FORQ-based Language Inclusion Formal Testing⋆

Kyveli Doveri[1,2][0000−0001−9403−2860], Pierre Ganty[1][0000−0002−3625−6003], and
Nicolas Mazzocchi[1,3][0000−0001−6425−5369]

CAV
Artifact
Evaluation
★
Available

[1] IMDEA Software Institute, Madrid, Spain
`firsname.lastname@imdea.org`
[2] Universidad Politécnica de Madrid, Spain
[3] Institute of Science and Technology Austria, Klosterneuburg, Austria

Abstract. We propose a novel algorithm to decide the language inclusion between (nondeterministic) Büchi automata, a PSPACE-complete problem. Our approach, like others before, leverage a notion of quasiorder to prune the search for a counterexample by discarding candidates which are subsumed by others for the quasiorder. Discarded candidates are guaranteed to not compromise the completeness of the algorithm. The novelty of our work lies in the quasiorder used to discard candidates. We introduce FORQs (family of right quasiorders) that we obtain by adapting the notion of family of right congruences put forward by Maler and Staiger in 1993. We define a FORQ-based inclusion algorithm which we prove correct and instantiate it for a specific FORQ, called the structural FORQ, induced by the Büchi automaton to the right of the inclusion sign. The resulting implementation, called FORKLIFT, scales up better than the state-of-the-art on a variety of benchmarks including benchmarks from program verification and theorem proving for word combinatorics.

Keywords: Language inclusion · Büchi automata · Well-quasiorders
Artifact: https://doi.org/10.5281/zenodo.6552870

1 Introduction

In verification [19,20] and theorem proving [31], Büchi automata have been used as the underlying formal model. In these settings, Büchi automata respectively encode 1) the behaviors of a system as well as properties about it; and 2) the set of valuations satisfying a predicate. Questions like asking whether a system complies with a specification naturally reduce to a language inclusion problem and so does proving a theorem of the form $\forall x\, \exists y,\, P(x) \Rightarrow Q(y)$.

In this paper we propose a new algorithm for the inclusion problem between ω-regular languages given by Büchi automata. The problem is PSPACE-complete [23] and significant effort has been devoted to the discovery of algorithms for inclusion that behave well in practice [8,10,14,18,22,25]. Each proposed algorithm is characterized by a set of techniques (e.g. Ramsey-based, rank-based) and heuristics (e.g. antichains, simulation relations). The algorithm we propose falls into the category of Ramsey-based algorithms and uses the antichain [11] heuristics: the search for counterexamples is pruned using quasiorders. Intuitively when two candidate counterexamples are comparable with respect to some considered quasiorder, the "higher" of the two can be discarded without compromising completeness of the search. In our setting, counterexamples to inclusion are ultimately periodic words, i.e., words of the form uv^ω, where u and v are called a *stem* and a *period*, respectively. Therefore pruning is done by comparing stems and periods of candidate counterexamples during the search.

In the work of Abdulla et al. [7,8] which was further refined by Clemente et al. [10] they use a single quasiorder to compare both stems and periods. Their effort has been focused on refining that single quasiorder by enhancing it with simulation relations. Others including some authors of this paper, followed an orthogonal line [13,22] that investigates the use of two quasiorders: one for the stems and another one, independent, for the periods. The flexibility of using different quasiorders yields more pruning when searching for a counterexample. In this paper, we push the envelope further by using an unbounded number of quasiorders: one for the stems and a family of quasiorders for the periods each of them depending on a distinct stem. We use the acronym FORQ, which stands for *family of right quasiorders*, to refer to these quasiorders. Using FORQs leads to significant algorithmic differences compared to the two quasiorders approaches. More precisely, the algorithms with two quasiorders [13,22] compute exactly two fixpoints (one for the stems and one for the periods) independently whereas the FORQ-based algorithm that we present computes two fixpoints for the stems and unboundedly many fixpoints for the periods (depending on the number of stems that belong to the first two fixpoints). Even though we lose the stem/period independence and we compute more fixpoints, in practice, the use of FORQs scales up better than the approaches based on one or two quasiorders.

We formalize the notion of FORQ by relaxing and generalizing the notion of *family of right congruences* introduced by Maler and Staiger [30] to advance the theory of recognizability of ω-regular languages and, in particular, questions related to minimal-state automata. Recently, families of right congruences have been used in other contexts like the learning of ω-regular languages (see [9] and references therein) and Büchi automata complementation [26].

Below, we describe how our contributions are organized:

- We define the notion of FORQs and leverage them to identify key finite sets of stems and periods that are sound and complete to decide the inclusion problem (Section 3).
- We introduce a FORQ called the structural FORQ which relies on the structure of a given Büchi automaton (Section 4).

- We formulate a FORQ-based inclusion algorithm that computes such key sets as fixpoints, and then use these key stems and periods to search for a counterexample to inclusion via membership queries (Section 5).
- We study the algorithmic complexity of the FORQ-based inclusion algorithm instantiated with structural FORQs (Section 6).
- We implement the inclusion algorithm with structural FORQs in a prototype called FORKLIFT and we conduct an empirical evaluation on a set of 674 benchmarks (Section 7).

2 Preliminaries

Languages. Let Σ be a finite and non-empty *alphabet*. We write Σ^* to refer to the set of finite words over Σ and we write ε to denote the empty word. Given $u \in \Sigma^*$, we denote by $|u|$ the length of u. In particular $|\varepsilon| = 0$. We also define $\Sigma^+ \triangleq \Sigma^* \setminus \{\varepsilon\}$, and $\Sigma^{\nabla n} \triangleq \{u \in \Sigma^* \mid |u| \nabla n\}$ with $\nabla \in \{\leq, \geq\}$, hence $\Sigma^* = \Sigma^{\geq 0}, \Sigma^+ = \Sigma^{\geq 1}$. We write Σ^ω to refer to the set of infinite words over Σ. An infinite word $\mu \in \Sigma^\omega$ is said to be *ultimately periodic* if it admits a decomposition $\mu = uv^\omega$ with $u \in \Sigma^*$ (called a *stem*) and $v \in \Sigma^+$ (called a *period*). We fix an alphabet Σ throughout the paper.

Order Theory. Let E be a set of elements and $\ltimes$ be a binary relation over E. The relation $\ltimes$ is said to be a *quasiorder* when it is *reflexive* and *transitive*. Given a subset X of E, we define its *upward closure* with respect to the quasiorder $\ltimes$ by $_\ltimes\!\uparrow X \triangleq \{e \in E \mid \exists x \in X, x \ltimes e\}$. Given two subsets $X, Y \subseteq E$ the set Y is said to be a *basis* for X with respect to $\ltimes$, denoted $\mathfrak{B}_\ltimes(Y, X)$, whenever $Y \subseteq X$ and $_\ltimes\!\uparrow X = {}_\ltimes\!\uparrow Y$. The quasiorder $\ltimes$ is a *well-quasiorder* iff for each set $X \subseteq E$ there exists a finite set $Y \subseteq E$ such that $\mathfrak{B}_\ltimes(Y, X)$. This property on bases is also known as the *finite basis property*. Other equivalent definitions of well-quasiorders can be found in the literature [27], we will use the followings:

1. For every $\{e_i\}_{i \in \mathbb{N}} \in E^\mathbb{N}$ there exists $i, j \in \mathbb{N}$ with $i < j$ such that $e_i \ltimes e_j$.
2. No sequence $\{X_i\}_{i \in \mathbb{N}} \in \wp(E)^\mathbb{N}$ is such that $_\ltimes\!\uparrow X_1 \subsetneq {}_\ltimes\!\uparrow X_2 \subsetneq \ldots$ holds.[4]

Automata. A (nondeterministic) *Büchi automaton* $\mathcal{B}$ (BA for short) is a tuple (Q, q_I, Δ, F) where Q is a finite set of states including q_I, the initial state, $\Delta \subseteq Q \times \Sigma \times Q$ is the transition relation, and, $F \subseteq Q$ is the set of accepting states. We lift Δ to finite words as expected. We prefer to write $\mathcal{B} \colon q_1 \xrightarrow{u} q_2$ instead of $(q_1, u, q_2) \in \Delta$. In addition, we write $\mathcal{B} \colon q_1 \xrightarrow{u}_F q_2$ when there exists a state $q_F \in F$ and two words u_1, u_2 such that $\mathcal{B} \colon q_1 \xrightarrow{u_1} q_F \xrightarrow{u_2} q_2$, and $u = u_1 u_2$.

A run π of $\mathcal{B}$ over $\mu = a_0 a_1 \cdots \in \Sigma^\omega$ is a function $\pi \colon \mathbb{N} \to Q$ such that $\pi(0) = q_I$ and for all position $i \in \mathbb{N}$, we have that $\mathcal{B} \colon \pi(i) \xrightarrow{a_i} \pi(i+1)$. A run is said to be *accepting* if $\pi(i) \in F$ for infinitely many values of $i \in \mathbb{N}$. The language $L(\mathcal{B})$ of words *recognized* by $\mathcal{B}$ is the set of ω-words for which $\mathcal{B}$ admits an accepting run. A language L is ω-*regular* if it is recognized by some BA.

---

[4] The notation $\wp(E)$ denotes the set of all subsets of E.

3   Foundations of our Approach

Let $\mathcal{A} \triangleq (P, p_I, \Delta_{\mathcal{A}}, F_{\mathcal{A}})$ be a Büchi automaton and M be an ω-regular language. The main idea behind our approach is to compute a finite subset $T_{\mathcal{A}}$ of ultimately periodic words of $L(\mathcal{A})$ such that:

$$T_{\mathcal{A}} \subseteq M \iff L(\mathcal{A}) \subseteq M \ . \qquad\qquad (\dagger)$$

Then $L(\mathcal{A}) \subseteq M$ holds iff each of the finitely many words of $T_{\mathcal{A}}$ belongs to M which is tested via membership queries.

First we observe that such a subset always exists: if the inclusion holds take $T_{\mathcal{A}}$ to be any finite subset of $L(\mathcal{A})$ (empty set included); else take $T_{\mathcal{A}}$ to contain some ultimately periodic word that is a counterexample to inclusion. In what follows, we will show that a finite subset $T_{\mathcal{A}}$ satisfying ($\dagger$) can be computed by using an ordering to prune the ultimately periodic words of $L(\mathcal{A})$. We will obtain such an ordering using a family of right quasiorders, a notion introduced below.

Definition 1 (FORQ). *A family of right quasiorders is a pair $\langle \lesssim, \{\lll_u\}_{u \in \Sigma^*} \rangle$ where $\lesssim \ \subseteq \Sigma^* \times \Sigma^*$ is a right-monotonic[5] quasiorder as well as every $\lll_u \subseteq \Sigma^* \times \Sigma^*$ where $u \in \Sigma^*$. Additionally, for all $u, u' \in \Sigma^*$, we require $u \lesssim u' \Rightarrow \lll_{u'} \subseteq \lll_u$ called the FORQ constraint.*

First, we observe that the above definition uses separate orderings for stems and periods. The definition goes even further, the ordering used for periods is depending on stems so that a period may or may not be discarded depending on the stem under consideration. The FORQ constraint tells us that if the periods v and w compare for a stem u', that is $v \lll_{u'} w$, then they also compare for every stem u subsuming u', that is $v \lll_u w$ if $u \lesssim u'$.

Expectedly, a FORQ needs to satisfy certain properties for $T_{\mathcal{A}}$ to be finite, computable and for ($\dagger$) to hold (in particular the left to right direction). The property of right-monotonicity of FORQs is needed so that we can iteratively compute $T_{\mathcal{A}}$ via a fixpoint computation (see Section 5).

Definition 2 (Suitable FORQ). *A FORQ $\mathcal{F} \triangleq \langle \lesssim, \{\lll_u\}_{u \in \Sigma^*} \rangle$ is said to be finite (resp. decidable) when $\lesssim$, its converse $\lesssim^{-1}$, and $\{\lll_u\}$ for all $u \in \Sigma^*$ are all well-quasiorders (resp. computable). Given $L \subseteq \Sigma^\omega$, $\mathcal{F}$ is said to preserve L when for all $u, \hat{u} \in \Sigma^*$ and all $v, \hat{v} \in \Sigma^+$ if $uv^\omega \in L$, $u \lesssim \hat{u}$, $v \lll_{\hat{u}} \hat{v}$ and $\hat{u}\hat{v} \lesssim \hat{u}$ then $\hat{u}\hat{v}^\omega \in L$. Finally, $\mathcal{F}$ is said to be L-suitable (for inclusion) if it is finite, decidable and preserves L.*

Intuitively, the "well" property on the quasiorders ensures finiteness of $T_{\mathcal{A}}$. The perservation property ensures completeness: a counterexample to $L(\mathcal{A}) \subseteq M$ can only be discarded (that is, not included in $T_{\mathcal{A}}$) if it is subsumed by another ultimately periodic word in $T_{\mathcal{A}}$ that is also a counterexample to inclusion.

Before defining $T_{\mathcal{A}}$ we introduce for each state $p \in P$ the sets of words

$$\mathtt{Stem}_p \triangleq \{u \in \Sigma^* \mid \mathcal{A} \colon p_I \xrightarrow{\ u\ } p\} \qquad \text{and} \qquad \mathtt{Per}_p \triangleq \{v \in \Sigma^+ \mid \mathcal{A} \colon p \xrightarrow{\ v\ } p\} \ .$$

---

[5] A quasiorder $\ltimes$ on Σ^* is *right-monotonic* when $u \ltimes v$ implies $u\,w \ltimes v\,w$ for all $w \in \Sigma^*$.

The set $\mathtt{Stem}_p$ is the set of stems of $L(\mathcal{A})$ that reach state p in $\mathcal{A}$ while the set $\mathtt{Per}_p$ is the set of periods read by a cycle of $\mathcal{A}$ on state p.

Given a M-suitable FORQ $\mathcal{F} \triangleq \langle \lesssim, \{\lessapprox_u\}_{u \in \Sigma^*} \rangle$, we let

$$T_{\mathcal{A}} \triangleq \left\{ uv^{\omega} \mid \exists s \in F_{\mathcal{A}} \colon u \in U_s, v \in V_s^w \text{ for some } w \in W_s \text{ with } u \lesssim w \right\} \qquad (\ddagger)$$

where for all $p \in P$, the set U_p is a basis of $\mathtt{Stem}_p$ with respect to $\lesssim$, that is $\mathfrak{B}_{\lesssim}(U_p, \mathtt{Stem}_p)$ holds. Moreover $\mathfrak{B}_{\lesssim^{-1}}(W_p, \mathtt{Stem}_p)$ holds and $\mathfrak{B}_{\lessapprox_w}(V_p^w, \mathtt{Per}_p)$ holds for all $w \in W_p$. Note that the quasiorder $\lessapprox_w$ used to prune the periods of $\mathtt{Per}_p$ depends on a maximal w.r.t. $\lesssim$ stem w of $\mathtt{Stem}_p$ since w belongs to the basis W_p for $\lesssim^{-1}$. The correctness argument for choosing $\lessapprox_w$ essentially relies on the FORQ constraint as the proof of (†) given below shows. In Section 8 we will show, that when w is not "maximal" the quasiorder $\lessapprox_w$ yields a set $T_{\mathcal{A}}$ for which (†) does not hold.

Furthermore, we conclude from the finite basis property of the quasiorders of $\mathcal{F}$ that U_p, W_p and $\{V_p^w\}_{w \in \Sigma^*}$ are finite for all $p \in P$, hence $T_{\mathcal{A}}$ is a finite subset of ultimately periodic words of $L(\mathcal{A})$. Next we prove the equivalence (†). The proof crucially relies on the preservation property of $\mathcal{F}$ which allows discarding candidate counterexamples without loosing completeness, that is, if inclusion does not hold a counterexample will be returned.

Proof (of (†)*).* Consider $\mathtt{Ultim}_{\mathcal{A}} \triangleq \{ uv^{\omega} \mid \exists s \in F_{\mathcal{A}} \colon u \in \mathtt{Stem}_s, v \in \mathtt{Per}_s, uv \lesssim u \}$. It is easy to show that $\mathtt{Ultim}_{\mathcal{A}} = \{ uv^{\omega} \mid \exists s \in F_{\mathcal{A}} \colon u \in \mathtt{Stem}_s, v \in \mathtt{Per}_s \}$ (same definition as $\mathtt{Ultim}_{\mathcal{A}}$ buth without the constraint $uv \lesssim u$) by reasoning on properties of well-quasi orders.[6] It is well-known that ω-regular language inclusion holds if and only if it holds for ultimately periodic words. Formally $L(\mathcal{A}) \subseteq M$ holds if and only if $\mathtt{Ultim}_{\mathcal{A}} \subseteq M$ holds. Therefore, to prove (†), we show that $T_{\mathcal{A}} \subseteq M \Leftrightarrow \mathtt{Ultim}_{\mathcal{A}} \subseteq M$.

To prove the implication $\mathtt{Ultim}_{\mathcal{A}} \subseteq M \Rightarrow T_{\mathcal{A}} \subseteq M$ we start by taking a word $uv^{\omega} \in T_{\mathcal{A}}$ such that, by definition ($\ddagger$), $u \in U_s$ and $v \in V_s^w$ for some $s \in F_{\mathcal{A}}$ and $w \in W_s$. We conclude from $\mathfrak{B}_{<}(U_s, \mathtt{Stem}_s)$ and $\mathfrak{B}_{\lessapprox_w}(V_s^w, \mathtt{Per}_s)$ that $u \in U_s \subseteq \mathtt{Stem}_s$ and $v \in V_s^w \subseteq \mathtt{Per}_s$. Thus, we find that $uv^{\omega} \in \mathtt{Ultim}_{\mathcal{A}}$ hence the assumption $\mathtt{Ultim}_{\mathcal{A}} \subseteq M$ shows that $uv^{\omega} \in M$ which proves the implication.

Next, we prove that $T_{\mathcal{A}} \subseteq M \Rightarrow \mathtt{Ultim}_{\mathcal{A}} \subseteq M$ holds as well. Let $uv^{\omega} \in \mathtt{Ultim}_{\mathcal{A}}$, i.e., such that there exists $s \in F_{\mathcal{A}}$ for which $u \in \mathtt{Stem}_s$ and $v \in \mathtt{Per}_s$, satisfying $uv \lesssim u$. Since $u \in \mathtt{Stem}_s$ and $v \in \mathtt{Per}_s$, there exist $u_0 \in U_s$, $w_0 \in W_s$ and $v_0 \in V_s^{w_0}$ such that $u_0 \lesssim u \lesssim w_0$ and $v_0 \lessapprox_{w_0} v$ thanks to the finite basis property. By definition we have $u_0 v_0^{\omega} \in T_{\mathcal{A}}$ and thus we find that $u_0 v_0^{\omega} \in M$ since $T_{\mathcal{A}} \subseteq M$. Next since $u \lesssim w_0$, the FORQ constraint shows that $\lessapprox_{w_0} \subseteq \lessapprox_u$ which, in turn, implies that $v_0 \lessapprox_u v$ holds. Finally, we deduce from $u_0 v_0^{\omega} \in M$, $u_0 \lesssim u$, $v_0 \lessapprox_u v$, $uv \lesssim u$ and the preservation of M by the FORQ $\mathcal{F}$ that $uv^{\omega} \in M$. We thus obtain that $\mathtt{Ultim}_{\mathcal{A}} \subseteq M$ and we are done.   □

---

[6] The case $\subseteq$ is trivial. For the case $\supseteq$, let uv^{ω} with $u \in \mathtt{Stem}_s$ and $v \in \mathtt{Per}_s$. If $uv \lesssim u$ then we are done for otherwise consider the sequence $\{uv^i\}_{i \in \mathbb{N}}$. Since $\lesssim^{-1}$ is a well-quasiorder, there exists $x, y \in \mathbb{N}$ such that $x < y$ and $uv^x \lesssim^{-1} uv^y$ (viz. $uv^y \lesssim uv^x$). Therefore we have $(uv^x)(v^{y-x})^{\omega} = uv^{\omega}$, $(uv^x) \in \mathtt{Stem}_s$, $(v^{y-x}) \in \mathtt{Per}_s$, and $(uv^x)(v^{y-x}) \lesssim (uv^x)$, hence $uv^{\omega} \in \mathtt{Ultim}_{\mathcal{A}}$.

Example 3. To gain more insights about our approach consider the BAs of Fig. 1 for which we want to check whether $L(\mathcal{A}) \subseteq L(\mathcal{B})$ holds. From the description of $\mathcal{A}$ it is routine to check that $\mathtt{Stem}_{p_I} = \Sigma^*$ and $\mathtt{Per}_{p_I} = \Sigma^+$. Let us assume the existence[7] of $\lesssim$ (hence $\lesssim^{-1}$), $\ll_\varepsilon$ and $\ll_{aa}$ such that $a \lesssim aa$ holds and so does $\mathfrak{B}_{\lesssim}(\{\varepsilon, a\}, \Sigma^*)$, $\mathfrak{B}_{\lesssim^{-1}}(\{\varepsilon, aa\}, \Sigma^*)$, $\mathfrak{B}_{\ll_\varepsilon}(\{b\}, \Sigma^+)$ and $\mathfrak{B}_{\ll_{aa}}(\{a\}, \Sigma^+)$. In addition, we set $U_{p_I} = \{\varepsilon, a\}$ since $\mathfrak{B}_{\lesssim}(\{\varepsilon, a\}, \Sigma^*)$ and $W_{p_I} = \{\varepsilon, aa\}$ since $\mathfrak{B}_{\lesssim^{-1}}(\{\varepsilon, aa\}, \Sigma^*)$. Moreover $V_{p_I}^\varepsilon = \{b\}$ since $\mathfrak{B}_{\ll_\varepsilon}(\{b\}, \Sigma^+)$, and $V_{p_I}^{aa} = \{a\}$ since $\mathfrak{B}_{\ll_{aa}}(\{a\}, \Sigma^+)$. Next by definition (‡) of $T_{\mathcal{A}}$ and from $a \lesssim aa$ we deduce that $T_{\mathcal{A}} = \{\varepsilon(b)^\omega, a(a)^\omega\}$. Finally, we conclude from (†) and $a^\omega \in T_{\mathcal{A}}$ that $a^\omega \in L(\mathcal{A})$ (since $T_{\mathcal{A}} \subseteq L(\mathcal{A})$) hence that $L(\mathcal{A}) \nsubseteq L(\mathcal{B})$ because $a^\omega \notin L(\mathcal{B})$. By checking membership of the two ultimately periodic words of $T_{\mathcal{A}}$ into $L(\mathcal{B})$ we thus have shown that $L(\mathcal{A}) \subseteq L(\mathcal{B})$ does not hold.

In the example above we did not detail how the FORQ was obtained let alone how to compute the finite bases. We fill that gap in the next two sections: we define FORQs based on the underlying structure of a given BA in Section 4 and show they are suitable; and we give an effective computation of the bases hence our FORQ-based inclusion algorithm in Section 5.

4   Defining **FORQs** from the Structure of an Automaton

In this section we introduce a type of FORQs called structural FORQs such that given a BA $\mathcal{B}$ the structural FORQ induced by $\mathcal{B}$ is $L(\mathcal{B})$-suitable.

Definition 4. *Let $\mathcal{B} \triangleq (Q, q_I, \Delta_{\mathcal{B}}, F_{\mathcal{B}})$ be a BA. The structural FORQ of $\mathcal{B}$ is the pair $\langle \lesssim^{\mathcal{B}}, \{\ll_u^{\mathcal{B}}\}_{u \in \Sigma^*} \rangle$ where the quasiorders are defined by:*

$$u_1 \lesssim^{\mathcal{B}} u_2 \stackrel{\triangle}{\Longleftrightarrow} \mathtt{Tgt}_{\mathcal{B}}(u_1) \subseteq \mathtt{Tgt}_{\mathcal{B}}(u_2)$$
$$v_1 \ll_u^{\mathcal{B}} v_2 \stackrel{\triangle}{\Longleftrightarrow} \mathtt{Cxt}_{\mathcal{B}}(\mathtt{Tgt}_{\mathcal{B}}(u), v_1) \subseteq \mathtt{Cxt}_{\mathcal{B}}(\mathtt{Tgt}_{\mathcal{B}}(u), v_2)$$

with $\mathtt{Tgt}_{\mathcal{B}} \colon \wp(Q) \times \Sigma^* \to \wp(Q)$ and $\mathtt{Cxt}_{\mathcal{B}} \colon \wp(Q) \times \Sigma^* \to \wp(Q^2 \times \{\bot, \top\})$

$$\mathtt{Tgt}_{\mathcal{B}}(u) \triangleq \{q' \in Q \mid \mathcal{B} \colon q_I \xrightarrow{u} q'\}$$
$$\mathtt{Cxt}_{\mathcal{B}}(X, v) \triangleq \{(q, q', k) \mid q \in X, \mathcal{B} \colon q \xrightarrow{v} q', (k = \top \Rightarrow \mathcal{B} \colon q \xrightarrow{v}_F q')\}$$

Given $u \in \Sigma^*$, the set $\mathtt{Tgt}_{\mathcal{B}}(u)$ contains states that u can "target" from the initial state q_I. A "context" (q, q', k) returned by $\mathtt{Cxt}_{\mathcal{B}}$, consists in a source state $q \in Q$, a sink state $q' \in Q$ and a boolean $k \in \{\top, \bot\}$ that keeps track whether an accepting state is visited. Note that, having $\bot$ as last component of a context does *not* mean that no accepting state is visited. When it is clear from the context, we often omit the subscript $\mathcal{B}$ from $\mathtt{Tgt}_{\mathcal{B}}$ and $\mathtt{Cxt}_{\mathcal{B}}$. Analogously, we omit the BA from the structural FORQ quasiorders when there is no ambiguity.

Lemma 5. *Given a BA $\mathcal{B}$, the pair $\langle \lesssim^{\mathcal{B}}, \{\ll_u^{\mathcal{B}}\}_{u \in \Sigma^*} \rangle$ of Definition 4 is a FORQ.*

---

[7] The definition of the orderings, needed to compute the bases, are given in Example 6.

Fig. 1: Büchi automata $\mathcal{A}$ and $\mathcal{B}$ over the alphabet $\Sigma = \{a, b\}$.

Proof. Let $\mathcal{B} \triangleq (Q, q_I, \Delta_\mathcal{B}, F_\mathcal{B})$ be a BA, we start by proving that the FORQ constraint holds: $u \lesssim^\mathcal{B} u' \implies \lll^\mathcal{B}_{u'} \subseteq \lll^\mathcal{B}_u$. First, we observe that, for all $Y \subseteq X \subseteq Q$ and all $v, v' \in \Sigma^*$, we have that $\mathtt{Cxt}(X, v) \subseteq \mathtt{Cxt}(X, v') \Rightarrow \mathtt{Cxt}(Y, v) \subseteq \mathtt{Cxt}(Y, v')$. Consider $u, u' \in \Sigma^*$ such that $u \lesssim^\mathcal{B} u'$ and $v, v' \in \Sigma^*$ such that $v \lll^\mathcal{B}_{u'} v'$. Let $X \triangleq \mathtt{Tgt}(u)$ and $X' \triangleq \mathtt{Tgt}(u')$, we have that $X \subseteq X'$ following $u \lesssim^\mathcal{B} u'$. Next, we conclude from $v \lll^\mathcal{B}_{u'} v'$ that $\mathtt{Cxt}(X', v) \subseteq \mathtt{Cxt}(X', v')$, hence that $\mathtt{Cxt}(X, v) \subseteq \mathtt{Cxt}(X, v')$ by the above reasoning using $X \subseteq X'$, and finally that $v \lll^\mathcal{B}_u v'$.

For the right monotonicity, Definition 4 shows that if $\mathtt{Tgt}(u) \subseteq \mathtt{Tgt}(v)$ then $\mathtt{Tgt}(ua) \subseteq \mathtt{Tgt}(va)$, hence we have $u \lesssim v$ implies $ua \lesssim va$ for all $a \in \Sigma$. The reasoning with the other quasiorders and $\mathtt{Cxt}$ proceeds analogously.   □

Example 6. Consider the BA $\mathcal{B}$ of Fig. 1 and let $\langle \lesssim, \{\lll_u\}_{u \in \Sigma^*}\rangle$ be its structural FORQ. More precisely, we have $\mathtt{Tgt}(\varepsilon) = \{q_I\}$; $\mathtt{Tgt}(a) = \mathtt{Tgt}(b) = \{q_1\}$; and $\mathtt{Tgt}(u) = \{q_1, q_2\}$ for all $u \in \Sigma^{\geq 2}$. In particular we conclude from $u_1 \lesssim u_2 \overset{\triangle}{\iff} \mathtt{Tgt}(u_1) \subseteq \mathtt{Tgt}(u_2)$ that $a \lesssim aa$, $a \lesssim b$ and $b \lesssim a$; ε and a are incomparable; and so are ε and aa. Since $\mathtt{Tgt}$ has only three distinct outputs, the set $\{\lll_u\}_{u \in \Sigma^*}$ contains three distinct quasiorders.

1. $v_1 \lll_\varepsilon v_2 \overset{\triangle}{\iff} \mathtt{Cxt}(\{q_I\}, v_1) \subseteq \mathtt{Cxt}(\{q_I\}, v_2)$ where
   - $\mathtt{Cxt}(\{q_I\}, \varepsilon) = \{(q_I, q_I, \bot)\}$
   - $\mathtt{Cxt}(\{q_I\}, a) = \mathtt{Cxt}(\{q_I\}, b) = \{(q_I, q_1, \bot)\}$
   - $\mathtt{Cxt}(\{q_I\}, v) = \{(q_I, q_1, \bot), (q_I, q_2, \bot), (q_I, q_2, \top)\}$ for all $v \in \Sigma^{\geq 2}$
2. $v_1 \lll_a v_2 \iff v_1 \lll_b v_2 \overset{\triangle}{\iff} \mathtt{Cxt}(\{q_1\}, v_1) \subseteq \mathtt{Cxt}(\{q_1\}, v_2)$ where
   - $\mathtt{Cxt}(\{q_1\}, \varepsilon) = \{(q_1, q_1, \bot)\}$
   - $\mathtt{Cxt}(\{q_1\}, v) = \{(q_1, q_1, \bot), (q_1, q_2, \bot), (q_1, q_2, \top)\}$ for all $v \in \Sigma^+$
3. $v_1 \lll_{u_1} v_2 \iff v_1 \lll_{u_2} v_2 \overset{\triangle}{\iff} \mathtt{Cxt}(\{q_1, q_2\}, v_1) \subseteq \mathtt{Cxt}(\{q_1, q_2\}, v_2)$ for all $u_1, u_2 \in \Sigma^{\geq 2}$ where
   - $\mathtt{Cxt}(\{q_1, q_2\}, \varepsilon) = \{(q_1, q_1, \bot), (q_2, q_2, \bot), (q_2, q_2, \top)\}$
   - $\mathtt{Cxt}(\{q_1, q_2\}, av') = \{(q_1, q_1, \bot), (q_1, q_2, \bot), (q_1, q_2, \top)\}$ for all $v' \in \Sigma^*$
   - $\mathtt{Cxt}(\{q_1, q_2\}, bv') = \mathtt{Cxt}(\{q_1, q_2\}, av') \cup \{(q_2, q_2, \bot), (q_2, q_2, \top)\}$ for all $v' \in \Sigma^*$

With the above definitions the reader is invited to check the following predicates $\mathfrak{B}_\lesssim(\{\varepsilon, a\}, \Sigma^*)$, $\mathfrak{B}_\lesssim(\{\varepsilon, b\}, \Sigma^*)$, $\mathfrak{B}_{\lesssim^{-1}}(\{\varepsilon, aa\}, \Sigma^*)$, $\mathfrak{B}_{\lll_\varepsilon}(\{b\}, \Sigma^+)$, $\mathfrak{B}_{\lll_a}(\{b\}, \Sigma^+)$ and $\mathfrak{B}_{\lll_{aa}}(\{a\}, \Sigma^+)$. Also observe that none of the above finite bases contains comparable words for the ordering thereof. We also encourage the reader to revisit Example 3.

As prescribed in Section 3, we show that for every BA $\mathcal{B}$ its structural FORQ is $L(\mathcal{B})$-suitable, namely it is finite, decidable and preserves $L(\mathcal{B})$.

Proposition 7. *Given a BA $\mathcal{B}$, its structural FORQ is $L(\mathcal{B})$-suitable.*

Proof. Let $\mathcal{B} \triangleq (Q, q_I, \Delta_{\mathcal{B}}, F_{\mathcal{B}})$ be a BA and $\mathcal{F} \triangleq \langle \lesssim, \{\lll_u\}_{u \in \Sigma^*}\rangle$ be its structural FORQ. The finiteness proof of $\mathcal{F}$ is trivial since Q is finite and so is the proof of decidability by Definition 4. For the preservation, given $u_0 v_0^\omega \in L(\mathcal{B})$, we show that for all $u \in \Sigma^*$ and all $v \in \Sigma^+$ such that $uv \lesssim u$ and $u_0 \lesssim u$ and $v_0 \lll_u v$ then $uv^\omega \in L(\mathcal{B})$ holds. Let a run $\pi_0 \triangleq q_I \xrightarrow{u_0} q_0 \xrightarrow{v_0} q_1 \xrightarrow{v_0} q_2 \ldots$ of $\mathcal{B}$ over $u_0 v_0^\omega$ which is accepting. Stated equivalently, we have $q_0 \in \mathtt{Tgt}(u_0)$ and $(q_i, q_{i+1}, x_i) \in \mathtt{Cxt}(\mathtt{Tgt}(u_0 v_0^i), v_0)$ for every $i \in \mathbb{N}$ with the additional constraint that $x_i = \top$ holds infinitely often.

We will show that $\mathcal{B}$ has an accepting run over uv^ω by showing that $q_0 \in \mathtt{Tgt}(u)$ holds; $(q_i, q_{i+1}, x_i) \in \mathtt{Cxt}(\mathtt{Tgt}(uv^i), v)$ holds for every $i \in \mathbb{N}$; and $x_i = \top$ holds infinitely often. Since $u_0 \lesssim u$ and $q_0 \in \mathtt{Tgt}(u_0)$ we find that $q_0 \in \mathtt{Tgt}(u)$ by definition of $\lesssim$. Next we show the remaining constraints by induction. The induction hypothesis states that for all $0 \le n$ we have $(q_n, q_{n+1}, x_n) \in \mathtt{Cxt}(\mathtt{Tgt}(uv^n), v)$. For the base case $(n = 0)$ we have to show that $(q_0, q_1, x_0) \in \mathtt{Cxt}(\mathtt{Tgt}(u), v)$. We conclude from $(q_0, q_1, x_0) \in \mathtt{Cxt}(\mathtt{Tgt}(u), v_0)$, $v_0 \lll_u v$ and the definition of $\lll_u$ that $\mathtt{Cxt}(\mathtt{Tgt}(u), v_0) \subseteq \mathtt{Cxt}(\mathtt{Tgt}(u), v)$ and finally that $(q_0, q_1, x_0) \in \mathtt{Cxt}(\mathtt{Tgt}(u), v)$. For the inductive case, assume $(q_n, q_{n+1}, x_n) \in \mathtt{Cxt}(\mathtt{Tgt}(uv^n), v)$. The definition of context shows that $q_{n+1} \in \mathtt{Tgt}(uv^{n+1})$. It takes an easy an induction to show that $uv^n \lesssim u$ for all n using $uv \lesssim u$ and right-monotonicity of $\lesssim$. We conclude from $uv^{n+1} \lesssim u$, the definition of $\lesssim$ and $q_{n+1} \in \mathtt{Tgt}(uv^{n+1})$ that $q_{n+1} \in \mathtt{Tgt}(u)$ also holds, hence that $(q_{n+1}, q_{n+2}, x_{n+1}) \in \mathtt{Cxt}(\mathtt{Tgt}(u), v_0)$ following the definition of contexts and that of π_0. Next, we find that $(q_{n+1}, q_{n+2}, x_{n+1}) \in \mathtt{Cxt}(\mathtt{Tgt}(u), v)$ following a reasoning analogous to the base case, this time starting with $(q_{n+1}, q_{n+2}, x_{n+1}) \in \mathtt{Cxt}(\mathtt{Tgt}(u), v_0))$. Finally, $q_{n+1} \in \mathtt{Tgt}(uv^{n+1})$ implies that $(q_{n+1}, q_{n+2}, x_{n+1}) \in \mathtt{Cxt}(\mathtt{Tgt}(uv^{n+1}), v)$. We have thus shown that $q_0 \in \mathtt{Tgt}(u)$ and $(q_i, q_{i+1}, x_i) \in \mathtt{Cxt}(\mathtt{Tgt}(uv^i), v)$ for every $i \in \mathbb{N}$ with the additional constraint that $x_i = \top$ holds infinitely often and we are done. $\square$

5   A **FORQ**-based Inclusion Algorithm

As announced at the end of Section 3 it remains, in order to formulate our FORQ-based algorithm deciding whether $L(\mathcal{A}) \subseteq M$ holds, to give an effective computation for the bases defining $T_{\mathcal{A}}$. We start with a fixpoint characterization of the stems and periods of BAs using the function $\mathtt{Rcat}_{\mathcal{A}} \colon \wp(\Sigma^*)^{|P|} \to \wp(\Sigma^*)^{|P|}$:

$$\mathtt{Rcat}_{\mathcal{A}}(\vec{X}).p \triangleq \vec{X}.p \cup \big\{wa \in \Sigma^* \mid w \in \vec{X}.p', a \in \Sigma, \mathcal{A} \colon p' \xrightarrow{a} p\big\}$$

where $\vec{S}.p$ denotes the p-th element of the vector $\vec{S} \in \wp(\Sigma^*)^{|P|}$. In Fig. 2, the repeat/until loops at lines 4 and 5 compute iteratively subsets of the stems of $\mathcal{A}$, while the loop at line 10 computes iteratively subsets of the periods of $\mathcal{A}$. The following lemma formalizes the above intuition.

Input: Büchi automaton $\mathcal{A} \triangleq (P, p_I, \Delta_{\mathcal{A}}, F_{\mathcal{A}})$
Input: ω-regular language M with procedure deciding $uv^\omega \in M$ given u, v
Input: M-suitable FORQ $\mathcal{F} \triangleq \langle \precsim, \{\lll_u\}_{u \in \Sigma^*} \rangle$
Output: Returns ok if $L(\mathcal{A}) \subseteq M$ and ko otherwise

1 **Function:**
2   let $\vec{U}_0 \in \wp(\Sigma^*)^{|P|}$ as $\vec{U}_0.p \triangleq \varnothing$ **with** $p \neq p_I$ **and** $\vec{U}_0.p_I \triangleq \{\varepsilon\}$
3    $\vec{W} := \vec{U} := \vec{U}_0$
4   **repeat** $\vec{W} := \mathrm{Rcat}_{\mathcal{A}}(\vec{W})$ **until** $\mathfrak{B}_{\precsim^{-1}}(\vec{W}.p, \mathrm{Rcat}_{\mathcal{A}}(\vec{W}).p)$ *for all $p \in P$*
5   **repeat** $\vec{U} := \mathrm{Rcat}_{\mathcal{A}}(\vec{U})$ **until** $\mathfrak{B}_{\precsim}(\vec{U}.p, \mathrm{Rcat}_{\mathcal{A}}(\vec{U}).p)$ *for all $p \in P$*
6   **for each** $s \in F_{\mathcal{A}}$ **do**
7   let $\vec{V}_1^s \in \wp(\Sigma^*)^{|P|}$ as $\vec{V}_1^s.p \triangleq \{a \in \Sigma \mid \mathcal{A}: s \xrightarrow{a} p\}$ **with** $p \in P$
8   **for each** $w \in \vec{W}.s$ **do**
9        $\vec{V}^s := \vec{V}_1^s$
10   **repeat** $\vec{V}^s := \mathrm{Rcat}_{\mathcal{A}}(\vec{V}^s)$ **until** $\mathfrak{B}_{\lll_w}(\vec{V}^s.p, \mathrm{Rcat}_{\mathcal{A}}(\vec{V}^s).p)$ *for all $p \in P$*
11   **for each** $v \in \vec{V}^s.s$ **do**
12   **for each** $u \in \vec{U}.s$ such that $u \precsim w$ **do**
13   **if** $uv^\omega \notin M$ **then return** ko
14   **return** ok

Fig. 2: FORQ-based algorithm

Lemma 8. *Consider $\vec{U}_0$ and $\vec{V}_1^s$ (with $s \in F_{\mathcal{A}}$) in the FORQ-based algorithm. The following holds for all $n \in \mathbb{N}$:*

$$\mathrm{Rcat}_{\mathcal{A}}^n(\vec{U}_0).p = \mathrm{Stem}_p \cap \Sigma^{\leq n} \text{ for all } p \in P, \text{ and } \mathrm{Rcat}_{\mathcal{A}}^n(\vec{V}_1^s).s = \mathrm{Per}_s \cap \Sigma^{\leq n+1} .$$

Prior to proving the correctness of the algorithm of Fig. 2 we need the following result which is key for establishing the correctness of the repeat/until loop conditions of lines 4, 5, and 10.

Lemma 9. *Let $\bowtie$ be a right-monotonic quasiorder over Σ^*. Given $\mathcal{A} \triangleq (P, p_I, \Delta_{\mathcal{A}}, F_{\mathcal{A}})$ and $\vec{S}, \vec{S}' \in \wp(\Sigma^*)^{|P|}$, if $\mathfrak{B}_{\bowtie}(\vec{S}'.p, \vec{S}.p)$ holds for all $p \in P$ then $\mathfrak{B}_{\bowtie}(\mathrm{Rcat}_{\mathcal{A}}(\vec{S}').p, \mathrm{Rcat}_{\mathcal{A}}(\vec{S}).p)$ holds for all $p \in P$.*

Proof. Consider $w \in \mathrm{Rcat}_{\mathcal{A}}(\vec{S}).p$ where $p \in P$, we show that there exists $w' \in \mathrm{Rcat}_{\mathcal{A}}(\vec{S}').p$ such that $w' \bowtie w$. Assume that $\mathfrak{B}_{\bowtie}(\vec{S}'.p, \vec{S}.p)$ holds for all $p \in P$. In particular, for all $w_1 \in \vec{S}.p$, there exists $w_1' \in \vec{S}'.p$ such that $w_1' \bowtie w_1$. In the case where $w_1 \in \mathrm{Rcat}_{\mathcal{A}}(\vec{S}).p \setminus \vec{S}.p$, by definition of $\mathrm{Rcat}_{\mathcal{A}}$ w_1 is of the form $w_2 a$ for some $a \in \Sigma$ and some $w_2 \in \vec{S}.\hat{p}$ such that $\mathcal{A}: \hat{p} \xrightarrow{a} p$. Since $\mathfrak{B}_{\bowtie}(\vec{S}'.\hat{p}, \vec{S}.\hat{p})$ and $w_2 \in \vec{S}.\hat{p}$, there exists $w_3 \in \vec{S}'.\hat{p}$ such that $w_3 \bowtie w_2$. We deduce that $w_3 a \bowtie w_2 a$ holds, hence $w_3 a \bowtie w_1$ holds as well from the right-monotonicity of $\bowtie$. Furthermore $w_3 a \in \mathrm{Rcat}_{\mathcal{A}}(\vec{S}').p$ by definition of $\mathrm{Rcat}_{\mathcal{A}}$ and since $\mathcal{A}: \hat{p} \xrightarrow{a} p$. Finally, we conclude that $\mathfrak{B}_{\bowtie}(\mathrm{Rcat}_{\mathcal{A}}(\vec{S}'), \mathrm{Rcat}_{\mathcal{A}}(\vec{S}))$ holds. $\square$

Theorem 10. *The FORQ-based algorithm decides the inclusion of BAs.*

Proof. We first show that every loop of the algorithm eventually terminates. First, we conclude from the definition of $\mathtt{Rcat}_{\mathcal{A}}$ and the initializations (lines 3 and 9) of each repeat/until loop (lines 4, 5, and 10) that each component of each vector holds a finite set of words. Observe that the halting conditions of the repeat/until loops are effectively computable since every quasiorder of $\mathcal{F}$ is decidable and because, in order to decide $\mathfrak{B}_{\bowtie}(Y, X)$ where X, Y are finite sets and $\bowtie$ is decidable, it suffices to check that $Y \subseteq X$ and that for every $x \in X$ there exists $y \in Y$ such that $y \bowtie x$. Next, we conclude from the fact that all the quasiorders of $\mathcal{F}$ used in the repeat/until loops are all well-quasiorders that there is no infinite sequence $\{X_i\}_{i \in \mathbb{N}}$ such that $_{\bowtie}\!\lceil X_1 \subsetneq {}_{\bowtie}\!\lceil X_2 \subsetneq \ldots$ Since $\mathfrak{B}_{\bowtie}(Y, X)$ is equivalent to $Y \subseteq X \wedge {}_{\bowtie}\!\lceil X \subseteq {}_{\bowtie}\!\lceil Y$ and since each time $\mathtt{Rcat}_{\mathcal{A}}$ updates a component its upward closure after the update includes the one before, we find that every repeat/until loop must terminate after finitely many iterations.

Next, we show that when the repeat/until loop of line 5 halts, $\mathfrak{B}_{\lesssim}(\vec{U}.p, \mathtt{Stem}_p)$ holds for all $p \in P$. It takes an easy induction on n together with Lemma 9 to show that if $\mathfrak{B}_{\lesssim}(\mathtt{Rcat}_{\mathcal{A}}^{n+1}(\vec{U_0}).p, \mathtt{Rcat}_{\mathcal{A}}^{n}(\vec{U_0}).p)$ holds for all $p \in P$ then $\mathfrak{B}_{\lesssim}(\mathtt{Rcat}_{\mathcal{A}}^{n}(\vec{U_0}).p, \mathtt{Rcat}_{\mathcal{A}}^{m}(\vec{U_0}).p)$ holds for all $m > n$. Hence Lemma 8 shows that $\mathfrak{B}_{\lesssim}(\mathtt{Rcat}_{\mathcal{A}}^{k}(\vec{U_0}).p, \mathtt{Stem}_p)$ holds for all $p \in P$ where k is the number of iterations of the repeat/until loop implying $\mathfrak{B}_{\lesssim}(\vec{U}.p, \mathtt{Stem}_p)$ holds when the loop of line 5 halts.

An analogue reasoning shows that $\mathfrak{B}_{\lesssim^{-1}}(\vec{W}.p, \mathtt{Stem}_p)$ holds for all $p \in P$, as well as $\mathfrak{B}_{\lessapprox_w}(\vec{V}^s.s, \mathtt{Per}_s)$ holds for all $w \in \vec{W}.s$ and all $s \in F_{\mathcal{A}}$ upon termination of the loops of lines 4 and 10.

To conclude, we observe that each time a membership query is performed at line 13, the ultimately periodic word uv^{ω} belongs to $T_{\mathcal{A}}$ defined by ($\ddagger$). This is ensured since $u \in \mathfrak{B}_{\lesssim}(\vec{U}.s, \mathtt{Stem}_s)$, $w \in \mathfrak{B}_{\lesssim^{-1}}(\vec{W}.s, \mathtt{Stem}_s)$, $v \in \mathfrak{B}_{\lessapprox_w}(\vec{V}^s.s, \mathtt{Per}_s)$ for some $s \in F_{\mathcal{A}}$ and, thanks to the test at line 12, the comparison $u \lesssim w$ holds. $\square$

Remark 11. The correctness of the FORQ-based algorithm still holds when, after every ":=" assignment (at lines 3, 4, 5, 9 and 10), we remove from the variable content zero or more subsumed words for the corresponding ordering. The effect of removing zero or more subsumed words from a variable can be achieved by replacing assignments like, for instance, $\vec{U} := \mathtt{Rcat}_{\mathcal{A}}(\vec{U})$ at line 5 with $\vec{U} := \mathtt{Rcat}_{\mathcal{A}}(\vec{U}); \vec{U} := \vec{U}_r$ where $\vec{U}_r$ satisfies $\mathfrak{B}_{\lesssim}(\vec{U}_r.p, \vec{U}.p)$ for all $p \in P$. The correctness of the previous modification follows from Lemma 9. Therefore, the sets obtained by discarding subsumed words during computations still satisfy the basis predicates of $T_{\mathcal{A}}$ given at ($\ddagger$).

It is worth pointing out that the correctness arguments developed above, do not depend on the specifics of the structural FORQs. The FORQ-based algorithm is sound as long as we provide a suitable FORQ. Next we study the algorithmic complexity of the algorithm of Fig. 2.

6   Complexity of the structural **FORQ**-based algorithm

In this Section, we establish an upper bound on the runtime of the algorithm of Fig. 2 when the input **FORQ** is the structural **FORQ** induced by a BA $\mathcal{B}$. Let $n_\mathcal{A}$ and $n_\mathcal{B}$ be respectively the number of states in the BA $\mathcal{A}$ and $\mathcal{B}$. We start by bounding the number of iterations in the repeat/until loops. In each repeat/until loop, each component of the vector holds a finite set of words the upward closure of which grows (for $\subseteq$) over time and when all the upward closures stabilize the loop terminates. In the worst case, an iteration of the repeat/until loop adds exactly one word to some component of the vector which keeps the halting condition falsified (the upward closure strictly increases). Therefore a component of the vector cannot be updated more than $2^{n_\mathcal{B}}$ times for otherwise its upward closure has stabilized. We thus find that the total number of iterations is bounded from above by $n_\mathcal{A} \cdot 2^{n_\mathcal{B}}$ for the loops computing $\vec{U}$ and $\vec{W}$. Using an analogous reasoning we conclude that each component of the $\vec{V}$ vector has no more than $2^{(2n_\mathcal{B}^2)}$ elements and the total number of iterations is upper-bounded by $n_\mathcal{A} \cdot 2^{(2n_\mathcal{B}^2)}$. To infer an upper bound on the runtime of each repeat/until loop we also need to multiply the above expressions by a factor $|\Sigma|$ since the number of concatenations in `Rcat` depends on the size of the alphabet.

Next, we derive an upper bound on the number of membership queries performed at line 13. The number of iterations of the loops of lines 6, 8, 10, 11 and 12 is $n_\mathcal{A}$, $2^{n_\mathcal{B}}$, $n_\mathcal{A} \cdot 2^{(2n_\mathcal{B}^2)}$, $2^{(2n_\mathcal{B}^2)}$ and $2^{n_\mathcal{B}}$, respectively. Since all loops are nested, we multiply these bounds to end up with $n_\mathcal{A}^2 \cdot 2^{\mathcal{O}(n_\mathcal{B}^2)}$ as an upper bound on the number of membership queries. The runtime for each ultimately periodic word membership query (with a stem, a period and $\mathcal{B}$ as input) is upper bounded by an expression polynomial in the size $n_\mathcal{B}$ of $\mathcal{B}$, $2^{n_\mathcal{B}}$ for the length of the stem and $2^{(2n_\mathcal{B}^2)}$ for the length of the period.

We conclude from the above that the runtime of the algorithm of Fig. 2 is at most $|\Sigma| \cdot n_\mathcal{A}^2 \cdot 2^{\mathcal{O}(n_\mathcal{B}^2)}$.

7   Implementation and experiments

We implemented the **FORQ**-based algorithm of Fig. 2 instantiated by the structural **FORQ** in a tool called FORKLIFT [2]. In this section, we provide algorithmic details about FORKLIFT and then analyze how it behaves in practice (Section 7.1).

Data structures. Comparing two words given a structural **FORQ** requires to compute the corresponding sets of target for stems (`Tgt`), and sets of context for periods (`Cxt`). A naïve implementation would be to compute `Tgt` and `Cxt` every time a comparison is needed. We avoid to compute this information over and over again by storing each stem together with its `Tgt` set and each period together with its `Cxt` set.

Moreover, the function `Rcat` inserts new words in the input vector by concatenating a letter on the right to some words already in the vector. In our

implementation, we do not recompute the associated set of targets nor context for the newly computed word from scratch. For all stem $u \in \Sigma^*$ and all letter $a \in \Sigma$, the set of states $\mathtt{Tgt}(ua)$ can be computed from $\mathtt{Tgt}(u)$ thanks to the following equality essentially stating that $\mathtt{Tgt}()$ can be computed inductively:

$$\mathtt{Tgt}(ua) = \left\{ q \in Q \mid q' \in \mathtt{Tgt}(u), \mathcal{B} \colon q' \xrightarrow{a} q \right\} \ .$$

Analogously, for all period $v \in \Sigma^+$, all $X \subseteq Q$ and all $a \in \Sigma$, the set of contexts $\mathtt{Cxt}(X, va)$ can be computed from $\mathtt{Cxt}(X, v)$ thanks to the following equality:

$$\mathtt{Cxt}(X, va) = \left\{ (q_0, q, k) \in Q^2 \times \{\bot, \top\} \ \middle| \ \begin{matrix} (q_0, q', k') \in \mathtt{Cxt}(X, v), \mathcal{B} \colon q' \xrightarrow{a} q \\ (k = \bot \vee k' = \top \vee \mathcal{B} \colon q' \xrightarrow{a}_F q) \end{matrix} \right\} \ .$$

Intuitively $\mathtt{Cxt}$ can be computed inductively as we did for $\mathtt{Tgt}$. The first part of the condition defines how new context are obtained by appending a transition to the right of an existing context while the second part defines the bit of information keeping record of whether an accepting state was visited.

Bases, Frontier and Membership test. We stated in Remark 11 that the correctness of the FORQ-based algorithm is preserved when removing, from the computed sets, zero or more subsumed words for the corresponding ordering. In FORKLIFT, we remove all the subsumed words from all the sets we compute which, intuitively, means each computed set is a basis that contains as few words as possible. To remove subsumed words we leverage the target or context sets kept along with the words. It is worth pointing out that the least fixpoint computations at lines 4, 5, and 10 are implemented using a frontier. Finally, the ultimately periodic word membership procedure is implemented as a classical depth-first search as described in textbooks [17, Chapter 13.1.1].

Technical details. FORKLIFT, a naïve prototype implemented by a single person over several weeks, implements the algorithm of Fig. 2 with the structural FORQ in less than 1 000 lines of Java code. One of the design goals of our tool was to have simple code that could be easily integrated in other tools. Therefore, our implementation relies solely on a few standard packages from the Java SE Platform (notably collections such as `HashSet` or `HashMap`).

7.1   Experimental Evaluation

Benchmarks. Our evaluation uses benchmarks stemming from various application domains including benchmarks from theorem proving, software verification, and from previous work on the ω-regular language inclusion problem. In this section, a *benchmark* means an ordered pair of BAs such that the "left"/"right" BAs refer, resp., to the automata on the left/right of the inclusion sign. The BAs of the Pecan [31] benchmarks encode sets of solutions of predicates, hence a logical implication between predicates reduces to a language inclusion problem between BAs. The benchmarks correspond to theorems of type $\forall x, \exists y, \ P(x) \implies Q(y)$

about Sturmian words [21]. We collected 60 benchmarks from Pecan for which inclusion holds, where the BAs have alphabets of up to 256 symbols and have up to 21 395 states.

The second collection of benchmarks stems from software verification. The Ultimate Automizer (UA) [19,20] benchmarks encode termination problems for programs where the left BA models a program and the right BA its termination proof. Overall, we collected 600 benchmarks from UA for which inclusion holds for all but one benchmark. The BAs have alphabets of up to 13 173 symbols and are as large as 6 972 states.

The RABIT benchmarks are BAs modeling mutual exclusion algorithms [8], where in each benchmark one BA is the result of translating a set of guarded commands defining the protocol while the other BA translates a modified set of guarded commands, typically obtained by randomly weakening or strengthening one guard. The resulting BAs are on a binary alphabet and are as large as 7 963 states. Inclusion holds for 9 out of the 14 benchmarks.

All the benchmarks are publicly available on GitHub [12]. We used all the benchmarks we collected, that is, we discarded no benchmarks.

Tools. We compared FORKLIFT with the following tools: SPOT 2.10.3, GOAL (20200822), RABIT 2.5.0, ROLL 1.0, and BAIT 0.1.

SPOT [15,16] decides inclusion problems by complementing the "right" BA via determinization to parity automata with some additional optimizations including simulation-based optimizations. It is invoked through the command line tool `autfilt` with the option `--included-in`. It is worth pointing out that SPOT works with symbolic alphabets where symbols are encoded using Boolean propositions, and sets of symbols are represented and processed using OBDDs. SPOT is written in C++ and its code is publicly available [6].

GOAL [34] contains several language inclusion checkers available with multiple options. We used the Piterman algorithm using the options `containment -m piterman` with and without the additional options `-sim -pre`. In our plots GOAL is the invocation with the additional options `-sim -pre` which compute and use simulation relations to further improve performance while GOAL$^-$ is the one without the additional options. Inclusion is checked by constructing on-the-fly the intersection of the "left" BA and the complement of the "right" BA which is itself built on-the-fly by the Piterman construction [32]. The Piterman check was deemed the "best effort" (cf. [10, Section 9.1] and [33]) among the inclusion checkers provided in GOAL. GOAL is written in Java and the source code of the release we used is not publicly available [3].

RABIT [10] performs the following operations to check inclusion: (1) Removing dead states and minimizing the automata with simulation-based techniques, thus yielding a smaller instance; (2) Witnessing inclusion by simulation already during the minimization phase; (3) Using a Ramsey-based method with antichain heuristics to witness inclusion or non-inclusion. The antichain heuristics of Step (3) uses a unique quasiorder leveraging simulation relations to discard candidate counterexamples. In our experiments we

ran RABIT with options `-fast -jf` which RABIT states as providing the "best performance". RABIT is written in Java and is publicly available [4].

ROLL [24,25] contains an inclusion checker that does a preprocessing analogous to that of RABIT and then relies on automata learning and word sampling techniques to decide inclusion. ROLL is written in Java and is publicly available [5].

BAIT [13] which shares authors with the authors of the present paper, implements a Ramsey-based algorithm with the antichain heuristics where two quasiorders (one for the stems and the other for the periods) are used to discard candidate counterexamples as described in Section 1. BAIT is written in Java and is publicly available [1].

As far as we can tell all the above implementations, including FORKLIFT, are sequential except for RABIT which, using the `-jf` option, performs some computations in a separate thread.

Experimental Setup. We ran our experiments on a server with 24 GB of RAM, 2 Xeon E5640 2.6 GHz CPUs and Debian Stretch 64-bit. We used open-JDK 11.0.12 2021-07-20 when compiling Java code and ran the JVM with default options. For RABIT, BAIT and FORKLIFT the execution time is computed using timers internal to their implementations. For ROLL, GOAL and SPOT the execution time is given by the "real" value of the `time(1)` command. We preprocessed the benchmarks passed to FORKLIFT and BAIT with a reduction of the set of final states of the "left" BA that does not alter the language it recognizes. This preprocessing aims to minimize the number of iterations of the loop at line 6 of Fig. 2 over the set of final states. It is carried out by GOAL using the `acc -min` command. Internally, GOAL uses a polynomial time algorithm that relies on computing strongly connected components. The time taken by this preprocessing is negligible.

Plots. We use survival plots for displaying our experimental results in Fig. 3. Let us recall how to obtain them for a family of benchmarks $\{p_i\}_{i=1}^n$: (1) run the tool on each benchmark p_i and store its runtime t_i; (2) sort the t_i's in increasing order and discard pairs corresponding to abnormal program termination like time out or memory out; (3) plot the points $(t_1, 1), (t_1 + t_2, 2), \ldots$, and in general $(\sum_{i=1}^k t_i, k)$; (4) repeat for each tool under evaluation.

Survival plots are effective at comparing how tools scale up on benchmarks: the further right and the flatter a plot goes, the better the tool thereof scales up. Also the closer to the x-axis a plot is, the less time the tool needs to solve the benchmarks.

Analysis. It is clear from Fig. 3a and 3b that FORKLIFT scales up best on both the Pecan and UA benchmarks. FORKLIFT's scalability is particularly evident on the PECAN benchmarks of Fig. 3a where its curve is the flattest and no other tool finishes on all benchmarks. Note that, in Fig. 3b, the plot for SPOT is missing because we did not succeed into translating the UA benchmarks in the

input format of SPOT. On the UA benchmarks, FORKLIFT, BAIT and GOAL scale up well and we expect SPOT to scale up at least equally well. On the other hand, RABIT and ROLL scaled up poorly on these benchmarks.

On the RABIT benchmarks at Fig. 3c both FORKLIFT and SPOT terminate 13 out of 14 times; BAIT terminates 9 out of 14 times; and GOAL, ROLL and RABIT terminate all the times. We claim that the RABIT benchmarks can all be solved efficiently by leveraging simulation relations which FORKLIFT does not use let alone compute. Next, we justify this claim. First observe at Fig. 3c how GOAL is doing noticeably better than GOAL$^-$ while we have the opposite situation for the Pecan benchmarks Fig. 3a and no noticeable difference for the UA benchmarks Fig. 3b. Furthermore observe how ROLL and RABIT, which both leverage simulation relations in one way or another, scale up well on the RABIT benchmarks but scale up poorly on the PECAN and UA benchmarks.

The reduced RABIT benchmarks at Fig. 3d are obtained by pre-processing every BA of every RABIT benchmark with the simulation-based reduction operation of SPOT given by `autfilt --high --ba`. This preprocessing reduces the state space of the BAs by more than 90% in some cases. The reduction significantly improves how FORKLIFT scales up (it now terminates on all benchmarks) while it has less impact on RABIT, ROLL and SPOT which, as we said above, already leverage simulation relation internally. It is also worth noting that GOAL has a regression (from 14/14 before the reduction to 13/14).

Overall FORKLIFT, even though it is a prototype implementation, is the tool that returns most often (673/674). Its unique failure disappears after a preprocessing using simulation relations of the two BAs. The FORKLIFT curve for the Pecan benchmarks shows FORKLIFT scales up best.
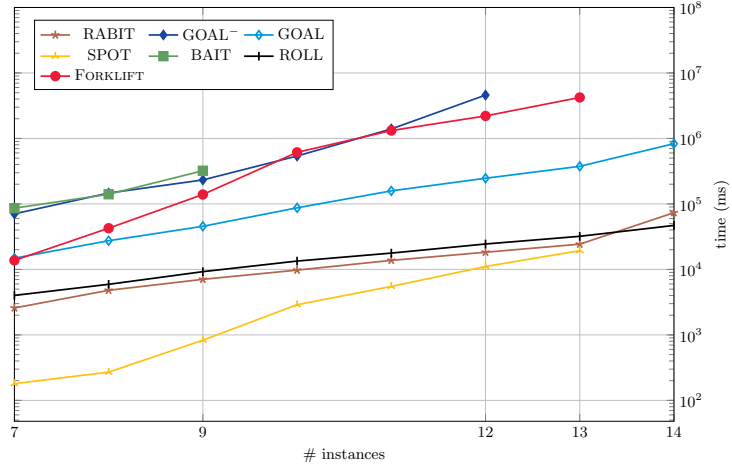
Our conclusion from the empirical evaluation is that, in practice FORKLIFT is competitive compared to the state-of-the-art in terms of scalability. Moreover


(a) Benchmarks from Pecan
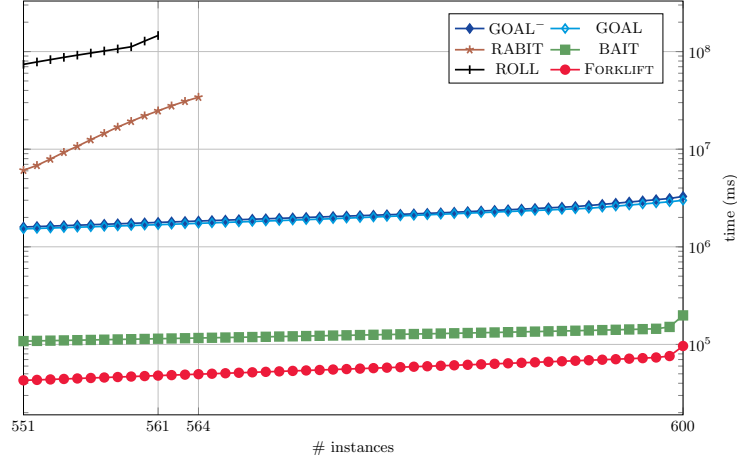
(b) Benchmarks from Ultimate Automizer

(c) Benchmarks from RABIT

(d) Benchmarks from RABIT (reduced)

Fig. 3: Survival plot with a logarithmic y axis and linear x axis. Each benchmark has a timeout value of 12h. Parts of the plots left out for clarity. A point is plotted for abscissa value x and tool r iff r returns with an answer for x benchmarks. All the failures of BAIT and the one of FORKLIFT are memory out.

the behavior of the FORQ-based algorithm in practice is far from its worst case exponential runtime.

8   Discussions

This section provides information that we consider of interest although not essential for the correctness of our algorithm or its evaluation.

Origin of FORQs. Our definition of FORQ and their suitability property (in particular the language preservation) are directly inspired from the definitions related to families of right congruences introduced by Maler and Staiger in 1993 [28] (revised in 2008 [30]). We now explain how our definition of FORQs generalizes and relaxes previous definitions [30, Definitions 5 and 6].

First we explain why the FORQ constraint does not appear in the setting of families of right congruences. In the context of congruences, relations are symmetric and thus, the FORQ constraint reduces to $u \lesssim u' \Rightarrow {\lll}_{u'} = {\lll}_u$. Therefore the FORQ constraint trivially holds if the set $\{{\lll}_u\}_{u \in \Sigma^*}$ is quotiented by the congruence relation $\lesssim$, which is the case in the definition [29, Definition 5].

Second, we point that the condition $v \lll_u v' \Rightarrow uv \lesssim uv'$ which appears in the definition for right families of congruences [30, Definition 5] is not needed in our setting. Nevertheless, this condition enables an improvement of the FORQ-based algorithm that we describe next.

Less membership queries. We put forward a property of structural FORQs allowing us to reduce the number of membership queries performed by FORK-LIFT. Hereafter, we refer to the *picky constraint* as the property of a FORQ stating $v \lll_u v' \Rightarrow uv \lesssim uv'$ where $u, v, v' \in \Sigma^*$. We first show how thanks to the picky constraint we can reduce the number of candidate counterexamples in the FORQ-based algorithm and then, we show that every structural FORQ satisfies the picky constraint.

In the algorithm of Fig. 2, periods are taken in a basis for the ordering $\lll_w$ where $w \in \Sigma^*$ belongs to a finite basis for the ordering $\lesssim^{-1}$. The only restriction on w is that of being comparable to the stem u, as ensured by the test at line 12. The following lemma formalizes the fact that we could consider a stronger restriction.

Lemma 12. *Let $\lesssim$ be a quasiorder over Σ^* such that $\lesssim^{-1}$ is a right-monotonic well-quasiorder. Let $S, S' \subseteq \Sigma^*$ be such that $\mathfrak{B}_{<^{-1}}(S', S)$ and S' contains no two distinct comparable words. For all $u \in \Sigma^*$ and $v \in \Sigma^+$ such that $u \in S$ and $\{wv \mid w \in S\} \subseteq S$, there exists $\mathring{w} \in S'$ such that $uv^i \lesssim \mathring{w}$ and $\mathring{w}v^j \lesssim \mathring{w}$ for some $i, j \in \mathbb{N} \setminus \{0\}$.*

As in Section 3, we show that the equivalence (†) holds but this time for an alternative definition of $T_{\mathcal{A}}$ we provide next. Given a M-suitable FORQ $\mathcal{F} \triangleq \langle \lesssim, \{{\lll}_u\}_{u \in \Sigma^*} \rangle$, let

$$\hat{T}_{\mathcal{A}} \triangleq \{uv^\omega \mid \exists s \in F_{\mathcal{A}} : u \in U_s, v \in V_s^w \text{ for some } w \in W_s \text{ with } u \lesssim w, wv \lesssim w\}$$

where for all $p \in P$ the sets U_p, W_p and $\{V_p^w\}_{w \in \Sigma^*}$ such that $\mathfrak{B}_{\lesssim}(U_p, \mathtt{Stem}_p)$, $\mathfrak{B}_{\lesssim^{-1}}(W_p, \mathtt{Stem}_p)$ and $\mathfrak{B}_{\lll_w}(V_p^w, \mathtt{Per}_p)$ for all $w \in \Sigma^*$. Since $\hat{T}_\mathcal{A} \subseteq T_\mathcal{A}$ by definition, it suffices to prove the implication $\hat{T}_\mathcal{A} \subseteq M \Rightarrow \mathtt{Ultim}_\mathcal{A} \subseteq M$. Let $uv^\omega \in \mathtt{Ultim}_\mathcal{A}$, i.e., such that there exists $s \in F_\mathcal{A}$ for which $u \in \mathtt{Stem}_s$ and $v \in \mathtt{Per}_s$, satisfying $uv \lesssim u$. In the context of Lemma 12, taking $S \triangleq \mathtt{Stem}_s$ and $S' \triangleq W_s$ fulfills the requirements $u \in S$ and $\{wv \mid w \in S\} \subseteq S$. We can thus apply the lemma and ensure the existence of some $w_0 \in W_s$ satisfying $uv^i \lesssim w_0$ and $w_0 v^j \lesssim w_0$ for some $i, j \in \mathbb{N} \setminus \{0\}$. Since $uv^i \in \mathtt{Stem}_s$ and $v^j \in \mathtt{Per}_s$ we find that there exist $u_0 \in U_s$ and $v_0 \in V_s^{w_0}$ such that $u_0 \lesssim uv^i$ and $v_0 \lll_{w_0} v^j$ thanks to the finite basis property. We conclude from above that $v_0 \lll_{w_0} v^j$, hence that $w_0 v_0 \lesssim w_0 v^j$ by the picky condition, and finally that $w_0 v_0 \lesssim w_0$ by Lemma 12 and transitivity. By definition $u_0 v_0^\omega \in \hat{T}_\mathcal{A}$ and the proof continues as the one in Section 3 for $T_\mathcal{A}$.

To summarize, if the considered FORQ fulfills the picky constraint then the algorithm of Fig. 2 remains correct when discarding the periods v at line 11 such that $wv \not\lesssim w$. Observe that discarding one period v possibly means skipping several membership queries $(u_1 v^\omega, u_2 v^\omega, \ldots)$. As proved below, the picky constraint holds for all structural FORQs.

Lemma 13. *Let $\mathcal{B} \triangleq (Q, q_I, \Delta_\mathcal{B}, F_\mathcal{B})$ be a BA and $\mathcal{F} \triangleq \langle \lesssim^\mathcal{B}, \{\lll_u^\mathcal{B}\}_{u \in \Sigma^*} \rangle$ its structural FORQ. For all $u \in \Sigma^*$ and all $v, v' \in \Sigma^+$ if $v \lll_u^\mathcal{B} v'$ then $uv \lesssim^\mathcal{B} uv'$.*

Proof. For all $q' \in \mathtt{Tgt}(uv)$, there exists $q \in Q$ such that $\mathcal{B} \colon q_I \xrightarrow{u} q \xrightarrow{v} q'$. Hence $(q, q', \bot) \in \mathtt{Cxt}(\mathtt{Tgt}(u), v)$. In fact $(q, q', \bot) \in \mathtt{Cxt}(\mathtt{Tgt}(u), v')$ holds as well since $v \lll_u^\mathcal{B} v'$. We deduce from the definition of $\mathtt{Cxt}$ that $\mathcal{B} \colon q_I \xrightarrow{u} q \xrightarrow{v'} q'$ which implies $q' \in \mathtt{Tgt}(uv')$. Thus $\mathtt{Tgt}(uv) \subseteq \mathtt{Tgt}(uv')$, i.e., $uv \lesssim^\mathcal{B} uv'$. $\square$

We emphasize that this reduction of the number of membership queries was not included in our experimental evaluation since (1) the proof of correctness is simpler and (2) FORKLIFT already scales up well without this optimization. We leave for future work the precise effect of such optimization.

Why a basis for $\lesssim^{-1}$ is computed? Taking periods in a basis for the ordering $\lll_w$ where $w \in \Sigma^*$ is picked in a basis for the ordering $\lesssim^{-1}$ may seem unnatural. In fact, the language preservation property of FORQs even suggests that an algorithm without computing a basis for $\lesssim^{-1}$ may exist. Here, we show that taking periods in a basis for the ordering $\lll_u$ where $u \in \Sigma^*$ is picked in a basis for the ordering $\lesssim$ is not correct. More precisely, redefining $T_\mathcal{A}$ as

$$\tilde{T}_\mathcal{A} \triangleq \{uv^\omega \mid \exists s \in F_\mathcal{A} \colon u \in U_s, v \in V_s^u\}$$

where for all $p \in P$ we have that $\mathfrak{B}_{\lesssim}(U_p, \mathtt{Stem}_p)$ and $\mathfrak{B}_{\lll_w}(V_p^w, \mathtt{Per}_p)$ for all $w \in \Sigma^*$, leads to an *incorrect* algorithm because the equivalence (†) given by $\tilde{T}_\mathcal{A} \subseteq M \iff L(\mathcal{A}) \subseteq M$ no longer holds as shown below in Example 14.

Example 14. Consider the BAs given by Fig. 1. We have that $L(\mathcal{A}) \not\subseteq L(\mathcal{B})$ and, in Example 3, we have argued that $T_\mathcal{A} = \{\varepsilon(b)^\omega, a(a)^\omega\}$ contains the ultimately

periodic a^ω which is a counterexample to inclusion. Recall from Example 3 and 6 that we can set $U_{p_I} = \{\varepsilon, a\}$ since $\mathfrak{B}_{\lesssim}(\{\varepsilon, a\}, \Sigma^*)$, and $V_{p_I}^a = V_{p_I}^\varepsilon = \{b\}$ since $\mathfrak{B}_{\lesssim_a}(\{b\}, \Sigma^+)$ and $\mathfrak{B}_{\lesssim_\varepsilon}(\{b\}, \Sigma^+)$. We conclude from the above definition that $\tilde{T}_\mathcal{A} = \{\varepsilon(b)^\omega, a(b)^\omega\}$, hence that $\tilde{T}_\mathcal{A} \subseteq L(\mathcal{B})$ which contradicts (†) since $L(\mathcal{A}) \nsubseteq L(\mathcal{B})$.

9 Conclusion and future work

We presented a novel approach to tackle in practice the language inclusion problem between Büchi automata. Our antichain heuristics is driven by the notion of FORQs that extends the notion of family of right congruences introduced in the nineties by Maler and Staiger [29]. We expect the notion of FORQs to have impact beyond the inclusion problem, e.g. in learning [9] and complementation [26]. A significant difference of our inclusion algorithm compared to other algorithms which rely on antichain heuristics, is the increased number of fixpoint computations that, counterintuitively, yield better scalability. Indeed our prototype FORKLIFT, which implements the FORQ-based algorithm, scales up well on benchmarks taken from real applications in verification and theorem proving.

In the future we want to increase further the search pruning capabilities of FORQs by enhancing them with simulation relations. We also plan to study whether FORQs can be extended to other settings like ω-visibly pushdown languages.

References

1. BAIT: An ω-regular language inclusion checker. https://github.com/parof/bait, accessed: 2022-01-17
2. FORKLIFT: Forq-based language inclusion formal testing. https://github.com/Mazzocchi/FORKLIFT, accessed: 2022-06-7
3. GOAL: Graphical tool for omega-automata and logics. http://goal.im.ntu.edu.tw/wiki/doku.php, accessed: 2022-01-17
4. RABIT/Reduce: Tools for language inclusion testing and reduction of nondeterministic Büchi automata and NFA. http://www.languageinclusion.org/doku.php?id=tools, accessed: 2022-01-17
5. ROLL library: Regular Omega Language Learning library. https://github.com/ISCAS-PMC/roll-library, accessed: 2022-01-17
6. Spot: a platform for LTL and ω-automata manipulation. https://spot.lrde.epita.fr/, accessed: 2022-01-17
7. Abdulla, P.A., Chen, Y., Clemente, L., Holík, L., Hong, C., Mayr, R., Vojnar, T.: Simulation subsumption in ramsey-based Büchi automata universality and inclusion testing. In: CAV: 22nd International Conference on Computer Aided Verification. LNCS (2010). https://doi.org/10.1007/978-3-642-14295-6_14
8. Abdulla, P.A., Chen, Y., Clemente, L., Holík, L., Hong, C., Mayr, R., Vojnar, T.: Advanced ramsey-based Büchi automata inclusion testing. In: CONCUR: 22nd International Conference on Concurrency Theory. LNCS (2011). https://doi.org/10.1007/978-3-642-23217-6_13

9. Angluin, D., Boker, U., Fisman, D.: Families of DFAs as Acceptors of ω-Regular Languages. Logical Methods in Computer Science **14** (2018). https://doi.org/10.23638/LMCS-14(1:15)2018

10. Clemente, L., Mayr, R.: Efficient reduction of nondeterministic automata with application to language inclusion testing. Log. Methods Comput. Sci. **15**(1) (2019). https://doi.org/10.23638/LMCS-15(1:12)2019

11. De Wulf, M., Doyen, L., Henzinger, T.A., Raskin, J.F.: Antichains: A new algorithm for checking universality of finite automata. In: CAV: 18th International Conference on Computer Aided Verification. LNCS (2006). https://doi.org/10.1007/11817963_5

12. Doveri, K., Ganty, P., Parolini, F., Ranzato, F.: Büchi Automata benchmarks for language inclusion. https://github.com/parof/buchi-automata-benchmark (2021)

13. Doveri, K., Ganty, P., Parolini, F., Ranzato, F.: Inclusion testing of Büchi automata based on well-quasiorders. In: CONCUR: 32nd International Conference on Concurrency Theory. LIPIcs (2021). https://doi.org/10.4230/LIPIcs.CONCUR.2021.3

14. Doyen, L., Raskin, J.F.: Antichains for the automata-based approach to model-checking. Logical Methods in Computer Science **5**(1) (2009). https://doi.org/10.2168/lmcs-5(1:5)2009

15. Duret-Lutz, A., Lewkowicz, A., Fauchille, A., Michaud, T., Renault, E., Xu, L.: Spot 2.0 - A framework for LTL and ω-automata manipulation. In: ATVA: 14th International Conference on Automated Technology for Verification and Analysis. LNCS (2016). https://doi.org/10.1007/978-3-319-46520-3_8

16. Duret-Lutz, A., Renault, E., Colange, M., Renkin, F., Aisse, A.G., Schlehuber-Cassier, P., Medioni, T., Martin, A., Dubois, J., Gillard, C., Lauko, H.: From Spot 2.0 to Spot 2.10: What's New? In: CAV: 34th International Conference on Computer Aided Verification. LNCS (2022)

17. Esparza, J.: Automata theory – An algorithmic approach (2017), https://www7.in.tum.de/~esparza/autoskript.pdf, Lecture Notes

18. Fogarty, S., Vardi, M.Y.: Efficient Büchi Universality Checking. In: TACAS: 16th International Conference on Tools and Algorithms for the Construction and Analysis of Systems. LNCS (2010). https://doi.org/10.1007/978-3-642-12002-2_17

19. Heizmann, M., Chen, Y., Dietsch, D., Greitschus, M., Hoenicke, J., Li, Y., Nutz, A., Musa, B., Schilling, C., Schindler, T., Podelski, A.: Ultimate automizer and the search for perfect interpolants – (competition contribution). In: TACAS: 24th International Conference on Tools and Algorithms for the Construction and Analysis of Systems. LNCS (2018). https://doi.org/10.1007/978-3-319-89963-3_30

20. Heizmann, M., Hoenicke, J., Podelski, A.: Software model checking for people who love automata. In: CAV: 25th International Conference on Computer Aided Verification. LNCS (2013). https://doi.org/10.1007/978-3-642-39799-8_2

21. Hieronymi, P., Ma, D., Oei, R., Schaeffer, L., Schulz, C., Shallit, J.: Decidability for Sturmian Words. In: CSL: 30th EACSL Annual Conference on Computer Science Logic. LIPIcs (2022). https://doi.org/10.4230/LIPIcs.CSL.2022.24

22. Kuperberg, D., Pinault, L., Pous, D.: Coinductive algorithms for Büchi automata. Fundam. Informaticae **180**(4) (2021). https://doi.org/10.3233/FI-2021-2046

23. Kupferman, O., Vardi, M.Y.: Verification of fair transition systems. In: CAV: International Conference on Computer Aided Verification. LNCS (1996). https://doi.org/10.1007/3-540-61474-5_84

24. Li, Y., Chen, Y.F., Zhang, L., Liu, D.: A novel learning algorithm for Büchi automata based on family of dfas and classification trees. Information and Computation (2020). https://doi.org/10.1016/j.ic.2020.104678

25. Li, Y., Sun, X., Turrini, A., Chen, Y., Xu, J.: ROLL 1.0: ω-regular language learning library. In: TACAS: 25th International Conference on Tools and Algorithms for the Construction and Analysis of Systems. LNCS (2019). https://doi.org/10.1007/978-3-030-17462-0_23

26. Li, Y., Tsay, Y.K., Turrini, A., Vardi, M.Y., Zhang, L.: Congruence relations for büchi automata. In: FM: Formal Methods. LNCS (2021)

27. de Luca, A., Varricchio, S.: Well quasi-orders and regular languages. Acta Informatica **31**(6) (1994). https://doi.org/10.1007/BF01213206

28. Maler, O., Staiger, L.: On syntactic congruences for ω-languages. In: STACS: 10th Annual Symposium on Theoretical Aspects of Computer Science. LNCS (1993). https://doi.org/10.1007/3-540-56503-5_58

29. Maler, O., Staiger, L.: On syntactic congruences for ω-languages. Theor. Comput. Sci. **183**(1) (1997). https://doi.org/10.1016/S0304-3975(96)00312-X

30. Maler, O., Staiger, L.: On syntactic congruences for ω-languages. Tech. rep., Verimag, France (2008), http://www-verimag.imag.fr/~maler/Papers/congr.pdf

31. Oei, R., Ma, D., Schulz, C., Hieronymi, P.: Pecan: An automated theorem prover for automatic sequences using Büchi automata. CoRR **abs/2102.01727** (2021), https://arxiv.org/abs/2102.01727

32. Piterman, N.: From nondeterministic Büchi and Streett automata to deterministic parity automata. Logical Methods in Computer Science **3**(3) (2007). https://doi.org/10.2168/lmcs-3(3:5)2007

33. Tsai, M., Fogarty, S., Vardi, M.Y., Tsay, Y.: State of Büchi complementation. Log. Methods Comput. Sci. **10**(4) (2014). https://doi.org/10.2168/LMCS-10(4:13)2014

34. Tsai, M., Tsay, Y., Hwang, Y.: GOAL for games, omega-automata, and logics. In: CAV: 25th International Conference on Computer Aided Verification. LNCS (2013). https://doi.org/10.1007/978-3-642-39799-8_62