

Lecture 1

*Lecturer: Pablo A. Parrilo**Scribe: Pablo A. Parrilo*

1 Introduction: what is this course about?

In this course we aim to understand the properties, both mathematical and computational, of sets defined by polynomial equations and inequalities. In particular, we want to work towards methods that will enable the solution (either exact or approximate) of optimization problems with feasible sets that are defined through polynomial systems. Needless to say (is it?), many problems in estimation, control, signal processing, etc., admit simple formulations in terms of polynomial equations and inequalities. However, these formulations can be tremendously difficult to solve, and thus our methods should try to exploit as many structural properties as possible.

The computational aspects of these sets are not very well understood at the moment. In the well-known case of polyhedra, for instance, there is a well defined relationship between the geometrical properties of the set (e.g., the number of facets, or the number of extreme points) and its algebraic representation. Furthermore, polyhedral sets are preserved by natural operations (e.g., projections). None of this will generally be true for (basic) semialgebraic sets, and this causes a very interesting interaction between their geometry and algebraic descriptions.

2 Topics

To understand better what is going on, we will embark in a journey to learn a wide variety of methods used to approach these problems. Some of our stops along the way will include:

- Linear optimization, second order cones, semidefinite programming
- Algebra: groups, fields, rings
- Univariate polynomials
- Resultants and discriminants
- Hyperbolic polynomials
- Sum of squares
- Ideals, varieties, Groebner bases, Nullstellensatz
- Quantifier elimination
- Real Nullstellensatz
- And much more...

We are interested in computational methods, and want to emphasize efficiency. Throughout, applications will play an important role, both as motivation and illustration of the techniques.

3 Review: convexity

A very important notion in modern optimization is that of *convexity*. To a large extent, modulo some (important) technicalities, there is a huge gap between the theoretical and practical solvability of optimization problems where the feasible set is convex, versus those where this property fails. Recommended presentations of convex optimization from a modern viewpoint are [BV04, BTN01, BNO03], with [Roc70] being the classical treatment of convex analysis.

Here are some relevant definitions:

Definition 1 A set S is *convex* if $x_1, x_2 \in S$ implies $\lambda x_1 + (1 - \lambda)x_2 \in S$ for all $0 \leq \lambda \leq 1$.

Definition 2 A set $S \subseteq \mathbb{R}^n$ is a *cone* if $\lambda \geq 0, x \in S \Rightarrow \lambda x \in S$.

Definition 3 The dual of a set S is $S^* = \{y \in \mathbb{R}^n : x \in S \Rightarrow \langle x, y \rangle \geq 0\}$.

A cone \mathcal{K} is *pointed* if $\mathcal{K} \cap (-\mathcal{K}) = \{0\}$, and *solid* if the interior of \mathcal{K} is not empty. A cone that is convex, closed, pointed and solid is called a *proper cone*. The dual set of a proper cone is also a proper cone, called the *dual cone*. An element x is in the interior of the cone K if and only if $\langle x, y \rangle > 0, \forall y \in K^*, y \neq 0$.

A proper cone induces a *partial order* in the space, via $x \preceq y$ if and only if $y - x \in \mathcal{K}$. We also use $x \prec y$ if $y - x$ is in the interior of \mathcal{K} . Important examples of proper cones are the nonnegative orthant, given by $\{x \in \mathbb{R}^n, x_i \geq 0\}$, and the set of symmetric positive semidefinite matrices.

4 Review: linear programming

Linear programming (LP) is the problem of minimizing a linear function, subject to linear inequality constraints. An LP in standard form is written as:

$$\min c^T x \quad \text{s.t.} \quad \begin{cases} Ax = b \\ x \geq 0 \end{cases} \quad (\text{P})$$

Every LP problem has a corresponding *dual* problem, which in this case is:

$$\max b^T y \quad \text{s.t.} \quad c - A^T y \geq 0. \quad (\text{D})$$

There are many important features of LP. Among them, we mention the following ones:

Geometry of the feasible set: The feasible set of linear programs are *polyhedra*. The geometry of polyhedra is quite well understood. In particular, the Minkowski-Weyl theorem (e.g., [BT97, Zie95]) states that every polyhedron P can be written as

$$P = \text{conv}(u_1, \dots, u_r) + \text{cone}(v_1, \dots, v_s),$$

where the u_i, v_i are the *extreme points* and *extreme rays* of P , respectively.

Weak duality: For any feasible solutions x, y of (P) and (D), respectively, it always holds that:

$$c^T x - b^T y = x^T c - (Ax)^T y = x^T (c - A^T y) \geq 0.$$

In other words, from any feasible dual solution we can obtain a lower bound on the primal. Conversely, primal feasible solutions give upper bounds on the value of the dual.

Strong duality: If both primal and dual are feasible, then they achieve exactly the same value, and there exist optimal feasible solutions x_*, y_* such that $c^T x_* = b^T y_*$.

Some of these properties (which ones?) will break down as soon as we leave LP and go the more general case of conic or semidefinite programming. These will cause some difficulties, although with the right assumptions, the resulting theory will closely parallel the LP case.

Remark The software codes `cdd` (Komei Fukuda, http://www.ifor.math.ethz.ch/~fukuda/cdd_home/index.html) and `lrs` (David Avis, <http://cgm.cs.mcgill.ca/~avis/C/lrs.html>) are very useful for polyhedral computations. In particular, both of them allow to convert an inequality representation of a polyhedron (usually called an H-representation) into extreme points/rays (V-representation), and viceversa.

References

- [BNO03] D.P. Bertsekas, A. Nedic, and A.E. Ozdaglar. *Convex analysis and optimization*. Athena Scientific, Belmont, MA, 2003.
- [BT97] D. Bertsimas and J. N. Tsitsiklis. *Introduction to Linear Optimization*. Athena Scientific, 1997.
- [BTN01] A. Ben-Tal and A. Nemirovski. *Lectures on modern convex optimization*. MPS/SIAM Series on Optimization. Society for Industrial and Applied Mathematics (SIAM), Philadelphia, PA, 2001.
- [BV04] S. Boyd and L. Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- [Roc70] R. T. Rockafellar. *Convex Analysis*. Princeton University Press, Princeton, New Jersey, 1970.
- [Zie95] G. M. Ziegler. *Lectures on polytopes*, volume 152 of *Graduate Texts in Mathematics*. Springer-Verlag, New York, 1995.

Lecture 2

Lecturer: Pablo A. Parrilo

Scribe: Pablo A. Parrilo

Notation: The set of real symmetric $n \times n$ matrices is denoted \mathcal{S}^n . A matrix $A \in \mathcal{S}^n$ is called *positive semidefinite* if $x^T A x \geq 0$ for all $x \in \mathbb{R}^n$, and is called *positive definite* if $x^T A x > 0$ for all nonzero $x \in \mathbb{R}^n$. The set of positive semidefinite matrices is denoted \mathcal{S}_+^n and the set of positive definite matrices is denoted by \mathcal{S}_{++}^n . The cone \mathcal{S}_+^n is a proper cone (i.e., closed, convex, pointed, and solid).

1 PSD matrices

There are several equivalent conditions for a matrix to be positive (semi)definite. We present below some of the most useful ones:

Proposition 1 *The following statements are equivalent:*

- The matrix $A \in \mathcal{S}^n$ is positive semidefinite ($A \succeq 0$).
- For all $x \in \mathbb{R}^n$, $x^T A x \geq 0$.
- All eigenvalues of A are nonnegative.
- All $2^n - 1$ principal minors of A are nonnegative.
- There exists a factorization $A = B^T B$.

For the definite case, we have a similar characterization:

Proposition 2 *The following statements are equivalent:*

- The matrix $A \in \mathcal{S}^n$ is positive definite ($A \succ 0$).
- For all nonzero $x \in \mathbb{R}^n$, $x^T A x > 0$.
- All eigenvalues of A are strictly positive.
- All n leading principal minors of A are positive.
- There exists a factorization $A = B^T B$, with B square and nonsingular.

Here are some useful additional facts:

- If T is nonsingular, $A \succ 0 \Leftrightarrow T^T A T \succ 0$.
- Schur complement. The following conditions are equivalent:

$$\begin{bmatrix} A & B \\ B^T & C \end{bmatrix} \succ 0 \Leftrightarrow \begin{cases} A \succ 0 \\ C - B^T A^{-1} B \succ 0 \end{cases} \Leftrightarrow \begin{cases} C \succ 0 \\ A - B C^{-1} B^T \succ 0 \end{cases}$$

2 Semidefinite programming

Semidefinite programming (SDP) is a specific kind of convex optimization problem (e.g., [VB96, Tod01, BV04]), with very appealing numerical properties. An SDP problem corresponds to the optimization of a linear function subject to matrix inequality constraints.

An SDP problem in standard primal form is written as:

$$\begin{aligned} & \text{minimize} && C \bullet X \\ & \text{subject to} && A_i \bullet X = b_i, \quad i = 1, \dots, m \\ & && X \succeq 0, \end{aligned} \tag{1}$$

where $C, A_i \in \mathcal{S}^n$, and $X \bullet Y := \text{Tr}(XY)$. The matrix $X \in \mathcal{S}^n$ is the variable over which the maximization is performed. The inequality in the second line means that the matrix X must be positive semidefinite, i.e., all its eigenvalues should be greater than or equal to zero. The set of feasible solutions, i.e., the set of matrices X that satisfy the constraints, is always a convex set. In the particular case in which $C = 0$, the problem reduces to whether or not the inequality can be satisfied for some matrix X . In this case, the SDP is referred to as a *feasibility problem*. The convexity of SDP has made it possible to develop sophisticated and reliable analytical and numerical methods to solve them.

A very important feature of SDP problems, from both the theoretical and applied viewpoints, is the associated *duality theory*. For every SDP of the form (1) (usually called the *primal problem*), there is another associated SDP, called the *dual problem*, that can be stated as

$$\begin{aligned} & \text{maximize} && b^T \mathbf{y} \\ & \text{subject to} && \sum_{i=1}^m A_i y_i \preceq C, \end{aligned} \tag{2}$$

where $b = (b_1, \dots, b_m)$, and the vector $\mathbf{y} = (y_1, \dots, y_m)$ contains the dual decision variables.

The key relationship between the primal and the dual problem is the fact that feasible solutions of one can be used to bound the values of the other problem. Indeed, let X and \mathbf{y} be any two feasible solutions of the primal and dual problems respectively. Then we have the following inequality:

$$C \bullet X - b^T \mathbf{y} = (C - \sum_{i=1}^m A_i y_i) \bullet X \geq 0, \tag{3}$$

where the last inequality follows from the fact that the two terms are positive semidefinite matrices. From (1) and (2) we can see that the left hand side of (3) is just the difference between the objective functions of the primal and dual problems. The inequality in (3) tells us that the value of the primal objective function evaluated at any feasible matrix X is always greater than or equal to the value of the dual objective function at any feasible vector \mathbf{y} . This property is known as *weak duality*. Thus, we can use any feasible X to compute an upper bound for the optimum of $b^T \mathbf{y}$, and we can also use any feasible \mathbf{y} to compute a lower bound for the optimum of $\text{Tr}(C \cdot X)$. Furthermore, in the case of feasibility problems (i.e., $C = 0$), the dual problem can be used to certify the nonexistence of solutions of the primal. This property will be crucial in our developments.

2.1 Conic duality

A general formulation, discussed briefly during the previous lecture, that unifies LP and SDP (as well as some other classes of optimization problems) is *conic programming*. We will be more careful than usual here (risking being a bit pedantic) in the definition of the respective spaces and mappings. It does not make much of a difference if we are working on \mathbb{R}^n (since we can identify a space and its dual), but it is “good hygiene” to keep these distinctions in mind, and also useful when dealing on more complicated spaces.

We will start with two real vector spaces, S and T , and a linear mapping $\mathcal{A} : S \rightarrow T$. Every real vector space has an associated dual space, which is the vector space of real-valued linear functionals. We will denote these dual spaces by S^* and T^* , respectively, and the pairing between an element of a vector space and one of the dual as $\langle \cdot, \cdot \rangle$ (i.e., $f(x) = \langle f, x \rangle$). The *dual mapping* of \mathcal{A} is the unique linear map $\mathcal{A}^* : T^* \rightarrow S^*$ defined through the property

$$\langle \mathcal{A}^* y, x \rangle_S = \langle y, \mathcal{A} x \rangle_T \quad \forall x \in S, y \in T^*.$$

Notice here that the brackets on the left-hand side of the equation represent the pairing in S , and those on the right-hand side correspond to the pairing in T . We can then define the primal-dual pair of (conic) optimization problems:

$$\min \langle c, x \rangle_S \quad \text{s.t.} \quad \begin{cases} \mathcal{A}x = b \\ x \in \mathcal{K} \end{cases} \quad \max \langle y, b \rangle_T \quad \text{s.t.} \quad c - \mathcal{A}^* y \in \mathcal{K}^*,$$

where $b \in T$, $c \in S^*$, $\mathcal{K} \subset S$ is a proper cone, and $\mathcal{K}^* \subset S^*$ is the corresponding dual cone. Notice that exactly the same proof presented earlier works here to show weak duality:

$$\begin{aligned} \langle c, x \rangle_S - \langle y, b \rangle_T &= \langle c, x \rangle_S - \langle y, \mathcal{A} x \rangle_T \\ &= \langle c, x \rangle_S - \langle \mathcal{A}^* y, x \rangle_S \\ &= \langle c - \mathcal{A}^* y, x \rangle_S \\ &\geq 0. \end{aligned}$$

In the usual cases (e.g., LP and SDP), the vector spaces are finite dimensional, and thus isomorphic to their duals. The specific correspondence between these is given through whatever inner product we use.

Among the classes of problems that can be interpreted as particular cases of the general conic formulation we have *linear programs*, *second-order cone programs* (SOCP), and SDP, when we take the cone \mathcal{K} to be the nonnegative orthant \mathbb{R}_+^n , the second order cone in n variables, or the PSD cone \mathcal{S}_+^n . We have then the following natural inclusion relationship among the different optimization classes.

$$\text{LP} \subseteq \text{SOCP} \subseteq \text{SDP}.$$

2.2 Geometric interpretation: separating hyperplanes

To be written.

2.3 Strong duality in SDP

We have seen that weak duality always holds for conic LP. As opposed to the LP case, strong duality can fail in SDP (and thus, in general conic programming). A nice example is given in [VB96, p. 65], where both the primal and dual problems are feasible, but their optimal values are different (i.e., there is a nonzero finite duality gap).

Nevertheless, under relatively mild constraint qualifications (Slater's condition, equivalent to the existence of strictly feasible primal and dual solutions) that usually satisfied in practice, SDP problems have strong duality, and thus zero duality gap.

There are several geometric interpretations of what causes the failure of strong duality for general SDP problems. A good one is based on the fact that the image of a proper cone under a linear transformation is not necessarily a proper cone. This fact seems quite surprising (or even wrong!) the first time one encounters it, but after a little while it starts being quite reasonable. Can you think of an example where this happens? What property will fail?

It should be mentioned that it is possible to formulate a more complicated SDP dual program (called the ‘‘Extended Lagrange-Slater Dual’’ in [Ram97]) for which strong duality always holds. For details, as well as a comparison with the more general ‘‘minimal cone’’ approach, we refer the reader to [Ram97, RTW97].

3 Applications

There have been *many* applications of SDP in a variety of areas of applied mathematics and engineering. We present here just a few, to give a flavor of what is possible. Many more will follow.

3.1 Lyapunov stability and control

Consider a linear difference equation (i.e., a discrete-time linear system) given by

$$x(k+1) = Ax(k), \quad x(0) = x_0.$$

It is well-known (and relatively simple to prove) that $x(k)$ converges to zero for all initial condition x_0 iff $|\lambda_i(A)| < 1, i = 1, \dots, n$.

There is a simple characterization of this spectral radius condition in terms of a quadratic *Lyapunov function* $V(x(k)) = x(k)^T P x(k)$:

$$|\lambda_i(A)| < 1 \quad \forall i \iff \exists P \succ 0 \quad A^T P A - P \prec 0$$

Proof

- (\Leftarrow) Let $Av = \lambda v$. Then,

$$0 > v^T (A^T P A - P) v = (|\lambda|^2 - 1) \underbrace{v^T P v}_{>0},$$

and therefore $|\lambda| < 1$

- (\Rightarrow) Let $P = \sum_{i=0}^{\infty} (A^i)^T Q A^i$, where $Q \succ 0$. The sum converges by the eigenvalue assumption. Then,

$$A^T P A - P = \sum_{i=1}^{\infty} (A^i)^T Q A^i - \sum_{i=0}^{\infty} (A^i)^T Q A^i = -Q \prec 0$$

■

Consider now the case where A is not stable, but we can use linear state feedback, i.e., $A(K) = A + BK$, where K is a fixed matrix. We want to find a matrix K such that $A + BK$ is stable, i.e., all its eigenvalues have absolute value smaller than one.

Use Schur complements to rewrite the condition:

$$(A + BK)^T P (A + BK) - P \prec 0, \quad P \succ 0$$

$$\Updownarrow$$

$$\begin{bmatrix} P & (A + BK)^T P \\ P(A + BK) & P \end{bmatrix} \succ 0$$

Condition is nonlinear in (P, K) . However, we can do a congruence transformation with $Q := P^{-1}$, and obtain:

$$\begin{bmatrix} Q & Q(A + BK)^T \\ (A + BK)Q & Q \end{bmatrix} \succ 0$$

Now, defining a new variable $Y := KQ$ we have

$$\begin{bmatrix} Q & QA^T + Y^T B^T \\ AQ + BY & Q \end{bmatrix} \succ 0.$$

This problem is now linear in (Q, Y) . In fact, it is an SDP problem. After solving it, we can recover the controller K via $K = Q^{-1}Y$.

3.2 Theta function

Given a graph $G = (V, E)$, a *stable set* is a subset of V with the property that the induced subgraph has no edges. In other words, none of the selected vertices are adjacent to each other.

The *stability number* of a graph, usually denoted by $\alpha(G)$, is the cardinality of the largest stable set. Computing the stability number of a graph is NP-hard. There are many interesting applications of the stable set problem. In particular, they can be used to provide bounds on the *Shannon capacity of a graph* [Lov79], a problem of importance in coding. In fact, this was one of the first appearances of what today is known as SDP.

The Lovász theta function is denoted by $\vartheta(G)$, and is defined as the solution of the SDP :

$$\max J \bullet X \quad \text{s.t.} \quad \begin{cases} \text{Tr}(X) = 1 \\ X_{ij} = 0 & (i, j) \in E \\ X \succeq 0 \end{cases} \quad (4)$$

where J is the matrix with all entries equal to one. The theta function is an upper bound on the stability number, i.e.,

$$\alpha(G) \leq \vartheta(G).$$

The inequality is easy to prove. Consider the indicator vector $\xi(S)$ of any stable set S , and define the matrix $X := \frac{1}{|S|} \xi \xi^T$. It is easy to see that this X is a feasible solution of the SDP, and thus the inequality follows.

4 Software

Remark There are many good software codes for semidefinite programming. Among the most well-known, we mention the following ones:

- SeDuMi, originally by Jos Sturm, now being maintained by the optimization group at McMaster: <http://sedumi.mcmaster.ca/>
- SDPT3, by Kim-Chuan Toh, Reha Tütüncü, and Mike Todd. <http://www.math.nus.edu.sg/~mattohkc/sdpt3.html>
- SDPA, by the research group of Masakazu Kojima, <http://grid.r.dendai.ac.jp/sdpa/>
- CSDP, by Brian Borchers, <http://infohost.nmt.edu/~borchers/csdp.html>

A very convenient way of using these (and other) SDP solvers under MATLAB is through the YALMIP parser/solver (Johan Löfberg, <http://control.ee.ethz.ch/~joloef/yalmip.php>).

References

- [BV04] S. Boyd and L. Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- [Lov79] L. Lovász. On the Shannon capacity of a graph. *IEEE Transactions on Information Theory*, 25(1):1–7, 1979.
- [Ram97] M. V. Ramana. An exact duality theory for semidefinite programming and its complexity implications. *Math. Programming*, 77(2, Ser. B):129–162, 1997.
- [RTW97] M. V. Ramana, L. Tunçel, and H. Wolkowicz. Strong duality for semidefinite programming. *SIAM J. Optim.*, 7(3):641–662, 1997.

- [Tod01] M. Todd. Semidefinite optimization. *Acta Numerica*, 10:515–560, 2001.
- [VB96] L. Vandenberghe and S. Boyd. Semidefinite programming. *SIAM Review*, 38(1):49–95, March 1996.

Lecture 3

Lecturer: Pablo A. Parrilo

Scribe: Pablo A. Parrilo

In this lecture, we will discuss one of the most important applications of semidefinite programming, namely its use in the formulation of convex relaxations of nonconvex optimization problems. We will present the results from several different, but complementary, points of view. These will also serve us as starting points for the generalizations to be presented later in the course.

We will discuss first the case of binary quadratic optimization, since in this case the notation is simpler, and perfectly illustrates many of the issues appearing in more complicated problems. Afterwards, a more general formulation containing arbitrary linear and quadratic constraints will be presented.

1 Binary optimization

Binary (or Boolean) quadratic optimization is a classical combinatorial optimization problem. In the version we consider, we want to minimize a quadratic function, where the decision variables can only take the values ± 1 . In other words, we are minimizing an (indefinite) quadratic form over the vertices of an n -dimensional hypercube. The problem is formally expressed as:

$$\begin{aligned} \text{minimize} \quad & x^T Q x \\ \text{s.t.} \quad & x_i \in \{-1, 1\} \end{aligned} \tag{1}$$

where $Q \in \mathcal{S}^n$. There are many well-known problems that can be naturally written in the form above. Among these, we mention the maximum cut problem (MAXCUT) discussed below, the 0-1 knapsack, the linear quadratic regulator (LQR) control problem with binary inputs, etc.

Notice that we can model the Boolean constraints using quadratic equations, i.e.,

$$x_i^2 - 1 = 0 \iff x_i \in \{-1, 1\}.$$

These n quadratic equations define a finite set, with an exponential number of elements, namely all the n -tuples with entries in $\{-1, 1\}$. There are exactly 2^n points in this set, so a direct enumeration approach to (1) is computationally prohibitive when n is large (already for $n = 30$, we have $2^n \approx 10^9$).

We can thus write the equivalent polynomial formulation:

$$\begin{aligned} \text{minimize} \quad & x^T Q x \\ \text{s.t.} \quad & x_i^2 = 1 \end{aligned} \tag{2}$$

We will denote the optimal value and optimal solution of this problem as f_\star and x_\star , respectively. It is well-known that the decision version of this problem is *NP-complete* (e.g., [GJ79]). Notice that this is true even if the matrix Q is positive definite (i.e., $Q \succ 0$), since we can always make Q positive definite by adding to it a constant multiple of the identity (this only shifts the objective by a constant).

Example 1 (MAXCUT) *The maximum cut (MAXCUT) problem consists in finding a partition of the nodes of a graph $G = (V, E)$ into two disjoint sets V_1 and V_2 ($V_1 \cap V_2 = \emptyset$, $V_1 \cup V_2 = V$), in such a way to maximize the number of edges that have one endpoint in V_1 and the other in V_2 . It has important practical applications, such as optimal circuit layout. The decision version of this problem (does there exist a cut with value greater than or equal to K ?) is NP-complete [GJ79].*

We can easily rewrite the MAXCUT problem as a binary optimization problem. A standard formulation (for the weighted problem) is the following:

$$\max_{y_i \in \{-1, 1\}} \frac{1}{4} \sum_{i,j} w_{ij} (1 - y_i y_j), \tag{3}$$

where w_{ij} is the weight corresponding to the (i, j) edge, and is zero if the nodes i and j are not connected. The constraints $y_i \in \{-1, 1\}$ are equivalent to the quadratic constraints $y_i^2 = 1$.

We can easily convert the MAXCUT formulation into binary quadratic programming. Removing the constant term, and changing the sign, the original problem is clearly equivalent to:

$$\min_{y_i^2=1} \sum_{i,j} w_{ij} y_i y_j. \quad (4)$$

1.1 Semidefinite relaxations

Computing “good” solutions to the binary optimization problem given in (2) is a quite difficult task, so it is of interest to produce accurate bounds on its optimal value. As in all minimization problems, *upper bounds* can be directly obtained from feasible points. In other words, if $x_0 \in \mathbb{R}^n$ has entries equal to ± 1 , it always holds that $f_* \leq x_0^T Q x_0$ (of course, for a poorly chosen x_0 , this upper bound may be very loose).

To prove *lower bounds*, we need a different technique. There are several approaches to do this, but as we will see in detail in the next sections, many of them will turn out to be exactly equivalent in the end. Indeed, many of these different approaches will yield a characterization of a lower bound in terms of the following primal-dual pair of semidefinite programming problems:

minimize $\text{Tr } QX$ s.t. $X_{ii} = 1$ $X \succeq 0$	maximize $\text{Tr } \Lambda$ s.t. $Q \succeq \Lambda$ Λ diagonal	(5)
--	---	-----

In the next sections, we will derive these SDPs several times, in a number of different ways. Let us notice here first that for this primal-dual pair of SDP, strong duality always holds, and both achieve their corresponding optimal solutions (why?).

1.2 Lagrangian duality

A general approach to obtain lower bounds on the value of general (non)convex minimization problems is to use Lagrangian duality. As we have seen the original Boolean minimization problem can be written as:

$$\begin{aligned} &\text{minimize} && x^T Q x \\ &\text{s.t.} && x_i^2 - 1 = 0 \end{aligned} \quad (6)$$

For notational convenience, let $\Lambda = \text{diag}(\lambda_1, \dots, \lambda_n)$. Then, the Lagrangian function can be written as:

$$L(x, \lambda) = x^T Q x - \sum_{i=1}^n \lambda_i (x_i^2 - 1) = x^T (Q - \Lambda) x + \text{Tr } \Lambda.$$

For the dual function $g(\lambda) := \inf_x L(x, \lambda)$ to be bounded below, we need the implicit constraint that the matrix $Q - \Lambda$ must be positive semidefinite. In this case, the optimal value of x is zero, and thus we obtain a lower bound given by the solution of the SDP:

$$\begin{aligned} &\text{maximize} && \text{Tr } \Lambda \\ &\text{s.t.} && Q - \Lambda \succeq 0 \end{aligned} \quad (7)$$

This is exactly the dual side of the SDP in (5).




Figure 1: The ellipsoids \mathcal{E}_1 and \mathcal{E}_2 .

1.3 Underestimator of the objective

A different but related interpretation of the SDP relaxation (5) is through the notion of an *underestimator* of the objective function. Indeed, the quadratic function $x^T \Lambda x$ is an “easily optimizable” function that is guaranteed to lie below the desired objective $x^T Q x$. To see this, notice that for any feasible x we have

$$x^T Q x \geq x^T \Lambda x = \sum_{i=1}^n \Lambda_{ii} x_i^2 = \text{Tr } \Lambda,$$

where

- The first inequality follows from $Q \succeq \Lambda$
- The second equation holds since the matrix Λ is diagonal
- Finally, the third one holds since $x_i \in \{+1, -1\}$

There is also a nice corresponding geometric interpretation. For simplicity, we assume without loss of generality that Q is positive definite. Then, the problem (2) can be interpreted as finding the largest value of γ for which the ellipsoid $\{x \in \mathbb{R}^n \mid x^T Q x \leq \gamma\}$ does not contain a vertex of the unit hypercube.

Consider now the two ellipsoids in \mathbb{R}^n defined by:

$$\begin{aligned} \mathcal{E}_1 &= \{x \in \mathbb{R}^n \mid x^T Q x \leq \text{Tr } \Lambda\} \\ \mathcal{E}_2 &= \{x \in \mathbb{R}^n \mid x^T \Lambda x \leq \text{Tr } \Lambda\}. \end{aligned}$$

The principal axes of ellipsoid \mathcal{E}_2 are aligned with the coordinates axes (since Λ is diagonal), and furthermore its boundary contains all the vertices of the unit hypercube. Also, it is easy to see that the condition $Q \succeq \Lambda$ implies $\mathcal{E}_1 \subseteq \mathcal{E}_2$.

With these facts, it is easy to understand the related problem that the SDP relaxation is solving: dilating \mathcal{E}_1 as much as possible, while ensuring the existence of another ellipsoid \mathcal{E}_2 with coordinate-aligned axes and touching the hypercube in all 2^n vertices; see Figure 1 for an illustration.

1.4 Probabilistic interpretation

To be written

ToDo




Figure 2: The three-dimensional “spectraplex.” This is the set of 3×3 positive semidefinite matrices, with unit diagonal.

1.5 Lifting and rank relaxation

We present yet another derivation of the SDP relaxations, this time focused on the primal side. Recall the original formulation of the optimization problem (2). Define now $X := xx^T$. By construction, the matrix $X \in \mathcal{S}^n$ satisfies $X \succeq 0$, $X_{ii} = x_i^2 = 1$, and has *rank one*. Conversely, any matrix X with

$$X \succeq 0, \quad X_{ii} = 1, \quad \text{rank } X = 1$$

necessarily has the form $X = xx^T$ for some ± 1 vector x (why?). Furthermore, by the cyclic property of the trace, we can express the objective function directly in terms of the matrix X , via:

$$x^T Q x = \text{Tr } x^T Q x = \text{Tr } Q x x^T = \text{Tr } Q X.$$

As a consequence, the original problem (2) can be exactly rewritten as:

$$\begin{aligned} & \text{minimize} && \text{Tr } Q X \\ & \text{s.t.} && X_{ii} = 1, \quad X \succeq 0 \\ & && \text{rank}(X) = 1 \end{aligned}$$

This is almost an SDP problem (all the constraints are either linear or conic), except for the rank one constraint on X . Since this is a minimization problem, a lower bound on the solution can be obtained by dropping the (nonconvex) rank constraint, which enlarges the feasible set.

A useful interpretation is in terms of a nonlinear *lifting* to a higher dimensional space. Indeed, rather than solving the original problem in terms of the n -dimensional vector x , we are instead solving for the $n \times n$ matrix X , effectively converting the problem from \mathbb{R}^n to \mathcal{S}^n (which has dimension $\binom{n+1}{2}$).

Observe that this line of reasoning immediately shows that if we find an optimal solution X of the SDP (5) that has rank one, then we have solved the original problem. Indeed, in this case the upper and lower bounds on the solution coincide.

As a graphical illustration, in Figure 2 we depict the set of 3×3 positive semidefinite matrices of unit diagonal. The rank one matrices correspond to the four “vertices” of this convex set, and are in (two-to-one) correspondence with the eight 3-vectors with ± 1 entries.

In general, it is not the case that the optimal solution of the SDP relaxation will be rank one. However, as we will see in the next section, it is possible to use *rounding schemes* to obtain “nearby” rank one solutions. Furthermore, in some cases, it is possible to do so while obtaining some approximation guarantees on the quality of the rounded solutions.

2 Bounds: Goemans-Williamson and Nesterov

So far, our use of the SDP relaxation (5) has been limited to providing only *a posteriori* bounds on the optimal solution of the original minimization problem. However, two desirable features are missing:

- Approximation guarantees: is it possible to prove general properties on the quality of the bounds obtained by SDP?
- Feasible solutions: can we (somehow) use the SDP relaxations to provide not just bounds, but actual feasible points with good (or optimal) values of the objective?

As we will see, it turns out that both questions can be answered in the positive. As it has been shown by Goemans and Williamson [GW95] in the MAXCUT case, and Nesterov in a more general setting, we can actually achieve both of these objectives by randomly “rounding” in an appropriate manner the solution X of this relaxation. We discuss these results below.

2.1 Goemans and Williamson rounding

In their celebrated MAXCUT paper, Goemans and Williamson developed the following randomized method for finding a “good” feasible cut from the solution of the SDP.

- Factorize X as $X = V^T V$, where $V = [v_1 \dots v_n] \in \mathbb{R}^{r \times n}$, where r is the rank of X .
- Then $X_{ij} = v_i^T v_j$, and since $X_{ii} = 1$ this factorization gives n vectors v_i on the unit sphere in \mathbb{R}^r .
- Instead of assigning either 1 or -1 to each variable, we have assigned to each a point on the unit sphere in \mathbb{R}^r .
- Now, choose a random hyperplane in \mathbb{R}^r , and assign to each variable x_i either a $+1$ or a -1 , depending on which side of the hyperplane the point v_i lies.

It turns out that this procedure gives a solution that, on average, is quite close to the value of the SDP bound. We will compute the expected value of the rounded solution in a slightly different form from the original G-W argument, but one that will be helpful later. The random hyperplane can be characterized by its normal vector p , which is chosen to be uniformly distributed on the unit sphere (e.g., by suitably normalizing a standard multivariate Gaussian random variable). Then, according to the description above, the rounded solution is given by $x_i = \text{sign}(p^T v_i)$. The expected value of this solution can then be written as:

$$\mathbf{E}_p[x^T Q x] = \sum_{ij} Q_{ij} \mathbf{E}_p[x_i x_j] = \sum_{ij} Q_{ij} \mathbf{E}_p[\text{sign}(p^T v_i) \text{sign}(p^T v_j)].$$

We can easily compute the value of this expectation. Consider the plane spanned by v_i and v_j , and let θ_{ij} be the angle between these two vectors. Then, it is easy to see that the desired expectation is equal to the probability that both points are on the same side of the hyperplane, minus the probability that they are on different sides. These probabilities are $1 - \frac{\theta_{ij}}{\pi}$ and $\frac{\theta_{ij}}{\pi}$, respectively. Thus, the expected value of the rounded solution is exactly:

$$\sum_{ij} Q_{ij} \left(1 - \frac{2\theta_{ij}}{\pi}\right) = \sum_{ij} Q_{ij} \left(1 - \frac{2}{\pi} \arccos(v_i^T v_j)\right) = \frac{2}{\pi} \sum_{ij} Q_{ij} \arcsin X_{ij}. \quad (8)$$

Notice that the expression is of course well-defined, since if X is PSD and has unit diagonal, all its entries are bounded in absolute value by 1. This result exactly characterizes the expected value of the rounding procedure, as a function of the optimal solution of the SDP. We would like, however, to directly relate this quantity to the optimal solution of the original optimization problem. For this, we will need additional assumptions on the matrix Q . We discuss next two of the most important results in this direction.




Figure 3: Bound on the inverse cosine function, for $\alpha \approx 0.878$.

2.2 MAXCUT bound

Recall from (3) that for the MAXCUT problem, the objective function does not only include the quadratic part, but there is actually a constant term:

$$\frac{1}{4} \sum_{ij} w_{ij} (1 - y_i y_j).$$

The expected value of the cut is then:

$$c_{\text{sdp-expected}} = \frac{1}{4} \sum_{ij} w_{ij} \left(1 - \frac{2}{\pi} \arcsin X_{ij} \right) = \frac{1}{4} \cdot \frac{2}{\pi} \sum_{ij} w_{ij} \arccos X_{ij}.$$

On the other hand, the solution of the SDP gives an upper bound on the cut capacity equal to:

$$c_{\text{sdp-upper-bound}} = \frac{1}{4} \sum_{ij} w_{ij} (1 - X_{ij}).$$

To relate these two quantities, we look for a constant α such that

$$\alpha (1 - t) \leq \frac{2}{\pi} \arccos(t) \quad \text{for all } t \in [-1, 1]$$

The best possible (i.e., largest) such constant is $\alpha = 0.878$; see Figure 3. So we have

$$c_{\text{sdp-upper-bound}} \leq \frac{1}{\alpha} \cdot \frac{1}{4} \cdot \frac{2}{\pi} \sum_{ij} w_{ij} \arccos X_{ij} = \frac{1}{\alpha} c_{\text{sdp-expected}}$$

Notice that here we have used the nonnegativity of the weights (i.e., $w_{ij} \geq 0$). Thus, so far we have the following inequalities:

- $c_{\text{sdp-upper-bound}} \leq \frac{1}{\alpha} c_{\text{sdp-expected}}$
- Also clearly $c_{\text{sdp-expected}} \leq c_{\text{max}}$
- And $c_{\text{max}} \leq c_{\text{sdp-upper-bound}}$

Putting it all together, we can sandwich the value of the relaxation as follows:

$$\alpha \cdot c_{\text{sdp-upper-bound}} \leq c_{\text{sdp-expected}} \leq c_{\text{max}} \leq c_{\text{sdp-upper-bound}}.$$

2.3 Nesterov's $\frac{2}{\pi}$ result

A result by Nesterov generalizes the MAXCUT bound described above, but for a larger class of problems. The original formulation is for the case of binary *maximization*, and applies to the case when the matrix A is *positive semidefinite*. Since the problem is homogeneous, the optimal value is guaranteed to be nonnegative.

As we have seen, the expected value of the solution after randomized rounding is given by (8). Since X is positive semidefinite, it follows from the nonnegativity of the Taylor series of $\arcsin(t) - t$ and the Schur product theorem that

$$\arcsin[X] \succeq X,$$

where the \arcsin function is applied componentwise. This inequality can be combined with (8) to give the bounds:

$$\frac{2}{\pi} \cdot f_{\text{sdp-upper-bound}} \leq f_{\text{sdp-expected}} \leq f_{\text{max}} \leq f_{\text{sdp-upper-bound}},$$

where $2/\pi \approx 0.636$. For more details, see [BTN01, Section 4.3.4]. Among others, the paper [Meg01] presents several new results, as well as a review of many of the available approximation schemes.

3 Linearly constrained problems

In this section we extend the earlier results, to general quadratic optimization problems under linear and quadratic constraints. For notational simplicity, we write the constraints in homogeneous form, i.e., in terms of the vector $x = [1 \ y^T]^T$.

The general primal form of the SDP optimization problems we are concerned with is

$$\begin{aligned} \min \quad & x^T Q x \\ \text{s.t.} \quad & x^T A_i x \geq 0 \\ & Bx \geq 0 \\ & x = \begin{bmatrix} 1 \\ y \end{bmatrix} \end{aligned}$$

The corresponding primal and dual SDP relaxations are given by

$\begin{aligned} \min \quad & Q \bullet X \\ \text{s.t.} \quad & A_i \bullet X \geq 0 \\ & BXB^T \geq 0 \\ & E_{11} \bullet X = 1 \\ & X \succeq 0 \end{aligned}$	$\begin{aligned} \max \quad & \gamma \\ \text{s.t.} \quad & Q \succeq \gamma E_{11} + \sum_i \lambda_i A_i + B^T N B \\ & \lambda_i \geq 0 \\ & N \geq 0 \\ & N_{ii} = 0 \end{aligned} \tag{9}$
---	---

Here E_{11} denotes the matrix with a 1 on the (1,1) component, and all the rest being zero. The dual variables λ_i can be interpreted as Lagrange multipliers associated to the quadratic constraints of the primal problem, while N corresponds to pairwise products of the linear constraints.

References

- [BTN01] A. Ben-Tal and A. Nemirovski. *Lectures on modern convex optimization*. MPS/SIAM Series on Optimization. Society for Industrial and Applied Mathematics (SIAM), Philadelphia, PA, 2001.

- [GJ79] M. R. Garey and D. S. Johnson. *Computers and Intractability: A guide to the theory of NP-completeness*. W. H. Freeman and Company, 1979.
- [GW95] M. X. Goemans and D. P. Williamson. Improved approximation algorithms for maximum cut and satisfiability problems using semidefinite programming. *Journal of the ACM*, 42(6):1115–1145, 1995.
- [Meg01] A. Megretski. Relaxations of quadratic programs in operator theory and system analysis. In *Systems, approximation, singular integral operators, and related topics (Bordeaux, 2000)*, volume 129 of *Oper. Theory Adv. Appl.*, pages 365–392. Birkhäuser, Basel, 2001.

Lecture 4

Lecturer: Pablo A. Parrilo

Scribe: Pablo A. Parrilo

In this lecture we will review some basic elements of abstract algebra. We also introduce and begin studying the main objects of our considerations, multivariate polynomials.

1 Review: groups, rings, fields

We present here standard background material on abstract algebra. Most of the definitions are from [Lan71, CLO97, DF91, BCR98].

Definition 1 A group consists of a set G and a binary operation “ \cdot ” defined on G , for which the following conditions are satisfied:

1. *Associative:* $(a \cdot b) \cdot c = a \cdot (b \cdot c)$, for all $a, b, c \in G$.
2. *Identity:* There exist $1 \in G$ such that $a \cdot 1 = 1 \cdot a = a$, for all $a \in G$.
3. *Inverse:* Given $a \in G$, there exists $b \in G$ such that $a \cdot b = b \cdot a = 1$.

For example, the integers \mathbb{Z} form a group under addition, but not under multiplication. Another example is the set $GL(n, \mathbb{R})$ of real nonsingular $n \times n$ matrices, under matrix multiplication.

If we drop the condition on the existence of an inverse, we obtain a *monoid*. Note that a monoid always has at least one element, the identity. As an example, given a set S , then the set of all strings of elements of S is a monoid, where the monoid operation is string concatenation and the identity is the empty string λ . Another example is given by \mathbb{N}_0 , with the operation being addition (in this case, the identity is the zero). Monoids are also known as *semigroups with identity*.

In a group we only have one binary operation (“multiplication”). We will introduce another operation (“addition”), and study the structure that results from their interaction.

Definition 2 A commutative ring (with identity) consists of a set k and two binary operations “ \cdot ” and “ $+$ ”, defined on k , for which the following conditions are satisfied:

1. *Associative:* $(a + b) + c = a + (b + c)$ and $(a \cdot b) \cdot c = a \cdot (b \cdot c)$, for all $a, b, c \in k$.
2. *Commutative:* $a + b = b + a$ and $a \cdot b = b \cdot a$, for all $a, b \in k$.
3. *Distributive:* $a \cdot (b + c) = a \cdot b + a \cdot c$, for all $a, b, c \in k$.
4. *Identities:* There exist $0, 1 \in k$ such that $a + 0 = a \cdot 1 = a$, for all $a \in k$.
5. *Additive inverse:* Given $a \in k$, there exists $b \in k$ such that $a + b = 0$.

A simple example of a ring are the integers \mathbb{Z} under the usual operations. After formally introducing polynomials, we will see a few more examples of rings.

If we add a requirement for the existence of multiplicative inverses, we obtain *fields*.

Definition 3 A field consists of a set k and two binary operations “ \cdot ” and “ $+$ ”, defined on k , for which the following conditions are satisfied:

1. *Associative:* $(a + b) + c = a + (b + c)$ and $(a \cdot b) \cdot c = a \cdot (b \cdot c)$, for all $a, b, c \in k$.
2. *Commutative:* $a + b = b + a$ and $a \cdot b = b \cdot a$, for all $a, b \in k$.
3. *Distributive:* $a \cdot (b + c) = a \cdot b + a \cdot c$, for all $a, b, c \in k$.

4. *Identities:* There exist $0, 1 \in k$, where $0 \neq 1$, such that $a + 0 = a \cdot 1 = a$, for all $a \in k$.
5. *Additive inverse:* Given $a \in k$, there exists $b \in k$ such that $a + b = 0$.
6. *Multiplicative inverse:* Given $a \in k, a \neq 0$, there exists $c \in k$ such that $a \cdot c = 1$.

Any field is obviously a commutative ring. Some commonly used fields are the rationals \mathbb{Q} , the reals \mathbb{R} and the complex numbers \mathbb{C} . There are also Galois or finite fields (the set k has a finite number of elements), such as \mathbb{Z}_p , the set of integers modulo p , where p is a prime. Another important field is given by $k(x_1, \dots, x_n)$, the set of *rational functions* with coefficients in the field k , with the natural operations.

2 Polynomials and ideals

Consider a given field k , and let x_1, \dots, x_n be indeterminates. We can then define *polynomials*.

Definition 4 A polynomial f in x_1, \dots, x_n with coefficients in a field k is a finite linear combination of monomials:

$$f = \sum_{\alpha} c_{\alpha} x^{\alpha} = \sum_{\alpha} c_{\alpha} x_1^{\alpha_1} \dots x_n^{\alpha_n}, \quad c_{\alpha} \in k, \quad (1)$$

where the sum is over a finite number of n -tuples $\alpha = (\alpha_1, \dots, \alpha_n)$, $\alpha_i \in \mathbb{N}_0$. The set of all polynomials in x_1, \dots, x_n with coefficients in k is denoted $k[x_1, \dots, x_n]$.

It follows from the previous definitions that $k[x_1, \dots, x_n]$, i.e., the set of polynomials in n variables with coefficients in k , is a commutative ring with identity. We also notice that it is possible (and sometimes, convenient) to define polynomials where the coefficients belong to a ring with identity, not necessarily to a field.

Definition 5 A form is a polynomial where all the monomials have the same degree $d := \sum_i \alpha_i$. In this case, the polynomial is homogeneous of degree d , since it satisfies $f(\lambda x_1, \dots, \lambda x_n) = \lambda^d f(x_1, \dots, x_n)$.

A commutative ring is called an *integral domain* if it has no zero divisors, i.e. $a \neq 0, b \neq 0 \Rightarrow a \cdot b \neq 0$. Every field is also an integral domain (why?). Two examples of rings that are not integral domains are the set of matrices $\mathbb{R}^{n \times n}$, and the set of integers modulo n , when n is a composite number (with the usual operations). If k is an integral domain, then so is $k[x_1, \dots, x_n]$.

Remark 6 Another important example of a ring (in this case, non-commutative) appears in systems and control theory, through the ring $\mathcal{M}(s)$ of stable proper rational functions. This is the set of matrices (of fixed dimension) whose entries are rational functions of s (i.e., in the field $\mathbb{C}(s)$), are bounded at infinity, and have all poles in the strict left-half plane. In this algebraic setting (usually called “co-prime factorization approach”), the question of finding a stabilizing controller is exactly equivalent to the solvability of a Diophantine equation $ax + by = 1$.

2.1 Algebraically closed and formally real fields

A very important property of a univariate polynomial p is the existence of a *root*, i.e., an element x_0 for which $p(x_0) = 0$. Depending on the solvability of these equations, we can characterize a particular nice class of fields.

Definition 7 A field k is algebraically closed if every nonconstant polynomial in $k[x]$ has a root in k .

If a field is algebraically closed, then it has an infinite number of elements (why?). What can we say about the most usual fields, \mathbb{C} and \mathbb{R} ? The Fundamental Theorem of Algebra (“every univariate polynomial has at least one complex root”) shows that \mathbb{C} is an algebraically closed field.

However, this is clearly *not* the case of \mathbb{R} , since for instance the polynomial $x^2 + 1$ does not have any real root. The lack of algebraic closure of \mathbb{R} is one of the main sources of complications when dealing with systems of polynomial equations and inequalities. To deal with the case when the base field is not algebraically closed, the *Artin-Schreier* theory of *formally real fields* was introduced.

The starting point is one of the intrinsic properties of \mathbb{R} :

$$\sum_{i=1}^n x_i^2 = 0 \implies x_1 = \dots = x_n = 0. \quad (2)$$

A field will be called *formally real* if it satisfies the above condition (clearly, \mathbb{R} and \mathbb{Q} are formally real, but \mathbb{C} is not). As we can see from the definition, the theory of formally real fields has very strong connections with sums of squares, a notion that will reappear in several forms later in the course. For example, an alternative (but equivalent) statement of (2) is to say that a field is formally real if and only if the element -1 is not a sum of squares.

A related important notion is that of an *ordered* field:

Definition 8 *A field k is said to be ordered if a relation $>$ is defined on k , that satisfies*

1. *If $a, b \in k$, then either $a > b$ or $a = b$ or $b > a$.*
2. *If $a > b$, $c \in k$, $c > 0$ then $ac > bc$.*
3. *If $a > b$, $c \in k$, then $a + c > b + c$.*

A crucial result relating these two notions is the following:

Lemma 9 *A field can be ordered if and only if it is formally real.*

For a field to be ordered (or equivalently, formally real), it necessarily must have an infinite number of elements. This is somewhat unfortunate, since this rules out several modular methods for dealing with real solutions to polynomial inequalities.

2.2 Ideals

We consider next *ideals*, which are subrings with an “absorbent” property:

Definition 10 *Let R be a commutative ring. A subset $I \subset R$ is an ideal if it satisfies:*

1. *$0 \in I$.*
2. *If $a, b \in I$, then $a + b \in I$.*
3. *If $a \in I$ and $b \in R$, then $a \cdot b \in I$.*

A simple example of an ideal is the set of even integers, considered as a subset of the integer ring \mathbb{Z} . Also, notice that if the ideal I contains the multiplicative identity 1, then $I = R$.

To introduce another important example of ideals, we need to define the concept of an algebraic variety as the zero set of a set of polynomial equations:

Definition 11 *Let k be a field, and let f_1, \dots, f_s be polynomials in $k[x_1, \dots, x_n]$. Let the set \mathbf{V} be*

$$\mathbf{V}(f_1, \dots, f_s) = \{(a_1, \dots, a_n) \in k^n : f_i(a_1, \dots, a_n) = 0 \quad \forall 1 \leq i \leq s\}.$$

We call $\mathbf{V}(f_1, \dots, f_s)$ the affine variety defined by f_1, \dots, f_s .

Then, the set of polynomials that vanish in a given variety, i.e.,

$$\mathbf{I}(V) = \{f \in k[x_1, \dots, x_n] : f(a_1, \dots, a_n) = 0 \quad \forall (a_1, \dots, a_n) \in V\},$$

is an ideal, called the *ideal of V*.

By Hilbert's Basis Theorem [CLO97], $k[x_1, \dots, x_n]$ is a *Noetherian* ring, i.e., every ideal $I \subset k[x_1, \dots, x_n]$ is finitely generated. In other words, there always exists a finite set $f_1, \dots, f_s \in k[x_1, \dots, x_n]$ such that for every $f \in I$, we can find $g_i \in k[x_1, \dots, x_n]$ that verify $f = \sum_{i=1}^s g_i f_i$.

We also define the *radical* of an ideal:

Definition 12 Let $I \subset k[x_1, \dots, x_n]$ be an ideal. The radical of I , denoted \sqrt{I} , is the set

$$\{f \mid f^k \in I \text{ for some integer } k \geq 1\}.$$

It is clear that $I \subset \sqrt{I}$, and it can be shown that \sqrt{I} is also a polynomial ideal. A very important result, that we will see later in some detail, is the following:

Theorem 13 (Hilbert's Nullstellensatz) If I is a polynomial ideal, then $\mathbf{I}(\mathbf{V}(I)) = \sqrt{I}$.

2.3 Associative algebras

Another important notion, that we will encounter at least twice later in the course, is that of an *associative algebra*.

Definition 14 An associative algebra \mathcal{A} over \mathbb{C} is a vector space with a \mathbb{C} -bilinear operation $\cdot : \mathcal{A} \times \mathcal{A} \rightarrow \mathcal{A}$ that satisfies

$$x \cdot (y \cdot z) = (x \cdot y) \cdot z, \quad \forall x, y, z \in \mathcal{A}.$$

In general, associative algebras do not need to be commutative (i.e., $x \cdot y = y \cdot x$). However, that is an important special case, with many interesting properties. We list below several examples of finite dimensional associative algebras.

- Full matrix algebra $\mathbb{C}^{n \times n}$, standard product.
- The subalgebra of square matrices with equal row and column sums.
- The n -dimensional algebra generated by a single $n \times n$ matrix.
- The group algebra: formal \mathbb{C} -linear combination of group elements.
- Polynomial multiplication modulo a zero dimensional ideal.
- The Bose-Mesner algebra of an association scheme.

We will discuss the last three in more detail later in the course.

3 Questions about polynomials

There are many natural questions that we may want to answer about polynomials, even in the univariate case. Among them, we mention:

- When does a univariate polynomial have *only* real roots?
- What conditions must it satisfy for *all* roots to be real?
- When does a polynomial satisfy $p(x) \geq 0$ for all x ?

We will answer many of these next week.

References

- [BCR98] J. Bochnak, M. Coste, and M-F. Roy. *Real Algebraic Geometry*. Springer, 1998.
- [CLO97] D. A. Cox, J. B. Little, and D. O'Shea. *Ideals, varieties, and algorithms: an introduction to computational algebraic geometry and commutative algebra*. Springer, 1997.
- [DF91] D. S. Dummit and R. M. Foote. *Abstract algebra*. Prentice Hall Inc., Englewood Cliffs, NJ, 1991.
- [Lan71] S. Lang. *Algebra*. Addison-Wesley, 1971.

Lecture 5

Lecturer: Pablo A. Parrilo

Scribe: Noah Stein

In this lecture we study univariate polynomials, particularly questions regarding the existence of roots and nonnegativity conditions.

- When does a univariate polynomial have *only* real roots?
- What conditions must it satisfy for *all* roots to be real?
- When does a polynomial satisfy $p(x) \geq 0$ for all x ?

1 Univariate polynomials

A univariate polynomial $p(x) \in \mathbb{R}[x]$ of degree n has the form:

$$p(x) = p_n x^n + p_{n-1} x^{n-1} + \cdots + p_1 x + p_0, \quad (1)$$

where the coefficients p_k are real. We normally assume $p_n \neq 0$, and occasionally we will normalize it to $p_n = 1$, in which case we say that $p(x)$ is *monic*.

As we have seen, the field \mathbb{C} of complex numbers is algebraically closed:

Theorem 1 (Fundamental theorem of algebra). *Every nonzero univariate polynomial of degree n has exactly n complex roots (counted with multiplicity). Furthermore, we have the unique factorization*

$$p(x) = p_n \prod_{k=1}^n (x - x_k),$$

where $x_k \in \mathbb{C}$ are the roots of $p(x)$.

If all the coefficients p_k are real, if x_k is a root, then so is its complex conjugate x_k^* . In other words, all complex roots appear in complex conjugate pairs.

2 Root bounds and Sturm sequences

To be completed

ToDo

3 Counting real roots

How many real roots does a polynomial have? There are many options, ranging from all roots being real (e.g., $(x-1)(x-2)\cdots(x-n)$), to all roots being complex (e.g., $x^{2d}+1$). We will give a couple of different characterizations of the location of the roots of a polynomial, both of them in terms of some associated symmetric matrices.

3.1 The companion matrix

A very well-known relationship between univariate polynomials and matrices is given through the so-called companion matrix.

Definition 2. The companion matrix \mathcal{C}_p associated with the polynomial $p(x)$ in (1) is the $n \times n$ real matrix

$$\mathcal{C}_p := \begin{bmatrix} 0 & 0 & \cdots & 0 & -p_0/p_n \\ 1 & 0 & \cdots & 0 & -p_1/p_n \\ 0 & 1 & \cdots & 0 & -p_2/p_n \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \cdots & 1 & -p_{n-1}/p_n \end{bmatrix}.$$

Lemma 3. The characteristic polynomial of \mathcal{C}_p is (up to a constant) equal to $p(x)$. Formally, $\det(xI - \mathcal{C}_p) = \frac{1}{p_n}p(x)$.

From this lemma, it directly follows that the eigenvalues of \mathcal{C}_p are exactly equal to the roots x_i of $p(x)$, including multiple roots the appropriate number of times. In other words, if we want to obtain the roots of a polynomial, we can do this by computing instead the eigenvalues of the associated (nonsymmetric) companion matrix. In fact, that is exactly the way that MATLAB computes roots of polynomials; see the source file `roots.m`.

For any $A \in \mathbb{C}^{n \times n}$, we always have $\text{Tr} A = \sum_{i=1}^n \lambda_i(A)$, and $\lambda_i(A^k) = \lambda_i(A)^k$. Therefore, it follows that $\sum_{i=1}^n x_i^k = \text{Tr} [\mathcal{C}_p^k]$. As a consequence of linearity, we have that if $q(x) = \sum_{j=0}^m q_j x^j$ is a univariate polynomial,

$$\sum_{i=1}^n q(x_i) = \sum_{i=1}^n \sum_{j=0}^m q_j x_i^j = \sum_{j=0}^m q_j \text{Tr} [\mathcal{C}_p^j] = \text{Tr} \left[\sum_{j=0}^m q_j \mathcal{C}_p^j \right] = \text{Tr} [q(\mathcal{C}_p)], \quad (2)$$

where the expression $q(\mathcal{C}_p)$ indicates the evaluation of the polynomial $q(x)$ on the companion matrix of $p(x)$. Note that if p is monic, then the final expression in (2) is a polynomial in the coefficients of p . This is an identity that we will use several times in the sequel.

Remark 4. Our presentation of the companion matrix has been somewhat unmotivated, other than noticing that “it just works.” After presenting some additional material on Gröbner basis, we will revisit this construction, where we will give a natural interpretation of \mathcal{C}_p as representing a well-defined linear operator in the quotient ring $\mathbb{R}[x]/\langle p(x) \rangle$. This will enable a very appealing generalization of companion matrices to multivariate polynomials, in the case where the underlying system has only a finite number of solutions (i.e., a “zero dimensional ideal”).

3.2 Inertia and signature

Definition 5. Consider a symmetric matrix A . The inertia of A , denoted $\mathcal{I}(A)$, is the triple (n_+, n_0, n_-) , where n_+, n_0, n_- are the number of positive, zero, and negative eigenvalues, respectively. The signature of A is equal to the number of positive eigenvalues minus the number of negative eigenvalues, i.e., the integer $n_+ - n_-$.

Notice that, with the notation above, the rank of A is equal to $n_+ + n_-$. A symmetric positive definite $n \times n$ matrix has inertia $(n, 0, 0)$, while a positive semidefinite one has $(n - k, k, 0)$ for some $k \geq 0$. The inertia is an important invariant of a quadratic form, since it holds that $\mathcal{I}(A) = \mathcal{I}(T^T A T)$, where T is nonsingular. This invariance of the inertia of a matrix under congruence transformations is known as *Sylvester’s law of inertia*; see for instance [HJ95].

3.3 The Hermite form

While the companion matrix is quite useful, we will present now a different characterization of the roots of a polynomial. Among others, an advantage of this formulation is the fact that we will be using *symmetric* matrices.

Let $q(x)$ be a fixed auxiliary polynomial. Consider the following $n \times n$ symmetric Hankel matrix $H_q(p)$ with complex entries defined by

$$[H_q(p)]_{jk} = \sum_{i=1}^n q(x_i) x_i^{j+k-2}. \quad (3)$$

Like every symmetric matrix, $H_q(p)$ defines an associated quadratic form via

$$\begin{aligned} f^T H_q(p) f &= \begin{bmatrix} f_0 \\ f_1 \\ \vdots \\ f_{n-1} \end{bmatrix}^T \begin{bmatrix} \sum_{i=1}^n q(x_i) & \sum_{i=1}^n q(x_i) x_i & \cdots & \sum_{i=1}^n q(x_i) x_i^{n-1} \\ \sum_{i=1}^n q(x_i) x_i & \sum_{i=1}^n q(x_i) x_i^2 & \cdots & \sum_{i=1}^n q(x_i) x_i^n \\ \vdots & \vdots & \ddots & \vdots \\ \sum_{i=1}^n q(x_i) x_i^{n-1} & \sum_{i=1}^n q(x_i) x_i^n & \cdots & \sum_{i=1}^n q(x_i) x_i^{2n-2} \end{bmatrix} \begin{bmatrix} f_0 \\ f_1 \\ \vdots \\ f_{n-1} \end{bmatrix} \\ &= \sum_{i=1}^n q(x_i) (f_0 + f_1 x_i + \cdots + f_{n-1} x_i^{n-1})^2 \\ &= \text{Tr}[(qf^2)(\mathcal{C}_p)]. \end{aligned}$$

Although not immediately obvious from the definition (3), the expression above shows that when $p(x)$ is monic, the entries of $H_q(p)$ are actually polynomials in the coefficients of $p(x)$. Notice that we have used (2) in the derivation of the last step.

Define now the $n \times n$ *Vandermonde matrix*

$$V = \begin{bmatrix} 1 & x_1 & \cdots & x_1^{n-1} \\ 1 & x_2 & \cdots & x_2^{n-1} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_n & \cdots & x_n^{n-1} \end{bmatrix}$$

where $x_1, \dots, x_n \in \mathbb{C}$. It can be shown that such a matrix is nonsingular if and only if all the x_i are distinct. Note that a Vandermonde matrix defines a linear transformation mapping the coefficients of a degree $n-1$ polynomial f to its values $[f(x_1), \dots, f(x_n)]$. Since this transformation is invertible, given any $y \in \mathbb{R}^n$ there always exists an f of degree $n-1$ such that $f(x_i) = y_i$ (i.e., there is always an interpolating polynomial). From the expressions above, we have the factorization

$$H_q(p) = V^T \text{diag}[q(x_1), \dots, q(x_n)] V.$$

This is almost a congruence transformation, except that there are complex terms if some of the x_i are complex. However, this can be easily resolved, to obtain the theorem below.

Theorem 6. *The signature of $H_q(p)$ is equal to the number of real roots x_j of p for which $q(x_j) > 0$, minus the number of real roots for which $q(x_j) < 0$.*

Proof. For simplicity, we assume all roots are distinct (this is easy to change, at the expense of slightly

more complicated notation). We have then

$$\begin{aligned}
f^T H_q(p) f &= \sum_{j=1}^n q(x_j) (f_0 + f_1 x_j + \cdots + f_{n-1} x_j^{n-1})^2 \\
&= \sum_{x_j \in \mathbb{R}} q(x_j) f(x_j)^2 + \sum_{x_j, x_j^* \in \mathbb{C} \setminus \mathbb{R}} q(x_j) f(x_j)^2 + q(x_j^*) f(x_j^*)^2 \\
&= \sum_{x_j \in \mathbb{R}} q(x_j) f(x_j)^2 + 2 \sum_{x_j, x_j^* \in \mathbb{C} \setminus \mathbb{R}} \begin{bmatrix} \Re f(x_j) \\ \Im f(x_j) \end{bmatrix}^T \begin{bmatrix} \Re q(x_j) & -\Im q(x_j) \\ -\Im q(x_j) & -\Re q(x_j) \end{bmatrix} \begin{bmatrix} \Re f(x_j) \\ \Im f(x_j) \end{bmatrix}.
\end{aligned}$$

Notice that an expression of the type $f(x_i)$ is a linear form in $[f_0, \dots, f_{n-1}]$. Because of the assumption that all the roots x_j are distinct, the linear forms $\{f(x_j)\}_{j=1, \dots, n}$ are linearly independent (the corresponding Vandermonde matrix is nonsingular), and thus so are $\{f(x_j)\}_{x_j \in \mathbb{R}} \cup \{\Re f(x_j), \Im f(x_j)\}_{x_j \in \mathbb{C} \setminus \mathbb{R}}$. Therefore, the expression above gives a congruence transformation of $H_q(p)$, and we can obtain its signature by adding the signatures of the scalar elements $q(x_j)$ and the 2×2 blocks. The signature of the 2×2 blocks is always zero (they have zero trace), and thus the result follows. \square

In particular, notice that if we want to count the number of roots, we can just use $q(x) = 1$. The matrix corresponding to this quadratic form (called the *Hermite form*) is:

$$H_1(p) = \begin{bmatrix} s_0 & s_1 & \cdots & s_{n-1} \\ s_1 & s_2 & \cdots & s_n \\ \vdots & \vdots & \ddots & \vdots \\ s_{n-1} & s_n & \cdots & s_{2n-2} \end{bmatrix}, \quad s_k = \sum_{j=1}^n x_j^k.$$

The s_k are known as the *power sums* and can be computed with 2. When $p(x)$ is monic, the s_k are polynomials of degree k in the coefficients of $p(x)$.

Example 7. Consider the monic cubic polynomial

$$p(x) = x^3 + p_2 x^2 + p_1 x + p_0.$$

Then, the first five power sums are:

$$\begin{aligned}
s_0 &= 3 \\
s_1 &= -p_2 \\
s_2 &= p_2^2 - 2p_1 \\
s_3 &= -p_2^3 + 3p_1 p_2 - 3p_0 \\
s_4 &= p_2^4 - 4p_1 p_2^2 + 2p_1^2 + 4p_0 p_2.
\end{aligned}$$

Lemma 8. The signature of $H_1(p)$ is equal to the number of real roots. The rank of $H_1(p)$ is equal to the number of distinct complex roots of $p(x)$.

Corollary 9. If $p(x)$ has odd degree, there is always at least one real root.

Example 10. Consider $p(x) = x^3 + 2x^2 + 3x + 4$. The corresponding Hermite matrix is:

$$H(p) = \begin{bmatrix} 3 & -2 & -2 \\ -2 & -2 & -2 \\ -2 & -2 & 18 \end{bmatrix}$$

This matrix has one negative and two positive eigenvalues, all distinct (i.e., its inertia is $(2, 0, 1)$). Thus, $p(x)$ has three simple roots, and exactly one of them is real.

Sylvester's law of inertia guarantees that this result is actually coordinate independent.

4 Nonnegativity

An important property of a polynomial is whether it only takes nonnegative values. As we will see, this is of interest in a wide variety of applications.

Definition 11. A univariate polynomial $p(x)$ is positive semidefinite or nonnegative if $p(x) \geq 0$ for all real values of x .

Clearly, if $p(x)$ is nonnegative, then its degree must be an even number. The set of nonnegative polynomials has very interesting properties. Perhaps the most appealing one for our purposes is the following:

Theorem 12. Consider the set P_n of nonnegative univariate polynomials of degree less than or equal to n (n is even). Then, identifying a polynomial with its $n + 1$ coefficients (p_n, \dots, p_0) , the set P_n is a proper cone (i.e., closed, convex, pointed, solid) in \mathbb{R}^{n+1} .

An equivalent condition for the (nonconstant) univariate polynomial (1) to be strictly positive, is that $p(x_0) > 0$ for some x_0 , and it has no real roots. Thus, we can use Theorem 6 to write explicit conditions for a polynomial $p(x)$ to be nonnegative in terms of the signature of the associated Hermite matrix $H_1(p)$.

5 Sum of squares

Definition 13. A univariate polynomial $p(x)$ is a sum of squares (SOS) if there exist $q_1, \dots, q_m \in \mathbb{R}[x]$ such that

$$p(x) = \sum_{k=1}^m q_k^2(x).$$

If a polynomial $p(x)$ is a sum of squares, then it obviously satisfies $p(x) \geq 0$ for all $x \in \mathbb{R}$. Thus, a SOS condition is a sufficient condition for global nonnegativity.

Interestingly, in the univariate case, the converse is also true:

Theorem 14. A univariate polynomial is nonnegative if and only if it is a sum of squares.

Proof. (\Leftarrow) Obvious. If $p(x) = \sum_k q_k^2(x)$ then $p(x) \geq 0$.

(\Rightarrow) Since $p(x)$ is univariate, we can factorize it as

$$p(x) = p_n \prod_j (x - r_j)^{n_j} \prod_k (x - a_k + ib_k)^{m_k} (x - a_k - ib_k)^{m_k},$$

where r_j and $a_k \pm ib_k$ are the real and complex roots, respectively. Because $p(x)$ is nonnegative, then $p_n > 0$ and the multiplicities of the real roots are even, i.e., $n_j = 2s_j$.

Notice that $(x - a + ib)(x - a - ib) = (x - a)^2 + b^2$. Then, we can write

$$p(x) = p_n \prod_j (x - r_j)^{2s_j} \prod_k ((x - a_k)^2 + b_k^2)^{m_k},$$

Since products of sums of squares are sums of squares, and all the factors in the expression above are SOS, it follows that $p(x)$ is SOS.

Furthermore, the two-squares identity $(\alpha^2 + \beta^2)(\gamma^2 + \delta^2) = (\alpha\gamma - \beta\delta)^2 + (\alpha\delta + \beta\gamma)^2$ allows us to combine every partial product as a sum of only two squares. \square

Notice that the proof shows that if $p(x)$ is SOS, then there exists a representation $p(x) = q_1^2(x) + q_2^2(x)$.

As we will see very soon, we can decide whether a univariate polynomial is a sum of squares (equivalently, if it is nonnegative) by solving a semidefinite optimization problem.

6 Positive semidefinite matrices

Recall from Lecture 2 the (apparent) disparity between the stated conditions for a matrix to be positive definite versus the semidefinite case. In the former, we could use a test (Sylvester's criterion) that required the calculation of only n minors, while for the semidefinite case apparently we needed a much larger number, $2^n - 1$.

If the matrix X is positive definite, Sylvester's criterion requires the positivity of the leading principal minors, i.e.,

$$\det(X_{1,1}) > 0, \quad \det X_{12,12} > 0, \quad \dots, \quad \det X > 0.$$

For positive semidefiniteness, it is not enough to replace strict positivity with the nonstrict inequality; a simple counterexample is the matrix

$$\begin{bmatrix} 0 & 0 \\ 0 & -1 \end{bmatrix},$$

for which the leading minors vanish, but is not PSD. As mentioned, an alternative approach is given by the following classical result:

Lemma 15. *Let $A \in \mathbb{R}^{n \times n}$ be symmetric. Then $A \succeq 0$ if and only if all $2^n - 1$ principal minors of A are nonnegative.*

Although the condition above requires the nonnegativity of $2^n - 1$ expressions, it is possible to do the same by checking only n inequalities:

Theorem 16 (e.g. [HJ95, p. 403]). *A real $n \times n$ symmetric matrix A is positive semidefinite if and only if all the coefficients c_i of its characteristic polynomial $p(\lambda) = \det(\lambda I - A) = \lambda^n + p_{n-1}\lambda^{n-1} + \dots + p_1\lambda + p_0$ alternate in sign, i.e., they satisfy $p_i(-1)^{n-i} \geq 0$.*

We prove this below, since we will use a slightly more general version of this result when discussing hyperbolic polynomials. Note that in the $n = 2$ case Theorem 16 is the familiar result that A is positive semidefinite if and only if $\det A \geq 0$ and $\text{Tr} A \geq 0$.

Lemma 17. *Consider a monic univariate polynomial $p(t) = t^n + \sum_{k=0}^{n-1} p_k t^k$, that has only real roots. Then, all roots are nonpositive if and only if all coefficients are nonnegative (i.e., $p_k \geq 0, k = 0, \dots, n-1$).*

Proof. Since all roots of $p(t)$ are real, this can be obtained from a direct application of Descartes' rules of signs; see e.g. [BPR03]. For completeness, we present here a direct proof.

If all roots t_i are nonpositive ($t_i \leq 0$), from the factorization

$$p(t) = \prod_{k=1}^n (t - t_i)$$

it follows directly that all coefficients p_k are nonnegative.

For the other direction, from the nonnegativity of the coefficients it follows that $p(0) \geq 0$ and $p(t)$ is nondecreasing. If there exists a $t_i > 0$ such that $p(t_i) = 0$, then the polynomial must vanish in the interval $[0, t_i]$, which is impossible since it is monic and hence nonzero. \square

Definition 18. *A set $S \subset \mathbb{R}^n$ is basic closed semialgebraic if it can be written as*

$$S = \{x \in \mathbb{R}^n \mid f_i(x) \geq 0, \quad h_j(x) = 0\}$$

for some finite number of polynomials f_i, h_j .

Theorem 19. *Both the primal and dual feasible sets of a semidefinite program are basic closed semialgebraic.*

Proof. The condition $X \succeq 0$ is equivalent to n nonstrict polynomial inequalities in the entries of X . This can be conveniently shown applying Lemma 17 to the characteristic polynomial of $-X$, i.e.,

$$p(\lambda) = \det(\lambda I + X) = \lambda^n + \sum_{k=0}^{n-1} p_k(X) \lambda^k.$$

where the $p_k(X)$ are homogeneous polynomials of degree $n - k$ in the entries of X . For instance, we have $p_0(X) = \det X$, and $p_{n-1}(X) = \text{Tr} X$.

Since X is symmetric, all its eigenvalues are real, and thus $p(\lambda)$ has only real roots. Positive semidefiniteness of X is equivalent to $p(\lambda)$ having no roots that are strictly positive. It then follows that the two following statements are equivalent:

$$X \succeq 0 \quad \Leftrightarrow \quad p_k(X) \geq 0 \quad k = 0, \dots, n-1.$$

□

Remark 20. *These inequalities correspond to the elementary symmetric functions e_k evaluated at the eigenvalues of the matrix X .*

As we will see in subsequent lectures, the same inequalities will reappear when we consider a class of optimization problems known as *hyperbolic programs*.

References

- [BPR03] S. Basu, R. Pollack, and M.-F. Roy. *Algorithms in real algebraic geometry*, volume 10 of *Algorithms and Computation in Mathematics*. Springer-Verlag, Berlin, 2003.
- [HJ95] R. A. Horn and C. R. Johnson. *Matrix Analysis*. Cambridge University Press, Cambridge, 1995.

Lecture 6

Lecturer: Pablo A. Parrilo

Scribe: ???

Last week we learned about explicit conditions to determine the number of real roots of a univariate polynomial. Today we will expand on these themes, and study two mathematical objects of fundamental importance: the *resultant* of two polynomials, and the closely related *discriminant*.

The resultant will be used to decide whether two univariate polynomials have common roots, while the discriminant will give information about the existence of multiple roots. Furthermore, we will see the intimate connections between discriminants and the boundary of the cone of nonnegative polynomials.

Besides the properties described above, a direct consequence of their definitions, there are many other interesting applications of resultants and discriminant. We describe a few of them below, and we will encounter them again in later lectures, when studying elimination theory and the construction of cylindrical algebraic decompositions. For much more information about resultants and discriminants, particularly their generalizations to the sparse and multipolynomial case, we refer the reader to the very readable introductory article [Stu98] and the book [CLO97].

1 Resultants

Consider two polynomials $p(x)$ and $q(x)$, of degree n, m , respectively. We want to obtain an easily checkable criterion to determine whether they have a common root, that is, there exists an $x_0 \in \mathbb{C}$ for which $p(x_0) = q(x_0) = 0$. There are several approaches, seemingly different at first sight, for constructing such a criterion:

- **Sylvester matrix:** If $p(x_0) = q(x_0) = 0$, then we can write the following $(n+m) \times (n+m)$ linear system:

$$\begin{bmatrix}
 p_n & p_{n-1} & \cdots & p_1 & p_0 & & & & \\
 & p_n & & \ddots & \ddots & \ddots & & & \\
 \vdots & & & & & & p_1 & p_0 & \\
 & & & & & & p_2 & p_1 & p_0 \\
 & & & & & & & & \\
 q_m & q_{m-1} & \cdots & q_0 & & & & & \\
 & q_m & & \ddots & \ddots & & & & \\
 & & & & \ddots & & & & \\
 & & & & & & q_1 & q_0 & \\
 & & & & & & q_2 & q_1 & q_0
 \end{bmatrix}
 \begin{bmatrix}
 x_0^{n+m-1} \\
 x_0^{n+m-2} \\
 \vdots \\
 x_0^n \\
 x_0^{n-1} \\
 \vdots \\
 x_0 \\
 1
 \end{bmatrix}
 =
 \begin{bmatrix}
 p(x_0)x_0^{m-1} \\
 p(x_0)x_0^{m-2} \\
 \vdots \\
 p(x_0)x_0 \\
 p(x_0) \\
 q(x_0)x_0^{n-1} \\
 q(x_0)x_0^{n-2} \\
 \vdots \\
 q(x_0)x_0 \\
 q(x_0)
 \end{bmatrix}
 = 0.$$

This implies that the matrix on the left-hand side, called the *Sylvester matrix* $\text{Syl}_x(p, q)$ associated to p and q , is singular and thus its determinant must vanish. It is not too difficult to show that the converse is also true; if $\det \text{Syl}_x(p, q) = 0$, then there exists a vector in the kernel of $\text{Syl}_x(p, q)$ of the form shown in the matrix equation above, and thus a common root x_0 .

- **Root products and companion matrices:** Let α_j, β_k be the roots of $p(x)$ and $q(x)$, respectively. By construction, the expression

$$\prod_{j=1}^n \prod_{k=1}^m (\alpha_j - \beta_k)$$

vanishes if and only if there exists a root of p that is equal to a root of q . Although the computation of this product seems to require explicit access to the roots, this can be avoided. Multiplying by

a convenient normalization factor, we have:

$$\begin{aligned}
 p_n^m q_m^n \prod_{j=1}^n \prod_{k=1}^m (\alpha_j - \beta_k) &= p_n^m \prod_{j=1}^n q(\alpha_j) = p_n^m \det q(\mathcal{C}_p) \\
 &= (-1)^{nm} q_m^n \prod_{k=1}^m p(\beta_k) = (-1)^{nm} q_m^n \det p(\mathcal{C}_q)
 \end{aligned} \tag{1}$$

- **Kronecker products:** Using a well-known connection to Kronecker products, we can also write (1) as

$$p_n^m q_m^n \det(\mathcal{C}_p \otimes I_m - I_n \otimes \mathcal{C}_q).$$

- **Bézout matrix**

To be completed

ToDo

If can be shown that all these constructions are equivalent. They define exactly the same polynomial, called the *resultant* of p and q , denoted as $\text{Res}_x(p, q)$:

$$\begin{aligned}
 \text{Res}_x(p, q) &= \det \text{Syl}_x(p, q) \\
 &= p_n^m \det q(\mathcal{C}_p) \\
 &= (-1)^{nm} q_m^n \det p(\mathcal{C}_q) \\
 &= p_n^m q_m^n \det(\mathcal{C}_p \otimes I_m - I_n \otimes \mathcal{C}_q).
 \end{aligned}$$

The resultant is a homogeneous multivariate polynomial, with integer coefficients, and of degree $n + m$ in the $n + m + 2$ variables p_j, q_k . It vanishes if and only if the polynomials p and q have a common root. Notice that the definition is not symmetric in its two arguments, $\text{Res}_x(p, q) = (-1)^{nm} \text{Res}(q, p)$ (of course, this does not matter in checking whether it is zero).

Remark 1. To compute the resultant of two polynomials $p(x)$ and $q(x)$ in Maple, you can use the command `resultant(p, q, x)`. In Mathematica, use instead `Resultant[p, q, x]`.

2 Discriminants

As we have seen, the resultant allow us to write an easily checkable condition for the simultaneous vanishing of two univariate polynomials. Can we use the resultant to produce a condition for a polynomial to have a double root? Recall that if a polynomial $p(x)$ as a double root at x_0 (which can be real or complex), then its derivative $p'(x)$ also vanishes at x_0 . Thus, we can check for the existence of a double root by computing the resultant between a polynomial and its derivative.

Definition 2. The discriminant of a univariate polynomial $p(x)$ is defined as

$$\text{Dis}_x(p) := (-1)^{n(n-1)/2} \frac{1}{p_n} \text{Res}_x \left(p(x), \frac{dp(x)}{dx} \right).$$

Similar to what we did in the resultant case, the discriminant can also be obtained by writing a natural condition in terms of the roots α_i of $p(x)$:

$$\text{Dis}_x(p) = p_n^{2n-2} \prod_{j < k} (\alpha_j - \alpha_k)^2.$$

If $p(x)$ has degree n , its discriminant is a homogeneous polynomial of degree $2n - 2$ in its $n + 1$ coefficients p_n, \dots, p_0 .

Example 3. Consider the quadratic univariate polynomial $p(x) = ax^2 + bx + c$. Its discriminant is:

$$\text{Dis}_x(p) = -\frac{1}{a}\text{Res}_x(ax^2 + bx + c, 2ax + b) = b^2 - 4ac.$$

For the cubic polynomial $p(x) = ax^3 + bx^2 + cx + d$ we have

$$\text{Dis}_x(p) = -27a^2d^2 + 18adcb + b^2c^2 - 4b^3d - 4ac^3.$$

3 Applications

3.1 Polynomial equations

One of the most natural applications of resultants is in the solution of polynomial equations in two variables. For this, consider a polynomial system

$$p(x, y) = 0, \quad q(x, y) = 0, \quad (2)$$

with only a finite number of solutions (which is generically the case). Consider a fixed value of y_0 , and the two univariate polynomials $p(x, y_0), q(x, y_0)$. If y_0 corresponds to the y -component of a root, then these two univariate polynomials clearly have a common root, hence their resultant vanishes.

Therefore, to solve (2), we can compute $\text{Res}_x(p, q)$, which is a univariate polynomial in y . Solving this univariate polynomial, we obtain a finite number of points y_i . Backsubstituting in p (or q), we obtain the corresponding values of x_i . Naturally, the same construction can be used by computing first the univariate polynomial in x given by $\text{Res}_y(p, q)$.

Example 4. Let $p(x, y) = 2xy + 3y^3 - 2x^3 - x - 3x^2y^2$, $q(x, y) = 2x^2y^2 - 4y^3 - x^2 + 4y + x^2y$. The resultant (in the x variable) is

$$\text{Res}_x(p, q) = y(y + 1)^3(72y^8 - 252y^7 + 270y^6 - 145y^5 + 192y^4 - 160y^3 + 28y + 4).$$

One particular root of this polynomial is $y_* \approx 1.6727$, with the corresponding value of $x_* \approx -1.3853$.

3.2 Implicitization of rational curves

To be completed

ToDo

3.3 Random matrices

To be completed

ToDo

4 The set of nonnegative polynomials

One of the main reasons why nonnegativity conditions about polynomials are difficult is because these sets can have a quite complicated structure, even though they are always convex.

Recall from last lecture that we have defined $P_n \subset \mathbb{R}^{n+1}$ as the set of nonnegative polynomials of degree n . It is easy to see that if $p(x)$ is in the boundary of the set P_n , then it must have a real root, of multiplicity at least two. Indeed, if there is no real root, then $p(x)$ is in the strict interior of P




Figure 1: The shaded region corresponds to the polynomial $x^4 + 2ax^2 + b$ being nonnegative. The numbers indicate how many real roots $p(x)$ has.




Figure 2: Region of nonnegativity of the polynomial $x^4 + 4ax^3 + 6bx^2 + 4ax + 1$, and number of real roots.

(small enough perturbations will not create a root), and if it has a simple real root it clearly cannot be nonnegative.

Thus, on the boundary of P_n , the discriminant of $p(x)$ must necessarily vanish. However, it turns out that $\text{Dis}_x(p)$ does not vanish *only* on the boundary, but it also vanishes at points inside the set. Why is this?

Example 5. Consider the univariate polynomial $p(x) = x^4 + 2ax^2 + b$. For what values of a, b does it hold that $p(x) \geq 0 \forall x \in \mathbb{R}$? Since the leading term x^4 has even degree and is strictly positive, $p(x)$ is strictly positive if and only if it has no real roots. The discriminant of $p(x)$ is equal to $256b(a^2 - b)^2$.

Here is a slightly different example, showing the same phenomenon.

Example 6. As another example, consider now $p(x) = x^4 + 4ax^3 + 6bx^2 + 4ax + 1$. Its discriminant, in factored form, is equal to $256(1 + 3b + 4a)(1 + 3b - 4a)(1 + 2a^2 - 3b)^2$. The corresponding nonnegativity region and number of real roots are presented in Figure 2.

As we can see, this creates some difficulties. For instance, even though we have a perfectly valid analytic expression for the boundary of the set, we cannot have a good sense of “how far we are” from the boundary by looking at the absolute value of the discriminant.

From the mathematical viewpoint, there are a couple of (unrelated?) reasons with these sets cannot be directly handled by “standard” optimization, at least if we want to keep the polynomial structure.




Figure 3: A three-dimensional convex set, described by one quadratic and one linear inequality, whose projection on the (a, b) plane is equal to the set in Figure 1.

One has to do with its algebraic structure, and the other one with convexity, and in particular its facial structure.

Lemma 7 (e.g., [And03]). *The set described in Figure 1 is not basic closed semialgebraic.*

Remark 8. *Notice that the convex sets described in Figures 1 and 2 both have an uncommon feature. They both have proper faces that are not exposed, i.e., they cannot be isolated by a supporting hyperplane¹. Indeed, in Figure 1 the origin $(0, 0)$ is a non-exposed zero-dimensional face, while in Figure 2 the point $(1, 1)$ has the same property. A non-exposed face is a known obstruction for a convex set to be the feasible set of a semidefinite program, see [RG95].*

Even though these sets have these complicating features, it turns out that we can often provide some “good” representations. These are normally given as a *projection* from higher dimensional spaces, where the object “upstairs” is much more smooth and well-behaved. For instance, as a graphical illustration, in Figure 3 we can see a three dimensional convex set, whose projection on the plane (a, b) is exactly the one discussed in Example 5 and Figure 1.

The presence of “extraneous” components of the discriminant inside the feasible set is an important roadblock for the availability of “easily computable” barrier functions. Indeed, every polynomial that vanishes on the boundary of the set P_n must necessarily have the discriminant as a factor. This is a striking difference with the case of the nonnegative orthant or the PSD cone, where the standard barriers are given (up to a logarithm) by products of the linear constraints or a determinant (which are polynomials). The way out of this problem is to produce non-polynomial barrier functions, either by partial minimization from a higher-dimensional barrier (i.e., projection) or using the “universal” barrier function introduced by Nesterov and Nemirovski.

In this direction, there have been several research efforts that aim at directly characterizing barrier functions for the set of nonnegative polynomials (or related modifications). Among them, we mention the work of Kao and Megretski [KM02] and Faybusovich [Fay02], both of which produce barriers that rely on the computation of one or more integral expressions. Given the fact that these integrals must be computed numerically, there is no clear consensus yet on how useful this approach is in practical optimization problems.

¹A *face* of a convex set S is a convex subset $F \subseteq S$, with the property that $x, y \in S, \frac{1}{2}(x + y) \in F \Rightarrow x, y \in F$. A face F is *exposed* if it can be written as $F = S \cap H$, where H is a supporting hyperplane of S .




Figure 4: The discriminant of the polynomial $x^4 + 4ax^3 + 6bx^2 + 4cx + 1$. The convex set inside the “bowl” corresponds to the region of nonnegativity. There is an additional one-dimensional component inside the set.

References

- [And03] C. Andradas. Characterization and description of basic semialgebraic sets. In *Algorithmic and quantitative real algebraic geometry (Piscataway, NJ, 2001)*, volume 60 of *DIMACS Ser. Discrete Math. Theoret. Comput. Sci.*, pages 1–12. Amer. Math. Soc., Providence, RI, 2003.
- [CLO97] D. A. Cox, J. B. Little, and D. O’Shea. *Ideals, varieties, and algorithms: an introduction to computational algebraic geometry and commutative algebra*. Springer, 1997.
- [Fay02] L. Faybusovich. Self-concordant barriers for cones generated by Chebyshev systems. *SIAM J. Optim.*, 12(3):770–781 (electronic), 2002.
- [KM02] C. Y. Kao and A. Megretski. A new barrier function for IQC optimization problems. In *Proceedings of the American Control Conference*, 2002.
- [RG95] M. Ramana and A. J. Goldman. Some geometric results in semidefinite programming. *J. Global Optim.*, 7(1):33–50, 1995.
- [Stu98] B. Sturmfels. Introduction to resultants. In *Applications of computational algebraic geometry (San Diego, CA, 1997)*, volume 53 of *Proc. Sympos. Appl. Math.*, pages 25–39. Amer. Math. Soc., Providence, RI, 1998.

Lecture 7

Lecturer: Pablo A. Parrilo

Scribe: ???

In this lecture we introduce a special class of multivariate polynomials, called *hyperbolic*. These polynomials were originally studied in the context of partial differential equations. As we will see, they have many surprising properties, and are intimately linked with convex optimization problems that have an algebraic structure. A few good references about the use of hyperbolic polynomials in optimization are [Gül97, BGLS01, Ren].

1 Hyperbolic polynomials

Consider a homogeneous multivariate polynomial $p \in \mathbb{R}[x_1, \dots, x_n]$ of degree d . Here *homogeneous of degree d* means that the sum of degrees of each monomial is constant and equal to d , i.e.,

$$p(x) = \sum_{|\alpha|=d} c_\alpha x^\alpha,$$

where $\alpha = (\alpha_1, \dots, \alpha_n)$, $\alpha_i \in \mathbb{N} \cup \{0\}$, and $|\alpha| = \alpha_1 + \dots + \alpha_n$. A homogeneous polynomial satisfies $p(tw) = t^d p(w)$ for all real t and vectors $w \in \mathbb{R}^n$. We denote the set of such polynomials by $\mathcal{H}_n(d)$. By identifying a polynomial with its vector of coefficients, we can consider $\mathcal{H}_n(d)$ as a normed vector space of dimension $\binom{n+d-1}{d}$.

Definition 1. Let e be a fixed vector in \mathbb{R}^n . A polynomial $p \in \mathcal{H}_n(d)$ is *hyperbolic with respect to e* if $p(e) > 0$ and, for all vectors $x \in \mathbb{R}^n$, the univariate polynomial $t \mapsto p(x - te)$ has only real roots.

A natural geometric interpretation is the following: consider the hypersurface in \mathbb{R}^n given by $p(x) = 0$. Then, hyperbolicity is equivalent to the condition that every line in \mathbb{R}^n parallel to e intersects this hypersurface at exactly d points (counting multiplicities), where d is the degree of the polynomial.

Example 2. The polynomial $x_1 x_2 \cdots x_n$ is hyperbolic with respect to the vector $(1, 1, \dots, 1)$, since the univariate polynomial $t \mapsto (x_1 - t)(x_2 - t) \cdots (x_n - t)$ has roots x_1, x_2, \dots, x_n .

Hyperbolic polynomials enjoy a very surprising property, that connects in an unexpected way algebra with convex analysis. Given a hyperbolic polynomial $p(x)$, consider the set defined as:

$$\Lambda_{++} := \{x \in \mathbb{R}^n : p(x - te) = 0 \Rightarrow t > 0\}.$$

Geometrically, this condition says that if we start at the point $x \in \mathbb{R}^n$, and slide along a line in the direction parallel to e , then we will never encounter the hypersurface $p(x) = 0$, while if we move in the opposite direction, we will cross it exactly n times. Figure 1 illustrates a particular hyperbolicity cone.

It is immediate from homogeneity and the definition above that $\lambda > 0, x \in \Lambda_{++} \Rightarrow \lambda x \in \Lambda_{++}$. Thus, we call Λ_{++} the *hyperbolicity cone* associated to p , and denote its closure by Λ_+ . As we will see shortly, it turns out that these cones are actually *convex cones*. We prove this following the arguments in Renegar [Ren]; the original results are due to Gårding [Går59].

Lemma 3. The hyperbolicity cone Λ_{++} is the connected component of $p(x) > 0$ that includes e .

Example 4. The hyperbolicity cone Λ_{++} associated with the polynomial $x_1 x_2 \cdots x_n$ discussed in Example 2 is the open positive orthant $\{x \in \mathbb{R}^n \mid x_i > 0\}$.

The first step is to show that we can replace e with any vector in the hyperbolicity cone.

Lemma 5. If $p(x)$ is hyperbolic with respect to e , then it is also hyperbolic with respect to every direction $v \in \Lambda_{++}$. Furthermore, the hyperbolicity cones are the same.




Figure 1: Hyperbolicity cone corresponding to the polynomial $p(x, y, z) = 4xyz + xz^2 + yz^2 + 2z^3 - x^3 - 3zx^2 - y^3 - 3zy^2$. This polynomial is hyperbolic with respect to $(0, 0, 1)$.

Proof. By Lemma 3 we have $p(v) > 0$. We need to show that for every $x \in \mathbb{R}^n$, the polynomial $\beta \mapsto p(\beta v + x)$ has only real roots if $v \in \Lambda_{++}$.

Let $\alpha > 0$ be fixed, and consider the polynomial $\beta \mapsto p(\alpha i e + \beta v + \gamma x)$, where i is the imaginary unit. We claim that if $\gamma \geq 0$, this polynomial has only roots in the lower half-plane. Let's look at the $\gamma = 0$ case first. It is clear that $\beta \mapsto p(\alpha i e + \beta v)$ cannot have a root at $\beta = 0$, since $p(\alpha i e) = (\alpha i)^d p(e) \neq 0$. If $\beta \neq 0$, we can write

$$p(\alpha i e + \beta v) = 0 \quad \Leftrightarrow \quad p(\alpha \beta^{-1} i e + v) = 0 \quad \Rightarrow \quad \alpha \beta^{-1} i < 0 \quad \Rightarrow \quad \beta \in i\mathbb{R}_-,$$

and thus the roots of this polynomial are on the strict negative imaginary axis (we have used $v \in \Lambda_{++}$ in the second implication). If by increasing γ there is ever a root in the upper half-plane, then there must exist a γ_* for which $\beta \mapsto p(\alpha i e + \beta v + \gamma_* x)$ has a real root β_* , and thus $p(\alpha i e + \beta_* v + \gamma_* x) = 0$. However, this contradicts hyperbolicity, since $\beta_* v + \gamma_* x \in \mathbb{R}^n$. Thus, for all $\gamma \geq 0$, the roots of $\beta \mapsto p(\alpha i e + \beta v + \gamma x)$ are in the lower half-plane.

The conclusion above was true for any $\alpha > 0$. Letting $\alpha \rightarrow 0$, by continuity of the roots we have that the polynomial $\beta \mapsto p(\beta v + \gamma x)$ must also have its roots in the lower closed half-plane. However, since it is a polynomial with real coefficients (and therefore its roots always appear in complex-conjugate pairs), then all the roots must actually be real. Taking now $\gamma = 1$, we have that $\beta \mapsto p(\beta v + x)$ has real roots for all x , or equivalently, p is hyperbolic in the direction v . \square

The following result shows that this set is actually convex:

Theorem 6 ([Gär59]). *The hyperbolicity cone Λ_{++} is convex.*

Proof. We want to show that $u, v \in \Lambda_{++}$, $\beta, \gamma > 0$ implies that $\beta u + \gamma v \in \Lambda_{++}$. The previous result implies that we can always assume $v = e$. But then the roots of $t \mapsto p(\beta u + \gamma e - te)$ are just a nonnegative affine scaling of the roots of $t \mapsto p(u - te)$, since

$$p(u - t_* e) = 0 \quad \Leftrightarrow \quad p(\beta u + \gamma e - (\beta t_* + \gamma)e) = 0,$$

and $u \in \Lambda_{++}$ implies that $t_* > 0$, hence $\beta t_* + \gamma > 0$, and as a consequence, $\beta u + \gamma e \in \Lambda_{++}$. \square

Hyperbolic polynomials are of interest in convex optimization, because they unify in a quite appealing way many facts about the most important tractable classes: linear, second order, and semidefinite programming.

Example 7 (SOCP). Let $p(x) = x_{n+1}^2 - \sum_{k=1}^n x_k^2$. This is a homogeneous quadratic polynomial, hyperbolic in the direction $e = (0, \dots, 0, 1)$, since

$$p(x - te) = (x_{n+1} - t)^2 - \sum_{k=1}^n x_k^2 = t^2 - 2tx_{n+1} + \left(x_{n+1}^2 - \sum_{k=1}^n x_k^2 \right),$$

and the discriminant of this quadratic equation is equal to

$$4x_{n+1}^2 - 4 \left(x_{n+1}^2 - \sum_{k=1}^n x_k^2 \right) = 4 \sum_{k=1}^n x_k^2,$$

which is always nonnegative, so the polynomial $t \mapsto p(x - te)$ has only real roots. The corresponding hyperbolicity cone is the Lorentz or second order cone given by

$$\Lambda_+ = \left\{ x \in \mathbb{R}^{n+1} \mid x_{n+1} \geq 0, \quad \sum_{k=1}^n x_k^2 \leq x_{n+1}^2 \right\}.$$

Example 8 (SDP). Consider the homogeneous polynomial

$$p(x) = \det(x_1 A_1 + \dots + x_n A_n),$$

where $A_i \in \mathcal{S}^d$ are given symmetric matrices, with $A_1 \succ 0$. The polynomial $p(x)$ is homogeneous of degree d . Letting $e = (1, 0, \dots, 0)$, we have

$$p(x - te) = \det \left(\sum_{k=1}^n x_k A_k - t A_1 \right) = \det A_1 \cdot \det \left(\sum_{k=1}^n x_k A_1^{-\frac{1}{2}} A_k A_1^{-\frac{1}{2}} - t I \right),$$

and as a consequence the roots of $p(x - te)$ are always real since they are the eigenvalues of a symmetric matrix. Thus, $p(x)$ is hyperbolic with respect to e . The corresponding hyperbolicity cone is

$$\Lambda_{++} = \{x \in \mathbb{R}_n \mid x_1 A_1 + \dots + x_n A_n \succ 0\}.$$

Thus, by Lemma 5, $p(x)$ is hyperbolic with respect to every $x \in \Lambda_{++}$.

Based on the results discussed earlier regarding the number of real roots of a univariate polynomial, we have the following lemma.

Lemma 9. The polynomial $p(x)$ is hyperbolic with respect to e if and only if the Hermite matrix $H_1(p) \in \mathcal{S}^n[x]$ is positive semidefinite for all $x \in \mathbb{R}^n$.

Lemma 10. The hyperbolicity cone Λ_+ is basic closed semialgebraic, i.e., it can be described by unquantified polynomial inequalities.

The two following results are of importance in optimization and the formulation of interior-point methods.

Theorem 11. A hyperbolic cone Λ_+ is facially exposed.

Theorem 12 ([Gül97]). The function $-\log p(x)$ is a logarithmically homogeneous self-concordant barrier¹ for the hyperbolicity cone Λ_{++} , with barrier parameter equal to d .

¹A function $f : \mathbb{R} \rightarrow \mathbb{R}$ is self-concordant if it satisfies $f''(x) \geq |\frac{1}{2}f'''(x)|^{\frac{2}{3}}$. A function $f : \mathbb{R}^n \rightarrow \mathbb{R}$ is self-concordant if the univariate function obtained when restricting to any line is self-concordant. Self-concordance implies strict convexity, and is a crucial property in the analysis of the polynomial-time global convergence of Newton's method; see [NN94] or [BV04, Section 9.6] for more details.

One of the main open issues regarding hyperbolic cones is about their generality. As Example 8 shows, the cone associated with a semidefinite program is a hyperbolic cone. An open question (known as the generalized Lax conjecture) is whether the converse holds, more specifically, whether every hyperbolic cone is a “slice” of the semidefinite cone, i.e., it can be represented as the intersection of an affine subspace and \mathcal{S}_+^n . As we will see in the next lecture, a special case of the conjecture has been settled recently.

2 SDP representability

Recall that in the previous lecture, we encountered a class of convex sets in \mathbb{R}^2 that lacked certain desirable properties (namely, being basic semialgebraic, and facially exposed). As we will see, hyperbolic polynomials will play a fundamental role in the characterization of the properties a set in \mathbb{R}^2 must satisfy for it to be the feasible set of a semidefinite program.

References

- [BGLS01] H. H. Bauschke, O. Güler, A. S. Lewis, and H. S. Sendov. Hyperbolic polynomials and convex analysis. *Canad. J. Math.*, 53(3):470–488, 2001.
- [BV04] S. Boyd and L. Vandenberghe. *Convex Optimization*. Cambridge University Press, 2004.
- [Går59] L. Gårding. An inequality for hyperbolic polynomials. *J. Math. Mech.*, 8:957–965, 1959.
- [Gül97] O. Güler. Hyperbolic polynomials and interior point methods for convex programming. *Math. Oper. Res.*, 22(2):350–377, 1997.
- [NN94] Y. E. Nesterov and A. Nemirovski. *Interior point polynomial methods in convex programming*, volume 13 of *Studies in Applied Mathematics*. SIAM, Philadelphia, PA, 1994.
- [Ren] J. Renegar. Hyperbolic programs, and their derivative relaxations. Preprint, 2005.

Lecture 8

Lecturer: Pablo A. Parrilo

Scribe: ???

1 SDP representability

A few lectures ago, when discussing the set of nonnegative polynomials, we encountered convex sets in \mathbb{R}^2 that lacked certain desirable properties (namely, being basic semialgebraic, and facially exposed). As we will see, hyperbolic polynomials will play a fundamental role in the characterization of the properties a set in \mathbb{R}^2 must satisfy for it to be the feasible set of a semidefinite program.

2 Convex sets in \mathbb{R}^2

In this lecture we will study conditions that a set $S \subset \mathbb{R}^2$ must satisfy for it to be *semidefinite representable*, i.e., to admit a characterization of the type

$$\{(x, y) \in \mathbb{R}^2 \mid I + xB + yC \succeq 0\}, \quad (1)$$

where $B, C \in \mathcal{S}^d$. Notice that we have assumed (without loss of generality) that $0 \in \text{int } S$, and normalized the first matrix in the matrix pencil to be an identity matrix (this can always be achieved by left- and right-multiplying by an appropriate factor).

Remark 1. *We should not confuse the notion of semidefinite representability described above, with the much more general lifted SDP representability, that allows the representation of the original set as a projection of a higher-dimensional SDP set. In other words, here we are not allowed to use additional variables.*

Clearly, from (1), we have the following necessary conditions for SDP representability:

- **Closed:** Every set of the form (1) is closed, in the standard topology.
- **Convex:** Every set of the form (1) is necessarily convex, since it is (the projection of) the intersection of an affine subspace and the convex set of PSD matrices. Of course, this is also easy to prove directly.
- **Basic semialgebraic:** As we have discussed, the boundary of the set (1) is defined by d unquantified polynomial inequalities of degree at most equal to d . In fact, the interior of this set exactly corresponds to the connected component of $\det(I + xB + yC) > 0$ that contains the origin.

There is a less obvious additional condition, which we have also seen already:

- **Exposed faces:** Every convex set of the form (1) has proper faces that are *exposed*. In other words, every face F must have a representation as $F = S \cap H$, where H is a supporting hyperplane of the convex set S .

A natural question, then, is the following: are the conditions listed above *sufficient* for SDP representability? If a set $S \subset \mathbb{R}^2$ satisfies these four conditions, do there always exist matrices B, C , for which the set (1) is exactly equal to S ? To ask a concrete question: does the set in Figure 1 admit an SDP representation? Before settling this issue, let us discuss first an apparently different question, involving hyperbolic polynomials.




Figure 1: Convex set defined by $x^4 + y^4 \leq 1$.

3 Hyperbolicity and the Lax conjecture

Recall from the previous lecture that a hyperbolic polynomial is a homogeneous polynomial $p(x)$ of degree d , with the property that when restricted to lines parallel to a particular direction e , the resulting univariate polynomial has all its d roots real.

Furthermore, we have also seen that every polynomial of the form

$$p(x) = \det(x_1 A_1 + \cdots + x_n A_n), \quad (2)$$

where $A_i \in \mathcal{S}^d$ and $A_1 \succ 0$, is hyperbolic with respect to the $(1, 0, \dots, 0)$ direction.

A 1958 conjecture by Peter Lax [Lax58], asks whether the converse is true in the case $n = 3$ (i.e., trivariate polynomials). In other words, is it true that for every hyperbolic polynomial $p(x)$ in three variables of degree d , there exist three symmetric matrices $\{A_1, A_2, A_3\} \subset \mathcal{S}^d$ for which (2) holds?

As a first step towards answering this question, let us verify that this at least makes sense in terms of dimension counting. As we have seen, the dimension of the set of hyperbolic polynomials in three variables ($n = 3$) and degree d is equal to $\binom{n+d-1}{d} = \binom{d+2}{2}$. On the other hand, for a polynomial of the form (2), by an appropriate similarity transform we can always assume without loss of generality $A_1 = a_0 I_d$, and $A_2 = \text{diag}(a_1, \dots, a_d)$. The total number of parameters is then $1 + d + \binom{d+1}{2}$, which is exactly equal to $\binom{d+2}{2}$. Of course, this by itself does not prove the result, but it shows that it is certainly possible.

4 Relating SDP-representable sets and hyperbolic polynomials

As we will see shortly, these two apparently different problems are in fact one and the same. Before showing this, let us consider one additional necessary condition for a set in \mathbb{R}^2 to be SDP-representable. For later reference, we first define the following notion:

Definition 2. A polynomial $p \in \mathbb{R}[x]$ is a real zero polynomial if for every $x \in \mathbb{R}^n$, $p(tx) = 0$ implies that t is real.

Recall that the boundary of a set described by (1) is determined by the zero set of the polynomial $\det(I + xB + yC)$. Consider now any line passing through the origin, i.e., of the form $(x, y) = (\beta t, \gamma t)$. We have then

$$\det[I + (\beta B + \gamma C)t] = 0,$$

and this univariate polynomial in t has exactly d real roots (namely, the negative inverse of the eigenvalues of $\beta B + \gamma C$). In terms of the notation just introduced, the polynomial defined by $\det(I + xB + yC)$ is

a real zero polynomial. Equivalently, for every set of the form (1), is it always the case that every line through the origin intersects (the Zariski closure¹ of) the boundary of the set exactly d times.

In the preceding, our starting point was directly a determinantal representation as in (1). It can be shown (see [HV]) that if we start directly from a given set that admits an SDP representation, we can precisely characterize a unique minimal polynomial that defines the boundary of the set.

Hence, this gives us an additional necessary condition ([HV]) for SDP representability:

- **Rigid convexity:** Consider a set in \mathbb{R}^2 , with the origin in the interior. Every line that passes through the origin must intersect the polynomial defining the boundary exactly d times (counting multiplicities, and points at infinity), where d is the degree of the boundary polynomial.

This additional requirement is quite strong, and immediately allows us to discard sets for which the previous conditions were satisfied.

Example 3. Consider the set described by $x^4 + y^4 \leq 1$; see Figure 1. It clearly satisfies the first four necessary conditions. However, if we consider any line through the origin, it will intersect the defining polynomial only two times, instead of the four required by the rigid convexity condition. Thus, this set is not rigidly convex, and hence does not admit a (non-lifted) semidefinite representation.

5 Characterization

It should be apparent that the rigid convexity condition looks very similar to the hyperbolicity property of a polynomial. In fact, they are exactly the *same* condition, provided we redefine things accordingly [LPR05]. As we will see, this equivalence will make explicit the connection between the Helton & Vinnikov characterization of SDP-representable sets and the Lax conjecture described earlier.

Theorem 4 ([LPR05]). *If $p \in \mathbb{R}[x, y, z]$ is a polynomial of degree d , hyperbolic with respect to $e = (0, 0, 1)$ and that satisfies $p(e) = 1$, then the polynomial in $\mathbb{R}[x, y]$ defined by $q(x, y) = p(x, y, 1)$ is a real zero polynomial of degree no more than d , and satisfying $q(0, 0) = 1$.*

Conversely, if $q \in \mathbb{R}[x, y]$ is a real zero polynomial of degree d satisfying $q(0, 0) = 1$, then the polynomial defined by

$$p(x, y, z) = z^d q\left(\frac{x}{z}, \frac{y}{z}\right)$$

is a hyperbolic polynomial of degree d with respect to $e = (0, 0, 1)$, and $p(e) = 1$.

In their paper [HV], Helton and Vinnikov proved that the rigid convexity condition fully characterizes the plane sets that are semidefinite representable.

Theorem 5 ([HV]). *If $p(x, y)$ is a real zero polynomial of degree d with $p(0) > 0$, then the closure of the connected component of $p(x) > 0$ containing the origin admits a representation as in (1).*

For hyperbolic cones, we have shown earlier that the specific hyperbolicity direction e does not matter too much (as long as it belong to the hyperbolicity cone). Similarly, it can be shown that when checking the real zero condition we can choose any point in the interior of the set, not necessarily the origin.

Combining these two results, the truth of the Lax conjecture follows:

Theorem 6. *Every hyperbolic polynomial in three variables admits a determinantal representation of the type (2). If coordinates are chosen so that $e = (1, 0, 0)$, then we can choose $A_1 = I$.*

¹The Zariski topology on \mathbb{C}^n can be defined in terms of its closed sets, which are the algebraic varieties, i.e., the vanishing set of a finite set of polynomial equations. The Zariski topology is a very weak topology, and is quite different from the usual topology in \mathbb{C}^n . For instance, the Zariski closure of the open interval $(0, 1)$ is equal to \mathbb{C} . The Zariski topology is not Hausdorff, i.e., distinct points do not always have disjoint neighborhoods.




Figure 2: Convex set defined by $\{3 + x - x^3 - 3x^2 - 2y^2 \geq 0, x \geq -1\}$. A semidefinite representation is given in (3).

An interesting issue concerns the possibility of a constructive approach. In other words, given a hyperbolic polynomial in three variables, how to effectively obtain matrices A_i that give a determinantal representation? While “explicit” formulae for these matrices are given in [HV] in terms of objects that are quite complicated to compute (namely, theta functions of Jacobian varieties), it seems likely that a more elementary formulation exists.

5.1 Example

As an illustration, consider the convex set shown in Figure 3, which corresponds to the “oval” of the elliptic curve given by $3 + x - x^3 - 3x^2 - 2y^2 = 0$. This set satisfies the real zero condition, since every line that passes through a point in the interior of the set intersects the polynomial defining the boundary at exactly three points (if the lines are vertical, then the corresponding intersections are at infinity).

Homogenizing this polynomial, we obtain $p(x, y, z) = 3z^3 + xz^2 - x^3 - 3x^2z - 2y^2z$; the corresponding zero set is given in Figure 3. As we can see (and proved earlier), the section corresponding to the plane $z = 1$ is exactly the zero set of the original polynomial. Furthermore, lines parallel to the hyperbolicity direction e are projectively mapped into lines in this plane that go through the origin. Hence, the number of intersections (and thus, real roots) is preserved.

The theorem presented above promises the existence of a semidefinite representation. In this case, one such representation is:

$$\begin{bmatrix} x+1 & 0 & y \\ 0 & 2 & -x-1 \\ y & -x-1 & 2 \end{bmatrix} \succeq 0, \quad (3)$$

with the corresponding determinantal representation of the hyperbolic polynomial being:

$$p(x, y, z) = \det \begin{bmatrix} x+z & 0 & y \\ 0 & 2z & -x-z \\ y & -x-z & 2z \end{bmatrix}. \quad (4)$$

References

- [HV] J.W. Helton and V. Vinnikov. Linear matrix inequality representation of sets. Preprint. Available from <http://arxiv.org/abs/math.OA/0306180>.




Figure 3: The polynomial $3z^3 + xz^2 - x^3 - 3x^2z - 2y^2z = 0$ and corresponding hyperbolicity cone.

- [Lax58] P. D. Lax. Differential equations, difference equations and matrix theory. *Comm. Pure Appl. Math.*, 11:175–194, 1958.
- [LPR05] A. S. Lewis, P. A. Parrilo, and M. V. Ramana. The Lax conjecture is true. *Proc. Amer. Math. Soc.*, 133(9):2495–2499, 2005.

Lecture 9

Lecturer: Pablo A. Parrilo

Scribe: ???

In this lecture, we study first a relatively simple type of polynomial equations, namely *binomial equations*. As we will see, in this case there exists a quite efficient solution method. We define next an important geometric and combinatorial object associated with every multivariate polynomial, called the *Newton polytope*. Finally, we put together these two notions in the formulation of a family of bounds on the number of solutions of systems of polynomial equations. Our presentation of the material here is inspired by [Stu02, Chapter 3].

1 Binomial equations

We introduce in this section a particular kind of polynomial equations, that have nice computational properties. A *binomial* system of polynomial equations is one where each equation has only two terms. We also assume that the system has only a finite number of solutions, i.e., the solution set is a finite set of points in \mathbb{C}^n . We are interested in determining the exact number of solutions, and in efficient computational procedures for solving the system.

Let's start with an example. Consider the binomial system given by

$$\begin{aligned} 8x^2y^3 - 1 &= 0 \\ 2x^3y^2 - yx &= 0. \end{aligned} \tag{1}$$

If we assume that the solutions satisfy $x \neq 0, y \neq 0$, then we can put these equations in the more symmetric form

$$\begin{aligned} 8x^2y^3 &= 1 \\ 2x^2y &= 1. \end{aligned} \tag{2}$$

Now, by dividing the first equation by the second one, we obtain $4y^2 = 1$, which has two solutions ($y = \frac{1}{2}$ and $y = -\frac{1}{2}$). Substituting into the resulting equations for every value of y we have two corresponding values of x , so the system has a big total of four complex solutions.

Let's try to understand in a bit more detail what manipulations we were performing here. For this, let's define the integer matrix

$$B = \begin{bmatrix} 2 & 3 \\ 2 & 1 \end{bmatrix}$$

corresponding to the exponents in (2). Notice that when we divided the two equations, that is equivalent to an elementary row operation in the matrix B , namely subtracting the second row of B from the first one. Thus, the operations we have done can be understood as the matrix multiplication $UB = C$, where

$$U = \begin{bmatrix} 0 & 1 \\ 1 & -1 \end{bmatrix}, \quad C = \begin{bmatrix} 2 & 1 \\ 0 & 2 \end{bmatrix}.$$

The fact that the matrix C is lower triangular, is what allows us to start solving the system for y , and then backsolving for the other variable.

It is not too difficult to understand from this example how to generalize this. Let $\mathbb{C}^* = \mathbb{C} \setminus \{0\}$, and consider a system of binomial equations in n variables, where we are interested in computing (or bounding the number of) solutions in $(\mathbb{C}^*)^n$. We can always put the system in the normalized form in (2). Notice that, in general, the entries of the integer B could be either positive or negative (i.e., we write polynomials in x_i and x_i^{-1} , which is fine since $x_i \neq 0$).

Then, a well-known result in integer linear algebra (the Hermite normal form of an integer matrix) guarantees the existence of a matrix $U \in SL_n(\mathbb{Z})$ (an integer matrix, with determinant equal to one),




Figure 1: Newton polytope of the polynomial $p(x, y) = 5 - xy - x^2y^2 + 3y^2 + x^4$.

such that $C = UB$ is a lower triangular matrix. We can then use this expression to obtain values for the last variable, and backsolve to obtain all solutions.

How can we determine the number of solutions from this factorization? When backsubstituting using C , at each step we have to solve an equation of the type $x_i^{c_{ii}} = d_i$, and thus the current number of possible solutions is multiplied by $|c_{ii}|$. Therefore, the total number of solutions in $(\mathbb{C}^*)^n$ will then be equal to $|\det(C)| = |\det(U) \det(B)| = |\det(B)|$.

Remark 1. To compute the Hermite normal form of an integer matrix in Maple, you can use the command *ihermite*. In Mathematica, use instead *HermiteNormalForm*.

2 Newton polytopes

Many of the polynomial systems that appear in practice are far from being “generic,” but rather present a number of structural features that, when properly exploited, allow for much more efficient computational techniques. This is quite similar to the situation in numerical linear algebra, where there is a big difference in performance between algorithms that take into account the sparsity structure of a matrix and those that do not. For matrices, the standard notion of sparsity is relatively straightforward, and relates mostly to the number of nonzero coefficients. In computational algebra, however, there exists a much more refined notion of sparsity that refers not only to the number of zero coefficients of a polynomial, but also to the underlying combinatorial structure.

This notion of sparsity for multivariate polynomials is usually presented in terms of the *Newton polytope* of a polynomial, defined below.

Definition 2. Consider a multivariate polynomial $p(x_1, \dots, x_n) = \sum_{\alpha} c_{\alpha} x^{\alpha}$. The Newton polytope of p , denoted by $\text{New}(f)$, is defined as the convex hull of the set of exponents α , considered as vectors in \mathbb{R}^n .

Thus, the Newton polytope of a polynomial always has integer extreme points, given by a subset of the exponents of the polynomial.

Example 3. Consider the polynomial $p(x, y) = 5 - xy - x^2y^2 + 3y^2 + x^4$. Its Newton polytope $\text{New}(f)$, displayed in Figure 1, is the convex hull of the points $(0, 0)$, $(1, 1)$, $(2, 2)$, $(0, 2)$, $(4, 0)$.

Example 4. Consider the polynomial $p(x, y) = 1 - x^2 + xy + 4y^4$. Its Newton polytope $\text{New}(p)$ is the triangle in \mathbb{R}^2 with vertices $\{(0, 0), (2, 0), (0, 4)\}$.

Newton polytopes are an essential tool when considering polynomial arithmetic because of the following fundamental identity:

$$\text{New}(g \cdot h) = \text{New}(g) + \text{New}(h),$$

where $+$ denotes the Minkowski addition of polytopes.

Example 5. Let $p(a, b, c, d) = (a^4 + 1)(b^4 + 1)(c^4 + 1)(d^4 + 1) + 2a + 3b + 4c + 5d$. Its Newton polytope is the hypercube in \mathbb{R}^4 of side length equal to 4, and with opposing vertices at $(0, 0, 0, 0)$ and $(4, 4, 4, 4)$.

It is a general theme in computational algebra that the complexity of many problems involving polynomials is directly related to some measure of the size of the corresponding Newton polytopes. We discuss an example below, in terms of the number of solutions of polynomial equations. We will encounter Newton polytopes again later in the course, when discussing the semidefinite characterization of polynomials that are sums of squares.

3 The Bézout and BKK bounds

Consider a system of two polynomial equations, $p(x, y) = 0$, $q(x, y) = 0$. As we have seen in previous lectures, we can solve this by computing the resultant of the polynomials p and q with respect to either variable, and then factorizing the corresponding univariate polynomial. If the degree of the polynomials is d_1 and d_2 , respectively, then the degree of the resultant is bounded by $d_1 \cdot d_2$, and thus the number of zeros of the system is at most this number.

However, when the polynomials p and q are sparse (in the sense defined earlier) then the number of solutions can be much smaller. For instance, the system

$$\begin{aligned} a + bx + cy + dy^2 &= 0 \\ ex + fy + gxy &= 0 \end{aligned} \tag{3}$$

has, for a generic choice of the coefficients $\{a, \dots, g\}$, exactly three complex roots, while the bound based on the individuals degrees (usually called the Bézout bound) will give a total of $2 \times 2 = 4$. As we will see, much sharper bounds can be obtained by considering the Newton polytopes of the individual equations.

To introduce the main theorem, we need to introduce the following concept, that generalizes the notion of volume of a polytope, to a collection of them.

Definition 6. Consider polytopes $P_1, \dots, P_n \subset \mathbb{R}^n$, nonnegative scalars $\lambda_1, \dots, \lambda_n$, and let $V(\lambda) = \text{Vol}(\lambda_1 P_1 + \dots + \lambda_n P_n)$. It can be shown that $V(\lambda)$ is a homogeneous polynomial of degree n . The mixed volume $MV(P_1, \dots, P_n)$ is the coefficient of this polynomial, corresponding to the monomial $\lambda_1 \lambda_2 \dots \lambda_n$.

The main result in this area, with different versions due to Bernstein, Kouchnirenko, and Khovanskii, relates the number of solutions of a sparse polynomial system with the mixed volume of the Newton polytopes of the individual equations. Formally, we have

Theorem 7. The number of solutions in $(\mathbb{C}^*)^n$ of a sparse polynomial system of n equations and n unknowns is less than or equal to the mixed volume of the n Newton polytopes. If the coefficients are “generic” enough, then the upper bound is achieved.

The basic idea behind the derivation of the theorem is to introduce an additional parameter t in the equations, in such a way that for $t = 1$ we have the original system, while for $t = 0$ the system is binomial, which as we have seen can be solved in an efficient manner. This process is usually called a *toric deformation*, and is somewhat similar in spirit to the homotopies used in interior point methods. To make our words a bit more precise, an important fact is that we will not deform to just one binomial system, but actually to a collection of them, given by what is called a *mixed subdivision* of the sum of Newton polytopes. The important fact is that the sum of the number of roots of all these binomial systems is exactly equal to the mixed volume of the collection of polytopes.

Example 8. Consider the univariate polynomial

$$p(x) = a_n x^n + a_{n-1} x^{n-1} + \dots + a_m x^m,$$

where $n \geq m$. It is clear that the Newton polytope is the line segment with endpoints in n and m . The mixed volume (in this case, just the volume) is equal to $n - m$. Thus, the BKK bound for this polynomial is equal to $n - m$, which is clearly exact for generic choices of the coefficients.

Example 9. *Let us consider again the example discussed in (1). The Newton polytope of the first polynomial is the line segment with endpoints $(0,0)$ and $(2,3)$, while the second one has endpoints $(1,1)$ and $(3,2)$. If we denote these by P_1 and P_2 , it is easy to see that*

$$\text{Vol}(\lambda_1 P_1 + \lambda_2 P_2) = 4\lambda_1 \lambda_2,$$

and thus the mixed volume of (P_1, P_2) is equal to 4, which is the number of solutions of (1).

4 Application: Nash equilibria

We can use the results described, to give a bound on the number of Nash equilibria of a game. For simplicity, consider the three player case, where each player has two pure strategies. We are interested here only in totally mixed equilibria, i.e., those where the players randomize among all their pure strategies with nonzero probability (if this is not the case, then by eliminating the never played strategies we can reduce the game to the totally mixed case). Thus, the mixed strategies can be parametrized in terms of three variables $a, b, c \in (0, 1)$, representing the probabilities with which they play their different strategies.

It can be shown that the Nash equilibrium condition result in a polynomial system of the structure

$$\begin{aligned} p_{11}bc + p_{12}b + p_{13}c + p_{14} &= 0 \\ p_{21}ca + p_{22}c + p_{23}a + p_{24} &= 0 \\ p_{31}ab + p_{32}a + p_{33}b + p_{34} &= 0, \end{aligned} \tag{4}$$

where the coefficients p_{ij} are explicit linear functions of the payoffs. The mixed volume of the Newton polytopes of these three equations is equal to 2, so the maximum number of totally mixed Nash equilibria that a three-player, two-strategy game can have is equal to two.

Theorem 10 ([Stu02, p.82]). *The maximum number of isolated totally mixed Nash equilibria for an n -person game where each player has two pure strategies is equal to the mixed volume of the n facets of the n cube. This number is the closest integer to $n!/e$.*

There are extensions of this result to the graphical case; see [Stu02] and the references therein for details.

References

[Stu02] B. Sturmfels. *Solving Systems of Polynomial Equations*. AMS, Providence, R.I., 2002.

Lecture 10

Lecturer: Pablo A. Parrilo

Scribe: ???

In this lecture we begin our study of one of the main themes of the course, namely the relationships between polynomials that are sums of squares and semidefinite programming.

1 Nonnegativity and sums of squares

Recall from a previous lecture the definition of a polynomial being a sum of squares.

Definition 1. A univariate polynomial $p(x)$ is a sum of squares (SOS) if there exist $q_1, \dots, q_m \in \mathbb{R}[x]$ such that

$$p(x) = \sum_{k=1}^m q_k^2(x). \quad (1)$$

If a polynomial $p(x)$ is a sum of squares, then it obviously satisfies $p(x) \geq 0$ for all $x \in \mathbb{R}$. Thus, a SOS condition is a sufficient condition for global nonnegativity.

As we have seen, in the univariate case, the converse is also true:

Theorem 2. A univariate polynomial is nonnegative if and only if it is a sum of squares.

As we will see, there is a very direct link between sum of squares conditions on polynomials and semidefinite programming. We study first the univariate case.

2 Sums of squares and semidefinite programming

Consider a polynomial $p(x)$ of degree $2d$ that is a sum of squares, i.e., it can be written as in (1). Notice that the degree of the polynomials q_k is at most equal to d , since the highest term of each q_k^2 is positive, and thus there cannot be any cancellation in the highest power of x . Then, we can write

$$\begin{bmatrix} q_1(x) \\ q_2(x) \\ \vdots \\ q_m(x) \end{bmatrix} = V \begin{bmatrix} 1 \\ x \\ \vdots \\ x^d \end{bmatrix}, \quad (2)$$

where $V \in \mathbb{R}^{m \times (d+1)}$, and its k th row contains the coefficients of the polynomial q_k . For future reference, let $[x]_d$ be the vector in the right-hand side of (2). Consider now the matrix $Q = V^T V$. We then have $p(x) = \sum_{k=1}^m q_k^2(x) = (V[x]_d)^T (V[x]_d) = [x]_d^T V^T V [x]_d = [x]_d^T Q [x]_d$.

Conversely, assume there exists a symmetric positive definite Q , for which $p(x) = [x]_d^T Q [x]_d$. Then, by factorizing $Q = V^T V$ (e.g., via Choleski, or square root factorization), we arrive at a SOS decomposition of p .

We formally express this in the following lemma, that gives a direct relation between positive semidefinite matrices and a sum of squares condition.

Lemma 3. Let $p(x)$ be a univariate polynomial of degree $2d$. Then, $p(x)$ is nonnegative (or SOS) if and only if there exists $Q \in \mathcal{S}_+^{d+1}$ that satisfies

$$p(x) = [x]_d^T Q [x]_d.$$

Indexing the rows and columns of Q by $\{0, \dots, d\}$, we have:

$$[x]_d^T Q [x]_d = \sum_{j=0}^d \sum_{k=0}^d Q_{jk} x^{j+k} = \sum_{i=0}^{2d} \left(\sum_{j+k=i} Q_{jk} \right) x^i$$

Thus, for this expression to be equal to $p(x)$, it should be the case that

$$p_i = \sum_{j+k=i} Q_{jk}, \quad i = 0, \dots, 2d. \quad (3)$$

This is a system of $2d + 1$ linear equations between the entries of Q and the coefficients of $p(x)$. Thus, since Q is simultaneously constrained to be positive semidefinite, and to belong to a particular affine subspace, a SOS condition is exactly equivalent to a semidefinite programming problem.

Lemma 4. *A polynomial $p(x) = \sum_{i=0}^{2d} p_i x^i$ is a sum of squares if and only if there exists $Q \in \mathcal{S}_+^{d+1}$ satisfying (3). This is a semidefinite programming problem.*

3 Applications and extensions

We discuss first a few applications of the SDP characterization of nonnegative polynomials, followed by several extensions.

3.1 Optimization

Our first application concerns the global optimization of a univariate polynomial $p(x)$. Rather than focusing on computing an x_* for which $p(x_*)$ is as small as possible, we attempt first to obtain a good (or the best) lower bound on its optimal value. It is easy to see that a number γ is a global lower bound of a polynomial $p(x)$, if and only if the polynomial $p(x) - \gamma$ is nonnegative, i.e.,

$$p(x) \geq \gamma \quad \forall x \in \mathbb{R} \quad \Longleftrightarrow \quad p(x) - \gamma \geq 0 \quad \forall x \in \mathbb{R}.$$

Notice that the polynomial $p(x) - \gamma$ has coefficients that depend affinely on γ . Consider now the optimization problem defined by

$$\max \gamma \quad \text{s.t.} \quad p(x) - \gamma \text{ is SOS.}$$

It should be clear that this is a *convex* problem, since the feasible set is defined by an infinite number of linear inequalities. Its optimal solution γ_* is equal to the global minimum of the polynomial, $p(x_*)$. Furthermore, using Lemma 4, we can easily write this as a semidefinite programming problem. We can thus obtain the global minimum of a univariate polynomial, by solving an SDP problem. Notice also that at optimality, we have $0 = p(x_*) - \gamma_* = \sum_{k=1}^m q_k^2(x_*)$, and thus all the q_k simultaneously vanish at x_* , which gives a way of computing the optimal solution x_* . As we shall see later, we can also obtain this solution directly from the dual problem, by using complementary slackness.

Notice that even though $p(x)$ may be highly nonconvex, we are nevertheless effectively computing its global minimum.

3.2 Nonnegativity on intervals

We have seen how to characterize a univariate polynomial that is nonnegative on $(-\infty, \infty)$ in terms of SDP conditions. But what if we are interested in polynomials that are nonnegative only in an interval (either finite, or semi-infinite)? As explained below, we can use very similar ideas, and two classical characterizations, usually associated to the names Pólya-Szegő, Fekete, or Markov-Lukacs. The basic results are the following:

Theorem 5. *The polynomial $p(x)$ is nonnegative on $[0, \infty)$, if and only if it can be written as*

$$p(x) = s(x) + x t(x),$$

where $s(x), t(x)$ are SOS. If $\deg(p) = 2d$, then we have $\deg(s) \leq 2d$, $\deg(t) \leq 2d - 2$, while if $\deg(p) = 2d + 1$, then $\deg(s) \leq 2d$, $\deg(t) \leq 2d$.

Theorem 6. *Let $a < b$. Then, $p(x)$ is nonnegative on $[a, b]$, if and only if it can be written as*

$$\begin{cases} p(x) = s(x) + (x - a)(b - x)t(x), & \text{if } \deg(p) \text{ is even} \\ p(x) = (x - a)s(x) + (b - x)t(x), & \text{if } \deg(p) \text{ is odd} \end{cases}$$

where $s(x), t(x)$ are SOS. In the first case, we have $\deg(p) = 2d$, and $\deg(s) \leq 2d$, $\deg(t) \leq 2d - 2$. In the second, $\deg(p) = 2d + 1$, and $\deg(s) \leq 2d$, $\deg(t) \leq 2d$.

Notice that in both of these results, one direction of the implication is evident.

3.3 Rational functions

What happens if we want to minimize a univariate rational function, rather than a polynomial? Consider a rational function given as a quotient of polynomials $p(x)/q(x)$, where $q(x)$ is strictly positive (why?). Then, we have

$$\frac{p(x)}{q(x)} \geq \gamma \quad \Leftrightarrow \quad p(x) - \gamma q(x) \geq 0,$$

and therefore we can find the global minimum of the rational function by solving

$$\max \gamma \quad \text{s.t.} \quad p(x) - \gamma q(x) \text{ is SOS.}$$

The constrained case (i.e., over finite or semi-infinite intervals) are very similar, and can be formulated using the results in the Section 3.2. The details are left for the exercises.

4 Multivariate polynomials

For polynomials in more than one variable, it is no longer true that nonnegativity is equivalent to a sum of squares condition. In fact, for polynomials of degree greater than or equal to four, deciding polynomial nonnegativity is an NP-hard problem (as a function of the number of variables).

More than a century ago, David Hilbert showed that equality between the set of nonnegative and SOS polynomials holds only in the following three cases:

- Univariate polynomials (i.e., $n = 1$)
- Quadratic polynomials ($2d = 2$)
- Bivariate quartics ($n = 2, 2d = 4$)

For all other cases, there always exist nonnegative polynomials that are *not* sums of squares. A classical counterexample is the bivariate sextic ($n = 2, 2d = 6$) due to Motzkin, given by (in dehomogenized form)

$$M(x, y) = x^4 y^2 + x^2 y^4 + 1 - 3x^2 y^2.$$

This polynomial is nonnegative, but is not a sum of squares. We will prove both facts later. An excellent account of much of the classical work in this area has been provided by Bruce Reznick [Rez00].

4.1 SDP formulation

Essentially the same construction we have seen in Lemma 4 applies to the multivariate case. In this case, we consider polynomials of degree $2d$ in n variables. In the dense case, i.e., when the polynomial is not sparse, the number of coefficients is equal to $\binom{n+2d}{2d}$. If we let $p(x) = \sum_{\alpha} p_{\alpha} x^{\alpha}$, and indexing the matrix Q by the $\binom{n+d}{d}$ monomials in n variables of degree d , we have the SDP conditions on $Q \in \mathcal{S}_+^{\binom{n+d}{d}}$:

$$Q \succeq 0, \quad p_{\alpha} = \sum_{\beta+\gamma=\alpha} Q_{\beta\gamma}. \quad (4)$$

We have exactly $\binom{n+2d}{2d}$ linear equations, one per each coefficient of $p(x)$. As before, these conditions are affine conditions relating the entries of Q and the coefficients of $p(x)$. Thus, we can decide membership to, or optimize over, the set of SOS polynomials by solving a semidefinite programming problem.

4.2 Using the Newton polytope

Recall that we have defined in a previous lecture the *Newton polytope* of a polynomial $p(x) \in \mathbb{R}[x_1, \dots, x_n]$ as the convex hull of the set of exponents appearing in p . This allowed us to introduce a notion of sparseness for a polynomial, related to the size of its Newton polytope. Sparsity (in this algebraic sense) allows a notable reduction in the computational cost of checking sum of squares conditions of multivariate polynomials. The reason is the following theorem due to Reznick:

Theorem 7 ([Rez78], Theorem 1). *If $p(x) = \sum q_i(x)^2$, then $\text{New}(q_i) \subseteq \frac{1}{2}\text{New}(p)$.*

In other words, this theorem allows us, without loss of generality, to restrict the set of monomials appearing in the representation (4) to those in the Newton polytope of p , scaled by a factor of $\frac{1}{2}$. This reduces the size of the corresponding matrix Q , thus simplifying the SDP problem.

Example 8. *Consider the following polynomial:*

$$p = (w^4 + 1)(x^4 + 1)(y^4 + 1)(z^4 + 1) + 2w + 3x + 4y + 5z.$$

The polynomial p has degree $2d = 16$, and four independent variables ($n = 4$). A naive approach, along the lines described earlier, would require a matrix Q of size $\binom{n+2d}{2d} = 495$. However, the Newton polytope of p is easily seen to be the four dimensional hypercube with vertices in $(0, 0, 0, 0)$ and $(4, 4, 4, 4)$. Therefore, the polynomials q_i in the SOS decomposition of p will have at most $3^4 = 81$ distinct monomials, and as a consequence the full decomposition can be computed by solving a much smaller SDP.

5 Duality and density

In the next lecture, we will revisit the sum of squares construction, but emphasizing this time the dual side, and its appealing measure-theoretic interpretation. We will also review some recent results on the relative density of the cones of nonnegative polynomials and SOS.

References

- [Rez78] B. Reznick. Extremal PSD forms with few terms. *Duke Mathematical Journal*, 45(2):363–374, 1978.
- [Rez00] B. Reznick. Some concrete aspects of Hilbert’s 17th problem. In *Contemporary Mathematics*, volume 253, pages 251–272. American Mathematical Society, 2000.

Lecture 11

Lecturer: Pablo A. Parrilo

Scribe: ???

In this lecture we continue our study of SOS polynomials. After presenting a couple of applications, we focus here on the dual side, and provide a natural probabilistic interpretation of the corresponding problem. We further present some recent results on the density of the cone of SOS polynomials relative to that of the nonnegative polynomials.

1 SOS applications

1.1 Lyapunov functions

The possibility of reformulating conditions for a polynomial to be a sum-of-squares as an SDP is very useful, since we can use the SOS property in a control context as a convenient sufficient condition for polynomial nonnegativity. Recent work has applied the sum-of-squares approach to the problem of finding a Lyapunov function for nonlinear systems [Par00, PP02]. This approach allows one to search over affinely parametrized polynomial or rational Lyapunov functions for systems with dynamics of the form

$$\dot{x}_i(t) = f_i(x(t)) \quad \text{for all } i = 1, \dots, n$$

where the functions f_i are polynomials or rational functions. Then the condition that the Lyapunov function be positive, and that its Lie derivative be negative, are both directly imposed as sum-of-squares constraints in terms of the coefficients of the Lyapunov function.

As an example, consider the following system:

$$\begin{aligned}\dot{x} &= -x + (1+x)y \\ \dot{y} &= -(1+x)x.\end{aligned}$$

Using SOSTOOLS [PPP05] we easily find a quartic polynomial Lyapunov function, which after rounding (for purely cosmetic reasons) is given by

$$V(x, y) = 6x^2 - 2xy + 8y^2 - 2y^3 + 3x^4 + 6x^2y^2 + 3y^4.$$

It can be readily verified that both $V(x, y)$ and $(-\dot{V}(x, y))$ are SOS, since

$$V = \begin{bmatrix} x \\ y \\ x^2 \\ xy \\ y^2 \end{bmatrix}^T \begin{bmatrix} 6 & -1 & 0 & 0 & 0 \\ -1 & 8 & 0 & 0 & -1 \\ 0 & 0 & 3 & 0 & 0 \\ 0 & 0 & 0 & 6 & 0 \\ 0 & -1 & 0 & 0 & 3 \end{bmatrix} \begin{bmatrix} x \\ y \\ x^2 \\ xy \\ y^2 \end{bmatrix}, \quad -\dot{V} = \begin{bmatrix} x \\ y \\ x^2 \\ xy \end{bmatrix}^T \begin{bmatrix} 10 & 1 & -1 & 1 \\ 1 & 2 & 1 & -2 \\ -1 & 1 & 12 & 0 \\ 1 & -2 & 0 & 6 \end{bmatrix} \begin{bmatrix} x \\ y \\ x^2 \\ xy \end{bmatrix},$$

and the matrices in the expression above are positive definite. Similar approaches may also be used for finding Lyapunov functionals for certain classes of hybrid systems.

1.2 Entangled states in quantum mechanics

The state of a finite-dimensional quantum system can be described in terms of a positive semidefinite Hermitian matrix, called the *density matrix*. An important property of a bipartite quantum state ρ is whether or not it is *separable*, which means that it can be written as a convex combination of tensor products of rank one matrices, i.e.,

$$\rho = \sum_i p_i (x_i x_i^T) \otimes (y_i y_i^T), \quad p_i \geq 0, \quad \sum_i p_i = 1,$$

where for simplicity we have restricted ρ, x_i, y_i to be real. Here $x_i \in \mathbb{R}^{n_1}$, $y_i \in \mathbb{R}^{n_2}$, and $\rho \in \mathcal{S}_+^{n_1 n_2}$.

How to recognize if a state is entangled or not?

Complete

ToDo

2 Moments

Consider a nonnegative measure μ on \mathbb{R} (or if you prefer, a real-valued random variable X). We can then define the *moments*, which are the expectation of powers of X .

$$\mu_k := E[X^k] = \int x^k d\mu \quad (1)$$

What constraints, if any, should the μ_k satisfy? Is it true that for any set of numbers $\mu_0, \mu_1, \dots, \mu_k$, there always exists a nonnegative measure having exactly these moments?

It should be apparent that some conditions are required. For instance, consider (1) for an even value of k . Since the measure μ is nonnegative, it is clear that in this case we have $\mu_k \geq 0$.

However, that's clearly not enough, and more restrictions should hold. A simple one can be derived by recalling the relationship between the first and second moments and the variance of a random variable, i.e., $\text{var}(X) = E[X^2] - E[X]^2 = \mu_2 - \mu_1^2$. Since the variance is always nonnegative, we should have $\mu_2 - \mu_1^2 \geq 0$.

How to systematically derive conditions of this kind? Notice that the previous inequality can be obtained by noticing that for all a, b ,

$$0 \leq E[(a + bX)^2] = a^2 + 2abE[X] + b^2E[X^2] = \begin{bmatrix} a \\ b \end{bmatrix}^T \begin{bmatrix} 1 & \mu_1 \\ \mu_1 & \mu_2 \end{bmatrix} \begin{bmatrix} a \\ b \end{bmatrix},$$

which implies that the 2×2 matrix above must be positive semidefinite. Interestingly, the inequality obtained earlier is exactly equal to the determinant of this matrix.

Exactly the same procedure can be done for higher-order moments. Proceeding this way, we have that the higher order moments must always satisfy:

$$\begin{bmatrix} 1 & \mu_1 & \mu_2 & \cdots & \mu_d \\ \mu_1 & \mu_2 & \mu_3 & \cdots & \mu_{d+1} \\ \mu_2 & \mu_3 & \mu_4 & \cdots & \mu_{d+2} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \mu_d & \mu_{d+1} & \mu_{d+2} & \cdots & \mu_{2d} \end{bmatrix} \succeq 0. \quad (2)$$

Notice that the diagonal elements correspond to the even-order moments, which should obviously be nonnegative.

Necessary and sufficient, multivariate case

ToDo

Remark 1. For unbounded intervals, the SDP conditions characterize the closure of the set of moments, but not necessarily the whole set. As an example, consider the set of moments given by $\mu = (1, 0, 0, 0, 1)$, corresponding to the Hankel matrix

$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}.$$

Although the matrix above is PSD, it is not hard to see that there is no nonnegative measure corresponding to those moments. However, the parametrized atomic measure given by

$$\mu_\varepsilon = \frac{\varepsilon^4}{2} \cdot \delta(x + \frac{1}{\varepsilon}) + (1 - \varepsilon^4) \cdot \delta(x) + \frac{\varepsilon^4}{2} \cdot \delta(x - \frac{1}{\varepsilon})$$

has as first five moments $(1, 0, \varepsilon^2, 0, 1)$, and thus as $\varepsilon \rightarrow 0$ the corresponding Hankel matrix is the one given above.

2.1 Nonnegative measures on intervals

Just like we did for the case of polynomials nonnegative on intervals, we can similarly obtain a necessary and sufficient characterization for moments. For simplicity, we present below only one particular case, corresponding to the interval $[-1, 1]$.

Lemma 2. *There exists a nonnegative measure in $[-1, 1]$ with moments $(\mu_0, \mu_1, \dots, \mu_{2d+1})$ if and only if*

$$\begin{bmatrix} \mu_0 & \mu_1 & \mu_2 & \cdots & \mu_d \\ \mu_1 & \mu_2 & \mu_3 & \cdots & \mu_{d+1} \\ \mu_2 & \mu_3 & \mu_4 & \cdots & \mu_{d+2} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \mu_d & \mu_{d+1} & \mu_{d+2} & \cdots & \mu_{2d} \end{bmatrix} \pm \begin{bmatrix} \mu_1 & \mu_2 & \mu_4 & \cdots & \mu_{d+1} \\ \mu_2 & \mu_3 & \mu_5 & \cdots & \mu_{d+2} \\ \mu_3 & \mu_4 & \mu_6 & \cdots & \mu_{d+3} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ \mu_{d+1} & \mu_{d+2} & \mu_{d+3} & \cdots & \mu_{2d+1} \end{bmatrix} \succeq 0. \quad (3)$$

Notice that the necessity is clear, since it follows from consideration of the quadratic form (in the a_i):

$$0 \leq E \left[(1 \pm X) \left(\sum_{i=0}^d a_i X^i \right)^2 \right] = \sum_{j=0}^d \sum_{k=0}^d (\mu_{j+k} \pm \mu_{j+k+1}) a_j a_k,$$

where the first inequality follows since $1 \pm X$ is always nonnegative, since X is supported on $[-1, 1]$. Notice the similarities (in fact, the duality) with the conditions for polynomial nonnegativity.

2.2 The moment curve

An appealing geometric interpretation of the set of valid moments is in terms of the so-called *moment curve*, which is the parametric curve in \mathbb{R}^{d+1} given by $t \mapsto (1, t, t^2, \dots, t^d)$. Indeed, it is easy to see that every point on the curve corresponds to a Dirac measure, where all the probability is concentrated on a given point. Thus, every finite (or infinite) measure on the interval corresponds to a point in the convex hull. In Figure 1 we present an illustration of the set of valid moments, for the case $d = 3$.

3 Bridging the gap

What to do in the cases where the set of nonnegative polynomials is no longer equal to the SOS ones? As we will see in much more detail later, it turns out that we can approximate *any* semialgebraic problem (including the simple case of a single polynomial being nonnegative) by sum of squares techniques.

As a preview, and a hint at some of the possibilities, let's consider how to prove nonnegativity of a particular polynomial which is not a sum of squares. Recall that the Motzkin polynomial was defined as:

$$M(x, y) = x^4 y^2 + x^2 y^4 + 1 - 3x^2 y^2.$$

and is a nonnegative polynomial that is not SOS. We can try multiplying it by another polynomial which is known to be positive, and check whether the resulting product is SOS. In this case, multiplying by the factor $(x^2 + y^2)$, we can find the decomposition

$$(x^2 + y^2) \cdot M(x, y) = y^2(1 - x^2)^2 + x^2(1 - y^2)^2 + x^2 y^2(x^2 + y^2 - 2)^2,$$




Figure 1: Set of valid moments (μ_1, μ_2, μ_3) of a probability measure on $[-1, 1]$. This is the convex hull of the moment curve (t, t^2, t^3) , for $-1 \leq t \leq 1$. An explicit SDP representation is given in (3).

which clearly certifies that $M(x, y) \geq 0$.

More details will follow...

References

- [Par00] P. A. Parrilo. *Structured semidefinite programs and semialgebraic geometry methods in robustness and optimization*. PhD thesis, California Institute of Technology, May 2000. Available at <http://resolver.caltech.edu/CaltechETD:etd-05062004-055516>.
- [PP02] A. Papachristodoulou and S. Prajna. On the construction of Lyapunov functions using the sum of squares decomposition. In *Proceedings of the 41th IEEE Conference on Decision and Control*, 2002.
- [PPP05] S. Prajna, A. Papachristodoulou, and P. A. Parrilo. *SOSTOOLS: Sum of squares optimization toolbox for MATLAB*, 2002-05. Available from <http://www.cds.caltech.edu/sostools> and <http://www.mit.edu/~parrilo/sostools>.

Lecture 12

Lecturer: Pablo A. Parrilo

Scribe: ???

1 Recovering a measure from its moments

We review next a classical method for producing a univariate atomic measure with a given set of moments. Essentially similar variations of this method are often used in signal processing, e.g., Pisarenko's harmonic decomposition method, where we are interested in producing a superposition of sinusoids with a given covariance matrix.

Consider the set of moments $(\mu_0, \mu_1, \dots, \mu_{2n-1})$ for which we want to find an associated nonnegative measure, supported on the real line. The resulting measure will be discrete, of the form $\sum_{i=1}^n w_i \delta(x - x_i)$. For this, consider the linear system

$$\begin{bmatrix} \mu_0 & \mu_1 & \cdots & \mu_{n-1} \\ \mu_1 & \mu_2 & \cdots & \mu_n \\ \vdots & \vdots & \ddots & \vdots \\ \mu_{n-1} & \mu_n & \cdots & \mu_{2n-2} \end{bmatrix} \begin{bmatrix} c_0 \\ c_1 \\ \vdots \\ c_{n-1} \end{bmatrix} = - \begin{bmatrix} \mu_n \\ \mu_{n+1} \\ \vdots \\ \mu_{2n-1} \end{bmatrix}. \quad (1)$$

The Hankel matrix on the left-hand side of this equation is the one that appeared earlier as a sufficient condition for the moments to represent a nonnegative measure. The linear system in (1) has a unique solution if the matrix is positive definite. In this case, we let x_i be the roots of the univariate polynomial

$$x^n + c_{n-1}x^{n-1} + \cdots + c_1x + c_0 = 0,$$

which are all real and distinct (why?). We can then obtain the corresponding weights w_i by solving the nonsingular Vandermonde system given by

$$\sum_{i=1}^n w_i x_i^j = \mu_j \quad (0 \leq j \leq n-1).$$

In the exercises, you will have to prove that this method actually works (i.e., the x_i are real and distinct, the w_i are nonnegative, and the moments are the correct ones).

Example 1. Let's find a nonnegative measure whose first six moments are given by $(1, 1, 2, 1, 6, 1)$. The solution of the linear system (1) yields the polynomial

$$x^3 - 4x^2 - 9x + 16 = 0,$$

whose roots are -2.4265 , 1.2816 , and 5.1449 . The corresponding weights are 0.0772 , 0.9216 , and 0.0012 , respectively.

2 A probabilistic interpretation

We also mention here an appealing probabilistic interpretation of the dual (2), commonly used in integer and quadratic programming, and developed by Lasserre in the polynomial case [Las01]. Consider as before the problem of minimizing a polynomial. Now, rather than looking for a minimizer x in \mathbb{R}^n , let's "relax" our notion of solution to allow for probabilities densities μ on \mathbb{R}^n , and replace the objective function by its natural generalization $\int p(x) d\mu$. It clearly holds that the new objective is never larger than the original one, since we are making the feasible set bigger.

This change makes the problem trivially convex, although infinite-dimensional. To produce a finite dimensional approximation (which may or may not be exact), we rewrite the objective function in terms of the moments of the measure μ , and write valid semidefinite constraints for the moments μ_k .

3 Duality and complementary slackness

What is the relationship between this classical method and semidefinite programming duality? Recall our approach to minimizing a polynomial $p(x)$ by computing

$$\max \gamma \quad \text{s.t.} \quad p(x) - \gamma \text{ is SOS.}$$

If this relaxation is exact (i.e., the optimal γ is equal to the optimal value of the polynomial) then at optimality, we necessarily have $p(x_\star) - \gamma_\star = \sum_i g_i^2(x_\star)$. This implies that all the g_i vanish at the optimal point. We can thus obtain the optimal value by looking at the roots of the polynomials $g_i(x)$.

However, it turns out that if we are simultaneously solving the primal and the dual SDPs (as most modern interior point solvers) this is unnecessary, since from complementary slackness we can extract almost all the information needed. In particular, notice that if we have

$$p(x) - \gamma = [x]_d^T Q [x]_d = 0$$

then necessarily $Q \cdot [x]_d = 0$.

Recall the SDP formulation is given by

$$\max \gamma \quad \text{s.t.} \quad \begin{cases} p_0 - \gamma = Q_{00} \\ p_i = \sum_{j+k=i} Q_{jk} \quad i = 1, \dots, 2d \\ Q \succeq 0 \end{cases}$$

and its dual

$$\min \sum_{i=0}^{2d} p_i \mu_i \quad \text{s.t.} \quad M(\mu) := \begin{bmatrix} \mu_0 & \mu_1 & \cdots & \mu_d \\ \mu_1 & \mu_2 & \cdots & \mu_{d+1} \\ \vdots & \vdots & \ddots & \vdots \\ \mu_d & \mu_{d+1} & \cdots & \mu_{2d} \end{bmatrix} \succeq 0, \quad \mu_0 = 1. \quad (2)$$

At optimality, complementarity slackness holds, i.e., the product of the primal and dual matrices vanishes. We have then $M(\mu) \cdot Q = 0$. Assume that the leading $k \times k$ submatrix of $M(\mu)$ is nonsingular. Then, the procedure described in Section 1 gives a k -atomic measure, with support in the minimizers of $p(x)$. Generically, this matrix $M(\mu)$ will be rank one, which will correspond to the case of a unique optimal solution.

Remark 2. *Unlike the univariate case, a multivariate polynomial that is bounded below may not achieve its minimum. A well-known example is $p(x, y) = x^2 + (1 - xy)^2$, which clearly satisfies $p(x, y) \geq 0$. Since $p(x, y) = 0$ would imply $x = 0$ and $1 - xy = 0$ (which is impossible), this value cannot be achieved. However, we can get arbitrarily close, since $p(\epsilon, 1/\epsilon) = \epsilon^2$, for any $\epsilon > 0$.*

4 Multivariate case

We have seen previously that in the multivariate case, it is no longer the case that nonnegative polynomials are always sums of squares. The corresponding result on the dual side is that the set of valid moments is no longer described by the “obvious” semidefinite constraints, obtained by considering the expected value of squares (even if we require strict positivity).

Example 3 (“Dual Motzkin”). *Consider the existence of a probability measure on \mathbb{R}^2 , that satisfies the moment constraints:*

$$\begin{aligned} E[1] &= E[X^4 Y^2] = E[X^2 Y^4] = 1, \\ E[X^2 Y^2] &= 2, \\ E[XY] &= E[XY^2] = E[X^2 Y] = E[X^2 Y^3] = E[X^3 Y^2] = E[X^3 Y^3] = 0. \end{aligned} \quad (3)$$

The “obvious” nonnegativity constraints are satisfied, since

$$E[(a + bXY + cXY^2 + dX^2Y)^2] = a^2 + 2b^2 + c^2 + d^2 \geq 0.$$

However, it turns out that these conditions are only necessary, but not sufficient. This can be seen by computing the expectation of the Motzkin polynomial (which is nonnegative), since in this case we have

$$E[X^4Y^2 + X^2Y^4 + 1 - 3X^2Y^2] = 1 + 1 + 1 - 6 = -3,$$

thus proving that no measure with the given moments can exist.

5 Density results

Recent results by Blekherman [Ble04] give quantitative bounds on the relative density of the cone of sum of squares versus the cone of nonnegative polynomials. Concretely, in [Ble04] it is proved that a suitably normalized section of the cone of positive polynomials $\tilde{P}_{n,2d}$ satisfies

$$c_1 n^{-\frac{1}{2}} \leq \left(\frac{\text{Vol } \tilde{P}_{n,2d}}{\text{Vol } B_M} \right)^{\frac{1}{D_M}} \leq c_2 n^{-\frac{1}{2}},$$

while the corresponding expression for the section of the cone of sum of squares $\tilde{\Sigma}_{n,2d}$ is

$$c_3 n^{-\frac{d}{2}} \leq \left(\frac{\text{Vol } \tilde{\Sigma}_{n,2d}}{\text{Vol } B_M} \right)^{\frac{1}{D_M}} \leq c_4 n^{-\frac{d}{2}},$$

where c_1, c_2, c_3, c_4 depend on d only (explicit expressions are available), $D_M = \binom{n+2d}{2d} - 1$, and B_M is the unit ball in \mathbb{R}^{D_M} .

These expressions show that for fixed d , as $n \rightarrow \infty$ the volume of the set of sum of squares becomes vanishingly small when compared to the nonnegative polynomials.

Show the values of the actual bounds, for reasonable dimensions

ToDo

References

- [Ble04] G. Blekherman. There are significantly more nonnegative polynomials than sums of squares. Preprint, 2004.
- [Las01] J. B. Lasserre. Global optimization with polynomials and the problem of moments. *SIAM J. Optim.*, 11(3):796–817, 2001.

Lecture 13

Lecturer: Pablo A. Parrilo

Scribe: ???

Today we introduce the first basic elements of algebraic geometry, namely ideals and varieties over the complex numbers. This dual viewpoint (ideals for the algebra, varieties for the geometry) is enormously powerful, and will help us later in the development of methods for solving polynomial equations. We also present the notion of quotient rings, which are very natural when considering functions defined on algebraic varieties (e.g., in polynomial optimization problems with equality constraints). Finally, we begin our study of Groebner bases, by defining the notion of term orders. A superb introduction to algebraic geometry, emphasizing the computational aspects, is the textbook of Cox, Little, and O'Shea [CLO97].

1 Polynomial ideals

For notational simplicity, we use $\mathbb{C}[\mathbf{x}]$ to denote the polynomial ring in n variables $\mathbb{C}[x_1, \dots, x_n]$. Specializing the general definition of an ideal to a polynomial ring, we have the following:

Definition 1. A subset $I \subset \mathbb{C}[\mathbf{x}]$ is an ideal if it satisfies:

1. $0 \in I$.
2. If $a, b \in I$, then $a + b \in I$.
3. If $a \in I$ and $b \in \mathbb{C}[\mathbf{x}]$, then $a \cdot b \in I$.

The two most important examples of polynomial ideals for our purposes are the following:

- The set of polynomials that vanish in a given set $S \subset \mathbb{C}^n$, i.e.,

$$\mathbf{I}(S) := \{f \in \mathbb{C}[\mathbf{x}] : f(a_1, \dots, a_n) = 0 \quad \forall (a_1, \dots, a_n) \in S\},$$

is an ideal, called the *vanishing ideal* of S .

- The ideal generated by a finite set of polynomials $\{f_1, \dots, f_s\}$, defined as

$$\langle f_1, \dots, f_s \rangle := \{f \mid f = g_1 f_1 + \dots + g_s f_s, \quad g_i \in \mathbb{C}[\mathbf{x}]\}. \quad (1)$$

An ideal is *finitely generated* if it can be written as in (1) for some finite set of polynomials $\{f_1, \dots, f_s\}$. An ideal is called *principal* if it can be generated by a single polynomial. The intersection of two ideals is again an ideal. What about the union of ideals?

Example 2. In the univariate case (i.e., the polynomial ring is $\mathbb{C}[x]$), every ideal is principal.

One of the most important facts about polynomial ideals is Hilbert's finiteness theorem:

Theorem 3 (Hilbert Basis Theorem). Every polynomial ideal in $\mathbb{C}[\mathbf{x}]$ is finitely generated.

We will present a proof of this after learning about Groebner bases. From the computational viewpoint, two very natural questions about ideals are the following:

- Given a polynomial $p(x)$, how to decide if it belongs to a given ideal?
- How to find a “convenient” representation of an ideal? What does “convenient” mean?




Figure 1: Two algebraic varieties. The one on the left is defined by the equation $(x^2 + y^2 - 1)(3x + 6y - 4) = 0$. The one on the right is a quartic surface, defined by $1 - x^2 - y^2 - 2z^2 + z^4 = 0$.

2 Algebraic varieties

An (affine) algebraic variety is the zero set of a finite collection of polynomials (see formal definition below). The word “affine” here means that we are working in the standard affine space, as opposed to projective space, where we identify $x, y \in \mathbb{C}^n$ if $x = \lambda y$ for some $\lambda \neq 0$.

Definition 4. Let f_1, \dots, f_s be polynomials in $\mathbb{C}[\mathbf{x}]$. Let the set \mathbf{V} be

$$\mathbf{V}(f_1, \dots, f_s) := \{(a_1, \dots, a_n) \in \mathbb{C}^n : f_i(a_1, \dots, a_n) = 0 \quad 1 \leq i \leq s\}.$$

We call $\mathbf{V}(f_1, \dots, f_s)$ the affine variety defined by f_1, \dots, f_s .

A simple example of a variety is a (complex) affine subspace, that corresponds to the vanishing of a finite collection of affine polynomials. A few additional examples of varieties are shown in Figure 1.

It is not too hard to show that *finite* unions and intersections of algebraic varieties are again algebraic varieties. What about the infinite case?

Remark 5. Recall our previous encounter with the Zariski topology, whose closed sets were defined to be the algebraic varieties, i.e., the vanishing set of a finite set of polynomial equations. To prove that this is actually a topology, we need to show that arbitrary intersections of closed sets are closed. Hilbert’s basis theorem precisely guarantees this fact.

Perhaps the most natural question about algebraic varieties is the following:

- Given a variety V , how to decide it is nonempty?

Let’s start connecting ideals and varieties. Consider a finite set of polynomials $\{f_1, \dots, f_s\}$. We already know how to generate an ideal, namely $\langle f_1, \dots, f_s \rangle$. However, we can also look at the corresponding variety $\mathbf{V}(f_1, \dots, f_s)$. Since this variety is a subset of \mathbb{C}^n , we can form the corresponding vanishing ideal, $\mathbf{I}(\mathbf{V}(f_1, \dots, f_s))$. How do these two ideals related to each other? Is it always the case that

$$\langle f_1, \dots, f_s \rangle = \mathbf{I}(\mathbf{V}(f_1, \dots, f_s)),$$

and if it is not, what are the reasons? The answer to these questions (and more) will be given by another famous result by Hilbert, known as the Nullstellensatz.

3 Quotient rings

Whenever we have an ideal in a ring, we can immediately define a notion of equivalence classes, where we identify two elements in the ring if and only if their difference is in the ideal.

Example 6. Recall that a simple example of an ideal in the ring \mathbb{Z} was the set of even integers. By identifying two integers if their difference is even, we partition \mathbb{Z} into two equivalence classes, namely the even and the odd numbers. More generally, if the ideal is given by the integer multiples of a given number m , then \mathbb{Z} can be partitioned into m equivalence classes.

We can do this for the polynomial ring $\mathbb{C}[\mathbf{x}]$, and any ideal I .

Definition 7. Let $I \subset \mathbb{C}[\mathbf{x}]$ be an ideal, and let $f, g \in \mathbb{C}[\mathbf{x}]$. We say f and g are congruent modulo I , written

$$f \equiv g \pmod{I},$$

if $f - g \in I$.

It is easy to show that this is an equivalence relation, i.e., it is reflexive, symmetric, and transitive. Thus, this partitions $\mathbb{C}[\mathbf{x}]$ into equivalence classes, where two polynomials are “the same” if their difference belongs to the ideal. This allows us to define the quotient ring:

Definition 8. The quotient $\mathbb{C}[\mathbf{x}]/I$ is the set of equivalence classes for congruence modulo I .

The quotient $\mathbb{C}[\mathbf{x}]/I$ inherits the ring structure of $\mathbb{C}[\mathbf{x}]$, with the natural operations. Thus, with these operations now defined between equivalence classes, $\mathbb{C}[\mathbf{x}]/I$ becomes a ring, known as the *quotient ring*.

Quotient rings are particularly useful when considering a polynomial function $p(x)$ over the algebraic variety defined by $g_i(x) = 0$. Notice that if we define the ideal $I = \langle g_i \rangle$, then *any* polynomial q that is congruent with p modulo I takes exactly the same values in the variety.

4 Monomial orderings

In order to begin studying “nice” bases for ideals, we need a way of ordering monomials. In the univariate case, this is straightforward, since we can define $x^a \succ x^b$ as being true if and only if $a > b$. In the multivariate case, there are a lot more options.

We also want the ordering structure to be consistent with polynomial multiplication. This is formalized in the following definition.

Definition 9. A monomial ordering on $\mathbb{C}[\mathbf{x}]$ is a relation \succ on \mathbb{Z}_+^n (i.e., the monomial exponents), such that:

1. The relation \succ is a total ordering.
2. If $\alpha \succ \beta$, and $\gamma \in \mathbb{Z}_+^n$, then $\alpha + \gamma \succ \beta + \gamma$.
3. The relation \succ is a well-ordering (every nonempty subset has a smallest element).

One of the simplest examples of a monomial ordering is the *lexicographic* ordering, where $\alpha \succ_{\text{lex}} \beta$ if the left-most nonzero entry of $\alpha - \beta$ is positive. We will see some other examples of monomial orderings later in the course.

References

- [CLO97] D. A. Cox, J. B. Little, and D. O’Shea. *Ideals, varieties, and algorithms: an introduction to computational algebraic geometry and commutative algebra*. Springer, 1997.

Lecture 14

Lecturer: Pablo A. Parrilo

Scribe: ???

After a brief review of monomial orderings, we develop the basic ideas of Groebner bases, followed by examples and applications. For background and much more additional material, we recommend the textbook of Cox, Little, and O'Shea [CLO97]. Other good, more specialized references are [AL94, BW93, KR00].

1 Monomial orderings

Recall from last lecture the notion of a monomial ordering:

Definition 1. A monomial ordering on $\mathbb{C}[\mathbf{x}]$ is a relation \succ on \mathbb{Z}_+^n (i.e., the monomial exponents), such that:

1. The relation \succ is a total ordering.
2. If $\alpha \succ \beta$, and $\gamma \in \mathbb{Z}_+^n$, then $\alpha + \gamma \succ \beta + \gamma$.
3. The relation \succ is a well-ordering (every nonempty subset has a smallest element).

There are several term orderings of interest in computational algebra. Among them, we mention:

- Lexicographic (“dictionary”). Here $\alpha \succ_{\text{lex}} \beta$ if the left-most nonzero entry of $\alpha - \beta$ is positive. Notice that a particular order of the variables is assumed, and by changing this, we obtain $n!$ nonequivalent lexicographic orderings.
- Graded lexicographic. Sort first by total degree, then lexicographic, i.e., $\alpha \succ_{\text{grlex}} \beta$ if $|\alpha| > |\beta|$, or if $|\alpha| = |\beta|$ and $\alpha \succ_{\text{lex}} \beta$.
- Graded reverse lexicographic. Here $\alpha \succ_{\text{grevlex}} \beta$ if $|\alpha| > |\beta|$, or if $|\alpha| = |\beta|$ and the right-most nonzero entry of $\alpha - \beta$ is negative. This ordering, although somewhat nonintuitive, has some desirable computational properties.
- General matrix orderings. Described by a weight matrix $W \in \mathbb{R}^{k \times n}$ ($k \leq n$), where $\alpha \succ_W \beta$ if $(W\alpha) \succ_{\text{lex}} (W\beta)$. For W to correspond to a monomial ordering as defined, the first nonzero entry on each column must be positive.

It turns out that every monomial ordering can be described by an associated matrix W , i.e., every monomial ordering is a matrix ordering. What are the matrices corresponding to the first three orderings described?

Example 2. Consider the polynomial ring $\mathbb{C}[x, y]$. In the lexicographic ordering (\prec_{lex}) discussed, we have:

$$1 \prec y \prec y^2 \prec \cdots \prec x \prec xy \prec xy^2 \prec \cdots \prec x^2 \prec x^2y \prec x^2y^2 \prec \cdots,$$

while for the other two orderings (\prec_{grlex} and \prec_{grevlex}), which in the special case of two variables coincide, we have:

$$1 \prec y \prec x \prec y^2 \prec xy \prec x^2 \prec y^3 \prec xy^2 \prec x^2y \prec x^3 \prec \cdots.$$

Picture comparing different orderings

ToDo

Example 3. Consider the monomials $\alpha = x^3y^2z^8$ and $\beta = x^2y^9z^2$. If the variables are ordered as (x, y, z) , we have

$$\alpha \succ_{\text{lex}} \beta, \quad \alpha \succ_{\text{grlex}} \beta, \quad \alpha \prec_{\text{grevlex}} \beta.$$

Notice that $x \succ y \succ z$ for all three orderings.

2 Groebner bases

2.1 Monomial ideals

Before studying general ideals, it is convenient to introduce first a special class, known as *monomial ideals*.

Definition 4. A monomial ideal is a polynomial ideal that can be generated by monomials.

What are the possible monomials that belong to a given monomial ideal? Since $x^\alpha \in I \Rightarrow x^{\alpha+\beta} \in I$ for $\beta \geq 0$, we have that these sets are “closed upwards.”

Picture of monomial ideals

ToDo

Furthermore, a polynomial belongs to a monomial ideal I if and only if all its terms are in I .

Theorem 5 (Dickson’s lemma). *Every monomial ideal is finitely generated.*

We consider next a special monomial ideal, associated to every polynomial ideal I . From now on, we assume a fixed monomial ordering (e.g., graded reverse lexicographic), and denote by $\text{in}(f)$ the “largest” monomial appearing in the polynomial $f \neq 0$.

Definition 6. Consider an ideal $I \subset \mathbb{C}[x]$, and a fixed monomial ordering. The initial ideal of I , denoted $\text{in}(I)$, is the monomial ideal generated by the leading terms of all the elements in I , i.e.,

$$\text{in}(I) := \langle \text{in}(f) : f \in I \setminus \{0\} \rangle.$$

A monomial x^α is called *standard*, if it does not belong to the initial ideal $\text{in}(I)$.

2.2 Groebner bases

Given an ideal $I = \langle f_1, \dots, f_s \rangle$, we can construct two monomial ideals associated with it. On the one hand, we have the initial ideal $\text{in}(I)$, previously defined. However, we can also consider the monomial ideal generated by the initial monomials of the generators, i.e., $\langle \text{in}(f_1), \dots, \text{in}(f_s) \rangle$. Although we always have $\langle \text{in}(f_1), \dots, \text{in}(f_s) \rangle \subset \text{in}(I)$, in general these two monomial ideals are distinct.

Example 7. Consider the ideal $I = \langle x^3 - 1, x^2 + 1 \rangle$. Since $1 = \frac{1}{2}(x-1)(x^3-1) - \frac{1}{2}(x^2-x-1)(x^2+1)$, we have $1 \in I$, and thus $\text{in}(I) = I = \mathbb{C}[x]$. On the other hand, $1 \notin \langle x^3, x^2 \rangle$.

However, it may be possible to produce a set of generators for which these two ideals are the same. This is exactly the notion of a *Groebner basis*.

Definition 8. Consider the polynomial ring $\mathbb{C}[x]$, with a fixed monomial ordering, and an ideal I . A finite set of polynomials $\{g_1, \dots, g_s\} \subset I$ is a Groebner basis of I if the initial ideal of I is generated by the leading terms of the g_i , i.e.,

$$\text{in}(I) = \langle \text{in}(g_1), \dots, \text{in}(g_s) \rangle. \quad (1)$$

Theorem 9. Every ideal I has a Groebner basis G . Furthermore, $I = \langle g_1, \dots, g_s \rangle$.

The previous theorem essentially establishes Hilbert’s finiteness result, and gives an explicit characterization of a finite generating set for the ideal I . Furthermore, since there are explicit algorithms to compute Groebner bases, this is a constructive version of this theorem.

Even though the monomial ordering is fixed, Groebner bases as defined are not unique (why?). This can be easily fixed, by refining the concept to the so-called *reduced* Groebner bases, which are uniquely defined.

There are several possible algorithms to effectively compute Groebner bases. The traditional one is *Buchberger's algorithm*, developed by Bruno Buchberger around 1965, and many variants have been proposed since. There are also several newer methods, based on sparse linear algebra, that in some instances can significantly outperform the Buchberger approach. Good specialized programs for Groebner bases calculations (and much more) are CoCoA[CoC], Macaulay2 [GS] and Singular [GPS05].

2.3 Quotients and normal forms

Recall that if we have an ideal $I \subset \mathbb{C}[x]$, we defined the quotient ring $\mathbb{C}[x]/I$ as the set of equivalence classes modulo the ideal. For computational purposes, we want a “good” representation of these classes, and in particular, a way to provide a “unique representative” to every polynomial. This can in fact be easily done once we have computed a Groebner basis. To each polynomial $p \in \mathbb{C}[x]$, we can associate a unique “normal form”, defined below.

Lemma 10. *Let G be a Groebner basis of the ideal $I \subset \mathbb{C}[x]$. Given any $p \in \mathbb{C}[x]$, there exists a unique polynomial \bar{p} , called the normal form of p , such that*

1. *The polynomials p and \bar{p} are congruent mod I , i.e., $p - \bar{p} \in I$.*
2. *Only standard monomials appear in \bar{p} .*

Notice that we have $p = q_1g_1 + \dots + q_sg_s + \bar{p}$. Thus, the normal form can be interpreted as the “remainder” after a division-like process by the generators g_i . The key property (1) guarantees that this remainder is uniquely defined.

As a consequence of this, we can solve the ideal membership problem: to check if a polynomial $p(x)$ is in a given ideal I , compute a Groebner basis G of I , and check if the normal form of $p(x)$ is the zero polynomial, i.e., $p \in I \Leftrightarrow \bar{p} = 0$.

3 Applications and examples

Groebner bases enable the algorithmic solution of many problems in computational algebraic geometry. We discuss some of these below.

- **Ideal membership.** As we have seen, given an ideal I and a polynomial p , we can check if $p \in I$ by computing the normal form of p .
- **Radical membership.** Consider an ideal $I = \langle f_1, \dots, f_s \rangle \subset \mathbb{C}[x]$, and a polynomial p , for which we want to check whether $p \in \sqrt{I}$. Since \sqrt{I} is also an ideal, we could compute a Groebner basis for it, and then reduce the problem to the previous one. However, it is often more efficient to instead use the following result (“Rabinowitch’s trick”):

$$p \in \sqrt{I} \quad \Leftrightarrow \quad 1 \in \langle f_1, \dots, f_s, 1 - yp \rangle,$$

where y is a (new) additional variable.

- **Consistency of polynomial equations.** Consider a finite set of polynomial equations $\{f_i = 0\}$, and let $I = \langle f_i \rangle$ be the corresponding ideal. By the Nullstellensatz, the given equations are infeasible if and only if $\{1\}$ is the reduced Groebner basis of I .
- **Elimination.** For notational simplicity, consider an ideal $I \subset \mathbb{C}[x, y, z]$. Suppose that we want to compute all the polynomials in I , that *do not* depend on the variable z , i.e., $I \cap \mathbb{C}[x, y]$. Geometrically, this elimination of variables corresponds to (the Zariski closure of) the projection of the corresponding variety into (x, y) . This intersection (or projection) can be easily obtained, by computing a Groebner basis G of I with respect to a lexicographic (or elimination) ordering. The corresponding ideal is then generated by $G \cap \mathbb{C}[x, y]$.

4 Zero-dimensional ideals

In practice, we are often interested in polynomial systems that have only a finite number of solutions (the “zero-dimensional” case), and many interesting things happen in this case. Among other properties, the quotient ring $\mathbb{C}[x]/I$ is now a finite dimensional vector space, with its dimension being equal to the number of standard monomials. Furthermore, Groebner bases can be used to fully reduce their solution to a classical eigenvalue problem, generalizing the “companion matrix” notion from the univariate case. All this, and much more, next time...

References

- [AL94] W.W. Adams and P. Loustaunau. *An introduction to Gröbner bases*, volume 3 of *Graduate Studies in Mathematics*. American Mathematical Society, Providence, RI, 1994.
- [BW93] T. Becker and V. Weispfenning. *Gröbner bases*, volume 141 of *Graduate Texts in Mathematics*. Springer-Verlag, New York, 1993.
- [CLO97] D. A. Cox, J. B. Little, and D. O’Shea. *Ideals, varieties, and algorithms: an introduction to computational algebraic geometry and commutative algebra*. Springer, 1997.
- [CoC] CoCoATeam. CoCoA: a system for doing Computations in Commutative Algebra. Available at <http://cocoa.dima.unige.it>.
- [GPS05] G.-M. Greuel, G. Pfister, and H. Schönemann. SINGULAR 3.0. A Computer Algebra System for Polynomial Computations, Centre for Computer Algebra, University of Kaiserslautern, 2005. <http://www.singular.uni-kl.de>.
- [GS] D.R. Grayson and M. E. Stillman. Macaulay 2, a software system for research in algebraic geometry. Available at <http://www.math.uiuc.edu/Macaulay2/>.
- [KR00] M. Kreuzer and L. Robbiano. *Computational commutative algebra. 1*. Springer-Verlag, Berlin, 2000.

Lecture 15

Lecturer: Pablo A. Parrilo

Scribe: ???

Today we will see a few more examples and applications of Groebner bases, and we will develop the zero-dimensional case.

1 Zero-dimensional ideals

In practice, we are often interested in polynomial systems that have only a finite number of solutions (the “zero-dimensional” case), and as we will see, many interesting things happen in this case.

Definition 1. An ideal I is zero-dimensional if the associated variety $V(I)$ is a finite set.

Given a system of polynomial equations, how to decide if it has a finite number of solutions (i.e., if the corresponding ideal is zero-dimensional)? We can state a simple criterion for this in terms of a Groebner basis.

Lemma 2. Let G be a Groebner basis of the ideal $I \subset \mathbb{C}[x_1, \dots, x_n]$. The ideal I is zero-dimensional if and only if for each i ($1 \leq i \leq n$), there exists an element in the Groebner basis whose initial term is a pure power of x_i .

Among other important consequences, when I is a zero-dimensional ideal the quotient ring $\mathbb{C}[x]/I$ is a finite dimensional vector space, with its dimension being equal to the number of standard monomials.

Furthermore, we can use Groebner bases to reduce the effective calculation of the solutions of a zero-dimensional polynomial system to an eigenvalue problem, generalizing the “companion matrix” notion from the univariate case. We sketch this below.

Recall that in this case, the quotient $\mathbb{C}[x]/I$ is a finite dimensional vector space. The main idea is to consider the homomorphisms given by the n linear maps $M_{x_i} : \mathbb{C}[x]/I \rightarrow \mathbb{C}[x]/I$, $f \mapsto \widehat{(x_i f)}$ (that is, multiplication by the coordinate variables, followed by normal form). Choosing as a basis the set of standard monomials, we can effectively compute a matrix representation of these linear maps. This defines n matrices M_{x_i} , that commute with each other (why?).

Assume for simplicity that all the roots have single multiplicity. Then, all the M_{x_i} can be simultaneously diagonalized by a single matrix V , and the k th diagonal entry of $VM_{x_i}V^{-1}$ contains the i th coordinate of the k th solution, for $1 \leq k \leq \#V(I)$.

(In general, we can block-diagonalize this commutative algebra, splitting into its semisimple and nilpotent components. The nilpotent part is trivial if and only if the ideal is radical.)

To understand these ideas a bit better, let’s recall the univariate case.

Example 3. Consider the ring $\mathbb{C}[x]$ of polynomials in a single variable x , and an ideal $I \subset \mathbb{C}[x]$. Since every ideal in this ring is principal, I can be generated by a single polynomial $p(x) = p_n x^n + \dots + p_1 x + p_0$. Then, we can write $I = \langle p(x) \rangle$, and $\{p(x)\}$ is a Groebner basis for the ideal (why?). The quotient $\mathbb{C}[x]/I$ is an n -dimensional vector, with a suitable basis given by the standard monomials $\{1, x, \dots, x^{n-1}\}$.

Consider as before the linear map $M_x : \mathbb{C}[x]/I \rightarrow \mathbb{C}[x]/I$. The matrix representation of this linear map in the given basis is given by

$$\begin{bmatrix} 0 & 0 & 0 & \cdots & -p_0/p_n \\ 1 & 0 & 0 & \cdots & -p_1/p_n \\ 0 & 1 & 0 & \cdots & -p_2/p_n \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \cdots & -p_{n-1}/p_n \end{bmatrix},$$

which is the standard companion matrix C_p associated with $p(x)$. Its eigenvalues are exactly the roots of $p(x)$.

We present next a multivariate example.

Example 4. Consider the ideal $I \subset \mathbb{C}[x, y, z]$ given by

$$I = \langle xy - z, yz - x, zx - y \rangle.$$

Choosing a term ordering (e.g., lexicographic, where $x \prec y \prec z$), we obtain the Groebner basis

$$G = \{x^3 - x, yx^2 - y, y^2 - x^2, z - yx\}.$$

We can directly see from this that I is zero-dimensional (why?). A basis for the quotient space is given by $\{1, x, x^2, y, yx\}$. Consider the maps M_x , M_y , and M_z , we have that their corresponding matrix representations are given by

$$M_x = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 \end{bmatrix}, \quad M_y = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 \end{bmatrix}, \quad M_z = \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 \\ 1 & 0 & 1 & 0 & 0 \end{bmatrix}.$$

It can be verified that these three matrices commute. A simultaneous diagonalizing transformation is given by the matrix:

$$V = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 1 & -1 & -1 \\ 1 & -1 & 1 & 1 & -1 \\ 1 & -1 & 1 & -1 & 1 \end{bmatrix}, \quad V^{-1} = \frac{1}{4} \begin{bmatrix} 4 & 0 & 0 & 0 & 0 \\ 0 & 1 & 1 & -1 & -1 \\ -4 & 1 & 1 & 1 & 1 \\ 0 & 1 & -1 & 1 & -1 \\ 0 & 1 & -1 & -1 & 1 \end{bmatrix}.$$

The corresponding transformed matrices are:

$$VM_xV^{-1} = \text{diag}(0, 1, 1, -1, -1)$$

$$VM_yV^{-1} = \text{diag}(0, 1, -1, 1, -1),$$

$$VM_zV^{-1} = \text{diag}(0, 1, -1, -1, 1)$$

from where the coordinates of the five roots can be read.

In the general (radical) case, the matrix V is a generalized Vandermonde matrix, with rows indexed by roots (points in the variety) and columns indexed by the standard monomials. The V_{ij} entry contains the j -th monomial evaluated at the i th root. Since $VV^{-1} = I$, we can also interpret the j th column of V^{-1} as giving the coefficients of a Lagrange interpolating polynomial $p_j(x)$, that vanishes at all the points in the variety, except at r_j , where it takes the value 1 (i.e., $p_j(r_k) = \delta_{jk}$).

Generalize Hermite form, etc

ToDo

2 Hilbert series

Consider an ideal $I \subset \mathbb{C}[x]$ and the corresponding quotient ring $\mathbb{C}[x]/I$. We have seen that, once a particular Groebner basis is chosen, we could associate to every element of $\mathbb{C}[x]/I$ a unique representative, namely a \mathbb{C} -linear combination of *standard monomials*, obtained as the remainder after division with the corresponding Groebner basis. We are interested in studying, for every integer k , the dimension of




Figure 1: A six-node graph.

the vector space of remainders of degree less than or equal to k . Expressed in a simpler way, we want to know how many standard monomials of degree k there are, for any given k .

Rather than studying this for different values of k separately, it is convenient to collect (or bundle) all these numbers together in a single object (this general technique is usually called “generating function”). The *Hilbert series* of I , denoted $H_I(t)$, is then defined as the generating function of the dimension of the space of residues of degree k , i.e.,

$$H_I(t) = \sum_{k=0}^{\infty} \dim(\mathbb{C}[x]/I \cap P_{n,k}) \cdot t^k, \quad (1)$$

where $P_{n,k}$ denotes the set of homogeneous polynomials in n variables of degree k .

Notice that, if the ideal is zero-dimensional, the corresponding Hilbert series is actually a finite sum, and thus a polynomial. The number of solutions is then equal to $H_I(1)$.

Example 5. For the ideal I in Example 4, the corresponding Hilbert function is $H_I(t) = 1 + 2t + 2t^2$.

In general, the Hilbert series *does* depend on the specific Groebner basis chosen, not only on the ideal I . However, almost all of the relevant algebraic and geometric properties (e.g., its degree, if it is a polynomial) are actually invariants associated only with the ideal.

3 Examples

3.1 Graph ideals

Consider a graph $G = (V, E)$, and define the associated *edge ideal* $I_G = \langle x_i x_j : (i, j) \in E \rangle$. Notice that I_G is a monomial ideal. For instance, for the graph in Figure 1, the corresponding ideal is given by:

$$I_G := \langle x_1 x_2, x_2 x_3, x_3 x_4, x_4 x_5, x_5 x_6, x_1 x_6, x_1 x_5, x_3 x_5 \rangle.$$

One of the motivations for studying this kind of ideals is that many graph-theoretic properties (e.g., bipartiteness, acyclicity, connectedness, etc) can be understood in terms of purely algebraic properties of the corresponding ideal. This enables the extension and generalization of these notions to much more abstract settings (e.g., simplicial complexes, resolutions, etc).

For our purposes here, rather than studying I_G directly, we will instead study the ideal obtained when restricting to zero-one solutions¹. For this, consider the ideal I_b defined as

$$I_b := \langle x_1^2 - x_1, \dots, x_n^2 - x_n \rangle. \quad (2)$$

¹There are more efficient ways of doing this, that would not require adding generators. We adopt this approach to keep the discussion relatively straightforward.

Clearly, this is a zero-dimensional radical ideal, with the corresponding variety having 2^n distinct points, namely $\{0, 1\}^n$. Its corresponding Hilbert series is $H_{I_b}(t) = (1 + t)^n = \sum_{k=0}^n \binom{n}{k} t^k$.

Since we want to study the intersection of the corresponding varieties, we must consider the *sum* of the ideals, i.e., the ideal $I := I_G + I_b$. It can be shown that the given set of generators (i.e., the ones corresponding to the edges, and the quadratic relations in (2)) are always a Groebner basis of the corresponding ideal. What are the standard monomials? How can they be interpreted in terms of the graph?

The *Hilbert function* of the ideal I can be obtained from the Groebner basis. In this case, the corresponding Hilbert function is given by

$$H_I(t) = 1 + 6t + 7t^2 + t^3,$$

and we can read from the coefficient of t^k the number of stable sets of size k . In particular, the degree of the Hilbert function (which is actually a polynomial, since the ideal is zero-dimensional) indicates the size of the maximum stable set, which is equal to three in this example (for the subset $\{x_2, x_4, x_6\}$).

3.2 Integer programming

Another interesting application of Groebner bases deals with integer programming. For more details, see the papers [CT91, ST97, TW97].

Consider the integer programming problem

$$\min c^T \mathbf{x} \quad \text{s.t.} \quad \begin{cases} A\mathbf{x} = b \\ \mathbf{x} \geq 0 \\ \mathbf{x} \in \mathbb{Z}^n \end{cases} \quad (3)$$

where $A \in \mathbb{Z}^{m \times n}$, $b \in \mathbb{Z}^m$, and $c \in \mathbb{Z}^n$. For simplicity, we assume that $A, c \geq 0$, and that we know a feasible solution \mathbf{x}_0 . These assumptions can be removed.

The main idea to solve (3) will be to interpret the nonnegative integer decision variables \mathbf{x} as the *exponents* of a monomial.

Complete

ToDo

Example 6. Consider the problem data given by

$$A = \begin{bmatrix} 4 & 5 & 6 & 1 \\ 1 & 2 & 7 & 3 \end{bmatrix}, \quad b = \begin{bmatrix} 750 \\ 980 \end{bmatrix}, \quad c^T = [1 \quad 2 \quad 3 \quad 4].$$

An initial feasible solution is given by $\mathbf{x}_0 = [0, 30, 80, 120]^T$. We will work on the ring $\mathbb{C}[z_1, z_2, w_1, w_2, w_3, w_4]$. Thus, we need to compute a Groebner basis G of the binomial ideal

$$\langle z_1^4 z_2 - w_1, z_1^5 z_2^2 - w_2, z_1^6 z_2^7 - w_3, z_1 z_2^3 - w_4 \rangle,$$

for a term ordering that combines elimination of the z_i with the weight vector c . To obtain the solution, we compute the normal form of the monomial given by the initial feasible point, i.e., $w_2^{30} w_3^{80} w_4^{120}$. This reduction process yields the result $w_2 w_3^{106} w_4^{74}$, and thus the optimal solution is $[0, 8, 106, 74]$. The corresponding costs of the initial and optimal solutions are $c^T \mathbf{x}_0 = 780$ and $c^T \mathbf{x}_{\text{opt}} = 630$.

We should remark that there are more efficient ways of implementing this than the one described. Also, although this basic method cannot currently compete with specialized techniques used in integer programming, there are some particular cases where it is very efficient, mostly related with the solution of parametric problems.

References

- [CT91] P. Conti and C. Traverso. Buchberger algorithm and integer programming. In *Applied algebra, algebraic algorithms and error-correcting codes (New Orleans, LA, 1991)*, volume 539 of *Lecture Notes in Comput. Sci.*, pages 130–139. Springer, Berlin, 1991.
- [ST97] B. Sturmfels and R.R. Thomas. Variation of cost functions in integer programming. *Math. Programming*, 77(3, Ser. A):357–387, 1997.
- [TW97] R. Thomas and R. Weismantel. Truncated Gröbner bases for integer programming. *Appl. Algebra Engrg. Comm. Comput.*, 8(4):241–256, 1997.

Lecture 16

Lecturer: Pablo A. Parrilo

Scribe: ???

1 Generalizing the Hermite matrix

Recall the basic construction of the Hermite matrix $H_q(p)$ in the univariate case, whose signature gave important information on the signs of the polynomial $q(x)$ on the real roots of $p(x)$.

In a very similar way to the extension of the companion matrix to the multivariate case, we can parallel the Hermite form to general zero-dimensional ideals. The basic idea is again to consider the zero-dimensional ideal $I \subset \mathbb{R}[x_1, \dots, x_n]$, and an associated basis of the quotient ring $B = \{x^{\alpha_1}, \dots, x^{\alpha_m}\}$, where the elements of B are standard monomials.

For simplicity, we assume first that I is radical. In this case, the corresponding finite variety is given by m distinct points, i.e., $V(I) = \{r_1, \dots, r_m\} \subset \mathbb{C}^n$. Notice first that, by the definition of the matrices M_{x_i} , we have $\sum_{i=1}^m r_i^\alpha = \text{Tr}[M_{x_1}^{\alpha_1} \cdots M_{x_n}^{\alpha_n}]$. Thus, in a similar way as we did in the univariate case, for any polynomial $q = \sum_{\beta} q_{\beta} x^{\beta}$ we have

$$\sum_{i=1}^m q(r_i) = \text{Tr}[q(M_{x_1}, \dots, M_{x_n})]. \quad (1)$$

Once again, this implies that if we have access to matrix representations M_{x_1}, \dots, M_{x_n} , then we can explicitly evaluate these expressions. Notice also that, if both q and the generators of the ideal have rational coefficients, then the expression above is also a rational number (even if the roots are not).

Example 1. Consider the system in Example 4 of the previous lecture, and the polynomial $p(x, y, z) = (x + y + z)^2$. To evaluate the sum of the values that this polynomial takes on the variety, we compute:

$$p(M_x, M_y, M_z) = \text{Tr}(M_x + M_y + M_z)^2 = \text{Tr} \begin{bmatrix} 0 & 0 & 0 & 0 & 0 \\ 2 & 3 & 2 & 2 & 2 \\ 3 & 2 & 3 & 2 & 2 \\ 2 & 2 & 2 & 3 & 2 \\ 2 & 2 & 2 & 2 & 3 \end{bmatrix} = 12.$$

As expected, the squares of the sum of the coordinates of each of the five roots are $\{0, 9, 1, 1, 1\}$, with the total sum being equal to 12.

Given any $q \in \mathbb{R}[x_1, \dots, x_n]$, we can then define a Hermite-like matrix $H_q(I)$ as

$$[H_q(I)]_{jk} := \sum_{i=1}^m q(r_i) r_i^{\alpha_j + \alpha_k}. \quad (2)$$

Notice that the rows and columns of $H_q(I)$ are indexed by standard monomials.

Consider now a vector $f = [f_1, \dots, f_m]^T$, and the quadratic form

$$\begin{aligned} f^T H_q(I) f &:= \sum_{j,k=1}^m \sum_{i=1}^m q(r_i) (f_j r_i^{\alpha_j}) (f_k r_i^{\alpha_k}) \\ &= \sum_{i=1}^m q(r_i) (f_1 r_i^{\alpha_1} + \cdots + f_m r_i^{\alpha_m})^2 \\ &= \text{Tr}[(q f^2)(M_{x_1}, \dots, M_{x_n})]. \end{aligned} \quad (3)$$

As we see, the matrix $H_q(I)$ is a specific representation, in a basis given by standard monomials, of a quadratic form $H_q : \mathbb{C}[x]/I \rightarrow \mathbb{C}$, with $H_q : f \rightarrow \sum_{i=1}^m (qf^2)(r_i)$. The expressions in (3) allow us to explicitly compute a matrix representation of this quadratic map. (What is the other “natural” representation of this map?)

The following theorem then generalizes the results of the univariate case, and enable, among other things, to do root counting.

Theorem 2. *The signature of the matrix $H_q(I)$ is equal to the number of real points r_i in $V(I)$ for which $q(r_i) > 0$, minus the number of real points for which $q(r_i) < 0$.*

Corollary 3. *Consider a zero dimensional ideal I . The signature of the matrix $H_1(I)$ is equal to the number of real roots, i.e., $|V(I) \cap \mathbb{R}^n|$.*

In the general (non-radical) case, we would take the property (3) as the definition of $H_q(I)$, instead of (2). Also, in Theorem 2, multiple real zeros are counted only once.

2 Parametric versions

One of the most appealing properties of Groebner-based eigenvalue methods is that they allow us to extend many of the results to the *parametric* case, i.e., when we are interested in obtaining all solutions of a polynomial system as a function of some additional parameters η_i .

Consider for simplicity the case of a single parameter η , and a polynomial system defined by $p_i(x, \eta) = 0$. In order to solve this for any fixed η , we need to compute a Groebner basis of the corresponding ideal. However, when η changes, it is possible that the resulting set of polynomials is no longer a GB. A way of fixing this inconvenience is to compute instead a *comprehensive Groebner basis*, which is a set of polynomials with the property that it remains a Groebner basis of I for all possible specializations of the parameters. Using the corresponding monomials as a basis for the quotient space, we can give an eigenvalue characterization of the solutions for all values of η .

3 SOS on quotients

For simplicity, we assume throughout that the ideal I is radical. We can interpret the previous result as essentially stating the fact that when a polynomial is nonnegative on a finite variety, then it is a sum of squares on the quotient ring; see [Par02].

Theorem 4. *Let $f(x)$ be nonnegative on $\{x \in \mathbb{R}^n | h_i(x) = 0\}$. If the ideal $I = \langle h_1, \dots, h_m \rangle$ is radical, then $f(x)$ is a sum of squares in the quotient ring $\mathbb{R}[x]/I$, i.e., there exist polynomials q_i, λ_i , such that*

$$f(x) = \sum_i q_i^2(x) + \sum_{i=1}^m \lambda_i(x) h_i(x).$$

Remark 5. *The assumption that I is radical (or a suitable local modification) is necessary when $f(x)$ is nonnegative but not strictly positive. For instance, the polynomial $f = x$ is nonnegative over the variety defined by the (non-radical) ideal $\langle x^2 \rangle$, although no decomposition of the form $x = s_0(x) + \lambda(x)x^2$ (where s_0 is SOS), can possibly exist.*

References

- [Par02] P. A. Parrilo. An explicit construction of distinguished representations of polynomials non-negative over finite sets. Technical Report IfA Technical Report AUT02-02. Available from <http://control.ee.ethz.ch/~parrilo>, ETH Zürich, 2002.

Lecture 17

Lecturer: Pablo A. Parrilo

Scribe: ???

One of our main goals in this course is to achieve a better understanding of the techniques available for polynomial systems over the real field. Today we discuss how to certify infeasibility for polynomial equations over the reals, and contrast these approaches with well-known results in linear algebra, linear programming, and complex algebraic geometry.

We will discuss the possible convergence of these schemes in the general case later in the course, concentrating today on an elementary proof of the finite convergence in the zero-dimensional case [Par02].

1 Infeasibility of real polynomial equations

Based on what we have learned in the past weeks, we have a quite satisfactory answer to the question of when a system of polynomial equations has solutions over the complex field. Indeed, as we have seen, given a system of polynomial equations $\{h_i(x) = 0, i = 1, \dots, m\}$, we can form the associated ideal $I = \langle h_1, \dots, h_m \rangle$. By the Nullstellensatz, the associated complex variety $V(I)$ (i.e., the solution set $\{x \in \mathbb{C}^n \mid h_i(x) = 0\}$) will be empty if and only if $I = \mathbb{C}[x]$, or equivalently, $1 \in I$. Computationally, this condition can be checked by computing a reduced Groebner basis of I (with respect to any term ordering), which will be equal to $\{1\}$ if this holds.

What happens, however, when we are interested in *real* solutions, and not just complex ones? Or, if not only we have equations, but also inequalities? Consider, for instance, the basic semialgebraic set given by

$$S = \{x \in \mathbb{R}^n \mid f_i(x) \geq 0, \quad h_i(x) = 0\}. \quad (1)$$

How to decide if the set S is empty? Can we give a Groebner-like criterion to that demonstrate the infeasibility of this system of equations? Even worse, do we even know that this question can be decided algorithmically¹?

Fortunately for us, a famous result, the Tarski-Seidenberg theorem, guarantees the algorithmic solvability of this problem (in fact, of a much larger class of problems, that may include quantifiers). We will discuss this powerful approach in more detail later, when presenting cylindrical algebraic decomposition (CAD) techniques, concentrating instead in a more direct way of tackling the feasibility problem.

2 Certificates

Discuss certificates: NP/co-NP, Linear algebra, LP, Nullstellensatz, P-satz

ToDo

3 The zero-dimensional case

What happens in the case where the equations in the system (1) define a zero dimensional ideal? It should be intuitively obvious that, in some sense, such a finite certificate exists. Indeed, *if* we had access to all the roots, of which there are a finite number, just by evaluating the corresponding expressions we could decide the feasibility or infeasibility. As we will see, we can actually “encode” this process in a set of polynomials, that prove the existence of these certificates.

¹There are certainly similar-looking problems that are *not* decidable. A famous one is the solvability of polynomial equations over the integers. This is Hilbert’s 10th problem, solved in 1970 by Matiyasevich; see [Dav73] for a full account of the solution and historical remarks. This result implies, in particular, the nonexistence of an algorithm to solve integer quadratic programming; see [Jer73].

Theorem 1. Consider the set S in (1), and assume the ideal $I = \langle h_1, \dots, h_m \rangle$ is radical. Then, S is empty if and only if there exists a decomposition

$$-1 = s_0(x) + \sum_{i=1} s_i(x)f_i(x) + \sum_{i=1} \lambda_i(x)h_i(x).$$

where the s_i are sums of squares.

Notice that we can equivalently write

$$-1 \equiv s_0(x) + \sum_{i=1} s_i(x)f_i(x) \pmod{I}.$$

It should be clear that one direction of the implication is obvious (which one?).

4 Optimization

Since optimization can be interpreted as a parametrized family of feasibility problems, we can directly apply these results towards optimization of polynomial or rational functions. For instance, we have the following result:

Theorem 2. Let $p(x)$ be nonnegative on $S = \{x \in \mathbb{R}^n \mid f_i(x) \geq 0, h_i(x) = 0\}$, and assume that the ideal $I = \langle h_1, \dots, h_m \rangle$ is radical. Consider the optimization problem

$$\max \gamma \quad \text{s.t.} \quad p(x) - \gamma = s_0(x) + \sum_{i=1} s_i(x)f_i(x) + \sum_{i=1} \lambda_i(x)h_i(x).$$

where the s_i are sums of squares, and the decision variables are γ and the coefficients of the polynomials $s_i(x)$, $\lambda_i(x)$. Then, the optimal value of γ is exactly equal to the minimum of $p(x)$ over S .

Notice that this is exactly a sum of squares program, since all the constraints are linear and/or sum of squares constraints.

Remark 3. The assumption that I is radical (or a suitable local modification) is necessary when $p(x)$ is nonnegative but not strictly positive. For instance, the polynomial $p = x$ is nonnegative over the variety defined by the (non-radical) ideal $\langle x^2 \rangle$, although no decomposition of the form $x = s_0(x) + \lambda(x)x^2$ (where s_0 is SOS), can possibly exist.

More details will follow...

References

- [Dav73] M. Davis. Hilbert's tenth problem is unsolvable. *Amer. Math. Monthly*, 80:233–269, 1973.
- [Jer73] R.G. Jeroslow. There cannot be any algorithm for integer programming with quadratic constraints. *Operations Res.*, 21:221–224, 1973.
- [Par02] P. A. Parrilo. An explicit construction of distinguished representations of polynomials non-negative over finite sets. Technical Report IfA Technical Report AUT02-02. Available from <http://control.ee.ethz.ch/~parrilo>, ETH Zürich, 2002.

Lecture 18

Lecturer: Pablo A. Parrilo

Scribe: ???

Quantifier elimination (QE) is a very powerful procedure for problems involving first-order formulas over real fields. The cylindrical algebraic decomposition (CAD) is a technique for the “efficient” implementation of QE, that effectively reduces an seemingly infinite problem into a finite (but potentially large) instance. For much more information about QE and CAD (including a reprint of Tarski’s original 1930 work), we recommend the book [CJ98].

1 Quantifier elimination

A quantifier-free formula is an expression consisting of polynomial equations ($f(x) = 0$) and inequalities ($f(x) \leq 0$) combined using the Boolean operators \wedge (and), \vee (or), and \Rightarrow (implies). We often also allow strict inequalities $f(x) > 0$ and inequations $f(x) \neq 0$, since these are just shorthands for particular boolean combinations of equations and inequalities.

In general, a *formula* (in prenex form) is an expression in the variables $x = (x_1, \dots, x_n)$ of the following type:

$$(Q_1 x_1) \dots (Q_s x_s) \quad \mathcal{F}(f_1(x), \dots, f_r(x)) \quad (1)$$

where Q_i is one of the quantifiers \forall (for all) and \exists (there exists). Furthermore, $\mathcal{F}(f_1(x), \dots, f_r(x))$ is assumed to be a quantifier-free formula. If there is a quantifier corresponding to the variable x_i , we say that x_i is *quantified*, or *free* otherwise.

Example 1. *The following are valid formulas*

$$\begin{aligned} &(\forall x) [(x \geq 0) \Rightarrow (x^2 + ax + b \geq 0)] \\ &(\forall x)(\exists y) [x > y^2] \\ &(\forall \delta)(\exists \epsilon) [(\epsilon^2 + \delta^2 \leq 1) \vee (\epsilon \neq 0)] \Rightarrow [\delta < 1]. \end{aligned}$$

The first formula has two free variables (since the variables a and b are unquantified), while for the other two all variables are quantified.

We will interpret the symbols in a formula as taking only real values. Notice that a formula without free variables (usually called a *closed* formula or a *sentence*) is either true or false. For instance, the last two expressions in Example 1 are sentences, with the first one being false and the second being true. Notice also that the truth value may depend on the order of the quantifiers.

Tarski showed that for every formula including quantifiers there is always an equivalent quantifier free formula. Obtaining the latter from the former is called quantifier elimination.

Theorem 2 (Tarski-Seidenberg). *For every first-order formula over the real field there exists an equivalent quantifier-free formula. Furthermore, there is an explicit algorithm to compute this quantifier-free formula.*

The Tarski-Seidenberg theorem is an extremely powerful result, since it provides a complete characterization and algorithmic technique for an extremely large collection of problems involving polynomials. Unfortunately, there are very serious computational barriers to the efficient practical implementation of these ideas, since the resulting algorithms have extremely poor scaling properties, with respect to the number of variables (towers of exponentials). Newer methods, such as the (partial) cylindrical algebraic decomposition (CAD) technique due to Collins and described below, or the critical point method, are by comparison much better. Nevertheless, by necessity they still behave exponentially (or worse) in terms of the number of variables.

2 Tarski-Seidenberg

Example 3. Consider the quantified first-order formula:

$$(\forall x)(\forall y) [(x^2 + ay^2 \leq 1) \Rightarrow (ax^2 - a^2xy + 2 \geq 0)]. \quad (2)$$

This formula is equivalent to the quantifier free expression:

$$(a \geq 0) \wedge (a^3 - 8a - 16 \leq 0),$$

which defines the interval $[0, a_\star]$, where $a_\star \approx 3.538$. Thus, the original expression (2) is true only for $a \in [0, a_\star]$.

2.1 Geometric interpretation

A geometric interpretation of the Tarski-Seidenberg theorem is the following:

Theorem 4. The projection of a semialgebraic set is semialgebraic.

2.2 Applications

Add many more

ToDo

Static output feedback An early application of Tarski-Seidenberg in control theory was the “solution” of the static output feedback stabilization problem in [ABJ75]. Given matrices $A \in \mathbb{R}^{n \times n}$, $B \in \mathbb{R}^{n \times m}$, we want to find a matrix $K \in \mathbb{R}^{m \times n}$ such that the matrix $A + BK$ is Hurwitz, i.e., all its eigenvalues are in the left-hand plane. Since the existence of such a matrix can be easily expressed as a formula in first order logic¹, the decidability and existence of an effective (but not efficient) algorithm immediately follows.

Simultaneous stabilization A very interesting result by Blondel [Blo94, BG93] shows that the simultaneous stabilization of three linear time-invariant systems is *not* decidable (and thus, cannot be semialgebraic). Notice however that, for any given bound on the degree of the controller, the problem is decidable.

3 Cylindrical Algebraic Decomposition (CAD)

There are a few approaches for effective implementation of the QE procedure. One of the most well-known, which is also relatively easy to understand, is the cylindrical algebraic decomposition (CAD) due to Collins [Col75]. We describe the elements of this approach below. We remark that much better algorithms (in the theoretical complexity sense) are known; see for instance the article by Renegar [Ren91] (also reprinted in [CJ98]) or [BPR03]. In particular, for CAD the number of operations usually scales in a doubly exponential fashion with the number of variables, while the newer methods are doubly exponential in the number of quantifier alternations.

¹For instance, $(\exists K)(\forall x)(\forall \lambda) [(A + BK)x = \lambda x \vee x \neq 0] \Rightarrow [\Re(\lambda) \leq 0]$. Notice that we are being a bit sloppy with notation, since for a fully real formulation, we should split x and λ into real and imaginary parts. There are many other equivalent expressions, using for instance a Lyapunov equation, or the Routh array.

3.0.1 Description

Given a set P of multivariate polynomials in n variables, a CAD is a special partition of \mathbb{R}^n into components, called *cells*, over which all the polynomials have constant signs. The algorithm for computing a CAD also provides a point in each cell, called *sample point*, which can be used to determine the sign of the polynomials in the cell.

A cell is called *cylindrical* if it has the form $S \times \mathbb{R}^k$, for some $k \leq n$. A decomposition of \mathbb{R}^n is a CAD if all polynomials have constant sign on each cell, and all cells are cylindrical.

The CAD associated to the formula (1) depends only on its quantifier free part $\mathcal{F}(f_1(x), \dots, f_r(x))$. Since all possible truth values of the formula are in correspondence with the values at the sample points, we can use the CAD to evaluate its truth value, and to perform quantifier elimination.

The basic CAD construction consists of two steps: *projection* and *lifting* (plus an additional third one, if formula construction is desired).

In the first projection phase, we compute successive sets of polynomials in $n-1, n-2, \dots, 1$ variables. The main idea is, given an input set of polynomials, to compute at each step a new set of polynomials obtained by eliminating one variable at a time. In general, the elimination order does matter and a good choice leads to lower computational complexity.

The second phase (lifting) constructs a decomposition of \mathbb{R} , at the lowest level of projection, after all but one variable have been eliminated. This decomposition of \mathbb{R} is successively extended to a decomposition of \mathbb{R}^n .

The basic operations necessary in the construction of CADs are (sub)resultants and (sub)discriminants.

Complete

ToDo

An implementation of (an improved version of) the CAD method for quantifier elimination is the software package QEPCAD [Bro03].

References

- [ABJ75] B. D. O. Anderson, N. K. Bose, and E. I. Jury. Output feedback stabilization and related problems—solution via decision methods. *IEEE Transactions on Automatic Control*, 20:53–66, 1975.
- [BG93] V. Blondel and M. Gevers. Simultaneous stabilizability of three linear systems is rationally undecidable. *Mathematics of Control, Signals, and Systems*, 6(2):135–145, 1993.
- [Blo94] V. Blondel. *Simultaneous stabilization of linear systems*, volume 191 of *Lecture Notes in Control and Information Sciences*. Springer-Verlag London Ltd., London, 1994.
- [BPR03] S. Basu, R. Pollack, and M.-F. Roy. *Algorithms in real algebraic geometry*, volume 10 of *Algorithms and Computation in Mathematics*. Springer-Verlag, Berlin, 2003.
- [Bro03] C.W. Brown. *QEPCAD - Quantifier Elimination by Partial Cylindrical Algebraic Decomposition*, 2003. Available from <http://www.cs.usna.edu/~qepcad/B/QEPCAD.html>.
- [CJ98] B. F. Caviness and J. R. Johnson, editors. *Quantifier elimination and cylindrical algebraic decomposition*, Texts and Monographs in Symbolic Computation, Vienna, 1998. Springer-Verlag.
- [Col75] G. E. Collins. Quantifier elimination for real closed fields by cylindrical algebraic decomposition. In *Automata theory and formal languages (Second GI Conf., Kaiserslautern, 1975)*, pages 134–183. Lecture Notes in Comput. Sci., Vol. 33. Springer, Berlin, 1975.

- [Ren91] J. Renegar. Recent progress on the complexity of the decision problem for the reals. In *Discrete and computational geometry (New Brunswick, NJ, 1989/1990)*, volume 6 of *DIMACS Ser. Discrete Math. Theoret. Comput. Sci.*, pages 287–308. Amer. Math. Soc., Providence, RI, 1991.

Lecture 19

Lecturer: Pablo A. Parrilo

Scribe: ???

Today we continue with some additional aspects of quantifier elimination. We will then recall the Positivstellensatz and its relations with semidefinite programming. After introducing copositive matrices, we present Pólya's theorem on positive forms on the simplex, and the associated relaxations. Finally, we conclude with an important result due to Schmüdgen about representation of positive polynomials on compact sets.

1 Certificates

Talk about certificates in QE

ToDo

2 Psatz revisited

Recall the statement of the Positivstellensatz.

Theorem 1 (Positivstellensatz). *Consider the set $S = \{x \in \mathbb{R}^n \mid f_i(x) \geq 0, h_i(x) = 0\}$. Then,*

$$S = \emptyset \quad \Leftrightarrow \quad \exists f, h \in \mathbb{R}[x] \text{ s.t. } \begin{cases} f + h &= -1 \\ f &\in \mathbf{cone}\{f_1, \dots, f_s\} \\ h &\in \mathbf{ideal}\{h_1, \dots, h_t\} \end{cases}$$

Once again, since the conditions on the polynomials f, h are convex and affine, respectively, by restricting their degree to be less than or equal to a given bound d we have a finite-dimensional semidefinite programming problem.

2.1 Hilbert 17th problem

As we have seen, in the general case nonnegative multivariate polynomials can fail to be a sum of squares (the Motzkin polynomial being the classical counterexample). As part of his famous list of twenty-three problems that he presented at the International Congress of Mathematicians in 1900, David Hilbert asked the following¹:

17. Expression of definite forms by squares. A rational integral function or form in any number of variables with real coefficient such that it becomes negative for no real values of these variables, is said to be definite. The system of all definite forms is invariant with respect to the operations of addition and multiplication, but the quotient of two definite forms in case it should be an integral function of the variables is also a definite form. The square of any form is evidently always a definite form. But since, as I have shown, not every definite form can be compounded by addition from squares of forms, the question arises which I have answered affirmatively for ternary forms whether every definite form may not be expressed as a quotient of sums of squares of forms. At the same time it is desirable, for certain questions as to the possibility of certain geometrical constructions, to know whether the coefficients of the forms to be used in the expression may always be taken from the realm of rationality given by the coefficients of the form represented.

¹ This text was obtained from <http://aleph0.clarku.edu/~djoyce/hilbert/>, and corresponds to Newson's translation of Hilbert's original German address. In that website you will also find links to the current status of the problems, as well as the original German text.

In other words, can we write every nonnegative polynomial as a sum of squares of *rational functions*? As we show next, this is a rather direct consequence of the Psatz. Of course, it should be clear (and goes without saying) that we are (badly) inverting the historical order! In fact, much of the motivation for the development of real algebra came from Hilbert's question.

How can we use the Psatz to prove that a polynomial $p(x)$ is nonnegative? Clearly, p is nonnegative if and only if the set $\{x \in \mathbb{R}^n \mid p(x) < 0\}$ is empty. Since our version of the Psatz does not allow for strict inequalities (there are slightly more general, though equivalent, formulations that do), we'll need a useful trick discussed earlier ("Rabinowitch's trick"). Introducing a new variable z , the nonnegativity of $p(x)$ is equivalent to the emptiness of the set described by

$$-p(x) \geq 0, \quad 1 - zp(x) = 0.$$

The Psatz can be used to show that this holds if and only if there exist polynomials $s_0, s_1, t \in \mathbb{R}[x, z]$ such that

$$s_0(x, z) - s_1(x, z) \cdot p + t(x, z) \cdot (1 - zp) = -1,$$

where s_0, s_1 are sums of squares. Replace now $z \rightarrow 1/p(x)$, and multiply by p^{2k} (where k is sufficiently large) to obtain

$$\tilde{s}_0 - \tilde{s}_1 \cdot p = -p^{2k},$$

where \tilde{s}_0, \tilde{s}_1 are sums of squares in $\mathbb{R}[x]$. Solving now for p , we have:

$$p(x) = \frac{\tilde{s}_0(x) + p(x)^{2k}}{\tilde{s}_1(x)} = \frac{\tilde{s}_1(x)(\tilde{s}_0(x) + p(x)^{2k})}{\tilde{s}_1^2(x)},$$

and since the numerator is a sum of squares, it follows that $p(x)$ is indeed a sum of squares of rational functions.

3 Copositive matrices and Pólya's theorem

An interesting class of matrices are the so-called *copositive matrices*, which are those for which the associated quadratic form is nonnegative on the nonnegative orthant.

Definition 2. A matrix $M \in \mathcal{S}^n$ is copositive if it satisfies

$$x^T M x \geq 0, \quad \text{for all } x_i \geq 0.$$

As opposed to positive semidefiniteness, which can be checked in polynomial time, the recognition problem for copositive matrices is an NP-hard problem. The set of copositive is a closed convex cone, for which checking membership is a difficult problem.

There are many interesting applications of copositive matrices. Among others, we mention:

- Consider a graph G , with A being its adjacency matrix. The stability number α of the graph G is equal to the cardinality of its largest stable set. By a result of Motzkin and Straus, it is known that it can be obtained as:

$$\frac{1}{\alpha(G)} = \min_{x_i \geq 0, \sum_i x_i = 1} x^T (I + A) x$$

This implies that $\alpha(G) \leq \gamma$ if and only if the matrix $\gamma \cdot (I + A) - ee^T$ is copositive.

- Another interesting application of copositive matrices is in the performance analysis of queueing networks; see e.g. [KM96]. Modulo some (important) details, the basic idea is to use a quadratic function $x^T M x$ as a Lyapunov function, where the matrix M is copositive and x represents the lengths of the queues.

An important related result is Pólya's theorem on positive forms on the simplex:

Theorem 3 (Pólya). *Consider a homogeneous polynomial in n variables of degree d , that is strictly positive in the unit simplex $\Delta_n := \{x \in \mathbb{R}^n \mid x_i \geq 0, \sum_{i=1}^n x_i = 1\}$. Then, for large enough k , the polynomial $(x_1 + \cdots + x_n)^k p(x)$ has nonnegative coefficients.*

We can provide a natural hierarchy of sufficient conditions for a matrix to be copositive. Completeness of this hierarchy follows directly from Pólya's theorem [Par00].

There are some very interesting connections between Pólya's result and a foundational theorem in probability known as De Finetti's exchangeability theorem.

4 Positive polynomials

The Positivstellensatz allows us to obtain certificates of the emptiness of a basic semialgebraic set, explicitly given by polynomials.

What if we want to apply this for optimization? As we have seen, it is relatively straightforward to convert an optimization problem to a family of feasibility problems, by considering the sublevel sets, i.e., the sets $\{x \in \mathbb{R}^n \mid f(x) \leq \gamma\}$.

In the general case of constrained problems, however, using the Psatz we will require conditions that are not linear in the unknown parameter γ (because we need products between the constraints), and this presents a difficulty to the direct use of SDP. Notice nevertheless, that the problem is certainly an SDP for any fixed value of γ , and it thus quasiconvex (which is almost as good, except for the fact that we cannot use "standard" SDP solvers to solve it directly, but rather rely on methods such as bisection).

Theorem 4 ([Sch91]). *If $p(x)$ is strictly positive on $K = \{x \in \mathbb{R}^n \mid f_i(x) \geq 0\}$, and K is compact, then $p(x) \in \text{cone}\{f_1, \dots, f_s\}$.*

In the next lecture we will describe the basic elements of Schmüdgen's proof. His approach combines both algebraic tools (using the Positivstellensatz to prove the boundedness of certain operators) and functional analysis (spectral measures of commuting families of operators and the Hahn-Banach theorem). We will also describe some alternative versions due to Putinar, as well as a related purely functional-analytic result due to Megretski.

For a comprehensive treatment and additional references, we mention [BCR98, Mar00, PD01] among others.

References

- [BCR98] J. Bochnak, M. Coste, and M-F. Roy. *Real Algebraic Geometry*. Springer, 1998.
- [KM96] P. R. Kumar and S. P. Meyn. Duality and linear programs for stability and performance analysis of queuing networks and scheduling policies. *IEEE Trans. Automat. Control*, 41(1):4–17, 1996.
- [Mar00] M. Marshall. *Positive polynomials and sums of squares*. Dottorato de Ricerca in Matematica. Dept. di Mat., Univ. Pisa, 2000.
- [Par00] P. A. Parrilo. *Structured semidefinite programs and semialgebraic geometry methods in robustness and optimization*. PhD thesis, California Institute of Technology, May 2000. Available at <http://resolver.caltech.edu/CaltechETD:etd-05062004-055516>.

- [PD01] A. Prestel and C. N. Delzell. *Positive polynomials: from Hilbert's 17th problem to real algebra*. Springer Monographs in Mathematics. Springer, 2001.
- [Sch91] K. Schmüdgen. The K -moment problem for compact semialgebraic sets. *Math. Ann.*, 289:203–206, 1991.

Lecture 20

Lecturer: Pablo A. Parrilo

Scribe: ???

In this lecture we introduce Schmüdgen's theorem about the K -moment problem (or equivalently, on the representation of positive polynomials) and describe the basic elements in his proof. This approach combines both algebraic tools (using the Positivstellensatz to prove the boundedness of certain operators) and functional analysis (spectral measures of commuting families of operators and the Hahn-Banach theorem). We will also describe some alternative versions due to Putinar, as well as a related purely functional-analytic result due to Megretski.

For a comprehensive treatment and additional references, we mention [BCR98, Mar00, PD01] among others.

1 Positive polynomials

As we have seen, the Positivstellensatz allows us to obtain certificates of the emptiness of a basic semialgebraic set, explicitly given by polynomials. When looking for bounded degree certificates, this provides a natural hierarchy of SDP-based conditions [Par00, Par03].

What if we want to apply this for the particular case of optimization? As we have seen, it is relatively straightforward to convert a polynomial optimization problem to a one-parameter family of feasibility problems, by considering the sublevel sets, i.e., the sets $\{x \in \mathbb{R}^n \mid f(x) \leq \gamma\}$.

In the general case of constrained problems, however, using the full power of the Psatz will yield conditions that are not linear in the unknown parameter γ (because we need products between the constraints and objective function), and in principle, this presents a difficulty to the direct use of SDP. Notice nevertheless, that the problem is certainly an SDP for any fixed value of γ , and is thus quasiconvex (which is almost as good, except for the fact that we cannot use "standard" SDP solvers to solve it directly, but rather rely on methods such as bisection).

Of course, we can always produce specific families of certificates that are linear in γ , and use them for optimization (e.g., like we did in the copositivity case). However, in general it is unclear whether the desired family is "complete," in the sense that we will be able to prove arbitrarily good bounds on the optimal value as the degree of the polynomials grows to infinity.

2 Schmüdgen's theorem

In 1991, Schmüdgen presented a characterization of the moment sequences of measures supported on a compact semialgebraic K (the K -moment problem). As in the one-dimensional case we studied earlier the question is, given an (infinite) sequence of moments, decide whether it actually corresponds to a nonnegative measure with support on a given set K .

His solution combined both real algebraic methods (the Psatz), with some functional analytic tools (reproducing kernel Hilbert spaces, bounded operators, and the spectral theorem).

This characterization of moment sequences can be used, in turn, to produce an explicit description of the set of strictly positive polynomials on a compact semialgebraic set:

Theorem 1 ([Sch91]). *If $p(x)$ is strictly positive on $K = \{x \in \mathbb{R}^n \mid f_i(x) \geq 0\}$, and K is compact, then $p(x) \in \text{cone}\{f_1, \dots, f_m\}$.*

expand

ToDo

There are several interesting ideas in the proof; a coarse description follows. The first step is to use the Positivstellensatz to produce an algebraic certificate of the compactness of the set K . Then the

given moment sequence (which is a positive definite function on the semigroup of monomials) is used to construct a particular pre-Hilbert space and its completion (namely, the associated reproducing kernel Hilbert space). In this Hilbert space, we consider linear operators T_{x_i} given by multiplication by the coordinate variables, and use the algebraic certificate of compactness to prove that these are bounded. Now, the T_{x_i} are a finite collection of pairwise commuting, bounded, self-adjoint operators, and thus there exists a spectral measure for the family, from which a measure, only supported in K , can be extracted. Finally, a Hahn-Banach (separating hyperplane) argument is used to prove the final result.

2.1 Putinar's approach

The theorem in the previous section requires (in principle) all $2^m - 1$ squarefree products of constraints¹. Putinar [Put93] presented a modified formulation (under stronger assumptions) for which the representation is *linear* in the constraints. We introduce the following concept:

Definition 2. Let $\{f_1, \dots, f_m\} \subset \mathbb{R}[x]$. The preprime generated by the f_i , and denoted by $\mathbf{preprime}\{f_1, \dots, f_m\}$ is the set of all polynomials of the form $s_0 + s_1 f_1 + \dots + s_m f_m$, where all the s_i are sums of squares.

Notice that $\mathbf{preprime}\{f_i\} \subset \mathbf{cone}\{f_i\}$, and that every element in the preprime takes only nonnegative values on $\{x \in \mathbb{R}^n, f_i(x) \geq 0\}$.

Theorem 3 ([Put93]). Consider a set $K = \{x \in \mathbb{R}^n \mid f_i(x) \geq 0\}$, such that there exists a $q \in \mathbf{preprime}\{f_1, \dots, f_m\}$ and $\{x \in \mathbb{R}^n, q(x) \geq 0\}$ is compact (this implies that K is compact). Then, $p(x) > 0$ on K if and only if $p(x) \in \mathbf{preprime}\{f_1, \dots, f_m\}$.

Notice that here, the polynomial q serves as an algebraic certificate of the compactness of K , so in this case the Psatz is not needed.

Putinar's theorem was used by Lasserre to present a hierarchy of semidefinite relaxations for polynomial optimization, based on the dual moment interpretation [Las01].

2.2 Tradeoffs

In principle (and often, in practice) there is a tradeoff between how “expressive” our family of certificates is, the quality of the resulting bounds, and the complexity of finding proofs.

On one extreme, the most general method is the Psatz, as it encapsulates pretty much every possible “algebraic deduction,” and will certainly provide the strongest bounds, since it includes the other techniques as special cases. For optimization, Schmüdgen's theorem provides the advantages of a linear representation, although (possibly) at the cost of having a large number of products between the constraints. Finally, the Putinar approach has a reduced number of constraints (and thus, SOS multipliers), although the obtained bounds can potentially be much weaker than the previous ones.

In the end, the decision concerning what approach to use should be dictated by the available computational resources, i.e., the size of the SDPs that we can solve in a reasonable time. It is not difficult to produce examples with significant gaps between the corresponding bounds; see for instance [Ste96] for a particularly simple example, that is trivial for the Psatz, but for which either the Schmüdgen or Putinar representations need large degree refutations.

Add examples

ToDo

¹Recall that in practice, this may not be an issue at all, since the restriction on the degree of the certificates imposes a strict limit on how many products can be included.

2.3 Trigonometric case

Recently, Megretski [Meg03] analyzed the trigonometric case. We introduce the following notation: let $\mathbb{T}_n = \{z \in \mathbb{C}^n, |z_i| = 1\}$ be the n -dimensional torus, P_n is the set of multivariate Laurent polynomials, and $RP_n \subset P_n$ are the Laurent polynomials that are real-valued on \mathbb{T}_n .

Theorem 4 ([Meg03]). *Let $\{F, Q_1, \dots, Q_m\} \subset RP_n$, such that $F(z) > 0$ for all $z \in \mathbb{T}_n$ satisfying $Q_1(z) = \dots = Q_m(z) = 0$. Then there exist $V_1, \dots, V_r \in P_n$, $H_1, \dots, H_m \in RP_n$, such that*

$$F(z) = \sum_{i=1}^r |V_i(z)|^2 + \sum_{j=1}^m H_j(z)Q_j(z).$$

Notice that, by splitting into real and imaginary part, this corresponds to a special kind of (standard) polynomials, and a compact semialgebraic set (so in principle, any of the previous theorems would apply). Of course, the result exploits the complex structure for a more concise representation.

In particular, Megretski's proof is purely functional-analytic, the main tools being Bochner's theorem and Hahn-Banach. Bochner's theorem is an important result in harmonic analysis, that characterizes a positive definite function on an Abelian group in terms of the nonnegativity of its Fourier transform.

Notice that the theorem above deals only with the equality case (no inequalities), and the feasible set is compact (since so is \mathbb{T}^n). It essentially states that a positive polynomial is a sum of squares modulo the ideal generated by the Q_i . Recall we have proved similar results in the zero-dimensional case, and this theorem naturally generalizes these.

In simplified terms, one reason why trigonometric (or Laurent) polynomials are somewhat “easier” than the general case is because in this case there is a *group* structure, as opposed to the *semigroup* structure of regular monomials. For the group case, the corresponding theory is the classical harmonic analysis on abelian groups (e.g., [Rud90]); while for semigroups there is the newer, but well-developed characterizations of positive functions on (Abelian) semigroups; see for instance [BCR84].

We also mention that there are “purely algebraic” versions of these theorems, that do not use functional analytic ideas (e.g., [Mar00]). Roughly, the role played by the compactness of K in proving the boundedness of the operators T_{x_i} is replaced with a property called *Archimedeanity* of the corresponding preorder.

References

- [BCR84] C. Berg, J. P. R. Christensen, and P. Ressel. *Harmonic analysis on semigroups*, volume 100 of *Graduate Texts in Mathematics*. Springer-Verlag, New York, 1984.
- [BCR98] J. Bochnak, M. Coste, and M-F. Roy. *Real Algebraic Geometry*. Springer, 1998.
- [Las01] J. B. Lasserre. Global optimization with polynomials and the problem of moments. *SIAM J. Optim.*, 11(3):796–817, 2001.
- [Mar00] M. Marshall. *Positive polynomials and sums of squares*. Dottorato de Ricerca in Matematica. Dept. di Mat., Univ. Pisa, 2000.
- [Meg03] A. Megretski. Positivity of trigonometric polynomials. In *Proceedings of the 42th IEEE Conference on Decision and Control*, pages 3814–3817, 2003.
- [Par00] P. A. Parrilo. *Structured semidefinite programs and semialgebraic geometry methods in robustness and optimization*. PhD thesis, California Institute of Technology, May 2000. Available at <http://resolver.caltech.edu/CaltechETD:etd-05062004-055516>.
- [Par03] P. A. Parrilo. Semidefinite programming relaxations for semialgebraic problems. *Math. Prog.*, 96(2, Ser. B):293–320, 2003.

- [PD01] A. Prestel and C. N. Delzell. *Positive polynomials: from Hilbert's 17th problem to real algebra*. Springer Monographs in Mathematics. Springer, 2001.
- [Put93] M. Putinar. Positive polynomials on compact semi-algebraic sets. *Indiana Univ. Math. J.*, 42(3):969–984, 1993.
- [Rud90] W. Rudin. *Fourier analysis on groups*. Wiley Classics Library. John Wiley & Sons Inc., New York, 1990.
- [Sch91] K. Schmüdgen. The K -moment problem for compact semialgebraic sets. *Math. Ann.*, 289:203–206, 1991.
- [Ste96] G. Stengle. Complexity estimates for the Schmüdgen Positivstellensatz. *J. Complexity*, 12(2):167–174, 1996.

Lecture 21

Lecturer: Pablo A. Parrilo

Scribe: ???

In this lecture we study techniques to exploit the symmetry that can be present in semidefinite programming problems, particularly those arising from sum of squares decompositions [GP04]. For this, we present the basic elements of the representation theory of finite groups. There are many possible applications of these ideas in different fields; for the case of Markov chains, see [BDPX05].

1 Groups and their representations

The representation theory of finite groups is a classical topic; good descriptions are given in [FS92, Ser77]. We concentrate here on the finite case; extensions to compact groups are relatively straightforward.

Definition 1. A group consists of a set G and a binary operation “ \cdot ” defined on G , for which the following conditions are satisfied:

1. *Associative:* $(a \cdot b) \cdot c = a \cdot (b \cdot c)$, for all $a, b, c \in G$.
2. *Identity:* There exist $1 \in G$ such that $a \cdot 1 = 1 \cdot a = a$, for all $a \in G$.
3. *Inverse:* Given $a \in G$, there exists $b \in G$ such that $a \cdot b = b \cdot a = 1$.

We consider a finite group G , and an n -dimensional vector space V . We define the associated (infinite) group $GL(V)$, which we can interpret as the set of invertible $n \times n$ matrices. A *linear representation* of the group G is a homomorphism $\rho : G \rightarrow GL(V)$. In other words, we have a mapping from the group into linear transformations of V , that respects the group structure, i.e.

$$\rho(st) = \rho(s)\rho(t) \quad \forall s, t \in G.$$

Example 2. Let $\rho(g) = 1$ for all $g \in G$. This is the trivial representation of the group.

Example 3. For a more interesting example, consider the symmetric group S_n , and the “natural” representation $\rho : S_n \rightarrow GL(\mathbb{C}^n)$, where $\rho(g)$ is a permutation matrix. For instance, for the group of permutations of two elements, $S_2 = \{e, g\}$, where $g^2 = e$, we have

$$\rho(e) = \begin{bmatrix} 1 & 0 \\ 0 & 1 \end{bmatrix}, \quad \rho(g) = \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}.$$

The representation given in Example 3 has an interesting property. The set of matrices $\{\rho(e), \rho(g)\}$ have common invariant subspaces (other than the trivial ones, namely $(0, 0)$ and \mathbb{C}^2). Indeed, we can easily verify that the (orthogonal) one-dimensional subspaces given by (t, t) and $(t, -t)$ are invariant under the action of these matrices. Therefore, the restriction of ρ to those subspaces also gives representations of the group G . In this case, the one corresponding to the subspace (t, t) is “equivalent” (in a well-defined sense) to the trivial representation described in Example 2. The other subspace $(t, -t)$ gives the one-dimensional *alternating* representation of S_2 , namely $\rho_A(e) = 1, \rho_A(g) = -1$. Thus, the representation ρ decomposes as $\rho = \rho_T \oplus \rho_A$, a direct sum of the trivial and the alternating representations.

The same ideas extend to arbitrary finite groups.

Definition 4. An irreducible representation of a group is a linear representation with no nontrivial invariant subspaces.

Theorem 5. Every finite group G has a finite number of nonequivalent irreducible representations ρ_i , of dimension d_i . The relation $\sum_i d_i^2 = |G|$ holds.




Figure 1: Two symmetric optimization problems, one non-convex and the other convex. For the latter, optimal solutions always lie on the fixed-point subspace.

Example 6. Consider the group S_3 (permutations in three elements). This group is generated by the two permutations $s : 123 \rightarrow 213$ and $c : 123 \rightarrow 312$ (“swap” and “cycle”), and has six elements $\{e, s, c, c^2, cs, sc\}$. Notice that $c^3 = e$, $s^2 = e$, and $s = csc$.

The group S_3 has three irreducible representations, two one-dimensional, and one two-dimensional (so $1^2 + 1^2 + 2^2 = |S_3| = 6$). These are:

$$\begin{aligned} \rho_T(s) &= 1, & \rho_T(c) &= 1 \\ \rho_A(s) &= -1, & \rho_A(c) &= 1 \\ \rho_S(s) &= \begin{bmatrix} 0 & 1 \\ 1 & 0 \end{bmatrix}, & \rho_S(c) &= \begin{bmatrix} \omega & 0 \\ 0 & \omega^2 \end{bmatrix} \end{aligned}$$

where $\omega = e^{\frac{2\pi i}{3}}$ is a cube root of 1. Notice that it is enough to specify a representation on the generators of the group.

1.1 Symmetry and convexity

A key property of symmetric *convex* sets is the fact that the “group average” $\frac{1}{|G|} \sum_{g \in G} \sigma(g)x$ always belongs to the set.

Therefore, in convex optimization we can always restrict the solution to the fixed-point subspace

$$\mathcal{F} := \{x | \sigma(g)x = x, \quad \forall g \in G\}.$$

In other words, for convex problems, no “symmetry-breaking” is ever necessary.

As another interpretation, that will prove useful later, the “natural” decision variables of a symmetric optimization problem are the *orbits*, not the points themselves. Thus, we may look for solutions in the quotient space.

1.2 Invariant SDPs

We consider a general SDP, described in geometric form. If \mathcal{L} is an affine subspace of \mathcal{S}^n , and $C, X \in \mathcal{S}^n$, an SDP is given by:

$$\min \langle C, X \rangle \quad \text{s.t.} \quad X \in \mathcal{X} := \mathcal{L} \cap \mathcal{S}_+^n.$$

Definition 7. Given a finite group G , and associated representation $\sigma : G \rightarrow GL(\mathcal{S}^n)$, a σ -invariant SDP is one where both the feasible set and the cost function are invariant under the group action, i.e.,

$$\langle C, X \rangle = \langle C, \sigma(g)X \rangle, \quad \forall g \in G, \quad X \in \mathcal{X} \Rightarrow \sigma(g)X \in \mathcal{X} \quad \forall g \in G$$




Figure 2: The cyclic graph C_n in n vertices (here, $n = 9$).

Example 8. Consider the SDP given by

$$\min a + c, \quad s.t. \quad \begin{bmatrix} a & b \\ b & c \end{bmatrix} \succeq 0,$$

which is invariant under the Z_2 action:

$$\begin{bmatrix} X_{11} & X_{12} \\ X_{12} & X_{22} \end{bmatrix} \rightarrow \begin{bmatrix} X_{22} & -X_{12} \\ -X_{12} & X_{11} \end{bmatrix}.$$

Usually in SDP, the group acts on \mathcal{S}^n through a congruence transformation, i.e., $\sigma(g)M = \rho(g)^T M \rho(g)$, where ρ is a representation of G on \mathbb{C}^n . In this case, the restriction to the fixed-point subspace takes the form:

$$\sigma(g)M = M \quad \implies \quad \rho(g)M - M\rho(g) = 0, \quad \forall g \in G. \quad (1)$$

The Schur lemma of representation theory exactly characterizes the matrices that commute with a group action.

As a consequence of an important structural result (Schur's lemma), it turns out that every representation can be written in terms of a finite number of primitive blocks, the *irreducible representations* of a group.

Theorem 9. Every group representation ρ decomposes as a direct sum of irreducible representations:

$$\rho = m_1 \vartheta_1 \oplus m_2 \vartheta_2 \oplus \dots \oplus m_N \vartheta_N$$

where m_1, \dots, m_N are the multiplicities.

This decomposition induces an isotypic decomposition of the space

$$\mathbb{C}^n = V_1 \oplus \dots \oplus V_N, \quad V_i = V_{i1} \oplus \dots \oplus V_{in_i}.$$

In the symmetry-adapted basis, the matrices in the SDP have a block diagonal form:

$$(I_{m_1} \otimes M_1) \oplus \dots \oplus (I_{m_N} \otimes M_N)$$

In terms of our symmetry-reduced SDPs, this means that not only the SDP block-diagonalizes, but there is also the possibility that many blocks are identical.

1.3 Example: symmetric graphs

Consider the MAXCUT problem on the cycle graph C_n with n vertices (see Figure 2). It is easy to see that the optimal cut has cost equal to n or $n - 1$, depending on whether n is even or odd, respectively.

What would the SDP relaxation yield in this case? If A is the adjacency matrix of the graph, then the SDP relaxations have essentially the form

$$\begin{array}{ll} \text{minimize} & \text{Tr } AX \\ \text{s.t.} & X_{ii} = 1 \\ & X \succeq 0 \end{array} \qquad \begin{array}{ll} \text{maximize} & \text{Tr } \Lambda \\ \text{s.t.} & A \succeq \Lambda \\ & \Lambda \text{ diagonal} \end{array} \quad (2)$$

By the symmetry of the graph, the matrix A is *circulant*, i.e., $A_{ij} = a_{i-j \bmod n}$.

We focus now on the dual form. It should be clear that the cyclic symmetry of the graph induces a cyclic symmetry in the SDP, i.e., if $\Lambda = \text{diag}(\lambda_1, \lambda_2, \dots, \lambda_n)$ is a feasible solution, then $\tilde{\Lambda} = \text{diag}(\lambda_n, \lambda_1, \lambda_2, \dots, \lambda_{n-1})$ is also feasible and achieves the same objective value. Thus, by averaging over the cyclic group, we can always restrict D to be a multiple of the identity matrix, i.e., $\Lambda = \lambda I$. Furthermore, the constraint $A \succeq \lambda I$ can be block-diagonalized via the Fourier matrix (i.e., the irreducible representations of the cyclic group), yielding:

$$A \succeq \lambda I \quad \Leftrightarrow \quad 2 \cos \frac{k\pi}{n} \geq \lambda \quad k = 0, \dots, n-1.$$

From this, the optimal solution of the relaxation can be directly computed, yielding the exact expressions for the upper bound on the size of the cut

$$mc(C_n) \leq SDP(C_n) = \begin{cases} n & n \text{ even} \\ n \cos^2 \frac{\pi}{2n} & n \text{ odd.} \end{cases}$$

Although this example is extremely simple, exactly the same techniques can be applied to much more complicated problems; see for instance [PP04, dKMP⁺, Sch05] for some recent examples.

1.4 Example: even polynomials

Another (but illustrative) example of symmetry reduction is the case of SOS decompositions of even polynomials. Consider a polynomial $p(x)$ that is *even*, i.e., it satisfies $p(x) = p(-x)$. Does this symmetry help in making the computations more efficient?

Complete

ToDo

1.5 Benefits

In the case of semidefinite programming, there are many benefits to exploiting symmetry:

- Replace one big SDP with smaller, coupled problems.
- Instead of checking if a big matrix is PSD, we use one copy of each repeated block (constraint aggregation).
- Eliminates multiple eigenvalues (numerical difficulties).
- For groups, the coordinate change depends only on the group, and not on the problem data.
- Can be used as a general preprocessing scheme. The coordinate change T is unitary, so well-conditioned.

As we will see in the next section, this approach can be extended to more general algebras that do not necessarily arise from groups.

1.6 Sum of squares

In the case of SDPs arising from sum of squares decompositions, a parallel theory can be developed by considering the symmetry-induced decomposition of the full polynomial ring $\mathbb{R}[x]$. Since the details involve some elements of invariant theory, we omit the details here; see [GP04] for the full story.

Include example

ToDo

2 Algebra decomposition

An alternative (and somewhat more general) approach can be obtained by focusing instead on the *associative algebra* generated by the matrices in a semidefinite program.

Definition 10. An associative algebra \mathcal{A} over \mathbb{C} is a vector space with a \mathbb{C} -bilinear operation $\cdot : \mathcal{A} \times \mathcal{A} \rightarrow \mathcal{A}$ that satisfies

$$x \cdot (y \cdot z) = (x \cdot y) \cdot z, \quad \forall x, y, z \in \mathcal{A}.$$

In general, associative algebras do not need to be commutative (i.e., $x \cdot y = y \cdot x$). However, that is an important special case, with many interesting properties. Important examples of finite dimensional associative algebras are:

- Full matrix algebra $\mathbb{C}^{n \times n}$, standard product.
- The subalgebra of square matrices with equal row and column sums.
- The n -dimensional algebra generated by a single $n \times n$ matrix.
- The group algebra: formal \mathbb{C} -linear combination of group elements.
- Polynomial multiplication modulo a zero dimensional ideal.
- The Bose-Mesner algebra of an association scheme.

We have already encountered some of these, when studying the companion matrix and its generalizations to the multivariate case. A particularly interesting class of algebras (for a variety of reasons) are the *semisimple* algebras.

Definition 11. The radical of an associative algebra \mathcal{A} , denoted $\text{rad}(\mathcal{A})$, is the intersection of all maximal left ideals of \mathcal{A} .

Definition 12. An associative algebra \mathcal{A} is semisimple if $\text{Rad}(\mathcal{A}) = 0$.

For a semidefinite programming problem in standard (dual) form

$$\max b^T y \quad \text{s.t.} \quad A_0 - \sum_{i=1}^m A_i y_i \succeq 0,$$

we consider the algebra generated by the A_i .

Theorem 13. Let $\{A_0, \dots, A_m\}$ be given symmetric matrices, and \mathcal{A} the generated associative algebra. Then, \mathcal{A} is a semisimple algebra.

Semisimple algebras have a very nice structure, since they are essentially the direct sum of much simpler algebras.

Theorem 14 (Wedderburn). *Every finite dimensional semisimple associative algebra over \mathbb{C} can be decomposed as a direct sum*

$$\mathcal{A} = \mathcal{A}_1 \oplus \mathcal{A}_2 \oplus \dots \oplus \mathcal{A}_k.$$

Each \mathcal{A}_i is isomorphic to a simple full matrix algebra.

Example 15. *A well-known example is the (commutative) algebra of circulant matrices, i.e., those of the form*

$$A = \begin{bmatrix} a_1 & a_2 & a_3 & a_4 \\ a_4 & a_1 & a_2 & a_3 \\ a_3 & a_4 & a_1 & a_2 \\ a_2 & a_3 & a_4 & a_1 \end{bmatrix}.$$

Circulant matrices are ubiquitous in many applications, such as signal processing. It is well-known that there exists a fixed coordinate change (the Fourier matrix) under which all matrices A are diagonal (with distinct scalar blocks).

Remark 16. *In general, any associative algebra is the direct sum of its radical and a semisimple algebra. For the n -dimensional algebra generated by a single matrix $A \in \mathbb{C}^{n \times n}$, we have that $A = S + N$, where S is diagonalizable, N is nilpotent, and $SN = NS$. Thus, this statement is essentially equivalent to the existence of the Jordan decomposition.*

References

- [BDPX05] S. Boyd, P. Diaconis, P. A. Parrilo, and L. Xiao. Symmetry analysis of reversible Markov chains. *Internet Math.*, 2(1):31–71, 2005.
- [dKMP⁺] E. de Klerk, J. Maharry, D.V. Pasechnik, R.B. Richter, and G. Salazar. Improved bounds for the crossing numbers of $K_{m,n}$ and K_n . <http://arxiv.org/abs/math.CO/0404142>.
- [FS92] A. Fässler and E. Stiefel. *Group Theoretical Methods and Their Applications*. Birkhäuser, 1992.
- [GP04] K. Gatermann and P. A. Parrilo. Symmetry groups, semidefinite programs, and sums of squares. *Journal of Pure and Applied Algebra*, 192(1-3):95–128, 2004.
- [PP04] P. A. Parrilo and R. Peretz. An inequality for circle packings proved by semidefinite programming. *Discrete and Computational Geometry*, 31(3):357–367, 2004.
- [Sch05] A. Schrijver. New code upper bounds from the Terwilliger algebra and semidefinite programming. *IEEE Transactions on Information Theory*, 51(8):2859–2866, 2005.
- [Ser77] J.-P. Serre. *Linear Representations of Finite Groups*. Springer-Verlag, 1977.

Sum of Squares Programs and Polynomial Inequalities

Pablo A. Parrilo
Laboratory for Information and Decision Systems
Massachusetts Institute of Technology
Cambridge, MA 02139-4307, USA
parrilo@mit.edu

Submitted to SIAG/OPT News and Views

1 Introduction

Consider a given system of polynomial equations and inequalities, for instance:

$$\begin{aligned} f_1(x_1, x_2) &:= x_1^2 + x_2^2 - 1 = 0, \\ g_1(x_1, x_2) &:= 3x_2 - x_1^3 - 2 \geq 0, \\ g_2(x_1, x_2) &:= x_1 - 8x_2^3 \geq 0. \end{aligned} \quad (1)$$

How can one find real solutions (x_1, x_2) ? How to *prove* that they do not exist? And if the solution set is nonempty, how to optimize a polynomial function over this set?

Until a few years ago, the default answer to these and similar questions would have been that the possible nonconvexity of the feasible set and/or objective function precludes any kind of analytic global results. Even today, the methods of choice for most practitioners would probably employ mostly local techniques (Newton's and its variations), possibly complemented by a systematic search using deterministic or stochastic exploration of the solution space, interval analysis or branch and bound.

However, very recently there have been renewed hopes for the efficient solution of specific instances of this kind of problems. The main reason is the appearance of methods that combine in a very interesting fashion ideas from *real algebraic geometry* and *convex optimization* [27, 30, 21]. As we will see, these methods are based on the intimate links between sum of squares decompositions for multivariate polynomials and semidefinite programming (SDP).

In this note we outline the essential elements of this new research approach as introduced in [30, 32], and provide pointers to the literature. The centerpieces will be the following two facts about multivariate polynomials and systems of polynomials inequalities:

Sum of squares decompositions can be computed using semidefinite programming.

The search for infeasibility certificates is a convex problem. For bounded degree, it is an SDP.

In the rest of this note, we define the basic ideas needed to make the assertions above precise, and explain the relationship with earlier techniques. For this, we will introduce sum of squares polynomials and the notion of *sum of squares programs*. We then explain how to use them to provide infeasibility certificates for systems of polynomial inequalities, finally putting it all together via the surprising connections with optimization.

On a related but different note, we mention a growing body of work also aimed at the integration of ideas from algebra and optimization, but centered instead on integer programming and toric ideals; see for instance [7, 42, 3] and the volume [1] as starting points.

2 Sums of squares and SOS programs

Our notation is mostly standard. The *monomial* x^α associated to the n -tuple $\alpha = (\alpha_1, \dots, \alpha_n)$ has the form $x_1^{\alpha_1} \dots x_n^{\alpha_n}$, where $\alpha_i \in \mathbb{N}_0$. The *degree* of a monomial x^α is the nonnegative integer $\sum_{i=1}^n \alpha_i$. A *polynomial* is a finite linear combination of monomials $\sum_{\alpha \in S} c_\alpha x^\alpha$, where the coefficients c_α are real. If all the monomials have the same degree d , we will call the polynomial *homogeneous* of degree d . We denote the ring of multivariate polynomials with real coefficients in the indeterminates $\{x_1, \dots, x_n\}$ as $\mathbb{R}[x]$.

A multivariate polynomial is a *sum of squares* (SOS) if it can be written as a sum of squares of

other polynomials, i.e.,

$$p(x) = \sum_i q_i^2(x), \quad q_i(x) \in \mathbb{R}[x].$$

If $p(x)$ is SOS then clearly $p(x) \geq 0$ for all x . In general, SOS decompositions are not unique.

Example 1 *The polynomial $p(x_1, x_2) = x_1^2 - x_1x_2^2 + x_2^4 + 1$ is SOS. Among infinite others, it has the decompositions:*

$$\begin{aligned} p(x_1, x_2) &= \frac{3}{4}(x_1 - x_2^2)^2 + \frac{1}{4}(x_1 + x_2^2)^2 + 1 \\ &= \frac{1}{9}(3 - x_2^2)^2 + \frac{2}{3}x_2^2 + \\ &\quad + \frac{1}{288}(9x_1 - 16x_2^2)^2 + \frac{23}{32}x_1^2. \end{aligned}$$

The sum of squares condition is a quite natural sufficient test for polynomial nonnegativity. Its rich mathematical structure has been analyzed in detail in the past, notably by Reznick and his coauthors [6, 38], but until very recently the computational implications have not been fully explored. In the last few years there have been some very interesting new developments surrounding sums of squares, where several independent approaches have produced a wide array of results linking foundational questions in algebra with computational possibilities arising from convex optimization. Most of them employ semidefinite programming (SDP) as the essential computational tool. For completeness, we present in the next paragraph a brief summary of SDP.

Semidefinite programming SDP is a broad generalization of linear programming (LP), to the case of symmetric matrices. Denoting by S^n the space of $n \times n$ symmetric matrices, the standard SDP primal-dual formulation is:

$$\min_X C \bullet X \quad \text{s.t.} \quad \begin{cases} A_i \bullet X = b_i, & i = 1, \dots, m \\ X \succeq 0 \end{cases}$$

$$\max_y b^T y, \quad \text{s.t.} \quad \sum_{i=1}^m A_i y_i \preceq C \quad (2)$$

where $A_i, C, X \in S^n$ and $b, y \in \mathbb{R}^m$. The matrix inequalities are to be interpreted in the partial order induced by the positive semidefinite cone, i.e., $X \succeq Y$ means that $X - Y$ is a positive semidefinite matrix. Since its appearance almost a decade ago (related ideas, such as eigenvalue optimization, have been around for decades) there has been a true “revolution” in computational methods, supported by an

astounding variety of applications. By now there are several excellent introductions to SDP; among them we mention the well-known work of Vandenberghe and Boyd [44] as a wonderful survey of the basic theory and initial applications, and the handbook [45] for a comprehensive treatment of the many aspects of the subject. Other survey works, covering different complementary aspects are the early work by Alizadeh [2], Goemans [15], as well as the more recent ones due to Todd [43], De Klerk [9] and Laurent and Rendl [25].

From SDP to SOS The main object of interest in semidefinite programming is

Quadratic forms, that are positive semi-definite.

When attempting to generalize this construction to homogeneous polynomials of higher degree, an unsurmountable difficulty that appears is the fact that deciding nonnegativity for quartic or higher degree forms is an NP-hard problem. Therefore, a computational tractable replacement for this is the following:

Even degree polynomials, that are sums of squares.

Sum of squares programs can then be defined as optimization problems over affine families of polynomials, subject to SOS constraints. Like SDPs, there are several possible equivalent descriptions. We choose below a free variables formulation, to highlight the analogy with the standard SDP dual form discussed above.

Definition 1 *A sum of squares program has the form*

$$\begin{aligned} \max_y \quad & b_1 y_1 + \dots + b_m y_m \\ \text{s.t.} \quad & P_i(x, y) \text{ are SOS}, \quad i = 1, \dots, p \end{aligned}$$

where $P_i(x, y) := C_i(x) + A_{i1}(x)y_1 + \dots + A_{im}(x)y_m$, and the C_i, A_{ij} are given polynomials in the variables x_i .

SOS programs are very useful, since they directly operate with polynomials as their basic objects, thus providing a quite natural modelling formulation for many problems. Among others, examples for this are the search for Lyapunov functions for nonlinear systems [30, 28], probability inequalities [4], as well as the relaxations in [30, 21] discussed below.

Interestingly enough, despite their apparently greater generality, sum of squares programs are in fact *equivalent* to SDPs. On the one hand, by choosing the polynomials $C_i(x), A_{ij}(x)$ to be quadratic

forms, we recover standard SDP. On the other hand, as we will see in the next section, it is possible to exactly embed every SOS program into a larger SDP. Nevertheless, the rich algebraic structure of SOS programs will allow us a much deeper understanding of their special properties, as well as enable customized, more efficient algorithms for their solution [26].

Furthermore, as illustrated in later sections, there are numerous questions related to some foundational issues in nonconvex optimization that have simple and natural formulations as SOS programs.

SOS programs as SDPs Sum of squares programs can be written as SDPs. The reason is the following theorem:

Theorem 1 *A polynomial $p(x)$ is SOS if and only if $p(x) = z^T Q z$, where z is a vector of monomials in the x_i variables, $Q \in \mathcal{S}^N$ and $Q \succeq 0$.*

In other words, every SOS polynomial can be written as a quadratic form in a set of monomials of cardinality N , with the corresponding matrix being positive semidefinite. The vector of monomials z (and therefore N) in general depends on the degree and sparsity pattern of $p(x)$. If $p(x)$ has n variables and total degree $2d$, then z can always be chosen as a subset of the set of monomials of degree less than or equal to d , of cardinality $N = \binom{n+d}{d}$.

Example 2 *Consider again the polynomial from Example 1. It has the representation*

$$p(x_1, x_2) = \frac{1}{6} \begin{bmatrix} 1 \\ x_2 \\ x_2^2 \\ x_1 \end{bmatrix}^T \begin{bmatrix} 6 & 0 & -2 & 0 \\ 0 & 4 & 0 & 0 \\ -2 & 0 & 6 & -3 \\ 0 & 0 & -3 & 6 \end{bmatrix} \begin{bmatrix} 1 \\ x_2 \\ x_2^2 \\ x_1 \end{bmatrix},$$

and the matrix in the expression above is positive semidefinite.

In the representation $f(x) = z^T Q z$, for the right- and left-hand sides to be identical, all the coefficients of the corresponding polynomials should be equal. Since Q is simultaneously constrained by linear equations and a positive semidefiniteness condition, the problem can be easily seen to be directly equivalent to an SDP feasibility problem in the standard primal form (2).

Given a SOS program, we can use the theorem above to construct an equivalent SDP. The conversion step is fully algorithmic, and has been implemented, for instance, in the SOSTOOLS [36] software package. Therefore, we can in principle directly apply all the available numerical methods for SDP to solve SOS programs.

SOS and convexity The connection between sum of squares decompositions and convexity can be traced back to the work of N. Z. Shor [39]. In this 1987 paper, he essentially outlined the links between Hilbert’s 17th problem and a class of convex bounds for unconstrained polynomial optimization problems. Unfortunately, the approach went mostly unnoticed for several years, probably due to the lack of the convenient framework of SDP.

3 Algebra and optimization

A central theme throughout convex optimization is the idea of *infeasibility certificates* (for instance, in LP via Farkas’ lemma), or equivalently, *theorems of the alternative*. As we will see, the key link relating algebra and optimization in this approach is the fact that infeasibility can *always* be certified by a particular algebraic identity, whose solution is found via convex optimization.

We explain some of the concrete results in Theorem 5, after a brief introduction to two algebraic concepts, and a comparison with three well-known infeasibility certificates.

Ideals and cones For later reference, we define here two important algebraic objects: the *ideal* and the *cone* associated with a set of polynomials:

Definition 2 *Given a set of multivariate polynomials $\{f_1, \dots, f_m\}$, let*

$$\mathbf{ideal}(f_1, \dots, f_m) := \{f \mid f = \sum_{i=1}^m t_i f_i, \quad t_i \in \mathbb{R}[x]\}.$$

Definition 3 *Given a set of multivariate polynomials $\{g_1, \dots, g_m\}$, let*

$$\begin{aligned} \mathbf{cone}(g_1, \dots, g_m) := & \{g \mid g = s_0 + \sum_{\{i\}} s_i g_i + \\ & + \sum_{\{i,j\}} s_{ij} g_i g_j + \sum_{\{i,j,k\}} s_{ijk} g_i g_j g_k + \dots\}, \end{aligned}$$

where each term in the sum is a squarefree product of the polynomials g_i , with a coefficient $s_\alpha \in \mathbb{R}[x]$ that is a sum of squares. The sum is finite, with a total of $2^m - 1$ terms, corresponding to the nonempty subsets of $\{g_1, \dots, g_m\}$.

These algebraic objects will be used for deriving *valid inequalities*, which are logical consequences of the given constraints. Notice that by construction, every polynomial in $\mathbf{ideal}(f_i)$ vanishes in the solution set

of $f_i(x) = 0$. Similarly, every element of $\mathbf{cone}(g_i)$ is clearly nonnegative on the feasible set of $g_i(x) \geq 0$.

The notions of *ideal* and *cone* as used above are standard in real algebraic geometry; see for instance [5]. In particular, the cones are also referred to as *preorders*. Notice that as geometric objects, ideals are affine sets, and cones are closed under convex combinations and nonnegative scalings (i.e., they are actually cones in the convex geometry sense). These convexity properties, coupled with the relationships between SDP and SOS, will be key for our developments in the next section.

Infeasibility certificates If a system of equations does not have solutions, how do we *prove* this fact? A very useful concept is that of *certificates*, which are formal algebraic identities that provide irrefutable evidence of the inexistence of solutions.

We briefly illustrate some well-known examples below. The first two deal with linear systems and polynomial equations over the complex numbers, respectively.

Theorem 2 (Range/kernel)

$$\begin{aligned} Ax = b \quad & \text{is infeasible} \\ \Updownarrow \\ \exists \mu \text{ s.t. } A^T \mu = 0, b^T \mu = -1. \end{aligned}$$

Theorem 3 (Hilbert’s Nullstellensatz) Let $f_i(z), \dots, f_m(z)$ be polynomials in complex variables z_1, \dots, z_n . Then,

$$\begin{aligned} f_i(z) = 0 \quad (i = 1, \dots, m) \quad & \text{is infeasible in } \mathbb{C}^n \\ \Updownarrow \\ -1 \in \mathbf{ideal}(f_1, \dots, f_m). \end{aligned}$$

Each of these theorems has an “easy” direction. For instance, for the first case, given the multipliers μ the infeasibility is obvious, since

$$Ax = b \quad \Rightarrow \quad \mu^T Ax = \mu^T b \quad \Rightarrow \quad 0 = -1,$$

which is clearly a contradiction.

The two theorems above deal only with the case of *equations*. The inclusion of inequalities in the problem formulation poses additional algebraic challenges, because we need to work on an *ordered* field. In other words, we need to take into account special properties of the *reals*, and not just the complex numbers.

For the case of *linear* inequalities, LP duality provides the following characterization:

Theorem 4 (Farkas lemma)

$$\begin{aligned} \begin{cases} Ax + b &= 0 \\ Cx + d &\geq 0 \end{cases} \quad & \text{is infeasible} \\ \Updownarrow \\ \exists \lambda \geq 0, \mu \text{ s.t. } \begin{cases} A^T \mu + C^T \lambda &= 0 \\ b^T \mu + d^T \lambda &= -1. \end{cases} \end{aligned}$$

Although not widely known in the optimization community until recently, it turns out that similar certificates do exist for *arbitrary* systems of polynomial equations and inequalities over the reals. The result essentially appears in this form in [5], and is due to Stengle [40].

Theorem 5 (Positivstellensatz)

$$\begin{aligned} \begin{cases} f_i(x) &= 0, \quad (i = 1, \dots, m) \\ g_i(x) &\geq 0, \quad (i = 1, \dots, p) \end{cases} \quad & \text{is infeasible in } \mathbb{R}^n \\ \Updownarrow \\ \exists F(x), G(x) \in \mathbb{R}[x] \text{ s.t. } \begin{cases} F(x) + G(x) &= -1 \\ F(x) &\in \mathbf{ideal}(f_1, \dots, f_m) \\ G(x) &\in \mathbf{cone}(g_1, \dots, g_p). \end{cases} \end{aligned}$$

The theorem states that for every infeasible system of polynomial equations and inequalities, there exists a simple algebraic identity that directly certifies the inexistence of real solutions. By construction, the evaluation of the polynomial $F(x) + G(x)$ at any feasible point should produce a nonnegative number. However, since this expression is identically equal to the polynomial -1 , we arrive at a contradiction. Remarkably, the Positivstellensatz holds under *no assumptions* whatsoever on the polynomials.

The use of the German word “Positivstellensatz” is standard in the field, and parallels the classical “Nullstellensatz” (roughly, “theorem of the zeros”) obtained by Hilbert in 1901 and mentioned above.

In the worst case, the degree of the infeasibility certificates $F(x), G(x)$ could be high (of course, this is to be expected, due to the NP-hardness of the original question). In fact, there are a few explicit counterexamples where large degree refutations are necessary [16]. Nevertheless, for many problems of practical interest, it is often the case that it is possible to prove infeasibility using relatively low-degree certificates. There is significant numerical evidence that this is the case, as indicated by the large number of practical applications where SDP relaxations based on these techniques have provided solutions of very high quality.

Degree \ Field	Complex	Real
Linear	<i>Range/Kernel</i> Linear Algebra	<i>Farkas Lemma</i> Linear Programming
Polynomial	<i>Nullstellensatz</i> Bounded degree: Linear Algebra Groebner bases	<i>Positivstellensatz</i> Bounded degree: SDP

Table 1: Infeasibility certificates and associated computational techniques.

Of course, we are concerned with the effective computation of these certificates. For the cases of Theorems 2–4, the corresponding refutations can be obtained using either linear algebra, linear programming, or Groebner bases techniques (see [8] for a superb introduction to Groebner bases).

For the Positivstellensatz, we notice that the cones and ideals as defined above are always *convex sets* in the space of polynomials. A key consequence is that the conditions in Theorem 5 for a certificate to exist are therefore *convex*, regardless of any convexity property of the original problem. Even more, the same property holds if we consider only bounded-degree sections, i.e., the intersection with the set of polynomials of degree less than or equal to a given number D . In this case, the conditions in the P-satz have *exactly* the form of a SOS program! Of course, as discussed earlier, this implies that we can find bounded-degree certificates, by solving semidefinite programs. In Table 1 we present a summary of the infeasibility certificates discussed, and the associated computational techniques.

Example 3 Consider again the system (1). We will show that it has no solutions $(x_1, x_2) \in \mathbb{R}^2$. By the P-satz, the system is infeasible if and only if there exist polynomials $t_1, s_0, s_1, s_2, s_{12} \in \mathbb{R}[x_1, x_2]$ that satisfy

$$\underbrace{f_1 \cdot t_1}_{\text{ideal}(f_1)} + \underbrace{s_0 + s_1 \cdot g_1 + s_2 \cdot g_2 + s_{12} \cdot g_1 \cdot g_2}_{\text{cone}(g_1, g_2)} \equiv -1, \quad (3)$$

where s_0, s_1, s_2 and s_{12} are SOS.

A SOS relaxation is obtained by looking for solutions where all the terms in the left-hand side have degree less than or equal to D . For each fixed integer $D > 0$ this can be tested by semidefinite programming.

For instance, for $D = 4$ we find the certificate


$$\begin{aligned} t_1 &= -3x_1^2 + x_1 - 3x_2^2 + 6x_2 - 2, \\ s_1 &= 3, \quad s_2 = 1, \quad s_{12} = 0, \\ s_0 &= 3x_1^4 + 2x_1^3 + 6x_1^2x_2^2 - 6x_1^2x_2 - x_1^2 - x_1x_2^2 + \\ &\quad + 3x_2^4 + 2x_2^3 - x_2^2 - 3x_2 + 3 \\ &= \frac{1}{2} z^T \begin{bmatrix} 6 & -3 & -3 & 0 & 0 & -3 \\ -3 & 4 & 2 & 0 & 1 & 1 \\ -3 & 2 & 6 & -2 & 0 & -3 \\ 0 & 0 & -2 & 4 & -7 & 2 \\ 0 & 1 & 0 & -7 & 18 & 0 \\ -3 & 1 & -3 & 2 & 0 & 6 \end{bmatrix} z, \end{aligned}$$

where

$$z = [1 \quad x_2 \quad x_2^2 \quad x_1 \quad x_1x_2 \quad x_1^2]^T.$$

The resulting identity (3) thus certifies the inconsistency of the system $\{f_1 = 0, g_1 \geq 0, g_2 \geq 0\}$.

As outlined in the preceding paragraphs, there is a direct connection going from general polynomial optimization problems to SDP, via P-satz infeasibility certificates. Pictorially, we have the following:



Even though we have discussed only feasibility problems, there are obvious straightforward connections with optimization. By considering the emptiness of the sublevel sets of the objective function, sequences of converging bounds indexed by certificate degree can be directly constructed.

4 Further developments and applications

We have covered only the core elements of the SOS/SDP approach. Much more is known, and even

more still remains to be discovered, both in the theoretical and computational ends. Some specific issues are discussed below.

Exploiting structure and numerical computation To what extent can the inherent structure in SOS programs be exploited for efficient computations? Given the algebraic origins of the formulation, it is perhaps not surprising to find that several intrinsic properties of the input polynomials can be profitably used [29]. In this direction, symmetry reduction techniques have been employed by Gatermann and Parrilo in [14] to provide novel representations for symmetric polynomials. Kojima, Kim and Waki [20] have recently presented some results for sparse polynomials. Parrilo [31] and Laurent [23] have analyzed the further simplifications that occur when the inequality constraints define a zero-dimensional ideal.

Other relaxations Lasserre [21, 22] has independently introduced a scheme for polynomial optimization dual to the one described here, but relying on Putinar’s representation theorem for positive polynomials rather than the P-satz. There are very interesting relationships between SOS-based methods and earlier relaxation and approximation schemes, such as Lovász-Schrijver and Sherali-Adams. Laurent [24] analyzes this in the specific case of 0-1 programming.

Implementations The software SOSTOOLS [36] is a free, third-party MATLAB¹ toolbox for formulating and solving general sum of squares programs. The related software Gloptipoly [17] is oriented toward global optimization problems. In their current version, both use the SDP solver SeDuMi [41] for numerical computations.

Approximation properties There are several important open questions regarding the provable quality of the approximations. In this direction, De Klerk and Pasechnik [11] have established some approximations guarantees of a SOS-based scheme for the approximation of the stability number of a graph. Recently, De Klerk, Laurent and Parrilo [10] have shown that a related procedure based on a result by Pólya provides a polynomial-time approximation scheme (PTAS) for polynomial optimization over simplices.

Applications There are many exciting applications of the ideas described here. The descriptions that follow are necessarily brief; our main objective

here is to provide the reader with some good starting points to this growing literature.

In systems and control theory, the techniques have provided some of the best available analysis and design methods, in areas such as nonlinear stability and robustness analysis [30, 28, 35], state feedback control [19], fixed-order controllers [18], nonlinear synthesis [37], and model validation [34]. Also, there have been interesting recent applications in geometric theorem proving [33] and quantum information theory [12, 13].

Acknowledgments: The author would like to thank Etienne de Klerk and Luis Vicente for their helpful comments and suggestions.

References

- [1] K. Aardal and R. Thomas (eds.), *Algebraic and geometric methods in discrete optimization*, Springer-Verlag, Heidelberg, 2003, Math. Program. **96** (2003), no. 2, Ser. B. MR 1 993 044
- [2] F. Alizadeh, *Interior point methods in semidefinite programming with applications to combinatorial optimization*, SIAM J. Optim. **5** (1995), no. 1, 13–51.
- [3] D. Bertsimas, G. Perakis, and S. Tayur, *A new algebraic geometry algorithm for integer programming*, Management Science **46** (2000), no. 7, 999–1008.
- [4] D. Bertsimas and I. Popescu, *Optimal inequalities in probability: A convex optimization approach*, INSEAD working paper, available at <http://www.insead.edu/facultyresearch/tm/popescu/>, 1999-2001.
- [5] J. Bochnak, M. Coste, and M-F. Roy, *Real algebraic geometry*, Springer, 1998.
- [6] M. D. Choi, T. Y. Lam, and B. Reznick, *Sums of squares of real polynomials*, Proceedings of Symposia in Pure Mathematics **58** (1995), no. 2, 103–126.
- [7] P. Conti and C. Traverso, *Buchberger algorithm and integer programming*, Applied algebra, algebraic algorithms and error-correcting codes (New Orleans, LA, 1991), Lecture Notes in Comput. Sci., vol. 539, Springer, Berlin, 1991, pp. 130–139. MR 1 229 314
- [8] D. A. Cox, J. B. Little, and D. O’Shea, *Ideals, varieties, and algorithms: an introduction to computational algebraic geometry and commutative algebra*, Springer, 1997.
- [9] E. de Klerk, *Aspects of semidefinite programming: Interior point algorithms and selected applications*, Applied Optimization, vol. 65, Kluwer Academic Publishers, 2002.
- [10] E. de Klerk, M. Laurent, and P. A. Parrilo, *A PTAS for the minimization of polynomials of fixed degree over the simplex*, Submitted, available at <http://www.mit.edu/~parrilo>, 2004.

¹A registered trademark of The MathWorks, Inc.

- [11] E. de Klerk and D.V. Pasechnik, *Approximating the stability number of a graph via copositive programming*, SIAM Journal on Optimization **12** (2002), no. 4, 875–892.
- [12] A. C. Doherty, P. A. Parrilo, and F. M. Spedalieri, *Distinguishing separable and entangled states*, Physical Review Letters **88** (2002), no. 18.
- [13] ———, *Complete family of separability criteria*, Physical Review A **69** (2004), 022308.
- [14] K. Gatermann and P. A. Parrilo, *Symmetry groups, semidefinite programs, and sums of squares*, Journal of Pure and Applied Algebra **192** (2004), no. 1-3, 95–128.
- [15] M. X. Goemans, *Semidefinite programming in combinatorial optimization*, Math. Programming **79** (1997), no. 1-3, 143–161. MR 98g:90028
- [16] D. Grigoriev and N. Vorobjov, *Complexity of Null and Positivstellensatz proofs*, Annals of Pure and Applied Logic **113** (2002), no. 1-3, 153–160.
- [17] D. Henrion and J.-B. Lasserre, *GloptiPoly - global optimization over polynomials with Matlab and SeDuMi*, Available from <http://www.laas.fr/~henrion/software/gloptipoly/>.
- [18] D. Henrion, M. Sebek, and V. Kucera, *Positive polynomials and robust stabilization with fixed-order controllers*, IEEE Transactions on Automatic Control **48** (2003), no. 7, 1178–1186.
- [19] Z. Jarvis-Wloszek, R. Feeley, W. Tan, K. Sun, and A. Packard, *Some controls applications of sum of squares*, Proceedings of the 42th IEEE Conference on Decision and Control, 2003, pp. 4676–4681.
- [20] M. Kojima, S. Kim, and H. Waki, *Sparsity in sum of squares of polynomials*, Research report B-391, Dept. of Mathematical and Computing Sciences, Tokyo Institute of Technology, 2003.
- [21] J. B. Lasserre, *Global optimization with polynomials and the problem of moments*, SIAM J. Optim. **11** (2001), no. 3, 796–817. MR 1 814 045
- [22] ———, *Polynomials nonnegative on a grid and discrete optimization*, Trans. Amer. Math. Soc. **354** (2002), no. 2, 631–649. MR 1 862 561
- [23] M. Laurent, *Semidefinite representations for finite varieties*, Preprint.
- [24] ———, *A comparison of the Sherali-Adams, Lovász-Schrijver, and Lasserre relaxations for 0-1 programming*, Math. Oper. Res. **28** (2003), no. 3, 470–496. MR 1 997 246
- [25] M. Laurent and F. Rendl, *Semidefinite programming and integer programming*, Tech. Report PNA-R0210, CWI, Amsterdam, April 2002.
- [26] J. Löfberg and P. A. Parrilo, *From coefficients to samples: a new approach in SOS optimization*, Proceedings of the 43th IEEE Conference on Decision and Control, 2004.
- [27] Y. Nesterov, *Squared functional systems and optimization problems*, High Performance Optimization (J.B.G. Frenk, C. Roos, T. Terlaky, and S. Zhang, eds.), Kluwer Academic Publishers, 2000, pp. 405–440.
- [28] A. Papachristodoulou and S. Prajna, *On the construction of Lyapunov functions using the sum of squares decomposition*, Proceedings of the 41th IEEE Conference on Decision and Control, 2002.
- [29] P. A. Parrilo, *Exploiting algebraic structure in sum of squares programs*, In *Positive Polynomials in Control*, D. Henrion and A. Garulli, eds., LNCIS, Springer, forthcoming.
- [30] ———, *Structured semidefinite programs and semialgebraic geometry methods in robustness and optimization*, Ph.D. thesis, California Institute of Technology, May 2000, Available at <http://resolver.caltech.edu/CaltechETD:etd-05062004-055516>.
- [31] ———, *An explicit construction of distinguished representations of polynomials non-negative over finite sets*, Tech. Report IfA Technical Report AUT02-02. Available from <http://control.ee.ethz.ch/~parrilo>, ETH Zürich, 2002.
- [32] ———, *Semidefinite programming relaxations for semialgebraic problems*, Math. Prog. **96** (2003), no. 2, Ser. B, 293–320.
- [33] P. A. Parrilo and R. Peretz, *An inequality for circle packings proved by semidefinite programming*, Discrete and Computational Geometry **31** (2004), no. 3, 357–367.
- [34] S. Prajna, *Barrier certificates for nonlinear model validation*, Proceedings of the 42th IEEE Conference on Decision and Control, 2003, pp. 2884–2889.
- [35] S. Prajna and A. Papachristodoulou, *Analysis of switched and hybrid systems – beyond piecewise quadratic methods*, Proceedings of the American Control Conference, 2003.
- [36] S. Prajna, A. Papachristodoulou, and P. A. Parrilo, *SOSTOOLS: Sum of squares optimization toolbox for MATLAB*, 2002-04, Available from <http://www.cds.caltech.edu/sostools> and <http://www.mit.edu/~parrilo/sostools>.
- [37] S. Prajna, P. A. Parrilo, and A. Rantzer, *Nonlinear control synthesis by convex optimization*, IEEE Transactions on Automatic Control **49** (2004), no. 2, 310–314.
- [38] B. Reznick, *Some concrete aspects of Hilbert’s 17th problem*, Contemporary Mathematics, vol. 253, American Mathematical Society, 2000, pp. 251–272.
- [39] N. Z. Shor, *Class of global minimum bounds of polynomial functions*, Cybernetics **23** (1987), no. 6, 731–734, (Russian orig.: Kibernetika, No. 6, (1987), 9–11).

- [40] G. Stengle, *A Nullstellensatz and a Positivstellensatz in semialgebraic geometry*, Math. Ann. **207** (1974), 87–97.
- [41] J. Sturm, *SeDuMi version 1.05*, October 2001, Available from <http://fewcal.uvt.nl/sturm/software/sedumi.html>.
- [42] R. Thomas and R. Weismantel, *Truncated Gröbner bases for integer programming*, Appl. Algebra Engrg. Comm. Comput. **8** (1997), no. 4, 241–256. MR 98i:90055
- [43] M. Todd, *Semidefinite optimization*, Acta Numerica **10** (2001), 515–560.
- [44] L. Vandenberghe and S. Boyd, *Semidefinite programming*, SIAM Review **38** (1996), no. 1, 49–95.
- [45] H. Wolkowicz, R. Saigal, and L. Vandenberghe (eds.), *Handbook of semidefinite programming*, Kluwer, 2000.