

Alternating Weak Automata from Universal Trees

Laure Daviaud

Marcin Jurdziński

Karoliina Lehtinen

Department of Computer Science
City, University of London
Laure.Daviaud@city.ac.uk

Department of Computer Science
University of Warwick
Marcin.Jurdzinski@warwick.ac.uk

Department of Computer Science
University of Liverpool
k.lehtinen@liverpool.ac.uk

Abstract

An improved translation from alternating parity automata on infinite words to alternating weak automata is given. The blow-up of the number of states is related to the size of the smallest universal ordered trees and hence it is quasi-polynomial, and only polynomial if the asymptotic number of priorities is logarithmic in the number of states. This is an exponential improvement on the translation of Kupferman and Vardi (2001) and a quasi-polynomial improvement on the translation of Boker and Lehtinen (2018). Any slightly better such translation would (if—like all presently known such translations—it is efficiently constructive) lead to algorithms for solving parity games that are asymptotically faster in the worst case than the current state of the art (Calude, Jain, Khoussainov, Li, and Stephan, 2017; Jurdziński and Lazić, 2017; and Fearnley, Jain, Schewe, Stephan, and Wojtczak, 2017), and hence it would yield a significant breakthrough.

I. INTRODUCTION

We give a new translation of alternating parity automata on infinite words into alternating weak automata. Alternating weak automata have the same expressivity as parity automata [14], yet enjoy simpler algorithmic properties. They are, however, less concise.

Understanding the exact trade-off between the complexity of the acceptance condition—parity or weak—and the size of the automaton is interesting for several algorithmic problems. For instance, the emptiness check for alternating parity automata reduces to checking the emptiness of an equivalent weak automaton, which, in the case of automata over the singleton alphabet, is linear. The complexity of this procedure then depends on the blow-up involved in the translation into weak automata. Deciding the winner in a parity game on the other hand is a special case of checking the emptiness of alternating parity word automata; a concise translation of alternating parity automata into weak automata immediately implies an algorithm for solving parity games of which the complexity matches the increase in the size of the automaton.

The necessary size blow-up when going from alternating parity automata to alternating weak automata is wide open: the best known lower bound is $\Omega(n \log n)$ [14], closely related to the $2^{\Omega(n \log n)}$ lower bound on Büchi complementation [17], while the upper bound is quasi-polynomial. Boker and Lehtinen’s quasi-polynomial translation [1] combines Lehtinen’s register technique [15], which reduces the number of priorities to logarithmic in the state space, with Kupferman and Vardi’s exponential (in the number of priorities) procedure [13], which eliminates the remaining priorities one by one. This results in a blow-up of the form $n^{\Theta(\lg n \cdot \lg(d/\lg n))}$, where n and d are the number of states and the number of priorities of the original parity automaton.

In contrast, we first establish a hierarchical decomposition of runs of parity automata as a generalization of the decomposition of runs of co-Büchi automata due to Kupferman and Vardi [14], and then we use the recently introduced universal trees [10], [5] to reduce the number of priorities to 2 (Büchi). The blow-up involved in the parity to Büchi translation is proportional to the size of the smallest universal trees, and the further blow-up involved in the Büchi to weak translation is only quadratic [14]. The composite blow-up is polynomial if $d = O(\log n)$, and it is of the form $n^{O(\lg(d/\lg n))}$ if $d = \omega(\log n)$ [10]. This way we obtain an exponential improvement on the translation of Kupferman and Vardi [14] (which has $n^{d+O(1)}$ blow-up) and a quasi-polynomial factor $n^{\Omega(\lg n)}$ improvement on the translation of Boker and Lehtinen [1].

The blow-up we obtain nearly matches the current state-of-the-art quasi-polynomial upper bounds on the complexity of solving parity games [2], [10], [8].

The exponential translation from alternating parity automata to alternating weak automata due to Kupferman and Vardi [13] is done by a rather involved induction on the number of priorities. For an automaton with highest priority d , it goes through a sequence of d intermediate automata of a generalized type, which they call *weak-parity alternating automata*. In contrast, our construction is significantly more streamlined and transparent; in particular, it avoids introducing a new class of hybrid parity-weak automata.

Our work is yet another application of the recently introduced notion of universal trees [10], [5]. Such applications typically focus on algorithms for solving games [10], [6], [5], [3]; our work is the first whose primary focus is on automata.

In addition to universal trees, we use a notion of lazy progress measure, which adapts classical parity progress measures, used to solve parity games, to infinite arenas. Unlike regular parity progress measures, which can be recognised by safety automata but require a definite bound to be known on the successive occurrences of odd priorities, lazy progress measures are recognised by Büchi automata and can deal with finite but unbounded occurrences of successive odd priorities. This Büchi automaton is similar to (but more concise than) the automaton used to characterise parity tree automata that recognise co-Büchi recognisable tree languages [16], itself a generalisation of automata used to decide the weak definability of tree languages given as Büchi automata [18], [4].

A similar concept to our lazy progress measures was already introduced by Klarlund for complementation of Büchi and Streett automata on words [11]. Klarlund indeed proves a result that is equivalent to one of our key lemmas on parity progress measures. Our proof, however, is more constructive, and it explicitly provides a hierarchical decomposition, which clearly describes the structure of accepting run dags of parity word automata. Moreover—unlike Klarlund’s proof, which relies on the result about Rabin measures [12]—our proof is self-contained. We suspect that the opaqueness of Klarlund’s paper [11] may have been responsible for attracting less attention and shallower absorption by the research community than it deserves. In particular, some of the techniques and results that he presents there have been rediscovered and refined by various authors, often much later [14], [13], [9], [19], [10], [5], including this work. We hope that our paper will help a wider and more thorough reception and appreciation of Klarlund’s work.

II. ALTERNATING AUTOMATA

For a finite set X , we write $\mathcal{B}^+(X)$ for the set of positive Boolean formulas over X . We say that a set $Y \subseteq X$ *satisfies* a formula $\varphi \in \mathcal{B}^+(X)$ if φ evaluates to true when all variables in Y are set to true and all variables in $X \setminus Y$ are set to false. For example, the sets $\{x, y\}$ and $\{x, z\}$ satisfy the positive Boolean formula $x \wedge (y \vee z)$, but the set $\{y, z\}$ does not. An *alternating automaton* has a finite set Q of *states*, an *initial state* $q_0 \in Q$, a finite *alphabet* Σ , and a transition function $\delta : Q \times \Sigma \rightarrow \mathcal{B}^+(Q)$. Alternating automata allow to combine both *non-deterministic* and *universal* transitions; disjunctions in transition formulas model the non-deterministic choices and conjunctions model the universal choices.

We consider alternating automata as acceptors of infinite words. Whether infinite sequences of states in runs of such automata are *accepting* or not is determined by an *acceptance condition*. Here, we consider *parity*, *Büchi*, *co-Büchi*, *weak*, and *safety* acceptance conditions. In a *parity* condition, given by a *state priority function* $\pi : Q \rightarrow \{0, 1, 2, \dots, d\}$ for some positive even integer d , an infinite sequence of states is accepting if the largest state priority that occurs infinitely many times is even. *Büchi* conditions are a special case of parity conditions in which all states have priorities 1 or 2, and *co-Büchi* conditions are parity conditions in which all states have priorities 0 or 1.

Let the *transition graph* of an alternating automaton have an edge $(q, r) \in Q \times Q$ if r occurs in $\delta(q, a)$ for some letter $a \in \Sigma$. We say that a parity automaton has a *weak* acceptance condition if it is *stratified*: in every cycle in the transition graph, all states have the same priority. Weak conditions are a special case of both Büchi and co-Büchi conditions in the following sense: if the transition graph of a parity automaton

is stratified, then every infinite path in the transition graph satisfies each of the following three conditions if and only if it satisfies the other two:

- the parity condition $\pi : Q \rightarrow \{0, 1, 2, \dots, d\}$;
- the co-Büchi condition $\pi' : Q \rightarrow \{0, 1\}$;
- the Büchi condition $\pi'' : Q \rightarrow \{1, 2\}$;

where $\pi'(q) = \pi(q) \bmod 2$ and $\pi''(q) = 2 - \pi'(q)$ for all $q \in Q$.

We say that a state is *absorbing* if its only successor in the transition graph is itself. A parity automaton has a *safety* acceptance condition if all of its states have priority 0, except for the additional absorbing reject state that has priority 1. An automaton with a safety acceptance condition is stratified, and hence safety conditions are a special case of weak conditions.

Whether an infinite word $w = w_0w_1w_2 \dots \in \Sigma^\omega$ is *accepted* or *rejected* by an alternating automaton \mathcal{A} is determined by the winner of the following *acceptance game* $\mathcal{G}(\mathcal{A}, w)$. The set of positions in the game is the set $Q \times \mathbb{N}$ and the two players, Alice and Elvis, play in the following way. The *initial position* is $(q_0, 0)$; for every *current position* (q_i, i) , first Elvis chooses a subset P of Q that satisfies $\delta(q_i, w_i)$, then Alice picks a state $q_{i+1} \in P$, and $(q_{i+1}, i+1)$ becomes the next current position. Note that Elvis can be thought of making the non-deterministic choices and Alice can be thought of making the universal choices in the transition function of the alternating automaton. This interaction of Alice and Elvis yields an infinite sequence of states q_0, q_1, q_2, \dots , and whether Elvis is declared the winner or not is determined by whether the sequence is *accepting* according to the acceptance condition of the automaton. Acceptance games for parity, Büchi, co-Büchi, and weak conditions are parity games, which are *determined* [7], [21]: in every acceptance game, either Alice or Elvis has a winning strategy. We say that an infinite word $w \in \Sigma^\omega$ is *accepted* by an alternating automaton \mathcal{A} if Elvis has a winning strategy in the acceptance game $\mathcal{G}(\mathcal{A}, w)$, and otherwise it is *rejected*.

A *run dag* of an alternating automaton \mathcal{A} on an infinite word w is a directed acyclic graph $G = (V, E, \rho : V \rightarrow Q)$, where $V \subseteq Q \times \mathbb{N}$ is the set of vertices; successors (according to the directed edge relation E) of every vertex (q, i) are of the form $(q', i+1)$; the following conditions hold:

- $(q_0, 0) \in V$,
- for every $(q, i) \in V$, the Boolean formula $\delta(q, w_i)$ is satisfied by the set of states p , such that $(p, i+1)$ is a successor of (q, i) ;

and ρ projects vertices onto the first component. Note that every vertex in a run dag has a successor, and hence every maximal path is infinite. We say that a run dag of an automaton \mathcal{A} is *accepting* if the sequence of states on every infinite path in the run dag is accepting according to the accepting condition of \mathcal{A} . The *positional* determinacy theorem for parity games [7], [21] implies that an infinite word w is accepted by an alternating automaton \mathcal{A} with a parity (or Büchi, co-Büchi, weak, or safety) condition if and only if there is an accepting run dag of \mathcal{A} on w . In other words, run dags are compact representations of (positional) winning strategies for Elvis in the acceptance games.

Run dags considered here are a special case of *layered dags*, whose vertices can be partitioned into sets L_0, L_1, L_2, \dots , such that every edge goes from some layer L_i to the next layer L_{i+1} . We define the *width* of a layered dag with an infinite number of layers L_0, L_1, L_2, \dots to be $\liminf_{i \rightarrow \infty} |L_i|$. Note that the width of a run dag of an alternating automaton is trivially upper-bounded by the number of states of the automaton.

III. FROM CO-BÜCHI AND BÜCHI TO WEAK

In this section we summarize the results of Kupferman and Vardi [14] who have given translations from alternating co-Büchi and Büchi automata to alternating weak automata with only a quadratic blow-up in the state space. We recall the decomposition of co-Büchi accepting run dags of Kupferman and Vardi in detail because it motivates and prepares the reader for our generalization of their result to accepting parity run dags. Our main technical result is a translation from alternating parity automata to alternating

Büchi automata with only a quasi-polynomial blow-up in the state space, but the ultimate goal is a quasi-polynomial translation from parity to weak automata. Therefore, we also recall how Kupferman and Vardi use their quadratic co-Büchi to weak translation in order to obtain a quadratic Büchi to weak translation.

A. Co-Büchi progress measures

The main technical concept that underlies Kupferman and Vardi's [14] translation from alternating co-Büchi automata to alternating weak automata is that of a *ranking function* for accepting run dags of alternating co-Büchi automata. As Kupferman and Vardi themselves point out, ranking functions can be seen as equivalent to Klarlund's *progress measures* [11]. We will adopt Klarlund's terminology because the theory of progress measures for certifying parity conditions is very well developed [7], [11], [12], [9], [10], [5] and our main goal in this paper is to use a version of parity progress measures to give a simplified, streamlined, and improved translation from alternating parity to alternating weak automata.

Let $G = (V, E, \pi : V \rightarrow \{0, 1\})$ be a layered dag with vertex priorities 0 or 1, and in which every vertex has a successor. Note that all run dags of an alternating co-Büchi automaton are such layered dags and if the automaton has n states then the width of the run dag is at most n . (Observe, however, that while formally the third component $\rho : V \rightarrow Q$ in a run dag maps vertices to states, here we instead consider the labeling $\pi : V \rightarrow \{0, 1\}$ that labels vertices by the priorities of the states $\pi(v) = \pi(\rho(v))$.)

A *co-Büchi progress measure* [7], [12], [9] is a mapping $\mu : V \rightarrow M$, where (M, \leq) is a well-ordered set, such that for every edge $(v, u) \in E$, we have

- 1) if $\pi(v) = 0$ then $\mu(v) \geq \mu(u)$,
- 2) if $\pi(v) = 1$ then $\mu(v) > \mu(u)$.

It is elementary to argue that existence of a co-Büchi progress measure on a graph is sufficient for every infinite path in the graph satisfying the co-Büchi condition. Importantly, it is also necessary, which can be, for example, deduced from the proof of positional determinacy for parity games due to Emerson and Jutla [7]. In other words, co-Büchi progress measures are witnesses for the property that all infinite paths in a graph satisfy the co-Büchi condition. The appeal of such witnesses stems from the property that while certifying a global and infinitary condition, it suffices to verify them locally by checking a simple inequality between the labels of the source and the target of each edge in the graph.

The disadvantage of progress measures as above is that on graphs of infinite size, such as run dags, the well-ordered sets of labels that are needed to certify co-Büchi conditions may be of unbounded (and possibly infinite) size. In order to overcome this disadvantage, and to enable automata-theoretic uses of progress measure certificates, Klarlund has proposed the following concept of lazy progress measures [11]. A *lazy (co-Büchi) progress measure* is a mapping $\mu : V \rightarrow M$, where (M, \leq) is a well-ordered set and $L \subset M$ is the set of *lazy-progress* elements, and such that:

- 1) for every edge $(v, u) \in E$, we have $\mu(v) \geq \mu(u)$;
- 2) if $\pi(v) = 1$ then $\mu(v) \in L$;
- 3) on every infinite path in G , there are infinitely many vertices v such that $\mu(v) \notin L$.

It is elementary to prove the following proposition.

Proposition 1. *If a graph has a lazy co-Büchi progress measure then all infinite paths in it satisfy the co-Büchi condition.*

The following converse establishes the attractiveness of lazy co-Büchi progress measures for certifying the co-Büchi conditions on layered dags of bounded width, and hence for certifying accepting run dags of alternating co-Büchi automata.

Lemma 2 (Klarlund [11]). *If all infinite paths in a layered dag $(V, E, \pi : V \rightarrow \{0, 1\})$ satisfy the co-Büchi condition and the width of the dag is at most n , then there is a lazy co-Büchi progress measure $\mu : V \rightarrow M$, where $M = \{1, 2, \dots, 2n\}$ and $L = \{2, 4, 6, \dots, 2n\}$.*

Proof. We summarize a proof given by Kupferman and Vardi [14] that provides an explicit decomposition of the accepting run dag into (at most) $2n$ parts from which a lazy co-Büchi progress measure can be straightforwardly defined. The proof by Klarlund [11] is more succinct, but the former is more constructive and hence more transparent.

Observe that if all infinite paths satisfy the co-Büchi condition then there must be a vertex v whose all descendants (i.e., vertices to which there is a—possibly empty—path from v) have priority 0; call such vertices *1-safe* in $G_1 = G$. Indeed, otherwise it would be easy to construct an infinite path with infinitely many occurrences of vertices of priority 1.

Let S_1 be the set of all the 1-safe vertices in G_1 , and let $G'_1 = G_1 \setminus S_1$ be the layered dag obtained from G_1 by removing all vertices in S_1 . Note that there is an infinite path in the subgraph of G_1 induced by S_1 , and hence the width of G'_1 is strictly smaller than the width of G_1 .

Let R_1 be the set of all vertices in G'_1 that have only finitely many descendants; call such vertices *transient* in G'_1 . Let G_2 be the layered dag obtained from G'_1 by removing all vertices in R_1 . Since G_2 is a subgraph of G'_1 , the width of G_2 is strictly smaller than the width of G_1 . Moreover, G_2 shares the key properties with G_1 : every vertex has a successor and hence all the maximal paths are infinite, and all infinite paths satisfy the co-Büchi condition.

By applying the same procedure to G_2 that we have described for G_1 above, we obtain the set S_2 of 1-safe vertices in G_2 and the set R_2 of vertices transient in G'_2 , and the layered dag G_3 —obtained from G_2 by removing all vertices in $S_2 \cup R_2$ —has the width that is strictly smaller than that of G_2 . We can continue in this fashion until the graph G_{k+1} , for some $k \geq 1$, is empty. Since the width of G is at most n , and the widths of graphs G_1, G_2, \dots, G_{k+1} are strictly decreasing, it follows that $k \leq n$.

We define $\mu : V \rightarrow \{1, 2, \dots, 2n\}$ by:

$$\mu(v) = \begin{cases} 2i - 1 & \text{if } v \in S_i, \\ 2i & \text{if } v \in R_i, \end{cases}$$

and note that it is routine to verify that if we let $L = \{2, 4, \dots, 2n\}$ be the set of lazy-progress elements then μ is a lazy co-Büchi progress measure. \square

B. From co-Büchi and Büchi to weak

In this section we present a proof of the following result.

Theorem 3 (Kupferman and Vardi [14]). *There is a translation that given an alternating co-Büchi automaton with n states yields an equivalent alternating weak automaton with $O(n^2)$ states.*

Proof. It suffices to argue that, given an alternating co-Büchi automaton $\mathcal{A} = (Q, q_0, \Sigma, \delta, \pi : Q \rightarrow \{0, 1\})$ with n states, we can design an alternating weak automaton with $O(n^2)$ states that guesses and certifies a dag run of \mathcal{A} together with a lazy co-Büchi progress measure on it as described in Lemma 2. First we construct a *safety* automaton \mathcal{S} with $O(n^2)$ states that simulates the automaton \mathcal{A} while guessing a lazy co-Büchi progress measure and verifying conditions 1) and 2) of its definition. Condition 3) will be later handled by turning the safety automaton \mathcal{S} into a weak automaton \mathcal{W} by appropriately assigning odd or even priorities to all states in \mathcal{S} . We split the design of \mathcal{W} into those two steps so that we can better motivate and explain the generalized constructions in Section V.

The safety automaton \mathcal{S} has the following set of states:

$$Q \times \{2, 4, \dots, 2n\} \cup (\pi^{-1}(0) \times \{1, 3, \dots, 2n - 1\}) \cup \{\text{reject}\};$$

its initial state is $(q_0, 2n)$; and its transition function δ' is obtained from the transition function δ of \mathcal{A} in the following way: for every state (q, i) , and for every $a \in \Sigma$, the formula $\delta'((q, i), a)$ is obtained from $\delta(q, a)$ by replacing every occurrence of state $q' \in Q$ by the disjunction (i.e., a non-deterministic choice)

$$(q', i) \vee (q', i - 1) \vee \dots \vee (q', 1) \tag{1}$$

where every occurrence (q', j) for which $\pi(q') = 1$ and j is odd stands for the state reject.

In other words, the safety automaton \mathcal{S} can be thought of as consisting of $2n$ copies $\mathcal{A}_{2n}, \mathcal{A}_{2n-1}, \dots, \mathcal{A}_1$ of \mathcal{A} , with the non-accepting states $\pi^{-1}(1)$ removed from the odd-indexed copies $\mathcal{A}_{2n-1}, \mathcal{A}_{2n-3}, \dots, \mathcal{A}_1$, and in whose acceptance games, Elvis always has the choice to stay in the current copy of \mathcal{A} or to move to one of the lower-indexed copies of \mathcal{A} . Since the transitions of the safety automaton \mathcal{S} always respect the transitions of the original co-Büchi automaton \mathcal{A} , an accepting run dag of \mathcal{S} yields a run dag of \mathcal{A} (obtained from the first components of the states (q, i)) and a labelling of its vertices by numbers in $\{1, 2, \dots, 2n\}$ (obtained from the second components of the states (q, i)). It is routine to verify that the design of the state set and of the transition function of the safety automaton \mathcal{S} guarantees that the latter labelling satisfies conditions 1) and 2) of the definition of a lazy co-Büchi progress measure, where the set of lazy-progress elements is $\{2, 4, \dots, 2n\}$.

By setting the state priority function $\pi' : (q, i) \mapsto i + 1$ for all non-reject states in \mathcal{S} , and $\pi' : \text{reject} \mapsto 1$, we obtain from \mathcal{S} an automaton \mathcal{W} whose acceptance condition is weak because—by design—the transition function is non-increasing w.r.t. the state priority function. One can easily verify that the addition of this weak acceptance condition to \mathcal{S} allows the resulting automaton \mathcal{W} , for every input word, to guess and verify a lazy progress measure—if one exists—on a run dag of automaton \mathcal{A} on the input word, while \mathcal{W} rejects the input word otherwise. This completes our summary of the proof of Theorem 3. \square

C. From Büchi to weak

Theorem 4 (Kupferman and Vardi [14]). *There is a translation that given an alternating Büchi automaton with n states yields an equivalent alternating weak automaton with $O(n^2)$ states.*

The argument of Kupferman and Vardi is simple and it exploits the ease with which alternating automata can be complemented. Given an alternating Büchi automaton \mathcal{A} with n states, first complement it with no state space blow-up, obtaining an alternating co-Büchi automaton with n states, next use the translation from Theorem 3 to obtain an equivalent alternating weak automaton with $O(n^2)$ states, and finally complement the latter again with no state space blow-up, hence obtaining an alternating weak automaton that is equivalent to \mathcal{A} and that has $O(n^2)$ states.

We also present an explicit and direct translation from alternating Büchi automata to alternating weak automata that foregoes the double complementation. While it is technically identical to the one described above, we hope that it gives a reader more immediate insights about winning strategies of Elvis and Alice in acceptance games for alternating automata.

Proof. (of Theorem 4) Let $\mathcal{B} = (Q, q_0, \Sigma, \delta, \pi : Q \rightarrow \{1, 2\})$ be an alternating Büchi automaton with n states. We now describe the design of an alternating weak automaton \mathcal{U} , with $O(n^2)$ states, that is equivalent to \mathcal{B} . Recall that the main idea in the design of the weak alternating automaton \mathcal{W} was that, in its acceptance games, Elvis guessed a run of the co-Büchi automaton \mathcal{A} together with a lazy progress measure on it, which certified that the run was indeed accepting. The main idea in the design of the weak alternating automaton \mathcal{U} is that, in its acceptance games, Elvis is responsible for guessing a run of the Büchi automaton \mathcal{B} , but it is Alice who is given the task of guessing a lazy co-Büchi progress measure on it. The goal of Alice is to certify that the run of the Büchi automaton \mathcal{B} guessed by Elvis is not accepting.

More specifically, we construct the alternating weak automaton \mathcal{U} from the alternating Büchi automaton \mathcal{B} in a way that mirrors the way that the alternating weak automaton \mathcal{W} is constructed from the alternating co-Büchi automaton \mathcal{A} in the proof of Theorem 3. The only modifications needed are as follows.

- Instead of an additional absorbing reject state, we have an additional absorbing accept state.
- The disjunction (1) is replaced by the conjunction:

$$(q', i) \wedge (q', i - 1) \wedge \dots \wedge (q', 1)$$

in which every occurrence (q', j) for which $\pi(q') = 2$ and j is odd stands for the state `accept`.

- The state priority function is defined by $\pi' : (q, i) \mapsto i$ for all non-accept states and $\pi' : \text{accept} \mapsto 0$. Note that the automaton \mathcal{U} can be thought of as consisting of $2n$ copies $\mathcal{B}_{2n}, \mathcal{B}_{2n-1}, \dots, \mathcal{B}_1$ of \mathcal{B} , with the accepting states $\pi^{-1}(2)$ in the odd-indexed copies replaced by the absorbing accepting state `accept`. Moreover, in the acceptance games for \mathcal{U} , Alice always has a choice to stay in the current copy of \mathcal{B} or to move to one of the lower-indexed copies of \mathcal{B} . Finally, all states in the even-indexed copies of \mathcal{B} are accepting according to the weak acceptance condition of \mathcal{U} , and all non-accept states in the odd-indexed copies are rejecting.

We show that the alternating weak automaton \mathcal{U} is equivalent to the alternating Büchi automaton \mathcal{B} .

Firstly, we argue that if Elvis has a winning strategy in the acceptance game $\mathcal{G}(\mathcal{B}, w)$ then he also has a winning strategy in the acceptance game $\mathcal{G}(\mathcal{U}, w)$. Consider a straightforward adaptation of a winning strategy for Elvis in $\mathcal{G}(\mathcal{B}, w)$: in each copy of \mathcal{B} in \mathcal{U} , Elvis plays in the same way as he plays in $\mathcal{G}(\mathcal{B}, w)$. As long as a play remains inside one of the even-indexed copies of \mathcal{B} in \mathcal{U} , Elvis wins because all states in such a copy are accepting in \mathcal{U} . On the other hand, all non-accept states in every odd-indexed copy of \mathcal{B} in \mathcal{U} are rejecting, so in order to show that Elvis wins, we argue that every play that stays in such an odd-indexed copy of \mathcal{B} eventually reaches the absorbing state `accept`. Indeed, this happens because Elvis uses a winning strategy in the Büchi game $\mathcal{G}(\mathcal{B}, w)$, and hence every play eventually reaches a state in \mathcal{B} of priority 2.

Next, we argue that if Alice has a winning strategy in the acceptance game $\mathcal{G}(\mathcal{B}, w)$ then she also has a winning strategy in the acceptance game $\mathcal{G}(\mathcal{U}, w)$. By positional determinacy of parity games [7], [21], Alice has a positional winning strategy in $\mathcal{G}(\mathcal{B}, w)$, whose strategy subgraph [9] is a layered dag of width at most $|Q|$, and in which all paths satisfy the co-Büchi condition induced by the function $\pi' : q \rightarrow \pi(q) - 1$. By Lemma 2, it follows that there is a lazy co-Büchi progress measure on this strategy subgraph that uses $2n$ labels $\{1, 2, \dots, 2n\}$, where the even numbers among them are the lazy-progress labels.

The strategy for Alice in $\mathcal{G}(\mathcal{U}, w)$ follows her winning strategy in $\mathcal{G}(\mathcal{B}, w)$, thus restricting the set of plays to those that are in the strategy subgraph above. However, Alice has more choice in $\mathcal{G}(\mathcal{U}, w)$ than in $\mathcal{G}(\mathcal{B}, w)$: in every move she has to decide whether to stay in the current copy of \mathcal{B} in \mathcal{U} or to move to another one with a lower index. We argue that Alice can win in $\mathcal{G}(\mathcal{U}, w)$ by always staying in the copy of \mathcal{B} whose index is equal to the label in the lazy co-Büchi progress measure. Indeed, every play according to this strategy eventually leaves any of the even-indexed copies of \mathcal{B} by condition (3) in the definition of a lazy co-Büchi progress measure. Moreover, no play according to this strategy ever reaches the absorbing state `accept` because of condition (2) in the definition of a lazy co-Büchi progress measure. It follows that every play according to Alice's strategy eventually forever stays in one of the odd-indexed copies of \mathcal{B} and hence it is winning for her. \square

IV. LAZY PARITY PROGRESS MEASURES

Before we introduce *lazy parity progress measures*, we recall the definition of (standard) parity progress measures [10], [5]. We define a *well-ordered tree* to be a finite prefix-closed set of sequences of elements of a well-ordered set. We call such sequences *nodes* of the tree, and their components are *branching directions*. We use the standard ancestor-descendant terminology to describe relative positions of nodes in a tree. For example, $\langle \rangle$ is the *root*; node $\langle x, y \rangle$ is the *child* of the node $\langle x \rangle$ that is reached from it via the branching direction y ; node $\langle x, y \rangle$ is the *parent* of node $\langle x, y, z \rangle$; nodes $\langle x, y \rangle$ and $\langle x, y, z \rangle$ are *descendants* of nodes $\langle \rangle$ and $\langle x \rangle$; nodes $\langle \rangle$ and $\langle x \rangle$ are *ancestors* of nodes $\langle x, y \rangle$ and $\langle x, y, z \rangle$; and a node is a *leaf* if it does not have any children. All nodes in a well-ordered tree are well-ordered by the *lexicographic order* that is induced by the well-order on the branching directions; for example, we have $\langle x \rangle < \langle x, y \rangle$, and $\langle x, y, z \rangle < \langle x, w \rangle$ if $y < w$. We define the *depth* of a node to be the number of elements in the eponymous sequence, the *height* of a tree to be the maximum depth of a node, and the *size* of a tree to be the number of its nodes.

Parity progress measures assign labels to vertices of graphs with vertex priorities, and the labels are nodes in a well-ordered tree. A *tree labelling* of a graph with vertex priorities that do not exceed a positive even integer d is a mapping from vertices of the graph to nodes in a well-ordered tree of height at most $d/2$. We write $\langle m_{d-1}, m_{d-3}, \dots, m_\ell \rangle$, for some odd ℓ , $1 \leq \ell < d$, to denote such nodes. We say that such a node has an (odd) *level* ℓ and an (even) level $\ell - 1$, and the root $\langle \rangle$ has level d . Moreover, for every priority p , $0 \leq p \leq d$, we define the p -truncation $\langle m_{d-1}, m_{d-3}, \dots, m_\ell \rangle|_p$ in the following way:

$$\langle m_{d-1}, m_{d-3}, \dots, m_\ell \rangle|_p = \begin{cases} \langle m_{d-1}, m_{d-3}, \dots, m_\ell \rangle & \text{for } p \leq \ell, \\ \langle m_{d-1}, m_{d-3}, \dots, m_{p+1} \rangle & \text{for even } p > \ell, \\ \langle m_{d-1}, m_{d-3}, \dots, m_p \rangle & \text{for odd } p > \ell. \end{cases}$$

We then say that a tree labelling μ of a graph $G = (V, E)$ with vertex priorities $\pi : V \rightarrow \{0, 1, 2, \dots, d\}$ is a *parity progress measure* if the following *progress condition* holds for every edge $(v, u) \in E$:

- 1) if $\pi(v)$ is even then $\mu(v)|_{\pi(v)} \geq \mu(u)|_{\pi(v)}$;
- 2) if $\pi(v)$ is odd then $\mu(v)|_{\pi(v)} > \mu(u)|_{\pi(v)}$.

It is well-known that satisfaction of such local conditions on every edge in a graph is sufficient for every infinite path in the graph satisfying the parity condition [9], [10]. Less obviously, it is also necessary, which can be, again, deduced from the proof of positional determinacy of parity games due to Emerson and Jutla [7]. In other words, parity progress measures are witnesses for the property that all infinite paths in a graph satisfy the parity condition. Like for the simpler co-Büchi condition, their appeal stems from the property that they certify conditions that are global and infinitary by verifying conditions that are local to every edge in the graph.

Similar to the simpler co-Büchi progress measures, parity progress measures may unfortunately require unbounded or even infinite well-ordered trees to certify parity conditions on infinite graphs, and hence we consider lazy parity progress measures, also inspired by Klarlund's pioneering work [11]. A *lazy tree* is a well-ordered tree with a distinguished subset of its nodes called *lazy nodes*. For convenience, we assume that only leaves may be lazy and the root never is.

A *lazy parity progress measure* is a tree labelling μ of a graph (V, E) , where the labels are nodes in a lazy tree T , such that:

- 1) for every edge $(v, u) \in E$, $\mu(v)|_{\pi(v)} \geq \mu(u)|_{\pi(v)}$;
- 2) if $\pi(v)$ is odd then node $\mu(v)$ is lazy and its level is at least $\pi(v)$;
- 3) on every infinite path in G , there are infinitely many vertices v , such that $\mu(v)$ is not lazy.

First we establish that existence of a lazy progress measure is sufficient for all infinite paths in a graph to satisfy the parity condition.

Lemma 5. *If a graph has a lazy parity progress measure then all infinite paths in it satisfy the parity condition.*

Proof. For the sake of contradiction, assume that there is an infinite path v_1, v_2, v_3, \dots in the graph for which the highest priority p that occurs infinitely often is odd. Let $i \geq 1$ be such that $\pi(v_j) \leq p$ for all $j \geq i$. By condition 1), we have:

$$\mu(v_i)|_p \geq \mu(v_{i+1})|_p \geq \mu(v_{i+2})|_p \geq \dots \quad (2)$$

Let $i \leq i_1 < i_2 < i_3 < \dots$ be such that $\pi(v_{i_k}) = p$ for all $k = 1, 2, 3, \dots$. By condition 2), all labels $\mu(v_{i_k})$, for $k = 1, 2, 3, \dots$, are lazy and their level in the tree is at least p . By condition 3), for infinitely many k , $\pi(v_k)$ is not lazy, so infinitely many of the inequalities in (2) must be strict, which contradicts the well-ordering of the tree T . \square

Now we argue that existence of lazy parity progress measure is also necessary for a graph to satisfy the parity condition. Moreover, we explicitly quantify the size of a lazy ordered tree the labels from which

are sufficient to give a lazy progress measure for a layered dag, as a function of the width of the dag. Before we do that, however, we introduce a simple operation that we call a lazification of a finite ordered tree. If T is a finite ordered tree, then its *lazification* $\text{lazi}(T)$ is a finite lazy tree that is obtained from T in the following way:

- all nodes in T are also nodes in $\text{lazi}(T)$ and they are not lazy;
- for every non-leaf node t in T , t has extra lazy children in the tree $\text{lazi}(T)$, one smaller and one larger than all the other children, and one in-between every pair of consecutive children.

It is routine to argue that if a tree has n leaves and it is of height at most h then its lazification $\text{lazi}(T)$ has $O(nh)$ nodes and it is also of height h .

Theorem 6 (Klarlund [11]). *If all infinite paths in a layered dag satisfy the parity condition and the width of the dag is at most n , then there is a lazy parity progress measure whose labels are nodes in a tree that is a lazification of an ordered tree with at most n leaves.*

Proof. Klarlund’s proof [11] is very succinct and it heavily relies on the result of Klarlund and Kozen on Rabin measures [12]. Our proof is not only self-contained but it also is more constructive and transparent. The hierarchical decomposition describes the fundamental structure of accepting run dags of alternating parity automata and it may be of independent interest. The argument presented here is a generalization of the proof of Lemma 2—given in Section III—from co-Büchi conditions to parity conditions.

Consider a layered dag $G = (V, E, \pi)$ where $\pi : V \rightarrow \{0, 1, 2, \dots, d\}$. For a priority p , $0 \leq p < d$, we write $G^{\leq p}$ for the subgraph induced by the vertices whose priority is at most p .

We describe the following decomposition of G . Let D be the set that consists of all vertices of the top even priority d in G , and R_0 all those vertices in the subgraph $G^{\leq d-1}$ that have finitely many descendants. We say that those vertices are $(d-1)$ -transient in $G^{\leq d-1}$. In other words, $D \cup R_0$ is the set of vertices from which every path reaches (possibly immediately) a vertex of priority d .

Let $G_1 = G \setminus (D \cup R_0)$ be the layered dag obtained from G by removing all vertices in $D \cup R_0$. Observe that every vertex in G_1 has at least one successor and hence—unless G_1 is empty—all maximal paths are infinite. W.l.o.g., assume henceforth that G_1 is not empty. We argue that there must be a vertex in G_1 whose all descendants have priorities at most $d-2$; call such vertices $(d-1)$ -safe in G_1 . Indeed, otherwise it would be easy to construct an infinite path with infinitely many occurrences of the odd priority $d-1$ and no occurrences of the top even priority d .

Let S_1 be the set of all the $(d-1)$ -safe vertices in G_1 . Let H_1 be the subgraph of G induced by S_1 , let n_1 be the width of H_1 , and note that $n_1 > 0$. Set $G'_1 = G_1 \setminus S_1$ to be the layered dag obtained from G_1 by removing all $(d-1)$ -safe vertices in G_1 .

Let R_1 be the set of all vertices in G'_1 that have only finitely many descendants; call such vertices $(d-1)$ -transient in G'_1 . Finally, let G_2 be the layered dag obtained from G'_1 by removing all the $(d-1)$ -transient vertices in G'_1 . Note that the width of G_2 is smaller than the width of G_1 by at least $n_1 > 0$.

Unless graph G_2 is empty, we can now apply the same steps to G_2 that we have described for G_1 , and obtain:

- the set S_2 of $(d-1)$ -safe vertices in G_2 ;
- the subgraph H_2 of G_2 induced by S_2 , which is of width $n_2 > 0$;
- the layered dag G'_2 , obtained from G_2 by removing all the vertices in S_2 ;
- the set R_2 of $(d-1)$ -transient vertices in G'_2 ;
- the layered graph G_3 , obtained from G_2 by removing all vertices in $S_2 \cup R_2$, and whose width is smaller than the width of G_2 by at least $n_2 > 0$.

We can continue in this fashion, obtaining graphs G_1, G_2, \dots, G_{k+1} , until the graph G_{k+1} , for some $k \geq 0$, is empty. Since the width of G is at most n and the widths of the graphs G_1, G_2, \dots, G_k are positive (unless $k = 0$), we have that $k \leq n$ and $\sum_{i=1}^k n_i \leq n$.

The proces described above yields a hierarchical decomposition of the layered dag; we now define—by induction on d —the tree that describes the shape of this decomposition. We then argue that the lazification of this tree provides the set of labels in a lazy parity progress measure.

In the base case $d = 0$, the shape of the decomposition is the well-ordered tree T of height $h = 0/2 = 0$ with only a root node $\langle \rangle$. It is straightforward to see that the function that maps every vertex onto the root is a (lazy) progress measure.

For $d \geq 2$, note that all vertices in dags H_1, H_2, \dots, H_k have priorities at most $d - 2$. By the inductive hypothesis, there are trees T_1, T_2, \dots, T_k , of heights at most $h - 1 = (d - 2)/2$ and with at most n_1, n_2, \dots, n_k leaves, respectively, which are the shapes of the hierarchical decompositions of dags H_1, H_2, \dots, H_k , respectively.

We now construct the finite ordered tree T of height at most $h = d/2$ that is the shape of the hierarchical decomposition of G : let T consist of the root node $\langle \rangle$ that has k children, which are the roots of the subtrees T_1, T_2, \dots, T_k , in that order. Note that the number of leaves of T is at most $\sum_{i=1}^k n_i \leq n$. Consider the following mapping from vertices in the graph onto nodes in the lazification $\text{lazi}(T)$ of tree T :

- vertices in set D are mapped onto the root of $\text{lazi}(T)$;
- vertices in transient sets $R_0, R_1, R_2, \dots, R_k$ are mapped onto the lazy children of the root of $\text{lazi}(T)$: those in R_0 onto the smallest lazy child, those in R_1 onto the lazy child between the roots of T_1 and T_2 , etc.;
- vertices in subgraphs H_1, H_2, \dots, H_k are inductively mapped onto the appropriate nodes in the lazy subtrees of $\text{lazi}(T)$ that are rooted in the k non-lazy children of the root.

It is easy to verify that this mapping satisfies conditions 1) and 2) of the definition of a lazy parity progress measure. Condition 3) is ensured by the fact that the root of T is not lazy and by the inductive hypothesis. Recall that every infinite path satisfies the parity condition, thus the highest priority p seen infinitely often on a given path is even. If $p = d$, the path visits infinitely often vertices labelled by the root of T . Otherwise, eventually the path contains only vertices in one of the sets S_i and we can use the inductive hypothesis. \square

V. FROM PARITY TO BÜCHI VIA UNIVERSAL TREES

In this section we complete the proof of the main technical result of the paper, which is a quasi-polynomial translation from alternating parity automata to alternating weak automata. The main technical tools that we use to design our translation are lazy progress measures and universal trees [10], [5], and the state space blow-up of the translation is merely quadratic in the size of the smallest universal tree. Nearly tight quasi-polynomial upper and lower bounds have recently been given for the size of the smallest universal trees [10], [5] and, in particular, they imply that if the number of priorities in a family of alternating parity automata is at most logarithmic in the number of states, then the state space blow-up of our translation is only polynomial.

Theorem 7. *There is a translation that given an alternating parity automaton with n states and d priorities yields an equivalent alternating weak automaton whose number of states is polynomial if $d = O(\log n)$ and it is $n^{O(\lg(d/\lg n))}$ if $d = \omega(\log n)$.*

Before we proceed to prove the theorem, we recall the notion of universal ordered trees. An (n, h) -universal (ordered) tree [5] is an ordered tree, such that every finite ordered tree of height at most h and with at most n leaves can be isomorphically embedded into it. In such an embedding, the root of the tree must be mapped onto the root of the universal tree, and the children of every node must be mapped—injectively and in an order-preserving way—onto the children of its image. In order to upper-bound the size of the blow-up in our parity to weak translation, we rely on the following upper bound on the size of the smallest universal trees.

Theorem 8 (Jurdziński and Lazić [10]). *For all positive integers n and h , there is an (n, h) -universal tree with at most quasi-polynomial number of leaves. More specifically, the number of leaves is polynomial in n if $h = O(\log n)$, and it is $n^{\lg(h/\lg n) + O(1)}$ if $h = \omega(\log n)$.*

We also note that Czerwiński et al. [5] have subsequently proved a nearly-matching quasi-polynomial lower bound, hence establishing that the smallest universal trees have quasi-polynomial size.

In order to prove Theorem 7, we establish the following lemma that provides a translation from alternating parity automata to alternating Büchi automata whose state-space blow-up is tightly linked to the size of universal trees.

Lemma 9. *There is a translation that given an alternating parity automaton with n states and d priorities yields an equivalent alternating Büchi automaton whose number of states is $O(ndL_U)$ where L_U is the number of leaves in an (n, d) -universal ordered tree U .*

Note that Theorem 7 follows from Lemma 9 by Theorem 8 and Corollary 4.

Proof. (of Lemma 9) Given an alternating parity automaton $\mathcal{A} = (Q, q_0, \Sigma, \delta, \pi : Q \rightarrow \{0, 1, \dots, d\})$ with n states, we now design an alternating Büchi automaton that guesses and certifies a dag run of \mathcal{A} together with a lazy parity progress measure on it. As for the co-Büchi to weak case, we first construct a *safety* automaton that simulates the automaton \mathcal{A} while guessing a lazy parity progress measure and verifying conditions 1) and 2) of its definition. Condition 3) will be later handled by turning the safety automaton into a Büchi automaton by appropriately assigning priorities 1 or 2 to all states in the safety automaton.

Below we give a general construction of an alternating Büchi automaton \mathcal{B}_T from any lazy well-ordered tree T , and then we argue that the alternating parity automaton \mathcal{A} is equivalent to the alternating Büchi automaton $\mathcal{B}_{\text{lazi}(U)}$, for every $(n, d/2)$ -universal tree U .

Let T be a lazy tree of width n and height $d/2$. The construction is by induction on d . The safety automaton \mathcal{S}_T has the following set of states, which are pairs of an element of Q and of a node in T .

- If $d = 0$, then the set of states of \mathcal{S}_T is $(Q \times \{\langle \rangle\}) \cup \{\text{reject}\}$.
- Otherwise, let $\langle x_1 \rangle, \langle x_2 \rangle, \dots, \langle x_k \rangle$ be the children of the root, and $1 \leq i_1 < i_2 < \dots < i_m \leq k$ are the indices of its leaves that are lazy.

For $i \notin \{i_1, i_2, \dots, i_m\}$, let T_i be the lazy subtrees of T of height at most $d/2 - 1$ rooted in $\langle x_i \rangle$. By induction, for all i , we obtain an alternating Büchi automaton that is obtained from the lazy tree T_i and from the alternating parity automaton \mathcal{A} restricted to the states of priority up to $d - 2$. Let Ω_i denote its set of non-reject states. They are pairs consisting of an element of Q and of a node in a tree of height $d/2 - 1$; the latter is a sequence $\langle m_{d-3}, m_{d-5}, \dots, m_\ell \rangle$ of at most $d/2 - 1$ branching directions. Let Γ_i be the set consisting of the pairs $(q, \langle x_i, m_{d-3}, m_{d-5}, \dots, m_\ell \rangle)$ for $(q, \langle m_{d-3}, m_{d-5}, \dots, m_\ell \rangle) \in \Omega_i$. Set $Q^{(d)}$ (resp. $Q^{(<d)}$) the set of states of priority d (resp. $< d$) in \mathcal{A} . The states of \mathcal{S}_T are defined as:

$$(Q^{(d)} \times \{\langle \rangle\}) \cup (Q^{(<d)} \times \{\langle x_{i_1} \rangle, \langle x_{i_2} \rangle, \dots, \langle x_{i_m} \rangle\}) \cup \bigcup_{i=1}^k \Gamma_i \cup \{\text{reject}\}.$$

The initial state is (q_0, t) where t is the largest tuple such that (q_0, t) is a state. Let us now define the transition function: for every state (q, t) , and for every $a \in \Sigma$, the formula $\delta'((q, t), a)$ is obtained from $\delta(q, a)$ by replacing every occurrence of state $q' \in Q$ by the disjunction (i.e., a non-deterministic choice)

$$\bigvee \{ (q', t') : t|_{\pi(q)} \geq t'|_{\pi(q)} \},$$

where every occurrence (q', t') which is not in the set of states stands for the state reject.

In other words, the safety automaton \mathcal{S}_T can be thought of as consisting of copies of \mathcal{A} , for each node of the tree T , in whose acceptance games Elvis always has the choice to stay in the current copy of \mathcal{A} or to move to one of a smaller node with respect to the priority of the current state. Since the transitions of the safety automaton \mathcal{S}_T always respect the transitions of the original parity automaton \mathcal{A} , an accepting run dag of \mathcal{S}_T yields a run dag of \mathcal{A} (obtained from the first components of the states (q, t)) and a labelling of its vertices by nodes in T (obtained from the second components of the states (q, t)). It is routine to verify that the design of the state set and of the transition function of the safety automaton \mathcal{S}_T guarantees that the latter labelling satisfies condition 1) and 2) of the definition of a lazy parity progress measure.

In order to obtain the Büchi automaton \mathcal{B}_T from the safety automaton \mathcal{S}_T , it suffices to appropriately assign priorities 1 and 2 to all states: we let the state reject and all states (q, t) such that t is a lazy node in tree T have priority 1, and we let all states (q, t) such that t is not a lazy node in tree T have priority 2. Note that this ensures that a run of \mathcal{B}_T is accepting if and only if the tree labelling of a run dag of \mathcal{A} that the underlying safety automaton \mathcal{S}_T guesses—in the second component of its states—satisfies condition 3) of the definition of a lazy parity progress measure.

We now argue that if U is an (n, d) -universal tree then the alternating Büchi automaton $\mathcal{B}_{\text{lazi}(U)}$ is equivalent to the alternating parity automaton \mathcal{A} . Firstly, all words accepted by \mathcal{B}_T for any finite lazy ordered tree T are also accepted by \mathcal{A} . This is because—as we have argued above—every accepting run dag of any such automaton \mathcal{B}_T yields both a run dag of \mathcal{A} (in the first state components) and a lazy parity progress measure on it (in the second state components), and the latter certifies that the former is accepting.

It remains to argue that every word accepted by \mathcal{A} is also accepted by $\mathcal{B}_{\text{lazi}(U)}$. By Theorem 6, for every accepting run dag of \mathcal{A} , there is a lazy progress measure whose labels are nodes in a tree $\text{lazi}(T)$, where T is an ordered tree of height at most $d/2$ and with at most n leaves. It is routine to verify that if an ordered tree can be isomorphically embedded in another, then the same holds for their lazifications. By $(n, d/2)$ -universality of U , it follows that $\text{lazi}(T)$ can be isomorphically embedded in $\text{lazi}(U)$. Therefore, for every word on which there is an accepting run dag of \mathcal{A} , the automaton $\mathcal{B}_{\text{lazi}(U)}$ has the capacity to guess the run dag of \mathcal{A} and to guess and certify a lazy progress measure on it.

In order to conclude the $O(ndL_U)$ upper bound on the number of states of $\mathcal{B}_{\text{lazi}(U)}$, it suffices to observe that if the number of leaves in an ordered tree T of height h is L then the number of nodes in $\text{lazi}(T)$ is $O(hL)$. \square

VI. CONCLUSION

Our new translation from alternating parity to weak automata nearly matches the current state of the art upper bounds for solving parity games [2], [10], [8]; any further significant improvements would imply a breakthrough in the complexity of parity games. Our use of universal trees to turn alternating parity automata into Büchi automata, like Boker and Lehtinen’s [1] register technique, does not exploit alternations (although the further Büchi to weak translation does): all transitions that are not copied from the original automaton are non-deterministic. Whether universal and non-deterministic choices could be combined to further improve these translations remains open. Conversely, perhaps there is hope that the long-standing $\Omega(n \log n)$ lower bound [17] could be improved for example by a combination of the full-automata technique of Yan [20] and the recent lower bound techniques for non-deterministic safety separating automata based on universal trees [5] and universal graphs [3].

ACKNOWLEDGMENT

This work has been supported by the EPSRC grant EP/P020992/1 (Solving Parity Games in Theory and Practice).

We thank Moshe Vardi for encouraging us to bring the state space blow-up of alternating parity to weak word automata translation in line with the state of the art complexity of solving parity games.

REFERENCES

- [1] U. Boker and K. Lehtinen, “On the way to alternating weak automata,” in *FSTTCS 2018*, ser. LIPIcs, vol. 122. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2018, pp. 21:1–21:22.
- [2] C. S. Calude, S. Jain, B. Khoussainov, W. Li, and F. Stephan, “Deciding parity games in quasipolynomial time,” in *STOC 2017*. ACM, 2017, pp. 252–263.
- [3] T. Colcombet and N. Fijalkow, “Parity games and universal graphs,” arXiv:1810.05106, 2018.
- [4] T. Colcombet, D. Kuperberg, C. Löding, and M. Vanden Boom, “Deciding the weak definability of Büchi definable tree languages,” in *CSL 2013*, ser. LIPIcs, vol. 23. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2013, pp. 215–230.
- [5] W. Czerwiński, L. Daviaud, N. Fijalkow, M. Jurdziński, R. Łazić, and P. Parys, “Universal trees grow inside separating automata: Quasi-polynomial lower bounds for parity games,” in *SODA 2019*. ACM/SIAM, 2019, pp. 2333–2349.
- [6] L. Daviaud, M. Jurdziński, and R. Łazić, “A pseudo-quasi polynomial algorithm for solving mean-payoff parity games,” in *LICS 2018*. IEEE, 2018, pp. 325–334.
- [7] E. A. Emerson and C. Jutla, “Tree automata, mu-calculus and determinacy,” in *FOCS 1991*. IEEE, 1991, pp. 368–377.
- [8] J. Fearnley, S. Jain, S. Schewe, F. Stephan, and D. Wojtczak, “An ordered approach to solving parity games in quasi polynomial time and quasi linear space,” in *SPIN 2017*, 2017, pp. 112–121.
- [9] M. Jurdziński, “Small progress measures for solving parity games,” in *STACS 2000*, 2000, pp. 290–301.
- [10] M. Jurdziński and R. Łazić, “Succinct progress measures for solving parity games,” in *LICS 2017*. IEEE, 2017, p. 9pp.
- [11] N. Klarlund, “Progress measures for complementation of ω -automata with applications to temporal logic,” in *FOCS 1991*. IEEE, 1991, pp. 358–367.
- [12] N. Klarlund and D. Kozen, “Rabin measures,” *Chicago J. Theor. Comput. Sci.*, pp. 1–24, 1995, article 3.
- [13] O. Kupferman and M. Y. Vardi, “Weak alternating automata and tree automata emptiness,” in *STOC 1998*. ACM, 1998, pp. 224–233.
- [14] —, “Weak alternating automata are not that weak,” *ACM Trans. Comput. Log.*, vol. 2, no. 3, pp. 408–429, 2001.
- [15] K. Lehtinen, “A modal μ perspective on solving parity games in quasi-polynomial time,” in *LICS 2018*. IEEE, 2018, pp. 639–648.
- [16] K. Lehtinen and S. Quickert, “ Σ_2^μ is decidable for Π_2^μ ,” in *CiE 2017: Unveiling Dynamics and Complexity*, 2017, pp. 292–303.
- [17] M. Michel, “Complementation is more difficult with automata on infinite words,” *CNET, Paris*, vol. 15, 1988.
- [18] M. Skrzypczak and I. Walukiewicz, “Deciding the topological complexity of Büchi languages,” in *ICALP 2016*, ser. LIPIcs, vol. 55. Schloss Dagstuhl–Leibniz-Zentrum für Informatik, 2016, pp. 99:1–99:13.
- [19] M. Y. Vardi, “The Büchi complementation saga,” in *STACS 2007*. Springer, 2007, pp. 12–22.
- [20] Q. Yan, “Lower bounds for complementation of ω -automata via the full automata technique,” *Logical Methods in Computer Science*, vol. 4, pp. 1–20, 2008.
- [21] W. Zielonka, “Infinite games on finitely coloured graphs with applications to automata on infinite trees,” *Theor. Comput. Sci.*, vol. 2000, no. 1–2, pp. 135–183, 1998.