

Winning regions of higher-order pushdown games*

A. Carayol[†]

M. Hague[‡]

A. Meyer[§]

C.-H. L. Ong[¶]

O. Serre^{||}

Abstract

In this paper we consider parity games defined by higher-order pushdown automata. These automata generalise pushdown automata by the use of higher-order stacks, which are nested “stack of stacks” structures. Representing higher-order stacks as well-bracketed words in the usual way, we show that the winning regions of these games are regular sets of words. Moreover a finite automaton recognising this region can be effectively computed.

A novelty of our work are abstract pushdown processes which can be seen as (ordinary) pushdown automata but with an infinite stack alphabet. We use the device to give a uniform presentation of our results.

From our main result on winning regions of parity games we derive a solution to the Modal Mu-Calculus Global Model-Checking Problem for higher-order pushdown graphs as well as for ranked trees generated by higher-order safe recursion schemes.

1 Introduction

Higher-order pushdown automata were introduced by Maslov [31] as a generalisation of pushdown automata and nested pushdown automata to extend the indexed languages introduced by Aho [1]. Whereas an ordinary (*i.e.* order-1) pushdown automaton works with a stack of symbols (*i.e.* order-1 stack), a pushdown automaton of order 2 works with a stack of (order-1) stacks. In addition to pushing a symbol onto and popping a symbol from the top-most order-1 stack, an order-2 pushdown automaton can duplicate or remove the entire top-most (order-1) stack. Pushdown automata of higher orders are defined in a similar

way and have been extensively studied as language acceptors [15, 17].

Recently, the infinite structures defined by these automata have received a lot of attention. In [28], the families of infinite terms defined by higher-order pushdown automata were shown to correspond to the solutions of safe higher-order recursion schemes. Subsequently, in [14, 12], the ε -closure of their configuration graphs were shown to be exactly those constructible from finite graphs using natural graph transformations (see [35] for a survey).

A remarkable property of these graphs is that we can decide the validity of any formula of monadic second-order (MSO) logic. Unfortunately the decision procedure is non-elementary (in the size of the formula) and this is already so in the case of pushdown graphs. In order to obtain an elementary decision procedure, we consider the μ -calculus: a weaker modal logic equi-expressive with MSO over trees [26]. The two main algorithmic problems in this setting are the (*local*) *model-checking problem* (*i.e.* to decide if a particular configuration satisfies a given μ -calculus formula) and the *global model-checking problem* (*i.e.* to compute a finite description of the set of configurations satisfying a given formula). To solve these problems, we consider the associated two-player parity game, whose size is polynomial in both the size of the automaton and of the formula [16]. The two versions of the model-checking problem above are respectively equivalent to deciding which player wins from a given configuration and to giving a finite description of the winning region for each player.

For parity games defined by pushdown automata, Walukiewicz has given an EXPTIME decision procedure to compute the winner from a given configuration [37]. In [6, 32], the winning region is shown to be regular when a configuration (q, w) is represented by the word qw . Note that this result can easily be derived from the results by Vardi [36]. For order- n pushdown automata, an n -EXPTIME decision procedure for the local version of the problem was given by Cachet [7] using techniques from [36]. In this article, we shall consider the global version of the problem *i.e.* the computation of a finite representation of the winning region, and we obtain, as a by-product, a new proof of Cachet’s result.

To give a finite description of a set of higher-order stacks,

*We direct readers to the (downloadable) long version [11] of this paper in which all proofs are presented.

[†]Arnaud.Carayol@univ-mlv.fr IGM (Université Paris Est & CNRS)

[‡]Matthew.Hague@comlab.ox.ac.uk Oxford University Computing Laboratory (OUEL)

[§]Antoine.Meyer@liafa.jussieu.fr LIAFA (Université Paris Diderot – Paris 7 & CNRS)

[¶]Luke.Ong@comlab.ox.ac.uk OUEL

^{||}Olivier.Serre@liafa.jussieu.fr LIAFA

we represent these stacks as well-bracketed words. Note that the depth of the bracketing is bounded by the order of the stack. By extension, we say that a set of higher-order stacks is regular if the set of associated well-bracketed words is regular. Our main result is that the winning regions of parity games over higher-order pushdown graphs are regular. Moreover we can construct in n -EXPTIME a finite deterministic automaton accepting it.

To simplify the presentation, we consider a more general notion of pushdown automata, called *abstract pushdown automata*, which work with a possibly infinite stack alphabet. Standard pushdown automata are abstract pushdown automata with a finite stack alphabet; order- $(k + 1)$ pushdown automata are abstract pushdown automata whose stack alphabet is the (infinite) set of order- k stacks. Our main technical result concerns parity games over the configuration graphs of abstract pushdown automata. From an abstract pushdown parity game, we construct a reduced parity game based on the stack alphabet (*i.e.* which does not make use of a stack structure) and we show that a finite description of the winning region of the original game can be derived from a finite description of the winning region of the reduced game. Applied to higher-order pushdown automata, this result allows us to reduce the problem of computing the winning region of an order- $(k + 1)$ pushdown parity game to that of computing the winning region of an order- k pushdown parity game. We also show that starting from a winning strategy in the reduced parity game one can effectively build an abstract pushdown automaton that realises a winning strategy in the original game and whose stack is synchronised with the one used in the game. Applied to higher-order pushdown parity games, this result allows us to prove that any such game always admits an effective winning strategy realised by a higher-order pushdown automata of the same order and whose stack is synchronised with the one used in the game.

As an application of these results, we solve the μ -calculus global model-checking problem for higher-order pushdown graphs, and for ranked trees generated by safe higher-order recursion schemes.

Related work. In [4], Bouajjani and Meyer considered simpler reachability games over higher-order pushdown automata with a single control state. They showed that the winning regions in these games are regular. In [25], Hague and Ong extended the result to arbitrary higher-order pushdown automata. The results presented here have been obtained by Hague and Ong, and also (separately and independently) by the other co-authors. The present article follows the latter approach whereas the proof in [23] generalises the saturation method developed in [25].

A similar result was obtained in [13] for a stronger notion of regularity (introduced in [10, 20]) which coincides

with MSO-definable sets of configurations. In contrast, the simpler notion presented here can only capture properties definable in μ -calculus. In particular, owing to the higher-order *push* operations, the set of configurations reachable from a given configuration is not regular in the sense of the present paper. (To our knowledge the results presented here do not appear to be derivable from Carayol’s work [10].)

2 Definitions

An *alphabet* A is a (possibly infinite) set of letters. In the sequel A^* denotes the set of *finite words* over A and A^ω the set of *infinite words* over A . The empty word is denoted by ε .

Infinite two-player games. Let $G = (V, E)$ be a (possibly infinite) graph with vertex-set V and edge-set $E \subseteq V \times V$. Let $V_E \cup V_A$ be a partition of V between two players, Éloïse and Abelard. A *game graph* is a tuple $\mathcal{G} = (V_E, V_A, E)$. An *infinite two-player game* on a game graph \mathcal{G} is a pair $\mathbb{G} = (\mathcal{G}, \Omega)$, where $\Omega \subseteq V^\omega$ is a *winning condition*.

Éloïse and Abelard play in \mathbb{G} by moving a token between vertices. A *play* from some initial vertex v_0 proceeds as follows: the player owning v_0 moves the token to a vertex v_1 such that $(v_0, v_1) \in E$. Then the player owning v_1 chooses a successor v_2 and so on. If at some point one of the players cannot move, she/he loses the play. Otherwise, the play is an infinite word $\Lambda \in V^\omega$ and is won by Éloïse if and only if $\Lambda \in \Omega$. Nevertheless, for all game graphs considered in this article one can always assume without loss of generality that they have no dead-ends. A *partial play* is any prefix of a play.

A strategy for Éloïse is a function assigning, to any partial play ending in some vertex $v \in V_E$, a vertex v' such that $(v, v') \in E$. Éloïse *respects a strategy* Φ during some play $\Lambda = v_0 v_1 v_2 \dots$ if $v_{i+1} = \Phi(v_0 \dots v_i)$, for all $i \geq 0$ such that $v_i \in V_E$. A strategy Φ for Éloïse is *winning* from some position $v \in V$ if she wins every play that starts from v and respects Φ . Finally, a vertex $v \in V$ is *winning* for Éloïse if she has a winning strategy from v , and the winning region for Éloïse consists of all winning vertices for her. Symmetrically, one defines the corresponding notions for Abelard.

A game \mathbb{G} is *determined* if, from any position, either Éloïse or Abelard has a winning strategy. Determinacy of all games considered here follows from Martin’s Theorem [30].

For more details and basic results on games, we refer the reader to [34, 39, 22, 38].

Higher-order pushdown processes. An *order- n pushdown process* is a tuple $\mathcal{P} = \langle Q, \Sigma, \perp, \delta \rangle$ where Q is a finite

set of control states, Σ is a finite stack alphabet containing a bottom-of-stack symbol $\perp \in \Sigma$, and Δ is a transition function (to be defined below). Let $Op_1 = \{pop_1, push_1^\sigma \mid \sigma \in \Sigma \setminus \{\perp\}\}$ (resp. $Op_k = \{pop_k, push_k\}$) be the set of order-1 (resp. order- k with $2 \leq k \leq n$) stack operations. Then $\Delta : Q \times \Sigma \rightarrow 2^{Q \times Op_{\leq n}}$ where $Op_{\leq i} = \bigcup_{1 \leq k \leq i} Op_k$.

An *order-0 stack* over Σ is an element of Σ , and for $k > 0$ an *order- k stack* over Σ is a finite sequence s_1, \dots, s_ℓ of order- $(k-1)$ stacks. Moreover we require that an order- k stack be non-empty whenever $k \geq 2$. Let $k > 0$ and take an order- k stack $s = s_1, \dots, s_\ell$, we define the following partial functions:

$$\begin{aligned} top_1(s) &= \begin{cases} s_\ell & \text{if } \ell \neq 0 \text{ and } k = 1, \\ top_1(s_\ell) & \text{if } k > 1. \end{cases} \\ pop_i(s) &= \begin{cases} s_1, \dots, s_{\ell-1} & \text{if } \ell > 1 \text{ and } k = i, \\ s_1, \dots, s_{\ell-1}, pop_i(s_\ell) & \text{if } k > i. \end{cases} \\ push_1^\sigma(s) &= \begin{cases} s_1, s_2, \dots, s_\ell, \sigma & \text{if } k = 1, \\ s_1, s_2, \dots, s_{\ell-1}, push_i(s_\ell), & \text{if } k > 1. \end{cases} \\ push_i(s) &= \begin{cases} s_1, s_2, \dots, s_\ell, s_\ell & \text{if } k = i, \\ s_1, s_2, \dots, s_{\ell-1}, push_i(s_\ell), & \text{if } k > i. \end{cases} \end{aligned}$$

A configuration of \mathcal{P} is a pair (q, s) where $q \in Q$ and s is an order- n stack. At a configuration (q, s) , \mathcal{P} can apply any transition $(q', op) \in \Delta(q, top_1(s))$, leading to the configuration $(q', op(s))$. The *configuration graph* of \mathcal{P} is defined to be the graph whose vertices are the configurations and whose edges are given by the transitions of \mathcal{P} .

In order to define a notion of regularity over higher-order stacks, we associate to each order- k stack a well-bracketed word of depth k as follows. Let $s = s_1, \dots, s_\ell$ be an order- k stack. We define $\tilde{s} \in (\Sigma \cup \{[,]\})^*$ as

$$\tilde{s} = \begin{cases} [\tilde{s}_1 \cdots \tilde{s}_\ell] & \text{if } k \geq 1 \\ s & \text{if } k = 0 \text{ (i.e. } s \in \Sigma) \end{cases}$$

A set S of order- k stacks is said to be *regular* if the language $\tilde{S} = \{\tilde{s} \mid s \in S\}$ is a regular subset of $(\Sigma \cup \{[,]\})^*$. By extension, a set C of configurations of an order- k pushdown automaton is said to be regular if the set $\tilde{C} = \{\tilde{s}q \mid (q, s) \in C\}$ is a regular subset of $(\Sigma \cup \{[,]\})^*Q$. For example $\{[\underbrace{[a] \cdots [a]}_n] \mid n \geq 0\}$ is a regular set of order-2

stacks, but $\{[\underbrace{[a \cdots a]}_n \underbrace{[a \cdots a]}_n] \mid n \geq 0\}$ is not.

In particular, the set of configurations reachable from a given configuration is not in general regular. Consider for example an order-2 pushdown automaton with two states q_0 and q_1 and two transitions $\Delta(q_0, a) = \{(q_0, push_1^a), (q_1, push_2)\}$. The language of words representing the configurations reachable from $(q_0, [[a]])$ is

the context-free language $\{[[a^n]]q_0 \mid n \geq 1\} \cup \{[[a^n][a^n]]q_1 \mid n \geq 1\}$.

When considering higher-order pushdown automata as acceptors of finite-word languages over a finite alphabet Σ , we attach to each transition a symbol in $\Sigma \cup \{\varepsilon\}$ and fix an initial state q_0 together with a set F of final states. We use the symbol ε to label silent transitions. We call such automata *labelled higher-order pushdown automata*.

A word $w \in \Sigma^*$ is accepted by the automaton if there exists a sequence of configurations $c_0 \xrightarrow{a_1} \dots \xrightarrow{a_n} c_n$ where c_0 is the initial configuration $(q_0, [{}^n\perp]^n)$, c_n contains a final state in $F \subseteq Q$ and w is the word obtained by removing all occurrences of ε in $a_1 \dots a_n$. (We say that w is the word *traced out* by the configuration sequence $c_0 \xrightarrow{a_1} \dots \xrightarrow{a_n} c_n$.)

Proposition 1. *Take an (resp. deterministic) order- k pushdown automaton. The set of words traced out by configuration sequences from the initial configuration c_0 to a configuration c , where c ranges over a regular set of configurations, is accepted by an (resp. deterministic) order- k pushdown automaton.*

Abstract Pushdown Processes. We situate the techniques developed here in a general and abstract framework of (order-1) pushdown processes whose stack alphabet is a possibly infinite set.

An *abstract pushdown process* is a tuple $\mathcal{P} = \langle Q, \Gamma, \Delta \rangle$ where Q is a finite set of states, Γ is a (possibly infinite) set called an *abstract pushdown alphabet* and containing a bottom-of-stack symbol denoted $\perp \in \Gamma$, and

$$\Delta : Q \times \Gamma \rightarrow 2^{Q \times \{rew(\gamma), pop, push(\gamma) \mid \gamma \in \Gamma\}}$$

is the transition relation. We additionally require that for all $\gamma \neq \perp$, $\Delta(q, \gamma)$ does not contain any element of the form $(q, push(\perp))$ or $(q', rew(\perp))$, and that $\Delta(q, \perp)$ does not contain any element of the form (q', pop) or $(q', rew(\gamma))$ with $\gamma \neq \perp$, i.e. the bottom-of-stack symbol can only occur at the bottom of the stack, and is never popped or rewritten.

An *abstract pushdown content* is a word in $St = \perp(\Gamma \setminus \{\perp\})^*$. A configuration of \mathcal{P} is a pair (q, σ) with $q \in Q$ and $\sigma \in St$. Note that the top stack symbol in some configuration (q, σ) is the rightmost symbol of σ .

Remark 1. In general an abstract pushdown process is not finitely describable, as the domain of Δ is infinite and no further assumption is made on Δ .

Example 1. A pushdown process is an abstract pushdown process whose stack alphabet is finite.

A abstract pushdown process \mathcal{P} induces a possibly infinite graph, called an *abstract pushdown graph*, denoted $G = (V, E)$, whose vertices are the configurations of \mathcal{P}

(i.e. pairs from $Q \times St$), and edges are defined by the transition relation Δ , i.e., from a vertex $(p, \sigma\gamma)$ one has edges to:

- $(q, \sigma\gamma')$ whenever $(q, \text{rew}(\gamma')) \in \Delta(p, \gamma)$.
- (q, σ) whenever $(q, \text{pop}) \in \Delta(p, \gamma)$.
- $(q, \sigma\gamma')$ whenever $(q, \text{push}(\gamma')) \in \Delta(p, \gamma)$.

Example 2. Higher-order pushdown processes are special cases of abstract pushdown processes. Let $n > 1$ and consider an order- n pushdown process $\mathcal{P} = \langle Q, \Sigma, \Delta \rangle$. Set Γ to be the set of all order- $(n-1)$ stacks over Σ , and for every $p \in Q$ and $\gamma \in \Gamma$ with $\sigma = \text{top}_1(\gamma)$, we define $\Delta'(p, \gamma)$ by

- $(q, \text{pop}) \in \Delta'(p, \gamma)$ iff $(q, \text{pop}_n) \in \Delta(p, \sigma)$;
- $(q, \text{push}(\gamma)) \in \Delta'(p, \gamma)$ iff $(q, \text{push}_n) \in \Delta(p, \sigma)$;
- $(q, \text{rew}(\sigma\gamma)) \in \Delta'(p, \gamma)$ iff $(q, \sigma) \in \Delta(p, \gamma)$ where $k < n$ and σ is an order- k action.

It follows that the abstract pushdown process $\langle Q, \Gamma, \Delta' \rangle$ and \mathcal{P} have isomorphic transition graphs.

Abstract Pushdown Parity Game. Consider a partition $Q_E \cup Q_A$ of Q between Éloïse and Abelard. It induces a natural partition $V_E \cup V_A$ of V by setting $V_E = Q_E \times St$ and $V_A = Q_A \times St$. The resulting game graph $\mathcal{G} = (V_E, V_A, E)$ is called an *abstract pushdown game graph*. Finally, an *abstract pushdown game* is a game played on such a game graph.

Let ρ be a colouring function from Q to a finite set of colours $C \subset \mathbb{N}$. This function is easily extended to a function from V to C by setting $\rho((q, \sigma)) = \rho(q)$. The parity condition is the winning condition defined by:

$$\Omega_{\text{par}} = \{v_0 v_1 \dots \mid \liminf((\rho(v_i))_{i \geq 0}) \text{ is even}\}$$

For an abstract pushdown parity game, the main questions are the following:

1. For a given vertex, decide who wins from it.
2. For a given vertex, compute a winning strategy for the winning player.
3. Compute a finite representation of the winning regions.

These three problems are polynomially equivalent to the following problems respectively: model-checking for μ -calculus, controller synthesis against μ -calculus specifications, and global model-checking for μ -calculus.

Automata with oracles. We now define a class of automata to accept the winning positions in an abstract pushdown game. An *automaton with oracles* is a tuple $\mathcal{A} = (\mathcal{S}, Q, \Gamma, \delta, s_{\text{in}}, \mathcal{O}_1 \dots \mathcal{O}_n, \text{Acc})$ where \mathcal{S} is a finite set of control states, Q is a set of input states, Γ is a (possibly infinite) input alphabet, $s_{\text{in}} \in \mathcal{S}$ is the initial state, \mathcal{O}_i are subsets of Γ (called *oracles*) and $\delta : \mathcal{S} \times \{0, 1\}^n \rightarrow \mathcal{S}$ is the transition function. Finally Acc is a function from \mathcal{S} to 2^Q . Such an automaton is designed to accept in a *deterministic* way configurations of an abstract pushdown process whose abstract pushdown content alphabet is Γ and whose control states are Q .

Let $\mathcal{A} = (\mathcal{S}, Q, \Gamma, \delta, s_{\text{in}}, \mathcal{O}_1 \dots \mathcal{O}_n, \text{Acc})$ be such an automaton. With every $\gamma \in \Gamma$ we associate a Boolean vector $\pi(\gamma) = (b_1, \dots, b_n)$ where

$$b_i = \begin{cases} 1 & \text{if } \gamma \in \mathcal{O}_i \\ 0 & \text{otherwise.} \end{cases}$$

The automaton reads a configuration $C = (q, \gamma_1 \gamma_2 \dots \gamma_\ell)$ from left to right. A *run* over C is the sequence $s_0, \dots, s_{\ell+1}$ such that $s_0 = s_{\text{in}}$ and $s_{i+1} = \delta(s_i, \pi(\gamma_i))$ for every $i = 0, \dots, \ell$. Finally the run is *accepting* if and only if $q \in \text{Acc}(s_{\ell+1})$.

Remark 2. When the input alphabet is finite, it is easily seen that automata with oracles behave as (standard) deterministic finite automata.

In this article, we are going to use automata with oracles to accept sets of configurations of higher-order pushdown automata. As seen in Example 2 for an order- $(k+1)$ pushdown automaton, we take Γ to be the set of all order- k stacks. The sets of regular configurations of an order- $(k+1)$ pushdown automaton are naturally captured by automata using, as oracles, regular sets of order- k stacks.

Proposition 2. Fix an order- $(k+1)$ pushdown automaton \mathcal{P} and consider an automaton \mathcal{A} with oracles $\mathcal{O}_1, \dots, \mathcal{O}_n$ respectively accepted by deterministic word automata $\mathcal{A}_1, \dots, \mathcal{A}_n$. Let C be the set of configurations of \mathcal{P} accepted by \mathcal{A} . Then we can construct a deterministic finite automaton, of size $\mathcal{O}(|\mathcal{A}| |\mathcal{A}_1| \dots |\mathcal{A}_n|)$, accepting the set \tilde{C} .

3 Preliminary results

From now on, let us fix an abstract pushdown process $\mathcal{P} = \langle Q, \Gamma, \Delta \rangle$ together with a partition $Q_E \cup Q_A$ of Q and a colouring function ρ using a finite set of colours C . Denote respectively by $\mathcal{G} = (V, E)$ and \mathbb{G} the associated abstract pushdown game graph and abstract pushdown parity game.

We can define an automaton with oracles that accepts Éloïse's winning region of the game \mathbb{G} . The oracles of this

automaton are defined using conditional games. For every subset R of Q the game $\mathbb{G}(R)$ played over \mathcal{G} is the *conditional game induced by R over \mathcal{G}* . A play Λ in $\mathbb{G}(R)$ is winning for Éloïse iff one of the following happens:

- In Λ no configuration with an empty stack (*i.e.* of the form (q, \perp)) is visited, and Λ satisfies the parity condition.
- In Λ a configuration with an empty stack is visited and the control state in the first such configuration belongs to R .

More formally, the winning condition in $\mathbb{G}(R)$ is

$$[\Omega_{par} \setminus V^*(Q \times \{\perp\})V^\omega] \cup V^*(R \times \{\perp\})V^\omega$$

For any state q , any stack letter $\gamma \neq \perp$, and any subset $R \subseteq Q$ it follows from Martin's Determinacy theorem that either Éloïse or Abelard has a winning strategy from $(q, \perp\gamma)$ in $\mathbb{G}(R)$. We denote by $\mathcal{R}(q, \gamma)$ the set of subsets R for which Éloïse wins in $\mathbb{G}(R)$ from $(q, \perp\gamma)$:

$$\mathcal{R}(q, \gamma) = \{R \subseteq Q \mid (q, \perp\gamma) \text{ is winning for Éloïse in } \mathbb{G}(R)\}$$

Proposition 3. *Let $\sigma \in (\Gamma \setminus \{\perp\})^*$, $q \in Q$ and $\gamma \in \Gamma \setminus \{\perp\}$. Then Éloïse has a winning strategy in \mathbb{G} from $(q, \perp\sigma\gamma)$ if and only if there exists some $R \in \mathcal{R}(q, \gamma)$ such that $(r, \perp\sigma)$ is winning for Éloïse in \mathbb{G} for every $r \in R$.*

Proof (sketch). Assume Éloïse has a winning strategy from $(q, \perp\sigma\gamma)$ in \mathbb{G} and call it φ . Define R to be the set of all $r \in Q$ such that there is a play $v_0 \cdots v_k(r, \perp\sigma)v_{k+1} \cdots$ where Éloïse respects φ and each v_i for $0 \leq i \leq k$ is of the form $(p, \perp\sigma\sigma')$ for some $\sigma' \neq \varepsilon$. Mimicking φ , Éloïse wins in $\mathbb{G}(R)$ from (q, γ) and hence $R \in \mathcal{R}(q, \gamma)$. Finally, for every $r \in R$ there is a partial play λ_r that starts from $(q, \perp\sigma\gamma)$, where Éloïse respects φ , and that ends in $(r, \perp\sigma)$. Hence $(r, \perp\sigma)$ is also winning for Éloïse in \mathbb{G} .

Conversely, let us assume that there is some $R \in \mathcal{R}(q, \gamma)$ such that $(r, \perp\sigma)$ is winning for Éloïse in \mathbb{G} for every $r \in R$. For every $r \in R$, let us denote by φ_r a winning strategy for Éloïse from $(r, \perp\sigma)$ in \mathbb{G} . Let φ_R be a winning strategy for Éloïse in $\mathbb{G}(R)$ from $(q, \perp\gamma)$. In order to win in \mathbb{G} from $(q, \perp\sigma\gamma)$, Éloïse first mimics φ_R and, if eventually some configuration $(r, \perp\sigma)$ is reached she follows, φ_r from that point onward. \square

Proposition 3 easily implies the following result.

Theorem 1. *Let \mathbb{G} be an abstract pushdown parity game induced by an abstract pushdown process $\mathcal{P} = \langle Q, \Gamma, \Delta \rangle$. Then the set of winning positions in \mathbb{G} for Éloïse (respectively for Abelard) is accepted by an automaton with oracles $\mathcal{A} = (\mathcal{S}, Q, \Gamma, \delta, s_i, \mathcal{O}_1 \cdots \mathcal{O}_n, \text{Acc})$ such that*

- $\mathcal{S} = 2^Q$;

- $s_i = \emptyset$.
- There is an oracle $\mathcal{O}_{p,R}$ for every $p \in Q$ and $R \subseteq Q$, and $\gamma \in \mathcal{O}_{p,R}$ iff $R \in \mathcal{R}(p, \gamma)$ and $\gamma \neq \perp$.
- There is an oracle \mathcal{O}_\perp and $\gamma \in \mathcal{O}_\perp$ iff $\gamma = \perp$.
- Using the oracles, δ is designed such that:
 - From state \emptyset on reading \perp , \mathcal{A} goes to $\{p \mid (p, \perp) \text{ is winning for Éloïse in } \mathbb{G}\}$.
 - From state S on reading γ , \mathcal{A} goes to $\{p \mid S \in \mathcal{R}(p, \gamma)\}$.
- Acc is the identity function.

We will later use Theorem 1 in combination with Remark 2 to prove that the set of winning positions in any higher-order pushdown parity games is regular (see Theorem 4).

4 Reducing the conditional games

The main purpose of this section is to build a new game whose winning region embeds all the information needed to determine the sets $\mathcal{R}(q, \gamma)$. Moreover the underlying game graph no longer uses a stack.

For an infinite play $\Lambda = v_0v_1 \cdots$, let Steps_Λ be the set of indices of positions where no configuration of strictly smaller stack height is visited later in the play. More formally, $\text{Steps}_\Lambda = \{i \in \mathbb{N} \mid \forall j \geq i \text{ } sh(v_j) \geq sh(v_i)\}$, where $sh((q, \perp\gamma_1 \cdots \gamma_n)) = n + 1$. Note that Steps_Λ is always infinite and hence induces a factorisation of the play Λ into finite pieces.

In the factorisation induced by Steps_Λ , a factor $v_i \cdots v_j$ is called a *bump* if $sh(v_j) = sh(v_i)$, called a *Stair* otherwise (that is, if $sh(v_j) = sh(v_i) + 1$).

For any play Λ with $\text{Steps}_\Lambda = \{n_0 < n_1 < \cdots\}$, we can define the sequence $(mcol_i^\Lambda)_{i \geq 0} \in \mathbb{N}^\mathbb{N}$ by setting $mcol_i^\Lambda = \min\{\rho(v_k) \mid n_i \leq k \leq n_{i+1}\}$. This sequence fully characterises the parity condition.

Proposition 4. *For a play Λ , $\Lambda \in \Omega_{par}$ iff $\liminf((mcol_i^\Lambda)_{i \geq 0})$ is even.*

In the sequel, we build a new parity game $\tilde{\mathbb{G}}$ over a new game graph $\tilde{\mathcal{G}} = (\tilde{V}, \tilde{E})$. This new game *simulates* the abstract pushdown game, in the sense that the sequence of visited colours during a correct simulation of some play Λ in \mathbb{G} is exactly the sequence $(mcol_i^\Lambda)_{i \geq 0}$. Moreover, a play in which a player does not correctly simulate the abstract pushdown game is losing for that player. We shall see that the winning regions in $\tilde{\mathbb{G}}$ allow us to compute the sets $\{\gamma \in \Gamma \mid R \in \mathcal{R}(q, \gamma)\}$.

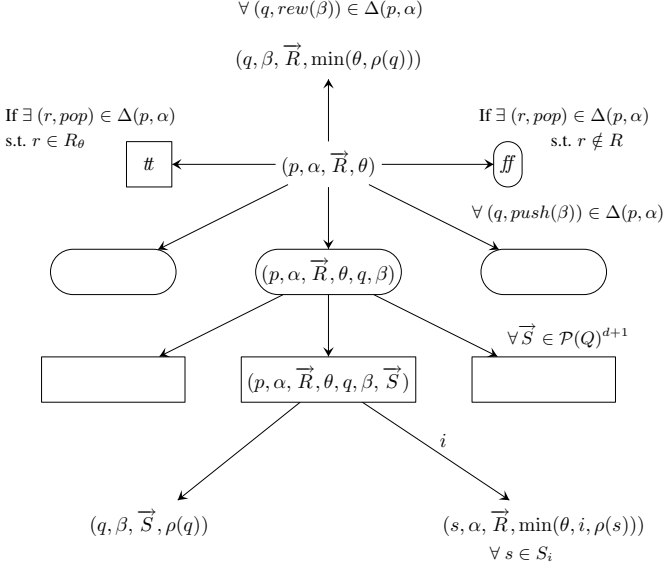


Figure 1. Local structure of $\tilde{\mathcal{G}}$.

Before providing a description of the game graph $\tilde{\mathcal{G}}$, let us consider the following informal description of this simulation game. We aim at simulating a play in the abstract pushdown game from some initial vertex (p_{in}, \perp) . In $\tilde{\mathcal{G}}$ we keep track of only the control state and the top stack symbol of the simulated configuration.

The interesting case is when it is in a control state p with top stack symbol α , and the player owning p wants to push a symbol β onto the stack and change the control state to q . For every strategy of Éloïse, there is a certain set of possible (finite) continuations of the play that will end with popping β (or actually a symbol into which β was rewritten in the meantime) from the stack. We require Éloïse to declare a vector $\vec{S} = (S_0, \dots, S_d)$ of $(d+1)$ subsets of Q , where S_i is the set of all states the game can be in after popping (possibly a rewriting of) β along those plays where in addition the smallest visited colour while (possibly a rewriting of) β was on the stack is i .

Abelard has two choices. He can continue the game by pushing β onto the stack and updating the state (we call this a *pursue* move). Otherwise, he can pick a set S_i and a state $s \in S_i$, and continue the simulation from that state s (we call this a *jump* move). If he does a pursue move, then he remembers the vector \vec{S} claimed by Éloïse; if later on, a pop transition is simulated, the play stops and Éloïse wins if and only if the resulting state is in S_θ where θ is the smallest colour seen in the current level (this information is encoded in the control state, reset after each pursue move and updated after each jump move). If Abelard does a jump

move to a state s in S_i , the currently stored value for θ is updated to $\min(\theta, i, \rho(s))$, which is the smallest colour seen since the current stack level was reached.

There are extra moves to simulate *rew* rules where the top stack element and the value of θ are updated.

Therefore the main vertices of this new game graph are of the form $(p, \alpha, \vec{R}, \theta)$, which are controlled by the player who controls p . Intermediate vertices are used to handle the previously described intermediate steps. The local structure is given in Figure 1 (circled vertices are those controlled by Éloïse). Two special vertices $\#$ and ff are used to simulate pop moves. This game graph is equipped with a colouring function on the vertices and on the edges: vertices of the form $(p, \alpha, \vec{R}, \theta)$ have colour $\rho(p)$, an edge leaving from a vertex $(p, \alpha, \vec{R}, \theta, q, \beta, \vec{S})$ has colour i where i is the colour of the simulated bump. Note that intermediate vertices could be introduced in order to have only colours on vertices. A precise description of the graph is given in the full proof of the following main result.

Theorem 2. *The following holds.*

1. A configuration (p_{in}, \perp) is winning for Éloïse in \mathbb{G} if and only if $(p_{in}, \perp, (\emptyset, \dots, \emptyset), \rho(p_{in}))$ is winning for Éloïse in $\tilde{\mathbb{G}}$.
2. For every $q \in Q$, $\gamma \in \Gamma$ and $R \subseteq Q$, $R \in \mathcal{R}(q, \gamma)$ if and only if $(q, \gamma, (R, \dots, R), \rho(q))$ is winning for Éloïse in $\tilde{\mathbb{G}}$.

Proof (sketch). We concentrate here on the first point, as the second is actually a part of the proof of the first one. For the direct implication, assume that the configuration (p_{in}, \perp) is winning for Éloïse in \mathbb{G} , and let Φ be a corresponding strategy for her.

Using Φ , we define a strategy φ for Éloïse in $\tilde{\mathbb{G}}$ from $(p_{in}, \perp, (\emptyset, \dots, \emptyset), \rho(p_{in}))$. This strategy stores a partial play Λ in \mathbb{G} , that is an element in V^* (where V denotes the set of vertices of \mathcal{G}). At the beginning Λ is initialised to the vertex (p_{in}, \perp) .

By a *round* we mean a factor of a play between two visits through vertices of the form $(p, \alpha, \vec{R}, \theta)$. Both the strategy φ and the update of Λ , are described for a round. When Éloïse has to play from some vertex $(p, \alpha, \vec{R}, \theta)$ she considers the value of $\Phi(\Lambda)$. If it is a pop move then she goes to $\#$ (one proves that this move is always possible). If it equals $(q, \text{rew}(\beta))$, she goes to $(q, \beta, \vec{R}, \min(\theta, \rho(q)))$. Finally, if it is equal to $(q, \text{push}(\beta))$, she goes to $(p, \alpha, \vec{R}, \theta, \beta)$.

From some vertex $(p, \alpha, \vec{R}, \theta, q, \beta)$, Éloïse has to provide a vector $\vec{S} \in \mathcal{P}(Q)^{d+1}$ that describes which states can be reached if β (or one of its successors by top rewriting) is popped, depending on the smallest visited colour in the meantime. In order to define \vec{S} , Éloïse considers the set

of all possible continuations of $\Lambda \cdot (q, \sigma\alpha\beta)$ (where $(p, \sigma\alpha)$ denotes the last vertex of Λ) where she respects her strategy Φ . For each such play, she checks whether some configuration of the form $(s, \sigma\alpha)$ is visited after $\Lambda \cdot (q, \sigma\alpha\beta)$, that is if the β is eventually popped. If it is the case, she considers the first such configuration and the smallest colour i seen in the meantime. For every $i \in \{0, \dots, d\}$, S_i is exactly the set of states $s \in Q$ such that the preceding case happens.

Let $(p, \sigma\alpha)$ be the last vertex in Λ . The memory Λ is updated after each visit to a vertex of the form $(p, \alpha, \vec{R}, \theta)$, we have three cases depending on the kind of the round. If it was simulating a $(q, \text{rew}(\beta))$ action then the updated memory is $\Lambda \cdot (q, \beta\sigma)$. If it was simulating a bump of colour i starting with some action $(q, \text{push}(\beta))$ and ending in a state $s \in S_i$ then the memory becomes Λ extended by $(q, \sigma\alpha\beta)$ followed by a sequence of moves, where Éloïse respects Φ , that ends by popping β and reaches $(s, \sigma\alpha)$ while having i as smallest colour. Finally, if it was simulating a stair starting with a $(q, \text{push}(\beta))$ action, then the updated memory is $\Lambda \cdot (q, \sigma\alpha\beta)$.

Therefore, any partial play λ in $\tilde{\mathbb{G}}$ — in which Eloïse respects her strategy φ — is associated with a partial play Λ in \mathbb{G} . One shows that Éloïse respects Φ in Λ . The same arguments work for an infinite play λ , and the corresponding play Λ is infinite, starts from (p_{in}, \perp) and Éloïse respects Φ in that play. Therefore it is a winning play. Relying on that fact one concludes that φ is winning.

For the converse implication one can reason in a rather similar way by constructing a winning strategy for Abelard in $\tilde{\mathbb{G}}$ from one in \mathbb{G} . \square

From a more constructive proof of Theorem 2 one can construct natural strategies in \mathbb{G} from strategies in $\tilde{\mathbb{G}}$.

Theorem 3. *Assume Éloïse has a winning strategy φ in $\tilde{\mathbb{G}}$ that uses a memory ranging from some set M . Then one can construct an abstract pushdown process \mathcal{T} with output that realises a winning strategy Φ for Éloïse in \mathbb{G} . Moreover the abstract pushdown alphabet used by \mathcal{T} is $\tilde{V} \times M$ and, at any moment in a play where Éloïse respects Φ , the abstract pushdown content of \mathcal{T} has exactly the same height as the one in the current position of the game graph. Finally, if φ is effective the same holds for Φ .*

Remark 3. In the special case of pushdown games, since Γ is finite so is $\tilde{\mathbb{G}}$. Hence in the previous statement, φ can be chosen to be memoryless. Therefore one concludes that for pushdown games one can construct a deterministic pushdown automaton that realises a winning strategy and whose stack is synchronised with the one in the game [37].

5 Uniform solution of higher-order pushdown parity games and strategies

In this section we prove that the winning regions in higher-order pushdown games are regular. The first step is to note the following property.

Property 1. *Let \mathbb{G} be a higher-order pushdown parity game of order- n and let $\tilde{\mathbb{G}}$ be as in Theorem 2. Then $\tilde{\mathbb{G}}$ is a higher-order pushdown parity game of order- $(n - 1)$.*

Proof. One simply needs to consider how the game graph $\tilde{\mathcal{G}}$ is defined. It suffices to make the following observations concerning the local structure given in Figure 1 when \mathbb{G} is played on the transition graph of a pushdown automaton of order- n .

1. For every vertex of the form $(p, \alpha, \vec{R}, \theta)$, $(p, \alpha, \vec{R}, \theta, q, \beta)$ or $(p, \alpha, \vec{R}, \theta, q, \beta, \vec{S})$, α is an order- $(n - 1)$ stack.
2. For every vertex of the form $(p, \alpha, \vec{R}, \theta, q, \beta)$ or $(p, \alpha, \vec{R}, \theta, q, \beta, \vec{S})$, it holds that $\alpha = \beta$.

This implies that any vertex in $\tilde{\mathcal{G}}$ can be seen as a pair formed by a state in a finite set and an order- $(n - 1)$ stack. Then one concludes the proof by checking that the edge relation is the one of an order- $(n - 1)$ pushdown automaton (for the transition to vertices t and ff one can introduce vertices (t, α) and (ff, α) for any order- $(n - 1)$ stack α). \square

Remark 4. The number of states of the higher-order pushdown automaton describing $\tilde{\mathbb{G}}$ is exponential in the number of states of the pushdown automaton describing \mathbb{G} but both games have the same number of colours.

Consider the order-1 case. As $\tilde{\mathbb{G}}$ is finite one can solve it and therefore effectively construct the automaton with oracles as in Theorem 1. As this automaton has a finite input alphabet, using Remark 2, we deduce the following result.

Property 2. [7, 32] *The set of winning position in a pushdown parity game is regular.*

Now, one can iterate this reasoning: applying inductively Property 1 together with Proposition 2 and Theorem 2 easily leads to the desired result.

Theorem 4. *The sets of winning positions in a higher-order pushdown parity game are regular and can be effectively computed. Computing these regions is an n -EXPTIME-complete problem for an order- n pushdown parity game.*

Starting with an order- n pushdown parity game, and applying n times the reduction of Proposition 1, one ends up with a parity game using the same number of colours

over a finite game graph whose size, using Remark 4, is n times exponential in the size of the original order- n pushdown automaton. As solving this latter game is exponential only on the number of colours [22] the global procedure is in n -EXPTIME. The lower bound follows from the fact that deciding the winner in two-player reachability games over order- n pushdown is already n -EXPTIME-hard. A self-contained proof can be found in [9]. The next remark sketches a much simpler proof of this result.

Remark 5. Note that using the reduction of Theorem 2, we can deduce n -EXPTIME-hardness by reduction to the $(n - 1)$ -EXPTIME-hardness of the emptiness problem for order- n pushdown automata [17]. Consider an order- $(n+1)$ pushdown automaton \mathcal{P} over a one-letter input alphabet. The emptiness problem for \mathcal{P} is polynomially equivalent to deciding the winner in the associated one-player reachability game. In the case of a one-player game, the reduction of Theorem 2 can be tailored to yield a reduced game \mathbb{G} which is a two-player game of *polynomial* size w.r.t to \mathcal{P} . Hence we establish that the emptiness problem for order- $(n+1)$ pushdown automata can be polynomially reduced to deciding the winner of two-player order- n game. This proves the n -EXPTIME-hardness of the latter problem.

Remark 6. Theorem 4 generalises the result obtained by Hague and Ong [25] for higher-order pushdown reachability games.

Concerning strategies, we already noted in Remark 3 that a winning strategy can be realised by a pushdown automaton whose stack is synchronised with the one from the game. Reasoning by induction and relying on Theorem 3, one can obtain a similar result for the general case. More precisely, if one defines the shape of an order- k stack to be the stack of obtained by rewriting the stack symbols by a fixed symbol \diamond , we get the following result.

Theorem 5. *Consider an order- k pushdown parity game. Then one can construct, for any player, a deterministic order- k pushdown automaton that realises a winning strategy and whose stack has always the same shape as the one in the game.*

Proof (sketch). The order-1 case was noted in Remark 3. For the general case, we reason by induction using Theorem 3 together with the base case (order-1).

Assume the result is proved at order- k and consider some order- $(k+1)$ game \mathbb{G} . Then using Property 1 we get that $\tilde{\mathbb{G}}$ is a game of order k . By induction hypothesis, there exists a strategy realised by a deterministic order- k pushdown automaton whose stack has always the same shape as the one in $\tilde{\mathbb{G}}$. Now apply Theorem 3: it leads to a strategy in \mathbb{G} realised by an abstract pushdown process whose abstract pushdown content is the product of the vertices in $\tilde{\mathbb{G}}$ together with the memory used to win in $\tilde{\mathbb{G}}$, i.e. the product

of two order- k stacks having the *same* shape. Hence, this product can be thought as a single order- k stack, and therefore the resulting strategy in \mathbb{G} is realised by an automaton using a stack of order- k stacks, in other terms by an order- $(k+1)$ automaton. Finally, the fact that this latter automaton always has the same shape as the current configuration in the game, follows from the induction hypothesis and the height property of the strategy of Theorem 3. \square

Remark 7. Theorem 5 means that the memory needed to win a higher-order pushdown parity game can actually be implemented in the underlying higher-order pushdown automaton defining the game by enriching its stack alphabet.

6 Global μ -calculus model-checking

Given a vertex (state) s in a state-transition graph \mathcal{K} and a formula φ , the *model-checking problem* asks whether $\mathcal{K}, s \models \varphi$ holds (in words, whether the formula φ holds at s in the structure \mathcal{K}). The *global model-checking problem* is the task of computing (if possible) a finite representation of the set $\|\mathcal{K}\|_\varphi = \{s \mid \mathcal{K}, s \models \varphi\}$ of vertices in \mathcal{K} that satisfy a given formula φ .

In the following we tackle the global model checking problem when φ is a formula of the modal μ -calculus [29, 2]. We consider two kinds of structures for \mathcal{K} : configuration graphs of higher-order pushdown automata, and trees generated by higher-order *safe* recursion schemes.

Global model-checking of configuration graphs of higher-order pushdown automata

The μ -calculus global model-checking problem for families of graphs closed under Cartesian product with finite graphs is well-known to be equivalent to the solvability of parity games over the same class. Hence Theorem 4 implies the following characterisation of μ -calculus definable sets over higher-order pushdown graphs:

Theorem 6. *The μ -calculus definable sets over configuration graphs of higher-order pushdown automata are exactly the regular sets of configurations.*

Proof (sketch). Given a μ -calculus formula φ , a classical construction [16] leads to a finite game graph \mathcal{G}_φ whose vertices are all possible subformulas of φ , together with a colouring function ρ . Then, for any structure \mathcal{K} and any vertex v in \mathcal{K} , the following holds: $\mathcal{K}, v \models \varphi$ iff Éloïse wins from (v, φ) in the parity game induced by $\mathcal{K} \times \mathcal{G}_\varphi$. The underlying game graph of this game is the cartesian product of \mathcal{K} and \mathcal{G}_φ , the partition of its vertices as well as the associated colouring function being inherited from \mathcal{G}_φ . Then, it follows that $\|\mathcal{K}\|_\varphi = \{v \mid (v, \varphi) \in W_{\mathbf{E}}\}$ where $W_{\mathbf{E}}$ denotes the winning region for Éloïse in the previous game.

The previous construction actually does not rely on \mathcal{K} being finite and can be used, for instance, to solve the μ -calculus model-checking problem against pushdown processes [37]. In the special case where \mathcal{K} is the transition graph of some order- k pushdown automaton, the game graph $\mathcal{K} \times \mathcal{G}_\varphi$ is also the transition graph of some order- k pushdown automaton and the partition of its vertices as well as its associated colouring function only depend on the control states. Hence, the solution of the global μ -calculus model-checking problem for such a graph can be deduced from the (regular) winning region of a higher-order parity pushdown game by a simple operation preserving regularity. Hence μ -calculus definable sets over configuration graphs of higher-order pushdown automata are regular. \square

Remark 8. The preceding Theorem generalises a result of Bouajjani and Meyer in [4] for a weaker logic (i.e. the $E(U, X)$ fragment of CTL) over weaker structures (i.e. higher-order pushdown automata with a single control state).

Global model-checking of trees generated by higher-order safe recursion schemes

Fix a (ranked) alphabet Σ . *Types* are generated from a base type using the arrow constructor \rightarrow . A higher-order (deterministic) recursion scheme is a finite set of equations of the form $F x_1 \cdots x_n = e$, where F is a typed non-terminal, each x_i is a typed variable, and e is an applicative term constructed from the non-terminals, terminals (which are the Σ -symbols), and variables x_1, \dots, x_n . The scheme is said to be *order- k* if the highest order of the non-terminals is k . We use recursion schemes here as generators of possibly infinite term-trees. The ranked *tree* generated by a recursion scheme is defined to be the (possibly infinite) term-tree built up from the terminal symbols by applying the equations *qua* rewrite rules, replacing the formal parameters by the actual parameters, starting from the initial non-terminal. We refer the reader to [28] for a precise description of the preceding, and for the meaning of *safety* which is a syntactic constraint; here we rely on the following characterisation:

Theorem 7. [28] *For each $k \geq 0$, the ranked trees generated by safe order- k recursion schemes coincide with the ε -closure of the unravelling of the configuration graphs of deterministic order- k pushdown automata.*

Consider a higher-order deterministic pushdown automaton labelled by $\Sigma \cup \{\varepsilon\}$. After unravelling its configuration graph from the initial configuration, the operation of ε -closure (see [33]) first adds an a -labelled edge from c_1 to c_2 whenever there is a path from c_1 to c_2 that traces out the word $a\varepsilon^*$ and c_2 is not the source-vertex of an ε -labelled edge, and then removes the source-vertex of every ε -labelled edge. The resultant graph is a tree.

A node in this tree is naturally represented by the word over Σ labelling the path from the root to this node. Hence, every μ -calculus formula over such a tree induces a language in Σ^* i.e. the set of words labelling a path from the root to a node satisfying the formula. These languages can be characterised as follows:

Theorem 8. *Let $k \geq 0$ and let \mathcal{T} be the ranked tree generated by a given order- k safe recursion scheme. The μ -calculus definable sets of \mathcal{T} -nodes are recognisable by a deterministic order- k pushdown automaton.*

Proof (sketch). Let \mathcal{C} be the configuration graph of a deterministic order- k $\Sigma \cup \{\varepsilon\}$ -labelled pushdown automaton \mathcal{P} , and let U be the unravelling of \mathcal{C} from the initial configuration c_0 , such that the ε -closure of U is isomorphic to \mathcal{T} . Let φ be a μ -calculus formula. A node in U satisfies φ if and only if it corresponds to a path ending in a vertex of \mathcal{C} that satisfies φ . By Theorem 6, the set F of configurations of \mathcal{C} satisfying φ in \mathcal{C} is regular. Moreover it is easy to see that the set E of configurations of \mathcal{C} which are not the source-vertex of a ε -labelled edge in \mathcal{C} is also regular. A node in \mathcal{T} at the end of a path from the root labelled by w satisfies φ in \mathcal{T} if and only if w is accepted by the labelled pushdown automaton \mathcal{P} from the initial configuration to the regular set $E \cap F$. By Proposition 1, the language consisting of such w is accepted by an order- k pushdown automaton. \square

7 Discussions

There are a number of further directions:

- Can the quite general class of stack data games together with Theorem 2 be used to prove the decidability of games over new structures (e.g. allowing arithmetic on the stack)?
- Regular stack properties allow assertions over the stack contents. Since information regarding regular tests may be encoded in the control states and stack of an abstract pushdown process, the definable sets are again regular. Properties of this kind have been shown to have applications in inter-procedural data-flow analysis [18] and security [27] for the case of order-1 pushdown systems. One such security property, implemented in Java and .Net [21, 5], allows the programmer to decorate code with permission checks. These checks require that all callers on the stack have sufficient privileges to proceed. Do regular stack properties have similar applications for other structures encodable as abstract pushdown processes?
- The notion of regularity used in this article can capture the μ -calculus definable sets of configurations. As mentioned previously, MSO-definable sets can be captured by a stronger notion of regularity introduced in [10, 20]. Hence a natural question is the decidability of the following problem: given a strong regular set (in the sense of [10, 20]),

decide whether it is regular (in the sense of this article). Moreover, it is known from [20] that positional winning strategies of higher-order pushdown parity game can be described using strong regular sets (*i.e.* for each transition, the set of configurations from which the strategy plays the transition is a strong regular set) and a k -Exptime algorithm was recently given in [13]. Is it possible to describe positional winning strategies using the weak notion of regularity of this article?

- *Order- k collapsible pushdown automata* (CPDA) are a generalisation of order- k pushdown automata in which each symbol in the k -stack has a link to a stack below it. As generators of ranked trees, they are equi-expressive with (arbitrary) order- k recursion schemes [24]. A natural question is to consider configuration graphs of CPDA, and compute a finite representation of the μ -calculus definable vertex-sets thereof.

An alternative method for computing the winning regions of a higher-order pushdown parity game develops the order-1 saturation techniques introduced by Bouajjani *et al.* [3] and Finkel *et al.* [19] and generalised to Büchi games by Cachat [8]. This algorithm uses a characterisation of a game's winning regions as a series of greatest and least fixed points. Following this characterisation, a small initial automaton, accepting higher-order stacks, is expanded until the winning region has been computed. Because this technique does not require an n -exponential reduction to a finite state game, it is possible that the state-explosion problem may be avoided in some, low-order, cases.

References

- [1] A. Aho. Indexed grammars, an extension of context-free grammars. *Journal of the Association for Computing Machinery (ACM)*, 15:647–671, 1968.
- [2] A. Arnold and D. Niwiński. *Rudiments of μ -calculus*, volume 146 of *Studies in Logic and the Foundations of Mathematics*. Elsevier, 2001.
- [3] A. Bouajjani, J. Esparza, and O. Maler. Reachability analysis of pushdown automata: Application to model-checking. In *International Conference on Concurrency Theory*, pages 135–150, 1997.
- [4] A. Bouajjani and A. Meyer. Symbolic reachability analysis of higher-order context-free processes. In *FST&TCS 2004: Foundations of Software Technology and Theoretical Computer Science, 24th International Conference, Proceedings*, volume 3328 of *Lecture Notes in Computer Science*, pages 135–147. Springer, 2004.
- [5] D. Box and T. Pattison. *Essential .NET: The Common Language Runtime*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 2002.
- [6] T. Cachat. Uniform solution of parity games on prefix-recognizable graphs. In Antonin Kucera and Richard Mayr, editors, *4th International Workshop on Verification of Infinite-State Systems, Infinity 2002*, volume 68 of *Electronic Notes in Theoretical Computer Science*. Elsevier Science Publishers, 2002.
- [7] T. Cachat. Higher order pushdown automata, the Caucal hierarchy of graphs and parity games. In *30th International Colloquium on Automata, Languages, and Programming, ICALP 2003, Proceedings*, volume 2719 of *Lecture Notes in Computer Science*, pages 556–569. Springer, 2003.
- [8] T. Cachat. *Games on Pushdown Graphs and Extensions*. PhD thesis, RWTH Aachen, 2003.
- [9] T. Cachat and I. Walukiewicz. The complexity of games on higher order pushdown automata, 2007. arXiv:0705.0262v1.
- [10] A. Carayol. Regular sets of higher-order pushdown stacks. In *Mathematical Foundations of Computer Science 2005, 30th International Symposium, MFCS 2005, Proceedings*, volume 3618 of *Lecture Notes in Computer Science*, pages 168–179. Springer, 2005.
- [11] A. Carayol, M. Hague, A. Meyer, C.-H. L. Ong, and O. Serre. Winning regions of higher-order pushdown games. Technical report, 2007. Preprint, 20 pages, downloadable from <http://www.liafa.jussieu.fr/~serre/>
- [12] A. Carayol and S. Wöhrle. The Caucal hierarchy of infinite graphs in terms of logic and higher-order pushdown automata. In *FST&TCS 2003: Foundations of Software Technology and Theoretical Computer Science*, volume 2914 of *Lecture Notes in Computer Science*, pages 112–123. Springer, 2003.
- [13] A. Carayol and M. Slaats. Positional strategies for higher-order pushdown parity games. Submitted, 2008.
- [14] D. Caucal. On infinite terms having a decidable monadic theory. In *Proceedings of the 27th International Symposium on Mathematical Foundations of Computer Science (MFCS 2002)*, volume 2420 of *Lecture Notes in Computer Science*, pages 165–176. Springer, 2002.

- [15] W. Damm. The IO- and OI-hierarchies. *Theoretical Computer Science*, 20:95–207, 1982.
- [16] E. Allen Emerson and Charanjit S. Jutla. Tree automata, mu-calculus and determinacy (extended abstract). In *32nd Annual Symposium on Foundations of Computer Science, FoCS 1991, Proceedings*, pages 368–377. IEEE Computer Society Press, 1991.
- [17] J. Engelfriet. Iterated stack automata and complexity classes. *Information and Computation*, pages 21–75, 1991.
- [18] J. Esparza, A. Kučera, and S. Schwoon. Model-checking LTL with regular valuations for pushdown systems. In *Proceedings of Tools and Algorithms for the Construction and Analysis of Systems, 7th International Conference, TACAS 2001*, number 2215 in Lecture Notes in Computer Science, pages 306–339, 2001.
- [19] A. Finkel, B. Willems, and P. Wolper. A direct symbolic approach to model checking pushdown systems. In *Proc. 2nd Int. Workshop on Verification of Infinite State Systems (INFINITY’97)*, Bologna, Italy, July 11–12, 1997, volume 9 of *Electronic Notes in Theor. Comp. Sci.* Elsevier, 1997.
- [20] S. Fratani. *Automates à pile de piles ... de piles*. PhD thesis, Université de Bordeaux I, 2005.
- [21] L. Gong. *Inside Java 2 platform security architecture, API design, and implementation*. Addison-Wesley Longman Publishing Co., Inc., Boston, MA, USA, 1999.
- [22] E. Grädel, W. Thomas, and T. Wilke, editors. *Automata, Logics, and Infinite Games: A Guide to Current Research [outcome of a Dagstuhl seminar, February 2001]*, volume 2500 of *Lecture Notes in Computer Science*. Springer, 2002.
- [23] M. Hague. *Global Model Checking of Higher-Order Pushdown Systems*. PhD thesis, University of Oxford, 2008.
- [24] M. Hague, A. S. Murawski, C.-H. L. Ong, and O. Serre. Collapsible pushdown automata and recursion schemes. In *Proceedings of the 23rd Annual IEEE Symposium on Logic in Computer Science (LICS’08)*. Computer Society Press, 2008. To appear.
- [25] M. Hague and C.-H. L. Ong. Symbolic backwards-reachability analysis for higher-order pushdown systems. In *Foundations of Software Science and Computational Structures, 10th International Conference, FOSSACS 2007, Proceedings*, volume 4423 of *Lecture Notes in Computer Science*, pages 213–227. Springer, 2007.
- [26] D. Janin and I. Walukiewicz. On the expressive completeness of the propositional mu-calculus with respect to monadic second order logic. In *Proceedings of the 7th International Conference on Concurrency Theory (CONCUR 1996)*, volume 1119 of *Lecture Notes in Computer Science*, pages 263–277. Springer, 1996.
- [27] T. P. Jensen, D. Le Métayer, and T. Thorn. Verification of control flow based security properties. In *IEEE Symposium on Security and Privacy*, pages 89–103, 1999.
- [28] T. Knapik, D. Niwiński, and P. Urzyczyn. Higher-order pushdown trees are easy. In *Proceedings of the 5th International Conference on Foundations of Software Science and Computational Structures (FOSSACS’02)*, volume 2303 of *Lecture Notes in Computer Sciences*, pages 205–222. Springer-Verlag, 2002.
- [29] D. Kozen. Results on the propositional mu-calculus. *Theoretical Computer Science*, 27:333–354, 1983.
- [30] D. A. Martin. Borel determinacy. *Annals of Mathematics*, 102(363–371), 1975.
- [31] A. N. Maslov. Multi-level stack automata. *Problems Information Transmission*, 12:38–43, 1976.
- [32] O. Serre. Note on winning positions on pushdown games with omega-regular winning conditions. *Information Processing Letters*, 85:285–291, 2003.
- [33] C. Stirling. Decidability of bisimulation equivalence for pushdown processes. Technical Report EDI-INF-RR-0005, School of Informatics, University of Edinburgh, 2000.
- [34] W. Thomas. On the synthesis of strategies in infinite games. In *8th Annual Symposium on Theoretical Aspects of Computer Science, STACS 1995, Proceedings*, volume 900 of *Lecture Notes in Computer Science*, pages 1–13. Springer, 1995.
- [35] W. Thomas. Constructing infinite graphs with a decidable mso-theory. In *Proceedings of Mathematical Foundations of Computer Science 2003, 28th International Symposium, MFCS 2003*, volume 2747 of LNCS, pages 113–124. Springer Verlag, 2003.
- [36] M. Y. Vardi. Reasoning about the past with two-way automata. In *Proceedings of the 25th International Colloquium on Automata, Languages and Programming*, pages 628–641, 1998.

- [37] I. Walukiewicz. Pushdown processes: games and model checking. *Information and Computation*, 157:234–263, 2000.
- [38] I. Walukiewicz. A landscape with games in the background. In *Proceedings of the 19th Annual IEEE Symposium on Logic in Computer Science (LICS'04)*, pages 356–366. Computer Society Press, 2004.
- [39] W. Zielonka. Infinite games on finitely coloured graphs with applications to automata on infinite trees. *Theoretical Computer Science*, 200(1-2):135–183, 1998.