# DReX: A Declarative Language for Efficiently Evaluating Regular String Transformations [*]

Rajeev Alur    Loris D'Antoni    Mukund Raghothaman

University of Pennsylvania

{alur, lorisdan, rmukund}@seas.upenn.edu

## Abstract

We present DReX, a declarative language that can express all regular string-to-string transformations, and can still be efficiently evaluated. The class of regular string transformations has a robust theoretical foundation including multiple characterizations, closure properties, and decidable analysis questions, and admits a number of string operations such as insertion, deletion, substring swap, and reversal. Recent research has led to a characterization of regular string transformations using a primitive set of function combinators analogous to the definition of regular languages using regular expressions. While these combinators form the basis for the language DReX proposed in this paper, our main technical focus is on the complexity of evaluating the output of a DReX program on a given input string. It turns out that the natural evaluation algorithm involves dynamic programming, leading to complexity that is cubic in the length of the input string. Our main contribution is identifying a consistency restriction on the use of combinators in DReX programs, and a single-pass evaluation algorithm for consistent programs with time complexity that is linear in the length of the input string and polynomial in the size of the program. We show that the consistency restriction does not limit the expressiveness, and whether a DReX program is consistent can be checked efficiently. We report on a prototype implementation, and evaluate it using a representative set of text processing tasks.

***Categories and Subject Descriptors*** D.3.2 [*Language Classifications*]: Specialized application languages; F.1.1 [*Theory of Computation*]: Models of Computation, Automata

***Keywords*** DReX, string transformations, declarative languages

## 1. Introduction

Programs that transform plain text are ubiquitous and used for many different tasks, from reformatting documents to translating data between different formats. String specific utilities such as sed, AWK, and Perl have been used to query and reformat text files for many years. Since these tools are Turing complete they can express very complex transformations, however this comes at the cost of not being amenable to algorithmic analysis. To address this issue, restricted languages have been proposed in the context of verification of string sanitizers [26] and string coders [13], and for the analysis and optimization of list-manipulating programs [15]. These languages build on variants of finite-state transducers, which are automata-based representations of programs mapping strings to strings, and each of these languages supports different algorithmic analyses that are enabled by the properties of the underlying transducer model. Due to the focus on analyzability, expressiveness is a limiting factor in all such languages and many programs, in particular those that reorder input chunks, cannot be represented. Moreover, these languages are not declarative and their semantics are tightly coupled to the transducer model, forcing the programmer to reason in terms of finite state machines and process the input left-to-right.

In the theory of string-to-string transformations the class of regular string transformations is a robust class that strikes a balance between decidability and expressiveness. In particular this class captures transformations that involve reordering of input chunks, it is closed under composition [9], it has decidable equivalence [21], and has several equivalent characterizations, such as one-way transducers with a finite set of write-only registers [1], two-way transducers [16] and monadic second-order definable graph transformations [10]. Recently Alur et al [3] proposed a set of combinators that captures the class of regular string transformations. In [3] the focus is on expressiveness and the paper does not provide an efficient procedure to evaluate programs written with these combinators. Efficient evaluation of such programs is the main focus of this paper.

Starting with the combinators presented in [3], we develop DReX, an expressive declarative language to describe string transformations. The base combinator of DReX, $\varphi \mapsto d$, maps any character $a$ that satisfies the predicate $\varphi$ to the string $d(a)$. This combinator symbolically extends the one proposed in [3] with predicates and can therefore succinctly model strings over large (and potentially infinite) alphabets, such as Unicode. The other combinators supported by DReX are: (a) $\mathtt{split}(f,g)$ that unambiguously splits the input string into two parts and outputs the concatenation of the results obtained using $f$ on the first part and $g$ on the second part; (b) $\mathtt{iterate}(f)$ that unambiguously splits the input string into multiple parts and outputs the concatenation of evaluating $f$ on each of such parts; (c) $\mathtt{combine}(f,g)$ that applies both $f$ and $g$ to the input string and concatenates the obtained results; (d) the conditional $f \mathtt{\ else\ } g$ that first tries to apply $f$ to the input and if $f$ cannot be applied it applies $g$; (e) $\mathtt{chain}(f,R)$ that unambiguously splits the string into multiple parts $\sigma = \sigma_1 \ldots \sigma_n$ each belonging to the language described by the regular expression $R$, applies $f$ to every two pair of adjacent chunks $\sigma_i \sigma_{i+1}$, and finally concatenates these results. In order to model operations such as reversing a string, the

operators split, iteration and chained sum also have a *left-additive* version in which the outputs computed on each split of the string are concatenated in reverse order.

A straightforward algorithm to evaluate DReX programs involves "operationalizing" the semantics, i.e. use dynamic programming and evaluate each sub-program on each substring of the input. Unfortunately, this algorithm takes time cubic in the length of the input string, and does not scale to strings longer than approximately a thousand characters. Because of the analogy between DReX operators (split sum, conditionals, iteration, etc.) and regular expressions (concatenation, union, Kleene-*, etc.) one approach is to construct an automaton model for evaluating DReX programs similarly to the approach taken to evaluate regular expressions. This is not simple because of various reasons, such as: (*a*) the conditional operator, $f$ `else` $g$, applies the transformation $g$ to the input only if the input string is not accepted by $f$. To check whether a string is in the complement of the domain of $f$ one needs to determinize the domain automaton and this is an exponential time operation; and (*b*) the operator `combine`$(f, g)$ is only defined on the intersection of the domains of $f$ and $g$. Repeating automata intersections multiple times also causes an exponential blow-up.

The main technical challenge is to identify a fragment of DReX which does not sacrifice expressiveness, and still permits "fast" evaluation algorithms. We call this subset of DReX the *consistent* fragment. Intuitively, we require each operator to admit unambiguous parsing, and limit the operators' ability to express the complement and intersection of languages. For example, `split`$(f, g)$ is consistent iff the domains of $f$ and of $g$ are unambiguously concatenable, i.e. there is no string with multiple viable splits. In the case of a conditional $f$ `else` $g$, the domains of $f$ and $g$ are required to be disjoint, making the complementation of the domain of $f$ unnecessary. Similarly, for the operator `combine`$(f, g)$, we require the domains of $f$ and $g$ to be identical, so that the domain of the entire program is equal to the domain of its sub-expressions and no language intersection is required. For the `chain`$(f, R)$ operator to be consistent, the language $R^*$ is required to be unambiguous, and $f$ is required to be a `split` operator for which both the left and right hand sides are exactly defined on the language described by $R$. We show that consistency of a DReX program can be efficiently determined.

We present an algorithm that evaluates a consistent DReX program $f$ on an input string $\sigma$, in time polynomial in the size of $f$ and linear in the length of $\sigma$. Intuitively, we construct a machine for each sub-program which reads the input in a single left-to-right pass. Each machine keeps track of potential parse trees of $\sigma$ as multiple *threads*, and updates the threads on reading each input symbol. The goal of the algorithm is to have a number of threads that is linear in the size of the program but does not depend on the length of the input string. This bound is achieved using the consistency requirements to eagerly kill threads whenever they become inactive. For example, the machine for `split`$(f, g)$ outputs a result as soon as it discovers a single viable split of the input string $\sigma$, since the consistency rules guarantee the absence of any other split. If the program were not consistent the machine would need to delay the output causing the number of alive threads to depend on the length of the string. Similarly, the machine for $f$ `else` $g$ can output the results of $g$ whenever this is defined, because of the requirement that $f$ and $g$ have disjoint domains and cannot simultaneously emit results. We also present a dynamic programming algorithm that can handle all DReX programs, but is limited by its cubic time complexity in the length of the input string.

We implemented our evaluation and consistency-checking algorithms and evaluated them on several text transformations: deletion of comments from a program, insertion of quotes around words, tag extraction from XML documents, reversing dictionaries, and the reordering and aligning of misplaced fields in BibTeX files. The evaluation algorithm for consistent DReX scales to large inputs (less than for 8 seconds for $100,000$ characters), while the dynamic programming algorithm, due to the cubic complexity in the size of the input, does not scale in practice (more than 60 seconds for $5,000$ characters) and therefore has limited applicability. Finally, the consistency-checking algorithm is very fast in practice (less than $0.6$ seconds for programs of size $\approx 3,600$ subexpressions), and it is also very helpful in identifying sources of ambiguity in the implemented programs.

In summary we offer the following contributions.

1. DReX, a language for describing string transformations that extends the combinators proposed in [3] to model strings over arbitrary sorts, and consistent DReX, a fragment of DReX that admits efficient evaluation without sacrificing expressiveness (section 2);

2. an algorithm for evaluating consistent DReX programs in a single left-to-right pass that is linear in the size of the input string and polynomial in the size of the program (section 3);

3. a dynamic programming algorithm for evaluating unrestricted DReX programs that has cubic time complexity in the length of the input string (section 4);

4. a proof that adding a composition operator to DReX causes the evaluation problem to become PSPACE-complete, and the dynamic programming algorithm to run in time exponential in the size of the program (section 4); and

5. an implementation of DReX together with an evaluation of our algorithms on practical string transformations (section 5).

## 2. The Syntax and Semantics of DReX

### 2.1 Regular combinators for string transformations

Given a character $a \in \Sigma$, and an output string $d \in \Gamma^*$, the function $a \mapsto d : \Sigma^* \to \Gamma^*_\perp$ maps the single-character input string $\sigma = a$ to the output $d$, and is undefined for all other inputs:[1]

$$[\![a \mapsto d]\!](\sigma) = \begin{cases} d & \text{if } \sigma = a, \text{ and} \\ \perp & \text{otherwise.} \end{cases}$$

Another basic function is $\epsilon \mapsto d$ which maps the empty string $\epsilon$ to the output $d \in \Gamma^*$, and is undefined everywhere else. The final basic function `bottom` is undefined for all input strings.

The split sum operators are the counterparts of concatenation in regular expressions. Given an input string $\sigma$, if there exists a unique split $\sigma = \sigma_1 \sigma_2$, such that both $[\![f_1]\!](\sigma_1)$ and $[\![f_2]\!](\sigma_2)$ are defined, then

$$[\![\texttt{split}(f_1, f_2)]\!](\sigma) = [\![f_1]\!](\sigma_1)[\![f_2]\!](\sigma_2), \text{ and}$$
$$[\![\texttt{left–split}(f_1, f_2)]\!](\sigma) = [\![f_2]\!](\sigma_2)[\![f_1]\!](\sigma_1).$$

For all other inputs (where there is either no split, or multiple viable splits), both functions are undefined. Note the insistence on a unique parse tree — this is so that programs define functions, rather than relations.

Given two DReX functions $f_1$ and $f_2$, the function $f_1$ `else` $f_2$ first tries to apply $f_1$, and if this fails, applies $f_2$:

$$[\![f_1 \texttt{ else } f_2]\!](\sigma) = \begin{cases} [\![f_1]\!](\sigma) & \text{if } [\![f_1]\!](\sigma) \neq \perp, \text{ and} \\ [\![f_2]\!](\sigma) & \text{otherwise.} \end{cases}$$

This is the unambiguous counterpart of the union operator of traditional regular expressions.

---

[1] We adopt the convention of saying $f(x) = \perp$ when $f$ is undefined for the input $x$, and write $A_\perp$ for $A \cup \{\perp\}$, when $\perp \notin A$.

Similarly, if both $[\![f_1]\!](\sigma)$ and $[\![f_2]\!](\sigma)$ are defined, then

$$[\![\texttt{combine}(f_1, f_2)]\!](\sigma) = [\![f_1]\!](\sigma)[\![f_2]\!](\sigma).$$

If either function is undefined for the input $\sigma$, $\texttt{combine}(f_1, f_2)$ is undefined as well. This combinator can be used to make multiple passes over the input string, and a typical example would be the function that copies the input string twice: $\sigma$ transformed into $\sigma\sigma$. In terms of the input domain, the operator $\texttt{combine}$ is the counterpart of intersection in regular languages, and is necessary to achieve expressive parity with regular string transformations because of the non-commutativity of string concatenation.

If $f$ is a DReX program, and the input string $\sigma$ can be uniquely split into substrings $\sigma = \sigma_1\sigma_2\dots\sigma_n$, with $n \geq 0$, and such that $[\![f]\!](\sigma_i) \neq \bot$, for each $i$, then

$$[\![\texttt{iterate}(f)]\!](\sigma) = [\![f]\!](\sigma_1)[\![f]\!](\sigma_2)\dots[\![f]\!](\sigma_n), \text{ and}$$
$$[\![\texttt{left--iterate}(f)]\!](\sigma) = [\![f]\!](\sigma_n)[\![f]\!](\sigma_{n-1})\dots[\![f]\!](\sigma_1).$$

Otherwise (if the input $\sigma$ cannot be split, or if multiple viable splits exist), then both iterated sums are undefined. This is the counterpart of Kleene-* of regular expressions.

The chained sum operator allows us to "mix" outputs produced by different parts of the input string. This is a new operator, without a regular expression counterpart, and is necessary for expressive completeness. Let $R$ be a regular expression that defines the language $[\![R]\!] = L$, and $f$ be a DReX program. Given an input $\sigma$, if there is a unique split $\sigma = \sigma_1\sigma_2\dots\sigma_n$, such that $\sigma_i \in L$ for each $i$, then, if $n \geq 2$

$$[\![\texttt{chain}(f, R)]\!](\sigma) = [\![f]\!](\sigma_1\sigma_2)[\![f]\!](\sigma_2\sigma_3)\dots$$
$$[\![f]\!](\sigma_{n-1}\sigma_n), \text{ and}$$
$$[\![\texttt{left--chain}(f, R)]\!](\sigma) = [\![f]\!](\sigma_{n-1}\sigma_n)[\![f]\!](\sigma_{n-2}\sigma_{n-1})\dots$$
$$[\![f]\!](\sigma_1\sigma_2).$$

For notational convenience, we treat $\sigma\bot = \bot\sigma = \bot$, and so if $[\![f]\!](\sigma_i\sigma_{i+1})$ is undefined for any $i$, both functions are undefined. Furthermore, if a unique split of the input string $\sigma$ does not exist, both the chained and left-chained sums are undefined. Notice that the regular expression $R$ in $\texttt{chain}(f, R)$ defines the split of the input string, and $f$ is applied to each pair of adjacent splits.

The final operator is function composition. If $f_1$ and $f_2$ are DReX programs such that $[\![f_1]\!] : \Sigma^* \to \Gamma^*_\bot$, and $[\![f_2]\!] : \Gamma^* \to \Lambda^*_\bot$, are partial functions, $\texttt{compose}(f_1, f_2)$ is defined as

$$[\![\texttt{compose}(f_1, f_2)]\!](\sigma) = [\![f_2]\!]([\![f_1]\!](\sigma)),$$

with the notational convention that $[\![f_1]\!](\bot) = \bot$.

Recall that regular string transformations can be defined in multiple equivalent ways: as two-way finite state transducers, as one-way streaming string transducers, and as MSO-definable graph transformations. We summarize the main result of [3]:

**Theorem 1** (Expressive completeness). *For every finite input alphabet $\Sigma$, and output alphabet $\Gamma$, every regular string transformation $f : \Sigma^* \to \Gamma^*_\bot$ can be expressed by a DReX program.*

More precisely, when we include the chained sum, function composition is unnecessary for expressive completeness, while the chained sum can itself be expressed using function composition, and so if composition is included, the chained sum is unnecessary.

## 2.2 Character sorts and predicates

Consider the basic combinator $a \mapsto d$ we described in the previous subsection, which maps the input $\sigma = a$ to the output $d$. For large alphabets, such as the set of all Unicode characters, this approach of explicitly mentioning each character does not scale. Basic transformations in DReX may therefore also reference symbolic predicates and character functions, as we will now describe. This is inspired by the recent development of symbolic transducers [26], which has proved to be useful in several practical applications.

Let $\Sigma, \Gamma, \dots$ be a collection of character sorts. For each character sort $\Sigma$, we pick a (possibly infinite) collection of predicates $P_\Sigma$ such that (a) $P_\Sigma$ is closed under the standard boolean operations: for each $\varphi, \psi \in P_\Sigma, \neg\varphi, \varphi \wedge \psi, \varphi \vee \psi \in P_\Sigma$, and (b) the satisfiability of predicates is decidable: given $\varphi \in P_\Sigma$, whether there exists an $x \in \Sigma$ such that $\varphi(x)$ holds is decidable.

A simple example is the sort $\Sigma_2 = \{a, b\}$ together with the set of predicates $P_{\Sigma_2} = \{x = a, x = b, \texttt{true}, \texttt{false}\}$. Another example is the set of all integers $\mathbb{Z}$, and with $P_\mathbb{Z} = \{\texttt{odd}(x), \texttt{even}(x), \texttt{true}, \texttt{false}\}$. We will write $\mathbb{U}$ for the set of all Unicode characters, with the various character properties $P_\mathbb{U} = \{\texttt{uppercase}(x), \texttt{digit}(x), \dots\}$.

If $\varphi \in P_\Sigma$ and $d = [d_1, d_2, \dots, d_k]$ is a list of character transformations, i.e. $d_i : \Sigma \to \Gamma$, then $\varphi \mapsto d$ is a basic transformation which maps every single-character string $\sigma$ which satisfies $\varphi(\sigma)$ to the output string $d_1(\sigma)d_2(\sigma)\dots d_k(\sigma)$, and is undefined for all other strings.

For example, the function $\texttt{uppercase}(x) \mapsto \texttt{tolowercase}(x)$ transforms every upper-case Unicode character to lower-case, while the function $\texttt{uppercase}(x) \mapsto xx$ outputs two copies of an upper-case character. The function $x \geq 0 \mapsto x - 1$ transforms a non-negative integer by subtracting one from it. Given an input digit $x \in [2\text{-}9]$, the function $x \in [2\text{-}9] \mapsto x - 2$ subtracts 2 from it.

Note that the basic symbolic transformations can still only operate on individual characters in isolation, and cannot relate properties of adjacent characters. For example, we do not allow transformations such as $[x > 0, y > x] \mapsto x, y$, which outputs two consecutive symbols $x$ and $y$, if $x > 0$ and $y > x$. It is known that allowing such "multi-character predicates" makes several analysis questions undecidable [12].

## 2.3 Consistent DReX programs

We now define consistent DReX, a restricted class which still captures all regular string transformations but for which we can provide an efficient evaluation algorithm (section 3). Intuitively, we restrict each operator to only allow unambiguous parsing, and limit the operators' ability to express expensive automata operations such as intersection and complement. Since the purpose of the consistency rules is for the correctness of the evaluation algorithm, we defer their motivation to subsection 3.5.

### 2.3.1 Consistent unambiguous regular expressions

The consistency rules we propose are based on the notion of consistent unambiguous regular expression (CURE). CUREs are similar to conventional regular expressions, but with the additional guarantee that all matched strings have unique parse trees. Unambiguous regular expressions have been studied in the literature [7, 8, 25] — we explicitly qualify them as consistent here to emphasize that there are no strings with multiple parse trees. They are defined inductively as follows:

1. $\bot$ and $\epsilon$ are CUREs. The language associated with $\bot$ is the empty set $[\![\bot]\!] = \emptyset$, and the language associated with $\epsilon$ is the singleton $[\![\epsilon]\!] = \{\epsilon\}$.

2. For each satisfiable predicate $\varphi \in P_\Sigma$, $\varphi$ is a CURE. The language $[\![\varphi]\!]$ associated with the CURE $\varphi$ is the set of all single-character strings $\{x \in \Sigma \mid \varphi(x) \text{ holds}\}$.

3. For each pair of non-empty CUREs $R_1$ and $R_2$, if the associated languages $L_1 = [\![R_1]\!]$ and $L_2 = [\![R_2]\!]$ are disjoint, then $R_1 \cup R_2$ is also a CURE, and $[\![R_1 \cup R_2]\!] = L_1 \cup L_2$.

4. Given a pair of non-empty CUREs $R_1$ and $R_2$, we say that they are *unambiguously concatenable*, if for each string $\sigma \in \Sigma^*$,

there is at most one split $\sigma = \sigma_1\sigma_2$ such that $\sigma_1 \in [\![R_1]\!]$ and $\sigma_2 \in [\![R_2]\!]$. If $R_1$ and $R_2$ are unambiguously concatenable, then $R_1 \cdot R_2$ is also a CURE, and $[\![R_1 \cdot R_2]\!] = [\![R_1]\!][\![R_2]\!]$.

5. A non-empty CURE $R$ is *unambiguously iterable* if for every string $\sigma$, there is at most one split $\sigma = \sigma_1\sigma_2\ldots\sigma_n$ into substrings such that $\sigma_i \in [\![R]\!]$ for each $i$. If $R$ is unambiguously iterable, then $R^*$ is also a CURE, and $[\![R^*]\!] = [\![R]\!]^*$.

For example, the regular expressions $\varphi$ and $(\neg\varphi)^*$ are unambiguously concatenable for every character predicate $\varphi$: every string $\sigma$ matching $\varphi \cdot (\neg\varphi)^*$ has to be split after the first character. On the other hand, $\Sigma^*$ is not unambiguously concatenable with itself: there are three ways to parse the string $aa$ in $\Sigma^* \cdot \Sigma^*$, because the left part of the concatenation can either match $\epsilon$, $a$, or $aa$. The regular expression $\Sigma^*$ is unambiguous — there is only one way to split each string $\sigma$ such that each substring is in $\Sigma$ — but $(\Sigma^*)^*$ is not unambiguous.[2]

We call two CUREs $R_1$ and $R_2$ *equivalent*, and write $R_1 \equiv R_2$, if $[\![R_1]\!] = [\![R_2]\!]$.

### 2.3.2 Consistency rules

A consistent DReX program is one that satisfies the rules defined in this section. One major effect of these rules is to guarantee that no string has multiple parse trees, so the word "unique" in the definitions of subsection 2.1 is unnecessary. The domain of a DReX program $f$ is the set containing every string $\sigma$ such that $[\![f]\!](\sigma)$ is defined. In the following rules, we assign each consistent DReX program a *domain type*, which is a representation of its domain as a CURE.

Most of the consistency rules are straight-forward, except for chained sum and `combine`. Recall from the definition of `chain` that for `chain(f, R)` to be defined on a string $\sigma$, we must have a split $\sigma = \sigma_1\sigma_2\ldots\sigma_n$ with $n \geq 2$. Therefore the domain type of `chain(f, R)` requires at least two matches of $R$: $R \cdot R \cdot R^*$. Next, we want expressions of the form `chain(combine(split(f_{11}, f_{12}), \ldots split(f_{k1}, f_{k2})), R \cdot R)`, where $R$ is unambiguously concatenable with itself, and the domain type of each $f_{ij}$ is $R_{ij} \equiv R$, to be consistent. We therefore pay special attention to the rule for `combine`.

1. All basic functions `bottom`, $\epsilon \mapsto d$, and $\varphi \mapsto d$ (where $\varphi$ is satisfiable) are consistent. Their domain types are $\bot$, $\epsilon$ and $\varphi$ respectively.

2. If $f_1$ and $f_2$ are both consistent and have unambiguously concatenable domain types $R_1$ and $R_2$ respectively, then `split(f_1, f_2)` and `left–split(f_1, f_2)` are also both consistent and have the domain $R_1 \cdot R_2$.

3. If $f$ is consistent and has domain type $R$, and $R$ is unambiguously iterable, then `iterate(f)` and `left–iterate(f)` are both consistent, with domain $R^*$.

4. If $f_1$ and $f_2$ are consistent with disjoint domain types $R_1$ and $R_2$ respectively, then $f_1$ `else` $f_2$ is also consistent with the domain $R_1 \cup R_2$.

5. If $f$ is consistent, and has domain type $R_1 \cdot R_2$, such that $R_1 \equiv R_2 \equiv R$, where $R$ is an unambiguously iterable CURE, then `chain(f, R)` and `left–chain(f, R)` are both consistent, and have the domain $R \cdot R \cdot R^*$.

6. If $f_1$ and $f_2$ are consistent with domain types $R_1$ and $R_2$ respectively, and $R_1 \equiv R_2$, then `combine(f_1, f_2)` is also consistent. Depending on the syntactic structure of the CUREs

$R_1$ and $R_2$, the domain type of `combine(f_1, f_2)` is defined as follows:

(a) If $R_1 = R_{11} \cdot R_{12}$, and $R_2 = R_{21} \cdot R_{22}$, with $R_{11} \equiv R_{12} \equiv R_{21} \equiv R_{22}$, then the domain type of `combine(f_1, f_2)` is $R_1$.

(b) Otherwise, if $R_1$ is *not* of the form $R_{11} \cdot R_{12}$ with $R_{11} \equiv R_{12}$, then the domain type is $R_1$.

(c) Otherwise, the domain type is $R_2$.

We now strengthen the claim originally made in theorem 1. While consistency was not an explicit goal in the original proof of theorem 1, it is the case that every expression constructed was actually consistent, and we can therefore state:

**Theorem 2.** *For every finite input alphabet $\Sigma$, and output alphabet $\Gamma$, every regular function $f : \Sigma^* \to \Gamma_\bot^*$ can be expressed by a consistent DReX program.*

The consistency and domain computation rules are syntax-directed, and straightforward to implement directly. We need to be able to answer the following basic questions about unambiguous regular expressions:

1. "Given CUREs $R_1$ and $R_2$, are $R_1$ and $R_2$ unambiguously concatenable?", "Given a CURE $R$, is it unambiguously iterable?", "Given CUREs $R_1$ and $R_2$, are they disjoint, or equivalently, is $R_1 \cup R_2$ also a CURE?" Observe that the traditional algorithm [24] to convert regular expressions to NFAs converts unambiguous regular expressions to unambiguous NFAs, where each accepted string has exactly one accepting path. Whether a regular expression $R$ is unambiguous can therefore be checked in polynomial time [8]: take the product of the corresponding ($\epsilon$-transition free) NFA $A_R$ with itself, and check for the presence of a reachable state $(q, q')$, with $q \neq q'$, which can itself reach a pair of accepting states $(q_f, q'_f) \in F \times F$, where $F$ is the set of accepting states of $A_L$. Thus, if the input alphabet $\Sigma$ is finite, these questions can be answered in polynomial time. Otherwise, the same problems for symbolic automata (representing $R_1$, $R_2$, etc.) are also decidable in polynomial time assuming that we can check in polynomial time whether a predicate is satisfiable.

2. "Given CUREs $R_1$ and $R_2$, is $R_1 \equiv R_2$?" If $\Sigma$ is finite, then from [25], we have that this can be checked in time $\mathcal{O}(poly(|R_1|, |R_2|, |\Sigma|))$. Otherwise, if CUREs are expressed using the symbolic notation of section 2.2, they can be translated into symbolic automata, and the equivalence of symbolic automata is decidable in polynomial time in the size of $R_1$ and $R_2$ and exponential[3] in the number of predicates appearing in $R_1$ and $R_2$.

**Theorem 3.** *Given a DReX program $f$ over an input alphabet $\Sigma$, checking whether $f$ is consistent is decidable. Furthermore, if the input alphabet $\Sigma$ is finite, then the consistency of $f$ can be determined in time $\mathcal{O}(poly(|f|, |\Sigma|))$.*

Note specifically that programs involving function composition are not consistent. In the rest of this paper, to distinguish the class of consistent DReX programs from the bigger class of *all* DReX programs, we will qualify the latter as the unrestricted class.

---

[2] To be consistent with the classic notation we write $\Sigma$ do denote the CURE that accepts all possible input characters, but formally in our definition the corresponding CURE is the predicate `true`.

[3] The algorithm proposed in [25] can check in polynomial time whether two unambiguous NFAs are equivalent. The algorithm requires the alphabet to be finite, and using the Minterm generation technique proposed in [14] one can make a symbolic alphabet finite by constructing the Boolean combinations of the predicates in the automaton. This operation however can cause an exponential blow-up.

## 2.4 Examples of consistent DReX programs

The simplest non-trivial DReX program is the identity function $id = \texttt{iterate}(\texttt{true} \mapsto x)$. Several variations of this program are also useful: $\texttt{iterate}(\texttt{lowercase}(x) \mapsto \texttt{touppercase}(x))$ maps strings of lower-case characters to upper-case, and $id_{\neg \texttt{space}} = \texttt{iterate}(\neg \texttt{space}(x) \mapsto x)$ is the identity function restricted to strings not containing a space.

More interesting functions can be constructed using the conditional operator: the function $sw\text{-}case = \texttt{uppercase}(x) \mapsto \texttt{tolowercase}(x) \, \texttt{else} \, \texttt{lowercase}(x) \mapsto \texttt{touppercase}(x)$ flips the case of a single input character, and so $\texttt{iterate}(sw\text{-}case)$ switches the case of each character in the input string.

Given a string of the form "*First-name Last-name*", the function $echo\text{-}first = \texttt{split}(id_{\neg \texttt{space}}, \texttt{space}(x) \mapsto \epsilon, \texttt{iterate}(\texttt{true} \mapsto \epsilon))$ outputs "*First-name*". Similarly, the function *echo-last* which outputs the last name could be written, and the two can be combined into $\texttt{combine}(echo\text{-}last, echo\text{-}first)$, which outputs "*Last-name First-name*". Note that the space in between is omitted — the expression $\texttt{combine}(\texttt{split}(echo\text{-}last, \epsilon \mapsto \texttt{" "}), echo\text{-}first)$ preserves this space. An example of the use of the left-additive operators is in string reversal: the function $\texttt{left-iterate}(\texttt{true} \mapsto x)$ reverses the input string.

Finally, to present an example of the chained sum combinator, we consider the situation of misaligned titles in BibTeX files. Assume that, by mistake, the title of the first entry appears in the second entry, the title of the second entry appears in the third entry, and so on. Let $R_{bib}$ be the unambiguous regular expression matching a BibTeX entry. Let $f_{bib}$ be the DReX program which examines pairs of adjacent BibTeX entries and outputs the title of the second entry and all other fields from the first entry. Then $\texttt{chain}(f_{bib}, R_{bib})$ corrects the misaligned text in the input file. We now outline the construction of $f_{bib}$. Let *echo-header* be the function which maps each BibTeX entry to its header (such as "@book{"), and *echo-title* be the function which maps each BibTeX entry to its title. Then $make\text{-}title = \texttt{split}(echo\text{-}header, echo\text{-}title)$ copies the header from the first BibTeX entry and the title from the second. Similarly, if *copy-body* echoes all the fields of a BibTeX entry except the title, and *delete-entry* maps an entire BibTeX entry to the empty string $\epsilon$, then $make\text{-}body = \texttt{split}(copy\text{-}body, delete\text{-}entry)$ completes the body of the output using the fields of the first BibTeX entry. We can then write $f_{bib} = \texttt{combine}(make\text{-}title, make\text{-}body)$.

The reader is referred to appendix A for a description of the consistent DReX programs used in our evaluation.

## 3. A Single-Pass Algorithm for Consistent DReX

In this section, we present the main technical contribution of this paper: a single-pass linear time algorithm to evaluate consistent DReX programs. We describe the intuition and present the idea of *function evaluators* in subsection 3.1, and then construct the evaluators for each DReX combinator. We conclude with subsection 3.5, a brief discussion of why this algorithm does not work with unrestricted DReX programs. Full proofs, and the omitted case of the chained sum will be included in the full version of this paper.

### 3.1 Intuition and the idea of function evaluators

Given a consistent DReX program $f$, we construct an evaluator $T$ which computes the associated function $[\![f]\!]$. The evaluator $T$ processes the input string from left-to-right, one character at a time. After reading each character, it outputs the value of $f$ on the string read so far, if it is defined.

To understand the input / output specifications of $T$, we consider the example program $\texttt{split}(f, g)$. In this case, $T$ is given the sequence of input signals $(Start, 0), (\sigma_1, 1), (\sigma_2, 2), \ldots, (\sigma_n, n)$. The first signal $(Start, 0)$ indicates the beginning of the string, and



Figure 3.1: Example run of the evaluator $T$ for $\texttt{split}(f, g)$ over a string $\sigma$. The evaluator $T_f$ emits a result at indices $i$ and $j$ of the input string. The evaluator $T_g$ for $g$ may simultaneously be processing multiple threads, corresponding to different potential parse trees of the input string $\sigma$. From the consistency rules, we know that at most one thread may return a result at each index, and so $T$ can safely emit a result in response to getting a result from $T_g$.

each character $\sigma_i$ is annotated with its index $i$ in the input string. After reading $(\sigma_i, i)$, $T$ responds with the value of $\texttt{split}(f, g)$ on $\sigma_1 \sigma_2 \ldots \sigma_i$, if it is defined.

Assume that $f$ and $g$ are consistent, and have unambiguously concatenable domain types $R_f$ and $R_g$ respectively. The evaluator $T$ maintains two sub-evaluators $T_f$ and $T_g$ for the functions $f$ and $g$ respectively. Each time $T$ receives the input $(a, i)$, it forwards this signal to both $T_f$ and $T_g$. Whenever $T_f$ reports a result, i.e. that $f$ is defined on the input string read so far, $T$ sends the signal *Start* to $T_g$ to start processing the suffix. Consider the situation in figure 3.1, where $f$ is defined for the prefixes $\sigma_1 \sigma_2 \ldots \sigma_i$ and $\sigma_1 \sigma_2 \ldots \sigma_j$. The input to the sub-evaluator $T_g$ is then the sequence $(\sigma_1, 1), (\sigma_2, 2), \ldots, (\sigma_i, i), (Start, i), (\sigma_{i+1}, i+1), \ldots, (\sigma_j, j), (Start, j), \ldots, (\sigma_n, n)$.

For each signal $(Start, i)$ occurring in the input string, we call the subsequent sequence of characters $\sigma_{i+1} \sigma_{i+2} \ldots$ the *thread* beginning at index $i$. Note that each thread corresponds to a potential parse tree of $\sigma$, and that $T_g$ may be processing multiple such threads simultaneously. The main challenge is to ensure that the number of active threads in $T_g$ is bound by $\mathcal{O}(|g|)$, and is independent of the length of the input string. After reading $\sigma_n$, $T_g$ reports a result to $T$, the evaluator for the $\texttt{split}(f, g)$. To uniquely identify the thread $j$ reporting the result, the result signal $(Result, j, \gamma_g)$ is annotated with the index $j$ at which the corresponding *Start* was received.

Note that the consistency rules guarantee that, after reading each input symbol, $T_g$ emits at most one result, for otherwise the prefix of the input string read so far would have multiple parse trees.

When $T$ receives this result signal from $T_g$, it combines it with the response $(Result, 0, \gamma_f)$ initially obtained from $f$ at position $j$, and itself emits the result $(Result, 0, \gamma_f \gamma_g)$. To do this, it maintains a set $th_g$ (for threads) of triples $(i_{0f}, i_{0g}, \gamma_f)$, where $i_{0f}$ is the index along the input string at which $T_f$ was started, $i_{0g}$ was the index at which $T_f$ reported a result and $T_g$ was started, and $\gamma_f$ was the result reported by $T_f$. In order to prevent this set $th_g$ from becoming too large, $T_g$ emits *kill signals*. Say that, at index $k$, $T_g$ discovers that for every possible suffix $\tau \in \Sigma^*$, $g$ will be undefined for the string $\sigma_{i+1} \sigma_{i+2} \ldots \sigma_k \tau$, and so the thread $(Start, i)$ of $T_g$ initiated at the input index $i$ can never return a result. It then emits $(Kill, i)$ to signal to $T$ that the relevant entries in the set $th_g$ can be deleted.

***Formal specification of the evaluator.*** The input alphabet to each evaluator is therefore $In = (\Sigma \cup \{Start\}) \times \mathbb{N}$ and the output alphabet is $Out = (\{Result\} \times \mathbb{N} \times \Gamma^*) \cup (\{Kill\} \times \mathbb{N})$, where $\Sigma$ is the input and $\Gamma$ is the output alphabet of the DReX program.

While constructing the evaluator $T$ for a DReX program $f$, we assume the following condition of *input validity*: for each prefix of

the input stream *in*, there is at most one thread for which $f$ is defined. Thus, for example, $T$ can never see two consecutive $(Start, i)$ signals for the same $i$. In return, we make the following guarantees:

**Correctness of results.** After reading each input signal $(\sigma_j, j)$ in *in*, we report the result $(Result, i, \gamma)$ exactly for that thread $(Start, i)$ such that $[\![f]\!](\sigma_{i+1}\sigma_{i+2}\ldots\sigma_j) = \gamma$, if it exists.

**Eagerness of kills.** Every thread $\sigma$ beginning at $(Start, i)$ of *in*, such that there is no suffix $\tau$ for which $f(\sigma\tau)$ is defined, is killed exactly once while reading *in*. Furthermore, there are always at most $\mathcal{O}(|f|)$ active threads, where $|f|$ is the size of $f$.

If an evaluator $T$ satisfies these requirements for $f$, then we say that the evaluator *computes* $f$. On the input $(Start, 0), (\sigma_1, 1), (\sigma_2, 2),$ $\ldots, (\sigma_n, n)$, the evaluator outputs a result $\gamma$ in exactly those cases when $f$ is defined, and in that case, $[\![f]\!](\sigma) = \gamma$. We will ensure that the evaluator $T$ processes each input signal in time $\mathcal{O}(poly(|f|))$.[4]

### 3.2 Basic evaluators

The simplest case is when $f = \texttt{bottom}$. The evaluator $T_\perp$ is defined by the following rules:

1. On input $(Start, i)$, respond with $(Kill, i)$.

2. On input $(a, i)$, for $a \in \Sigma$, do nothing.

Next, we consider the evaluator $T_{\epsilon \mapsto d}$, for the case when $f = \epsilon \mapsto d$, for some $d \in \Gamma^*$. Intuitively, this evaluator returns a result immediately on receiving a start signal, but can only kill the thread after reading the next symbol. It therefore maintains a set $th \subseteq \mathbb{N}$ of currently active threads, which are to be killed on reading the next input symbol. The set $th$ is initialized to $\emptyset$.

1. On input $(Start, i)$, respond with $(Result, i, d)$. Update $th := th \cup \{i\}$.

2. On input $(a, i)$, for $a \in \Sigma$, respond with $(Kill, j)$, for each thread start index $j \in th$. Update $th := \emptyset$.

Observe that by the condition of input validity, we can never observe two consecutive start signals in the input stream. Therefore, $|th| \leq 1$, and the response time of $T_{\epsilon \mapsto d}$ to each input signal is bounded by a constant.

The final basic function is $f = \varphi \mapsto d$ for some character predicate $\varphi$ and $d : (\Sigma \to \Gamma)^*$. The evaluator $T_{\varphi \mapsto d}$ maintains two sets $th, th' \subseteq \mathbb{N}$ of thread start indices, initialized to $th = th' = \emptyset$. $th$ is the set of threads for which no symbol has yet been seen, while $th'$ is the set of threads for which one input symbol has been seen, and that input symbol satisfied the predicate $\varphi$.

1. On input $(Start, i)$, update $th := th \cup \{i\}$.

2. On input $(a, i)$, for $a \in \Sigma$:

    (a) Emit $(Kill, j)$, for each thread $j \in th'$.

    (b) If $a$ satisfies the predicate $\varphi$, for each thread $j \in th$, emit $(Result, j, d(a))$. Update $th' := th$, and $th := \emptyset$.

    (c) If $a$ does not satisfy the predicate $\varphi$, then for each thread $j \in th$, emit $(Kill, j)$. Update $th := \emptyset$, and $th' := \emptyset$.

Just as in the case of $\epsilon \mapsto d$, we have $|th|, |th'| \leq 1$, and so $T_{\varphi \mapsto d}$ responds to each input signal in time bounded by some constant.

---

[4] We assume a representation for strings with concatenation requiring only constant time. Specifically, strings are only concatenated symbolically using a pointer representation. Such "lazily" represented strings can be converted into the traditional sequence-of-characters representation in time linear in the string length.

### 3.3 State-free evaluators: combination and conditionals

The simplest non-trivial evaluator is for $\texttt{combine}(f, g)$. Recall that, by the consistency requirements, we have $R_f \equiv R_g$ for the domain types $R_f$ and $R_g$ of the sub-expressions. Thus, all state can be maintained by the sub-evaluators $T_f$ and $T_g$ and $T_{\texttt{combine}(f,g)}$ can be entirely state-free. It has the following behavior:

1. On input $(Start, i)$, send the signal $(Start, i)$ to both sub-evaluators $T_f$ and $T_g$.

2. On input $(a, i)$, send the signal $(a, i)$ to both $T_f$ and $T_g$.

3. On receiving the result $(Result, i, \gamma_f)$ from $T_f$ and the result $(Result, i, \gamma_g)$ from $T_g$ (which, according to the consistency requirements for $\texttt{combine}(f, g)$, have to occur simultaneously), respond with $(Result, i, \gamma_f\gamma_g)$.

4. On receiving the kill signals $(Kill, i)$ from $T_f$ and $T_g$ (by the consistency rules, necessarily simultaneously), emit the kill signal $(Kill, i)$.

The evaluator $T_{f\,\texttt{else}\,g}$ maintains two sub-evaluators $T_f$ and $T_g$. In addition, it maintains two sets $th_f, th_g \subseteq \mathbb{N}$ of threads currently active in $T_f$ and $T_g$ respectively. Both sets are initialized to $\emptyset$. The behavior of $T_{f\,\texttt{else}\,g}$ is defined as follows:

1. On receiving the input $(Start, i)$, update $th_f := th_f \cup \{i\}$, and $th_g := th_g \cup \{i\}$. Send the start signal $(Start, i)$ to both $T_f$ and $T_g$.

2. On receiving the input $(a, i)$ for some $a \in \Sigma$, send the input $(a, i)$ to both $T_f$ and $T_g$.

3. When either $T_f$ or $T_g$ respond with the result $(Result, i, \gamma)$ (by the consistency rules, we know that the other sub-evaluator is not responding with a result), emit the result $(Result, i, \gamma)$.

4. If a kill signal $(Kill, i)$ is received from $T_f$ (resp. $T_g$), update $th_f := th_f \setminus \{i\}$ (resp. $th_g := th_g \setminus \{i\}$). If $i \notin th_f$ and $i \notin th_g$, then kill the thread by emitting $(Kill, i)$.

The sizes of $th_f$ and $th_g$ are bounded by the number of active threads of $T_f$ and $T_g$ respectively, and hence it follows that $T_{f\,\texttt{else}\,g}$ responds to each input signal in time $\mathcal{O}(|f| + |g|) + t_f + t_g$, where $t_f$ and $t_g$ are the response times of $T_f$ and $T_g$ respectively.

### 3.4 Stateful evaluators: iteration and split sum

We now construct evaluators for $\texttt{iterate}(f)$ and $\texttt{split}(f, g)$. The evaluators for $\texttt{left-iterate}(f)$ and $\texttt{left-split}(f, g)$ are symmetric with respect to concatenation and can be constructed similarly.

First, we build the evaluator $T_{\texttt{iterate}(f)}$, where $f$ is consistent and has the unambiguously iterable domain type $R_f$. Whenever $T_{\texttt{iterate}(f)}$ receives a start signal $(Start, i)$, or an input signal $(a, i)$, this is passed to $T_f$. Consider a sequence of input signals $\sigma$, as shown in figure 3.2. After reading each input symbol, say $(\sigma_n, n)$, $T_f$ may report that $f$ is defined for a suffix of the input stream $(Start, k), (\sigma_{k+1}, k+1), \ldots, (\sigma_n, n)$ seen so far. The evaluator $T_{\texttt{iterate}(f)}$ responds by initiating a new thread of $T_f$ by sending it the start signal $(Start, n)$. Furthermore, it has to record the result $(Result, k, \gamma_3)$ just reported by $T_f$. It does this by adding the entry $(i, n, \gamma_1\gamma_2\gamma_3)$ to the set $th$. Each entry $(i_0, j_0, \gamma) \in th$ refers to an active thread $j_0$ of $T_f$, the index of the signal $(Start, i_0)$ received by $T_{\texttt{iterate}(f)}$, and the cumulative result $\gamma$ obtained so far.

Formally, the set $th \subseteq \mathbb{N} \times \mathbb{N} \times \Gamma^*$ is initialized to $\emptyset$. The evaluator $T_{\texttt{iterate}(f)}$ does the following:

1. On input $(Start, i)$:

    (a) Update $th := th \cup \{(i, i, \epsilon)\}$.

$$1 \quad (\textit{Start}, i) \qquad j \qquad k \qquad n$$

Figure 3.2: For each thread $(\textit{Start}, i_0)$ of the evaluator $T_{\texttt{iterate}(f)}$, there may be multiple potential parse trees. The evaluator $T_{\texttt{iterate}(f)}$ maps individual threads $(\textit{Start}, i)$ of $T_f$ to the corresponding start signal $(\textit{Start}, i_0)$ in $T_{\texttt{iterate}(f)}$ through the entry $(i_0, i, \gamma)$ in the set $th$. Thus, after obtaining the response from $T_f$ at index $n$, $T_{\texttt{iterate}(f)}$ updates $th := th \cup \{(i, n, \gamma_1 \gamma_2 \gamma_3)\}$.

    (b) Send $(\textit{Start}, i)$ to $T_f$. Assert that $T_f$ does not respond with a result $(\textit{Result}, i, \gamma)$, because by the consistency rules, $f(\epsilon)$ is undefined for $R_f$ to be unambiguously iterable.

    (c) Respond with the result $(\textit{Result}, i, \epsilon)$.

2. On input $(a, i)$, send the signal $(a, i)$ to $T_f$. For each response of $T_f$, do the following:

    (a) If the response is a result, $(\textit{Result}, j, \gamma_f)$, then, find the corresponding entry $(j_0, j, \gamma) \in th$, for some values of $j_0$ and $\gamma$. Assert (by the invariant that $th$ records the active threads of $T_f$) that this entry exists, and is unique.

       i. Update $th := th \cup \{(j_0, i, \gamma \gamma_f)\}$.

       ii. Send the signal $(\textit{Start}, i)$ to $T_f$. Confirm that $T_f$ does not respond with a result $(\textit{Result}, i, 0, \gamma'_f)$, for that would violate the consistency requirements.

       iii. Respond with the result $(\textit{Result}, j_0, \gamma \gamma_f)$.

    (b) If the response is a kill signal, $(\textit{Kill}, j)$:

       i. Let *kill-ring* be the set of all tuples $(j_0, j, \gamma) \in th$, for some values of $j_0$ and $\gamma$. By the consistency requirements, *kill-ring* is asserted to be a singleton set.

       ii. Update $th := th \setminus \textit{kill-ring}$.

       iii. For every entry $(j_0, j, \gamma) \in \textit{kill-ring}$ if there is no entry of the form $(j_0, j', \gamma') \in th$, then emit the kill signal $(\textit{Kill}, j_0)$.

Observe that an element is added to $th$ exactly when it is sent a start signal, and an entry is deleted exactly when $T_{\texttt{iterate}(f)}$ receives a kill signal. Thus, the entries of $th$ correspond to the active threads of $T_f$, and its size is bounded by $\mathcal{O}(|f|)$. The response time of $T_{\texttt{iterate}(f)}$ to each input signal is therefore $\mathcal{O}(|f|) + t_f$, where $t_f$ is the response time of $T_f$.

The evaluator $T_{\texttt{split}(f,g)}$ for $\texttt{split}(f, g)$ is similar, except that it maintains two sets: the first set $th_f \subseteq \mathbb{N}$ is the set of thread start indices which are still active in $T_f$, and the second set $th_g \subseteq \mathbb{N} \times \mathbb{N} \times \Gamma^*$ is the set of triples $(i_0, i, \gamma_f)$ which indicates, for each active thread in $T_g$, the index $i$ at which $T_g$ was signaled to start, the index $i_0$ of the original start received by $T_{\texttt{split}(f,g)}$, when $T_f$ was started, and the value $\gamma_f$ returned by $T_f$ on the prefix. Both sets are initialized to $\emptyset$, and $T_{\texttt{split}(f,g)}$ follows the following rules:

1. On input $(\textit{Start}, i)$:

    (a) Update $th_f := th_f \cup \{i\}$.

    (b) Send $(\textit{Start}, i)$ to $T_f$. Let $Rsp_f$ be the responses from $T_f$. Let $Rsp_g = \emptyset$.

2. On input $(a, i)$:

    (a) Send $(a, i)$ to $T_f$. Let $Rsp_f$ be the set of responses from $T_f$.

    (b) Send $(a, i)$ to $T_g$. Let $Rsp_g$ be the set of responses from $T_g$.

3. For each response $r \in Rsp_f \cup Rsp_g$, do the following:

    (a) If $r$ is a result $(\textit{Result}, j_0, \gamma_f)$ from $T_f$,

       i. Update $th_g := th_g \cup \{(j_0, i, \gamma_f)\}$.

       ii. Send the signal $(\textit{Start}, i)$ to $T_g$. Update $Rsp_g := Rsp_g \cup Rsp'_g$, where $Rsp'_g$ is the set of responses from $T_g$.

    (b) If $r$ is a result $(\textit{Result}, j, \gamma_g)$ from $T_g$, let $(j_0, j, \gamma_f)$ be the (by the consistency rules, necessarily unique) corresponding record in $th_g$. Respond with $(\textit{Result}, j_0, \gamma_f \gamma_g)$.

    (c) If $r$ is a kill signal $(\textit{Kill}, j_0)$ from $T_f$,

       i. Update $th_f := th_f \setminus \{j_0\}$.

       ii. If there is no element of the form $(j_0, j, \gamma_f) \in th_g$, for some values of $j$, $\gamma_f$, kill the thread: $(\textit{Kill}, j_0)$.

    (d) Finally, if $r$ is a kill signal $(\textit{Kill}, j)$ from $T_g$,

       i. Let *kill-ring* be the set of tuples $(j_0, j, \gamma_f) \in th_g$ for some values of $j_0$, $\gamma_f$.

       ii. Update $th_g := th_g \setminus \textit{kill-ring}$.

       iii. For every record $(j_0, j, \gamma_f) \in \textit{kill-ring}$, if there is no longer a record of the form $(j_0, j', \gamma'_f) \in th_g$, and $j_0 \notin th_f$, kill the thread beginning at $j_0$: $(\textit{Kill}, j_0)$.

Observe that $|th_f|$ is bound by the number of active threads of $T_f$ and $|th_g|$ is bound by the number of active threads of $T_g$. Thus, $T_{\texttt{split}(f,g)}$ responds to each input signal in time $\mathcal{O}(|f| + |g|) + t_f + t_g$, where $t_f$ and $t_g$ are the response times of $T_f$ and $T_g$ respectively.

### 3.5 What breaks with unrestricted DReX programs?

First, notice that the function composition operator is un-amenable to the evaluator model we just described. We wish to process each character in bounded time, regardless of the length of the input string. Consider the case of $\texttt{compose}(f, g)$: when the evaluator $T_f$ returns a result $(\textit{Result}, i, \gamma)$, we have to pass the *entire* intermediate result string $\gamma$ to $T_g$, and this is possibly as long as the input seen so far. Notice that this limitation should be unsurprising, because nested function compositions — such as the transformation $\sigma \mapsto \sigma \sigma$ composed with itself $k$ times — can cause an exponential blowup in the length of the output string.

Next, consider a potential evaluator $T$ for $\texttt{split}(f, g)$, in the absence of any consistency requirement. Thus, there might exist strings $\sigma = \sigma_1 \sigma_2 \ldots \sigma_n$ which admit two splits $\sigma = \tau_1 \tau_2 = \tau'_1 \tau'_2$, such that all of $[\![f]\!](\tau_1)$, $[\![f]\!](\tau'_1)$, $[\![g]\!](\tau_2)$, and $[\![g]\!](\tau'_2)$ are defined. In this case, $\texttt{split}(f, g)$ is undefined for the entire string $\sigma$. We have to drop the requirement of input validity, because the nested evaluator $T_g$ emits *two Result* signals after reading $\sigma_n$. We could conceivably modify $T$ to emit an output signal when exactly one thread of $T_g$ returns a result. Unfortunately, it turns out that this modification is insufficient, as the induction hypothesis now breaks — the evaluator $T$ has to perform additional book-keeping to report results correctly. The consistency rules provide an easy way to avoid this non-trivial book-keeping.

Finally, note that unrestricted use of the chained sum operator does not satisfy the output bound $|[\![f]\!](\sigma)| = \mathcal{O}(poly(|f|) |\sigma|)$, while every consistent DReX program obeys $|[\![f]\!](\sigma)| \leq |f| |\sigma|$.

## 4. The Complexity of Unrestricted DReX

In this section we first describe the dynamic programming algorithm to evaluate DReX programs. We show that it has time complexity cubic in the size of the input string, and when function composition is allowed, requires time exponential in the size of the program. We then show that the evaluation problem for DReX programs with composition is PSPACE-complete, and is thus computationally hard. Finally, we argue that for unrestricted DReX programs, even

when disallowing composition, there is no evaluation algorithm for which the complexity is linear in the length of the input string and polynomial in the size of the program.

## 4.1 Evaluation by dynamic programming

We now describe the dynamic programming algorithm to evaluate unrestricted DReX programs. This algorithm also works on programs containing the function composition operator. The algorithm mimics the semantics of DReX by computing the following functions (represented as lookup tables). Given a program $f$, and a string $\sigma$, for any two numbers $i$ and $j$ the algorithm computes the function $\mathrm{OUT}(f, \sigma, i, j)$ representing the output of $f$ on the substring[5] $\sigma[i, j]$, $[\![f]\!](\sigma[i, j])$. To evaluate the operator $\mathtt{iterate}(f)$ we also need to compute the function $\mathrm{COUNT}(\mathtt{iterate}(f), \sigma, i, j)$ that counts the number of possible ways to split $\sigma[i, j]$ so that each split is accepted by $f$, and to evaluate the operator $\mathtt{chain}(f, R)$ we compute the function $\mathrm{BEL}(R, \sigma, i, j)$ which checks whether a substring $\sigma[i, j]$ belongs to the language $[\![R]\!]$. Every function $\mathrm{OUT}$, $\mathrm{COUNT}$, and $\mathrm{BEL}$ is represented by a table and for each string $\sigma$, and each sub-program $g$ of $f$, each table will have $\mathcal{O}(|\sigma|^2)$ entries, corresponding to the substrings of $\sigma$. The final output of the algorithm is $\mathrm{OUT}(f, \sigma, 1, |\sigma| + 1)$.

We explain the intuition of the algorithm by showing how the entries are computed for the iteration and composition operators. The value $\mathrm{OUT}(\mathtt{iterate}(f), \sigma, i, j)$ corresponding to the output of $\mathtt{iterate}(f)$ on the string $\sigma' = \sigma[i, j]$ is defined iff there is a unique way to split the string $\sigma'$ into multiple chunks so that $f$ is defined on each chunk, i.e. iff $\mathrm{COUNT}(\mathtt{iterate}(f), \sigma, i, j) = 1$. If this is the case, then we know that there is a unique value $k$, such that $i \leq k < j$, for which both $\gamma_{pre} = \mathrm{OUT}(\mathtt{iterate}(f), \sigma, i, k)$ and $\gamma_{post} = \mathrm{OUT}(f, \sigma, k, j)$ are defined, and $\gamma_{pre}\gamma_{post}$ is the output of $\mathtt{iterate}(f)$ on $\sigma[i, j]$. Looking for this witness $k$ takes at most $|\sigma|$ steps if all the required table entries have already been computed. Similarly, the entry $\mathrm{COUNT}(\mathtt{iterate}(f), \sigma, i, j)$ can be computed in at most $|\sigma|$ steps by counting for how many values of $l$ the function $\mathrm{COUNT}(\mathtt{iterate}(f), \sigma, i, l)$ is greater than 0 and the functions $\mathrm{OUT}(f, \sigma, l, j)$ is defined.

The rule for computing $\mathrm{OUT}(\mathtt{compose}(f_1, f_2), \sigma, i, j)$ is what causes an exponential blow-up in the evaluation time. To compute the output of $\mathtt{compose}(f_1, f_2)$ on $\sigma[i, j]$, we first need to compute the output $\tau = \mathrm{OUT}(f_1, \sigma, i, j)$ of the program $f_1$ on the string $\sigma[i, j]$ and then the output $\mathrm{OUT}(f_2, \tau, 1, |\tau| + 1)$ of the program $f_2$ on the string $\tau$. As we will show, when using the composition operator the size of the output $\tau$ may grow exponentially, and since computing each entry of the table $\mathrm{OUT}(f_2, \tau, \cdot, \cdot)$ requires $\mathcal{O}(\tau)$ steps, the resulting complexity is also exponential.

We now bound the output length of programs containing the composition operator.

**Lemma 4** (Output size). *Given a program $f$ and an input string $\sigma \in \Sigma^*$ such that $\tau = [\![f]\!](\sigma)$ is defined, we have $|\tau| \leq |f|^{d+1} |\sigma|$. Here $d$ is the number of composition operators appearing in $f$.*

We now state the complexity of the dynamic programming routine to evaluate DReX programs.

**Theorem 5** (Complexity of dynamic programming). *Given a program $f$, and an input string $\sigma \in \Sigma^*$, the output $[\![f]\!]$ can be computed in time $\mathcal{O}(|f|^{2d^2+5d+4} |\sigma|^{2d+3})$ where $d$ is the number of composition operators in $f$. If $d = 0$ we can show that the algorithm has complexity $\mathcal{O}(|f| |\sigma|^3)$.*

*Proof sketch.* For a particular string $\sigma$ and if the program $f$ that does not contain any composition operators ($d = 0$), computing

---

[5] We define the substring $\sigma[i, j]$ of $\sigma$ as $\sigma_i \sigma_{i+1} \ldots \sigma_{j-1}$, so that $\sigma[i, i] = \epsilon$, for each $i$, and $\sigma[1, |\sigma| + 1]$ is the entire string $\sigma$.

each entry of the tables $\mathrm{OUT}$, $\mathrm{COUNT}$, and $\mathrm{BEL}$ takes time $\mathcal{O}(|\sigma|)$ and since there are $|f| |\sigma|^2$ entries the algorithm has complexity $\mathcal{O}(|f| |\sigma|^3)$. In the presence of composition operators which can produce intermediate results, for each intermediate string $\tau$, a new table of size $|\tau|^2$ must be created. By lemma 4 we know that an intermediate result has size at most $\mathcal{O}(|f|^{d+1} |\sigma|)$, hence the complexity is exponential in $d$. □

## 4.2 DReX evaluation with composition is PSPACE-complete

While the main appeal of the algorithm in subsection 4.1 is ease of implementation, it can use exponential space. It turns out that, even in the presence of composition, DReX programs can be evaluated in PSPACE. First observe that, since the output computed by a program has at most exponentially many characters (lemma 4), the index of each character in the output is only polynomially many bits long. We therefore adopt an implicit representation of strings with the following operations (in contrast with the traditional explicit list-of-characters representation of strings): (*a*) check whether $\sigma$ is defined; (*b*) compute the length of $\sigma$; and (*c*) given an index $i$, compute the $i$-th character of $\sigma$. Finally, by structural induction on the DReX program $f$, and given an implicit representation of the input string $\sigma$, we build an implicit representation of $[\![f]\!](\sigma)$ using only polynomial space. For example, the implicit representation of $[\![f \mathtt{ else } g]\!](\sigma)$ would function as follows: (*a*) to check whether the output is defined, simply determine whether either $[\![f]\!](\sigma)$ or $[\![g]\!](\sigma)$ is defined; (*b*) to compute the length of the output, if $[\![f]\!](\sigma)$ is defined, return the length of $[\![f]\!](\sigma)$, and otherwise, return the length of $[\![g]\!](\sigma)$; and (*c*) to compute the $i$-th character of $[\![f \mathtt{ else } g]\!](\sigma)$, if $\tau_1 = [\![f]\!](\sigma)$ is defined, then return the $i$-th character of $\tau_1$, and otherwise return the $i$-th character of $\tau_2 = [\![g]\!](\tau)$. Observe that since both nested implicit representations $[\![f]\!](\sigma)$ and $[\![g]\!](\sigma)$ consume only polynomial space, $[\![f \mathtt{ else } g]\!](\sigma)$ is itself evaluated in polynomial space. The most interesting case is $\mathtt{compose}(f, g)$ where we simply connect the implicit representation of the output of $[\![f]\!](\sigma)$ to the input of the function $g$. The only non-trivial case is when $f = \mathtt{iterate}(g)$. To check whether $f$ is defined on the input $\sigma$, we need to determine whether there is exactly one way to split $\sigma$ such that $g$ is defined on each split. Consider each position in the string $\sigma$ as a vertex in a graph, with an edge between two vertices iff $g$ is defined on the substring between them. Then each path from the initial to the final node of this graph corresponds to a viable split of $\sigma$, and thus $f$ is defined on $\sigma$ iff there is a unique path from the initial node to the final node in this implicitly represented graph of potentially exponential size. This problem can be solved in PSPACE.

**Theorem 6.** *Given a DReX program $f$, and strings $\sigma \in \Sigma^*$ and $\tau \in \Gamma^*$, the problem of determining whether $[\![f]\!](\sigma) = \tau$ is in PSPACE.*

Finally, we show that when we allow the use of composition operators polynomial space is required, and the problem of evaluating a DReX program is PSPACE-complete.

**Theorem 7.** *The following problems are PSPACE-complete:*

1. *given a program $f$ in DReX check whether $[\![f]\!](\epsilon)$ is defined;*
2. *given a program $f$ in DReX and a string $\sigma \in \Sigma^*$ check whether $[\![f]\!](\sigma)$ is defined;*
3. *given a program $f$ in DReX, a string $\sigma \in \Sigma^*$, and a string $\tau \in \Gamma^*$, check whether $[\![f]\!](\sigma) = \tau$;*

*Proof sketch of PSPACE-hardness.* To show that the first problem is PSPACE-hard we reduce from the validity problem for quantified Boolean formulas (QBF). Intuitively given a QBF $\Phi = \forall x_1 \exists x_2 \ldots \exists x_n \varphi(x_1, \ldots, x_n)$ we construct a DReX program $f_\Phi$ such that $[\![f_\Phi]\!](\epsilon)$ is defined iff $\Phi$ is valid. The program $f_\Phi$ is the composition of three programs $f_{01}$, $f_{3\mathrm{CNF}}$ and $f_Q$ where:

1. $f_{01}$ takes as input $\epsilon$ and outputs all the strings in $\{0,1\}^n$ in lexicographic order and separated by a $\#$; this program generates all the possible assignments of the boolean variables.

2. $f_{3\text{CNF}}$ takes as input the string of all the assignments produced by $f_{01}$ and replaces each assignment in $a \in \{0,1\}^n$ with $T$ if the assignment $a$ satisfies the 3SAT formula $\varphi$ and $F$ otherwise.

3. $f_Q$ takes as input the string over $(\{T,F\}\#)^*$ and checks whether such a sequence of satisfying assignments is valid for the quantified formula $\Phi$. If it is valid it outputs $\epsilon$ and otherwise it is undefined.

The other problems can be reduced to the first problem using the composition operator. □

### 4.3 Single-pass algorithms for unrestricted DReX

In the proof of theorem 5 we showed that in the absence of the function composition combinator the dynamic programming algorithm has complexity $\mathcal{O}(f|\sigma|^3)$ where $\sigma$ is the input string and $f$ the program. In this section we argue that, if one wants to obtain an algorithm that is linear in the size of the input, it is necessary to pay at least an exponential complexity in the size of the program.

DReX operators are similar to those offered by regular expressions (iteration, split, etc. are the broad analogues of Kleene-*, concatenation, etc.). Since one can evaluate regular expressions efficiently by transforming them into nondeterministic finite automata one can try to construct an automaton model corresponding to DReX programs. Unfortunately as we discussed in section 2, DReX combinators can also express language intersection and other complex operations. In the presence of such operations, directly constructing an automaton model from the program seems hard (see [23], where the author summarizes the state-of-the-art on matching regular expressions extended with an intersection operator). The following is currently the best claim we can make about evaluating unrestricted DReX programs with a single linear time pass over the input string:

**Theorem 8.** *Given a DReX program $f$ and an input string $\sigma \in \Sigma^*$, we can compute the output $[\![f]\!](\sigma)$ in time linear in the length of the input string $\sigma$, and a tower of exponentials with height $\mathcal{O}(|f|)$.*

The algorithm compiles $f$ into an equivalent streaming string transducer (SST) [1] using the procedure described in [3]. An SST is a finite state machine that reads the input in a left-to-right fashion and stores intermediate results inside variables. The final output is a combination of such variables. An SST $A$ can be evaluated on an input string $\sigma$ in one pass in time $\mathcal{O}(|A| |\sigma|)$. Since SSTs are deterministic, operations such as concatenation and iteration (even in the absence of function composition) cause an exponential blow-up, making the overall complexity non-elementary.

## 5. Evaluation

We implemented the algorithms described in this paper, and evaluated their performance on a representative set of text and BibTeX file transformations. We show that

1. the evaluation algorithm for consistent DReX scales to inputs with more than 100,000 characters (subsection 5.2.1), and

2. the dynamic programming algorithm presented in section 4.1 *does not* scale for inputs with more than 3,000 characters (subsection 5.2.2).

Finally, we remark on the subjective experience of expressing string transformations using DReX (subsection 5.3).

### 5.1 Implementation details

The prototype implementation of DReX was written in Java and uses the recently released Java SE 8. We used the symbolic automata

| Program name | Size | CC (ms) |
|---|---|---|
| *delete_comm* | 28 | 12 |
| *insert_quotes* | 28 | 6 |
| *get_tags* | 31 | 6 |
| *reverse* | 5 | 1 |
| *swap_bibtex* | 1663 | 262 |
| *align_bibtex* | 3652 | 537 |

Table 1: Evaluated programs with sizes and time to check consistency.



Figure 5.1: Evaluation time for the single-pass algorithm.

library SVPAlib [11] to implement the symbolic operations required by the consistency-checking algorithm (theorem 3). The set of characters were all 16-bit UTF-16 code units, and the predicates were unions of character intervals (such as `[a-z, A-Z, 0-9]`).

The experiments were run on regular contemporary hardware: Windows 7 running on a 64-bit quad-core Intel Core i7-2600 CPU, at 3.40 GHz with 8 GB of RAM. Each experiment was run 10 times and the results reported are the mean of the obtained running times.

The dynamic programming (DP) algorithm for the extended version of DReX (theorem 5) is implemented lazily. Each entry $\text{OUT}(f, \sigma, i, j)$ is only computed and allocated when its value is required by another entry. Without this technique the algorithm runs out of memory for inputs of length smaller than 1,000. We also optimize the DP algorithm to take advantage of the consistency-checking; for consistent programs the algorithm does not need to check whether there is more than one way to match a string.

### 5.2 Benchmark programs

Table 1 shows the programs we considered in our evaluation together with their sizes and the running time of the consistency-checking algorithm. These programs are described in appendix A. Observe that for every program the consistency-checking algorithm terminates in less than 600 ms.

We evaluated the first four programs on randomly generated text files of size varying between 1 and 100,000 characters, and evaluated the more complex functions *swap_bibtex* and *align_bibtex* on actual BibTeX files of size also varying between 1 and 100,000 characters. We set a timeout of 60 seconds for each operation.

### 5.2.1 Single-pass algorithm for consistent DReX

Figure 5.1 shows how the running time for the algorithm presented in section 3 depends linearly on the size of the input. For inputs up to 100,000 characters each program takes less than 8 seconds to compute the output. Also observe that the evaluation algorithm can successfully handle reasonably large programs such as *align_bibtex* which has an AST with 3652 nodes.

Figure 5.2: Evaluation time for dynamic programming algorithm. Note that the $x$-axis in log-scale.

### 5.2.2 Dynamic programming for unrestricted DReX

Figure 5.2 shows the running time for the dynamic programming algorithm presented in section 4.1. The $x$-axis is shown in log scale to better appreciate the difference between the different programs. From the figure we can see that the running time depends polynomially on the length of the input and, as a consequence, all the considered programs time out for inputs with more than 70,000 characters. In some cases large programs such as *swap_bibtex* execute faster than smaller programs such as *get_tags*. This is due to the fact that our implementation uses some optimizations that may depend on the shape of the program: for example in the case of programs that are only defined on strings of a fixed length $k$, the algorithm is evaluated only for those $i$ and $j$ such that $j - i = k$. In conclusion, although all DReX programs can be evaluated using the dynamic programming algorithm presented in section 4.1, the procedure does not scale to large inputs.

### 5.3 User experience and comparison to existing tools

We were able to easily program several non-trivial string-to-string transformations without having to worry about efficiency. The main restrictions of the streaming evaluation algorithm are that programs cannot use compositions and have to be consistent. In our case study we did not find instances where composition was required and all the natural implementations of our programs were consistent. Moreover, in many cases the consistency algorithm helped us in identifying sources of ambiguity that caused our program to be incorrect. As an anecdotal account, we had mistakenly concatenated the sub-program $copy\_spaces = \mathtt{iterate}(\mathtt{space}(x) \mapsto x)$, which copies all whitespace characters (including tabs and newlines), with itself. In this case, the type-checker warned us that this concatenation was ambiguous on the input string "\n". This was clearly a bug in our script, and would have led to unexpected behavior even if we had used the dynamic programming evaluation algorithm.

We also programmed the benchmark transformations in sed, AWK, and Perl. Regular-expression based substitution, as present in all of these tools, is very efficient and usable to substitute or delete substrings based on patterns. This includes the benchmark programs *delete_comm* and *insert_quotes*, for which the sed implementations were $\approx 6$ times faster than the DReX ones. On the other hand, *reverse*, *swap_bibtex* and *align_bibtex* were hard to express in line-based tools such as sed and AWK. The Perl implementations of these functions were $\approx 2$ times faster than ours.

## 6. Related Work

***Regular string transformations.*** This class of string-to-string transformations is robust and has many equivalent characterizations, including deterministic two-way string transducers [16, 18], streaming string transducers [1], transducers with origin information [6], and MSO-definable string transformations [10]. Regular string transformations are also closed under composition [9], and

enjoy decidable equivalence [21]. Alur et al proposed a set of combinators that captures the set of regular string-to-string transformations [3] and this paper builds on it. The results in [3] only focus on expressiveness and do not try to answer questions about complexity. In particular the transformation to streaming string transducers proposed in [3] has non-elementary complexity in the size of the program. In this paper, however, we are primarily driven by issues related to the complexity of evaluation.

***DSLs for string transformations.*** DSLs for string transformations mainly fall into two classes: string specific utilities such as sed, AWK, and Perl, and transducer-based languages [13, 15, 26].

Utilities such as sed, AWK, and Perl provide the programmer with powerful programming constructs to manipulate strings. These languages are Turing-complete and in general cannot be efficiently compiled into fast executable code and are not amenable to algorithmic analysis. We also argued in section 5.3 that some of the programs that can be naturally expressed in DReX are actually hard to define using these tools.

All the existing transducer-based languages simply act as frontends to an underlying transducer model that they use to reason about the implemented programs. BEK uses symbolic finite transducers [26] and it has been used to analyze string sanitization functions. BEX is a frontend for extended symbolic finite transducers [12, 13] and it has been used to prove the correctness of string encoders and decoders such as Base64. FAST is based on symbolic tree transducers with regular look-ahead [15] and it is used to reason about programs that manipulate strings and trees over arbitrary domains. While these languages enable powerful analysis and verification techniques, (*a*) their semantics are tightly coupled to the transducer model, forcing the programmer to think in terms of a finite state machine, and a left-to-right reading of the input string, and (*b*) they only capture a strict subset of the class of regular string transformations; none of these models can reverse a string.

Another language based on automata is Boomerang, a bidirectional programming language for string editing [5]. Bidirectional programs contain combinators for extracting a view from a concrete input and then reconstructing an updated input from the updated view. Boomerang also supports extractions where each record is associated with a "key". Although Boomerang has a similar type-system to that of DReX (it forces unambiguous operations), we are not aware of a complexity analysis for the problem of evaluating a Boomerang program. The goals of Boomerang and DReX are orthogonal: the former focuses on bidirectional transformations, while the latter focuses on efficiently evaluating all regular string transformations. Despite this difference we believe that Boomerang could benefit from the evaluation techniques proposed in this paper.

***Efficient string manipulation.*** Little effort has been devoted to design languages and algorithms to efficiently evaluate string transformations in linear time. The general approach has been to identify an automaton model that processes the string in a single left-to-right pass and can express interesting programs. However, all existing tools that use this approach [13, 15, 26] take advantage of composition or combination operators that make the compilation to transducers exponential in the size of the program. Streaming string transducers (SST) [1] capture all programs that can be written in DReX and can be executed in a single left-to-right pass over the input. However, transforming DReX programs into SSTs also causes an super-exponential blow-up in the size of the input program.

In the context of XML processing numerous languages or fragments have been proposed for efficiently querying (XPath, XQuery), stream processing (STX [4]), and manipulating (XSLT) XML trees. Some of these languages particularly focus on efficiently processing the input in a linear time left-to-right pass. Although in the case of XML documents with bounded depth some XML

transformations can be described in DReX, the main goal of DReX remains that of providing a well-defined fragment (regular) of *string* (and not tree) transformations that can be efficiently executed.

***Future Directions.*** A major motivation for choosing the class of regular string transformations was the decidability of analysis questions. In particular, consider regular type-checking: given a program $f$ and two regular languages $I$ and $O$, is it the case that for every input $\sigma \in I$, $[\![p]\!](\sigma) \in O$? Such a tool would be helpful to audit string sanitizers against specific kinds of attacks. Implementing these procedures is an open research direction.

In FlashFill [20] simple string transformations can be synthesized from examples. The expressiveness of the combinators used in FlashFill has not been characterized. Can DReX programs be efficiently learnt or synthesized from input-output examples?

Recently Mytkowicz et al. [22] proposed new techniques for evaluating finite automata in a data-parallel fashion. Can these techniques be used to parallelize the evaluation algorithm proposed in this paper?

Extending our techniques to tree transformations is another open problem. Streaming tree transducers (STTs) [2] are to regular tree transformations (or equivalently macro tree transducers [19], and MSO-definable tree transformations [17]) as SSTs are related to regular string transformations. Can we design a similar declarative language to express regular tree transformations?

Finally, certain operations such as counting the number of substring matches, sorting the elements of a dictionary, or deleting duplicate entries in a list are *not* regular string transformations. Extending DReX with non-regular primitives but which can still be efficiently evaluated is an interesting direction of future work.

## 7. Conclusion

We presented DReX, a declarative language for describing regular string transformations. The basic transformers are symbolic, so DReX can succinctly express transformations even over large alphabets such as Unicode. We demonstrated that the evaluation problem for unrestricted DReX is PSPACE-complete, and so we considered a restricted fragment, *consistent* DReX which permits a fast single-pass evaluation algorithm, and still retains expressive completeness. In experiments over representative string transformations such as BibTeX file manipulations the evaluation algorithm for consistent DReX scaled to process thousands of characters per second.

## References

[1] R. Alur and P. Černý. Streaming transducers for algorithmic verification of single-pass list-processing programs. In *Proceedings of the 38th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 599–610. ACM, 2011.

[2] R. Alur and L. D'Antoni. Streaming tree transducers. In A. Czumaj, K. Mehlhorn, A. Pitts, and R. Wattenhofer, editors, *Automata, Languages, and Programming*, volume 7392 of *Lecture Notes in Computer Science*, pages 42–53. Springer, 2012.

[3] R. Alur, A. Freilich, and M. Raghothaman. Regular combinators for string transformations. In *Proceedings of the Joint Meeting of the 23rd EACSL Annual Conference on Computer Science Logic (CSL) and the 29th Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*, CSL-LICS '14, pages 9:1–9:10. ACM, 2014.

[4] O. Becker. Streaming transformations for xml-stx. In *XMIDX*, volume 24 of *LNI*, pages 83–88. GI, 2003.

[5] A. Bohannon, N. Foster, B. Pierce, A. Pilkiewicz, and A. Schmitt. Boomerang: Resourceful lenses for string data. In *Proceedings of the 35th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 407–419. ACM, 2008.

[6] M. Bojańczyk. Transducers with origin information. In *Automata, Languages, and Programming*, volume 8573 of *Lecture Notes in Computer Science*, pages 26–37. Springer, 2014.

[7] R. Book, S. Even, S. Greibach, and G. Ott. Ambiguity in graphs and expressions. *IEEE Transactions on Computers*, 20(2):149–153, February 1971.

[8] A. Brüggemann-Klein. Regular expressions into finite automata. In *LATIN '92*, volume 583 of *Lecture Notes in Computer Science*, pages 87–98. Springer, 1992.

[9] M. Chytil and V. Jákl. Serial composition of 2-way finite-state transducers and simple programs on strings. In *Automata, Languages, and Programming*, volume 52 of *Lecture Notes in Computer Science*, pages 135–147. Springer, 1977.

[10] B. Courcelle. Monadic second-order definable graph transductions: a survey. *Theoretical Computer Science*, 126(1):53 – 75, 1994.

[11] L. D'Antoni and R. Alur. Symbolic visibly pushdown automata. In *Computer Aided Verification*, volume 8559 of *Lecture Notes in Computer Science*, pages 209–225. Springer, 2014.

[12] L. D'Antoni and M. Veanes. Equivalence of extended symbolic finite transducers. In *Computer Aided Verification*, volume 8044 of *Lecture Notes in Computer Science*, pages 624–639. Springer, 2013.

[13] L. D'Antoni and M. Veanes. Static analysis of string encoders and decoders. In *Verification, Model Checking, and Abstract Interpretation*, volume 7737 of *Lecture Notes in Computer Science*, pages 209–228. Springer, 2013.

[14] L. D'Antoni and M. Veanes. Minimization of symbolic automata. In *Proceedings of the 41st ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 541–553, New York, NY, USA, 2014. ACM.

[15] L. D'Antoni, M. Veanes, B. Livshits, and D. Molnar. Fast: A transducer-based language for tree manipulation. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*, pages 384–394. ACM, 2014.

[16] J. Engelfriet and H. J. Hoogeboom. MSO definable string transductions and two-way finite-state transducers. *ACM Transactions on Computational Logic*, 2(2):216–254, April 2001.

[17] J. Engelfriet and S. Maneth. Macro tree transducers, attribute grammars, and MSO definable tree translations. *Information and Computation*, 154(1):34–91, 1999.

[18] J. Engelfriet, G. Rozenberg, and G. Slutzki. Tree transducers, L systems, and two-way machines. *Journal of Computer and System Sciences*, 20(2):150 – 202, 1980.

[19] J. Engelfriet and H. Vogler. Macro tree transducers. *Journal of Computer and System Sciences*, 31(1):71–146, 1985.

[20] S. Gulwani. Automating string processing in spreadsheets using input-output examples. In *Proceedings of the 38th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 317–330. ACM, 2011.

[21] E. Gurari. The equivalence problem for deterministic two-way sequential transducers is decidable. In *21st Annual Symposium on Foundations of Computer Science*, pages 83–85, 1980.

[22] T. Mytkowicz, M. Musuvathi, and W. Schulte. Data-parallel finite-state machines. In *Proceedings of the 19th International Conference on Architectural Support for Programming Languages and Operating Systems*, pages 529–542. ACM, 2014.

[23] G. Rosu. An effective algorithm for the membership problem for extended regular expressions. In *Foundations of Software Science and Computational Structures*, volume 4423 of *Lecture Notes in Computer Science*, pages 332–345. Springer, 2007.

[24] M. Sipser. *Introduction to the Theory of Computation*. Cengage Learning, 3rd edition, 2012.

[25] R. Stearns and H. Hunt. On the equivalence and containment problems for unambiguous regular expressions, grammars, and automata. In *Proceedings of the 22nd Annual Symposium on Foundations of Computer Science*, pages 74–81. IEEE Computer Society, 1981.

[26] M. Veanes, P. Hooimeijer, B. Livshits, D. Molnar, and N. Bjorner. Symbolic finite state transducers: Algorithms and applications. In *Proceedings of the 39th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, pages 137–150. ACM, 2012.

## A. Examples of Consistent DReX Programs

This appendix contains simplified descriptions of the consistent DReX programs that we used in our experiments in section 5. We will use the following macros in our definitions.

$$\mathtt{copy}(\varphi) = \varphi(x) \mapsto [x]$$
$$\mathtt{del}(\varphi) = \varphi(x) \mapsto [\,]$$
$$\mathtt{iterate\text{-}plus}(f) = \mathtt{split}(\mathtt{iterate}(f), f)$$
$$\mathtt{split}(f_1, \ldots, f_n) = \mathtt{split}(f_1, \mathtt{split}(\ldots, f_n))$$

$\mathtt{copy}(\varphi)$ and $\mathtt{del}(\varphi)$ respectively copy or delete a character matching $\varphi$. When the predicate is of the form $x = c$ we simply write $c$. $\mathtt{iterate\text{-}plus}(f)$ repeats the function $f$ one or more times. The last macro simply repeats the split sum operator on all the arguments.

### A.1 Delete one-line comments from a program

This program deletes all one-line comments from a file (i.e. the lines starting with //). We define each component of the program separately. The program *del_comm_line* deletes strings of the form "//$\sigma$\n", where $\sigma$ does not contain any occurrence of the new line character "\n".

$$del\_slashes = \mathtt{split}(\mathtt{del}('/'), \mathtt{del}('/')),$$
$$del\_non\_nl = \mathtt{iterate}(\mathtt{del}(x \neq '\backslash\mathtt{n}')),$$
$$del\_comm = \mathtt{split}(del\_slashes, del\_non\_nl),$$
$$del\_comm\_line = \mathtt{split}(del\_comm, \mathtt{del}('\backslash\mathtt{n}')).$$

The program *copy_line* copies strings of the form "$\sigma$\n", where the string $\sigma$ is either the empty string, or it does not start with the character '/'.

$$copy\_non\_nl = \mathtt{iterate}(\mathtt{copy}(x \neq '\backslash\mathtt{n}')),$$
$$copy\_txt = \mathtt{split}(\mathtt{copy}(x \neq '/'), copy\_non\_nl),$$
$$copy\_line = \mathtt{copy}('\backslash\mathtt{n}')\,\mathtt{else}\,\mathtt{split}(copy\_txt, \mathtt{copy}('\backslash\mathtt{n}')).$$

The program *process_line* reads a line and deletes it if it is a comment line and copies it otherwise. The last line might not end with a '\n' and the program *process_last_line* deals with this exception. Finally the program *delete_comm* repeats *process_line* and at the end processes the last line, therefore deleting all the one-line comments in the input.

$$process\_line = del\_comm\_line\,\mathtt{else}\,copy\_line,$$
$$last\_line = del\_comm\,\mathtt{else}\,copy\_txt,$$
$$process\_last\_line = process\_line\,\mathtt{else}\,last\_line,$$
$$process\_lines = \mathtt{iterate}(process\_line),$$
$$delete\_comm = \mathtt{split}(process\_lines, process\_last\_line).$$

### A.2 Insert quotes around words

This program inserts quotes ('"') around every alphabetic substring appearing in the input. The program *add_qts_skip*, given a string $\sigma_1\sigma_2$ where $\sigma_1$ is alphabetic and $\sigma_2$ does not contain any letter, outputs the string ""$\sigma_1$"$\sigma_2$".

$$copy\_astr = \mathtt{iterate\text{-}plus}(\mathtt{copy}([\mathtt{a\text{-}zA\text{-}Z}])),$$
$$copy\_nastr = \mathtt{iterate\text{-}plus}(\mathtt{copy}(x \notin [\mathtt{a\text{-}zA\text{-}Z}])),$$
$$add\_qt = \epsilon \mapsto ['"'],$$
$$add\_qts = \mathtt{split}(add\_qt, copy\_astr, add\_qt),$$
$$add\_qts\_skip = \mathtt{split}(add\_qts, copy\_nastr).$$

The program *insert_quotes* repeats the *add_qts* function. Since the string might start with a symbol that is not alphabetic the program *start* deals with this case. Similarly the program *ending* checks whether the string does not end with an alphabetic sequence. Finally,

the program *insert_quotes* inserts quotes around every alphabetic substring in the input.

$$start = \mathtt{iterate}(\mathtt{copy}(x \notin [\mathtt{a\text{-}zA\text{-}Z}])),$$
$$ending = add\_qts\,\mathtt{else}\,\epsilon \mapsto [\,],$$
$$repeat\_add = \mathtt{iterate}(add\_qts\_skip),$$
$$insert\_quotes = \mathtt{split}(start, repeat\_add, ending).$$

### A.3 Extracting tags from a malformed XML file

This program extracts and concatenates all the substrings of the form <$\sigma$> where $\sigma$ does not contain any character < or > (this is a generalization of a program shown in [26]). For simplicity we assume that the string does not contain occurrences of the substring "<>". The program *copy_match* copies any string of the form <$\sigma$> where $\sigma$ does not contain any character < or >.

$$copy\_ntag = \mathtt{iterate\text{-}plus}(\mathtt{copy}(x \notin [\mathtt{<>}])),$$
$$copy\_match = \mathtt{split}(\mathtt{copy}('<'), copy\_ntag, \mathtt{copy}('>')).$$

The program *del_not_match* deletes any string that does not contain a substring of the form <$s$> where $s$ does not contain any character < or >. The program *del_not_match* looks for the following pattern: $s = \sigma_1 > \cdots > \sigma_i < \cdots < \sigma_n$ (with $i \geq 1$). Its sub-program *del_close_op* deletes all the string of the described form containing at least one open character ('<').

$$del\_not\_opn = \mathtt{iterate}(\mathtt{del}(x \neq '<')),$$
$$del\_not\_cls = \mathtt{iterate}(\mathtt{del}(x \neq '>')),$$
$$del\_opn\_not\_cls = \mathtt{split}(\mathtt{del}('<'), \mathtt{iterate}(\mathtt{del}(x \neq '>'))),$$
$$del\_close\_op = \mathtt{split}(del\_not\_opn, del\_opn\_not\_cls),$$
$$del\_not\_match = del\_not\_opn\,\mathtt{else}\,del\_close\_op.$$

The program *find_match* keeps looking for tags and removes eventual non-matches at the end of the string. Finally, *get_tags* repeats *find_match*, and therefore outputs all the substrings of the form <$\sigma$>.

$$find\_match = \mathtt{split}(del\_not\_match, copy\_match),$$
$$repeat\_get\_tags = \mathtt{iterate}(find\_match),$$
$$get\_tags = \mathtt{split}(repeat\_get\_tags, del\_not\_match).$$

### A.4 Reversing a dictionary

Given a dictionary of the form $\sigma_1; \sigma_2; \cdots; \sigma_n;$ we want to output the reverse $\sigma_n; \sigma_{n-1}; \cdots; \sigma_1;$. The program *copy_stretch* copies a string of the form $\sigma;$ such that $\sigma$ does not contain a ';'. The program *reverse* implements the final transformation by left iterating *copy_stretch*.

$$copy\_stretch = \mathtt{split}(\mathtt{iterate}(\mathtt{copy}(x \neq ';')), \mathtt{copy}(';')),$$
$$reverse = \mathtt{left\text{-}iterate}(copy\_stretch).$$

### A.5 Reformatting BibTeX files

In this section we define two functions that operate over BibTeX files. The first function, *swap_bibtex*, reorders attributes within each BibTeX entry by moving the title to the top. The second function, *align_bibtex*, rearranges a *misaligned* file by moving the title of each entry inside the previous entry.

To do so we first define a few auxiliary functions that are used for copying and deleting alphabetic strings, spaces, and delimiters.

$$copy\_astr = \mathtt{iterate\text{-}plus}(\mathtt{copy}([\mathtt{a\text{-}zA\text{-}Z}])),$$
$$copy\_anum = \mathtt{iterate\text{-}plus}(\mathtt{copy}([\mathtt{a\text{-}zA\text{-}Z0\text{-}9}])),$$
$$copy\_spaces = \mathtt{iterate}(\mathtt{copy}([\backslash\mathtt{n}\ \backslash\mathtt{r}\ \backslash\mathtt{b}\ \backslash\mathtt{t}])),$$
$$del\_astr = \mathtt{iterate\text{-}plus}(\mathtt{del}([\mathtt{a\text{-}zA\text{-}Z}])).$$

```
@book{Gal1638,                          @book{Gal1638,
    publisher = {Elzevir},                  title = {Two New Sciences},
    place = {Leiden},                       publisher = {Elzevir},
    year = {1638},                          place = {Leiden},
    title = {Two New Sciences},             year = {1638},
    author = {Galileo},                     author = {Galileo},
}                                       }
```

Figure A.1: Example application of the transformation *swap_bibtex* from the entry on the left to the entry on the right.

It is easy to see how given a DReX program that copies a pattern, one can easily define a DReX program that deletes the same pattern. To simplify the presentation in the following we assume that for every program of the form *copy_something* there is an analogous program *del_something* that is defined on the same input as *copy_something* but always outputs the empty string.

The program *copy_header* copies the header of an entry. In the example of figure A.1, it copies the string "`@book{Gal1638`" along with the following spaces.

$$copy\_header = \texttt{split}(\texttt{copy}(`@'), copy\_astr, \texttt{copy}(`\{'),$$
$$copy\_anum, \texttt{copy}(`,'), copy\_spaces).$$

We now define macros for copying and deleting a particular string $s$ or a set of strings.

$$\texttt{copy}(s) = \texttt{split}(\texttt{copy}(s[1]), \dots, \texttt{copy}(s[|s|])),$$
$$\texttt{copy}(\{\sigma_1, \dots, \sigma_n\}) = \texttt{copy}(\sigma_1) \, \texttt{else} \dots \texttt{else} \, \texttt{copy}(\sigma_n).$$

The program *copy_title* copies the string "`title`", and the program *copy_non_title* copies every attribute name different from "`title`". We omit the full list of attributes for readability.

$$copy\_title = \texttt{copy}(`\texttt{title}'),$$
$$copy\_non\_title = \texttt{copy}(\{`\texttt{author}', `\texttt{year}', `\texttt{place}', \dots\}).$$

The program *copy_att_value* copies the value of an attribute along with the surrounding parentheses (i.e. "` = { Elzevir},`").

$$copy\_non\_par = \texttt{iterate}(\texttt{copy}(x \notin [\}\{])),$$
$$copy\_att\_value = \texttt{split}(copy\_spaces, \texttt{copy}(`='),$$
$$copy\_spaces, \texttt{copy}(`\{'), copy\_non\_par,$$
$$\texttt{copy}(`\}'), \texttt{copy}(`,'), copy\_spaces).$$

The program *copy_title_att* copies a complete title attribute (i.e. "`title = {Two New Sciences},`"), while *copy_non_title_att* copies a complete non-title attribute.

$$copy\_title\_att = \texttt{split}(copy\_title, copy\_att\_value),$$
$$copy\_non\_title\_att = \texttt{split}(copy\_non\_title, copy\_att\_value).$$

Given a list of attributes, the program *title_only* copies the title and deletes all the non-title attributes, the program *all_but_title* deletes the title and copies all the other attributes, and the program *copy_attrs* copies the entire list.

$$title\_only = \texttt{iterate}(copy\_title \, \texttt{else} \, del\_non\_title),$$
$$all\_but\_title = \texttt{iterate}(del\_title \, \texttt{else} \, copy\_non\_title),$$
$$all\_but\_title = \texttt{iterate}(copy\_title \, \texttt{else} \, copy\_non\_title).$$

### A.5.1 Polishing a BibTeX file

The following function defines a typical transformation a paper author may perform on BibTeX files. The program $swap\_bibtex$ moves the title attribute to the top of each entry of an input BibTeX file. Figure A.1 shows the result of applying $swap\_bibtex$ to a particular BibTeX entry. In this presentation we assume that every attribute is followed by a comma.[6] The program *title_first* copies the title first and then all the other attributes, therefore obtaining the desired attribute reordering.

$$title\_first = \texttt{combine}(title\_only, all\_but\_title).$$

The program *swap_entry* performs the operation of figure A.1 for a single BibTeX entry.

$$swap\_entry = \texttt{split}(copy\_header, title\_first,$$
$$\texttt{copy}(`\}'), copy\_spaces).$$

Finally, the program *swap_bibtex* applies the transformation to all the entries in the file.

$$swap\_bibtex = \texttt{iterate}(swap\_entry).$$

### A.5.2 Aligning titles in a misaligned BibTeX file

Assume that by mistake the titles of a BibTeX file have been misaligned: the title of the first entry now appears in the second one, the second title appears in the third entry and so on. The function *align_bibtex*, given such a misaligned file, moves the title of each entry $i + 1$ to the correct entry $i$, and since the last entry would now not have a title, it gets deleted. This program requires the use of the chained sum operator in order to process two BibTeX entries at a time (first and second, second and third, and so on). The program *header_only_entry*, given a single BibTeX entry, copies only the header and deletes the rest of the entry.

$$header\_only\_entry = \texttt{split}(copy\_header, del\_attrs,$$
$$\texttt{del}(`\}'), del\_spaces).$$

The program *title_only_entry*, given a single BibTeX entry, copies only the title and deletes the rest of the entry, while the program *all_but_title_entry*, deletes the header, the title, and copies the rest of the entry.

$$title\_only\_entry = \texttt{split}(del\_header, title\_only,$$
$$\texttt{del}(`\}'), del\_spaces),$$
$$all\_but\_title\_entry = \texttt{split}(del\_header, all\_but\_title,$$
$$\texttt{copy}(`\}'), copy\_spaces).$$

The program *delete_entry* deletes a single BibTeX entry.

$$delete\_entry = \texttt{split}(del\_header, del\_attrs,$$
$$\texttt{del}(`\}'), del\_spaces).$$

The program *move_title*, given two BibTeX entries outputs the first entry where the title has been replaced with the title of the second entry: it first copies the header of the first entry, then the title of the second entry, and then the remaining attributes of the first entry.

$$move\_title = \texttt{combine}($$
$$\texttt{split}(header\_only\_entry, title\_only\_entry),$$
$$\texttt{split}(all\_but\_title\_entry, delete\_entry)).$$

Finally, the program *align_bibtex* applies the transformation to all the entries in the file. We omit the regular expression describing the format of an entry in the chain operator since DReX can infer it automatically.

$$align\_bibtex = \texttt{chain}(move\_title).$$

---

[6] The actual program used for the experiments is able to deal with missing commas.