

Stably Computable Properties of Network Graphs

Dana Angluin, James Aspnes, Melody Chan, Michael J. Fischer,
Hong Jiang, and René Peralta

Yale University

Abstract. We consider a scenario in which anonymous, finite-state sensing devices are deployed in an ad-hoc communication network of arbitrary size and unknown topology, and explore what properties of the network graph can be stably computed by the devices. We show that they can detect whether the network has degree bounded by a constant d , and, if so, organize a computation that achieves asymptotically optimal linear memory use. We define a model of stabilizing inputs to such devices and show that a large class of predicates of the multiset of final input values are stably computable in any weakly-connected network. We also show that nondeterminism in the transition function does not increase the class of stably computable predicates.

1 Introduction

In some applications, a large number of sensors will be deployed without fine control of their locations and communication patterns in the target environment. To enable the distributed gathering and processing of information, the sensors must constitute themselves into an ad-hoc network and use it effectively. A fundamental question in this context is whether there are protocols that determine enough about the topological properties of this network to exploit its full potential for distributed computation, when sensors are severely limited in their computational power.

We consider a model introduced in [1] in which communication is represented by pairwise interactions of anonymous finite-state devices in networks of finite but unbounded size. These systems correspond to sensor networks with intermittent two-way communications between nearby nodes. A **communication graph** describes which pairs of nodes are close enough to eventually interact. Our goal is to explore what graph-theoretic properties of the communication graph are **stably computable**, where a property is stably computable if all sensors eventually converge to the correct answer. The model assumes that the devices have no identifiers and only a few bits of memory each. Because each device is so limited, their collective ability to achieve nontrivial computation must be based on their capacity to organize a distributed computation in the network.

In this setting, the structure of the network has a profound influence on its computational potential. If n is the number of vertices in the communication graph, previous results show that $O(\log n)$ bits of memory are sufficient for a nondeterministic Turing machine to simulate any protocol in the all-pairs communication graph, in which every pair of vertices is joined by an edge [1]. In contrast, we give a protocol that can determine whether the communication graph is a directed cycle and if so, use it as a linear memory of $O(n)$ bits, which is asymptotically optimal in terms of memory capacity. More generally, we show that for every d there is a protocol that can organize any communication graph of maximum degree d into a linear memory of $O(n)$ bits, also asymptotically optimal.

For general communication graphs, we show that any property that is determined by the existence of a fixed finite subgraph is stably computable, as are Boolean combinations of such properties. In addition, there are protocols to compute the following graph properties: whether the communication graph G is a directed star, whether G is a directed arborescence, whether G contains a directed cycle, and whether G contains a directed cycle of odd length. Furthermore, for any positive integer d , there is a protocol that stabilizes to a proper d -coloring of any d -colorable graph, but does not stabilize if the graph is not d -colorable.

In the model of [1], the sensor readings were all assumed to be available as inputs at the start of the computation. In this paper, we extend the model to allow for a more realistic scenario of stabilizing inputs, that may change finitely many times before attaining a final value. In addition to allowing fluctuations in the inputs, these results allow composition of stabilizing protocols: a protocol that works with stabilizing inputs can use the stabilizing outputs of another protocol. We generalize two fundamental theorems to the case of stabilizing inputs: all the **Presburger-definable** predicates are stably computable in the all-pairs graph, and any predicate stably computable in the all-pairs graph is stably computable in any weakly connected graph with the same number of nodes.

Another powerful tool for the design of protocols is to permit a nondeterministic transition function; we give a simulation to show that this does not increase the class of stably computable predicates.

1.1 Other Related Work

Population protocols and similar models as defined in [1–3] have connections to a wide range of theoretical models and problems involving automata of various kinds, including Petri nets [4–7], semilinear sets and Presburger expressions [8], vector addition systems [9], the Chemical Abstract Machine [10, 11] and other models of communicating automata [12, 13]. See [1] for a more complete discussion of these connections.

The potential for distributed computation during aggregation of sensor data is studied in [14, 15], and distributed computation strategies for conserving resources in tracking targets in [16, 17]. Issues of random mobility in a wireless packet network are considered in [18].

Passively mobile sensor networks have been studied in several practical application contexts. The **smart dust** project [19,20] designed a cloud of tiny wireless microelectromechanical sensors (MEMS) with wireless communication capacity, where each sensor or “mote” contains sensing units, computing circuits, bidirectional wireless capacity, and a power supply, while being inexpensive enough to deploy massively.

2 The Model of Stable Computation

We represent a network communication graph by a directed graph $G = (V, E)$ with n vertices numbered 1 through n and no multi-edges or self-loops. Each vertex represents a finite-state sensing device, and an edge (u, v) indicates the possibility of a communication between u and v in which u is the initiator and v is the responder.¹ We assume G is *weakly connected*, that is, between any pair of nodes there is a path (disregarding the direction of the edges in the path). The **all-pairs graph** contains all edges (u, v) such that $u \neq v$.

We first define protocols without inputs, which is sufficient for our initial results on graph properties, and extend the definition to allow stabilizing inputs in Section 4.

A **protocol** consists of a finite set of **states** Q , an **initial state** $q_0 \in Q$, an **output function** $O : Q \rightarrow Y$, where Y is a finite set of output symbols, and a **transition function** δ that maps every pair of states (p, q) to a nonempty set of pairs of states. If $(p', q') \in \delta(p, q)$, we call $(p, q) \mapsto (p', q')$ a **transition**.

The transition function, and the protocol, is **deterministic** if $\delta(p, q)$ always contains just one pair of states. In this case we write $\delta(p, q) = (p', q')$ and define $\delta_1(p, q) = p'$ and $\delta_2(p, q) = q'$.

A **configuration** is a mapping $C : V \rightarrow Q$ specifying the state of each device in the network. We assume that there is a global start signal transmitted simultaneously to all the devices, e.g., from a base station, that puts them all in the initial state and starts the computation. The **initial configuration** assigns the initial state q_0 to every device.

Let C and C' be configurations, and let u, v be distinct nodes. We say that C goes to C' via **pair** $e = (u, v)$, denoted $C \xrightarrow{e} C'$, if the pair $(C'(u), C'(v))$ is in $\delta(C(u), C(v))$ and for all $w \in V - \{u, v\}$ we have $C'(w) = C(w)$. We say that C can go to C' in one step, denoted $C \rightarrow C'$, if $C \xrightarrow{e} C'$ for some edge $e \in E$. We write $C \xrightarrow{*} C'$ if there is a sequence of configurations $C = C_0, C_1, \dots, C_k = C'$, such that $C_i \rightarrow C_{i+1}$ for all i , $0 \leq i < k$, in which case we say that C' is **reachable** from C .

A **computation** is a finite or infinite sequence of population configurations C_0, C_1, \dots such that for each i , $C_i \rightarrow C_{i+1}$. An infinite computation is **fair** if for every pair of population configurations C and C' such that $C \rightarrow C'$, if C

¹ The distinct roles of the two devices in an interaction is a fundamental assumption of asymmetry in our model; symmetry-breaking therefore does not arise as a problem within the model.

occurs infinitely often in the computation, then C' also occurs infinitely often in the computation. We remark that by induction, if $C \xrightarrow{*} C'$ and C occurs infinitely often in a fair computation, then C' also occurs infinitely often in a fair computation.

2.1 Leader Election

As an example, we define a simple deterministic leader election protocol that succeeds in any network communication graph. The states of the protocol are $\{1, 0\}$ where 1 is the initial state. The transitions are defined by

$$\begin{aligned} (1) \quad (1, 1) &\mapsto (0, 1) \\ (2) \quad (1, 0) &\mapsto (0, 1) \\ (3) \quad (0, 1) &\mapsto (1, 0) \\ (4) \quad (0, 0) &\mapsto (0, 0) \end{aligned}$$

We think of 1 as the leader mark. In every infinite fair computation of this protocol starting with the initial configuration in any communication graph, after some finite initial segment of the computation, every configuration has just one vertex labeled 1 (the leader), and every vertex has label 1 in infinitely many different configurations of the computation. Thus, eventually there is one “leader” mark that hops incessantly around the graph, somewhat like an ancient English king visiting the castles of his lords. Note that in general the devices have no way of knowing whether a configuration with just one leader mark has been reached yet.

2.2 The Output of a Computation

Our protocols are not designed to halt, so there is no obvious fixed time at which to view the output of the computation. Rather, we say that the output of the computation stabilizes if it reaches a point after which no device can subsequently change its output value, no matter how the computation proceeds thereafter. Stability is a global property of the graph configuration, so individual devices in general do not know when stability has been reached.²

We define an **output assignment** y as a mapping from V to the output symbols Y . We extend the output map O to take a configuration C and produce an output assignment $O(C)$ defined by $O(C)(v) = O(C(v))$. A configuration C is said to be **output-stable** if $O(C') = O(C)$ for all C' reachable from C . Note that we do not require that $C = C'$, only that their output assignments be equal. An infinite computation **output-stabilizes** if it contains an output-stable configuration C , in which case we say that it **stabilizes to output assignment** $y = O(C)$. Clearly an infinite computation stabilizes to at most one output assignment.

² With suitable stochastic assumptions on the rate at which interactions occur, it is possible to bound the expected number of interactions until the output stabilizes, a direction explored in [1].

The output of a finite computation is the output assignment of its last configuration. The output of an infinite computation that stabilizes to output assignment y is y ; the output is undefined if the computation does not stabilize to any output assignment. Because of the nondeterminism inherent in the choice of encounters, the same initial configuration may lead to different computations that stabilize to different output assignments.

2.3 Graph Properties

We are interested in what properties of the network communication graph can be stably computed by protocols in this model. A **graph property** is a function P from graphs G to the set $\{0, 1\}$ where $P(G) = 1$ if and only if G has the corresponding property. We are interested in families of graphs, $\mathcal{G}_1, \mathcal{G}_2, \dots$, where \mathcal{G}_n is a set of network graphs with n vertices. The **unrestricted** family of graphs contains all possible network communication graphs. The **all-pairs** family of graphs contains for each n just the all-pairs graph with n vertices. For every d , the family of **d -degree-bounded** graphs contains all the network communication graphs in which the in-degree and out-degree of each vertex is bounded by d . Similarly, the family of **d -colorable** graphs contains all the network communication graphs properly colorable with at most d colors.

We say that a protocol \mathcal{A} stably computes the graph property P in the family of graphs $\mathcal{G}_1, \mathcal{G}_2, \dots$ if for every graph G in the family, every infinite fair computation of \mathcal{A} in G starting with the initial configuration stabilizes to the constant output assignment equal to $P(G)$. Thus the output of every device stabilizes to the correct value of $P(G)$.

3 Example: Is G a Directed Cycle?

In this section, we show that the property of G being a directed cycle is stably computable in the unrestricted family of graphs. Once a directed cycle is recognized, it can be organized (using leader-election techniques) to simulate a Turing machine tape of n cells for the processing of inputs, which vastly increases the computational power over the original finite-state devices, and is optimal with respect to the memory capacity of the network.

Theorem 1. *Whether G is a directed cycle is stably computable in the unrestricted family of graphs.*

Proof sketch: A weakly connected directed graph G is a directed cycle if and only if the in-degree and out-degree of each vertex is 1.

Lemma 1. *Whether G has a vertex of in-degree greater than 1 is stably computable in the unrestricted family of graphs.*

We give a protocol to determine whether some vertex in G has in-degree at least 2. The protocol is deterministic and has 4 states: $\{-, I, R, Y\}$, where $-$ is

the initial state, I and R stand for “initiator” and “responder” and Y indicates that there is a vertex of in-degree at least 2 in the graph. The transitions are as follows, where x is any state and unspecified transitions do not change their inputs:

- (1) $(-, -) \mapsto (I, R)$
- (2) $(I, R) \mapsto (-, -)$
- (3) $(-, R) \mapsto (Y, Y)$
- (4) $(Y, x) \mapsto (Y, Y)$
- (5) $(x, Y) \mapsto (Y, Y)$

The output map takes Y to 1 and the other states to 0. State Y is contagious, spreading to all states if it ever occurs. If every vertex has in-degree at most 1 then only transitions of types (1) and (2) can occur, and if some vertex has in-degree at least 2, then in any fair computation some transition of type (3) must eventually occur.

An analogous four-state protocol detects whether any vertex has out-degree at least 2. We must also guarantee that every vertex has positive in-degree and out-degree. The following deterministic two-state protocol stably labels each vertex with Z if it has in-degree 0 and P if it has in-degree greater than 0. The initial state is Z . The transitions are given by the following, where x and y are any states:

- (1) $(x, y) \mapsto (x, P)$

To detect whether all states are labeled P in the limit, we would like to treat the outputs of this protocol as the inputs to a another protocol to detect any occurrences of Z in the limit. However, to do this, the protocol to detect any occurrences of Z must cope with inputs that may change before they stabilize to their final values. In the next section we show that all the Presburger predicates (of which the problem of Z detection is a simple instance) are stably computable with such “stabilizing inputs” in the unrestricted family of graphs, establishing the existence of the required protocol.

Similarly, there is a protocol that stably computes whether every vertex has out-degree at least 1. By running all four protocols in parallel, we may stably compute the property: does every vertex have in-degree and out-degree exactly 1? Thus, there is a protocol that stably computes the property of G being a directed cycle for the unrestricted family of graphs, proving Theorem 1. ■

These techniques generalize easily to recognize other properties characterized by conditions on vertices having in-degrees or out-degrees of zero or one. A directed line has one vertex of in-degree zero, one vertex of out-degree zero, and all other vertices have in-degree and out-degree one. An out-directed star has one vertex of in-degree zero and all other vertices of in-degree one and out-degree zero, and similarly for an in-directed star. An out-directed arborescence has one vertex of in-degree zero and all other vertices have in-degree one, and similarly for an in-directed arborescence.

Theorem 2. *The graph properties of being a directed line, a directed star, or a directed arborescence are stably computable in the unrestricted family of graphs.*

In a later section, we generalize Theorem 1 to show that for any d there is a protocol to recognize whether the network communication graph has in-degree and out-degree bounded by d and to organize it as $O(n)$ cells of linearly ordered memory if so. We now turn to the issue of inputs.

4 Computing with Stabilizing Inputs

We define a model of stabilizing inputs to a network protocol, in which the input to each node may change finitely many times before it stabilizes to a final value. We are interested in what predicates of the final input assignment are stably computable. An important open question is whether any predicate stably computable in a family of graphs is stably computable with stabilizing inputs in the same family of graphs.

Though we do not fully answer this question, we show that a large class of protocols can be adapted to stabilizing inputs by generalizing two theorems from [1] to the case of stabilizing inputs: all Presburger predicates are stably computable with stabilizing inputs in the all-pairs family of graphs, and any predicate that can be stably computed with stabilizing inputs in the all-pairs family of graphs can be stably computed with stabilizing inputs in the unrestricted family of graphs.

We assume that there is a finite set of input symbols, and each device has a separate input port at which its current input symbol (representing a sensed value) is available at every computation step. Between any two computation steps, the inputs to any subset of the devices may change arbitrarily.

In the full paper, we define formally an extension of the basic stable computation model that permits protocols with stabilizing inputs. This extension is used to prove several results concerning such protocols, including:

Lemma 2. *Any Boolean combination of a finite set of predicates stably computable with stabilizing inputs in a family of graphs $\mathcal{G}_1, \mathcal{G}_2, \dots$ is stably computable with stabilizing inputs in the same family of graphs.*

Theorem 3. *For every nondeterministic protocol \mathcal{A} there exists a deterministic protocol \mathcal{B} such that if \mathcal{A} stably computes a predicate P with stabilizing inputs in a family of graphs, then \mathcal{B} also stably computes P with stabilizing inputs in the same family of graphs.*

4.1 The Presburger Predicates

The Presburger graph predicates form a useful class of predicates on the multiset of input symbols, including such things as “all the inputs are a’s”, “at least 5 inputs are a’s”, “the number of a’s is congruent to 3 modulo 5”, and “twice the number of a’s exceeds three times the number of b’s” and Boolean combinations of such predicates. Every Presburger graph predicate is stably computable with unchanging inputs in the all-pairs family of graphs [1].

We define the Presburger graph predicates as those expressible in the following expression language.³ For each input symbol $a \in X$, there is a variable $\#(a)$ that represents the number of occurrences of a in the input. A **term** is a linear combination of variables with integer coefficients, possibly modulo an integer. An **atom** is two terms joined by one of the comparison operators: $<, \leq, =, \geq, >$. An **expression** is a Boolean combination of atoms.

Thus, if the input alphabet $X = \{a, b, c\}$, the predicate “all the inputs are a’s” can be expressed as $\#(b) + \#(c) = 0$, the predicate that the number of a’s is congruent to 3 modulo 5 can be expressed as $(\#(a) \bmod 5) = 3$, and the predicate that twice the number of a’s exceeds three times the number of b’s by $2\#(a) > 3\#(b)$.

In the full paper, we show:

Theorem 4. *Every Presburger graph predicate is stably computable with stabilizing inputs in the family of all-pairs graphs.*

Theorem 5. *For any protocol \mathcal{A} there exists a protocol \mathcal{B} such that for every n , if \mathcal{A} stably computes predicate P with stabilizing inputs in the all-pairs interaction graph with n vertices and G is any communication graph with n vertices, protocol \mathcal{B} stably computes predicate P with stabilizing inputs in G .*

The following corollary is an immediate consequence of Theorems 4 and 5.

Corollary 1. *The Presburger predicates are stably computable with stabilizing inputs in the unrestricted family of graphs.*

5 Computing in Bounded-Degree Networks

In Section 3 we showed that there is a protocol that stably computes whether G is a directed cycle. If G is a directed cycle, another protocol can organize a Turing machine computation of space $O(n)$ in the graph to determine properties of the inputs. Thus, certain graph structures can be recognized and exploited to give very powerful computational capabilities.

In this section we generalize this result to the family of graphs with degree at most d .

Lemma 3. *For every positive integer d , there is a protocol that stably computes whether G has maximum degree less than or equal to d .*

Proof sketch: We describe a protocol to determine whether G has any vertex of out-degree greater than d ; in-degree is analogous, and the “nor” of these properties is what is required. The protocol is nondeterministic and based on leader election. Every vertex is initially a leader and has output 0. Each leader repeatedly engages in the following searching behavior. It hops around

³ These are closely related to the Presburger integer predicates defined in [1]. Details will be given in the full paper.

the vertices of G and decides to check a vertex by marking it and attempting to place $d + 1$ markers of a second type at the ends of edges outgoing from the vertex being checked. It may decide to stop checking, collect a set of markers corresponding to the ones it set out, and resume its searching behavior. If it succeeds in placing all its markers, it changes its output to 1, hops around the graph to collect a set of markers corresponding to the ones it set out, and resumes its searching behavior.

When two leaders meet, one becomes a non-leader, and the remaining leader collects a set of markers corresponding to the ones that both had set out, changes its output back to 0, and resumes its searching behavior. Non-leaders copy the output of any leader they meet.

After the leader becomes unique and collects the set of markers that it and the last deposed leader had set out, then the graph is clear of markers and the unique leader resumes the searching behavior with its output set to 0. If there is no vertex of out-degree greater than d , the output will remain 0 (and will eventually be copied by all the non-leaders.) If there is some vertex of out-degree greater than d , the searching behavior eventually finds it and sets the output to 1, which will eventually be copied by all the non-leaders. ■

Note that this technique can be generalized (using a finite collection of distinguishable markers) to determine the existence of any fixed subgraph in G . (For the lemma, the fixed subgraph is the out-directed star on $d + 2$ vertices.) Another variant of this idea (mark a vertex and try to reach the mark by following directed edges) gives protocols to determine whether G contains a directed cycle or a directed odd cycle.

Theorem 6. *There are protocols that stably compute whether G contains a fixed subgraph, or a directed cycle, or a directed cycle of odd length in the unrestricted family of graphs.*

For bounded-degree graphs, we can organize the nodes into a spanning tree rooted at a leader, which can then simulate a Turing machine tape of size $O(n)$ distributed across the nodes. This allows a population with a bounded-degree interaction graph to compute any function of the graph structure that is computable in linear space.

Theorem 7. *For every positive integer d , there is a protocol that for any d -bounded graph G stably constructs a spanning tree structure in G that can be used to simulate n cells of Turing machine tape.*

Proof sketch: The protocol is rather involved; details are deferred to the full paper. We give an outline of the protocol here.

The starting point is to label the vertices of G in such a way that no two neighbors of any vertex have the same label. A vertex can then send messages to a specific neighbor by waiting to encounter another vertex with the appropriate label.

To see that such a labeling exists, consider the graph G' obtained from G by ignoring the direction of edges and including an edge between any two nodes at

distance 2 in G . A proper coloring of G' will give the required labeling of G , and the degree of G' is at most $4d^2$, so G' can be colored with $4d^2 + 1$ colors.

One part of the protocol eventually constructs a labeling of the desired kind by using leader election and a searching behavior that attempts to find two vertices at distance two that have the same label and nondeterministically relabel both. Eventually there will be no more relabeling.

The other part of the protocol attempts to use the constructed labeling to build a spanning tree structure in G . This is also based on leader election. Each leader begins building a spanning tree from the root, recording in each vertex the labels of the known neighbors of the vertex, and designating (by label) a parent for each non-root vertex. The leader repeatedly traverses its spanning tree attempting to recruit new vertices. When it meets another leader, one becomes a non-leader and the other begins building a new spanning tree from scratch.

The labeling portion of the protocol is running in parallel with the tree-construction, and it produces a “restart” marker whenever it relabels vertices. When a leader encounters a “restart” marker, it deletes the marker and again begins building a new spanning tree from scratch.

Eventually all the relabeling will be completed, and only one leader will remain, and the final spanning tree construction will not be restarted. However, it is important that the spanning tree construction be able to succeed given a correct labeling but otherwise arbitrary states left over from preceding spanning tree construction attempts. The leader repeatedly traverses the constructed spanning tree, setting a phase indicator (0 or 1) at each pass to detect vertices that are not yet part of the tree. An arbitrary ordering on the labels gives a fixed traversal order for the spanning tree, and a portion of each state can be devoted to simulating a Turing machine tape cell.

Because a leader cannot determine when it is unique, when the labeling is stable, or when the spanning tree is complete, it simply restarts computation in the tree each time it begins rebuilding a spanning tree. Eventually, however, the computation has access to n simulated cells of tape. ■

Itkis and Levin use a similar construction in their self-stabilizing protocols for identical nameless nodes in an asynchronous communication network (represented by an undirected, connected, reflexive graph G) with worst-case transient faults [21]. In their model, in addition to a finite number of bits of storage, each node x maintains a finite set of pointers to itself and its neighbors, and can detect whether a neighbor y has a pointer to x and/or y , and can set a pointer to point to the first (in a fixed ordering) neighbor with a given property. This additional information about the graph structure permits them to exploit storage $O(|V|)$ in all cases. By comparison, in our model there are no pointers and each device truly has only a constant amount of memory regardless of the size and topology of the network. For this reason, in an all-pairs graph of size n , only memory proportional to $\log n$ is achievable. So the bounded-degree restriction of Theorem 7, or some other limitation that excludes the all-pairs graph, appears to be necessary.

6 Discussion and Open Problems

It is open whether every predicate stably computable with unchanging inputs in a family of graphs is stably computable with stabilizing inputs in the same family of graphs. For the family of all-pairs graphs, both classes contain the Presburger predicates. It follows that if, in this family, some predicate is stably computable with unchanging inputs but not with stabilizing inputs, that predicate would not be Presburger. The existence of such a predicate would disprove a conjecture from [1].

A natural measure of the information-theoretic memory capacity of a graph G with a finite set of vertex labels L is the log of number of different isomorphism classes of G with vertex labels drawn from L . If L has at least two elements, the all-pairs graphs have memory capacity $\Theta(\log n)$, which can be exploited for computation in the setting of randomized interactions and computation with errors [1]. When the cardinality of L is large enough compared to d ($O(d^2)$ suffices), the memory capacity of d -degree bounded graphs is $\Theta(n)$, which we have shown above can be exploited by protocols for stable computation. Are there protocols to exploit the full-information theoretic memory capacity of arbitrary network communication graphs?

An important future direction is to study the running time of protocols in this model, perhaps under a stochastic model of pairwise interactions, as in the model of conjugating automata [1].

Acknowledgments

James Aspnes was supported in part by NSF grants CCR-9820888, CCR-0098078, CNS-0305258, and CNS-0435201. Michael J. Fischer and René Peralta were supported in part by NSF grant CSE-0081823. Hong Jiang was supported by NSF grant ITR-0331548.

The authors would like to thank David Eisenstat for helpful comments on the paper. The authors would also like to thank the anonymous referees for DCOSS 2005 for their many useful suggestions.

References

1. Angluin, D., Aspnes, J., Diamadi, Z., Fischer, M.J., Peralta, R.: Computation in networks of passively mobile finite-state sensors. In: PODC '04: Proceedings of the twenty-third annual ACM symposium on Principles of distributed computing, ACM Press (2004) 290–299
2. Angluin, D., Aspnes, J., Diamadi, Z., Fischer, M.J., Peralta, R.: Urn automata. Technical Report YALEU/DCS/TR-1280, Yale University Department of Computer Science (2003)
3. Diamadi, Z., Fischer, M.J.: A simple game for the study of trust in distributed systems. Wuhan University Journal of Natural Sciences **6** (2001) 72–82 Also appears as Yale Technical Report TR-1207, January 2001, available at URL <ftp://ftp.cs.yale.edu/pub/TR/tr1207.ps>.

4. Bernardinello, L., Cindio, F.D.: A survey of basic net models and modular net classes. In Rozenberg, G., ed.: *Advances in Petri Nets: The DEMON Project*. Volume 609 of *Lecture Notes in Computer Science*. Springer-Verlag (1992) 304–351
5. Esparza, J., Nielsen, M.: Decidability issues for Petri nets - a survey. *Journal of Informatik Processing and Cybernetics* **30** (1994) 143–160
6. Esparza, J.: Decidability and complexity of Petri net problems-an introduction. In Rozenberg, G., Reisig, W., eds.: *Lectures on Petri Nets I: Basic models*. Springer Verlag (1998) 374–428 Published as LNCS 1491.
7. Mayr, E.W.: An algorithm for the general Petri net reachability problem. *SIAM J. Comput.* **13** (1984) 441–460
8. Ginsburg, S., Spanier, E.H.: Semigroups, Presburger formulas, and languages. *Pacific Journal of Mathematics* **16** (1966) 285–296
9. Hopcroft, J., Pansiot, J.: On the reachability problem for 5-dimensional vector addition systems. *Theoretical Computer Science* **8** (1978) 135–159
10. Berry, G., Boudol, G.: The Chemical Abstract Machine. *Theoretical Computer Science* **96** (1992) 217–248
11. Ibarra, O.H., Dang, Z., Egecioglu, O.: Catalytic p systems, semilinear sets, and vector addition systems. *Theor. Comput. Sci.* **312** (2004) 379–399
12. Brand, D., Zafropulo, P.: On communicating finite-state machines. *Journal of the ACM* **30** (1983) 323–342
13. Milner, R.: *Biographical reactive systems: basic theory*. Technical report, University of Cambridge (2001) UCAM-CL-TR-523.
14. Intanagonwiwat, C., Govindan, R., Estrin, D.: Directed diffusion: a scalable and robust communication paradigm for sensor networks. In: *Proceedings of the 6th Annual International Conference on Mobile computing and networking*, ACM Press (2000) 56–67
15. Madden, S.R., Franklin, M.J., Hellerstein, J.M., Hong, W.: TAG: a Tiny AGgregation service for ad-hoc sensor networks. In *OSDI 2002: Fifth Symposium on Operating Systems Design and Implementation* (December, 2002)
16. Fang, Q., Zhao, F., Guibas, L.: Lightweight sensing and communication protocols for target enumeration and aggregation. In: *Proceedings of the 4th ACM International Symposium on Mobile ad hoc networking & computing*, ACM Press (2003) 165–176
17. Zhao, F., Liu, J., Liu, J., Guibas, L., Reich, J.: Collaborative signal and information processing: An information directed approach. *Proceedings of the IEEE* **91** (2003) 1199–1209
18. Grossglauser, M., Tse, D.N.C.: Mobility increases the capacity of ad hoc wireless networks. *IEEE/ACM Transactions on Networking* **10** (2002) 477–486
19. Kahn, J.M., Katz, R.H., Pister, K.S.J.: Next century challenges: mobile networking for “Smart Dust”. In: *MobiCom ’99: Proceedings of the 5th annual ACM/IEEE international conference on Mobile computing and networking*, ACM Press (1999) 271–278
20. Chatzigiannakis, I., Nikolettseas, S., Spirakis, P.: Smart dust protocols for local detection and propagation. In: *POMC ’02: Proceedings of the second ACM international workshop on Principles of mobile computing*, ACM Press (2002) 9–16
21. Itkis, G., Levin, L.A.: Fast and lean self-stabilizing asynchronous protocols. In: *Proceeding of 35th Annual Symposium on Foundations of Computer Science*, IEEE Press (1994) 226–239