C-SHORe: A Collapsible Approach
to Verifying Higher-Order Programs

C. Broadbent[a,b]   A. Carayol[c]   M. Hague[a,c,d]   O. Serre[a]

Abstract

Higher-order recursion schemes (HORS) have recently received much attention as a useful abstraction of higher-order functional programs with a number of new verification techniques employing HORS model-checking as their centrepiece. This paper contributes to the ongoing quest for a truly scalable model-checker for HORS by offering a different, automata theoretic perspective. We introduce the first practical model-checking algorithm that acts on a generalisation of pushdown automata equi-expressive with HORS called *collapsible pushdown systems* (CPDS). At its core is a substantial modification of a recently studied saturation algorithm for CPDS. In particular it is able to use information gathered from an approximate forward reachability analysis to guide its backward search. Moreover, we introduce an algorithm that prunes the CPDS prior to model-checking and a method for extracting counter-examples in negative instances. We compare our tool with the state-of-the-art verification tools for HORS and obtain encouraging results. In contrast to some of the main competition tackling the same problem, our algorithm is fixed-parameter tractable, and we also offer significantly improved performance over the only previously published tool of which we are aware that also enjoys this property. The tool and additional material are available from `http://cshore.cs.rhul.ac.uk`.

1. Introduction

Functional languages such as Haskell, OCaML and Scala strongly encourage the use of higher-order functions. This represents a challenge for software verification, which usually does not model recursion accurately, or models only first-order calls (e.g. SLAM [2] and Moped [33]). However, there has recently been much interest in a model called *higher-order recursion schemes (HORS)* (see e.g. [29]), which offers a way of abstracting functional programs in a manner that precisely models higher-order control-flow.

The execution trees of HORS enjoy decidable μ-calculus theories [29], which testifies to the good algorithmic properties of the model. Even 'reachability' properties (subsumed by the μ-calculus) are very useful in practice. As a simple example, the safety of

[a] LIAFA (Université Paris Diderot – Paris 7 & CNRS).

[b] University of Tokyo.

[c] LIGM (Université Paris-Est & CNRS).

[d] Royal Holloway University of London

incomplete pattern matching clauses could be checked by asking whether the program can 'reach a state' where a pattern match failure occurs. More complex 'reachability' properties can be expressed using a finite automaton and could, for example, specify that the program respects a certain discipline when accessing a particular resource (see [22]). Despite even reachability being $(n-1)$-EXPTIME complete, recent research has revealed that useful properties of HORS can be checked in practice.

Kobayashi's TRecS [23] tool, which checks properties expressible by a deterministic trivial Büchi automaton (all states accepting), was the first to achieve this. It works by determining whether a HORS is typable in an intersection-type system characterising the property to be checked [22]. In a bid to improve scalability, a number of other algorithms have subsequently been designed and implemented such as Kobayashi *et al.*'s GTRecS(2) [25, 26] and Neatherway *et al.*'s TravMC [28] tools (appearing in ICFP 2012), all of which remain based on intersection type inference.

This work is the basis of various techniques for verifying functional programs. Kobayashi *et al.* have developed MoCHi [27] that checks safety properties of (OCaML) programs, and EHMTT Verifier [38] for tree processing programs. Both use a model-checker for recursion schemes as a central component. Similarly, Ramsay and Ong [30] provide a verification procedure for programs with pattern matching employing recursion schemes as an abstraction.

Despite much progress, even the state-of-the-art TRecS does not scale to recursion schemes big enough to model realistically sized programs; achieving scalability while accurately tracking higher-order control-flow is a challenging problem. This paper offers an automata-theoretic perspective on this challenge, providing a fresh set of tools that contrast with previous intersection-type approaches. Techniques based on pushdown automata have previously visited ICFP, such as the approximate higher-order control-flow analysis CFA2 [40], but our aims are a bit different in that we wish to match the expressivity of HORS. Consequently we require a more sophisticated notion of pushdown automaton.

Collapsible pushdown systems (CPDS) [16] are an alternative representation of the class of execution trees that can be generated by recursion schemes (with linear-time mutual-translations between the two formalisms [10, 16]). While pushdown systems augment a finite-state machine with a stack and provide an ideal model for first-order programs [20], collapsible pushdown systems model higher-order programs by extending the stack of a pushdown system to a nested "stack-of-stacks" structure. The nested stack structure enables one to represent closures. Indeed the reader might find it helpful to view a CPDS as being Krivine's Abstract Machine in a guise making it amenable to the generalisation of techniques for pushdown model-checking. Salvati and Walukiewicz have studied in detail the connection with the Krivine abstract machine [32].

For ordinary ('order-1') pushdown systems, a model-checking approach called *saturation* has been successfully implemented by tools such as Moped [33] and PDSolver [15]. Given a regular set

of configurations of the pushdown system (represented by a finite automaton A acting on stacks), saturation can solve the 'backward reachability problem' by computing another finite automaton recognising a set of configurations from which a configuration in $\mathcal{L}(A)$ can be reached. This is a fixed-point computation that gradually adds transitions to A until it is 'saturated'. If A recognises a set of error configurations, one can determine whether the pushdown system is 'safe' by checking if its initial configuration is recognised by the automaton computed by saturation.

We recently extended the saturation method to a backward reachability analysis of collapsible pushdown systems [7]. This runs in PTIME when the number of control states is bounded. Crucially, this condition is satisfied when translating from recursion schemes of bounded arity with properties represented by automata of bounded size [16]. Whilst the HORS/intersection-type based tool GTRecS(2) also enjoys this fixed-parameter tractability (others do not), it times out on many benchmarks that our tool solves quickly.

Motivated by these facts, we revisit the foundations of higher-order verification tools and introduce C-SHORe — the first model-checking tool for the (direct) analysis of collapsible pushdown systems. Whilst the tool is based on the ICALP 2012 result, some substantial modifications and additions are made to the algorithm, leading to several novel practical and theoretical contributions:

1. An approximate *forward* reachability algorithm providing data

    (a) …allowing for the CPDS to be pruned so that saturation receives a smaller input.

    (b) …employed by a modified saturation algorithm to guide its *backward* search.

    This is essential for termination on most of our benchmarks.

2. A method for extracting witnesses to reachability.

3. A complete reworking of the saturation algorithm that speeds up the fixed-point computation.

4. Experimental results showing our approach compares well with TRecS, GTRecS(2) and TravMC.

It is worth remarking that the other type-based tools mentioned above all work by propagating information in a forward direction with respect to the evaluation of the model. In contrast, the raw saturation algorithm works backwards, but we also show here how forward and backward propagation can be combined.

In Sections 5 to 8 we describe the original contributions of this paper. These sections can be understood independently of one another, and hence the reader does not need to fully grasp each section before continuing to the next. The remaining sections describe the background, related work and conclusions.

In Section 3 we recall the basic structures used in the paper as well as recapping the ICALP 2012 algorithm. In Section 5 we describe the approximate forwards-reachability analysis and how it is exploited. In Section 6 we show how to generate witnesses to reachability. In Section 7, we then consider how to restructure the saturation algorithm to more efficiently compute the fixed-point. We provide experimental results in Section 8. Note that we do not discuss in formal detail the translation from HORS model-checking to reachability for CPDS, which essentially follows [10]. However, we do give an informal overview in Section 2, which we hope serves to demonstrate how closures can be accurately modelled.

The tool is available at http://cshore.cs.rhul.ac.uk.

2. Modelling Higher-Order Programs

In this section we give an informal introduction to the process of modelling higher-order programs for verification. In particular, we show how a simple example program can be modelled using a higher-order recursion scheme, and then we show how this scheme is evaluated using a collapsible pushdown system. For a more systematic approach to modelling higher-order programs with recursion schemes, we refer the reader to work by Kobayashi [27]. This section is for background only, and can be safely skipped.

For this section, consider the toy example below.

```
main = Play toy
Play x = if * (playWith x) else (Upgrade Play x)
Upgrade f y = if * (f null) else (f (up y))
```

In this example, * represents a non-deterministic choice (that may, for example, be a result of some input by the user). Execution begins at the main function, where we begin playing with toy (where toy is a suitable constant value). To play with a toy we have a choice of either immediately playing, or attempting to upgrade. The upgrade function may be malicious and force us to play with a null toy (causing an error), or may be benevolent and allow us to play with the upgraded toy.

2.1 Higher-Order Recursion Schemes

As a first step in modelling this program, we introduce, informally, higher-order recursion schemes. These are rewrite systems that generate the computation tree of a functional program. A rewrite rule takes the form

$$N \, \phi \, x \, \hookrightarrow \, t$$

where N is a typed non-terminal with (possibly higher-order) arguments ϕ and x. A term $N \, t_\phi \, t_x$ rewrites to t with t_ϕ substituted for ϕ and t_x substituted for x. Note that recursion schemes require t to be of ground type. We will illustrate the behaviour of a recursion scheme and its use in analysis using the toy example from above.

We can directly model our example with the scheme

$$
\begin{aligned}
main &\hookrightarrow P \, toy \\
P \, x &\hookrightarrow or \, (playWith \, x) \, (U \, P \, x) \\
U \, \phi \, y &\hookrightarrow or \, (\phi \, null) \, (\phi \, (up \, y))
\end{aligned}
$$

where P is the non-terminal associated with the $Play$ function, and U is the non-terminal associated with the $Upgrade$ function; toy, or, $playWith$ and up are terminal symbols of arity 0, 2, 1 and 1 respectively (e.g. in the second rule, or takes the two arguments $(playWith \ x)$ and $(U \ P \ x)$). The scheme above begins with the non-terminal $main$ and, through a sequence of rewrite steps, generates a tree representation of the evolution of the program. Figure 1 shows such a sequence and is and described below.

Beginning with the non-terminal $main$, we apply the first rewrite rule to obtain the tree representing the term $P \ toy$. We then apply the second rewrite rule, instantiating x with toy to obtain the next tree in the sequence. This continues *ad infinitum* to produce a possibly infinite tree labelled only by terminals.

We are interested in ensuring the correctness of the toy program. In our case, this means ensuring that the program never attempts to $playWith$ a $null$ toy. By inspecting the rightmost tree in Figure 1, we can identify a branch labelled $or, or, or, playWith, null$. This is an error situation because $playWith$ is being called with a $null$ argument. In general we can define the regular language $\mathcal{L}_{err} = or^* playWith \ or^* null$. If the tree generated by the recursion scheme contains a branch labelled by a word appearing in $\mathcal{L}_{err}$, then we have identified an error in the program.

2.2 Collapsible Pushdown Automata

Previous research into the verification of recursion schemes has used an approach based on *intersection types* (e.g. [24, 28]). In this work we investigate a radically different approach exploiting the connection between higher-order recursion schemes and an automata model called *collapsible pushdown automata* (CPDA). These two formalisms are, in fact, equivalent.

$$main \quad \hookrightarrow \quad P \quad \hookrightarrow \quad or \qquad \hookrightarrow \cdots \hookrightarrow \qquad or \qquad \hookrightarrow \cdots$$

main → P → or (playWith toy | toy, U P toy) → ⋯ → or (playWith (toy, or (playWith | null, U P toy)), or (U P toy)) → ⋯

Figure 1: The behaviour of a toy recursion scheme.

Theorem 2.1 (Equi-expressivity [16])**.** *For each order-n recursion scheme, there is an order-n collapsible pushdown automaton generating the same tree, and vice-versa. Furthermore, the translations in both directions are linear.*

We describe at a high level the structure of a CPDA and how they can be used to evaluate recursion schemes. In our case, this means outputting a sequence of non-terminals representing each path in the tree. More formal definitions are given in Section 3. At any moment, a CPDA is in a *configuration* $\langle p, w \rangle$, where p is a control state taken from a finite set $\mathcal{P}$, and w is a higher-order collapsible stack. In the following we will focus on the stack. Control states are only needed to ensure that sequences of stack operations occur in the correct order and are thus elided for clarity.

In the case of our toy example, we have an order-2 recursion scheme and hence an order-2 stack. An order-1 stack is a stack of characters a from a finite alphabet Σ. An order-2 stack is a stack of order-1 stacks. Thus we can write $[[main]]$ to denote the order-2 stack containing only the order-1 stack $[main]$; $[main]$ is an order-1 stack containing only the character $main$. In general Σ will contain all subterms appearing in the original statement of our toy example recursion scheme. The evolution of the CPDA stack is given in Figure 2 and explained below.

The first step is to rewrite $main$ using $main \hookrightarrow P\ toy$. Since $(P\ toy)$ is a subterm of our recursion scheme, we have $(P\ toy) \in \Sigma$ and we simply rewrite the stack $[[main]]$ to $[[P\ toy]]$.

The next step is to call the function P. As is typical in the execution of programs, a function call necessitates a new stack frame. In particular, this means pushing the body of P (that is $or\ (playWithx)\ (U\ P\ x)$) onto the stack, resulting in the third stack in Figure 2. Note that we do not instantiate the variable x, hence we use only the subterms appearing in the recursion scheme.

Recall that we want to obtain a CPDA that outputs a sequence of terminals representing each path in the tree. To evaluate the term $or\ (\cdots)\ (\cdots)$ we have to output the terminal or and then (non-deterministically) choose a branch of the tree to follow. Let us choose $(U\ P\ x)$. Hence, the CPDA outputs the terminal or and rewrites the top term to $(U\ P\ x)$. Next we make a call to the U function, pushing its body on to the stack, and then pick out the $(\phi\ null)$ branch of the or terminal. This takes us to the beginning of the second row of Figure 2.

To proceed, we have to evaluate $(\phi\ null)$. To be able to do this, we have to know the value of ϕ. We can obtain this information by inspecting the stack and seeing that the first argument of the call of U is P. However, since we can only see the top of a stack, we would have to remove the character $(\phi\ null)$ to be able to determine that $\phi = P$, thus losing our place in the computation.

This is where we use the power of order-2 stacks. An order-2 stack is able — via a $push_2$ operation — to create a copy of its topmost order-1 stack. Hence, we perform this copy (note that the top of the stack is written on the left) and delve into the copy of the stack to ascertain the value of ϕ. While doing this we also create a *collapse link*, pictured as an arrow from P to the term $(\phi\ null)$. This collapse link is a pointer from P to the context in which P will be evaluated. In particular, if we need to know the value

of x in the body of P, we will need to know that P was called with the $null$ argument, within the term $(\phi\ null)$; the collapse link provides a pointer to this information (in other words we have encoded a closure in the stack). We can access this information via a *collapse* operation. These are the two main features of a higher-order collapsible stack, described formally in the next section.

To continue the execution, we push the body of P on to the stack, output the or symbol and choose the $(playWith\ x)$ branch. Since $playWith$ is a terminal, we output it and pick out x for evaluation. To know the value of x, we have to look into the stack and follow the collapse link from P to $(\phi\ null)$. Note that we do not need to create a copy of the stack here because x is an order-0 variable and thus represents a self-contained execution. Since $null$ is the value of the argument we are considering, we pick it out and then output it before terminating. This completes the execution corresponding to the error branch that we identified in Figure 1.

2.3 Collapsible Pushdown Systems

During the execution above, we see that the CPDA has output the sequence $or, or, or, playWith, null$. According to our definition of $\mathcal{L}_{err}$, this is an error sequence and should be flagged as such. Our general approach is to take the finite automaton A representing the regular language $\mathcal{L}_{err}$ and form a product of this automaton and the CPDA described above. This results in a CPDA that does not output any symbols, but instead keeps track in its control state of the progression of A. Thus we are interested not in the language that such a CPDA outputs, but whether it is able to reach an accepting state of A. We call a CPDA without output symbols a *collapsible pushdown system* (CPDS), and the question of whether a CPDS can reach a given state is the reachability problem. This is the subject of the remainder of the paper.

3. Preliminaries

3.1 Collapsible Pushdown Systems

We first introduce higher-order collapsible stacks and their operations, before giving the definition of collapsible pushdown systems.

3.1.1 Higher-Order Collapsible Stacks and Their Operations

Higher-order collapsible stacks are built from a stack alphabet Σ and form a nested "stack-of-stacks" structure. Using an idea from *panic automata* [21], each stack character contains a pointer — called a "link" — to a position lower down in the stack. Operations updating stacks (defined below) may create copies of sub-stacks. The link is intuitively a pointer to the context in which the stack character was first created. In the sequel, we fix the maximal order to n, and use k to range between 1 and n. In the definition below, we defer the meaning of *collapse link* to Definition 3.2.

Definition 3.1 (Order-n Collapsible Stacks)**.** *Given a finite set of stack characters Σ, an order-0 stack is simply a character $a \in \Sigma$. An order-n stack is a sequence $w = [w_\ell \ldots w_1]_n$ such that each w_i is an order-$(n-1)$ stack and each character a on the stack is augmented with a collapse link. The top-most stack is w_ℓ. Let $Stacks_n$ denote the set of order-n stacks.*
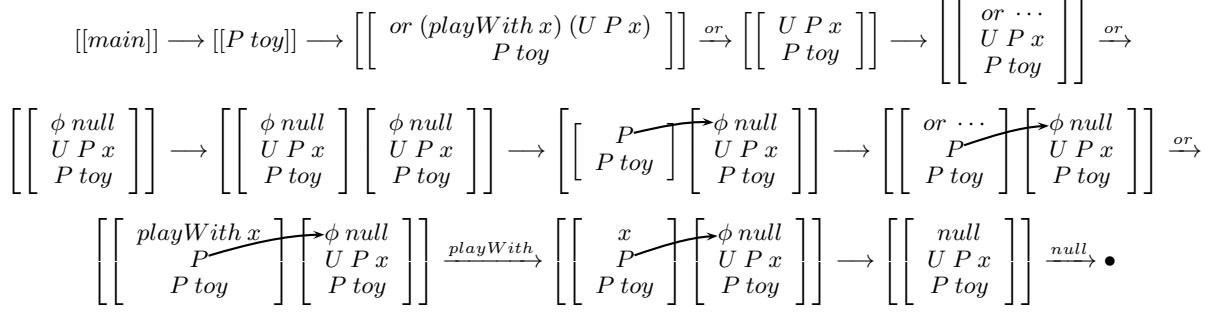
$$[[main]] \longrightarrow [[P\ toy]] \longrightarrow \left[\left[\begin{array}{c} or\ (playWith\ x)\ (U\ P\ x) \\ P\ toy \end{array}\right]\right] \xrightarrow{or} \left[\left[\begin{array}{c} U\ P\ x \\ P\ toy \end{array}\right]\right] \longrightarrow \left[\left[\begin{array}{c} or\ \cdots \\ U\ P\ x \\ P\ toy \end{array}\right]\right] \xrightarrow{or}$$

$$\left[\left[\begin{array}{c} \phi\ null \\ U\ P\ x \\ P\ toy \end{array}\right]\right] \longrightarrow \left[\left[\begin{array}{c} \phi\ null \\ U\ P\ x \\ P\ toy \end{array}\right]\left[\begin{array}{c} \phi\ null \\ U\ P\ x \\ P\ toy \end{array}\right]\right] \longrightarrow \left[\left[\begin{array}{c} P \\ P\ toy \end{array}\right.\!\!\!\left.\begin{array}{c} \phi\ null \\ U\ P\ x \\ P\ toy \end{array}\right]\right] \longrightarrow \left[\left[\begin{array}{c} or\ \cdots \\ P \\ P\ toy \end{array}\right.\!\!\!\left.\begin{array}{c} \phi\ null \\ U\ P\ x \\ P\ toy \end{array}\right]\right] \xrightarrow{or}$$

$$\left[\left[\begin{array}{c} playWith\ x \\ P \\ P\ toy \end{array}\right.\!\!\!\left.\begin{array}{c} \phi\ null \\ U\ P\ x \\ P\ toy \end{array}\right]\right] \xrightarrow{playWith} \left[\left[\begin{array}{c} x \\ P \\ P\ toy \end{array}\right.\!\!\!\left.\begin{array}{c} \phi\ null \\ U\ P\ x \\ P\ toy \end{array}\right]\right] \longrightarrow \left[\left[\begin{array}{c} null \\ U\ P\ x \\ P\ toy \end{array}\right]\right] \xrightarrow{null} \bullet$$

Figure 2: A stack evaluating the toy example.

$$\bullet \rightarrow \bullet \xrightarrow{[_2} \bullet \xrightarrow{[_1} \bullet \xrightarrow{a} \bullet \xrightarrow{b} \bullet \xrightarrow{]_1} \bullet \xrightarrow{]_2} \bullet \xrightarrow{[_2} \bullet \xrightarrow{[_1} \bullet \xrightarrow{c} \bullet \xrightarrow{]_1} \bullet \xrightarrow{[_1} \bullet \xrightarrow{d} \bullet \xrightarrow{e} \bullet \xrightarrow{]_1} \bullet \xrightarrow{]_2} \bullet$$

Figure 3: A graph representation of a stack.

Collapse links point to positions in the stack. Before describing them formally, we give an informal description and some basic definitions. An order-n stack can be represented naturally as an edge-labelled word-graph over the alphabet $\{[_{n-1}, \ldots, [_1,]_1, \ldots,]_{n-1}\} \uplus \Sigma$, with additional collapse-links pointing from a stack character in Σ to the beginning of the graph representing the target of the link. For technical convenience we do not use $[_n$ or $]_n$ symbols (these appear uniquely at the beginning and end of the stack). An example order-3 stack is given in Figure 3, with only a few collapse links shown, ranging from order-3 to order-1 respectively.

Stacks are written with the top on the left. Given an order-n stack $[w_\ell \ldots w_1]_n$, we define

$$\begin{array}{rcll} top_{n+1}([w_\ell \ldots w_1]_n) &=& [w_\ell \ldots w_1]_n & \\ top_n([w_\ell \ldots w_1]_n) &=& w_\ell & \text{if } \ell > 0 \\ top_n([\ ]_n) &=& [\ ]_{n-1} & \text{otherwise} \\ top_k([w_\ell \ldots w_1]_n) &=& top_k(w_\ell) & \text{if } k < n \text{ and } \ell > 0 \end{array}$$

noting that $top_k(w)$ is undefined if $top_{k'}(w)$ is empty for any $k' > k$. We also remove the top portion of a top_k stack using

$$bot_n^i([w_\ell \ldots w_1]_n) = [w_i \ldots w_1]_n$$

when $i \leq \ell$ and $\ell > 0$, and

$$bot_k^i([w_\ell \ldots w_1]_n) = [bot_k^i(w_\ell)w_{\ell-1} \ldots w_1]_n$$

when $k < n$ and $\ell > 0$. We are now ready to define collapse links.

Definition 3.2 (Collapse Links). *An order-k collapse link is a pair (k, i) where $1 \leq k \leq n$ and $i > 0$.*

For $top_1(w) = a$ where a has the link (k, i), the destination of the link is $bot_k^i(w)$. We disallow collapse links where bot_k^i does not lead to a valid stack. The example stack in Figure 3 is thus $[[[a^{(3,1)}b]_1]_2[[c^{(2,1)}]_1]_2[[d^{(1,1)}e]_1]_2]_3$, where collapse links are denoted as superscripts. Often (as we have done for the stack characters b and e) we will omit these superscripts for readability.

Finally the following notation appends a stack on top of another. Given an order-k stack $v = [v_\ell \ldots v_1]_k$ and an order-$(k'-1)$ stack u (with $k' \leq k$), we define $u :_{k'} v = [u\ v_\ell \ldots v_1]_k$ if $k' = k$ and $u :_{k'} v = [(u :_{k'} v_\ell) \ldots v_1]_k$ if $k' < k$.

The following operations apply to an order-n collapsible stack.

$$\begin{array}{rcl} \mathcal{O}_n &=& \{pop_1, \ldots, pop_n\} \cup \{push_2, \ldots, push_n\}\ \cup \\ && \{collapse_2, \ldots, collapse_n\}\ \cup \\ && \{\,push_a^2, \ldots, push_a^n, rew_a \mid a \in \Sigma\,\} \end{array}$$

We define each stack operation for an order-n stack w. Collapse links are created by $push_a^k$, which add a character to the top of a

given stack w with a link pointing to $top_{k+1}(pop_k(w))$. This gives a access to the context in which it was created. We set

1. $pop_k(u :_k v) = v$,

2. $push_k(u :_k v) = u :_k (u :_k v)$,

3. $collapse_k(w) = bot_k^i(w)$ where $top_1(w) = a^{(k,i)}$ for some i,

4. $push_a^k(w) = b^{(k,\ell-1)} :_1 w$ where $top_{k+1}(w) = [w_\ell \ldots w_1]_k$,

5. $rew_b\left(a^{(k,i)} :_1 v\right) = b^{(k,i)} :_1 v$.

Note that, for a $push_k$ operation, links outside of $u = top_k(w)$ point to the same destination in both copies of u, while links pointing within u point within the respective copies of u. Since $collapse_1$ is always equivalent to pop_1, we neither create nor follow order-1 links. For a more detailed introduction see [10].

3.1.2 Collapsible Pushdown Systems

We are now ready to define collapsible pushdown systems.

Definition 3.3 (Collapsible Pushdown Systems). *An order-n collapsible pushdown system (n-CPDS) is a tuple $\mathcal{C} = (\mathcal{P}, \Sigma, \mathcal{R})$ where $\mathcal{P}$ is a finite set of control states, Σ is a finite stack alphabet, and $\mathcal{R} \subseteq \mathcal{P} \times \Sigma \times \mathcal{O}_n \times \mathcal{P}$ is a set of rules.*

A *configuration* of a CPDS is a pair $\langle p, w \rangle$ where $p \in \mathcal{P}$ and $w \in Stacks_n$. We denote by $\langle p, w \rangle \longrightarrow \langle p', w' \rangle$ a transition from a rule (p, a, o, p') with $top_1(w) = a$ and $w' = o(w)$. A run of a CPDS is a finite sequence $\langle p_0, w_0 \rangle \longrightarrow \cdots \longrightarrow \langle p_\ell, w_\ell \rangle$.

3.2 Representing Sets of Stacks

Our algorithm represents sets of configurations using order-n stack automata. These are a kind of alternating automata with a nested structure that mimics the nesting in a higher-order collapsible stack. We recall the definition below.

Definition 3.4 (Order-n Stack Automata). *An order-n stack automaton $A = (\mathcal{Q}_n, \ldots, \mathcal{Q}_1, \Sigma, \Delta_n, \ldots, \Delta_1, \mathcal{F}_n, \ldots, \mathcal{F}_1)$ is a tuple where Σ is a finite stack alphabet, and*

1. *for all $n \geq k \geq 2$, we have $\mathcal{Q}_k$ is a finite set of states, $\mathcal{F}_k \subseteq \mathcal{Q}_k$ is a set of accepting states, and $\Delta_k \subseteq \mathcal{Q}_k \times \mathcal{Q}_{k-1} \times 2^{\mathcal{Q}_k}$ is a transition relation such that for all q and Q there is at most one q' with $(q, q', Q) \in \Delta_k$, and*

2. *$\mathcal{Q}_1$ is a finite set of states, $\mathcal{F}_1 \subseteq \mathcal{Q}_1$ a set of accepting states, and $\Delta_1 \subseteq \bigcup_{2 \leq k \leq n} \left(\mathcal{Q}_1 \times \Sigma \times 2^{\mathcal{Q}_k} \times 2^{\mathcal{Q}_1}\right)$ a transition relation.*

The sets $\mathcal{Q}_k$ are disjoint and their states recognise order-k stacks. Stacks are read from "top to bottom". A transition $(q, q', Q) \in \Delta_k$, written $q \xrightarrow{q'} Q$, from q to Q for some $k > 1$ requires that the top_{k-1} stack is accepted from $q' \in \mathcal{Q}_{(k-1)}$ and the rest of the stack is accepted from each state in Q. At order-1, a transition (q, a, Q_{col}, Q) has the additional requirement that the stack linked to by a is accepted from Q_{col}. A stack is accepted if a subset of $\mathcal{F}_k$ is reached at the end of each order-k stack. We write $w \in \mathcal{L}_q(A)$ to denote the set of all w accepted from q. Note that a transition to the empty set is distinct from having no transition. Figure 4 shows part of a run over the stack in Figure 3 where each node in the graph is labelled by the states from which the remainder of the stack containing it (as well as the stacks linked to) must be accepted. Note, e.g., that since Q_2 appears at the bottom of an order-2 stack, we must have $Q_2 \subseteq \mathcal{F}_2$ for the run to be accepting. The transitions used are $q_3 \xrightarrow{q_2} Q_3 \in \Delta_3$, $q_2 \xrightarrow{q_1} Q_2 \in \Delta_2$, and $q_1 \xrightarrow[Q_{col}]{a} Q_1 \in \Delta_1$. See Section 4 for further examples.

$$q_3 \xrightarrow{[_2} q_2 \xrightarrow{[_1} q_1 \xrightarrow{a} Q_1 \xrightarrow{b} \cdots \longrightarrow Q_2 \xrightarrow{]_2} Q_3 \cup Q_{col} \xrightarrow{[_2} \cdots$$

Figure 4: Part of a run over an example stack.

3.2.1 Representing Transitions and States

We use a *long-form* notation (defined below) that captures nested sequences of transitions. For example, we may write $q_3 \xrightarrow[Q_{col}]{a} (Q_1, Q_2, Q_3)$ to capture the transitions shown in Figure 4. Together, these indicate that after starting from the beginning of the stack and reading only the topmost stack character, the remainder of the stack must be accepted by Q_{col}, Q_1, Q_2, and Q_3. More generally, we may also use $q_3 \xrightarrow{q_1} (Q_2, Q_3)$, and $q_3 \xrightarrow{q_2} (Q_3)$.

Formally, when $q \in \mathcal{Q}_k$, $q' \in \mathcal{Q}_{k'}$, $Q_i \subseteq \mathcal{Q}_i$ for all $k \geq i \geq 1$, and there is some i with $Q_{col} \subseteq \mathcal{Q}_i$, we write

$$q \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k) \text{ and } q \xrightarrow{q'} (Q_{k'+1}, \ldots, Q_k).$$

In the first case, there exist $q_{k-1}, \ldots, q_1$ such that $q \xrightarrow{q_{k-1}} Q_k \in \Delta_k$, $q_{k-1} \xrightarrow{q_{k-2}} Q_{k-1} \in \Delta_{k-1}, \ldots, q_1 \xrightarrow[Q_{col}]{a} Q_1 \in \Delta_1$. In the second case there exist $q_{k-1}, \ldots, q_{k'+1}$ with $q \xrightarrow{q_{k-1}} Q_k \in \Delta_k$, $q_{k-1} \xrightarrow{q_{k-2}} Q_{k-1} \in \Delta_{k-1}, \ldots, q_{k'+2} \xrightarrow{q_{k'+1}} Q_{k'+2} \in \Delta_{k'+2}$ and $q_{k'+1} \xrightarrow{q'} Q_{k'+1} \in \Delta_{k'+1}$.

Remark 3.1. *We may also write $q_{Q_k, \ldots, Q_{k'+1}}$ for the q' above (which is uniquely determined by $Q_k, \ldots, Q_{k'+1}$).*

Note that our definitions mean that we have, e.g., $q \xrightarrow[Q_{col}]{a} (Q_1, Q_2, Q_3)$ if and only if we have $q_{Q_3, Q_2} \xrightarrow[Q_{col}]{a} Q_1$ in Δ_1.

3.2.2 Representing Sets of Transitions

Let Δ_k^S denote the set of all order-k *long-form* transitions $q \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k)$ of order-k.

Given a set $T = \{t_1, \ldots, t_\ell\} \subseteq \Delta_k^S$, we say T is of the form

$$Q \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k)$$

whenever $Q = \{q_1, \ldots, q_\ell\}$ and for all $1 \leq i \leq \ell$ we have $t_i = q_i \xrightarrow[Q_{col}^i]{a} (Q_1^i, \ldots, Q_k^i)$ and $Q_{col} = \bigcup_{1 \leq i \leq \ell} Q_{col}^i$ and for all k, $Q_k = \bigcup_{1 \leq i \leq \ell} Q_k^i$. Because a link can only be of one order, we insist that $Q_{col} \subseteq \mathcal{Q}_k$ for some k.

3.3 Representing Sets of Configurations

We define a notion of $\mathcal{P}$-multi-automata [4] for representing sets of configurations of collapsible pushdown systems.

Definition 3.5 ($\mathcal{P}$-Multi Stack Automata)**.** *Given an order-n CPDS with control states $\mathcal{P}$, a $\mathcal{P}$-multi stack automaton is an order-n stack automaton $A = (\mathcal{Q}_n, \ldots, \mathcal{Q}_1, \Sigma, \Delta_n, \ldots, \Delta_1, \mathcal{F}_n, \ldots, \mathcal{F}_1)$ such that for each $p \in \mathcal{P}$ there exists a state $q_p \in \mathcal{Q}_n$.*

A state is *initial* if it is of the form $q_p \in Q_n$ for some control state p or if it is a state $q_k \in Q_k$ for $k < n$ such that there exists a transition $q_{k+1} \xrightarrow{q_k} Q_{k+1}$ in Δ_{k+1}. The language of a $\mathcal{P}$-multi stack automaton A is the set $\mathcal{L}(A) = \{ \langle p, w \rangle \mid w \in \mathcal{L}_{q_p}(A) \}$.

3.4 Basic Saturation Algorithm

Our algorithm computes the set $Pre_{\mathcal{C}}^*(A_0)$ of a collapsible pushdown system $\mathcal{C}$ and a $\mathcal{P}$-multi stack automaton A_0. We assume without loss of generality that initial states of A_0 do not have incoming transitions and are not final. To accept empty stacks from initial states, a bottom-of-stack symbol can by used.

We define $Pre_{\mathcal{C}}^*(A_0)$ as the smallest set such that $Pre_{\mathcal{C}}^*(A_0) \supseteq \mathcal{L}(A_0)$, and

$$Pre_{\mathcal{C}}^*(A_0) \supseteq \{ \langle p, w \rangle \mid \exists \langle p, w \rangle \longrightarrow \langle p', w' \rangle \in Pre_{\mathcal{C}}^*(A_0) \} .$$

We begin with A_0 and iterate a saturation function Π — adding new transitions to A_0 — until a 'fixed point' is reached; that is, we cannot find any more transitions to add.

Notation for Adding Transitions During saturation we designate transitions $q_n \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$ to be added to the automaton. Recall this represents $q \xrightarrow{q_{k-1}} Q_k \in \Delta_k$, $q_{k-1} \xrightarrow{q_{k-2}} Q_{k-1} \in \Delta_{k-1}, \ldots, q_1 \xrightarrow[Q_{col}]{a} Q_1 \in \Delta_1$. Hence, we first, for each $n \geq k > 1$, add $q_k \xrightarrow{q_{k-1}} Q_k$ to Δ_k if it does not already exist. Then, we add $q_1 \xrightarrow[Q_{col}]{a} Q_1$ to Δ_1.

Justified Transitions In this paper, we extend the saturation function to add *justifications* to new transitions that indicate the provenance of each new transition. This permits counter example generation. To each $t = q \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$ we will define the justification $J(t)$ to be either 0 (indicating the transition is in A_0), a pair (r, i), a tuple (r, t', i) or a tuple (r, t', T, i) where r is a rule of the CPDS, i is the number of iterations of the saturation function required to introduce the transition, t' is a long-form transition and T is a set of such transitions. This information will be used in Section 6 for generating counter examples. Note that we apply J to the long-form notation. In reality, we associate each justification with the unique order-1 transition $q_1 \xrightarrow[Q_{col}]{a} Q_1$ associated to each t.

The Saturation Function We are now ready to recall the saturation function Π for a given $\mathcal{C} = (\mathcal{P}, \Sigma, \mathcal{R})$. As described above, we apply this function to A_0 until a fixed point is reached. First set $J(t) = 0$ for all transitions of A_0. The intuition behind the saturation rules can be quickly understood via a rewrite rule (p, a, rew_b, p') which leads to the addition of a transition $q_p \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$ whenever there already existed a transition $q_{p'} \xrightarrow[Q_{col}]{b} (Q_1, \ldots, Q_n)$. Because the rewrite can change the control state from p to p' and the top character from a to b, we must have an accepting run from q_p with a on top whenever we had an accepting run from $q_{p'}$ with b on top. We give examples and intuition of the more complex steps in Section 4, which may be read alongside the definition below.

Definition 3.6 (The Saturation Function Π). *Given an order-n stack automaton A_i we define $A_{i+1} = \Pi(A_i)$. The state-sets of A_{i+1} are defined implicitly by the transitions which are those in A_i plus, for each $r = (p, a, o, p') \in \mathcal{R}$,*

1. *when $o = pop_k$, for each $q_{p'} \xrightarrow{q_k} (Q_{k+1}, \ldots, Q_n)$ in A_i, add to A_{i+1}*

$$t = q_p \xrightarrow[\emptyset]{a} (\emptyset, \ldots, \emptyset, \{q_k\}, Q_{k+1}, \ldots, Q_n)$$

   and set $J(t) = (r, i+1)$ whenever t is not already in A_{i+1},

2. *when $o = push_k$, for each $t = q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k, \ldots, Q_n)$ and T of the form $Q_k \xrightarrow[Q'_{col}]{a} (Q'_1, \ldots, Q'_k)$ in A_i, add to A_{i+1} the transition*

$$t' = q_p \xrightarrow[Q_{col} \cup Q'_{col}]{a} \begin{pmatrix} Q_1 \cup Q'_1, \ldots, Q_{k-1} \cup Q'_{k-1}, \\ Q'_k, \\ Q_{k+1}, \ldots, Q_n \end{pmatrix}$$

   and set $J(t') = (r, t, T, i+1)$ if t' is not already in A_{i+1},

3. *when $o = collapse_k$, when $k = n$, add $t = q_p \xrightarrow[\{q_{p'}\}]{a} (\emptyset, \ldots, \emptyset)$ if it does not exist, and when $k < n$, for each transition $q_{p'} \xrightarrow{q_k} (Q_{k+1}, \ldots, Q_n)$ in A_i, add to A_{i+1} the transition $t = q_p \xrightarrow[\{q_k\}]{a} (\emptyset, \ldots, \emptyset, Q_{k+1}, \ldots, Q_n)$ if t does not already exist. In all cases, if t is added, set $J(t) = (r, i+1)$,*

4. *when $o = push_b^k$ for all transitions $t = q_{p'} \xrightarrow[Q_{col}]{b} (Q_1, \ldots, Q_n)$ and $T = Q_1 \xrightarrow[Q'_{col}]{a} Q'_1$ in A_i with $Q_{col} \subseteq \mathcal{Q}_k$, add to A_{i+1} the transition*

$$t' = q_p \xrightarrow[Q'_{col}]{a} (Q'_1, Q_2, \ldots, Q_k \cup Q_{col}, \ldots, Q_n) ,$$

   and set $J(t') = (r, t, T, i+1)$ if t' is not already in A_{i+1},

5. *when $o = rew_b$ for each transition $t = q_{p'} \xrightarrow[Q_{col}]{b} (Q_1, \ldots, Q_n)$ in A_i, add to A_{i+1} the transition $t' = q_p \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$, setting $J(t') = (r, t, i)$ when t' is not already in A_{i+1}.*

From A_0, we iterate $A_{i+1} = \Pi(A_i)$ until $A_{i+1} = A_i$. Generally, we terminate in n-EXPTIME. When all $q \xrightarrow{q'} Q \in \Delta_n$ in A_0 have at most one element in Q, we can restrict Π to only add transitions where Q_n has at most one element, giving $(n-1)$-EXPTIME complexity. In all cases saturation is linear in the size of Σ.

4. Examples of Saturation

We give an example of how the saturation method solves the backwards reachability problem. Consider a CPDS with the run

$$\langle p_1, [b]\,[c]\,[d]\rangle \xrightarrow{push_a^2} \langle p_2, [\overset{\frown}{ab}]\,[c]\,[d]\rangle \xrightarrow{push_2}$$

$$\langle p_3, [\overset{\frown}{ab}]\,[\overset{\frown}{ab}]\,[c]\,[d]\rangle \xrightarrow{collapse_2} \langle p_4, [c]\,[d]\rangle \xrightarrow{pop_2} \langle p_5, [d]\rangle .$$

Figure 5 shows the sequence of saturation steps, beginning with an accepting run of $\langle p_5, [d]\rangle$ and finishing with an accepting run of $\langle p_1, [b]\,[c]\,[d]\rangle$. The individual steps are explained below.

Initial Automaton We begin at the top of Figure 5 with a stack automaton containing the transitions $q_{p_5} \xrightarrow{q_1} \emptyset$ and $q_1 \xrightarrow[\emptyset]{d} \emptyset$, which we write $q_{p_5} \xrightarrow[\emptyset]{d} (\emptyset, \emptyset)$. This gives the pictured run over $\langle p_5, [d]\rangle$.

Rule (p_4, c, pop_2, p_5) When the saturation step considers such a pop rule, it adds $q_{p_4} \xrightarrow[\emptyset]{c} (\emptyset, \{q_{p_5}\})$. We add such a transition because we only require the top order-1 stack (removed by pop_2) to have the top character c (hence $\emptyset$ is the next order-1 label), and after the pop_2 the remaining stack needs to be accepted from q_{p_5} (hence $\{q_{p_5}\}$ is the next order-2 label). This new transition allows us to construct the next run over $\langle p_4, [c]\,[d]\rangle$ in Figure 5.

Rule $(p_3, a, collapse_2, p_4)$ Similarly to the pop rule above, the saturation step adds the transition $q_{p_3} \xrightarrow[\{q_{p_4}\}]{a} (\emptyset, \emptyset)$. The addition of such a transition allows us to construct the pictured run over $\langle p_3, [ab]\,[ab]\,[c]\,[d]\rangle$ (collapse links omitted), recalling that $\emptyset \xrightarrow[\emptyset]{\emptyset} \emptyset$, $\emptyset \xrightarrow[\emptyset]{a} \emptyset$ and $\emptyset \xrightarrow[\emptyset]{b} \emptyset$ transitions are always possible due to the empty initial set. Note that the labelling of $\{q_{p_4}\}$ in the run comes from the collapse link on the topmost a character on the stack.

Rule $(p_2, a, push_2, p_3)$ Consider the run from q_{p_3} in Figure 5. The initial transition of the run accepting the first order-1 stack is $q_{p_3} \xrightarrow[\{q_{p_4}\}]{a} (\emptyset, \emptyset)$. We also have $\emptyset \xrightarrow{\emptyset} \emptyset$ (trivially) accepting the second order-1 stack. Any $push_2$ predecessor of this stack must have a top order-1 stack that could have appeared twice at the top of the stack from q_{p_3}. Thus, the saturation step makes the intersection of the initial order-1 transitions of first two order-1 stacks. This results in the transition $q_{p_2} \xrightarrow[\{q_{p_4}\} \cup \emptyset]{a} (\emptyset \cup \emptyset, \emptyset)$, which is used to form the shown run over $\langle p_2, [ab]\,[c]\,[d]\rangle$ (collapse links omitted).

Rule $(p_1, b, push_a^2, p_2)$ The run from q_{p_2} in Figure 5 begins with $q_{p_2} \xrightarrow[\{q_{p_4}\}]{a} (\emptyset, \emptyset)$ and $\emptyset \xrightarrow[\emptyset]{b} \emptyset$. Note that the $push_a^2$ gives a stack with ab on top. Moreover, the collapse link on a should point to the order-1 stack just below the current top one. Since the transition from q_{p_2} requires that the linked-to stack is accepted from q_{p_4}, we need this requirement in the preceding stack (accepted from q_{p_1} and without the a on top). Thus, we move the target of the collapse link into the order-2 destination of the new transition. That is, the saturation step for $push_a^2$ rules creates $q_{p_1} \xrightarrow[\emptyset]{b} (\emptyset, \emptyset \cup \{q_{p_4}\})$. This can be used to construct an accepting run over $\langle p_1, [b]\,[c]\,[d]\rangle$.

5. Initial Forward Analysis

First recall that Sections 5, 6, and 7 may be understood independently of each other. In this section we are interested only in whether $\mathcal{C}$ can reach a configuration of the form $\langle p_{err}, w\rangle$ where p_{err} is a distinguished *error state* from a designated initial configuration $\langle p_0, [\cdots [a]_1 \cdots]_n\rangle$. This suffices to capture the same safety (reachability) properties of recursion schemes as TRecS. We fix a stack-automaton $\mathcal{E}$ recognising all error configurations (those with the state p_{err}). We write $Post_{\mathcal{C}}^*$ for the set of configurations reachable by $\mathcal{C}$ from the initial configuration. This set is irregular and cannot be represented precisely by a stack automaton [5]. We first summarise our approach, giving details in Sections 5.1 and 5.2.

It is generally completely impractical to compute $Pre_{\mathcal{C}}^*(\mathcal{E})$ in full (most non-trivial examples considered in our experiments would time-out). For our saturation algorithm to be usable in practice, it is therefore essential that the search space is restricted, which we achieve by means of an initial forward analysis of the CPDS. Ideally we would compute only $Pre_{\mathcal{C}}^*(\mathcal{E}) \cap Post_{\mathcal{C}}^*$. Since this is irregular, we instead compute a sufficient approximation T where:

$$Pre_{\mathcal{C}}^*(\mathcal{E}) \cap Post_{\mathcal{C}}^* \subseteq T \subseteq Pre_{\mathcal{C}}^*(\mathcal{E}).$$

The initial configuration will belong to T iff it can reach a configuration recognised by $\mathcal{E}$. Computing such a T is much more feasible.

$$(1) \quad q_{p_5} \xrightarrow{[_1} q_1 \xrightarrow{d} \emptyset \xrightarrow{]_1} \emptyset \qquad\qquad (2) \quad q_{p_4} \xrightarrow{[_1} q_2 \xrightarrow{c} \emptyset \xrightarrow{]_1} \{q_{p_5}\} \xrightarrow{[_1} q_1 \xrightarrow{d} \emptyset \xrightarrow{]_1} \emptyset$$

$$(3) \quad q_{p_3} \xrightarrow{[_1} q_3 \xrightarrow{a} \emptyset \xrightarrow{b} \emptyset \xrightarrow{]_1} \emptyset \xrightarrow{[_1} \emptyset \xrightarrow{a} \emptyset \xrightarrow{b} \emptyset \xrightarrow{]_1} \{q_{p_4}\} \xrightarrow{[_1} q_2 \xrightarrow{c} \emptyset \xrightarrow{]_1} \{q_{p_5}\} \xrightarrow{[_1} q_1 \xrightarrow{d} \emptyset \xrightarrow{]_1} \emptyset$$

$$(4) \quad q_{p_2} \xrightarrow{[_1} q_4 \xrightarrow{a} \emptyset \xrightarrow{b} \emptyset \xrightarrow{]_1} \{q_{p_4}\} \xrightarrow{[_1} q_2 \xrightarrow{c} \emptyset \xrightarrow{]_1} \{q_{p_5}\} \xrightarrow{[_1} q_1 \xrightarrow{d} \emptyset \xrightarrow{]_1} \emptyset$$

$$(5) \quad q_{p_1} \xrightarrow{[_1} q_5 \xrightarrow{b} \emptyset \xrightarrow{]_1} \{q_{p_4}\} \xrightarrow{[_1} q_2 \xrightarrow{c} \emptyset \xrightarrow{]_1} \{q_{p_5}\} \xrightarrow{[_1} q_1 \xrightarrow{d} \emptyset \xrightarrow{]_1} \emptyset$$

Figure 5: A sequence of saturation steps.

We first compute an over-approximation of $Post_\mathcal{C}^*$ that may be a superset of $Pre_\mathcal{C}^*(\boldsymbol{\mathcal{E}})$. For this we use a *summary algorithm* [34] (that happens to be order-1) from which we extract an over-approximation of the set of CPDS rules that may be used on a run to p_{err}. Let $\mathcal{C}'$ be the (smaller) CPDS containing only these rules. That is, we remove all rules that we know cannot appear on a run to p_{err}. We could thus take $T = Pre_{\mathcal{C}'}^*(\boldsymbol{\mathcal{E}})$ (computable by the saturation algorithm for $\mathcal{C}'$) since it satisfies the conditions above.

However, we further improve performance by computing an even smaller T. We extract contextual information from our over-approximation of $Post_\mathcal{C}^*$ about how pops and collapses might be used during a run to p_{err}. Our $\mathcal{C}'$ is then restricted to a model $\mathcal{C}''$ that 'guards' its rules by these contextual constraints. Taking $T = Pre_{\mathcal{C}''}^*(\boldsymbol{\mathcal{E}})$ we have a T smaller than $Pre_{\mathcal{C}'}^*(\boldsymbol{\mathcal{E}})$, but still satisfying our sufficient conditions. In fact, $\mathcal{C}''$ will be a '*guarded CPDS*' (defined in the next subsection). We cannot compute $Pre_{\mathcal{C}''}^*(\boldsymbol{\mathcal{E}})$ precisely for a guarded CPDS, but we can adjust the saturation algorithm to compute T such that $Pre_{\mathcal{C}''}^*(\boldsymbol{\mathcal{E}}) \subseteq T \subseteq Pre_{\mathcal{C}'}^*(\boldsymbol{\mathcal{E}})$. This set will thus also satisfy our sufficient conditions.

5.1 Guarded Destruction

An *order-n* guarded *CPDS (n-GCPDS)* is an n-CPDS where conventional pop_k and $collapse_{k'}$ operations are replaced by *guarded operations* of the form pop_k^S and $collapse_{k'}^S$ where $S \subseteq \Sigma$. These operations may only be fired if the stack resulting from the operation has a member of S on top. That is, for o *collapse* or *pop*:

$$o^S(u) := \begin{cases} o(u) & \text{if } o(u) \text{ defined and } top_1(o(u)) \in S \\ \text{undefined} & \text{otherwise .} \end{cases}$$

Note, we do not guard the other stack operations since these themselves guarantee the symbol on top of the new stack.

Given a GCPDS $\mathcal{C}$, we write $\mathbf{Triv}(\mathcal{C})$ for the ordinary CPDS that is the *trivialisation* of $\mathcal{C}$, obtained by replacing each pop_k^S (resp. $collapse_k^S$) in the rules of $\mathcal{C}$ with pop_k (resp. $collapse_k$).

We modify the saturation algorithm to use 'guarded' saturation steps for pop and collapse rules. Other saturation steps are unchanged. Non-trivial guards reduce the size of the stack-automaton constructed by avoiding certain additions that are only relevant for unreachable (and hence uninteresting) configurations in the pre-image. This thus improves performance.

1. when $o = pop_k^S$, for each $q_{p'} \xrightarrow{q_k} (Q_{k+1}, \ldots, Q_n)$ in A such that there is a transition of the form $q_k \xrightarrow{b} (\_, \ldots, \_)$ in A such that $b \in S$, add $q_p \xrightarrow{a}_{\emptyset} (\emptyset, \ldots, \emptyset, \{q_k\}, Q_{k+1}, \ldots, Q_n)$ to A',

3. when $o = collapse_k^S$, for each $q_{p'} \xrightarrow{q_k} (Q_{k+1}, \ldots, Q_n)$ in A such that there is a transition of the form $q_k \xrightarrow{b} (\_, \ldots, \_)$ in A with $b \in S$, add $q_p \xrightarrow{a}_{\{q_k\}} (\emptyset, \ldots, \emptyset, Q_{k+1}, \ldots, Q_n)$ to A'.

Lemma 5.1. *The revised saturation algorithm applied to $\boldsymbol{\mathcal{E}}$ gives a stack automaton recognising T such that $Pre_\mathcal{C}^*(\boldsymbol{\mathcal{E}}) \subseteq T \subseteq Pre_{\mathbf{Triv}(\mathcal{C})}^*(\boldsymbol{\mathcal{E}})$*

Remark 5.1. *The algorithm may result in a stack-automaton recognising configurations that do not belong to $Pre_\mathcal{C}^*(\boldsymbol{\mathcal{E}})$ (albeit still in $Pre_{\mathbf{Triv}(\mathcal{C})}^*(\boldsymbol{\mathcal{E}})$). This is because a stack-automaton state q_k containing a transition $q_k \xrightarrow{b} (\_, \ldots, \_)$ may also contain another transition $q_k \xrightarrow{b'} (\_, \ldots, \_)$ with $b \neq b'$ (and so it might recognise a stack from which a pop_k, say, guarded by b cannot be performed).*

Remark 5.2. *The above modification to the naive saturation algorithm can also be easily incorporated into the efficient fixed point algorithm described in Section 7.*

5.2 Approximate Reachability Graphs

We now give an overview of the algorithm used to obtain an over-approximation of $Post_\mathcal{C}^*$ and thus compute the GCPDS $\mathcal{C}''$ mentioned previously. We refer the reader to Appendix C for details. To keep things simpler, we will assume that a stack symbol uniquely determines the order of any link that it emits (which is the case for a CPDS obtained from a HORS).

An *approximate reachability graph* for $\mathcal{C}$ is a structure (H, E) describing an over-approximation of the reachable configurations of $\mathcal{C}$. The set of nodes of the graph H consists of *heads* of the CPDS, where a head is a pair $(p, a) \in \mathcal{P} \times \Sigma$ and describes configurations of the form $\langle p, u \rangle$ where $top_1(u) = a$. The set E contains directed edges $((p, a), r, (p', a'))$ labelled by rules of $\mathcal{C}$. Such edges over-approximate the transitions that $\mathcal{C}$ might make using a rule r from a configuration described by (p, a) to one described by (p', a'). For example, suppose that $\mathcal{C}$ is order-2 and has, amongst others, the rules $r_1 := (p_1, b, push_2, p_2)$, $r_2 := (p_2, b, push_c^2, p_3)$ and $r_3 := (p_3, c, pop_1, p_4)$ so that it can perform transitions:

$$\left\langle p_1, \left[\left[\begin{array}{c} b \\ a \end{array} \right]\right]\right\rangle \xrightarrow{r_1} \left\langle p_2, \left[\left[\begin{array}{c} b \\ a \end{array} \right]\left[\begin{array}{c} b \\ a \end{array} \right]\right]\right\rangle$$

$$\xrightarrow{r_2} \left\langle p_3, \left[\left[\begin{array}{c} c \\ b \\ a \end{array} \right]\left[\begin{array}{c} b \\ a \end{array} \right]\right]\right\rangle \xrightarrow{r_3} \left\langle p_4, \left[\left[\begin{array}{c} b \\ a \end{array} \right]\left[\begin{array}{c} b \\ a \end{array} \right]\right]\right\rangle$$

where the first configuration mentioned here is reachable. We should then have edges $((p_1, b), r_1, (p_2, b))$, $((p_2, b), r_2, (p_3, c))$ and $((p_3, c), r_3, (p_4, b))$ in E. Below we refer to the configurations above as C_1, C_2, C_3 and C_4 respectively, with respective stacks s_1, s_2, s_3, s_4.

Such a graph can be computed using an *approximate summary algorithm*, which builds up an object (H, E, B, U) consisting of an approximate reachability graph together with two additional components. B is a map assigning each head h in the graph a set $B(h)$ of *stack descriptors*, which are $(n + 1)$-tuples $(h_n, \ldots, h_1, h_c)$ of

heads. Roughly speaking, h_i describes at which head the new top_i-stack resulting from a pop_i operation (applied to a configuration with head h) may have been created, and h_c does likewise for a *collapse* operation. (The algorithm will use $\perp$ in place of a head to indicate when pop_i or $collapse_i$ is undefined.)

Consider $C_3 = \langle p_3, s_3 \rangle$ from the example above. This has control-state p_3 and top stack symbol c and so is associated with the head (p_3, c). Thus $B((p_3, c))$ should contain the stack-descriptor $((p_1, b), (p_2, b), (p_1, b))$, which describes s_3. The first component is because $top_2(s_3)$ was created by a $push_2$ operation from a configuration with head (p_1, b), and the second component is because the top symbol was created via an order-1 push from (p_2, b). Moreover, the order-2 link from the top of s_3 points to a stack occurring on top of a configuration at the head (p_1, b), giving rise to the final component describing the collapse link.

Tracking this information allows the summary algorithm to process the rule r_3 to obtain a description of C_4 from the description of C_3. Since this rule performs a pop_1, it can look at the second component of the stack descriptor to see the head (p_2, b), telling us that pop_1 results in b being on top of the stack. Since the rule r_3 moves into control-state p_4, this tells us that the new head should be (p_4, b). It also tells us that certain pieces of information in $B((p_2, b))$ are relevant to the description of $top_2(s_4)$ contained in $B((p_4, b))$. This aspect is handled by *summary edges*, which describe how information contained in stack descriptors should be shared between heads.

The component U is a set of *approximate higher-order summary edges*. An *order-k summary edge* from a head h to a head h' is a triple of the form $(h, (h'_n, \ldots, h'_{k+1}), h')$ where each h'_i is a head. Intuitively such a summary edge means that if $(h_n, \ldots, h_{k+1}, h_k, \ldots, h_1, h_c) \in B(h)$, then we should also have $(h'_n, \ldots, h'_{k+1}, h_k, \ldots, h_1, h_c) \in B(h')$. When $n = k = 1$ (so that h_c is also unnecessary since there would be no links) note that $(h, (), h')$ behaves like a summary edge in a standard order-1 summary algorithm [34], which is complete at order-1.

To continue our example, the r_3 rule from C_3 to C_4 means U should contain an order-1 summary edge $((p_2, b), ((p_1, b)), (p_4, b))$. This is because $pop_2(s_3) = pop_2(s_4)$ and so (p_1, b) (the first component of the stack descriptor for s_3) should also be the first component of a stack descriptor for s_4. However, since $top_1(s_4)$ was created at a configuration with head (p_2, b), the second and third components of such a stack descriptor for s_4 should be inherited from a stack descriptor in $B((p_2, b))$.

The construction of the approximate reachability graph is described in algorithms 16, 17, 18 and 19. Note that summary edges are created when processing *pop* and *collapse* rules since it is such rules that cause stacks created earlier in the run to re-emerge on top of the stack. In such cases it is necessary to refer to information provided by stack descriptors for stacks in earlier parts of the run to create stack descriptors for the new configuration.

5.3 Extracting the Guarded CPDA

Let $\mathcal{G} = (H, E)$ be an approximate reachability graph for $\mathcal{C}$. Let $\mathbf{Heads}(\mathcal{E})$ be the set of heads of error configurations, *i.e.* $\mathbf{Heads}(\mathcal{E}) := \{ (p_{err}, a) \mid a \in \Sigma \}$. We do a simple backwards reachability computation on the finite graph $\mathcal{G}$ to compute $\mathbf{BackRules}(\mathcal{G})$, defined to be the smallest set satisfying:

$$\mathbf{BackRules}(\mathcal{G}) = \left\{ e \in E \;\middle|\; \begin{array}{l} e = (h, r, h') \in E \text{ for some} \\ h' \in \mathbf{Heads}(\mathcal{E}) \end{array} \right\}$$
$$\cup \left\{ e \in E \;\middle|\; \begin{array}{l} e = (h, r, h') \in E \text{ for some} \\ (h', \_, \_) \in \mathbf{BackRules}(\mathcal{G}) \end{array} \right\}$$

The CPDS rules occurring in the triples in $\mathbf{BackRules}(\mathcal{G})$ can be used to define a pruned CPDS $\mathcal{C}'$ that reaches an error state if and only if the original also does. However, the approximate reacha-

bility graph provides enough information to construct a guarded CPDS $\mathcal{C}''$ whose guards are non-trivial. It should be clear that the following set $\mathbf{BackRulesG}(\mathcal{G})$ of *guarded* rules can be computed:

$$\left\{ (p, a, o', p') \;\middle|\; \begin{array}{l} (\_, (p, a, o, p'), \_) \in \mathbf{BackRules}(\mathcal{G}) \text{ and} \\ o' = \begin{cases} o^S & \text{if } o \text{ is a pop or a collapse and } S \text{ is} \\ & \left\{ b \in \Sigma \;\middle|\; \begin{array}{l} ((p, a), r, (p', b)) \\ \in E \end{array} \right\} \\ & \text{with } r = (p, a, o, p') \\ o & \text{if } o \text{ is a rewrite or push} \end{cases} \end{array} \right\}$$

These rules define a GCPDS on which C-SHORe finally performs saturation.

Lemma 5.2. *The GCPDS $\mathcal{C}''$ defined using the rules $\mathbf{BackRulesG}(\mathcal{G})$ satisfies:* $Post^*_{\mathcal{C}} \cap Pre^*_{\mathcal{C}}(\mathcal{E}) \subseteq Pre^*_{\mathcal{C}''}(\mathcal{E}) \subseteq Pre^*_{\mathcal{C}}(\mathcal{E})$

Algorithm 1 The Approximate Summary Algorithm

Require: An n-CPDS with rules $\mathcal{R}$ and heads $\mathcal{P} \times \Sigma$ and initial configuration $\langle p_0, [\cdots [a_0]_1 \cdots]_n \rangle$
Ensure: The creation of a structure (H, E, B, U) where (H, E) is an approximate reachability graph and U is a set of approximate higher-order summary edges.

Set $H := \{(p_0, a_0)\}$ and set E, B and U to be empty
Call AddStackDescriptor$((p_0, a_0), (\perp, \ldots, \perp, \perp))$
return Done, (H, E, B, U) will now be as required

Algorithm 2 AddStackDescriptor$(h, (h_n, \ldots, h_1, h_c))$

Require: A head $h \in H$ and a stack descriptor $(h_n, \ldots, h_1, h_c)$
Ensure: $(h_n, \ldots, h_1, h_c) \in B(h)$ and that any further additions to $B(h')$ for each $h' \in H$ necessary to respect summary edges are made.

if $(h_n, \ldots, h_1, h_c) \in B(h)$ **then**
  return Done (Nothing to do)
Add $(h_n, \ldots, h_1, h_c)$ to $B(h)$
Call ProcessHeadWithStackDescriptor$(h, (h_n, \ldots, h_1, h_c))$
for $h' \in H$ such that $(h, (h'_n, \ldots, h'_{k+1}), h') \in U$ **do**
  Call AddStackDescriptor$(h', (h'_n, \ldots, h'_{k+1}, h_k, \ldots, h_1, h_c))$
return Done

---

6. Counter Example Generation

In this section, we describe an algorithm that given a CPDS $\mathcal{C}$ and a stack automaton A_0 such that a configuration $\langle p, w \rangle$ of $\mathcal{C}$ belongs $Pre^*_{\mathcal{C}}(A_0)$, constructs a sequence of rules of $\mathcal{C}$ which when applied from $\langle p, w \rangle$ leads to a configuration in $\mathcal{L}(A_0)$. In practice, we use the algorithm with A_0 accepting the set of all configurations starting with some error state p_{err}. The output is a counter-example showing how the CPDS can reach this error state.

The algorithm itself is a natural one, and can be understood by considering the example in Figure 5 described in Section 4. To construct a trace from $\langle p_1, [b] [c] [d] \rangle$ to $\langle p_5, [d] \rangle$ we first note that, when adding the initial transition of the pictured run from q_{p_1}, the saturation step marked that the transition was added due to the rule $(p_1, b, push^2_a, p_2)$. If we apply this rule to $\langle p_1, [b] [c] [d] \rangle$ we obtain $\langle p_2, [ab] [c] [d] \rangle$ (collapse links omitted) and, from the workings of the saturation step, we know that the pictured run from q_{p_2} is a run of the stack automaton. In this way, we can follow the justifications back to a run of A_0, obtaining a counter example.

The main technical difficulty lies in proving that the reasoning outlined above leads to a terminating algorithm. For example, we need to prove that following the justifications does not result us fol-

Algorithm 3 ProcessHeadWithStackDescriptor($h, (h_n, \ldots, h_1, h_c)$)

Require: A head $h := (p, a) \in H$ and a stack descriptor $(h_n, \ldots, h_1, h_c) \in B(h)$

Ensure: All necessary modifications to the graph are made so that it is consistent with the presence of $(h_n, \ldots, h_1, h_c) \in B(h)$. In particular this is the procedure that processes the CPDS rules from h (with respect to a stack described by h and the stack descriptor).

> **for** o and p' such that $r = (p, a, o, p') \in \mathcal{R}$ **do**
> > **if** o of form rew_b **then**
> > > Add (p', b) to H and $((p, a), r, (p', b))$ to E
> > > Call AddStackDescriptor($(p', b), (h_n, \ldots, h_1, h_c)$)
> > **else if** o of form $push_b^k$ **then**
> > > Add (p', b) to H and $((p, a), r, (p', b))$ to E
> > > Call AddStackDescriptor($(p', b), (h_n, \ldots, h_2, (p, a), h_k)$)
> > **else if** o of form $push_k$ **then**
> > > Add (p', a) to H and $((p, a), r, (p', a))$ to E
> > > Call AddStackDescriptor($(p', a), (h_n, \ldots, h_{k+1}, (p, a),$
> > > $\quad h_{k-1}, \ldots, h_1, h_c)$)
> > **else if** o of form pop_k with $h_k = (p_k, a_k)$ where $a_k \neq \perp$
> > **then**
> > > Add (p', a_k) to H and $((p, a), r, (p', a_k))$ to E
> > > Call AddSummary($(p_k, a_k), (h_n, \ldots, h_{k+1}), (p', a_k)$)
> > **else if** o of form $collapse_k$ with $h_c = (p_c, a_c)$ where $a_c \neq \perp$
> > **then**
> > > Add (p_c, a_c) to H and $((p, a), r, (p', a_c))$ to E
> > > Call AddSummary($(p_c, a_c), (h_n, \ldots, h_{k+1}), (p', a_c)$)
> **return** Done

Algorithm 4 AddSummary($h, (h'_n, \ldots, h_{k+1}), h'$)

Require: An approximate higher-order summary edge $(h, (h'_n, \ldots, h_{k+1}), h')$

Ensure: $(h, (h'_n, \ldots, h'_{k+1}), h') \in U$ and that all necessary stack descriptors are added to the appropriate $B(h'')$ for $h'' \in H$ so that all summary edges (including the new one) are respected.

> **if** $(h, (h'_n, \ldots, h'_{k+1}), h') \in U$ **then**
> > **return** Done (Nothing to do)
> Add $(h, (h'_n, \ldots, h'_{k+1}), h')$ to U
> **for** $(h_n, \ldots, h_{k+1}, h_k, \ldots, h_1, h_c) \in B(h)$ **do**
> > AddStackDescriptor($h', (h'_n, \ldots, h'_{k+1}, h_k, \ldots, h_1, h_c)$)
> **return** Done

---

lowing a loop indefinitely. This is a non-trivial property requiring a subtle relation on runs over higher-order collapsible stacks.

First, to describe how to construct counter examples, we need an alternative, more manipulable, definition of stack automaton runs.

6.1 Alternative definition of the runs

We give an alternative definition of a run of a stack automaton that is more appropriate to perform the *run surgery* needed below. The definition of an accepting run requires two intermediary notions.

A run of A on an order-$\leq n$ stack v is an annotation of each symbol of this stack by a subset of Δ_1, the order-1 transitions of A. Formally a run over an order-k stack v is an order-k stack over the alphabet $\Sigma \times 2^{\Delta_1}$ such that when projecting on the Σ-component we retrieve the stack v.

Let w be an order-k run of A. For a set $Q \subseteq \mathcal{Q}_k$ of order-k states, we say that w *is Q-valid* if the following holds. If the run w is empty, Q must be a subset of $\mathcal{F}_k$. Assume now that w is not empty. If $k = 1$ and $w = (a, T) :_1 w'$, there must exist $Q' \subseteq \mathcal{Q}_1$ such that w' is Q'-valid and for all $q \in Q$, there exists a transition in T of the form $q \xrightarrow[Q_{col}]{a} Q''$ with $Q'' \subseteq Q'$. If $k > 1$

and $w = u :_k w'$ then there must exist a subset $Q' \subseteq \mathcal{Q}_k$ of order-k states such that w' is Q'-valid and for all $q \in Q$, there exists a transition $q \xrightarrow{q_{Q''}} Q'' \in \Delta_k$ such that u is $\{q_{Q''}\}$-valid.

Note that the notion of Q-validity does not check the constraint imposed by the Q_{col} component appearing in order-1 transitions. This is done by the notion of *link-validity* which is only meaningful on order-n runs: An order-n run w is *link-valid* if for every substack of w of the form $w' = (a, T)^{k,i} : w''$ and every transition $q \xrightarrow[Q_{col}]{a} Q$ then $top_{k+1}(collapse_k(w'))$ is Q_{col}-valid.

For $q \in \mathcal{Q}_n$, an order-n run w is q-accepting if it is both $\{q\}$-valid and link-valid. In addition, we require that if w is non empty and hence of the form $(a, T) :_1 w'$ then T is reduced to a singleton $\{t\}$ and we refer to t as the head transition of the run.

6.2 Extracting a Counter Example

Fix a CPDS $\mathcal{C}$ and a stack automaton A_0. Let A be the result of the saturation algorithm for $\mathcal{C}$ and A_0. The key idea is that, thanks to the transition justifications, we can define a well-founded relation on runs of A. We define by induction a relation $\hookrightarrow_k$ on the order-k runs of A (runs are represented by annotated stacks as explained in Section 6.1). Note that this is not an order relation as it is not always transitive. There are several cases to $\hookrightarrow_k$.

1. For $k = 1$, we say $w' \hookrightarrow_1 w$ if for some $i \geq 0$, w contains strictly fewer transitions in Δ_1 justified at step i than w' and that for all $j > i$ they both contain the same number of transitions in Δ_1 justified at step j.

2. For $k > 1$, we say $u = [u_\ell \ldots u_1]_k \hookrightarrow_k v = [v_{\ell'} \ldots v_1]_k$ if

   (a) $\ell' < \ell$ and $u_i = v_i$ for $i \in [1, \ell' - 1]$ and either $u_{\ell'} = v_{\ell'}$ or $u_{\ell'} \hookrightarrow_{k-1} v_{\ell'}$, or

   (b) $\ell' \geq \ell$ and $u_i = v_i$ for $i \in [1, \ell - 1]$ and $u_\ell \hookrightarrow_{k-1} v_i$ for all $i \in [\ell, \ell']$.

Lemma 6.1. *For all $k \in [1, n]$, the relation $\hookrightarrow_k$ is well-founded. Namely there is no infinite sequence $w_0 \hookrightarrow_k w_1 \hookrightarrow_k w_2 \hookrightarrow_k \cdots$.*

Algorithm 5 shows how we construct counter examples. Variable w contains a run of A which is q_p-accepting for some state p. The initial value of w, denoted w_0, is an accepting run for the initial configuration $\langle p_0, u_0 \rangle$. We update w at the end of each iteration of the while-loop. Let w_i be the value of w at the end of the i-th iteration which we assume to be an accepting run for a configuration $\langle p_i, u_i \rangle$. Moreover, let t_i denote the head transition of w_i.

The invariant of the algorithm is that t_{i-1} has a justification containing a rule of the form $(p_{i-1}, a_{i-1}, o, p_i)$ with $top_1(u_{i-1}) = a_{i-1}$ and $u_i = o(u_{i-1})$. Moreover and crucially for termination, $w_{i-1} \hookrightarrow_n w_i$. As $\hookrightarrow_n$ is well-founded, we eventually exit the while-loop after N iterations with t_N justified by 0. It is then possible to prune the last run w_N to form a run that consists entirely of transitions already belonging to A_0 by the assumption that initial states at every level of A have no incoming transitions. It follows that the last configuration reached $\langle p_N, u_N \rangle$ belongs to $\mathcal{L}(A_0)$.

Proposition 6.1. *Algorithm 5 is correct.*

7. Efficient Fixed Point Computation

We introduce an efficient method of computing the fixed point in Section 3, inspired by Schwoon *et al.*'s algorithm for alternating (order-1) pushdown systems [36]. Rather than checking all CPDS rules at each iteration, we fully process *all* consequences of each new transition at once. New transitions are kept in a set Δ_{new} (implemented as a stack), processed, then moved to a set Δ_{done}, which forms the transition relation of the final stack automaton. We assume w.l.o.g. that a character's link order is determined by the character. This is true of all CPDSs obtained from HORS.

Algorithm 5 Counter-example Extraction

Require: A stack-automaton A generated by saturating A_0.
Ensure: Print a finite sequence of CPDS rules that when executed will lead from the initial
configuration $\langle p_0, [\cdots[a]_1 \cdots]_n\rangle$ to one in $\mathcal{L}(A_0)$.
$\boldsymbol{w} := [\cdots[\{t_0\}]_1 \cdots]_n$ where $t_0 := q_{p_0} \xrightarrow[\emptyset]{a} (\emptyset, \ldots, \emptyset)$

    while the head transition t of $\boldsymbol{w}$ is not justified by 0 **do**
        Print $\boldsymbol{r} = (p, a, o, p')$, the CPDS transition appearing in the justification of t.
        if $o = pop_k$ for some $1 \le k \le n$ **then**
            The transition t is of the form $q_p \xrightarrow[\emptyset]{a} (\emptyset, \ldots, \emptyset, \{q_{p',Q_n,\ldots,Q_{k+1}}\}, Q_{k+1}, \ldots, Q_n)$

            Pick $Q_k, \ldots, Q_1, Q_{col}$ such that $t' := q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k, Q_{k+1}, \ldots, Q_n)$ is in

            $top_1(pop_k(\boldsymbol{w}))$
            $\boldsymbol{w} := rew_{\{t'\}}(pop_k(\boldsymbol{w}))$
        else if $o = collapse_k$ for $2 \le k \le n$ **then**
            The transition t is of the form $t = q_p \xrightarrow[\{q_{p',Q_n,\ldots,Q_{k+1}}\}]{a} (\emptyset, \ldots, \emptyset, Q_{k+1}, \ldots, Q_n)$

            Pick $Q_k, \ldots, Q_1, Q_{col}$ such that $t' := q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k, Q_{k+1}, \ldots, Q_n)$ is in

            $top_1(collapse_k(\boldsymbol{w}))$
            $\boldsymbol{w} := rew_{\{t'\}}(collapse_k(\boldsymbol{w}))$
        else if $o = rew_b$ for some $b \in \Sigma$ **then**
            $J(t)$ must be of form (r, t', i), set $\boldsymbol{w} := rew_{\{t'\}}(rew_b(\boldsymbol{w}))$
        else if $o = push_k$ or $o = push_b^{k'}$ **then**
            $J(t)$ must be of the form (r, t', T, i), set $\boldsymbol{w} := rew_{\{t'\}}(o(rew_T(\boldsymbol{w})))$

In most cases, new transitions only depend on a single existing transition, hence processing the consequences of a new transition is straightforward. The key difficulty is the push rules, which depend on *sets* of existing transitions. Given a rule $(p, a, push_k, p')$, processing $t = q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k, \ldots, Q_n)$ 'once and once only' must somehow include adding a new transition whenever there is a set of transitions of the form $Q_k \xrightarrow[Q'_{col}]{a} (Q'_1, \ldots, Q'_k)$ in A_i either now *or in the future*. When t is processed, we therefore create a *trip-wire*, consisting of a *source* and *target*. A *target* collects transitions from a given set of states (such as Q_k above), whilst a *source* describes how such a collection could be used to form a new transition according to a *push* saturation step.

Definition 7.1. *An order-k source for $k \ge 1$ is defined as a tuple* (q_k, q_{k-1}, a, Q_k) *in* $\mathcal{Q}_k \times \mathcal{Q}_{k-1}^\perp \times \Sigma \times 2^{\mathcal{Q}_k}$ *where* $\mathcal{Q}_0^\perp := \{\perp\}$ *and* $\mathcal{Q}_i^\perp = \mathcal{Q}_i \cup \{\perp\}$ *for* $i \ge 1$. *An order-k target is a tuple*

$$(Q_k, Q_k^C, Q_{lbl}, Q'_k) \in 2^{\mathcal{Q}_k} \times 2^{\mathcal{Q}_k} \times 2^{\mathcal{Q}_{k-1}} \times 2^{\mathcal{Q}_k}$$

if $k \ge 2$, and if $k = 1$

$$(Q_1, Q_1^C, a, Q_{col}, Q'_1) \in \bigcup_{k'=2}^{n} \left(2^{\mathcal{Q}_1} \times 2^{\mathcal{Q}_1} \times \Sigma \times 2^{\mathcal{Q}_{k'}} \times 2^{\mathcal{Q}_1}\right).$$

The set Q_k^C is a *countdown* containing states in Q_k still awaiting a transition. We always have $Q_k^C \subseteq Q_k$ and $(Q_k \setminus Q_k^C) \xrightarrow{Q_{lbl}} Q'_k$. Likewise, an order-1 target $(Q_1, Q_1^C, a, Q_{col}, Q_1)$ will satisfy $(Q_1 \setminus Q_1^C) \xrightarrow[Q_{col}]{a} Q'_1$. A target is *complete* if $Q_k^C = \emptyset$ or $Q_1^C = \emptyset$.

A *trip-wire* of order-k is an order-k source-target pair of the form $((\_, \_, \_, Q_k), (Q_k, \_, \_, \_))$ when $k \ge 2$ or $((\_, \_, a, Q_k), (Q_k, \_, a, \_, \_))$ when $k = 1$. When the target in a trip-wire is *complete*, the action specified by its source is triggered, which we now sketch.

An order-k source for $k \ge 2$ describes how an order-$(k-1)$ source should be created from a complete target, propagating the computation to the level below, and an order-1 source describes

how a new long-form transition should be created from a complete target. That is, when we have $(q_k, \_, a, Q_k)$ (we hide the second component for simplicity of description) and $(Q_k, \emptyset, Q_{lbl}, Q'_k)$ this means we've found a set of transitions witnessing $Q_k \xrightarrow{Q_{lbl}} Q'_k$ and should now look for transitions from Q_{lbl}. Hence the algorithm creates a new source and target for the order-$(k-1)$ state-set Q_{lbl}. When this process reaches order-1, a new transition is created. This results in the construction of the t' from a *push* saturation step.

Algorithm 6 gives the main loop and introduces the global sets Δ_{done} and Δ_{new}, and two arrays $\mathcal{U}_{src}[k]$ and $\mathcal{U}_{targ}[k]$ containing sources and targets for each order. Omitted are loops processing pop_n and $collapse_n$ rules like the naive algorithm. Algorithm 7 gives the main steps processing a new transition. We present only two CPDS rule cases here. In most cases a new transition is created, however, for *push* rules we create a trip-wire. Remaining algorithms, definitions, justification handling, and proofs are given in Appendix B. We describe some informally below.

In create_trip_wire we create a trip-wire with a new target $(Q_k, Q_k, \emptyset, \emptyset)$. This is added using an add_target procedure which also checks Δ_{done} to create further targets. E.g., a new target $(Q_k, Q_k^C, Q_{lbl}, Q'_k)$ combines with an existing $q_k \xrightarrow{q_{k-1}} Q''_k$ to create a new target $(Q, Q_k^C \setminus \{q_k\}, Q_{lbl} \cup \{q_{k-1}\}, Q'_k \cup Q''_k)$. (This step corrects a bug in the algorithm of Schwoon *et al.*) Similarly update_trip_wires updates existing targets by new transitions. In all cases, when a source and matching complete target are created, we perform the propagations described above.

Proposition 7.1. *Given a CPDS $\mathcal{C}$ and stack automaton A_0, let A be the result of Algorithm 6. We have $\mathcal{L}(A) = Pre_{\mathcal{C}}^*(A_0)$.*

8. Experimental Results

We compared C-SHORe with the current state-of-the-art verification tools for higher-order recursion schemes (HORS): TRecS [23], GTRecS2 [26] (the successor of [25]), and TravMC [28]. Benchmarks are from the TRecS and TravMC benchmark suites. The majority of the TravMC benchmarks were translated into HORS from

Algorithm 6 Computing $Pre^*_{\mathcal{C}}(A_0)$

Let $\Delta_{done} = \emptyset$, $\Delta_{new} = \bigcup_{n \geq k \geq 1} \Delta_k$, $\mathcal{U}_{src}[k] = \emptyset$, $\mathcal{U}_{targ}[k] = \{(\emptyset, \emptyset, \emptyset, \emptyset)\}$ for each $n \geq k > 1$ and $\mathcal{U}_{targ}[1] = \{(\emptyset, \emptyset, a, \emptyset, \emptyset) \mid a \in \Sigma\}$.
. . .
while $\exists t \in \Delta_{new}$ **do**
    update_rules(t); update_trip_wires(t); move t from Δ_{new} to Δ_{done}

Algorithm 7 update_rules(t)

if t is an order-k transition for $2 \leq k \leq n$ of the form $q_{p'Q_n\cdots Q_{k+1}} \longrightarrow Q_k$ **then**
    for $p \in \mathcal{P}$ and $a \in \Sigma$ such that $r := (p, a, pop_k, p') \in \mathcal{R}$ **do**
      add_to_worklist$\left(q_p \xrightarrow[\mathcal{Q}_{col}]{a} (\emptyset, \ldots, \emptyset, \{q_{p'\mathcal{Q}_n\cdots\mathcal{Q}_{k+1}}\}, \mathcal{Q}_{k+1}, \ldots, \mathcal{Q}_n), r \right)$
    for $p \in \mathcal{P}$ and $a \in \Sigma$ such that $r := (p, a, push_k, p') \in \mathcal{R}$ **do**
      create_trip_wire$\left(q_{p,Q_n,\ldots,Q_{k+1}}, q_{p',Q_n,\ldots,Q_{k+1},Q_k}, a, Q_k, (r, t) \right)$
. . .

---

an extended formalism, HORS with Case statements (HORSC), using a script by Kobayashi. For fairness, all tools in our experiments took a pure HORS as input. However, the authors of TravMC report that TravMC performs faster on the original HORSC examples than on their HORS translations.

The experiments were run on a Dell Latitude e6320 laptop with 4Gb of RAM and four 2.7GHz Intel i7-2620M cores. We ran C-SHORe on OpenJDK 7.0 with IcedTea7 replacing binary plugs, using the argument "-Xmx" to limit RAM usage to 2.5Gb. As advised by the TravMC developers, we ran TravMC on the Mono JIT Compiler version 3.0.3 with no command line arguments. Finally TRecS and GTRecS2 were compiled with the OCaml version 4.00.1 compilers. On negative examples, GTRecS2 was run with its -neg argument. We used the "ulimit" command to limit memory usage to 2.5Gb and set a CPU timeout of 600 seconds (per benchmark). The given runtimes were reported by the respective tools and are the means of three separate runs on each example.

Table 1 shows trials where at least one tool took over 1s. This is because virtual machine "warm-up" and HORS to CPDS conversion can skew the results on small benchmarks. Full results are in Appendix D. Examples violating their property are marked "(bug)". The order (Ord) and size (Sz) of the schemes were reported by TRecS. We show reported times in seconds for TRecS (T), GTRecS2 (G), TravMC (TMC) and C-SHORe (C) where "—" means analysis failed. For C-SHORe, we report the times for HORS to CPDS translation (Ctran), CPDS analysis (Ccpds), and building the approximation graph (Capprox). Capprox is part of Ccpds, and the full time (C) is the sum of Ctran and Ccpds.

Of 12 benchmarks, C-SHORe performed best on the order-2 filepath. In 4 cases, C-SHORe was the slowest. The only example that did not terminate is g41. The examples filter-nonzero and map-plusone are the smallest of the families of examples; C-SHORe performed well on the larger family members. The remaining slowest example safe-init only required a short verification time. Overall, TRecS appears to be the fastest; however, it is also the most mature and hence most highly optimised. C-SHORe does outperform TravMC in a number of cases, by a large margin for filter-nonzero-1.

It is also very important to note that C-SHORe and GTRecS2 are the only implemented fixed-parameter tractable algorithms in the literature for HORS model-checking of which we are aware (both TRecS and TravMC have worst-case run-times non-elementary in the size of the recursion scheme). Moreover, C-SHORe generally performs much better than GTRecS2. Thus not only does C-SHORe's performance seem promising when compared to the competition, there is also theoretical reason to suggest that the approach could in principle be scalable, in contrast to some of the alternatives. Thus initial work justifies further investigation into saturation based algorithms for higher-order model-checking.

Finally, we remark that without the forwards analysis described in Section 5, all shown examples except filepath timed out. We also note that we did not implement a naive version of the saturation algorithm, where after each change to the stack automaton, each

rule of the CPDS is checked for further updates. However, experience implementing PDSolver [15] (for order-1 pushdown systems) indicates that the naive approach is at least an order of magnitude slower than the techniques [36] we generalised in Section 7.

9. Related Work

The saturation technique has proved popular in the literature. It was introduced by Bouajjani *et al.* [4] and Finkel *et al.* [13] and based on a string rewriting algorithm by Benois [3]. It has since been extended to Büchi games [8], parity and μ-calculus conditions [15], and concurrent systems [1, 37], as well as weighted pushdown systems [31]. In addition to various implementations, efficient versions of these algorithms have also been developed [12, 36].

The saturation algorithm for CPDS that we introduced in [7], extending and improving [14] (and [5]), follows a number of papers solving parity games on the configuration graphs of higher-order automata [6, 9, 11, 16]. While only handling reachability, saturation lends itself well to implementation. This paper describes such a practical incarnation and a number of significant optimisations, such as using a forwards analysis to guide the backward search.

This latter point is an important way in which C-SHORe differs from previous model-checkers for HORS, which employ intersection types and propagate information purely in a forward direction. This is related to the fact that the latter accept 'co-reachability properties' (represented by trivial Büchi automata) as input, expressing the complement of properties taken by C-SHORe.

Indeed it would be interesting to investigate in more detail how approximate forward and backward analyses of varying degrees of accuracy could be combined for efficiency. It would also be helpful to more closely analyse the relationship between CPDS and type-based algorithms allowing a transfer of ideas. In any case, this paper shows that saturation-based algorithms for HORS/CPDS perform sufficiently well in practice to warrant further study.

To finish, we briefly mention several approaches to analysing higher-order programs with differing aims to ours. In static analysis, k-CFA [35] and CFA2 [39] perform an over-approximative analysis of higher-order languages with at-most first-order granularity. Similarly Jhala *et al.* use refinement types to analyse OCaml programs by reducing the problem to first-order model-checking, which is thus incomplete [19]. Finally, Hopkins *et al.* have produced tools for equivalence checking fragments of ML and Idealized Algol up to order-3 [17, 18].

10. Conclusion

We have considered the problem of verifying safety properties of a model that can be used to precisely capture control-flow in the presence of higher-order recursion. Whilst previous approaches to such an analysis are based on higher-order recursion schemes and intersection types, our approach is based on automata and saturation techniques previously only applied in practice to the first-order case. At a more conceptual level, our algorithm works by propagating information backwards from error states towards the initial state. Moreover, it combines this with an approximate

Benchmark file	Ord	Sz	T	TMC	G	C	Ctran	Ccpds	Capprox
`order5`	5	52	0.010	0.045	—	**0.531**	0.098	0.434	0.196
`order5-2`	5	40	0.017	0.091	—	**0.317**	0.082	0.235	0.128
`order5-variant`	5	55	0.017	0.042	**1.470**	**0.525**	0.095	0.430	0.170
`filepath`	2	5956	210.767	—	—	**0.657**	0.426	0.231	0.192
`filter-nonzero (bug)`	5	484	0.008	0.126	**0.163**	**2.150**	0.177	1.973	1.149
`filter-nonzero-1`	5	890	0.170	216.842	—	**6.879**	0.399	6.480	4.120
`map-head-filter-1`	3	880	0.137	1.359	—	**0.469**	0.205	0.263	0.239
`map-plusone`	5	302	0.013	0.126	**0.126**	**1.282**	0.162	1.120	0.871
`map-plusone-1`	5	459	0.023	0.760	—	**1.771**	0.223	1.547	1.031
`map-plusone-2`	5	704	1.325	14.242	—	**3.559**	0.323	3.236	2.184
`safe-init`	3	680	0.031	0.228	**0.300**	**1.233**	0.171	1.062	0.175
`g41`	4	31	—	0.107	**0.054**	—	0.068	—	0.244

Table 1: Comparison of model-checking tools. Shown in bold are the two fixed-parameter tractable algorithms, GTRecS2 and C-SHORe.

forward analysis to gather information that guides the backward search. In contrast, the preceding type-based algorithms all work by propagating information purely in a forward direction.

Our preliminary work brings new techniques to the table for tackling a problem, which in contrast to its first-order counterpart, has proven difficult to solve in a scalable manner. Our algorithm has the advantage that it accurately models higher-order recursion whilst also being fixed-parameter tractable, therefore giving a theoretical reason for hope that it could scale. In contrast TRecS and TravMC have worst-case run-times non-elementary in the size of the recursion scheme. Our tool also seems to work significantly better in practice than GTRecS2, the only other HORS model-checker in the literature that does enjoy fixed-parameter tractability.

We therefore believe that a C-SHORe-like approach shows much promise and warrants further investigation.

References

[1] M. F. Atig. Global model checking of ordered multi-pushdown systems. In *FSTTCS*, 2010.

[2] T. Ball and S. K. Rajamani. The SLAM project: Debugging system software via static analysis. In *POPL*, 2002.

[3] M. Benois. Parties rationnelles du groupe libre. *Comptes-Rendus de l'Acamdémie des Sciences de Paris, Série A*, 269:1188–1190, 1969.

[4] A. Bouajjani, J. Esparza, and O. Maler. Reachability analysis of pushdown automata: Application to model-checking. In *CONCUR*, 1997.

[5] A. Bouajjani and A. Meyer. Symbolic Reachability Analysis of Higher-Order Context-Free Processes. In *FSTTCS*, 2004.

[6] C. H. Broadbent, A. Carayol, C.-H. L. Ong, and O. Serre. Recursion schemes and logical reflection. In *LICS*, 2010.

[7] C. H. Broadbent, A. Carayol, M. Hague, and O. Serre. A saturation method for collapsible pushdown systems. In *ICALP*, 2012.

[8] T. Cachat. *Games on Pushdown Graphs and Extensions*. PhD thesis, RWTH Aachen, 2003.

[9] T. Cachat. Higher order pushdown automata, the Caucal hierarchy of graphs and parity games. In *ICALP*, 2003.

[10] A. Carayol and O. Serre. Collapsible pushdown automata and labeled recursion schemes: Equivalence, safety and effective selection. In *LICS*, 2012.

[11] A. Carayol, M. Hague, A. Meyer, C.-H. L. Ong, and O. Serre. Winning Regions of Higher-Order Pushdown Games. In *LICS*, 2008.

[12] J. Esparza, D. Hansel, P. Rossmanith, and S. Schwoon. Efficient algorithms for model checking pushdown systems. In *CAV*, 2000.

[13] A. Finkel, B. Willems, and P. Wolper. A direct symbolic approach to model checking pushdown systems. In *INFINITY*, 1997.

[14] M. Hague and C.-H. L. Ong. Symbolic backwards-reachability analysis for higher-order pushdown systems. *Logical Methods in Computer Science*, 4(4), 2008.

[15] M. Hague and C.-H. L. Ong. Analysing mu-calculus properties of pushdown systems. In *SPIN*, 2010.

[16] M. Hague, A. S. Murawski, C.-H. L. Ong, and O. Serre. Collapsible pushdown automata and recursion schemes. In *LICS*, 2008.

[17] D. Hopkins and C.-H. L. Ong. Homer: A higher-order observational equivalence model checker. In *CAV*, 2009.

[18] D. Hopkins, A. S. Murawski, and C.-H. L. Ong. Hector: An equivalence checker for a higher-order fragment of ml. In *CAV*, 2012.

[19] R. Jhala, R. Majumdar, and A. Rybalchenko. Hmc: Verifying functional programs using abstract interpreters. In *CAV*, 2011.

[20] N. D. Jones and S. S. Muchnick. Even simple programs are hard to analyze. *J. ACM*, 24:338–350, April 1977.

[21] T. Knapik, D. Niwinski, P. Urzyczyn, and I. Walukiewicz. Unsafe grammars and panic automata. In *ICALP*, 2005.

[22] N. Kobayashi. Types and higher-order recursion schemes for verification of higher-order programs. In *POPL*, 2009.

[23] N. Kobayashi. Model-checking higher-order functions. In *PPDP*, 2009.

[24] N. Kobayashi. Higher-order model checking: From theory to practice. In *LICS*, 2011.

[25] N. Kobayashi. A practical linear time algorithm for trivial automata model checking of higher-order recursion schemes. In *FOSSACS*, 2011.

[26] N. Kobayashi. GTRECS2: A model checker for recursion schemes based on games and types. A tool available at `http://www-kb.is.s.u-tokyo.ac.jp/~koba/gtrecs2/`, 2012.

[27] N. Kobayashi, R. Sato, and H. Unno. Predicate abstraction and cegar for higher-order model checking. In *PLDI*, 2011.

[28] R. P. Neatherway, S. J. Ramsay, and C.-H. L. Ong. A traversal-based algorithm for higher-order model checking. In *ICFP*, 2012.

[29] C.-H. L. Ong. On model-checking trees generated by higher-order recursion schemes. In *LICS*, 2006.

[30] C.-H. L. Ong and S. J. Ramsay. Verifying higher-order functional programs with pattern-matching algebraic data types. In *POPL*, 2011.

[31] T. W. Reps, S. Schwoon, S. Jha, and D. Melski. Weighted pushdown systems and their application to interprocedural dataflow analysis. *Sci. Comput. Program.*, 58(1-2):206–263, 2005.

[32] S. Salvati and I. Walukiewicz. Recursive schemes, krivine machines, and collapsible pushdown automata. In *RP*, 2012.

[33] S. Schwoon. *Model-checking Pushdown Systems*. PhD thesis, Technical University of Munich, 2002.

[34] M. Sharir and A. Pnueli. *Two approaches to interprocedural data flow analysis*, chapter 7, pages 189–234. Prentice-Hall, 1981.

[35] O. Shivers. *Control-Flow Analysis of Higher-Order Languages*. PhD thesis, Carnegie-Mellon University, 1991.

[36] D. Suwimonteerabuth, S. Schwoon, and J. Esparza. Efficient algorithms for alternating pushdown systems with an application to the computation of certificate chains. In *ATVA*, 2006.

[37] D. Suwimonteerabuth, J. Esparza, and S. Schwoon. Symbolic context-bounded analysis of multithreaded java programs. In *SPIN*, 2008.

[38] H. Unno, N. Tabuchi, and N. Kobayashi. Verification of tree-processing programs via higher-order model checking. In *APLAS*, 2010.

[39] D. Vardoulakis. *CFA2: Pushdown-Flow Analysis for Higher-Order Languages*. PhD thesis, Northeastern University, Boston, 2012.

[40] D. Vardoulakis and O. Shivers. Pushdown flow analysis of first-class control. In *ICFP*, 2011.

A. Counter Example Extraction

In this section, we establish the correctness of Algorithm 5 and give omitted proofs. We start with the proof of Lemma 6.1.

Proof. For $k = 1$, consider for any order-1 run w the tuple $|w| = (n_m, \ldots, n_0)$ where m is the step at which the saturation algorithm terminates and for all $i \in [0, m]$, n_i is the number of occurrences in w of transitions in Δ_1 justified at step i. The relation $\hookrightarrow_1$ can be equivalently defined as $w \hookrightarrow_1 w'$ if $|r'|$ is lexicographically smaller than $|r|$. It immediately follows that $\hookrightarrow_1$ is well-founded.

For $k + 1 > 1$ assuming the property holds for $\hookrightarrow_k$. Suppose for contradiction that $\hookrightarrow_{k+1}$ is not well-founded. Then there must be an infinite chain of runs of the form:

$$w_1 \hookrightarrow_{k+1} w_2 \hookrightarrow_{k+1} w_3 \hookrightarrow_{k+1} \cdots$$

Now pick an index i such that for every $j > i$ it is the case that w_j is at least as long (w.r.t the number of order-$(k-1)$ stacks) as the run w_i (infinitely many such indices must clearly exist since comparing runs by their lengths is a well-founded relation.). If $w_i = u : w'_i$, it is a straightforward induction to see that for every $j > i$ w_j is of the form $w''_j w'_i$ with $u \hookrightarrow_k^+ v$ for all order-k run v occuring in w''_j where $\hookrightarrow_k^+$ designates the transitive closure of $\hookrightarrow_k$.

So in particular if we pick infinitely many positions in the chain i_ℓ such that the run $w_{i_\ell} = u_{i_\ell} : w'_{i_\ell}$ is at least as long as the sequence w_j for all $j > i_\ell$ it must be the case that:

$$u_{i_1} \hookrightarrow_k^+ u_{i_2} \hookrightarrow_k^+ u_{i_3} \hookrightarrow_k^+ \cdots$$

This in turn contradicts the fact that $\hookrightarrow_k$ is well-founded. $\qquad\square$

The following lemma describes two simple sufficient conditions condition for $r \hookrightarrow_k r'$ to hold.

Lemma A.1. *The following properties hold:*

1. *Let w and w' be two order-n runs such that for some $1 \leq k < n$, $top_{k+1} \hookrightarrow_k top_{k+1}(w')$ and $pop_{k+1}(w) = pop_{k+1}(w')$ then $r \hookrightarrow_n r'$.*

2. *Let w be an order-k run and let T be a set of transitions that is smaller that some transition appearing in $top_1(w)$, $r \hookrightarrow_k rew_T(w)$.*

Proof. For the first property, we will show by induction on k' that for all $k' \in [k, n]$, $top_{k'+1}(w) \hookrightarrow_{k'} top_{k'+1}(w')$. The case $k' = k$ is assumed to hold in the hypothesis. Assume that the property holds for k'. We have $top_{k'+2}(w) = top_{k'+1}(w) :_{k'} top_{k'+2}(pop_{k'+1}(w))$ and $top_{k'+2}(w') = top_{k'+1}(w') :_{k'} top_{k'+2}(pop_{k'+1}(w'))$. Remark that $pop_{k'+1}(w) = pop_{k'+1}(w') = u$. This is by assumption for $k = k'$ and if $k' > k$ $pop_{k'+2}(w) = pop_{k'+2}(pop_{k+1}(w)) = pop_{k'+2}(pop_{k+1}(w')) = pop_{k'+2}(pop_{k+1}(w'))$. Hence $top_{k'+2}(w) = top_{k'+1}(w) :_{k'} u$ and $top_{k'+2}(w') = top_{k'+1}(w') :_{k'} u$ with $top_{k'+1} \hookrightarrow_{k'} top_{k'+1}(w')$. By definition of $\hookrightarrow_{k'+1}$, we have $top_{k'+2} \hookrightarrow_{k'+1} top_{k'+2}(w')$.

For the second property, we have $top_2(w) \hookrightarrow_1 top_2(rew_T(w))$ (by definition of $\hookrightarrow_1$) and $pop_2(w) = pop_2(w')$. Hence by the first-property $r \hookrightarrow_n rew_T(w)$. $\qquad\square$

We now prove the correctness of Algorithm 5.

As it is often the case, we restrict our attention to runs containing only "useful" transitions. A run w is *trimmed* if for any $o_1, \ldots, o_j$ of *pop* operations producing a subrun $w' = o_j(\ldots o_1(w) \ldots)$, for any order-1 transition

$$q \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$$

appearing in $top_1(w')$, we have for all $i \in [1, m-1]$, that $top_{i+1}(pop_i(w'))$ is Q_i-valid where m is the smallest index such that pop_m appears in the sequence $o_1, \ldots, o_j$.

We now give the proof of Proposition 6.1.

Proof. The initial value of w, denoted w_0, is an accepting run for the initial configuration $\langle p_0, u_0 \rangle$. The value of w is updated at the end of each iteration of the while-loop. We denote by w_i the value of w at the end of the i-th iteration. Let N be the total number of the iteration of the while-loop. Strictly speaking we have not yet proved that the algorithm terminates so N could be equal to ∞.

We are going to prove by induction on the iteration step i that w_i is a trimmed q_{p_i}-accepting run on some stack s_i. Furthermore for $i > 0$, the head transition has a justification containing a transition of the CPDS of the form $(p_{i-1}, top_1(s_{i-1}), o, p_i)$ and $s_i = o(s_{i-1})$. Furthermore we have $w_{i-1} \hookrightarrow_n w_i$.

For $i = 0$ the property is immediate as w_0 which contains only one transition is necessarily trimmed. Assume that the property holds for $i \geq 0$, let us prove it for $i + 1$. To simplify the writing let us write w for w_i and w' for w_{i+1}. Similarly, we write p for p_i. By the induction hypothesis, w is a trimmed q_p-accepting run on a stack s. This implies that its head transition is of the form:

$$t = q_p \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n) .$$

Hence its justification contains a transition of the CPDS of the form (p, a, o, p'). We take p_{i+1} equal to p'.

We now reason by case distinction on the operation o.

If $o = rew_b$ **for some** $b \in \Sigma$**.** The transition t is of the form

$$q_p \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$$

with a justification of the form $J(t) = (r, t', i)$ with t' of the form

$$q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$$

Note that t' was introduced before t.

The run w' is equal to $rew_{\{t'\}}(rew_b(w))$ It is clear that w' is a trimmed $q_{p'}$-accepting run on the stack $rew_b(s)$. By the second property of Lemma A.1, $w \hookrightarrow_n w'$.

If $o = pop_k$ **for some** $k \in [1, n]$**.** The transition t is of the form

$$q_p \xrightarrow{a}_{\emptyset} (\emptyset, \ldots, \emptyset, \{q_{p', Q_n, \ldots, Q_{k+1}}\}, Q_{k+1}, \ldots, Q_n).$$

As by w is q_p-accepting, it follows that for all $j \in [k+1, n]$, $top_{j+1}(pop_j(w))$ is Q_j-valid and that $top_{k+1}(pop_k(w))$ is $\{q_*\}$-valid for $q_* = q_{p', Q_n, \ldots, Q_{k+1}}$. By unfolding the notion of $\{q_*\}$-validity, we obtain that $top_1(pop_k(w))$ contains at least one transition t' of the form:

$$q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k, Q_{k+1}, \ldots, Q_n)$$

Let t' the transition of this form picked by the algorithm. As w is trimmed it follows that for all $j \in [1, k]$, $top_{j+1}(pop_j(pop_k(w)))$ is Q_j-valid.

We have $w' = rew_{\{t'\}}(pop_k(w))$. As w' is a subrun of w (which is link-valid and trimmed), it is link-valid and trimmed. It is also $q_{p'}$-valid it is enought to show that for all $i \in [1, n]$, we have $top_{i+1}(pop_i(w')$ is Q_i-valid. For $i \in [k+1, n]$, we have seen that $top_{i+1}(pop_i(w')) = top_{i+1}(pop_i(w))$ is Q_i-valid. For $i \in [1, k]$, we have seen that $top_{i+1}(pop_i(w')) = top_{i+1}(pop_i(pop_k(w)))$ is Q_i-valid.

It only remains to show that $r \hookrightarrow_n r'$. By the first property of Lemma A.1, it is enought to show that $top_{k+1}(w) \hookrightarrow_k top_k(w')$ (as $pop_{k+1}(w') = pop_{k+1}(w)$ if $k < n$). First consider the case when $k = 1$. It follows from the fact that the set of order-1 transitions appearing in $top_2(w')$ is stricly included in $top_2(w)$. Now assume that $k > 1$. The run $top_{k+1}(w)$ can be written as $u :_{k+1} u' :_{k+1} v$

and $top_{k+1}(w) = rew_T(u') :_{k+1} v$. By the second property of Lemma A.1, $u' \hookrightarrow k - 1rew_T(u')$ and by defintion of $\hookrightarrow k$, $top_{k+1}(w) \hookrightarrow_k top_k(w')$.

If $o = collapse_k$ for some $k \in [2, n]$. This case is similar to the *pop* case.

If $o = push_k$ for some $k \in [2, n]$.

The transition t is of the form

$$t = q_p \xrightarrow[Q_{col} \cup Q'_{col}]{a} \left(Q_1 \cup Q'_1, \dots, Q_{k-1} \cup Q'_{k-1}, Q'_k, Q_{k+1}, \dots, Q_n\right)$$

with $J(t) = (r, t', T, i + 1)$ where

$$t' = q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \dots, Q_k, \dots, Q_n)$$

and T is a set of transitions of the form:

$$Q_k \xrightarrow[Q'_{col}]{a} (Q'_1, \dots, Q'_k)$$

The run w' is equal to $rew_{\{t'\}}(o(rew_T(\boldsymbol{w})))$. Let $w = u :_k v$. The run w' is then equal to $rew_{\{t'\}}(u) :_k rew_T(u) :_k v$.

Let us first show that w' is $\{q_{p'}\}$-valid. For this it is enough to show that:

- for all $k' \in [k + 1, n]$, $top_{k'+1}(pop_{k'}(w')) = top_{k'+1}(pop_{k'}(w))$ is $Q_{k'}$-valid. This immediately follows from the fact that w is q-accepting with head transition t.

- $top_k(pop_k(w')) = rew_T(u) :_k v = rew_T(w)$ is Q_k-valid. As T has the form $Q_k \xrightarrow[Q'_{col}]{a} (Q'_1, \dots, Q'_k)$, it is enough for us to show that for all $k' \in [1, k]$, $top_{k'+1}(pop_{k'}(rew_T(w))) = top_{k'+1}(pop_{k'}(w))$ is $Q_{k'}$-valid. This immediately follows from the fact that w is q_p-accepting with head transition t.

- for all $k' \in [1, k - 1]$, $top_{k'+1}(pop_{k'}(w')) = top_{k'+1}(pop_{k'}(w))$ is $Q_{k'}$-valid.This immediately follows from the fact that w is q-accepting with head transition t.

We now show that w' is link-valid. We only need to check the validity for the substack $rew_T(u) :_k v$ and the substacks of the form $u' :_k rew_T(u) :_k v$ where u' is a substack of $rew_{\{t'\}}(u)$. Let us first consider the stack $rew_T(u) :_k v$ and let h be a transition in T of the form

$$q_h \xrightarrow[Q^h_{col}]{a} \left(Q^h_1, \dots, Q^h_n\right)$$

We have that Q^h_{col} is a subset of Q'_{col}. Let k' be the order of the link on top of $rew_T(u) :_k v$. As w is link-valid, we now that $top_{k'+1}(collapse_{k'}(w)) = top_{k+1}(collapse_{k'}(rew_T(u) :_k v))$ is $Q_{col} \cup Q'_{col}$-valid hence it is also Q^h_{col}-valid. We now move on to the case of $x = rew_{\{t'\}}(u) :_k rew_T(u) :_k v$. Let k' be the order of the link on top of x. We have that $top_{k'+1}(collapse_{k'}(x)) = top_{k'+1}(collapse_{k'}(w))$. By link-validity of w, it is the case that $top_{k'+1}(collapse_{k'}(w))$ is $Q_{col} \cup Q'_{col}$-valid and in particular Q_{col}-valid.

Finally let u' be a strict substack of $rew_{\{t'\}}(u)$. Let k' be the order of the link appearing on top of $x = u' :_k rew_T(u) :_k v$ and let h be a transition attached to the top of x of the form:

$$q_h \xrightarrow[Q^h_{col}]{a} \left(Q^h_1, \dots, Q^h_n\right)$$

We have that $top_{k'+1}(collapse_{k'}(x)) = top_{k'+1}(collapse_{k'}(w))$. By link-validity of w, it is the case that $top_{k'+1}(collapse_{k'}(w))$ is Q^h_{col}-valid.

Algorithm 8 Computing $Pre_{\mathcal{C}}^*(A_0)$

Let $\Delta_{done} = \emptyset$, $\Delta_{new} = \bigcup_{n \geq k > 1} \Delta_k$, $\mathcal{U}_{src}[k] = \emptyset$, $\mathcal{U}_{targ}[k] = \{(\emptyset, \emptyset, \emptyset, \emptyset)\}$ for
each $n \geq k > 1$ and $\mathcal{U}_{targ}[1] = \{ (\emptyset, \emptyset, a, \emptyset, \emptyset) \mid a \in \Sigma \}$.
for $r := (p, a, pop_n, p') \in \mathcal{R}$ **do**

    add_to_worklist$\left(q_p \xrightarrow[\emptyset]{a} (\emptyset, \ldots, \emptyset, \{q_{p'}\}) , r \right)$

for $r := (p, a, collapse_n, p') \in \mathcal{R}$ **do**

    add_to_worklist$\left(q_p \xrightarrow[\{q_{p'}\}]{a} (\emptyset, \ldots, \emptyset) , r \right)$

while $\exists t \in \Delta_{new}$ **do**

    update_rules(t); update_trip_wires(t); move t from Δ_{new} to Δ_{done}

It now remains to show that w' is trimmed. The only interesting case is that of the substack $rew_T(u) :_k v$ which is reach by a pop_k operation. Any transition $h \in T$, is of the form

$$q_h \xrightarrow[Q_{col}^h]{a} \left(Q_1^h, \ldots, Q_n^h \right)$$

with for all $k' \in [1, k]$, $Q_{k'}^h \subseteq Q_{k'}'$. Hence it is enough for us to show that for all $k' \in [1, k]$, $top_{k'+1}(pop_{k'}(rew_T(u) :_k v)) = top_{k'+1}(pop_{k'}(w))$ is Q_k'-valid. This immediately follows from the fact that w is q-accepting with head transition t.

It only remains to show that $r \hookrightarrow_n r'$. First remark that $u \hookrightarrow_{k-1} rew_{\{t'\}}(u)$ and $u \hookrightarrow_{k-1} rew_T(u)$ as in both cases t is replaced by one or several transition with a smaller timestamp (cf. second property of Lemma A.1). By definition of $\hookrightarrow_k$, we have $top_{k+1}(w) \hookrightarrow_k top_{k+1}(w')$. The first property of Lemma A.1 then implies that $w \hookrightarrow_k w'$.

If $o = push_b^k$ **for some** $b \in \Sigma$ **and** $k \in [2, n]$. This case is similar to the $push_k$ case. $\quad\square$

B. Efficient Fixed Point Computation

B.1 Omitted Algorithms

We present the full definitions of the subroutines used in the efficient fixed point computation described in Section 7. Let the function **extract_short_forms** obtain from a long-form transition its (unique) corresponding set of (*short-form*) transitions.

B.2 Correctness

We prove Proposition 7.1 that states the fast algorithm is correct. The proposition is proved in two parts in the following sub-sections. In particular in Lemma B.8 and Lemma B.3.

In the sequel, we fix the following notation. Let $(A_i)_{i \geq 0}$ be the sequence of automata constructed by the naive fixed point algorithm. Then, let $\left(\Delta_{done}^j \right)_{j \geq 0}$ be the sequence of sets of transitions such that Δ_{done}^j is Δ_{done} after j iterations of the main loop of Algorithm 6. Similarly, define $\mathcal{U}_{src}^j[k]$ and $\mathcal{U}_{targ}^j[k]$.

B.3 Soundness

We prove that the algorithm is sound. First, we show two preliminary lemmas about the data-structures maintained by the algorithm.

Algorithm 9 update_rules(t)

Require: A transition t to be processed against Δ_{done}

  if t is an order-k transition for $2 \leq k \leq n$ of the form $q_{p'Q_n\cdots Q_{k+1}} \longrightarrow Q_k$ **then**

    for $p \in \mathcal{P}$ and $a \in \Sigma$ such that $r := (p, a, pop_{k-1}, p') \in \mathcal{R}$ **do**

      add_to_worklist$\left(q_p \xrightarrow[\mathcal{Q}_{col}]{a} (\emptyset, \ldots, \emptyset, \{q_{p'\mathcal{Q}_n\cdots\mathcal{Q}_k}\}, \mathcal{Q}_k, \ldots, \mathcal{Q}_n), r \right)$

    for $p \in \mathcal{P}$ and $a \in \Sigma$ such that $r := (p, a, collapse_{k-1}, p') \in \mathcal{R}$ **do**

      add_to_worklist$\left(q_p \xrightarrow[\{q_{p'Q_n\cdots Q_k}\}]{a} (\emptyset, \ldots, \emptyset, Q_k, \ldots, Q_n), r \right)$

    for $p \in \mathcal{P}$ and $a \in \Sigma$ such that $r := (p, a, push_k, p') \in \mathcal{R}$ **do**

      create_trip_wire$\left(q_{p,Q_n,\ldots,Q_{k+1}}, q_{p',Q_n,\ldots,Q_{k+1},Q_k}, a, Q_k, (r, t) \right)$

  else if t is an order-1 transition of the form $q_{p'Q_n\cdots Q_2} \xrightarrow[Q_{col}]{b} Q_1$ **then**

    for $p \in \mathcal{P}$ and $a \in \Sigma$ such that $r := (p, a, rew_b, p') \in \mathcal{R}$ **do**

      add_to_worklist$\left(q_p \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n), (r, t) \right)$

    for $p \in \mathcal{P}$ and $a \in \Sigma$ such that $r := \left(p, a, push_b^k, p' \right) \in \mathcal{R}$ **do**

      create_trip_wire$\left(q_{p,Q_n,\ldots,Q_{k+1},Q_k\cup Q_{col},Q_{k-1},\ldots,Q_2}, \perp, a, Q_1, (r, t) \right)$

Algorithm 10 update_trip_wires$\left(t = q_p \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n) \right)$
---

  for $t_k = q_k \longrightarrow Q_k \in$ **extract_short_forms**(t) **do**

    for $targ \in \mathcal{U}_{targ}[k]$ with $targ = (\_, Q_{k'}^C, \_, \_)$ or $(\_, Q_{k'}^C, a, \_, \_)$ and $q_k \in Q_{k'}^C$ **do**

      proc_targ_against_tran$(targ, t_k)$

Algorithm 11 create_trip_wire$(q_k, q_{k-1}, a, Q_k, (r, t))$
---

  if $(q_k, q_{k-1}, a, Q_k) \notin \mathcal{U}_{src}[k]$ **then**

    Add $src := (q_k, q_{k-1}, a, Q_k)$ to $\mathcal{U}_{src}[k]$, set $J(src) := (r, t)$

    Let $targ := (Q_k, Q_k, \emptyset, \emptyset)$ if $k > 1$ or $(Q_k, Q_k, a, \emptyset, \emptyset)$ if $k = 1$

    if $targ \in \mathcal{U}_{targ}[k]$ **then**

      for each *complete* target $targ$ matching src **do**

        proc_source_complete_targ$(src, targ)$

    else

      add_target$(targ, k)$; set $J(targ) := \emptyset$

Algorithm 12 proc_targ_against_tran$(targ, t)$

Suppose $\begin{cases} t = q_k \longrightarrow Q_k'' \text{ and } targ = (Q_k, Q_k^C, Q_{lbl}, Q_k') & \text{if } k \geq 2 \\ t = q_1 \xrightarrow[Q_{col}]{a} Q_1'' \text{ and } targ = (Q_1, Q_1^C, a, Q_{lbl}, Q_1') & \text{if } k = 1 \end{cases}$

Let $targ' := \begin{cases} \left(Q_k, Q_k^C \setminus \{q_k\}, Q_{lbl} \cup \left\{q_{kQ_k''}\right\}, Q_k' \cup Q_k''\right) & \text{if } k \geq 2 \\ \left(Q_1, Q_1^C \setminus \{q_1\}, a, Q_{lbl} \cup Q_{col}, Q_1' \cup Q_1''\right) & \text{if } k = 1 \end{cases}$

if $q_k \in Q_k^C$ and $targ' \notin \mathcal{U}_{targ}[k]$ **then**

    add_target$(targ', k)$; if $k = 1$, set $J(targ') := J(targ) \cup \{t\}$

    if $Q_k^C \setminus \{q_k\} = \emptyset$ **then**

        for each source $src \in \mathcal{U}_{src}[k]$ of form $(\_, \_, \_, Q_k)$ **do**

            proc_source_complete_targ$(src, targ')$

Algorithm 13 proc_source_complete_targ$(src, comp\_targ)$

Require: An order-k source of the form $src = (q_k, q_{k-1}, a, Q_k)$ and an order-k complete target of the form $comp\_targ = (Q_k, \emptyset, Q_{lbl}, Q_k')$ when $k \geq 2$ and $(Q_1, \emptyset, a, Q_{lbl}, Q_1')$ when $k = 1$

if $k \geq 2$ **then**

    Let $S := \{ q_{k-1} \mid q_{k-1} \neq \bot \}$

    create_trip_wire$\left(q_{kQ_k'}, \bot, a, Q_{lbl} \cup S, J(src)\right)$

else if $k = 1$ **then**

    Suppose $q_1 = q_{p, Q_n, \ldots, Q_2}$ and $J(src) = (r, t)$ and $J(comp\_targ) = T$

    add_to_worklist$\left(q_p \xrightarrow[Q_{lbl}]{a} (Q_1', Q_2, \ldots, Q_n), (r, t, T)\right)$

Algorithm 14 add_to_worklist$(t, justif)$

Require: A long form transition t and justification $justif$.

    for $u \in$ **extract_short_forms**(t) such that $u \notin \Delta_{done} \cup \Delta_{new}$ **do**

        Add u to Δ_{new} and set $J(u) := (justif, |\Delta_{new} \cup \Delta_{done}|)$ if u is order-1

Algorithm 15 add_target$(targ, k)$

---

if $targ \notin \mathcal{U}_{targ}[k]$ **then**

    Add $targ$ to $\mathcal{U}_{targ}[k]$

    for $t' \in \Delta_{done}$ **do**

        proc_targ_against_tran$(targ, t')$

Lemma B.1. *For all $j \geq 0$ and $n \geq k > 1$, if $\left(Q_k, Q_k \setminus Q_k^T, Q_{k-1}^T, Q_k^{T'}\right) \in \mathcal{U}_{targ}^j[k]$, then we have $T \subseteq \Delta_{done}^j$ that witnesses $Q_k^T \xrightarrow{Q_{k-1}^T} Q_k^{T'}$.*

Proof. We proceed by induction over j and the order in which targets are created. In the base case we only have $(\emptyset, \emptyset, \emptyset, \emptyset) \in \mathcal{U}_{targ}^0[k]$. Setting $T = \emptyset$ witnesses $\emptyset \xrightarrow{\emptyset} \emptyset$.

In the inductive case, consider the location of the call to add_target. This is either in create_trip_wire or proc_targ_against_tran. When the call location is create_trip_wire, we have a target of the form $(Q_k, Q_k, \emptyset, \emptyset)$, hence $Q_k^T = \emptyset$ and we trivially have $T = \emptyset \subset \Delta_{done}^j$ witnessing $\emptyset \xrightarrow{\emptyset} \emptyset$.

Otherwise the call is from proc_targ_against_tran against a transition $t = q_k \longrightarrow Q_k''$ and a target $targ = \left(Q_k, Q_k \setminus Q_k^T, Q_{k-1}^T, Q_k^{T'}\right)$ already in $\mathcal{U}_{targ}$. Hence, by induction, we know that there is some $T \subseteq \Delta_{done}^j[k]$ witnessing $Q_k^T \xrightarrow{Q_{k-1}^T} Q_k^{T'}$. The transition t is either already in Δ_{done}^j or will be moved there at the end of the jth iteration. Combining t with T we have $T \cup \{t\} \subseteq \Delta_{done}^j$ witnessing $Q_k^T \cup \{q_k\} \xrightarrow{Q_{k-1}^T \cup \left\{q_k Q_k''\right\}} Q_k^{T'} \cup Q_k''$. Since the new target added is $\left(Q_k, Q_k \setminus \left(Q_k^T \cup \{q_k\}\right), Q_{col}^T \cup \left\{q_k Q_k''\right\}, Q_k^{T'} \cup Q_k''\right)$ we are done. $\square$

Lemma B.2. *For all $j \geq 0$, if $\left(Q_1, Q_1 \setminus Q_1^T, a, Q_{col}^T, Q_1^{T'}\right) \in \mathcal{U}_{targ}^j[1]$, then we have $T \subseteq \Delta_{done}^j$ that witnesses $Q_1^T \xrightarrow[Q_{col}^T]{a} Q_1^{T'}$.*

Proof. The proof is essentially the same as the order-k case above. We proceed by induction over j and the order in which targets are created. In the base case we only have $(\emptyset, \emptyset, a, \emptyset, \emptyset) \in \mathcal{U}_{targ}^0[1]$. Setting $T = \emptyset$ witnesses $\emptyset \xrightarrow[\emptyset]{a} \emptyset$.

In the inductive case, consider the location of the call to add_target. This is either in create_trip_wire or proc_targ_against_tran. When the call location is create_trip_wire, we have a target of the form $(Q_1, Q_1, a, \emptyset, \emptyset)$, hence $Q_1^T = \emptyset$ and we trivially have $T = \emptyset \subset \Delta_{done}^j$ witnessing $\emptyset \xrightarrow[\emptyset]{a} \emptyset$.

Otherwise the call is from proc_targ_against_tran against a transition $t = q_1 \xrightarrow[Q_{col}]{a} Q_1''$ and a target $targ = \left(Q_1, Q_1 \setminus Q_1^T, Q_{col}^T, Q_1^{T'}\right)$ already in $\mathcal{U}_{targ}$. Hence, by induction, we know that there is some $T \subseteq \Delta_{done}^j[1]$ witnessing $Q_1^T \xrightarrow{Q_{col}^T} Q_1^{T'}$. The transition t is either already in Δ_{done}^j or will be moved there at the end of the jth iteration. Combining t with T we have $T \cup \{t\} \subseteq \Delta_{done}^j$ witnessing $Q_1^T \cup \{q_1\} \xrightarrow[Q_{col}^T \cup Q_{col}]{a} Q_1^{T'} \cup Q_1''$. Since the new target is $\left(Q_1, Q_1 \setminus \left(Q_1^T \cup \{q_1\}\right), a, Q_{col} \cup Q_{col}^T, Q_1^{T'} \cup Q_1''\right)$ we are done. $\square$

We are now ready to prove the algorithm sound.

Lemma B.3. *Given a CPDS $\mathcal{C}$ and stack automaton A_0, let A be the result of Algorithm 6. We have $\mathcal{L}(A) \subseteq Pre_{\mathcal{C}}^*(A_0)$.*

Proof. We proceed by induction over j and show every transition appearing in Δ_{done}^j appears in A_i for some i. This implies the lemma.

When $j = 0$ the property is immediate, since the only transitions added are already in A_0, or added to A_1 during the first processing of the pop_n and $collapse_n$ rules.

In the inductive step, we consider some t first appearing in Δ_{new}^j (and thus, eventually in $\Delta_{done}^{j'}$ for some j'). There are several cases depending on how t was added to Δ_{new} (i.e. from where

add_to_worklist was called). We consider the simple cases first. In all the following cases, t was added during update_rules against a transition t' appearing in Δ_{new}^{j-1}.

- If $t' = q_{p'Q_n\cdots Q_{k+1}} \longrightarrow Q_k$ and t was added as part of

$$t_1 = q_p \xrightarrow[\mathcal{Q}_{col}]{a} \left(\emptyset, \ldots, \emptyset, \left\{q_{p'\mathcal{Q}_n\cdots\mathcal{Q}_k}\right\}, \mathcal{Q}_k, \ldots, \mathcal{Q}_n\right)$$

  during the processing of t' against a pop_{k-1} rule. By induction t' appears in A_i for some i, and hence t_1 (which includes t) is present in A_{i+1}.

- If $t' = q_{p'Q_n\cdots Q_{k+1}} \longrightarrow Q_k$ and t was added as part of

$$t_1 = q_p \xrightarrow[\left\{q_{p'Q_n\cdots Q_k}\right\}]{a} \left(\emptyset, \ldots, \emptyset, Q_k, \ldots, Q_n\right)$$

  during the processing of t' against a $collapse_{k-1}$ rule. By induction t' appears in A_i for some i, and hence t_1 (which includes t) is present in A_{i+1}.

- If $t' = q_{p'Q_n\cdots Q_2} \xrightarrow[Q_{col}]{b} Q_1$ and t was added as part of

$$t_1 = q_p \xrightarrow[Q_{col}]{a} \left(Q_1, \ldots, Q_n\right)$$

  during the processing of t' against a rew_b rule. By induction t' appears in A_i for some i, and hence t_1 (which includes t) is present in A_{i+1}.

In the final case, add_to_worklist is called during proc_source_complete_targ. There are two cases depending on the provenance of the source. In the first case, the source was added by a call to create_trip_wire from update_rules while processing a $push_b^k$ rule against $t' = q_{p'Q_n\cdots Q_2} \xrightarrow[Q_{col}]{b} Q_1$. Therefore, t was added as part of

$$q_p \xrightarrow[Q'_{col}]{a} \left(Q'_1, Q_2, \ldots, Q_{k-1}, Q_k \cup Q_{col}, Q_{k+1}, Q_n\right)$$

from a source $(q_1, \bot, a, Q_1) \in \mathcal{U}_{src}^j[1]$ with

$$q_1 = q_{p, Q_n, \ldots, Q_{k+1}, Q_k \cup Q_{col}, Q_{k-1}, \ldots, Q_2} \,.$$

By induction, from t' we know that

$$q_{p'} \xrightarrow[Q_{col}]{b} \left(Q_1, \ldots, Q_n\right)$$

appears in A_i for some i. Now, consider the target $(Q_1, \emptyset, a, Q'_{col}, Q'_1) \in \mathcal{U}_{targ}^j[1]$ that was combined with the source to add the new transition. By Lemma B.2 we have $Q_1 \xrightarrow[Q'_{col}]{a} Q'_1$ in Δ_{done}^j and hence (since all transitions in Δ_{done}^j passed through Δ_{new}) by induction we have $Q_1 \xrightarrow[Q'_{col}]{a} Q'_1$ in $A_{i'}$ for some i'. Hence, in $A_{\max(i,i')+1}$ we have t as required.

In the second case we have a source $(q_1, \bot, a, Q_1^s) \in \mathcal{U}_{src}^j[1]$ and a complete target of the form $\left(Q_1^s, \emptyset, a, Q_{col}^t, Q_1^{t'}\right) \in \mathcal{U}_{targ}^j[1]$ and the source derived from a call to create_trip_wire in proc_source_complete_targ. Note, by Lemma B.2 we have $Q_1^s \xrightarrow[Q_{col}^t]{a} Q_1^{t'}$ in Δ_{done}^j. The call to create_trip_wire implies we have a source $(q_2, q_1', a, Q_2^s) \in \mathcal{U}_{src}^j[2]$ and complete target of the form

$\left(Q_2^s, \emptyset, Q_1^t, Q_2^{t'}\right) \in \mathcal{U}_{targ}^j[2]$, with $Q_1^s = Q_1^t \cup S_1$ where $S_1 = \{ q_1' \mid q_1' \neq \perp \}$ and $q_1 = q_{2Q_2^{t'}}$. The proof will now iterate from $k = 2$ upwards until a source is discovered that was added during a call to create_trip_wire from update_rules while processing some $push_{k'}$ rule. Note that sources not added by $push_b^k$ rules can only be added in this way and, for all $k < k'$, the second component of the source (q_{k-1}') will be $\perp$.

Hence, inductively, we have a source $src = \left(q_k, q_{k-1}', a, Q_k^s\right) \in \mathcal{U}_{src}^j[k]$ and complete target $\left(Q_k^s, \emptyset, Q_{k-1}^t, Q_k^{t'}\right) \in \mathcal{U}_{targ}^j[k]$ with $Q_{k-1}^s = Q_{k-1}^t \cup S_{k-1}$ where $S_{k-1} = \{ q_{k-1}' \mid q_{k-1}' \neq \perp \}$ and $q_{k-1} = q_{kQ_k^{t'}}$. Furthermore, by Lemma B.1 we have $Q_k^s \xrightarrow{Q_{k-1}^t} Q_k^{t'}$ in Δ_{done}^j.

In the first case, suppose src was added due to a call to create_trip_wire in proc_source_complete_targ. The call to create_trip_wire implies we have a source $\left(q_{k+1}, q_k', a, Q_{k+1}^s\right) \in \mathcal{U}_{src}^j[k+1]$ and complete target $\left(Q_{k+1}^s, \emptyset, Q_k^t, Q_{k+1}^{t'}\right) \in \mathcal{U}_{targ}^j[k+1]$, with $Q_k^s = Q_k^t \cup S_k$ where $S_k = \{ q_k' \mid q_k' \neq \perp \}$ and $q_k = q_{k+1Q_{k+1}^{t'}}$.

For the final case, suppose that src was added due to a call to create_trip_wire in update_rules from a $push_k$ rule. Then we were processing a new transition $q_{p'Q_n \cdots Q_{k+1}} \longrightarrow Q_k$, and we have $q_k = q_{p,Q_n,\ldots,Q_{k+1}}$ and $q_{k-1}' = q_{p',Q_n,\ldots,Q_k}$ and $Q_k^s = Q_k$. From the induction and since $q_{k'}' = \perp$ for all $k' < k$, we have $Q_{k-1}^t \cup \{q_{k-1}'\} \xrightarrow[Q_{col}^t]{a} \left(Q_1^{t'}, \ldots, Q_{k-1}^{t'}\right)$ in Δ_{done}^j which can be split into

$Q_{k-1}^t \xrightarrow[Q_{col}']{a} \left(Q_1', \ldots, Q_{k-1}'\right)$ and $q_{k-1}' \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_{k-1})$. Thus, because $q_{k-1}' = q_{p',Q_n,\ldots,Q_k}$ and $Q_k^s = Q_k$ and letting $Q_k' = Q_k^{t'}$, we have

$$q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n) \quad \text{and} \quad Q_k \xrightarrow[Q_{col}']{a} (Q_1', \ldots, Q_k')$$

in Δ_{done}^j and thus by induction in A_i for some i. Since we have

$$q_1 = q_{kQ_k', Q_{k-1} \cup Q_{k-1}', \ldots, Q_1 \cup Q_1'} = q_{p,Q_n,\ldots,Q_{k+1},Q_k',Q_{k-1} \cup Q_{k-1}',\ldots,Q_2 \cup Q_2'}$$

we added t as part of a transition

$$q_p \xrightarrow[Q_{col} \cup Q_{col}']{a} \left(Q_1 \cup Q_1', \ldots, Q_{k-1} \cup Q_{k-1}', Q_k', Q_{k+1}, \ldots, Q_n\right)$$

which is the transition added by the naive saturation algorithm from the $push_k$ rule and the transitions in A_i. Hence, we satisfy the lemma. $\square$

B.4 Completeness

We prove that the algorithm is complete. For this we need some preliminary lemmas stating properties of the data-structures maintained by the algorithm.

Lemma B.4. *For all $k \geq 2$ and $j \geq 0$, all $T \subseteq \Delta_{done}^j$ witnessing $Q_k^T \xrightarrow{Q_{k-1}^T} Q_k'$, and all $(q_k, q_{k-1}, a, Q_k) \in \mathcal{U}_{src}^j[k]$ such that $Q_k^T \subseteq Q_k$, we have the target $(Q_k, Q_k \setminus Q_k^T, Q_{k-1}^T, Q_k')$ in $\mathcal{U}_{targ}^j[k]$.*

Proof. Let j_1 be the iteration of Algorithm 6 where (q_k, q_{k-1}, a, Q_k) was first added to $\mathcal{U}_{src}^{j_1}[k]$. We perform an induction over j_1. In the base case the lemma is trivially true. In the inductive case, the only position where a source may be added is in the create_trip_wire procedure. After adding the source, the induction hypothesis needs to be re-established. There are two cases.

Let $targ = (Q_k, Q_k, \emptyset, \emptyset)$. If $targ$ is already in $\mathcal{U}_{targ}^{j_1}$ then we observe that a target of the form $(Q_k, Q_k, \ldots)$ is only created in create_trip_wire (targets are also created in proc_targ_against_tran, but

these targets are obtained by removing a state from the second component of an existing target, hence the two first components cannot be equal). This implies the existence of a source $(\_,\_,\_,Q_k) \in \mathcal{U}_{src}^{j'}[k]$ for some $j' < j_1$. This implies the result by induction since neither T nor the desired target depend any but the final component of the source.

If $targ$ is not in $\mathcal{U}_{targ}^{j_1}[k]$, then we add it. Next, split $T = T_1 \cup T_2$ such that T_1 contains all $t \in T$ appearing in $\Delta_{done}^{j_1-1}$. The balance is contained in T_2. The algorithm proceeds to call proc_targ_against_tran on $targ$ and all $t \in \Delta_{done}^{j_1}$. In particular, this includes all $t \in T_1$.

We aim to prove that, after the execution of this loop, we have $(Q_k, Q_k \setminus Q_k^1, Q_{k-1}^1, Q_k^{1'}) \in \mathcal{U}_{targ}^{j_1}[k]$ when T_1 witnesses $Q_k^1 \xrightarrow{Q_{k-1}^1} Q_k^{1'}$.

Let $t_1, \ldots, t_\ell$ be a linearisation of T_1 in the order they appear in iterations over Δ_{done} (we assume a fixed order here for convenience, though the proof can generalise if the order changes between iterations). Additionally, let $T_z = \{t_1, \ldots, t_z\}$ witness $Q_k^{t_z} \xrightarrow{Q_{k-1}^{t_z}} Q_k^{t_z'}$. We show after T_z has been processed, we have $(Q_k, Q_k \setminus Q_k^{t_z}, Q_{k-1}^{t_z}, Q_k^{t_z'}) \in \mathcal{U}_{targ}^{j_1}[k]$. This gives us the property once $z = \ell$. In the base case $z = 0$ and we are done. Otherwise, we know $targ_z = (Q_k, Q_k \setminus Q_k^{t_z}, Q_{k-1}^{t_z}, Q_k^{t_z'}) \in \mathcal{U}_{targ}^{j_1}[k]$ and prove the case for $(z+1)$. Consider the call to add_target that added $targ_z$. Now take the iteration against Δ_{done} that processes t_{z+1}. This results in the addition of $(Q_k, Q_k \setminus Q_k^{t_{z+1}}, Q_{k-1}^{t_{z+1}}, Q_k^{t_{z+1}'})$ as required.

Hence, we have $(Q_k, Q_k \setminus Q_k^1, Q_{k-1}^1, Q_k^{1'}) \in \mathcal{U}_{targ}^{j_1}[k]$. Now, let $t_1, \ldots, t_\ell$ be a linearisation of T_2 in the order they are added to Δ_{done}. Additionally, we write $Q_k^{t_z} \xrightarrow{Q_{k-1}^{t_z}} Q_k^{t_z'}$ for the state-sets and transitions witnessed by $T_1 \cup \{t_1, \ldots, t_z\}$.

We show after t_z has been added to Δ_{done} on the j'th iteration, we have $(Q_k, Q_k \setminus Q_k^{t_z}, Q_{k-1}^{t_z}, Q_k^{t_z'}) \in \mathcal{U}_{targ}^{j'}[k]$ for some j'. In the base case $z = 0$ and we are done by the argument above. Otherwise, we know $targ_z = (Q_k, Q_k \setminus Q_k^{t_z}, Q_{k-1}^{t_z}, Q_k^{t_z'}) \in \mathcal{U}_{targ}^{j'}[k]$ and prove the case for $(z+1)$. Consider the call to update_trip_wires with t_{z+1}. This results in the addition of the target $(Q_k, Q_k \setminus Q_k^{t_{z+1}}, Q_{k-1}^{t_{z+1}}, Q_k^{t_{z+1}'})$ via the call to proc_targ_against_tran. When $z = \ell$, we have the lemma as required. $\square$

Lemma B.5. *For all $j \geq 0$, all $T \subseteq \Delta_{done}^j$ witnessing $Q_1^T \xrightarrow[Q_{col}^T]{a} Q_1'$, and all $(q_1, \bot, a, Q_1) \in \mathcal{U}_{src}^j[1]$ such that $Q_k^T \subseteq Q_1$, we have $(Q_1, Q_1 \setminus Q_1^T, a, Q_{col}^T, Q_1')$ in $\mathcal{U}_{targ}^j[1]$.*

Proof. The proof is essentially the same as the proof when $k \geq 2$. Let j_1 be the iteration of Algorithm 6 where (q_1, q_{k-1}, a, Q_1) was first added to $\mathcal{U}_{src}^{j_1}[1]$. We perform an induction over j_1. In the base case the lemma is trivially true. In the inductive case, the only position where a source may be added is in the create_trip_wire procedure. After adding the source, the induction hypothesis needs to be re-established. There are two cases.

Let $targ = (Q_1, Q_1, a, \emptyset, \emptyset)$. If $targ$ is already in $\mathcal{U}_{targ}^{j_1}$ then we observe that a target of the form $(Q_1, Q_1, \ldots)$ is only created in create_trip_wire. This implies the existence of a source $(\_, \_, \_, Q_1) \in \mathcal{U}_{src}^{j'}[1]$ for some $j' < j_1$. This implies the result by induction since neither T nor the desired target depend any but the final component of the source.

If $targ$ is not in $\mathcal{U}_{targ}^{j_1}[1]$, then we add it. Next, split $T = T_1 \cup T_2$ such that T_1 contains all $t \in T$ appearing in $\Delta_{done}^{j_1-1}$. The balance is contained in T_2. The algorithm proceeds to call proc_targ_against_tran on $targ$ and all $t \in \Delta_{done}^{j_1}$. In particular, this includes all $t \in T_1$.

We aim to prove that, after the execution of this loop, we have $(Q_1, Q_1 \setminus Q_1^1, a, Q_{col}^1, Q_1^{1'}) \in \mathcal{U}_{targ}^{j_1}[1]$ when T_1 witnesses $Q_1^1 \xrightarrow[Q_{col}^1]{a} Q_1^{1'}$.

Let $t_1, \ldots, t_\ell$ be a linearisation of T_1 in the order they appear in iterations over Δ_{done}. Additionally, let $T_z = \{t_1, \ldots, t_z\}$ witness $Q_1^{t_z} \xrightarrow[Q_{col}^{t_z}]{a} Q_1^{t'_z}$. We show after T_z has been processed, we have

$(Q_1, Q_1, \backslash Q_1^{t_z}, a, Q_{col}^{t_z}, Q_1^{t'_z}) \in \mathcal{U}_{targ}^{j_1}[1]$. This gives us the property once $z = \ell$. In the base case $z = 0$ and we are done. Otherwise, we know $targ_z = (Q_1, Q_1 \backslash Q_1^{t_z}, aQ_{col}^{t_z}, Q_1^{t'_z}) \in \mathcal{U}_{targ}^{j_1}[1]$ and prove the case for $(z + 1)$. Consider the call to add_target that added $targ_z$. Now take the iteration against Δ_{done} that processes t_{z+1}. This results in the addition of $(Q_1, Q_1 \backslash Q_1^{t_{z+1}}, a, Q_{col}^{t_{z+1}}, Q_1^{t'_{z+1}})$ as required.

Hence, we have $(Q_1, Q_1 \backslash Q_1^1, a, Q_{col}^1, Q_1^{1'}) \in \mathcal{U}_{targ}^{j_1}[1]$. Now, let $t_1, \ldots, t_\ell$ be a linearisation of T_2 in the order they are added to Δ_{done}. Additionally, we write $Q_1^{t_z} \xrightarrow[Q_{col}^{t_z}]{a} Q_1^{t'_z}$ for the state-sets and

transitions witnessed by $T_1 \cup \{t_1, \ldots, t_z\}$.

We show after t_z is added to Δ_{done} on the j'th iteration, we have $(Q_1, Q_1 \backslash Q_1^{t_z}, a, Q_{col}^{t_z}, Q_1^{t'_z}) \in \mathcal{U}_{targ}^{j'}[1]$ for some j'. In the base case $z = 0$ and we are done by the argument above. Otherwise, we know $targ_z = (Q_1, Q_1 \backslash Q_1^{t_z}, a, Q_{col}^{t_z}, Q_1^{t'_z}) \in \mathcal{U}_{targ}^{j'}[1]$ and prove the case for $(z + 1)$. Consider the call to update_trip_wires with t_{z+1}. This results in the addition of the target $(Q_1, Q_1 \backslash Q_1^{t_{z+1}}, aQ_{col}^{t_{z+1}}, Q_1^{t'_{z+1}})$ via the call to proc_targ_against_tran. When $z = \ell$, we have the lemma as required. $\square$

Lemma B.6. *For all $k > 1$ and $j \geq 0$, if we have $(q_k, q_{k-1}, a, Q_k) \in \mathcal{U}_{src}^j[k]$ and $(Q_k, \emptyset, Q_{k-1}, Q'_k) \in \mathcal{U}_{targ}^j[k]$, then it is the case that there exists $j' \geq 0$ such that $\left(q_{kQ'_k}, \bot, a, Q_{k-1} \cup S\right) \in \mathcal{U}_{src}^{j'}[k-1]$ where $S = \{ q_{k-1} \mid q_{k-1} \neq \bot \}$.*

Proof. Let j_1 be the smallest such that $(q_k, q_{k-1}, a, Q_k) \in \mathcal{U}_{src}^{j_1}[k]$ and j_2 be the smallest such that $(Q_k, \emptyset, Q_{k-1}, Q'_k) \in \mathcal{U}_{targ}^{j_2}[k]$.

In the case $j_1 \leq j_2$, we consider the j_2th iteration of Algorithm 6 at the moment where the target is added to $\mathcal{U}_{targ}^{j_2}[k]$. This has to be a result of the call to add_target during Algorithm 12. The only other place add_target may be called is during Algorithm 11; however, this implies the target is of the form $(Q_k, Q_k, \emptyset, \emptyset)$ and hence, for the target to be complete, it must be $(\emptyset, \emptyset, \emptyset, \emptyset)$ and hence $j_2 = 0$, and since $j_1 > 0$ (since there are initially no sources) we have a contradiction. Hence, the target is added during Algorithm 12 and the procedure goes on to call proc_source_complete_targ against each matching source in $\mathcal{U}_{src}^{j_2}[k]$, including (q_k, q_{k-1}, a, Q_k). This results in the addition of $\left(q_{kQ'_k}, \bot, a, Q_{k-1} \cup S\right)$ to $\mathcal{U}_{src}^{j_2}[k-1]$, if it is not there already, satisfying the lemma.

In the case $j_1 > j_2$, we consider the j_1th iteration of Algorithm 6 at the moment where the source is added. This is necessarily in the create_trip_wire procedure. Since $(Q_k, \emptyset, Q_{k-1}, Q'_k) \in \mathcal{U}_{targ}^{j_1}[k]$ and since this target must have been obtained from a target of the form $(Q_k, Q_k, \emptyset, \emptyset)$, we know $(Q_k, Q_k, \emptyset, \emptyset) \in \mathcal{U}_{targ}^{j_1}[k]$ and thus the procedure calls proc_source_complete_targ against each complete target including $(Q_k, \emptyset, Q_{k-1}, Q'_k)$. This results in the addition of $\left(q_{kQ'_k}, \bot, a, Q_{k-1} \cup S\right)$ to $\mathcal{U}_{src}^{j_2}[k-1]$, if it is not there already, satisfying the lemma. $\square$

Lemma B.7. *For all $j \geq 0$, if $(q_1, \bot, a, Q_1) \in \mathcal{U}_{src}^j[1]$ and $(Q_1, \emptyset, Q_{col}, Q'_1) \in \mathcal{U}_{targ}^j[1]$, if $q_1 = q_{p, Q_n, \ldots, Q_2}$, then for each t in*

$$\text{extract\_short\_forms}\left(q_p \xrightarrow[Q_{col}]{a} (Q'_1, Q_2, \ldots, Q_n)\right)$$

there exists some $j' \geq 0$ such that $t \in \Delta_{done}^{j'}$.

Proof. As before, the proof of this order-1 case is very similar to the order-k proof.

Let j_1 be the smallest such that $(q_1, \bot, a, Q_1) \in \mathcal{U}^{j_1}_{src}[1]$ and j_2 be the smallest such that $(Q_1, \emptyset, a, Q_{col}, Q'_1) \in \mathcal{U}^{j_2}_{targ}[1]$.

In the case $j_1 \leq j_2$, we consider the j_2th iteration of Algorithm 6 at the moment where the target is added to $\mathcal{U}^{j_2}_{targ}[1]$. This has to be a result of the call to add_target during Algorithm 12. The only other place add_target may be called is during Algorithm 11; however, this implies the target is of the form $(Q_1, Q_1, \emptyset, \emptyset)$ and hence, for the target to be complete, it must be $(\emptyset, \emptyset, \emptyset, \emptyset)$ and hence $j_2 = 0$, and since $j_1 > 0$ (since there are initially no sources) we have a contradiction. Hence, the target is added during Algorithm 12 and the procedure goes on to call proc_source_complete_targ against each matching source in $\mathcal{U}^{j_2}_{src}[1]$, including $(q_1, \bot, a, Q_1)$. This results in the addition of $q_p \xrightarrow[Q_{col}]{a} (Q'_1, Q_2, \ldots, Q_n)$ satisfying the lemma.

In the case $j_1 > j_2$, we consider the j_1th iteration of Algorithm 6 at the moment where the source is added. This is necessarily in the create_trip_wire procedure. Since $(Q_1, \emptyset, a, Q_{col}, Q'_1) \in \mathcal{U}^{j_1}_{targ}[1]$ and since this target must have been obtained from a target of the form $(Q_1, Q_1, a, \emptyset, \emptyset)$, we know $(Q_1, Q_1, a, \emptyset, \emptyset) \in \mathcal{U}^{j_1}_{targ}[1]$ and thus the procedure calls proc_source_complete_targ against each complete target including $(Q_1, \emptyset, a, Q_{col}, Q'_1)$. This results in the addition of the source $q_p \xrightarrow[Q_{col}]{a} (Q'_1, Q_2, \ldots, Q_n)$ satisfying the lemma. $\qquad\square$

We are now ready to prove completeness.

Lemma B.8. *Given a CPDS $\mathcal{C}$ and stack automaton A_0, let A be the result of Algorithm 6. We have $\mathcal{L}(A) \supseteq Pre^*_{\mathcal{C}}(A_0)$.*

Proof. We know (from ICALP [7]) that the fixed point of $(A_i)_{i \geq 0}$ is an automaton recognising $Pre^*_{\mathcal{C}}(A_0)$. We prove, by induction, that for each transition t appearing in A_i for some i, there exists some j such that t appears in Δ^j_{done}.

We first prove the only if direction. In the base case we have all transitions in A_0 in Δ_{new} at the beginning of Algorithm 6. Since the main loop continues until Δ_{new} has been completely transferred to Δ_{done}, the result follows.

Now, let t be an order-k transition appearing for the first time in A_i $(i > 0)$. We perform a case split on the pushdown operation that led to the introduction of the new transition. Let $r = (p, a, o, p')$ be the rule that led to the new transition. We first deal with the simple cases.

- When $o = pop_k$, then there was $q_{p'} \xrightarrow{q_k} (Q_{k+1}, \ldots, Q_n)$ in A_i and we added to A_{i+1}

$$q_p \xrightarrow[\emptyset]{a} (\emptyset, \ldots, \emptyset, \{q_k\}, Q_{k+1}, \ldots, Q_n)$$

  of which t is a transition. By induction we have j such that $t' = q_{p' Q_n \cdots Q_{k+2}} \xrightarrow{q_k} Q_{k+1}$ appears in Δ^j_{done}. Consider the jth iteration of Algorithm 6 when update_rules is called on t'. The $pop_{k'}$ loop immediately adds

$$q_p \xrightarrow[\emptyset]{a} (\emptyset, \ldots, \emptyset, \{q_k\}, Q_{k+1}, \ldots, Q_n)$$

  and hence t to Δ_{new}, giving us some $j' > j$ such that t appears in $\Delta^{j'}_{done}$.

- When $o = collapse_k$, when $k = n$, we added t as part of $q_p \xrightarrow[\{q_{p'}\}]{a} (\emptyset, \ldots, \emptyset)$. In this case we also added t to Δ_{new} as part of the initialisation steps of Algorithm 6. Otherwise, $n > k$ and from a transition $q_{p'} \xrightarrow{q_k} (Q_{k+1}, \ldots, Q_n)$ we added t as part of

$$q_p \xrightarrow[\{q_k\}]{a} (\emptyset, \ldots, \emptyset, Q_{k+1}, \ldots, Q_n)$$

By induction we have j such that $t' = q_{p'Q_n\cdots Q_{k+2}} \xrightarrow{q_k} Q_{k+1}$ appears in Δ_{done}^j. Consider the jth iteration of Algorithm 6 when update_rules is called on t'. The $collapse_k$ loop immediately adds

$$q_p \xrightarrow[\{q_k\}]{a} (\emptyset, \ldots, \emptyset, Q_{k+1}, \ldots, Q_n)$$

and hence t to Δ_{new}, giving us some $j' > j$ such that t appears in $\Delta_{done}^{j'}$.

- when $o = rew_b$ then from a transition $q_{p'} \xrightarrow[Q_{col}]{b} (Q_1, \ldots, Q_n)$ we added t as part of a transition

$q_p \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$. By induction, we know that $t' = q_{p',Q_n,\ldots,Q_2} \xrightarrow[Q_{col}]{b} Q_1$ appears in Δ_{done}^j for some j. Consider the jth iteration of the main loop of Algorithm 6. During this iteration t' is passed to update_rules, and the loop handling rules containing rew_b adds $q_p \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_n)$ to Δ_{new}.

Since this transition contains t, there must be some j' such that t appears in $\Delta_{done}^{j'}$.

We now consider the push rules, which require more intricate reasoning.

- when $o = push_k$, we had $q_{p'} \xrightarrow[Q_{col}]{a} (Q_1, \ldots, Q_k, \ldots, Q_n)$ and T of the form $Q_k \xrightarrow[Q'_{col}]{a} (Q'_1, \ldots, Q'_k)$ in A_i, and we added to the transition

$$q_p \xrightarrow[Q_{col} \cup Q'_{col}]{a} (Q_1 \cup Q'_1, \ldots, Q_{k-1} \cup Q'_{k-1}, Q'_k, Q_{k+1}, \ldots, Q_n)$$

which contains t. By induction, there exists some j where $t_1 = q_{p'Q_n,\ldots,Q_{k+1}} \longrightarrow Q_k$ first appears in Δ_{done}^j. Also by induction, for each $t' \in T$, there is some j' such that t' first appears in $\Delta_{done}^{j'}$. We divide T into $T_1 \cup \cdots \cup T_k$, where $T_{k'}$ contains all order-k' transitions in T.

Consider the jth iteration where t_1 is added to Δ_{done}. During the call to update_rules we call create_trip_wire in the loop handling push rules with $q_k = q_{p,Q_n,\ldots,Q_{k+1}}$, $q_{k-1} = q_{p',Q_n,\ldots,Q_{k+1}}$, $a = a$ and $Q_k = Q_k$.

The call ensures $(q_k, q_{k-1}, a, Q_k) \in \mathcal{U}_{src}^j[k]$. Now take j' such that all $T_k \in \Delta_{done}^{j'}$. We know T_k witnesses $Q_k \xrightarrow{Q_{k-1}} Q'_k$. By Lemma B.4 we know that we have $(Q_k, \emptyset, Q_{k-1}, Q'_k) \in \mathcal{U}_{targ}^{j'}[k]$, and then by Lemma B.6 that we have $\left(q_{kQ'_k}, \perp, a, Q_{k-1} \cup \{q_{k-1}\}\right) \in \mathcal{U}_{src}^{j''}[k-1]$ for some j''.

We essentially iterate the above argument from $k' = k - 1$ down to $k' = 1$. Begin with j' such that $\left(q_{k'}, \perp, a, Q_{k'}^{lbl}\right) \in \mathcal{U}_{src}^{j'}[k']$ and all $T_{k'} \in \Delta_{done}^{j'}$. We know $T_{k'}$ witnesses $Q_{k'}^{lbl} \xrightarrow{Q_{k'-1}^{lbl}} Q_{k'} \cup Q'_{k'}$. By Lemma B.4 we know it to be the case that $\left(Q_{k'}^{lbl}, \emptyset, Q_{k'-1}^{lbl}, Q_{k'} \cup Q'_{k'}\right) \in \mathcal{U}_{targ}^{j'}[k]$, and then by Lemma B.6 that we have $\left(q_{k'Q_{k'} \cup Q'_{k'}}, \perp, a, Q_{k'-1}^{lbl}\right) \in \mathcal{U}_{src}^{j''}[k-1]$ for some j''.

Finally, when $k' = 1$, we have some j' such that $\left(q_1, \perp, a, Q_1^{lbl}\right) \in \mathcal{U}_{src}^{j'}[1]$ and $T_1 \in \Delta_{done}^{j'}$. Note that

$$q_1 = q_{p,Q_n,\ldots,Q_{k+1},Q'_k,Q_{k-1} \cup Q'_{k-1},\ldots,Q_2 \cup Q'_2} .$$

We know $T_{k'}$ witnesses $Q_1^{lbl} \xrightarrow[Q_{col} \cup Q'_{col}]{a} Q_1 \cup Q'_1$. By Lemma B.5 we know it to be the case that $\left(Q_1^{lbl}, \emptyset, a, Q_{col} \cup Q'_{col}, Q_1 \cup Q'_1\right) \in \mathcal{U}_{targ}^{j'}[k]$, and then by Lemma B.7 we have j'' such that we have all t' in

$$q_p \xrightarrow[Q_{col} \cup Q'_{col}]{a} \left(Q_1 \cup Q'_1, \ldots, Q_{k-1} \cup Q'_{k-1}, Q'_k, Q_{k+1}, \ldots, Q_n\right)$$

in $\Delta_{done}^{j''}$. This, in particular, includes t.

- when $o = push^k_b$ we had transitions $q_{p'} \xrightarrow[Q_{col}]{b} (Q_1, \ldots, Q_n)$ and $T = Q_1 \xrightarrow[Q'_{col}]{a} Q'_1$ in A_i with $Q_{col} \subseteq \mathcal{Q}_k$ and added the transitions

$$q_p \xrightarrow[Q'_{col}]{a} (Q'_1, Q_2, \ldots, Q_{k-1}, Q_k \cup Q_{col}, Q_{k+1}, \ldots, Q_n)$$

which include t.

By induction, there exists some j where $t_1 = q_{p'Q_n,\ldots,Q_2} \xrightarrow[Q_{col}]{b} Q_1$ first appears in Δ^j_{done}. Also by induction, for each $t' \in T$, there is some j' such that t' first appears in $\Delta^{j'}_{done}$.

Consider the jth iteration where t_1 is added to Δ_{done}. During the call to update_rules we call create_trip_wire in the loop handling push rules with $q_1 = q_k = q_{p,Q_n,\ldots,Q_{k+1},Q_k \cup Q_{col},Q_{k-1},\ldots,Q_2}$, $q_{k-1} = \bot$, $a = b$ and $Q_k = Q_1$.

The call ensures $(q_1, \bot, b, Q_1) \in \mathcal{U}^j_{src}[1]$. Now take j' such that all $T_1 \in \Delta^{j'}_{done}$. We know T witnesses $Q_1 \xrightarrow[Q_{col}]{a} Q'_1$. By Lemma B.5 we know that we have $(Q_1, \emptyset, a, Q_{col}, Q'_1) \in \mathcal{U}^{j'}_{targ}[1]$, and then by Lemma B.7 we have j'' such that we have all t' in

$$q_p \xrightarrow[Q'_{col}]{a} (Q_1 \cup Q'_1, Q_2, \ldots, Q_{k-1}, Q_k \cup Q_{col}, Q_{k+1}, \ldots, Q_n)$$

in $\Delta^{j''}_{done}$. This, in particular, includes t.

This completes the proof. $\qquad\square$

C. Initial Forward Analysis

C.1 The variant of the saturation algorithm for guarded CPDS

Let $\mathcal{C}$ be a guarded CPDS.

Lemma 5.1 If the revised saturation algorithm is applied to a stack automaton $\mathcal{E}$, then it will output a stack automaton recognising a set T such that:

$$Pre^*_{\mathcal{C}}(\mathcal{E}) \subseteq T \subseteq Pre^*_{\mathsf{Triv}(\mathcal{C})}(\mathcal{E})$$

Proof. We can see that $T \subseteq Pre^*_{\mathsf{Triv}(\mathcal{C})}(\mathcal{E})$ since every time we can add a transition during the modified saturation algorithm we could have added the corresponding guard-free rule in the original, and the original is already known to be sound.

Checking that $Pre^*_{\mathcal{C}}(\mathcal{E}) \subseteq T$ is an easy modification of the completeness proof for the original algorithm in [7]. This works by induction on the length of a path from a configuration in $Pre^*_{\mathcal{E}}(\mathcal{C})$ to one in $\mathcal{E}$. Suppose we have a stack-automaton A recognising a configuration (p', u') together with a rule (p, a, o^S, p') of $\mathcal{C}$ where o is either a *pop* or a *collapse* operation. Suppose that (p, u) can reach (p', u') in a single step via this rule. By definition it must then be the case that $top_1(u') = b$ for some $b \in S$ (and also that $u' = o(u)$). But then the run recognising (p', u') must begin with a transition of the form $q_{p'} \xrightarrow[Q_{col}]{b} (Q_1, \ldots, Q_n)$. Thus in particular $q_{p',Q_n,\ldots,Q_{k+1}} \xrightarrow[Q_{col}]{b} (Q_1, \ldots, Q_k)$ is the first long-form order-k transition in this run. But then taking $q_k := q_{p',Q_n,\ldots,Q_{k+1}}$ we can see that applying the step for the operation o^S in the revised saturation algorithm will create a stack-automaton recognising u. $\qquad\square$

The reason that the algorithm may result in a stack-automaton recognising configurations that do not belong to $Pre^*_{\mathcal{C}}(\mathcal{E})$ (albeit still in $Pre^*_{\mathsf{Triv}(\mathcal{C})}(\mathcal{E})$) is that a stack-automaton state q_k emitting a

transition $q_k \xrightarrow{b} (\_, \ldots, \_)$ may also emit another transition $q_k \xrightarrow{b'} (\_, \ldots, \_)$ with $b \neq b'$. We could obtain a precise algorithm by taking level-n stack-automaton states of the form $\mathcal{P} \times \Sigma$ so that they represent the top stack-character of a configuration as well as its control-state. However, since Σ is usually large compared to $\mathcal{P}$ and since the worst-case size of the stack-automaton is n-exponential in the number of level-n states this would potentially come at a large practical cost and in any case destroy fixed-parameter tractability. We leave it for future work to investigate how this potential for accuracy could be balanced with the inevitable cost.

C.2 The Approximate Reachability Graph and Approximate Summary Algorithm

Let us fix an ordinary n-CPDS with rules $\mathcal{R}$ and initial configuration $c_0 := (p_0, [\cdots [a_0] \cdots])$. A *head* is an element $(p, a) \in \mathcal{P} \times \Sigma$ and should be viewed as describing *stacks* u such that there is a *reachable* configuration of the form (p, u) where $top_1(u) = a$. Formally we define:

$$\llbracket (p, a) \rrbracket := \{ \, u \in Stacks_n \mid top_1(u) = a \text{ and } (p, u) \in Post^*_\mathcal{C} \, \}$$

A *stack descriptor* is an $(n + 1)$-tuple $(h_n, \ldots, h_1, h_c)$ where for each $1 \leq i \leq n$, each of h_i and h_c is either a head or $\bot$. We write $\mathbf{SDesc} := (\mathcal{P} \times \Sigma)^{\bot^{n+1}}$ for the set of stack descriptors and it will also be useful to have $\mathbf{SDesc}_k := (\mathcal{P} \times \Sigma)^{\bot^{n-k}}$ for the set of *order-k stack-descriptor prefixes*. Note that $\mathbf{SDesc}_n = \{()\}$—*i.e.* consists only of the empty tuple. Assuming a map $B : (\mathcal{P} \times \Sigma) \to \mathbf{SDesc}$ a stack descriptor describes a set of stacks $\llbracket (h_n, \ldots, h_1, h_c) \rrbracket :=$

$$\left\{ u \in Stacks_n \; \middle| \; \begin{array}{l} pop_k(u) \text{ is undefined if } h_k = \bot \text{ otherwise} \\ top_1(pop_k(u)) = b_k \text{ where } h_k = (\_, b_k) \text{ and} \\ pop_k(u) \in \llbracket (h_n, \ldots, h_{k+1}, h'_k, \ldots, h'_1, h'_c) \rrbracket \\ \quad \text{for some } (\_, \ldots, \_, h'_k, \ldots, h'_1, h'_c) \in B(h_k), \\ \quad \text{for every } 1 \leq k \leq n, \\ \text{and } top_1(u) \text{ emits no link if } h_c = \bot \text{ otherwise} \\ top_1(collapse_k(u)) = b_c \text{ where } h_c = (\_, b_c) \text{ and} \\ collapse_k(u) \in \llbracket (h_n, \ldots, h_{k+1}, h'_k, \ldots, h'_1, h'_c) \rrbracket \\ \quad \text{for some } (\_, \ldots, \_, h'_k, \ldots, h'_1, h'_c) \in B(h_c), \\ \quad \text{for every } 1 \leq k \leq n \end{array} \right\}$$

We now define an approximate reachability graph. In the main body of the paper we defined this to be a pair (H, E), and then later consider an extended structure (H, E, B, U). This is because H and E contain all of the information used to extract the GCPDS. However, for the soundness proof we will express an invariant in terms of the 'semantics' of an approximate reachability graph, and it is helpful to have B as part of this semantics. We thus add the B component to the approximate reachability graph for the purposes of the appendix (so it is now a triple (H, E, B).

Definition C.1. *An* approximate reachability graph *for the CPDS $\mathcal{C}$ is a triple (H, E, B) such that (i) $H \subseteq \mathcal{P} \times \Sigma$ is a set of heads such that $(p, u) \in Post^*_\mathcal{C}$ implies that $(p, top_1(u)) \in H$, (ii) $E \subseteq H \times \mathcal{R} \times H$ is a set of triples such that if $(p, u) \in Post^*_\mathcal{C}$ and $r = (p, top_1(u), o, p') \in \mathcal{R}$ for which $o(u)$ is defined, then $((p, top_1(u)), r, (p', top_1(o(u)))) \in E$, (iii) B is a map $B : H \to \mathbf{SDesc}$ such that for every $h \in H$ we have $\llbracket h \rrbracket \subseteq \{ \, \llbracket d \rrbracket \mid d \in B(h) \, \}$.*

Let $\mathcal{G} = (H, E, B)$ be an approximate reachability graph for $\mathcal{C}$. Let $\mathbf{Heads}(\mathcal{E})$ be the set of heads of error configurations, *i.e.* $\mathbf{Heads}(\mathcal{E}) := \{ \, (p_{err}, a) \mid a \in \Sigma \, \}$. We do a simple backwards reachability computation on the finite graph $\mathcal{G}$ to compute $\mathbf{BackRules}(\mathcal{G})$, defined to be the smallest set satisfying:

$$\mathbf{BackRules}(\mathcal{G}) = \{ \, e \in E \mid e = (h, r, h') \in E \text{ for some } h' \in \mathbf{Heads}(\mathcal{E}) \, \}$$
$$\cup \, \{ \, e \in E \mid e = (h, r, h') \in E \text{ for some } (h', \_, \_) \in \mathbf{BackRules}(\mathcal{G}) \, \}$$

The CPDS rules occurring in the triples in **BackRules**($\mathcal{G}$) can be used to define a pruned CPDS that is safe if and only if the original also is. However, the approximate reachability graph provides enough information to construct a guarded CPDS whose guards are non-trivial. It should be clear that the following set **BackRulesG**($\mathcal{G}$) of *guarded* rules can be computed:

$$\left\{ (p, a, o', p') \ \middle| \ \begin{array}{l} (\_, (p, a, o, p'), \_) \in \textbf{BackRules}(\mathcal{G}) \text{ and} \\ o' = \begin{cases} o^S & \text{if } o \text{ is a pop or a collapse and } S \text{ is} \\ & \left\{ b \in \Sigma \ \middle| \ \begin{array}{l} ((p, a), r, (p', b)) \\ \in E \end{array} \right\} \\ & \text{with } r = (p, a, o, p') \\ o & \text{if } o \text{ is a rewrite or push} \end{cases} \end{array} \right\}$$

These rules define a GCPDS on which C-SHORe finally performs saturation.

Lemma C.1. *The GCPDS $\mathcal{C}'$ defined using the rules* **BackRulesG**($\mathcal{G}$) *satisfies:*

$$Post^*_{\mathcal{C}} \ \cap \ Pre^*_{\mathcal{C}}(\mathcal{E}) \ \subseteq \ Pre^*_{\mathcal{C}'}(\mathcal{E}) \ \subseteq \ Pre^*_{\mathcal{C}}(\mathcal{E})$$

Proof. $Pre^*_{\mathcal{C}'}(\mathcal{E}) \subseteq Pre^*_{\mathcal{C}}(\mathcal{E})$ is trivial since the rules of **Triv**($\mathcal{C}'$) are a subset of the rules for $\mathcal{C}$.

Now suppose that $(p_1, u_1) \in Post^*_{\mathcal{C}} \cap Pre^*_{\mathcal{C}}(\mathcal{E})$. By (i) in the definition of approximate reachability graph it must be the case that $(p_1, top_1(u_1)) \in H$ (since $(p_1, u_1) \in Post^*_{\mathcal{C}}$).

Since $(p, u) \in Pre^*_{\mathcal{C}}(\mathcal{E})$ we must also have a finite sequence of $\mathcal{C}$-rules $r_i = (p_i, a_i, o_i, p_{i+1})$ for $1 \leq i < m$, and a finite sequence of configurations such that $(p_{i+1}, u_{i+1}) = (p_i, o_i(u_i))$, and $p_m = p_{err}$. But then by (i) and (ii) in the definition of approximate reachability graph, $h_i := (p_i, top_1(u_i)) \in H$ for every $1 \leq i \leq m$ and $(h_i, r_i, h_{i+1}) \in E$ for every $1 \leq i < m$.

Thus when o_i is neither a *pop* nor *collapse* operation $r'_i := r_i$ will itself occur as a rule of $\mathcal{C}'$. Otherwise $r'_i := (p_i, a_i, o_i^S, p_{i+1})$ will be in $\mathcal{C}'$ where $a_{i+1} \in S$. Thus $r'_1, \ldots, r'_{m-1}$ witnesses $(p_1, u_1) \in Pre^*_{\mathcal{C}'}(\mathcal{E})$, as required. □

A non-trivial approximate reachability graph is computed using an algorithm that works in the *forwards* direction (unlike saturation which works backwards), and which resembles a summary algorithm.

C.3 The Approximate Summary Algorithm

The approximate summary algorithm computes an approximate reachability graph (H, E, B) 'as accurately as possible based on an order-1 approximation'. In order to do this, the algorithm builds up an object (H, E, B, U) where the additional component U is a set of *approximate higher-order summary edges*. An *order-k summary edge* is a triple in $H \times \textbf{SDesc}_k \times H$. Intuitively such a summary $(h, (h'_n, \ldots, h'_{k+1}), h')$ indicates that if $(h_n, \ldots, h_{k+1}, h_k, \ldots, h_1, h_c) \in B(h)$, then we should also have $(h'_n, \ldots, h'_{k+1}, h_k, \ldots, h_1, h_c) \in B(h')$. When $n = k = 1$ (so that h_c is also unnecessary since there would be no links) note that $(h, (), h')$ behaves like a summary edge in a standard order-1 summary algorithm [34], which is complete at order-1.

The algorithm is presented as Algorithm 16.

Lemma C.2. *Algorithm 16 terminates and the resulting structure (H, E, B, U) gives an approximate reachability graph (H, E, B).*

Proof. For termination note that the respective procedures in Algorithms 17 and 19 will immediately return if the stack-descriptor (respectively summary) that they are called with is already contained in a particular set. If it does not belong to this set, then it is added. Since there are only finitely many possible arguments for these functions, they can thus only be called finitely many times without immediately returning. From this fact it is easy to see that the entire algorithm must always terminate.

Algorithm 16 The Approximate Summary Algorithm

Require: An n-CPDS with rules $\mathcal{R}$ and heads $\mathcal{P} \times \Sigma$ and initial configuration $(p_0, [\cdots [a_0] \cdots])$

Ensure: The creation of a structure (H, E, B, U) where (H, E, B) is an approximate reachability graph and U is a set of approximate higher-order summary edges.

Set $H := \{(p_0, a_0)\}$ and set E, B and U to be empty
Call AddStackDescriptor$((p_0, a_0), (\bot, \ldots, \bot, \bot))$
return Done, (H, E, B, U) will now be as required

Algorithm 17 AddStackDescriptor$(h, (h_n, \ldots, h_1, h_c))$

Require: A head $h \in H$ and a stack descriptor $(h_n, \ldots, h_1, h_c)$

Ensure: $(h_n, \ldots, h_1, h_c) \in B(h)$ and that any further additions to $B(h')$ for each $h' \in H$ necessary to respect summary edges are made.

if $(h_n, \ldots, h_1, h_c) \in B(h)$ **then**
   return Done (Nothing to do)
Add $(h_n, \ldots, h_1, h_c)$ to $B(h)$
Call ProcessHeadWithStackDescriptor$(h, (h_n, \ldots, h_1, h_c))$
for $h' \in H$ such that $(h, (h'_n, \ldots, h'_{k+1}), h') \in U$ **do**
   Call AddStackDescriptor$(h', (h'_n, \ldots, h'_{k+1}, h_k, \ldots, h_1, h_c))$
return Done

Now we show that (H, E, B) is an approximate reachability graph. Recursively define $Post_{\mathcal{C}}^0 := \{c_0\}$ and

$$Post_{\mathcal{C}}^{i+1} := Post_{\mathcal{C}}^i \cup \left\{ c \mid \exists c^- \in Post_{\mathcal{C}}^i \text{ s.t. } c^- \to_{\mathcal{C}} c \text{ in one step} \right\}$$

That is $Post_{\mathcal{C}}^i$ is the set of configurations that can be reached from the initial configuration in at most i steps. For a head $(p, a) \in \mathcal{P} \times \Sigma$, define

$$[\![(p, a)]\!]_i := \left\{ (p, u) \mid (p, u) \in Post_{\mathcal{C}}^i \text{ and } top_1(u) = a \right\}$$

We can now define an *i-partial approximate reachability graph* to be a version of an approximate reachability graph defined for 'reachability up to depth i'.

Definition C.2. *An i-partial approximate reachability graph for the CPDS $\mathcal{C}$ is a triple (H, E, B) such that (i) $H \subseteq \mathcal{P} \times \Sigma$ is a set of heads such that $(p, u) \in Post_{\mathcal{C}}^i$ implies that $(p, top_1(u)) \in H$, (ii) $E \subseteq H \times \mathcal{R} \times H$ is a set of triples such that if $i > 0$ and $(p, u) \in Post_{\mathcal{C}}^{i-1}$ and $r = (p, top_1(u), o, p') \in \mathcal{R}$ for which $o(u)$ is defined, then $((p, top_1(u)), r, (p', top_1(o(u)))) \in E$, (iii) B is a map $B : H \to$ **SDesc** such that for every $h \in H$ we have $[\![h]\!]_i \subseteq \{ [\![d]\!] \mid d \in B(h) \}$.*

Observe that a structure (H, E, B) is an approximate reachability graph if and only if it is an i-partial approximate reachability graph for every $i \geq 0$.

Now observe that the algorithm monotonically grows the sets making up (H, E, B, U) (it only adds to the sets, it never removes from them). We may thus argue by induction to show that the (H, E, B) after termination is an i-partial approximate reachability graph for every $i \geq 0$ (and hence

Algorithm 18 ProcessHeadWithStackDescriptor($h, (h_n, \ldots, h_1, h_c)$)

Require: A head $h := (p, a) \in H$ and a stack descriptor $(h_n, \ldots, h_1, h_c) \in B(h)$

Ensure: All necessary modifications to the graph are made so that it is consistent with the presence of $(h_n, \ldots, h_1, h_c) \in B(h)$. In particular this is the procedure that processes the CPDS rules from h (with respect to a stack described by h and the stack descriptor)

> **for** o and p' such that $r = (p, a, o, p') \in \mathcal{R}$ **do**
> > **if** o of form rew_b **then**
> > > Add (p', b) to H
> > > Add $((p, a), r, (p', b))$ to E
> > > Call AddStackDescriptor($(p', b), (h_n, \ldots, h_1, h_c)$)
> >
> > **else if** o of form $push_b^k$ **then**
> > > Add (p', b) to H
> > > Add $((p, a), r, (p', b))$ to E
> > > Call AddStackDescriptor($(p', b), (h_n, \ldots, h_2, (p, a), h_k)$)
> >
> > **else if** o of form $push_k$ **then**
> > > Add (p', a) to H
> > > Add $((p, a), r, (p', a))$ to E
> > > Call AddStackDescriptor($(p', a), (h_n, \ldots, h_{k+1}, (p, a),$
> > > $h_{k-1}, \ldots, h_1, h_c)$)
> >
> > **else if** o of form pop_k with $h_k = (p_k, a_k)$ where $a_k \neq \perp$ **then**
> > > Add (p', a_k) to H
> > > Add $((p, a), r, (p', a_k))$ to E
> > > Call AddSummary($(p_k, a_k), (h_n, \ldots, h_{k+1}), (p', a_k)$)
> >
> > **else if** o of form $collapse_k$ with $h_c = (p_c, a_c)$ where $a_c \neq \perp$ **then**
> > > Add (p_c, a_c) to H
> > > Add $((p, a), r, (p', a_c))$ to E
> > > Call AddSummary($(p_c, a_c), (h_n, \ldots, h_{k+1}), (p', a_c)$)
> **return** Done

Algorithm 19 AddSummary($h, (h'_n, \ldots, h_{k+1}), h'$)

Require: An approximate higher-order summary edge $(h, (h'_n, \ldots, h_{k+1}), h')$

Ensure: $(h, (h'_n, \ldots, h'_{k+1}), h') \in U$ and that all necessary stack descriptors are added to the appropriate $B(h'')$ for $h'' \in H$ so that all summary edges (including the new one) are respected.

> **if** $(h, (h'_n, \ldots, h'_{k+1}), h') \in U$ **then**
> > **return** Done (Nothing to do)
>
> Add $(h, (h'_n, \ldots, h'_{k+1}), h')$ to U
> **for** $(h_n, \ldots, h_{k+1}, h_k, \ldots, h_1, h_c) \in B(h)$ **do**
> > AddStackDescriptor($h', (h'_n, \ldots, h'_{k+1}, h_k, \ldots, h_1, h_c)$)
> **return** Done

an approximate reachability graph). First note that the opening statements of Algorithm 16 (including the call to add $(\bot, \ldots, \bot, \bot)$ as a stack descriptor to $B(p_0, a_0)$) guarantees that (H, E, B) is a 0-partial approximate reachability graph.

Now suppose that (H, E, B) is an i-partial approximate reachability graph. We show that it is also an $(i + 1)$-partial approximate reachability graph. Let $(p, u) \in Post_C^i$ and let $r := (p, a, o, p') \in \mathcal{R}$ be such that $o(u)$ is defined and $top_1(u) = a$ so that $(p', o(u)) \in Post_C^{i+1}$. Let $a' := top_1(o(u))$. It suffices to show that (i) $h' := (p', a') \in H$, (ii) $e := ((p, a), r, (p', a')) \in E$ and that (iii) Some $d' := (h'_n, \ldots, h'_1, h'_c) \in B(p', a')$ with $o(u) \in [\![d']\!]$.

By the induction hypothesis (that the structure is an i-partial approximate reachability graph) we must have $h := (p, top_1(u)) \in H$ and $d = (h_n, \ldots, h_1, h_c) \in B(h)$ such that $u \in [\![d]\!]$. Inspection of the algorithm shows that the addition of d to $B(h)$ is only possible if **AddStackDescriptor**(h, d) was called at some point during its execution. However, this also implies that **ProcessHeadWithStackDescriptor**(h, d) must have been called.

Note also that when o is a rewrite operation we must have $pop_j(o(u)) = pop_j(u)$ and $collapse_j(o(u)) = collapse_j(u)$ for all j. When $o = push_k$ for $k \geq 2$ we must have $top_{j+1}(pop_j(o(u))) = top_{j+1}(pop_j(u))$ and $top_{j+1}(collapse_j(o(u))) = top_{j+1}(collapse_j(u))$ for all $j \neq k$ and $pop_k(o(u)) = u$. When $o = push_{b'}^k$ we must have $pop_j(o(u)) = pop_j(u)$ for all $j \geq 2$, but $pop_1(o(u)) = u$ and $collapse_k(o(u)) = pop_k(u)$.

Thus if o is any operation other than pop_k or $collapse_k$ it can be seen that **AddStackDescriptor**(h', d') must be called for a d' such that $u \in [\![d']\!]$. Also, e is added to E. Since the algorithm never deletes elements from sets, this ensures that (H, E, B) must satisfy the constraints (i), (ii) and (iii) above.

Now consider the case when o is either pop_k or $collapse_k$. Suppose again that $top_1(o(u)) = a'$. Since $u \in [\![d]\!]$ we must have:

- For some control-state p^- we have: $h_k = (p^-, a')$ if $o = pop_k$ and $h_c = (p^-, a')$ if $o = collapse_k$ such that. . .

- . . . there exists $(\_, \ldots, \_, h'_k, \ldots, h'_1, h'_c) \in B((p^-, a'))$ such that $o(u) \in [\![(h_n, \ldots, h_{k+1}, h'_k, \ldots, h'_1, h'_c)]\!]$.

Thus a suitable d' is $d' = (h_n, \ldots, h_{k+1}, h'_k, \ldots, h'_1, h'_c)$.

The call to **ProcessHeadWithStackDescriptor**(h, d) guarantees that (i) $h' = (p', a') \in H$ and (ii) $e := ((p, a), r, (p', a')) \in E$. It just remains to check that $d' \in B((p', a'))$.

Note that the above call must also ensure a call to **AddSummary**$((p^-, a'), (h_n, \ldots, h_{k+1}), (p', a'))$. We are thus guaranteed the existence of a summary edge s:

$$((p^-, a'), (h_n, \ldots, h_{k+1}), (p', a')) \in U$$

(although it may have been added at an earlier point in the algorithm).

There are two cases to consider:

- If the summary edge s was created *after* a stack-descriptor of the form $(\_, \ldots, \_, h'_k, \ldots, h'_1, h'_c)$ was added to $B((p^-, a'))$, then the call to **AddSummary** creating s must add d' to $B((p', a'))$.

- If the summary edge s was created *before* a stack-descriptor of the form $d^- = (\_, \ldots, \_, h'_k, \ldots, h'_1, h'_c)$ was added to $B((p^-, a'))$, then the call to **ProcessHeadWithStackDescriptor**$((p^-, d^-))$ creating this stack-descriptor must result in d' being added to $B((p', a'))$.

Either way, (iii) must also be satisfied. $\qquad\blacksquare$

C.4 A Remark On Complexity

The approximate summary algorithm runs in time polynomial in the size of the CPDS. Since the graph constructed must also be of polynomial size, it follows that the rules for the guarded CPDS $\mathcal{C}'$ can also

be extracted in polynomial time. Since the raw saturation algorithm is also PTIME when the number of control-states is fixed, it follows that the C-SHORe algorithm as a whole is fixed-parameter tractable.

We sketch here how to see that the approximate summary algorithm runs in polynomial time. First note that an approximate reachability graph can contain at most $|Q|.|\Sigma|$ heads and at most $|Q|.|\Sigma|.|\mathcal{R}|.|Q|.|\Sigma|$ edges (recalling that $\mathcal{R}$ is the set of CPDS rules). Moreover the maximum size of the function B (viewed as a relation $\{(h, d) \in (Q \times \Sigma) \times (Q \times \Sigma)^{n+1} \mid d \in B(h)\}$) is $|Q|.|\Sigma|.(|Q|.|\Sigma|)^{n+1}$. The maximum number of summary edges is

$$\sum_{i=2}^{n} |Q|.|\Sigma|.(|Q|.|\Sigma|)^{n-i}.|Q|.|\Sigma|$$

It follows that the size of the structure (H, E, B, U) constructed by algorithm is at most polynomial in the size of the original CPDS. Moreover, since the algorithm only *adds* to the structure and never removes elements previously added, it will perform at most polynomially many additions. Let Z be this polynomial bound on the size of the structure.

Moreover, recall that the procedures for adding summaries and heads/stack-descriptors are guarded so that the procedure only processes the new object if it had not already been added; if it had already been added, the procedure in question will return after constant time.

So we consider the cases when the object being created is new. For each *new* head/stack-descriptor pair, **ProcessHeadWithStackDescriptor** will check it against every rule and for each rule may attempt to create a new object. Disregarding the result of the calls to create *new* objects (with calls to create old objects returning in constant time), the run-time of this procedure will thus be bounded by $O(|\mathcal{R}|)$. Likewise each time a *new* stack descriptor is added, **AddStackDescriptor** will compare it against existing summary edges and so run in time $O(Z)$.

Similarly the run-time of a call to **AddSummary** on a new summary edge (disregarding run-times to calls from this procedure that create *new* objects) is $O(Z)$ since the new summary edge will, at worst, be compared against every possible stack-descriptor.

Thus creating a new object takes at most $O(Z.|\mathcal{R}|)$ time and new objects are created only during the call to a procedure that itself is creating a new object. Thus the overall run-time is bounded by $O(Z.Z.|\mathcal{R}|)$ and so is polynomial.

D. Full Experimental Results

The full experimental results are given in Table 2. In the main paper we restricted attention to trials where at least one tool took over 1s. This is because virtual machine "warm-up" and HORS to CPDS conversion can skew the results on small benchmarks.

Benchmark file	Ord	Sz	T	G	TMC	C	Ctran	Ccpds	Capprox
example2-1	1	7	0.003	0.002	0.029	0.189	0.040	0.150	0.017
example2-3	1	13	0.002	0.003	0.027	0.053	0.037	0.016	0.010
example3-1 (bug)	1	8	0.000	—	0.061	0.070	0.041	0.029	0.014
exception	1	18	0.003	0.004	0.028	0.053	0.038	0.015	0.009
file	1	8	0.000	0.006	0.028	0.063	0.041	0.022	0.014
fileocamlc	4	111	0.011	0.062	0.049	0.975	0.117	0.858	0.226
filewrong (bug)	4	45	0.000	0.017	0.062	0.247	0.086	0.161	0.080
flow	4	16	0.001	0.004	0.029	0.093	0.047	0.046	0.023
lock1	4	38	0.008	0.007	0.032	0.124	0.070	0.055	0.043
lock2	4	45	0.016	0.268	0.062	0.466	0.079	0.386	0.217
order5	5	52	0.010	—	0.045	0.531	0.098	0.434	0.196
order5-2	5	40	0.017	—	0.091	0.317	0.082	0.235	0.128
order5-variant	5	55	0.017	1.470	0.042	0.525	0.095	0.430	0.170
twofiles	4	47	0.010	0.087	0.039	0.269	0.088	0.181	0.095
twofilesexn	4	56	0.011	0.040	0.041	0.268	0.094	0.174	0.086
checknz	2	93	0.004	0.009	0.032	0.117	0.074	0.043	0.028
checkpairs (bug)	2	251	0.004	0.031	0.064	0.287	0.102	0.185	0.064
filepath	2	5956	210.767	—	—	0.657	0.426	0.231	0.192
filter-nonzero (bug)	5	484	0.008	0.163	0.126	2.150	0.177	1.973	1.149
filter-nonzero-1	5	890	0.170	—	216.842	6.879	0.399	6.480	4.120
last	2	193	0.012	0.014	0.041	0.157	0.094	0.063	0.054
map-head-filter (bug)	3	370	0.007	0.121	0.085	0.647	0.128	0.519	0.122
map-head-filter-1	3	880	0.137	—	1.359	0.469	0.205	0.263	0.239
map-plusone	5	302	0.013	0.126	0.126	1.282	0.162	1.120	0.871
map-plusone-1	5	459	0.023	—	0.760	1.771	0.223	1.547	1.031
map-plusone-2	5	704	1.325	—	14.242	3.559	0.323	3.236	2.184
mkgroundterm	2	379	0.038	0.039	0.082	0.252	0.125	0.127	0.116
risers	2	563	0.043	0.058	0.098	0.625	0.148	0.477	0.117
safe-head	3	354	0.017	0.054	0.047	0.500	0.132	0.368	0.086
safe-init	3	680	0.031	0.300	0.228	1.233	0.171	1.062	0.175
safe-tail	3	468	0.028	0.087	0.062	0.756	0.143	0.613	0.123
tails	3	259	0.019	0.447	0.056	0.191	0.108	0.082	0.074
g41	4	31	—	0.107	0.054	—	0.068	—	0.244
merge4	2	141	0.066	0.491	0.236	0.896	0.416	0.480	0.164
stress	1	35	0.019	0.004	0.152	0.073	0.043	0.030	0.022

Table 2: Comparison of model-checking tools on all benchmarks.