A Fixed-Parameter Algorithm for the Directed Feedback Vertex Set Problem

JIANER CHEN, YANG LIU, AND SONGJIAN LU

Texas A&M University, College Station, Texas

AND

BARRY O'SULLIVAN AND IGOR RAZGON

University College Cork, Cork, Ireland

Abstract. The (parameterized) FEEDBACK VERTEX SET problem on directed graphs (i.e., the DFVS problem) is defined as follows: given a directed graph G and a parameter k, either construct a feedback vertex set of at most k vertices in G or report that no such a set exists. It has been a well-known open problem in parameterized computation and complexity whether the DFVS problem is fixed-parameter tractable, that is, whether the problem can be solved in time $f(k)n^{O(1)}$ for some function f. In this article, we develop new algorithmic techniques that result in an algorithm with running time $4^k k! n^{O(1)}$ for the DFVS problem. Therefore, we resolve this open problem.

Categories and Subject Descriptors: F.2.0 [**Analysis of Algorithms and Problem Complexity**]: General

General Terms: Algorithms, Theory

Additional Key Words and Phrases: Deadlock, feedback vertex set, parameterized computation

Authors' addresses: J. Chen, Y. Liu, and S. Lu, Department of Computer Science, Texas A&M University, College Station, TX 77843, e-mail: {chen,yangliu,sjlu}@cs.tamu.edu; B. O'Sullivan and I. Razgon, Computer Science Department, University College Cork, Ireland, e-mail: {b.osullivan, i.razgon}@cs.ucc.ie.

1. *Introduction*

Let G be a directed graph. A *feedback vertex set F* (briefly, FVS) for G is a set of vertices in G such that every directed cycle in G contains at least one vertex in F, or equivalently, that the removal of F from the graph G leaves a directed acyclic graph (i.e., a DAG). The (parameterized) FEEDBACK VERTEX SET problem on directed graphs (briefly, the DFVS problem) is defined as follows: given a directed graph G and a parameter k, either construct an FVS of at most k vertices for G or report that no such set exists.

The DFVS problem is a classic NP-complete problem that appeared in the first list of NP-complete problems in Karp's seminal paper [Karp 1972], and has a variety of applications in areas such as operating systems [Silberschatz and Galvin 1994], database systems [Gardarin and Spaccapietra 1976], and circuit testing [Leiserson and Saxe 1991]. In particular, the DFVS problem has played an essential role in the study of *deadlock recovery* in database systems and in operating systems [Silberschatz and Galvin 1994; Gardarin and Spaccapietra 1976]. In such a system, the status of system resource allocations can be represented as a directed graph G (i.e., the *system resource-allocation graph*), and a directed cycle in G represents a deadlock in the system. Therefore, in order to recover from deadlocks, we need to abort a set of processes in the system, that is, to remove a set of vertices in the graph G, so that all directed cycles in G are broken. Equivalently, we need to find an FVS in the graph G. In practice, one may expect and desire that the number of vertices removed from the graph G, which is the number of processes to be aborted in the system, be small. This motivates the study of *parameterized algorithms* for the DFVS problem that find an FVS of k vertices in a directed graph of n vertices and run in time $f(k)n^{O(1)}$ for a fixed function f; thus, the algorithms become practically efficient when the value k is small.

This work has been part of a systematic study of the *theory of fixed-parameter tractability* [Downey and Fellows 1999], which has received considerable attention in recent years. A problem Q is a *parameterized problem* if each instance of Q contains a specific integral parameter k. A parameterized problem is *fixed-parameter tractable* if it can be solved in time $f(k)n^c$ for a function $f(k)$ and a constant c, where the function $f(k)$ is independent of the instance size n. A large number of NP-hard parameterized problems, such as the VERTEX COVER problem [Chen et al. 2001] and the ML TYPE-CHECKING problem [Lichtenstein and Pnueli 1985], have been shown to be fixed-parameter tractable. On the other hand, strong evidence has been given that another group of well-known parameterized problems, including the INDEPENDENT SET problem and the DOMINATING SET problem, are not fixed-parameter tractable [Downey and Fellows 1999]. The study of fixed-parameter tractability of parameterized problems has become increasingly interesting, for both theoretical research and practical computation.

The fixed-parameter tractability of the DFVS problem was posted as an open problem in the very first papers on the study of fixed-parameter tractability [Downey and Fellows 1992, 1995]. After numerous significant efforts, however, the problem still remained open. In the past fifteen years, the problem has been constantly and explicitly posted as an open problem in a large number of publications in the literature (see [Gutin and Yeo 2008] for a recent survey on this study). The problem has become a well-known and outstanding open problem in parameterized computation and complexity.

In this article, we develop new algorithmic techniques that lead to the conclusion that the DFVS problem is fixed-parameter tractable, and thus resolve the above open problem in parameterized computation and complexity. We first show that the DFVS problem can be reduced in time $f(k)n^{O(1)}$ for some function f to a special version of the multi-cut problem, which will be called the SKEW SEPARATOR problem. We then develop an algorithm that shows the fixed-parameter tractability of the SKEW SEPARATOR problem. The combination of these two results gives an algorithm with running time $4^k k! n^{O(1)}$ for the DFVS problem, which proves its fixed-parameter tractability.

The relationship between the DFVS problem and multi-cut problems has been studied in the research of approximation algorithms for the FEEDBACK VERTEX SET problem [Even et al. 1998; Leighton and Rao 1999]. However, our problem formulations and the corresponding techniques are significantly different from those studied in the approximation algorithms. In particular, our formulations and techniques seem especially suitable for developing faster and more effective exact algorithms (of exponential-time) for NP-hard multi-cut problems. First of all, instead of seeking a multi-cut that separates a given set of *terminal vertices*, as formulated in most multi-cut problems, our problem is more general: we wish to construct a multi-cut that separates a *collection of terminal vertex-subsets*. This more general version of the multi-cut problem enables us to effectively reduce the search space size when we are searching for an *optimal* solution of a given problem instance. Second, unlike most multi-cut problems whose solutions are multi-cuts that are in general symmetric to the given terminal vertices, the multi-cuts for the SKEW SEPARATORproblem are asymmetric to the terminal vertex-subsets. Third, we develop an (exponential-time) reduction that effectively reduces the problem of multi-cuts for multiple terminal vertex-subsets to the problem of minimum cuts from a single source vertex to a single sink vertex. Note that the latter is solvable in polynomial time via algorithms for the MAXIMUM FLOW problem. Such an exponential-time reduction is obviously very different from the polynomial-time processes used in the development of polynomial-time approximation algorithms. Finally, unlike most parameterized algorithms that are focused on effectively decreasing the sole parameter value k, our algorithm for the SKEW SEPARATOR problem adds another dimension of bounds in terms of the size of a minimum cut between two properly chosen terminal vertex-subsets. This dimension of bounds has become crucial in our development of the algorithm for the SKEW SEPARATOR problem because it effectively bounds the number of branches in which the parameter value k is not decreased.

Before we move to the technical discussion of our algorithms, we remark that the FEEDBACK VERTEX SET problem on undirected graphs (briefly, the UFVS problem) has also been an interesting and active research topic in parameterized computation and complexity. Since the first fixed-parameter tractable algorithm for the UFVS problem was published fifteen years ago [Bodlaender 1993], there has been an impressive list of improved algorithms for the problem. Currently the best algorithm for the UFVS problem runs in time $O(5^k k n^2)$ [Chen et al. 2007a]. The FEEDBACK VERTEX SET problem on directed graphs (i.e., the DFVS problem) seems very different from the problem on undirected graphs (i.e., the UFVS problem). This fact has also been reflected in the study of approximation algorithms for the problems. The FEEDBACK VERTEX SET problem on undirected graphs is polynomial-time approximable with a ratio 2. This holds true even for weighted graphs [Bafna et al. 1999]. On the other

hand, it still remains open whether the FEEDBACK VERTEX SET problem on directed graphs has a constant-ratio polynomial-time approximation algorithm. The current best polynomial-time approximation algorithm for the problem on directed graphs has a ratio $O(\log \tau \log \log \tau)$, where τ is the size of a minimum FVS for the input graph [Even et al. 1998].

2. *Preliminaries*

Let $G = (V, E)$ be a directed graph and let $e = [u, v]$ be a (directed) edge in G. We say that the edge *e goes out* from the vertex u and *comes into* the vertex v. The edge e is called an *outgoing edge* of the vertex u, and an *incoming edge* of the vertex v. These concepts can be extended from single vertices to general vertex sets. Thus, for two vertex sets S_1 and S_2, we can say that an edge goes out from S_1 and comes into S_2 if the edge goes out from a vertex in S_1 and comes into a vertex in S_2. Moreover, we say that an edge *goes out from* S_1 if the edge goes out from a vertex in S_1 and comes into a vertex not in S_1, and that an edge *comes into* S_2 if the edge goes out from a vertex not in S_2 and comes into a vertex in S_2.

A *path P* from a vertex v_1 to a vertex v_h in the graph G is a sequence $\{v_1, v_2, \ldots, v_h\}$ of vertices in G such that $[v_i, v_{i+1}]$ is an edge in G for all $1 \leq i \leq h - 1$. The path P is *simple* if no vertex is repeated in P. The path P is a *cycle* if $v_1 = v_h$, and the cycle is *simple* if no other vertices are repeated. We say that a path is from a vertex set S_1 to a vertex set S_2 if the path is from a vertex in S_1 to a vertex in S_2. The graph G is a *DAG* (i.e., directed acyclic graph) if it contains no cycles.

For a vertex subset $V' \subseteq V$ in the directed graph $G = (V, E)$, we denote by $G[V']$ the subgraph of G that is induced by the vertex subset V'. Without any ambiguity, we will denote by $G - V'$ the induced subgraph $G[V - V']$, and by $G - w$, where w is a vertex in G, the induced subgraph $G[V - \{w\}]$.

A vertex subset F in the directed graph G is a *feedback vertex set* (*FVS*) if the graph $G - F$ is a DAG. Since a vertex v with a self-loop (i.e., an edge that both goes out from and comes into v) must be included in every FVS for the graph G, we will assume, without loss of generality, that the graphs in our discussion have no self-loops.

Definition 2.1.   Let $[S_1, \ldots, S_l]$ and $[T_1, \ldots, T_l]$ be two collections of l vertex subsets in a directed graph $G = (V, E)$. A *skew separator* X for $([S_1, \ldots, S_l],$ $[T_1, \ldots, T_l])$ is a vertex subset in $V - \bigcup_{i=1}^{l}(S_i \cup T_i)$ such that for any pair of indices i and j satisfying $l \geq i \geq j \geq 1$, there is no path from S_i to T_j in the graph $G - X$.

The subsets $S_1, \ldots, S_l$ will be called the *source sets* and the subsets $T_1, \ldots, T_l$ will be called the *sink sets*. A vertex is a *non-terminal vertex* if it is not in $\bigcup_{i=1}^{l}$ $(S_i \cup T_i)$. Note that by definition, all vertices in a skew separator must be non-terminal vertices. Moreover, a skew separator X is asymmetric to the source sets and the sink sets: a path from S_i to T_j with $i < j$ may exist in the graph $G - X$.

When there is only one source set S_1 and one sink set T_1, a skew separator for the pair $([S_1], [T_1])$ becomes a regular cut for S_1 and T_1, i.e., a vertex set whose removal leaves a graph in which there is no path from the set S_1 to the set T_1. Therefore, a skew separator for the pair $([S_1], [T_1])$ is also called a *cut from S_1 to T_1*. A cut from

S_1 to T_1 is a *min-cut* (i.e., a minimum cut) if it has the smallest cardinality over all cuts from S_1 to T_1.

The following lemma can be easily derived based on standard maximum flow techniques [Schrijver 2003]. Thus, we omit its proof.

LEMMA 2.2. *There is an $O(kn^2)$ time algorithm that for two given vertex subsets S and T in a directed graph G of n vertices, and a parameter k, either constructs a min-cut from S to T whose size is bounded by k, or reports that the min-cut from S to T has a size larger than k.*

The algorithm for the DFVS problem is obtained through careful development of algorithms for a series of problems. In the following, we give the formal definitions of these problems.

> SKEW SEPARATOR: Given $(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$, where G is a directed graph, $[S_1, \ldots, S_l]$ is a collection of l source sets and $[T_1, \ldots, T_l]$ is a collection of l sink sets in G, and a parameter k, such that
>
> (1) all sets $S_1, \ldots, S_l, T_1, \ldots, T_l$ are pairwise disjoint;
> (2) for each i, $1 \le i \le l - 1$, there is no edge coming into the set S_i; and
> (3) for each j, $1 \le j \le l$, there is no edge going out from the set T_j,
>
> either construct a skew separator of at most k vertices for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$, or report that no such separator exists.

Note that in an instance of the SKEW SEPARATOR problem, condition (2) on source sets and condition (3) on sink sets are not completely symmetric. Although the first $l - 1$ source sets are not allowed to have incoming edges, the last source set S_l is allowed to have incoming edges. On the other hand, *all* sink sets are not allowed to have outgoing edges.

We remark that conditions (1)–(3) in the definition of the SKEW SEPARATOR problem (plus the restriction that the skew separator can consist of only non-terminal vertices) may be relaxed, and our techniques for the problem may still be applicable. However, the above formulation of the problem will make our discussion simpler, and will also be sufficient for our solution to the DFVS problem, which is the focus of the current paper. We leave the investigation of the separator problems of more general forms to later research.

Let $G = (V, E)$ be a directed graph, and let (D_1, D_2) be a bi-partition of the vertex set V of G, that is, $D_1 \cup D_2 = V$ and $D_1 \cap D_2 = \emptyset$. The bi-partition (D_1, D_2) is a *DAG-bipartition* for the graph G if both induced subgraphs $G[D_1]$ and $G[D_2]$ are DAGs. A vertex subset F in the graph G is a D_1-*FVS* if F is an FVS for G and $F \subseteq D_1$.

> DAG-BIPARTITION FVS: Given (G, D_1, D_2, k), where G is a directed graph, (D_1, D_2) is a DAG-bipartition for G, and k is the parameter, either construct a D_1-FVS of size bounded by k for the graph G, or report that no such D_1-FVS exists.

We will be also interested in a special version of the FEEDBACK VERTEX SET problem.

DFVS REDUCTION: Given a triple (G, F, k), where G is a directed graph and F is an FVS of size $k + 1$ for G, either construct an FVS of size bounded by k for G, or report that no such FVS exists.

Finally, our central problem in this article is as follows.

DFVS: Given a pair (G, k), where G is a directed graph and k is the parameter, either construct an FVS of size bounded by k for G, or report that no such FVS exists.

3. Solving the SKEW SEPARATOR Problem

In this section, we study the complexity of the SKEW SEPARATOR problem.

Let $(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$ be an instance of the SKEW SEPARATOR problem. Define $T_{all} = \bigcup_{1 \le i \le l} T_i$. There are a few cases in which we can directly reduce the instance size:

Rule R1. There is no path from S_l to T_{all}, i.e., the size of a min-cut from S_l to T_{all} is 0: then we only need to find a skew separator of size k that separates S_i from T_j for all indices i and j satisfying $l - 1 \ge i \ge j \ge 1$, i.e., we can work instead on the instance $(G, [S_1, \ldots, S_{l-1}], [T_1, \ldots, T_{l-1}], k)$. Note that in this case, by definition, if $l = 1$, then the solution to the instance $(G, [S_1, \ldots, S_{l-1}], [T_1, \ldots, T_{l-1}], k)$ is simply the empty set $\emptyset$;

Rule R2. There is an edge from S_l to T_{all}: then there is no way to even separate S_l from T_{all} – we can simply stop and claim that the given instance is a "No" instance;

Rule R3. There exists a non-terminal vertex w, an edge from S_l to w, and an edge from w to T_{all}: then the vertex w must be included in the skew separator in order to separate S_l and T_{all} – we can simply work on the instance $(G - w, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k - 1)$ and recursively find a skew separator of size $k - 1$.

Note in Rules R1 and R3, the reduced instances $(G, [S_1, \ldots, S_{l-1}], [T_1, \ldots, T_{l-1}], k)$ and $(G - w, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k - 1)$ are still valid instances of the SKEW SEPARATOR problem.

In the following, we assume that for the instance $(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$, none of the rules above is applicable. In particular, since Rule R1 is not applicable, a min-cut from S_l to T_{all} has size larger than 0. Because Rules R1–R3 are not applicable, there must be a nonterminal vertex u_0 such that (1) there is an edge from S_l to u_0; and (2) there is no edge from u_0 to T_{all}. Such a vertex u_0 will be called an S_l-extended vertex. Fix an S_l-extended vertex u_0, let $S'_l = S_l \cup \{u_0\}$.

We start with the following simple but important lemma. The proof of this lemma is straightforward. Thus, we leave it to the reader.

LEMMA 3.1. *Let X be a subset of vertices in the graph G that does not contain the S_l-extended vertex u_0. Then X is a skew separator for $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$ if and only if X is a skew separator for $([S_1, \ldots, S_{l-1}, S'_l], [T_1, \ldots, T_{l-1}, T_l])$.*

Lemma 3.1 also directly implies the following two useful corollaries.

COROLLARY 3.2. *A skew separator for $([S_1, \ldots, S_{l-1}, S'_l], [T_1, \ldots, T_{l-1}, T_l])$ is also a skew separator for $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$.*

COROLLARY 3.3. *The size of a min-cut from S_l' to T_{all} in the graph G is at least as large as the size of a min-cut from S_l to T_{all} in G.*

Now we are ready for our main theorem in this section.

THEOREM 3.4. *If the size of a min-cut from S_l to T_{all} is equal to the size of a min-cut from S_l' to T_{all}, then the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$ has a skew separator of size bounded by k if and only if the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$ has a skew separator of size bounded by k.*

PROOF

$\Leftarrow$: Suppose that the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$ has a skew separator X' of size bounded by k. By Corollary 3.2, X' is also a skew separator for $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$. In consequence, $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$ has a skew separator of size bounded by k.

$\Rightarrow$: Suppose that the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$ has a skew separator X of size bounded by k. If the skew separator X does not contain the S_l-extended vertex u_0, then by Lemma 3.1, X is also a skew separator of size bounded by k for the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$, and the theorem is proved. Therefore, we can assume that the set X contains the S_l-extended vertex u_0. We will define another set X' that does not contain u_0. We will show that $|X'| \leq |X|$ and that X' is a skew separator for $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$. Then the theorem will immediately follow.

Let Y be a min-cut from S_l' to T_{all}. Then Y does not contain the S_l-extended vertex u_0. Moreover, since there is no edge coming into S_i from outside of S_i for all $i \leq l - 1$, the set Y does not contain any vertex in $\bigcup_{i=1}^{l-1} S_i$. In consequence, the set Y consists of only nonterminal vertices. By Corollary 3.2, Y is also a cut from S_l to T_{all}. Moreover, by the assumption of the theorem that the size of a min-cut from S_l to T_{all} is equal to the size of a min-cut from S_l' to T_{all}, Y is actually also a min-cut from S_l to T_{all}. Let $R_Y(S_l)$ be the set of vertices v such that either $v \in S_l$ or there is a path from S_l to v in the subgraph $G - Y$. In particular, $u_0 \in R_Y(S_l)$ because Y does not contain u_0 and there is an edge from S_l to u_0.

We introduce a number of sets as follows.

$$Z = X \cap Y;$$
$$X_{in} = X \cap R_Y(S_l);$$
$$X_{out} = X - (X_{in} \cup Z).$$

That is, the skew separator X for $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$ is decomposed into three disjoint subsets Z, X_{in}, and X_{out} (note that by definitions, $R_Y(S_l)$ and Y do not intersect).

Let Y_T be the set of vertices v in the min-cut Y such that there is a path from v to T_{all} in the subgraph $G - X$. By definition, we have $Y_T \cap Z = \emptyset$. Let

$$Y_S = Y - (Y_T \cup Z).$$

Thus, the min-cut Y from S_l to T_{all} is decomposed into three disjoint subsets Z, Y_T, and Y_S. Figure 1 gives an intuitive illustration of the sets Z, X_{in}, X_{out}, Y_T, Y_S, and $R_Y(S_l)$.

We first show that the set $Y' = Y_S \cup Z \cup X_{in}$ is also a cut from S_l to T_{all}. If by contradiction Y' is not a cut from S_l to T_{all}, then there is a path P_1 from S_l to T_{all} in

Y: a min-cut from S_l' to T_{all}

X: a skew separator for
$([S_1, \ldots, S_l], [T_1, \ldots, T_l])$

$R_Y(S_l)$: vertices reachable
from S_l in $G - Y$

$Z = X \cap Y$

$X_{in} = X \cap R_Y(S_l)$

$X_{out} = X - (X_{in} \cup Z)$

Y_T: vertices in Y from which
T_{all} is reachable in $G - X$

$Y_S = Y - (Y_T \cup Z)$

FIG. 1.   Sets in the proof of Theorem 3.4.

the subgraph $G - Y'$. The path P_1 must contain vertices in the set Y since Y is a cut from S_l to T_{all}. Let w be the first vertex on the path P_1 that is in Y when we traverse from S_l to T_{all} along the path P_1. Then, w must be in Y_T since Y' contains both Y_S and Z. Now the partial path P_1' of P_1 from S_l to w (not including w) must be entirely contained in $R_Y(S_l)$ (note that the path P_1 does not intersect $Y_S \cup Z$). Moreover, the path P_1' contains neither vertices in $X_{in} \cup Z$ (by the definition of the set Y') nor vertices in X_{out} (since the sets X_{out} and $R_Y(S_l)$ are disjoint). In summary, the subpath P_1' from S_l to w contains no vertex in the set X. Moreover, by the definition of the set Y_T, and $w \in Y_T$, there is a path P_1'' from w to T_{all} in the subgraph $G - X$. Now the concatenation of the paths P_1' and P_1'' would result in a path from S_l to T_{all} in the graph $G - X$, contradicting the fact that X is a skew separator for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$. This contradiction shows that the set Y' must be a cut from S_l to T_{all}.

Since Y is a min-cut from S_l to T_{all}, we have $|Y| \leq |Y'|$. By definition, $Y = Y_S \cup Z \cup Y_T$ and $Y' = Y_S \cup Z \cup X_{in}$. Also note that Y_S, Z, and Y_T are pairwise disjoint, and that Y_S, Z, and X_{in} are also pairwise disjoint. Therefore, we must have $|Y_T| \leq |X_{in}|$.

Consider the set $X' = X_{out} \cup Z \cup Y_T$. The set X' has the following properties: (1) X' consists of only non-terminal vertices (because both X and Y consist of only non-terminal vertices); (2) $|X'| \leq |X|$ (because $|Y_T| \leq |X_{in}|$), so the size of X' is bounded by k; and (3) the set X' does not contain the S_l-extended vertex u_0 (this is because u_0 is in X_{in} and Y does not contain u_0). Therefore, if we can prove that X' is a skew separator for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$, then by Lemma 3.1, X' is also a skew separator of size bounded by k for the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$. This will complete the proof of the theorem.

Therefore, what remains is to prove that the set $X' = X_{out} \cup Z \cup Y_T$ is a skew separator for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$. Let $R_Y(T_{all})$ be the set of vertices v such that either $v \in T_{all}$, or there is a path from v to T_{all} in the subgraph $G - Y$.

Suppose by contradiction that the set X' is not a skew separator for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$. Then there is a path P_2 in the subgraph $G - X'$ from S_i to T_j for some $i \geq j$. The path P_2 has the following properties:

(1) The path P_2 must contain a vertex in $R_Y(S_l)$: since X is a skew separator for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$, the path P_2 from S_i to T_j with $i \geq j$ must contain at least one vertex w_1 in $X = X_{in} \cup Z \cup X_{out}$. Now since the path P_2 is in the subgraph $G - X'$, where $X' = X_{out} \cup Z \cup Y_T$, the vertex w_1 must be in X_{in}, which is a subset of $R_Y(S_l)$;

(2) The path P_2 must contain a vertex in Y_S: by Property (1), P_2 contains a vertex w_1 in $R_Y(S_l)$. From the vertex w_1 to T_{all} along the path P_2, there must be a vertex w_2 in $Y = Y_S \cup Z \cup Y_T$ since Y is a cut from S_l to T_{all} while w_1 is reachable from S_l in the subgraph $G - Y$. Now since $X' = X_{out} \cup Z \cup Y_T$, and the path P_2 is in the subgraph $G - X'$, the vertex w_2 on the path P_2 must be in the set Y_S;

(3) The path P_2 must end at a vertex in $R_Y(T_{all})$; this is simply because P_2 is ended in T_{all}. Note that by definition, no vertex in Y_S can be in $R_Y(T_{all})$.

By Properties (2) and (3), the path P_2 contains a vertex not in $R_Y(T_{all})$ and ends at a vertex in $R_Y(T_{all})$. Thus, there must be an internal vertex w in the path such that w is not in $R_Y(T_{all})$ but all vertices after w along the path P_2 (from S_i to T_j) are in $R_Y(T_{all})$. Note that no vertex w' after the vertex w along the path P_2 can be in the set X: w' in X would imply w' in X_{in} (since P_2 is a path in the subgraph $G - X'$), which would imply that there is another vertex after w' that is in Y thus is not in $R_Y(T_{all})$. Moreover, the vertex w must be in the set Y (otherwise, w would be in $R_Y(T_{all})$). Since P_2 is a path in $G - X'$ and $X' = X_{out} \cup Z \cup Y_T$, the vertex w must be in the set Y_S. However, this derives a contradiction: the subpath of P_2 from w to T_{all} shows that the vertex w should belong to the set Y_T (note that all vertices after w on the path are not in X), and the sets Y_S and Y_T are disjoint. This contradiction proves that the set X' must be a skew separator for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$. Since the size of the set X' is bounded by k and X' does not contain the S_l-extended vertex u_0, by Lemma 3.1, the set X' is also a skew separator for the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$, and the size of X' is bounded by k.

This completes the proof of the theorem. □

Theorem 3.4 enables us to develop a parameterized algorithm for the SKEW SEPARATOR problem. The algorithm is presented in Figure 2.

THEOREM 3.5. *The algorithm* **SMC**$(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$ *solves the problem* SKEW SEPARATOR *in time* $O(4^k k n^3)$*, where n is the number of vertices in the input graph G.*

PROOF. For the correctness of the algorithm, let $(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$ be an input to the algorithm, which is an instance of the SKEW SEPARATOR problem, where $G = (V, E)$ is a directed graph, $[S_1, \ldots, S_l]$ and $[T_1, \ldots, T_l]$ are the source sets and the sink sets, respectively, and k is the upper bound of the size of the skew separator we are looking for.

If $l = 1$, then the problem becomes the construction of a min-cut of size bounded by k from S_1 to T_1, which can be solved in $O(k n^2)$ time by Lemma 2.2. Steps 2–4 were justified in the discussions of Rules R2, R1, R3, respectively, at the

Algorithm SMC$(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$
input: an instance $(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$ of the SKEW SEPARATOR problem.
output: a skew separator of size bounded by k for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$,
       or report "No" (i.e., no such separator exists).

1.   **if** $l = 1$ **then** solve the problem in time $O(kn^2)$;
2.   **if** Rule R2 applies or $k < 0$ **then** return "No";
3.   **if** Rule R1 applies **then** return **SMC**$(G, [S_1, \ldots, S_{l-1}], [T_1, \ldots, T_{l-1}], k)$;
4.   **if** Rule R3 applies on a vertex w
       then return $\{w\} \cup$ **SMC**$(G - w, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k - 1)$;§
5.   pick an S_l-extended vertex u_0;   let $S_l' = S_l \cup \{u_0\}$;
6.   let m be the size of a min-cut from S_l to $T_{all} = \bigcup_{i=1}^{l} T_i$;
7.   **if** $m > k$ **then** return "No";
8.   let m' be the size of a min-cut from S_l' to T_{all};
9.   **if** $(m = m')$
9.1. **then** return **SMC**$(G, [S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l], k)$;
9.2. **else** $X = \{u_0\} \cup$ **SMC**$(G - u_0, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k - 1)$;
       if $X \neq$ "No" **then** return X;
9.3.   **else** return **SMC**$(G, [S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l], k)$.

§To simplify our description, we assume that a "No" plus anything gives a "No".

FIG. 2.   An algorithm for the SKEW SEPARATOR problem.

beginning of this section (note that we have also consistently defined that an instance is a "No" instance if the parameter k has a negative value). Therefore, if the algorithm reaches step 5, then none of the Rules 1-3 are applicable. In particular, since Rule 1 is not applicable and the sets S_l and T_{all} are disjoint, there must be an edge $[v, w]$, where $v \in S_l$ and $w \notin S_l$. Since Rule 2 is not applicable, the vertex w is not in the set T_{all}. The vertex w also cannot be in any source set S_i for $i < l$ because there is no edge coming into S_i from outside of S_i. Therefore, the vertex w is a non-terminal vertex. Finally, since Rule 3 is not applicable, there is no edge from w to T_{all}. Thus, w must be an S_l-extended vertex. This proves that at step 5, the algorithm can always find an S_l-extended vertex u_0.

In the case $m > k$ in step 7, that is, the size m of a min-cut from S_l to T_{all} is larger than the parameter k, then even separating a single source set S_l from the sink sets $T_{all} = \bigcup_{j=1}^{l} T_j$ requires more than k vertices. Thus, no skew separator of size bounded by k can exist to separate S_i from T_j for all $l \geq i \geq j \geq 1$. Step 7 correctly handles this case by returning "No".

In the case $m = m'$ in step 9, that is, the size m of a min-cut from S_l to T_{all} is equal to the size m' of a min-cut from S_l' to T_{all}, by Theorem 3.4, the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$ has a skew separator of size bounded by k if and only if the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$ has a skew separator of size bounded by k. Moreover, by Corollary 3.2, a skew separator of size bounded by k for the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$ is also a skew separator for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$. Therefore, in this case, we can recursively call the algorithm **SMC**$(G, [S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l], k)$, and look instead for a skew separator of size bounded by k for the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$, as handled by step 9.1.

In the case $m \neq m'$, then the algorithm branches into two subcases: step 9.2 includes the S_l-extended vertex u_0 in the skew separator and recursively looks for a skew separator of size bounded by $k - 1$ in the remaining graph $G - u_0$ for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$; and step 9.3 excludes the S_l-extended vertex u_0 from the

skew separator and recursively looks for a skew separator that does not contain u_0 and is of size bounded by k in the graph G for the pair $([S_1, \ldots, S_l], [T_1, \ldots, T_l])$ (which, by Lemma 3.1, is a skew separator of size bounded by k for the pair $([S_1, \ldots, S_{l-1}, S_l'], [T_1, \ldots, T_{l-1}, T_l])$). This completes the verification of the correctness of the algorithm. Now we analyze its complexity.

The recursive execution of the algorithm can be described as a search tree $\mathcal{T}$. We first count the number of leaves in the search tree $\mathcal{T}$. Note that only steps 9.2–9.3 of the algorithm correspond to branches in the search tree $\mathcal{T}$. Let $D(k, m)$ be the total number of leaves in the search tree $\mathcal{T}$ for the algorithm $\mathbf{SMC}(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$, where m is the size of a min-cut from S_l to T_{all}. Then steps 9.2–9.3 induce the following recurrence relation:

$$D(k, m) \le D(k - 1, m_1) + D(k, m_2), \qquad (1)$$

where m_1 is the size of a min-cut from S_l to T_{all} in the graph $G - u_0$ as given in step 9.2, and m_2 is the size of a min-cut from S_l' to T_{all} in the graph G as given in step 9.3. Note that $m - 1 \le m_1 \le m$ because removing the vertex u_0 from the graph G cannot increase the size of a min-cut from S_l to T_{all}, and can decrease the size of a min-cut for the two sets by at most 1. Moreover, by Corollary 3.3, in step 9.3 we must have $m_2 \ge m + 1$. Summarizing these, we have

$$m - 1 \le m_1 \le m \quad \text{and} \quad m_2 \ge m + 1. \qquad (2)$$

We prove, by induction on $t = 2k - m$, that $D(k, m) \le 2^{2k-m}$. First, note that we always have $t = 2k - m \ge 0$ because by the definitions of k and m we always have $k \ge m \ge 0$. In particular, in the initial case when $t = 2k - m = 0$, we must have $k = m = 0$; in this case, the algorithm can solve the instance without further branching. Therefore, we have $D(k, m) = 1$ when $t = 2k - m = 0$. For the inductive step, note that by Inequalities (2), we have

$$t_1 = 2(k - 1) - m_1 \le 2(k - 1) - (m - 1) = 2k - m - 1,$$

and

$$t_2 = 2k - m_2 \le 2k - (m + 1) = 2k - m - 1.$$

Therefore, we can apply the inductive hypothesis on Inequality (1), which gives

$$\begin{aligned} D(k, m) &\le D(k - 1, m_1) + D(k, m_2) \\ &\le 2^{2(k-1)-m_1} + 2^{2k-m_2} \\ &\le 2^{2k-m-1} + 2^{2k-m-1} \\ &= 2^{2k-m}. \end{aligned} \qquad (3)$$

This completes the inductive proof. Moreover, we also note that certain non-branching steps (i.e., steps 3, 4, and 9.1) may also change the values of k and m, thus changing the value $t = 2k - m$. However, none of these steps increases the value $t = 2k - m$: (i) step 3 keeps the value k unchanged and does not decrease the value m (because in this case the size of a min-cut from S_l to T_{all} is 0 that cannot be larger than the size of a min-cut from S_{l-1} to $\bigcup_{j=1}^{l-1} T_j$); (ii) step 4 decreases the value k by 1 and the value m by at most 1 (because removing a vertex from G can reduce the size of a min-cut from S_l to T_{all} by at most 1), which as a total will decrease the value $t = 2k - m$ by at least 1; (iii) by the condition assumed, step 9.1 keeps both the values k and m unchanged, thus unchanging the value $t = 2k - m$.

As a result, the value $t = 2k - m$ after a branching step to the next branching step can never be increased.

Summarizing the above discussion, we conclude that the total number of leaves, $D(k, m)$, in the search tree $\mathcal{T}$ for the algorithm **SMC**$(G, [S_1, \ldots, S_l], [T_1, \ldots, T_l], k)$, where m is the size of a min-cut from S_l to T_{all}, satisfies the following inequality

$$D(k, m) \leq 2^{2k-m} \leq 4^k.$$

The running time of each execution of the algorithm **SMC**, not counting the time for the recursive calls in the execution, is bounded by $O(kn^2)$, where n is the number of vertices in the input graph. In particular, by Lemma 2.2, step 1 that looks for a min-cut of size bounded by k from S_1 to T_1, steps 6–7 that determine if the size m of a min-cut from S_l to T_{all} is bounded by k, and steps 8–9 that determine if the size of a min-cut from S'_l to T_{all} is equal to m ($m \leq k$ at this point), all have their running time bounded by $O(kn^2)$.

Observe that for each recursive call in an execution of the algorithm **SMC**, either the number of source-sink pairs in the instance is decreased by 1 (step 3), or the number of non-terminal vertices in the instance is decreased by 1 (steps 4, 9.1, 9.2, and 9.3). When the number of source-sink pairs is equal to 1, the problem is solved in time $O(kn^2)$ by step 1, and when the number of nonterminal vertices is equal to 0, either step 2 or step 3 can be applied directly. In conclusion, along each root-leaf path in the search tree $\mathcal{T}$, there are at most $O(n)$ recursive calls to the algorithm **SMC**. Therefore, the running time along each root-leaf path in the search tree $\mathcal{T}$ is bounded by $O(kn^3)$.

Summarizing the above discussions, we conclude that the running time of the algorithm **SMC** is bounded by $O(4^k kn^3)$. This completes the proof of the theorem. □

4. Solving the DAG-BIPARTITION FVS Problem

In this section, we describe how to use the results in the previous section to solve the DAG-BIPARTITION FVS problem.

Recall that an instance of DAG-BIPARTITION FVS is given as a tuple (G, D_1, D_2, k), where G is a directed graph, (D_1, D_2) is a DAG-bipartition of G, and k is the parameter, with the objective of finding an FVS X for the graph G such that $X \subseteq D_1$ (recall that such an FVS is called a D_1-FVS) and that the size of X is bounded by k.

Let $\pi = \{v_1, v_2, \ldots, v_h\}$ be a topologically sorted order of the vertices in the induced DAG $G[D_2]$. We construct an instance of the SKEW SEPARATOR problem as follows:

(1) Let G' be the graph obtained from G by removing all edges in $G[D_2]$.
(2) In the graph G', replace each vertex v_i in D_2 by a pair (t_i, s_i) of vertices such that all incoming edges into v_i are now coming into the vertex t_i, and that all outgoing edges from v_i are now going out from the vertex s_i. Let the resulting graph be G_π.

Note that in the resulting graph G_π, the vertices s_i, $1 \leq i \leq h$, have no incoming edges, and the vertices t_j, $1 \leq j \leq h$, have no outgoing edges. Moreover, since we have removed all edges between the vertices in $G[D_2]$, every edge going out from a vertex s_i must come into a vertex in the set D_1, and every edge coming into a vertex t_j must go out from a vertex in the set D_1. In particular,

$(G_\pi, [\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}], k)$ is a valid instance for the SKEW SEPARATOR problem, which will be called an instance of the SKEW SEPARATOR *induced by the instance* (G, D_1, D_2, k) *of* DAG-BIPARTITION FVS *and the topologically sorted order* π *of the vertices in* $G[D_2]$.

Thus, each vertex v_i in the set D_2 in the graph G is now "split" into the two vertices s_i and t_i in the graph G_π. Moreover, there is a one-to-one mapping between the vertices in the set D_1 in the graph G and the non-terminal vertices in the graph G_π. Thus, in case of no ambiguity, we will use the same vertex name to refer to both a non-terminal vertex in the graph G_π and a vertex in the set D_1 in the graph G. In particular, a skew separator for the pair $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ in the graph G_π corresponds to a subset of D_1 in the graph G. We have the following important theorem.

THEOREM 4.1. *Let* (G, D_1, D_2, k) *be an instance of the* DAG-BIPARTITION FVS *problem, and let* X *be a* D_1*-FVS for the graph* G. *Then there is a topologically sorted order* $\pi = \{v_1, \ldots, v_h\}$ *of the vertices in* $G[D_2]$ *such that in the instance* $(G_\pi, [\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}], k)$ *induced by* (G, D_1, D_2, k) *and* π: (1) X *is a skew separator for the pair* $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ *in the graph* G_π; *and* (2) *every skew separator for the pair* $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ *in* G_π *is a* D_1*-FVS for the graph* G.

PROOF. As assumed in the theorem, let (G, D_1, D_2, k) be an instance of the DAG-BIPARTITION FVS problem, and let X be a D_1-FVS for the graph G. Consider the subgraph $G - X$. Since X is an FVS for G, the graph $G - X$ is a DAG. Therefore, the vertices in $G - X$ can be topologically sorted into an ordered list π' such that there is no edge in $G - X$ that goes out from a later vertex in π' and comes into an earlier vertex in π'. Let $\pi = \{v_1, \ldots, v_h\}$ be the order of the vertices in D_2 that is induced from the order π' (i.e., π is obtained from π' by removing the vertices not in D_2. Note that all vertices in X are in D_1). The order π is obviously a topologically sorted order for the DAG $G[D_2]$. We show that this order π of the vertices in D_2 and the corresponding instance $(G_\pi, [\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}], k)$ induced by (G, D_1, D_2, k) and π satisfy the conclusions of the theorem.

We first show that the set X is a skew separator for $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ in the graph G_π. If this were not the case, then there would be a path P in the graph $G_\pi - X$ that starts from a vertex s_i and ends at a vertex t_j with $i \geq j$. Since no vertex in $\{s_1, \ldots, s_h\}$ has incoming edges and no vertex in $\{t_1, \ldots, t_h\}$ has outgoing edges, all internal vertices on the path P are nonterminal vertices in G_π. In consequence, all internal vertices on P are vertices in the set D_1 in the graph G. Therefore, the path P in $G_\pi - X$ corresponds to a path P' in the graph $G - X$ that starts from the vertex v_i and ends at the vertex v_j, where $i \geq j$, with all internal vertices of P' in the set D_1. But it is impossible: (1) if $i = j$ then the path P' would be a cycle in the graph $G - X$, contradicting the assumption that X is an FVS for the graph G; and (2) if $i > j$, then P' would become a path from v_i to v_j with $i > j$ in the graph $G - X$, contradicting the assumption that $\pi = \{v_1, \ldots, v_h\}$ is an order of the vertices in D_2 that is induced from the topologically sorted order π' of the vertices in the DAG $G - X$. In conclusion, the path P does not exist, and the set X is a skew separator for the pair $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ in the graph G_π.

Now we prove that every skew separator X' for $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ in the graph G_π is a D_1-FVS for the graph G. First of all, by definition, a skew separator consists of only nonterminals, thus, all vertices in X' are in the set D_1.

Algorithm DBF(G, D_1, D_2, k)
input: an instance (G, D_1, D_2, k) of the DAG-BIPARTITION FVS problem.
output: a D_1-FVS of size bounded by k for G, or report "No" (i.e., no such
$\qquad D_1$-FVS exists).

1.   **for** each topologically sorted order $\pi = \{v_1, \ldots, v_h\}$ of the vertices in $G[D_2]$ **do**
1.1.   construct the instance $(G_\pi, [\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}], k)$ of the SKEW
$\qquad$ SEPARATOR problem induced by (G, D_1, D_2, k) and π;
1.2.   let $X = \mathbf{SMC}(G_\pi, [\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}], k)$;
1.3.   **if** X is a D_1-FVS of size bounded by k for G
$\qquad$ **then** return(X); stop;
2.   return("No").

FIG. 3.   An algorithm for the DAG-BIPARTITION FVS problem.

Suppose for a contradiction that X' is not a D_1-FVS for the graph G. Then there is a cycle C in the graph $G - X'$. Without loss of generality, we can assume that C is a simple cycle. Since both the induced subgraphs $G[D_1]$ and $G[D_2]$ are DAGs, the cycle C must contain both vertices in D_1 and vertices in D_2. We consider two different cases.

Case 1. The cycle C contains a single vertex v_i in the set D_2. Then, all other vertices in the cycle C are in the set D_1. But then the cycle C would correspond to a path P_1 in the graph $G_\pi - X'$ that starts with the vertex s_i and ends at the vertex t_i (with all internal vertices being non-terminal vertices). But this contradicts the assumption that X' is a skew separator for the pair $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ that should have cut all paths from s_i to t_i.

Case 2. The cycle C contains more than one vertex in D_2. Let $\{v_{i_1}, v_{i_2}, \ldots, v_{i_d}, v_{i_1}\}$ be the order of the vertices in D_2 that we encounter when traversing along the cycle C (starting from an arbitrary vertex v_{i_1} in D_2), where $d > 1$. Then there must be an index j such that $i_j > i_{j+1}$ (where we take $i_{j+1} = i_1$ if $j = d$). Now consider the subpath P_2 of C that starts from the vertex v_{i_j} and ends at the vertex $v_{i_{j+1}}$. The path P_2 cannot be a single edge from v_{i_j} to $v_{i_{j+1}}$ since $\pi = \{v_1, v_2, \ldots, v_h\}$ is a topologically sorted order for the vertices in the DAG $G[D_2]$ and $i_j > i_{j+1}$. Thus, the path P_2 contains at least one internal vertex. Since all internal vertices on the path P_2 are not in D_2 thus correspond to non-terminal vertices in the graph $G_\pi - X'$, the path P_2 would correspond to a path P_2' in the graph $G_\pi - X'$ that starts from the vertex s_{i_j} and ends at the vertex $t_{i_{j+1}}$, with $i_j > i_{j+1}$. Again this contradicts the assumption that X' is a skew separator for the pair $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$, which should have cut all paths from s_{i_j} to $t_{i_{j+1}}$ when $i_j > i_{j+1}$.

This proves that the skew separator X' for $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ in the graph G_π must be a D_1-FVS for the graph G. This completes the proof of the theorem. $\square$

Theorem 4.1 enables us to reduce the DAG-BIPARTITION FVS problem to the SKEW SEPARATOR problem. An algorithm for the DAG-BIPARTITION FVS problem is given in Figure 3.

THEOREM 4.2. *The algorithm* **DBF**(G, D_1, D_2, k) *solves the* DAG-BIPARTITION FVS *problem in time* $O(4^k k n^3 h!)$, *where h is the number of vertices in the set D_2, and n is the number of vertices in the input graph G.*

PROOF.   The running time of the algorithm is obvious: the **for**-loop in step 1 is executed at most $h!$ times, and the time for each execution is dominated by the

subroutine call to the algorithm **SMC** in step 1.2. By Theorem 3.5, the running time of each execution of step 1.2 is bounded by $O(4^k k n^3)$.

For the correctness of the algorithm, first note that the algorithm always returns "No" unless it actually constructs a D_1-FVS of size bounded by k for G in step 1.3. In particular, if the input instance (G, D_1, D_2, k) contains no D_1-FVS of size bounded by k for the graph G, then the algorithm always correctly reports "No".

On the other hand, suppose that there is a D_1-FVS X_0 of size bounded by k for the graph G. Then by Theorem 4.1, there is a topologically sorted order $\pi = \{v_1, \ldots, v_h\}$ of the vertices in the DAG $G[D_2]$ such that in the instance $(G_\pi, [\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}], k)$ of the SKEW SEPARATOR problem induced by (G, D_1, D_2, k) and π, X_0 is a skew separator for $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ in the graph G_π, and every skew separator for $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ in G_π is a D_1-FVS for the graph G. In particular, $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ has at least one skew separator of size bounded by k (e.g., X_0) in the graph G_π. Therefore, step 1.2 of the algorithm **DBF** must return a skew separator X of size bounded by k for the pair $([\{s_1\}, \ldots, \{s_h\}], [\{t_1\}, \ldots, \{t_h\}])$ in the graph G_π (the set X may be different from the set X_0), and this set X is a D_1-FVS for the graph G. In conclusion, if there is a D_1-FVS of size bounded by k for the graph G, then the algorithm **DBF**(G, D_1, D_2, k) will correctly return a D_1-FVS of size bounded by k in step 1.3. □

5. Solving the DFVS Problem

We present our algorithm for the DFVS problem. We start with a more restricted version of the problem, the DFVS REDUCTION problem, defined as follows.

> DFVS REDUCTION: Given a triple (G, F, k), where G is a directed graph and F is an FVS of size $k + 1$ for G, either construct an FVS of size bounded by k for G, or report that no such FVS exists.

LEMMA 5.1. *The* DFVS REDUCTION *problem on a triple* (G, F, k) *is solvable in time* $O(n^3 4^k k^3 k!)$, *where n is the number of vertices in the input graph G.*

PROOF. Let $G = (V, E)$ be the input directed graph with $n = |V|$ vertices, and let F be the input FVS of size $k+1$ for the graph G. Every FVS F' of size bounded by k for G can be split into two disjoint subsets F_1 and F_2, where F_2 consists of j vertices in F for some integer j, $0 \le j \le k$, and F_1 consists of at most $k - j$ vertices in $V - F$. Note that since we assume that no vertex in $F - F_2$ is in the FVS F', the induced subgraph $G[F - F_2]$ must be a DAG. Therefore, for each j, $0 \le j \le k$, we enumerate all subsets of j vertices in F. For each such subset F_2 of F such that $G[F - F_2]$ is a DAG, we seek a subset F_1 of at most $k - j$ vertices in $V - F$ such that $F_1 \cup F_2$ makes an FVS for the graph G.

Fix a subset F_2 of F, such that $|F_2| = j$ and that the induced subgraph $G[F - F_2]$ is a DAG. Note that the graph G has an FVS $F_1 \cup F_2$ of size bounded by k, where $F_1 \subseteq V - F$, if and only if the subset F_1 of $V - F$ is an FVS for the graph $G - F_2$ and the size of F_1 is bounded by $k - j$. Therefore, to solve the original problem, we can instead consider how to construct an FVS F_1 for the graph $G - F_2$ such that $|F_1| \le k - j$ and $F_1 \subseteq V - F$.

Since F is an FVS for G, we have that the induced subgraph $G[V - F] = G - F$ is a DAG. Moreover, by our assumption, the induced subgraph $G[F - F_2]$ is also

a DAG. Note that $(V - F) \cup (F - F_2) = V - F_2$, which is the vertex set for the graph $G' = G - F_2$. Therefore, $(V - F, F - F_2)$ is a DAG-bipartition of the graph G'. Thus, an FVS F_1 for the graph G' such that $|F_1| \leq k - j$ and $F_1 \subseteq V - F$, is actually a $(V - F)$-FVS of size bounded by $k - j$ for the graph G' with the DAG-bipartition $(V - F, F - F_2)$. Therefore, the set F_1 can be constructed by the algorithm **DBF**$(G', V - F, F - F_2, k - j)$.

Since $|F| = k + 1$ and $|F_2| = j$, we have $|F - F_2| = k + 1 - j$. Therefore, the DAG $G[F - F_2]$ contains exactly $k + 1 - j$ vertices. By Theorem 4.2, the running time of the algorithm **DBF**$(G', V - F, F - F_2, k - j)$ is bounded by $O(4^{k-j}(k - j)n^3(k + 1 - j)!)$. Now for all integers $j, 0 \leq j \leq k$, we enumerate all subsets F_2 of j vertices in F and apply the algorithm **DBF**$(G', V - F, F - F_2, k - j)$ for those F_2 such that $G[F - F_2]$ is a DAG. As we discussed above, the graph G has an FVS of size bounded by k if and only if for some F_2 of j vertices in F, where $0 \leq j \leq k$, the algorithm **DBF**$(G', V - F, F - F_2, k - j)$ produces an FVS F_1 of size bounded by $k - j$ for the graph G'. The running time of this process is bounded by the order of

$$\sum_{j=0}^{k} \binom{k + 1}{j} \left(4^{k-j}(k - j)n^3(k + 1 - j)!\right) = O(n^3 4^k k^3 k!).$$

This completes the proof of the lemma. $\square$

The rest of our process for solving the original DFVS problem is to apply the *iterative compression* method. The method was proposed by Reed et al. [2004] and has been used for solving the FEEDBACK VERTEX SET problem on undirected graphs [Dehne et al. 2007; Guo et al. 2005, 2006]. Here we extend the method and apply it to solve the DFVS problem.

THEOREM 5.2.   *The* DFVS *problem is solvable in time* $O(n^4 4^k k^3 k!)$.

PROOF.   Let (G, k) be an instance of the DFVS problem, where $G = (V, E)$ is a directed graph with $n = |V|$ vertices, and k is the parameter. Pick any subset V_0 of $k + 1$ vertices in G, and let F_0 be any subset of k vertices in V_0. Note that the set F_0 is an FVS of k vertices for the induced subgraph $G_0 = G[V_0]$ since the graph $G_0 - F_0$ consists of a single vertex (note that by our assumption, the graph G contains no self-loops).

Let $V - V_0 = \{v_1, v_2, \ldots, v_{n-k-1}\}$. Let $V_i = V_0 \cup \{v_1, \ldots, v_i\}$, and let $G_i = G[V_i]$ be the subgraph induced by V_i, for $i = 0, 1, \ldots, n - k - 1$. Inductively, suppose that for an integer $i, 0 \leq i < n - k - 1$, we have constructed an FVS F_i of size bounded by k for the induced subgraph G_i (this has been the case for $i = 0$). Without loss of generality, we can assume that the set F_i consists of exactly k vertices – otherwise we simply pick $k - |F_i|$ vertices (arbitrarily) from $G_i - F_i$ and add them to the set F_i. Now consider the set $F_{i+1}' = F_i + v_{i+1}$. Since $G_{i+1} - F_{i+1}' = G_i - F_i$ and F_i is an FVS for G_i, the set F_{i+1}' is an FVS of size $k + 1$ for the induced subgraph G_{i+1}. In particular, the triple (G_{i+1}, F_{i+1}', k) is a valid instance for the DFVS REDUCTION problem.

Apply Lemma 5.1 to the instance (G_{i+1}, F_{i+1}', k), which either returns an FVS F_{i+1} of size bounded by k for the graph G_{i+1}, or claims that no such FVS exists. It is easy to see that if the induced subgraph $G_{i+1} = G[V_{i+1}]$ does not have an FVS of size bounded by k, then the original graph G cannot have an FVS of size

bounded by k. Therefore, in this case, we can simply stop and conclude that there is no FVS of size bounded by k for the original input graph G. On the other hand, suppose that an FVS F_{i+1} of size bounded by k is constructed for the graph G_{i+1} in the above process, then the induction successfully proceeds from i to $i+1$ with a new pair (G_{i+1}, F_{i+1}).

In conclusion, the above process either stops at some point and correctly reports that the input graph G has no FVS of size bounded by k, or eventually ends with an FVS F_{n-k-1} of size bounded by k for the graph $G_{n-k-1} = G[V_{n-k-1}] = G$.

This process is involved in solving at most $n - k - 1$ instances (G_i, F_i, k) of the DFVS REDUCTION problem, for $0 \leq i \leq n - k - 2$. By Lemma 5.1, the running time of the process is bounded by $O(n^3 4^k k^3 k!(n-k-1)) = O(n^4 4^k k^3 k!)$, and the process correctly solves the DFVS problem. $\square$

Remark 5.3. The running time of the algorithm in Theorem 5.2 can be further improved by taking advantage of existing approximation algorithms for the FEEDBACK VERTEX SET problem on directed graphs. Even et al. [1998] have developed a polynomial time approximation algorithm for the FEEDBACK VERTEX SET problem that for a given directed graph G, produces an FVS F of size bounded by $c \cdot \tau \log \tau \log \log \tau$ in time $O(n^2 M(n) \log^2 n)$, where c is a constant, τ is the size of a minimum FVS for the graph G, and $M(n) = O(n^{2.376})$ is the complexity of the multiplication of two $n \times n$ matrices. Therefore, for a given instance (G, k) of the DFVS problem, we can first apply the approximation algorithm in Even et al. [1998] to construct an FVS F for the graph G. If $|F| > c \cdot k \log k \log \log k$, then we know that the graph G has no FVS of size bounded by k. On the other hand, suppose that $|F| \leq c \cdot k \log k \log \log k$. Then we pick a subset F_0 of arbitrary k vertices in F, and let $G_0 = G - (F - F_0)$. The set F_0 is an FVS of size k for the graph G_0. Now we can proceed exactly the same way as we did in Theorem 5.2: let $F - F_0 = \{v_1, v_2, \ldots, v_h\}$, where $h \leq c \cdot k \log k \log \log k - k$, and let $V_i = V_0 \cup \{v_1, \ldots, v_i\}$, and $G_i = G[V_i]$, for $i = 0, 1, \ldots, h$. By repeatedly applying the algorithm in Lemma 5.1, we can either stop with a certain index i where the induced subgraph G_{i+1} has no FVS of size bounded by k (thus the original input graph G has no FVS of size bounded by k), or eventually construct an FVS F_h of size bounded by k for the graph $G_h = G[V_h] = G$. This process calls for the execution of the algorithm in Lemma 5.1 at most $h = O(k \log k \log \log k)$ times, and each execution takes time $O(n^3 4^k k^3 k!)$. In conclusion, the DFVS problem can be solved in time $O(n^3 4^k k^4 k! \log k \log \log k + n^{4.376} \log^2 n)$, where the second term in the complexity is due to the approximation algorithm given in Even et al. [1998].

6. *Remarks and Future Research*

We presented a parameterized algorithm of running time $O(n^4 4^k k^3 k!)$ for the DFVS problem, which shows that the problem is fixed-parameter tractable, and resolves an outstanding open problem in parameterized computation and complexity. Before we close the paper, we give a few remarks on our results and on directions for future research.

There is an edge version of the FEEDBACK SET problem, which is called the FEEDBACK ARC SET problem (briefly, the DFAS problem): given a directed graph G and a parameter k, either construct a set of at most k edges in G whose removal leaves a DAG, or report that no such edge set exists. The DFAS problem is also a

well-known NP-complete problem [Garey and Johnson 1979]. As shown by Even et al. [1998], the DFAS problem and the DFVS problem can be reduced in linear time from one to the other with the same parameter. Therefore, our results also imply an $O(n^4 4^k k^3 k!)$ time algorithm for the DFAS problem.

The techniques developed in this article for solving the SKEW SEPARATOR problem seem to be powerful and generally useful in the study of a variety of separator problems. For example, it has been used recently in developing improved algorithms for a multi-cut problem on undirected graphs in which a separator is sought to (uniformly) separate a set of given terminals [Chen et al. 2007b]. It will be interesting to identify the conditions for the multi-cut problems under which these techniques (and their variations and generalizations) are applicable. In particular, it will be interesting to see if the techniques are applicable to derive the fixed-parameter tractability of the FEEDBACK VERTEX SET problem on *weighted and directed graphs*. Note that the fixed-parameter tractability of the problem on weighted and *undirected* graphs has been derived recently [Chen et al. 2007a].

It will be interesting to develop new techniques that lead to faster parameterized algorithms for the DFVS problem and other related problems. For example, is it possible that the DFVS problem can be solved in time $O(c^k n^{O(1)})$ for a constant c? Another direction is to look at the *kernelization* of the DFVS problem, by which we refer to a polynomial-time algorithm that on an instance (G, k) of the DFVS problem, produces a (smaller) instance (G', k') of the problem, such that the size of the graph G' (the kernel) is bounded by a function $g(k)$ of k (but independent of the size of the original graph G), that $k' \leq k$, and that the graph G has an FVS of size bounded by k if and only if the graph G' has an FVS of size bounded by k'. Since now it is known that the DFVS problem is fixed-parameter tractable, by a general theorem in parameterized complexity theory [Downey and Fellows 1999], such a kernelization algorithm exists for the DFVS problem. However, how small can the size of the kernel G' be? In particular, can the kernel G' have its size bounded by a polynomial of the parameter k? We note that recently there has been progress in the study of kernelization for the FEEDBACK VERTEX SET problem on undirected graphs. Bodlaender [2007] was able to give a kernel of size $O(k^3)$ for the FEEDBACK VERTEX SET problem on undirected graphs, and Bodlaender and Penninkx [2008] have shown that the FEEDBACK VERTEX SET problem on undirected planar graphs has a kernel of size $O(k)$.

REFERENCES

BAFNA, V., BERMAN, P., AND FUJITO, T. 1999.   A 2-approximation algorithm for the undirected feedback vertex set problem. *SIAM J. Disc. Math. 12*, 289–297.

BODLAENDER, H. 1993.   On linear time minor tests with depth first search. *J. Algorithms 14*, 1–23.

BODLAENDER, H. 2007.   A cubic kernel for feedback vertex set. In *Proceedings of the 24th Annual Symposium on Theoretical Aspects of Computer Science (STACS 2007),* Lecture Notes in Computer Science, vol. 4393. Springer-Verlag, New York, 320–331.

BODLAENDER, H., AND PENNINKX, E. 2008. A linear kernel for planar feedback vertex set. In *Proceedings of the 3rd International Workshop on Parameterized and Exact Computation (IWPEC 2008),* Lecture Notes in Computer Science, vol. 5018. Springer-Verlag, New York, 160–171.

CHEN, J., FOMIN, F., LIU, Y., LU, S., AND VILLANGER, Y. 2007a. Improved algorithms for the feedback vertex set problems. In *Proceedings of the 10th Workshop on Algorithms and Data Structures (WADS'07),* Lecture Notes in Computer Science, vol. 4619. Springer-Verlag, New York, 422–433.

CHEN, J., KANJ, I., AND JIA, W. 2001. Vertex cover: further observations and further improvements. *J. Algorithms 41*, 280–301.

CHEN, J., LIU, Y., AND LU, S. 2007b. An improved parameterized algorithm for the minimum node multi-way cut problem. In *Proceedings of the 10th Workshop on Algorithms and Data Structures (WADS'07),* Lecture Notes in Computer Science, vol. 4619. Springer-Verlag, New York, 495–506.

DEHNE, F., FELLOWS, M., LANGSTON, M., ROSAMOND, F., AND STEVENS, K. 2007. An $O(2^{O(k)}n^3)$ FPT algorithm for the undirected feedback vertex set problem. *Theory Comput. Syst. 41*, 479–492.

DOWNEY, R., AND FELLOWS, M. 1992. Fixed parameter intractability. In *Proceedings of the 7th Annual Structural Complexity Conference*. IEEE Computer Society Press, Los Alamitos, CA, 36–49.

DOWNEY, R., AND FELLOWS, M. 1995. Fixed-parameter tractability and completeness I: Basic results. *SIAM J. Comput. 24*, 873–921.

DOWNEY, R., AND FELLOWS, M. 1999. *Parameterized Complexity*. Springer-Verlag, New York.

EVEN, G., NAOR, J., SCHIEBER, B., AND SUDAN, M. 1998. Approximating minimum feedback sets and multicuts in directed graphs. *Algorithmica 20*, 151–174.

GARDARIN, G., AND SPACCAPIETRA, S. 1976. Integrity of databases: A general lockout algorithm with deadlock avoidance. In *Modeling in Data Base Management System,* G, Nijsssen, Ed. North-Holland, Amsterdam, The Netherlands, 395–411.

GAREY, M., AND JOHNSON, D. 1979. *Computers and Intractability: A Guide to the Theory of NP-Completeness*. W.H. Freeman, New York.

GUO, J., GRAMM, J., HÜFFNER, F., NIEDERMEIER, R., AND WERNICKE, S. 2005. Improved fixed-parameter algorithms for two feedback set problems. In *Proceedings of the 9th Workshop on Algorithms and Data Structures (WADS'05),* Lecture Notes in Computer Science, vol. 3608. Springer-Verlag, New York, 158–168.

GUO, J., GRAMM, J., HÜFFNER, F., NIEDERMEIER, R., AND WERNICKE, S. 2006. Compression-based fixed-parameter algorithms for feedback vertex set and edge bipartization. *J. Comput. Syst. Sci. 72*, 1386–1396.

GUTIN, G., AND YEO, A. 2008. Some parameterized problems on digraphs. *Comput. J. 51*, 363–371.

KARP, R. 1972. Reducibility among combinatorial problems. In *Complexity of Computer Computations,* R. Miller and J. Thatcher, Eds. Plenum Press, New York, 85–103.

LEIGHTON, T., AND RAO, S. 1999. An approximation max-flow min-cut theorem for uniform multi-commodity flow problems with applications to approximation algorithms. *J. ACM 46*, 787–832.

LEISERSON, C., AND SAXE, J. 1991. Retiming synchronous circuitry. *Algorithmica 6*, 5–35.

LICHTENSTEIN, O., AND PNUELI, A. 1985. Checking that finite state concurrent programs satisfy their linear specification. In *Proceedings of the 12th ACM Symposium on Principles of Programming Languages*. ACM, New York, 97–107.

REED, B., SMITH, K., AND VETTA, A. 2004. Finding odd cycle transversals. *Oper. Res. Lett. 32*, 299–301.

SCHRIJVER, A. 2003. *Combinatorial Optimization*. Springer-Verlag, Berlin, Germany.

SILBERSCHATZ, A., AND GALVIN, P. 1994. *Operating System Concepts,* 4th ed. Addison Wesley, Reading, MA.