

Rule-Based Modelling, Symmetries, Refinements

Vincent Danos^{1,3}, Jérôme Feret², Walter Fontana²,
Russell Harmer³, and Jean Krivine²

¹ University of Edinburgh

² Harvard Medical School

³ CNRS, Université Paris Diderot

Abstract. Rule-based modelling is particularly effective for handling the highly combinatorial aspects of cellular signalling. The dynamics is described in terms of interactions between partial complexes, and the ability to write rules with such partial complexes -*i.e.*, not to have to specify all the traits of the entities partaking in a reaction but just those that matter- is the key to obtaining compact descriptions of what otherwise could be nearly infinite dimensional dynamical systems. This also makes these descriptions easier to read, write and modify.

In the course of modelling a particular signalling system it will often happen that more traits matter in a given interaction than previously thought, and one will need to strengthen the conditions under which that interaction may happen. This is a process that we call *rule refinement* and which we set out in this paper to study. Specifically we present a method to refine rule sets in a way that preserves the implied stochastic semantics.

This stochastic semantics is dictated by the number of different ways in which a given rule can be applied to a system (obeying the mass action principle). The refinement formula we obtain explains how to refine rules and which choice of refined rates will lead to a neutral refinement, *i.e.*, one that has the same global activity as the original rule had (and therefore leaves the dynamics unchanged). It has a pleasing mathematical simplicity, and is reusable with little modification across many variants of stochastic graph rewriting. A particular case of the above is the derivation of a maximal refinement which is equivalent to a (possibly infinite) Petri net and can be useful to get a quick approximation of the dynamics and to calibrate models. As we show with examples, refinement is also useful to understand how different subpopulations contribute to the activity of a rule, and to modulate differentially their impact on that activity.

1 Semi-liquid Computing

To the eye of the computational scientist, cellular signalling looks like an intriguing computational medium. Various types of agents (proteins) of limited means interact in what, at first sight, may seem to be a liquid universe of chance encounters where there is little causality. But in fact a rich decentralized choreography of bindings (complex formation) and mutual modifications (post-translational modifications) can be observed. Transient devices (complexes) are built by agents to

integrate, convey, and amplify signals and channel them to the appropriate outputs (transcriptional regulation). The intricate pathways of the response to the epidermal growth factor (EGF) sketched in Fig. 1 are a well-studied and well-modelled example [1]. This universe of semi-liquid computing is brought about by a surprisingly small number of elementary interactions. It sits somewhere in between the worlds of the random graphs of statistical physics [2] which perhaps lack expressivity, and the solid colliding sphere models of chemical kinetics [3] which perhaps lack programmability.

The generativity of those systems, that is to say the number of different non-isomorphic combinations (aka complexes or species) that may come to exist along different realizations of the implied stochastic process, may well be enormous, but this does not say how complex those systems really are. A lot fewer rules than there are reactions (interactions between complete complexes) may be good enough to describe some of them. For instance the sketch of Fig. 1 once properly formalized uses about 300 rules whereas it produces about 10^{40} unique combinations. One sees that the number of rules is a more meaningful estimate of its inherent complexity.

Rule-based languages [4,5,6,7,8,9,10,11], and more generally process algebraic approaches to modelling [12,13,14,15,16,17,18,19], because they can express such generic interactions, can work around this apparent descriptive complexity and achieve compact descriptions. Let us also mention, although we will not treat this aspect of the question here, that another benefit of rule-based modelling is that one can trace the evolution of a system at the level of agents (or individuals) and explore the causal relationships between events occurring in a system [6].

The difference between an assembly of agents with random uncorrelated encounters and a signalling system is that there is a causal structure channelling the interactions towards a particular response. Typically a binding will not happen before one or both of the bindees has been modified. Combining those micro-causal constraints into a coherent pathway is a programming art that we don't master or even understand yet, but one that signalling systems have been honing for a considerable time. Rule-based modelling incorporates such causality constraints in the rules themselves by using partial complexes: not everything needs to be described in a rule, only the aspects of the state of a complex which matter for an event to happen need to be specified. That is the difference between a reaction between complete entities, and a rule between partial ones. As said, this reliance on partial complexes allows to capture compact descriptions and work around the huge numbers of combinations one would have to contemplate (or neglect) otherwise. The more detailed the partial complex, that is to say the less partial, the more conditions must be met for a particular event to happen.

The purpose of the present paper is to understand better the mechanics of *refinement*, that is to say the process by which one can make a complex less partial, or equivalently a rule more demanding. We specifically consider the problem of replacing a rule with a family of refined rules which will exhibit the same collective activity, and will therefore generate an identical stochastic behaviour. Note that there are really two questions in one: one is to define what




Fig. 1. A informal sketch of the many interactions involved in the ERK/AKT pathway responding to an EGF stimulus. The corresponding rule-based model generates about 10^{40} different species.

constitutes a good family of refined rules, another is to define their rates so as to preserve the underlying dynamics. It turns out that the latter question has an intimate connexion with the notion of symmetry, and what becomes of a symmetry group of a partial complex under refinement. The solution we propose to the former question can certainly be made to cover more cases (of which more later).

Seeing how the notion of partial complex is central to rule-based modelling¹ it certainly makes sense to try to theorize around it, as we start doing in this paper. But there are also very concrete reasons to do so. First it will often happen that in a modelling situation a rule has to be revised because people come to believe that its rate depends on more information about the context than the rule actually provides. A typical example would be that a post-translational modification increases or decreases the likelihood that an agent will bind another one. Replacing a rule with a bunch of more specific ones, in order to express those

¹ An aside: the name rule-based is a little unfortunate since just about any computational formalism is rule-based but that is the name under which this approach has become known in the biological modelling community.

context-dependent modulations of the rule activity, is a transformation which we call a kinetic refinement of the rule. It can be usefully decomposed as first introducing a neutral refinement -as defined in this paper- and second changing those base rates to achieve the modulation of interest. In this application the neutral refinement, that is to say the choice of refined rates that will *not* change the behaviour, serves as a baseline. One needs it to know where to start from. In fact even when one does not actually modulate the rates of the refined rules and keeps the refinement neutral, the procedure allows one to peek into the relative contributions of the various subpopulations of complexes that can intervene in an instance of the original unrefined rule (see the examples at the end of the next section). So for both reasons it is important to understand how to compute this baseline which is the question we address here.

Maximal refinements are of special interest. This is the case where one replaces a rule with all its ground instances (in general an infinite set) where only complete complexes take part. Such a transformation when applied to all rules in a rule set will obtain a set of multiset rewriting rules, that is to say a (possibly infinite) Petri net. This transformation will be unfeasible in general, owing to the combinatorial explosion mentioned earlier, because the obtained Petri net, even if finite in principle, will be simply too large to be written (this is not even a problem of computational complexity but of mere size of the output). However it is easy to imagine running truncated versions of a complete expansion using an ODE semantics. That could be useful for model calibration, and similar exploration mechanisms that are particularly demanding in terms of the number of simulations required while not necessarily needing the accuracy an exact expansion would provide.

We start with a brief presentation of Kappa which is the rule-based language we shall use in this paper. This is an occasion to get familiar with some of the notations, but is in no way a formal presentation. Then we turn to two simple examples of refinement to get a more concrete sense of what the notion of refinement is trying to achieve and how it is relevant to practical modelling questions. Explanations given in this paper, further than the ones given above, about the relevance of Kappa for the actual practical modelling are all to be found in the next Section. The reader interested in more can consult Refs. [6] and [7].

After this presentation the mathematical development reintroduces a simplified Kappa, this time in a completely formal and algebraic way which is conducive to a study of refinement which will be of general import and not tied in specific syntactical details. In fact the refinement formula we obtain is of general validity and assumes nothing about the arity of the rule to be refined, and actually assumes little about the rewriting framework itself. We are conscious that this incurs some cost to the reader unacquainted with basic category-theoretic notions in that some heavy-looking machinery is involved. However, mathematically it is natural, and hopefully at this stage the preliminary informal explanations will have clarified what is achieved in the mathematical development. After the derivation of the refinement formula for partial complexes first, and then for rules, we explain how to sharpen this result by considering more inclusive notions

of refinements. Only the first refinement formula is treated in detail; its extensions are just sketched.

2 A Brief Guide to Kappa

Let us start with an informal and brief account of our modelling language. In Kappa, agents (think of them as idealized proteins) have sites and their sites can be used to bind other sites (at most once), and can also hold an internal state. The former possibility accounts for domain-mediated complex formation, while the latter accounts for post-translational modifications. Accordingly one distinguishes three types of (atomic) rules for binding, unbinding, and modification. In the full language one also considers agent creation and deletion (see later the formal presentation), and it is possible to combine actions in a single rule.


Note that a binding rule requires two distinct agents, each with a free (*i.e.*, not already bound) site, which bind via those sites. In other words, it is not possible to bind a site more than once.

A Kappa *model* consists of (i) an *initial solution* that declares the names and all sites (with default state values for all sites we wish to carry a state) of the relevant agents; and (ii) a *rule set* specifying how the initial solution may evolve. We will see an example very soon.

The behaviour of a model is stochastic. Given a global state of the system one assigns to each rule a likelihood to be applied which is proportional to the number of ways in which this rule can be applied, and its intrinsic rate (the rate is a measure of how efficient a rule is at turning a chance encounter of reagents into an actual reaction). In the particular case where agents have no sites at all, one has a Petri net, and the dynamics is none other than the mass action law put in Gillespie form [3].


2.1 A Simple Cascade

As a way of getting more familiar with the notation we can consider a simple and yet ubiquitous motif of cellular biology that consists of one protein (typically an enzyme or kinase) covalently modifying another. Let us call them S (as signal) and X which we are going to assume have each a unique site s . That situation can be easily expressed by the following rule triplet:



where we represent a binding between two sites by a shared exponent, s^1 , and the internal state of a site as a subscript to this site, s_u or s_p , where (u) p is a mnemonic for (un-) phosphorylated.

One can add a second and similar triplet for X and a new agent Y :






Fig. 2. Low off rate ($k = 0.1$): the activated Y s tend to stay attached to their activators X (the XYp ‘sleep’ curve dominates the XYu ‘ready’ one); as a consequence the production of activated Y is slowed down




Fig. 3. Medium off rate ($k = 10$): most of the XY complexes have now an inactivated Y (the XYu ‘ready’ curve now dominates); the production of activated Y is visibly faster

Note that X has to be activated and Y inactivated for the binding and subsequent activation of Y to happen. This ensures in particular that no Y is activated

in the absence of a signal S (an example of the causal constraints we were alluding to earlier). Such *cascades* regularly arise in real signalling networks. Referring back to the actual EGFR pathway in Fig. 1 we see the famous examples of the Ras, Raf, MEK and ERK cascade, and the PIP3, PDK1, and AKT one.

What about the dynamical behaviour of such a simplified cascade? We would like to understand how the cascade *throughput*, that is to say the rate of production of the active form of Y , depends on the rate k at which X detaches from Y (hereafter the XY *off-rate*), namely the rate of the rule $r := X(s^1), Y(s^1) \rightarrow X(s), Y(s)$ defined above.

Well intuitively, with too small an off-rate (high affinity binding), X will tend to remain bound to Y even after the Y has been activated. Whereas with too large an off-rate (low affinity binding), X will often detach from Y before having activated Y . Somewhere in between, an optimal choice of k will strike the right balance and maximise the rate of activation of Y .

This is something that we can verify numerically. Suppose one starts with $15S(s) + 60X(s_u) + 120Y(s_u)$ as an initial state, and suppose further all other rules have a rate of 1. As expected, we see in Fig. 2 where $k = 0.1$, that the activation of Y is slower than in Fig. 2, where $k = 10$; with an even higher $k = 1000$ the activation rate goes down again (Fig. 4).

This demonstrates the tension between binding loosely, and “not always getting the job done”, and binding tightly which amounts to “sleeping on the job”. It also nicely shows that the cascade throughput depends on a lot more than just the rate attached to the rule performing that activation.

Now our off-rate k measures how likely it is that X and Y will detach - *independently* of their respective internal states. If we were to optimize the cascade throughput it would be natural to let the off-rate depend on the state of Y . In terms of rules, all we have to do is to split the unbinding rule under consideration into two subcases:

$$\begin{aligned} r_u &:= X(s^1), Y(s_u^1) \rightarrow_{k_u} X(s), Y(s_u) \\ r_p &:= X(s^1), Y(s_p^1) \rightarrow_{k_p} X(s), Y(s_p) \end{aligned}$$

with respective rates k_u, k_p . One calls the substitution of r with such more specific rules r_u, r_p a *refinement*. If in addition the new rates are both taken equal to k , then in this simple case evidently the behaviour of the system will be unchanged. That special case, where nothing changes in the dynamics, is what we call a neutral refinement.

If one favours unbinding from active Y , then clearly this allows X both to bind long enough to Y to activate it and then to unbind quickly to maximize throughput. In particular the combination $k_p = \infty$ (detach as soon as activated), and $k_u = 0$ (never detach before activation) leads to the best possible throughput, all other things being equal (see Fig. 5).

2.2 A Less Obvious Refinement

Here is a second example which shows that choosing the rates of the refined rules and obtaining a neutral refinement may require some ingenuity.




Fig. 4. High off rate ($k = 1000$): the production of activated Y has gone down again




Fig. 5. Split rate ($k_u = 0, k_p = \infty$): there is no XYp anymore; the production of activated Y is optimal

Consider two agent types B, C each with only one site x , and define a family of systems $x(n_1, n_2)$ consisting of n_1 single C s and n_2 dimers $C(x^1), B(x^1)$. In other words set $x(n_1, n_2) := n_1 C(x) + n_2 (C(x^1), B(x^1))$. Now consider the rule $r := C(), B() \rightarrow C()$ with rate 1. Note that r does not mention x at all (we

say that both agents have an empty signature in this case). This means that r applies irrespective of the binding state of x in B and C . Both agents could be free, or bound, or even bound together. Whichever is the case, the effect of the rule will be the same, namely to delete a B and to bring $x(n_1, n_2)$ to a new state $x(n_1 + 1, n_2 - 1)$. This supposes $n_2 > 0$. If on the other hand $n_2 = 0$ then there is no B left in the system and no further event is possible (deadlock).

Now we would like to refine r into mutually exclusive sub-cases depending on the relationship in which C and B stand; specifically we want to use the following three refined rules:

$$\begin{aligned} r_1 &:= C(x^1), B(x^1) \rightarrow_1 C(x) \\ r_2 &:= C(x^1), B(x^1), C(x^2), B(x^2) \rightarrow_2 C(x), C(x^2), B(x^2) \\ r_3 &:= C(x^1), B(x^1), C(x) \rightarrow_1 C(x), C(x) \end{aligned}$$

Each of them is a particular case of r in the sense that their left hand sides embed (sometimes in more than one way) that of r (see below the notion of morphism). Intuitively, r_1 is the sub-case where B, C are bound together, r_2 is the sub-case where they are both bound but not together, and r_3 is the sub-case where B is bound but C is free. Given the particular family of states $x(n_1, n_2)$ we are dealing with, those seem to cover all possible cases, and to be indeed exclusive.

Define the *activity* of a rule as the number of possible ways to apply the rule multiplied by its rate. This determines its likelihood to apply next and only depends on the current state of the system. Now we have chosen for each refined rule a rate (indicated as a subscript to the reaction arrow), and in particular r_2 was assigned a rate of 2. We claim this is the unique correct choice if one wants the stochastic behaviour of the system to be preserved by the refinement. Figure 6 shows a run of the refined system with $x(0, 100)$ as the initial state. The y axis traces the activity of all rules including the base one r .

We see that indeed at all times the refined activities add up to the original one (the top curve).

There are other things worth noticing. Firstly, r_1 keeps a low probability that decreases linearly during the simulation since its activity is exactly the number of dimers n_2 ; so suppressing r_1 would change very little to the behaviour of the system. Secondly, r_2 dominates the early events, since near the initial state there are only dimers, and no free C s yet; however, as time passes there will be more of those free C s, and the corresponding rule r_3 will come to dominate. Hence we see that the relative importance of the sub-cases changes over time, and that refinement can be used as a way of profiling the contribution various subpopulations of agents make to a given type of event.

The corrective factor applied to r_2 accounts for two opposite effects: on the one hand r_2 embeds r in more than one way which tends to scale the rate of r_2 upwards, on the other hand r_2 is more symmetric than r and that would tend to scale the rate of r_2 downwards.

What we are interested in is to handle the general case, *i.e.*, to explain what constitutes a good set of refined rules as r_1, r_2 , and r_3 above, and how one can




Fig. 6. The activities of the refined rules r_1 , r_2 , and r_3 add up exactly to that of the initial rule r (top curve)

choose the refined rates in a way that the global activity is preserved. We will return to the example once we have a general solution.

3 Rule-Based Modelling

To give proper generality and portability to our study, we will frame it into some simple categorical language where a system is seen as an object x and the various ways a rule r may apply to x are identified using a notion of morphism f from r 's left hand side to x .

As said we shall also simplify the Kappa syntax in two respects. First, we suppose agents have no internal states. Second, we suppose no wildcards are used in left hand sides, e.g., expressions like $A(x^*)$ meaning x is bound to some unspecified other site, are not considered. The former simplification is only a matter of readability, as internal values offer no difficulty. The latter simplification is more significant, and we will see later in our development that reintroducing wildcards allows us to strengthen our main result. With these simplifications we can give a syntax-less presentation of Kappa that will facilitate the derivation of the refinement formula.

We suppose given two sets \mathcal{A} and \mathcal{S} of agent names and sites.

A *matching* over a set X is an irreflexive and symmetric binary relation on X such that no element is in relation with more than one element.

Definition 1. An object is a quadruple $(V, \lambda, \sigma, \mu)$ where:

- V is a set of nodes,
- $\lambda \in \mathcal{A}^V$ assigns names to nodes,
- $\sigma \in \mathcal{P}(\mathcal{S})^V$ assigns sets of sites to nodes,
- μ is a matching over $\sum_{v \in V} \sigma(v)$.

The matching represents bindings, and hence any given site can be bound at most once. A node however can be bound many times via different sites.

We define $(u, x) \in \mu$ as shorthand for $\exists(v, y) : (u, x, y, v) \in \mu$, and say u, x is free when $(u, x) \notin \mu$, bound when $(u, x) \in \mu$.

The simplest non-empty object is a single node named A with no sites and therefore no binding. In the preceding section we wrote $A()$ for this object. There we also introduced a textual notation to designate objects where bindings are indicated by exponents.

Note that we sometimes use the same family of symbols x, y , etc. for sites and objects. Hopefully this will not cause any confusion since they are entities of a very different nature.

We define a *signature* as a map $\Sigma : \mathcal{A} \rightarrow \mathcal{S}$; this can be used to constrain the set of sites per agent type. We write $x \leq \Sigma$ if for all $v \in V_x$, $\sigma_x(v) \subseteq \Sigma(\lambda_x(v))$; likewise we write $\Sigma \leq x$ if for all $v \in V_x$, $\Sigma(\lambda_x(v)) \subseteq \sigma_x(v)$, and $x : \Sigma$ when $x \leq \Sigma \leq x$.

When $x : \Sigma$ for some Σ , we say x is *homogeneous*, which means all agents of the same type in x use exactly the same set of sites.

Definition 2. An arrow $(V, \lambda, \sigma, \mu) \rightarrow (V', \lambda', \sigma', \mu')$ is a map $f : V \rightarrow V'$ such that

- 1) f preserve names: $\lambda' \circ f = \lambda$
- 2) f preserve sites: $\sigma' \circ f \supseteq \sigma$
- 3a) f preserve edges: $(u, x, y, v) \in \mu \Rightarrow (f(u), x, y, f(v)) \in \mu'$
- 3b) f reflects edges: $(f(u), x) \in \mu', x \in \sigma_x(u) \Rightarrow (u, x) \in \mu$
- 4) f is a monomorphism

This then is the category of *graphs with sites* we shall work with. We also call arrows morphisms sometimes; we write $[x, y]$ for the arrows from x to y ; $iso[x, y]$ for the isomorphisms (meaning invertible arrows), and therefore $[x, x] = iso[x, x]$ denotes the set of *automorphisms* (or symmetries) of x ; we say that y *embeds* x when $[x, y] \neq \emptyset$.

Define the *image* of $f \in [x, y]$ as $Im(f) := \{f(v), x; v \in V, x \in \sigma(v)\}$.

Note that $Im(f)$ is but a subset of $\sum_{v \in V} \sigma'(f(v))$, and only sites in $Im(f)$ are mentioned in the arrow-defining clauses above.

One has obviously a forgetful functor to the category of graphs and graph morphisms, and that allows us to import the usual graph-theoretical vocabulary of connected components and paths, which we will freely use in the sequel. Note that, from the point of view of graphs, the reflectivity condition 3b) above does not really make sense, one really needs sites to express edge reflection. Moreover the rather stringent notion of arrow constrains the homsets $[x, y]$:

Lemma 1 (rigidity). *Suppose x is connected, then any non-empty partial injection f from V_x to V_y extends to at most one morphism in $[x, y]$.*

Proof: If f is strictly partial, that is to say $V_x \setminus \text{dom}(f)$ is not empty, pick a v in there such that for some node $w \in \text{dom}(f)$, and some sites x, y , $(w, y, v, x) \in \mu_x$. This is always possible because x is connected. Then, either $(f(w), y, v', x) \in \mu_y$ for some $v' \in V_y$, and by 3a) one must extend f as $f(v) = v'$, or there is no such extension. \square

Clearly being a monomorphism, *i.e.*, being post-cancellable, is equivalent to being a one-one map. On the other hand there are far more epimorphisms than surjections:

Lemma 2 (epis). *A map $h \in [x, y]$ is an epimorphism iff every connected component of y intersects $f(x)$; that is to say for all connected component $c_y \subseteq y$, $h^{-1}(c_y) \neq \emptyset$.*

Proof: Suppose $f_1 h = f_2 h$ for $h \in [x, y]$, $f_i \in [y, z]$, and let $c_y \subseteq y$ be a connected component of y such that $h^{-1}(c_y) \neq \emptyset$. Pick u such that $h(u) \in c_y$, then $f_1(h(u)) = f_2(h(u))$ and by the preceding lemma $f_1/c_y = f_2/c_y$. \square

We write $[x, y]^e \subseteq [x, y]$ for the epis from x to y .

4 Object Refinements

Now that we have our basics in place we turn to the first question of what constitutes a refinement of a (partial) object s . As we have seen in the example, a refinement of s is intuitively a collection of objects t_i that embed s and such that any embedding in an object of higher signature x (*i.e.*, that has more sites everywhere) can be unambiguously attributed to one t_i . We first make this intuition into a real definition and then proceed to define the refinements of rules.

Definition 3 (factorisation). *One says an object t factors $f \in [s, x]$ if $f = \gamma\phi$ for some $\phi, \gamma \in [s, t]^e \times [t, x]$; ϕ, γ is called a factorisation of f via t .*

The first thing to notice is that one cannot ask for unique factorisations.

Suppose given a factorisation $\phi, \gamma \in [s, t]^e \times [t, x]$ of f via t and an isomorphism $\alpha \in [t, t']$. Define $\phi', \gamma' := \alpha\phi, \gamma\alpha^{-1} \in [s, t']^e \times [t', x]$; this new pair verifies $\gamma\phi = \gamma'\phi'$, and since ϕ' is clearly an epimorphism, the pair is also a factorisation of f via t' .

$$\begin{array}{ccc}
 s & \xrightarrow{\phi} & t \\
 \phi' \downarrow & \nearrow \alpha & \downarrow \gamma \\
 t' & \xrightarrow{\gamma'} & x
 \end{array} \tag{1}$$

In this case we will say that ϕ, γ and ϕ', γ' are *conjugate* under α , and write $\phi, \gamma \simeq_{t'} \phi', \gamma'$. We also write $[s, t] \times_{[t, t]} [t, x]$ for the quotient of $[s, t] \times [t, x]$ under \simeq_{tt} ; this notation is justified by the following:

Lemma 3 (conjugates). *The equivalence relation \simeq_{tt} has $||[s, t] \times [t, x]||/||[t, t]||$ classes.*

Proof: Suppose, using the notations of (1), that $\phi, \gamma \simeq_{tt'} \phi', \gamma'$, then this uniquely determines α since $\gamma\alpha^{-1} = \gamma'\alpha'^{-1}$ implies $\alpha = \alpha'$ by γ being a monomorphism. In particular the set of conjugates of ϕ, γ over the same t is in one-one correspondence with $[t, t]$. \square

Note that being an epimorphism is stable by conjugation so we can say that a class is an epimorphism, and we can restrict the equivalence to $[s, t]^e \times [t, x]$, so the version of the Lemma relative to $[s, t]^e$ also holds.²

Unicity of factorization is then to be understood up to isomorphisms; furthermore, even if we select one representative t_i per isomorphism class, unicity is up to automorphisms of each of the representative t_i .

Definition 4 (object refinement). *Given s, Σ such that $s \leq \Sigma$, a refinement of s under Σ , written $\Sigma(s)$, is a set of objects obtained by selecting one representative in each isomorphism class defined by $\{t \mid t : \Sigma, [s, t]^e \neq \emptyset\}$.*

Note that the actual choice of representatives does not matter, but we do have to choose one for our counting purposes.

Another noteworthy fact is that $\Sigma(s)$ in general will be infinite. However in practice one may get information about the reachables of the system which will allow to control the size of the expansion [4]; indeed it is not necessary to include ts which are not reachable, and we took advantage of this in the example of the first section.

Lemma 4 (injectivity). *Given Σ, s, x such that $s \leq \Sigma \leq x$ the composition map from the disjoint sum $\sum_{t \in \Sigma(s)} [s, t]^e \times_{[t, t]} [t, x]$ to $[s, x]$ is injective.*

Proof: Suppose given two factorisations $f = \gamma\phi = \gamma'\phi'$ via t and t' as in (1).

$$\begin{array}{ccc}
 v \in s & \xrightarrow{\phi} & t \supseteq c \ni \phi(v) \\
 \phi' \downarrow & \nearrow \alpha & \downarrow \gamma \\
 \phi'(v) \in c' \subseteq t' & \xrightarrow{\gamma'} & x
 \end{array} \tag{2}$$

Pick a connected component $c \subseteq t$, such that $\phi(v) \in c$ for some $v \in s$. Call $c' \subseteq t'$ the connected component of $\phi'(v)$ in t' . By construction $\gamma(c)$ and $\gamma'(c')$ intersect

² This prompts a more general argument for the restriction of that Lemma to the case of $[s, t]^e$. Suppose given ϕ, γ as above, one has a map from $[t, t]$ to the class $[\phi, \gamma]$ in $[s, t] \times [t, x]$: $\alpha \mapsto \alpha \cdot (\phi, \gamma) := \phi\alpha, \alpha^{-1}\gamma$ (an action of the group $[t, t]$ on $[s, t] \times [t, x]$); this map is surjective by definition of conjugation; it is also injective because γ is a mono. Now if in addition ϕ is an epi, it is injective for a second reason, namely ϕ is an epi ($\alpha_1\phi = \alpha_2\phi$ implies $\alpha_1 = \alpha_2$). This seems to indicate that one can relax the monomorphism requirement in the ambient category and still develop the same theory.

at $\gamma\phi(v) = f(v) = \gamma'\phi'(v)$. It is easy to see that they both are Σ -homogeneous. This means they must be equal.

Indeed suppose $w \in \gamma(c)$ is a node which is directly connected to $\gamma(c) \cap \gamma'(c')$, meaning w is such that $(u, x, y, w) \in \mu_x$, for some $u \in \gamma(c) \cap \gamma'(c')$ and $(u, x), (w, y) \in \text{Im}(\gamma)$. Because c' is Σ -homogeneous, $u, x \in \text{Im}(\gamma')$, ie x is also a site of the (unique) antecedent of u in c' , which we can write $x \in \sigma_{t'}\gamma'^{-1}(u)$. By condition 3b) this site cannot be free, and by 3a) it must be bound to $\gamma'^{-1}(w), y$, so $w \in \gamma'(c')$. Since $\gamma(c)$ is connected, $\gamma(c) \cap \gamma'(c')$ must contain $\gamma(c)$, and by symmetry $\gamma'(c')$.

Hence $\gamma(c) = \gamma'(c')$, therefore c and c' are isomorphic. In fact, since ϕ is an epi, we can repeat the above for any connected component in t , and therefore t embeds in t' (it is readily seen that the assignment of a c' to a c above is injective), and by symmetry they must be isomorphic under a certain isomorphism α . By definition of $\Sigma(s)$ we have picked exactly one representative in each isomorphism class, therefore $t = t'$, $\alpha \in [t, t]$, and the two factorizations are conjugate under α . \square

Theorem 1. *Given Σ, s, x such that $s \leq \Sigma$ and $x : \Sigma$:*

$$[s, x] \simeq \sum_{t \in \Sigma(s)} [s, t]^e \times_{[t, t]} [t, x]$$

Proof: From the preceding lemma we know the composition map is injective, so all there remains to prove is that it is surjective.

Consider $f \in [s, x]$, define $f(s) := \{u \mid \exists x : (u, x) \in \text{Im}(f)\} \subseteq V_x$, and write $[f(s)]$ for the connected closure of $f(s)$ in x . We claim there is a $t \in \Sigma(s)$ which is isomorphic to $[f(s)]$. Indeed every node in $[f(s)]$ has a signature in accordance with Σ because $x : \Sigma$, and $[f(s)]$ embeds s since $f(s)$ does (via f). \square

Using Lemma 3 in addition we can use the above theorem to obtain:

Corollary 1. *Given Σ, s, x such that $s \leq \Sigma$ and $x : \Sigma$, one has:*

$$|[s, x]| = \sum_{t \in \Sigma(s)} |[s, t]^e| / |[t, t]| \cdot |[t, x]| \quad (3)$$

There are several things worth noticing about the theorem and its numerical form as a corollary.

First, the $|[s, t]^e| / |[t, t]|$ is a *static* term that can be computed once and for all, and which we shall use to determine the rule rates. The positive contribution $[s, t]^e$ is rather intuitive since the more copies of s one finds in t the higher the contribution of t to the number of instances of s should be; the negative contribution $|[t, t]|$ is less intuitive however.

Second, one cannot relax the homogeneity condition on x and ask only $\Sigma \leq x$. That would break the easy part of the proof, namely that of surjectivity. Here is an example; set $s := A(x) < \Sigma := A \mapsto \{x, y\} < A(x, y^1, z), A(x, y, z^1) =: x$. Choose f to be the ‘left’ morphism mapping s ’s unique A to $A(x, y^1, z)$ in x ; then $[f(s)] = x$ and no $t \in \Sigma(s)$ can factorize f because the (y, z) binding is not reproducible in t , because $z \notin \Sigma(A)$.

However, one can modify the notion of object (and accordingly that of arrow) by introducing new partial objects such as $t = A(x, y^{-\Sigma})$, meaning y binds an otherwise unspecified non- Σ site (ie A, y is bound to some B, z such that $z \notin \Sigma(B)$). This t is homogeneous and factorizes the f above. This variant allows to recover surjectivity and extend our decomposition theorem above. Similar wildcard expressions are already present in the actual syntax of Kappa, and it is amusing to see that those convenient notations also have a theoretical status.

This begs a last remark, namely that we are the ones choosing how to relate the base object s and its refinement. For example, here, we are using epis to relate them. Below we will allude to a finer-grained correspondence based on using a pointed version of the ambient category that will allow us to go beyond the homogeneity requirement in another way. But before we do that we will return to the example of the first section.

4.1 Example Continued

We can now reconsider our initial example. Set $s := C(), B()$, for the left hand side of the base rule r , and t_i for that of the refined rule r_i :

$$\begin{aligned} t_1 &:= C(x^1), B(x^1) \\ t_2 &:= C(x^1), B(x^1), C(x^2), B(x^2) \\ t_3 &:= C(x^1), B(x^1), C(x) \end{aligned}$$

Set also $\Sigma := B, C \mapsto \{x\}$. Clearly $s < \Sigma$ and t_i , and $x(n_1, n_2)$ are Σ -homogeneous. Besides due to the particular form of $x(n_1, n_2)$, the t_i s are the only elements in $\Sigma(s)$ that $x(n_1, n_2)$ embed. Using Lemma 3 we get:

$$\begin{aligned} |[s, x(n_1, n_2)]| &= n_2(n_1 + n_2) \\ |[s, t_1]^e \times_{[t_1, t_1]} [t_1, x(n_1, n_2)]| &= |[s, t_1]^e [t_1, x(n_1, n_2)]| / |[t_1, t_1]| = 1.n_2/1 = n_2 \\ |[s, t_2]^e \times_{[t_2, t_1]} [t_2, x(n_1, n_2)]| &= |[s, t_2]^e [t_2, x(n_1, n_2)]| / |[t_2, t_2]| = 2.n_2(n_2 - 1)/2 \\ |[s, t_3]^e \times_{[t_3, t_1]} [t_3, x(n_1, n_2)]| &= |[s, t_3]^e [t_3, x(n_1, n_2)]| / |[t_2, t_2]| = 1.n_1 n_2 / 1 = n_1 n_2 \end{aligned}$$

and the corollary correctly predicts $n_1 n_2 + n_2(n_2 - 1) + n_2 = n_1(n_1 + n_2)$.

4.2 Pointed Refinements

Let us look at an example which breaks injectivity. This is the kind of complication the theorem is staying cautiously away from by asking the t s to be homogeneous.

The set of nodes $V_s = \{1, 2\}$ is represented as subscripts to agents below; the subscripts to the y sites, y_0 and y_1 , denote bindings to agents with only one site and different names (to save space):

$$\begin{array}{ccc} s = A(x^1)_1, A(x^1)_2 & \xrightarrow{I} & t_0 = A(x^1, y_0)_1, A(x^1)_2 \\ \downarrow I & \swarrow I \text{ (dotted)} & \downarrow I \\ t_1 = A(x^1)_1, A(x^1, y_1)_2 & \xrightarrow{I} & x = A(x^1, y_0)_1, A(x^1, y_1)_2 \end{array}$$

If one refers to the situation of (1), the unique possible candidate conjugating α , i.e., the unique diagonal that makes both triangle commute, fails to be a morphism. That means that t_0, t_1 provide really distinct extensions of $f(s)$ in x and form an ambiguous decomposition of s . Indeed, applying (wrongly since the t_i s are not homogeneous) the refinement formula (3) betrays this redundancy problem since $||[s, x]|| = 2$ while $||[s, t_i]||/||[t_i, t_i]||[t_i, x]|| = 2$.

To deal with a case such as this one, one needs to break the symmetry. To do this, a possibility is to work out the static part of the refinement formula in a *pointed* subcategory where objects have in addition to their usual structure a distinguished node per connected component, and arrows are asked to preserve them. Then one can replace homogeneity by a weaker requirement, namely that across all expansions of s no two agents with the same coordinates with respect to a distinguished node differ in their signature. In the example above, that would force to decide whether the additional binding is to sit on the distinguished node or not, and *then* both extensions would become truly distinct and unambiguous. Obviously a little more work is needed to say with complete confidence that this will work, but it seems it will.

5 Rule Refinements

Now that we know how to refine objects, we will proceed to the case of rules.

5.1 Action, Rules, Events

An atomic action on s is one of the following:

- an edge addition $+(u, x, y, v)$
- an edge deletion $-(u, x, y, v)$
- an agent addition $+(A, \sigma)$ with A a name, σ a set of free sites
- an agent deletion $-(u)$ with $u \in V_s$, $v \in V_s$, $x \in \sigma_s(u)$, and $y \in \sigma_s(v)$.

An action on s is a finite sequence of atomic actions on s . An atomic action is well defined on s :

- if $\alpha = +(u, x, y, v)$, when both (u, x) and (v, y) are free in s ,
- if $\alpha = -(u, x, y, v)$, when $(u, x, y, v) \in \mu_s$.

This notion extends readily to non-atomic actions; we consider only well-defined actions hereafter.

Definition 5. A rule is a triple $r = s, \alpha, \tau$ where:

- s in an object,
- α is an action on s ,
- and τ a rate which can be any positive real number.

We write $\alpha \cdot s$ for the effect of the action α on s .

Given $f \in [s, x]$ and α there is an obvious definition of the transport of α along f , written $f(\alpha)$, and it is easy to verify that $f(\alpha)$ is itself a well-defined action on x if α is a well-defined action on s (condition 3b) is crucial though).

Definition 6. A set R of rules defines a labelled transition relation:

$$x \xrightarrow{f}^{s, \alpha, \tau} f(\alpha) \cdot x \quad (4)$$

where $s, \alpha, \tau \in R$, and $f \in [s, x]$.

The labelled transition system just defined can be enriched quantitatively in a way that generalizes the notion of stochastic Petri nets [20] (Petri nets correspond to the case of a uniformly empty signature $\Sigma = \emptyset$).

To do this we need to define the activity of a rule.

Definition 7. Given an object x and a rule $r = s, \alpha, \tau$, the activity of r at x is $a(x, r) := \tau[s, x]$, and the global activity of a set of rules R at x is $a(x) := \sum_{r \in R} a(x, r)$.

Supposing $a(x) > 0$, the probability at x that the next event is $f \in [s, x]$ is $p(x, f) := \tau/a(x)$, and the subsequent time advance is a random variable $\delta t(x)$ such that $p(\delta t(x) > t) := e^{-a(x)t}$. For our present purposes, all we need to remember is that the quantitative structure of the transition system is entirely determined by the activities of its rules. In fact this means our result will hold for a larger class of stochastic system -for what it is worth.

5.2 The Main Result

Given a rule $r = s, \alpha, \tau$ and $\theta \in [s, t]$, we define $\theta(r) := \theta(s), \theta(\alpha), \tau$.

We say r, r' are isomorphic rules, written $r \simeq r'$, if there is an isomorphism $\theta \in [s, s']$ such that $r' = \theta(r)$. If that is the case then r and $\theta(r)$ have isomorphic transitions:

$$x \xrightarrow{f}^r f(\alpha) \cdot x \Leftrightarrow x \xrightarrow{f\theta^{-1} \in [\theta(s), x]}^{\theta(r)} f\theta^{-1}(\theta(\alpha)) \cdot x$$

and in particular the same activity $a(r, x) = a(\theta(r), x)$.

Definition 8 (rule refinement). Given s, Σ such that $s \leq \Sigma$ and $r = s, \alpha, \tau$, the refinement of r under Σ is the following family of rules:

$$\Sigma(s, \alpha, \tau) := (t, \phi(\alpha), \tau; t \in \Sigma(s), \phi \in [s, t]^e/[t, t]) \quad (5)$$

where the notation $\phi \in [s, t]^e/[t, t]$ means that for each t , one selects one $\phi \in [s, t]^e$ per symmetry class on t (the equivalence relation $\exists \theta \in [t, t] : \phi = \theta\phi'$).

It is easily seen that the particular selection made is irrelevant, but one has to choose one to define refinement as a syntactic transformation.

Note also that the above family can have isomorphic or even identical rules, it is important to have them all, *i.e.*, *stricto sensu* the expansion is a multiset of rules not a set. However one can always pack n isomorphic copies together by choosing a representative and multiplying its rate by n so we carry on with our slight abuse of terminology.

Given R a rule set, r a rule in R , we write $R[r \setminus \Sigma(r)]$ for the rule set obtained by replacing r with $\Sigma(r)$.

We write $r = s, \alpha, \tau \leq \Sigma$ if $s \leq \Sigma$, and $R \leq \Sigma$ if for all $r \in R$, $r \leq \Sigma$.

Theorem 2. *Given R, Σ , such that $R \leq \Sigma$, one has $R[r \setminus \Sigma(r)] \leq \Sigma$, and R and $R[r \setminus \Sigma(r)]$ determine the same stochastic transition system over Σ -homogeneous objects.*

Proof: By Th. 1 events $f \in [s, x]$ associated to rule $r = s, \alpha, \tau$ are in one-one correspondence with factorizations $f = \gamma\phi$ via some t , and therefore determine a unique matching refined event γ . This refined event has the same effect as f since:

$$x \xrightarrow[\gamma \in [t = \phi(s), x]]{t, \phi, \tau} \gamma\phi(\alpha) \cdot x = f(\alpha) \cdot x$$

so r and its refinements are equally likely and have the same effect on the underlying state x ; hence their stochastic transition systems are the same. \square

Note that the activity of t, ϕ, τ in the refined system is $\tau|[t, x]|$ so the cumulative activity of the refined rules is:

$$\sum_{t \in \Sigma(s)} \sum_{\phi \in [s, t]^e / [t, t]} \tau|[t, x]| = \sum_{t \in \Sigma(s)} \tau|[s, t]^e| / |[t, t]| |[t, x]| = a(r, x)$$

by Coro. 1, so we can directly derive the fact that the refined rules have the same activity, but we also needed to prove they have the same effect.

5.3 Example Concluded

We can now conclude our initial example.

There we had $s := C(), B()$, and:

$$\begin{aligned} t_1 &:= C(x^1), B(x^1) \\ t_2 &:= C(x^1), B(x^1), C(x^2), B(x^2) \\ t_3 &:= C(x^1), B(x^1), C(x) \end{aligned}$$

Since $|[s, t_2]^e| = 2$ (recall that epis must have images in all connected components), the refinement of r via t_2 will contribute two rules to $\Sigma(r)$ -according to Def. 8. In this particular case the action of the rule to be refined is $\alpha(r) = -B$, and both epimorphisms $\phi \in [s, t_2]$ lead to the same transported action $\phi(-B)$ up to isomorphism. One can then pack them into one rule r_2 , as we did intuitively when we considered the example, and as a consequence the rates must be added. This explains why r_2 has a rate of 2.

6 Conclusion

We have presented in this article the beginning of a theory of refinements for rule-based modelling. Specifically we have defined what constitutes a notion of a good set of refined rules and how, given such a set, one can compute the new refined rates in such a way that the overall activity of the system is preserved and the underlying stochastic semantics therefore unchanged. We have suggested two improvements to extend the type of refinement one can consider.

We have also shown how one can use such refinements to obtain a complete expansion (at least in principle), a construction which could be useful in practice

to get cheap and fast approximations of a system. We have further shown by examples that refinements can be useful to modulate the influence of the context in which a rule is applied.

A point worth commenting in this conclusion is that the formulas obtained in our two main results, Th. 1 and 2, are couched in rather general terms and are likely to be of a larger relevance than the particular case of graph-rewriting we were contemplating here. In particular the epi-mono factorization system which we rely on implicitly for the concrete case we have treated would point to a more abstract approach. That in itself is valuable since such combinatorial results as we have presented here can become nearly intractable if looked at in a too concrete way. This in fact is one of the reasons why we framed our results in a categorical language which has revealed the pervasiveness of symmetries (the other reason is that the syntax is simpler to deal with). It would be particularly interesting to recast the theory in the axiomatic framework of adhesive categories [21], with a view on understanding the formula as a traditional partition formula (which it is, at least intuitively).

A longer term goal that this preliminary work might help to reach is that of finding exact model reduction techniques. This needs to lift a key assumption made here, namely that refinements are made of mutually exclusive sub-cases.

References

1. Orton, R.J., Sturm, O.E., Vyshemirsky, V., Calder, M., Gilbert, D.R., Kolch, W.: Computational modelling of the receptor tyrosine kinase activated MAPK pathway. *Biochemical Journal* 392(2), 249–261 (2005)
2. Söderberg, B.: General formalism for inhomogeneous random graphs. *Physical Review E* 66(6), 66121 (2002)
3. Gillespie, D.T.: Exact stochastic simulation of coupled chemical reactions. *J. Phys. Chem.* 81, 2340–2361 (1977)
4. Danos, V., Feret, J., Fontana, W., Krivine, J.: Abstract interpretation of cellular signalling networks. In: Logozzo, F., et al. (eds.) *VMCAI 2008*. LNCS, vol. 4905, pp. 83–97. Springer, Heidelberg (2008)
5. Danos, V., Feret, J., Fontana, W., Krivine, J.: Scalable simulation of cellular signaling networks. In: Shao, Z. (ed.) *APLAS 2007*. LNCS, vol. 4807, pp. 139–157. Springer, Heidelberg (2007)
6. Danos, V., Feret, J., Fontana, W., Harmer, R., Krivine, J.: Rule-based modelling of cellular signalling. In: Caires, L., Vasconcelos, V.T. (eds.) *CONCUR*. LNCS, vol. 4703. Springer, Heidelberg (2007)
7. Danos, V.: Agile modelling of cellular signalling. In: *Proceedings of ICCMSE (2007)*
8. Blinov, M.L., Faeder, J.R., Goldstein, B., Hlavacek, W.S.: A network model of early events in epidermal growth factor receptor signaling that accounts for combinatorial complexity. *BioSystems* 83, 136–151 (2006)
9. Hlavacek, W.S., Faeder, J.R., Blinov, M.L., Posner, R.G., Hucka, M., Fontana, W.: Rules for Modeling Signal-Transduction Systems. *Science's STKE* 2006(344) (2006)
10. Blinov, M.L., Yang, J., Faeder, J.R., Hlavacek, W.S.: Graph theory for rule-based modeling of biochemical networks. In: Abadi, M., de Alfaro, L. (eds.) *CONCUR 2005*. LNCS, vol. 3653. Springer, Heidelberg (2005)

11. Faeder, J.R., Blinov, M.L., Goldstein, B., Hlavacek, W.S.: Combinatorial complexity and dynamical restriction of network flows in signal transduction. *Systems Biology* 2(1), 5–15 (2005)
12. Regev, A., Silverman, W., Shapiro, E.: Representation and simulation of biochemical processes using the π -calculus process algebra. In: Altman, R.B., Dunker, A.K., Hunter, L., Klein, T.E. (eds.) *Pacific Symposium on Biocomputing*, vol. 6, pp. 459–470. World Scientific Press, Singapore (2001)
13. Priami, C., Regev, A., Shapiro, E., Silverman, W.: Application of a stochastic name-passing calculus to representation and simulation of molecular processes. *Information Processing Letters* (2001)
14. Regev, A., Shapiro, E.: Cells as computation. *Nature* 419 (September 2002)
15. Priami, C., Quaglia, P.: Beta binders for biological interactions. In: Danos, V., Schachter, V. (eds.) *CMSB 2004. LNCS (LNBI)*, vol. 3082, pp. 20–33. Springer, Heidelberg (2005)
16. Danos, V., Krivine, J.: Formal molecular biology done in CCS. In: *Proceedings of BIO-CONCUR 2003*, Marseille, France. *Electronic Notes in Theoretical Computer Science*, vol. 180, pp. 31–49. Elsevier, Amsterdam (2003)
17. Regev, A., Panina, E.M., Silverman, W., Cardelli, L., Shapiro, E.: BioAmbients: an abstraction for biological compartments. *Theoretical Computer Science* 325(1), 141–167 (2004)
18. Cardelli, L.: Brane calculi. In: *Proceedings of BIO-CONCUR 2003*. *Electronic Notes in Theoretical Computer Science*, vol. 180. Elsevier, Amsterdam (2003)
19. Calder, M., Gilmore, S., Hillston, J.: Modelling the influence of RKIP on the ERK signalling pathway using the stochastic process algebra PEPA. In: Priami, C., Ingólfssdóttir, A., Mishra, B., Riis Nielson, H. (eds.) *Transactions on Computational Systems Biology VII. LNCS (LNBI)*, vol. 4230, pp. 1–23. Springer, Heidelberg (2006)
20. Gillespie, D.T.: A general method for numerically simulating the stochastic time evolution of coupled chemical reactions. *J. Comp. Phys.* 22, 403–434 (1976)
21. Lack, S., Sobocinski, P.: Adhesive and quasiadhesive categories. *Theoretical Informatics and Applications* 39(3), 511–546 (2005)