

Product Lines of Theorems

Benjamin Delaware William R. Cook Don Batory

Department of Computer Science
University of Texas at Austin
{bendy,wcook,batory}@cs.utexas.edu

Abstract

Mechanized proof assistants are powerful verification tools, but proof development can be difficult and time-consuming. When verifying a family of related programs, the effort can be reduced by proof reuse. In this paper, we show how to engineer product lines with theorems and proofs built from feature modules. Each module contains proof fragments which are composed together to build a complete proof of correctness for each product. We consider a product line of programming languages, where each variant includes metatheory proofs verifying the correctness of its semantic definitions. This approach has been realized in the Coq proof assistant, with the proofs of each feature independently certifiable by Coq. These proofs are composed for each language variant, with Coq mechanically verifying that the composite proofs are correct. As validation, we formalize a core calculus for Java in Coq which can be extended with any combination of casts, interfaces, or generics.

Categories and Subject Descriptors D.3.1 [Programming Languages]: D.3.1 Formal Definitions and Theory

General Terms Design, Theory, Verification

Keywords Feature-Orientation, Mechanized Metatheory, Product Line Verification

1. Introduction

Mechanized theorem proving is hard: large-scale proof developments [13, 16] take multiple person-years and consist of tens of thousand lines of proof scripts. Given the effort invested in formal verification, it is desirable to reuse as much of the formalization as possible when developing similar proofs. The problem is compounded when verifying members of a *product line* – a family of related systems [2, 5] –

in which the prospect of developing and maintaining individual proofs for each member is untenable.

Product lines can be decomposed into *features* – units of functionality. By selecting and composing different features, members of a product line can be synthesized. The challenge of feature modules for software product lines is that their contents cut across normal object-oriented boundaries [5, 25]. The same holds for proofs. Feature modularization of proofs is an open, fundamental, and challenging problem.

Surprisingly, the programming language literature is replete with examples of product lines which include proofs. These product lines typically only have two members, consisting of a core language such as *Featherweight Java (FJ)* [14], and an updated one with modified syntax, semantics, and proofs of correctness. Indeed, the original FJ paper also presents *Featherweight Generic Java (FGJ)*, a modified version of FJ with support for generics. An integral part of any type system are the metatheoretic proofs showing *type soundness* – a guarantee that the type system statically enforces the desired run-time behavior of a language, typically preservation and progress [24].

Typically, each research paper only adds a single new feature to a core calculus, and this is accomplished manually. Reuse of existing syntax, semantics, and proofs is achieved by copying existing rules, and in the case of proofs, following the structure of the original proof with appropriate updates. As more features are added, this manual process grows increasingly cumbersome and error prone. Further, the enhanced languages become more difficult to maintain. Adding a feature requires changes that cut across the normal structural boundaries of a language – its syntax, operational semantics, and type system. Each change requires arduously rechecking existing proofs by hand.

Using theorem provers to mechanically formalize languages and their metatheory provides an interesting testbed for studying the modularization of product lines which include proofs. By implementing an extension in the proof assistant as a *feature module*, which includes updates to existing definitions and proofs, we can compose feature modules to build a completely mechanized definition of an enhanced language, with the proofs mechanically checked by the theorem prover. Stepwise development is enabled, and it

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

OOPSLA'11, October 22–27, 2011, Portland, Oregon, USA.
Copyright © 2011 ACM 978-1-4503-0940-0/11/10...\$10.00

FJ Expression Syntax			FGJ Expression Syntax	
$ \begin{aligned} e &::= x \\ & e.f \\ & e.m(\bar{e}) \\ & \text{new } C(\bar{e}) \\ & (C) e \end{aligned} $		\Rightarrow	$ \begin{aligned} e &::= x \\ & e.f \\ & e.m \langle \bar{T} \rangle^\beta (\bar{e}) \\ & \text{new } C \langle \bar{T} \rangle^\beta (\bar{e}) \\ & (C \langle \bar{T} \rangle^\beta) e \end{aligned} $	
FJ Subtyping	$T <: T$		FGJ Subtyping	$\Delta^\delta \vdash T <: T$
$ \frac{S <: T \quad T <: V}{S <: V} \quad (\text{S-TRANS}) $		\Rightarrow	$ \frac{\Delta \vdash X <: \Delta(X) \text{ (GS-VAR)}^\alpha \quad \Delta^\delta \vdash S <: T \quad \Delta^\delta \vdash T <: V}{\Delta^\delta \vdash S <: V} \quad (\text{GS-TRANS}) $	
$T <: T \quad (\text{S-REFL})$		\Rightarrow	$\Delta^\delta \vdash T <: T \quad (\text{GS-REFL})$	
$ \frac{\text{class } C \text{ extends } D \{ \dots \}}{C <: D} \quad (\text{S-DIR}) $		\Rightarrow	$ \frac{\text{class } C \langle \bar{X} < \bar{N} \rangle^\beta \text{ extends } D \langle \bar{V} \rangle^\beta \{ \dots \}}{\Delta^\delta \vdash C \langle \bar{T} \rangle^\beta <: [\bar{T}/\bar{X}] D \langle \bar{V} \rangle^\beta} \quad (\text{GS-DIR}) $	
FJ New Typing	$\Gamma \vdash e : T$		FGJ New Typing	$\Delta;^\delta \Gamma \vdash e : T$
$ \frac{\text{fields}(C) = \bar{V} \bar{f} \quad \Gamma \vdash \bar{e} : \bar{U} \quad \bar{U} <: \bar{V}}{\Gamma \vdash \text{new } C(\bar{e}) : C} \quad (\text{T-NEW}) $		\Rightarrow	$ \frac{\Delta \vdash C \langle \bar{T} \rangle^\gamma \quad \text{fields}(C \langle \bar{T} \rangle^\beta) = \bar{V} \bar{f} \quad \Delta;^\delta \Gamma \vdash \bar{e} : \bar{U} \quad \Delta^\delta \vdash \bar{U} <: \bar{V}}{\Delta;^\delta \Gamma \vdash \text{new } C \langle \bar{T} \rangle^\beta (\bar{e}) : C} \quad (\text{GT-NEW}) $	

Figure 1: Selected FJ Definitions with FGJ Changes Highlighted

is possible to start with a core language and add features to progressively build a family or product line of more detailed languages with tool support and less difficulty.

In this paper, we present a methodology for feature-oriented development of a language using a variant of FJ as an example. We implement feature modules in Coq [8] and demonstrate how to build mechanized proofs that can adapt to new extensions. Each module is a separate Coq file which includes inductive definitions formalizing a language and proofs over those definitions. A feature's proofs can be independently checked by Coq, with no need to recheck existing proofs after composition. We validate our approach through the development of a family of feature-enriched languages, culminating in a generic version of FJ with generic interfaces. Though our work is geared toward mechanized metatheory in Coq, the techniques should apply to different formalizations in other higher-order proof assistants.

2. Background

2.1 On the Importance of Engineering

Architecting product lines (sets of similar programs) has long existed in the software engineering community [18, 23]. So too has the challenge of achieving object-oriented code reuse in this context [26, 30]. The essence of reusable designs – be they code or proofs – is engineering. There is no

magic bullet, but rather a careful trade-off between flexibility and specialization. A spectrum of common changes must be explicitly anticipated in the construction of a feature and its interface. This is no different from using abstract classes and interfaces in the design of OO frameworks [7]. The plug-compatibility of features is not an after-thought but is essential to their design and implementation, allowing the easy integration of new features *as long as* they satisfy the assumptions of existing features. Of course, unanticipated features do arise, requiring a refactoring of existing modules. Again, this is no different than typical software development. Exactly the same ideas hold for modularizing proofs. It is against this backdrop that we motivate our work.

2.2 A Motivating Example

Consider adding generics to the calculus of FJ [14] to produce the FGJ calculus. The required changes are woven throughout the syntax and semantics of FJ. The left-hand column of Figure 1 presents a subset of the syntax of FJ, the rules which formalize the subtyping relation that establish the inheritance hierarchy, and the typing rule that ensures expressions for object creation are well-formed. The corresponding definitions for FGJ are in the right-hand column.

The categories of changes are tagged in Figure 1 with Greek letters:

FJ Fields of a Supertype Lemma		FGJ Fields of a Supertype Lemma
Lemma 2.1. If $S \leq T$ and $\text{fields}(T) = T \bar{f}$, then $\text{fields}(S) = \bar{S} \bar{g}$ and $S_i = T_i$ and $g_i = f_i$ for all $i \leq \#(f)$.	\Rightarrow	Lemma 2.2. If $\Delta^\delta \vdash S \leq T$ and $\text{fields}(\text{bound}_\Delta(T)^\eta) = \bar{T} \bar{f}$, then $\text{fields}(\text{bound}_\Delta(S)^\eta) = \bar{S} \bar{g}$, $S_i = T_i$ and $g_i = f_i$ for all $i \leq \#(f)$.
<i>Proof.</i> By induction on the derivation of $S \leq T$		<i>Proof.</i> By induction on the derivation of $\Delta^\delta \vdash S \leq T$
		Case GS-VAR $^\alpha S = X$ and $T = \Delta(X)$. Follows immediately from the fact that $\text{bound}_\Delta(\Delta(X)) = \Delta(X)$ by the definition of bound .
Case S-REFL $S = T$. Follows immediately.	\Rightarrow	Case GS-REFL $S = T$. Follows immediately.
Case S-TRANS $S \leq V$ and $V \leq T$. By the inductive hypothesis, $\text{fields}(V) = \bar{V} \bar{h}$ and $V_i = T_i$ and $h_i = f_i$ for all $i \leq \#(f)$. Again applying the inductive hypothesis, $\text{fields}(S) = \bar{S} \bar{g}$ and $S_i = V_i$ and $g_i = h_i$ for all $i \leq \#(h)$. Since $\#(f) \leq \#(h)$, the conclusion is immediate.	\Rightarrow	Case GS-TRANS $\Delta^\delta \vdash S \leq V$ and $\Delta^\delta \vdash V \leq T$. By the inductive hypothesis, $\text{fields}(\text{bound}_\Delta(V)^\eta) = \bar{V} \bar{h}$ and $V_i = T_i$ and $h_i = f_i$ for all $i \leq \#(f)$. Again applying the inductive hypothesis, $\text{fields}(\text{bound}_\Delta(S)^\eta) = \bar{S} \bar{g}$ and $S_i = V_i$ and $g_i = h_i$ for all $i \leq \#(h)$. Since $\#(f) \leq \#(h)$, the conclusion is immediate.
Case S-DIR $S = C$, $T = D$, class C extends $D \{\bar{S} \bar{g}; \dots\}$. By the rule F-CLASS, $\text{fields}(C) = \bar{U} \bar{f}; \bar{S} \bar{g}$, where $\bar{U} \bar{f} = \text{fields}(D)$, from which the conclusion is immediate.	\Rightarrow	Case GS-DIR $S = C \langle \bar{T} \rangle^\beta$, $T = [\bar{T}/\bar{X}]^\eta D \langle \bar{V} \rangle^\beta$, class $C \langle \bar{X} \triangleleft N \rangle^\beta$ extends $D \langle \bar{V} \rangle^\beta \{\bar{S} \bar{g}; \dots\}$. By the rule F-CLASS, $\text{fields}(C \langle \bar{T} \rangle^\beta) = \bar{U} \bar{f}; [\bar{T}/\bar{X}]^\eta \bar{S} \bar{g}$, where $\bar{U} \bar{f} = \text{fields}([\bar{T}/\bar{X}]^\eta D \langle \bar{V} \rangle^\beta)$. By definition, $\text{bound}_\Delta(V) = V$ for all non-variable types V , from which the conclusion is immediate.

Figure 2: An Example FJ Proof with FGJ Changes Highlighted

- α . *Adding new rules or pieces of syntax.* FGJ adds type variables to parameterize classes and methods. The subtyping relation adds the GS-VAR rule for this new kind of type.
- β . *Modifying existing syntax.* FGJ adds type parameters to method calls, object creation, casts, and class definitions.
- γ . *Adding new premises to existing typing rules.* The updated GT-NEW rule includes a new premise requiring that the type of a new object must be well-formed.
- δ . *Extending judgment signatures.* The added rule GS-VAR looks up the bound of a type variable using a typing context, Δ . This context must be added to the signature of the subtyping relation, transforming all occurrences to a new ternary relation.
- η . *Modifying premises and conclusions in existing rules.* The type parameters used for the parent class D in a class definition are instantiated with the parameters used for the child in the conclusion of GS-DIR.

In addition to syntax and semantics, the definitions of FJ and FGJ include proofs of progress and preservation for their type systems. With each change to a definition, these proofs must also be updated. As with the changes to definitions in Figure 1, these changes are threaded throughout existing proofs. Consider the related proofs in Figure 2 of a lemma used in the proof of progress for both languages. These lem-

mas are used in the same place in the proof of progress and are structurally similar, proceeding by induction on the derivation of the subtyping judgment. The proof for FGJ has been adapted to reflect the changes that were made to its definitions. These changes are highlighted in Figure 2 and marked with the kind of definitional change that triggered the update. Throughout the lemma, the signature of the subtyping judgment has been altered include a context for type variables ^{δ} . The statement of the lemma now uses the auxiliary bound function, due to a modification to the premises of the typing rule for field lookup ^{η} . These changes are not simply syntactic: both affect the applications of the inductive hypothesis in the GS-TRANS case. The proof must now include a case for the added GS-VAR subtyping rule ^{α} . The case for GS-DIR requires the most drastic change, as the existing proof for that case is modified to include an additional statement about the behavior of bound .

As more features are added to a language, its metatheoretic proofs of correctness grow in size and complexity. In addition, each different selection of features produces a new language with its own syntax, type system, operational semantics. While the proof of type safety is similar for each language, (potentially subtle) changes occur throughout the proof depending on the features included. By modularizing the type safety proof into distinct features, each language variant is able to build its type safety proof from a com-

mon set of proofs. There is no need to manually maintain separate proofs for each language variant. As we shall see, this allows us to add new features to an existing language in a structured way, exploiting existing proofs to build more feature-rich languages that are semantically correct.

We demonstrate in the following sections how each kind of extension to a language’s syntax and semantics outlined above requires a structural change to a proof. Base proofs can be updated by filling in the pieces required by these changes, enabling reuse of potentially complex proofs for a number of different features. Further, we demonstrate how this modularization can be achieved within the Coq proof assistant. In our approach, each feature has a set of assumptions that serve as variation points, allowing a feature’s proofs to be checked independently. As long as an extension provides the necessary proofs to satisfy these assumptions, the composite proof is guaranteed to hold for any composed language. Generating proofs for a composed language is thus a straightforward check that all dependencies are satisfied, with no need to recheck existing proofs.

3. The Structure of Features

Features impose a mathematical structure on the universe of programming languages (including type systems and proofs of correctness) that are to be synthesized. In this section, we review concepts that are essential to our work.

3.1 Features and Feature Compositions

We start with a *base* language or *base* feature to which extensions are added. It is modeled as a constant or zero-ary function. For our study, the *core Featherweight Java* cFJ language is a cast-free variant of FJ. (This omission is not without precedent, as other core calculi for Java [28] omit casts). There are also *optional* features, which are unary functions, that extend the base or other features:

cFJ	core Featherweight Java
Cast	adds casts to expressions
Interface	adds interfaces
Generic	adds type parameters

Assuming no feature interactions, features are composed by function composition (\cdot). Each expression corresponds to a composite feature or a distinct language. Composing Cast with cFJ builds the original version of FJ:

```

cFJ      // Core FJ
Cast · cFJ  // Original FJ [14]
Interface · cFJ  // Core FJ with Interfaces
Interface · Cast · cFJ  // Original FJ with Interfaces
Generic · cFJ  // Core Featherweight Generic Java
Generic · Cast · cFJ  // Original FGJ
Generic · Interface · cFJ  // core Generic FJ with
                          // Generic Interfaces
Generic · Interface
  · Cast · cFJ  // FGJ with
                // Generic Interfaces

```

3.2 Feature Models

Not all compositions of features are meaningful. Some features require the presence or absence of other features. An *if* statement, for example, requires a feature that introduces some notion of booleans to use in test conditions. Feature models define the compositions of features that produce meaningful languages. A *feature model* is a context sensitive grammar, consisting of a context free grammar whose sentences define a superset of all legal feature expressions, and a set of constraints (the context sensitive part) that eliminates nonsensical sentences [6]. The grammar of feature model P (below) defines eight sentences (features *k*, *i*, *j* are optional; *b* is mandatory). Constraints limit legal sentences to those that have at least one optional feature, and if feature *k* is selected, so too must *j*.

```

P : [k] [i] [j] b;      // context free grammar
    k ∨ j ∨ i;          // additional constraints
    k ⇒ j;

```

Given a sentence of a feature model (kjb) a dot-product is taken of its terms to map it to an expression ($k \cdot j \cdot b$). A language is synthesized by evaluating the expression. The feature model L that used in our study is context free:

```
L : [Generic] [Interface] [Cast] cFJ;
```

3.3 Multiple Representations of Languages

Every base language (cFJ) has multiple representations: its syntax s_{cFJ} , operational semantics o_{cFJ} , type system t_{cFJ} , and metatheory proofs p_{cFJ} . A base language is a tuple of representations $cFJ = [s_{cFJ}, o_{cFJ}, t_{cFJ}, p_{cFJ}]$. An optional feature *i* extends each representation: the language’s syntax is extended with new productions Δs_i , its operational semantics are extended by modifying existing rules and adding new rules to handle the updated syntax Δo_i , etc. Each of these changes is modeled by a unary function. Feature *i* is a tuple of such functions $i = [\Delta s_i, \Delta o_i, \Delta t_i, \Delta p_i]$ that update each representation of a language.

The representations of a language are computed by composing tuples element-wise. The tuple for language FJ = Cast · cFJ is:

```

FJ = Cast · cFJ
    = [ $\Delta s_C, \Delta o_C, \Delta t_C, \Delta p_C$ ] · [ $s_{cFJ}, o_{cFJ}, t_{cFJ}, p_{cFJ}$ ]
    = [ $\Delta s_C \cdot s_{cFJ}, \Delta o_C \cdot o_{cFJ}, \Delta t_C \cdot t_{cFJ}, \Delta p_C \cdot p_{cFJ}$ ]

```

That is, the syntax of FJ is the syntax of the base s_{cFJ} composed with extension Δs_C , the semantics of FJ are the base semantics o_{cFJ} composed with extension Δo_C , and so on. In this way, all parts of a language are updated lock-step when features are composed. See [5, 12] for generalizations of these ideas.

3.4 Feature Interactions

Feature interactions are ubiquitous. Consider the Interface feature which introduces syntax for interface declarations:

```
J ::= interface I {Mty}
```

This declaration may be changed by other features. When *Generic* is added, the syntax of an interface declaration must be updated to include type parameters:

$$J ::= \text{interface } I \quad \langle \overline{X \triangleleft N} \rangle \quad \{\overline{Mty}\}$$

Similarly, any proofs in *Generic* that induct over the derivation of the subtyping judgement must add new cases for the subtyping rule introduced by the *Interface* feature. Such proof updates are necessary only when *both* features are present. The set of additional changes made across all representations is the *interaction* of these features, written *Generic#Interface*.¹

Until now, features were composed by only one operation (dot or \cdot). Now we introduce two additional operations: product (\times) and interaction ($\#$). When designers want a set of features, they really want the \times -product of these features, which includes the dot-product of these features *and* their interactions. The \times -product of features f and g is:

$$f \times g = (f\#g) \cdot f \cdot g \quad (1)$$

where $\#$ distributes over dot and $\#$ takes precedence over dot:

$$f\#(g \cdot h) = (f\#g) \cdot (f\#h) \quad (2)$$

That is, the interaction of a feature with a dot-product is the dot-product of their interactions. \times is right-associative and $\#$ is commutative and associative.²

The connection of \times and $\#$ to prior discussions is simple. *To allow for feature interactions*, a sentence of a feature model ($\langle kjb \rangle$) is mapped to an expression by a \times -product of its terms ($k \times j \times b$). Equations (1) and (2) are used to reduce an expression with \times operations to an expression with only dot and $\#$, as below:

$$\begin{aligned} p &= k \times j \times b && // \text{def of } p \\ &= k \times (j\#b \cdot j \cdot b) && // (1) \\ &= k\#(j\#b \cdot j \cdot b) \cdot k \cdot (j\#b \cdot j \cdot b) && // (1) \\ &= k\#j\#b \cdot k\#j \cdot k\#b \cdot k \cdot j\#b \cdot j \cdot b && // (2) \end{aligned} \quad (4)$$

¹ Our *Generic#Interface* example is isomorphic to the classical example of fire and flood control [15]. Let b denote the design of a building. The flood control feature adds water sensors to every floor of b . If standing water is detected, the water main to b is turned off. The fire control feature adds fire sensors to every floor of b . If fire is detected, sprinklers are turned on. Adding flood or fire control to the building (e.g. $\text{flood} \cdot b$ and $\text{fire} \cdot b$) is straightforward. However, adding both ($\text{flood} \cdot \text{fire} \cdot b$) is problematic: if fire is detected, the sprinklers turn on, standing water is detected, the water main is turned off, and the building burns down. This is not the intended semantics of the composition of the flood, fire, and b features. The fix is to apply an additional extension, labeled $\text{flood}\#\text{fire}$, which is the interaction of flood and fire. $\text{flood}\#\text{fire}$ represents the changes (extensions) that are needed to make the flood and fire features work correctly together. The correct building design is $\text{flood}\#\text{fire} \cdot \text{flood} \cdot \text{fire} \cdot b$.

² A more general algebra has operations \times , $\#$, and \cdot that are all commutative and associative [4]. This generality is not needed for this paper.

Language p is synthesized by evaluating expression (4). Interpreting modules for individual features like k , j , and b as 1-way feature interactions (where $k\#j$ denotes a 2-way interaction and $k\#j\#b$ is 3-way), the universe of modules in a feature-oriented construction are exclusively those of feature interactions.

An \times -product of n features results in $O(2^n)$ interactions (i.e. all possible feature combinations). Fortunately, the *vast* majority of feature interactions are empty, meaning that they correspond to the identity transformation 1, whose properties are:

$$1 \cdot f = f \cdot 1 = f \quad (3)$$

Most non-empty interactions are pairwise (2-way). Occasionally higher-order interactions arise. The \times -product of *cFJ*, *Interface*, and *Generic* is:

$$\begin{aligned} &\text{Generic} \times \text{Interface} \times \text{cFJ} \\ &= \text{Generic}\#\text{Interface}\#\text{cFJ} \cdot \text{Generic}\#\text{Interface} \\ &\quad \cdot \text{Generic}\#\text{cFJ} \cdot \text{Generic} \cdot \text{Interface}\#\text{cFJ} \\ &\quad \cdot \text{Interface} \cdot \text{cFJ} \\ &= \text{Generic}\#\text{Interface} \cdot \text{Generic} \cdot \text{Interface} \cdot \text{cFJ} \end{aligned}$$

which means that all 2- and 3-way interactions, except *Generic#Interface*, equal 1. In our case study, the complete set of interaction modules that are not equal to 1 is:

Module	Description
cFJ	core Featherweight Java
Cast	cast
Interface	interfaces
Generic	generics
Generic#Interface	generic and interface interactions
Generic#Cast	generic and cast interactions

Each of these interaction modules is represented by a tuple of definitions or a tuple of changes to these definitions.

4. Decomposing a Language into Features

We designed features to be monotonic: what was true before a feature is added remains valid after composition, although the scope of validity may be qualified. This is standard in feature-based designs, as it simplifies reasoning with features [2].

All representations of a language (syntax, operational semantics, type system, proofs) are themselves written in distinct languages. Language syntax uses BNF, operational semantics and type systems use standard rule notations, and metatheoretic proofs are formal proofs in Coq.

Despite these different representations, there are only two kinds of changes that a feature makes to a document: new definitions can be added and existing definitions can be modified. Addition is just the union of definitions. Modification requires definitions to be engineered for change.

In the following sections, we explain how to accomplish addition and modification. We alert readers that our tech-

niques for extending language syntax are identical to extension techniques for the other representations. The critical contribution of our approach is how we guarantee the correctness of composed proofs, the topic of Section 5.1.

4.1 Language Syntax

We use BNF to express language syntax. Figure 3a shows the BNF for expressions in cFJ, Figure 3b the production that the Cast feature adds to cFJ's BNF, and Figure 3c the composition (union) of these productions, that defines the expression grammar of the FJ = Cast · cFJ language (Figure 1).




Figure 3: Union of Grammars

Modifying existing productions requires foresight to anticipate how productions may be changed by other features. (This is no different from object-oriented refactorings that prepare source code for extensions – visitors, frameworks, strategies, etc. – as discussed in Section 2.) Consider the impact of adding the Generics feature to cFJ and Cast: type parameters must be added to the expression syntax of method calls and class types now have type parameters. What we do is to insert *variation points* (VP), a standard concept in product line designs [1], to allow new syntax to appear in a production. For syntax rules, a VP is simply the name of an (initially) empty production.

Figure 4a-b shows the VPs TP_m added to method calls in the cFJ expression grammar and TP_t added to class types in the cFJ and Cast expression grammars. Figure 4c shows the composition (union) of the revised Cast and cFJ expression grammars. Since TP_m and TP_t are empty, Figure 4c can be inlined to produce the grammar in Figure 3c.

Now consider the changes that Generic makes to expression syntax: it redefines TP_m and TP_t to be lists of type parameters, thereby updating all productions that reference these VPs. Figure 5a shows this definition. Figure 5b shows the productions of Figure 4c with these productions inlined, building the expression grammar for Generic · Cast · cFJ.

Replacing an empty production with a non-empty one is a standard programming practice in frameworks (e.g. EJB [19]). Framework hook methods are initially empty and users can override them with a definition that is specific to their context. We do the same here.

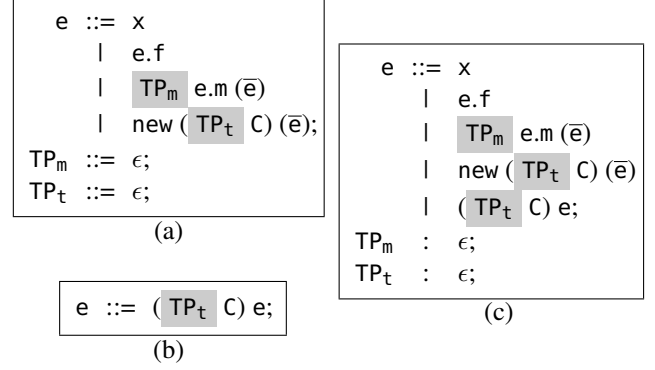


Figure 4: Modification of Grammars




Figure 5: The Effect of Adding Generics to Expressions

These are simple and intuitively appealing techniques for defining and composing language extensions. As readers will see, these same ideas apply to rules and proofs as well.

4.2 Reduction and Typing Rules

The judgments that form the operational semantics and type system of a language are defined by rules. Figure 6a shows the typing rules for cFJ expressions, Figure 6b the rule that the Cast feature adds, and Figure 6c the composition (union) of these rules, defining the typing rules for FJ.

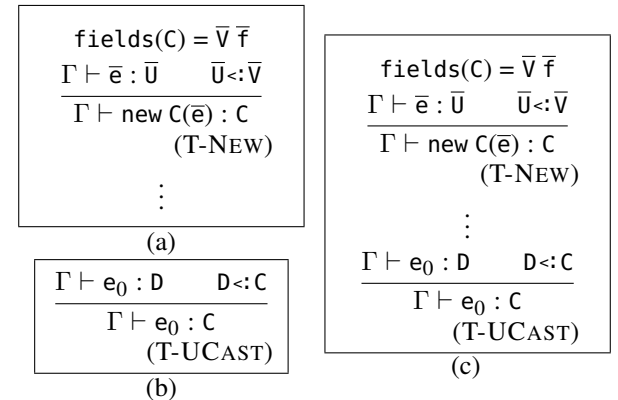


Figure 6: Union of Typing Rules

Modifying existing rules is analogous to language syntax. There are three kinds of VPs for rules: (a) predicates that extend the premise of a rule, (b) relational holes which

To build the typing rules for FGJ, the Generic feature adds non-empty definitions for the $\text{WF}_C(\text{D}, \text{TP}_t \text{ C})$ predicate and for the D relational hole in the cFJ typing definitions. (Compare Figure 6a to its VP-extended counterpart in Figure 7a). Figure 7b shows the non-empty definitions for these VPs introduced by the Generic feature, with Figure 7c showing the T-NEW rule with these definitions inlined.




Variation points also appear in the statements of lemmas and theorems, enabling the construction of feature-extensible proofs. Consider the lemma in Figure 8 with its seven VPs.



Without an accompanying proof, feature-extensible theorem statements are uninteresting. Ideally, a proof should adapt to any VP instantiation or case introduction, allowing the proof to be reused in any target language variant. Of course, proofs must rule out broken extensions which do not guarantee progress and preservation, and admit only “correct” new cases or VP instantiations. This is the key challenge in crafting modular proofs.

The syntax, operational semantics, and typing rules of a language are embedded in Coq as standard inductive data types. The metatheoretic proofs of a language are then written over these encodings. Figure 10a-b gives the Coq definitions for the syntax of Figure 3a and the typing rules of Figure 7a. A feature module in Coq is realized as a Coq file containing its definitions and proofs. The target language is itself a Coq file which combines the definitions and proofs from a set of Coq feature modules.



601

$ \begin{array}{c} TP_t : \epsilon; \quad TP_m : \epsilon; \quad \mu : \epsilon; \\ D := \epsilon \quad \frac{\mathbb{T}}{WF_{mc}(\epsilon, \epsilon, \bar{T})} \\ \frac{\mathbb{T}}{WF_{ne}(\epsilon, C)} \quad \Phi_M(\epsilon, \epsilon, \bar{T}) := \bar{T} \end{array} $	<p>Lemma 4.1 (cFJ Well-Formed MBody). <i>If $mtype(m, C) = \bar{V} \rightarrow V$ and \mathbb{T} with $mbody(m, C) = \bar{x}.e$ where \mathbb{T}, then there exists some N and S such that $\vdash C <: N$ and $\vdash S <: V$ and $\bar{x} : \bar{V}, this : N \vdash e : S$.</i></p>
$ \begin{array}{c} TP_t : \bar{T}; \quad TP_m : \bar{T}; \quad \mu : \langle \bar{Y} \triangleleft \bar{P} \rangle; \\ D := \Delta \quad \frac{\Delta \vdash \bar{U} \text{ ok} \quad \Delta \vdash \bar{U} <: [\bar{U}/\bar{Y}]\bar{P}}{WF_{mc}(\Delta, \langle \bar{Y} \triangleleft \bar{P} \rangle, \bar{U})} \\ \frac{\Delta \vdash \langle \bar{T} \rangle C \text{ ok}}{WF_{ne}(\Delta, \langle \bar{T} \rangle C)} \quad \Phi_M(\langle \bar{T} \rangle, \langle \bar{Y} \triangleleft \bar{P} \rangle, \bar{U}) := [\bar{T}/\bar{Y}]\bar{U} \end{array} $	<p>Lemma 4.1 (FGJ Well-Formed MBody). <i>If $mtype(m, \langle \bar{T} \rangle C) = \langle \bar{Y} \triangleleft \bar{P} \rangle \bar{V} \rightarrow V$ and $\Delta \vdash \langle \bar{T} \rangle C \text{ ok}$ with $mbody(\langle \bar{U} \rangle, m, \langle \bar{T} \rangle C) = \bar{x}.e$, where $\Delta \vdash \bar{U} \text{ ok}$ and $\Delta \vdash \bar{U} <: [\bar{U}/\bar{Y}]\bar{P}$, then there exists some N and S such that $\Delta \vdash \langle \bar{T} \rangle C <: N$ and $\Delta \vdash S <: [\bar{U}/\bar{Y}]V$ and $\Delta; \bar{x} : [\bar{U}/\bar{Y}]\bar{V}, this : N \vdash e : S$</i></p>

Figure 9: VP Instantiations for cFJ and Generic and the resulting statements of Lemma 4.1 for cFJ and FGJ

the definitions from each of the features. For abstract predicates, the target predicate is the conjunction of all the VP definitions. The Coq encoding of expressions the Cast, and Generic features and the result of their composition with cFJ is given in Figure 11.

<p>Inductive e : Set := cast : C → e → e.</p> <p style="text-align: center;">Cast</p>
<p>Definition TP_m := list Type. Definition TP_t := list Type.</p> <p style="text-align: center;">Generic</p>
<p>Definition TP_m := (list Type, unit). Definition TP_t := (list Type, unit). Inductive C : Set := ty : TP_t → Name → e. Inductive e : Set := e_var : Var → e fd_access : e → F → e m_call : TP_m → e → M → List e → e new : C → List e → e cast : C → e → e.</p> <p style="text-align: center;">Cast · Generic · cFJ</p>

Figure 11: Coq Encoding of Fig. 3b and Fig. 5a-b.

In the discussion so far, composition has been strictly syntactic: definitions are directly unioned together or defaults are replaced. Modular reasoning within a feature requires a more semantic form of composition that is supported by Coq. OO frameworks are implemented using inheritance and mixin layers [3], techniques that are not available in most proof assistants. Our feature modules instead rely on the higher-order parameterization mechanisms of the Coq theorem prover to support case extension and VPs. Modules can now be composed within Coq by instantiating parameterized definitions. Using Coq's native abstraction mechanism enables independent certification of each of the feature modules.

Figure 12 shows a concrete example of crafting an extensible inductive definition in Coq. The target language of $FJ = \text{Cast} \cdot \text{cFJ}$ is built by importing the Coq modules for features cFJ and Cast. The target syntax is defined as a new data type, e, with data constructors cFJ and Cast from each feature. Each constructor wraps the syntax definitions from their corresponding features, closing the inductive loop by instantiating the abstract parameter e' with e, the data type for the syntax of target language.

<p>Inductive e (e' : Set) : Set := e_var : Var → e fd_access : e' → F → e m_call : e' → M → List e' → e new : Ty → List e' → e.</p> <p style="text-align: center;">cFJ.v</p>
<p>Inductive e (e' : Set) : Set := e_cast : Ty → e' → e.</p> <p style="text-align: center;">cast.v</p>
<p>Require Import cFJ. Require Import cast. Inductive e : Set := cFJ : cFJ.e e → e cast : cast.e e → e.</p> <p style="text-align: center;">FJ.v</p>

Figure 12: Syntax from cFJ and Cast Features and their Union.

These parameters also affect data types which reference open inductive definitions. In particular, the signature of typing rules and the transition relation are now over the parameter used for the final language. Exp_WF from Fig. 10b ranges over the complete set of expressions from the final language, so its signature becomes $\forall e' : \text{Set}, \text{Context} \rightarrow e' \rightarrow \text{Ty} \rightarrow \text{Prop}$. Of course, within a feature module these rules are written over the actual syntax definitions it provides. In order for the signatures to sync up, these

rules are parameterized by a function that injects the syntax defined in the feature module into the syntax of the final language. Since the syntax of a module is always included alongside its typing and reduction rules in the target language, such an injection always exists.

Parameterization also allows feature modules to include VPs, as shown in Figure 13. The VPs in each module are explicitly represented as abstract sets/predicates/functions, as with the parameter TP_m used to extend the expression for method calls in $cFJ.v$. Other features can provide appropriate instantiations for this parameter. In Figure 13, for example, $FGJ.v$ builds the syntax for the target language by instantiating this VP with the definition of TP_m given in $Generic.v$. Alternatively, the syntax of cFJ can be built from the same inductive definition from cFJ using the default definition of TP_m it provides.

```

Definition TP_m := unit.
Inductive cFJ_e (e : Set) (TP_m : Set): Set :=
| e_var : Var → cFJ_e
| fd_access : e → F → cFJ_e
| m_call : TP_m → e → M → List e → cFJ_e
| new : C → List e → cFJ_e.

```

cFJ.v

```

Definition TP_m := List Ty.

```

Generic.v

```

Require Import cFJ.
Require Import Generic.
Definition TP_m := Generic.TP_m.
Inductive e : Set :=
| cFJ : cFJ_e e TP_m → e

```

FGJ.v

Figure 13: Coq Syntax from cFJ with a Variation Point, and its Instantiation in FGJ .

5.1 Crafting Modular Proofs

Rather than writing multiple related proofs, our goal is to create a single proof for a generic statement of a theorem. This proof is then specialized for the target language by instantiating the variation points appropriately. Instead of separately proving the two lemmas in Figure 2, the cFJ feature has a single proof of the generic Lemma 5.1 (Figure 14). This lemma is then specialized to the variants FJ and FGJ shown in Figure 2. The proof now reasons over the generic subtyping rules with variation points, as in the case for **S-Dir** in Figure 14. The definition of these holes depends on the set of features included in the final language, so from the (human or computer) theorem prover's point of view, they are opaque. Thus, this proof becomes stuck when it requires knowledge about behavior of Φ_f .

In order to proceed, the lemma must constrain possible VP instantiations to those that have the properties required by the proof. In the case of Lemma 5.1, this behavior is that

Lemma 5.1. *If $\Delta \vdash S \leq T$ and $\text{fields}(\Phi_f(\Delta, T)) = \bar{T} \bar{f}$, then $\text{fields}(\Phi_f(\Delta, S)) = \bar{S} \bar{g}$, $S_i = T_i$ and $g_i = f_i$ for all $i \leq \#(f)$.*

Case S-Dir

$S = TP_0 C$, CP_0 class C extends $TP_1 D \{ \bar{S} \bar{g}; \dots \}$,
 $T = \Phi_{SD}(TP_0, CP_0, TP_1 D)$.

By the rule F-CLASS, $\text{fields}(\Phi_{SD}(TP_0, CP_0, TP_1 D)) = \bar{U} \bar{h}$ with $\text{fields}(TP_0 C) = \bar{U} \bar{h}$; $\Phi_{SD}(TP_0, CP_0, \bar{S}) \bar{g}$. Assuming that for all class types $TP_2 D'$, $\Phi_f(\Delta, TP_2 D') = TP_2 D'$ and $\Phi_{SD}(TP_0, CP_0, TP_2 D')$ returns a class type, $\Phi_f(\Delta, \Phi_{SD}(TP_0, CP_0, TP_1 D)) = \Phi_{SD}(TP_0, CP_0, TP_1 D)$. It follows that $\bar{T} \bar{f} = \text{fields} \Phi_{SD}(TP_0, CP_0, TP_1 D) = \bar{U} \bar{h}$ from which the conclusion is immediate.

Figure 14: Generic Statement of Lemmas 2.2 and 2.1 and Proof for **S-Dir** Case.

Φ_f must be the identity function for non-variable types and that Φ_{SD} maps class types to class types. For this proof to hold for the target language, the instantiations of Φ_f and Φ_{SD} must have this property. More concretely, the proof assumes this behavior for all instantiations of Φ_f and Φ_{SD} , producing the new generic Lemma 5.2. In order to produce the desired lemma, the target language instantiates the VPs and provides proofs of all the assumed behaviors. Each feature which supplies a concrete realization of a VP also provides the necessary proofs about its behavior. The assumptions of a proof form an explicit interface against which it is written. The interface of a feature module is the union of all these assumptions together with the the set of lemmas about the behavior of its VP instantiations and definitions it provides. As long as the features included in the target language provide proofs satisfying this interface, a feature's generic proofs can be specialized and reused in their entirety.

Lemma 5.2. *As long as $\Phi_f(\Delta, V) = V$ for all non-variable types V and Φ_{SD} maps class types to class types, if $\Delta \vdash S \leq T$ and $\text{fields}(\Phi_f(\Delta, T)) = \bar{T} \bar{f}$, then $\text{fields}(\Phi_f(\Delta, S)) = \bar{S} \bar{g}$, $S_i = T_i$ and $g_i = f_i$ for all $i \leq \#(f)$.*

We also have to deal with new cases. Whenever a new rule or production is added, a new case must be added to proofs which induct over or case split on the original production or rule. For FGJ , this means that a new case must be added to Lemma 5.2 for **GS-Var**. When writing an inductive proof, a feature provides cases for each of the rules or productions it introduces. To build the proof for the target language, a new skeleton proof by induction is started. Each of the cases is discharged by the proof given in the introducing feature.

5.2 Engineering Extensible Proofs in Coq

Each Coq feature module contains proofs for the extensible lemmas it provides. To get a handle on the behavior of opaque parameters, Coq feature modules make explicit assumptions about their behavior. Just as definitions were parameterized on variation points, proofs are now parameterized on a set of lemmas that define legal extensions. These assumptions enable separate certification of feature modules. Coq certifies that a proof is correct for all instantiations or case introductions that satisfy its assumptions, enabling proof reuse for all compatible features.

As a concrete example, consider the Coq proof of Lemma 5.3 given in Figure 16. The cFJ feature provides the statement of the lemma, which is over the abstract subtype relation. Both the Generic and cFJ features give proofs for their definitions of the subtype relation. Notably, the Generic feature assumes that if a type variable is found in a Context Gamma, it will have the same value in app_context Gamma Delta for any Context Delta. Any compatible extension of Context and app_Context can thus reuse this proof.

Lemma 5.3 (Subtype Weakening). *For all contexts Γ and Δ , if $\Gamma \vdash S <: T$, $\Gamma; \Delta \vdash S <: T$.*

Figure 15: Weakening lemma for subtyping.

To build the final proof, the target language inducts over subtype, as shown in the final box of Figure 16. For each constructor, the lemma dispatches to the proofs from the corresponding feature module. To reuse those proofs, each of their assumptions has to be fulfilled by a theorem (e.g. TLookup_app' satisfies TLookup_app). The inductive hypothesis is provided to cFJ_Weaken_subtype_app for use on its subterms. As long as every assumption is satisfied for each proof case, Coq will certify the composite proof. There is one important caveat: proofs which use the inductive hypothesis can only do so on subterms or subjudgements. By using custom induction schemes to build proofs, features can ensure that this check will always succeed. The cFJ_subtype_ind induction scheme used to combine cFJ's cases in the first box of Figure 16 is an example.

5.3 Feature Composition in Coq

Each feature module is implemented as a Coq file which contains the inductive definitions, variation points, and proofs provided by that feature. These modules are certified independently by Coq. Once the feature modules have been verified, a target language is built as a new Coq file. This file imports the files for each of the features included in the language, e.g. “Require Import cFJ.” in Figure 12. First, each target language definition is built as a new inductive type using appropriately instantiated definitions from the included feature modules, as shown in Figures 12 and 13. Proofs for the target language are then built using the proofs

```

Variables (app_context : Context → Context → Context)
(FJ_subtype Wrap : forall gamma S T,
  FJ_subtype gamma S T → subtype gamma S T).
Definition Weaken_Subtype_app_P
  delta S T (sub_S_T : subtype delta S T) :=
  forall gamma, subtype (app_context delta gamma) S T.

Lemma cFJ_Weaken_Subtype_app_H1 :
  forall (ty : Ty) (gamma : Context),
  Weaken_Subtype_app_P _ _ (sub_refl ty gamma).
Lemma cFJ_Weaken_Subtype_app_H2 : forall c d e gamma
  (sub_c : subtype gamma c d) (sub_d : subtype gamma d e),
  Weaken_Subtype_app_P _ _ sub_c →
  Weaken_Subtype_app_P _ _ sub_d →
  Weaken_Subtype_app_P _ _ (sub_trans _ _ _ sub_c sub_d).
Lemma cFJ_Weaken_Subtype_app_H3 :
  forall ce c d fs k' ms te te' delta CT_c
  bld_te, Weaken_Subtype_app_P _ _ _
  (sub_dir ce c d fs
    k' ms te te' delta CT_c bld_te).
Definition cFJ_Weaken_Subtype_app :=
  cFJ_subtype_ind _ cFJ_Weaken_Subtype_app_H1
  cFJ_Weaken_Subtype_app_H2 cFJ_Weaken_Subtype_app_H3.

Variables (app_context:Context → Context → Context)
(TLookup_app : forall gamma delta X ty,
  TLookup gamma X ty →
  TLookup (app_context gamma delta) X ty).
(GJ_subtype Wrap : forall gamma S T,
  GJ_subtype gamma S T → subtype gamma S T).
Definition Weaken_Subtype_app_P :=
  cFJ_Pinitions.Weaken_Subtype_app_P _ _ subtype app_context.

Lemma GJ_Weaken_Subtype_app : forall gamma
  S T (sub_S_T : GJ_subtype gamma S T),
  Weaken_Subtype_app_P _ _ sub_S_T.
cbv beta delta; intros; apply GJ_subtype Wrap.
inversion sub_S_T; subst.
econstructor; eapply TLookup_app; eauto.
Qed.

Fixpoint Weaken_Subtype_app gamma S T
  (sub_S_T : subtype gamma S T) :
  Weaken_Subtype_app_P _ _ sub_S_T :=
  match sub_S_T return Weaken_Subtype_app_P _ _ sub_S_T with
  | cFJ_subtype Wrap gamma S' T' sub_S_T' =>
    cFJ_Weaken_Subtype_app _ _ _ cFJ_Ty.Wrap _ _ CT
    _ subtype GJ_Phi_sb cFJ_subtype Wrap app_context _ _
    sub_S_T' Weaken_Subtype_app
  | GJ_subtype Wrap gamma S' T' sub_S_T' =>
    GJ_Weaken_Subtype_app _ _ Gty _ TLookup subtype
    GJ_subtype Wrap app_context TLookup_app' _ _ sub_S_T'
  end.

```

Figure 16: Coq proofs of Lemma 5.3 for the cFJ and Generic features and the composite proof.

from the constituent feature modules per the above discussion. Proof composition requires a straightforward check by Coq that the assumptions of each feature module are satisfied, i.e. that a feature's interface is met by the target language. Currently each piece of the final language is composed by hand in this straightforward manner; future work includes automating feature composition directly.

6. Implementation of the FJ Product Line

We have implemented the six feature modules of Section 3.4 in the Coq proof assistant. Each contains pieces of syntax, semantics, type system, and metatheoretic proofs needed by

that feature or interaction. Using them, we can build the seven variants on Featherweight Java listed in Section 3.2³.

Module	Lines of Code in Coq
cFJ	2612 LOC
Cast	463 LOC
Interface	499 LOC
Generic	6740 LOC
Generic#Interfaces	1632 LOC
Generic#Cast	296 LOC

Figure 17: Feature Module Sizes

While we achieve feature composition by manually instantiating these modules, the process is straightforward and should be mechanizable. Except for some trivial lemmas, the proofs for a final language are assembled from proof pieces from its constituent features by supplying them with lemmas which satisfy their assumptions. Importantly, once the proofs in each of the feature modules have been certified by Coq, they do not need to be rechecked for the target language. A proof is guaranteed to be correct for any language which satisfies the interface formed by the set of assumptions for that lemma. This has a practical effect as well: certifying larger feature modules takes a non-trivial amount of time. Figure 18 lists the certification times for feature modules and all the possible language variants built from their composition. By checking the proofs of each feature in isolation, Coq is able to certify the entire product line in roughly the same amount of time as the cFJ feature module. Rechecking the work of each feature for each individual product would quickly become expensive. Independent certification is particularly useful when modifying a single feature. Recertifying the product line is a matter of rechecking the proofs of the modified features and then performing a quick check of the products, without having to recheck the independent features.




Figure 18: Certification Times for Feature Modules and All Language Variants.

³ The source for these feature modules and language variants can be found at <http://www.cs.utexas.edu/users/bendy/MMMDevelopment.php>.

7. Discussion and Future Work

Relying on parameterization for feature composition allows feature modules to be built and independently certified by Coq “out of the box” with the same level of assurance. With this approach, a familiar set of issues is encountered when a new feature is added to a product line. Ideally, a new feature would be able to simply update the existing definitions and proofs, allowing language designers to leverage all the hard work expended on formalizing the original language. Some foresight, called *domain analysis* [22], allows language designers to predict VPs in advance, thus enabling a smooth and typically painless addition of new features. What our work shows is a path for the structured evolution of languages. But of course, when unanticipated features are added using this style of composition, additional engineering may be required.

Existing definitions can be extended and reused *as long as* they already have the appropriate VPs and their inductive definitions are left open. For example, once class definitions have a variation point inserted for interfaces, the same VP can also be extended with type parameters for generics. Similarly, once the definition of subtyping has been left open, both interfaces and generics can add new rules for the target language.

Proof reuse is almost as straightforward: as long as an extension is compatible with the set of assumptions in a proof’s interface, the proof can be reused directly in the final language. A new feature is responsible for providing the proofs that its extension satisfies the assumptions of the original base language.

Refactoring is necessary when a new feature requires VPs that are not in existing features. A feature which makes widespread changes throughout the base language (i.e. Generic), will probably make changes in places that the original feature designer did not anticipate. In this situation, as mentioned in Section 2.1, existing features have to be refactored to allow the new kind of extension by inserting variation points or breaking open the recursion on inductive definitions. Any proofs over the original definition may have to be updated to handle the new extensions, possibly adding new assumptions to its interface.

Feature modules tend to be inoculated from changes in another, unless they reference the definitions of another feature. This only occurs when two features must appear together: modules which resolve feature interactions, for example, only appear when their base features are present. Thus, it is possible to develop an enhanced language incrementally: starting with the base and then iteratively refactoring other features, potentially creating new modules to handle interactions. Once a new feature has been fully integrated into the feature set, all of the previous languages in the product line should still be derivable. If two features F and G commute (i.e. $F \cdot G = G \cdot F$) their integration comes for free as their interaction module is empty (i.e. $F\#G = 1$).

A new feature can also invalidate the assumptions of an existing feature's proofs. In this case, assumptions might have to be weakened and the proof refactored to permit the new extension. Alternatively, if an extension breaks the assumption of an existing proof, the offending feature can simply build a new proof of that lemma. This proof can then be utilized in any other proofs which used that lemma as an assumption, replacing the original lemma and allowing the other proofs to be reused. In this manner, each proof is insulated from features which break the assumptions of other lemmas. Again, all of this is just another variation of the kinds of problems that are encountered when one generalizes and refactors typical object-oriented code bases.

Future work includes creating a new module-level composition operator that eases the burden of integrating new features. Ideally, this operator will allow subsequent features to extend a feature's definitions with new cases or variations without modifying the feature and to provide patches to allow existing proofs to work with the extended definitions. As alluded to earlier, by operating at the module level, this operator would automate the tedious piece-by-piece composition currently employed to build target languages.

8. Related Work

The Tinkertype project [17] is a framework for modularly specifying formal languages. Features consist of a set of variants of inference rules with a feature model determining which rule is used in the final language. An implementation of these ideas was used to format the language variants used in Pierce's Types and Programming Languages [24]. This corresponds to our notion of case introduction. Importantly, our approach uses variations points to allow variations on a single definition. This allows us to construct of a single generic proof which can be specialized for each variant, as opposed to maintaining a separate proof for each variation. Levin et al. consider using their tool to compose handwritten proofs, but these proofs must be rechecked after composition. In contrast, we have crafted a systematic approach to proof extension that permits the creation of machine-checkable proofs. After a module's proofs are certified, they can be reused without needing to be rechecked. As long as the module's assumptions hold, the proofs are guaranteed to hold for the final language.

Stärk et. al [27] develop a complete Java 1.0 compiler through incremental refinement of a set of Abstract State Machines. Starting with ExpI, a core language of imperative Java expressions which contains a grammar, interpreter, and compiler, the authors add features which incrementally update the language until an interpreter and compiler are derived for the full Java 1.0 specification. The authors then write a monolithic proof of correctness for the full language. Later work casts this approach in the calculus of features [2], noting that the proof could have been developed incrementally. While we present the incremental development of the

formal specification of a language here, many of the ideas are the same. An important difference is that our work focuses on structuring languages and proofs for mechanized proof assistants, while the development proposed by [2] is completely by hand.

Thüm et. al [29] consider proof composition in the verification of a Java-based software product line. Each product is annotated with invariants from which the Krakatoa/Why tool generates proof obligations to be verified in Coq. To avoid maintaining these proofs for each product, the authors maintain proof pieces in each feature and compose the pieces for an individual product. Their notion of composition is strictly syntactic: proof scripts are copied together to build the final proofs and have to be rechecked for each product. Importantly, features only add new premises and conjunctions to the conclusions of the obligations generated by Krakatoa/Why, allowing syntactic composition to work well for this application. As features begin to apply more subtle changes to definitions and proofs, it is not clear how to effectively syntactically glue together Coq's proof scripts. Using the abstraction mechanisms provided by Coq to implement features, as we have, enables a more semantic notion of composition.

The modular development of reduction rules are the focus of Mosses' Modular Structural Operational Semantics [20]. In this paradigm, rules are written with an abstract label which effectively serves as a repository for all effects, allowing rules to be written once and reused with different instantiations depending on the effects supported by the final language. Effect-free transitions pass around the labels of their subexpressions:

$$\frac{d \xrightarrow{X} d'}{\text{let } d \text{ in } e \xrightarrow{X} \text{let } d' \text{ in } e} \quad (\text{R-LETB})$$

Those rules which rely on an effectual transition specify that the final labeling supports effects:

$$\frac{e \xrightarrow{\{p=p_1[p_0] \dots\}} e'}{\text{let } p_0 \text{ in } e \xrightarrow{\{p=p_1 \dots\}} \text{let } p_0 \text{ in } e} \quad (\text{R-LETE})$$

These abstract labels correspond to the abstract contexts used by the cFJ subtyping rules to accommodate the updates of the Generic feature. In the same way that R-LETE depends on the existence of a store in the final language, S-VAR requires the final context to support a type lookup operation. Similarly, both R-LETB and S-TRANS pass along the abstract labels / contexts from their subjudgements.

Both Boite [9] and Mulhern [21] consider how to extend existing inductive definitions and reuse related proofs in the Coq proof assistant. Both only consider adding new cases and rely on the critical observation that proofs over the extended language can be patched by adding pieces for the new cases. The latter promotes the idea of 'proof weaving' for merging inductive definitions of two languages which

merges proofs from each by case splitting and reusing existing proof terms. An unimplemented tool is proposed to automatically weave definitions together. The former extends Coq with a new `Extend` keyword that redefines an existing inductive type with new cases and a `Reuse` keyword that creates a partial proof for an extended datatype with proof holes for the new cases which the user must interactively fill in. These two keywords explicitly extend a concrete definition and thus modules which use them cannot be checked by Coq independently of those definitions. This presents a problem when building a language product line: adding a new feature to a base language can easily break the proofs of subsequent features which are written using the original, fixed language. Interactions can also require updates to existing features in order to layer them onto the feature enhanced base language, leading to the development of parallel features that are applied depending on whether the new feature is included. These keyword extensions were written for a previous version of Coq and are not available for the current version of the theorem prover. As a result of our formulation, it is possible to check the proofs in each feature module independently, with no need to recheck proof terms when composing features.

Chlipala [10] proposes a using adaptive tactics written in Coq's tactic definition language LTac [11] to achieve proof reuse for a certified compiler. The generality of the approach is tested by enhancing the original language with let expressions, constants, equality testing, and recursive functions, each of which required relatively minor updates to existing proof scripts. In contrast to our approach, each refinement was incorporated into a new monolithic language, with the new variant having a distinct set of proofs to maintain. Our feature modules avoid this problem, as each target language derives its proofs from a uniform base, with no need to recheck the proofs in existing feature modules when composing them with a new feature. Adaptive proofs could also be used within our feature modules to make existing proofs robust in to the addition of new syntax and semantic variation points.

9. Conclusion

Mechanically verifying artifacts using theorem provers can be hard work. The difficulty is compounded when verifying all the members of a product line. Features, transformations which add a new piece of functionality, are a natural way of decomposing a product line. Decomposing proofs along feature boundaries enables the reuse of proofs from a common base for each target product. These ideas have a natural expression in the evolution of formal specification of programming languages, using the syntax, semantics, and metatheoretic proofs of a language as the core representations. In this paper, we have shown how introductions and variation points can be used to structure product lines of formal language specifications.

As a proof of concept, we used this approach to implement features modules that enhance a variant of Featherweight Java in the Coq proof assistant. Our implementation uses the standard facilities of Coq to build the composed languages. Coq is able to mechanically check the proofs of progress and preservation for the composed languages, which reuse pieces of proofs defined in the composed features. Each extension allows for the structured evolution of a language from a simple core to a fully-featured language. Harnessing these ideas in a mechanized framework transforms the mechanized formalization of a language from a rigorous check of correctness into an important vehicle for reuse of definitions and proofs across a family of related languages.

Acknowledgments. This work was supported by NSF's Science of Design Project CCF 0724979. Also we appreciate comments from Thomas Thüm and the referees on earlier drafts of this paper.

References

- [1] Paul Bassett. Frame-based software engineering. *IEEE Software*, 4(4), 1987.
- [2] D. Batory and E. Börger. Modularizing theorems for software product lines: The jbook case study. *Journal of Universal Computer Science*, 14(12):2059–2082, 2008.
- [3] D. Batory, Rich Cardone, and Y. Smaragdakis. Object-oriented frameworks and product-lines. In *SPLC*, 2000.
- [4] D. Batory, J. Kim, and P. Höfner. Feature interactions, products, and composition. In *GPCE*, 2011.
- [5] D. Batory, J.N. Sarvela, and A. Rauschmayer. Scaling Step-Wise Refinement. *IEEE TSE*, 30, June 2004.
- [6] Don Batory. Feature models, grammars, and propositional formulas. *Software Product Lines*, pages 7–20, 2005.
- [7] Don Batory, Rich Cardone, and Yannis Smaragdakis. Object-oriented framework and product lines. In *SPLC*, pages 227–247, 2000.
- [8] Yves Bertot and Pierre Castéran. *Interactive Theorem Proving and Program Development*. Springer-Verlag, Berlin, 2004.
- [9] Olivier Boite. Proof reuse with extended inductive types. In *Theorem Proving in Higher Order Logics*, pages 50–65, 2004.
- [10] Adam Chlipala. A verified compiler for an impure functional language. In *POPL 2010*, January 2010.
- [11] David Delahaye. A tactic language for the system coq. In *Proceedings of Logic for Programming and Automated Reasoning (LPAR), Reunion Island, volume 1955 of LNCS*, pages 85–95. Springer, 2000.
- [12] Feature oriented programming. http://en.wikipedia.org/wiki/Feature_Oriented_Programming, 2008.
- [13] Georges Gonthier. In Deepak Kapur, editor, *Computer Mathematics*, chapter The Four Colour Theorem: Engineering of a Formal Proof, pages 333–333. Springer-Verlag, Berlin, Heidelberg, 2008.

- [14] Atsushi Igarashi, Benjamin C. Pierce, and Philip Wadler. Featherweight java: a minimal core calculus for java and gj. *ACM Trans. Program. Lang. Syst.*, 23(3):396–450, 2001.
- [15] K.C. Kang. Private Correspondence, 2005.
- [16] Xavier Leroy. Formal verification of a realistic compiler. *Commun. ACM*, 52:107–115, July 2009.
- [17] Michael Y. Levin and Benjamin C. Pierce. Tinkertype: A language for playing with formal systems. *Journal of Functional Programming*, 13(2), March 2003. A preliminary version appeared as an invited paper at the *Logical Frameworks and Metalanguages Workshop (LFM)*, June 2000.
- [18] M. D. McIlroy. Mass-produced software components. *Proc. NATO Conf. on Software Engineering, Garmisch, Germany*, 1968.
- [19] R. Monson-Haefel. *Enterprise Java Beans*. O’Reilly, 3rd edition, 2001.
- [20] Peter D. Mosses. Modular structural operational semantics. *J. Log. Algebr. Program.*, 60-61:195–228, 2004.
- [21] Anne Mulhern. Proof weaving. In *Proceedings of the First Informal ACM SIGPLAN Workshop on Mechanizing Metatheory*, September 2006.
- [22] J. Neighbors. The draco approach to constructing software from reusable components. *IEEE TSE*, September 1984.
- [23] D.L. Parnas. On the design and development of program families. *IEEE TSE*, SE-2(1):1 – 9, March 1976.
- [24] Benjamin C. Pierce. *Types and Programming Languages*. MIT Press, 2002.
- [25] Y. Smaragdakis and D. Batory. Mixin Layers: An Object-Oriented Implementation Technique for Refinements and Collaboration-Based Designs. *ACM TOSEM*, December 2001.
- [26] Yannis Smaragdakis and Don Batory. Implementing reusable object-oriented components. In *In the 5th Int. Conf. on Software Reuse (ICSR 98)*, pages 36–45. Society Press, 1998.
- [27] Robert Stärk, Joachim Schmid, and Egon Börger. Java and the java virtual machine - definition, verification, validation, 2001.
- [28] Rok Strnisa, Peter Sewell, and Matthew J. Parkinson. The Java module system: core design and semantic definition. In *OOPSLA*, pages 499–514, 2007.
- [29] T. Thüm, I. Schaefer, M. Kuhlemann, and S. Apel. Proof composition for deductive verification of software product lines. In *Software Testing, Verification and Validation Workshops (ICSTW) 2011*, pages 270 –277, march 2011.
- [30] Michael VanHilst and David Notkin. Decoupling change from design. *SIGSOFT Softw. Eng. Notes*, 21:58–69, October 1996.