

LOWER BOUNDS FOR DEPTH-4 FORMULAS COMPUTING ITERATED MATRIX MULTIPLICATION*

HERVÉ FOURNIER[†], NUTAN LIMAYE[‡], GUILLAUME MALOD[†], AND SRIKANTH
SRINIVASAN[§]

Abstract. We study the arithmetic complexity of iterated matrix multiplication. We show that any multilinear homogeneous depth-4 arithmetic formula computing the product of d generic matrices of size $n \times n$, $\text{IMM}_{n,d}$, has size $n^{\Omega(\sqrt{d})}$ as long as $d = n^{O(1)}$. This improves the result of Nisan and Wigderson [*Comput. Complexity*, 6 (1997), pp. 217–234] for depth-4 set-multilinear formulas. We also study $\Sigma\Pi^{[O(d/t)]}\Sigma\Pi^{[t]}$ formulas, which are depth-4 formulas with the stated bounds on the fan-ins of the Π gates. A recent depth reduction result of Tavenas [*Lecture Notes in Comput. Sci.* 8087, 2013, pp. 813–824] shows that any n -variate degree $d = n^{O(1)}$ polynomial computable by a circuit of size $\text{poly}(n)$ can also be computed by a depth-4 $\Sigma\Pi^{[O(d/t)]}\Sigma\Pi^{[t]}$ formula of top fan-in $n^{O(d/t)}$. We show that any such formula computing $\text{IMM}_{n,d}$ has top fan-in $n^{\Omega(d/t)}$, proving the optimality of Tavenas’ result. This also strengthens a result of Kayal, Saha, and Saptharishi [*Proceedings of STOC*, 2014, pp. 146–153], which gives a similar lower bound for an explicit polynomial in VNP.

Key words. arithmetic circuits, depth four, iterated matrix product, shifted partial derivative measure, lower bound

AMS subject classifications. 68Q17, 68Q25, 68W30

DOI. 10.1137/140990280

1. Introduction. Arithmetic circuits are a convenient way to model computation when considering objects of an algebraic nature such as the determinant. Thanks to the work of Valiant [Val79, Val82], they are also the basis of a clean theoretical framework to study the complexity of such objects.

In particular, Valiant defined two classes: the class **VP** of tractable polynomials and the larger class **VNP**, which contains polynomials thought to be intractable. He then showed the completeness of the permanent polynomial for the class VNP. This contrasts with the determinant polynomial, whose expression is very close to that of the permanent, but which is efficiently computable. Indeed, a slight restriction of tractable computations [Tod92, MP08] yields a class, VP_s , for which the determinant is complete. As a result, the major open question of the equality of the classes VP_s and VNP can be stated as the question of whether a permanent can always be expressed as a “not too big” determinant, without mention of a computation model.

Many other questions remain open in arithmetic circuit complexity. One of them is whether the determinant, or other polynomials from the associated class VP_s , can be computed efficiently by weaker models. Among these models are formulas,

*Received by the editors October 7, 2014; accepted for publication (in revised form) May 20, 2015; published electronically September 8, 2015. This research was conducted as a part of Indo-French Project 4702-1(A) funded by IFCPAR/CEFIPRA.

<http://www.siam.org/journals/sicomp/44-5/99028.html>

[†]Université Paris Diderot, Sorbonne Paris Cité, Institut de Mathématiques de Jussieu, UMR 7586 CNRS, F-75205 Paris, France (fournier@math.univ-paris-diderot.fr, malod@math.univ-paris-diderot.fr).

[‡]Department of Computer Science and Engineering, Indian Institute of Technology, Bombay, Mumbai, India (nutan@cse.iitb.ac.in).

[§]Department of Mathematics, Indian Institute of Technology, Bombay, Mumbai, India (srikanth@math.iitb.ac.in).

which define a class VP_E , where partial results cannot be reused, and constant-depth circuits.

In this paper we will focus on **iterated matrix multiplication**, another fundamental computation which is **complete for the class VP_s** , and whether it can be computed by depth-4 formulas, with alternating sum and product gates (so-called $\Sigma\Pi\Sigma\Pi$ formulas).

Motivation and results. Interest in depth-4 formulas for arithmetic computation was sparked by the result of Agrawal and Vinay [AV08], showing that, for certain lower bound questions, it was enough to consider this depth-4 case. This was later pursued by Koiran [Koi12] and Tavenas [Tav13], the latter showing that a polynomial of degree d over N variables computed by a circuit of size s can also be computed by a formula of depth 4 and size $\exp(O(\sqrt{d} \log(ds) \log N))$. In particular, every polynomial p of degree $d = \text{poly}(N)$ that has a circuit of size $\text{poly}(N)$ has a depth-4 $\Sigma\Pi\Sigma\Pi$ formula C of size $\exp(O(\sqrt{d} \log N))$. The formula C , additionally, has the property that its Π -gates have fan-in $O(\sqrt{d})$; such formulas are called $\Sigma\Pi^{[O(\sqrt{d})]}\Sigma\Pi^{[O(\sqrt{d})]}$ formulas. In the special case that p is homogeneous, then C is also homogeneous.

There has been recent progress toward proving strong lower bounds for $\Sigma\Pi^{O(\sqrt{d})}\Sigma\Pi^{O(\sqrt{d})}$ formulas as well: in a breakthrough result, Gupta et al. [GKKS13] give $\exp(\Omega(\sqrt{n}))$ lower bounds for $\Sigma\Pi^{[O(\sqrt{n})]}\Sigma\Pi^{[O(\sqrt{n})]}$ formulas computing the $n \times n$ permanent and determinant polynomials. Note that this gives a lower bound of $\exp(\Omega(\sqrt{d}))$, for a polynomial in VP of degree d , which is off by a factor of $\log N$ (here $N = n^2$) in the exponent as compared to the upper bound given by Tavenas. More recently, Kayal, Saha, and Saptharishi [KSS14] gave a polynomial of degree d in N variables (for $d = \sqrt{N}$) in VNP such that any $\Sigma\Pi^{[O(\sqrt{d})]}\Sigma\Pi^{[O(\sqrt{d})]}$ formula computing the polynomial has size $\exp(\Omega(\sqrt{d} \log N))$. Thus, improving either the result of Tavenas or the lower bound techniques of [GKKS13, KSS14] a little further could yield the desired separation between VP and VNP .

Here, we look at the first approach and consider the question of whether the result of Tavenas [Tav13] can be strengthened. Formally, we ask,

Is it possible to show that any polynomial (respectively, homogeneous polynomial) of degree d over N variables that has a $\text{poly}(N)$ -sized circuit has a $\Sigma\Pi^{[O(\sqrt{d})]}\Sigma\Pi^{[O(\sqrt{d})]}$ (respectively, homogeneous $\Sigma\Pi\Sigma\Pi$) formula of size $\exp(o(\sqrt{d} \log N))$?

Our results indicate that the answer to this question is negative: Theorem 16 implies that for all d there is an explicit polynomial $f \in \text{VP}$ of degree d on N variables such that any $\Sigma\Pi^{[O(\sqrt{d})]}\Sigma\Pi^{[O(\sqrt{d})]}$ formula computing it has size $\exp(\Omega(\sqrt{d} \log N))$. Thus, we strengthen the result of [KSS14] by obtaining a similar lower bound (for all d) for a polynomial in VP .

Moreover, the polynomial f is the iterated matrix multiplication polynomial, which computes a single entry in the product of d many $n \times n$ matrices whose entries are all distinct variables. We will denote this polynomial by $\text{IMM}_{n,d}$. Its complexity is by itself of great interest. It occupies a central position both in algebraic complexity theory—being complete for VP_s as mentioned above—and in complexity theory in general, since it is closely related to the Boolean and counting versions of the canonical NL-complete problem of deciding reachability in a directed graph. In particular, showing that $\text{IMM}_{n,d}$ does not have polynomial-sized formulas is equivalent to showing a separation between VP_E and VP_s .

It is easy to see that for any $r \in \mathbb{N}$, the polynomial $\text{IMM}_{n,d}$ has a formula of size $n^{O(rd^{1/r})}$ and product depth $r \leq \log d$ (i.e., at most r Π -gates on any root to leaf path).

This formula, constructed using a simple divide-and-conquer technique that requires r levels of recursion, is furthermore a *set-multilinear* formula (see section 2). In particular, for $r = \log d$, this technique yields a set-multilinear formula of size $n^{O(\log d)}$ for $\text{IMM}_{n,d}$, which is the best known formula upper bound for this polynomial.

In a seminal work, Nisan and Wigderson [NW97] showed lower bounds on the size of product depth- r set multilinear formulas computing $\text{IMM}_{n,d}$. For the case $r = 1$, [NW97] prove an optimal lower bound of n^{d-1} . For $r \geq 2$, however, they prove a lower bound of $\exp(\Omega(d^{1/r}))$. Note that the dimension n of the matrices does not appear in the lower bound: indeed, we get the same lower bound for any $n \geq 2$. We rectify this situation for $r = 2$ with the following theorem.

THEOREM 15. *For any large enough $n, d \in \mathbb{N}$, any set-multilinear $\Sigma\Pi\Sigma\Pi$ formula computing $\text{IMM}_{n,d}$ has size $n^{\Omega(\sqrt{d})}$.*

In fact, our lower bound holds in the more general setting of homogeneous multilinear $\Sigma\Pi\Sigma\Pi$ formulas.

THEOREM 29. *For any large enough $n, d \in \mathbb{N}$ such that $d = n^{O(1)}$, any homogeneous multilinear $\Sigma\Pi\Sigma\Pi$ formula computing $\text{IMM}_{n,d}$ has size $n^{\Omega(\sqrt{d})}$.*

As a final consequence of our technical theorems, we also obtain an optimal lower bound for *regular formulas* (see section 6) for the same polynomial f , answering a question raised in [KSS14].

THEOREM 18. *For large enough $n, d \in \mathbb{N}$, any regular formula for $\text{IMM}_{n,d}$ has size at least $n^{\Omega(\log d)}$.*

Related work. As mentioned above, the iterated matrix multiplication polynomial has been considered before in a work of Nisan and Wigderson [NW97], which also introduced the important technique of using *partial derivatives* to prove lower bounds in arithmetic complexity. We use a recent strengthening of this technique due to Kayal [Kay12] and Gupta et al. [GKKS13], which uses *shifted partial derivatives*. We briefly survey some results that use this technique but refer the reader to [GKKS13] for a more thorough account.

Kayal [Kay12] used the shifted partial derivatives technique to show a lower bound for expressing the monomial $x_1 x_2 \cdots x_n$ as a sum of powers of bounded degree polynomials in x_1, \dots, x_n . Gupta et al. [GKKS13] showed lower bounds for $\Sigma\Pi\Sigma\Pi$ formulas (with fan-in bounds on the Π -gates) computing the permanent and determinant polynomials. More recently, the shifted partial derivative method has been used by Kumar and Saraf [KS13] to prove lower bounds for homogeneous $\Sigma\Pi\Sigma\Pi$ formulas (see section 2) with bounded fan-in at the top Σ gate computing the permanent and by Kayal, Saha, and Saptharishi [KSS14] to prove stronger lower bounds for bounded Π -gate fan-in $\Sigma\Pi\Sigma\Pi$ formulas computing a certain explicit polynomial in VNP.

It is interesting to note that the result of [GKKS13] itself implies a (superpolynomially weaker) lower bound for $\Sigma\Pi^{[O(\sqrt{d})]}\Sigma\Pi^{[O(\sqrt{d})]}$ formulas computing the iterated matrix multiplication polynomial. This is because of the well-known fact (see, for instance, [MV97]) that an $m \times m$ determinant is a projection of $\text{IMM}_{n,d}$, where $n = O(m^2)$ and $d = m$. Thus, for this setting of parameters, the lower bound of [GKKS13] for the determinant gives a lower bound of $\exp(\Omega(\sqrt{d}))$ for $\text{IMM}_{n,d}$.

There has also been a considerable amount of research into lower bounds for set-multilinear and, more generally, multilinear formulas. Nisan and Wigderson [NW97] proved lower bounds on the size of small-depth set-multilinear formulas for the iterated matrix multiplication polynomial. Building on their techniques, the breakthrough work of Raz [Raz09] proved superpolynomial lower bounds for multilinear formulas computing the determinant and permanent polynomials. Follow-up work of

Raz [Raz06] (see also Raz and Yehudayoff [RY08]) showed a superpolynomial separation between VP_E and VP in the multilinear setting. This was recently strengthened by Dvir et al. [DMPY12] to a superpolynomial separation between VP_E and VP_s in the multilinear setting.

A result that is closely related to ours is the work of Raz and Yehudayoff [RY09], who also prove strong exponential lower bounds for constant-depth multilinear formulas. More precisely, they give an explicit multilinear polynomial of degree N over N variables that has no multilinear $\Sigma\Pi\Sigma\Pi$ formulas of size less than $\exp(\Omega(\sqrt{N \log N}))$. Their results are somewhat incomparable to ours, for the following reasons:

- Our lower bound is stronger in that it matches Tavenas' upper bound [Tav13] for $\Sigma\Pi\Sigma\Pi$ formulas for any degree- d polynomial with $\text{poly}(N)$ -sized circuits. The above lower bound is slightly weaker.
- The results of Raz and Yehudayoff apply not just to $\Sigma\Pi\Sigma\Pi$ formulas but to all *small-depth* (up to $o(\log N / \log \log N)$) multilinear formulas, without homogeneity restrictions. (The bounds get weaker with larger depth.)
- As far as we are aware, their techniques—or indeed, any of the general techniques used to prove multilinear formula lower bounds—are not applicable to the iterated matrix multiplication polynomial.

Subsequent work. In a recent work, building upon [GKKS13, KSS14, CM14], Kumar and Saraf [KS14a] prove the tightness of Tavenas' depth reduction result even for homogeneous formulas, thus strengthening our Theorem 16. Specifically, they give a polynomial of degree d on N variables computable by a $\text{poly}(N)$ size homogeneous $\Sigma\Pi\Sigma\Pi$ formula with top fan-in $O(\log d)$ such that any homogeneous $\Sigma\Pi\Sigma\Pi^{[t]}$ formula computing it has size $\exp(\Omega(\frac{d}{t} \log N))$. In the same work, they also give a family of polynomials $\{f_n\}_n$ in VNP such that any $\Sigma\Pi\Sigma\Pi$ homogeneous formula with top fan-in $o(\log n)$ computing f_n requires superpolynomial size (specifically, any $\Sigma\Pi\Sigma\Pi$ homogeneous formula with top fan-in bounded r needs size $\exp(\Omega(n^{1/r} \log n))$).

The work of [KLSS14b] proves that any homogeneous $\Sigma\Pi\Sigma\Pi$ formula (with *no* restriction on the top fan-in) computing $\text{IMM}_{n,d}$ requires size $n^{\Omega(\log n)}$. In a more recent work [KLSS14a] gives a polynomial in VNP such that any homogeneous $\Sigma\Pi\Sigma\Pi$ formula computing it requires size $2^{\Omega(\sqrt{d} \log N)}$. Finally, the most recent work of [KS14b] proves that any homogeneous $\Sigma\Pi\Sigma\Pi$ formula computing $\text{IMM}_{n,d}$ requires size $2^{\Omega(\sqrt{d} \log N)}$. This result subsumes the lower bound proved in the current work. We would like to note that [KLSS14b] and [KS14b] use the random restriction technique which we have used to lower bound the dimension of the shifted partial derivative space of $\text{IMM}_{n,d}$ in this paper.

2. Definitions and notation. Let X be a set of variables and let $\mathbb{F}[X]$ denote the set of polynomials over variables X and field \mathbb{F} .

Arithmetic circuits and branching programs. An *arithmetic circuit* is a finite simple directed acyclic graph. The vertices of in-degree 0 are called *input gates* and are labeled by constants from \mathbb{F} or variables from X . The vertices of in-degree at least 2 are labeled by $+$ or \times . The *output gate* is a vertex with out-degree 0. The polynomial computed by a node is defined in an obvious inductive way. The polynomial computed by the arithmetic circuit is the polynomial computed by the output gate.

The size of the circuit is the number of nodes in the graph. The depth of the circuit is the length of the longest input gate to output gate path. The in-degree (out-degree) of a node/gate is often called its *fan-in* (*fan-out*, respectively). We do not assume any bound on the fan-ins or fan-outs of the nodes unless stated otherwise. A circuit is called layered if the underlying graph is layered.

An *algebraic branching program* (ABP) over the set of variables X and field \mathbb{F} is a tuple (G, s, t) , where G is a weighted simple directed acyclic graph and s and t are special vertices in G . The *weight* of an edge in a branching program is a linear form in $\mathbb{F}[X]$. The *weight* of a path is the product of the weights of its edges. The polynomial computed by a branching program G is the sum of the weights of all the paths from s to t in G . The *size* of a branching program is the number of its vertices. The *length* of a branching program is the length of the longest s to t path. If we can partition the vertices of a branching program in levels so that there are only edges between vertices in successive levels, we say that the branching program is *layered*.

Arithmetic formulas and variants. An *arithmetic formula* is an arithmetic circuit which is a simple directed tree. The size, depth, fan-ins, fan-outs, and layers for formulas are defined similarly to that of circuits. We fix the convention that in a layered circuit/formula, the layers are numbered in increasing order with input gates getting the smallest number (0) and output gates getting the largest number.

A $\Sigma\Pi\Sigma\Pi$ formula is a layered formula in which gates at layers 1 and 3 are labeled \times and gates at layers 2 and 4 are labeled $+$. We will also use the notation $\Sigma\Pi^{[\alpha]}\Sigma\Pi^{[\beta]}$ to indicate that the fan-in of gates on the first and third layers is bounded by β and α , respectively.

Recall that a polynomial is called *homogeneous* if each monomial in it has the same degree. A formula is called homogeneous if each of its gates computes a homogeneous polynomial.

Fix a partition X_1, X_2, \dots, X_d of X . For a subset $T \subseteq [d]$ we say that a monomial over the variables in X is *T -multilinear* if it is a product of variables such that exactly one variable comes from each X_i ($i \in T$). A polynomial is called *T -multilinear* if it is a linear combination of T -multilinear monomials. We say that a polynomial is *set-multilinear* if it is T -multilinear for some $T \subseteq [d]$.

A formula is called set-multilinear if every node in the formula computes a set-multilinear polynomial. Note that a set-multilinear formula is by definition homogeneous.

We also consider *multilinear polynomials*, which are a slight generalization of set-multilinear polynomials. A monomial over a set of variables X is called multilinear if each variable in X has degree at most one in the monomial. A polynomial is called multilinear if it is a linear combination of multilinear monomials. A formula is called multilinear if each node in the formula computes a multilinear polynomial. It is called *homogeneous multilinear* if it is simultaneously homogeneous and multilinear.

For any node g in the formula, let X_g denote the set of variables in the polynomial computed by g . A formula is called *syntactic multilinear* if, for each \times node g in the formula, the sets $X_{g_1}, X_{g_2}, \dots, X_{g_k}$ are mutually disjoint, where g_1, g_2, \dots, g_k are the children of g .

It is known from [SY10] that if there is a multilinear formula F of size s computing a multilinear polynomial $p \in \mathbb{F}[X]$, then there exists a syntactic multilinear formula of size at most s computing p ; similar statements are also true for $\Sigma\Pi\Sigma\Pi$ and homogeneous $\Sigma\Pi\Sigma\Pi$ formulas. Therefore, we assume without loss of generality that the formulas computing multilinear polynomials are syntactic multilinear.

It will be convenient for us to blur the distinction between multilinear monomials over the set of variables X and subsets of X . Thus, we freely apply reasonable set-theoretic operations to multilinear monomials. For example, for multilinear monomials m_1 and m_2 , $m_1 \cup m_2$ is the *multilinear* monomial that contains exactly the variables that occur in either m_1 or m_2 ; we can similarly define $m_1 \cap m_2$ and $m_1 \setminus m_2$; $|m|$ will denote the degree of a multilinear monomial m .

The iterated matrix multiplication polynomial. Throughout, let $n, d \geq 2$ be fixed parameters. We consider polynomials defined on variable sets X_1, \dots, X_d . For $i \in [d] \setminus \{1, d\}$, let X_i be the set of variables $x_{j,k}^{(i)}$ for $j, k \in [n]$; for $i \in \{1, d\}$, let X_i be the set of variables $x_j^{(i)}$ for $j \in [n]$. Let $X = \bigcup_{i \in [d]} X_i$. We will use N to denote $|X| = (d-2)n^2 + 2n$.

The iterated matrix multiplication polynomial on X , denoted $\text{IMM}_{n,d}$, is defined to be

$$\text{IMM}_{n,d} = \sum_{j_1, \dots, j_{d-1}} x_{j_1}^{(1)} x_{j_1, j_2}^{(2)} x_{j_2, j_3}^{(3)} \cdots x_{j_{d-2}, j_{d-1}}^{(d-1)} x_{j_{d-1}}^{(d)}.$$

Note that the polynomial $\text{IMM}_{n,d}$ is the sole entry of the product of d generic matrices (of dimensions $1 \times n$, $n \times n$ ($d-2$ times), and $n \times 1$), the i th matrix having entries from the variable set X_i . Hence in the remainder of the paper we refer to “the matrix X_i .”

Another way to define this polynomial is to see it as a generic layered ABP with $d+1$ layers V_0, \dots, V_d , where $V_i = \{v_1^{(i)}, \dots, v_n^{(i)}\}$ for $0 < i < d$ and $V_i = \{v^{(i)}\}$ for $i \in \{0, d\}$. The graph contains all possible edges from V_i to V_{i+1} for $i \in \{0, \dots, d-1\}$. The edge from $v_j^{(i-1)}$ to $v_k^{(i)}$ is labeled with the variable $x_{j,k}^{(i)}$ for $0 < i < d-1$ and the edges from $v^{(0)}$ to $v_j^{(1)}$ and $v_j^{(d-1)}$ to $v^{(d)}$ are labeled $x_j^{(1)}$ and $x_j^{(d)}$, respectively. Then, $\text{IMM}_{n,d}$ is the polynomial computed by the branching program, i.e., the sum of the weights of all the paths from the vertex $v^{(0)}$ to the vertex $v^{(d)}$ (see Figure 1).




FIG. 1. ABP for $\text{IMM}_{n,d}$.

We denote by \mathcal{A} the canonical ABP defined above. Given a path ρ in the ABP \mathcal{A} , we will also denote by ρ the product of all the variables that occur along the edges in the path ρ .

The dimension of the shifted partial derivatives. As in [Kay12, GKKS13], we will use the dimension of shifted partial derivatives as our complexity measure. For $k, \ell \in \mathbb{N}$ and a multivariate polynomial $f \in \mathbb{F}[x_1, \dots, x_n]$, we define

$$\langle \partial_k f \rangle_{\leq \ell} = \text{span} \left\{ x_1^{j_1} \cdots x_n^{j_n} \cdot \frac{\partial^k f}{\partial x_1^{i_1} \cdots \partial x_n^{i_n}} \mid i_1 + \cdots + i_n = k, j_1 + \cdots + j_n \leq \ell \right\}.$$

The complexity measure we use is $\dim(\langle \partial_k f \rangle_{\leq \ell})$.

3. Preliminaries. In this section we give a few technical lemmas and definitions which will be used in the subsequent sections.

The derivatives of $\text{IMM}_{n,d}$. The derivatives of $\text{IMM}_{n,d}$ have a simple form that is easily described. Since we will be interested in lower bounding the size of the partial derivative space of this polynomial, we only choose a subset of all partial derivatives available to us. Let k denote a parameter which we will choose later. Let r denote $\lfloor \frac{d}{k+1} \rfloor - 1$. We fix k matrices among X_1, \dots, X_d that are placed evenly apart. Formally, choose k matrices X_{p_1}, \dots, X_{p_k} such that $p_q - (p_{q-1} + 1) \geq r$ for all $1 \leq q \leq k+1$, where $p_0 = 0$ and $p_{k+1} = d+1$. We then choose one variable from each of these chosen matrices, say, $x_{i_1, j_1}^{(p_1)}, \dots, x_{i_k, j_k}^{(p_k)}$, and take derivatives with respect to these variables. We denote this derivative by $\partial_{\mathcal{I}} \text{IMM}_{n,d}$, where \mathcal{I} denotes $(i_1, j_1, \dots, i_k, j_k) \in [n]^{2k}$ (see Figure 2).

Note that $\partial_{\mathcal{I}} \text{IMM}_{n,d}$ can be written as a sum of monomials m such that $m = \rho_1 \rho_2 \dots \rho_{k+1}$, where ρ_q is a path from $v_{j_{q-1}}^{(p_{q-1})}$ to $v_{i_q}^{(p_q)}$ in \mathcal{A} for all $2 \leq q \leq k$, ρ_1 is a path from vertex $v^{(0)}$ to $v_{i_1}^{(p_1)}$ in \mathcal{A} , and ρ_{k+1} is a path from $v_{j_k}^{(p_k)}$ to vertex $v^{(d)}$ in \mathcal{A} . Clearly, $\partial_{\mathcal{I}} \text{IMM}_{n,d}$ is a homogeneous polynomial of degree $d - k$.




FIG. 2. The derivatives of $\text{IMM}_{n,d}$.

Restrictions. By a *restriction* of the variable set X , we will mean a function $\sigma : X \rightarrow \{0, *\}$. Given $f \in \mathbb{F}[X]$ and a restriction σ on X , we denote by $f|_{\sigma}$ the polynomial $g \in \mathbb{F}[X]$ obtained by setting all the variables $x \in \sigma^{-1}(0)$ to 0 (the other variables remain as they are).

Given polynomials $f, g \in \mathbb{F}[X]$, we say that g is a *restriction* of f if there exists a restriction σ on X such that $g = f|_{\sigma}$.

Given a formula C over the variables X and a restriction σ , we define $C|_{\sigma}$ to be the circuit obtained by replacing all the variables $x \in \sigma^{-1}(0)$ with 0 then simplifying the formula accordingly, by suppressing any Π gate receiving a variable set to 0. Clearly, if C computes the polynomial $f \in \mathbb{F}[X]$, then $C|_{\sigma}$ computes the polynomial $f|_{\sigma}$.

We will mostly be interested in restrictions of $\text{IMM}_{n,d}$. In this setting, the following basic observation helps simplify many arguments. In section 2, we defined the $\text{IMM}_{n,d}$ polynomial to be the polynomial computed by an ABP \mathcal{A} such that for each edge of \mathcal{A} , the linear form labeling that edge was a distinct variable from X . Restrictions F of $\text{IMM}_{n,d}$ are polynomials obtained when we set certain variables of X to 0 in $\text{IMM}_{n,d}$; equivalently, we may see F as the polynomial computed by the ABP \mathcal{A}_F obtained when we delete the edges corresponding to the variables that are set to 0 by the restriction.

The shifted partial derivative space of $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas. We need an upper bound on the dimension of the shifted partial derivative space of polynomials computed by small $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas. The following is implicit in the work of Gupta et al. [GKKS13] and is stated explicitly in [KSS14].

LEMMA 1 (see [KSS14, Lemma 4]). *Let $D, t, k, \ell \in \mathbb{N}$ be arbitrary parameters. Let C be a $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula with at most s Π gates at layer 3 computing a polynomial in N variables. Then, we have*

$$\dim(\langle \partial_k C \rangle_{\leq \ell}) \leq s \cdot \binom{D}{k} \cdot \binom{N + \ell + (t-1)k}{\ell + (t-1)k}.$$

We also need the following technical facts.

Fact 2. For any integers N, ℓ, r such that $r < \ell$, we have

$$\left(\frac{N + \ell}{\ell}\right)^r \leq \frac{\binom{N + \ell}{\ell}}{\binom{N + \ell - r}{\ell - r}} \leq \left(\frac{N + \ell - r}{\ell - r}\right)^r.$$

Fact 3. For any integers $n, d \geq 2$, $N = (d-2)n^2 + 2dn$ and $t \geq 1$, there exists an integer $\ell > d$ such that $n^{1/16} \leq \left(\frac{N + \ell}{\ell}\right)^t \leq n^{1/4}$.

Proof. We choose ℓ to be the least positive integer such that $f(\ell) := \left(\frac{N + \ell}{\ell}\right)^t \leq n^{1/4}$. Note that such an ℓ exists since $f(1) = (N+1)^t > n^{1/4}$ and $\lim_{\ell \rightarrow \infty} f(\ell) = 1$. We must also have $\ell > d$ since for $\ell \leq d$, we have $f(\ell) \geq ((N/d) + 1)^t \geq (n+1)^t > n^{1/4}$.

The only thing left to show is that for this choice of ℓ , we have $\left(\frac{N + \ell}{\ell}\right)^t \geq n^{1/16}$. To prove this, we claim that it suffices to prove the following inequality for any $\ell' \geq 1$:

$$(1) \quad \sqrt{f(\ell')} \leq f(\ell' + 1).$$

To see this, note that assuming the above inequality, we have $f(\ell) \geq \sqrt{f(\ell-1)} \geq n^{1/16}$, where the last inequality follows from the fact that $f(\ell-1) \geq n^{1/4}$.

The proof of inequality (1) is elementary. We need to show that

$$\begin{aligned} \left(\frac{N + \ell'}{\ell'}\right)^{t/2} &\leq \left(\frac{N + \ell' + 1}{\ell' + 1}\right)^t \\ \Leftrightarrow \left(\frac{N + \ell'}{\ell'}\right)^{1/2} &\leq \frac{N + \ell' + 1}{\ell' + 1}. \end{aligned}$$

Squaring both sides and cross multiplying we see that (1) is equivalent to

$$\begin{aligned} \frac{(\ell' + 1)^2}{\ell'} &\leq \frac{(N + \ell' + 1)^2}{N + \ell'} \\ \Leftrightarrow \ell' + \frac{1}{\ell'} + 2 &\leq N + \ell' + \frac{1}{N + \ell'} + 2 \end{aligned}$$

which is easily verified for $N \geq 1$. \square

4. Proof overview. In this section, we briefly describe the outline of the proof of the main theorems.

Our theorems prove strong lower bounds on variants of $\Sigma\Pi\Sigma\Pi$ formulas computing $\text{IMM}_{n,d}$. Recall that we already have tight lower bounds on $\Sigma\Pi\Sigma$ set-multilinear formulas computing $\text{IMM}_{n,d}$ due to Nisan and Wigderson [NW97]. A natural first step for us, therefore, would be to prove an optimal lower bound for set-multilinear formulas which are sums of products of quadratics, i.e., set-multilinear $\Sigma\Pi\Sigma\Pi^{[2]}$, or more generally, sums of products of low-degree polynomials. To do this, we use the *shifted partial derivative* method of Gupta et al. [GKKS13], who introduced this technique to prove that any $\Sigma\Pi^{[O(n/t)]}\Sigma\Pi^{[t]}$ formula (not necessarily homogeneous) computing the permanent or the determinant polynomial (on n^2 variables) must have size $\exp(\Omega(n/t))$. Their proof was made up of two steps.

- First, they observed that the shifted partial derivative space of $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas, for suitable D and t , has *small* dimension.
- Then, they showed that the dimension of $\langle \partial_k F \rangle_{\leq \ell}$ is quite large for suitable k and ℓ , where F is any one of the determinant or permanent polynomials.

We prove a strong lower bound on dimension of the shifted partial derivative space of $\text{IMM}_{n,d}$, thereby proving a lower bound of $n^{\Omega(d/t)}$ for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas computing $\text{IMM}_{n,d}$, as long as D is small enough compared to n (the formal statement and proof can be found in section 6). In fact, we manage to prove something slightly stronger. We prove that some carefully chosen *restrictions* (see section 3) of the $\text{IMM}_{n,d}$ polynomial have shifted partial derivative spaces of large dimension. Putting this together with Lemma 1 implies strong lower bounds for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas computing even these restrictions of $\text{IMM}_{n,d}$.

In order to prove our next result, a lower bound for set-multilinear $\Sigma\Pi\Sigma\Pi$ formulas and homogeneous multilinear $\Sigma\Pi\Sigma\Pi$ formulas of possibly unbounded bottom fan-in computing $\text{IMM}_{n,d}$ (the formal statement and proof can be found in section 5), we reduce to the case of formulas with bounded bottom fan-in using the idea of *random restrictions*. This is motivated by and reminiscent of some arguments in [FSS84, Hås87, NW97]; our restrictions themselves, however, look quite different.

We force the fan-in of the bottom Π gates to less than some threshold t by using random restrictions. This is quite intuitive, since a random restriction that sets any variable to 0 with good probability should set any high degree (multilinear) monomial to 0 with probability close to 1. Importantly for us, though, we can devise such a set of restrictions with the additional property that these restrictions remain hard to compute for homogeneous $\Sigma\Pi\Sigma\Pi^{[t]}$ formulas, by the ideas used to prove the lower bound for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas.

We consider two different sets of restrictions. The first set of restrictions is simpler, but works only to reduce the fan-in of *set-multilinear* $\Sigma\Pi\Sigma\Pi$ formulas. By considering a second, slightly more involved, family of restrictions, we prove a lower bound for homogeneous multilinear $\Sigma\Pi\Sigma\Pi$ formulas as well (the formal statement and proof can be found in section 7). Note that the lower bound result for $\Sigma\Pi\Sigma\Pi$ homogeneous multilinear formulas subsumes the lower bound result for set-multilinear $\Sigma\Pi\Sigma\Pi$ formulas and, indeed, there is a good amount of overlap between the two proofs. However, for the sake of clarity of exposition, we give detailed proofs for both.

5. Lower bounds for set-multilinear formulas. We start by defining a set \mathcal{R} of restrictions of $\text{IMM}_{n,d}$, then we prove a lower bound for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas computing them, and finally we show that for any set-multilinear depth-4 formula Φ computing $\text{IMM}_{n,d}$, there is a restriction $\sigma \in \mathcal{R}$ such that $\Phi|_{\sigma}$ is in fact a multilinear $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula.

5.1. Nice restrictions of $\text{IMM}_{n,d}$. Our restrictions are related to the evenly spaced matrices chosen before: recall that we have chosen indices p_1, \dots, p_k and also set $p_0 = 0$ and $p_{k+1} = d + 1$. Also, recall that r equals $\lfloor \frac{d}{k+1} \rfloor - 1$. We will now choose new indices p'_q , where p'_q is roughly in the middle between p_{q-1} and p_q : for each $q \in [k+1]$, we choose a p'_q such that $p_{q-1} \leq p'_q \leq p_q$ and $\min\{p'_q - (p_{q-1} + 1), p_q - (p'_q + 1)\} \geq \lfloor \frac{r-1}{2} \rfloor$. We define P' to be $\{p_q \mid q \in [k]\} \cup \{p'_q \mid q \in [k+1]\}$.

We then consider the set \mathcal{R} of restrictions which

- keep only one variable in the first (row) matrix,
- for each index $p \notin P' \cup \{1, d\}$, keep only the variables in X_p of the form $x_{i, \pi_p(i)}$ for some permutation π_p of $[n]$,

- for each index $p \in P'$, leave the variables of X_p untouched,
- keep only one variable in the last (column) matrix.

Restrictions in \mathcal{R} are called *nice*.

More formally, let \mathcal{R} be the set of restrictions τ defined below, for any choice of integers j_1 and j_d in $[n]$ and any set $\{\pi_p \mid p \notin P' \cup \{1, d\}\}$ of permutations of $[n]$:

$$\tau(x) = \begin{cases} 0 & \text{if } x = x_j^{(1)} \text{ for some } j \neq j_1, \text{ or} \\ & \text{if } x = x_j^{(d)} \text{ for some } j \neq j_d, \text{ or} \\ & \text{if } x = x_{i,j}^{(p)} \text{ for } j \neq \pi_p(i) \text{ and } p \notin P' \cup \{1, d\}, \\ * & \text{otherwise.} \end{cases}$$

For example, we can choose to keep the first variable of X_1 and X_d and, in each matrix X_p for $p \notin P'$, only variables on the diagonal (which corresponds to choosing π_p to be the identity permutation for all $p \notin P'$), thus defining a restriction σ :

$$\sigma(x) = \begin{cases} 0 & \text{if } x = x_j^{(1)} \text{ for some } j \neq 1, \text{ or} \\ & \text{if } x = x_j^{(d)} \text{ for some } j \neq 1, \text{ or} \\ & \text{if } x = x_{i,j}^{(p)} \text{ for } i \neq j \text{ and } p \notin P' \cup \{1, d\}, \\ * & \text{otherwise.} \end{cases}$$

Let F be the polynomial $\text{IMM}_{n,d}|_{\sigma}$. As we saw in section 3, we can also define the polynomial F in the language of ABPs. Consider the ABP \mathcal{A} defined in section 2 above. Construct a new ABP \mathcal{A}' by removing edges from \mathcal{A} as follows:

- Remove all edges from $v_j^{(0)}$ to $v_j^{(1)}$ for $j \neq 1$.
- For $p \notin P' \cup \{1, d\}$, remove all edges between V_{p-1} to V_p except for those of the form $(v_j^{(p-1)}, v_j^{(p)})$ for $j \in [n]$.
- Remove all edges from $v_j^{(d-1)}$ to $v_j^{(d)}$ for $j \neq 1$.

The ABP \mathcal{A}' computes exactly the polynomial F .

Intuition. The reasons for the choice of the restrictions in \mathcal{R} are the following. To simplify the analysis of $\dim(\langle \partial_k \text{IMM}_{n,d} \rangle_{\leq \ell})$, it is helpful to have a basis for $\langle \partial_k \text{IMM}_{n,d} \rangle_{\leq \ell}$ that is entirely made up of *monomials* (since then the problem of analyzing the dimension reduces to a simpler monomial counting problem). If we consider a k th order derivative of $\text{IMM}_{n,d}|_{\tau}$ ($\tau \in \mathcal{R}$) by variables $x_{i,j}^{(p_q)}$ ($q \in [k]$), it can be checked that this derivative is indeed a monomial. Since we have not set any of the variables in the layers corresponding to $p \in P'$, the derivative space of $\text{IMM}_{n,d}|_{\tau}$ is seen to have a *large* set of monomials. We will be able to analyze the number of the shifts of these monomials and hence give a lower bound on $\dim(\langle \partial_k \text{IMM}_{n,d} \rangle_{\leq \ell})$. This can then be used to give a lower bound on the size of $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas computing $\text{IMM}_{n,d}|_{\tau}$.

Finally, also note that the restrictions in \mathcal{R} are quite sparse: i.e., most of the variables are set to 0 in the restrictions. This will be useful when we analyze general depth-4 set-multilinear formulas (with no restriction on bottom fan-in) computing $\text{IMM}_{n,d}$. By choosing a random restriction $\tau \in \mathcal{R}$ and setting variables to 0 accordingly, we will be able to reduce the formula to a $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula (for suitable D and t) with high probability, and then appeal to our lower bound for this model in order to conclude the set-multilinear formula lower bound.

5.2. A lower bound for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas for nice restrictions of $\text{IMM}_{n,d}$. We will work with F for most of the lower bound proof and then show that

the lower bound holds for all the polynomials obtained from $\text{IMM}_{n,d}$ by a restriction in \mathcal{R} .

5.2.1. The dimension of the space of shifted partial derivatives of F .

As with $\text{IMM}_{n,d}$ in section 3, we will consider the derivatives of F with respect to a tuple of variables $x_{i_1, j_1}^{(p_1)}, \dots, x_{i_k, j_k}^{(p_k)}$ and denote this derivative by $\partial_{\mathcal{I}} F$, where \mathcal{I} denotes $(i_1, j_1, \dots, i_k, j_k)$ as before. It can be observed from the restriction defining F that $\partial_{\mathcal{I}} F$ is now a *single monomial* of degree $d - k$, which we denote $p(\mathcal{I})$. Indeed, as illustrated in Figure 3, the choice of $(i_1, j_1, \dots, i_k, j_k)$ fixes some edges a path must go through. Consider a possible path between the chosen edges (i_1, j_1) and (i_2, j_2) . It must leave at the level j_1 and stay there (as we have projected the matrices to the identity) until it reaches $X_{p'_2}$. Conversely the path must reach level i_2 in X_{p_2} , so it must have been at this level when leaving $X_{p'_2}$, and therefore it used the edge (j_1, i_2) in $X_{p'_2}$. The unicity of the path between two consecutive chosen edges holds everywhere, so there is a unique path through the chosen edges and the partial derivative is a single monomial. This argument also explains why we had to chose “intermediate” matrices $X_{p'}$ and leave them untouched by restrictions: they are used to connect paths coming from or going to X_p layers.

In fact, we can write $p(\mathcal{I}) = \rho_1 \rho_2 \dots \rho_{k+1}$, where

$$\begin{aligned} \rho_1 &= \underbrace{\left(x_1^{(1)} \cdot \prod_{1 < p < p'_1} x_{1,1}^{(p)} \right)}_{g_1^{\mathcal{I}}} \cdot \underbrace{x_{1,i_1}^{(p'_1)} \cdot \left(\prod_{p'_1 < p < p_1} x_{i_1,i_1}^{(p)} \right)}_{h_1^{\mathcal{I}}}, \\ \rho_q &= \underbrace{\left(\prod_{p_{q-1} < p < p'_q} x_{j_{q-1}, j_{q-1}}^{(p)} \right)}_{g_q^{\mathcal{I}}} \cdot \underbrace{x_{j_{q-1}, i_q}^{(p'_q)} \cdot \left(\prod_{p'_q < p < p_q} x_{i_q, i_q}^{(p)} \right)}_{h_q^{\mathcal{I}}} \quad (\text{for } 1 < q < k+1), \\ \rho_{k+1} &= \underbrace{\left(\prod_{p_k < p < p'_{k+1}} x_{j_k, j_k}^{(p)} \right)}_{g_{k+1}^{\mathcal{I}}} \cdot \underbrace{\left(x_{j_k, 1}^{(p'_{k+1})} \cdot \left(\prod_{p'_{k+1} < p < d} x_{1,1}^{(p)} \right) \cdot x_1^{(d)} \right)}_{h_{k+1}^{\mathcal{I}}}. \end{aligned}$$




FIG. 3. The partial derivative of F with respect to $x_{i_1, j_1}^{(p_1)}, \dots, x_{i_k, j_k}^{(p_k)}$.

We would like to lower bound the dimension of the vector space generated by the shifted k -partial derivatives of F . Clearly, we have

$$\dim(\langle \partial_k F \rangle_{\leq \ell}) \geq \dim(\text{span}(\mathcal{M})),$$

where $\mathcal{M} = \{m \cdot \partial_{\mathcal{I}} F \mid m \text{ a monomial of degree at most } \ell \text{ and } \mathcal{I} \in [n]^{2k}\}$. Since \mathcal{M} is a set of *monomials*, the dimension of the span of \mathcal{M} is exactly $|\mathcal{M}|$.

Another way of looking at \mathcal{M} is as follows. For $\mathcal{I} \in [n]^{2k}$, define

$$\mathcal{M}_{\mathcal{I}} := \{m' \mid m' \text{ a monomial of degree at most } \ell + d - k \text{ and } p(\mathcal{I}) \text{ divides } m'\}.$$

Since by definition, $p(\mathcal{I})$ is exactly $\partial_{\mathcal{I}} F$, we have

$$\begin{aligned} \mathcal{M} &= \{m \cdot p(\mathcal{I}) \mid m \text{ a monomial of degree at most } \ell \text{ and } \mathcal{I} \in [n]^{2k}\} \\ &= \{m' \mid m' \text{ of degree at most } \ell + d - k \text{ and } \exists \mathcal{I} \in [n]^{2k} \text{ such that } p(\mathcal{I}) \mid m'\} \\ &= \bigcup_{\mathcal{I} \in [n]^{2k}} \mathcal{M}_{\mathcal{I}}. \end{aligned}$$

We have shown the following.

Claim 4. For F and $\mathcal{M}_{\mathcal{I}}$ ($\mathcal{I} \in [n]^{2k}$) as defined above, we have $\dim(\langle \partial_k F \rangle_{\leq \ell}) \geq |\mathcal{M}|$, where $\mathcal{M} = \bigcup_{\mathcal{I} \in [n]^{2k}} \mathcal{M}_{\mathcal{I}}$.

Our aim is now to prove a lower bound on $|\mathcal{M}|$. We will need the following simple technical claim, the intuitive content of which is that any two distinct monomials $p(\mathcal{I})$ and $p(\mathcal{I}')$ are quite different. Recall that we do not distinguish between multilinear monomials over the variable set X and subsets of X .

Claim 5. For any $\mathcal{I}, \mathcal{I}' \in [n]^{2k}$, we have

$$|p(\mathcal{I}') \setminus p(\mathcal{I})| \geq \Delta(\mathcal{I}, \mathcal{I}') \cdot \left\lfloor \frac{r-1}{2} \right\rfloor,$$

where $\Delta(\mathcal{I}, \mathcal{I}')$ denotes the Hamming distance between \mathcal{I} and \mathcal{I}' and r equals $\lfloor \frac{d}{k+1} \rfloor - 1$.

Proof. Consider any $\mathcal{I}, \mathcal{I}' \in [n]^{2k}$. Say $\mathcal{I} = (i_1, j_1, \dots, i_k, j_k)$ and $\mathcal{I}' = (i'_1, j'_1, \dots, i'_k, j'_k)$. Then, using the notation from the definition of $p(\mathcal{I})$, we have

$$\begin{aligned} p(\mathcal{I}') \setminus p(\mathcal{I}) &\supseteq \bigcup_{q \in [k]} (g_{q+1}^{\mathcal{I}'} \setminus g_{q+1}^{\mathcal{I}}) \dot{\cup} \bigcup_{q \in [k]} (h_q^{\mathcal{I}'} \setminus h_q^{\mathcal{I}}) \\ &\supseteq \bigcup_{q \in [k]: j_q \neq j'_q} (g_{q+1}^{\mathcal{I}'} \setminus g_{q+1}^{\mathcal{I}}) \dot{\cup} \bigcup_{q \in [k]: i_q \neq i'_q} (h_q^{\mathcal{I}'} \setminus h_q^{\mathcal{I}}). \end{aligned}$$

(Recall that $A \dot{\cup} B$ denotes the union of *disjoint* sets A and B .)

Note that when $j_q \neq j'_q$, then the monomials $g_{q+1}^{\mathcal{I}}$ and $g_{q+1}^{\mathcal{I}'}$ do not share *any* variables and hence $|g_{q+1}^{\mathcal{I}'} \setminus g_{q+1}^{\mathcal{I}}| = |g_{q+1}^{\mathcal{I}'}| \geq \lfloor \frac{r-1}{2} \rfloor$. (Recall that by construction, there are at least r matrices between p_q and p_{q+1} and thus at least $\lfloor \frac{r-1}{2} \rfloor$ matrices between p_q and p'_{q+1} .) Similarly, when $i_q \neq i'_q$, we have $|h_q^{\mathcal{I}'} \setminus h_q^{\mathcal{I}}| \geq \lfloor \frac{r-1}{2} \rfloor$.

$$\begin{aligned} |p(\mathcal{I}') \setminus p(\mathcal{I})| &\geq \sum_{q \in [k]: j_q \neq j'_q} |g_{q+1}^{\mathcal{I}'} \setminus g_{q+1}^{\mathcal{I}}| + \sum_{q \in [k]: i_q \neq i'_q} |h_q^{\mathcal{I}'} \setminus h_q^{\mathcal{I}}| \\ &\geq \Delta(\mathcal{I}, \mathcal{I}') \cdot \left\lfloor \frac{r-1}{2} \right\rfloor, \end{aligned}$$

which completes the proof of the claim. \square

Claim 6. For any $\mathcal{I} \in [n]^{2k}$, we have $|\mathcal{M}_{\mathcal{I}}| = \binom{N+\ell}{\ell}$.

Proof. A monomial $m \in \mathcal{M}_{\mathcal{I}}$ iff there is a monomial m' of degree at most ℓ such that $m = m' \cdot p(\mathcal{I})$. Thus, $|\mathcal{M}_{\mathcal{I}}|$ is equal to the number of monomials of degree at most ℓ , which is $\binom{N+\ell}{\ell}$. \square

Claim 7. For any $\mathcal{I}, \mathcal{I}' \in [n]^{2k}$, we have $|\mathcal{M}_{\mathcal{I}} \cap \mathcal{M}_{\mathcal{I}'}| = \binom{N+\ell-|p(\mathcal{I}') \setminus p(\mathcal{I})|}{\ell-|p(\mathcal{I}') \setminus p(\mathcal{I})|}$.

Proof. Fix any $\mathcal{I}, \mathcal{I}'$ as above. Any monomial $m \in \mathcal{M}_{\mathcal{I}} \cap \mathcal{M}_{\mathcal{I}'}$ may be factored as $m = m' \cdot p(\mathcal{I}) \cdot (p(\mathcal{I}') \setminus p(\mathcal{I}))$. Note that the degree of m' can be bounded by $\ell + d - k - (d - k) - |p(\mathcal{I}') \setminus p(\mathcal{I})| = \ell - |p(\mathcal{I}') \setminus p(\mathcal{I})|$.

Thus, $|\mathcal{M}_{\mathcal{I}} \cap \mathcal{M}_{\mathcal{I}'}|$ is equal to the number of monomials of degree at most $\ell - |p(\mathcal{I}') \setminus p(\mathcal{I})|$, from which the claim follows. \square

Claim 8. Fix any $k, n \in \mathbb{N}$. Then there exists an $\mathcal{S} \subseteq [n]^{2k}$ such that

- $|\mathcal{S}| = \lfloor \left(\frac{n}{4}\right)^k \rfloor$,
- For all distinct $\mathcal{I}, \mathcal{I}' \in \mathcal{S}$, we have $\Delta(\mathcal{I}, \mathcal{I}') \geq k$.

Proof. Greedily pick vectors which have pairwise Hamming distance at least k . A standard volume argument (see, e.g., [Gur10]) shows that the set picked has size at least $\frac{n^{2k}}{\text{Vol}_n(2k, k)}$, where $\text{Vol}_n(2k, k)$ stands for the volume of the Hamming ball of radius k for strings of length $2k$ over an alphabet of size n . We can upper bound $\text{Vol}_n(2k, k)$ by $n^k \binom{2k}{k}$. This shows that there exists a set \mathcal{S} of size at least $n^k / \binom{2k}{k} \geq (n/4)^k$. We choose \mathcal{S} such that it has size exactly $\lfloor \left(\frac{n}{4}\right)^k \rfloor$. Hence the lemma follows. \square

Now we are ready to prove lower bound on the dimension of the space of shifted partial derivatives of F .

LEMMA 9. Let $k, \ell \in \mathbb{N}$ be arbitrary parameters such that $20k < d < \ell$ and $k \geq 2$. Then,

$$\dim(\langle \partial_k F \rangle_{\leq \ell}) \geq M \cdot \binom{N+\ell}{\ell} - M^2 \cdot \binom{N+\ell-d/10}{\ell-d/10},$$

where $M = \lfloor \left(\frac{n}{4}\right)^k \rfloor$.

Proof. Fix \mathcal{S} as guaranteed by Claim 8. By Claim 4, it suffices to lower bound $|\mathcal{M}|$. For this, we use inclusion-exclusion. Since $\mathcal{M} = \bigcup_{\mathcal{I} \in \mathcal{S}} \mathcal{M}_{\mathcal{I}}$, we have

$$\begin{aligned} |\mathcal{M}| &\geq \left| \bigcup_{\mathcal{I} \in \mathcal{S}} \mathcal{M}_{\mathcal{I}} \right| \\ (2) \quad &\geq \sum_{\mathcal{I} \in \mathcal{S}} |\mathcal{M}_{\mathcal{I}}| - \sum_{\mathcal{I} \neq \mathcal{I}' \in \mathcal{S}} |\mathcal{M}_{\mathcal{I}} \cap \mathcal{M}_{\mathcal{I}'}|. \end{aligned}$$

By Claim 6, we know that $|\mathcal{M}_{\mathcal{I}}| = \binom{N+\ell}{\ell}$. By Claims 7 and 5 and our choice of \mathcal{S} , we see that for any distinct $\mathcal{I}, \mathcal{I}' \in \mathcal{S}$, we have

$$|\mathcal{M}_{\mathcal{I}} \cap \mathcal{M}_{\mathcal{I}'}| \leq \binom{N+\ell-k \cdot \lfloor (r-1)/2 \rfloor}{\ell-k \cdot \lfloor (r-1)/2 \rfloor} \leq \binom{N+\ell-d/10}{\ell-d/10},$$

where the last inequality follows since $\lfloor (r-1)/2 \rfloor \geq d/10k$ for $k \leq d/20$ (recall that r denotes $\lfloor \frac{d}{k+1} \rfloor - 1$).

Plugging the above into (2), we obtain

$$|\mathcal{M}| \geq |\mathcal{S}| \cdot \binom{N+\ell}{\ell} - |\mathcal{S}|^2 \cdot \binom{N+\ell-d/10}{\ell-d/10}.$$

Since $|\mathcal{S}| = \lfloor \left(\frac{n}{4}\right)^k \rfloor$, the lemma follows. \square

The above lemma can be used to obtain a lower bound on the dimension of the shifted partial derivative space of $\text{IMM}_{n,d}$; see Appendix A.

5.2.2. A lower bound for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas computing F . We now prove the main lemma for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas.

LEMMA 10. *Let $n, d, D, t, k \in \mathbb{N}$ be such that $n \geq 10$, and $2 \leq k \leq d/160t$. Then, any $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula for F has top fan-in at least $\Omega((\frac{n^{3/4}k}{15D})^k)$.*

Proof. Recall that $N = (d-2)n^2 + 2dn = |X|$. By Fact 3, we can choose ℓ to be a positive integer such that $n^{1/16} \leq (\frac{N+\ell}{\ell})^t \leq n^{1/4}$ and $\ell \geq d$. We now analyze $\dim(\langle \partial_k F \rangle_{\leq \ell})$. By our choice of parameters, we have that $20k < d$. By Lemma 9, we have

$$\dim(\langle \partial_k F \rangle_{\leq \ell}) \geq \underbrace{M \cdot \binom{N+\ell}{\ell}}_{T_1} - \underbrace{M^2 \cdot \binom{N+\ell-d/10}{\ell-d/10}}_{T_2},$$

where $M = \lfloor (\frac{n}{4})^k \rfloor$.

However, for our choice of parameters, we have

$$\begin{aligned} \frac{T_1}{T_2} &= \frac{\binom{N+\ell}{\ell}}{M \cdot \binom{N+\ell-d/10}{\ell-d/10}} \\ &\geq \frac{1}{M} \cdot \left(\frac{N+\ell}{\ell} \right)^{d/10} && \text{(by Fact 2)} \\ &\geq \frac{n^{d/160t}}{M} && \text{(by our choice of } \ell) \\ &\geq \frac{n^k}{M} \geq 4^k \geq 2. \end{aligned}$$

Hence, we have

$$(3) \quad \dim(\langle \partial_k F \rangle_{\leq \ell}) \geq T_1 - T_2 \geq T_1/2 = \frac{M}{2} \cdot \binom{N+\ell}{\ell}.$$

Now, let C be a $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula for F of top fan-in s . Then, by Lemma 1 and using Stirling's approximation, we have

$$\dim(\langle \partial_k C \rangle_{\leq \ell}) \leq s \cdot \binom{D}{k} \cdot \binom{N+\ell+(t-1)k}{\ell+(t-1)k} \leq s \cdot \left(\frac{De}{k} \right)^k \cdot \binom{N+\ell+(t-1)k}{\ell+(t-1)k}.$$

An application of inequality (3) implies that we must have

$$s \cdot \left(\frac{De}{k} \right)^k \cdot \binom{N+\ell+(t-1)k}{\ell+(t-1)k} \geq \frac{M}{2} \cdot \binom{N+\ell}{\ell}.$$

Therefore,

$$\begin{aligned}
 s &\geq \frac{M}{2 \left(\frac{De}{k}\right)^k} \cdot \frac{\binom{N+\ell}{\ell}}{\binom{N+\ell+(t-1)k}{\ell+(t-1)k}} \\
 &\geq \frac{M}{2 \left(\frac{De}{k}\right)^k} \cdot \left(\frac{\ell}{N+\ell}\right)^{(t-1)k} \quad (\text{by Fact 2}) \\
 &\geq \frac{1}{4 \left(\frac{De}{k}\right)^k} \cdot \left(\frac{n}{4}\right)^k \cdot \left(\frac{\ell}{N+\ell}\right)^{(t-1)k} \quad (\text{by our choice of } M) \\
 &= \frac{1}{4} \cdot \left(\frac{nk}{4eD} \cdot \left(\frac{\ell}{N+\ell}\right)^{(t-1)}\right)^k \\
 &\geq \frac{1}{4} \cdot \left(\frac{nk}{15D \cdot \left(\frac{N+\ell}{\ell}\right)^t}\right)^k \geq \frac{1}{4} \cdot \left(\frac{n^{3/4}k}{15D}\right)^k \quad (\text{by our choice of } \ell, \text{ and } 4e < 15).
 \end{aligned}$$

This proves the lemma. \square

5.2.3. A lower bound for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas computing nice restrictions of $\text{IMM}_{n,d}$. We will show that any restriction in \mathcal{R} , when applied to $\text{IMM}_{n,d}$, yields a polynomial whose complexity is equivalent to that of F .

DEFINITION 11. For a polynomial $g \in \mathbb{F}[X]$ and a permutation ϕ of X , define $\phi(g)$ as the polynomial obtained by replacing in g each variable $x \in X$ by $\phi(x)$. Two polynomials $f, g \in \mathbb{F}[X]$ are said to be equivalent if there exists a permutation ϕ of X such that $f = \phi(g)$.

Note that if two polynomials $f, g \in \mathbb{F}[X]$ are equivalent, then their complexity with regard to $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas is the same. That is, there exists a $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula of size s for f iff there exists a $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula of size s for g . The following claim is unsurprising and directly implies Lemma 13 below. The proof of the claim is given in Appendix C.

Claim 12. For any restriction $\tau \in \mathcal{R}$ as defined above, $\text{IMM}_{n,d}|_{\tau}$ is equivalent to F .

LEMMA 13. Let $n, d, D, t, k \in \mathbb{N}$ be such that $n \geq 10$ and $2 \leq k \leq d/160t$. Let τ be a restriction in \mathcal{R} . Then any $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ for $\text{IMM}_{n,d}|_{\tau}$ has top fan-in at least $\Omega\left(\left(\frac{n^{3/4}k}{15D}\right)^k\right)$.

5.3. From set-multilinear formulas to $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas. In this section we reduce the case of a depth-4 set-multilinear formula to the case of bounded bottom fan-in by finding a suitable nice restriction.

LEMMA 14. Let n, d be large enough integers, and let $k, t \in \mathbb{N}$ be such that $t \geq 4k$. Let C be a set-multilinear $\Sigma\Pi\Sigma\Pi$ formula of size $s < n^{t/10}$. Then there exists a restriction $\tau \in \mathcal{R}$ such that $C|_{\tau}$ is a $\Sigma\Pi\Sigma\Pi^{[t]}$ formula.

Proof. We consider the uniform distribution over restrictions in \mathcal{R} and prove that with high probability the property in Lemma 14 holds.

Let us fix any bottom level Π gate G in C that has fan-in $t' > t$. Let m be the (set-multilinear) monomial computed by this Π gate. We can write m as a product of t' variables, each coming from a different variable set. That is, there exists a set

$S \subseteq [d]$, $|S| = t'$ such that $m = \prod_{i \in S} y^{(i)}$, where $y^{(i)}$ is a variable from X_i . We claim the following:

$$(4) \quad \Pr_{\tau}[G \text{ not set to 0 in } C|_{\tau}] \leq \frac{1}{n^{t/3}}.$$

To see this, first note that for all $p \in P'$, each variable in X_p survives with probability 1, i.e., the restriction does not set any variable in X_p to 0. But from the definition of our restriction and choice of t , $|P'| = 2k + 1$ and $t' \geq t \geq 4k$. Therefore, the monomial m has at least $t/3$ variables coming from matrices X_p ($p \notin P'$). And for all $p \notin P'$, the probability over τ that a variable survives in X_p is exactly $1/n$. Therefore, the probability that the monomial m survives is at most $(1/n)^{t/3}$. Therefore we have (4).

Since there are at most $s < n^{t/10}$ bottom level Π gates of fan-in greater than t , by a union bound, the probability that any such Π gate survives is at most $n^{t/10} \cdot \frac{1}{n^{t/3}} = o(1)$. \square

We can now bring everything together.

THEOREM 15. *For any large enough $n, d \in \mathbb{N}$, any set-multilinear $\Sigma\Pi\Sigma\Pi$ formula computing $\text{IMM}_{n,d}$ has size $n^{\Omega(\sqrt{d})}$.*

Proof. Let us choose k, t such that $d/320 \leq kt \leq d/160$ and $t = 4k$. Let C be a $\Sigma\Pi\Sigma\Pi$ set-multilinear formula of size s computing $\text{IMM}_{n,d}$ and say $s < n^{t/10}$. Let σ be the restriction guaranteed by Lemma 14. Therefore, we have that $C|_{\sigma}$ is a $\Sigma\Pi\Sigma\Pi^{[t]}$ formula of size at most s computing $\text{IMM}_{n,d}|_{\sigma}$, which is equivalent to F .

We can further convert $C|_{\sigma}$ into a $\Sigma\Pi^{[3d/t]}\Sigma\Pi^{[t]}$ formula, say, C' , so that the top fan-ins of both $C|_{\sigma}$ and C' are the same. This can be done (as explained in Remark 11 of [GKKS13]) by multiplying out polynomials of degree less than $t/2$ feeding into the Π gates at layer 3 so that at most one of them has degree less than $t/2$ and all others have degree between $t/2$ and t . This will imply that the fan-in of Π gates at layer 3 is at most $2d/t + 1$ which is bounded by $3d/t$. From Theorem 10, we have that for $kt \leq d/160$ any $\Sigma\Pi^{[3d/t]}\Sigma\Pi^{[t]}$ formula computing F has top fan-in at least $\frac{1}{4} \cdot (\frac{n^{3/4}kt}{15(3d)})^k$. Therefore, for the above choice of k, t , we get that F has size $n^{\Omega(k)}$. Therefore, we get that $s > \min\{n^{\Omega(t)}, n^{\Omega(k)}\}$. Since $kt = \Theta(d)$ and $t = 4k$, we have proved $s = n^{\Omega(\sqrt{d})}$. \square

Raz and Yehudayoff [RY09] proved that any $\Sigma\Pi\Sigma\Pi$ multilinear formula computing the determinant polynomial has size $\exp(\Omega(n^{1/27}))$. Using a technique similar to the one used above allows us to prove a stronger lower bound in the set-multilinear case (see Appendix B).

6. Lower bounds for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas and regular formulas. In this section, we derive our lower bounds for some flavors of $\Sigma\Pi\Sigma\Pi$ formulas. We start with a specific case that has been the focus of a few recent results ([GKKS13, KSS14]), the $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ model, where the Π gates at layers 3 and 1 have fan-ins bounded by D and t , respectively.

THEOREM 16. *Let $n, d, D, t \in \mathbb{N}$ be such that $n \geq 10$ and $1 \leq t \leq d/160$. Then, any $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula computing $\text{IMM}_{n,d}$ has top fan-in at least $(\frac{n^{3/4}d}{Dt})^{\Omega(d/t)}$. In particular, if $D = O(d/t)$, then any $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula computing $\text{IMM}_{n,d}$ has top fan-in at least $n^{\Omega(d/t)}$.*

Proof. Fix $k = \lfloor d/160t \rfloor$. We will show that any $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula C for $\text{IMM}_{n,d}$ has top fan-in at least $\Omega((\frac{n^{3/4}k}{15D})^k)$, which will prove the theorem.

But this follows from Lemma 10 and the simple fact that if $\text{IMM}_{n,d}$ has a $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula with top fan-in at most s , then so does any of its restrictions, and in particular F (defined in section 5.1) does. \square

For the next result of this section we need a few additional definitions. The *degree* of any node is defined to be the degree of the polynomial computed by it. The degree of the circuit (formula) is the degree of the output node. The *syntactic degree* is defined inductively. The syntactic degree of an input node is 1. The syntactic degree of $+$ gate is the maximum of the syntactic degrees of its children. The syntactic degree of \times gate is the sum of the syntactic degrees of its children. The syntactic degree of the circuit (formula) is the syntactic degree of its output node.

A formula is called *regular* if it is a layered formula, the alternate layers in the formula are labeled by $+$ and \times , for every layer the fan-in of all the gates at that layer is the same, and the syntactic degree of the formula is at most twice the degree of the formula.

Regular formulas were defined and studied recently by Kayal, Saha, and Saptharishi [KSS14]. They show the existence of a certain polynomial in VNP of degree d over N variables that has no regular formula of size less than $N^{\Omega(\log d)}$.

They also explicitly ask the following: is it true that any degree d polynomial in N variables that has a polynomial-sized ABP also has a regular formula of size $N^{o(\log d)}$? Here, we answer this question in the negative by showing that $\text{IMM}_{n,d}$ has no regular formulas of size less than $n^{\Omega(\log d)}$. We will need the following theorem of [KSS14].

THEOREM 17 (see [KSS14, Theorem 9]). *Let X be any set of N variables and let $F \in \mathbb{F}[X]$ be a polynomial of degree d with the property that there exists a $\delta > 0$ such that for any $t < d/100$, any $\Sigma\Pi^{[O(d/t)]}\Sigma\Pi^{[t]}$ formula computing the polynomial F has top fan-in at least $\exp(\delta(\frac{d}{t})\log N)$. Then, any regular formula computing F must be of size $N^{\Omega(\log d)}$.*

Though the theorem above is stated for $t < d/100$, it holds for $t < d/C$ for any constant C . We leave this check to the interested reader. Putting the above theorem together with Theorem 16, we have the next theorem.

THEOREM 18. *For large enough $n, d \in \mathbb{N}$, any regular formula for $\text{IMM}_{n,d}$ has size at least $n^{\Omega(\log d)}$.*

Note that the above is tight, up to the constant in the exponent, since the standard construction of an $n^{O(\log d)}$ sized formula for $\text{IMM}_{n,d}$ yields a regular formula.

7. Homogeneous multilinear depth-4 formulas. We will follow the same strategy as in section 5, first defining a set of nice restrictions, then proving a lower bound for $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas computing them, and finally showing that there exists a restriction in the set which changes a homogeneous multilinear formula to a $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula.

7.1. Nice restrictions of $\text{IMM}_{n,d}$. Unfortunately, we are unable to use the same family of restrictions for homogeneous multilinear formulas that we used in the set-multilinear formula case (see section 5.1 for the definition of these restrictions). The reason for this is that these restrictions leave variables $x_{i,j}^{(p)}$ ($p \in P'$) untouched. In particular, this implies that a monomial entirely made up of these variables will not reduce at all in degree upon applying such a restriction, which makes these restrictions unsuitable for the purpose of the random restriction argument that is used to reduce the bottom fan-in of formulas. In the set-multilinear case, this turns out to not be a problem, since each monomial can contain at most one variable from each X_p and consequently, a monomial of the above form can have degree at most $|P'|$, which turns out to be low enough for our setting of parameters.

However, in the multilinear case, this is an issue. Hence, we need to redo our argument with a slightly more complicated family of restrictions, which we now define. Our restrictions are once again related to the evenly spaced matrices chosen before. We will now choose new indices p''_q in a slightly different way. For $q \in [k+1]$, let $p''_q \in [d]$ be defined so that $p_{q-1} < p''_q < p_q$ and $\min\{p''_q - (p_{q-1} + 1), p_q - (p''_q + 2)\} \geq \lfloor \frac{r}{2} \rfloor - 1$. Let P''_1 denote the set $\{p_q \mid q \in [k]\}$ and P''_2 denote $\{p''_q \mid q \in [k+1]\}$. We define P'' to be $P''_1 \cup \{p, p+1 \mid p \in P''_2\}$.

DEFINITION 19. Let \mathcal{R} be the set of restrictions τ such that

1. for $p \in \{1, d\}$, there is a unique $j_p \in [n]$ such that $\tau(x_{j_p}^{(p)}) = *$,
2. for $p \notin P'' \cup \{1, d\}$, there is a permutation π_p of $[n]$ such that for any $i, j \in [n]$, $\tau(x_{i,j}^{(p)}) = *$ iff $j = \pi_p(i)$,
3. for $p \in P''_2$ and for all $i, j \in [n]$, there is at least one h in $[n]$ such that $\tau(x_{i,h}^{(p)}) = \tau(x_{h,j}^{(p+1)}) = *$,
4. for $p \in P''_1$, $|X_p \cap \tau^{-1}(*)| \geq n^{1.7}$.

In our final random restriction argument, we restrict the variables $x_{i,j}^{(p)}$ ($p \notin P''$) as in the set-multilinear case, and set all the variables $x_{i,j}^{(p)}$ ($p \in P''$) independently to $*$ with probability $1/n^{0.2}$. This ensures that any large degree multilinear monomial is set to 0 with high probability.¹

However, we also need to show that the dimension of the shifted partial derivative space is large after applying this restriction. We show that the random restriction that is sampled lies in \mathcal{R} with high probability. In particular, condition 4 (which holds with high probability since the expected number of surviving variables in X_p is $n^{1.8} \gg n^{1.7}$) in the definition of \mathcal{R} will ensure that we still have a large number of possible derivatives to work with, and condition 3 ensures that all these derivatives are nonzero. This is enough to recover the lower bound on the shifted partial derivative space of the restricted polynomial.

7.2. A lower bound for nice restrictions of $\text{IMM}_{n,d}$. We will not proceed exactly as we did in section 5, where we chose a specific nice restriction and showed the lower bound for the resulting polynomial. Instead, we study the polynomial obtained from a nice restriction in general.

7.2.1. The dimension of the space of shifted partial derivatives of nice restrictions of $\text{IMM}_{n,d}$. Let σ be a restriction in \mathcal{R} and F the polynomial $\text{IMM}_{n,d}|_\sigma$. Let \mathcal{A}_σ be the ABP corresponding to F .

As in section 5.1, we first analyze $\partial_{\mathcal{I}} F$ for $\mathcal{I} \in [n]^{2k}$. Let $\mathcal{I} = (i_1, j_1, \dots, i_k, j_k)$. Clearly, if there is a $q \in [k]$ such that $\sigma(x_{i_q, j_q}^{(p_q)}) = 0$, then we have $\partial_{\mathcal{I}} F = 0$. Hence we define

$$\mathcal{T} = \left\{ \mathcal{I} \in [n]^{2k} \mid \sigma(x_{i_q, j_q}^{(p_q)}) = * \text{ for all } q \in [k] \right\}.$$

By property 4 in Definition 19, we have $|\mathcal{T}| = |\prod_{q=1}^k (X_{p_q} \cap \sigma^{-1}(*))| \geq n^{(1.7) \cdot k}$.

For any $\mathcal{I} = (i_1, j_1, \dots, i_k, j_k) \in \mathcal{T}$, the polynomial $\partial_{\mathcal{I}} F$ is the sum of all monomials m such that $m = \rho_1 \rho_2 \dots \rho_{k+1}$, where ρ_q is a path from $v_{j_{q-1}}^{(p_{q-1})}$ to $v_{i_q}^{(p_q-1)}$ in \mathcal{A}_σ for all $2 \leq q \leq k$, ρ_1 is a path from vertex $v^{(0)}$ to $v_{i_1}^{(p_1)}$ in \mathcal{A}_σ , and ρ_{k+1} is a path from $v_{j_k}^{(p_k)}$ to vertex $v^{(d)}$ in \mathcal{A}_σ . Clearly, $\partial_{\mathcal{I}} F$ is a homogeneous polynomial of degree $d - k$.

¹There is nothing special about the constant 0.2; any other small enough constant would do as well.

Moreover, given any $\mathcal{I} \in \mathcal{T}$, the polynomial $\partial_{\mathcal{I}}F$ is nonzero; that is, there is a path from $v^{(0)}$ to $v^{(d)}$ in \mathcal{A}_{σ} which contains each edge $(v_{i_q}^{(p_q-1)}, v_{j_q}^{(p_q)})$ for $q \in [k]$ by properties 1, 2, and 3 in Definition 19.

We would like to lower bound $\dim(\langle \partial_k F \rangle_{\leq \ell})$, which is at least $\dim(\mathcal{V})$, where

$$\mathcal{V} = \text{span} \{m \cdot \partial_{\mathcal{I}}F \mid \mathcal{I} \in \mathcal{T} \text{ and } m \text{ a monomial of degree at most } \ell\}.$$

To lower bound $\dim(\mathcal{V})$, we will use the monomial-ordering technique as in [GKKS13] (see also [CLO97]). Let \geq be an arbitrary linear ordering of the variables in X and extend it to the lexicographic ordering on the set of all monomials in $\mathbb{F}[X]$ —the linear order on monomials is also denoted \geq . Given this monomial ordering \geq , for any polynomial $f \in \mathbb{F}[X]$, we denote by $\text{LM}(f)$ the leading monomial of f under this ordering (the ordering will be clear from context). The following fact will be useful.

Fact 20. Let \geq be any ordering as described above. Let $m_1, m_2 \in \mathbb{F}[X]$ be arbitrary monomials such that $m_1 \geq m_2$. Then, for any monomial m , we have

$$m_1 \cdot m \geq m_2 \cdot m.$$

This immediately implies that for $f, g \in \mathbb{F}[X]$, we have $\text{LM}(f \cdot g) = \text{LM}(f) \cdot \text{LM}(g)$.

Now, to lower bound $\dim(\mathcal{V})$, note that by Gaussian elimination, we know that

$$\begin{aligned} \dim(\mathcal{V}) &= |\{\text{LM}(f) \mid f \in \mathcal{V}\}| \\ &\geq |\{\text{LM}(m \cdot \partial_{\mathcal{I}}F) \mid \mathcal{I} \in \mathcal{T} \text{ and } m \text{ a monomial of degree at most } \ell\}| \\ &= |\{m \cdot \text{LM}(\partial_{\mathcal{I}}F) \mid \mathcal{I} \in \mathcal{T} \text{ and } m \text{ a monomial of degree at most } \ell\}|, \end{aligned}$$

where the last equality follows from Fact 20. We denote by $p(\mathcal{I})$ the monomial $\text{LM}(\partial_{\mathcal{I}}F)$. From the above, we see that $\dim(\mathcal{V}) \geq |\mathcal{M}|$, where

$$\mathcal{M} = \left\{m' \mid m' \text{ a monomial of degree at most } \ell + d - k \text{ and } \exists \mathcal{I} \in \mathcal{T} \text{ such that } p(\mathcal{I}) \mid m'\right\}.$$

Also, for $\mathcal{I} \in \mathcal{T}$, if we let

$$\mathcal{M}_{\mathcal{I}} = \{m' \mid m' \text{ a monomial of degree at most } \ell + d - k \text{ such that } p(\mathcal{I}) \mid m'\},$$

then we have $\mathcal{M} = \bigcup_{\mathcal{I} \in \mathcal{T}} \mathcal{M}_{\mathcal{I}}$.

The above arguments prove the following claim.

Claim 21. $\dim(\langle \partial_k F \rangle_{\leq \ell}) \geq |\mathcal{M}|$.

We need some technical claims, which are analogous to the claims proved in section 5.

Claim 22. For any $\mathcal{I}, \mathcal{I}' \in \mathcal{T}$, we have

$$|p(\mathcal{I}') \setminus p(\mathcal{I})| \geq \Delta(\mathcal{I}, \mathcal{I}') \cdot \left(\left\lfloor \frac{r}{2} \right\rfloor - 1\right),$$

where $\Delta(\mathcal{I}, \mathcal{I}')$ denotes the Hamming distance between \mathcal{I} and \mathcal{I}' .

Proof. Let $\mathcal{I}, \mathcal{I}' \in \mathcal{T}$. Say $\mathcal{I} = (i_1, j_1, \dots, i_k, j_k)$ and $\mathcal{I}' = (i'_1, j'_1, \dots, i'_k, j'_k)$. In the ABP \mathcal{A}_{σ} , let $g_{\mathcal{I}}^{\mathcal{I}}$ be the unique path from $v^{(0)}$ to $V_{p_{\mathcal{I}}-1}$ and for all $q \in [k]$, let $g_{q+1}^{\mathcal{I}}$ be the unique path from $v_{j_q}^{(p_q)}$ to $V_{p_{q+1}-1}$. For all $q \in [k]$, let $h_q^{\mathcal{I}}$ be the unique path from $V_{p_q'+1}$ to $v_{i_q}^{(p_q-1)}$, and let $h_{k+1}^{\mathcal{I}}$ be the unique path from $V_{p_{k+1}'+1}$ to $v^{(d)}$. (These paths are unique by property 2 in Definition 19.) For $q \in [k+1]$, define $g_q^{\mathcal{I}'}$

and $h_q^{\mathcal{I}'}$ in the same way for \mathcal{I}' . We have $p(\mathcal{I}) = m \cdot \prod_{q \in [k+1]} g_q^{\mathcal{I}} h_q^{\mathcal{I}}$ and $p(\mathcal{I}') = m' \cdot \prod_{q \in [k+1]} g_q^{\mathcal{I}'} h_q^{\mathcal{I}'}$, where m and m' are monomials on the variables $\bigcup_{q \in P_2''} (X_q \cup X_{q+1})$. Hence $|p(\mathcal{I}') \setminus p(\mathcal{I})| \geq \sum_{q \in [k+1]} (|g_q^{\mathcal{I}'} \setminus g_q^{\mathcal{I}}| + |h_q^{\mathcal{I}'} \setminus h_q^{\mathcal{I}}|)$.

If $i_q \neq i'_q$, the paths $h_q^{\mathcal{I}}$ and $h_q^{\mathcal{I}'}$ are edge disjoint (again by property 2 in Definition 19). In the same way, $g_{q+1}^{\mathcal{I}}$ and $g_{q+1}^{\mathcal{I}'}$ are edge disjoint if $j_q \neq j'_q$. Now all paths $g_q^{\mathcal{I}}$, $h_q^{\mathcal{I}}$ (and $g_q^{\mathcal{I}'}$, $h_q^{\mathcal{I}'}$) are of length at least $\lfloor r/2 \rfloor - 1$ by the choice of p_q and p'_q . \square

Claim 23. For any $\mathcal{I} \in \mathcal{T}$, we have $|\mathcal{M}_{\mathcal{I}}| = \binom{N+\ell}{\ell}$.

Claim 24. For any $\mathcal{I}, \mathcal{I}' \in \mathcal{T}$, we have $|\mathcal{M}_{\mathcal{I}} \cap \mathcal{M}_{\mathcal{I}'}| = \binom{N+\ell - |p(\mathcal{I}') \setminus p(\mathcal{I})|}{\ell - |p(\mathcal{I}') \setminus p(\mathcal{I})|}$.

The proofs of the above claims are identical to those of Claims 6 and 7 in section 5.2.1 and are hence omitted.

Claim 25. Fix any $k, n \in \mathbb{N}$. Then there exists an $\mathcal{S} \subseteq \mathcal{T}$ such that

- $|\mathcal{S}| = \lfloor (\frac{\sqrt{n}}{4})^k \rfloor$,
- for all distinct $\mathcal{I}, \mathcal{I}' \in \mathcal{S}$, we have $\Delta(\mathcal{I}, \mathcal{I}') \geq k$.

Proof. As in the proof of Claim 8, a volume argument shows that we can pick $\lfloor \frac{n^{1.7k}}{n^k \binom{2k}{k}} \rfloor \geq \frac{n^{k/2}}{2^{2k}}$ elements in \mathcal{T} with pairwise Hamming distance at least k . \square

Now we are ready to prove the lower bound on the dimension of the space of shifted partial derivatives of F .

LEMMA 26. *Let $k, \ell \in \mathbb{N}$ be arbitrary parameters such that $20k < d < \ell$ and $k \geq 2$. Then,*

$$\dim(\langle \partial_k F \rangle_{\leq \ell}) \geq M \cdot \binom{N+\ell}{\ell} - M^2 \cdot \binom{N+\ell-d/10}{\ell-d/10},$$

where $M = \lfloor (\frac{\sqrt{n}}{4})^k \rfloor$.

Proof. By Claim 21, it suffices to lower bound $|\mathcal{M}|$. Since $\mathcal{M} = \bigcup_{\mathcal{I} \in \mathcal{S}} \mathcal{M}_{\mathcal{I}}$, we have

$$\begin{aligned} |\mathcal{M}| &\geq \left| \bigcup_{\mathcal{I} \in \mathcal{S}} \mathcal{M}_{\mathcal{I}} \right| \\ (5) \quad &\geq \sum_{\mathcal{I} \in \mathcal{S}} |\mathcal{M}_{\mathcal{I}}| - \sum_{\mathcal{I} \neq \mathcal{I}' \in \mathcal{S}} |\mathcal{M}_{\mathcal{I}} \cap \mathcal{M}_{\mathcal{I}'}|. \end{aligned}$$

By Claim 23, we know that $|\mathcal{M}_{\mathcal{I}}| = \binom{N+\ell}{\ell}$. By Claims 24 and 22 and our choice of \mathcal{S} (Claim 25), we see that for any distinct $\mathcal{I}, \mathcal{I}' \in \mathcal{S}$, we have

$$|\mathcal{M}_{\mathcal{I}} \cap \mathcal{M}_{\mathcal{I}'}| \leq \binom{N+\ell-k(\lfloor r/2 \rfloor - 1)}{\ell-k(\lfloor r/2 \rfloor - 1)} \leq \binom{N+\ell-d/10}{\ell-d/10},$$

where the last inequality follows since $\lfloor r/2 \rfloor - 1 \geq d/10k$ for $k \leq d/20$.

Plugging the above into (5), we obtain

$$|\mathcal{M}| \geq |\mathcal{S}| \cdot \binom{N+\ell}{\ell} - |\mathcal{S}|^2 \cdot \binom{N+\ell-d/10}{\ell-d/10}.$$

Since $|\mathcal{S}| = \lfloor (\frac{\sqrt{n}}{4})^k \rfloor$, the lemma follows. \square

7.2.2. A lower bound for $\Sigma \Pi^{[D]} \Sigma \Pi^{[t]}$ formulas computing nice restrictions of $\text{IMM}_{n,d}$.

LEMMA 27. *Let $n, D, k, t, d \in \mathbb{N}$ be such that $2 \leq k \leq d/160t$. Let σ be a restriction in \mathcal{R} . Then any $\Sigma \Pi^{[D]} \Sigma \Pi^{[t]}$ formula C for $\text{IMM}_{n,d}|_{\sigma}$ has top fan-in $\Omega((\frac{n^{1/4}k}{15D})^k)$.*

Proof. Let $F = \text{IMM}_{n,d}|_{\sigma}$. We proceed as in Lemma 10. Fix $\ell \in \mathbb{N}$ such that $n^{1/16} \leq (\frac{N+\ell}{\ell})^t \leq n^{1/4}$, which exists by Fact 3. By Lemma 26, we have

$$\dim(\langle \partial_k F \rangle_{\leq \ell}) \geq \underbrace{M \cdot \binom{N+\ell}{\ell}}_{T_1} - \underbrace{M^2 \cdot \binom{N+\ell-d/10}{\ell-d/10}}_{T_2},$$

where $M = \lfloor (\frac{\sqrt{n}}{4})^k \rfloor$.

However, for our choice of parameters, we have

$$\begin{aligned} \frac{T_1}{T_2} &= \frac{\binom{N+\ell}{\ell}}{M \cdot \binom{N+\ell-d/10}{\ell-d/10}} \\ &\geq \frac{1}{M} \cdot \left(\frac{N+\ell}{\ell} \right)^{d/10} && \text{(by Fact 2)} \\ &\geq \frac{n^{d/160t}}{M} && \text{(by our choice of } \ell \text{)} \\ &\geq \frac{n^k}{M} \geq 4^k \geq 2. \end{aligned}$$

Hence, we have

$$(6) \quad \dim(\langle \partial_k F \rangle_{\leq \ell}) \geq T_1 - T_2 \geq T_1/2 = \frac{M}{2} \cdot \binom{N+\ell}{\ell}.$$

Now, let C be a $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formula for F of top fan-in s . Then, by Lemma 1 and using Stirling's approximation, we have

$$\dim(\langle \partial_k C \rangle_{\leq \ell}) \leq s \cdot \binom{D}{k} \cdot \binom{N+\ell+(t-1)k}{\ell+(t-1)k} \leq s \cdot \left(\frac{De}{k} \right)^k \cdot \binom{N+\ell+(t-1)k}{\ell+(t-1)k}.$$

An application of inequality (6) implies that we must have

$$s \cdot \left(\frac{De}{k} \right)^k \cdot \binom{N+\ell+(t-1)k}{\ell+(t-1)k} \geq \frac{M}{2} \cdot \binom{N+\ell}{\ell}.$$

Therefore,

$$\begin{aligned} s &\geq \frac{M}{2 \left(\frac{De}{k} \right)^k} \cdot \frac{\binom{N+\ell}{\ell}}{\binom{N+\ell+(t-1)k}{\ell+(t-1)k}} \\ &\geq \frac{M}{2 \left(\frac{De}{k} \right)^k} \cdot \left(\frac{\ell}{N+\ell} \right)^{(t-1)k} && \text{(by Fact 2)} \\ &\geq \frac{1}{4 \left(\frac{De}{k} \right)^k} \cdot \left(\frac{\sqrt{n}}{4} \right)^k \cdot \left(\frac{\ell}{N+\ell} \right)^{(t-1)k} && \text{(by our choice of } M \text{)} \\ &= \frac{1}{4} \cdot \left(\frac{\sqrt{n}k}{4eD} \cdot \left(\frac{\ell}{N+\ell} \right)^{(t-1)} \right)^k \\ &\geq \frac{1}{4} \cdot \left(\frac{\sqrt{n}k}{15D \cdot \left(\frac{N+\ell}{\ell} \right)^t} \right)^k \\ &\geq \frac{1}{4} \cdot \left(\frac{n^{1/4}k}{15D} \right)^k && \text{(by our choice of } \ell \text{ and because } 4e < 15). \quad \square \end{aligned}$$

7.3. From homogeneous multilinear formulas to $\Sigma\Pi^{[D]}\Sigma\Pi^{[t]}$ formulas.

LEMMA 28. *The following holds for any large enough $n \in \mathbb{N}$ and any k, t, d such that $1 \leq k, t \leq d \leq n^{O(1)}$. Let C be a homogeneous multilinear $\Sigma\Pi\Sigma\Pi$ formula of size $s < n^{t/10}$; there is a restriction $\tau \in \mathcal{R}$ such that $C|_\tau$ is $\Sigma\Pi\Sigma\Pi^{[t]}$.*

Proof. As in section 5, we will use a probabilistic argument. We define a suitable distribution \mathcal{D} over restrictions $\sigma : X \rightarrow \{0, *\}$ in general and show that with high probability over the choice of $\sigma \sim \mathcal{D}$, the restriction $F := \text{IMM}_{n,d}|_\sigma$ belongs to \mathcal{R} and satisfies the required property. We specify the distribution \mathcal{D} by describing how to sample a single restriction σ .

- For $p \in \{1, d\}$, pick $j_p \in [n]$ uniformly at random. Set $\sigma(x_j^{(p)}) = *$ if $j = j_p$ and 0 otherwise.
- For $p \notin P'' \cup \{1, d\}$, pick an independent and uniformly random permutation π_p of $[n]$. Set $\sigma(x_{i,j}^{(p)}) = *$ if $j = \pi_p(i)$ and 0 otherwise.
- For each $x \in \bigcup_{p \in P''} X_p$, set $\sigma(x) = *$ independently with probability $\frac{1}{n^{0.2}}$.

We denote by \mathcal{A}_σ the ABP corresponding to this restriction.

σ belongs to \mathcal{R} with high probability. From the description of \mathcal{D} above, it follows that for any $\sigma \sim \mathcal{D}$, the restriction σ always satisfies properties 1 and 2 of Definition 19. So we only need to consider properties 3 and 4.

Let \mathcal{E}_3 denote the event that property 3 is not satisfied. Fix any $p \in P_2''$ and any $i, j \in [n]$. For each $v \in V_p$, define the $\{0, 1\}$ -valued random variable $Y_{v,i,j}$ so that $Y_{v,i,j} = 1$ iff the edges $(v_i^{(p-1)}, v)$ and $(v, v_j^{(p+1)})$ both survive in the ABP \mathcal{A}_σ and let $Y'_{p,i,j} = \sum_{v \in V_p} Y_{v,i,j}$. Since each of the edges $(v_i^{(p-1)}, v)$ and $(v, v_j^{(p+1)})$ survives independently with probability $1/n^{0.2}$, it follows that $\Pr_\sigma[Y_{v,i,j} = 1] = 1/n^{0.4}$.

Note that the random variables $Y_{v,i,j}$ are mutually independent. Hence, we have

$$\begin{aligned} \Pr_\sigma[\text{There is no path from } v_i^{(p-1)} \text{ to } v_j^{(p+1)}] &= \Pr_\sigma[Y'_{p,i,j} = 0] = \Pr_\sigma\left[\bigwedge_v Y_{v,i,j} = 0\right] \\ &= \prod_v \Pr_\sigma[Y_{v,i,j} = 0] \\ &= \left(1 - \frac{1}{n^{0.4}}\right)^n = \exp(-n^{\Omega(1)}). \end{aligned}$$

Union bounding over the choices of p, i, j , we see that $\Pr[\mathcal{E}_3] = \exp(-n^{\Omega(1)}) = o(1)$.

Let \mathcal{E}_4 denote the event that property 4 is not satisfied. Fix any $p \in P_1''$ and for each $x \in X_p$, define the $\{0, 1\}$ -valued random variable $Z(x)$ that is 1 iff $\sigma(x) = *$; let $Z_p = \sum_{x \in X_p} Z(x)$.

We have $\mathbf{E}_\sigma[Z_p] = \sum_{x \in X_p} \Pr_\sigma[Z(x) = 1] = |X_p| \cdot n^{-0.2} = n^{1.8}$. Moreover, the random variables $Z(x)$ are mutually independent and hence, by a Chernoff bound (see, e.g., [DP09, Chapter 1]), we have for large enough n ,

$$\Pr_\sigma[Z_p < n^{1.7}] \leq \Pr_\sigma[Z_p < \mathbf{E}[Z_p]/2] = \exp(-n^{\Omega(1)}).$$

Thus, union bounding over the at most $d = n^{O(1)}$ choices of $p \in P_1''$, we have $\Pr[\mathcal{E}_4] \leq n^{O(1)} \cdot \exp(-n^{\Omega(1)}) = \exp(-n^{\Omega(1)}) = o(1)$.

Thus, we have

$$\Pr_\sigma[\sigma \text{ not nice}] = \Pr_\sigma[\mathcal{E}_3 \vee \mathcal{E}_4] = o(1).$$

$C|_\sigma$ is $\Sigma\Pi\Sigma\Pi^{[t]}$ with high probability. We need to show that, with high probability, the fan-in of bottom Π gates of $C|_\sigma$ is at most t . In other words, we need to show that with high probability, all bottom Π gates in C that have fan-in greater than t are set to 0 by σ .

Let us fix any bottom level Π gate G in C that has fan-in greater than t . Let m be the (multilinear) monomial computed by this Π gate.² Write $m = m_1 \cdot m_2 \cdots m_d$, where $m_p = \prod_{x \in X_p: x|m} x$. Let t_p denote $\deg(m_p)$. We have $\sum_{p \in [d]} t_p = \deg(m) > t$. We claim that

$$(7) \quad \Pr_\sigma[G \text{ not set to 0 in } C|_\sigma] \leq \frac{1}{n^{t/5}}.$$

To see this, note that by the independence of the random restriction σ across the different X_p ($p \in [d]$), we have

$$(8) \quad \Pr_\sigma[G \text{ not set to 0 in } C|_\sigma] = \prod_{p \in [d]} \underbrace{\Pr_\sigma[\text{No variable in } m_p \text{ set to 0}]}_{\alpha_p}.$$

Now, fix any $p \in [d]$. We upper bound α_p based on a case analysis.

- If $p \in \{1, d\}$, $\alpha_p = 0$ if $t_p > 1$ and $\alpha_p = 1/n^{t_p}$ otherwise.
- If $p \in P''$, $\alpha_p = 1/n^{t_p/5}$.
- If $p \notin P'' \cup \{1, d\}$, then $\alpha_p = 0$ if the monomial m_p contains at least two variables from any row or column of X_p . Otherwise, $\alpha_p = \prod_{z=0}^{t_p-1} \frac{1}{n-z} \leq \left(\frac{1}{n-t_p}\right)^{t_p} \leq \left(\frac{1}{\sqrt{n}}\right)^{t_p/2} \leq \frac{1}{n^{t_p/4}}$.

Thus, we see that in all cases, we have $\alpha_p \leq \frac{1}{n^{t_p/5}}$. Substituting in (8), we have $\Pr_\sigma[G \text{ not set to 0}] \leq 1/n^{\sum_{p \in [d]} (t_p/5)} \leq 1/n^{t/5}$, which proves (7).

Since there are at most $s < n^{t/10}$ bottom level Π gates of fan-in greater than t , by a union bound, the probability that any such Π gate survives is at most $n^{t/10} \cdot \frac{1}{n^{t/5}} = o(1)$.

Finally, another union bound shows that the probability that either $\sigma \notin \mathcal{R}$ or $C|_\sigma$ not a $\Sigma\Pi\Sigma\Pi^{[t]}$ formula is $o(1)$. This proves the lemma. \square

THEOREM 29. *For any large enough $n, d \in \mathbb{N}$ such that $d = n^{O(1)}$ any homogeneous multilinear $\Sigma\Pi\Sigma\Pi$ formula computing $\text{IMM}_{n,d}$ has size $n^{\Omega(\sqrt{d})}$.*

Proof. Let us choose k, t such that $d/320 \leq kt \leq d/160$ and $t = 4k$. Let C be a $\Sigma\Pi\Sigma\Pi$ homogeneous multilinear formula of size s computing $\text{IMM}_{n,d}$ and say $s < n^{t/10}$. Let σ be the nice restriction guaranteed by Lemma 28. Therefore, we have that $C|_\sigma$ is a $\Sigma\Pi\Sigma\Pi^{[t]}$ formula computing $F = \text{IMM}_{n,d}|_\sigma$.

As in the proof of Theorem 15 we can further convert $C|_\sigma$ into a $\Sigma\Pi^{[3d/t]}\Sigma\Pi^{[t]}$ formula, say, C' , by multiplying out polynomials of degree less than $t/2$ feeding into the Π gates at layer 3. This can be done keeping the top fan-ins of $C|_\sigma$ and C' the same. From Lemma 27, we have that for $kt \leq d/160$ any $\Sigma\Pi^{[3d/t]}\Sigma\Pi^{[t]}$ formula computing F has top fan-in at least $\Omega\left(\left(\frac{n^{1/4}kt}{15(3d)}\right)^k\right)$. Therefore, for the above choice of k, t , we get that F has size $n^{\Omega(k)}$. Therefore, we get that $s > \min\{n^{\Omega(t)}, n^{\Omega(k)}\}$. Since $kt = \Theta(d)$ and $t = 4k$, we have proved $s = n^{\Omega(\sqrt{d})}$. \square

Remark 30. Note that we used multilinearity only in the proof of Lemma 28. Even there, multilinearity was not strictly necessary. We only needed that the $\Sigma\Pi\Sigma\Pi$ formula C has the property that all the Π gates on layer 1, just above the input variables, are multilinear.

²This is the only place where multilinearity is necessary.

8. Discussion. Our aims in this paper were twofold: to explore the limits of depth reduction and to understand better the arithmetic circuit complexity of $\text{IMM}_{n,d}$. We have made progress on both fronts, but many interesting questions remain unanswered.

We have shown that Tavenas' result [Tav13] is optimal up to a polynomial (i.e., up to constant factors in the exponent) even for polynomials in the class VP_s , by showing that any $\Sigma\Pi^{[O(\sqrt{d})]}\Sigma\Pi^{[O(\sqrt{d})]}$ formulas for $\text{IMM}_{n,d}$ has size $\exp(\Omega(\sqrt{d} \log N))$. Our results also answer a question of Kayal, Saha, and Saptharishi [KSS14] regarding the simulation of polynomial-sized circuits by regular formulas. Thus, in order to use depth reduction based techniques to prove a separation between VP and VNP , we will need to exhibit a polynomial in VNP that requires $\Sigma\Pi^{[O(\sqrt{d})]}\Sigma\Pi^{[O(\sqrt{d})]}$ formulas of size $\exp(\omega(\sqrt{d} \log N))$ to compute it.

One might also wonder whether lower bounds for weaker models, such as arithmetic formulas, might follow from either the shifted partial derivative technique or by depth reduction. Can one show nontrivial upper bounds on the dimension of the shifted partial derivative space of polynomials computed by small formulas? The result of Kumar and Saraf [KS14a] shows that the answer to the question as stated is probably no, but perhaps a modification of the shifted partial derivative measure (e.g., [KLSS14b]) could work.

Coming to the question of the complexity of $\text{IMM}_{n,d}$, we have been able to pin down almost exactly the $\Sigma\Pi\Sigma\Pi$ complexity of $\text{IMM}_{n,d}$ in the set-multilinear and, more generally, in the homogeneous multilinear setting. Can we extend these results to show that, in general, that set-multilinear formulas of product depth r (for constant r) computing $\text{IMM}_{n,d}$ must have size $\exp(\Omega(d^{1/r} \log n))$?

This would count as tangible progress toward the goal of showing that set-multilinear formulas for $\text{IMM}_{n,d}$ must have size $n^{\Omega(\log d)}$. Raz [Raz10] has shown that, for $d = o(\log n / \log \log n)$, the set-multilinear formula complexity of $\text{IMM}_{n,d}$ and the formula complexity of $\text{IMM}_{n,d}$ are polynomially related, and hence, a superpolynomial lower bound for set-multilinear formulas in this regime would immediately imply a separation between VP_s and VP_E .

Appendix A. Lower bound on the dimension of the shifted partial derivative space of $\text{IMM}_{n,d}$. We now prove a proposition which states a lower bound on the dimension of the shifted partial derivative space of $\text{IMM}_{n,d}$. Although this is not necessary to obtain any of our lower bounds, it may be of general interest. We first need the following intuitive statement.

LEMMA 31. *Let X be a set of variables with $|X| = N$ and $f, g \in \mathbb{F}[X]$ be such that g is a restriction of f . Then, for any $k, \ell \geq 1$ we have $\dim(\langle \partial_k f \rangle_{\leq \ell}) \geq \dim(\langle \partial_k g \rangle_{\leq \ell})$.*

Proof. Given $x \in X$, let $S_x : \mathbb{F}[X] \rightarrow \mathbb{F}[X]$ be the linear map that maps any polynomial f' to the polynomial g' obtained by setting x to 0 in f' . Similarly, we use $S_\sigma : \mathbb{F}[X] \rightarrow \mathbb{F}[X]$ to denote the linear map that maps f' to the polynomial g' obtained by setting all $x \in \sigma^{-1}(0)$ to 0 in f' . Note that S_σ is the composition of all the S_x for $x \in \sigma^{-1}(0)$ (the order of composition is insignificant).

We want to show that for any $f \in \mathbb{F}[X]$, we have $\dim(\langle \partial_k f \rangle_{\leq \ell}) \geq \dim(\langle \partial_k (S_\sigma f) \rangle_{\leq \ell})$. From the reasoning in the previous paragraph, it suffices to show that for any $x \in X$, we have $\dim(\langle \partial_k f \rangle_{\leq \ell}) \geq \dim(\langle \partial_k (S_x f) \rangle_{\leq \ell})$. We can write f as $f = \sum_{j=0}^a x^j f_j$, where $f_j \in \mathbb{F}[X \setminus \{x\}]$ for $j \leq a$. Using this notation, $S_x f$ is simply the polynomial f_0 . Thus, what we want to show is that $\dim(\langle \partial_k f \rangle_{\leq \ell}) \geq \dim(\langle \partial_k (f_0) \rangle_{\leq \ell})$.

We introduce some useful notation here. Let the variable in X be denoted x_1, \dots, x_N . For $\mathbf{i} \in \mathbb{N}^N$ such that $i_1 + \dots + i_N = k$ and $g \in \mathbb{F}[X]$, we denote by $\mathbf{x}^{\mathbf{i}}$ the monomial $x_1^{i_1} \dots x_N^{i_N}$ and by $\frac{\partial^k g}{\partial \mathbf{x}^{\mathbf{i}}}$ the polynomial $\partial^k g / (\partial^{i_1} x_1 \dots \partial^{i_N} x_N)$.

Let L denote the dimension of $\langle \partial_k(f_0) \rangle_{\leq \ell}$. Choose an arbitrary basis B for this space. Such a basis may be written as

$$B = \left\{ m_1 \cdot \frac{\partial^k f_0}{\partial \mathbf{x}^{\mathbf{i}^{(1)}}}, m_2 \cdot \frac{\partial^k f_0}{\partial \mathbf{x}^{\mathbf{i}^{(2)}}}, \dots, m_L \cdot \frac{\partial^k f_0}{\partial \mathbf{x}^{\mathbf{i}^{(L)}}} \right\}$$

for some monomials m_1, \dots, m_L of degree at most ℓ and $\mathbf{i}^{(1)}, \dots, \mathbf{i}^{(L)} \in \mathbb{N}^N$ such that for each $r \in [L]$, $i_1^{(r)} + \dots + i_N^{(r)} = k$. In particular, we must have $\frac{\partial^k f_0}{\partial \mathbf{x}^{\mathbf{i}^{(r)}}} \neq 0$ for each $r \in [L]$. As $f_0 \in \mathbb{F}[X \setminus \{x\}]$, this implies that $x \nmid \mathbf{x}^{\mathbf{i}^{(r)}}$.

We claim that the elements

$$B' = \left\{ m_1 \cdot \frac{\partial^k f}{\partial \mathbf{x}^{\mathbf{i}^{(1)}}}, m_2 \cdot \frac{\partial^k f}{\partial \mathbf{x}^{\mathbf{i}^{(2)}}}, \dots, m_L \cdot \frac{\partial^k f}{\partial \mathbf{x}^{\mathbf{i}^{(L)}}} \right\}$$

of $\langle \partial_k f \rangle_{\leq \ell}$ are linearly independent. This would prove that $\dim(\langle \partial_k f \rangle_{\leq \ell}) \geq L$ and finish the proof of the lemma.

To see that the elements of B' are linearly independent, we partition the monomials m_r ($r \in [L]$) depending on the highest power of x dividing them. For $j \leq \ell$, let $T_j = \{r \in [L] \mid x^j \mid m_r \text{ but } x^{j+1} \nmid m_r\}$. Now, fix any nonzero linear combination of the elements of B' , say,

$$H = \sum_{r \in [L]} \alpha_r \cdot m_r \cdot \frac{\partial^k f}{\partial \mathbf{x}^{\mathbf{i}^{(r)}}},$$

where $\alpha_r \in \mathbb{F}$ for each $r \in [L]$. Let j be the least element of $\{0\} \cup [\ell]$ such that there is an $r \in T_j$ with $\alpha_r \neq 0$. Consider the coefficient H_j of x^j in H as a polynomial in $\mathbb{F}[X \setminus \{x\}]$. Since $x \nmid \mathbf{x}^{\mathbf{i}^{(r)}}$ for any $r \in [L]$, it can be seen that

$$x^j H_j = \sum_{r \in T_j} \alpha_r \cdot m_r \cdot \frac{\partial^k f_0}{\partial \mathbf{x}^{\mathbf{i}^{(r)}}}.$$

(Notice that f has been replaced by f_0 in the equation above.) But by the linear independence of the elements in B , we have $H_j \neq 0$. Hence so is H . Thus, the elements of B' are linearly independent. \square

PROPOSITION 32. *Let $k, \ell \in \mathbb{N}$ be arbitrary parameters such that $20k < d < \ell$ and $k \geq 2$. Then,*

$$\dim(\langle \partial_k \text{IMM}_{n,d} \rangle_{\leq \ell}) \geq M \cdot \binom{N + \ell}{\ell} - M^2 \cdot \binom{N + \ell - d/10}{\ell - d/10},$$

where $M = \left\lfloor \left(\frac{n}{4}\right)^k \right\rfloor$.

Proof. The proof straightaway follows from Lemmas 31 and 9, since the polynomial F from Lemma 9 is a restriction of $\text{IMM}_{n,d}$. \square

Appendix B. Lower bounds for set-multilinear formulas computing the determinant.

PROPOSITION 33. *A $\Sigma\Pi\Sigma\Pi$ set-multilinear formula computing the determinant of a generic $n \times n$ matrix has size $\exp(\Omega(n^{1/2}))$.*

Proof sketch. We consider the determinant polynomial of a generic $n \times n$ matrix which is set-multilinear with respect to its columns. Let $S, T \subseteq [n]$ and $|S| = |T| =$

$n/2$ be chosen uniformly at random. Let ϕ be a random bijection from $[n] \setminus S$ to $[n] \setminus T$. Now consider a restriction σ as given below: $\sigma(x_{i,j}) = *$ if $i \in S, j \in T$, $\sigma(x_{i,\phi(i)}) = 1$ if $i \in [n] \setminus S$, and $\sigma(x_{i,j}) = 0$ otherwise.³

Under this restriction, an $n \times n$ determinant reduces to an $n/2 \times n/2$ determinant of the matrix defined by S, T . Also, under this restriction a $\Sigma\Pi\Sigma\Pi$ set-multilinear formula of size $s = 2^{o(\sqrt{n})}$ computing the determinant reduces to a $\Sigma\Pi\Sigma\Pi^{[t]}$ set-multilinear formula, where $t = \Theta(\sqrt{n})$, with high probability. To see this, we first note that any set-multilinear monomial m of degree larger than t is set to 0 with probability $1 - \exp(-\Omega(\sqrt{n}))$.

We sketch a proof of this bound. Let C be the set of columns such that variables from these columns appear in the monomial m : note that $|C| \geq t$. Since T is chosen uniformly at random from $\binom{[n]}{n/2}$, it follows from standard concentration bounds for negatively associated random variables (see, e.g., [DP09, Chapter 3]) that the probability that C does not contain $t/4$ columns *outside* T is at most $\exp(-\Omega(t))$. We condition on a choice of T such that $|C \setminus T| \geq t/4$. Now, the monomial m contains at least $t/4$ variables outside T . The probability that none of these variables is set to 0 by the restriction is exactly the probability that for each such variable $x_{i,j}$ appearing in m , we have $j = \phi(i)$; the probability of this event (over the choice of the random bijection ϕ) can be seen to be $n^{-\Omega(t)}$. Hence, we see that the probability that any monomial m of degree at least t is not set to 0 by our random restriction is at most $\exp(-\Omega(t)) + n^{-\Omega(t)} = \exp(-\Omega(t))$.

Hence by a union bound over all monomials m of degree at least t appearing in the formula (there are at most s of them), we see that with high probability over the choice of the restriction, we get a $\Sigma\Pi\Sigma\Pi^{[t]}$ formula. It can be easily checked that this formula is set-multilinear as well and hence, in particular, homogeneous. Fix any such restriction. Now, applying the main lower bound of [GKKS13]—which shows that any homogeneous $\Sigma\Pi\Sigma\Pi^{[t]}$ formula for the $r \times r$ determinant must have size at least $\exp(\Omega(r/t))$ —gives a lower bound of $\exp(\Omega(n^{1/2}))$ for $\Sigma\Pi\Sigma\Pi$ set-multilinear formula computing the determinant polynomial. \square

Appendix C. Miscellaneous proofs.

Proof of Claim 12. Recall that the polynomial F was obtained from $\text{IMM}_{n,d}$ by the restriction σ :

$$\sigma(x) = \begin{cases} 0 & \text{if } x = x_j^{(1)} \text{ for some } j \neq 1, \text{ or} \\ & \text{if } x = x_j^{(d)} \text{ for some } j \neq 1, \text{ or} \\ & \text{if } x = x_{i,j}^{(p)} \text{ for } i \neq j \text{ and } p \notin P' \cup \{1, d\}, \\ * & \text{otherwise.} \end{cases}$$

Consider a restriction τ obtained by picking j_1, j_d and permutations π_p for $p \notin P' \cup \{1, d\}$. We wish to define a permutation ϕ such that $\text{IMM}_{n,d}|_\tau = \phi(F)$. We start necessarily by defining $\phi(x_1^{(1)}) = x_{j_1}^{(1)}$.

We then define ϕ for the matrices X_p for $p \in \{2, \dots, p'_1 - 1\}$. Once again viewing our polynomials as ABPs: let \mathcal{A}_σ be the graph corresponding to F and \mathcal{A}_τ the graph corresponding to $\text{IMM}_{n,d}|_\tau$. To lighten notation, we only write the index of the vertex at each layer; the path $1, 1, \dots, 1$ from layer p to layer q is thus the path $v_1^{(p)}, v_1^{(p+1)}, \dots, v_1^{(q)}$.

³This restriction differs slightly in form from the restrictions we used for $\text{IMM}_{n,d}$ since variables can be set to any of 0, 1, or $*$ instead of simply 0 or $*$.

Note that there are n pairwise edge-disjoint paths i, i, \dots, i going from layer 1 to layer $p'_1 - 1$ in \mathcal{A}_σ , one path for each $i \in [n]$. There are also n pairwise edge-disjoint paths going from layer 1 to layer $p'_1 - 1$ in \mathcal{A}_τ , these paths being defined by the composition of the permutations π_p for $p \in \{2, \dots, p'_1 - 1\}$. We define ϕ successively in each matrix $X_2, X_3, \dots, X_{p'_1-1}$ so that it sends the path $1, 1, \dots, 1$ to the path starting from vertex j_1 . Let i' be the end-vertex of this path at layer $p'_1 - 1$. We then define ϕ over the variables in $X_{p'_1}$ by requiring only that $\phi(x_{1,j}^{p'_1})$ be equal to $x_{i',j}^{p'_1}$ for all $j \in [n]$.

We will then define the permutation separately over the following intervals of matrix indices: $\{p'_1 + 1, \dots, p_1\}, \{p_1 + 1, \dots, p'_2\}, \dots, \{p_k + 1, \dots, p'_{k+1} - 1\}$. We will describe the case of the interval $\{p'_1 + 1, \dots, p_1\}$ in some detail; other intervals can be treated in a similar fashion, with a slight exception for the last one.

Once again in \mathcal{A}_σ there are n pairwise edge-disjoint paths from layer p'_1 to layer $p_1 - 1$. Similarly, there is a set of n pairwise edge-disjoint paths from layer p'_1 to layer $p_1 - 1$ in the graph \mathcal{A}_τ . We define ϕ such that it sends the path i, i, \dots, i to the path starting at vertex i of layer p'_1 in \mathcal{A}_τ . For example, we send the path $1, 1, \dots, 1$ in \mathcal{A}_σ to the path $1, \pi_{p'_1+1}(1), \pi_{p'_1+2} \circ \pi_{p'_1+1}(1), \dots, \pi_{p_1-1} \circ \pi_{p_1-2} \circ \dots \circ \pi_{p'_1+2} \circ \pi_{p'_1+1}(1)$. Next we define the effect of ϕ on the matrix X_{p_1} . Define the permutation α by setting $\alpha(i)$ to be the index of the end-vertex of the path starting at vertex i of layer p'_1 in \mathcal{A}_τ , i.e., $\alpha(i) = \pi_{p_1-1} \circ \pi_{p_1-2} \circ \dots \circ \pi_{p'_1+2} \circ \pi_{p'_1+1}(i)$. We then let $\phi(x_{i,j}^{(p_1)}) = x_{\alpha(i),j}^{(p_1)}$ for all $i \in [n]$. A path i, i, \dots, i in \mathcal{A}_σ is sent by ϕ to the path starting at i and ending at $\alpha(i)$ in \mathcal{A}_τ . This path i, i, \dots, i in \mathcal{A}_σ could then be extended by any edge (i, j) in X_{p_1} . Since we have sent the path i, i, \dots, i to the path ending in $\alpha(i)$, by setting $\phi(x_{i,j}^{(p_1)}) = x_{\alpha(i),j}^{(p_1)}$ we ensure that any path from vertex i of layer p'_1 to vertex j of layer p_1 in \mathcal{A}_σ is sent to a path from vertex i of layer p'_1 to vertex j of layer p_1 in \mathcal{A}_τ . We can then start the process again with the next interval, with the exception of the last one, where we do not set ϕ for the variables in $X_{p'_{k+1}}$ yet. Let α be the permutation obtained as above for the interval $\{p_k + 1, \dots, p'_{k+1} - 1\}$.

To define ϕ on the end of the graph, we will start from the end. Clearly, we must set $\phi(x_1^{(d)}) = x_{j_d}^{(d)}$. We then need to do the interval $\{p'_{k+1} + 1, \dots, d - 1\}$. We define ϕ on the interval $\{p'_{k+1} + 1, \dots, d - 1\}$ by sending the path $1, 1, \dots, 1$ to the unique path ending at vertex j_d of layer $d - 1$ in \mathcal{A}_τ . Let j' be the index at layer p'_{k+1} of the starting vertex of this path. To define ϕ on $X_{p'_{k+1}}$, we only require of ϕ that it send $x_{i,1}^{(p'_{k+1})}$ to $x_{\alpha(i),j'}^{(p'_{k+1})}$ for all $i \in [n]$.

Then $\text{IMM}_{n,d}|_\tau = \phi(F)$. \square

Acknowledgments. The authors wish to thank the funding agency. NL and SS would like to thank the Université Paris Diderot for hosting them for the period during which this research was conducted. The authors also wish to thank Sylvain Perifel, Yann Strozecki, and Meena Mahajan for useful discussions and feedback. In the previous version of the paper, Theorems 15 and 16 were stated for $d \leq n^{1/10}$. The authors wish to thank Chandan Saha for pointing out that these statements can be proved without any restriction on d .

REFERENCES

- [AV08] M. AGRAWAL AND V. VINAY, *Arithmetic circuits: A chasm at depth four*, in Proceedings of FOCS, 2008, pp. 67–75.

- [CLO97] D. A. COX, J. LITTLE, AND D. O'SHEA, *Ideals, Varieties, and Algorithms: An Introduction to Computational Algebraic Geometry and Commutative Algebra*, 2nd ed., Undergrad Texts Math., Springer, New York, 1997.
- [CM14] S. CHILLARA AND P. MUKHOPADHYAY, *Depth-4 lower bounds, determinantal complexity: A unified approach*, in Proceedings of STACS, 2014, pp. 239–250.
- [DMPY12] Z. DVIR, G. MALOD, S. PERIFEL, AND A. YEHUDAYOFF, *Separating multilinear branching programs and formulas*, in Proceedings of STOC, H. J. Karloff and T. Pitassi, eds., ACM, New York, 2012, pp. 615–624.
- [DP09] D. P. DUBHASHI AND A. PANCONESI, *Concentration of Measure for the Analysis of Randomized Algorithms*, Cambridge University Press, Cambridge, UK, 2009.
- [FSS84] M. L. FURST, J. B. SAXE, AND M. SIPSER, *Parity, circuits, and the polynomial-time hierarchy*, Math. Systems Theory, 17 (1984), pp. 13–27.
- [GKKS13] A. GUPTA, P. KAMATH, N. KAYAL, AND R. SAPTHARISHI, *Approaching the chasm at depth four*, in Proceedings of the Conference on Computational Complexity (CCC), 2013.
- [Gur10] V. GURUSWAMI, *Introduction to Coding Theory, Lecture 2: Gilbert-Varshamov Bound*, <http://www.cs.cmu.edu/~venkatg/teaching/codingtheory> (2010).
- [Hås87] J. HÅSTAD, *Computational Limitations of Small-Depth Circuits*, MIT Press, Cambridge, MA, 1987.
- [Kay12] N. KAYAL, *An exponential lower bound for the sum of powers of bounded degree polynomials*, Electron. Colloquium Comput. Complexity, 19 (2012), p. 81.
- [KLSS14a] N. KAYAL, N. LIMAYE, C. SAHA, AND S. SRINIVASAN, *An exponential lower bound for homogeneous depth four arithmetic formulas*, Electron. Colloquium Comput. Complexity, 21 (2014), p. 5.
- [KLSS14b] N. KAYAL, N. LIMAYE, C. SAHA, AND S. SRINIVASAN, *Super-polynomial lower bounds for depth-4 homogeneous arithmetic formulas*, in Proceedings of STOC, 2014, pp. 119–127.
- [Koi12] P. KOIRAN, *Arithmetic circuits: The chasm at depth four gets wider*, Theoret. Comput. Sci., 448 (2012), pp. 56–65.
- [KS13] M. KUMAR AND S. SARAF, *Lower bounds for depth 4 homogenous circuits with bounded top fanin*, Electron. Colloquium Comput. Complexity, 20 (2013), p. 68.
- [KS14a] M. KUMAR AND S. SARAF, *The limits of depth reduction for arithmetic formulas: It's all about the top fan-in*, in Proceedings of STOC, 2014, pp. 136–145.
- [KS14b] M. KUMAR AND S. SARAF, *On the power of homogeneous depth 4 arithmetic circuits*, Electron. Colloquium Comput. Complexity, 21 (2014), p. 45.
- [KSS14] N. KAYAL, C. SAHA, AND R. SAPTHARISHI, *A super-polynomial lower bound for regular arithmetic formulas*, in Proceedings of STOC, 2014, pp. 146–153.
- [MP08] G. MALOD AND N. PORTIER, *Characterizing valiant's algebraic complexity classes*, J. Complex., 24 (2008), pp. 16–38.
- [MV97] M. MAHAJAN AND V. VINAY, *Determinant: Combinatorics, algorithms, and complexity*, Chic. J. Theoret. Comput. Sci., 1997 (1997).
- [NW97] N. NISAN AND A. WIGDERSON, *Lower bounds on arithmetic circuits via partial derivatives*, Comput. Complexity, 6 (1997), pp. 217–234.
- [Raz06] R. RAZ, *Separation of multilinear circuit and formula size*, Theory Comput., 2 (2006), pp. 121–135.
- [Raz09] R. RAZ, *Multi-linear formulas for permanent and determinant are of super-polynomial size*, J. ACM, 56 (2009).
- [Raz10] R. RAZ, *Tensor-rank and lower bounds for arithmetic formulas*, in Proceedings of STOC, L. J. Schulman, ed., ACM, New York, 2010, pp. 659–666.
- [RY08] R. RAZ AND A. YEHUDAYOFF, *Balancing syntactically multilinear arithmetic circuits*, Comput. Complexity, 17 (2008), pp. 515–535.
- [RY09] R. RAZ AND A. YEHUDAYOFF, *Lower bounds and separations for constant depth multilinear circuits*, Comput. Complexity, 18 (2009), pp. 171–207.
- [SY10] A. SHPILKA AND A. YEHUDAYOFF, *Arithmetic circuits: A survey of recent results and open questions*, Found. Trends Theor. Comput. Sci., 5 (2010), pp. 207–388.
- [Tav13] S. TAVENAS, *Improved bounds for reduction to depth 4 and 3*, in Mathematical Foundations of Computer Science, Lecture Notes in Comput. Sci. 8087, Springer, New York, 2013, pp. 813–824.
- [Tod92] S. TODA, *Classes of arithmetic circuits capturing the complexity of computing the determinant*, IEICE Trans. Inform. Systems, E75-D (1992), pp. 116–124.

- [Val79] L. G. VALIANT, *Completeness classes in algebra*, in Proceedings of STOC '79, 1979, ACM, New York, pp. 249–261.
- [Val82] L. G. VALIANT, *Reducibility by algebraic projections*, in Logic and Algorithmic: An International Symposium Held in Honor of Ernst Specker, Monogr. Enseign. Math. 30, Enseignement Mathématique, Geneva, 1982, pp. 365–380.