# Automating Regression Verification

Dennis Felsing[†]
dennis.felsing@student.kit.edu

Sarah Grebing[†]
sarah.grebing@kit.edu

Vladimir Klebanov[†]
klebanov@kit.edu

Philipp Rümmer[‡]
philipp.ruemmer@it.uu.se

Mattias Ulbrich[†]
ulbrich@kit.edu

[†]Karlsruhe Institute of Technology, Germany
[‡]Uppsala University, Sweden

## ABSTRACT

Regression verification is an approach complementing regression testing with formal verification. The goal is to formally prove that two versions of a program behave either equally or differently in a precisely specified way. In this paper, we present a novel automatic approach for regression verification that reduces the equivalence of two related imperative integer programs to Horn constraints over uninterpreted predicates. Subsequently, state-of-the-art SMT solvers are used to solve the constraints. We have implemented the approach, and our experiments show non-trivial integer programs that can now be proved equivalent without further user input.

## Categories and Subject Descriptors

F.3.1 [**Logics and Meanings of Programs**]: Specifying and Verifying and Reasoning about Programs; D.2.4 [**Software Engineering**]: Software/Program Verification

## Keywords

Regression verification; program equivalence; invariant generation; formal methods

## 1. INTRODUCTION

One of the main concerns during software evolution is to prevent the introduction of unwanted behavior, commonly known as *regressions,* when implementing new features, fixing defects, or during optimization. Undetected regressions can have severe consequences and incur high cost, in particular in late stages of development, or in software that is already deployed. Currently, the main quality assurance measure during software evolution is *regression testing* [3]. Regression testing uses a carefully crafted test suite to check that a modified version of a program is equivalent to the original one in relevant behavioral aspects.

*Regression verification* is a complementary approach that attempts to achieve the same goals as regression testing with techniques from formal verification. This means establishing a formal proof of equivalence of two program versions. In its basic form, we are trying to prove that the two versions produce the same output for all inputs. In more sophisticated scenarios, we want to verify that the two versions are equivalent only on some inputs (conditional equivalence) or differ in a formally specified way (relational equivalence).

Regression verification is not intended to replace testing, as testing has unique capabilities. Tests can, for instance, validate non-functional aspects of software (e.g., performance) or its interactions with the underlying software (and even hardware) layers. On the other hand, regression verification—especially if automated—is an attractive additional instrument of software quality assurance. If successful, it offers guaranteed coverage, while not requiring additional expenses to develop and maintain a test suite.

At the same time, regression verification offers a more favorable pragmatics than standard verification of functional properties of individual programs. For regression verification, one does not need to write and maintain complex specifications (which can be a significant bottleneck in the verification process). Furthermore, given two program versions that are both complex but similar to each other, much less effort is required to prove their equivalence than to prove that they satisfy an—also complex—functional specification. The effort for proving equivalence mainly depends on the difference between the programs and not on their overall size and complexity. Regression verification can exploit the fact that modifications are often local and only affect a small portion of a program.

A number of approaches and tools for regression verification exist already (see Section 6), but the majority of them are not automatic and require the user to supply inductive invariants (e.g., [9, 22, 32]). We present an approach and a tool for *automatic* regression verification of imperative programs with integer variables. We use automatic invariant generation techniques to infer sufficiently strong *coupling predicates*[1] between programs—and thus prove behavior equivalence.

Our approach is targeted towards showing equivalence of programs with complex arithmetic and control flow. This

---

[1]A coupling predicate is an inductive two-program invariant that relates the two programs throughout their execution. We are typically interested in coupling predicates that imply result equality upon termination of both programs.

**Figure 1: Architecture of our approach**

```
int g(int n) {          int g(int n, int s) {
  int r = 0;              int r = 0;

  if (n <= 0) {           if (n <= 0) {
    r = 0;                  r = s;
  } else {                } else {
    r = g(n-1) + n;         r = g(n-1, n+s);
  }                       }
  return r;               return r;
}                       }
```

   (a) basic version $P_1$        (b) optimized version $P_2$

**Figure 2: Computing the $n$-th triangular number**

kind of programs is poorly supported by existing automatic approaches, as these either require static (i.e., known at compile time) control flow [31], employ coarse abstractions on program computations [17,31], or are overly restrictive (e.g., require small bounds on loops or that equivalent unbounded loops have equivalent bodies) [26].

Our method works well whenever sufficiently "simple" coupling predicates over linear arithmetic exist that prove program equivalence. This is often the case in practice, as we argue throughout the paper. In particular, in Section 5 we demonstrate the effectiveness of our technique using a collection of small but non-trivial benchmarks.

In detail, the contributions of this paper are:

- A method for automatic regression verification for programs employing complex arithmetic on integer variables

- As part of the above, a method for computing efficient verification conditions for program equivalence

- A tool implementing the approach (available at `http://formal.iti.kit.edu/improve/`).

The architecture of our approach is shown in Figure 1 and can be described as follows: a frontend translates the two programs into efficient logical verification conditions (VC) for program equivalence using the algorithm presented in Section 3. The translation is completely automatic; the user does not have to supply the coupling predicates, loop invariants, or function summaries. Instead, placeholders for these entities are inserted into the VC formulas. The produced VC are in Horn normal form and are passed to an SMT solver for Horn constraints (such as Z3 [23] or ELDARICA [27]), as presented in Section 4. The solver tries to find a solution for the placeholders that would make the VC true. If the solver succeeds in finding a solution and thus inferring, among other things, a coupling predicate, the programs are equivalent. Alternatively, the solver may show that no solution exists (i.e., disprove equivalence) or time out.

## 1.1 Illustration

EXAMPLE 1. *Consider the function* $g_1$ *in Figure 2(a).*[2] *The function recursively computes the sum of integers in the interval* $[1..n]$ *(also known as the* $n$-th *triangular number). The function* $g_2$ *in Figure 2(b) computes essentially the same result, but it has been optimized to employ tail recursion.*

---

[2]Our approach requires that the two programs which we prove equivalent have disjoint variable and function names. To distinguish equally named identifiers from the two programs, we add subscripts indicating the program to which they belong. We may also concurrently use the original identifiers without a subscript as long as the relation is clear from the context.

*The advantage of tail-recursive functions is that they can be executed without growing the stack. To enable this optimization, an accumulator parameter* s *has been added to the signature of* $g_2$ *for collecting and passing on intermediate results. As another consequence,* $g_1$ *performs summation from the end of the interval, while* $g_2$ *starts from the beginning.*

It is our goal to prove *automatically* using program verification technology that the two functions are equivalent in the sense that

$$g_1(n) = g_2(n, 0) \text{ for any } n. \tag{1}$$

The *conceptually* simplest way to achieve this goal is to infer and compare a complete logical specification of each function, such as $g_1(n) = \frac{n(n+1)}{2}$ for $n \geqslant 0$, and 0 otherwise. Yet, such a "brute-force" approach is clearly infeasible at the moment or in the foreseeable future.

Instead, we exploit the similarities between program versions and attempt to infer coupling predicates, i.e., a *relative* specification that only states how the executions of two versions relate to each other, but not what they compute in general. Of course, we demand that the coupling predicate is such that the two executions terminate in the same state, but in general, it will be a stronger assertion. To deal with unbounded recursion and loops, the predicate must be *inductive*: assuming that it holds at a key point of the computation (i.e., loop iteration, recursive call) should be sufficient to show that it also holds at the next key point. Conveniently, the complexity of such a relative specification is often proportional to the difference between the two program versions and not the program size.

Suppose, we want to calculate g(5). We call $g_1$(5) resp. $g_2$(5,0). The functions descend recursively, proceeding to compute $g_1$(4)+5 resp. $g_2$(4,5), and then $g_1$(3)+4+5 resp. $g_2$(3,9). At this point, one could suspect that at *every* recursion step

$$g_1(n) + s = g_2(n, s) . \tag{2}$$

A simple induction proof can establish that our suspicion is indeed correct: assuming this relation for the callees in both functions allows us to prove the relation for the callers. Thus, the formula (2) is a valid coupling predicate. Fortunately, (2) also implies the desired equivalence for the top-level call (1) with $s = 0$. Indeed, if one knows or guesses the formula (2), then the fact that it is a valid coupling predicate and that it implies equivalence can be proved automatically with existing verification technology (cf., e.g., [9, 22, 32]).

In this paper, we show that it is actually in many cases possible to automatically *infer* coupling predicates that imply program equivalence. The Horn encoding of the VCs for

the illustrative example is discussed in Section 4, and can be solved by ELDARICA [27] in roughly three seconds, inferring the coupling predicate $n_1 = n_2 \rightarrow r_1 + s_2 = r_2$, where $n_1$ is the argument of $\mathsf{g}_1$, $n_2$ and $s_2$ are the arguments of $\mathsf{g}_2$, and $r_i$ denote the respective return values. We note that the coupling predicate is linear even though the mathematical function computed by the two programs is non-linear.

## 2. PROGRAM EQUIVALENCE

This section introduces the considered programming language, and formalizes our notion of program equivalence in terms of Dijkstra's weakest preconditions [12]. The resulting program equivalence condition can be reduced to a program-free verification condition by applying reduction rules for weakest preconditions. A set of reduction rules optimized for equivalence proofs is defined in Section 3. Automation of the procedure is discussed in Section 4.

### The Programming Language.

We consider deterministic imperative programs with un-bounded integer variables (mathematical integers), written in ANSI C notation. Determinism means that program runs starting in the same state also terminate in the same state. Sequential real-world programs are deterministic, provided that all variables are initialized before they are used (which can be efficiently checked by a compiler). Furthermore, we require that all considered programs terminate for all inputs; this can be checked with one of the existing termination checkers for imperative programs, such as, e.g., [14, 15]. To simplify presentation, we assume that every function ends with a `return` statement, and that `return` is always the last statement in a function. Further program features, for instance heap or arrays, are discussed in Section 5.2, but not the main focus of the present paper.

We also assume that all programs have a distinguished function that is the entry point of the program. The entry point of the programs in our examples below is clear from the context.

### Syntactical Conventions.

For reasons of presentation, we require that the programs $P_1$ and $P_2$ checked for equivalence have disjoint sets of variables. To distinguish equally named variables from the two programs, we add subscripts indicating the program version (1 or 2) to which they belong. We also establish the syntactic convention that program inputs (i.e., formal function parameters) are designated as $\bar{i}_1$ resp. $\bar{i}_2$, returned result variables as $r_1$ resp. $r_2$, and the vectors of all variables occurring in the programs as $\bar{x}_1$ resp. $\bar{x}_2$.

### Background: Weakest Precondition Calculus.

Our reasoning about programs is formulated in terms of Dijkstra's weakest precondition calculus [12]. The *weakest precondition* predicate $wp(P, \varphi)$ denotes the weakest condition that needs to hold before an execution of statement list[3] $P$ such that the execution terminates and the postcondition $\varphi$ holds in the final state. The termination requirement is often considered optional. Relinquishing it, one obtains the *weakest liberal precondition* predicate $wlp(P, \varphi)$,

---

[3]To simplify presentation, we will use the terms "statement list" and "program" interchangeably. The exact relation will be clear from the context.

which only demands that $\varphi$ holds after the execution of $P$ *if* $P$ terminates. Thus, the formula $pre \rightarrow wlp(P, post)$ has the same intuitive meaning as the Floyd-Hoare triple $\{pre\}P\{post\}$.

A weakest precondition calculus is a set of rules which allow the resolution of $wp/wlp$ predicates into formulas in pure first-order logic. Figure 3(a) lists a calculus for the $wlp$ predicate for the considered programming language; the rules are standard, except that, for technical reasons, our calculus performs rewriting from the *beginning* of the statement list to its end, while a presentation with rules operating in the opposite direction is more customary. Reduction in forward direction is more convenient, however, for identifying structural similarity between the programs whose equivalence is verified. The calculus in Figure 3(a) is *complete* in the sense that every $wlp$-expression can be reduced to a pure first-order formula.

The rules (5), (6) and (7) allow the direct resolution of assignments, conditional statements and return statements (remember that the latter may only appear at the end of function bodies). The rule (8) for while loops is parametrized by a *loop invariant* $I(\bar{x}_1)$, a formula which needs to hold before the loop and must be preserved by the loop body under assumption of the loop condition. Likewise, the rule (9) for a (recursive) invocation $b = \mathsf{f}_1(\bar{a})$ of the function $\mathsf{f}_1$ is parametrized by a *function summary predicate* $S_{\mathsf{f}1}(\bar{a}, b)$ that relates the arguments $\bar{a}$ to the result value assigned to variable $b$. When the function summary $S_{\mathsf{f}1}$ is used as abstraction for the behavior of $\mathsf{f}_1$, the correctness of the summary has to be justified globally by an additional verification condition

$$wlp(P_1, S_{\mathsf{f}1}(\bar{i}_1, r_1)) \ , \qquad (3)$$

in which $P_1$ is the function body of $\mathsf{f}_1$.

The invariant rule (8) and the recursive invocation rule (9) may approximate loop or function behavior depending on the chosen invariant or function summary. In this case, the formula derived by applying the rules will still be a correct precondition, but not necessarily the weakest one. Even when approximating, finding suitable loop invariants and summaries is in general a difficult task.

### Stating Program Equivalence.

We consider two statement lists (usually the bodies of two functions) $P_1$ and $P_2$ *equivalent,* in writing

$$pre \rightarrow P_1 \simeq P_2 \ ,$$

when they behave equally (return the same value) for all inputs for which the precondition *pre* holds. We lift this notion to whole programs, by defining it as equivalence of the two program entry functions. The precondition *pre*, which can speak about variables from both $P_1$ and $P_2$, makes our notion of equivalence *conditional*. It is also possible to relax the equality between results to some other specified relation, yielding *relational equivalence*.

These notions can be formalized using the $wlp$ predicate introduced above. Since we assume that $P_1$ and $P_2$ have disjoint vocabulary, their code can simply be combined sequentially. We define:

$$pre \rightarrow P_1 \simeq P_2 :=$$
$$\forall \bar{i}_1, \bar{i}_2. \ (\bar{i}_1 = \bar{i}_2 \wedge pre \rightarrow wlp(P_1 \,;\, P_2, \ r_1 = r_2)) \ . \quad (4)$$

This kind of construction is known as *self-composition* [10, 11]. The weakest *liberal* precondition predicate has been used in this definition, since we deliberately abstract from termination issues in this paper.

# 3. EFFICIENT CONDITIONS FOR PROGRAM EQUIVALENCE

At this point, one could in theory directly resolve the *wlp* predicate in (4) by applying the rules from Figure 3(a) to obtain a first-order verification condition for equivalence of $P_1$ and $P_2$. However, the sequential composition of the two programs would require that they be analyzed individually without exploiting structural similarities between them.

Instead, we devise additional rules for the *wlp* predicate for the case that the program code given as the first argument is composed of two pieces with disjoint vocabulary. The disjointness allows us to use more rules than would be sound otherwise, as the statements with disjoint data cannot interfere with each other. The additional rules make use of two forms of *coupling predicates* that relate the states of the compared programs: *mutual invariants C,* which describe reachable states of two loops in the respective programs, iterating in a synchronized manner, and *mutual function summaries R* that express the relative behavior of two functions in the programs. The result of applying the new rules is a much more efficient first-order verification condition for equivalence.

In Figure 3(b), we present the additional rules. To make the composition of two programs with disjoint vocabulary explicit, we use ⨟ instead of ; as separator between them. Semantically, both are equivalent. In particular, it is always sound to replace $P_1 ⨟ P_2$ with $P_1 ; P_2$. Conversely, it is sound to replace $P_1 ; P_2$ with $P_1 ⨟ P_2$ whenever $P_1$ and $P_2$ have disjoint vocabulary.

Rule (10) allows us to swap the two programs, thus enabling resolution of statements from both programs in an alternating fashion. The rule is sound since the statements of the two programs cannot possibly interfere; they have no common variables to refer to.

Together with the rules (5) and (6) of Figure 3(a), the swap rule allows us to resolve all statements but loops or recursion. These are the difficult cases since they require finding a suitable loop invariant or a function summary. The next two sections therefore introduce efficient rules for pairwise loops and function calls. The *wlp* calculus can isolate the relevant loop pairs from within their programs even if they are embedded into enclosing conditionals or loops.

PROPOSITION 1    (SOUNDNESS AND COMPLETENESS). *Let $\Phi$ be a purely first-order formula derived from the condition $wlp(P_1 ⨟ P_2, \varphi)$ by rules from Figure 3(a) and (b). If the program $P_1 ; P_2$ is started in a state satisfying the precondition $\Phi$, and terminates, then $\varphi$ holds in its final state. Furthermore, it is possible to choose suitable mutual invariants and summaries such that the derived formula is the weakest such precondition.*

We give a justification for the validity of the proposition in the following.

## 3.1 While loops

We first consider equivalence of programs with loops, but without recursive function invocations. The loop rule for program equivalence is different from the rules discussed so far in that it talks about both programs at the same time and actually connects the two:

$$wlp(\texttt{while}(\psi_1) \ B_1 ; P_1 ⨟ \texttt{while}(\psi_2) \ B_2 ; P_2, \varphi) \rightsquigarrow$$
$$C(\bar{x}_1, \bar{x}_2) \wedge \forall \bar{x}_1, \bar{x}_2. \big($$
$$(C(\bar{x}_1, \bar{x}_2) \wedge \psi_1 \wedge \psi_2 \rightarrow wlp(B_1 ⨟ B_2, C(\bar{x}_1, \bar{x}_2)) \wedge$$
$$(C(\bar{x}_1, \bar{x}_2) \wedge \neg\psi_1 \wedge \psi_2 \rightarrow wlp(B_2, C(\bar{x}_1, \bar{x}_2))) \wedge$$
$$(C(\bar{x}_1, \bar{x}_2) \wedge \psi_1 \wedge \neg\psi_2 \rightarrow wlp(B_1, C(\bar{x}_1, \bar{x}_2))) \wedge$$
$$(C(\bar{x}_1, \bar{x}_2) \wedge \neg\psi_1 \wedge \neg\psi_2 \rightarrow wlp(P_1 ⨟ P_2, \psi))\big) \ .$$

The rule is parametrized by the mutual loop invariant $C(\bar{x}_1, \bar{x}_2)$, which is part of the coupling predicate that we are interested in. Unlike the invariant rule for a single program (8), which has two cases (loop condition holds or does not hold), this rule has four possible evaluations of the two loop conditions to consider.

For the justification of this rule, let us look at a particular reordering of the statements in the two loops. The central idea behind the rearrangement is that the two loops can be subject to a loop fusion resulting in the following program equivalence:

$$\texttt{while}(\psi_1) \ B_1 ⨟ \texttt{while}(\psi_2) \ B_2 \quad \simeq$$
$$\texttt{while} \ (\psi_1 || \psi_2) \ \{ \ \texttt{if}(\psi_1) \ B_1 \ ; \ \texttt{if}(\psi_2) \ B_2 \ \} \quad . \tag{14}$$

Why is the single loop equivalent to the sequential execution of the separate loops? Running the two loops sequentially results in running the sequence of statements

$$(\underbrace{B_1, B_1, \dots, B_1}_{n \text{ times}}, \ \underbrace{B_2, B_2, \dots, B_2}_{m \text{ times}}) \ ,$$

in which the first loop body $B_1$ is repeated $n$ times followed by $m$ repetitions of the second body $B_2$. Let w.l.o.g. the second loop be executed more often than the first in this schematic example (i.e., $m > n$). Due to disjoint vocabulary, loop body executions from different programs may be swapped. The run may hence be rearranged to

$$(\underbrace{B_1, B_2, \ B_1, B_2, \ \dots, B_1, B_2}_{n \text{ times}}, \ \underbrace{B_2, \dots, B_2}_{m-n \text{ times}}) \tag{15}$$

without changing the semantics. One can make out $m$ iterations now, of which the first $n$ execute both loop bodies $B_1, B_2$, while the remaining $m - n$ rounds only execute the second loop body $B_2$. The sequence (15) is a run for the fused loop from (14). It is the additional `if`-statements that ensure that bodies are only executed as often as they would be executed in a sequential execution. The disjunction in the guard ensures that the fused loop is iterated precisely as often as the maximum iterations of the individual loops.

Applying the traditional while *wlp* rule (8) to the fused loop from (14), has the same effect as applying the two-program rule (12). Since the traditional *wlp* calculus is sound and complete, our extension thus inherits these properties.

Mutual loop invariants are simpler than full functional invariants if the two programs are related. To show equivalence between a while loop and (a copy of) itself, for instance, the simple invariant $\bar{x}_1 = \bar{x}_2$ is sufficient regardless of what the loop computes.

$$wlp(\mathtt{x = t} \,; P, \, \varphi) \rightsquigarrow \text{let } x = t \text{ in } wlp(P, \, \varphi) \tag{5}$$

$$wlp(\mathtt{if(}\psi\mathtt{)} \ T \ \mathtt{else} \ E \,; P, \, \varphi) \rightsquigarrow \text{if } \psi \text{ then } wlp(T \,; P, \, \varphi) \text{ else } wlp(E \,; P, \, \varphi) \tag{6}$$

$$wlp(\mathtt{return} \ r, \, \varphi) \rightsquigarrow \varphi \tag{7}$$

$$wlp(\mathtt{while(}\psi\mathtt{)} \ B \,; P, \, \varphi) \rightsquigarrow I(\bar{x}_1) \wedge \forall \bar{x}_1.( \quad (I(\bar{x}_1) \wedge \psi \rightarrow wlp(B, \, I(\bar{x}_1))) \tag{8}$$
$$\wedge \, (I(\bar{x}_1) \wedge \neg\psi \rightarrow wlp(P, \, \varphi)))$$

$$wlp(r = \mathtt{f_1}(\bar{t}) \,; P, \, \varphi) \rightsquigarrow \forall r. \ S_{\mathtt{f1}}(\bar{t}, s) \rightarrow wlp(P, \, \varphi) \tag{9}$$

**(a) Conventional $wlp$ calculus rules**

$$wlp(P_1 \,\mathring{;}\, P_2, \, \varphi) \rightsquigarrow wlp(P_2 \,\mathring{;}\, P_1, \, \varphi) \tag{10}$$

$$wlp(\mathtt{return} \ r \,\mathring{;}\, P_2, \, \varphi) \rightsquigarrow wlp(P_2, \, \varphi) \tag{11}$$

$$wlp(\mathtt{while(}\psi_1\mathtt{)} \ B_1 \,; P_1 \,\mathring{;}\, \mathtt{while(}\psi_2\mathtt{)} \ B_2 \,; P_2, \, \varphi) \rightsquigarrow C(\bar{x}_1, \bar{x}_2) \tag{12}$$
$$\wedge \, \forall \bar{x}_1, \bar{x}_2.( \quad (C(\bar{x}_1, \bar{x}_2) \wedge \psi_1 \wedge \psi_2 \rightarrow wlp(B_1 \,\mathring{;}\, B_2, \, C(\bar{x}_1, \bar{x}_2))$$
$$\wedge \, (C(\bar{x}_1, \bar{x}_2) \wedge \neg\psi_1 \wedge \psi_2 \rightarrow wlp(B_2, \, C(\bar{x}_1, \bar{x}_2)))$$
$$\wedge \, (C(\bar{x}_1, \bar{x}_2) \wedge \psi_1 \wedge \neg\psi_2 \rightarrow wlp(B_1, \, C(\bar{x}_1, \bar{x}_2)))$$
$$\wedge \, (C(\bar{x}_1, \bar{x}_2) \wedge \neg\psi_1 \wedge \neg\psi_2 \rightarrow wlp(P_1 \,\mathring{;}\, P_2, \, \psi)))$$

$$wlp(r_1 = \mathtt{f_1}(\bar{t}_1) \,; P_1 \,\mathring{;}\, r_2 = \mathtt{f_2}(\bar{t}_2) \,; P_2, \, \varphi) \rightsquigarrow \forall r_1, r_2. \ R_{\mathtt{f1/f2}}(\bar{t}_1, r_1, \bar{t}_2, r_2) \rightarrow wlp(P_1 \,\mathring{;}\, P_2, \, \varphi) \tag{13}$$

**(b) Additional $wlp$ calculus rules for independent programs**

**Figure 3: Weakest precondition calculus**

## 3.2 Recursion

We now consider programs that have (recursive) function calls but no loops. Recursive calls of related functions in both programs can be abstracted by a single predicate, a *mutual function summary* (a term originated in [21]), that describes the relation between the arguments and result values of both invocations simultaneously and in relation to one another. The calculus rule to handle simultaneous function invocations is

$$wlp(r_1 = \mathtt{f_1}(\bar{t}_1) \,; P_1 \,\mathring{;}\, r_2 = \mathtt{f_2}(\bar{t}_2) \,; P_2, \, \varphi) \rightsquigarrow$$
$$\forall r_1, r_2. \ R_{\mathtt{f1/f2}}(\bar{t}_1, r_1, \bar{t}_2, r_2) \rightarrow wlp(P_1 \,\mathring{;}\, P_2, \, \varphi) \ .$$

The rule is parametrized by the mutual function summary $R_{\mathtt{f1/f2}}(\bar{t}_1, r_1, \bar{t}_2, r_2)$.

Abstracting function invocations with a mutual summary requires a (global) justification that the summary is a faithful abstraction, and we need to add the proof obligation

$$\forall \bar{i}_1, \bar{i}_2. \ wlp(P_1 \,\mathring{;}\, P_2, \, R_{\mathtt{f1/f2}}(\bar{i}_1, r_1, \bar{i}_2, r_2)) \tag{16}$$

to the verification conditions of equivalence. Here, $P_1$ and $P_2$ are the statement lists from the function bodies of the invoked functions $\mathtt{f_1}$ and $\mathtt{f_2}$.

The justification of rule (13) is as follows. Due to the disjointness of the program vocabulary, the statements in the rule can be reordered:

$$r_1 = \mathtt{f_1}(\bar{t}_1) \,; P_1 \,\mathring{;}\, r_2 = \mathtt{f_2}(\bar{t}_2) \,; P_2 \quad \simeq$$
$$r_1 = \mathtt{f_1}(\bar{t}_1) \,; r_2 = \mathtt{f_2}(\bar{t}_2) \,; P_1 \,; P_2 \ .$$

Condition (16) guarantees that $R_{\mathtt{f1/f2}}(\bar{t}_1, r_1, \bar{t}_2, r_2)$ is a faithful abstraction of $P_1 \,; P_2$. Just as in the single-program case, it is thus sound to overapproximate the two recursive invocations with the mutual function summary.

As with mutual loop invariants, mutual function summaries are simpler than individual function summaries if

the two programs are related. In case a recursive function is verified against (a copy of) itself, the simple mutual function summary $\bar{i}_1 = \bar{i}_2 \rightarrow r_1 = r_2$ can be used.

Note that the same mutual summary $R_{\mathtt{f1/f2}}(\bar{t}_1, r_1, \bar{t}_2, r_2)$ is used for every occurrence of the pair $\mathtt{f_1/f_2}$ of functions; this is in contrast to the coupling invariant rule (12) for loops, where it is possible to choose different mutual invariants $C(\bar{x}_1, \bar{x}_2)$ for every application. While our calculus could in principle be extended to support multiple mutual summaries per $\mathtt{f_1/f_2}$ pair, the use of only a single such summary minimizes the number of required proof obligations (16).

## 4. AUTOMATIC EQUIVALENCE PROOFS

The application of the $wlp$ rules in Figure 3 requires knowledge of specific predicates, namely loop invariants $I(\bar{i}_1, \bar{x}_1)$ in rule (8), mutual loop invariants $C(\bar{x}_1, \bar{x}_2)$ in (12), function summaries $S_{\mathtt{f1}}(\bar{t}, s)$ in (9), and mutual function summaries $R_{\mathtt{f1/f2}}(\bar{t}_1, r_1, \bar{t}_2, r_2)$ in (13). Together, those formulas represent the coupling predicate that witnesses program equivalence. Derivation of summaries and invariants is in general a complicated process and can require creativity and manual intervention. Thanks to the specialized $wlp$-rules for program equivalence, however, it is often possible to carry out equivalence proofs with comparatively simple predicates. In Section 1.1, for instance, it is possible to show the equivalence of programs computing non-linear functions with the help of just linear predicates; our experiments (Section 5) show that such simple predicates are sufficient for a wide range of realistic cases from regression verification.

We leverage recent methods for solving fixed-point constraints in order to compute required predicates fully automatically [19, 23, 27]. Such methods are in principle incomplete, but they are effective for deriving predicates in practical cases arising from equivalence proofs.

## Recursive Horn Clauses.

In order to derive invariants and coupling predicates, verification conditions are represented in form of *Horn constraints* over (uninterpreted) relation symbols, including $I$, $C, S_{\mathtt{f}1}, R_{\mathtt{f}}$, and then solved with the help of model checking techniques like predicate abstraction and Craig interpolation. More generally, we fix a set $\mathcal{R}$ of uninterpreted fixed-arity *relation symbols,* and consider *Horn clauses* of the form $H \leftarrow \varphi \wedge B_1 \wedge \cdots \wedge B_n$, where:

- $\varphi$ is a constraint over variables occurring in the clause; in our experiments, $\varphi$ is always a formula in quantifier-free Presburger arithmetic, but extension to other theories (e.g., arrays) is possible;

- each $B_i$ is an application $p(t_1, \ldots, t_k)$ of a relation symbol $p \in \mathcal{R}$ to first-order terms;

- $H$ is similarly either an application $p(t_1, \ldots, t_k)$ of a symbol $p \in \mathcal{R}$ to first-order terms, or *false*.

$H$ is called the *head* of the clause, $\varphi \wedge B_1 \wedge \cdots \wedge B_n$ the *body*. In case $\varphi = true$, we usually leave out $\varphi$ and just write $H \leftarrow B_1 \wedge \cdots \wedge B_n$. First-order variables in a clause are considered implicitly universally quantified; relation symbols represent set-theoretic relations over the universe of a first-order semantic structure. A set of Horn clauses $HC$ over predicates $\mathcal{R}$ is called *solvable* if there is an interpretation of the predicates $\mathcal{R}$ as set-theoretic relations such the universal closure of every clause $h \in HC$ holds.

EXAMPLE 2 (EXAMPLE 1 CONTINUED). *Figure 4 shows the equivalence VC for the programs from the illustration example (Figure 2) as Horn clauses. Here, R is the uninterpreted predicate symbol (placeholder) for the coupling predicate (mutual function summary) of $g_1$ and $g_2$ introduced by application of rule (13). The uninterpreted predicates $S_{g1}$ and $S_{g2}$ are the function summaries for the respective individual functions and are introduced by (9). Clauses with head false result from equivalence proof obligations (4), whereas the clauses with a head different from false are due to justification conditions (3) and (16).*

## Verification Conditions as Horn Clauses.

For the encoding of verification conditions as Horn clauses, we assume that the set $\mathcal{R}$ contains symbols that can act as summaries for individual functions and function pairs (of appropriate arity), as well as relation symbols $I_1, I_2, I_3, \ldots$ and $C_1, C_2, C_3, \ldots$ to represent loop invariants:

$$\mathcal{R} = \{S_{\mathtt{f}}, R_{\mathtt{f}1/\mathtt{f}2} \mid \mathtt{f}, \mathtt{f}_1, \mathtt{f}_2 \text{ functions}\} \cup$$
$$\{I_1, I_2, I_3, \ldots, C_1, C_2, C_3, \ldots\} .$$

We then consider the conjunction of the equivalence statement $pre \rightarrow P_1 \simeq P_2$ and the correctness of the summaries for all functions reachable from $P_1$ or $P_2$:

$$\forall \bar{i}_1, \bar{i}_2. \, (\bar{i}_1 = \bar{i}_2 \wedge pre \rightarrow wlp(P_1 \,;;\, P_2, \, r_1 = r_2))$$

$$\wedge \bigwedge_{\substack{\mathtt{f} \text{ a} \\ \text{function}}} \forall \bar{i}_{\mathtt{f}}. \, wlp(P_{\mathtt{f}}, \, S_{\mathtt{f}}(\bar{i}_{\mathtt{f}}, r_{\mathtt{f}}))$$

$$\wedge \bigwedge_{\substack{\mathtt{f}1, \mathtt{f}2 \\ \text{functions}}} \forall \bar{i}_{\mathtt{f}1}, \bar{i}_{\mathtt{f}2}. \, wlp(P_{\mathtt{f}1} \,;;\, P_{\mathtt{f}2}, \, R_{\mathtt{f}1/\mathtt{f}2}(\bar{i}_{\mathtt{f}1}, r_{\mathtt{f}1}, \bar{i}_{\mathtt{f}2}, r_{\mathtt{f}2})) .$$

$$(17)$$

$$false \leftarrow \begin{array}{l} n_1 \leqslant 0 \wedge n_2 \leqslant 0 \wedge \\ n_1 = n_2 \wedge s_2 = 0 \wedge 0 \neq s_2 \end{array}$$

$$false \leftarrow \begin{array}{l} n_1 > 0 \wedge n_2 > 0 \wedge \\ n_1 = n_2 \wedge s_2 = 0 \wedge r_1 \neq r_2 \wedge \\ R(n_1 - 1, r_1, n_2 - 1, n_2 + s_2, r_2) \end{array}$$

$$false \leftarrow \begin{array}{l} n_1 > 0 \wedge n_2 \leqslant 0 \wedge \\ n_1 = n_2 \wedge s_2 = 0 \wedge r_1 \neq s_2 \wedge \\ S_{g1}(n_2 - 1, r_1) \end{array}$$

$$false \leftarrow \begin{array}{l} n_1 \leqslant 0 \wedge n_2 > 0 \wedge \\ n_1 = n_2 \wedge s_2 = 0 \wedge 0 \neq r_2 \wedge \\ S_{g2}(n_2 - 1, n_2 + s_2, r_2) \end{array}$$

$$S_{g1}(n_1, 0) \leftarrow n_1 \leqslant 0$$
$$S_{g1}(n_1, r_1 + n_1) \leftarrow n_1 > 0 \wedge S_{g2}(n_1 - 1, r_1)$$
$$S_{g2}(n_2, s_2, s_2) \leftarrow n_2 \leqslant 0$$
$$S_{g2}(n_2, s_2, r_2) \leftarrow n_2 > 0 \wedge S_{g2}(n_2 - 1, n_2 + s_2, r_2)$$

$$R(n_1, 0, n_2, s_2, s_2) \leftarrow n_1 \leqslant 0 \wedge n_2 \leqslant 0$$

$$R(n_1, r_1 + n_1, n_2, s_2, r_2) \leftarrow \begin{array}{l} n_1 > 0 \wedge n_2 > 0 \wedge \\ R(n_1 - 1, r_1, n_2 - 1, n_2 + s_2, r_2) \end{array}$$

**Figure 4: Program equivalence VC as Horn clauses**

Intuitively, any valuation of the relation symbols $\mathcal{R}$ that makes (17) valid is a witness for the equivalence of $P_1$ and $P_2$, assuming *pre* holds initially.

The next step is the elimination of the *wlp* transformer from (17), by means of exhaustive application of the rules in Figure 3. When applying (9) or (13) to replace function calls $\mathtt{f}, \mathtt{f}_1, \mathtt{f}_2$ with the corresponding summary, the relation symbol $S_{\mathtt{f}}$ or $R_{\mathtt{f}1/\mathtt{f}2}$ is inserted in the formula; similarly, when applying the loop rules (8) or (12), a fresh relation symbol $I_k$ or $C_k$ is introduced. We explain one possible strategy for applying the reduction rules below. Once application of the *wlp* rules to (17) has terminated, Horn clauses can be extracted from the reduct $VC$ (a pure first-order formula), thanks to the following lemma:

LEMMA 1. *Suppose VC resulted from exhaustive application of rules in Figure 3 to (17). Then the clause normal form $VC_H$ of VC is Horn.*

The clause normal form $VC_H$ is derived by first distributing negations (*negation normal form*) in $VC$, then pulling out all universal quantifiers $\forall$ (*prenex normal form*), and finally transforming to *conjunctive normal form* [20]. To see that the clause normal form $VC_H$ is Horn, observe that (17) only contains *wlp* in positive positions, and that any two positive occurrences of relation symbols are separated by a conjunction; both properties are preserved by application of *wlp* rules, and entail that each clause in the clause normal form contains at most one positive relation symbol.

## Reduction Strategy.

In some situations, it can happen that more than one rule in Figure 3 is applicable to a *wlp* expression, so that in principle more than one verification condition $VC$ can be derived from (17). Different VCs can represent different ways to match up loops and corresponding function calls in the two programs checked for equivalence, and can therefore make

the subsequent solving of the Horn constraints $VC_H$ more or less difficult.

At the moment, we resolve such choice points using a greedy application strategy:

1. as long as possible, rules (5), (6), (7), (11) to eliminate assignments, conditionals, and return statements of the individual programs, possibly together with (10) to change the order of programs.

2. if no further rules from point 1 are applicable, try to use (12) or (13) for synchronous handling of loops or function calls; if this succeeds, go back to 1.

3. if no further rules from point 1 or 2 are applicable, use (8) or (9) to eliminate single loops or function calls; if this succeeds, go back to 1.

This strategy matches up loops and function calls in the order in which they occur in the considered programs. The strategy produces good results in our experiments, but can clearly be refined to take more sophisticated similarity measures into account. Further discussion is given in Section 5.

### *Solving Horn Clauses.*

A number of algorithms exist to solve the Horn clauses $VC_H$, including *predicate abstraction* [19, 27] and *property-directed reachability* (PDR, also known as IC3) implemented in Z3 [23]. The procedures attempt to construct a symbolic solution of $VC_H$ in a decidable logic, for instance in (quantifier-free) Presburger arithmetic; such a solution maps every $n$-ary relation symbol in $\mathcal{R}$ to a symbolic predicate over $n$ variables.

EXAMPLE 3 (EXAMPLE 2 CONTINUED). *For the clauses in Example 2, the following predicates are found for the uninterpreted symbols:*

$$R(n_1, r_1, n_2, s_2, r_2) \;\mapsto\; (n_1 = n_2 \rightarrow r_1 + s_2 = r_2)$$
$$S_{\mathtt{g1}}(n_1, r_1) \;\mapsto\; true$$
$$S_{\mathtt{g2}}(n_2, s_2, r_2) \;\mapsto\; true$$

*which is the solution already discussed in Section 1.1. The function summaries $S_{\mathtt{g1}}$ and $S_{\mathtt{g2}}$ can be trivially chosen to be true since the Horn clauses in which they occur in the body are already valid without them.*

In general, if it terminates, a Horn solver will produce one of two possible results: (i) a symbolic *solution* of the processed Horn clauses, or (ii) a concrete *counterexample tree* that witnesses that no solution of the Horn clauses exists. The leaves in a counterexample tree correspond to entry clauses (clauses without relation symbols in the body), the root of the tree to an assertion clause with head *false*; the counterexample shows that every attempt to satisfy the Horn clauses has to lead to one of the assertion clauses being violated.[4] Through additional bookkeeping and labeling, counterexamples can be translated back to runs of the programs $P_1, P_2$ that are checked for equivalence; the counterexample specifies the path taken through each program, as well as the values of all program variables.

---

[4]Due to reasons of computability, sets of Horn clauses exist for which neither (i) nor (ii) can be returned: those are clauses that are solvable in a set-theoretic sense, but no solution can be expressed in the decidable language used for predicates. In such cases, usually non-termination occurs.

We summarize by stating the correctness of our procedure. It is important to note that the procedure is correct independently of the order in which *wlp* rules are applied for translating (17) to $VC_H$; in particular, counterexamples are always genuine, and point to an actual case of non-equivalence. Good strategies when applying the rules can, however, improve efficiency and prevent non-termination of the Horn solver.

THEOREM 1 (CORRECTNESS). *If a Horn solver applied to $VC_H$ terminates, then one of the following holds:*

- *a solution is found for $VC_H$, and in this case the considered equivalence pre $\rightarrow P_1 \simeq P_2$ holds;*

- *a counterexample is found, and the programs are not equivalent.*

## 5. IMPLEMENTATION AND EXPERIMENTS

### *Implementation.*

We have implemented our approach for a language close to a subset of ANSI C in a tool named RÊVE. Program data is limited to local variables and function parameters of type `int`, which is interpreted as unbounded (i.e., mathematical) integers. Bounded integers can be simulated by instrumenting programs with modulo operations, at the cost of increased reasoning complexity. Supported control structures are if-then-else and while statements, function calls and returns. For simplicity, the return statement must always be the last statement of a function and must return a local variable. Recursive function calls may not occur within the conditions of if or while statements. Checking conditional and relational equivalence of programs is supported.

The tool (i.e., the *wlp* calculus) is implemented in Standard ML. As Horn constraint solvers we used Z3 (unstable branch as of 2013-11-27) and ELDARICA (as of 2014-04-16).

### *Experiments.*

We have evaluated the effectiveness and performance of our tool on a collection of benchmarks. The benchmarks vary in size from 16–53 lines of code (for both programs together) and are available with the tool at the URL given in the introduction. Benchmark results are summarized in Table 1. We also give results from the only automatic tool that is directly comparable to ours (due to scope, cf. Section 6), the Regression Verification Tool (RVT) by Strichman and Godlin [17].

The programs in the first group in Table 1 are recursive, while the ones in the second group contain loops. Benchmarks where the two programs were not equivalent are in the third group, and their names end with a bang (`!`). All other benchmarks contain equivalent programs; the ✗ outcome is in this case a false negative.

Benchmarks `limit1` to `limit3` were given by Strichman and Godlin as beyond the limits of their approach to regression verification. Benchmarks `barthe2-big` and `barthe2-big2` embed the benchmark `barthe2` into a larger program that is syntactically identical in both versions. We could not prove equivalent the `ackermann` benchmark, as the result of a recursive function call is used as the argument to another recursive function call. Furthermore, we originally could not

prove the `limit1` benchmark, as two steps of the first loop are equivalent to one step of the second loop, an issue that we solve in the next section and illustrate with the larger `digits10` benchmark.

The `triangular-mod` benchmark corresponds to the illustrating example instrumented with modulo operations to simulate integer overflow.

As far as we are aware, RVT does not supply additional information to assist the user in case of a failed proof attempt. While, in theory, the model checker underlying RVT produces a counterexample, such a counterexample can be spurious due to the fixed abstraction employed. The ELDARICA solver that we use, in contrast, returns a genuine counterexample for many failed proofs (cf. Section 4). We found these counterexamples useful in diagnosing problems with the programs, even though we currently do not translate these counterexamples into source code terms.

## 5.1  An Example for Loop Equivalence

We consider a real-world example from [1]. The program $P_1$ in Figure 5(a) computes the number of digits in the decimal expansion of n through a series of integer divisions by 10. The program $P_2$ in Figure 5(c) computes the same result but (asymptotically) about seven times faster. This speedup is accomplished by reducing the strength of operations. The loop has been unrolled four times[5] *and* the majority of divisions have been replaced by pure comparisons.

Unsurprisingly, $P_1$ and $P_2$ cannot be proved equivalent automatically. To do so, the tool would in the least need to figure out the (very complex) relation between one iteration of the loop in $P_1$ and four iterations of the same loop. To overcome this barrier, the software engineer needs to supply to the tool the knowledge that an unrolling transformation took place. At the moment, we achieve this transfer by manually carrying out the unrolling on $P_1$ and producing the intermediate program $P_1'$ shown in Figure 5(b). We then prove automatically that $P_1'$ and $P_2$ are equivalent. Note that $P_1'$ is still significantly different from $P_2$, as unrolling is not the only optimization that has been performed originally. The program $P_1'$ still performs four times as many divisions as $P_2$. The if-conditions directly follow the divisions and depend on them, which slows the program down, while the four if-conditions in $P_2$ are all dependent on the same division result.

After 11.3 seconds, RÊVE with ELDARICA succeeds in proving equivalence with the following automatically inferred coupling predicate:

$$(b_2 = 1 \wedge r_1 = r_2 \wedge 10n_1 \leqslant n_2 \wedge n_2 \leqslant 10n_1 + 9)$$
$$\vee (b_2 = 0 \wedge r_1 = v_2 \wedge n_2 \geqslant 10n_1 \wedge n_1 \leqslant 0) \ .$$

Here, $n_1$ and $r_1$ denote the variables of $P_1'$, and $n_2$, $b_2$, $r_2$ the variables of $P_2$. The variable $b_2$ indicates whether the loop will ($b_2 = 1$) or will not ($b_2 = 0$) be executed once more. The coupling predicate is hence a disjunction over these two cases: While the loop is iterated, $r_1$ and $r_2$ hold the same value and $n_1$ is one division by 10 ahead of $n_2$, i.e., $n_1 = n_2$ DIV 10. Exactly this fact is expressed by the linear constraint $10n_1 \leqslant n_2 \wedge n_2 \leqslant 10n_1 + 9$. When the loop of $P_2$

---

[5]Loop unrolling is a simple transformation, in which the loop body is replicated within the loop and guarded by the loop guard. This transformation preserves the semantics of the program.

Table 1: Benchmark results

| Benchmark | LOC | Run time (seconds) RVT | Run time (seconds) RÊVE+ Z3 | Run time (seconds) RÊVE+ ELDARICA | Source |
|---|---|---|---|---|---|
| **Recursion** | | | | | |
| ackermann | 30 | 0.8 | − | − | [17] |
| mccarthy91 | 22 | 1.1 | 1.8 | 1.7 | [17] |
| limit1 | 22 | ✗ | − | − | [17] |
| limit2 | 22 | ✗ | − | 5.2 | [17] |
| limit3 | 24 | ✗ | − | 4.5 | [17] |
| add-horn | 26 | ✗ | − | 4.3 | |
| triangular | 23 | ✗ | − | 3.3 | |
| triangular-mod | 53 | ✗ | − | − | |
| inlining | 20 | ✗ | − | 5.7 | |
| **Loops** | | | | | |
| simple-loop | 16 | 0.8 | 0.1 | 5.3 | |
| loop | 22 | ✗ | − | 2.8 | |
| loop2 | 22 | ✗ | − | 3.2 | |
| loop3 | 28 | ✗ | − | 5.4 | |
| loop4 | 22 | ✗ | − | 27.4 | |
| loop5 | 22 | ✗ | − | 26.4 | |
| while-if | 22 | ✗ | − | 3.8 | |
| digits10 | 32 | ✗ | − | 11.3 | [1] |
| barthe | 28 | ✗ | − | 4.8 | [9] |
| barthe2 | 22 | 0.5 | 10.2 | 3.9 | [9] |
| barthe2-big | 32 | 1.6 | − | 5.7 | |
| barthe2-big2 | 42 | 1.7 | − | 8.0 | |
| bug15 | 26 | 1.0 | 0.1 | 1.8 | [17] |
| nested-while | 28 | 1.5 | − | 5.2 | [17] |
| **Not equivalent** | | | | | |
| ackermann! | 30 | ✗ | 0.1 | 4.9 | |
| limit1! | 22 | ✗ | 0.0 | 1.4 | |
| limit2! | 25 | ✗ | 0.7 | 10.9 | |
| add-horn! | 28 | ✗ | 0.1 | 1.9 | |
| triangular-mod! | 49 | ✗ | 0.5 | 28.1 | |
| inlining! | 20 | ✗ | 0.1 | 2.7 | |
| loop5! | 22 | ✗ | 0.0 | 2.2 | |
| barthe! | 31 | ✗ | 1.8 | 21.9 | |
| nested-while! | 28 | ✗ | 0.1 | 5.2 | |

LOC=non-empty, non-comment lines of code in both programs together. Dash (−) denotes timeout at 600 seconds, cross (✗) denotes that the tool terminates but cannot prove equivalence. All times have been measured on a 2.5 GHz Intel Core2 Quad machine, using only one core.

has finished, its negated loop guard $n_1 \leqslant 0$ holds and the final results are stored in $r_1$ and $v_2$.

## 5.2  Discussion

Our method exploits structural similarities between the compared programs, and can generally be expected to perform well when applied to programs with a high degree of similarity. In other situations, for instance when exchanging complete algorithms (e.g., replacing a bubble sort procedure

```
int f(int n) {
  int r = 1;
  n = n/10;

  while (n > 0) {
    r++;
    n = n / 10;
  }
  return r;
}
```

```
int f(int n) {
  int r = 1;
  n = n/10;

  while (n > 0) {
    r++;
    n = n / 10;
    if (n > 0) {
      r++;
      n = n / 10;
      if (n > 0) {
        r++;
        n = n / 10;
        if (n > 0) {
          r++;
          n = n / 10;
        }
      }
    }
  }
  return r;
}
```

```
int f(int n) {
  int r = 1;
  int b = 1;
  int v = -1;

  while (b != 0) {
    if      (n<   10) { v = r;   b = 0; }
    else if (n<  100) { v = r+1; b = 0; }
    else if (n< 1000) { v = r+2; b = 0; }
    else if (n<10000) { v = r+3; b = 0; }
    else {
      n = n / 10000;
      r = result + 4;
    }
  }
  return v;
}
```

|  (a) basic version $P_1$  |  (b) intermediate version $P_1'$  |  (c) optimized version $P_2$  |

The programs $P_1$ and $P_2$ shown above are reformulations of those given in [1] in order to comply with the input requirements of our tool. The `do-while` and `for` loops have been replaced by `while` loops. The boolean flag `b` and the temporary storage variable `v` in $P_2$ have been introduced to avoid premature returns from the function.

**Figure 5: Computing the number of digits (`digits10`) from [1]**

with quicksort), or when changing the design of a system in a fundamental way, it is less likely that an equivalence proof can be found automatically.

Our method works well whenever sufficiently "simple" coupling predicates exist that prove program equivalence. This applies in a number of important cases:

- as a baseline, our procedure will always be able to prove that a program is equivalent to itself, by applying the greedy reduction strategy from Section 4, and choosing the coupling predicates $\bar{x}_1 = \bar{x}_2$.

- the procedure is also complete when applied to two programs with the same control structure and locally equivalent (though not necessarily identical) loop and function bodies. In this case, the same coupling predicates $\bar{x}_1 = \bar{x}_2$ can be chosen for the entry points of the bodies.

- the procedure is complete for program transformations that correspond to affine mappings of program states; this includes renaming or exchanging variables, shifting the value of a variable by a constant offset, or changing the sign of some variable.

We currently do not consider equivalence of programs that use arrays or heap data structures, but we intend to work on lifting these limitations in the future. It is, for instance, known that assertions about arrays can be encoded in Horn clauses [13]. An approach to regression verification of programs with tree-shaped heap structures can be presumably adapted from [17].

## 6. RELATED WORK

Research on proving program equivalence is driven by a variety of applications, including security verification, com-

piler optimizations, backwards compatibility and refactoring, cryptographic algorithms, hardware design, and general-purpose regression verification.

Godlin and Strichman [16–18] present an approach for automatic general-purpose regression verification. In this approach, loops in the programs are transformed to recursive procedures, and matching recursive calls are abstracted by an uninterpreted function. The equivalence of functions (that no longer contain recursion) is then checked by the CBMC model checker. In our vernacular, the approach can be described as an attempt to verify equivalence with the fixed coupling predicate $\bar{i}_1 = \bar{i}_2 \rightarrow r_1 = r_2$ for every related pair of recursive functions. This abstraction imposes the limitation that function calls with different arguments or a different number of recursions of two matching recursive functions are not supported. The technique is implemented in the RVT tool and supports a subset of ANSI C.

Verdoolaege et al. [30,31] have developed an automatic approach to prove equivalence of static affine programs. The approach focuses on programs with array-manipulating `for` loops and can automatically deal with complex loop transformations such as loop interchange, reversal, skewing, tiling, and others. It is implemented in the ISA tool for the static affine subset of ANSI C. Initially, dataflow analysis is applied to build a dependence graph abstraction of each of the two programs. Then the equivalence hypothesis for outputs is propagated through the graphs towards the inputs, in a manner resembling verification condition generation. The static control flow requirement means that the control flow of the program must be known already at compile time. Furthermore, arithmetical operations in the loop/function bodies are abstracted. Addition is, e.g., replaced by an associative and commutative uninterpreted function. The abstraction prevents proving equivalence of such programs as `x=x+1; x=x+1;` and `x=x+2;`.

Barthe et al. [9] present a calculus for reasoning about relations between programs that is based on pure program transformation. The calculus offers rules to merge two programs into a single *product program*. The merging process is guided by the user and facilitates proving relational properties with the help of existing verification technology (the WHY tool, in that particular case). The verification process still requires user-supplied annotations though.

Almeida et al. [2] have verified the correctness of the Open-SSL implementation of the RC4 cipher w.r.t. a reference implementation. The authors use self-composition of programs together with interactively verified lemmas about particular program transformations and optimizations.

Sinz and Post [26] prove equivalence of two AES cipher implementations by means of bounded model checking. The approach unrolls resp. inlines all loops and recursive calls. Such reasoning is only feasible if the program admits small bounds on loops or depth of recursive calls. In the case of AES, a complete unrolling of the main loop was not possible, so the authors proved equivalence of loop bodies instead.

Backes et al. [5] propose to leverage slicing and impact analysis to improve scalability of regression verification. The idea is to subject both program versions to a dependency analysis, then to remove the code present in both versions that has no data or control dependencies on the introduced change, and to apply an existing technique (e.g., bounded symbolic execution) to show equivalence of the reduced programs.

Mutual function summaries have been prominently put forth by Hawblitzel et al. in [21] and later developed in [22]. The concept is implemented in the equivalence checker SYM-DIFF [25], where the user supplies the mutual summary, and the verification conditions are discharged by BOOGIE. Loops are encoded as recursion. The BCVERIFIER tool for proving backwards compatibility of Java class libraries by Welsch and Poetzsch-Heffter [32] has a similar pragmatics.

Banerjee and Naumann [6, 7] study equivalence of Java-like programs from the perspective of data encapsulation. They develop a programming discipline and a static analysis ensuring that changes in an object-oriented data structure's implementation are confined and cannot affect its clients other than through specified public methods.

Several relational program logics (e.g., [4,8,28]) have been developed for security applications. Proving in these logics requires user-supplied inductive invariants.

A large body of work also exists on equivalence checking of hardware logic circuits; see [24] for an overview. The approaches fall into two major groups. One group builds the product machine of two circuits and exhaustively traverses the state space to ensure that the corresponding outputs of the two circuits are identical in every reachable state. The other group recognizes that the incremental nature of the design process induces structural similarity between the circuit variants under verification and tries to exploit them. The techniques to do so include functional equivalences, indirect implications, permissible functions, and others (see e.g., [29]).

## 7. CONCLUSION AND FUTURE WORK

In this paper, we have presented a novel approach that uses invariant inference techniques to automatically conduct regression proofs for two imperative integer programs. To this end, the two versions of the program are trans-formed into Horn clauses over uninterpreted predicate symbols. These clauses constrain equivalence-witnessing coupling predicates that connect the states of the two programs at key points. A Horn constraint solver is used to find a solution for the coupling predicates, if one exists.

The approach is implemented and we have demonstrated its effectiveness on integer programs with non-trivial arithmetic and control flow. Future work includes an extension to programs with arrays and heap structures, as well as development of more fine-grained coupling schemes.

## 8. REFERENCES

[1] A. Alexandrescu. Three optimization tips for C++, 2012. A presentation at Facebook NYC. Available at `www.facebook.com/notes/facebook-engineering/three-optimization-tips-for-c/10151361643253920`.

[2] J. Almeida, M. Barbosa, J. Sousa Pinto, and B. Vieira. Verifying cryptographic software correctness with respect to reference implementations. In M. Alpuente, B. Cook, and C. Joubert, editors, *Formal Methods for Industrial Critical Systems*, volume 5825 of *Lecture Notes in Computer Science*, pages 37–52. Springer Berlin / Heidelberg, 2009.

[3] P. Ammann and J. Offutt. *Introduction to Software Testing*. Cambridge University Press, New York, NY, USA, first edition, 2008.

[4] T. Amtoft, S. Bandhakavi, and A. Banerjee. A logic for information flow in object-oriented programs. In *Conference record of the 33rd ACM SIGPLAN-SIGACT symposium on Principles of programming languages*, POPL '06, pages 91–102, New York, NY, USA, 2006. ACM.

[5] J. Backes, S. Person, N. Rungta, and O. Tkachuk. Regression verification using impact summaries. In E. Bartocci and C. Ramakrishnan, editors, *Model Checking Software*, volume 7976 of *Lecture Notes in Computer Science*, pages 99–116. Springer Berlin Heidelberg, 2013.

[6] A. Banerjee and D. A. Naumann. Ownership confinement ensures representation independence for object-oriented programs. *J. ACM*, 52(6):894–960, 2005.

[7] A. Banerjee and D. A. Naumann. State based ownership, reentrance, and encapsulation. In *Proceedings of the 19th European Conference on Object-Oriented Programming*, ECOOP'05, pages 387–411, Berlin, Heidelberg, 2005. Springer-Verlag.

[8] G. Barthe, J. Crespo, B. Grégoire, C. Kunz, and S. Zanella Béguelin. Computer-aided cryptographic proofs. In L. Beringer and A. Felty, editors, *Interactive Theorem Proving*, volume 7406 of *Lecture Notes in Computer Science*, pages 11–27. Springer Berlin Heidelberg, 2012.

[9] G. Barthe, J. M. Crespo, and C. Kunz. Relational verification using product programs. In M. Butler and W. Schulte, editors, *Proceedings, 17th International Symposium on Formal Methods (FM)*, volume 6664 of *Lecture Notes in Computer Science*, pages 200–214. Springer, 2011.

[10] G. Barthe, P. R. D'Argenio, and T. Rezk. Secure information flow by self-composition. In *17th IEEE Computer Security Foundations Workshop, CSFW-17, Pacific Grove, CA, USA*, pages 100–114. IEEE Computer Society, 2004.

[11] A. Darvas, R. Hähnle, and D. Sands. A theorem proving approach to analysis of secure information flow. In *Proceedings of the Second International Conference on Security in Pervasive Computing*, SPC'05, pages 193–209, Berlin, Heidelberg, 2005. Springer-Verlag.

[12] E. W. Dijkstra. Guarded commands, nondeterminacy and formal derivation of programs. *Communications of the ACM*, 18(8):453–457, Aug. 1975.

[13] I. Dillig, T. Dillig, and A. Aiken. Fluid updates: Beyond strong vs. weak updates. In *Proceedings of the 19th European Conference on Programming Languages and Systems*, ESOP'10, pages 246–266, Berlin, Heidelberg, 2010. Springer-Verlag.

[14] S. Falke, D. Kapur, and C. Sinz. Termination analysis of imperative programs using bitvector arithmetic. In *Proceedings of the 4th International Conference on Verified Software: Theories, Tools, Experiments (VSTTE'12)*, pages 261–277, Berlin, Heidelberg, 2012. Springer-Verlag.

[15] J. Giesl, R. Thiemann, P. Schneider-Kamp, and S. Falke. Automated termination proofs with AProVE. In V. van Oostrom, editor, *Rewriting Techniques and Applications, 15th International Conference (RTA 2004), Proceedings*, volume 3091 of *Lecture Notes in Computer Science*, pages 210–220. Springer, 2004.

[16] B. Godlin and O. Strichman. Inference rules for proving the equivalence of recursive procedures. *Acta Inf.*, 45(6):403–439, 2008.

[17] B. Godlin and O. Strichman. Regression verification. In *Proceedings of the 46th Annual Design Automation Conference*, DAC '09, pages 466–471. ACM, 2009.

[18] B. Godlin and O. Strichman. Regression verification: proving the equivalence of similar programs. *Software Testing, Verification and Reliability*, 23(3):241–258, 2013.

[19] S. Grebenshchikov, N. P. Lopes, C. Popeea, and A. Rybalchenko. Synthesizing software verifiers from proof rules. In *Proceedings of the 33rd ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI '12, pages 405–416. ACM, 2012.

[20] J. Harrison. *Handbook of Practical Logic and Automated Reasoning*. Cambridge University Press, 2009.

[21] C. Hawblitzel, M. Kawaguchi, S. K. Lahiri, and H. Rebêlo. Mutual summaries: Unifying program comparison techniques. In *Proceedings, First International Workshop on Intermediate Verification Languages (BOOGIE)*, 2011. Available at

http://research.microsoft.com/en-us/um/people/moskal/boogie2011/boogie2011_pg40.pdf.

[22] C. Hawblitzel, M. Kawaguchi, S. K. Lahiri, and H. Rebêlo. Towards modularly comparing programs using automated theorem provers. In M. P. Bonacina, editor, *Automated Deduction - CADE-24 - 24th International Conference on Automated Deduction, Lake Placid, NY, USA, June 9-14, 2013. Proceedings*, volume 7898 of *Lecture Notes in Computer Science*, pages 282–299. Springer, 2013.

[23] K. Hoder and N. Bjørner. Generalized property directed reachability. In *Proceedings of the 15th International Conference on Theory and Applications of Satisfiability Testing*, SAT'12, pages 157–171, Berlin, Heidelberg, 2012. Springer-Verlag.

[24] S.-Y. Huang and K.-T. Cheng. *Formal Equivalence Checking and Design DeBugging*. Kluwer Academic Publishers, Norwell, MA, USA, 1998.

[25] S. K. Lahiri, C. Hawblitzel, M. Kawaguchi, and H. Rebêlo. SymDiff: A language-agnostic semantic diff tool for imperative programs. In *Proceedings of the 24th International Conference on Computer Aided Verification*, CAV'12, pages 712–717, Berlin, Heidelberg, 2012. Springer-Verlag.

[26] H. Post and C. Sinz. Proving functional equivalence of two AES implementations using bounded model checking. In *Proceedings of the 2009 International Conference on Software Testing Verification and Validation*, ICST '09, pages 31–40. IEEE Computer Society, 2009.

[27] P. Rümmer, H. Hojjat, and V. Kuncak. Disjunctive interpolants for Horn-clause verification. In *Proceedings of the 25th International Conference on Computer Aided Verification*, CAV'13, pages 347–363, Berlin, Heidelberg, 2013. Springer-Verlag.

[28] C. Scheben and P. H. Schmitt. Efficient self-composition for weakest precondition calculi. In C. B. Jones, P. Pihlajasaari, and J. Sun, editors, *Proceedings, 19th International Symposium on Formal Methods (FM)*, volume 8442 of *Lecture Notes in Computer Science*, pages 579–594. Springer, 2014.

[29] C. van Eijk. Sequential equivalence checking based on structural similarities. *Computer-Aided Design of Integrated Circuits and Systems, IEEE Transactions on*, 19(7):814–819, 2000.

[30] S. Verdoolaege, G. Janssens, and M. Bruynooghe. Equivalence checking of static affine programs using widening to handle recurrences. *ACM Trans. Program. Lang. Syst.*, 34(3):11:1–11:35, 2012.

[31] S. Verdoolaege, M. Palkovic, M. Bruynooghe, G. Janssens, and F. Catthoor. Experience with widening based equivalence checking in realistic multimedia systems. *J. Electronic Testing*, 26(2):279–292, 2010.

[32] Y. Welsch and A. Poetzsch-Heffter. Verifying backwards compatibility of object-oriented libraries using Boogie. In *Proceedings of the 14th Workshop on Formal Techniques for Java-like Programs*, FTfJP '12, pages 35–41. ACM, 2012.