

From Two-Way to One-Way Finite State Transducers

Emmanuel Filiot Olivier Gauwin Pierre-Alain Reynier Frédéric Servais
 LACL LaBRI, LIF, Aix-Marseille Univ. Hasselt University and
 University Paris-Est Créteil University of Bordeaux & CNRS, UMR 7279 Transnational University of Limburg

Abstract—Any two-way finite state automaton is equivalent to some one-way finite state automaton. This well-known result, shown by Rabin and Scott and independently by Shepherdson, states that two-way finite state automata (even non-deterministic) characterize the class of regular languages. It is also known that this result does not extend to finite string transductions: (deterministic) two-way finite state transducers strictly extend the expressive power of (functional) one-way transducers. In particular deterministic two-way transducers capture exactly the class of MSO-transductions of finite strings.

In this paper, we address the following definability problem: given a function defined by a two-way finite state transducer, is it definable by a one-way finite state transducer? By extending Rabin and Scott's proof to transductions, we show that this problem is decidable. Our procedure builds a one-way transducer, which is equivalent to the two-way transducer, whenever one exists.

I. INTRODUCTION

In formal language theory, the importance of a class of languages is often supported by the number and the diversity of its characterizations. One of the most famous example is the class of regular languages of finite strings, which enjoys, for instance, computational (automata), algebraic (syntactic congruence) and logical (monadic second order (MSO) logic with one successor) characterizations. The study of regular languages has been very influential and several generalizations have been established. Among the most notable ones are the extensions to infinite strings [1] and trees [2]. On finite strings, it is well-known that both deterministic and non-deterministic finite state automata define regular languages. It is also well-known that the expressive power of finite state automata does not increase when the reading head can move left and right, even in presence of non-determinism. The latter class is known as non-deterministic *two-way finite state automata* and it is no more powerful than (one-way) finite state automata. The proof of this result was first shown in the seminal paper of Rabin and Scott [3], and independently by Shepherdson [4].

The picture of automata models over finite strings changes substantially when, instead of languages, string *transductions*, i.e. relations from strings to strings, are considered. *Transducers* generalize automata as they are equipped with a one-way output tape. At each step they read an input symbol, they can append several symbols to the output tape. Their transition systems can be either deterministic or non-deterministic. *Functional* transducers are transducers that define functions instead

of relations. For instance, deterministic transducers are always functional. In this paper, we are interested in transducers that define functions, but that can be non-deterministic.

As for automata, the reading head of transducers can move one-way (left-to-right) or two-way. (*One-way*) *finite state transducers* have been extensively studied [5], [6]. Non-deterministic (even functional) one-way transducers (*NFTs*) strictly extend the expressive power of deterministic one-way transducers (*DFTs*), because non-determinism allows one to express local transformations that depend on properties of the future of the input string.

Two-way finite state transducers define regular transformations that are beyond the expressive power of one-way transducers [7]. They can for instance reverse an input string, swap two substrings or copy a substring. The transductions defined by two-way transducers have been characterized by other logical and computational models. Introduced by Courcelle, monadic second-order definable transductions are transformations from graphs to graphs defined with the logic MSO [8]. Engelfriet and Hoozeboom have shown that the monadic second-order definable functions are exactly the functions definable by deterministic two-way finite state transducers (*2DFTs*) when the graphs are restricted to finite strings [9]. Recently, Alur and Černý have characterized *2DFT*-definable transductions by a deterministic one-way model called *streaming string transducers* [10] and shown how they can be applied to the verification of list-processing programs [11]. Streaming string transducers extend *DFTs* with a finite set of output string variables. At each step, their content can be reset or updated by either prepending or appending a finite string, or the content of another variable, in a copyless manner. Extending *2DFTs* with non-determinism does not increase their expressive power when they define functions: non-deterministic two-way finite state transducers (*2NFTs*) that are functional define exactly the class of functions definable by *2DFTs* [9], [12]. To summarize, there is a strict hierarchy between *DFT*-, functional *NFT*- and *2DFT*-definable transductions.

Several important problems are known to be decidable for one-way transducers. The *functionality* problem for *NFT*, decidable in PTime [13], [14], asks whether a given *NFT* is functional. The *determinizability* problem, also decidable in PTime [15], [14], asks whether a given functional *NFT* can be determinized, i.e. defines a *subsequential* function. Subsequential functions are those functions that can be defined by *DFTs* equipped with an additional output function from final states to finite strings, which is used to append a last string

This work has been partly supported by the project ECSPER funded by the french agency for research (ANR-09-JCJC-0069), by the project SOSP funded by the CNRS, and by the Faculty of Sciences of University Paris-Est Créteil.

to the output when the computation terminates successfully in some final state. Over strings that always end with a unique end marker, subsequential functions are exactly the functions definable by *DFTs*. For *2NFTs*, the functionality problem is known to be decidable [16]. Therefore the determinizability problem is also decidable for *2NFTs*, since functional *2NFTs* and *2DFTs* have the same expressive power. In the same line of research, we address a definability problem in this paper. In particular we answer the fundamental question of *NFT*-definability of transductions defined by functional *2NFTs*.

Theorem 1. *For all functional 2NFTs T , it is decidable whether the transduction defined by T is definable by an NFT.*

The proof of Theorem 1 extends the proof of Rabin and Scott [3] from automata to transducers¹. The original proof of Rabin and Scott is based on the following observation about the runs of two-way automata. Their shapes have a nesting structure: they are composed of many zigzags, each zigzag being itself composed of simpler zigzags. Basic zigzags are called *z*-motions as their shapes look like a *Z*. Rabin and Scott prove that for automata, it is always possible to replace a *z*-motion by a single pass. Then from a two-way automaton A it is possible to construct an equivalent two-way automaton B (called the squeeze of A) which is simpler in the following sense: accepting runs of B are those of A in which some *z*-motions have been replaced by single pass runs. Last, they argue² that after a number of applications of this construction that depends only on the number of states of A , every zigzag can be removed, yielding an equivalent one-way automaton.

The extension to *2NFTs* faces the following additional difficulty: it is not always possible to replace a *z*-motion of a transducer by a single pass. Intuitively, this is due to the fact that *2NFTs* are strictly more expressive than *NFTs*. As our aim is to decide when a *2NFT* T is *NFT*-definable, we need to prove that the *NFT*-definability of T implies that of every *z*-motion of T , to be able to apply the squeeze construction. The main technical contribution of this paper is thus the study of the *NFT*-definability of *z*-motions of transducers. We show that this problem is decidable, and identify a characterization which allows one to prove that the *NFT*-definability of T implies that of every *z*-motion of T .

This characterization expresses requirements about the output strings produced along loops of *z*-motions. We show that when *z*-motions are *NFT*-definable, the output strings produced by the three passes on a loop are not arbitrary, but conjugates. This allows us to give a precise characterization of the form of these output strings. We show that it is decidable to check whether all outputs words have this form. Last, we present how to use this characterization to simulate an *NFT*-definable *z*-motion by a single pass.

Applications By Theorem 1 and since functionality is decid-

able for *2NFTs*, it is also decidable, given a *2NFT*, whether the transduction it defines is definable by a functional *NFT*. Another corollary of Theorem 1 and the fact that functionality of *2NFTs* and determinizability of *NFTs* are both decidable is the following theorem:

Theorem 2. *For all 2NFTs T , it is decidable whether the transduction defined by T is a subsequential function.*

A practical application of this result lies in the static analysis of memory requirements for evaluating (textual and functional) document transformations in a streaming fashion. In this scenario, the input string is received as a left-to-right stream. When the input stream is huge, it should not be entirely loaded in memory but rather processed on-the-fly. Similarly, the output string should not be stored in memory but produced as a stream. The remaining amount of memory needed to evaluate the transformation characterizes its streaming space complexity. *Streamable* transformations are those transformations for which the required memory is bounded by a constant, and therefore is independent on the length of the input stream. It is known that streamable transformations correspond to transformations definable by subsequential (functional) *NFTs* [18]. The *streamability* problem asks, given a transformation defined by some transducer, whether it is streamable. Therefore for transformations defined by functional *NFTs*, streamability coincides with determinizability, and is decidable in PTime [15], [14]. Theorem 2 is a generalization of this latter result to regular transformations, i.e. transformations defined by functional *2NFTs*, MSO transducers or streaming string transducers [10]. Other streamability problems have been studied for XML validation [19], [20], XML queries [21] and XML transformations [18]. However the XML tree transformations of [18] are incomparable with the regular string transformations studied in this paper.

Related work Most of the related work has already been mentioned. To the best of our knowledge, it is the first result that addresses a definability problem between two-way and one-way transducers. In [22], two-way transducers with a two-way output tape are introduced with a special output policy: each time a cell at position i of the input tape is processed, the output is written in the cell at position i of the output tape. With that restriction, it is shown that two-way and one-way transducers (*NFTs*) define the same class of functions. In [23], the result of Rabin and Scott, and Shepherdson, is extended to two-way automata with multiplicities. In this context, two-way automata strictly extend one-way automata.

Organization of the paper Section II introduces necessary preliminary definitions. In Section III, we describe the general decision procedure for testing *NFT*-definability of functional *2NFTs*. We introduce *z*-motion transductions induced by *2NFTs* and show that their *NFT*-definability is necessary. The decidability of this necessary condition as well as the construction from *z*-motion transducers to *NFTs* are the most technical results of this paper and are the subject of Section IV. We finally discuss side results and further questions in Section V.

¹Shepherdson [4] and then Vardi [17] proposed arguably simpler constructions for automata. It is however not clear to us how to extend these constructions to transducers.

²To our knowledge, there is no published proof of this result, thus we prove it in this paper as we use it for transducers.

All technical details and omitted proofs can be found in the full version of this paper [24].

II. ONE-WAY AND TWO-WAY FINITE STATE MACHINES

Words, Languages and Transductions Given a finite alphabet Σ , we denote by Σ^* the set of finite words over Σ , and by ϵ the empty word. The length of a word $u \in \Sigma^*$ is its number of symbols, denoted by $|u|$. For all $i \in \{1, \dots, |u|\}$, we denote by $u[i]$ the i -th letter of u . Given $1 \leq i \leq j \leq |u|$, we denote by $u[i..j]$ the word $u[i]u[i+1] \dots u[j]$ and by $u[j..i]$ the word $u[j]u[j-1] \dots u[i]$. We say that $v \in \Sigma^*$ is a *factor* of u if there exist $u_1, u_2 \in \Sigma^*$ such that $u = u_1vu_2$. By \bar{u} we denote the *mirror* of u , i.e. the word of length $|u|$ such that $\bar{u}[i] = u[|u| - i + 1]$ for all $1 \leq i \leq |u|$.

The *primitive root* of $u \in \Sigma^*$ is the shortest word v such that $u = v^k$ for some integer $k \geq 1$, and is denoted by $\mu(u)$. Two words u and v are *conjugates*, denoted by \sim , if there exist $x, y \in \Sigma^*$ such that $u = xy$ and $v = yx$, i.e. u can be obtained from v by a cyclic permutation. Note that \sim is an equivalence relation. We will use this fundamental lemma:

Lemma 1 ([25]). *Let $u, v \in \Sigma^*$. If there exists $n \geq 0$ such that u^n and v^n have a common factor of length at least $|u| + |v| - \gcd(|u|, |v|)$, then $\mu(u) \sim \mu(v)$.*

Note that if $\mu(u) \sim \mu(v)$, then there exist $x, y \in \Sigma^*$ such that $u \in (xy)^*$ and $v \in (yx)^*$.

A *language* over Σ is a set $L \subseteq \Sigma^*$. A *transduction* over Σ is a relation $R \subseteq \Sigma^* \times \Sigma^*$. Its domain is denoted by $\text{dom}(R)$, i.e. $\text{dom}(R) = \{u \mid \exists v, (u, v) \in R\}$, while its image $\{v \mid \exists u, (u, v) \in R\}$ is denoted by $\text{img}(R)$. A transduction R is *functional* if it is a function.

Automata A *non-deterministic two-way finite state automaton*³ (2NFA) over a finite alphabet Σ is a tuple $A = (Q, q_0, F, \Delta)$ where Q is a finite set of states, $q_0 \in Q$ is the initial state, $F \subseteq Q$ is a set of final states, and Δ is the transition relation, of type $\Delta \subseteq Q \times \Sigma \times Q \times \{+1, -1\}$. It is *deterministic* if for all $(p, a) \in Q \times \Sigma$, there is at most one pair $(q, m) \in Q \times \{+1, -1\}$ such that $(p, a, q, m) \in \Delta$. In order to see how words are evaluated by A , it is convenient to see the input as a right-infinite input tape containing the word (starting at the first cell) followed by blank symbols. Initially the head of A is on the first cell in state q_0 (the cell at position 1). When A reads an input symbol, depending on the transitions in Δ , its head moves to the left (-1) if the head was not in the first cell, or to the right ($+1$) and changes its state. A stops as soon as it reaches a blank symbol (therefore at the right of the input word), and the word is accepted if the current state is final.

A *configuration* of A is a pair $(q, i) \in Q \times (\mathbb{N} - \{0\})$ where q is a state and i is a position on the input tape. A *run* ρ of A is a finite sequence of configurations. The run $\rho = (p_1, i_1) \dots (p_m, i_m)$ is a run on an input word $u \in \Sigma^*$ of length n if $p_1 = q_0$, $i_1 = 1$, $i_m \leq n + 1$, and for all

$k \in \{1, \dots, m-1\}$, $1 \leq i_k \leq n$ and $(p_k, u[i_k], p_{k+1}, i_{k+1} - i_k) \in \Delta$. It is *accepting* if $i_m = n + 1$ and $p_m \in F$. The language of a 2NFA A , denoted by $L(A)$, is the set of words u such that there exists an accepting run of A on u .

A *non-deterministic (one-way) finite state automaton* (NFA) is a 2NFA such that $\Delta \subseteq Q \times \Sigma \times Q \times \{+1\}$, therefore we will often see Δ as a subset of $Q \times \Sigma \times Q$. Any 2NFA is effectively equivalent to an NFA. It was first proved by Rabin and Scott, and independently by Shepherdson [3], [4].

Transducers *Non-deterministic two-way finite state transducers* (2NFTs) over Σ extend NFAs with a one-way left-to-right output tape. They are defined as 2NFAs except that the transition relation Δ is extended with outputs: $\Delta \subseteq Q \times \Sigma \times \Sigma^* \times Q \times \{-1, +1\}$. If a transition (q, a, v, q', m) is fired on a letter a , the word v is appended to the right of the output tape and the transducer goes to state q' . Wlog we assume that for all $p, q \in Q$, $a \in \Sigma$ and $m \in \{+1, -1\}$, there exists at most one $v \in \Sigma^*$ such that $(p, a, v, q, m) \in \Delta$. We also denote v by $\text{out}(p, a, q, m)$.

A run of a 2NFTs is a run of its underlying automaton, i.e. the 2NFAs obtained by ignoring the output. A run ρ may be simultaneously a run on a word u and on a word $u' \neq u$. However, when the underlying input word is given, there is a unique sequence of transitions associated with ρ . Given a 2NFT T , an input word $u \in \Sigma^*$ and a run $\rho = (p_1, i_1) \dots (p_m, i_m)$ of T on u , the output of ρ on u , denoted by $\text{out}^u(\rho)$, is the word obtained by concatenating the outputs of the transitions followed by ρ , i.e. $\text{out}^u(\rho) = \text{out}(p_1, u[i_1], p_2, i_2 - i_1) \dots \text{out}(p_{m-1}, u[i_{m-1}], p_m, i_m - i_{m-1})$. If ρ contains a single configuration, we let $\text{out}^u(\rho) = \epsilon$. When the underlying input word u is clear from the context, we may omit the exponent u . The transduction defined by T is the relation $R(T) = \{(u, \text{out}^u(\rho)) \mid \rho \text{ is an accepting run of } T \text{ on } u\}$. We may often just write T when it is clear from the context. A 2NFT T is *functional* if the transduction it defines is functional. The class of functional 2NFTs is denoted by $f2NFT$. In this paper, we mainly focus on $f2NFT$ s. The *domain* of T is defined as $\text{dom}(T) = \text{dom}(R(T))$. The domain $\text{dom}(T)$ is a regular language that can be defined by the 2NFA obtained by projecting away the output part of the transitions of T , called the *underlying input automaton*. A *deterministic two-way finite state transducer* (2DFT) is a 2NFT whose underlying input automaton is deterministic. Note that 2DFTs are always functional, as there is at most one accepting run per input word. A *non-deterministic (one-way) finite state transducer* (NFT) is a 2NFT whose underlying automaton is an NFA⁴. It is deterministic (written DFT) if the underlying automaton is a DFA.

We say that two transducers T, T' are equivalent, denoted by $T \equiv T'$, whenever they define the same transduction, i.e. $R(T) = R(T')$. For all transducer classes \mathcal{C} , we say that a transduction $R \subseteq \Sigma^* \times \Sigma^*$ is \mathcal{C} -definable if there exists $T \in \mathcal{C}$

³We follow the definition of Vardi [17], but without stay transitions. This is without loss of generality though.

⁴This definition implies that there is no ϵ -transitions that can produce outputs, which may cause the image of an input word to be an infinite language. Those NFTs are sometimes called *real-time* in the literature.

such that $R=R(T)$. Given two classes $\mathcal{C}, \mathcal{C}'$ of transducers, and a transducer $T \in \mathcal{C}$, we say that T is (effectively) \mathcal{C}' -definable if one can construct an equivalent transducer $T' \in \mathcal{C}'$.

The $(\mathcal{C}, \mathcal{C}')$ -definability problem takes as input a transducer $T \in \mathcal{C}$ and asks to decide whether T is \mathcal{C}' -definable. If so, one may want to construct an equivalent transducer $T' \in \mathcal{C}'$. In this paper, we prove that $(f2NFT, NFT)$ -definability is decidable.

It is known that whether an NFT T is functional can be decided in PTime [13]. The class of functional NFT s is denoted by $fNFT$. Functional NFT s are strictly more expressive than DFT s. For instance, the function that maps any word $u \in \{a, b\}^+$ to $a^{|u|}$ if $u[|u|] = a$, and to $b^{|u|}$ otherwise, is $fNFT$ -definable but not DFT -definable. This result does not hold for $2NFT$ s: functional $2NFT$ s and $2DFT$ s define the same class of transductions (Theorem 22 of [9]).

Examples Let $\Sigma = \{a, b\}$ and $\# \notin \Sigma$, and consider the transductions

- 1) $R_0 = \{(u, a^{|u|}) \mid u \in \Sigma^+, u[|u|] = a\}$
- 2) $R_1 = \{(u, b^{|u|}) \mid u \in \Sigma^+, u[|u|] = b\} \cup R_0$
- 3) $R_2 = \{(\#u\#, \#u\#) \mid u \in \Sigma^*\}$.

R_0 is DFT -definable: it suffices to replace each letter by a and to accept only if the last letter is a . Therefore it can be defined by the DFT $T_0 = (\{q_a, q_b\}, q_b, \{q_a\}, \{(q_x, y, a, q_y) \mid x, y \in \Sigma\})$.

R_1 is $fNFT$ -definable but not DFT -definable: similarly as before we can define a DFT $T'_0 = (\{p_a, p_b\}, p_a, \{p_b\}, \{(p_x, y, b, p_y) \mid x, y \in \Sigma\})$ that defines the transduction $\{(u, b^{|u|}) \mid u \in \Sigma^+, u[|u|] = b\}$, and construct an NFT T_1 as follows: its initial state is some fresh state p_0 , and when reading $x \in \Sigma$ the first time, it non-deterministically goes to T_0 or T'_0 by taking the transition (p_0, x, a, q_x) or (p_0, x, b, p_x) , and proceeds in either T_0 or T'_0 . Even if R_1 is functional, it is not DFT -definable, as the transformation depends on the property of the last letter, which can be arbitrarily far away from the beginning of the string.

R_2 is $2DFT$ -definable: it suffices to go to the end of the word by producing ϵ each time a letter is read, to go back to the beginning while copying each input letter, and return to the end without outputting anything, and to accept. Hence it is defined by $T_2 = (\{q_0, q_1, q_2, q_3, q_f\}, q_0, \{q_f\}, \delta_2)$ where states q_1, q_2, q_3 denote passes, and δ_2 is made of the transitions $(q_0, \#, \epsilon, q_1, +1)$, $(q_1, x \in \Sigma, \epsilon, q_1, +1)$ (during the first pass, move to the right), $(q_1, \#, \epsilon, q_2, -1)$, $(q_2, x \in \Sigma, x, q_2, -1)$, $(q_2, \#, \#, q_3, +1)$, $(q_3, x \in \Sigma, \epsilon, q_3, +1)$, $(q_3, \#, \#, q_f, +1)$.

Crossing Sequences, Loops and Finite-Crossing $2NFT$ s
The notion of crossing sequence is a useful notion in the theory of two-way automata [4], [26], that allows one to pump runs of two-way automata. Given a $2NFA$ A , a word $u \in \Sigma^*$ and a run ρ of A on u , the *crossing sequence* at position i , denoted by $CS(\rho, i)$ is given by the sequence of states q such that (q, i) occurs in ρ . The order of the sequence is given by the order in which the pairs of the form (q, i) occur in ρ . E.g. if $\rho = (q_1, 1)(q_2, 2)(q_3, 1)(q_4, 2)(q_5, 1)(q_6, 2)(q_7, 3)$ then $CS(\rho, 1) = q_1q_3q_5$, $CS(\rho, 2) = q_2q_4q_6$ and $CS(\rho, 3) = q_7$. We write $CS(\rho)$ the sequence $CS(\rho, 1), \dots, CS(\rho, |u| + 1)$.

Crossing sequences allow one to define the loops of a run. Given a run ρ of the $2NFA$ A on some word u of length n , a pair of positions (i, j) is a *loop*⁵ in ρ if (i) $1 \leq i \leq j \leq n$, (ii) $CS(\rho, i) = CS(\rho, j)$ and (iii) $u[i] = u[j]$. Let $u_1 = u[1..(i-1)]$, $u_2 = u[i..(j-1)]$ and $u_3 = u[j..n]$. If (i, j) is a loop in ρ and $u \in L(A)$, then $u_1(u_2)^k u_3 \in L(A)$ for all $k \geq 0$. We say that a loop (i, j) is *empty* if $i = j$, in this case we have $u_2 = \epsilon$. The notions of crossing sequence and loop carry over to transducers through their underlying input automata.

Given a $2NFT$ T , $N \in \mathbb{N}$ and a run ρ of T on a word of length n , ρ is said to be N -crossing if $|CS(\rho, i)| \leq N$ for all $i \in \{1, \dots, n\}$. The transducer T is *finite-crossing* if there exists $N \in \mathbb{N}$ such that for all $(u, v) \in R(T)$, there is an accepting N -crossing run ρ on u such that $out(\rho) = v$. In that case, T is said to be N -crossing. It is easy to see that if T is N -crossing, then for all $(u, v) \in R(T)$ there is an accepting run ρ on u such that $out(\rho) = v$ and no states repeat in $CS(\rho, i)$ for all $i \in \{1, \dots, |u|\}$. Indeed, if some state q repeats in some $CS(\rho, i)$, then it is possible to pump the subrun between the two occurrences of q on $CS(\rho, i)$. This subrun has an empty output, otherwise T would not be functional.

Proposition 1. Any $f2NFT$ with N states is N -crossing.

III. FROM TWO-WAY TO ONE-WAY TRANSDUCERS

In this section, we prove the main result of this paper, i.e. the decidability of $(f2NFT, NFT)$ -definability.

A. Rabin and Scott's Construction for Automata

The proof of Theorem 1 relies on the same ideas as Rabin and Scott's construction for automata [3]. It is based on the following key observation: Any accepting run is made of many zigzags, and those zigzags are organized by a nesting hierarchy: zigzag patterns may be composed of simpler zigzag patterns. The simplest zigzags of the hierarchy are those that do not nest any other zigzag: they are called z -motions. Rabin and Scott described a procedure that removes those zigzags by iterating a construction that removes z -motions.

A *one-step sequence* is an indexed sequence $s = a_1, \dots, a_n$ of positions such that $a_i \in \{1, 2, \dots, m\}$, $a_1 = 1$, $a_n = m$, and $|a_{i+1} - a_i| = 1$. The sequence s is N -crossing if for all $x \in \{1, 2, \dots, m\}$ we have $|\{i \mid a_i = x\}| \leq N$. The reversals of s are the indexes $1 < r_1 < r_2 < \dots < r_l < n$ such that $a_{r_i+1} = a_{r_i-1}$. In the sequel we let $r_0 = 1$ and $r_{l+1} = n$.

A *z-motion* z in s is a subsequence a_e, a_{e+1}, \dots, a_f such that there is $0 < i < l$ with $r_{i-1} \leq e < r_i < r_{i+1} < f \leq r_{i+2}$, and $a_e = a_{r_{i+1}}$ and $a_f = a_{r_i}$. We may denote z by the pair of reversals (r_i, r_{i+1}) . E.g. the sequences $z_1 = 1, 2, 3, 2, 1, 2, 3$ and $z_2 = 4, 3, 2, 3, 4, 3, 2$ are z -motions. The *shape* of a run ρ is defined as the second projection of ρ , written $shape(\rho)$. A run ρ is a *z-motion run* if $shape(\rho)$ is a z -motion. When there is no ambiguity, z -motion runs are just called z -motions.

⁵Observe that we include the input letter in the notion of loop. We use this to avoid technical difficulties due to backward transitions (which do not read the local symbol, but its successor).




Fig. 1. Zigzags removal by applications of squeeze.

If T is a 2NFA, it is possible to construct a new automaton denoted by $\text{squeeze}(T)$ such that, for all accepting runs ρ of T on some input word u , there exists a “simpler” accepting run of $\text{squeeze}(T)$ on u , obtained from ρ by replacing some z -motions by one-way runs that simulate three passes in parallel. It is illustrated by Fig. 1. For instance at the first step, there are two z -motions from q_1 to q_2 and from q_3 to q_4 respectively. Applying $\text{squeeze}(T)$ consists in non-deterministically guessing those z -motions and simulating them by one-way runs. This is done by the NFA $R_T(q_1, q_2)$ and $R_T(q_3, q_4)$ respectively. Depending on whether the z -motions enter from the left or the right, z -motions are replaced by runs of NFAs $R_T(\cdot, \cdot)$ (that read the input backwardly) or $L_T(\cdot, \cdot)$, as illustrated by the second iteration of squeeze on Fig. 1.

An N -crossing run ρ can be simplified into a one-way run after a constant number of applications of squeeze . This result is unpublished so we prove it in this paper. In particular, we show that if ρ is N -crossing, then its zigzag nesting depth decreases after N steps. Moreover, if ρ is N -crossing, then its zigzag nesting depth is also bounded by N . Therefore after N^2 applications of squeeze , ρ is transformed into a simple one-way run. It is sufficient to prove those results at the level of integer sequences. In particular, one can define $\text{squeeze}(s)$ the set of sequences obtained from a one-step sequence s by replacing *some* z -motions of s by strictly increasing or decreasing subsequences (for more details see [24]).

Lemma 2. *Let s be an N -crossing one-step sequence over $\{1, \dots, m\}$. Then $1, 2, \dots, m$ is in $\text{squeeze}^{N^2}(s)$.*

At the automata level, it is known that for all words u accepted by a 2NFA T with N states, there exists an N -crossing accepting run on u . Therefore it suffices to apply squeeze N^2 times to T . One gets an equivalent 2NFA T^* from which the backward transitions can be removed while




Fig. 2. Output decomposition in property \mathcal{P} .

preserving equivalence with T^* , and so T .

B. Extension to transducers: overview

The construction used to show decidability of $f2NFT$ follows the same ideas as Rabin and Scott’s construction. The main difference relies in the transformation of the local transducers defined by z -motion runs (that we call $ZNFTs$) into $NFTs$. Our procedure is built over a $ZNFT$ -to- NFT procedure. It is seen as a black-box in this section, but is the subject of the next section.

Compared to two-way automata, one faces an extra difficulty caused by the fact that $2NFTs$ (and $ZNFTs$) are not always NFT -definable. Therefore one defines a necessary condition that has to be tested each time we want to apply squeeze . Let us consider again Fig. 1 when T is a $2NFT$. One defines from T the transductions induced by local z -motion runs from a starting state q_1 to an ending state q_2 , and show that those local transductions must be NFT -definable.

Once this necessary condition is satisfied, the construction squeeze can be applied and works as for Rabin and Scott’s construction: the new transducer $\text{squeeze}(T)$ simulates T and non-deterministically may guess that the next zigzag of T is a z -motion run from some state q_1 to some state q_2 , and thus can be simulated by a run of some NFT $R_T(q_1, q_2)$ or $L_T(q_1, q_2)$, depending on whether it enters from the left or the right. Then $\text{squeeze}(T)$ switches to $R_T(q_1, q_2)$ (if it entered from the right) and once $R_T(q_1, q_2)$ reaches an accepting state, it may come back to its normal mode.

C. z -motion transducers

z -motion transducers are defined like $2NFTs$ except that they must define **functions** and to be accepting, a run on a word of length n must be of the form $\rho.(q_f, n+1)$ where ρ is a z -motion run and q_f is an accepting state. Note that it implies that $\text{shape}(\rho)$ is always of the form $1, \dots, n, n-1, \dots, 1, \dots, n$. The class of z -motion transducers is denoted by $ZNFTs$. Note that z -motion transducers are incomparable with $f2NFTs$. Indeed, z -motion transducers can define the transduction $u \in \Sigma^* \mapsto \bar{u}$, which is not $f2NFT$ -definable as there are no end markers.

Let $T \in ZNFT$ and $\rho = (p_1, 1) \dots (p_n, n) (q_{n-1}, n-1) \dots (q_1, 1) (r_2, 2) \dots (r_{n+1}, n+1)$ be a run of T on a word of length n . We let $q_n = p_n$ and $r_1 = q_1$ and define the following shortcuts: for $1 \leq i \leq j \leq n$, $\text{out}_1[i, j] = \text{out}((p_i, i) \dots (p_j, j))$, and $\text{out}_2[i, j] = \text{out}((q_j, j) \dots (q_i, i))$ and $\text{out}_3[i, j] = \text{out}((r_i, i) \dots (r_j, j))$, and $\text{out}_3[i, n+1] = \text{out}((r_i, i) \dots (r_{n+1}, n+1))$.

We characterize the NFT -definability of a $ZNFT$ by a property that we prove to be decidable. Intuitively, this property

requires that the outputs produced by loops can be produced by a single forward pass:

Definition 1 (\mathcal{P} -property). *Let T be a ZNFT. We say that T satisfies the property \mathcal{P} , denoted by $T \models \mathcal{P}$, if for all words $u \in \text{dom}(T)$, for all accepting runs ρ on u , and for all pairs of loops (i_1, j_1) and (i_2, j_2) of ρ such that $j_1 \leq i_2$, there exist $\beta_1, \beta_2, \beta_3, \beta_4, \beta_5 \in \Sigma^*$, $f, g : \mathbb{N}^2 \rightarrow \Sigma^*$ and constants $c_1, c'_1, c_2, c'_2 \geq 0$ such that $c_1, c_2 \neq 0$ and for all $k_1, k_2 \geq 0$,*

$$f(k_1, k_2) x_0 v_1^{\eta_1} x_1 w_1^{\eta_2} x_2 w_2^{\eta_2} x_3 v_2^{\eta_1} x_4 v_3^{\eta_1} x_5 w_3^{\eta_2} x_6 g(k_1, k_2) \\ = \beta_1 \beta_2^{k_1} \beta_3 \beta_4^{k_2} \beta_5$$

where $\eta_i = k_i c_i + c'_i$, $i \in \{1, 2\}$, and, x_i 's, v_i 's and w_i 's are words defined as depicted in Fig. 2.

The following key lemma is proved in Section IV.

Lemma 3. *Let $T \in \text{ZNFT}$. $T \models \mathcal{P}$ iff T is NFT-definable. Moreover, \mathcal{P} is decidable and if $T \models \mathcal{P}$, one can (effectively) construct an equivalent NFT.*

Definition 2 (z -motion transductions induced by a $f2\text{NFT}$). *Let $T = (Q, q_0, F, \Delta)$ be a $f2\text{NFT}$ and $q_1, q_2 \in Q$. The transduction $\mathcal{L}_T(q_1, q_2)$ (resp. $\mathcal{R}_T(q_1, q_2)$) is defined as the set of pairs (u_2, v_2) such that there exist $u \in \Sigma^*$, two positions $i_1 < i_2$ (resp. $i_2 < i_1$), an accepting run ρ of T on u which can be decomposed as $\rho = \rho_1(q_1, i_1)\rho_2(q_2, i_2)\rho_3$ such that $u_2 = u[i_1 \dots i_2]$ and*

- $(q_1, i_1)\rho_2(q_2, i_2)$ is a z -motion run
- $\text{out}((q_1, i_1)\rho_2(q_2, i_2)) = v_2$

z -motions can be of two forms: either they start from the left and end to the right, or start from the right and end to the left. In order to avoid considering these two cases each time, we introduce the notation \bar{T} that denotes the mirror of T : it is T where the moves $+1$ are replaced by -1 and the moves -1 by $+1$. Moreover, the way \bar{T} reads the input tape is slightly modified: it starts in position n and a run is accepting if it reaches position 0 in some accepting state. All the notions defined for 2NFT s carry over to their mirrors. In particular, $(u, v) \in R(T)$ iff $(\bar{u}, v) \in R(\bar{T})$. The z -motion transductions $\mathcal{R}_T(q_1, q_2)$ and $\mathcal{L}_T(q_1, q_2)$ are symmetric in the following sense: $\mathcal{R}_T(q_1, q_2) = \mathcal{L}_{\bar{T}}(q_1, q_2)$ and $\mathcal{L}_T(q_1, q_2) = \mathcal{R}_{\bar{T}}(q_1, q_2)$.

Proposition 2. *The transductions $\mathcal{R}_T(q_1, q_2)$ and $\mathcal{L}_T(q_1, q_2)$ are ZNFT-definable.*

Proof: We only consider the case $\mathcal{L}_T(q_1, q_2)$, the other case being solved by using the equality $\mathcal{R}_T(q_1, q_2) = \mathcal{L}_{\bar{T}}(q_1, q_2)$. We first construct from T a ZNFT $Z'_T(q_1, q_2)$ which is like T but its initial state is q_1 , and it can move to an accepting state whenever it is in q_2 . However $Z'_T(q_1, q_2)$ may define input/output pairs (u_2, v_2) that cannot be embedded into some pair $(u, v) \in R(T)$ as required by the definition of $\mathcal{L}_T(q_1, q_2)$. Based on Shepherdson's construction, we modify $Z'_T(q_1, q_2)$ in order to take this constraint into account. ■

In the next subsection, we show that $\mathcal{R}_T(q_1, q_2)$ and $\mathcal{L}_T(q_1, q_2)$ must necessarily be NFT-definable for T to be NFT-definable. For that purpose, it is crucial in Definition 2

to make sure that the z -motion $(q_1, i_1)\rho_2(q_2, i_2)$ can be embedded into a global accepting run of T . Without that restriction, it might be the case that $\mathcal{L}_T(q_1, q_2)$ or $\mathcal{R}_T(q_1, q_2)$ is not NFT-definable although the 2NFT T is. Indeed, the domain of $\mathcal{L}_T(q_1, q_2)$ or $\mathcal{R}_T(q_1, q_2)$ would be too permissive and accept words that would be otherwise rejected by other passes of global runs of T . This is another difficulty when lifting Rabin and Scott's proof to transducers, as for automata, the context in which a z -motion occurs is not important.

D. Decision procedure and proof of Theorem 1

We show that the construction $\text{squeeze}(T)$ can be applied if the following necessary condition is satisfied.

Lemma 4. *If T is NFT-definable, then so are the transductions $\mathcal{R}_T(q_1, q_2)$ and $\mathcal{L}_T(q_1, q_2)$ for all states q_1, q_2 . Moreover, it is decidable whether the transductions $\mathcal{R}_T(q_1, q_2)$ and $\mathcal{L}_T(q_1, q_2)$ are NFT-definable.*

Sketch of proof: We have seen in Lemma 3 that NFT-definability of an ZNFT is characterized by Property \mathcal{P} . Let $Z \in \text{ZNFT}$ that defines $\mathcal{L}_T(q_1, q_2)$ for some q_1, q_2 , we thus sketch the proof that $Z \models \mathcal{P}$.

Consider two loops (i_1, j_1) , (i_2, j_2) of a run ρ of Z on some word u , as in the premises of Property \mathcal{P} . They induce a decomposition of u as $u = u_1 u_2 u_3 u_4 u_5$ with $u_2 = u[i_1 \dots j_1 - 1]$ and $u_4 = u[i_2 \dots j_2 - 1]$. By definition of the transduction $\mathcal{L}_T(q_1, q_2)$, any word in $\text{dom}(Z)$ can be extended into a word in $\text{dom}(T)$. By hypothesis, T is NFT-definable, thus there exists an equivalent NFT T' . As T' has finitely many states, it is possible, by iterating the loops (i_1, j_1) and (i_2, j_2) , to identify an input word of the form

$$u' = \alpha u_1 u_2^{c_1} u_2^{c_2} u_2^{c_3} u_3 u_4^{c'_1} u_4^{c'_2} u_4^{c'_3} u_5 \alpha'$$

and a run ρ' of T' on this word which has two loops on the input subwords $u_2^{c_2}$ and $u_4^{c'_2}$. It is then easy to conclude. ■

Construction of $\text{squeeze}(T)$ Assuming that the necessary condition is satisfied, we now explain how to construct the $f2\text{NFT}$ $\text{squeeze}(T)$. By hypothesis, the transductions $\mathcal{L}_T(q_1, q_2)$ and $\mathcal{R}_T(q_1, q_2)$ are NFT-definable for all q_1, q_2 by NFT $L_T(q_1, q_2)$ and $R_T(q_1, q_2)$ respectively (they exist by Proposition 2 and Lemma 3). As already said before, the main idea to define $\text{squeeze}(T)$ is to non-deterministically (but repeatedly) apply $L_T(q_1, q_2)$, $R_T(q_1, q_2)$, or T , for some $q_1, q_2 \in Q$. However when applying $R_T(q_1, q_2)$, the head of $\text{squeeze}(T)$ should move from the right to the left, so that we have to mirror the transitions of $R_T(q_1, q_2)$.

The transducer $\text{squeeze}(T)$ has two modes, Z-mode or T-mode. In T-mode, it works as T until it non-deterministically decides that the next zigzag is a z -motion from some state q_1 to some state q_2 . Then it goes in Z-mode and runs $L_T(q_1, q_2)$ or $\bar{R}_T(q_1, q_2)$, in which transitions to an accepting state have been replaced by transitions from q_2 in T , so that $\text{squeeze}(T)$ returns in T-mode. From those transitions we also add transitions from the initial states of $L_T(q_2, q_3)$ and $\bar{R}_T(q_2, q_3)$ for all $q_3 \in Q$, in case $\text{squeeze}(T)$ guesses that




Fig. 3. From ZNFT to NFT.

the next z -motion starts immediately at the end of the previous z -motion (for more details see [24]).

Proposition 3. *Let $T \in f2NFT$ such that T is NFT-definable. Then $\text{squeeze}(T)$ is defined and equivalent to T .*

Let $T \in f2NFT$. If T is NFT-definable, then the operator squeeze can be iterated on T while preserving equivalence with T , by the latter proposition. By Proposition 1 T is N -crossing, and therefore, based on Lemma 2, it suffices to iterate squeeze N^2 times to remove all zigzags from accepting runs of T , as stated by the following lemma:

Lemma 5. *Let T be a $f2NFT$ with N states. If T is $fNFT$ -definable, then $\text{squeeze}^{N^2}(T)$ is defined and equivalent to T , and moreover, for all $(u, v) \in R(T)$, there exists an accepting run ρ of $\text{squeeze}^{N^2}(T)$ on u such that $\text{out}(\rho) = v$ and ρ is made of forward transitions only.*

Proof of Theorem 1 In order to decide whether a $f2NFT$ T is NFT-definable, it suffices to test whether squeeze can be applied N^2 times. More precisely, it suffices to set T_0 to T , i to 0, and, while T_i satisfies the necessary condition (which is decidable by Lemma 4) and $i \leq N^2$, to increase i and set T_i to $\text{squeeze}(T_{i-1})$. If the procedure exits the loops before reaching N^2 , then T is not NFT-definable, otherwise it is NFT-definable by the NFT obtained by removing from T_{N^2} all its backward transitions.

IV. FROM ELEMENTARY ZIGZAGS TO LINES

This section is devoted to the proof of Lemma 3 that characterizes NFT-definable ZNFT by the property \mathcal{P} and states its decidability. Moreover, we give a ZNFT-to-NFT construction when \mathcal{P} is satisfied.

We first prove that Property \mathcal{P} is a necessary condition for NFT-definability. To prove the converse, we proceed in two steps. First, we define a procedure that tests whether a given ZNFT T is equivalent to a ZNFT that does not output anything on its backward pass (called ϵZNFT), and then define another procedure that tests whether the latter ZNFT is equivalent to an NFT. We show that it is always true whenever $T \models \mathcal{P}$. This approach is depicted in Fig. 3. The two steps are similar, therefore we mainly focus on the first step.

A. Property \mathcal{P} is a necessary condition

We show that Property \mathcal{P} only depends on transductions.

Lemma 6. *Let $T, T' \in \text{ZNFT}$. If $T \models \mathcal{P}$ and $T \equiv T'$ then $T' \models \mathcal{P}$.*

Proof: Consider two loops $(i_1, j_1), (i_2, j_2)$ as in Property \mathcal{P} in a run of T' on some word u . They induce a decomposition of u as $u = u_1 u_2 u_3 u_4 u_5$ where $u_2 = u[i_1 \dots (j_1 - 1)]$ and




Fig. 4. Decomposition of the output according to Property \mathcal{P}_1 .

$u_4 = u[i_2 \dots (j_2 - 1)]$, with $u_1 u_2^{k_1} u_3 u_4^{k_2} u_5 \in \text{dom}(T')$ for all $k_1, k_2 \geq 0$.

As T is equivalent to T' and has finitely many states, there exist iterations of the loops on u_2 and u_4 which constitute loops in T on powers of u_2 and u_4 . Formally, there exist integers $d_1, e_1, h_1, d_2, e_2, h_2$ with $e_1, e_2 > 0$ such that T has a run ρ on the input word $u_1 u_2^{d_1} u_2^{e_1} u_2^{h_1} u_3 u_4^{d_2} u_4^{e_2} u_4^{h_2} u_5$ which contains a loop on the input subwords $u_2^{e_1}$ and $u_4^{e_2}$.

We conclude easily by using the fact that $T \models \mathcal{P}$. ■

As a consequence, we obtain that Property \mathcal{P} is a necessary condition for NFT-definability.

Lemma 7. *Let $T \in \text{ZNFT}$. If T is NFT-definable, then $T \models \mathcal{P}$.*

Proof: Let T' be an NFT equivalent to T . It is easy to turn T' into a ZNFT T'' that performs two additional backward and forward passes which output ϵ . Consider two loops (i_1, j_1) and (j_1, j_2) in a run of T'' , and let us write the output of this run as depicted on Fig. 2. These loops are also loops of T' , and thus we can define β_1 (resp. $\beta_2, \beta_3, \beta_4$ and β_5) as x_0 (resp. v_1, x_1, w_1 and x_2), and f, g as the constant mappings equal to ϵ . Hence $T'' \models \mathcal{P}$, and we conclude by Lemma 6. ■

B. From ZNFT to ϵZNFT

The goal is to devise a procedure that tests whether the first and second passes (forward and backward) of the run can be done with a single forward pass, and constructs an NFT that realizes this single forward pass. Then, in order to obtain an ϵZNFT , it suffices to replace the first pass of T by the latter NFT and add a backward pass that just comes back to the beginning of the word and outputs ϵ all the time. The procedure constructs an ϵZNFT , and tests whether it is equivalent to T . It is based on the following key property that characterizes the form of the output words of the two first passes of any ZNFT satisfying \mathcal{P} . Intuitively, when these words are long enough, they can be decomposed as words whose primitive roots are conjugate.

Definition 3 (\mathcal{P}_1 -property). *Let $T \in \text{ZNFT}$ with m states, and let $(u, v) \in R(T)$ where u has length n . Let $K = 2 \cdot o \cdot m^3 \cdot |\Sigma|$ where $o = \max\{|v| \mid (p, a, v, q, m) \in \Delta\}$. The pair (u, v) satisfies the property \mathcal{P}_1 , denoted by $(u, v) \models \mathcal{P}_1$, if for all accepting runs ρ on u , there exist a position $1 \leq \ell \leq n$ and $w, w', t_1, t_2, t_3 \in \Sigma^*$ such that $v \in w t_1 t_2^* t_3 w'$ and:*

$$\begin{aligned} \text{out}_1[1, \ell] &= w & \text{out}_2[1, \ell] &= t_3 & \text{out}_1[\ell, n] \text{out}_2[\ell, n] &\in t_1 t_2^* \\ \text{out}_3[1, n+1] &= w' & |t_i| &\leq 2K, \forall i \in \{1, 2, 3\} \end{aligned}$$

This decomposition is depicted in Fig. 4. T satisfies property \mathcal{P}_1 , denoted $T \models \mathcal{P}_1$, if all $(u, v) \in R(T)$ satisfy it.




Fig. 5. Decomposition of the two first passes of a z -motion run with loop.

Proposition 4. *Let $T \in \text{ZNFT}$. If $T \models \mathcal{P}$, then $T \models \mathcal{P}_1$.*

Proof: • If $|\text{out}_2[1, n-1]| \leq K$, then clearly, it suffices to take $\ell = n$, $t_1 = \text{out}_2[n-1, n]$, $t_2 = \varepsilon$, $t_3 = \text{out}_2[0, n-1]$, $w = \text{out}_1[1, n]$ and $w' = \text{out}_3[1, n+1]$.

• Otherwise, $|\text{out}_2[1, n-1]| > K$. Therefore u is of length $2.m^3 \cdot |\Sigma|$ at least and there exists a (non-empty) loop (i, j) in ρ . We can always choose this loop such that $|\text{out}_2[1, i]| \leq K$ and $1 \leq |\text{out}_2[i, j]| \leq K$.

The loop partitions the input and output words into factors that are depicted in Fig. 5 (only the two first passes are depicted). Formally, let $u = u_1 u_2 u_3$ such that $u_2 = u[i \dots (j-1)]$. Let $x_0 = \text{out}_1[1, i]$, $v_1 = \text{out}_1[i, j]$, $x_1 = \text{out}_1[j, n] \text{out}_2[j, n]$, $v_2 = \text{out}_2[i, j]$, $x_2 = \text{out}_1[1, i]$, $x_3 = \text{out}_3[1, i]$, $v_3 = \text{out}_3[i, j]$ and $x_4 = \text{out}_3[j, n+1]$. In particular, we have $|x_2| \leq K$, $1 \leq |v_2| \leq K$ and $x_0 v_1 x_1 v_2 x_2 x_3 v_3 x_4 \in T(u)$. Since (i, j) is a loop we also get $x_0 v_1^k x_1 v_2^k x_2 x_3 v_3^k x_4 \in T(u_1 u_2^k u_3)$ for all $k \geq 0$. We then distinguish two cases:

1) If $v_1 \neq \epsilon$. We can apply Property \mathcal{P} by taking the second loop empty. We get that for all $k \geq 0$

$$f(k) x_0 v_1^{kc+c'} x_1 v_2^{kc+c'} x_2 x_3 v_3^{kc+c'} x_4 g(k) = \beta_1 \beta_2^k \beta_3$$

where $f, g : \mathbb{N} \rightarrow \Sigma^*$, $c \in \mathbb{N}_{>0}$, $c' \in \mathbb{N}$, and $\beta_1, \beta_2, \beta_3 \in \Sigma^*$. Since the above equality holds for all $k \geq 0$, we can apply Lemma 1 and we get $\mu(v_1) \sim \mu(\beta_2)$ and $\mu(\beta_2) \sim \mu(v_2)$, and therefore $\mu(v_1) \sim \mu(v_2)$. So there exist $x, y \in \Sigma^*$ such that $v_1 \in (xy)^*$ and $v_2 \in (yx)^*$. One can show (see [24]) that $v_1 x_1 v_2 \in x(yx)^*$. Then it suffices to take $\ell = i$, $w = x_0$, $t_1 = x$, $t_2 = yx$ and $t_3 = x_2$.

2) The second case ($v_1 = \epsilon$) is more complicated as it requires to use the full Property \mathcal{P} , using two non-empty loops. First, we distinguish two cases whether $|\text{out}_1[j, n]| \leq K$ or not. For the latter case, we identify a second loop and then apply Property \mathcal{P} . ■

Construction of an ϵZNFT from a ZNFT We construct an ϵZNFT T' from a ZNFT T such that $R(T') = \{(u, v) \in R(T) \mid (u, v) \models \mathcal{P}_1\}$. Intuitively, the main idea is to perform the two first passes in a single forward pass, followed by a non-producing backward pass, and the final third pass is exactly as T does. Therefore, T' guesses the words t_1, t_2 and t_3 and makes sure that the output v is indeed of the form characterized by \mathcal{P}_1 . This can be done in a one-way fashion while simulating the forward and backward passes in parallel and by guessing non-deterministically the position ℓ . In addition, the output mechanism of T' exploits the special

form of v : the idea is to output powers of t_2 while simulating the two first passes.

First, let us describe how T' simulates the forward and backward passes in parallel during the first forward pass. It guesses both the state of the backward pass, and the current symbol (this is needed as the symbol read by the backward transition is the next symbol). The first state (q^*) guessed for the backward pass needs to be stored, as the last (forward) pass should start from q^* . The transducer can go from state (p, q, σ) to state (p', q', σ') if the current symbol is σ and there is a (forward) transition $(p, \sigma, x, p', +1)$ and a (backward) transition $(q', \sigma', y, q, -1)$. Therefore if Q is the set of states of T , T' uses, on the first pass, elements of $Q \times Q \times \Sigma$ in its states. The transducer T' can non-deterministically decide to perform the backward and non-producing backward pass whenever it is in some state (q, q, σ) and the current symbol is σ . This indeed happens precisely when the forward and backward passes are in the same state q . If the current symbol is not the last of the input word, then the whole run of T' is not a z -motion and therefore it is not accepting.

Second, we describe how the ϵZNFT T' , with the guess of t_1, t_2, t_3 , verifies during its first forward pass that the output has the expected form, and how it produces this output. During the first pass, T' can be in two modes: In mode 1 (before the guess ℓ), T' verifies that the output on the simulated backward pass is t_3 and proceeds as T in the first forward pass (it outputs what T outputs on the forward pass). Mode 2 starts when the guess ℓ has been made. In this mode, T' first outputs t_1 and then verifies that the output of the forward/backward run from and to position ℓ is of the form $t_1 t_2^*$. It can be done by using pointers on t_1 and t_2 . There are two cases (guessed by T'): either t_1 ends during the forward pass or during the backward pass (using notations of Fig.4, either t_1 is a prefix of x , or x is a prefix of t_1).

In the first case, T' needs a pointer on t_1 to make sure that the output of T in the forward pass starts with t_1 . It also needs a pointer on t_2 , initially at the end of t_2 , to make sure that the output of T on the simulated backward pass is a suffix of t_2^* (the pointer moves backward, coming back to the last position of t_2 whenever it reaches the first position of t_2). Once the verification on t_1 is done, T' starts, by using a pointer initially at the first position in t_2 , to verify that the output of T in the forward pass is a prefix of t_2^* . Once the forward and the simulated backward passes merge, the two pointers on t_2 must be at the same position, otherwise the run is rejected.

During this verification, T' also has to output a power of t_2 (remind that it has already output t_1). However the transitions of T may not output exactly one t_2 , nor a power of t_2 , but may cut t_2 before its end. Therefore T' needs another pointer h to know where it is in t_2 . Initially this pointer is at the first position of t_2 ($h = 1$). Suppose that T' simulates T using the (forward) transition $(p, \sigma, x, p', +1)$ and the (backward) transition $(q', \sigma', y, q, -1)$. If this step occurs before the end of t_1 , then T' outputs $t_2^\omega[h \dots (h + |y|)]$ (t_2^ω is the infinite concatenation of t_2), and the pointer h is updated to $1 + ((h + |y| - 1) \bmod |t_2|)$. Otherwise, T' outputs $t_2^\omega[h \dots (h + |x| + |y|)]$




Fig. 6. Decomposition of the output according to Property \mathcal{P}_2 .

and h is updated to $1 + ((h + |x| + |y| - 1) \bmod |t_2|)$.

The second case (when T' guesses that t_1 ends during the backward pass) is similar. T' has to guess exactly the position in the output where t_1 ends. On the first pass it verifies that the output is a prefix of t_1 , and on the simulated backward pass, it checks that the output is a suffix of t_2^* (and outputs as many t_2 as necessary, like before), until the end of t_1 is guessed to occur. From that moment it enters a verification mode on both passes.

The main property of this construction is that no wrong output words are produced by T' , due to the verification and the way the output words are produced, i.e. for all $(u, v) \in R(T')$, we have $(u, v) \in R(T)$.

Proposition 5. *Let $T \in \text{ZNFT}$. $R(T') = \{(u, v) \in R(T) \mid (u, v) \models \mathcal{P}_1\}$.*

Lemma 8. *Let $T \in \text{ZNFT}$. If $T \models \mathcal{P}$, then T is equivalent to the ϵZNFT T' . Moreover, the latter is decidable.*

Proof: If $T \models \mathcal{P}$, then by Proposition 4, $T \models \mathcal{P}_1$. Therefore by Proposition 5, T and T' are equivalent.

We know that $R(T') \subseteq R(T)$, and since T and T' are both functional, they are equivalent iff $\text{dom}(T) \subseteq \text{dom}(T')$. Both domains can be defined by NFAs. Those NFAs simulate the three passes in parallel and make sure that those passes define a z -motion. Therefore testing the equivalence of T and T' amounts to test the equivalence of two NFAs. ■

C. From ϵZNFT to NFT

We have seen how to go from a ZNFT to an ϵZNFT . We now briefly sketch how to go from an ϵZNFT to a (functional) NFT . Given an ϵZNFT T' , we define an $f\text{NFT}$ T'' such that T' and T'' are equivalent as soon as $T' \models \mathcal{P}$. The ideas are very similar to the previous construction therefore we do not give all the details here.

We exhibit a property on the form of output words produced by an ϵZNFT that verifies \mathcal{P} . Intuitively, apart from the beginning of the first pass, and the end of the second pass, if the two passes produce long enough outputs, then these outputs can be decomposed so as to exhibit conjugate primitive roots.

Definition 4 (\mathcal{P}_2 -property). *Let $T' \in \epsilon\text{ZNFT}$ with m states, and let $(u, v) \in R(T')$ where u has length n . Let $K = 2om^3|\Sigma|$ where $o = \max\{|v| \mid (p, a, v, q, m) \in \Delta\}$. The pair (u, v) satisfies the property \mathcal{P}_2 , denoted by $(u, v) \models \mathcal{P}_2$, if for all accepting runs ρ on u , there exist two positions*

$1 \leq \ell_1 \leq \ell_2 \leq n$ and $w, w', t_1, t_2, t_3 \in \Sigma^*$ such that:

$$\begin{array}{ll} \text{out}_1[1, \ell_1] = w & |t_i| \leq 3.K, \forall i \in \{1, 2, 3\} \\ \text{out}_3[\ell_2, n+1] = w' & |\text{out}_1[\ell_2, n]| \leq 3.K \\ \text{out}_1[\ell_1, n]\text{out}_3[1, \ell_2] \in t_1 t_2^* t_3 & |\text{out}_3[1, \ell_1]| \leq 3.K \end{array}$$

This decomposition is depicted in Fig. 6. T' satisfies property \mathcal{P}_2 , denoted $T \models \mathcal{P}_2$, if all $(u, v) \in R(T')$ satisfy it.

The proof of the following proposition uses the same structure and techniques as that of Proposition 4. Using a (long) case analysis, we identify loops in runs, and apply Property \mathcal{P} to show that output words have the expected form.

Proposition 6. *Let $T' \in \epsilon\text{ZNFT}$. If $T' \models \mathcal{P}$, then $T' \models \mathcal{P}_2$.*

We can now sketch the construction of an $f\text{NFT}$ T'' which recognizes the subrelation of T' defined as $\{(u, v) \in R(T') \mid (u, v) \models \mathcal{P}_2\}$. Again, the construction is rather similar and uses the same techniques to that of T' starting from T .

The transducer T'' simulates, in a single forward pass, the three passes of T' . Hence it also checks that the run of the ZNFT T' it simulates is a z -motion run, which is a semantic restriction of accepting runs of ZNFTs . The $f\text{NFT}$ T'' also guesses positions ℓ_1 and ℓ_2 , and uses three modes accordingly. It also guesses the words t_1 , t_2 and t_3 , and words for $\text{out}_3[1, \ell_1]$ and $\text{out}_1[\ell_2, n]$, which are all of bounded length (see Property \mathcal{P}_2). The output of T'' is produced according to the mode, using pointers to check the guesses, similarly to T' .

If all the guesses happen to be verified, it outputs the correct output word, otherwise the input word is rejected. As a consequence, T'' recognizes a subrelation of T' and thus checking the equivalence of T' and T'' amounts to checking the equivalence of their domains (as the two transducers are functional), which is decidable. From Proposition 6 we get:

Lemma 9. *Let $T' \in \epsilon\text{ZNFT}$. If $T' \models \mathcal{P}$, then T' is equivalent to the $f\text{NFT}$ T'' . Moreover, the latter property is decidable.*

Proof of Lemma 3. Lemma 7 states that if T is NFT -definable, then $T \models \mathcal{P}$. Conversely, if $T \models \mathcal{P}$, then by Lemma 8, the first construction outputs an equivalent ϵZNFT T' . By Lemma 6, we have $T' \models \mathcal{P}$. By Lemma 9, the second construction outputs an equivalent NFT T'' . Therefore T is NFT -definable by T'' . In order to decide whether $T \models \mathcal{P}$, it suffices to construct T' , check that T and T' are equivalent, and then construct T'' and check whether T' and T'' are equivalent. Both problems are decidable by Lemma 8 and 9.

V. DISCUSSION

Complexity The procedure to decide $(f\text{ZNFT}, \text{NFT})$ -definability is non-elementary exponential time and space. This is due to the ZNFT -to- NFT construction which outputs an NFT of doubly exponential size. Indeed, the first step of this construction transforms any ZNFT with n states into an ϵZNFT with at least $|\Sigma|^{4on^3|\Sigma|}$ states, as the ϵZNFT has to guess words of length $4on^3|\Sigma|$, where o is the maximal length of an output word of a transition. The ϵZNFT -to- NFT construction also outputs an exponentially bigger transducer.

Therefore the squeeze operation outputs a transducer which is doubly exponentially larger. Since this operation has to be iterated N^2 times in the worst case, where N is the number of states of the initial $f2NFT$, this leads to a non-elementary procedure. On the other hand, the best lower bound we have for this problem is PSpace (by a simple proof that reduces the emptiness problem of the intersection of n DFAs is given in [24]).

Succinctness It is already known that $2DFAs$ are exponentially more succinct than $NFAs$ [27]. Therefore this result carries over to transducers, already for transducers defining identity relations on some particular domains. However we show here a stronger result: the succinctness of $2NFTs$ also comes from the transduction part and not only from the domain part. We can indeed exhibit a family of NFT -definable transductions $(R_n)_n$ that can be defined by $2DFTs$ that are exponentially more succinct than their smallest equivalent NFT , and such that the family of languages $(\text{dom}(R_n))_n$ does not show an exponential blow up between $2DFAs$ and $NFAs$. For all $n \geq 0$, we define R_n whose domain is the set of words $\#u\#$ for all $u \in \{a, b\}^*$ of length n , and the transduction is the mirror transduction, i.e. $R_n(\#u\#) = \#\bar{u}\#$. Clearly, R_n is definable by a $2DFT$ with $O(n)$ states that counts up to n the length of the input word by a forward pass, and then mirrors it by a backward pass. It is also definable by an NFT with $O(2^n)$ states: the NFT guesses a word u of length n (so it requires $O(2^n)$ states), outputs its reverse, and then verifies that the guess was correct. It is easy to prove that any NFT defining R_n needs at least 2^n states by a pumping argument. On the other hand, the domain of R_n can be defined by a DFA with $O(n)$ states that counts the length of the input word up to n . Note that the alphabet does not depend on n .

Further Questions We have shown that $(f2NFT, NFT)$ -definability is decidable, however with a non-elementary procedure. We would like to characterize precisely the complexity of this problem. Our procedure works for functional $2NFTs$, which are equivalent to $2DFTs$. Therefore we could have done our proof directly for $2DFTs$. However (functional) non-determinism was added with no cost in the proof so we rather did it in this more general setting. The extension of our results to relations instead of functions is still open.

Our proof is an adaptation of the proof of Rabin and Scott [3] to transducers. Alternative constructions based on the proofs of Shepherdson [4] or Vardi [17], and alternative models such as streaming string transducers [10] or MSO transformations [8], [9], could lead to better complexity results or refined results. In particular, we believe that our results are highly related to the problem of minimizing the number of variables in a streaming string transducer.

Finally, we plan to study extensions of our results to infinite string transformations, defined for instance by streaming string transducers [28], and to tree transformations, following our initial motivation from XML applications.

Acknowledgements We warmly thank Sebastian Maneth and Julien Tierny for interesting discussions.

REFERENCES

- [1] J. R. Büchi, "On a decision method in restricted second order arithmetic," in *Proc. of the International Congress on Logic, Methodology, and Philosophy of Science*. Stanford University Press, 1962, pp. 1–11.
- [2] J. W. Thatcher and J. B. Wright, "Generalized finite automata theory with an application to a decision problem of second-order logic," *Mathematical Systems Theory*, vol. 2, no. 1, pp. 57–81, 1968.
- [3] M. O. Rabin and D. Scott, "Finite automata and their decision problems," *IBM Journal of Research and Development*, vol. 3, no. 2, pp. 114–125, 1959.
- [4] J. C. Shepherdson, "The reduction of two-way automata to one-way automata," *IBM Journal of Research and Development*, vol. 3, no. 2, pp. 198–200, 1959.
- [5] J. Berstel, *Transductions and context-free languages*. Teubner, 1979.
- [6] J. Sakarovich, *Elements of Automata Theory*. Cambridge University Press, 2009.
- [7] B. Courcelle, "The expression of graph properties and graph transformations in monadic second-order logic," in *Handbook of Graph Transformation*. World Scientific, 1996, vol. 1, Foundations.
- [8] —, "Monadic second-order definable graph transductions: a survey," *Theoretical Computer Science*, vol. 126, no. 1, pp. 53–75, 1994.
- [9] J. Engelfriet and H. J. Hoogeboom, "MSO definable string transductions and two-way finite-state transducers," *ACM Transactions on Computational Logic (TOCL)*, vol. 2, no. 2, pp. 216–254, 2001.
- [10] R. Alur and P. Černý, "Expressiveness of streaming string transducers," in *FSTTCS*, vol. 8. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik, 2010, pp. 1–12.
- [11] —, "Streaming transducers for algorithmic verification of single-pass list-processing programs," in *POPL*, 2011, pp. 599–610.
- [12] R. de Souza, "Uniformisation of two-way transducers," in *LATA*, ser. LNCS, vol. 7810. Springer, 2013, pp. 547–558.
- [13] E. M. Gurari and O. H. Ibarra, "A note on finite-valued and finitely ambiguous transducers," *Mathematical Systems Theory*, vol. 16, no. 1, pp. 61–66, 1983.
- [14] M.-P. Béal, O. Carton, C. Prieur, and J. Sakarovitch, "Squaring transducers: an efficient procedure for deciding functionality and sequentiality," *Theoretical Computer Science*, vol. 292, no. 1, pp. 45–63, 2003.
- [15] A. Weber and R. Klemm, "Economy of description for single-valued transducers," *Information and Computation*, vol. 118, no. 2, pp. 327–340, 1995.
- [16] K. Culik and J. Karhumäki, "The equivalence problem for single-valued two-way transducers (on NPDTOL languages) is decidable," *SIAM Journal on Computing*, vol. 16, no. 2, pp. 221–230, 1987.
- [17] M. Y. Vardi, "A note on the reduction of two-way automata to one-way automata," *Information Processing Letters*, vol. 30, no. 5, pp. 261–264, 1989.
- [18] E. Filiot, O. Gauwin, P.-A. Reynier, and F. Servais, "Streamability of nested word transductions," in *FSTTCS*, vol. 13. Schloss Dagstuhl–Leibniz-Zentrum fuer Informatik, 2011, pp. 312–324.
- [19] L. Segoufin and C. Sirangelo, "Constant-memory validation of streaming XML documents against DTDs," in *ICDT*, ser. LNCS, vol. 4353. Springer, 2007, pp. 299–313.
- [20] V. Bárány, C. Löding, and O. Serre, "Regularity problems for visibly pushdown languages," in *STACS*, ser. LNCS, vol. 3884. Springer, 2006, pp. 420–431.
- [21] O. Gauwin, J. Niehren, and S. Tison, "Queries on XML streams with bounded delay and concurrency," *Information and Computation*, vol. 209, no. 3, pp. 409–442, 2011.
- [22] O. Carton, "Two-way transducers with a two-way output tape," in *DLT*, ser. LNCS, vol. 7410. Springer, 2012, pp. 263–272.
- [23] M. Anselmo, "Two-way automata with multiplicity," in *ICALP*, ser. LNCS. Springer, 1990, vol. 443, pp. 88–102.
- [24] E. Filiot, O. Gauwin, P.-A. Reynier, and F. Servais, "From two-way to one-way finite state transducers," *CoRR*, vol. abs/1301.5197, 2013.
- [25] C. Choffrut and J. Karhumäki, *Combinatorics on words*. Springer-Verlag, 1997, vol. 1, pp. 329–438.
- [26] J. Hopcroft and J. Ullman, *Introduction to Automata Theory*. Addison-Wesley, 1979.
- [27] J.-C. Birget, "State-complexity of finite-state devices, state compressibility and incompressibility," *Mathematical Systems Theory*, vol. 26, no. 3, pp. 237–269, 1993.
- [28] R. Alur, E. Filiot, and A. Trivedi, "Regular transformations of infinite strings," in *LICS*. IEEE, 2012, pp. 65–74.