

Profile Trees for Büchi Word Automata, with Application to Determinization*

Seth Fogarty

Computer Science Department
Trinity University

Moshe Y. Vardi

Department of Computer Science
Rice University

Orna Kupferman

School of Computer Science and Engineering
Hebrew University of Jerusalem

Thomas Wilke

Institut für Informatik
Christian-Albrechts-Universität zu Kiel

The determinization of Büchi automata is a celebrated problem, with applications in synthesis, probabilistic verification, and multi-agent systems. Since the 1960s, there has been a steady progress of constructions: by McNaughton, Safra, Piterman, Schewe, and others. Despite the proliferation of solutions, they are all essentially ad-hoc constructions, with little theory behind them other than proofs of correctness. Since Safra, all optimal constructions employ trees as states of the deterministic automaton, and transitions between states are defined operationally over these trees. The operational nature of these constructions complicates understanding, implementing, and reasoning about them, and should be contrasted with complementation, where a solid theory in terms of automata run DAGs underlies modern constructions.

In 2010, we described a *profile*-based approach to Büchi complementation, where a profile is simply the history of visits to accepting states. We developed a structural theory of profiles and used it to describe a complementation construction that is deterministic in the limit. Here we extend the theory of profiles to prove that every run DAG contains a *profile tree* with at most a finite number of infinite branches. We then show that this property provides a theoretical grounding for a new determinization construction where macrostates are doubly preordered sets of states. In contrast to extant determinization constructions, transitions in the new construction are described declaratively rather than operationally.

1 Introduction

Büchi automata were introduced in the context of decision problems for second-order arithmetic [3]. These automata constitute a natural generalization of automata over finite words to languages of infinite words. Whereas a run of an automaton on finite words is accepting if the run ends in an accepting state, a run of a Büchi automaton is accepting if it visits an accepting state infinitely often.

Determinization of nondeterministic automata is a fundamental problem in automata theory, going back to [19]. Determinization of Büchi automata is employed in many applications, including synthesis of reactive systems [18], verification of probabilistic systems [4, 25], and reasoning about multi-agent systems [2]. Nondeterministic automata over finite words can be determinized with a simple, although exponential, *subset construction* [19], where a state in the determinized automaton is a set of states of the input automaton. Nondeterministic Büchi automata, on the other hand, are not closed under determinization, as deterministic Büchi automata are strictly less expressive than their nondeterministic

*Work supported in part by NSF grants CNS 1049862 and CCF-1139011, by NSF Expeditions in Computing project “ExCAPE: Expeditions in Computer Augmented Program Engineering”, by a gift from Intel, by BSF grant 9800096, and by a stipend from Trinity University. A full version, with appendices and missing proofs, is available at <http://www.cs.trinity.edu/~sfogarty/papers/gandalf13rj.pdf>

counterparts [13]. Thus, a determinization construction for Büchi automata must result in automata with a more powerful acceptance condition, such as Muller [15], Rabin [20], or parity conditions [9, 17].

The first determinization construction for Büchi automata was presented by McNaughton, with a doubly-exponential blowup [15]. In 1988, Safra introduced a singly exponential construction [20], matching the lower bound of $n^{O(n)}$ [14]. Safra’s construction encodes a state of the determinized automaton as a labeled tree, now called a *Safra tree*, of sets of states of the input Büchi automaton. Subsequently, Safra’s construction was improved by Piterman, who simplified the use of tree-node labels [17], and by Schewe, who moved the acceptance conditions from states to edges [22]. In a separate line of work, Muller and Schupp proposed in 1995 a different singly exponential determinization construction, based on *Muller-Schupp trees* [16], which was subsequently simplified by Kähler and Wilke [9].

Despite the proliferation of Büchi determinization constructions, even in their improved and simplified forms all constructions are essentially ad-hoc, with little theory behind them other than correctness proofs. These constructions rely on the encoding of determinized-automaton states as finite trees. They are operational in nature, with transitions between determinized-automaton states defined “horticulturally,” as a sequence of operations that grow trees and then prune them in various ways. The operational nature of these constructions complicates understanding, implementing, and reasoning about them [1, 23], and should be contrasted with complementation, where an elegant theory in terms of automata run DAGs underlies modern constructions [8, 11, 21]. In fact, the difficulty of determinization has motivated attempts to find determinization-free decision procedures [12] and works on determinization of fragments of LTL [10].

In a recent work [6], we introduced the notion of *profiles* for nodes in the run DAG. We began by labeling accepting nodes of the DAG by 1 and non-accepting nodes by 0, essentially recording visits to accepting states. The profile of a node is the lexicographically *maximal* sequence of labels along paths of the run DAG that lead to that node. Once profiles and a lexicographic order over profiles were defined, we removed from the run DAG edges that do not contribute to profiles. In the pruned run DAG, we focused on lexicographically maximal runs. This enabled us to define a novel, profile-based Büchi complementation construction that yields *deterministic-in-the-limit* automata: one in which every accepting run of the complementing automaton is eventually deterministic [6]. A state in the complementary automaton is a set of states of the input nondeterministic automaton, augmented with the preorder induced by profiles. Thus, this construction can be viewed as an augmented subset construction.

In this paper, we develop the theory of profiles further, and consider the equivalence classes of nodes induced by profiles, in which two nodes are in the same class if they have the same profile. We show that profiles turn the run DAG into a *profile tree*: a binary tree of bounded width over the equivalence classes. The profile tree affords us a novel singly exponential Büchi determinization construction. In this profile-based determinization construction, a state of the determinized automaton is a set of states of the input automaton, augmented with *two* preorders induced by profiles. Note that while a Safra tree is finite and encodes a single level of the run DAG, our profile tree is infinite and encodes the entire run DAG, capturing the accepting or rejecting nature of all paths. Thus, while a state in a traditional determinization construction corresponds to a Safra tree, a state in our deterministic automaton corresponds to a single level in the profile tree.

Unlike previous Büchi determinization constructions, transitions between states of the determinized automaton are defined declaratively rather than operationally. We believe that the declarative character of the new construction will open new lines of research on Büchi determinization. For Büchi complementation, the theory of run DAGs [11] led not only to tighter constructions [8, 21], but also to a rich body of work on heuristics and optimizations [5, 7]. We foresee analogous developments in research on Büchi determinization.

2 Preliminaries

This section introduces the notations and definitions employed in our analysis.

2.1 Relations on Sets

Given a set R , a binary relation \leq over R is a *preorder* if \leq is reflexive and transitive. A *linear preorder* relates every two elements: for every $r_1, r_2 \in R$ either $r_1 \leq r_2$, or $r_2 \leq r_1$, or both. A relation is *antisymmetric* if $r_1 \leq r_2$ and $r_2 \leq r_1$ implies $r_1 = r_2$. A preorder that is antisymmetric is a *partial order*. A linear partial order is a *total order*. Consider a partial order \leq . If for every $r \in R$, the set $\{r' \mid r' \leq r\}$ of smaller elements is totally ordered by \leq , then we say that \leq is a *tree order*. The equivalence class of $r \in R$ under \leq , written $[r]$, is $\{r' \mid r' \leq r \text{ and } r \leq r'\}$. The equivalence classes under a linear preorder form a totally ordered partition of R . Given a set R and linear preorder \leq over R , define the minimal elements of R as $\min_{\leq}(R) = \{r_1 \in R \mid r_1 \leq r_2 \text{ for all } r_2 \in R\}$. Note that $\min_{\leq}(R)$ is either empty or an equivalence class under \leq . Given a non-empty set R and a total order \leq , we instead define $\min_{\leq}(R)$ as the the unique minimal element of R .

Given two finite sets R and R' where $|R| \leq |R'|$, a linear preorder \leq over R , and a total order $<'$ over R' , define the $\langle \leq, <' \rangle$ -*minjection* from R to R' to be the function mj that maps all the elements in the k -th equivalence class of R to the k -th element of R' . The number of equivalence classes is at most $|R|$, and thus at most $|R'|$. If \leq is also a total order, then the $\langle \leq, <' \rangle$ -minjection is also an injection.

Example 2.1. Let $R = \mathbb{Q}$ and $R' = \mathbb{Z}$ be the sets of rational numbers and integers, respectively. Define the linear preorder \leq_1 over \mathbb{Q} by $x \leq_1 x'$ iff $\lfloor x \rfloor \leq \lfloor x' \rfloor$, and the total order $<_2$ over \mathbb{Z} by $x <_2 x'$ if $x < x'$. Then, the $\langle \leq_1, <_2 \rangle$ -minjection from \mathbb{Q} to \mathbb{Z} maps a rational number x to $\lfloor x \rfloor$.

2.2 ω -Automata

A *nondeterministic ω -automaton* is a tuple $\mathcal{A} = \langle \Sigma, Q, Q^{\text{in}}, \rho, \alpha \rangle$, where Σ is a finite alphabet, Q is a finite set of states, $Q^{\text{in}} \subseteq Q$ is a set of initial states, $\rho: Q \times \Sigma \rightarrow 2^Q$ is a nondeterministic transition relation, and α is an acceptance condition defined below. An automaton is *deterministic* if $|Q^{\text{in}}| = 1$ and, for every $q \in Q$ and $\sigma \in \Sigma$, we have $|\rho(q, \sigma)| = 1$. For a function $\delta: Q \times \Sigma \rightarrow 2^Q$, we lift δ to sets R of states in the usual fashion: $\delta(R, \sigma) = \bigcup_{r \in R} \delta(r, \sigma)$. Further, we define the inverse of δ , written δ^{-1} , to be $\delta^{-1}(r, \sigma) = \{q \mid r \in \delta(q, \sigma)\}$.

A *run* of an ω -automaton \mathcal{A} on a word $w = \sigma_0 \sigma_1 \dots \in \Sigma^\omega$ is an infinite sequence of states $q_0, q_1, \dots \in Q^\omega$ such that $q_0 \in Q^{\text{in}}$ and, for every $i \geq 0$, we have that $q_{i+1} \in \rho(q_i, \sigma_i)$. Correspondingly, a *finite run* of \mathcal{A} to q on $w = \sigma_0 \dots \sigma_{n-1} \in \Sigma^*$ is a finite sequence of states p_0, \dots, p_n such that $p_0 \in Q^{\text{in}}$, $p_n = q$, and for every $0 \leq i < n$ we have $p_{i+1} \in \rho(p_i, \sigma_i)$.

The acceptance condition α determines if a run is *accepting*. If a run is not accepting, we say it is *rejecting*. A word $w \in \Sigma^\omega$ is accepted by \mathcal{A} if there exists an accepting run of \mathcal{A} on w . The words accepted by \mathcal{A} form the *language* of \mathcal{A} , denoted by $L(\mathcal{A})$. For a *Büchi automaton*, the acceptance condition is a set of states $F \subseteq Q$, and a run q_0, q_1, \dots is accepting iff $q_i \in F$ for infinitely many i 's. For convenience, we assume $Q^{\text{in}} \cap F = \emptyset$. For a *Rabin automaton*, the acceptance condition is a sequence $\langle G_0, B_0 \rangle, \dots, \langle G_k, B_k \rangle$ of pairs of sets of states. Intuitively, the sets G are “good” conditions, and the sets B are “bad” conditions. A run q_0, q_1, \dots is accepting iff there exists $0 \leq j \leq k$ so that $q_i \in G_j$ for infinitely many i 's, while $q_i \in B_j$ for only finitely many i 's. Our focus in this paper is on nondeterministic Büchi automata on words (NBW) and deterministic Rabin automata on words (DRW).

2.3 Safra's Determinization Construction

This section presents Safra's determinization construction, using the exposition in [17]. Safra's construction takes an NBW and constructs an equivalent DRW. Intuitively, a state in this construction is a tree of subsets. Every node in the tree is labeled by the states it follows. The label of a node is a strict superset of the union of labels of its descendants, and the labels of siblings are disjoint. Children of a node are ordered by "age". Let $\mathcal{A} = \langle \Sigma, Q, Q^{in}, \rho, F \rangle$ be an NBW, $n = |Q|$, and $V = \{0, \dots, n-1\}$.

Definition 2.2. [17] A *Safra tree* over \mathcal{A} is a tuple $t = \langle N, r, p, \psi, l, G, B \rangle$ where:

- $N \subseteq V$ is a set of nodes.
- $r \in N$ is the root node.
- $p : (N \setminus \{r\}) \rightarrow N$ is the parent function over $N \setminus \{r\}$.
- ψ is a partial order defining 'older than' over siblings.
- $l : N \rightarrow 2^Q$ is a labeling function from nodes to non-empty sets of states. The label of every node is a proper superset of the union of the labels of its sons. The labels of two siblings are disjoint.
- $G, B \subseteq V$ are two disjoint subsets of V .

The only way to move from one Safra tree to the next is through a sequence of "horticultural" operations, growing the tree and then pruning it to ensure that the above invariants hold.

Definition 2.3. Define the DRW $D^S(\mathcal{A}) = \langle \Sigma, Q_S, \rho_S, t_0, \alpha \rangle$ where:

- Q_S is the set of Safra trees over \mathcal{A} .
- $t_0 = \langle \{0\}, 0, \emptyset, \emptyset, l_0, \emptyset, \{1, \dots, n-1\} \rangle$ where $l_0(0) = Q^{in}$
- For $t = \langle N, r, p, \psi, l, G, B \rangle \in Q_S$ and $\sigma \in \Sigma$, the tree $t' = \rho_S(t, \sigma)$ is the result of the following sequence of operations. We temporarily use a set V' of names disjoint from V . Initially, let $t' = \langle N', r', p', \psi', l', G', B' \rangle$ where $N' = N$, $r' = r$, $p' = p$, $\psi' = \psi$, l' is undefined, and $G' = B' = \emptyset$.
 - (1) For every $v \in N'$, let $l'(v) = \rho(l(v), \sigma)$.
 - (2) For every $v \in N'$ such that $l'(v) \cap F \neq \emptyset$, create a new node $v' \in V'$ where: $p(v') = v$; $l'(v') = l'(v) \cap F$; and for every $w' \in V'$ where $p(w') = v$ add (w', v') to ψ .
 - (3) For every $v \in N'$ and $q \in l'(v)$, if there is a $w \in N'$ such that $(w, v) \in \psi$ and $q \in l'(w)$, then remove q from $l'(v)$ and, for every descendant v' of v , remove q from $l'(v')$.
 - (4) Remove all nodes with empty labels.
 - (5) For every $v \in N'$, if $l'(v) = \bigcup \{l'(v') \mid p'(v') = v\}$ remove all children of v , add v to G .
 - (6) Add all nodes in $V \setminus N'$ to B .
 - (7) Change the nodes in V' to unused nodes in V .
- $\alpha = \{\langle G_0, B_0 \rangle, \dots, \langle G_{n-1}, B_{n-1} \rangle\}$, where:
 - $G_i = \{\langle N, r, p, \psi, l, G, B \rangle \in Q_S \mid i \in G\}$
 - $B_i = \{\langle N, r, p, \psi, l, G, B \rangle \in Q_S \mid i \in B\}$

Theorem 2.4. [20] For an NBW \mathcal{A} with n states, $L(D^S(\mathcal{A})) = L(\mathcal{A})$ and $D^S(\mathcal{A})$ has $n^{O(n)}$ states.

3 From Run DAGs to Profile Trees

In this section, we present a framework for simultaneously reasoning about all runs of a Büchi automaton on a word. We use a DAG to encode all possible runs, and give each node in this DAG a profile based on its history. The lexicographic order over profiles induces a preorder \preceq_i over the nodes on level i of the run DAG. Using \preceq_i , we prune the edges of the run DAG, and derive a binary tree of bounded width. Throughout this paper we fix an NBW $\mathcal{A} = \langle \Sigma, Q, Q^{in}, \rho, F \rangle$ and an infinite word $w = \sigma_0 \sigma_1 \dots$.

3.1 Run DAGs and Profiles

The runs of \mathcal{A} on w can be arranged in an infinite DAG $G = \langle V, E \rangle$, where

- $V \subseteq Q \times \mathbb{N}$ is such that $\langle q, i \rangle \in V$ iff there is a finite run of \mathcal{A} to q on $\sigma_0 \cdots \sigma_{i-1}$.
- $E \subseteq \bigcup_{i \geq 0} (Q \times \{i\}) \times (Q \times \{i+1\})$ is such that $E(\langle q, i \rangle, \langle q', i+1 \rangle)$ iff $\langle q, i \rangle \in V$ and $q' \in \rho(q, \sigma_i)$.

The DAG G , called the *run DAG of \mathcal{A} on w* , embodies all possible runs of \mathcal{A} on w . We are primarily concerned with *initial paths* in G : paths that start in $Q^{in} \times \{0\}$. A node $\langle q, i \rangle$ is an *F-node* if $q \in F$, and a path in G is *accepting* if it is both initial and contains infinitely many *F-nodes*. An accepting path in G corresponds to an accepting run of \mathcal{A} on w . If G contains an accepting path, we say that G is *accepting*; otherwise it is *rejecting*. Let G' be a sub-DAG of G . For $i \geq 0$, we refer to the nodes in $Q \times \{i\}$ as *level i* of G' . Note that a node on level $i+1$ has edges only from nodes on level i . We say that G' has *bounded width of degree c* if every level in G' has at most c nodes. By construction, G has bounded width of degree $|Q|$.

Consider the run DAG $G = \langle V, E \rangle$ of \mathcal{A} on w . Let $f: V \rightarrow \{0, 1\}$ be such that $f(\langle q, i \rangle) = 1$ if $q \in F$ and $f(\langle q, i \rangle) = 0$ otherwise. Thus, f labels *F-nodes* by 1 and all other nodes by 0. The *profile* of a path in G is the sequence of labels of nodes in the path. We define the profile of a node to be the lexicographically maximal profile of all initial paths to that node. Formally, the profile of a finite path $b = v_0, v_1, \dots, v_n$ in G , written h_b , is $f(v_0)f(v_1) \cdots f(v_n)$, and the profile of an infinite path $b = v_0, v_1, \dots$ is $h_b = f(v_0)f(v_1) \cdots$. Finally, the profile of a node v , written h_v , is the lexicographically maximal element of $\{h_b \mid b \text{ is an initial path to } v\}$.

The lexicographic order of profiles induces a linear preorder over nodes on every level of G . We define a sequence of linear preorders \preceq_i over the nodes on level i of G as follows. For nodes u and v on level i , let $u \prec_i v$ if $h_u < h_v$, and $u \approx_i v$ if $h_u = h_v$. We group nodes by their equivalence classes under \preceq_i . Since the final element of a node's profile is 1 if and only if the node is an *F-node*, all nodes in an equivalence class agree on membership in F . Call an equivalence class an *F-class* when all members are *F-nodes*, and a *non-F-class* when none of its members are *F-nodes*. When a state can be reached by two finite runs, a node will have multiple incoming edges in G . We now remove from G all edges that do not contribute to profiles. Formally, define the pruned run DAG $G' = \langle V, E' \rangle$ where $E' = \{\langle u, v \rangle \in E \mid \text{for every } u' \in V, \text{ if } \langle u', v \rangle \in E \text{ then } u' \preceq_{|u|} u\}$. Note that the set of nodes in G and G' are the same, and that an edge is removed from E' only when there is another edge to its destination.

Lemma 3.1 states that, as we have removed only edges that do not contribute to profiles, nodes derive their profiles from their parents in G' .

Lemma 3.1. [6] *For two nodes u and u' in V , if $\langle u, u' \rangle \in E'$, then $h_{u'} = h_u 0$ or $h_{u'} = h_u 1$.*

While nodes with different profiles can share a child in G , Lemma 3.2 precludes this in G' .

Lemma 3.2. *Consider nodes u and v on level i of G' and nodes u' and v' on level $i+1$ of G' . If $\langle u, u' \rangle \in E'$, $\langle v, v' \rangle \in E'$, and $u' \approx_{i+1} v'$, then $u \approx_i v$.*

Proof: Since $u' \approx_{i+1} v'$, we have $h_{u'} = h_{v'}$. If u' is an *F-node*, then v' is an *F-node* and the last letter in both $h_{u'}$ and $h_{v'}$ is 1. By Lemma 3.1 we have $h_u 1 = h_{u'} = h_{v'} = h_v 1$. If u' and v' are *non-F-nodes*, then we have $h_u 0 = h_{u'} = h_{v'} = h_v 0$. In either case, $h_u = h_v$ and $u \approx_i v$. \square

Finally, we have that G' captures the accepting or rejecting nature of G . This result was employed to provide deterministic-in-the-limit complementation in [6]

Theorem 3.3. [6] *The pruned run DAG G' of an NBW \mathcal{A} on a word w is accepting iff \mathcal{A} accepts w .*

3.2 The Profile Tree

Using profiles, we define the *profile tree* T , which we show to be a binary tree of bounded width that captures the accepting or rejecting nature of the pruned run DAG G' . The nodes of T are the equivalence classes $\{[u] \mid u \in V\}$ of $G' = \langle V, E' \rangle$. To remove confusion, we refer to the nodes of T as *classes* and use u and U and W for classes in T , while reserving u and v for nodes in G or G' . The edges in T are induced by those in G' as expected: for an edge $\langle u, v \rangle \in E'$, the class $[v]$ is the child of $[u]$ in T . A class W is a *descendant* of a class U if there is a, possibly empty, path from U to W .

Theorem 3.4. *The profile tree T of an n -state NBW A on an infinite word w is a binary tree whose width is bounded by n .*

Proof: That T has bounded width follows from the fact that a class on level i contains at least one node on level i of G , and G is of bounded width of degree $|Q|$. To prove that every class has one parent, for a class W let $U = \{u \mid \text{there is } v \in W \text{ such that } \langle u, v \rangle \in E'\}$. Lemma 3.2 implies that U is an equivalence class, and is the sole parent of W . To show that T has a root, note that as $Q^{\text{in}} \cap F = \emptyset$, all nodes on the first level of G have profile 0, and every class descends from this class of nodes with profile 0. Finally, as noted Lemma 3.1 entails that a class U can have at most two children: the class with profile $h_U 1$, and the class with profile $h_U 0$. Thus T is binary. \square

A *branch* of T is a finite or infinite initial path in T . Since T is a tree, two branches share a prefix until they *split*. An infinite branch is *accepting* if it contains infinitely many F -classes, and *rejecting* otherwise. An infinite rejecting branch must reach a suffix consisting only of non- F -classes. A class U is called *finite* if it has finitely many descendants, and a finite class U *dies out* on level k if it has a descendant on level $k - 1$, but none on level k . Say T is *accepting* if it contains an accepting branch, and *rejecting* if all branches are rejecting.

As all members of a class share a profile, we define the profile h_U of a class U to be h_u for some node $u \in U$. We extend the function f to classes, so that $f(U) = 1$ if U is an F -class, and $f(U) = 0$ otherwise. We can then define the profile of an infinite branch $b = U_0, U_1, \dots$ to be $h_b = f(U_0)f(U_1)\dots$. For two classes U and W on level i , we say that $U \prec_i W$ if $h_U < h_W$. For two infinite branches b and b' , we say that $b \prec b'$ if $h_b < h_{b'}$. Note that \prec_i is a total order over the classes on level i , and that \prec is a total order over the set of infinite branches.

As proven above, a class U has at most two children: the class of F -nodes with profile $h_U 1$, and the class of non- F -nodes with profile $h_U 0$. We call the first class the F -child of U , and the second class the non- F -child of U . While the DAG G' can have infinitely many infinite branches, bounding the width of a tree also bounds the number of infinite branches it may have.

Corollary 3.5. *The profile tree T of an NBW A on an infinite word w has a finite number of infinite branches.*

Example 3.6. Consider, for example, the NBW in Figure 1.(a) and the first four levels of a tree of equivalence classes in Figure 1.(b). This tree corresponds to all runs of the NBW on the word ab^ω . There is only one infinite branch, $\{\langle q, 0 \rangle\}, \{\langle p, 1 \rangle\}, \{\langle p, 2 \rangle\}, \dots$, which is accepting. The set of labels and the global labeling gl are explained below, in Section 4.1.

We conclude this section with Theorem 3.7, which enables us to reduce the search for an accepting path in G' to a search for an accepting branch in T .




Figure 1: An automaton and tree of classes. Each class is a singleton set, brackets are omitted for brevity. F -classes are circled twice. Each class is labeled with its profile h , as well as the set `labels` and the global label gl as defined in Section 4.1.

Theorem 3.7. *The profile tree T of an NBW \mathcal{A} on an infinite word w is accepting iff \mathcal{A} accepts w .*

Proof: If $w \in L(\mathcal{A})$, then by Theorem 3.3 we have that G' contains an accepting path u_0, u_1, \dots . This path gives rise to an accepting branch $[u_0], [u_1], \dots$ in T . In the other direction, if T has an accepting branch U_0, U_1, \dots , consider the infinite subgraph of G' consisting only of the nodes in U_i , for $i > 0$. For every $i > 0$ there exists $u_i \in U_i$ and $u_{i+1} \in U_{i+1}$ so that $E'(u_i, u_{i+1})$. Because no node is orphaned in G' , Lemma 3.2 implies that every node in U_{i+1} has a parent in U_i , thus this subgraph is connected. As each node has degree of at most n , König's Lemma implies that there is an infinite initial path u_0, u_1, \dots through this subgraph. Further, at every level i where U_i is an F -class, we have that $u_i \in F$, and thus this path is accepting and $w \in L(\mathcal{A})$. \square

4 Labeling

In this section we present a method of deterministically labeling the classes in T with integers, so we can determine if T is accepting by examining the labels. Each label m represents the proposition that the lexicographically minimal infinite branch through the first class labeled with m is accepting. On each level we give the label m to the lexicographically minimal descendant, on any branch, of this first class labeled with m . We initially allow the use of global information about T and an unbounded number of labels. We then show how to determine the labeling using bounded information about each level of T , and how to use a fixed set of labels.

4.1 Labeling T

We first present a labeling that uses an unbounded number of labels and global information about T . We call this labeling the *global labeling*, and denote it with gl . For a class U on level i of T , and a class W on level j , we say that W is *before* U if $j < i$ or $j = i$ and $W \prec_i U$. For each label m , we refer to the first class labeled m as $\text{first}(m)$. Formally, $U = \text{first}(m)$ if U is labeled m and, for all classes W before U , the label of W is not m . We define the labeling function gl inductively over the nodes of T . For the initial class $U_0 = \{\langle q, 0 \rangle \mid q \in Q^m\}$ with profile 0, let $gl(U_0) = 0$.

Each label m follows the lexicographically minimal child of $\text{first}(m)$ on every level. When a class with label m has two children, we are not certain which, if either, is part of an infinite branch. We are thus conservative, and follow the non- F -child. If the non- F -child dies out, we revise our guess and move to a descendant of the F -child. For a label m and level i , let the *lexicographically minimal descendant* of m on level i , written $\text{lmd}(m, i)$, be $\min_{\preceq}(\{W \mid W \text{ is a descendant of } \text{first}(m) \text{ on level } i\})$: the class with the minimal profile among all the descendants of $\text{first}(m)$ on level i . For a class U on level i , define $\text{labels}(U) = \{m \mid U = \text{lmd}(m, i)\}$ as the set of valid labels for U . When labelling U , if U has more than one valid label, we give it the smallest label, which corresponds to the earliest ancestor. If $\text{labels}(U)$ is empty, U is given an unused label one greater than the maximum label occurring earlier in T .

Definition 4.1. $gl(U) = \begin{cases} \min(\text{labels}(U)) & \text{if } \text{labels}(U) \neq \emptyset, \\ \max(\{gl(W) \mid W \text{ is before } U\}) + 1 & \text{if } \text{labels}(U) = \emptyset. \end{cases}$

Lemma 4.2 demonstrates that every class on a level gets a unique label, and that despite moving between nephews the labeling adheres to branches in the tree.

Lemma 4.2. *For classes U and W on level i of T , it holds that:*

- (1) *If $U \neq W$ then $gl(U) \neq gl(W)$.*
- (2) *U is a descendant of $\text{first}(gl(U))$.*
- (3) *If U is a descendant of $\text{first}(gl(W))$, then $W \preceq_i U$. Consequently, if $U \prec_i W$, then U is not a descendant of $\text{first}(gl(W))$.*
- (4) *$\text{first}(gl(U))$ is the root or an F -class with a sibling.*
- (5) *If $U \neq \text{first}(gl(U))$, then there is a class on level $i - 1$ that has label $gl(U)$.*
- (6) *If $gl(U) < gl(W)$ then $\text{first}(gl(U))$ is before $\text{first}(gl(W))$.*

As stated above, the label m represents the proposition that the lexicographically minimal *infinite* branch going through $\text{first}(m)$ is accepting. Every time we pass through an F -child, this is evidence towards this proposition. Recall that when a class with label m has two children, we initially follow the non- F -child. If the non- F -child dies out, we revise our guess and move to a descendant of the F -child. Thus revising our guess indicates that at an earlier point the branch did visit an F -child, and also provides evidence towards this proposition. Formally, we say that a label m is *successful on level i* if there is a class U on level $i - 1$ and a class U' on level i such that $gl(U) = gl(U') = m$, and either U' is the F -child of U , or U' is not a child of U at all.

Example 4.3. In Figure 1.(b), the only infinite branch $\{\langle q, 0 \rangle\}, \{\langle p, 1 \rangle\}, \dots$ is accepting. At level 0 this branch is labeled with 0. At each level $i > 0$, we conservatively assume that the infinite branch beginning with $\langle q, 0 \rangle$ goes through $\{\langle q, i \rangle\}$, and thus label $\{\langle q, i \rangle\}$ by 0. As $\{\langle q, i \rangle\}$ is proven finite on level $i + 1$, we revise our assumption and continue to follow the path through $\{\langle p, i \rangle\}$. Since $\{\langle p, i \rangle\}$ is an F -class, the label 0 is successful on every level $i + 1$. Although the infinite branch is not labeled 0 after the first level, the label 0 asymptotically approaches the infinite branch, checking along the way that the branch is lexicographically minimal among the infinite branches through the root.

Theorem 4.4 demonstrates that the global labeling captures the accepting or rejecting nature of T . Intuitively, at each level the class U with label m is on the lexicographically minimal branch from $\text{first}(m)$. If U is on the lexicographically minimal *infinite* branch from $\text{first}(m)$, the label m is waiting for the branch to next reach an F -class. If U is not on the lexicographically minimal infinite branch from $\text{first}(m)$, then U is finite and m is waiting for U to die out.

Theorem 4.4. *A profile tree T is accepting iff there is a label m that is successful infinitely often.*

Proof: In one direction, assume there is a label m that is successful infinitely often. The label m can be successful only when it occurs, and thus m occurs infinitely often, $\text{first}(m)$ has infinitely many descendants, and there is at least one infinite branch through $\text{first}(m)$. Let $b = U_0, U_1, \dots$ be the lexicographically minimal infinite branch that goes through $\text{first}(m)$. We demonstrate that b cannot have a suffix consisting solely of non- F -classes, and therefore is an accepting branch. By way of contradiction, assume there is an index j so that for every $k > j$, the class U_k is a non- F -class. By Lemma 4.2.(4), $\text{first}(m)$ is an F -class or the root and thus occurs before level j .

Let $\mathcal{U} = \{W \mid W \prec_j U_j, W \text{ is a descendant of } \text{first}(m)\}$ be the set of descendants of $\text{first}(m)$, on level j , that are lexicographically smaller than U_j . Since b is the lexicographically minimal infinite branch through $\text{first}(m)$, every class in \mathcal{U} must be finite. Let $j' \geq j$ be the level at which the last class in \mathcal{U} dies out. At this point, $U_{j'}$ is the lexicographically minimal descendant of $\text{first}(m)$. If $gl(U_{j'}) \neq m$, then there is no class on level j' with label m , and, by Lemma 4.2.(5), m would not occur after level j' . Since m occurs infinitely often, it must be that $gl(U_{j'}) = m$. On every level $k > j'$, the class U_k is a non- F -child, and thus U_k is the lexicographically minimal descendant of $U_{j'}$ on level k and so $gl(U_k) = m$. This entails m cannot be not successful after level j' , and we have reached a contradiction. Therefore, there is no such rejecting suffix of b , and b must be an accepting branch.

In the other direction, if there is an infinite accepting branch, then let $b = U_0, U_1, \dots$ be the lexicographically minimal infinite accepting branch. Let B' be the set of infinite branches that are lexicographically smaller than b . Every branch in B' must be rejecting, or b would not be the minimal infinite accepting branch. Let j be the first index after which the last branch in B' splits from b . Note that either $j = 0$, or U_{j-1} is part of an infinite rejecting branch $U_0, \dots, U_{j-1}, W_j, W_{j+1}, \dots$ smaller than b . In both cases, we show that U_j is the first class for a new label m that occurs on every level $k > j$ of T .

If $j = 0$, then let $m = 0$. As m is the smallest label, and there is a descendant of U_j on every level of T , it holds that m will occur on every level. In the second case, where $j > 0$, then W_j must be the non- F -child of U_{j-1} , and so U_j is the F -child. Thus, U_j is given a new label m where $U_j = \text{first}(m)$. For every label $m' < m$ and level $k > j$, since for every descendant U' of U_j it holds that $W_k \preceq_k U'$, it cannot be that $\text{lmd}(m', k)$ is a descendant of U_j . Thus, on every level $k > j$, the lexicographically minimal descendant of U_j will be labeled m , and m occurs on every level of T .

We show that m is successful infinitely often by defining an infinite sequence of levels, j_0, j_1, j_2, \dots so that m is successful on j_i for all $i > 0$. As a base case, let $j_0 = j$. Inductively, at level j_i , let U' be the class on level j_i labeled with m . We have two cases. If $U' \neq U_{j_i}$, then as all infinite branches smaller than b have already split from b , U' must be finite in T . Let j_{i+1} be the level at which U' dies out. At level j_{i+1} , m will return to a descendant of U_{j_0} , and m will be successful. In the second case, $U' = U_{j_i}$. Take the first $k > j_i$ so that U_k is an F -class. As b is an accepting branch, such a k must exist. As every class between U_j and U_k is a non- F -class, $gl(U_{k-1}) = m$. If U_k is the only child of U_{k-1} then let $j_{i+1} = k$: since $gl(U_k) = m$ and U_k is not the non- F -child of U_{k-1} , it holds that m is successful on level k . Otherwise let U'_k be the non- F -child of U_{k-1} , so that $gl(U'_k) = m$. Again, U'_k is finite. Let j_{i+1} be the level at which U'_k dies out. At level j_{i+1} , the label m will return to a descendant of U_k , and m will be successful. \square

4.2 Determining Lexicographically Minimal Descendants

Recall that the definition of the labeling gl involves the computation of $\text{lmd}(m, i)$, the class with the minimal profile among all the descendants of $\text{first}(m)$ on level i . Finding $\text{lmd}(m, i)$ requires knowing the descendants of $\text{first}(m)$ on level i . We show how to store this information with a partial order, denoted \leq_i , over classes that tracks which classes are minimal cousins of other classes. Using this partial order, we can determine the class $\text{lmd}(m, i+1)$ for every label m that occurs on level i , using only information about levels i and $i+1$ of T . Lemma 4.2.(5) implies that we can safely restrict ourselves to labels that occur on level i .

Definition 4.5. For two classes U and W on level i of T , say that U is a *minimal cousin* of W , written $U \leq_i W$, iff W is a descendant of $\text{first}(gl(U))$. Say $U <_i W$ when $U \leq_i W$ and $U \neq W$.

For a label m and level i , we can determine $\text{lmd}(m, i+1)$ given only the classes on levels i and $i+1$ and the partial order $<_i$. Let U be a class U on level i . Because labels can move between branches, the minimal descendant of $\text{first}(gl(U))$ on level $i+1$ may be a nephew of U , not necessarily a direct descendant. Define the \leq_i -nephew of U as $\text{neph}_i(U) = \min_{\leq_{i+1}}(\{W' \mid W \text{ is the parent of } W' \text{ and } U \leq_i W\})$.

Lemma 4.6. For a class U on level i of T , it holds that $\text{lmd}(gl(U), i+1) = \text{neph}_i(U)$.

Proof: We prove that $\{W' \mid W \text{ is the parent of } W' \text{ and } U \leq_i W\}$ contains every descendant of $\text{first}(gl(U))$ on level $i+1$, and thus that its minimal element is $\text{lmd}(gl(U), i+1)$. Let W' be a class on level $i+1$, with parent W on level i . If $U \leq_i W$, then W is a descendant of $\text{first}(gl(U))$ and W' is likewise a descendant of $\text{first}(gl(U))$. Conversely, as $gl(U)$ exists on level i , if W' is a descendant of $\text{first}(gl(U))$, then its parent W must also be a descendant of $\text{first}(gl(U))$ and $U \leq_i W$. \square

By using neph_i , we can in turn define the set of valid labels for a class U' on level $i+1$. Formally, define the \leq_i -uncles of U' as $\text{unc}_i(U') = \{U \mid U' = \text{neph}_i(U)\}$. Lemma 4.7 demonstrates how unc_i corresponds to labels.

Lemma 4.7. Consider a class U' on level $i+1$. The following hold:

- (1) $\text{labels}(U') \cap \{gl(W) \mid W \text{ on level } i\} = \{gl(U) \mid U \in \text{unc}_i(U')\}$.
- (2) $\text{labels}(U') = \emptyset$ iff $\text{unc}_i(U') = \emptyset$.

Proof:

- (1) Let U be a class on level i . By definition, $gl(U) \in \text{labels}(U')$ iff $U' = \text{lmd}(gl(U), i+1)$. By Lemma 4.6, it holds that $\text{lmd}(gl(U), i+1) = \text{neph}_i(U)$. By the definition of unc_i , we have that $U' = \text{neph}_i(U)$ iff $U \in \text{unc}_i(U')$. Thus every label in $\text{labels}(U')$ that occurs on level i labels some node in $\text{unc}_i(U')$.
- (2) If $\text{unc}_i(U') \neq \emptyset$, then part (1) implies $\text{labels}(U') \neq \emptyset$. In other direction, let $m = \min(\text{labels}(U'))$. By Lemma 4.2.(5), there is a U on level i so that $gl(U) = m$, and by part (1) $U \in \text{unc}_i(U')$. \square

Finally, we demonstrate how to compute \leq_{i+1} only using information about the level i of T and the labeling for level $i+1$. As the labeling depends only on \leq_i , this removes the final piece of global information used in defining gl .

Lemma 4.8. Let U' and W' be two classes on level $i+1$ of T , where $U' \neq W'$. Let W be the parent of W' . We have that $U' \leq_{i+1} W'$ iff there exists a class U on level i so that $gl(U) = gl(U')$ and $U \leq_i W$.

Proof: If there is no class U on level i so that $gl(U) = gl(U')$, then $U' = \text{first}(gl(U'))$. Since W' is not a descendant of U' , it cannot be that $U' \leq_{i+1} W'$. If such a class U exists, then $U \leq_i W$ iff W is a descendant of $\text{first}(gl(U))$, which is true iff W' is a descendant of $\text{first}(gl(U'))$: the definition of $U' \leq_{i+1} W'$. \square

4.3 Reusing Labels

As defined, the labeling function gl uses an unbounded number of labels. However, as there are at most $|Q|$ classes on a level, there are at most $|Q|$ labels in use on a level. We can thus use a fixed set of labels by reusing dead labels. For convenience, we use $2|Q|$ labels, so that we never need reuse a label that was in use on the previous level. The full version demonstrates how to use $|Q| - 1$ labels. There are two barriers to reusing labelings. First, we can no longer take the numerically minimal element of $\text{labels}(U)$ as the label of U . Instead, we calculate which label is the oldest through \preceq . Second, we must ensure that a label that is good infinitely often is not reused infinitely often. To do this, we introduce a Rabin condition to reset each label before we reuse it.

We inductively define a sequence of labelings, l_i , each from the i th level of T to $\{0, \dots, 2|Q|\}$. As a base case, there is only one equivalence class U on level 0 of T , and define $l_0(U) = 0$. Inductively, given the set of classes \mathcal{U}_i on level i , the function l_i , and the set of classes \mathcal{U}_{i+1} on level $i+1$, we define l_{i+1} as follows. Define the set of unused labels $\text{FL}(l_i)$ to be $\{m \mid m \text{ is not in the range of } l_i\}$. As T has bounded width $|Q|$, we have that $|Q| \leq |\text{FL}(l_i)|$. Let mj_{i+1} be the $\langle \preceq_{i+1}, < \rangle$ -minjection from $\{U' \text{ on level } i+1 \mid \text{unc}_i(U') = \emptyset\}$ to $\text{FL}(l_i)$. Finally, define the labeling l_{i+1} as

$$l_{i+1}(U') = \begin{cases} l_i(\min_{\preceq_i}(\text{unc}_i(U'))) & \text{if } \text{unc}_i(U') \neq \emptyset, \\ \text{mj}_{i+1}(U') & \text{if } \text{unc}_i(U') = \emptyset. \end{cases}$$

Because we are reusing labels, we need to ensure that a label that is good infinitely often is not reused infinitely often. Say that a label m is *bad* in l_i if $m \notin \text{FL}(l_{i-1})$, but $m \in \text{FL}(l_i)$. We say that a label m is *good* in l_i if there is a class U on level $i-1$ and a class U' on level i such that $l_{i-1}(U) = l_i(U') = m$ and U' is either the F -child of U or is not a child of U at all.

Theorem 4.9 demonstrates that the Rabin condition of a label being good infinitely often, but bad only finitely often, is a necessary and sufficient condition to T being accepting. The proof, omitted for brevity, associates each label m in gl with the label $l_i(\text{first}(m))$.

Theorem 4.9. *A profile tree T is accepting iff there is a label m where $\{i \mid m \text{ is bad in } l_i\}$ is finite, and $\{i \mid m \text{ is good in } l_i\}$ is infinite.*


5 A New Determinization Construction for Büchi Automata

In this section we present a determinization construction for \mathcal{A} based on the profile tree T . For clarity, we call the states of our deterministic automaton *macrostates*.

Definition 5.1. Macrostates over \mathcal{A} are six-tuples $\langle S, \preceq, l, \leq, G, B \rangle$ where:

- $S \subseteq Q$ is a set of states.
- \preceq is a linear preorder over S .
- $l: S \rightarrow \{0, \dots, 2|Q|\}$ is a labeling.
- $\leq \subseteq \preceq$ is another preorder over S .
- G, B are sets of good and bad labels used for the Rabin condition.

For two states q and r in Q , we say that $q \approx r$ if $q \preceq r$ and $r \preceq q$. We constrain the labeling l so that it characterizes the equivalence classes of S under \preceq , and the preorder \leq to be a partial order over the equivalence classes of \preceq . Let \mathbf{Q} be the set of macrostates.

(a) An automaton \mathcal{B}

$$\mathbf{q}_0 = \langle \{q\}^0, \emptyset, G = \emptyset, B = \emptyset \rangle$$

$$\mathbf{q}_1 = \langle \{q\}^0 \prec \{p\}^1, q \prec p, G = \emptyset, B = \emptyset \rangle$$

$$\mathbf{q}_2 = \langle \{q\}^0 \prec \{p\}^2, q \prec p, G = \{0\}, B = \{1\} \rangle$$

$$\mathbf{q}_3 = \langle \{q\}^0 \prec \{p\}^1, q \prec p, G = \{0\}, B = \{2\} \rangle$$

(b) The first four macrostates in the run of $D^R(\mathcal{B})$ on ab^ω .

Figure 2: An automaton and four macrostates. For each macrostate $\langle S, \preceq, l, \preceq, G, B \rangle$, we first display the equivalence classes of S under \preceq in angle brackets, superscripted with the labels of l . We then display the \preceq relation, and finally the sets G and B .

Before defining transitions between macrostates, we reproduce the pruning of edges from G' by restricting the transition function ρ with respect to S and \preceq . For a state $q \in S$ and $\sigma \in \Sigma$, let $\rho_{S, \preceq}(q, \sigma) = \{q' \in \rho(q, \sigma) \mid \text{for every } r \in \rho^{-1}(q', \sigma) \cap S, r \preceq q\}$. Thus, when a state has multiple incoming σ -transitions from S , the function $\rho_{S, \preceq}$ keeps only the transitions from states maximal under the \preceq relation. For every state $q' \in \rho(S, \sigma)$, the set $\rho_{S, \preceq}^{-1}(q', \sigma) \cap S$ is an equivalence class under \preceq . We note that $\rho(S, \sigma) = \rho_{S, \preceq}(S, \sigma)$.

Example 5.2. Figure 2 displays the first four macrostates in a run of this determinization construction. Consider the state $\mathbf{q}_1 = \langle \{q, p\}, \preceq, l, \preceq, \emptyset, \emptyset \rangle$ where $q \prec p$, $q \preceq p$, $l(q) = 0$, and $l(p) = 1$. We have $\rho(q, a) = \{p, q\}$. However, $p \in \rho(p, a)$ and $q \prec p$. Thus we discard the transition from q to p , and $\rho_{S, \preceq}(q, a) = \{q\}$. In contrast, $\rho_{S, \preceq}(p, a) = \rho(p, a) = \{p\}$, because while $p \in \rho(q, a)$, it holds that $q \prec p$.

For $\sigma \in \Sigma$, we define the σ -successor of $\langle S, \preceq, l, \preceq, G, B \rangle$ to be $\langle S', \preceq', l', \preceq', G', B' \rangle$ as follows. First, $S' = \rho(S, \sigma)$. Second, define \preceq' as follows. For states $q', r' \in S'$, let $q \in \rho_{S, \preceq}^{-1}(q', \sigma)$ and $r \in \rho_{S, \preceq}^{-1}(r', \sigma)$. As the parents of q' and r' under $\rho_{S, \preceq}$ are equivalence classes the choice of q and r is arbitrary.

- If $q \prec r$, then $q' \prec' r'$.
- If $q \approx r$ and $q' \in F$ iff $r' \in F$, then $q' \approx' r'$.
- If $q \approx r$, $q' \notin F$, and $r' \in F$, then $q' \prec' r'$.

Example 5.3. As a running example we detail the transition from $\mathbf{q}_1 = \langle \{q, p\}, \preceq, l, \preceq, \emptyset, \emptyset \rangle$ to $\mathbf{q}_2 = \langle S', \preceq', l', \preceq', G', B' \rangle$ on b . We have $S' = \rho(\{q, p\}, b) = \{q, p\}$. To determine \preceq' , we note that $p \in S$ is the parent of both $q \in S'$ and $p \in S'$. Since $q \notin F$, and $p \in F$, we have $q \prec' p$.

Third, we define the labeling l' as follows. As in the profile tree T , on each level we give the label m to the minimal descendants, under the \preceq relation, of the first equivalence class to be labeled m . For a state $q \in S$, define the *nephews of q* to be $\text{neph}(q, \sigma) = \min_{\preceq'}(\rho_{S, \preceq}(\{r \in S \mid q \preceq r\}, \sigma))$. Conversely, for a state $r' \in S'$ we define the *uncles of r'* to be $\text{unc}(r', \sigma) = \{q \mid r' \in \text{neph}(q, \sigma)\}$.

Each state $r' \in S'$ inherits the oldest label from its uncles. If r' has no uncles, it gets a fresh label. Let $\text{FL}(l) = \{m \mid m \text{ not in the range of } l\}$ be the free labels in l , and let mj be the $\langle \preceq', < \rangle$ -minjection from $\{r' \in S' \mid \text{unc}(r', \sigma) = \emptyset\}$ to $\text{FL}(l)$, where $<$ is the standard order on $\{0, \dots, 2|Q|\}$. Let

$$l'(r') = \begin{cases} l(q), & \text{for some } q \in \min_{\preceq}(\text{unc}(r', \sigma)) & \text{if } \text{unc}(r', \sigma) \neq \emptyset, \\ \text{mj}(r') & \text{if } \text{unc}(r', \sigma) = \emptyset. \end{cases}$$

Example 5.4. The nephews of $q \in S$ is the \preceq' -minimal subset of the set $\rho_{s,\preceq}(\{r \in S \mid q \preceq r\}, \sigma)$. Since $q \preceq q$ and $q \preceq p$, we have that $\text{neph}(q, b) = \min_{\preceq'}(\{q, p\}) = \{q\}$. Similarly, for $p \in S$ we have $p \preceq p$ and $\text{neph}(p, b) = \min_{\preceq'}(\{p, q\}) = \{q\}$. Thus for $q \in S'$, we have $\min_{\preceq}(\text{unc}(q, b)) = \min_{\preceq}(\{p, q\}) = \{q\}$ and we set $l'(q) = l(q) = 0$. For $p \in S'$, we have $\text{unc}(p, b) = \emptyset$ and $l'(p)$ is the first unused label: $l'(p) = 2$.

Fourth, define the preorder \preceq' as follows. For states $q', r' \in S'$, define $q' \preceq' r'$ iff $q' \approx' r'$ or there exist $q, r \in S$ so that: $r' \in \rho_{s,\preceq}(r, \sigma)$; $q \in \text{unc}(q', \sigma)$; and $q \preceq r$. The labeling l' depends on recalling which states descend from the first equivalence class with a given label, and \preceq' tracks these descendants.

Finally, for a label m let $S_m = \{r \in S \mid l(r) = m\}$ and $S'_m = \{r' \in S' \mid l'(r') = m\}$ be the states in S , resp S' , labeled with m . Recall that a label m is good either when the branch it is following visits F -states, or the branch dies and it moves to another branch. Thus say m is *good* when: $S_m \neq \emptyset$; $S'_m \neq \emptyset$; and either $S'_m \subseteq F$ or $\rho_{s,\preceq}(S_m, \sigma) \cap S'_m = \emptyset$. G' is then $\{m \mid m \text{ is good}\}$. Conversely, a label is *bad* when it occurs in S , but not in S' . Thus the set of *bad* labels is $B' = \{m \mid S_m \neq \emptyset, S'_m = \emptyset\}$.

Example 5.5. As $p \in \rho_{s,\preceq}(p, b)$; $q \in \text{unc}(q, b)$; and $q \preceq p$, we have $q \preceq' p$. Since $l(q) = 0$ and $l'(q) = 0$, but $q \notin \rho_{s,\preceq}(q, b)$, we have $0 \in G'$, and as nothing is labeled 1 in l' , we have $1 \in B'$.

Lemma 5.6, proven in the full version, states that $\langle S', \preceq', l', \preceq', G', B' \rangle$ is a valid macrostate.

Lemma 5.6. *For a macrostate $\mathbf{q} \in \mathbf{Q}$ and $\sigma \in \Sigma$, the σ -successor of \mathbf{q} is a macrostate.*

Definition 5.7. Define the DRW automaton $D^R(\mathcal{A})$ to be $\langle \Sigma, \mathbf{Q}, \mathbf{Q}^{in}, \rho_{\mathbf{Q}}, \alpha \rangle$, where:

- $\mathbf{Q}^{in} = \{\langle Q^{in}, \preceq_0, l_0, \preceq_0, \emptyset, \emptyset \rangle\}$, where:
 - $\preceq_0 = \preceq_0 = Q^{in} \times Q^{in}$
 - $l_0(q) = 0$ for all $q \in Q^{in}$
- For $\mathbf{q} \in \mathbf{Q}$ and $\sigma \in \Sigma$, let $\rho_{\mathbf{Q}}(\mathbf{q}, \sigma) = \{\mathbf{q}'\}$, where \mathbf{q}' is the σ -successor of \mathbf{q}
- $\alpha = \langle G_0, B_0 \rangle, \dots, \langle G_{2|Q|}, B_{2|Q|} \rangle$, where for a label $m \in \{0, \dots, 2|Q|\}$:
 - $G_m = \{\langle S, \preceq, l, \preceq, G, B \rangle \mid m \in G\}$
 - $B_m = \{\langle S, \preceq, l, \preceq, G, B \rangle \mid m \in B\}$

Theorem 5.8, proven in the full version, asserts the correctness of the construction and says that its blowup is comparable with known determinization constructions.

Theorem 5.8. *For an NBW \mathcal{A} with n states, $L(D^R(\mathcal{A})) = L(\mathcal{A})$ and $D^R(\mathcal{A})$ has $n^{O(n)}$ states.*

There are two simple improvements to the new construction, detailed in the full version. First, we do not need $2|Q|$ labels: it is sufficient to use $|Q| - 1$ labels. Second, Piterman's technique of dynamic renaming can reduce the Rabin condition to a parity condition.

6 Discussion

In this paper we extended the notion of profiles from [6] and developed a theory of profile trees. This theory affords a novel determinization construction, where determinized-automaton states are sets of input-automaton states augmented with two preorders. In the future, a more thorough analysis could likely improve the upper bound on the size of our construction. We hope to see heuristic optimization techniques developed for this construction, just as heuristic optimization techniques were developed for Safra's construction [24].

More significantly, profile trees afford us the first theoretical underpinnings for determinization. Decades of research on Büchi determinization have resulted in a plethora of constructions, but a paucity of mathematical structures underlying their correctness. This is the first new major line of research in Büchi determinization since [16], and we expect it to lead to further research in this important area.

One important question is to understand better the connection between profile trees and Safra's construction. A key step in the transition between Safra trees is to remove states if they appear in more than one node. This seems analogous to the pruning of edges from G' . The second preorder in our construction, namely the relation \leq_i , seems to encode the order information embedded in Safra trees. Perhaps our approach could lead to declarative definition of constructions based on Safra and Muller-Schupp trees. In any case, it is our hope that profile trees will encourage the development of new methods to analyze and optimize determinization constructions.

References

- [1] C.S. Althoff, W. Thomas & N. Wallmeier (2005): *Observations on determinization of Büchi automata*. In: *ICALP*, doi:10.1016/j.tcs.2006.07.026.
- [2] R. Alur, T. A. Henzinger & O. Kupferman (2002): *Alternating-time temporal logic*. *J. ACM*, doi:10.1145/585265.585270.
- [3] J.R. Büchi (1962): *On a decision method in restricted second order arithmetic*. In: *ICLMPS*.
- [4] C. Courcoubetis & M. Yannakakis (1995): *The complexity of probabilistic verification*. *J. ACM*, doi:10.1145/210332.210339.
- [5] L. Doyen & J.-F. Raskin (2007): *Improved algorithms for the automata-based approach to model-checking*. In: *TACAS*, doi:10.1007/978-3-540-71209-1_34.
- [6] S. Fogarty, O. Kupferman, M.Y. Vardi & Th. Wilke (2011): *Unifying Büchi complementation constructions*. In: *CSL*, doi:10.4230/LIPIcs.CSL.2011.248.
- [7] S. Fogarty & M.Y. Vardi (2010): *Efficient Büchi universality checking*. In: *TACAS*, doi:10.1007/978-3-642-12002-2_17.
- [8] E. Friedgut, O. Kupferman & M.Y. Vardi (2006): *Büchi complementation made tighter*. *IJFCS*, doi:10.1142/S0129054106004145.
- [9] D. Kähler & Th. Wilke (2008): *Complementation, disambiguation, and determinization of Büchi automata unified*. In: *ICALP*, doi:10.1007/978-3-540-70575-8_59.
- [10] J. Kretínský & J. Esparza (2012): *Deterministic automata for the (F, G)-fragment of LTL*. In: *CAV*, doi:10.1007/978-3-642-31424-7_7.
- [11] O. Kupferman & M.Y. Vardi (2001): *Weak alternating automata are not that weak*. *TOCL*, doi:10.1145/377978.377993.
- [12] O. Kupferman & M.Y. Vardi (2005): *Safraless decision procedures*. In: *FOCS*, doi:10.1109/SFCS.2005.66.
- [13] L.H. Landweber (1969): *Decision problems for ω -automata*. *MST*, doi:10.1007/BF01691063.
- [14] C. Löding (1999): *Optimal bounds for the transformation of omega-automata*. In: *FSTTCS*, doi:10.1007/3-540-46691-6_8.
- [15] R. McNaughton (1966): *Testing and generating infinite sequences by a finite automaton*. *ICONT*, doi:10.1016/S0019-9958(66)80013-X.
- [16] D.E. Muller & P.E. Schupp (1995): *Simulating alternating tree automata by nondeterministic automata: new results and new proofs of theorems of Rabin, McNaughton and Safra*. *TCS*, doi:10.1016/0304-3975(94)00214-4.
- [17] N. Piterman (2006): *From nondeterministic Büchi and Streett automata to deterministic parity automata*. In: *LICS*, doi:10.2168/LMCS-3(3:5)2007.
- [18] A. Pnueli & R. Rosner (1989): *On the synthesis of a reactive module*. In: *POPL*, doi:10.1145/75277.75293.
- [19] M.O. Rabin & D. Scott (1959): *Finite automata and their decision problems*. *IBM JRD*, doi:10.1147/rd.32.0114.

- [20] S. Safra (1988): *On the complexity of ω -automata*. In: FOCS, doi:10.1109/SFCS.1988.21948.
- [21] S. Schewe (2009): *Büchi complementation made tight*. In: STACS, doi:10.4230/LIPIcs.STACS.2009.1854.
- [22] S. Schewe (2009): *Tighter bounds for the determinisation of Büchi automata*. In: FOSSACS, doi:10.1007/978-3-642-00596-1_13.
- [23] S. Tasiran, R. Hojati & R.K. Brayton (1995): *Language containment of non-deterministic omega-automata*. In: CHARME, doi:10.1007/3-540-60385-9_16.
- [24] M.-H. Tsai, S. Fogarty, M. Y. Vardi & Y.-K. Tsay (2010): *State of Büchi complementation*. In: CIAA, doi:10.1007/978-3-642-18098-9_28.
- [25] M.Y. Vardi (1985): *Automatic verification of probabilistic concurrent finite-state programs*. In: FOCS, doi:10.1109/SFCS.1985.12.