

# Faster Polynomial Multiplication over Finite Fields

DAVID HARVEY, University of New South Wales

JORIS VAN DER HOEVEN and GRÉGOIRE LECERF, CNRS, École Polytechnique

Polynomials over finite fields play a central role in algorithms for cryptography, error correcting codes, and computer algebra. The complexity of multiplying such polynomials is still a major open problem. Let  $p$  be a prime, and let  $M_p(n)$  denote the bit complexity of multiplying two polynomials in  $\mathbb{F}_p[X]$  of degree less than  $n$ . For  $n$  large compared to  $p$ , we establish the bound  $M_p(n) = O(n \log n 8^{\log^* n} \log p)$ , where  $\log^* n = \min\{k \in \mathbb{N} : \log^{\times k} \log n \leq 1\}$  stands for the iterated logarithm. This improves on the previously best known bound  $M_p(n) = O(n \log n \log \log n \log p)$ , which essentially goes back to the 1970s.

CCS Concepts: • **Mathematics of computing** → **Computations in finite fields**; **Computations on polynomials**; *Mathematical software performance*; • **Theory of computation** → *Algebraic complexity theory*; • **Computing methodologies** → *Algebraic algorithms*;

Additional Key Words and Phrases: Polynomial multiplication, finite field, algorithm, complexity bound, FFT

## ACM Reference Format:

David Harvey, Joris van der Hoeven, and Grégoire Lecerf. 2017. Faster polynomial multiplication over finite fields. *J. ACM* 63, 6, Article 52 (January 2017), 23 pages.

DOI: <http://dx.doi.org/10.1145/3005344>

## 1. INTRODUCTION

Given a ring  $R$ , a basic problem in complexity theory is to find an upper bound for the cost of multiplying two polynomials in  $R[X]$  of degree less than  $n$ . Several complexity models may be considered. In algebraic complexity models, such as the straight-line program model (see Chapter 4 of Bürgisser et al. [1997]), we count the number of ring operations in  $R$ , denoted by  $M_R(n)$ . In this model, the best currently known bound  $M_R(n) = O(n \log n \log \log n)$  was obtained by Cantor and Kaltofen [1991]. More precisely, their algorithm performs  $O(n \log n \log \log n)$  additions and subtractions and  $O(n \log n)$  multiplications in  $R$ . Their bound may be viewed as an algebraic analogue of the Schönhage–Strassen result for integer multiplication [Schönhage and Strassen 1971] and generalises previous work by Schönhage [1977]. These algorithms all rely on suitable incarnations of the fast Fourier transform (FFT) [Cooley and Tukey 1965]. For details, we refer the reader to Chapter 8 of von zur Gathen and Gerhard [2002].

In this article, we are mainly interested in the case that  $R$  is the finite field  $\mathbb{F}_p = \mathbb{Z}/p\mathbb{Z}$  for some prime  $p$ . The standard bit complexity model based on deterministic multitape Turing machines [Papadimitriou 1994] is more realistic in this setting, as it takes into account the dependence on  $p$ . We write  $M_p(n)$  for the bit complexity of multiplying two polynomials in  $\mathbb{F}_p[X]$  of degree less than  $n$ .

---

Authors' addresses: D. Harvey, School of Mathematics and Statistics, University of New South Wales, Sydney NSW 2052, Australia; email: [d.harvey@unsw.edu.au](mailto:d.harvey@unsw.edu.au); J. van der Hoeven and G. Lecerf, Laboratoire d'informatique, UMR 7161 CNRS, École Polytechnique, 91128 Palaiseau Cedex, France; emails: [vdhoeven@lix.polytechnique.fr](mailto:vdhoeven@lix.polytechnique.fr), [lecerf@lix.polytechnique.fr](mailto:lecerf@lix.polytechnique.fr).

Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies show this notice on the first page or initial screen of a display along with the full citation. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, to republish, to post on servers, to redistribute to lists, or to use any component of this work in other works requires prior specific permission and/or a fee. Permissions may be requested from Publications Dept., ACM, Inc., 2 Penn Plaza, Suite 701, New York, NY 10121-0701 USA, fax +1 (212) 869-0481, or [permissions@acm.org](mailto:permissions@acm.org).

© 2017 ACM 0004-5411/2017/01-ART52 \$15.00

DOI: <http://dx.doi.org/10.1145/3005344>

The complexity of polynomial multiplication is fundamental in the analysis of many other basic operations on mathematical objects, such as division with remainder, finding greatest common divisors, and computing power series expansions of algebraic functions. Faster multiplication algorithms over finite fields therefore admit many applications in areas such as cryptography and error correcting codes, at least from a theoretical point of view.

Two basic approaches are known for obtaining good asymptotic bounds for  $M_p(n)$ . If  $n$  is not too large compared to  $p$ , then one may use *Kronecker substitution*, which converts the problem to integer multiplication by packing the coefficients into large integers. Let  $l(n)$  denote the bit complexity of  $n$ -bit integer multiplication. Throughout the article, we make the customary assumption that  $l(n)/n$  is increasing, and for convenience we define  $l(x) := l(\lceil x \rceil)$  for  $x \in \mathbb{R}$ ,  $x > 0$ . According to our recent sharpening [Harvey et al. 2016a] of Fürer's bound [Fürer 2009], we may take

$$l(n) = O(n \log n 8^{\log^* n}), \quad (1)$$

where  $\log^* x$  denotes the iterated logarithm for  $x \in \mathbb{R}$ —that is,

$$\begin{aligned} \log^* x &:= \min\{k \in \mathbb{N} : \log^{\circ k} x \leq 1\}, \\ \log^{\circ k} &:= \log \circ \dots \circ \log \quad (k \text{ times}). \end{aligned} \quad (2)$$

If  $\log n = O(\log p)$ , then Kronecker substitution yields  $M_p(n) = O(l(n \log p))$  (see Section 2.6), so we have

$$M_p(n) = O(n \log p \log(n \log p) 8^{\log^*(n \log p)}).$$

However, Kronecker substitution leads to inferior bounds when  $n$  is large compared to  $p$ , due to coefficient growth in the lifted product in  $\mathbb{Z}[X]$ . In this situation, the best existing method is the algebraic Schönhage–Strassen algorithm (i.e., the Cantor–Kaltofen algorithm). In the Turing model, this yields the bound

$$M_p(n) = O(n \log n \log \log n \log p + n \log n l(\log p)).$$

The first term dominates for large enough  $n$ , say for  $\log \log p 8^{\log^* p} = O(\log \log n)$ , and we simply get

$$M_p(n) = O(n \log n \log \log n \log p). \quad (3)$$

Since Fürer's result [2009], there has been a gap between what is known for integer multiplication and for multiplication in  $\mathbb{F}_p[X]$ . Namely, it has remained an open question whether the factor  $\log \log n$  appearing in (3) can be improved to  $K^{\log^* n}$  for some  $K > 1$ . It seems that Fürer's construction does not work in  $\mathbb{F}_p[X]$ . Instead, we rely on recent techniques from Harvey et al. [2016a] and several new ideas that are specific to the finite field setting.

### 1.1. Our Contribution

The main result of this article is a new algorithm that completely closes the gap mentioned earlier.

**THEOREM 1.1.** *The bound*

$$M_p(n) = O(n \log p \log(n \log p) 8^{\log^*(n \log p)})$$

*holds uniformly for all  $n \geq 2$  and all primes  $p$ .*

This bound follows immediately via Kronecker substitution if  $\log n = O(\log p)$ , as was essentially already pointed out in Corollary 7.7 of Fürer [2009]. However, in Fürer's

statement, the term  $8^{\log^*(n \log p)}$  is replaced by  $K^{\log^*(n \log p)}$  for some unspecified  $K > 1$ . The improvement to  $K = 8$  in this case is a consequence of our previous result (1) for integer multiplication.

The bulk of the article concerns the reverse case where  $n$  is large compared to  $p$ . For example, if  $\log \log p = O(\log n)$ , the bound simply becomes

$$M_p(n) = O(n \log n 8^{\log^* n} \log p).$$

As promised, this replaces the factor  $\log \log n$  in (3) by  $K^{\log^* n}$  with  $K = 8$ . In this regime, the new bound does not follow directly from any results on integer multiplication.

The basic idea of the new algorithm is as follows. We first construct a special extension  $\mathbb{F}_{p^\kappa}$  of  $\mathbb{F}_p$ , where  $\kappa$  is exponentially smaller than  $n$ , and for which  $p^\kappa - 1$  has many small divisors. These divisors are themselves exponentially smaller than  $n$ , yet they are so numerous that their product is comparable to  $n$ . We now convert the given multiplication problem in  $\mathbb{F}_p[X]$  to multiplication in  $\mathbb{F}_{p^\kappa}[Y]$  by cutting up the polynomials into small pieces, and we then multiply in  $\mathbb{F}_{p^\kappa}[Y]$  by using FFTs over  $\mathbb{F}_{p^\kappa}$ . Applying the Cooley–Tukey method to the small divisors of  $p^\kappa - 1$ , we decompose the FFTs into transforms of exponentially shorter lengths (not necessarily powers of two). We use Bluestein’s method to convert each short transform to a convolution over  $\mathbb{F}_{p^\kappa}$  and then use Kronecker substitution (from bivariate to univariate polynomials) to convert this to multiplication in  $\mathbb{F}_p[X]$ . These latter multiplications are then handled recursively. We continue the recursion until  $n$  is comparable to  $p$ , at which point we switch to using Kronecker substitution (from polynomials to integers).

In many respects, this approach is very similar to the algorithm for integer multiplication introduced in Harvey et al. [2016a], and this partly explains why the final complexity bound has such a similar form. However, in the polynomial case, there are many additional technical issues to address. We recommend Harvey et al. [2016a] as a gentle introduction to the main ideas.

## 1.2. Outline of the Article

The article is structured as follows. In Section 2, we start with a survey of relevant basic techniques: discrete Fourier transforms (DFTs) and FFTs, Bluestein’s chirp transform, and Kronecker substitution. Most of this material is repeated from Section 2 of Harvey et al. [2016a] for the convenience of the reader; however, Sections 2.5 and 2.6 differ substantially.

In Section 3, we recall basic complexity results for arithmetic in finite fields. In particular, we consider the construction of irreducible polynomials in  $\mathbb{F}_p[X]$ , algorithms for finding roots of unity in  $\mathbb{F}_{p^\kappa}$ , and the cost of arithmetic in  $\mathbb{F}_{p^\kappa}$  and  $\mathbb{F}_{p^\kappa}[Y]$ .

In Section 4, we show how to construct special extensions of  $\mathbb{F}_p$  whose multiplicative group has a large subgroup of highly smooth order (i.e., is divisible by many small integers). The main tool is Theorem 3 of Adleman et al. [1983].

Section 5 gives complexity bounds for functions that satisfy recurrence inequalities involving postcompositions with “logarithmically slow” functions. The prototype of such an inequality is  $T(n) \leq KT(\log n) + L$ , where  $K$  and  $L$  are constants. The definitions and theorems are duplicated from Section 5 of Harvey et al. [2016a]; for the proofs, see Harvey et al. [2016a].

To minimize the constant  $K$  in the bound  $M_p(n) = O(n \log n K^{\log^* n} \log p)$  (for  $n$  large relative to  $p$ ), we need one more tool: in Section 6, we present a polynomial analogue of the Crandall–Fagin convolution algorithm [Crandall and Fagin 1994]. This allows us to convert a cyclic convolution over  $\mathbb{F}_p$  of length  $n$ , where  $n$  is arbitrary, into a cyclic convolution over  $\mathbb{F}_{p^\kappa}$  of somewhat smaller length  $N$ , where  $N$  is prescribed and where  $\kappa \approx 2n/N$ . (We can still obtain  $K = 16$  without using this Crandall–Fagin analogue,

but we do not know how to reach  $K = 8$  without it.) Section 7 is devoted to the proof of Theorem 1.1. It is noteworthy that  $K = 4$  can be achieved under “plausible number theoretic assumptions”: we refer the reader to Harvey et al. [2014] for details.

Section 8 offers some final notes and suggested directions for generalization. We first quickly dispense with the bit complexity of multiplication in  $\mathbb{F}_q[X]$  for prime powers  $q$  and in  $(\mathbb{Z}/m\mathbb{Z})[X]$  for arbitrary integers  $m \geq 1$ . We then sketch some algebraic complexity bounds for polynomial multiplication over  $\mathbb{F}_p$ -algebras and  $(\mathbb{Z}/m\mathbb{Z})$ -algebras, especially in the straight-line program model. Our techniques also give rise to new strategies for polynomial evaluation interpolation over  $\mathbb{F}_q$ . This may for instance be applied to the efficient multiplication of polynomial matrices over  $\mathbb{F}_q$ .

Implementing the new algorithms from this article is a work in progress. In particular, the techniques from Section 4 turn out to be very useful in practice. The field  $\mathbb{F}_{2^{60}}$  is one example of a field with primitive roots of unity of high smooth order  $N = 2^{60} - 1 = 3^2 \cdot 5^2 \cdot 7 \cdot 11 \cdot 13 \cdot 31 \cdot 41 \cdot 61 \cdot 151 \cdot 331 \cdot 1321$ . For finite fields  $\mathbb{F}_q$  with  $q = 2^K$ , it turns out [Harvey et al. 2016b] that polynomial multiplication over  $\mathbb{F}_{2^{60}}$  can be used as a building block for developing efficient polynomial arithmetic over  $\mathbb{F}_q$ . The gains are most important for our implementation of polynomial matrix multiplication over  $\mathbb{F}_q$ .

*Notations.* We use Hardy’s notations  $f \prec g$  for  $f = o(g)$ , and  $f \asymp g$  for  $f = O(g)$  and  $g = O(f)$ . The “soft-Oh” notation  $f(n) = \tilde{O}(g(n))$  means that  $f(n) = g(n)(\log(3+g(n)))^{O(1)}$  (see Section 7 of Chapter 25 in von zur Gathen and Gerhard [2002] for details). The symbol  $\mathbb{R}^{\geq}$  denotes the set of nonnegative real numbers, and  $\mathbb{N}$  denotes  $\{0, 1, 2, \dots\}$ . For a ring  $R$  and  $n \in \mathbb{N}$ , we write  $R[X]_n := \{P \in R[X] : \deg P < n\}$ .

We will write  $\lg n := \lceil \log n / \log 2 \rceil$ . In expressions like  $\log \log p$  or  $\lg \lg p$ , we tacitly assume that the function is adjusted so as to take positive values for small primes such as  $p = 2$ .

## 2. SURVEY OF CLASSICAL TOOLS

This section recalls basic facts on Fourier transforms and related techniques used in subsequent sections. For more details and historical references, we refer the reader to standard books on the subject, such as those of Aho et al. [1974], Bürgisser et al. [1997], von zur Gathen and Gerhard [2002], and Rao et al. [2010].

### 2.1. Arrays and Sorting

In the Turing model, we have a fixed number of linear tapes available. An  $n_1 \times \dots \times n_d$  array  $M_{i_1, \dots, i_d}$  of  $b$ -bit elements is stored as a linear array of  $n_1 \dots n_d b$  bits. We generally assume that the elements are ordered lexicographically by  $(i_1, \dots, i_d)$ , although this is just an implementation detail.

What is significant from a complexity point of view is that occasionally we must switch representations to access an array (say two dimensional) by “rows” or by “columns.” In the Turing model, we may transpose an  $n_1 \times n_2$  matrix of  $b$ -bit elements in time  $O(bn_1n_2 \lg \min(n_1, n_2))$  using the algorithm in the appendix to Bostan et al. [2007]. Briefly, the idea is to split the matrix into two halves along the “short” dimension and transpose each half recursively.

We will also require more complex rearrangements of data, for which we resort to sorting. Suppose that  $X$  is a totally ordered set, whose elements are represented by bit strings of length  $b$ , and suppose that we can compare elements of  $X$  in time  $O(b)$ . Then an array of  $n$  elements of  $X$  may be sorted in time  $O(bn \lg n)$  using merge sort [Knuth 1998], which can be implemented efficiently on a Turing machine.

## 2.2. Discrete Fourier Transforms

Let  $R$  be a commutative ring with identity, and let  $n \geq 1$ . An element  $\omega \in R$  is said to be a *principal  $n$ -th root of unity* if  $\omega^n = 1$  and

$$\sum_{k=0}^{n-1} (\omega^i)^k = 0 \quad (4)$$

for all  $i \in \{1, \dots, n-1\}$ . In this case, we define the *discrete Fourier transform* (or DFT) of an  $n$ -tuple  $a = (a_0, \dots, a_{n-1}) \in R^n$  with respect to  $\omega$  to be  $\text{DFT}_\omega(a) = \hat{a} = (\hat{a}_0, \dots, \hat{a}_{n-1}) \in R^n$ , where

$$\hat{a}_i := a_0 + a_1\omega^i + \dots + a_{n-1}\omega^{(n-1)i}.$$

In other words,  $\hat{a}_i$  is the evaluation of the polynomial  $A(X) := a_0 + a_1X + \dots + a_{n-1}X^{n-1}$  at  $\omega^i$ .

If  $\omega$  is a principal  $n$ -th root of unity, then so is its inverse  $\omega^{-1} = \omega^{n-1}$ , and we have

$$\text{DFT}_{\omega^{-1}}(\text{DFT}_\omega(a)) = na.$$

Indeed, writing  $b := \text{DFT}_{\omega^{-1}}(\text{DFT}_\omega(a))$ , the relation (4) implies that

$$b_i = \sum_{j=0}^{n-1} \hat{a}_j \omega^{-ji} = \sum_{j=0}^{n-1} \sum_{k=0}^{n-1} a_k \omega^{j(k-i)} = \sum_{k=0}^{n-1} a_k \sum_{j=0}^{n-1} \omega^{j(k-i)} = \sum_{k=0}^{n-1} a_k (n\delta_{i,k}) = na_i,$$

where  $\delta_{i,k} = 1$  if  $i = k$  and  $\delta_{i,k} = 0$  otherwise.

*Remark 2.1.* In all of the new algorithms introduced in this article, we actually work over a field whose characteristic does not divide  $n$ . In this setting, the concept of principal root of unity coincides with the more familiar primitive root of unity.

## 2.3. The Cooley–Tukey FFT

Let  $\omega$  be a principal  $n$ -th root of unity, and let  $n = n_1 n_2$ , where  $1 < n_1 < n$ . Then  $\omega^{n_1}$  is a principal  $n_2$ -th root of unity and  $\omega^{n_2}$  is a principal  $n_1$ -th root of unity. Moreover, for any  $i_1 \in \{0, \dots, n_1 - 1\}$  and  $i_2 \in \{0, \dots, n_2 - 1\}$ , we have

$$\begin{aligned} \hat{a}_{i_1 n_2 + i_2} &= \sum_{k_1=0}^{n_1-1} \sum_{k_2=0}^{n_2-1} a_{k_2 n_1 + k_1} \omega^{(k_2 n_1 + k_1)(i_1 n_2 + i_2)} \\ &= \sum_{k_1=0}^{n_1-1} \omega^{k_1 i_2} \left( \sum_{k_2=0}^{n_2-1} a_{k_2 n_1 + k_1} (\omega^{n_1})^{k_2 i_2} \right) (\omega^{n_2})^{k_1 i_1}. \end{aligned} \quad (5)$$

If  $\mathcal{A}_1$  and  $\mathcal{A}_2$  are algorithms for computing DFTs of length  $n_1$  and  $n_2$ , we may use (5) to construct an algorithm  $\mathcal{A}_1 \odot \mathcal{A}_2$  for computing DFTs of length  $n$  as follows.

For each  $k_1 \in \{0, \dots, n_1 - 1\}$ , the sum inside the brackets corresponds to the  $i_2$ -th coefficient of a DFT of the  $n_2$ -tuple  $(a_{0n_1+k_1}, \dots, a_{(n_2-1)n_1+k_1}) \in R^{n_2}$  with respect to  $\omega^{n_1}$ . Evaluating these inner DFTs requires  $n_1$  calls to  $\mathcal{A}_2$ . Next, we multiply by the *twiddle factors*  $\omega^{k_1 i_2}$ , at a cost of  $n$  operations in  $R$ . (Actually, fewer than  $n$  multiplications are required, as some of the twiddle factors are equal to 1. This optimization, although important in practice, has no asymptotic effect on the algorithms discussed in this article.) Finally, for each  $i_2 \in \{0, \dots, n_2 - 1\}$ , the outer sum corresponds to the  $i_1$ -th coefficient of a DFT of an  $n_1$ -tuple in  $R^{n_1}$  with respect to  $\omega^{n_2}$ . These outer DFTs require  $n_2$  calls to  $\mathcal{A}_1$ .

Denoting by  $F_R(n)$  the number of ring operations needed to compute a DFT of length  $n$ , and assuming that we have available a precomputed table of twiddle factors, we obtain

$$F_R(n_1 n_2) \leq n_1 F_R(n_2) + n_2 F_R(n_1) + n.$$

For a factorisation  $n = n_1 \cdots n_d$ , this recursively yields

$$F_R(n) \leq \sum_{i=1}^d \frac{n}{n_i} F_R(n_i) + (d-1)n. \quad (6)$$

The corresponding algorithm is denoted by  $\mathcal{A}_1 \odot \cdots \odot \mathcal{A}_d$ . The  $\odot$  operation is neither commutative nor associative; the preceding expression will always be taken to mean  $(\cdots((\mathcal{A}_1 \odot \mathcal{A}_2) \odot \mathcal{A}_3) \odot \cdots) \odot \mathcal{A}_d$ .

The preceding discussion requires several modifications in the Turing model. Assume that elements of  $R$  are represented by  $b$  bits.

First, for  $\mathcal{A}_1 \odot \mathcal{A}_2$ , we must add a rearrangement cost of  $O(bn \lg \min(n_1, n_2))$  to efficiently access the rows and columns for the recursive subtransforms (see Section 2.1). For the general case  $\mathcal{A}_1 \odot \cdots \odot \mathcal{A}_d$ , the total rearrangement cost is bounded by  $O(\sum_i bn \lg n_i) = O(bn \lg n)$ .

Second, we will sometimes use nonalgebraic algorithms to compute the subtransforms, so it may not make sense to express their cost in terms of  $F_R$ . The relation (6) therefore becomes

$$F(n) \leq \sum_{i=1}^d \frac{n}{n_i} F(n_i) + (d-1)nm_R + O(bn \lg n), \quad (7)$$

where  $F(n)$  is the (Turing) cost of a transform of length  $n$  over  $R$ , and where  $m_R$  is the cost of a single multiplication in  $R$ .

Finally, we point out that  $\mathcal{A}_1 \odot \mathcal{A}_2$  requires access to a table of twiddle factors  $\omega^{i_1 i_2}$ , ordered lexicographically by  $(i_1, i_2)$ , for  $0 \leq i_1 < n_1$ ,  $0 \leq i_2 < n_2$ . Assuming that we are given as input a precomputed table of the form  $1, \omega, \dots, \omega^{n-1}$ , we must show how to extract the required twiddle factor table in the correct order. We first construct a list of triples  $(i_1, i_2, i_1 i_2)$ , ordered by  $(i_1, i_2)$ , in time  $O(n \lg n)$ ; then sort by  $i_1 i_2$  in time  $O(n \lg^2 n)$  (see Section 2.1); then merge with the given root table to obtain a table  $(i_1, i_2, \omega^{i_1 i_2})$ , ordered by  $i_1 i_2$ , in time  $O(n(b + \lg n))$ ; and finally sort by  $(i_1, i_2)$  in time  $O(n \lg n(b + \lg n))$ . The total cost of the extraction is thus  $O(n \lg n(b + \lg n))$ .

The corresponding cost for  $\mathcal{A}_1 \odot \cdots \odot \mathcal{A}_d$  is determined as follows. Assuming that the table  $1, \omega, \dots, \omega^{n-1}$  is given as input, we first extract the subtables of  $(n_1 \cdots n_i)$ -th roots of unity for  $i = d-1, \dots, 2$  in time  $O((n_1 \cdots n_d + \cdots + n_1 n_2)(b + \lg n)) = O(n(b + \lg n))$ . Extracting the twiddle factor table for the decomposition  $(n_1 \cdots n_{i-1}) \times n_i$  then costs  $O(n_1 \cdots n_i \lg n(b + \lg n))$ ; the total over all  $i$  is again  $O(n \lg n(b + \lg n))$ .

*Remark 2.2.* An alternative approach is to compute the twiddle factors directly in the correct order. When working over  $\mathbb{C}$ , as in Section 3 of Harvey et al. [2016a], this requires a slight increase in the working precision. Similar comments apply to the root tables used in Bluestein's algorithm in Section 2.5.

## 2.4. Fast Fourier Multiplication

Let  $\omega$  be a principal  $n$ -th root of unity in  $R$ , and assume that  $n$  is invertible in  $R$ . Consider two polynomials  $A = a_0 + \cdots + a_{n-1}X^{n-1}$  and  $B = b_0 + \cdots + b_{n-1}X^{n-1}$  in  $R[X]$ .



Let  $C = c_0 + \dots + c_{n-1}X^{n-1}$  be the polynomial defined by

$$c := \frac{1}{n} \text{DFT}_{\omega^{-1}}(\text{DFT}_{\omega}(a) \text{DFT}_{\omega}(b)),$$

where the product of the DFTs is taken pointwise. By construction, we have  $\hat{c} = \hat{a}\hat{b}$ , which means that  $C(\omega^i) = A(\omega^i)B(\omega^i)$  for all  $i \in \{0, \dots, n-1\}$ . The product  $S = s_0 + \dots + s_{n-1}X^{n-1}$  of  $A$  and  $B$  modulo  $X^n - 1$  also satisfies  $S(\omega^i) = A(\omega^i)B(\omega^i)$  for all  $i$ . Consequently,  $\hat{s} = \hat{a}\hat{b}$ ,  $s = \text{DFT}_{\omega^{-1}}(\hat{s})/n = c$ , from which  $C = S$ .

For polynomials  $A, B \in R[X]$  with  $\deg A < n$  and  $\deg B < n$ , we thus obtain an algorithm for the computation of  $AB$  modulo  $X^n - 1$  using at most  $3F_R(n) + O(n)$  operations in  $R$ . Modular products of this type are also called *cyclic convolutions*. If  $\deg(AB) < n$ , then we may recover the product  $AB$  from its reduction modulo  $X^n - 1$ . This multiplication method is called *FFT multiplication*.

If one of the arguments (say  $B$ ) is fixed and we want to compute many products  $AB$  (or cyclic convolutions) for different  $A$ , then we may precompute  $\text{DFT}_{\omega}(b)$ , after which each new product  $AB$  can be computed using only  $2F_R(n) + O(n)$  operations in  $R$ .

## 2.5. Bluestein's Chirp Transform

We have shown earlier how to multiply polynomials using DFTs. Inversely, it is possible to reduce the computation of DFTs—of arbitrary length, not necessarily a power of two—to polynomial multiplication [Bluestein 1970], as follows.

Let  $\omega$  be a principal  $n$ -th root of unity, and consider an  $n$ -tuple  $a \in R^n$ . First consider the case that  $n$  is odd. Let

$$f_i := \omega^{(i^2-i)/2}, \quad f'_i := \omega^{(i^2+i)/2}, \quad g_i := \omega^{(-i^2-i)/2}.$$

Note that  $(i^2 - i)/2$  and  $(i^2 + i)/2$  are integers, and that

$$g_{i+n} = \omega^{(-(i+n)^2-(i+n))/2} = \omega^{(-i^2-i)/2} \omega^{n(-i-(n+1)/2)} = g_i.$$

Let  $F := f_0a_0 + \dots + f_{n-1}a_{n-1}X^{n-1}$ ,  $G := g_0 + \dots + g_{n-1}X^{n-1}$ , and  $C := c_0 + \dots + c_{n-1}X^{n-1} \equiv FG$  modulo  $X^n - 1$ . Then

$$f'_i c_i = \sum_{j=0}^{n-1} f'_i (f_j a_j) g_{i-j} = \sum_{j=0}^{n-1} \omega^{(i^2+i)/2} \omega^{(j^2-j)/2} \omega^{(-(i-j)^2-(i-j))/2} a_j = \sum_{j=0}^{n-1} \omega^{ij} a_j = \hat{a}_i.$$

In other words, the computation of a DFT of odd length  $n$  reduces to a cyclic convolution product of the same length  $n$ , together with  $O(n)$  additional operations in  $R$ . Notice that the polynomial  $G$  is fixed and independent of  $a$  in this product.

Now suppose that  $n$  is even. In this case, we require that 2 be invertible in  $R$  and that  $\omega^{n/2} = -1$ . Let  $\sigma := (-1)^{n/2}$ , and put

$$f_i := \omega^{i^2}, \quad f'_i := \omega^{i^2+i}, \quad g_i := \omega^{-i^2} + \omega^{-i^2-i}.$$

Then

$$g_{i+\frac{n}{2}} = \sigma(\omega^{-i^2} - \omega^{-i^2-i}), \quad g_{i+n} = g_i,$$

and

$$\begin{aligned} \frac{1}{2} (g_i + \sigma g_{i+\frac{n}{2}}) &= \omega^{-i^2}, \\ \frac{1}{2} (g_i - \sigma g_{i+\frac{n}{2}}) &= \omega^{-i^2-i}. \end{aligned}$$

As earlier, let  $F := f_0 a_0 + \dots + f_{n-1} a_{n-1} X^{n-1}$ ,  $G := g_0 + \dots + g_{n-1} X^{n-1}$ , and  $C := c_0 + \dots + c_{n-1} X^{n-1} \equiv FG \text{ modulo } X^n - 1$ . Then for  $i \in \{0, \dots, \frac{n}{2} - 1\}$ , we have

$$\begin{aligned} \frac{1}{2} f_i (c_i + \sigma c_{i+\frac{n}{2}}) &= \frac{1}{2} f_i \sum_{j=0}^{n-1} f_j a_j (g_{i-j} + \sigma g_{i-j+\frac{n}{2}}) \\ &= \sum_{j=0}^{n-1} \omega^{i^2} \omega^{j^2} \omega^{-(i-j)^2} a_j = \sum_{j=0}^{n-1} \omega^{2ij} a_j = \hat{a}_{2i}, \end{aligned}$$

and similarly

$$\frac{1}{2} f'_i (c_i - \sigma c_{i+\frac{n}{2}}) = \sum_{j=0}^{n-1} \omega^{i^2+i} \omega^{j^2} \omega^{-(i-j)^2-(i-j)} a_j = \sum_{j=0}^{n-1} \omega^{(2i+1)j} a_j = \hat{a}_{2i+1}.$$

Again, the DFT reduces to a convolution of length  $n$ , together with  $O(n)$  additional operations in  $R$ .

The only complication in the Turing model is the cost of extracting the  $f_i$ ,  $f'_i$ , and  $g_i$  in the correct order. For example, consider the  $f_i$  in the case that  $n$  is odd. Given as input a precomputed table  $1, \omega, \omega^2, \dots, \omega^{n-1}$ , we may extract the  $f_i$  in time  $O(n \lg n (b + \lg n))$  by applying the strategy from Section 2.3 to the pairs  $(i, \frac{1}{2}(i^2 - i) \bmod n)$  for  $0 \leq i < n$ . The other sequences are handled similarly. For the  $g_i$  in the case that  $n$  is even, we also need to perform  $O(n)$  additions in  $R$ ; assuming that additions in  $R$  have cost  $O(b)$ , this is already covered by the preceding bound.

*Remark 2.3.* Bluestein's original formulation used the weights  $f_i := \omega^{i^2/2}$ ,  $g_i := \omega^{-i^2/2}$ . This has two drawbacks in our setting. First, it requires the presence of  $2n$ -th roots of unity, which may not exist in  $R$ . Second, if  $n$  is odd, it leads to negacyclic convolutions rather than cyclic convolutions. The variants used here avoid both of these problems.

*Remark 2.4.* It is also possible to give variants of the new multiplication algorithms in which Bluestein's transform is replaced by a different method for converting DFTs to convolutions, such as Rader's algorithm [Rader 1968].

## 2.6. Kronecker Substitution

Multiplication in  $\mathbb{F}_p[X]$  may be reduced to multiplication in  $\mathbb{Z}$  using the classical technique of Kronecker substitution (see Corollary 8.27 in von zur Gathen and Gerhard [2002]). More precisely, let  $n > 0$ , and suppose that we are given two polynomials  $A, B \in \mathbb{F}_p[X]$  of degree less than  $n$ . Let  $\tilde{A}, \tilde{B} \in \mathbb{Z}[X]$  be lifts of these polynomials, with coefficients  $A_i$  and  $B_i$  satisfying  $0 \leq A_i, B_i < p$ . Then for the product  $\tilde{C} = \tilde{A}\tilde{B}$ , we have  $0 \leq C_i < np^2$ , and the coefficients of  $\tilde{C}$  may be read off the integer product  $\tilde{C}(2^N) = \tilde{A}(2^N)\tilde{B}(2^N)$ , where  $N := 2 \lg p + \lg n$ . We deduce the coefficients of  $C := AB$  by dividing each  $C_i$  by  $p$ . This shows that  $M_p(n) = O(\lg(n(\lg p + \lg n)) + n \lg p)$ . Using the assumption that  $\lg(n)/n$  is increasing, we obtain

$$M_p(n) = O(\lg(n(\lg p + \lg n))).$$

If we also assume that  $\lg n = O(\lg p)$  (i.e., that the degree is not too large), then this simply becomes

$$M_p(n) = O(\lg(n \lg p)). \quad (8)$$



A second type of Kronecker substitution reduces bivariate polynomial multiplication to the univariate case. Indeed, let  $n \geq 1$  and  $\kappa \geq 1$ , and suppose that  $A, B \in \mathbb{F}_p[X, Z]$ , where  $\deg_X A, \deg_X B < n$  and  $\deg_Z A, \deg_Z B < \kappa$ . We may then recover  $C := AB$  from the univariate product  $C(Y^{2\kappa}, Y) = A(Y^{2\kappa}, Y)B(Y^{2\kappa}, Y)$  in  $\mathbb{F}_p[Y]$ . Note that  $A(Y^{2\kappa}, Y)$  and  $B(Y^{2\kappa}, Y)$  have degree less than  $2n\kappa$ , so the cost of the bivariate product is  $M_p(2n\kappa) + O(n\kappa \lg p)$ .

The same method works for computing cyclic convolutions: to multiply  $A, B \in \mathbb{F}_p[X, Z]/(X^n - 1)$ , the same substitution leads to a product in  $\mathbb{F}_p[Y]/(Y^{2n\kappa} - 1)$ . The cost is thus  $M'_p(2n\kappa) + O(n\kappa \lg p)$ , where  $M'_p(d)$  denotes the cost of a multiplication in  $\mathbb{F}_p[X]/(X^d - 1)$ .

### 3. ARITHMETIC IN FINITE FIELDS

Let  $p$  be a prime, and let  $\kappa \geq 1$ . In this section, we review basic results concerning arithmetic in  $\mathbb{F}_{p^\kappa}$  and  $\mathbb{F}_{p^\kappa}[Y]$ .

We assume throughout that  $\mathbb{F}_{p^\kappa}$  is represented as  $\mathbb{F}_p[Z]/P$  for some irreducible monic polynomial  $P \in \mathbb{F}_p[Z]$  of degree  $\kappa$ . Thus, an element of  $\mathbb{F}_{p^\kappa}$  is represented uniquely by its residue modulo  $P$ —that is, by a polynomial  $F \in \mathbb{F}_p[Z]$  of degree less than  $\kappa$ .

**LEMMA 3.1.** *Let  $p$  be a prime, and let  $\kappa \geq 1$ . We may compute a monic irreducible polynomial  $P \in \mathbb{F}_p[Z]$  of degree  $\kappa$  in time  $\tilde{O}(\kappa^4 p^{1/2})$ .*

**PROOF.** See Theorem 3.2 in Shoup [1990].  $\square$

The preceding complexity bound is very pessimistic in practice. Better complexity bounds are known if we allow randomized algorithms or unproved hypotheses. For instance, assuming GRH, the bound reduces to  $(\kappa \log p)^{O(1)}$  [Adleman and Lenstra 1986].

We now consider the cost of arithmetic in  $\mathbb{F}_{p^\kappa}$ , assuming that  $P$  is given. Addition and subtraction in  $\mathbb{F}_{p^\kappa}$  have cost  $O(\kappa \lg p)$ . For multiplication, we will always use the Schönhage–Strassen algorithm. Denote by  $S_p(\kappa)$  the cost of multiplying polynomials in  $\mathbb{F}_p[Z]_\kappa$  by this method—that is, using the polynomial variant [Schönhage 1977] for the polynomials themselves, followed by the integer version [Schönhage and Strassen 1971] to handle the coefficient multiplications. Then we have

$$S_p(\kappa) = O(\kappa \lg \kappa \lg \lg \kappa \lg p \lg \lg p \lg \lg \lg p).$$

Of course, we could use the new multiplication algorithm recursively for these products, but it turns out that Schönhage–Strassen is fast enough and leads to a simpler complexity analysis in Section 7. Let  $D_p(\kappa)$  denote the cost of dividing a polynomial in  $\mathbb{F}_p[Z]_{2\kappa}$  by  $P$ , returning the quotient and remainder. Using Newton's method (see Chapter 9 in von zur Gathen and Gerhard [2002]), we have  $D_p(\kappa) = O(S_p(\kappa))$ . Thus, elements of  $\mathbb{F}_{p^\kappa}$  may be multiplied in time  $O(S_p(\kappa))$ .

Let  $N \geq 1$  be a divisor of  $p^\kappa - 1$ . To compute a DFT over  $\mathbb{F}_{p^\kappa}$  of length  $N$ , we must first have access to a primitive  $N$ -th root of unity in  $\mathbb{F}_{p^\kappa}$ . In general, it is very difficult to find a primitive root of  $\mathbb{F}_{p^\kappa}$ , as it requires knowledge of the factorisation of  $p^\kappa - 1$ . However, to find a primitive  $N$ -th root, Lemma 3.3 in the following shows that it is enough to know the factorisation of  $N$ . The construction relies on the following existence result.

**LEMMA 3.2.** *Let  $\ell$  be a positive integer such that  $p^\ell > c_1 r^4 (\log r + 1)^4 \kappa^2$ , where  $r$  is the number of distinct prime divisors of  $p^\kappa - 1$ , and where  $c_1 > 0$  is a certain absolute constant. Then there exists a monic irreducible polynomial  $\Theta \in \mathbb{F}_p[Z]$  of degree  $\ell$  such that  $\Theta$  modulo  $P$  is a primitive root of unity for  $\mathbb{F}_{p^\kappa} = \mathbb{F}_p[Z]/P$ .*

**PROOF.** This is Theorem 1.1 in Shoup [1992].  $\square$

**LEMMA 3.3.** *Assume that  $N$  divides  $p^\kappa - 1$  and that the factorisation of  $N$  is given. Then we may compute a primitive  $N$ -th root of unity in  $\mathbb{F}_p[Z]/P$  in time  $\tilde{O}(\kappa^9 p)$ .*

**PROOF.** Testing whether a given  $\alpha \in \mathbb{F}_{p^\kappa}$  is a primitive  $N$ -th root of unity reduces to checking that  $\alpha^N = 1$  and  $\alpha^{N/s} \neq 1$  for every prime divisor  $s$  of  $N$ . According to Lemma 3.2, it suffices to apply this test to  $\alpha := (\Theta \bmod P)^{(p^\kappa - 1)/N}$  for  $\Theta \in \mathbb{F}_p[Z]$  running over all monic polynomials of degree  $\ell := \lceil \log(c_1 r^4 (\log r + 1)^4 \kappa^2) / \log p \rceil$ , where  $r = O(\kappa \lg p)$  is as in the lemma. The number of candidates is at most  $p^\ell \leq cr^4 (\log r + 1)^4 \kappa^2 p = \tilde{O}(\kappa^6 p)$ , and each candidate can be tested in time  $\tilde{O}(r \lg N(\kappa \lg p)) = \tilde{O}(\kappa^3 \lg^3 p)$  using binary exponentiation.  $\square$

Of course, we can do much better if randomized algorithms are allowed, as  $\zeta^{(p^\kappa - 1)/N}$  is reasonably likely to be a primitive  $N$ -th root for randomly selected  $\zeta$ .

Finally, we consider polynomial multiplication over  $\mathbb{F}_{p^\kappa}$ . Let  $n \geq 1$ , and let  $A, B \in \mathbb{F}_{p^\kappa}[X]_n$ . Let  $\tilde{A}, \tilde{B} \in \mathbb{F}_p[X, Z]$  be their natural lifts (i.e., of degree less than  $\kappa$  with respect to  $Z$ ). The bivariate product  $\tilde{C} := \tilde{A}\tilde{B}$  may be computed using Kronecker substitution (Section 2.6) in time  $M_p(2n\kappa) + O(n\kappa \lg p)$ . Writing  $\tilde{C} = \sum_{i=0}^{2n-2} \tilde{C}_i(Z)X^i$ , to recover  $AB$  we must divide each  $\tilde{C}_i$  by  $P$ . Denoting by  $M_{p^\kappa}(n)$  the complexity of the original multiplication problem, we obtain

$$\begin{aligned} M_{p^\kappa}(n) &= M_p(2n\kappa) + O(nD_p(\kappa)) \\ &= M_p(2n\kappa) + O(nS_p(\kappa)). \end{aligned}$$

As in Section 2.6, the same method may be used for cyclic convolutions. If  $M'_{p^\kappa}(n)$  denotes the cost of multiplication in  $\mathbb{F}_{p^\kappa}[X]/(X^n - 1)$ , then we get

$$M'_{p^\kappa}(n) = M'_p(2n\kappa) + O(nS_p(\kappa)).$$

**Remark 3.4.** In the other direction, we can also reduce multiplication in  $\mathbb{F}_p[X]$  to multiplication in  $\mathbb{F}_{p^\kappa}[Y]$  by splitting the inputs into chunks of size  $\lfloor \kappa/2 \rfloor$ . This leads to the bound  $M_p(n) \leq M_{p^\kappa}(\lceil n/\lfloor \kappa/2 \rfloor \rceil) + O(n \lg p)$  for  $n \geq \kappa$ . A variant of this procedure is developed in detail in Section 6.

#### 4. PREPARING FOR DFTS OF LARGE SMOOTH ORDERS

The aim of this section is to prove the following theorem, which allows us to construct “small” extensions  $\mathbb{F}_{p^\lambda}$  of  $\mathbb{F}_p$  containing “many” roots of unity of “low” order. Recall that a positive integer is said to be  $y$ -smooth if all of its prime divisors are less than or equal to  $y$ .

**THEOREM 4.1.** *There exist computable absolute constants  $c_3 > c_2 > 0$  and  $n_0 \in \mathbb{N}$  with the following properties. Let  $p$  be a prime, and let  $n \geq n_0$ . Then there exists an integer  $\lambda$  in the interval*

$$(\lg n)^{c_2 \lg \lg \lg n} < \lambda < (\lg n)^{c_3 \lg \lg \lg n},$$

*and a  $(\lambda + 1)$ -smooth integer  $M \geq n$ , such that  $M \mid p^\lambda - 1$ . Moreover, given  $p$  and  $n$ , we may compute  $\lambda$  and the prime factorisation of  $M$  in time  $O((\lg n)^{\lg \lg n})$ .*

The proof depends on a series of preparatory lemmas. For  $\lambda \geq 2$ , define

$$H_\lambda := \prod_{\substack{q \text{ prime,} \\ q-1 \mid \lambda}} q.$$

For example,  $H_{36} = 1919190 = 2 \cdot 3 \cdot 5 \cdot 7 \cdot 13 \cdot 19 \cdot 37$  and  $H_{37} = 2$ . Note that  $H_\lambda$  is always squarefree and  $(\lambda + 1)$ -smooth. We are interested in finding  $\lambda$  for which  $H_\lambda$  is unusually large—that is, for which  $\lambda$  has many divisors  $d$  such that  $d + 1$  happens to be prime. Most of the heavy lifting is done by the following remarkable result.

**LEMMA 4.2.** *Define  $\lambda_0(k) := \min\{\lambda \in \mathbb{N} : H_\lambda \geq \sqrt{k}\}$ . There exist computable absolute constants  $c_5 > c_4 > 0$  such that for any integer  $k > 100$ , we have*

$$(\log k)^{c_4 \log \log \log k} < \lambda_0(k) < (\log k)^{c_5 \log \log \log k}.$$

**PROOF.** This is part of Theorem 3 in Adleman et al. [1983]. (The threshold  $\sqrt{k}$  could of course be replaced by any fixed power of  $k$ . It is stated this way in Adleman et al. [1983] because that work is concerned with primality testing.)  $\square$

The link between  $H_\lambda$  and  $\mathbb{F}_{p^\lambda}$  is as follows.

**LEMMA 4.3.** *Let  $p$  be a prime and let  $\lambda \geq 2$ . Then there exists a  $(\lambda + 1)$ -smooth integer  $M \geq H_\lambda/(\lambda + 1)$  such that  $M \mid p^\lambda - 1$ .*

**PROOF.** We take  $M := H_\lambda/p$  if  $p$  divides  $H_\lambda$ , and otherwise  $M := H_\lambda$ . In the former case we must have  $p \leq \lambda + 1$ , so in both cases  $M \geq H_\lambda/(\lambda + 1)$ . To see that  $M \mid p^\lambda - 1$ , consider any prime divisor  $q \neq p$  of  $H_\lambda$ . Then  $q - 1 \mid \lambda$ , so  $q \mid p^\lambda - 1$  by Fermat's little theorem.  $\square$

**Remark 4.4.** The integer  $M$  constructed in Lemma 4.3 only takes into account the structure of  $H_\lambda$  and ignores  $p^\lambda - 1$  itself. In practice,  $p^\lambda - 1$  will often have small factors other than those in  $H_\lambda$ , possibly including repeated factors (which are never detected by  $H_\lambda$ ). For example,  $H_6 = 2 \cdot 3 \cdot 7$ , but  $19^6 - 1 = 2^3 \cdot 3^3 \cdot 5 \cdot 7^3 \cdot 127$ . In an implementation, one should always incorporate these highly valuable “accidental” factors into  $M$ . We will ignore them in our theoretical discussion.

Next we give a simple sieving algorithm that tabulates approximations of  $\log H_\lambda$  for all  $\lambda$  up to a prescribed bound.

**LEMMA 4.5.** *Let  $m \geq 2$ . In time  $O(m^2)$ , we may compute integers  $\ell_1, \dots, \ell_m$  with*

$$|\ell_\lambda - \log H_\lambda| \leq 1, \quad (1 \leq \lambda \leq m). \quad (9)$$

**PROOF.** Initialise a table of  $\ell_\lambda$  with  $\ell_\lambda := 0$  for  $\lambda = 1, \dots, m$ . Since  $\log H_\lambda \leq \sum_{q \leq \lambda+1} \log q = O(\lambda)$ , it suffices to set aside  $O(\lg m)$  bits for each  $\ell_\lambda$ .

For each integer  $q = 2, \dots, m$ , perform the following steps. First test whether  $q$  is prime, discarding it if not. Using trial division, the cost is  $q^{1/2+o(1)}$ , so  $m^{3/2+o(1)} = O(m^2)$  overall. Now assume that  $q$  is found to be prime. Using a fast algorithm for computing logarithms [Brent 1976], compute an integer  $L_q$  such that  $|L_q - 2^r \log q| \leq 1$ , where  $r := 1 + \lg m$ , in time  $O((\lg m)^{1+o(1)})$ . Scan through the table, replacing  $\ell_\lambda$  by  $\ell_\lambda + L_q$  for those  $\lambda$  divisible by  $q - 1$  (i.e., every  $(q - 1)$ -th entry) in time  $O(m \lg m)$ . The total cost for each prime  $q$  is  $O(m \lg m)$ , and there are  $O(m/\lg m)$  primes, so the overall cost is  $O(m^2)$ . At the end, we divide each  $\ell_\lambda$  by  $2^r$ , yielding the required approximations satisfying (9).  $\square$

**PROOF OF THEOREM 4.1.** We are given  $p$  and  $n$  as input. Applying Lemma 4.2 with  $k := n^3$ , we find that for large enough  $n$  there exists some  $\lambda$  with

$$(\log(n^3))^{c_4 \log \log \log(n^3)} < \lambda < (\log(n^3))^{c_5 \log \log \log(n^3)}$$

and such that  $H_\lambda \geq n^{3/2}$ . Choose any positive  $c_2 < c_4 \log 2$  and any  $c_3 > c_5 \log 2$ . Since

$$\log \log(n^3) \log \log \log(n^3) \sim (\log 2) \log \lg n \lg \lg \lg n,$$

we have

$$(\lg n)^{c_2 \lg \lg \lg n} < (\log(n^3))^{c_4 \log \log \log(n^3)} < \lambda < (\log(n^3))^{c_5 \log \log \log(n^3)} < (\lg n)^{c_3 \lg \lg \lg n}$$

for large  $n$ . Using Lemma 4.5, we may find one such  $\lambda$  in time  $O(m^2)$ , where  $m := (\lg n)^{c_3 \lg \lg \lg n}$ .

Let  $M$  be as in Lemma 4.3. Then  $M$  divides  $p^\lambda - 1$ , and

$$M \geq \frac{H_\lambda}{\lambda + 1} \geq \frac{n^{3/2}}{m + 1} \geq n$$

for large  $n$ . We may compute the prime factorisation of  $M$  by simply enumerating the primes  $q \leq \lambda + 1$ ,  $q \neq p$  and checking whether  $q - 1$  divides  $\lambda$  for each  $q$ . This can be done in time  $O(\lambda^2) = O(m^2)$ .  $\square$

The main multiplication algorithm depends on a reduction of a “long” DFT to many “short” DFTs. It is essential to have some control over the long and short transform lengths. The following result packages together the prime divisors of the previously mentioned  $M$  to obtain a long transform length  $N$  near a given target  $L$  and short transform lengths  $N_i$  near a given target  $S$ .

**THEOREM 4.6.** *Let  $p, n, \lambda, M$  be as in Theorem 4.1. Let  $L$  and  $S$  be positive integers such that  $\lambda < S < L < M$ . Then there exist  $(\lambda + 1)$ -smooth integers  $N_1, \dots, N_d$  with the following properties:*

- (a)  $N := N_1 \cdots N_d$  divides  $M$  (and hence divides  $p^\lambda - 1$ ).
- (b)  $L \leq N \leq (\lambda + 1)L$ .
- (c)  $S \leq N_i \leq S^3$  for all  $i$ .

Given  $\lambda, S, L$ , and the prime factorisation of  $M$ , we may compute such  $N_1, \dots, N_d$  (and their factorisations) in time  $\tilde{O}(\lambda^3)$ .

**PROOF.** Let  $M = N_1 \cdots N_s$  be the prime decomposition of  $M$ . Taking  $d$  minimal with  $N_1 \cdots N_d \geq L$ , we ensure that (a) and (b) are satisfied. At this stage, we have  $N_i \leq \lambda + 1 \leq S$  for all  $i$ . As long as the tuple  $(N_1, \dots, N_d)$  contains an entry  $N_i$  with  $N_i < S$ , we pick the two smallest entries  $N_i$  and  $N_j$  and replace them with a single entry  $N_i N_j$ . Clearly, this does not alter the product  $N$  of all entries, so (a) and (b) continue to hold. Furthermore, as long as there exist two entries  $N_i$  and  $N_j$  with  $N_i < S$  and  $N_j < S$ , new entries  $N_i N_j$  will always be smaller than  $S^2$ . Only at the very last step of the loop might the second smallest entry  $N_j$  be larger than  $S$ . In that case, the product of the entry  $N_i$  with  $N_i < S$  with  $N_j$  is still bounded by  $S \cdot S^2 = S^3$ . This shows that condition (c) is satisfied at the end of the loop.

Determining  $d$  requires at most  $s = O(\lambda)$  multiplications of integers less than  $M$ . There are at most  $s$  iterations of the replacement loop. Each iteration must scan through at most  $s$  integers of bit size  $O(\log M) = O(\lambda)$  and perform one multiplication on such integers.  $\square$

## 5. LOGARITHMICALLY SLOW RECURRENCE INEQUALITIES

In this section, we briefly recall some facts about “logarithmically slow functions” from Harvey et al. [2016a].

Let  $\Phi : (x_0, \infty) \rightarrow \mathbb{R}$  be a smooth increasing function for some  $x_0 \in \mathbb{R}$ . We say that  $\Phi^* : (x_0, \infty) \rightarrow \mathbb{R}^{\geq}$  is an *iterator* of  $\Phi$  if  $\Phi^*$  is increasing and if

$$\Phi^*(x) = \Phi^*(\Phi(x)) + 1 \tag{10}$$

for all sufficiently large  $x$ .

For instance, the standard iterated logarithm  $\log^*$  defined in (2) is an iterator of  $\log$ . An analogous iterator may be defined for any smooth increasing function  $\Phi : (x_0, \infty) \rightarrow \mathbb{R}$  for which there exists some  $\sigma \geq x_0$  such that  $\Phi(x) \leq x - 1$  for all  $x > \sigma$ . Indeed, in that case,

$$\Phi^*(x) := \min\{k \in \mathbb{N} : \Phi^{\circ k}(x) \leq \sigma\}$$

is well defined and satisfies (10) for all  $x > \sigma$ . It will sometimes be convenient to increase  $x_0$  so that  $\Phi(x) \leq x - 1$  is satisfied on the whole domain of  $\Phi$ .

We say that  $\Phi$  is *logarithmically slow* if there exists an  $\ell \in \mathbb{N}$  such that

$$(\log^{\circ \ell} \circ \Phi \circ \exp^{\circ \ell})(x) = \log x + O(1) \quad (11)$$

for  $x \rightarrow \infty$ . For example, the functions  $\log(2x)$ ,  $2 \log x$ ,  $(\log x)^2$  and  $(\log x)^{\log \log x}$  are logarithmically slow, with  $\ell = 0, 1, 2, 3$ , respectively.

**LEMMA 5.1.** *Let  $\Phi : (x_0, \infty) \rightarrow \mathbb{R}$  be a logarithmically slow function. Then there exists  $\sigma \geq x_0$  such that  $\Phi(x) \leq x - 1$  for all  $x > \sigma$ . Consequently, all logarithmically slow functions admit iterators.*

In this work, the main role played by logarithmically slow functions is to measure *size reduction* in multiplication algorithms. In other words, multiplication of objects of size  $n$  will be reduced to multiplication of objects of size  $n'$ , where  $n' \leq \Phi(n)$  for some logarithmically slow function  $\Phi(x)$ . The following result asserts that, from the point of view of iterators, such functions are more or less interchangeable with  $\log x$ .

**LEMMA 5.2.** *For any iterator  $\Phi^*$  of a logarithmically slow function  $\Phi$ , we have*

$$\Phi^*(x) = \log^* x + O(1).$$

The next result is our main tool for converting recurrence inequalities into actual asymptotic bounds for solutions.

**PROPOSITION 5.3.** *Let  $K > 1$ ,  $B \geq 0$ , and  $\ell \in \mathbb{N}$ . Let  $x_0 \geq \exp^{\circ \ell}(1)$ , and let  $\Phi : (x_0, \infty) \rightarrow \mathbb{R}$  be a logarithmically slow function such that  $\Phi(x) \leq x - 1$  for all  $x > x_0$ . Then there exists a positive constant  $C$  (depending on  $x_0$ ,  $\Phi$ ,  $K$ ,  $B$ , and  $\ell$ ) with the following property.*

*Let  $\sigma \geq x_0$  and  $L > 0$ . Let  $S \subseteq \mathbb{R}$ , and let  $T : S \rightarrow \mathbb{R}^{\geq}$  be any function satisfying the following recurrence. First,  $T(y) \leq L$  for all  $y \in S$ ,  $y \leq \sigma$ . Second, for all  $y \in S$ ,  $y > \sigma$ , there exist  $y_1, \dots, y_d \in S$  with  $y_i \leq \Phi(y)$ , and weights  $\gamma_1, \dots, \gamma_d \geq 0$  with  $\sum_i \gamma_i = 1$ , such that*

$$T(y) \leq K \left( 1 + \frac{B}{\log^{\circ \ell} y} \right) \sum_{i=1}^d \gamma_i T(y_i) + L.$$

*Then we have  $T(y) \leq CLK^{\log^* y - \log^* \sigma}$  for all  $y \in S$ ,  $y > \sigma$ .*

## 6. THE CRANDALL–FAGIN ALGORITHM

Consider the problem of computing the product of two  $n$ -bit integers modulo  $m := 2^n - 1$ . The naive approach is to compute their ordinary  $2n$ -bit product and then reduce modulo  $m$ . The reduction cost is negligible because of the special form of  $m$ . If  $n$  is divisible by a high power of two, one can save a factor of two by using the fact that FFTs naturally compute cyclic convolutions.

An ingenious algorithm of Crandall and Fagin [1994] allows for the gain of this precious factor of two for *arbitrary*  $n$ . Their algorithm is routinely used in the extreme case where  $n$  is prime, in the large-scale GIMPS search for Mersenne primes (see <http://www.mersenne.org/>).

A variant of the Crandall–Fagin algorithm was a key ingredient of our algorithm for integer multiplication that conjecturally achieves the bound  $l(n) = O(n \log n 4^{\log^* n})$  (see Section 9 of Harvey et al. [2016a]). In this section, we present yet another variant for computing products in  $\mathbb{F}_p[X]/(X^n - 1)$ .

### 6.1. The Crandall–Fagin Reduction for Polynomials

Let  $n$ ,  $N$ , and  $\kappa$  be positive integers with  $1 \leq N \leq n$  and  $\kappa \geq 2\lceil n/N \rceil$ . Our aim is to reduce multiplication in  $\mathbb{F}_p[X]/(X^n - 1)$  to multiplication in  $\mathbb{F}_{p^\kappa}[Y]/(Y^N - 1)$ . In the applications in subsequent sections,  $N$  will be a divisor of  $p^\kappa - 1$  with many small factors so that we can multiply efficiently in  $\mathbb{F}_{p^\kappa}[Y]/(Y^N - 1)$  using FFTs over  $\mathbb{F}_{p^\kappa}$ .

Suppose that  $\mathbb{F}_{p^\kappa}$  is represented as  $\mathbb{F}_p[Z]/P$  for some irreducible  $P \in \mathbb{F}_p[Z]$  of degree  $\kappa$ . The reduction relies on the existence of an element  $\theta \in \mathbb{F}_{p^\kappa}$  such that  $\theta^N = Z$ . (This is the analogue of the real  $N$ -th root of 2 appearing in the usual Crandall–Fagin algorithm, which was originally formulated over  $\mathbb{C}$ .) It is easy to see that such  $\theta$  may not always exist. For example, there is no cube root of  $Z$  in  $\mathbb{F}_{16} = \mathbb{F}_2[Z]/(Z^4 + Z + 1)$ . Nevertheless, the next result shows that if  $\kappa$  is large enough, then we can always find some modulus  $P$  for which a suitable  $\theta$  exists.

**PROPOSITION 6.1.** *Suppose that  $p^{\kappa/2} > N$ . Then we may compute an irreducible polynomial  $P \in \mathbb{F}_p[Z]$  of degree  $\kappa$ , and an element  $\theta \in \mathbb{F}_{p^\kappa} = \mathbb{F}_p[Z]/P$  such that  $\theta^N = Z$ , in time  $\tilde{O}(\kappa^9 p)$ .*

**PROOF.** We first observe that if  $\zeta \in \mathbb{F}_{p^\kappa}$  is a primitive root, then  $\zeta^N$  cannot lie in a proper subfield of  $\mathbb{F}_{p^\kappa}$ . (This property is independent of  $P$ .) Indeed, if  $\zeta^N \in \mathbb{F}_{p^\ell}$  for some proper divisor  $\ell$  of  $\kappa$ , then every  $N$ -th power in  $\mathbb{F}_{p^\kappa}$  lies in  $\mathbb{F}_{p^\ell}$ . This contradicts the fact that the number of  $N$ -th powers in  $\mathbb{F}_{p^\kappa}$  is at least  $1 + (p^\kappa - 1)/N \geq p^\kappa/N > p^{\kappa/2} \geq p^\ell$ .

Now we give an algorithm for computing  $P$  and  $\theta$ . We start by using Lemma 3.1 to compute an irreducible  $\tilde{P} \in \mathbb{F}_p[U]$  of degree  $\kappa$  in time  $\tilde{O}(\kappa^4 p^{1/2})$ , and we temporarily represent  $\mathbb{F}_{p^\kappa}$  as  $\mathbb{F}_p[U]/\tilde{P}$ .

Using Lemma 3.2, we may construct a list of  $\tilde{O}(\kappa^6 p)$  candidates for primitive roots of  $\mathbb{F}_{p^\kappa}$  (see the proof of Lemma 3.3). For each candidate  $\zeta$ , we compute  $\zeta^N$  and its powers  $1, \zeta^N, \dots, \zeta^{(\kappa-1)N}$ . If these are linearly dependent over  $\mathbb{F}_p$ , then  $\zeta^N$  belongs to a proper subfield (and so, as shown earlier,  $\zeta$  is not actually a primitive root). Thus, we must eventually encounter some  $\zeta$  for which they are linearly independent. We take  $P$  to be the minimal polynomial of this  $\zeta^N$ , which may be computed by finding a linear relation among  $1, \zeta^N, \dots, \zeta^{\kappa N}$ . Let  $\varphi: \mathbb{F}_p[Z]/P \rightarrow \mathbb{F}_p[U]/\tilde{P}$  be the isomorphism that sends  $Z$  to  $\zeta^N$ . The matrix of  $\varphi$  with respect to the standard bases  $1, Z, \dots, Z^{\kappa-1}$  and  $1, U, \dots, U^{\kappa-1}$  is given by the coefficients of  $1, \zeta^N, \dots, \zeta^{(\kappa-1)N}$ . The inverse of this matrix yields the matrix of  $\varphi^{-1}: \mathbb{F}_p[U]/\tilde{P} \rightarrow \mathbb{F}_p[Z]/P$ . We then set  $\theta = \varphi^{-1}(\zeta)$  so that  $\theta^N = \varphi^{-1}(\zeta^N) = Z$ .

For each candidate  $\zeta$ , the cost of computing the necessary powers of  $\zeta$  is  $\tilde{O}((\kappa + \lg N)\kappa \lg p) = \tilde{O}(\kappa^2 \lg^2 p)$ , and the various linear algebra steps require time  $\tilde{O}(\kappa^3 \lg p)$  using classical matrix arithmetic.  $\square$

In the remainder of this section, we fix some  $P$  and  $\theta$  as in Proposition 6.1 and assume that  $\mathbb{F}_{p^\kappa}$  is represented as  $\mathbb{F}_p[Z]/P$ . Suppose that we wish to compute the product of  $u, v \in \mathbb{F}_p[X]/(X^n - 1)$ . The presentation here closely follows that of Section 9.2 in Harvey et al. [2016a]. We decompose  $u$  and  $v$  as

$$u = \sum_{i=0}^{N-1} u_i X^{ei}, \quad v = \sum_{i=0}^{N-1} v_i X^{ei}, \quad (12)$$



where

$$e_i := \lceil ni/N \rceil, \\ u_i, v_i \in \mathbb{F}_p[X]_{e_{i+1}-e_i}.$$

Notice that  $e_{i+1} - e_i$  takes only two possible values:  $\lfloor n/N \rfloor$  or  $\lceil n/N \rceil$ .

For  $0 \leq i < N$ , let

$$c_i := Ne_i - ni \quad (13)$$

so that  $0 \leq c_i < N$ . For any  $0 \leq i_1, i_2 < N$ , define  $\delta_{i_1, i_2} \in \mathbb{Z}$  as follows. Choose  $\sigma \in \{0, 1\}$  so that  $i := i_1 + i_2 - \sigma N$  lies in the interval  $0 \leq i < N$ , and put

$$\delta_{i_1, i_2} := e_{i_1} + e_{i_2} - e_i - \sigma n.$$

From (13), we have

$$c_{i_1} + c_{i_2} - c_i = N(e_{i_1} + e_{i_2} - e_i) - n(i_1 + i_2 - i) = N\delta_{i_1, i_2}.$$

Since the left-hand side lies in the interval  $(-N, 2N)$ , this shows that  $\delta_{i_1, i_2} \in \{0, 1\}$ . Now, since  $e_{i_1} + e_{i_2} = e_i + \delta_{i_1, i_2} \pmod{n}$ , we have

$$uv = \sum_{i_1=0}^{N-1} \sum_{i_2=0}^{N-1} u_{i_1} v_{i_2} X^{e_{i_1}+e_{i_2}} = \sum_{i=0}^{N-1} w_i X^{c_i} \pmod{X^n - 1},$$

where

$$w_i := \sum_{i_1+i_2=i \pmod{N}} X^{\delta_{i_1, i_2}} u_{i_1} v_{i_2}.$$

Since  $u_{i_1} \in \mathbb{F}_p[X]_{\lceil n/N \rceil}$  and similarly for  $v_{i_2}$ , we have  $w_i \in \mathbb{F}_p[X]_{2\lceil n/N \rceil}$ . Note that we may recover  $uv$  from  $w_0, \dots, w_{N-1}$  in time  $O(n \lg p)$  by a standard overlap-add procedure (provided that  $N = O(n/\lg n)$ ).

Now, regarding  $u_i$  and  $v_i$  as elements of  $\mathbb{F}_{p^k}$  by sending  $X$  to  $Z$ , define polynomials  $U, V \in \mathbb{F}_{p^k}[Y]/(Y^N - 1)$  by  $U_i := \theta^{c_i} u_i$  and  $V_i := \theta^{c_i} v_i$  for  $0 \leq i < N$ , and let  $W = W_0 + \dots + W_{N-1} Y^{N-1} := UV$  be their (cyclic) product. Then

$$\tilde{w}_i := \theta^{-c_i} W_i = \sum_{i_1+i_2=i \pmod{N}} \theta^{-c_i} U_{i_1} V_{i_2} = \sum \theta^{c_{i_1}+c_{i_2}-c_i} u_{i_1} v_{i_2} = \sum Z^{\delta_{i_1, i_2}} u_{i_1} v_{i_2}$$

coincides with the reinterpretation of  $w_i$  as an element of  $\mathbb{F}_{p^k}$ . Moreover, we may recover  $w_i$  unambiguously from  $\tilde{w}_i$ , as  $\kappa \geq 2\lceil n/N \rceil$  and  $w_i \in \mathbb{F}_p[X]_{2\lceil n/N \rceil}$ . All together, this shows how to reduce multiplication in  $\mathbb{F}_p[X]/(X^n - 1)$  to multiplication in  $\mathbb{F}_{p^k}[Y]/(Y^N - 1)$ .

**Remark 6.2.** The pair  $(e_{i+1}, c_{i+1})$  can be computed from  $(e_i, c_i)$  in  $O(\lg n)$  bit operations, so we may compute the sequences  $e_0, \dots, e_{N-1}$  and  $c_0, \dots, c_{N-1}$  in time  $O(N \lg n)$ . Moreover, since  $c_{i+1} - c_i$  takes on only two possible values, we may compute the sequence  $\theta^{c_0}, \dots, \theta^{c_{N-1}}$  using  $O(N)$  multiplications in  $\mathbb{F}_{p^k}$ .

## 7. THE MAIN ALGORITHM

Consider the problem of computing  $t \geq 1$  products  $u_1 v, \dots, u_t v$  with  $u_1, \dots, u_t, v \in \mathbb{F}_p[X]/(X^n - 1)$  (i.e.,  $t$  products with one fixed operand). Denote the cost of this operation by  $C_{p,t}(n)$ . Our algorithm for this problem will perform  $t+1$  forward DFTs and  $t$  inverse DFTs, so it is convenient to introduce the normalization

$$C_p(n) := \sup_{t \geq 1} \frac{C_{p,t}(n)}{2t+1}.$$

This is well defined since clearly  $C_{p,t}(n) \leq tC_{p,1}(n)$ . Roughly speaking,  $C_p(n)$  may be thought of as the notional cost of a single DFT.

The problem of multiplying two polynomials in  $\mathbb{F}_p[X]$  of degree less than  $k$  may be reduced to the preceding problem by using zero padding (i.e., taking  $n := 2k$  and  $t := 1$ ). Since  $C_{p,1}(n) \leq 3C_p(n)$ , we obtain  $M_p(k) \leq 3C(2k) + O(k \lg p)$ . Thus, it suffices to obtain a good bound for  $C_p(n)$ .

The next theorem gives the core of the new algorithm for the case that  $n$  is large relative to  $p$ .

**THEOREM 7.1.** *There exist  $x_0 \geq 2$  and a logarithmically slow function  $\Phi : (x_0, \infty) \rightarrow \mathbb{R}$  with the following property. For all integers  $n > x_0$ , there exist integers  $n_1, \dots, n_d \leq \Phi(n)$  and weights  $\gamma_1, \dots, \gamma_d \geq 0$  with  $\sum_i \gamma_i = 1$  such that*

$$\frac{C_p(n)}{n \lg p \lg(n \lg p)} \leq \left(8 + O\left(\frac{1}{\lg \lg n}\right)\right) \sum_{i=1}^d \gamma_i \frac{C_p(n_i)}{n_i \lg p \lg(n_i \lg p)} + O(1) \quad (14)$$

uniformly for  $n > \max(x_0, p^2)$ .

**PROOF.** We wish to compute  $t \geq 1$  products  $u_1 v, \dots, u_t v$  with  $u_1, \dots, u_t, v \in \mathbb{F}_p[X]/(X^n - 1)$  for some sufficiently large  $n$ . We assume throughout that  $p^2 < n$ .

*Choose parameters.* Using Theorem 4.1, we obtain integers  $\lambda$  and  $M$  with

$$(\lg n)^{c_2 \lg \lg \lg n} < \lambda < (\lg n)^{c_3 \lg \lg \lg n}$$

so that  $M \geq n$ ,  $M \mid p^\lambda - 1$ , and  $M$  is  $(\lambda + 1)$ -smooth. We choose long and short target transform lengths  $L := \lceil n/\lambda^3 \rceil$  and  $S := \lambda^{(\lg \lg n)^2}$ . For large enough  $n$ , we then have  $\lambda < S < L < n \leq M$ , so we may apply Theorem 4.6. This yields  $(\lambda + 1)$ -smooth integers  $N_1, \dots, N_d$ , with known factorisations, such that  $N := N_1 \cdots N_d$  divides  $p^\lambda - 1$  and lies in the range  $L \leq N \leq (\lambda + 1)L$ , and such that  $S \leq N_i \leq S^3$  for all  $i$ . Finally, we set  $\kappa := \lceil 2\lceil n/N \rceil / \lambda \rceil \lambda$  so that  $N$  also divides  $p^\kappa - 1$ , and we put  $n_i := 2N_i \kappa$ . All of these parameters may be computed in time  $O((\lg n)^{\lg \lg n}) = O(n)$ .

For large  $n$ , we observe that the following estimates hold. First, we have

$$\frac{n}{\lambda^3} \leq N \leq \frac{2n}{\lambda^2}.$$

Since  $\log \lambda \asymp \lg \lg n \lg \lg \lg n$ , it follows that

$$\log_2 N = \left(1 + O\left(\frac{1}{\lg \lg n}\right)\right) \lg n,$$

and also that  $\log \kappa \asymp \lg \lg n \lg \lg \lg n$ . We have  $N\kappa = 2n + O(N\lambda) = 2n + O(n/\lambda)$ , so

$$N\kappa = \left(2 + O\left(\frac{1}{\lg n}\right)\right) n.$$

To estimate  $d$ , note that  $N \geq S^d$  and  $\log S \asymp (\lg \lg n)^3 \lg \lg \lg n$ , so

$$d = O\left(\frac{\lg n}{(\lg \lg n)^3 \lg \lg \lg n}\right).$$

Since  $p \leq n$ , we have

$$\lg(n \lg p) = \left(1 + O\left(\frac{1}{\lg \lg n}\right)\right) \lg n, \quad (15)$$

and as  $n_i \geq N_i \geq S \geq (\lg n)^{\lg \lg n}$ , we similarly obtain

$$\lg(n_i \lg p) = \left(1 + O\left(\frac{1}{\lg \lg n}\right)\right) \lg n_i. \quad (16)$$

*Crandall–Fagin reduction.* Let us check that the hypotheses of Section 6.1 are satisfied to enable the reduction to multiplication in  $\mathbb{F}_{p^\kappa}[Y]/(Y^N - 1)$ . We certainly have  $1 \leq N \leq n$  and  $\kappa \geq 2\lceil n/N \rceil$ . For Proposition 6.1, observe that  $\kappa \geq \lambda^2 > \lg n \sim \lg N$ , so  $p^{\kappa/2} \geq 2^{\kappa/2} > N$  for large  $n$ . Thus, we obtain an irreducible  $P \in \mathbb{F}_p[Z]$  of degree  $\kappa$ , and an element  $\theta \in \mathbb{F}_{p^\kappa} = \mathbb{F}_p[Z]/P$  such that  $\theta^N = Z$ , in time  $\tilde{O}(\kappa^9 p) = \tilde{O}(\lambda^{27} n^{1/2}) = O(n)$ . Each multiplication in  $\mathbb{F}_{p^\kappa}$  has cost  $O(S_p(\kappa))$  (see Section 3).

Computing the sequences  $e_i$  and  $c_i$  costs  $O(N \lg n) = O(n \lg n)$ , and computing the sequence  $\theta^{e_i}$  costs  $O(NS_p(\kappa))$ . The initial splitting and final overlap-add phases require time  $O(tn \lg p)$ , and the multiplications by  $\theta^{e_i}$  and  $\theta^{-c_i}$  cost  $O(tNS_p(\kappa))$ .

*Long transforms.* The factorisation of  $N$  is known, and  $N$  divides  $p^\kappa - 1$ , so by Lemma 3.3 we may compute a primitive  $N$ -th root of unity  $\omega \in \mathbb{F}_{p^\kappa}$  in time  $\tilde{O}(\kappa^9 p) = O(n)$ .

We will multiply in  $\mathbb{F}_{p^\kappa}[Y]/(Y^N - 1)$  by using DFTs with respect to  $\omega$ . The table of roots  $1, \omega, \dots, \omega^{N-1}$  may be computed in time  $O(NS_p(\kappa))$ . In a moment, we will describe an algorithm  $\mathcal{A}_i$  for computing a “short” DFT of length  $N_i$  with respect to  $\omega_i := \omega^{N/N_i}$ ; we then use the algorithm  $\mathcal{A} := \mathcal{A}_1 \odot \dots \odot \mathcal{A}_d$  for the main transform of length  $N$  (see Section 2.3). The corresponding twiddle factor tables may be extracted in time  $O(N \lg N(\kappa \lg p + \lg N)) = O(n \lg n \lg p)$ .

Let  $D$  denote the complexity of  $\mathcal{A}$ , and for  $t' \geq 1$  let  $D_{i,t'}$  denote the cost of performing  $t'$  independent DFTs of length  $N_i$  using  $\mathcal{A}_i$ . Then by (7), we have

$$D \leq \sum_{i=1}^d D_{i,N/N_i} + O(dNS_p(\kappa)) + O((\kappa \lg p)N \lg N).$$

The last term is simply  $O(n \lg n \lg p)$ .

*Bluestein conversion.* We now begin constructing  $\mathcal{A}_i$ , assuming that  $t' \geq 1$  independent transforms are sought. We first use Bluestein’s algorithm (Section 2.5) to convert each DFT of length  $N_i$  to a multiplication in  $\mathbb{F}_{p^\kappa}[X]/(X^{N_i} - 1)$ . We must check that 2 is invertible in  $\mathbb{F}_{p^\kappa}$  if  $N_i$  is even; indeed, if  $N_i$  is even, then so is  $p^\kappa - 1$ , so  $p \neq 2$ . The Bluestein conversion contributes  $O(t'N_iS_p(\kappa))$  to the cost of  $\mathcal{A}_i$ .

We must also compute a suitable table of roots, once at the top level. We first extract the table  $1, \omega_i, \dots, \omega_i^{N_i-1}$  from the top-level table in time  $O(N\kappa \lg p) = O(n \lg p)$  and then sort them into the correct order (and perform any necessary additions) in time  $O(N_i \lg N_i(\kappa \lg p + \lg N)) = O(S^3 \lg N_i \kappa \lg p) = O(\lg N_i(n \lg p))$ . Over all  $i$ , the cost is  $O(\lg N(n \lg p)) = O(n \lg n \lg p)$ .

*Kronecker substitution.* We finally convert each multiplication in  $\mathbb{F}_{p^\kappa}[X]/(X^{N_i} - 1)$  to a multiplication in  $\mathbb{F}_p[X]/(X^{N_i} - 1)$  using Kronecker substitution (see Section 2.6). The latter multiplications have cost  $C_{p,t'}(n_i)$ , since one argument is fixed. After the multiplications, we must also perform  $t'N_i$  divisions by  $P$  to recover the results in  $\mathbb{F}_{p^\kappa}[X]/(X^{N_i} - 1)$  at cost  $O(t'N_iS_p(\kappa))$ . Consolidating the estimates for the Bluestein conversion and Kronecker substitution, we have

$$\begin{aligned} D_{i,t'} &\leq C_{p,t'}(n_i) + O(t'N_iS_p(\kappa)) \\ &\leq (2t' + 1)C_p(n_i) + O(t'N_iS_p(\kappa)). \end{aligned}$$

For  $t' = N/N_i \geq N/S^3 > \lg \lg n$ , this becomes

$$D_{i,N/N_i} \leq (2 + O(1/\lg \lg n))(N/N_i)C_p(n_i) + O(NS_p(\kappa)).$$

*Conclusion.* Summing over  $i$  yields

$$D \leq \left(2 + O\left(\frac{1}{\lg \lg n}\right)\right) \sum_{i=1}^d \frac{N}{N_i} C_p(n_i) + O(dNS_p(\kappa)) + O(n \lg n \lg p).$$

Since

$$\begin{aligned} dNS_p(\kappa) &= O\left(\frac{\lg n}{(\lg \lg n)^3 \lg \lg n} N \kappa \lg \kappa \lg \lg \kappa \lg p \lg \lg p \lg \lg \lg p\right) \\ &= O\left(n \lg n \lg p \frac{(\lg \lg \lg n)^2}{\lg \lg n}\right) \\ &= O(n \lg n \lg p) \end{aligned}$$

and

$$\frac{N}{N_i} = \frac{2N\kappa}{2N_i\kappa} = \left(4 + O\left(\frac{1}{\lg n}\right)\right) \frac{n}{n_i},$$

this becomes

$$D \leq \left(8 + O\left(\frac{1}{\lg \lg n}\right)\right) \sum_{i=1}^d \frac{n}{n_i} C_p(n_i) + O(n \lg n \lg p).$$

Let  $\gamma_i := \log N_i / \log N$  so that  $\sum_i \gamma_i = 1$ . Since

$$\lg n_i = \log_2 N_i + O(\log \kappa) = \left(1 + O\left(\frac{\log \lambda}{\log S}\right)\right) \log_2 N_i = \left(1 + O\left(\frac{1}{\lg \lg n}\right)\right) \log_2 N_i$$

and  $\lg n = (1 + O(1/\lg \lg n)) \log_2 N$ , we get

$$D \leq \left(8 + O\left(\frac{1}{\lg \lg n}\right)\right) \sum_{i=1}^d \gamma_i \frac{n \lg n}{n_i \lg n_i} C_p(n_i) + O(n \lg n \lg p).$$

To compute the desired  $t$  products, we must execute  $t + 1$  forward transforms and  $t$  inverse transforms. Each product also requires  $O(N)$  pointwise multiplications in  $\mathbb{F}_{p^r}$  and  $O(N)$  multiplications by  $1/N$ . These have cost  $O(NS_p(\kappa))$ , which is absorbed by the  $O(n \lg n \lg p)$  term. Thus, we obtain

$$C_{p,t}(n) \leq (2t + 1)D + O(tn \lg n \lg p).$$

Dividing by  $(2t + 1)n \lg n \lg p$ , taking suprema over  $t \geq 1$ , and using (15) and (16), yields the bound (14).

Finally, since

$$n_i = O(S^3 \lambda^3) = O((\lg n)^{3c_3 \lg \lg n (\lg \lg n^2 + 1)}) = O((\lg n)^{(\lg \lg n)^3}),$$

we have  $\log n_i = O((\log \log n)^4)$  and hence  $\log \log \log n_i \leq \log \log \log \log n + C$  for some constant  $C$  and all large  $n$ . We may then take  $\Phi(x) := \exp^{o^3}(\log^{o^4} x + C)$ .  $\square$

Now we may prove the main theorem announced in Section 1.

**PROOF OF THEOREM 1.1.** For  $n \geq 2$ , define

$$T_p(n) := \frac{C_p(n)}{n \lg p \lg(n \lg p)}.$$

It suffices to prove that  $T_p(n) = O(8^{\log^*(n \lg p)})$  uniformly in  $n$  and  $p$ .

Let  $x_0$  and  $\Phi(x)$  be as in Theorem 7.1. Increasing  $x_0$  if necessary, we may assume that  $\Phi(x) \leq x - 1$  for  $x > x_0$  and that  $x_0 \geq \exp(\exp(1))$ .

Let  $\sigma_p := \max(x_0, p^2)$  for each  $p$ . We will consider the regions  $n \leq \sigma_p$  and  $n > \sigma_p$  separately. First consider the case  $n \leq \sigma_p$ . There are only finitely many primes  $p < (x_0)^{1/2}$ , so we may assume that  $p^2 \geq x_0$  and that  $n \leq \sigma_p = p^2$ . In this region, we use Kronecker substitution: by (8) and (1), we have

$$M_p(n) = O(\lg(n \lg p)) = O(n \lg p \lg(n \lg p) 8^{\log^*(n \lg p)}),$$

and since

$$C_p(n) = \sup_{t \geq 1} \frac{C_{p,t}(n)}{2t+1} \leq \sup_{t \geq 1} \frac{tC_{p,1}(n)}{2t+1} = O(C_{p,1}(n)) = O(M_p(n)),$$

we get  $T_p(n) = O(8^{\log^*(n \lg p)})$  uniformly for  $n \leq \sigma_p$ . In fact, this even shows that  $T_p(n) = O(8^{\log^* p})$  uniformly for  $n \leq \sigma_p$ .

Now consider the case  $n > \sigma_p$ . Here we invoke Theorem 7.1 to obtain absolute constants  $B, L > 0$  such that for every  $n > \sigma_p$ , there exist  $n_1, \dots, n_d \leq \Phi(n)$  and  $\gamma_1, \dots, \gamma_d$  such that

$$T_p(n) \leq 8 \left( 1 + \frac{B}{\log \log n} \right) \sum_{i=1}^d \gamma_i T_p(n_i) + L.$$

Set  $L_p := \max(L, \max_{2 \leq n \leq \sigma_p} T_p(n))$ . Applying Proposition 5.3 with  $\ell := 2$ ,  $K := 8$ ,  $S := \{2, 3, \dots\}$ , we find that  $T_p(n) = O(L_p 8^{\log^* n - \log^* \sigma_p})$  uniformly for  $n > \sigma_p$ . Since  $\log^* \sigma_p = \log^* p + O(1)$  and  $L_p = O(8^{\log^* p})$ , we conclude that  $T_p(n) = O(8^{\log^* n}) = O(8^{\log^*(n \lg p)})$  uniformly for  $n > \sigma_p$ .  $\square$

## 8. NOTES AND GENERALIZATIONS

In this section, we outline some directions along which the results in this article can be extended. We also provide some hints concerning the practical usefulness of the new ideas. Our treatment is more sketchy, and we plan to provide more details in a forthcoming work.

### 8.1. Multiplication of Polynomials over $\mathbb{F}_{p^\kappa}$

Recall that  $M_{p^\kappa}(n)$  denotes the cost of multiplying polynomials in  $\mathbb{F}_{p^\kappa}[X]$  of degree less than  $n$ , where we assume that some model  $\mathbb{F}_p[Z]/P$  for  $\mathbb{F}_{p^\kappa}$  has been fixed in advance.

**THEOREM 8.1.** *We have*

$$M_{p^\kappa}(n) = O(n\kappa \lg p \lg(n\kappa \lg p) 8^{\log^*(n\kappa \lg p)})$$

*uniformly for all  $n, \kappa \geq 1$  and all primes  $p$ .*

Indeed, we saw in Section 3 that

$$M_{p^\kappa}(n) = M_p(2n\kappa) + O(nD_p(\kappa)),$$

where  $D_p(\kappa)$  denotes the cost of dividing a polynomial of degree less than  $2\kappa$  by  $P$ . Having established the bound  $M_p(n) = O(n \lg p \lg(n \lg p) 8^{\log^*(n \lg p)})$ , it is now permissible to assume that  $M_p(n)/n$  is increasing, so the usual argument for Newton iteration shows that  $D_p(\kappa) = O(M_p(\kappa))$ . Using again that  $M_p(n)/n$  is increasing, we obtain  $nD_p(\kappa) = O(M_p(n\kappa))$ .

### 8.2. Multiplication of Polynomials over $\mathbb{Z}/p^\alpha\mathbb{Z}$

For any prime  $p$  and any integer  $\alpha \geq 1$ , denote by  $M_{p,\alpha}(n)$  the bit complexity of multiplying polynomials in  $(\mathbb{Z}/p^\alpha\mathbb{Z})[X]_n$ .

**THEOREM 8.2.** *We have*

$$M_{p,\alpha}(n) = O(n\alpha \lg p \lg(n\alpha \lg p) 8^{\log^*(n\alpha \lg p)})$$

*uniformly for all  $n, \alpha \geq 1$  and all primes  $p$ .*

For  $\lg n = O(\alpha \lg p)$ , we may simply use Kronecker substitution. For  $\alpha \lg p = O(\lg n)$ , we must modify the algorithm of Section 7. Recall that in that algorithm we reduced multiplication in  $\mathbb{F}_p[X]/(X^n - 1)$  to multiplication  $\mathbb{F}_{p^\kappa}[Y]/(Y^N - 1)$ , where  $N$  is a transform length dividing  $p^\kappa - 1$ . Our task is to define a ring  $R$ , analogous to  $\mathbb{F}_{p^\kappa}$ , so that multiplication in  $(\mathbb{Z}/p^\alpha\mathbb{Z})[X]/(X^n - 1)$  can be reduced to multiplication in  $R[Y]/(Y^N - 1)$ .

Let  $\pi$  be the natural projection  $\mathbb{Z}/p^\alpha\mathbb{Z} \rightarrow \mathbb{F}_p$ , and write  $\pi$  also for the corresponding map  $(\mathbb{Z}/p^\alpha\mathbb{Z})[Z] \rightarrow \mathbb{F}_p[Z]$ . Let  $P \in \mathbb{F}_p[Z]$  be any monic irreducible polynomial of degree  $\kappa$ , and let  $\tilde{P} \in (\mathbb{Z}/p^\alpha\mathbb{Z})[Z]$  be an arbitrary lift via  $\pi$ , also monic of degree  $\kappa$ . Then  $\tilde{P}$  is irreducible, and we will take  $R := (\mathbb{Z}/p^\alpha\mathbb{Z})[Z]/\tilde{P}$ .

Any primitive  $N$ -th root of unity  $\omega \in \mathbb{F}_{p^\kappa}$  has a unique lift to a principal  $N$ -th root of unity  $\tilde{\omega} \in R$ . (This can be seen, for example, by observing that  $\tilde{\omega}$  must be the image in  $R$  of a Teichmüller lift of  $\omega$  in the  $p$ -adic field whose residue field is  $\mathbb{F}_{p^\kappa}$ .) Given  $\omega$ , we may compute  $\tilde{\omega}$  efficiently using fast Newton lifting (see Section 12.3 in Cohen et al. [2006]).

Moreover, if we choose  $P$  and  $\theta \in \mathbb{F}_{p^\kappa} = \mathbb{F}_p[Z]/P$  as in Proposition 6.1, so that  $\theta^N = Z$ , then fast Newton lifting can also be used to obtain  $\tilde{\theta} \in R$  such that  $\tilde{\theta}^N = Z$  in  $R$ .

### 8.3. Multiplication of Polynomials over $\mathbb{Z}/m\mathbb{Z}$

For any  $m \geq 1$ , denote by  $M_m(n)$  the bit complexity of multiplying polynomials in  $(\mathbb{Z}/m\mathbb{Z})[X]_n$ .

**THEOREM 8.3.** *We have*

$$M_m(n) = O(n \lg m \lg(n \lg m) 8^{\log^*(n \lg m)})$$

*uniformly for all  $n, m \geq 1$ .*

We use the isomorphism  $\mathbb{Z}/m\mathbb{Z} \cong \prod_i (\mathbb{Z}/p_i^{\alpha_i}\mathbb{Z})$ , where  $m = p_1^{\alpha_1} \cdots p_i^{\alpha_i}$  is the prime decomposition of  $m$ . The cost of converting between the  $\mathbb{Z}/m\mathbb{Z}$  and  $\prod_i (\mathbb{Z}/p_i^{\alpha_i}\mathbb{Z})$  representations is  $O(l(\lg m) \lg l) = O(l(\lg m) \lg \lg m)$  (see Section 10.3 in von zur Gathen and Gerhard [2002]). Using Theorem 8.2, it follows that

$$M_m(n) = O(n \lg m \lg(n \lg m) 8^{\log^*(n \lg m)} + nl(\lg m) \lg \lg m).$$

The first term dominates if  $n \geq m$ . If  $n < m$ , we may simply use Kronecker substitution.

### 8.4. Complexity Bounds for Straight-Line Programs

Until now, we have considered only complexity bounds in the Turing model. The new techniques also lead to improved bounds in algebraic complexity models. In what follows,  $\mathcal{A}$  is always a ring with identity.

For the simplest case, first assume that  $\mathcal{A}$  is an  $\mathbb{F}_p$ -algebra for some prime  $p$ . In the straight-line program model (see Chapter 4 in Bürgisser et al. [1997]), we count the number of additions and subtractions in  $\mathcal{A}$ , the number of scalar multiplications (multiplications by elements of  $\mathbb{F}_p$ ), and the number of nonscalar multiplications (multiplications in  $\mathcal{A}$ ).



**THEOREM 8.4.** *Let  $\mathcal{A}$  be an  $\mathbb{F}_p$ -algebra. We may multiply two polynomials in  $\mathcal{A}[X]$  of degree less than  $n$  using  $O(n \lg n 8^{\lg^* n})$  additions, subtractions, and scalar multiplications, and  $O(n 4^{\lg^* n})$  nonscalar multiplications. These bounds are uniform over all primes  $p$  and all  $\mathbb{F}_p$ -algebras  $\mathcal{A}$ .*

The idea of the proof is to use the same algorithm as in Section 7, but instead of switching to Kronecker substitution when we reach  $n \approx p$ , we simply recurse all the way down to  $n = 1$ . The role of the extension  $\mathbb{F}_{p^k}$  is played by  $\mathcal{A} \otimes_{\mathbb{F}_p} \mathbb{F}_{p^k} \cong \mathcal{A}[Z]/P$ , where  $P \in \mathbb{F}_p[Z]$  is monic and irreducible of degree  $k$  (here we have used the fact that  $\mathbb{F}_p[X]$  may be viewed as a subring of  $\mathcal{A}[X]$ , since  $\mathcal{A}$  contains an identity element and hence a copy of  $\mathbb{F}_p$ ). Thus, we reduce multiplication in  $\mathcal{A}[X]/(X^n - 1)$  to multiplication in  $(\mathcal{A} \otimes_{\mathbb{F}_p} \mathbb{F}_{p^k})[Y]/(Y^N - 1)$ . The latter multiplication may be handled by DFTs over  $\mathcal{A} \otimes_{\mathbb{F}_p} \mathbb{F}_{p^k}$ , as any primitive  $N$ -th root of unity in  $\mathbb{F}_{p^k}$  corresponds naturally to a principal  $N$ -th root of unity in  $\mathcal{A} \otimes_{\mathbb{F}_p} \mathbb{F}_{p^k}$ .

The  $O(n \lg n 8^{\lg^* n})$  bound covers the cost of the DFTs, which are accomplished entirely using additions, subtractions, and scalar multiplications. Nonscalar multiplications are needed only for the pointwise multiplication step. To explain the  $O(n 4^{\lg^* n})$  bound, we observe that at each recursion level the total “data size” grows by a factor of four: one factor of two arises from the Crandall–Fagin splitting and another factor of two from the Kronecker substitution.

Note that in the straight-line program model, we give a different algorithm for each  $n$ . Thus, all precomputed objects, such as the defining polynomial  $P$  and the required roots of unity, are obtained free of cost (they are built directly into the structure of the algorithm). The uniformity in  $p$  follows from the fact that the bounds for  $\lambda$  and  $M$  in Theorem 4.1 are independent of  $p$ , although of course the algorithm will be different for each  $p$ . For each fixed  $p$ , we may use the same straight-line program for all  $\mathbb{F}_p$ -algebras  $\mathcal{A}$ .

Theorem 8.4 can be generalized to  $(\mathbb{Z}/p^\alpha\mathbb{Z})$ -algebras, along the same lines as in Section 8.2. It is also possible to handle  $(\mathbb{Z}/m\mathbb{Z})$ -algebras for any integer  $m \geq 1$  (i.e., rings of finite characteristic), but we cannot proceed exactly as in Section 8.3 because the straight-line model has no provision for a “reduction modulo  $p^\alpha$ ” operation. Instead, we use a device introduced in Cantor and Kaltofen [1991].

Suppose that  $m = p_1^{\alpha_1} \cdots p_l^{\alpha_l}$  and that  $\mathcal{A}$  is now a  $(\mathbb{Z}/m\mathbb{Z})$ -algebra. For each  $i$ , we may construct a straight-line program  $\mathcal{M}_i$  that takes as input polynomials  $f, g \in \mathcal{B}[X]$  of degree less than  $n$ , where  $\mathcal{B}$  is any  $(\mathbb{Z}/p_i^{\alpha_i}\mathbb{Z})$ -algebra, and computes  $fg$ . By replacing every constant in  $\mathbb{Z}/p_i^{\alpha_i}\mathbb{Z}$  by a compatible constant in  $\mathbb{Z}/m\mathbb{Z}$ , we obtain a straight-line program  $\mathcal{M}'_i$  that takes as input  $f, g \in \mathcal{A}[X]$  of degree less than  $n$  and computes  $h_i \in \mathcal{A}[X]$  such that  $h_i - fg \in p_i^{\alpha_i}\mathcal{A}[X]$ . By the Chinese remainder theorem, we may choose  $e_i \in \mathbb{Z}/m\mathbb{Z}$  such that  $e_i = 0 \pmod{m/p_i^{\alpha_i}}$  and  $e_i = 1 \pmod{p_i^{\alpha_i}}$  for each  $i$ . The linear combination  $\sum_i e_i h_i$  is then equal to  $fg$  in  $\mathcal{A}[X]$ . We conclude that we may multiply polynomials in  $\mathcal{A}[X]$  using  $O(n \lg n 8^{\lg^* n})$  additions, subtractions, and scalar multiplications (by elements of  $\mathbb{Z}/m\mathbb{Z}$ ), and  $O(n 4^{\lg^* n})$  nonscalar multiplications (i.e., multiplications in  $\mathcal{A}$ ). These bounds are not uniform in  $m$ , but for each  $m$  they are uniform over all  $(\mathbb{Z}/m\mathbb{Z})$ -algebras.

## 8.5. Other Algebraic Models

As pointed out in Section 1, the best known result prior to Theorem 8.4 was the Cantor–Kaltofen algorithm: assuming a somewhat different bilinear complexity model, it was proved in Cantor and Kaltofen [1991] that polynomials of degree less than  $n$  over an arbitrary effective ring  $\mathcal{A}$  can be multiplied using  $O(n \lg n \lg \lg n)$  additions and subtractions, and  $O(n \lg n)$  multiplications. If we replace the effective  $\mathbb{F}_p$ -algebra  $\mathcal{A}$  from Theorem 8.4 by a more general effective ring  $\mathcal{A}$  with  $p\mathcal{A} = 0$ , then multiplying

two polynomials in  $\mathcal{A}[X]$  of degree less than  $n$  requires  $O(n \lg n 8^{\lg^* n} \lg p)$  additions and subtractions, and  $O(n 4^{\lg^* n})$  multiplications. Indeed, each multiplication with a scalar in  $\mathbb{F}_p$  can be emulated using  $O(\lg p)$  additions.

Furthermore, if we wish to take into account the cost of precomputations, and give a single algorithm that works uniformly for all  $p$  and  $n$ , then we may use a more refined complexity model such as the BSS model [Blum et al. 1998]. In this model, a *machine over  $\mathcal{A}$*  is, roughly speaking, a Turing machine in which the tape cells hold elements of  $\mathcal{A}$ . Actually, we need a multitape version of the model described in Blum et al. [1998]. The machine can perform arithmetic operations on elements of  $\mathcal{A}$  in unit time, but can also manipulate data such as index variables in the same way as a Turing machine, and must deal with data layout in the same way as a Turing machine. In this model, we obtain similar bounds to Theorem 8.4 but without uniformity in  $p$ . Alternatively, we could obtain bounds uniform in  $p$  if we added extra terms to account for the precomputations.

Another point of view toward Theorem 8.4 is that we have described a new evaluation-interpolation strategy for polynomials over an  $\mathbb{F}_p$ -algebra  $\mathcal{A}$ . We refer the reader to Sections 2.1 through 2.4 in van der Hoeven [2010] for classical examples of evaluation-interpolation schemes and to Gao and Mateer [2010] for algorithms specific to finite fields. Such schemes are characterized by two quantities: the evaluation/interpolation complexity  $E(n)$  and the number  $N(n)$  of evaluation points. The new algorithms yield the bounds  $E(n) = O(n \lg n 8^{\lg^* n})$  and  $N(n) = O(n 4^{\lg^* n})$ , and these can be used to prove complexity bounds for problems more general than polynomial multiplication. For example, we can multiply  $r \times r$  matrices with entries in  $\mathcal{A}[X]_n$  using  $O(r^2 n \lg n 8^{\lg^* n} + r^\omega n 4^{\lg^* n})$  ring operations, where  $\omega$  is an exponent of matrix multiplication. One of the main advantages of our algorithms is that  $N(n)$  is almost linear, contrary to synthetic FFT methods [Schönhage 1977; Cantor and Kaltofen 1991] derived from Schönhage–Strassen multiplication [Schönhage and Strassen 1971], which achieve only  $N(n) = O(n \lg n)$ . By working out a mixed-radix generalization of the truncated Fourier transform [van der Hoeven 2004, 2005], it should also be possible to smooth any occasional jumps in the function  $n \mapsto N(n)$ .

In the setting of bilinear complexity (see Chapter 14 in Bürgisser et al. [1997]), the new algorithms do not improve asymptotically on the best known bounds. For example, it is known [Pieltant and Randriam 2015] that for any  $n$  there exist  $\mathbb{F}_2$ -linear maps  $a_i, b_i : \mathbb{F}_2[X]_n \rightarrow \mathbb{F}_2$  and polynomials  $c_i \in \mathbb{F}_2[X]_{2n}$  for  $i \in \{1, \dots, k(n)\}$ , with  $k(n) = (189/22 + o(1))n$ , such that  $uv = \sum_{i=1}^{k(n)} a_i(u)b_i(v)c_i$  for all  $u, v \in \mathbb{F}_2[X]_n$ . The new method yields the asymptotically inferior bound  $k(n) = O(n 4^{\lg^* n})$ . Nevertheless, these bounds are asymptotic and do not take into account that our new algorithm would rather work over (say)  $\mathbb{F}_{2^{60}}$  instead of  $\mathbb{F}_2$ . In practice, the new method might therefore outperform the bilinear algorithms from Pieltant and Randriam [2015].

## ACKNOWLEDGMENTS

We would like to thank Karim Belabas for directing us to the work of Adleman et al. [1983] and Igor Shparlinski for answering some analytic number theory questions.

## REFERENCES

- L. M. Adleman and H. W. Lenstra Jr. 1986. Finding irreducible polynomials over finite fields. In *Proceedings of the 18th Annual ACM Symposium on Theory of Computing (STOC'86)*. ACM, New York, NY, 350–355.
- L. M. Adleman, C. Pomerance, and R. S. Rumely. 1983. On distinguishing prime numbers from composite numbers. *Annals of Mathematics* 117, 173–206.
- A. V. Aho, J. E. Hopcroft, and J. D. Ullman. 1974. *The Design and Analysis of Computer Algorithms*. Addison-Wesley.

- L. I. Bluestein. 1970. A linear filtering approach to the computation of discrete Fourier transform. *IEEE Transactions on Audio and Electroacoustics* 18, 4, 451–455.
- L. Blum, F. Cucker, M. Shub, and S. Smale. 1998. *Complexity and Real Computation*. Springer-Verlag.
- A. Bostan, P. Gaudry, and É. Schost. 2007. Linear recurrences with polynomial coefficients and application to integer factorization and Cartier-Manin operator. *SIAM Journal on Computing* 36, 1777–1806.
- R. P. Brent. 1976. Fast multiple-precision evaluation of elementary functions. *Journal of the ACM* 23, 2, 242–251.
- P. Bürgisser, M. Clausen, and M. A. Shokrollahi. 1997. *Algebraic Complexity Theory*. Springer-Verlag.
- D. G. Cantor and E. Kaltofen. 1991. On fast multiplication of polynomials over arbitrary algebras. *Acta Informatica* 28, 693–701.
- H. Cohen, G. Frey, R. Avanzi, Ch. Doche, T. Lange, K. Nguyen, and F. Vercauteren (Eds.). 2006. *Handbook of Elliptic and Hyperelliptic Curve Cryptography*. Chapman & Hall/CRC, Boca Raton, FL.
- J. W. Cooley and J. W. Tukey. 1965. An algorithm for the machine calculation of complex Fourier series. *Mathematics of Computation* 19, 90, 297–301.
- R. Crandall and B. Fagin. 1994. Discrete weighted transforms and large-integer arithmetic. *Mathematics of Computation* 62, 205, 305–324.
- M. Fürer. 2009. Faster integer multiplication. *SIAM Journal on Computing* 39, 3, 979–1005.
- S. Gao and T. Mateer. 2010. Additive fast Fourier transforms over finite fields. *IEEE Transactions on Information Theory* 56, 12, 6265–6272.
- J. von zur Gathen and J. Gerhard. 2002. *Modern Computer Algebra* (2nd ed.). Cambridge University Press.
- D. Harvey, J. van der Hoeven, and G. Lecerf. 2014. *Faster Polynomial Multiplication over Finite Fields*. Technical Report. arXiv:1407.3361. <http://arxiv.org/abs/1407.3361>.
- D. Harvey, J. van der Hoeven, and G. Lecerf. 2016a. Even faster integer multiplication. *Journal of Complexity* 36, 1–30.
- D. Harvey, J. van der Hoeven, and G. Lecerf. 2016b. Fast polynomial multiplication over  $\mathbb{F}_{2^{60}}$ . In *Proceedings of the ACM International Symposium on Symbolic and Algebraic Computation (ISSAC'16)*. ACM, New York, NY, 255–262.
- J. van der Hoeven. 2004. The truncated Fourier transform and applications. In *Proceedings of the 2004 International Symposium on Symbolic and Algebraic Computation (ISSAC'04)*. ACM, New York, NY, 290–296.
- J. van der Hoeven. 2005. *Notes on the Truncated Fourier Transform*. Technical Report 2005-5. Université Paris-Sud, Orsay, France.
- J. van der Hoeven. 2010. Newton's method and FFT trading. *Journal of Symbolic Computation* 45, 8, 857–878.
- D. E. Knuth. 1998. *The Art of Computer Programming*. Vol. 3: Sorting and Searching. Addison-Wesley, Reading, MA.
- C. H. Papadimitriou. 1994. *Computational Complexity*. Addison-Wesley.
- J. Pielant and H. Randriam. 2015. New uniform and asymptotic upper bounds on the tensor rank of multiplication in extensions of finite fields. *Mathematics of Computation* 84, 294, 2023–2045.
- C. M. Rader. 1968. Discrete Fourier transforms when the number of data samples is prime. *Proceedings of the IEEE* 56, 6, 1107–1108.
- K. R. Rao, D. N. Kim, and J. J. Hwang. 2010. *Fast Fourier Transform: Algorithms and Applications*. Springer-Verlag.
- A. Schönhage. 1977. Schnelle Multiplikation von Polynomen über Körpern der Charakteristik 2. *Acta Informatica* 7, 395–398.
- A. Schönhage and V. Strassen. 1971. Schnelle Multiplikation großer Zahlen. *Computing* 7, 281–292.
- V. Shoup. 1990. New algorithms for finding irreducible polynomials over finite fields. *Mathematics of Computation* 54, 189, 435–447.
- V. Shoup. 1992. Searching for primitive roots in finite fields. *Mathematics of Computation* 58, 197, 369–380.

Received August 2015; revised October 2016; accepted October 2016