

# Sooner is safer than later

Thomas A. Henzinger

Computer Science Department, Cornell University, Ithaca, NY 14853, USA

Communicated by F.B. Schneider

Received 29 May 1991

Revised 11 June 1992

## Abstract

Henzinger, T.A., Sooner is safer than later, Information Processing Letters (43) (1992) 135–141.

It has been observed repeatedly that the standard safety-liveness classification for properties of reactive systems does not fit for real-time properties. This is because the implicit “liveness” of time shifts the spectrum towards the safety side. While, for example, response—that “something good” will happen eventually—is a classical liveness property, bounded response—that “something good” will happen soon, within a certain amount of time—has many characteristics of safety. We account for this phenomenon formally by defining safety and liveness *relative* to a given condition, such as the progress of time.

**Keywords:** Safety, liveness, real time, topology, concurrency, semantics

## 1. Safety, liveness, and operationality

The behavior of a discrete reactive system can be described as an infinite string

$\sigma: \sigma_0 \sigma_1 \sigma_2 \sigma_3 \sigma_4 \dots$

over an alphabet  $\Sigma$ , which represents the states of the system. A *property*  $\Pi$  is a subset of  $\Sigma^\omega$ , the set of all infinite strings over  $\Sigma$ ; a reactive system has property  $\Pi$  iff all of its possible behaviors are contained in  $\Pi$ .

It is useful to classify properties of reactive systems into two categories, because they require qualitatively different means for their specification and verification [13]:

- A *safety* property stipulates that “nothing bad” will happen, ever, during the execution of a system. If “something bad” were to happen

during the execution, it would have to happen within a finite number of states. Thus we can formalize safety as follows:

$\Pi \subseteq \Sigma^\omega$  is a *safety* property iff for all  $\sigma \in \Sigma^\omega$ , whenever every finite prefix of  $\sigma$  can be extended to a string in  $\Pi$ , then  $\sigma \in \Pi$  [3].

- A *liveness* property stipulates that “something good” will happen, eventually, during the execution of a system. Even if “nothing good” were to happen within a finite prefix of the execution, “something good” could still happen in a later state; only if an irretrievably bad situation is reached within a finite number of states, “nothing good” will happen during the entire execution. Thus we can formalize liveness as follows:

$\Pi \subseteq \Sigma^\omega$  is a *liveness* property iff every finite prefix of a string in  $\Sigma^\omega$  can be extended to a string in  $\Pi$  [4].

There is a natural topology on  $\Sigma^\omega$  – the Cantor topology – in which the safety properties are

Correspondence to: T.A. Henzinger, Department of Computer Science, 4130 Upson Hall, Cornell University, Ithaca, NY 14853-7501, USA.

exactly the closed sets, and the liveness properties are exactly the dense sets. It follows that (1) only  $\Sigma^\omega$  itself is both a safety and a liveness property and (2) every property is the intersection of a safety property and a liveness property. Hence any correctness proof for a reactive system can be decomposed into a safety part and a liveness part.

Let us briefly sketch the standard topological construction for showing observation (2) [4], because we shall generalize it later. The construction is well-known to prove a strong formulation of the observation that is based on the following definition. We say that a safety property  $\Pi_S$  and a liveness property  $\Pi_L$  specify the property  $\Pi = \Pi_S \cap \Pi_L$  *congruously* iff every finite prefix of a string in  $\Pi_S$  can be extended to a string in  $\Pi$ . In other words, the safety part of a congruous specification is complete: the liveness part does not preclude any safe prefixes. A congruous pair  $(\Pi_S, \Pi_L)$  is called *machine closed* in [1], *feasible* in [8], and  $\Pi_L$  is called *live with respect to  $\Pi_S$*  in [9].

**Theorem 1.** (Existence of congruous specifications.) *Every property has a congruous specification.*

**Proof** (sketch). Since safety properties are closed under intersection, we can define the *closure*  $\bar{\Pi}$  of  $\Pi \subseteq \Sigma^\omega$  as the smallest safety property containing  $\Pi$ . Given a property  $\Pi$ , let  $\Pi_S$  be  $\bar{\Pi}$ . For  $\Pi_L$  take the complement of  $\Pi_S - \Pi$ . Then  $(\Pi_S, \Pi_L)$  specifies  $\Pi$  congruously.  $\square$

Congruous specifications are *operational*: a machine that incrementally generates safe execution sequences will never reach an irremedial situation from which the liveness conditions cannot be satisfied. On the other hand, a machine trying to execute an incongruous specification without look-ahead may “paint itself into a corner” from which no legal continuation is possible [8]. Examples of congruous specifications are fair transition systems; examples of formalisms that admit incongruous specifications are temporal logic and finite automata (see [17] and [19] for surveys of these formalisms).

## 2. Relative safety and liveness

Instead of looking at all strings in  $\Sigma^\omega$ , it is often useful to have a concept of safety and liveness under the assumption that, a priori, only a certain subset  $\Psi \subseteq \Sigma^\omega$  of strings are possible behaviors of a system. We call these notions *safety and liveness relative to the property  $\Psi$* :

- $\Pi \subseteq \Psi$  is a *safety property relative to  $\Psi \subseteq \Sigma^\omega$*  iff for all  $\sigma \in \Psi$ , whenever every finite prefix of  $\sigma$  can be extended to a string in  $\Pi$ , then  $\sigma \in \Pi$ .
- $\Pi \subseteq \Psi$  is a *liveness property relative to  $\Psi \subseteq \Sigma^\omega$*  iff every finite prefix of a string in  $\Psi$  can be extended to a string in  $\Pi$ .

Thus unconditional safety and liveness are safety and liveness relative to  $\Sigma^\omega$ .

The Cantor topology on  $\Sigma^\omega$  induces a topological subspace on  $\Psi \subseteq \Sigma^\omega$ , which is called the *relativization* of the  $\Sigma^\omega$ -topology to  $\Psi$  [11]. We show that the properties that are safe relative to  $\Psi$  are exactly the closed sets of the relative topology, and the properties that are live relative to  $\Psi$  are exactly the dense sets of the relative topology.

**Proposition 2.** (Relative safety.)  $\Pi \subseteq \Psi$  is a *safety property relative to  $\Psi \subseteq \Sigma^\omega$*  iff  $\bar{\Pi} \cap \Psi \subseteq \Pi$ .

**Proposition 3.** (Relative liveness.)  $\Pi \subseteq \Psi$  is a *liveness property relative to  $\Psi \subseteq \Sigma^\omega$*  iff  $\Psi \subseteq \bar{\Pi}$ .

**Proof of Propositions 2 and 3.** First observe that a string  $\sigma \in \Sigma^\omega$  is in the closure of a property  $\Pi \subseteq \Sigma^\omega$  (that is,  $\sigma \in \bar{\Pi}$ ) iff every finite prefix of  $\sigma$  can be extended to a string in  $\Pi$ . Then apply this observation to the definitions of relative safety and relative liveness.  $\square$

It follows that  $\Pi$  is safe relative to  $\Psi$  iff  $\Pi = \Pi_S \cap \Psi$  for some unconditional safety property  $\Pi_S$ . In particular, if the property  $\Pi = \Pi_S \cap \Pi_L$  is specified by a safety property  $\Pi_S$  and a liveness property  $\Pi_L$ , then  $\Pi$  is safe relative to  $\Pi_L$ . Furthermore, if the specification  $(\Pi_S, \Pi_L)$  is congruous, then  $\Pi$  is live relative to  $\Pi_S$ .

It is convenient to extend the notions of safety and liveness relative to a property  $\Psi$  to properties that are not necessarily subsets of  $\Psi$ : we say that  $\Pi \subseteq \Sigma^\omega$  is a safety (liveness) property relative to  $\Psi \subseteq \Sigma^\omega$  iff  $\Pi \cap \Psi$  is safe (live) relative to  $\Psi$ . Clearly, unconditional safety properties are, in this sense, safe relative to any property  $\Psi$ . More generally:

**Proposition 4.** (Downward preservation of safety.) *Suppose that  $\Psi_1 \subseteq \Psi_2$ . If  $\Pi$  is a safety property relative to  $\Psi_2$ , then it is also a safety property relative to  $\Psi_1$ .*

**Proof.** Let  $\Psi_1 \subseteq \Psi_2$ . First observe that the closure operator is monotonic; that is,  $\Pi \subseteq \Psi$  implies  $\overline{\Pi} \subseteq \overline{\Psi}$  for all  $\Pi, \Psi \in \Sigma^\omega$ . In particular, we have  $\overline{\Pi \cap \Psi_1} \subseteq \overline{\Pi \cap \Psi_2}$ .

By Proposition 2, we may assume that

$$(\overline{\Pi \cap \Psi_2}) \cap \Psi_2 \subseteq \Pi \cap \Psi_2$$

and need to show that, then,

$$(\overline{\Pi \cap \Psi_1}) \cap \Psi_1 \subseteq \Pi \cap \Psi_1.$$

The derivation is simple.  $\square$

The converse of Proposition 4 holds only in a very restricted case:

**Proposition 5.** (Upward preservation of safety.) *Suppose that  $\Pi \subseteq \Psi_1 \subseteq \Psi_2$ . If  $\Pi$  is a safety property relative to  $\Psi_1$  and  $\Psi_1$  is a safety property relative to  $\Psi_2$ , then  $\Pi$  is a safety property relative to  $\Psi_2$ .*

**Proof.** Again, use Proposition 2 and the monotonicity of the closure operator.  $\square$

In general, properties become “safer” when they are viewed relative to stronger (i.e., more restrictive) properties: a property that is not an unconditional safety property may be safe relative to another property.

Indeed, there are natural properties relative to which all properties are safety properties. Let  $z \in \Sigma$  be a symbol that signals the termination of

a reactive system. Let  $\Psi_{fin} \subseteq \Sigma^\omega$  contain all infinite strings that are of the form that a finite prefix over the alphabet  $\Sigma - \{z\}$  is followed by an infinite suffix over the alphabet  $\{z\}$ ; that is, the property  $\Psi_{fin}$  of a reactive system asserts that “the system terminates.” It is not difficult to see that every property  $\Pi$  is safe relative to  $\Psi_{fin}$  (which itself is neither a safety property nor a liveness property). For suppose that every finite prefix of a string  $\sigma \in \Psi_{fin}$  can be extended to a string in  $\Pi \cap \Psi_{fin}$ . Then we can choose a sufficiently long prefix of  $\sigma$  that contains a  $z$ ; any extension of this prefix must be  $\sigma$  itself, which implies that  $\sigma \in \Pi \cap \Psi_{fin}$ .

In other words, under the assumption that all systems under consideration terminate, every property of a reactive system is a safety property. In the final section, we will present a less stringent assumption about reactive systems that, nonetheless, shifts interesting properties “towards safety”.

### 3. Operability and verification of relative specifications

We say that a pair  $(\Pi_S, \Pi_L)$  specifies the property  $\Pi \subseteq \Psi$  congruently relative to  $\Psi \subseteq \Sigma^\omega$  iff

- (1)  $\Pi = \Pi_S \cap \Pi_L \cap \Psi$ ,
- (2)  $\Pi_S$  is safe relative to  $\Psi$  and  $\Pi_L$  is live relative to  $\Psi$ , and
- (3) every finite string that is both a prefix of a string in  $\Pi_S$  and a prefix of a string in  $\Psi$  can be extended to a string in  $\Pi$ .

Thus a specification is unconditionally congruent iff it is congruent relative to  $\Sigma^\omega$ . The following theorem generalizes the main result about the unconditional safety-liveness classification (Theorem 1).

**Theorem 6.** (Existence of relatively congruent specifications.) *For all  $\Psi \subseteq \Sigma^\omega$ , every property  $\Pi \subseteq \Psi$  has a specification that is congruent relative to  $\Psi$ .*

**Proof.** Let  $\Pi_S = \overline{\Pi}$  and  $\Pi_L = \neg((\Pi_S \cap \Psi) - \Pi)$ ; then  $\Pi_S$  is unconditionally safe. Alternatively, let  $\Pi_S = \overline{\Pi} \cap \Psi$  and  $\Pi_L = \neg(\Pi_S - \Pi)$ ; then  $\Pi_S \subseteq \Psi$ .

We show that  $(\Pi_S, \Pi_L)$  specifies  $\Pi$  congruously relative to  $\Psi$  in either case.

(1) It is not hard to check that  $\Pi = \Pi_S \cap \Pi_L \cap \Psi$ .

(2) The unconditional safety property  $\Pi_S = \bar{\Pi}$  is safe relative to  $\Psi$ , and so is  $\Pi_S = \bar{\Pi} \cap \Psi$ . To see that  $\Pi_L$  is live relative to  $\Psi$ , by Proposition 3 it suffices to show that

$$\Psi \subseteq \overline{\neg((\bar{\Pi} \cap \Psi) - \Pi)} \cap \Psi.$$

Since  $\Pi \subseteq \Psi$ , this condition is equivalent to

$$\Psi \subseteq \bar{\Pi} \cup (\Psi - \bar{\Pi}).$$

We can derive both

$$\bar{\Pi} \cap \Psi \subseteq \bar{\Pi} \cup (\Psi - \bar{\Pi})$$

and

$$\neg \bar{\Pi} \cap \Psi \subseteq \bar{\Pi} \cup (\Psi - \bar{\Pi}),$$

using the monotonicity of the closure operator.

(3) Since  $\Pi_S \subseteq \bar{\Pi}$ , every finite prefix of a string in  $\Pi_S$  can be extended to a string in  $\Pi$ .  $\square$

Our definition of relative congruity ensures again operationality: a machine that incrementally generates prefixes in  $\Pi_S$  that are also prefixes of  $\Psi$  will never reach an irremedial situation from which the liveness conditions of  $\Pi_L \cap \Psi$  cannot be satisfied. Next we shall see that the relative congruity of system descriptions is desirable also from a verification point of view.

The notion of relative safety has ramifications for both the specification and the verification of reactive systems. Suppose that a property  $\Pi$  is safe relative to an assumption  $\Psi$ . We can take advantage of this fact in two ways:

1. The property  $\Pi$  can be *specified* by an unconditional safety property, namely,  $\bar{\Pi} \cap \Psi$ . This is because  $(\bar{\Pi} \cap \Psi) \cap \Psi = \Pi \cap \Psi$  by Proposition 2.
2. The property  $\Pi$  can be *verified* by safety reasoning. Suppose that the possible behaviors of a reactive system  $\hat{\Pi}$  are given by the congruous pair  $(\hat{\Pi}_S, \hat{\Pi}_L)$ . In order to verify that the system  $\hat{\Pi}$  has the property  $\Pi$ , it suffices to show that the safety component  $\hat{\Pi}_S$  of the system  $\hat{\Pi}$  satisfies the safety property  $\bar{\Pi} \cap \Psi$ .

This verification strategy is justified by the following theorem; the strategy is complete, provided that (1) we may also use the safety component  $\Psi_S$  of the assumption  $\Psi$  in the verification process, and (2) the system specification  $(\hat{\Pi}_S, \hat{\Pi}_L)$  is congruous relative to the assumption  $\Psi$ .

**Theorem 7.** (Verification of relative safety properties.) *Let  $(\Psi_S, \Psi_L)$  be a congruous specification of  $\Psi \subseteq \Sigma^\omega$ , let  $(\hat{\Pi}_S, \hat{\Pi}_L)$  be a specification of  $\hat{\Pi} \subseteq \Psi$  that is congruous relative to  $\Psi$ , and let  $\Pi \subseteq \Sigma^\omega$  be safe relative to  $\Psi$ . Then  $\hat{\Pi} \subseteq \Pi$  iff  $\hat{\Pi}_S \cap \Psi_S \subseteq \bar{\Pi} \cap \Psi$ .*

**Proof.** First, assume that  $\hat{\Pi}_S \cap \Psi_S \subseteq \bar{\Pi} \cap \Psi$ . Then

$$\hat{\Pi} \subseteq (\bar{\Pi} \cap \Psi) \cap \Psi$$

and, since  $\Pi$  is safe relative to  $\Psi$ , we have

$$(\bar{\Pi} \cap \Psi) \cap \Psi \subseteq \Pi$$

by Proposition 2. By transitivity,  $\hat{\Pi} \subseteq \Pi$  follows.

Second, assume that  $\hat{\Pi} \subseteq \Pi$ . Since the pair  $(\Psi_S, \Psi_L)$  is congruous,  $\hat{\Pi}_S \cap \Psi_S \subseteq \hat{\Pi}_S \cap \bar{\Psi}$ . As the specification  $(\hat{\Pi}_S, \hat{\Pi}_L)$  is congruous relative to  $\Psi$ , we have

$$\hat{\Pi}_S \cap \bar{\Psi} \subseteq \bar{\Pi}.$$

By our assumption,  $\hat{\Pi} \subseteq \Pi \cap \Psi$  and, by the monotonicity of the closure operator,

$$\bar{\Pi} \subseteq \bar{\Pi} \cap \Psi.$$

By transitivity,  $\hat{\Pi}_S \cap \Psi_S \subseteq \bar{\Pi} \cap \Psi$  as desired.  $\square$

Now let us illustrate the application of this result with the termination assumption  $\Psi_{fin}$ . Consider the liveness property  $\Pi_{\diamond p}$  that contains all infinite strings with at least one occurrence of the symbol  $p$ . Since every property is safe relative to  $\Psi_{fin}$ , so is in particular  $\Pi_{\diamond p}$ . Thus Theorem 7 tells us that, over terminating systems,  $\Pi_{\diamond p}$  can be specified and verified as the safety property  $\bar{\Pi}_{\diamond p} \cap \Psi_{fin}$ . This property consists of all infinite strings such that (1) each occurrence of  $z$  is followed by a  $z$  and (2) there is an occurrence of  $p$  before the first occurrence of  $z$  (including all strings that contain neither a  $p$  nor a  $z$ ). Note

that, indeed, if all runs of a system satisfy the safety property  $\Pi_{\Diamond p} \cap \Psi_{fin}$ , then all terminating runs of the system satisfy the desired property  $\Pi_{\Diamond p}$ .

#### 4. Real-time safety and liveness

The behavior of a discrete real-time system can be described by an infinite sequence of pairs

$$\rho: (\sigma_0, \tau_0) \rightarrow (\sigma_1, \tau_1) \rightarrow (\sigma_2, \tau_2) \rightarrow \dots$$

of states  $\sigma_i \in \Sigma$ , for  $i \geq 0$ , and corresponding times  $\tau_i \in \mathcal{T}$ . While we do not commit to any particular time domain  $\mathcal{T}$ , we assume that there is a real-valued distance function  $d$  on  $\mathcal{T}^2$  with  $d(x, x) = 0$  for all  $x \in \mathcal{T}$ . The sequence  $\rho = (\sigma, \tau)$  is called a *timed state sequence*.

A *real-time property*  $\Pi$  is a subset of  $\Psi_{all}$ , the set of all timed state sequences. It is straightforward to extend the definitions of unconditional and relative safety and liveness to real-time properties. All results of the previous sections carry over. In particular, any trivial one-element time domain yields a model that is isomorphic to the original untimed setup.

Different models of time and computation put vastly different requirements on the time component  $\tau$  of legal behaviors  $\rho = (\sigma, \tau)$  of a real-time system. For instance:

- *Interval* models of time associate with every state its duration over time, while *clock* models stamp observations of the system state with time instants. Intervals of the real line are a suitable time domain for the former model, points for the latter.
- *Analog-clock* models of time record the exact time of every state, while *digital-clock* models measure the time of a state only with finite precision. The reals are a suitable time domain for the former model, the integers for the latter.
- In *synchronous* models of computation, all concurrent activity happens in lock-step, while *asynchronous (interleaving)* models sequentialize simultaneous actions nondeterministically. Strictly monotonic time is appropriate for the

former model, while instantaneous actions are required by the latter.

(See [7] for a survey of various models of time that have been proposed for the verification of real-time systems.)

Given a particular choice of model, we consider, by definition, only a subset  $\Psi \subseteq \Psi_{all}$  of timed state sequences as possible behaviors of a real-time system; that is, the specification of property  $\Pi$  really defines  $\Pi \cap \Psi$ . Thus we can specify  $\Pi$  by describing any property  $\Pi'$  with  $\Pi' \cap \Psi = \Pi \cap \Psi$ , possibly even using a safety property  $\Pi'$  to specify a liveness property  $\Pi$ . Precisely this phenomenon is captured formally by the concept of safety and liveness relative to the *timing assumption*  $\Psi$ .

There are two particularly important model-independent timing assumptions:

1. All “reasonable” models of time require that time must not decrease. A timed state sequence  $(\sigma, \tau)$  is called *monotonic* iff time increases (weakly) monotonically:

$$d(\tau_i, \tau_j) \leq d(\tau_i, \tau_k) \quad \text{for all } 0 \leq i \leq j \leq k.$$

The set  $\Psi_{mon} \subseteq \Psi_{all}$  of all monotonic timed state sequences is clearly a safety property.

2. The behavior of a continuous system that may change its state infinitely often between any two points in time cannot be modeled adequately by an  $\omega$ -sequence of states. Thus, given our choice of a timed state sequence semantics, we may “reasonably” demand that time diverges. A timed state sequence  $(\sigma, \tau)$  is called *divergent* iff time eventually proceeds to any point:

$$\text{for all } i \geq 0 \text{ and } x \in \mathcal{T}, \text{ there is some } j \geq i \text{ such that } d(\tau_i, \tau_j) \geq d(\tau_i, x).$$

It can be checked that the set  $\Psi_{div} \subseteq \Psi_{all}$  of all divergent timed state sequences is a liveness property.

It follows that typical timing assumptions are subsets of  $\Psi_{time} = \Psi_{mon} \cap \Psi_{div}$ .

Therefore we are especially interested in safety, liveness and operationality *relative to*

*monotonic divergence* (i.e., relative to  $\Psi_{time}$ ). The class of properties that are safe relative to monotonic divergence includes many important real-time properties that are unconditional liveness properties; that is, all the liveness they stipulate is subsumed by the divergence of time.

Bounded response is the standard example of a real-time property that is unconditionally live and becomes safe under strong enough timing assumptions [10,14,15,18]. Let  $p, q \in \Sigma$  and let  $\delta$  be a nonnegative real. The *bounded-response* property  $\Pi_{p \rightarrow q}^\delta$  contains a timed state sequence  $(\sigma, \tau)$  iff for all  $i \geq 0$ , whenever  $\sigma_i = p$ , then  $\sigma_j = q$  and  $d(\tau_i, \tau_j) \leq \delta$  for some  $j \geq i$ ; that is, every  $p$ -state is followed by a  $q$ -state within time  $\delta$ . Since any finite prefix of a timed state sequence containing  $(p, x)$  can be extended with the pair  $(q, x)$ , the property  $\Pi_{p \rightarrow q}^\delta$  is an unconditional liveness property.

Now let us consider  $\Pi_{p \rightarrow q}^\delta$  relative to monotonicity, and then relative to monotonic divergence. Provided that  $p$  and  $q$  are different states,  $\Pi_{p \rightarrow q}^\delta$  is not safe relative to  $\Psi_{mon}$ , because it contains all monotonic timed state sequences of the form

$$(p, x) \rightarrow \cdots \rightarrow (p, x) \rightarrow (q, x) \rightarrow \cdots,$$

without containing the monotonic sequence

$$(p, x) \rightarrow (p, x) \rightarrow (p, x) \rightarrow \cdots.$$

Provided that there are two times  $x, y \in \mathcal{T}$  with  $d(x, y) > \delta$ , the property  $\Pi_{p \rightarrow q}^\delta$  is not live relative to  $\Psi_{mon}$  either, because the finite prefix

$$(p, x) \rightarrow (p, y)$$

cannot be extended to a monotonic sequence in  $\Pi_{p \rightarrow q}^\delta$ . Finally, suppose that for all  $x \in \mathcal{T}$  there is some  $y \in \mathcal{T}$  such that  $d(x, y) > \delta$ . Then it is not hard to check that the bounded-response property  $\Pi_{p \rightarrow q}^\delta$  is a safety property relative to monotonic divergence; the “bad thing” that is not supposed to happen is that, after a  $p$ -state,  $\delta$  time units pass without a  $q$ -state occurring.

Specifications that are congruous relative to monotonic divergence are called *nonZeno* [2], because they cannot define Zeno machines that

force time to converge. Real-time transition systems [10] and extended state machines [16] are examples of specifications that are nonZeno, and thus operational descriptions of real-time systems. So are the timed automata of [15], which specify only properties that are safe relative to monotonic divergence. On the other hand, real-time temporal logics such as [6,12,16] and the timed automata of [5] permit, relative to monotonic divergence, incongruous specifications of real-time systems. A machine trying to execute such a specification without look-ahead may find itself in a situation from which time cannot diverge without violating the specification.

For nonZeno specifications we can apply Theorem 7. If a system is given congruously relative to monotonic divergence, then the bounded-response property  $\Pi_{p \rightarrow q}^\delta$  can be verified as the safety property

$$\overline{\Pi_{p \rightarrow q}^\delta \cap \Psi_{time}}$$

[10]. This property states that (1) time does not decrease and (2) whenever  $\sigma_i = p$ , then either  $\sigma_j = q$  and  $d(\tau_i, \tau_j) \leq \delta$  for some  $j \geq i$  or  $d(\tau_i, \tau_j) \leq \delta$  for all  $j \geq i$ .

## Acknowledgment

The author thanks Martín Abadi, Rajeev Alur, David Dill, Leslie Lamport, Zohar Manna, Amir Pnueli, Fred Schneider, and two anonymous referees for many valuable suggestions and improvements.

## References

- [1] M. Abadi and L. Lamport, The existence of refinement mappings, In: *Proc. Third Ann. Symp. on Logic in Computer Science* (IEEE Computer Society Press, Silver Spring, MD, 1988) 165–175.
- [2] M. Abadi and L. Lamport, An old-fashioned recipe for real time, in: *Proc. REX Workshop, Real-time: Theory in Practice*, Lecture Notes in Computer Science **600** (Springer, Berlin, 1992).
- [3] B. Alpern, A.J. Demers and F.B. Schneider, Safety without stuttering, *Inform. Process. Lett.* **23** (1986) 177–180.

- [4] B. Alpern and F.B. Schneider, Defining liveness, *Inform. Process. Lett.* **21** (1985) 181–185.
- [5] R. Alur and D.L. Dill, Automata for modeling real-time systems, in: *Proc. 17th Internat. Coll. on Automata, Languages, and Programming*, Lecture Notes in Computer Science **443** (Springer, Berlin, 1990) 322–335.
- [6] R. Alur and T.A. Henzinger, A really temporal logic, in: *Proc. 30th Ann. IEEE Symp. on Foundations of Computer Science* (1989) 164–169.
- [7] R. Alur and T.A. Henzinger, Logics and models of real time: a survey, In: *Proceedings of the REX Workshop, Real-time: Theory in Practice*, Lecture Notes in Computer Science **600** (Springer, Berlin, 1992).
- [8] K.R. Apt, N. Francez and S. Katz, Appraising fairness in languages for distributed programming, *Distributed Comput.* **2** (1988) 226–241.
- [9] F. Dederichs and R. Weber, Safety and liveness from a methodological point of view, *Inform. Process. Lett.* **36** (1) (1990) 25–30.
- [10] T.A. Henzinger, Z. Manna and A. Pnueli, Temporal proof methodologies for real-time systems, in: *Proc. 18th Ann. ACM Symp. on Principles of Programming Languages* (ACM Press, New York, 1991) 353–366.
- [11] J.L. Kelley, *General Topology* (Springer, Berlin, 1955).
- [12] R. Koymans, Specifying real-time properties with metric temporal logic, *Real-time Systems* **2** (1990) 255–299.
- [13] L. Lamport, Proving the correctness of multiprocess programs, *IEEE Trans. Software Engineering* **3** (1977) 125–143.
- [14] L. Lamport, The temporal logic of actions, Tech. Rept., DEC Systems Research Center, February 1991.
- [15] N.A. Lynch and H. Attiya, Using mappings to prove timing properties, in: *Proc. Ninth Ann. ACM Symp. on Principles of Distributed Computing* (ACM Press, New York, 1990) 265–280.
- [16] J.S. Ostroff, *Temporal Logic of Real-time Systems* (Research Studies Press, 1990).
- [17] A. Pnueli, Applications of temporal logic to the specification and verification of reactive systems: a survey of current trends, in: *Current Trends in Concurrency*, Lecture Notes in Computer Science **224** (Springer, Berlin, 1986) 510–584.
- [18] F.B. Schneider, Private communication, February 1991.
- [19] W. Thomas, Automata on infinite objects, in: Jan van Leeuwen, ed., *Handbook of Theoretical Computer Science, Vol. B* (Elsevier, Amsterdam and MIT Press, Cambridge, MA, 1990) 133–191.