

The Theory of Hybrid Automata*

Thomas A. Henzinger[†]

Electrical Engineering and Computer Sciences
University of California at Berkeley

Abstract. We summarize several recent results about hybrid automata. Our goal is to demonstrate that concepts from the theory of discrete concurrent systems can give insights into partly continuous systems, and that methods for the verification of finite-state systems can be used to analyze certain systems with uncountable state spaces.

1 Hybrid Automata

A hybrid automaton is a formal model for a dynamical system with discrete and continuous components.

1.1 Syntax

A paradigmatic example of a mixed discrete-continuous system is a digital controller of an analog plant. The discrete state of the controller is modeled by the vertices of a graph (control modes), and the discrete dynamics of the controller is modeled by the edges of the graph (control switches). The continuous state of the plant is modeled by points in \mathbb{R}^n , and the continuous dynamics of the plant is modeled by flow conditions such as differential equations. The behavior of the plant depends on the state of the controller: each control mode determines a flow condition, and each control switch may cause a discrete change in the state of the plant, as determined by a jump condition. Dually, the behavior of the controller depends on the state of the plant: each control mode continuously observes an invariant condition of the plant state, and by violating the invariant condition, a continuous change in the plant state will cause a control switch.

*This research was supported in part by the Office of Naval Research Young Investigator award N00014-95-1-0520, by the National Science Foundation CAREER award CCR-9501708, by the National Science Foundation grant CCR-9504469, by the Air Force Office of Scientific Research contract F49620-93-1-0056, and by the Advanced Research Projects Agency grant NAG2-892.

[†]Email: tah@eecs.berkeley.edu.

Definition 1.1 [Hybrid automata] [5, 36, 3] A *hybrid automaton* H consists of the following components.

Variables. A finite set $X = \{x_1, \dots, x_n\}$ of real-numbered variables. The number n is called the *dimension* of H . We write \dot{X} for the set $\{\dot{x}_1, \dots, \dot{x}_n\}$ of dotted variables (which represent first derivatives during continuous change), and we write X' for the set $\{x'_1, \dots, x'_n\}$ of primed variables (which represent values at the conclusion of discrete change).

Control graph. A finite directed multigraph (V, E) . The vertices in V are called *control modes*. The edges in E are called *control switches*.

Initial, invariant, and flow conditions. Three vertex labeling functions *init*, *inv*, and *flow* that assign to each control mode $v \in V$ three predicates. Each initial condition *init*(v) is a predicate whose free variables are from X . Each invariant condition *inv*(v) is a predicate whose free variables are from X . Each flow condition *flow*(v) is a predicate whose free variables are from $X \cup \dot{X}$.

Jump conditions. An edge labeling function *jump* that assigns to each control switch $e \in E$ a predicate. Each jump condition *jump*(e) is a predicate whose free variables are from $X \cup X'$.

Events. A finite set Σ of events, and an edge labeling function *event* : $E \rightarrow \Sigma$ that assigns to each control switch an event. \square

Example 1.1 [Temperature control] The hybrid automaton of Figure 1 models a thermostat. The variable x represents the temperature. In control mode *Off*, the heater is off, and the temperature falls according to the flow condition $\dot{x} = -0.1x$. In control mode *On*, the heater is on, and the temperature rises according to the flow condition $\dot{x} = 5 - 0.1x$. Initially, the heater is off and the temperature is 20 degrees. According to the jump condition $x < 19$, the heater




Figure 1: Thermostat automaton

may go on as soon as the temperature falls below 19 degrees. According to the invariant condition $x \geq 18$, at the latest the heater will go on when the temperature falls to 18 degrees. \square

1.2 Safe Semantics

The execution of a hybrid automaton results in continuous change (flows) and discrete change (jumps). The mixed discrete-continuous dynamics can be abstracted by a fully discrete transition system.

Definition 1.2 [Labeled transition systems] A *labeled transition system* S consists of the following components.

State space. A (possibly infinite) set Q of *states*, and a subset $Q^0 \subseteq Q$ of *initial states*.

Transition relations. A (possibly infinite) set A of *labels*, and for each label $a \in A$, a binary relation \xrightarrow{a} on the state space Q . Each triple $q \xrightarrow{a} q'$ is called a *transition*.

A subset $R \subseteq Q$ of the state space is called a *region*. Given a region R and a label $a \in A$, we write $\text{post}_a(R) = \{q' \mid \exists q \in R. q \xrightarrow{a} q'\}$ for the region of a -successors of R , and we write $\text{pre}_a(R) = \{q \mid \exists q' \in R. q \xrightarrow{a} q'\}$ for the region of a -predecessors of R . \square

For a given hybrid automaton, we define two labeled transition systems. Both transition systems represent discrete jumps by transitions. The timed transition system abstracts continuous flows by transitions, retaining only information about the source, the target, and the duration of each flow. The time-abstract transition system abstracts also the duration of flows.

Definition 1.3 [Transition semantics of hybrid automata] The *timed transition system* S_H^t of the hybrid automaton H is the labeled transition system with the components Q , Q^0 , A , and \xrightarrow{a} for each $a \in A$, defined as follows.

- Define $Q, Q^0 \subseteq V \times \mathbb{R}^n$ such that $(v, \mathbf{x}) \in Q$ iff the closed predicate $\text{inv}(v)[X := \mathbf{x}]$ is true, and $(v, \mathbf{x}) \in Q^0$ iff both $\text{init}(v)[X := \mathbf{x}]$ and

$\text{inv}(v)[X := \mathbf{x}]$ are true. The set Q is called the *state space* of H , and the subsets of Q are called *H-regions*.

- $A = \Sigma \cup \mathbb{R}_{\geq 0}$.
- For each event $\sigma \in \Sigma$, define $(v, \mathbf{x}) \xrightarrow{\sigma} (v', \mathbf{x}')$ iff there is a control switch $e \in E$ such that (1) the source of e is v and the target of e is v' , (2) the closed predicate $\text{jump}(e)[X, X' := \mathbf{x}, \mathbf{x}']$ is true, and (3) $\text{event}(e) = \sigma$.
- For each nonnegative real $\delta \in \mathbb{R}_{\geq 0}$, define $(v, \mathbf{x}) \xrightarrow{\delta} (v', \mathbf{x}')$ iff $v = v'$ and there is a differentiable function $f : [0, \delta] \rightarrow \mathbb{R}^n$, with the first derivative $\dot{f} : (0, \delta) \rightarrow \mathbb{R}^n$, such that (1) $f(0) = \mathbf{x}$ and $f(\delta) = \mathbf{x}'$, and (2) for all reals $\varepsilon \in (0, \delta)$, both $\text{inv}(v)[X := f(\varepsilon)]$ and $\text{flow}(v)[X, X' := f(\varepsilon), \dot{f}(\varepsilon)]$ are true. The function f is called a *witness* for the transition $(v, \mathbf{x}) \xrightarrow{\delta} (v', \mathbf{x}')$.

The *time-abstract transition system* S_H^a of H is the labeled transition system with the components Q , Q^0 , B , and \xrightarrow{b} for each $b \in B$, defined as follows.

- Q and Q^0 are defined as above.
- $B = \Sigma \cup \{\tau\}$, for some event $\tau \notin \Sigma$.
- For each event $\sigma \in \Sigma$, define $\xrightarrow{\sigma}$ as above.
- Define $(v, \mathbf{x}) \xrightarrow{\tau} (v', \mathbf{x}')$ iff there is a nonnegative real $\delta \in \mathbb{R}_{\geq 0}$ such that $(v, \mathbf{x}) \xrightarrow{\delta} (v', \mathbf{x}')$.

The time-abstract transition system S_H^a is called the *time abstraction* of the timed transition system S_H^t . \square

Remark. [Definition 1.3] The state space Q and the timed label set A are infinite. The time-abstract label set B is finite. For all states q of a hybrid automaton, $q \xrightarrow{0} q$. Sequences of event transitions and time transitions with duration (label) 0 are permitted, which generalizes the interleaving semantics for discrete concurrent systems [7]. \square

Remark. [Time vs. phase view] The time-abstract transition system S_H^a , which projects away the time dimension, can be viewed as the phase portrait of the timed transition system S_H^t [21]. \square

Remark. [Time-silent transition semantics] None of the results presented in this paper change if the τ -transitions of time-abstract transition systems are considered silent [25]. \square

1.3 Live Semantics

If we consider the infinite behavior of a hybrid automaton, then we are interested only in infinite sequences of transitions which do not converge in time. The divergence of time is a liveness assumption, and it is the only liveness assumption we need to consider [20, 27]. A hybrid automaton is nonzeno if it cannot prevent time from diverging. Clearly, only nonzeno designs of real-time systems can be realized.

Definition 1.4 [Live transition systems] Consider a labeled transition system S and a state q_0 of S . A q_0 -rooted trajectory of S is a finite or infinite sequence of pairs $\langle a_i, q_i \rangle_{i \geq 1}$ of labels $a_i \in A$ and states $q_i \in Q$ such that $q_{i-1} \xrightarrow{a_i} q_i$ for all $i \geq 1$. If q_0 is an initial state of S , then $\langle a_i, q_i \rangle_{i \geq 1}$ is an *initialized trajectory* of S . A *live transition system* (S, L) is a pair consisting of a labeled transition system S and a set L of infinite initialized trajectories of S . The set L of infinite initialized trajectories is *machine-closed* for S if every finite initialized trajectory of S is a prefix of some trajectory in L .¹ If (S, L) is a live transition system, and $\langle a_i, q_i \rangle_{i \geq 1}$ is either a finite initialized trajectory of S or a trajectory in L , then the corresponding sequence $\langle a_i \rangle_{i \geq 1}$ of labels is called a (finite or infinite) *trace* of (S, L) . \square

Definition 1.5 [Trace semantics of hybrid automata] We associate with each transition of the timed transition system S_H^t a *duration* in $\mathbb{R}_{\geq 0}$. For events $\sigma \in \Sigma$, the duration of $q \xrightarrow{\sigma} q'$ is 0. For reals $\delta \in \mathbb{R}_{\geq 0}$, the duration of $q \xrightarrow{\delta} q'$ is δ . An infinite trajectory $\langle a_i, q_i \rangle_{i \geq 1}$ of the timed transition system S_H^t *diverges* if the infinite sum $\sum_{i \geq 1} \delta_i$ diverges, where each δ_i is the duration of the corresponding transition $q_{i-1} \xrightarrow{a_i} q_i$. An infinite trajectory $\langle b_i, q_i \rangle_{i \geq 1}$ of the time-abstract transition system S_H^a *diverges* if there is a divergent trajectory $\langle a_i, q_i \rangle_{i \geq 1}$ of S_H^t such that for all $i \geq 1$, either $a_i = b_i$ or $a_i, b_i \notin \Sigma$. Let L_H^t be the set of divergent initialized trajectories of the timed transition system S_H^t , and let L_H^a be set of divergent initialized trajectories of the time-abstract transition system S_H^a . The hybrid automaton H is *nonzeno* if L_H^t is machine-closed for S_H^t (or equivalently, L_H^a is machine-closed for S_H^a). Each trace of the live transition system (S_H^t, L_H^t) is called a *timed trace* of H , and each trace of the live transition system (S_H^a, L_H^a) is called a *time-abstract trace* of H . \square

1.4 Composition

For two hybrid automata H_1 and H_2 , we define the timed semantics and the time-abstract semantics of

the parallel composition $H_1 \parallel H_2$. The two hybrid automata H_1 and H_2 interact via joint events: if event a is both an event of H_1 and an event of H_2 , then H_1 and H_2 must synchronize on a -transitions; if a is an event of H_1 but not an event of H_2 , then each a -transition of H_1 synchronizes with a 0-duration time transition of H_2 , and vice versa. For each real $\delta > 0$, a time transition of H_1 with duration δ must synchronize with a time transition of H_2 with the same duration.

Definition 1.6 [Product of transition systems] A *consistency check* for two labeled transition systems S_1 and S_2 is an associative partial function \otimes on pairs consisting of a transition from S_1 and a transition from S_2 . The *product* $S_1 \times S_2$ with respect to the consistency check \otimes is the labeled transition system with the state space $Q_1 \times Q_2$, the set $Q_1^0 \times Q_2^0$ of initial states, the label set $\text{range}(\otimes)$, and the following transition relations: for each label $a \in \text{range}(\otimes)$, define $(q_1, q_2) \xrightarrow{a} (q'_1, q'_2)$ iff there is a label $a_1 \in A_1$ and a label $a_2 \in A_2$ such that a is the (defined) result of applying \otimes to the two transitions $q_1 \xrightarrow{a_1} q'_1$ and $q_2 \xrightarrow{a_2} q'_2$. \square

Definition 1.7 [Composition of hybrid automata] Consider two hybrid automata H_1 and H_2 . A transition $q_1 \xrightarrow{a_1} q'_1$ of $S_{H_1}^t$ and a transition $q_2 \xrightarrow{a_2} q'_2$ of $S_{H_2}^t$ are *consistent* if one of the following three conditions is true.

1. $a_1 = a_2$. In this case, the consistency check $\otimes =$ applied to the transitions $q_1 \xrightarrow{a_1} q'_1$ and $q_2 \xrightarrow{a_2} q'_2$ yields a_1 .
2. $a_1 \in \Sigma_1 \setminus \Sigma_2$ and $a_2 = 0$. In this case, the consistency check $\otimes =$ yields a_1 .
3. $a_1 = 0$ and $a_2 \in \Sigma_2 \setminus \Sigma_1$. In this case, the consistency check $\otimes =$ yields a_2 .

The *timed transition system* $S_{H_1 \parallel H_2}^t$ is defined to be the product $S_{H_1}^t \times S_{H_2}^t$ with respect to the consistency check $\otimes =$. The *time-abstract transition system* $S_{H_1 \parallel H_2}^a$ is defined to be the time abstraction of $S_{H_1 \parallel H_2}^t$. \square

Example 1.2 [Railroad gate control] The hybrid automaton of Figure 2 models a train on a circular track with a gate. The variable x represents the distance of the train from the gate. Initially, the speed of the train is between 40 and 50 meters per second. At the distance of 1000 meters from the gate, the train issues an *approach* event and may slow down to 30 meters per second. At the distance of 100 meters past the

¹ Assuming that every initial state of S has a successor state.




Figure 2: Train automaton




Figure 3: Controller automaton

gate, the train issues an *exit* event. The circular track is between 2000 and 5000 meters long. We write jump conditions as guarded commands, which allows us to suppress conjuncts of the form $x' = x$. In particular, the jump condition of the control switch from *Near* to *Past* is $x = 0 \wedge x' = x$, and the jump condition from *Past* to *Far* is $x = -100 \wedge 1900 \leq x' \leq 4900$. The hybrid automaton of Figure 3 models the gate controller. The variable u is a symbolic constant that represents the reaction delay of the controller. The variable z is a clock for measuring elapsed time. When an *approach* event is received, the controller issues a *lower* event within u seconds, and when an *exit* event is received, the controller issues a *raise* event within u seconds. The hybrid automaton of Figure 4 models the gate. The variable y represents the position of the gate in degrees. Initially, the gate is open ($y = 90$). When a *lower* event is received, the gate starts closing at the rate of 9 degrees per second, and when a *raise* event is received, the gate starts opening at the same rate. Which values of the symbolic constant u ensure that the gate is fully closed ($y = 0$) whenever the train is within 10 meters of the gate ($-10 \leq x \leq 10$)? \square




Figure 4: Gate automaton

Remark. [Shared variables] The consistency check $\otimes =$ depends only on the transition labels, and not on the source and target states of transitions. Alternative consistency checks can be used to model read-shared and even write-shared variables of hybrid automata [8]. \square

Remark. [Time-abstract hybrid automata] The time-abstract transition system $S_{H_1 \parallel H_2}^a$ is generally different from the product $S_{H_1}^a \times S_{H_2}^a$ of the time-abstract component systems. This is not the case for *time-abstract hybrid automata* [21]. Time-abstract design is desirable, because many useful properties of time-abstract component systems are inherited by the compound system. \square

Remark. [Receptiveness] The composition of two nonzeno hybrid automata is not necessarily nonzeno. It is an interesting modeling problem for real-time systems to guarantee the liveness of compound designs [1, 33, 32]. \square

2 On the Trace Languages of Hybrid Automata

We identify which requirements on the traces of a hybrid automaton can be checked algorithmically, and which cannot.

2.1 Verification Tasks

We study four paradigmatic questions about the traces of a hybrid automaton. The reachability problem is a fundamental subtask for the verification of safety requirements, and the emptiness problem is a fundamental subtask for the verification of liveness requirements. The timed trace inclusion problem compares the traces of a hybrid automaton against

a timed specification, and the time-abstract trace inclusion problem compares the traces of a hybrid automaton against a time-abstract specification.

Definition 2.1 [Reachability, emptiness, and trace inclusion] The *reachability problem* for a class \mathcal{H} of hybrid automata asks, given a hybrid automaton H from \mathcal{H} and a control mode v of H , if there is an initialized trajectory of S_H^t (or equivalently, S_H^a) that visits a state of the form (v, \mathbf{x}) . The *emptiness problem* for \mathcal{H} asks, given a hybrid automaton H from \mathcal{H} , if there is a divergent initialized trajectory of S_H^t (or equivalently, S_H^a). The *(finitary) timed trace inclusion problem* for \mathcal{H} asks, given two hybrid automata H_1 and H_2 from \mathcal{H} , if every (finite) timed trace of H_1 is also a timed trace of H_2 . The *(finitary) time-abstract trace inclusion problem* for \mathcal{H} asks, given two hybrid automata H_1 and H_2 from \mathcal{H} , if every (finite) time-abstract trace of H_1 is also a time-abstract trace of H_2 . \square

Remark. [Definition 2.1] Some of these problems are harder than others. In particular, reachability can be reduced to finitary time-abstract trace inclusion, and emptiness can be reduced to time-abstract trace inclusion. Also, finitary trace inclusion can be reduced to trace inclusion. \square

2.2 Rectangular Automata

A hybrid automaton is rectangular if the flow conditions are independent of the control modes, and the variables are pairwise independent. Specifically, in each control mode of a rectangular automaton, the first derivative of each variable is given a range of possible values, and that range does not change with control switches. With each control switch of a rectangular automaton, the value of each variable is either left unchanged, or changed nondeterministically to a new value within a given range of possibilities. The behaviors of the variables are decoupled, because the ranges of possible values and derivative values for one variable cannot depend on the value or derivative value of another variable.

Definition 2.2 [Rectangular automata] [38, 26] A *rectangle* $\mathbf{I} = \prod_{1 \leq i \leq n} \mathbf{I}_i$ of *dimension* n is the product of n intervals $\mathbf{I}_i \subseteq \mathbb{R}$ of the real line, each with rational or infinite endpoints. The rectangle \mathbf{I} is *bounded* (a *singleton*) if each interval \mathbf{I}_i , for $1 \leq i \leq n$, is bounded (a singleton). A hybrid automaton H is a *rectangular automaton* if the following three restrictions are met. Let $X = \{x_1, \dots, x_n\}$ be the set of variables of H .

1. For each control mode v of H , the initial condition $\text{init}(v)$ has the form $X \in \mathbf{I}^{\text{init}(v)}$ for a bounded n -dimensional rectangle $\mathbf{I}^{\text{init}(v)}$, and the invariant condition $\text{inv}(v)$ has the form $X \in \mathbf{I}^{\text{inv}(v)}$ for an n -dimensional rectangle $\mathbf{I}^{\text{inv}(v)}$.
2. There is a bounded n -dimensional rectangle \mathbf{I}^{flow} such that for each control mode v of H , the flow condition $\text{flow}(v)$ has the form $\dot{X} \in \mathbf{I}^{\text{flow}}$.
3. For each control switch e of H , the jump condition $\text{jump}(v)$ has the form $X \in \mathbf{I}^{\text{pre}(e)} \wedge Y' = Y \wedge X' \in \mathbf{I}^{\text{post}(e)}$ for two n -dimensional rectangles $\mathbf{I}^{\text{pre}(e)}$ and $\mathbf{I}^{\text{post}(e)}$, and a set $Y \subseteq X$ of variables. The control switch e is said to *reinitialize* the variables in $X \setminus Y$. For all $1 \leq i \leq n$, if the variable x_i is reinitialized by e , then the interval $\mathbf{I}_i^{\text{post}(e)}$ must be bounded.

The rectangular automaton H is a *singular automaton* if the flow rectangle \mathbf{I}^{flow} is a singleton. The singular automaton H is a *timed automaton* if $\mathbf{I}^{\text{flow}} = [1, 1]^n$. \square

Remark. [Clocks and drifting clocks] A clock can be modeled by a variable x_i with the flow interval $\mathbf{I}_i^{\text{flow}} = [1, 1]$. All variables of a timed automaton are clocks [6]. A clock with drift ε , for $\varepsilon \in \mathbb{Q}_{\geq 0}$, can be modeled by a variable with the flow interval $[1 - \varepsilon, 1 + \varepsilon]$ [12, 30]. \square

Remark. [Composition] Timed, singular, and rectangular automata are closed under composition: for two timed (singular; rectangular) automata H_1 and H_2 , we can construct a timed (singular; rectangular) automaton H such that $S_H^t = S_{H_1 \parallel H_2}^t$ (and therefore, $S_H^a = S_{H_1 \parallel H_2}^a$). If the dimension of H_1 is n_1 and the dimension of H_2 is n_2 , then the dimension of H is $n_1 + n_2$. \square

We define two generalizations of rectangular automata. Multirectangular automata allow flow conditions that vary with control switches, and triangular automata allow the comparison of variables.

Definition 2.3 [Multirectangular and triangular automata] A hybrid automaton H is a *multirectangular automaton* if the restrictions of Definition 2.2 are met, except that different control modes v and v' of H may have different flow rectangles $\mathbf{I}^{\text{flow}(v)}$ and $\mathbf{I}^{\text{flow}(v')}$. The multirectangular automaton H is a *multisingular automaton* if all flow rectangles of H are singletons. The intersection of an n -dimensional rectangle with any number of half-spaces of \mathbb{R}^n that are defined by inequalities of the form $x_i \leq x_j$, for $1 \leq i, j \leq n$, is called a *triangle of dimension* n . A hybrid automaton

is a *triangular automaton* if the restrictions of Definition 2.2 are met, except that every rectangle may be a triangle. \square

Remark. [Stopwatches and symbolic constants] A stopwatch can be modeled by a multisingular variable x_i with the two flow intervals $I_i^{flow(v)} = [1, 1]$ (the stopwatch is turned on) and $I_i^{flow(v')} = [0, 0]$ (the stopwatch is turned off). Stopwatches are useful for measuring accumulated durations, such as the cumulative amount of time that is spent in control mode v [4, 31]. An unknown system constant can be modeled by a singular variable x_j with the flow interval $I_j^{flow} = [0, 0]$ such that (1) x_j is not reinitialized by any control switch, and (2) the behaviors of other variables may depend on the (unknown but constant) value of x_j [9]. \square

Remark. [Initialized multirectangular automata] Some multirectangular automata can be translated to rectangular automata by increasing the dimension. In particular, this is the case for *initialized multirectangular automata*, where for each variable x_i and each control switch e , if the flow interval $I_i^{flow(v)}$ of the source v of e is different from the flow interval $I_i^{flow(v')}$ of the target v' of e , then x_i is reinitialized by e [38]. \square

Example 2.1 [Railroad gate control] The train automaton of Figure 2 is an initialized 1D multirectangular automaton and can be translated to a 2D rectangular automaton with the same timed traces. The controller automaton of Figure 3 is a 2D triangular automaton with a clock z and a symbolic constant u . If the reaction delay u of the controller is known (say, 5 seconds), then the controller can be modeled by a 1D timed automaton. The gate automaton of Figure 4 is a 1D multisingular automaton (not initialized). If the direction of the gate cannot be reversed midway, then the gate can be modeled by a singular automaton. \square

Remark. [Abstract interpretation] Nonsingular flow intervals permit the conservative approximation of complex continuous behavior with arbitrary accuracy [22]: we may split the state space into regions and within each region, use lower and upper bounds on the first derivatives of all variables. \square

2.3 Verification Results

The following theorem ensures the verifiability of rectangular automata against time-abstract finite-state specifications.

Theorem 2.1 [Time-abstract traces] [26] *For every rectangular automaton H , the set of finite time-abstract traces of H is regular, and the set of infinite time-abstract traces of H is ω -regular.*

Proof. Given an rectangular automaton H of dimension n , we can construct a singular automaton H' of dimension $2n$ such that H and H' have the same timed traces. The construction replaces each variable x_i of H by a variable x_i^l of H' that tracks the smallest possible value of x_i , and a variable x_i^u of H' that tracks the largest possible value of H' . In particular, if x_i has the flow interval $[\ell, u]$, then x_i^l has the flow interval $[\ell, \ell]$ and x_i^u has the flow interval $[u, u]$. Alur and Dill have shown that for every timed automaton H' one can construct a Büchi automaton H'' whose traces are exactly the time-abstract traces of H' (see Theorem 3.2 below). Their construction can be generalized to singular automata. \square

Corollary 2.1 [Time-abstract trace inclusion] *The time-abstract trace inclusion problem for rectangular automata can be decided in EXPSpace.*

Remark. [Emptiness] The emptiness problem for rectangular automata is in PSPACE, and the additional exponential for time-abstract trace inclusion is caused by an intermediate complementation step. PSPACE emptiness checking is optimal, because already the reachability problem for timed automata (and other real-time formalisms) is PSPACE-hard [6]. \square

Rectangular automata characterize an exact boundary between the decidability and undecidability of verification problems. If the flow conditions are allowed to vary with control switches (multirectangular automata), or if the values of different variables may be related (triangular automata), then already the reachability problem cannot be decided.

Theorem 2.2 [Reachability] [5] *The reachability problems for multisingular automata and for triangular automata are undecidable.*

Proof. Reduction from the halting problem for 2-counter machines. \square

Remark. [Theorem 2.2] Theorem 2.2 can be sharpened to more specific statements [9, 26]. For example, the combination of clocks with a single stopwatch causes undecidability, and so does the combination of clocks with symbolic constants. \square

We have focused on time-abstract trace inclusion, because there is no hope for deciding timed trace inclusion.

Theorem 2.3 [Timed trace inclusion] [6] *The finitary timed trace inclusion problem for timed automata is undecidable.*

Remark. [Complementation] Theorem 2.3 does not contradict the decidability of the emptiness problem for timed automata (which follows from Theorem 2.1), because the (finitary) timed trace sets of timed automata are not closed under complement [6]. \square

3 On the State Spaces of Hybrid Automata

Since the state space of a nontrivial hybrid automaton is infinite, it cannot be explored by enumerating states. We analyze the state space of a hybrid automaton by computing with finite symbolic representations of infinite regions. For example, if x is a real-numbered variable, then the predicate $1 \leq x \leq 5$ is a finite symbolic representation of an infinite set of real numbers.

3.1 Symbolic Analysis of Transition Systems

A labeled transition system can be analyzed using symbolic representations of regions if there are algorithms for performing certain operations on the symbolic representations.

Definition 3.1 [Theories for transition systems] Consider a labeled transition system S with the state space Q . A *theory* T for S is a set of predicates that are assigned truth values by the states in Q . Given a predicate p of T , we write $\llbracket p \rrbracket$ for the set of states in Q that satisfy p , and we say that p *defines* the region $\llbracket p \rrbracket \subseteq Q$. A set Φ of predicates from T *induces* an equivalence relation \approx^Φ on Q : for all states q and r of S , define $q \approx^\Phi r$ iff q and r satisfy the same predicates in Φ . The theory T is *decidable* if for each predicate p of T , it can be decided whether $\llbracket p \rrbracket$ is empty. The theory T is *effectively closed under boolean operations* if for all predicates p_1 and p_2 of T , one can construct a predicate $Or(p_1, p_2)$ of T that defines the region $\llbracket p_1 \rrbracket \cup \llbracket p_2 \rrbracket$, and a predicate $Not(p_1)$ that defines the region $Q \setminus \llbracket p_1 \rrbracket$. The theory T is *effectively closed under transitions* if for each predicate p of T , and each label a of S , one can construct a predicate $Post(p, a)$ of T that defines the region $post_a(\llbracket p \rrbracket)$,² and a predicate $Pre(p, a)$ that defines the region $pre_a(\llbracket p \rrbracket)$. The theory T *permits the*

²Effective closure under *post* is not required for the results presented in this paper.

symbolic analysis of S if (1) T is decidable, (2) T is effectively closed under boolean operations and transitions, and (3) there is a predicate of T that defines the set of initial states of S . \square

Definition 3.2 [Similarity, bisimilarity, and trace equivalence] Consider a labeled transition system S with the state space Q , and an equivalence relation \approx on Q . A \approx -*simulation* \preceq of S is a binary relation on Q such that $q \preceq r$ implies (1) $q \approx r$ and (2) for each label a of S , if $q \xrightarrow{a} q'$, then there exists a state r' such that $r \xrightarrow{a} r'$ and $q' \preceq r'$. A symmetric \approx -simulation is called a \approx -*bisimulation*. Two states q and r of S are \approx -*similar* if there is a \approx -simulation \preceq such that $q \preceq r$ and $r \preceq q$. The two states q and r are \approx -*bisimilar* if there is a \approx -bisimulation \simeq such that $q \simeq r$. The two states q and r are *trace equivalent* if the labeled transition systems $S[Q^0 := \{q\}]$ and $S[Q^0 := \{r\}]$ have the same finite traces. If T is a theory for S , and Φ is a set of predicates from T , then the \approx^Φ -(bi)similarity relation of S is called Φ -(bi)similarity. \square

Remark. [Definition 3.2] We remind the reader of some well-known facts from concurrency theory. If two states q and r are \approx -bisimilar, then q and r are also \approx -similar. If there is an equivalence relation \approx such that the two states q and r are \approx -similar, then q and r are trace equivalent. The converse of either statement is not necessarily true. \square

Bisimilarity and similarity relations of a labeled transition system S can be defined as greatest fixpoints of a monotonic operator, and approximated by iterating the operator. The iteration can be performed in a theory that permits the symbolic analysis of S . The iteration terminates iff the approximated equivalence relation has finitely many equivalence classes.

Definition 3.3 [Finitary equivalences] An equivalence relation \simeq is called *finitary* if there are finitely many \simeq -equivalence classes. \square

Proposition 3.1 [Symbolic bisimilarity approximation] *Let S be a labeled transition system with a finite label set, let T be a theory that permits the symbolic analysis of S , and let Φ be a finite set of predicates from T . Each step of the procedure *BisimApprox* (Figure 5) is effective, and upon termination the procedure returns the Φ -bisimilarity relation of S . Furthermore, the procedure *BisimApprox* terminates iff the Φ -bisimilarity relation of S is finitary. \square*

Proposition 3.2 [Symbolic similarity approximation] [19] *Let S , T , and Φ be as in Proposition 3.1. Each step of the procedure *SimApprox* (Figure 6) is effective, and upon termination the procedure returns*

procedure *BisimApprox*:

Input: a labeled transition system S with label set A , and a set Φ of predicates.

Output: the set Π of equivalence classes of the Φ -bisimilarity relation of S .

let Π be the set of \approx^Φ -equivalence classes;

while there are two regions $R, R' \in \Pi$ and a label $a \in A$ such that

both $R \cap \text{pre}_a(R')$ and $R \setminus \text{pre}_a(R')$ are nonempty do

$(R_1, R_2) := (R \cap \text{pre}_a(R'), R \setminus \text{pre}_a(R'))$;

$\Pi = (\Pi \setminus \{R\}) \cup \{R_1, R_2\}$

od;

return Π .

Figure 5: Symbolic bisimilarity computation

the Φ -similarity relation of S . Furthermore, the procedure *SimApprox* terminates iff the Φ -similarity relation of S is finitary. \square

Remark. [Proposition 3.2] For two states q and r of S , there is a \approx -simulation \preceq with $q \preceq r$ iff upon termination of the procedure *SimApprox*, $q \in R$ and $r \in \text{sim}(R)$ for some region $R \in \Pi$. \square

If a finitary bisimilarity or similarity relation of an infinite-state transition system S can be computed, then many verification problems for S can be reduced to finite-state problems. Alternatively, if a verification task can be stated in the μ -calculus, then we may compute directly on the infinite state space without computing a finitary reduction. The μ -calculus defines monotonic operators on regions, and the iteration of these operators can be performed in a theory that permits the symbolic analysis of S . The iteration is guaranteed to terminate if a suitable finitary reduction of the state space exists.

Definition 3.4 [The μ -calculus] Consider a labeled transition system S with the state space Q , and a theory T for S . The formulas of the (S, T) -based μ -calculus are generated by the grammar

$$\phi ::= p \mid \phi_1 \vee \phi_2 \mid \neg\phi \mid \exists @\phi \mid \mu R. \phi \mid R$$

for predicates p of T , labels a of S , and region variables R . A formula of the (S, T) -based μ -calculus is *legal* if each occurrence of a region variable is bound by a μ -quantifier and separated from the quantifier by an even number of negations. Given a map F that assigns to each region variable a region of S , every subformula of a legal formula of the (S, T) -based μ -calculus defines a region $[\phi]^F \subseteq Q$:

$$\begin{aligned} [p]^F &= [p] \\ [\phi_1 \vee \phi_2]^F &= [\phi_1]^F \cup [\phi_2]^F \\ [\neg\phi]^F &= Q \setminus [\phi]^F \end{aligned}$$

$$\begin{aligned} [\exists @\phi]^F &= \text{pre}_a([\phi]^F) \\ [\mu R. \phi]^F &= \bigcap \{Q' \subseteq Q \mid Q' = [\phi]^F[R:=Q']\} \\ [R]^F &= F(R) \end{aligned}$$

The legal formula ϕ of the (S, T) -based μ -calculus defines the region $[\phi] = [\phi]^F$, for some map F . The formula ϕ is *existential* if every occurrence of an $\exists \bigcirc$ connective lies within an even number of negations, and ϕ is *universal* if every occurrence of an $\exists \bigcirc$ connective lies within an odd number of negations. \square

Remark. [Negation-free μ -formulas] Let $\forall @\phi$ stand for $\neg \exists @\neg\phi$, and let $\nu R. \phi$ stand for $(\neg \mu R. \neg\phi[R := \neg R])$. Using the defined connectives \wedge , $\forall \bigcirc$, and ν , every formula of the μ -calculus can be translated into an equivalent formula in negation-free form, where all \neg connectives occur in front of predicates. If ϕ is a formula in negation-free form, then ϕ is existential iff ϕ contains no occurrence of the $\forall \bigcirc$ connective, and ϕ is universal iff ϕ contains no occurrence of the $\exists \bigcirc$ connective. \square

Remark. [Reachability and controllability] We mention two of the many system requirements that can be checked using the μ -calculus. Let S be a labeled transition system with a finite label set A , let p^0 be a predicate that defines the set of initial states of S , and let p^* be a predicate that defines a set of error states of S . There is no trajectory from an initial state to an error state iff the existential μ -formula $p^0 \wedge (\mu R. p^* \vee \bigvee_{a \in A} \exists @R)$ defines the empty region. Suppose that the behavior of S can be influenced by applying a control map that maps each state of S to a label in A [35]: during the execution of S , in each state q , the control map chooses a label a , and the system proceeds to an a -successor of q . There is a control map that keeps the system from entering an error state iff the existential μ -formula $p^0 \wedge (\mu R. p^* \vee \bigwedge_{a \in A} \exists @R)$ defines the empty region. \square

procedure *SimApprox*:

Input: a labeled transition system S with label set A , and a set Φ of predicates.
Output: the set Π of equivalence classes of the Φ -similarity relation of S .

let Π be the set of \approx^Φ -equivalence classes;
for each region $R \in \Pi$ do $\text{sim}(R) := R$ od;
while there are two regions $R, R' \in \Pi$ and a label $a \in A$ such that
both $R \cap \text{pre}_a(\text{sim}(R'))$ and $\text{sim}(R) \setminus \text{pre}_a(\text{sim}(R'))$ are nonempty do
 $(R_1, R_2) := (R \cap \text{pre}_a(\text{sim}(R')), R \setminus \text{pre}_a(\text{sim}(R')))$;
 $\Pi = (\Pi \setminus \{R\}) \cup \{R_1\}$;
 $\text{sim}(R_1) := \text{sim}(R) \cap \text{pre}_a(\text{sim}(R'))$;
 if R_2 is nonempty then $\Pi := \Pi \cup \{R_2\}$; $\text{sim}(R_2) := \text{sim}(R)$ fi
od;
return Π .

Figure 6: Symbolic similarity computation

procedure *MuApprox*(S, ϕ):

Input: a labeled transition system S , and a legal formula ϕ of the μ -calculus.
Output: a predicate p such that p and ϕ define the same region of S .

case ϕ is a predicate p :
 return p
case ϕ has the form $\phi_1 \vee \phi_2$:
 return $\text{Or}(\text{MuApprox}(S, \phi_1), \text{MuApprox}(S, \phi_2))$
case ϕ has the form $\neg\phi'$:
 return $\text{Not}(\text{MuApprox}(S, \phi'))$
case ϕ has the form $\exists @ \phi'$:
 return $\text{Pre}(\text{MuApprox}(S, \phi'), a)$
case ϕ has the form $\mu R. \phi'$:
 $p_1 := \text{false}$;
 repeat
 $p_2 := p_1$; $p_1 := \text{MuApprox}(S, \phi'[R := p_2])$
 until $\llbracket p_1 \rrbracket = \llbracket p_2 \rrbracket$;
 return p_1
end case.

Proposition 3.3 [Symbolic μ -calculus approximation] *Let T be a theory that permits the symbolic analysis of the labeled transition system S , let ϕ be a legal formula of the (S, T) -based μ -calculus, and let Φ be the set of predicates that occur in ϕ . Each step of the procedure *MuApprox* is effective, and upon termination the procedure returns a predicate p of T such that $\llbracket p \rrbracket = \llbracket \phi \rrbracket$. Furthermore, the procedure *MuApprox* is guaranteed to terminate if the Φ -bisimilarity relation of S is finitary. If ϕ is existential or universal, then the procedure *MuApprox* is guaranteed to terminate if the Φ -similarity relation of S is finitary. \square*

Proof. For termination, observe that each predicate that is generated by the procedure *MuApprox* defines a union of equivalence classes of the Φ -bisimilarity relation of S . Furthermore, if ϕ is existential, then the operator *Pre* is applied only to regions R that

are closed under simulators; that is, if $q \in R$ and $q \preceq r$ for some Φ -simulation \preceq , then $r \in R$. \square

Remark. [Termination] If the procedure *BisimApprox* terminates, then the procedure *SimApprox* terminates also. The converse is not necessarily true. For any given input formula, the procedure *MuApprox* may terminate even if the procedures *SimApprox* and *BisimApprox* do not terminate. This encourages practical experimentation, especially since in practice, with concrete time and space constraints, strong termination guarantees are always elusive. \square

3.2 Linear Hybrid Automata

The hybrid automata that can be analyzed symbolically in the theory of the reals with addition are called linear.

Definition 3.5 [Linear hybrid automata] A *linear term* is an expression of the form $k_0 + k_1x_1 + \dots + k_mx_m$, for real-numbered variables x_1, \dots, x_m and integer constants k_0, \dots, k_m . If t_1 and t_2 are linear terms, then $t_1 \leq t_2$ is a *linear inequality*. A hybrid automaton H is *linear* if the following two restrictions are met.

1. The initial, invariant, flow, and jump conditions of H are boolean combinations of linear inequalities.
2. If X is the set of variables of H , then the flow conditions of H contain free variables from \dot{X} only. \square

Remark. [Definition 3.5] The linear hybrid automata are closed under composition. All (multi)-rectangular automata and all triangular automata are

linear hybrid automata. The use of general linear hybrid automata for approximating complex continuous behavior can be more efficient than the use of rectangular automata [29, 37]. \square

Definition 3.6 [Theories for hybrid automata] Consider a hybrid automaton H with the set X of variables and the set V of control modes. An H -predicate is a predicate whose free variables are boolean-valued variables from V and real-valued variables from X . A state (v, \mathbf{x}) of H satisfies the H -predicate p if the closed predicate $p[v, V \setminus \{v\}, X := \text{true}, \text{false}, \mathbf{x}]$ is true. The H -predicate p is *linear* if p is a boolean combination of (1) boolean-valued variables from V and (2) linear inequalities whose (real-valued) variables are from X . The linear H -predicate p is a *rectangular H -predicate* if each linear inequality in p can be written in the form $x \leq k$ or $x \geq k$, for a variable x and a rational constant k . An H -formula of the *rectangular μ -calculus* is a formula of the (S_H^a, T) -based μ -calculus, for the set T of rectangular H -predicates. \square

Theorem 3.1 [Symbolic analysis of linear hybrid automata] [5, 8] *For every linear hybrid automaton H , the set of linear H -predicates permits the symbolic analysis of the time-abstract transition system S_H^a .*

Proof. The linear H -predicates are quantifier-free formulas of the first-order theory $(\mathbb{R}, +, \leq, 0, 1)$ of the reals with addition, comparison, and integer constants. Each time transition of a linear hybrid automaton has a witness that can be decomposed into a finite sequence of straight lines. Then, using quantification over the reals, the *Post* and *Pre* operations can be expressed in the theory $(\mathbb{R}, +, \leq, 0, 1)$. The proof concludes with the observation that the first-order theory $(\mathbb{R}, +, \leq, 0, 1)$ admits quantifier elimination. \square

Remark. [Timed automata] The symbolic analysis of singular automata does not require the full theory of linear predicates, which leads to more efficient implementations [18, 10, 11, 17]. For a hybrid automaton A , a linear H -predicate p is a *triangular H -predicate* if each linear inequality in p can be written in the form $x \leq k$, $x \geq k$, $x \leq y + k$, or $x \geq y + k$, for variables x and y and a rational constant k . For every timed automaton H , the set of triangular H -predicates permits the symbolic analysis of the time-abstract transition system S_H^a [28]. \square

Remark. [Polynomial hybrid automata] Since the theory of the reals with addition and multiplication is decidable, it is possible to define a class of hybrid

automata that are more general than linear hybrid automata and can be analyzed symbolically in the more powerful theory. The practicality of such a generalization has not been studied. \square

3.3 Bisimilarity and Similarity Relations

From Theorem 2.1 it follows that for every rectangular automaton H , the trace equivalence relation of the time-abstract transition system S_H^a is finitary. If we can identify subclasses of rectangular automata whose time-abstract transition systems have finitary similarity or bisimilarity relations, then we obtain termination guarantees for the procedure *MuApprox* applied to hybrid automata.

Definition 3.7 [Timed and time-abstract (bi)similarity] Consider a hybrid automaton H and a set Φ of H -predicates. The Φ -(bi)similarity relation of the timed transition system S_H^t is called the *timed Φ -(bi)similarity relation* of H , and the Φ -(bi)similarity relation of the time-abstract transition system S_H^a is called the *time-abstract Φ -(bi)similarity relation* of H . \square

The fundamental theorem of timed automata shows the existence of finitary time-abstract bisimilarity relations for timed automata. This result can be extended to singular automata.

Theorem 3.2 [Time-abstract bisimilarity of singular automata] [6, 5] *If H is a singular automaton, and Φ is a finite set of rectangular H -predicates, then the time-abstract Φ -bisimilarity relation of H is finitary.*

Proof. A rectangular automaton H and a finite set Φ of rectangular H -predicates are *normalized* if all non-flow interval endpoints in H and all constants in Φ are integers. Normalization can be achieved by multiplying all non-flow interval endpoints in H and all constants in Φ by a suitably chosen integer constant. Assuming that H and Φ are normalized, let K be the largest integer constant that occurs in H and Φ . If H is a timed automaton, then the (finite) set Φ of triangular H -predicates with integer constants no larger than K induces a finitary Φ -bisimulation on the time-abstract transition system S_H^a . The first panel of Figure 7 shows the induced bisimulation on the unit square $[0, 1]^2$ for a 2D timed automaton H . For instance, if $v = v'$ and $0 < \mathbf{x}_1 < \mathbf{x}_2 < 1$ and $0 < \mathbf{x}'_1 < \mathbf{x}'_2 < 1$, then the two states (v, \mathbf{x}) and (v', \mathbf{x}') of H are time-abstract Φ -bisimilar. If H is a singular automaton, then a slight




Figure 7: Four finitary equivalence relations on the unit square

extension of triangular H -predicates needs to be considered. For example, the second and third panels of Figure 7 show the induced bisimulations on the unit square for 2D singular automata with the flow rectangles $[2, 2] \times [1, 1]$ and $[1, 1] \times [3, 3]$, respectively. \square

Corollary 3.1 [Symbolic μ -calculus model checking for singular automata] *The procedure MuApprox terminates if given the time-abstract transition system of a singular automaton H and an H -formula ϕ of the rectangular μ -calculus.*

Remark. [Nonsingular automata and nonrectangular predicates] Singular automata with rectangular predicates identify a boundary between the existence and nonexistence of finitary bisimilarity relations. In fact, for the following three generalizations, bisimilarity degenerates to equality on infinite state spaces. Let $\Phi_1 = \{x_1 = 1, x_2 = 1, x_1 \leq x_2\}$ and $\Phi_2 = \{x_1 = 1, x_2 = 1\}$.

1. There is an (infinite-state) singular automaton H such that the time-abstract Φ_1 -bisimilarity relation of H is equality.
2. There is an (infinite-state) multisingular automaton H such that the time-abstract Φ_2 -bisimilarity relation of H is equality.
3. There is an (infinite-state) 2D rectangular automaton H such that the time-abstract Φ_2 -bisimilarity relation of H is equality. \square

The boundary between the existence and nonexistence of finitary similarity relations lies at 2D rectangular automata.

Theorem 3.3 [Time-abstract similarity of 2D rectangular automata] [19] *If H is a 2D rectangular automaton, and Φ is a finite set of rectangular H -predicates, then the time-abstract Φ -similarity relation of H is finitary.*

Proof. The structure of the finitary time-abstract similarity relation is best illustrated with an example. Let H be a 2D rectangular automaton with the flow rectangle $[1, 2] \times [1, 3]$. Assuming that H and Φ are normalized, the fourth panel of Figure 7 shows a finitary kernel of a time-abstract Φ -simulation on the unit square. The simulation is obtained by intersecting the bisimulations for the two cases of extremal flow: maximal \dot{x}_1 and minimal \dot{x}_2 , and vice versa. \square

Corollary 3.2 [Symbolic μ -calculus model checking for 2D rectangular automata] *The procedure MuApprox terminates if given the time-abstract transition system of a 2D rectangular automaton H and an existential or universal H -formula ϕ of the rectangular μ -calculus.*

Theorem 3.3 does not generalize to rectangular automata of arbitrary dimension.

Theorem 3.4 [Time-abstract similarity of 3D rectangular automata] [25] *Let $\Phi = \{x_1 = 1, x_2 = 1, x_3 = 1\}$. There is an (infinite-state) 3D rectangular automaton H such that the time-abstract Φ -similarity relation of H is equality.*

In summary, rectangular automata are a maximal class of hybrid automata with finitary time-abstract trace equivalence relations, 2D rectangular automata are a maximal class of hybrid automata with finitary time-abstract similarity relations, and singular automata are a maximal class of hybrid automata with finitary time-abstract bisimilarity relations.

Remark. [Context-free equivalences] We have restricted ourselves to decidability results that can be obtained by relating hybrid automata to finite automata. Additional decidability results can be obtained by relating hybrid automata to pushdown automata [13, 15]. Little is known, however, about which classes of hybrid automata are time-abstract trace equivalent (similar; bisimilar) to pushdown automata. \square

Remark. [Timed (bi)similarity] We have focused on time-abstract state space equivalences, because the timed counterparts are infinitary already for nontrivial timed automata. From Theorem 2.3 it follows that for timed automata, the timed trace equivalence of two states cannot be decided. It should be noted, however, that timed similarity and timed bisimilarity can be decided for timed automata. Specifically, if H is a timed automaton, Φ is a finite set of rectangular H -predicates, and q and r are two states of H , then it can be decided if q and r are timed Φ -(bi)similar [16, 39]. \square

3.4 Computation Tree Logics

We have studied the structure of the timed and time-abstract transition systems of hybrid automata. These transition systems, however, may not be directly useful for (dis)proving assertions about the behavior of a hybrid automaton H , because each trajectory of S_H^t and S_H^a only samples a piecewise-continuous trajectory of H at certain discrete points. In the following, we restrict ourselves to the time-abstract view. Since each time transition of S_H^a abstracts all information about intermediate states that are visited, by looking only at a trajectory of S_H^a , it is impossible to check if the corresponding piecewise-continuous trajectory of H visits any given state or region. We solve this problem by defining (time-abstract) observational transition systems, where each time transition is labeled with a region: the time transition t is labeled with the region R iff all intermediate states and the target state of t lie within R . Thus, an observational transition system results from the continuous observation of a hybrid automaton, with limited observational power: for a given set \mathcal{R} of regions, it can be observed whether a continuous trajectory fragment stays within any of the regions from \mathcal{R} .

Definition 3.8 [Piecewise-continuous semantics of hybrid automata] Consider a hybrid automaton H and a set \mathcal{R} of H -regions. The \mathcal{R} -observational transition system $S_H^{\mathcal{R}}$ of H is the labeled transition system with the components Q , Q^0 , A , and \xrightarrow{c} for each $c \in C$, defined as follows.³

- Q and Q^0 are defined as in Definition 1.3.
- $C = \Sigma \cup \mathcal{R}$.
- For each event $\sigma \in \Sigma$, define $\xrightarrow{\sigma}$ as in Definition 1.3.

³If $\mathcal{R} = \{R\}$ for a single H -region R , then we write S_H^R for the \mathcal{R} -observational transition system $S_H^{\mathcal{R}}$.

- For each region $R \in \mathcal{R}$, define $(v, \mathbf{x}) \xrightarrow{R} (v', \mathbf{x}')$ iff there is a nonnegative real $\delta \in \mathbb{R}_{\geq 0}$ and a witness f for the transition $(v, \mathbf{x}) \xrightarrow{\delta} (v', \mathbf{x}')$ such that for all reals $\varepsilon \in (0, \delta]$, $(v, f(\varepsilon)) \in R$. The real δ is a *possible duration* of the transition $(v, \mathbf{x}) \xrightarrow{R} (v', \mathbf{x}')$.

An infinite trajectory $\langle c_i, q_i \rangle_{i \geq 1}$ of the \mathcal{R} -observational transition system $S_H^{\mathcal{R}}$ *diverges* if there is an infinite sequence $\langle \delta_i \rangle_{i \geq 0}$ of reals such that (1) the infinite sum $\sum_{i \geq 1} \delta_i$ diverges, and (2) for all $i \geq 0$, either $a_i \in \Sigma$ and $\delta_i = 0$, or $a_i \in \mathcal{R}$ and δ_i is a possible duration of the corresponding transition $q_{i-1} \xrightarrow{a_i} q_i$. A set Φ of H -predicates *permits the observational symbolic analysis* of the hybrid automaton H if Φ permits the symbolic analysis of the observational transition system $S_H^{\mathcal{R}}$, where \mathcal{R} is the set of H -regions that are definable by predicates in Φ . An equivalence relation \simeq on Q is an *observational Φ -(bi)similarity relation* of H if \simeq is the Φ -(bi)similarity relation of the observational transition system $S_H^{\mathcal{R}}$, where \mathcal{R} is the set of \simeq -equivalence classes. \square

Since observational transition systems are defined in a time-abstract way, the results of Theorems 3.1, 3.2, and 3.3 carry over from time-abstract to observational transition systems.

Proposition 3.4 [Observational symbolic analysis of linear hybrid automata] *For every linear hybrid automaton H , the set of linear H -predicates permits the observational symbolic analysis of H . Consider a finite set Φ of rectangular H -predicates. If H is a singular automaton, then there is a finitary observational Φ -bisimilarity relation of H . If H is a 2D rectangular automaton, then there is a finitary observational Φ -similarity relation of H .*

For (dis)proving assertions about the infinite behavior of hybrid automata, we need to take into account the liveness assumption that time diverges. Computation tree logics for hybrid automata are branching-time temporal logics for stating requirements about divergent piecewise-continuous trajectories.

Definition 3.9 [Computation tree logics] Consider a linear hybrid automaton H with the state space Q . The H -formulas of linear CTL are generated by the grammar

$$\phi ::= p \mid \phi_1 \vee \phi_2 \mid \neg \phi \mid \phi_1 \exists \mathcal{U} \phi_2 \mid \exists \Box \phi$$

for linear H -predicates p . Every H -formula of linear CTL defines an H -region $\llbracket \phi \rrbracket \subseteq Q$:

$$\llbracket \phi_1 \vee \phi_2 \rrbracket = \llbracket \phi_1 \rrbracket \cup \llbracket \phi_2 \rrbracket$$

$\llbracket \neg\phi \rrbracket = Q \setminus \llbracket \phi \rrbracket$
 $q \in \llbracket \phi_1 \exists \mathcal{U} \phi_2 \rrbracket$ iff the observational transition system $S_H^{\llbracket \phi_1 \vee \phi_2 \rrbracket}$ has a (finite) q -rooted trajectory that visits a state in $\llbracket \phi_2 \rrbracket$
 $q \in \llbracket \exists \square \phi \rrbracket$ iff the observational transition system $S_H^{\llbracket \phi \rrbracket}$ has a divergent q -rooted trajectory

The H -formula ϕ of linear CTL is an H -formula of *rectangular* CTL if all predicates that occur in ϕ are rectangular H -predicates. The H -formula ϕ of rectangular CTL is an H -formula of *rectangular* \exists CTL if each occurrence of the $\exists \mathcal{U}$ connective in ϕ and each occurrence of the $\exists \square$ connective in ϕ lies within the scope of an even number of negations, and ϕ is an H -formula of *rectangular* \forall CTL if all occurrences of $\exists \mathcal{U}$ and $\exists \square$ lie within the scope of odd numbers of negations. \square

Remark. [Definition 3.9] The semantics of linear CTL is defined to be piecewise-continuous (it refers to trajectories of observational rather than time-abstract transition systems), live (it refers to divergent trajectories), and strict (the temporal connectives do not impose requirements on the current state; for example, a nonstrict version of $\phi_1 \exists \mathcal{U} \phi_2$ can be defined as $\phi_2 \vee (\phi_1 \exists \mathcal{U} \phi_2)$). The disjunction in the semantic clause for the $\exists \mathcal{U}$ connective is necessary, because a switch from ϕ_1 being true to ϕ_2 being true may occur from a right-closed to a left-open interval. In the following, we write $\exists \diamond \phi$ for $\text{true} \exists \mathcal{U} \phi$, and $\forall \square \phi$ for $\neg \exists \diamond \neg \phi$. \square

If a system requirement is given as a formula of a computation tree logic, then the corresponding verification task is a model-checking problem.

Definition 3.10 [Model checking] The *model-checking problem* for a class \mathcal{H} of hybrid automata and a computation tree logic \mathcal{L} asks, given a hybrid automaton H from \mathcal{H} and an H -formula ϕ from \mathcal{L} , if $\llbracket \phi \rrbracket$ contains all initial states of H . \square

For a linear hybrid automaton H , an H -formula ϕ of linear CTL can be translated into a formula ϕ' of the μ -calculus. The piecewise-continuity of linear CTL is taken care of by interpreting ϕ' over an observational transition system $S_H^{\mathcal{R}}$. The liveness of linear CTL is taken of by interpreting ϕ' over an extension of $S_H^{\mathcal{R}}$ with a clock variable that can observe the divergence of time. The translation leads to a model-checking procedure for linear hybrid automata and linear CTL.

Definition 3.11 [Clock extension] A *clock automaton* H_z is a timed automaton with a single variable, z ,

a single control mode with the initial condition $z = 0$ and the invariant condition *true*, and a single control switch with the jump condition $z' = 0$. If H is a hybrid automaton and z is not a variable of H , then the composition $H \parallel H_z$ is called a *clock extension* of H . A procedure is an *effective procedure for the observational symbolic analysis* of the hybrid automaton H if each step of the procedure is either effective or a subroutine call of the form $\text{MuApprox}(S_H^{\llbracket p \rrbracket}, \phi)$, for a clock extension H' of H , a linear H' -predicate p , and a formula ϕ of the $(S_H^{\llbracket p \rrbracket}, T)$ -based μ -calculus, where T is the set of linear H' -predicates. \square

Theorem 3.5 [From CTL to the μ -calculus] [28, 8] *Let H be a linear hybrid automaton and let ϕ be an H -formula of linear CTL. There is an effective procedure for the observational symbolic analysis of H which, upon termination, returns a linear H -predicate p with $\llbracket p \rrbracket = \llbracket \phi \rrbracket$. Furthermore, the procedure is guaranteed to terminate if H is a singular automaton and ϕ is a formula of rectangular CTL, and if H is a 2D rectangular automaton and ϕ is a formula of rectangular \exists CTL or of rectangular \forall CTL.*

Proof. The CTL formula $\phi_1 \exists \mathcal{U} \phi_2$ can be translated to the μ -formula $(\mu R. \exists \diamond (\phi_2 \vee R))$, for $c = \llbracket \phi_1 \vee \phi_2 \rrbracket$. The H -formula $\exists \square \phi$ can be translated to the H' -formula $(\nu R. \phi \exists \mathcal{U} (\phi \wedge z = 1 \wedge \phi \exists \mathcal{U} (\phi \wedge z = 0 \wedge R)))$, where $H' = H \parallel H_z$ is a clock extension of H . The latter formula asserts that ϕ is true throughout some infinite trajectory along which $z = 1$ is true infinitely often and $z = 0$ is true infinitely often. This can be the case if and only if the trajectory diverges. \square

Corollary 3.3 [CTL model checking] *The model-checking problem for rectangular CTL is PSPACE-decidable for singular automata. The model-checking problems for rectangular \exists CTL and rectangular \forall CTL are PSPACE-decidable for 2D rectangular automata.*

Remark. [HYTECH] The procedure of Theorem 3.5 for checking linear CTL requirements of linear hybrid automata has been implemented in the tool HYTECH [23, 24]. The procedure has been found to terminate on several examples of practical interest that do not fall into any of the classes for which a priori termination guarantees can be given [30]. \square

Example 3.1 [Railroad gate control] Recall the safety requirement for the railroad gate controller from Example 1.2, namely, that the gate is fully closed whenever the train is within 10 meters of the gate. This requirement is expressed by the formula $(\text{Far} \wedge \text{Idle} \wedge \text{Open}) \rightarrow \forall \square (-10 \leq x \leq 10 \rightarrow \text{Closed})$ of rectangular \forall CTL. HYTECH simplifies this CTL

formula, fully automatically, to a linear predicate whose projection onto the u -dimension is $5u < 49$. It follows that the safety requirement is met if and only if the reaction delay u of the controller is less than 9.8 seconds. \square

Remark. [Nonzenoness] The semantics of the $\exists\mathcal{U}$ connective of linear CTL is defined over finite trajectories. The alternative interpretation of $\exists\mathcal{U}$ over finite prefixes of divergent trajectories requires that the underlying hybrid automaton is nonzeno. For a linear hybrid automaton H and a clock extension $H||H_z$, the rectangular existential formula $\phi_{nz} = (\nu R. \exists\Diamond(z = 1 \wedge \exists\Diamond(z = 0 \wedge R)))$ defines the set of states q with divergent q -rooted trajectories. Hence, if ϕ_{nz} can be simplified to a linear H -predicate p_{nz} , then the addition of p_{nz} as a conjunct to all invariant conditions of H results in a nonzeno linear hybrid automaton H_{nz} such that H and H_{nz} have the same divergent timed traces. From Theorem 3.5 it follows that this is always possible for singular and 2D rectangular automata. \square

Remark. [CTL with clocks, stopwatches, and symbolic constants] The H -formulas of linear CTL can be generalized to permit real-numbered variables that are not variables of the hybrid automaton H . In this way, linear CTL has been extended to include clocks (TCTL) [7, 2], stopwatches [14, 8], and symbolic constants [40]. The symbolic-analysis result of Theorem 3.5 continues to hold for these logics, and the decidability results of Corollary 3.3 continue to hold for TCTL. Isolated decidability results are known also for computation tree logics with a limited use of stopwatches or symbolic constants [4, 40]. \square

Acknowledgments. My view of mixed discrete-continuous systems has been shaped in collaborations with Rajeev Alur, Costas Courcoubetis, Nicolas Halbwachs, Pei-Hsin Ho, Peter Kopke, Zohar Manna, Xavier Nicollin, Amir Pnueli, Anuj Puri, Joseph Sifakis, Howard Wong-Toi, Pravin Varaiya, Moshe Vardi, and Sergio Yovine. My special thanks go to Oded Maler, Zohar, and Amir for introducing me to the problem domain [34], to Pei-Hsin and Howard for implementing HyTECH, and to Peter for many valuable comments on several drafts of this manuscript.

References

- [1] M. Abadi and L. Lamport. An old-fashioned recipe for real time. In J.W. de Bakker, K. Huizing, W.-P. de Roever, and G. Rozenberg, eds., *Real Time: Theory in Practice*, Lecture Notes in Computer Science (LNCS) 600, pp. 1–27. Springer-Verlag, 1992.
- [2] R. Alur, C. Courcoubetis, and D.L. Dill. Model checking in dense real time. *Information and Computation*, 104(1):2–34, 1993.
- [3] R. Alur, C. Courcoubetis, N. Halbwachs, T.A. Henzinger, P.-H. Ho, X. Nicollin, A. Olivero, J. Sifakis, and S. Yovine. The algorithmic analysis of hybrid systems. *Theoretical Computer Science*, 138:3–34, 1995.
- [4] R. Alur, C. Courcoubetis, and T.A. Henzinger. Computing accumulated delays in real-time systems. In C. Courcoubetis, ed., *Computer-aided Verification*, LNCS 697, pp. 181–193. Springer-Verlag, 1993.
- [5] R. Alur, C. Courcoubetis, T.A. Henzinger, and P.-H. Ho. Hybrid automata: an algorithmic approach to the specification and verification of hybrid systems. In R.L. Grossman, A. Nerode, A.P. Ravn, and H. Rischel, eds., *Hybrid Systems I*, LNCS 736, pp. 209–229. Springer-Verlag, 1993.
- [6] R. Alur and D.L. Dill. A theory of timed automata. *Theoretical Computer Science*, 126:183–235, 1994.
- [7] R. Alur and T.A. Henzinger. Logics and models of real time: a survey. In J.W. de Bakker, K. Huizing, W.-P. de Roever, and G. Rozenberg, eds., *Real Time: Theory in Practice*, LNCS 600, pp. 74–106. Springer-Verlag, 1992.
- [8] R. Alur, T.A. Henzinger, and P.-H. Ho. Automatic symbolic verification of embedded systems. *IEEE Trans. Software Engineering*, 22(3):181–201, 1996.
- [9] R. Alur, T.A. Henzinger, and M.Y. Vardi. Parametric real-time reasoning. In *Proc. 25th Annual Symp. Theory of Computing*, pp. 592–601. ACM Press, 1993.
- [10] R. Alur and R.P. Kurshan. Timing analysis in COSPAN. In R. Alur, T.A. Henzinger, and E.D. Sontag, eds., *Hybrid Systems III*, LNCS 1066, pp. 220–231. Springer-Verlag, 1996.
- [11] J. Bengtsson, K.G. Larsen, F. Larsson, P. Pettersson, and W. Yi. UPPAAL: a tool-suite for automatic verification of real-time systems. In R. Alur, T.A. Henzinger, and E.D. Sontag, eds., *Hybrid Systems III*, LNCS 1066, pp. 232–243. Springer-Verlag, 1996.
- [12] D. Bosscher, I. Polak, and F. Vaandrager. Verification of an audio-control protocol. In H. Langmaack, W.-P. de Roever, and J. Vytupil, eds., *Formal Techniques in Real-time and Fault-tolerant Systems*, LNCS 863, pp. 170–192. Springer-Verlag, 1994.
- [13] A. Bouajjani, R. Echahed, and R. Robbana. Verification of context-free timed systems using linear hybrid observers. In D.L. Dill, ed., *Computer-aided Verification*, LNCS, pp. 118–131. Springer-Verlag, 1994.
- [14] A. Bouajjani, R. Echahed, and J. Sifakis. On model checking for real-time properties with durations. In *Proc. 8th Annual Symp. Logic in Computer Science*, pp. 147–159. IEEE Computer Society Press, 1993.

- [15] A. Bouajjani and R. Robbana. Verifying ω -regular properties for subclasses of linear hybrid systems. In P. Wolper, ed., *Computer-aided Verification*, LNCS 939, pp. 437–450. Springer-Verlag, 1995.
- [16] K. Cer  ns. Decidability of bisimulation equivalence for parallel timer processes. In G. von Bochmann and D.K. Probst, eds., *Computer-aided Verification*, LNCS 663, pp. 302–315. Springer-Verlag, 1992.
- [17] C. Daws, A. Olivero, S. Tripakis, and S. Yovine. The tool KRONOS. In R. Alur, T.A. Henzinger, and E.D. Sontag, eds., *Hybrid Systems III*, LNCS 1066, pp. 208–219. Springer-Verlag, 1996.
- [18] D.L. Dill. Timing assumptions and verification of finite-state concurrent systems. In J. Sifakis, ed., *Computer-aided Verification*, LNCS 407, pp. 197–212. Springer-Verlag, 1989.
- [19] M.R. Henzinger, T.A. Henzinger, and P.W. Kopke. Computing simulations on finite and infinite graphs. In *Proc. 36rd Annual Symp. Foundations of Computer Science*, pp. 453–462. IEEE Computer Society Press, 1995.
- [20] T.A. Henzinger. Sooner is safer than later. *Information Processing Letters*, 43:135–141, 1992.
- [21] T.A. Henzinger. Hybrid automata with finite bisimulations. In Z. F  l  p and F. G  cseg, eds., *Automata, Languages, and Programming*, LNCS 944, pp. 324–335. Springer-Verlag, 1995.
- [22] T.A. Henzinger and P.-H. Ho. Algorithmic analysis of nonlinear hybrid systems. In P. Wolper, ed., *Computer-aided Verification*, LNCS 939, pp. 225–238. Springer-Verlag, 1995.
- [23] T.A. Henzinger, P.-H. Ho, and H. Wong-Toi. HYTECH: the next generation. In *Proc. 16th Annual Real-time Systems Symp.*, pp. 56–65. IEEE Computer Society Press, 1995.
- [24] T.A. Henzinger, P.-H. Ho, and H. Wong-Toi. A user guide to HYTECH. In E. Brinksma, W.R. Cleaveland, K.G. Larsen, T. Margaria, and B. Steffen, eds., *Tools and Algorithms for the Construction and Analysis of Systems*, LNCS 1019, pp. 41–71. Springer-Verlag, 1995.
- [25] T.A. Henzinger and P.W. Kopke. State equivalences for rectangular hybrid automata. In U. Montanari, ed., *Concurrency Theory*, Lecture Notes in Computer Science. Springer-Verlag, 1996.
- [26] T.A. Henzinger, P.W. Kopke, A. Puri, and P. Varaiya. What’s decidable about hybrid automata? In *Proc. 27th Annual Symp. Theory of Computing*, pp. 373–382. ACM Press, 1995.
- [27] T.A. Henzinger, P.W. Kopke, and H. Wong-Toi. The expressive power of clocks. In Z. F  l  p and F. G  cseg, eds., *Automata, Languages, and Programming*, LNCS 944, pp. 417–428. Springer-Verlag, 1995.
- [28] T.A. Henzinger, X. Nicollin, J. Sifakis, and S. Yovine. Symbolic model checking for real-time systems. *Information and Computation*, 111(2):193–244, 1994.
- [29] T.A. Henzinger and H. Wong-Toi. Linear phase-portrait approximations for nonlinear hybrid systems. In R. Alur, T.A. Henzinger, and E.D. Sontag, eds., *Hybrid Systems III*, LNCS 1066, pp. 377–388. Springer-Verlag, 1996.
- [30] P.-H. Ho and H. Wong-Toi. Automated analysis of an audio control protocol. In P. Wolper, ed., *Computer-aided Verification*, LNCS 939, pp. 381–394. Springer-Verlag, 1995.
- [31] Y. Kesten, A. Pnueli, J. Sifakis, and S. Yovine. Integration graphs: a class of decidable hybrid systems. In R.L. Grossman, A. Nerode, A.P. Ravn, and H. Rischel, eds., *Hybrid Systems I*, LNCS 736, pp. 179–208. Springer-Verlag, 1993.
- [32] N.A. Lynch, R. Segala, F. Vaandrager, and H.B. Weinberg. Hybrid I/O Automata. In R. Alur, T.A. Henzinger, and E.D. Sontag, eds., *Hybrid Systems III*, LNCS 1066, pp. 496–510. Springer-Verlag, 1996.
- [33] N.A. Lynch and F. Vaandrager. Forward and backward simulations for timing-based systems. In J.W. de Bakker, K. Huizing, W.-P. de Roever, and G. Rozenberg, eds., *Real Time: Theory in Practice*, LNCS 600, pp. 397–446. Springer-Verlag, 1992.
- [34] O. Maler, Z. Manna, and A. Pnueli. From timed to hybrid systems. In J.W. de Bakker, K. Huizing, W.-P. de Roever, and G. Rozenberg, eds., *Real Time: Theory in Practice*, LNCS 600, pp. 447–484. Springer-Verlag, 1992.
- [35] O. Maler, A. Pnueli, and J. Sifakis. On the synthesis of discrete controllers for timed systems. In E.W. Mayr and C. Puech, eds., *Theoretical Aspects of Computer Science*, LNCS 900, pp. 229–242. Springer-Verlag, 1995.
- [36] X. Nicollin, A. Olivero, J. Sifakis, and S. Yovine. An approach to the description and analysis of hybrid systems. In R.L. Grossman, A. Nerode, A.P. Ravn, and H. Rischel, eds., *Hybrid Systems I*, LNCS 736, pp. 149–178. Springer-Verlag, 1993.
- [37] A. Puri, V. Borkar, and P. Varaiya. ϵ -approximation of differential inclusions. In R. Alur, T.A. Henzinger, and E.D. Sontag, eds., *Hybrid Systems III*, LNCS 1066, pp. 362–376. Springer-Verlag, 1996.
- [38] A. Puri and P. Varaiya. Decidability of hybrid systems with rectangular differential inclusions. In D.L. Dill, ed., *Computer-aided Verification*, LNCS 818, pp. 95–104. Springer-Verlag, 1994.
- [39] S. Tasiran, R. Alur, R.P. Kurshan, and R.K. Brayton. Verifying abstractions of timed systems. In U. Montanari, ed., *Concurrency Theory*, Lecture Notes in Computer Science. Springer-Verlag, 1996.
- [40] F. Wang. Timing behavior analysis for real-time systems. In *Proc. 10th Annual Symp. Logic in Computer Science*, pp. 112–122. IEEE Computer Society Press, 1995.