

Introduction to Automata Theory, Languages, and Computation, 2nd Edition

John E. Hopcroft, Rajeev Motwani, and Jeffrey D. Ullman

Cornell, Stanford, and Stanford

© 2001 Addison-Wesley

ISBN 0-201-44124-1

Preface (Abridged)

In the preface from the 1979 predecessor to this book, Hopcroft and Ullman marveled at the fact that the subject of automata had exploded, compared with its state at the time they wrote their first book, in 1969. Truly, the 1979 book contained many topics not found in the earlier work and was about twice its size. If you compare this book with the 1979 book, you will find that, like the automobiles of the 1970's, this book is "larger on the outside, but smaller on the inside." That sounds like a retrograde step, but we are happy with the changes for several reasons.

First, in 1979, automata and language theory was still an area of active research. A purpose of that book was to encourage mathematically inclined students to make new contributions to the field. Today, there is little direct research in automata theory (as opposed to its applications), and thus little motivation for us to retain the succinct, highly mathematical tone of the 1979 book.

Second, the role of automata and language theory has changed over the past two decades. In 1979, automata was largely a graduate-level subject, and we imagined our reader was an advanced graduate student, especially those using the later chapters of the book. Today, the subject is a staple of the undergraduate curriculum. As such, the content of the book must assume less in the way of prerequisites from the student, and therefore must provide more of the background and details of arguments than did the earlier book.

A third change in the environment is that Computer Science has grown to an almost unimaginable degree in the past two decades. While in 1979 it was often a challenge to fill up a curriculum with material that we felt would survive the next wave of technology, today very many subdisciplines compete for the limited amount of space in the undergraduate curriculum.

Fourthly, CS has become a more vocational subject, and there is a severe pragmatism among many of its students. We continue to believe that aspects of automata theory are essential tools in a variety of new disciplines, and we believe that the theoretical, mind-expanding exercises embodied in the typical automata course retain their value, no matter how much the student prefers to learn only the most immediately monetizable technology. However, to assure a continued place for the subject on the menu of topics available to the computer science student, we believe it is necessary to emphasize the applications along with the mathematics. Thus, we have replaced a number of the more abstruse topics in the earlier book with examples of how the ideas are used today. While applications of automata and language theory to compilers are now so well understood that they are normally covered in a compiler course, there are a variety of more recent uses, including model-checking algorithms to verify protocols and document-description languages that are patterned on context-free grammars.

A final explanation for the simultaneous growth and shrinkage of the book is that we were today able to take advantage of the \TeX and \LaTeX typesetting systems developed by Don Knuth and Les Lamport. The latter, especially, encourages the "open" style of typesetting that makes books larger, but easier to read. We appreciate the efforts of both men.

Use of the Book

This book is suitable for a quarter or semester course at the Junior level or above. At Stanford, we have used the notes in CS154, the course in automata and language theory. It is a one-quarter course, which both Rajeev and Jeff have taught. Because of the limited time available, Chapter II is not covered, and some of the later material, such as the more difficult polynomial-time reductions in Section 10.4 are omitted as well. The book's Web site (see below) includes notes and syllabi for several offerings of CS154.

Some years ago, we found that many graduate students came to Stanford with a course in automata theory that did not include the theory of intractability. As the Stanford faculty believes that these ideas are essential for every computer scientist to know at more than the level of "NP-complete means it takes too long," there is another course, CS154N, that students may take to cover only Chapters 8, 9, and 10. They actually participate in roughly the last third of CS154 to fulfill the CS154N requirement. Even today, we find several students each quarter availing themselves of this option. Since it requires little extra effort, we recommend the approach.

Prerequisites

To make best use of this book, students should have taken previously a course covering discrete mathematics, e.g., graphs, trees, logic, and proof techniques. We assume also that they have had several courses in programming, and are familiar with common data structures, recursion, and the role of major system components such as compilers. These prerequisites should be obtained in a typical freshman-sophomore CS program.

Exercises

The book contains extensive exercises, with some for almost every section. We indicate harder exercises or parts of exercises with an exclamation point. The hardest exercises have a double exclamation point.

Some of the exercises or parts are marked with a star. For these exercises, we shall endeavor to maintain solutions accessible through the book's Web page. These solutions are publicly available and should be used for self-testing. Note that in a few cases, one exercise B asks for modification or adaptation of your solution to another exercise A. If certain parts of A have solutions, then you should expect the corresponding parts of B to have solutions as well.

Support on the World Wide Web

The book's home page is

<http://www-db.stanford.edu/~ullman/ialc>

Here are solutions to starred exercises, errata as we learn of them, and backup materials. We hope to make available the notes for each offering of CS154 as we teach it, including homeworks, solutions, and exams.

Contents

1. Automata: The Methods and the Madness	1
1.1 Why Study Automata Theory? 2 / 1.1.1 Introduction to Finite Automata 2 / 1.1.2 Structural Representations 4 / 1.1.3 Automata and Complexity 5 / 1.2 Introduction	

to Formal Proof 5 / 1.2.1 Deductive Proofs 6 / 1.2.2 Reduction to Definitions 8 / 1.2.3 Other Theorem Forms 10 / 1.2.4 Theorems That Appear Not to Be If-Then Statements 13 / 1.3 Additional Forms of Proof 13 / 1.3.1 Proving Equivalences About Sets 14 / 1.3.2 The Contrapositive 14 / 1.3.3 Proof by Contradiction 16 / 1.3.4 Counterexamples 17 / 1.4 Inductive Proofs 19 / 1.4.1 Inductions on Integers 19 / 1.4.2 More General Forms of Integer Inductions 22 / 1.4.3 Structural Inductions 23 / 1.4.4 Mutual Inductions 26 / 1.5 The Central Concepts of Automata Theory 28 / 1.5.1 Alphabets 28 / 1.5.2 Strings 29 / 1.5.3 Languages 30 / 1.5.4 Problems 31 / 1.6 Summary of Chapter 1 34 / 1.7 References for Chapter 1 35

2. Finite Automata 37

2.1 An Informal Picture of Finite Automata 38 / 2.1.1 The Ground Rules 38 / 2.1.2 The Protocol 39 / 2.1.3 Enabling the Automata to Ignore Actions 41 / 2.1.4 The Entire System as an Automaton 43 / 2.1.5 Using the Product Automaton to Validate the Protocol 45 / 2.2 Deterministic Finite Automata 45 / 2.2.1 Definition of a Deterministic Finite Automaton 46 / 2.2.2 How a DFA Processes Strings 46 / 2.2.3 Simpler Notations for DFA's 48 / 2.2.4 Extending the Transition Function to Strings 49 / 2.2.5 The Language of a DFA 52 / 2.2.6 Exercises for Section 2.2 53 / 2.3 Nondeterministic Finite Automata 55 / 2.3.1 An Informal View of Nondeterministic Finite Automata 56 / 2.3.2 Definition of Nondeterministic Finite Automata 57 / 2.3.3 The Extended Transition Function 58 / 2.3.4 The Language of an NFA 59 / 2.3.5 Equivalence of Deterministic and Nondeterministic Finite Automata 60 / 2.3.6 A Bad Case for the Subset Construction 65 / 2.3.7 Exercises for Section 2.3 66 / 2.4 An Application: Text Search 68 / 2.4.1 Finding Strings in Text 68 / 2.4.2 Nondeterministic Finite Automata for Text Search 69 / 2.4.3 A DFA to Recognize a Set of Keywords 70 / 2.4.4 Exercises for Section 2.4 72 / 2.5 Finite Automata With Epsilon-Transitions 72 / 2.5.1 Uses of epsilon-Transitions 72 / 2.5.2 The Formal Notation for an epsilon-NFA 74 / 2.5.3 Epsilon-Closures 75 / 2.5.4 Extended Transitions and Languages for epsilon-NFA's 76 / 2.5.5 Eliminating epsilon-Transitions 77 / 2.5.6 Exercises for Section 2.5 80 / 2.6 Summary of Chapter 2 80 / 2.7 References for Chapter 2 81

3. Regular Expressions and Languages 83

3.1 Regular Expressions 83 / 3.1.1 The Operators of Regular Expressions 84 / 3.1.2 Building Regular Expressions 85 / 3.1.3 Precedence of Regular-Expression Operators 88 / 3.1.4 Exercises for Section 3.1 89 / 3.2 Finite Automata and Regular Expressions 90 / 3.2.1 From DFA's to Regular Expressions 91 / 3.2.2 Converting DFA's to Regular Expressions by Eliminating States 96 / 3.2.3 Converting Regular Expressions to Automata 101 / 3.2.4 Exercises for Section 3.2 106 / 3.3 Applications of Regular Expressions 108 / 3.3.1 Regular Expressions in UNIX 108 / 3.3.2 Lexical Analysis 109 / 3.3.3 Finding Patterns in Text 111 / 3.3.4 Exercises for Section 3.3 113 / 3.4 Algebraic Laws for Regular Expressions 114 / 3.4.1 Associativity and Commutativity 114 / 3.4.2 Identities and Annihilators 115 / 3.4.3 Distributive Laws 115 / 3.4.4 The Idempotent Law 116 / 3.4.5 Laws Involving Closures 117 / 3.4.6 Discovering Laws for Regular Expressions 117 / 3.4.7 The Test for a Regular-Expression Algebraic Law 119 / 3.4.8 Exercises for Section 3.4 120 / 3.5 Summary of Chapter 3 122 / 3.6 References for Chapter 3 122

4. Properties of Regular Languages 125

4.1 Proving Languages not to be Regular 126 / 4.1.1 The Pumping Lemma for Regular Languages 126 / 4.1.2 Applications of the Pumping Lemma 127 / 4.1.3 Exercises for Section 4.1 129 / 4.2 Closure Properties of Regular Languages 131 / 4.2.1 Closure of Regular Languages Under Boolean Operations 131 / 4.2.2 Reversal 137 / 4.2.3 Homomorphisms 139 / 4.2.4 Inverse Homomorphisms 140 / 4.2.5 Exercises for Section 4.2 145 / 4.3 Decision Properties of Regular Languages 149 / 4.3.1 Converting Among Representations 149 / 4.3.2 Testing Emptiness of Regular Languages 151 / 4.3.3 Testing Membership in a Regular Language 153 / 4.3.4 Exercises for Section 4.3 153 / 4.4 Equivalence and Minimization of Automata 154 / 4.4.1 Testing Equivalence of States 154 / 4.4.2 Testing Equivalence of Regular Languages 157 / 4.4.3 Minimization of DFA's 159 / 4.4.4 Why the Minimized DFA Can't Be Beaten 162 / 4.4.5 Exercises for Section 4.4 164 / 4.5 Summary of Chapter 4 165 / 4.6 References for Chapter 4 166

5. Context-Free Grammars and Languages 169

5.1 Context-Free Grammars 169 / 5.1.1 An Informal Example 170 / 5.1.2 Definition of Context-Free Grammars 171 / 5.1.3 Derivations Using a Grammar 173 / 5.1.4 Leftmost and Rightmost Derivations 175 / 5.1.5 The Language of a Grammar 177 / 5.1.6 Sentential Forms 178 / 5.1.7 Exercises for Section 5.1 179 / 5.2 Parse Trees 181 / 5.2.1 Constructing Parse Trees 181 / 5.2.2 The Yield of a Parse Tree 183 / 5.2.3 Inference, Derivations, and Parse Trees 184 / 5.2.4 From Inferences to Trees 185 / 5.2.5 From Trees to Derivations 187 / 5.2.6 From Derivations to Recursive Inferences 190 / 5.2.7 Exercises for Section 5.2 191 / 5.3 Applications of Context-Free Grammars 191 / 5.3.1 Parsers 192 / 5.3.2 The YACC Parser-Generator 194 / 5.3.3 Markup Languages 196 / 5.3.4 XML and Document-Type Definitions 198 / 5.3.5 Exercises for Section 5.3 204 / 5.4 Ambiguity in Grammars and Languages 205 / 5.4.1 Ambiguous Grammars 205 / 5.4.2 Removing Ambiguity From Grammars 207 / 5.4.3 Leftmost Derivations as a Way to Express Ambiguity 211 / 5.4.4 Inherent Ambiguity 212 / 5.4.5 Exercises for Section 5.4 214 / 5.5 Summary of Chapter 5 215 / 5.6 References for Chapter 5 216

6. Pushdown Automata 219

6.1 Definition of the Pushdown Automaton 219 / 6.1.1 Informal Introduction 219 / 6.1.2 The Formal Definition of Pushdown Automata 221 / 6.1.3 A Graphical Notation for PDA's 223 / 6.1.4 Instantaneous Descriptions of a PDA 224 / 6.1.5 Exercises for Section 6.1 228 / 6.2 The Languages of a PDA 229 / 6.2.1 Acceptance by Final State 229 / 6.2.2 Acceptance by Empty Stack 230 / 6.2.3 From Empty Stack to Final State 231 / 6.2.4 From Final State to Empty Stack 234 / 6.2.5 Exercises for Section 6.2 236 / 6.3 Equivalence of PDA's and CFG's 237 / 6.3.1 From Grammars to Pushdown Automata 237 / 6.3.2 From PDA's to Grammars 241 / 6.3.3 Exercises for Section 6.3 245 / 6.4 Deterministic Pushdown Automata 246 / 6.4.1 Definition of a Deterministic PDA 247 / 6.4.2 Regular Languages and Deterministic PDA's 247 / 6.4.3 DPDA's and Context-Free Languages 249 / 6.4.4 DPDA's and Ambiguous Grammars 249 / 6.4.5 Exercises for Section 6.4 251 / 6.5 Summary of Chapter 6 252 / 6.6 References for Chapter 6 253

7. Properties of Context-Free Languages 255

7.1 Normal Forms for Context-Free Grammars 255 / 7.1.1 Eliminating Useless Symbols 256 / 7.1.2 Computing the Generating and Reachable Symbols 258 / 7.1.3 Eliminating epsilon-Productions 259 / 7.1.4 Eliminating Unit Productions 262 / 7.1.5 Chomsky Normal Form 266 / 7.1.6 Exercises for Section 7.1 269 / 7.2 The Pumping Lemma for Context-Free Languages 274 / 7.2.1 The Size of Parse Trees 274 / 7.2.2 Statement of the Pumping Lemma 275 / 7.2.3 Applications of the Pumping Lemma for CFL's 276 / 7.2.4 Exercises for Section 7.2 280 / 7.3 Closure Properties of Context-Free Languages 281 / 7.3.1 Substitutions 282 / 7.3.2 Applications of the Substitution Theorem 284 / 7.3.3 Reversal 285 / 7.3.4 Intersection With a Regular Language 285 / 7.3.5 Inverse Homomorphism 289 / 7.3.6 Exercises for Section 7.3 291 / 7.4 Decision Properties of CFL's 293 / 7.4.1 Complexity of Converting Among CFG's and PDA's 294 / 7.4.2 Running Time of Conversion to Chomsky Normal Form 295 / 7.4.3 Testing Emptiness of CFL's 296 / 7.4.4 Testing Membership in a CFL 298 / 7.4.5 Preview of Undecidable CFL Problems 302 / 7.4.6 Exercises for Section 7.4 302 / 7.5 Summary of Chapter 7 303 / 7.6 References for Chapter 7 304

8. Introduction to Turing Machines 307

8.1 Problems That Computers Cannot Solve 307 / 8.1.1 Programs that Print "Hello, World" 308 / 8.1.2 The Hypothetical "Hello, World" Tester 310 / 8.1.3 Reducing One Problem to Another 313 / 8.1.4 Exercises for Section 8.1 316 / 8.2 The Turing Machine 316 / 8.2.1 The Quest to Decide All Mathematical Questions 317 / 8.2.2 Notation for the Turing Machine 318 / 8.2.3 Instantaneous Descriptions for Turing Machines 320 / 8.2.4 Transition Diagrams for Turing Machines 323 / 8.2.5 The Language of a Turing Machine 326 / 8.2.6 Turing Machines and Halting 327 / 8.2.7 Exercises for Section 8.2 328 / 8.3 Programming Techniques for Turing Machines 329 / 8.3.1 Storage in the State 330 / 8.3.2 Multiple Tracks 331 / 8.3.3 Subroutines 333 / 8.3.4 Exercises for Section 8.3 334 / 8.4 Extensions to the Basic Turing Machine 336 / 8.4.1 Multitape Turing Machines 336 / 8.4.2 Equivalence of One-Tape and Multitape TM's 337 / 8.4.3 Running Time and the Many-Tapes-to-One Construction 339 / 8.4.4 Nondeterministic Turing Machines 340 / 8.4.5 Exercises for Section 8.4 342 / 8.5 Restricted Turing Machines 345 / 8.5.1 Turing Machines With Semi-infinite Tapes 345 / 8.5.2 Multistack Machines 348 / 8.5.3 Counter Machines 351 / 8.5.4 The Power of Counter Machines 352 / 8.5.5 Exercises for Section 8.5 354 / 8.6 Turing Machines and Computers 355 / 8.6.1 Simulating a Turing Machine by Computer 355 / 8.6.2 Simulating a Computer by a Turing Machine 356 / 8.6.3 Comparing the Running Times of Computers and Turing Machines 361 / 8.7 Summary of Chapter 8 363 / 8.8 References for Chapter 8 365

9. Undecidability 367

9.1 A Language That Is Not Recursively Enumerable 368 / 9.1.1 Enumerating the Binary Strings 369 / 9.1.2 Codes for Turing Machines 369 / 9.1.3 The Diagonalization Language 370 / 9.1.4 Proof that L_d is not Recursively Enumerable 372 / 9.1.5 Exercises for Section 9.1 372 / 9.2 An Undecidable Problem That is RE 373 / 9.2.1 Recursive Languages 373 / 9.2.2 Complements of Recursive and RE languages 374 / 9.2.3 The Universal Language 377 / 9.2.4 Undecidability of the Universal Language 379 / 9.2.5 Exercises for Section 9.2 381 / 9.3 Undecidable Problems About Turing Machines 383

/ 9.3.1 Reductions 383 / 9.3.2 Turing Machines That Accept the Empty Language 384 / 9.3.3 Rice's Theorem and Properties of the RE Languages 387 / 9.3.4 Problems about Turing-Machine Specifications 390 / 9.3.5 Exercises for Section 9.3 390 / 9.4 Post's Correspondence Problem 392 / 9.4.1 Definition of Post's Correspondence Problem 392 / 9.4.2 The "Modified" PCP 394 / 9.4.3 Completion of the Proof of PCP Undecidability 397 / 9.4.4 Exercises for Section 9.4 403 / 9.5 Other Undecidable Problems 403 / 9.5.1 Problems About Programs 403 / 9.5.2 Undecidability of Ambiguity for CFG's 404 / 9.5.3 The Complement of a List Language 406 / 9.5.4 Exercises for Section 9.5 409 / 9.6 Summary of Chapter 9 410 / 9.7 References for Chapter 9 411

10. Intractable Problems 419

10.1 The Classes P and NP 414 / 10.1.1 Problems Solvable in Polynomial Time 414 / 10.1.2 An Example: Kruskal's Algorithm 414 / 10.1.3 Nondeterministic Polynomial Time 419 / 10.1.4 An NP Example: The Traveling Salesman Problem 419 / 10.1.5 Polynomial-Time Reductions 421 / 10.1.6 NP-Complete Problems 422 / 10.1.7 Exercises for Section 10.1 423 / 10.2 An NP-Complete Problem 426 / 10.2.1 The Satisfiability Problem 426 / 10.2.2 Representing SAT Instances 427 / 10.2.3 NP-Completeness of the SAT Problem 428 / 10.2.4 Exercises for Section 10.2 434 / 10.3 A Restricted Satisfiability Problem 435 / 10.3.1 Normal Forms for Boolean Expressions 436 / 10.3.2 Converting Expressions to CNF 437 / 10.3.3 NP-Completeness of CSAT 440 / 10.3.4 NP-Completeness of 3SAT 445 / 10.3.5 Exercises for Section 10.3 446 / 10.4 Additional NP-Complete Problems 447 / 10.4.1 Describing NP-complete Problems 447 / 10.4.2 The Problem of Independent Sets 448 / 10.4.3 The Node-Cover Problem 452 / 10.4.4 The Directed Hamilton-Circuit Problem 453 / 10.4.5 Undirected Hamilton Circuits and the TSP 460 / 10.4.6 Summary of NP-Complete Problems 461 / 10.4.7 Exercises for Section 10.4 462 / 10.5 Summary of Chapter 10 466 / 10.6 References for Chapter 10 467

11. Additional Classes of Problems 469

11.1 Complements of Languages in NP 470 / 11.1.1 The Class of Languages Co-NP 470 / 11.1.2 NP-Complete Problems and Co-NP 471 / 11.1.3 Exercises for Section 11.1 472 / 11.2 Problems Solvable in Polynomial Space 473 / 11.2.1 Polynomial-Space Turing Machines 473 / 11.2.2 Relationship of PS and NPS to Previously Defined Classes 474 / 11.2.3 Deterministic and Nondeterministic Polynomial Space 476 / 11.3 A Problem That Is Complete for PS 478 / 11.3.1 PS-Completeness 478 / 11.3.2 Quantified Boolean Formulas 479 / 11.3.3 Evaluating Quantified Boolean Formulas 480 / 11.3.4 PS-Completeness of the QBF Problem 482 / 11.3.5 Exercises for Section 11.3 487 / 11.4 Language Classes Based on Randomization 487 . 11.4.1 Quicksort: an Example of a Randomized Algorithm 488 . 11.4.2 A Turing-Machine Model Using Randomization 489 11.4.3 The Language of a Randomized Turing Machine 490 / 11.4.4 The Class RP 493 / 11.4.5 Recognizing Languages in RP 495 / 11.4.6 The Class ZPP 495 / 11.4.7 Relationship Between RP and ZPP 496 / 11.4.8 Relationships to the Classes P and NP 497 / 11.5 The Complexity of Primality Testing 499 / 11.5.1 The Importance of Testing Primality 499 / 11.5.2 Introduction to Modular Arithmetic 501 / 11.5.3 The Complexity of Modular-Arithmetic Computations 503 / 11.5.4 Random-Polynomial Primality Testing 504 / 11.5.5 Nondeterministic Primality Tests 505 / 11.5.6 Exercises for Section 11.5 508 / 11.6 Summary of Chapter 11 508 / 11.7 References for Chapter 11 510