

Regular Expression Types for XML

Haruo Hosoya

Benjamin C. Pierce

Jérôme Vouillon

University of Pennsylvania

(with Peter Buneman and Phil Wadler)

XML

- ◆ a standard format for exchange of structured data
- ◆ huge corporate and* academic buy-in

*hence?

XML “Types”

A major selling point for XML is that documents come with descriptions of their structure

- ◆ DTDs (old standard)
- ◆ XML-Schema (proposed new standard)
- ◆ Relax, etc. (other similar proposals)
- ◆ Growing literature on refinements and extensions (e.g. keys)

But...

XML “Types”

A major selling point for XML is that documents come with descriptions of their structure

- ◆ DTDs (old standard)
- ◆ XML-Schema (proposed new standard)
- ◆ Relax, etc. (other similar proposals)
- ◆ Growing literature on refinements and extensions (e.g. keys)

But... these “types” are not actually used for typechecking **programs!**

The standard approach

A typical XML processing program...

- ◆ reads an XML tree
- ◆ verifies (during parsing) that it matches some DTD (or Schema)
- ◆ ignores the DTD from this point on, treating the input as a generic labeled tree (e.g., a DOM structure)
- ◆ constructs new trees that may or may not conform to an intended DTD

The standard approach

A typical XML processing program...

- ◆ reads an XML tree
- ◆ verifies (during parsing) that it matches some DTD (or Schema)
- ◆ ignores the DTD from this point on, treating the input as a generic labeled tree (e.g., a DOM structure)
- ◆ constructs new trees that may or may not conform to an intended DTD

We can do better!

A Better Approach

There have been a number of proposals for using the existing features of popular statically typed languages to represent XML structures.

E.g.,

- ◆ Data-binding for Java [Sun]
- ◆ HaXML [Wallace and Runciman, ICFP99]

This is good!

Static type safety \longrightarrow increased robustness

- ◆ A well-typed program cannot fail because an intermediate stage generates a structure that another stage is not expecting
- ◆ A well-typed program will always produce output that conforms to the expected DTD

But...

Union has lost its flexibility

The problem with such embeddings is that they introduce **spurious structure**, mapping the **non-disjoint union** operator of DTDs to the **disjoint union** of the host-language type system.

XML	Java/C#	ML
$S \mid T$	<pre>class SorT class S extends SorT class T extends SorT</pre>	<pre>datatype SorT = left of S right of T</pre>

Consequences:

- ◆ An element of S cannot simply be viewed as an element of $S|T$ — an explicit coercion is required
- ◆ If S and T have fields in common, there is no way to access them directly in an element of $S|T$ — we must first perform a tag-test

Consequences:

- ◆ An element of S cannot simply be viewed as an element of $S|T$ — an explicit coercion is required
- ◆ If S and T have fields in common, there is no way to access them directly in an element of $S|T$ — we must first perform a tag-test

We can do even better!

Regular expression types

Our proposal: Use DTDs* themselves as types.

*suitably cleaned up and generalized.

XDuce

- ◆ An experimental programming language for writing “recursive tree transducers” for XML structures
- ◆ Static typechecking based on regular expression types
- ◆ Flexible subtyping relation supporting common patterns of evolution and data integration
- ◆ Generalization of ML-style pattern matching with regular expression patterns
- ◆ Similar “feel” to XSLT

Outline

- ◆ Introduction ✓
- ◆ Regular expression types
- ◆ Regular expression patterns
- ◆ Current work: Xtatic

Regular Expression Types

An Example (in XDuce notation)

```
type Addrbook = addrbook[Person*]
type Person   = person[Name,Email*,Tel?]
type Name     = name[String]
type Email    = email[String]
type Tel      = tel[String]

val mybook = addrbook[person[name["Haruo Hosoya"],
                             email["hahosoya@upenn"],
                             email["haruo@u-tokyo"]],
                     person[name["Jerome Vouillon"],
                             email["vouillon@upenn"],
                             tel["215-123-4567"]]]
```


Semantics of Types

Each type denotes a **set of sequences**.

E.g.:

$$\llbracket \text{Email*} \rrbracket = \{ \begin{array}{l} () \\ \text{email}[\dots] \\ \text{email}[\dots], \text{email}[\dots] \\ \text{etc.} \end{array} \}$$

Comma denotes concatenation of sequences:

$$\begin{aligned} \llbracket \text{Name}, \text{Email}^*, \text{Tel} \rrbracket &= \{ \text{name}[\dots], \text{tel}[\dots] \\ &\quad \text{name}[\dots], \text{email}[\dots], \text{tel}[\dots] \\ &\quad \text{name}[\dots], \text{email}[\dots], \text{email}[\dots], \text{tel}[\dots] \\ &\quad \text{etc.} \} \end{aligned}$$

Subtyping (Examples)

? means “optional”

Name <: Name, Email?

Name, Email <: Name, Email?

* means “zero or more”

Email, Email, Email <: Email*

Email, Email* <: Email*

| means “or”

Email <: Email|Tel

Tel <: Email|Tel

“forget ordering” subtyping

Email*, Tel* <: (Email|Tel)*

Subtyping (Examples)

? means “optional”

Name <: Name, Email?

Name, Email <: Name, Email?

* means “zero or more”

Email, Email, Email <: Email*

Email, Email* <: Email*

| means “or” Note that | is a non-disjoint union!

Email <: Email|Tel

Tel <: Email|Tel

“forget ordering” subtyping

Email*, Tel* <: (Email|Tel)*

Subtyping Recursive Types

Recursive types describe arbitrarily deep tree structures:

```
type BinTree = Leaf[String]  
             | node[BinTree, BinTree]
```

```
type UBinTree = Leaf[String]  
              | node[BinTree, UBinTree?]
```

```
BinTree <: UBinTree
```

Example: Data Integration

Challenge: Suppose we are given two data sources that have almost but not exactly the same types (e.g., two databases that have evolved separately from a common origin)

Task:

1. Integrate (combine) the two sources
2. Write a simple program that performs some operation on the common fields

Original data sources:

$$\text{src1} \in (\text{Name}, \text{Email})^*$$
$$\text{src2} \in (\text{Name}, \text{Tel})^*$$

Integration = concatenation

$$\text{src1}, \text{src2} \in (\text{Name}, \text{Email})^*, (\text{Name}, \text{Tel})^*$$

Next, we want to do something with the common part (i.e., the Name fields).

Since there are two *'s, we need to use two separate loops, right?

No! Just use subtyping between regular expression types to factor out the common part.

Step 1: Use “forget ordering” subtyping to merge *’s:

$$\begin{aligned} & (\text{Name}, \text{Email})^*, (\text{Name}, \text{Tel})^* \\ \prec: & ((\text{Name}, \text{Email}) \mid (\text{Name}, \text{Tel}))^* \end{aligned}$$

Step 2: Distribute union over concatenation to factor out Name:

$$\begin{aligned} & ((\text{Name}, \text{Email}) \mid (\text{Name}, \text{Tel}))^* \\ \prec: & (\text{Name}, (\text{Email} \mid \text{Tel}))^* \end{aligned}$$

(Distributivity is the most important advantage of real union types over disjoint unions a la ML or Java.)

Step 3: Write a single loop to extract the Names.

```
fun names : (Name, (Email|Tel))* -> Name* =  
  name[n], (Email,Tel), rest  
    ->  n, names(rest)  
| ()  
  ->  ()
```

Subtyping

[insert additional slides]

Efficiency of Subtyping

The algorithm we've sketched gives us a decision procedure for the full “semantic” subtyping relation.

How fast is it?

Efficiency of Subtyping

The algorithm we've sketched gives us a decision procedure for the full “semantic” subtyping relation.

How fast is it?

In the worst case, not very fast: subtyping between regular expression types is essentially inclusion-testing for regular tree languages, which is exptime-complete.

However, we can go a long way with heuristics...

Subtyping Heuristics

- ◆ low-level tricks (hash-consing and memoization)
- ◆ special treatment of empty types
- ◆ a variety of “set-theoretic” optimizations to avoid using the general union rule

cf. [Hosoya, Pierce, & Vouillon, ICFP 2000]

Some Preliminary Measurements

	Bookmark	Html2Latex
Size (code)	310 lines	312 lines
Size (types)	1242 lines	1217 lines
# of subtype checks	61	123
Type checking time	0.48 secs	0.88 secs

on a 300Mhz Ultrasparc

Regular Expression Patterns

Example

Recall the address book types from earlier:

```
type AddrBook = Person*  
type Person   = person[Name,Email*,Tel?]  
type Name     = name[String]  
type Email    = email[String]  
type Tel      = tel[String]
```

A simple pattern match:

```
match p with
  person[name[n], Email*, tel[t]]
    -> (* do some stuff involving n and t *)
| person[p]
  -> (* do other stuff *)
```

Note how the **type** `Email*` is used in the first pattern to match a **variable-length** sequence of email nodes.

A complete XDuce function

```
fun tels : Person* -> (Name,Tel)* =  
  person[name[n], Email*, tel[t]], rest  
    -> name[n], tel[t], tels(rest)  
| person[p], rest -> tels(rest)  
| ()              -> ()
```

- ◆ header is explicitly annotated with both argument and result types
- ◆ other type annotations (in particular, on pattern variables) are omitted

A More Interesting Example

Using regular expression patterns, we can extract the subcomponents of an HTML table with a single `match`...

```
match t with
  table[cap as Caption?,
        col as (Col*|Colgroup*),
        hd as Thead,
        ft as Tfoot?,
        bd as (Tbody+|Tr+)]
-> ...
```

Issues

- ◆ Type inference for pattern variables
 - ◆ complicated by...
 - recursive types and patterns
 - complex control flows arising from “first-match” policy

[Cf. Hosoya & Pierce , POPL 2001]

- ◆ Optimization of patterns to avoid type membership testing at run time whenever possible

Status

XDuce Status

- ◆ Prototype implementation
 - ◆ Interpreted runtime
 - ◆ Fairly fast typechecker
- ◆ Several small (but nontrivial) applications
 - ◆ Web browser bookmark formatter
 - ◆ Html2latex translator
 - ◆ Tree-diff [Chawathe, VLDB '99]
 - ◆ Prototype XML-Schema validator [Renneberg '00]

Ongoing work: the Xtatic project

A new language design and implementation based on XDuce

- ◆ Emphasis on inter-operability with a mainstream “host language” (Java or C#)
- ◆ Compiling to a common runtime system (JVM or MS .net common runtime)

Xtatic: New design issues

- ◆ unordered record types (based on “interleave” operator for regular trees)
- ◆ integration of “horizontal” (XSLT-, ML-, or XDuce-style) and “vertical” (DB query language) patterns
- ◆ object types (and/or higher-order functions)
- ◆ polymorphism?

Finishing up...

Related work

Other XML languages with static type systems:

- ◆ XML [Meijer & Shields, 2000]
- ◆ YAT [Cluet & Simeon, WebDB1998]

Similar aims. Limited support for subtyping (e.g., no distributivity laws).

Related work

Typechecking for XML query languages [Milo, Suciu, and Vianu, PODS 2000, Papakonstantinou & Vianu, PODS 2000]

- ◆ Complexity studies of typechecking problems
- ◆ Similar notions of types and subtyping
- ◆ Applications / implementation issues not considered

Related work

“Union Types for Semi-structured Data” [Buneman & Pierce, DBPL 99]

- ◆ Powerful subtyping (including similar distributive laws)
- ◆ Unordered records (rather than XML's ordered sequences)
- ◆ No recursive types

Related work

XML Algebra [Fernandez, Simeon, and Wadler, 2000]

- ◆ Proposed core for XML query languages
- ◆ Draft W3C standard
- ◆ Type system based directly on XDuce

Related work

“Set inclusion constraints” [Aiken&Murphy, FPCA91, Aiken&Wimmers, LICS92]

- ◆ Closely related algorithmic problem
- ◆ See our ICFP 2000 paper for a detailed comparison

If you want to play...

The XDuce home page

www.cis.upenn.edu/~hahosoya/xduce

contains papers, talks, and our prototype implementation.