Linear parametric model checking of timed automata☆

Thomas Hune[a], Judi Romijn[b,*], Mariëlle Stoelinga[b],
Frits Vaandrager[b]

[a] *Basic Research in Computer Science, University of Århus, Århus, Denmark*[1]
[b] *Computing Science Institute, University of Nijmegen, Nijmegen, The Netherlands*

Abstract

We present an extension of the model checker UPPAAL, capable of synthesizing linear parameter constraints for the correctness of parametric timed automata. A symbolic representation of the (parametric) state space in terms of parametric difference bound matrices is shown to be correct. A second contribution of this paper is the identification of a subclass of parametric timed automata (L/U automata), for which the emptiness problem is decidable, contrary to the full class where it is known to be undecidable. Also, we present a number of results that reduce the verification effort for L/U automata in certain cases. We illustrate our approach by deriving linear parameter constraints for a number of well-known case studies from the literature (exhibiting a flaw in a published paper). © 2002 Published by Elsevier Science Inc.

Keywords: Model checking; Parameter synthesis; Real-time; Timed automata; Verification

1. Introduction

Model checking is emerging as a practical tool for automated debugging of complex reactive systems such as embedded controllers and network protocols. In model checking, a high-level description of a system is compared against a logical correctness requirement to discover inconsistencies. The first techniques for model checking did not admit an explicit modeling of time, and are thus unsuitable for analysis of real-time systems whose correctness depends on relative magnitudes of different delays. Consequently, Alur and

[1] Basic Research in Computer Science (http://www.brics.dk) funded by the Danish National Research Foundation.

* Corresponding author. Address: P.O. Box 513, 5600 MB Eindhoven, The Netherlands.

E-mail addresses: baris@brics.dk (T. Hune), judi@cs.kun.nl, jromijn@win.tue.nl (J. Romijn), marielle@cs.kun.nl (M. Stoelinga), fvaan@cs.kun.nl (F. Vaandrager).

Dill [3] proposed timed automata as a formal notion to model the behavior of real-time systems. Timed automata are state-transition diagrams annotated with timing constraints using finitely many real-valued clock variables. During the last decade, there has been enormous progress in the area of timed model checking. We refer to [7,10,21,28] for overviews of the underlying theory and references to applications. Timed automata tools such as UPPAAL [21], KRONOS [8], and PMC [22] are now routinely used for industrial case studies.

A disadvantage of the traditional approaches is, however, that they can only be used to verify concrete timing properties: one has to provide the values of all timing parameters that occur in the system. Typical examples of such parameters are upper and lower bounds on computation times, message delays and timeouts. For practical purposes, one is often interested in deriving the (symbolic) constraints on the parameters that ensure correctness. The process of manually finding and proving such results is very time consuming and error prone (we have discovered minor errors in the examples we have been looking at). Therefore tool support for deriving the constraints *automatically* is very important.

In this paper, we study a parametric extension of timed automata, called *parametric timed automata* (PTAs), and present an extension to PTAs of the (forward) state-space exploration algorithm for timed automata. We show the theoretical correctness of our approach, and its feasibility by application to three non-trivial case studies. For this purpose, we have implemented a prototype extension of UPPAAL, an efficient real-time model checking tool [21]. The algorithm we propose and have implemented fundamentally relies on parametric difference bound matrices (PDBMs) and operations on these. PDBMs constitute a data type that extends the difference bound matrices (DBMs [13]) in a natural way. The latter are used for recording clock differences when model checking (non-parametric) timed automata. PDBMs are basically DBMs, where the matrix entries are parameter expressions rather than constants. Our algorithm is a semi-decision algorithm which will not terminate in all cases. In [5], the problem of synthesizing values for parameters such that a property is satisfied was shown to be undecidable, so this is the best we can hope for.

A second contribution of this paper is the identification of a subclass of parametric timed automata, called *lower bound/upper bound* (L/U) *automata*, which appears to be sufficiently expressive from a practical perspective, while it also has nice theoretical properties. Most importantly, we show that the emptiness problem, in [5] shown to be undecidable for parametric timed automata, is decidable for L/U automata. We also establish a number of results which allow one to reduce the number of parameters when tackling specific verification questions for L/U automata. The application of these lemmas has already reduced the verification effort drastically in several of our experiments.

Related work. There are currently several other tools available that can do parametric model checking, namely LPMC, HYTECH and TReX.

LPMC [22] is a parametric extension of the timed model checker PMC [9]. The model checking algorithm implemented in LPMC differs from ours: it represents the state space of a system as an unstructured set of constraints, whereas we use PDBMs. Moreover, LPMC implements a partition refinement technique, whereas we use forward reachability. Other differences with our approach are that LPMC also allows for the comparison of non-clock variables to parameter constraints and for more general specification properties (full TCTL with fairness assumptions).

The model checker HYTECH [15] is a tool for linear hybrid automata. These are more general than parametric timed automata, since they allow the modeling of continuous behavior via linear differential equations. The HYTECH implementation uses polyhedra as

its basic data type. It can explore the state space by using either forward reachability, as we do, or partition refinement, as in LPMC. The tool has been applied successfully to relatively small examples such as a railway gate controller. Experience so far has shown that HYTECH cannot cope with larger examples, such as the ones considered in this paper, see the results in [12].

The tool TReX [1,2] is currently the only one that can deal with non-linear parameter constraints. Moreover, TReX has a clever method for guessing the effect of control loops in a model, based on widening principles, which increases chances of termination. Independently, Annichini et al. [1] developed the same data structure as we did (PDBMs) and implemented some similar operations on these. However, the underlying theory was not worked out in the research literature. Hence, we believe that our contribution over [1] consists of the following. Our work presents an extensive elaboration of the theory behind our implementation. In particular, we present a correctness proof of the model checking algorithm we implemented. That is, we prove that the symbolic semantics of a PTA in terms of PDBMs is equivalent to its concrete semantics in terms of single states and transitions. These proofs rely on a number of non-trivial generalizations of results for DBMs.

Each of the tools above has been applied to the IEEE 1394 Root Contention Protocol [9,12]. We refer the reader to [26] for a comparison of the results. An important conclusion was that each of the verifications has it own merits, where our approach was the fastest.

Overview. The rest of this paper is organized as follows. Section 2 introduces the notion of parametric timed automata. Section 3 gives the symbolic semantics in terms of PDBMs and is the basis for the model checking algorithm presented in Section 3.5. In Section 4, we introduce the class of L/U automata. Section 5 reports on several experiments with our tool. Finally, Section 6 presents some conclusions.

2. Parametric timed automata

Parametric timed automata were first defined in [5]. They generalize the timed automata of [3]. The definition of parametric timed automata that we present in this section is very similar to the definition in [5], except that progress is ensured via location invariants rather than via accepting states. This difference is not essential.

2.1. Parameters and constraints

Throughout this paper, we assume a fixed set of *parameters* $P = \{p_1, \ldots, p_n\}$.

Definition 2.1 (*Constraints*). A *linear expression* e is either an expression of the form $t_1 p_1 + \cdots + t_n p_n + t_0$, where $t_0, \ldots, t_n \in \mathbb{Z}$, or ∞. We write E to denote the set of all linear expressions. A *constraint* is an inequality of the form $e \sim e'$, with e, e' linear expressions and $\sim \in \{<, \leqslant, >, \geqslant\}$. The *negation* of constraint c, denoted $\neg c$, is obtained by replacing relation symbols $<, \leqslant, >, \geqslant$ by $\geqslant, >, \leqslant, <$, respectively. A (*parameter*) *valuation* is a function $v : P \to \mathbb{R}^{\geqslant 0}$ assigning a non-negative real value to each parameter. There is a one-to-one correspondence between valuations and points in $(\mathbb{R}^{\geqslant 0})^n$. In fact we often identify a valuation v with the point $(v(p_1), \ldots, v(p_n)) \in (\mathbb{R}^{\geqslant 0})^n$.

If e is a linear expression and v is a valuation, then $e[v]$ denotes the expression obtained by replacing each parameter p in e with $v(p)$. Likewise, we define $c[v]$ for c a constraint.

Valuation v *satisfies* constraint c, denoted $v \models c$, if $c[v]$ evaluates to true. The *semantics* of a constraint c, denoted $[\![c]\!]$, is the set of valuations that satisfy c. A finite set of constraints C is called a *constraint set*. A valuation *satisfies* a constraint set if it satisfies each constraint in the set. The *semantics* of a constraint set C is given by $[\![C]\!] := \bigcap_{c \in C} [\![c]\!]$. We say that C is *satisfiable* if $[\![C]\!]$ is non-empty.

Constraint c *covers* constraint set C, denoted $C \models c$, iff $[\![C]\!] \subseteq [\![c]\!]$. Constraint set C is *split by* constraint c iff neither $C \models c$ nor $C \models \neg c$.

During the analysis questions arise of the kind: given a constraint set C and a constraint c, does c hold, i.e., does constraint c cover C? There are three possible answers to this, *yes*, *no*, and *split*. A split occurs when c holds for some valuations in the semantics of C and $\neg c$ holds for some other valuations. Here, we will not discuss in detail methods for answering such questions: in the remainder of this paper we just assume the presence of the following "oracle" function.

Definition 2.2 (*Oracle*).

$$\mathcal{O}(c, C) = \begin{cases} \text{yes} & \text{if } C \models c, \\ \text{no} & \text{if } C \models \neg c, \\ \text{split} & \text{otherwise.} \end{cases}$$

The oracle function can be computed in polynomial time using linear programming (LP) solvers. Suppose we want to compute $\mathcal{O}(c, C)$, where c takes the form $e \leqslant e'$. Then we first maximize the linear function $e' - e$ subject to the set C of linear inequalities. This is a linear programming problem. If the outcome is negative, then $\mathcal{O}(c, C) = $ no. Otherwise we maximize $e - e'$ subject to C. If the outcome is less than or equal to 0, then $\mathcal{O}(c, C) = $ yes. Otherwise $\mathcal{O}(c, C) = $ split. In our implementation we use an LP solver that was kindly provided to us by the authors of [9], who built it for their model checking tool LPMC. This LP solver is geared to perform well on small, simple sets of constraints rather than large, complicated ones.

Observe that using the oracle, we can easily decide semantic inclusion between constraint sets: $[\![C]\!] \subseteq [\![C']\!]$ iff $\forall c' \in C' : \mathcal{O}(c', C) = $ yes.

2.2. Parametric timed automata

Throughout this paper, we assume a fixed set of clocks $X = \{x_0, \ldots, x_m\}$ and a fixed set of actions $A = \{a_1, \ldots, a_k\}$. The special clock x_0, which is called the *zero clock*, always has the value 0 (and hence does not increase with time).

A *simple guard* is an expression f of the form $x_i - x_j \prec e$, where x_i, x_j are clocks, $\prec \in \{<, \leqslant\}$, and e is a linear expression. We say that f is *proper* if $i \neq j$. We define a *guard* to be a (finite) conjunction of simple guards. We let g range over guards and write G to denote the set of guards. A *clock valuation* is a function $w : X \to \mathbb{R}^{\geqslant 0}$ assigning a non-negative real value to each clock such that $w(x_0) = 0$. We will identify a clock valuation w with the point $(w(x_0), \ldots, w(x_m)) \in (\mathbb{R}^{\geqslant 0})^{m+1}$. Let g be a guard, v a parameter valuation, and w a clock valuation. Then $g[v, w]$ denotes the expression obtained by replacing each parameter p with $v(p)$, and each clock x with $w(x)$. A pair (v, w) of a parameter valuation and a clock valuation *satisfies* a guard g, denoted $(v, w) \models g$, if $g[v, w]$ evaluates to true. The *semantics* of a guard g, denoted $[\![g]\!]$, is the set of pairs

Fig. 1. A parametric timed automaton.

(v, w) such that $(v, w) \models g$. Given a parameter valuation v, we write $[\![g]\!]_v$ for the set of clock valuations $\{w \mid (v, w) \models g\}$.

A *reset* is an expression of the form, $x_i := b$, where $i \neq 0$ and $b \in \mathbb{N}$. A *reset set* is a set of resets containing at most one reset for each clock. The set of reset sets is denoted by R.

We now define an extension of timed automata [4,28] called parametric timed automata. Similar models have been presented in [1,5,9].

Definition 2.3 (*PTA*). A *parametric timed automaton* (PTA) over set of clocks X, set of actions A, and set of parameters P, is a quadruple $\mathscr{A} = (Q, q_0, \rightarrow, I)$, where Q is a finite set of *locations*, $q_0 \in Q$ is the *initial location*, $\rightarrow \subseteq Q \times A \times G \times R \times Q$ is a finite *transition relation*, and function $I : Q \rightarrow G$ assigns an *invariant* to each location. We abbreviate a $(q, a, g, r, q') \in \rightarrow$ consisting of a source location q, an action a, a guard g, a reset set r, and a target location q' as $q \xrightarrow{a,g,r} q'$. For a simple guard $x_i - x_j \prec e$ to be used in an invariant it must be the case that $j = 0$, that is, the simple guard represents an upper bound on a clock.[2]

Example 2.4. A PTA with clocks x, y and parameters p, q can be seen in Fig. 1. The initial location is S_0 and has invariant $x \leqslant p$. There is a transition from the initial location to S_1, which has guard $y \geqslant q$ and reset set $\{x := 0\}$. There are no actions on the transitions. The transition from S_0 to S_1 can only become enabled if $p \leqslant q$, otherwise the system will end up in a deadlock.

To define the semantics of PTAs, we require two auxiliary operations on clock valuations. For clock valuation w and non-negative real number d, $w + d$ is the clock valuation that adds to each clock (except x_0) a delay d. For clock valuation w and reset set r, $w[r]$ is the clock valuation that resets clocks according to r:

$$(w + d)(x) = \begin{cases} 0 & \text{if } x = x_0, \\ w(x) + d & \text{otherwise,} \end{cases} \qquad (w[r])(x) = \begin{cases} b & \text{if } x := b \in r, \\ w(x) & \text{otherwise.} \end{cases}$$

Definition 2.5 (*LTS*). A *labeled transition system (LTS)* over a set of symbols Σ is a triple $\mathscr{L} = (S, S_0, \rightarrow)$, with S a set of *states*, $S_0 \subseteq S$ a set of *initial states*, and $\rightarrow \subseteq S \times \Sigma \times S$ a *transition relation*. We write $s \xrightarrow{a} s'$ for $(s, a, s') \in \rightarrow$. A *run* of $\mathscr{L}$ is a finite alternating sequence $s_0 a_1 s_1 a_2 \cdots s_n$ of states $s_i \in S$ and symbols $a_i \in \Sigma$ such that $s_0 \in S_0$ and, for all $i < n$, $s_i \xrightarrow{a_{i+1}} s_{i+1}$. A state is *reachable* if it is the last state of some run.

---

[2] There is no fundamental reason to impose this restriction on invariants; our whole theory can be developed without it. However, technically the restriction makes our lives a bit easier, see for instance Proposition 3.17. In practice the condition is (as far as we are aware) always met.

Definition 2.6 (*Concrete semantics*). Let $\mathscr{A} = (Q, q_0, \rightarrow, I)$ be a PTA and v be a parameter valuation. The *concrete semantics of $\mathscr{A}$ under v*, denoted $[\![\mathscr{A}]\!]_v$, is the LTS $(S, S_0, \rightarrow)$ over $A \cup \mathbb{R}^{\geqslant 0}$, where

$$S = \left\{ (q, w) \in Q \times (X \rightarrow \mathbb{R}^{\geqslant 0}) \mid w(x_0) = 0 \wedge (v, w) \models I(q) \right\},$$
$$S_0 = \{ (q, w) \in S \mid q = q_0 \wedge w = \lambda x.0 \},$$

and the transition predicate $\rightarrow$ is specified by the following two rules. For all (q, w), $(q', w') \in S$, $d \geqslant 0$ and $a \in A$,

- $(q, w) \xrightarrow{d} (q', w')$ if $q = q'$ and $w' = w + d$.
- $(q, w) \xrightarrow{a} (q', w')$ if $\exists g, r : q \xrightarrow{a,g,r} q'$ and $(v, w) \models g$ and $w' = w[r]$.

Note that the LTS $[\![\mathscr{A}]\!]_v$ has at most one initial state. It has no initial state if the invariant assigned to the initial location of $\mathscr{A}$ is unsatisfiable.

2.3. The parametric model checking problem

In its current version, UPPAAL is able to check for reachability properties, in particular whether certain combinations of locations and constraints on clock variables are reachable from the initial configuration. Our parametric extension of UPPAAL handles exactly the same properties. However, rather than just telling whether a property holds or not, our tool looks for constraints on the parameters which ensure that the property holds.

Definition 2.7 (*Properties*). Let $\mathscr{A} = (Q, q_0, \rightarrow, I)$ be a PTA. The sets of *system properties* and *state formulas* for $\mathscr{A}$ are defined by, respectively,

$$\psi ::= \forall \Box \phi \mid \exists \Diamond \phi, \quad \phi ::= x - y \prec b \mid q \mid \neg \phi \mid \phi \wedge \phi \mid \phi \vee \phi,$$

where $x, y \in X$, $b \in \mathbb{N}$ and $q \in Q$. Let $\mathscr{A}$ be a PTA, v a parameter valuation, s a state of $[\![\mathscr{A}]\!]_v$, and ϕ a state formula. We write $s \models_v \phi$ if ϕ holds in state s of $[\![\mathscr{A}]\!]_v$, we write $[\![\mathscr{A}]\!]_v \models \forall \Box \phi$ if ϕ holds in all reachable states of $[\![\mathscr{A}]\!]_v$, and we write $[\![\mathscr{A}]\!]_v \models \exists \Diamond \phi$ if ϕ holds for some reachable state of $[\![\mathscr{A}]\!]_v$.

The problem that we address in this paper can now be stated as follows:

> *Given a parametric timed automaton $\mathscr{A}$ and a system property ψ, compute the set of parameter valuations v for which $[\![\mathscr{A}]\!]_v \models \psi$.*

Remark 2.8. Timed automata [4,28] arise as a special case of PTAs for which the set P of parameters is empty. If $\mathscr{A}$ is a PTA and v is a parameter valuation, then the structure $\mathscr{A}[v]$ that is obtained by replacing all linear expressions e that occur in $\mathscr{A}$ by $e[v]$ is a timed automaton.[3] It is easy to see that in general $[\![\mathscr{A}]\!]_v = [\![\mathscr{A}[v]]\!]$. Since the reachability problem for timed automata is decidable [4], this implies that, for any $\mathscr{A}$, integer valued v and ψ, $[\![\mathscr{A}]\!]_v \models \psi$ is decidable.

---

[3] Strictly speaking, $\mathscr{A}[v]$ is only a timed automaton if v assigns an integer to each parameter.

FISCHER (P_i)
lock := 0
repeat
  while lock $\neq 0$ **do skip od**
  lock := i
  delay
until lock = i
critical section
lock := 0

Fig. 2. Fischer's mutual exclusion algorithm.

2.4. Example: Fischer's mutual exclusion algorithm

Fig. 3 shows a PTA model of Fischer's mutual exclusion algorithm [19]. The purpose of this algorithm is to guarantee mutually exclusive access to a critical section among n competing processes $P_1, P_2, \ldots, P_n$. The algorithm, where each process P_i (perpetually) runs the code of Fig. 2, uses a shared variable lock for communication between the processes.

The correctness of this algorithm crucially depends on the timing of the operations. The key idea is that any process P_i that sets lock := i is made to wait long enough before checking lock = i to ensure that any other process P_j that tested lock = 0, before P_i set lock to its index, has already set lock to its index j, when P_i finally checks lock = i.

Assume that read/write access to the global variable (in the operations lock = i and lock := 0) takes between *min_rw* and *max_rw* time units and assume that the delay operation (including the time needed for the assignment lock := i) takes between *min_delay* and *max_delay* time units. If we assume the basic constraints $0 \leqslant min\_rw < max\_rw \wedge 0 \leqslant min\_delay < max\_delay$, then mutual exclusion is guaranteed if and only if max_rw $\leqslant$ min_delay.

Now consider the PTA in Fig. 3, which is represented in UPPAAL syntax. (Several different models of this algorithm exist [5,6,18,23]; our model is closest to the one in [23].) It consists of four locations *start* (which is initial), *set*, *try_enter* and *cs*; four parameters, *min_rw*, *max_rw*, *min_delay* and *max_delay*; one clock x and a shared variable *lock*. By convention, x and *lock* are initially 0. Note that the process can remain in the locations *start* and *set* for at least *min_rw* and strictly less than *max_rw* time units. Similarly, the process can remain in *try_enter* for any time in the interval [*min_delay*, *max_delay*).

The shared variable, which is not a part of the definition of PTAs, is syntactic sugar which allows for an efficient encoding of the algorithm as a PTA. Also the notion of parallel composition for PTAs is standard, see for instance [21] for their definitions.


Fig. 3. A PTA model of Fischer's mutual exclusion algorithm.

3. Symbolic state-space exploration

Our aim is to use basically the same algorithm for parametric timed model checking as for timed model checking. We represent sets of states symbolically in a similar way and support the same operations used for timed model checking. In the non-parametric case, sets of states can be efficiently represented using matrices [13]. Similarly, in this paper we represent sets of states symbolically as (*constrained*) *parametric difference bound matrices*.

3.1. Parametric difference bound matrices

In the non-parametric case, a difference bound matrix is a $(m + 1) \times (m + 1)$ matrix whose entries are elements from $(\mathbb{Z} \cup \{\infty\}) \times \{0, 1\}$. An entry $(c, 1)$ for D_{ij} denotes a non-strict bound $x_i - x_j \leqslant c$, whereas an entry $(c, 0)$ denotes a strict bound $x_i - x_j < c$. In the parametric case, instead of using integers in the entries, we will use linear expressions over the parameters. Also, we find it convenient to view the matrix slightly more abstractly as a set of guards.

Definition 3.1 (*PDBM*). A *parametric difference bound matrix* (PDBM) is a set D which contains, for all $0 \leqslant i, j \leqslant m$, a simple guard D_{ij} of the form $x_i - x_j \prec_{ij} e_{ij}$. We require that, for all i, D_{ii} is of the form $x_i - x_i \leqslant 0$. Given a parameter valuation v, the *semantics* of D is given by $[\![D]\!]_v = [\![\bigwedge_{i,j} D_{ij}]\!]_v$. PDBM D is *satisfiable* for v if $[\![D]\!]_v$ is non-empty. If f is a guard of the form $x_i - x_j \prec e$ with $i \neq j$ (i.e., a proper guard), then $D[f]$ denotes the PDBM obtained from D by replacing D_{ij} by f. If i, j are indices, then D^{ij} denotes the pair $(e_{ij}, \prec_{ij})$; we call D^{ij} a *bound* of D. Clearly, a PDBM is fully determined by its bounds.

Definition 3.2 (*Constrained PDBM*). A *constrained PDBM* is a pair (C, D), where C is a constraint set and D is a PDBM. We require that $C \models p \geqslant 0$, for each p, and $C \models e_{0i} \leqslant 0$, for each i. The *semantics* of (C, D) is given by $[\![C, D]\!] = \{(v, w) \mid v \in [\![C]\!] \land w \in [\![D]\!]_v\}$. We call (C, D) *satisfiable* if $[\![C, D]\!]$ is non-empty.

Condition $C \models p \geqslant 0$ expresses that parameter p may only take non-negative values. The condition $C \models e_{0i} \leqslant 0$ ensures a non-negative lower bound on the value of clock x_i. Such a condition is required since clocks in a PTA only take non-negative values. A similar condition occurs in [28]. In the setting of [13] the condition of non-negative lower bounds is not needed since in this paper clocks (called timers) may take values in $\mathbb{R}$. In [1,7,10,20] the condition (or something similar) is needed but not mentioned.[4]

The PDBMs with the tightest possible bounds are called *canonical*. To formalize this notion, we define an addition operation on linear expressions by

$$(t_1 p_1 + \cdots + t_n p_n + t_0) + (t'_1 p_1 + \cdots + t'_n p_n + t'_0)$$
$$\stackrel{\Delta}{=} (t_1 + t'_1) p_1 + \cdots + (t_n + t'_n) p_n + (t_0 + t'_0).$$

---

[4] For instance, in [10] it is claimed on page 289: "If the clock zone is empty or unsatisfiable, there will be at least one negative entry in the main diagonal." This claim is incorrect. A counterexample is the canonical form of a DBM that contains as the only non-trivial guard $x_1 - x_0 \leqslant -1$.

Also, we view Boolean connectives as operations on relation symbols $\leq$ and $<$ by identifying $\leq$ with 1 and $<$ with 0. For instance, $(\leq \wedge \leq) = \leq$, $(\leq \wedge <) = <$, $\neg \leqslant = <$, and $(\leq \Rightarrow <) = <$.

Our definition of a canonical form of a constrained PDBM is essentially equivalent to the one for standard DBMs.

Definition 3.3 (*Canonical form*). A constrained PDBM (C, D) is *in canonical form* iff for all i, j, k, $C \models e_{ij} (\prec_{ij} \Rightarrow \prec_{ik} \wedge \prec_{kj}) e_{ik} + e_{kj}$.

The proof of the following technical result is immediate from the definitions.

Lemma 3.4.
1. If $v \models e \prec e'$ and $v \models e' \prec' e''$, then $v \models e (\prec \wedge \prec') e''$.
2. If $(v, w) \models x - y \prec e$ and $v \models e \prec' e'$, then $(v, w) \models x - y (\prec \wedge \prec') e'$.
3. If $v \models e (\prec \wedge \prec') e'$, then $v \models e \prec e'$.
4. If $(v, w) \models x - y (\prec \wedge \prec') e$, then $(v, w) \models x - y \prec e$.
5. If $(v, w) \models x - y \prec e$ and $(v, w) \models y - z \prec' e'$, then $(v, w) \models x - z (\prec \wedge \prec') e + e'$.
6. $v \models \neg(e \prec e')$ iff $v \models e' (\neg \prec) e$.

The next lemma states that canonicity of a constrained PDBM guarantees satisfiability.

Lemma 3.5. *Suppose (C, D) is a constrained PDBM in canonical form and $v \in [\![C]\!]$. Then D is satisfiable for v.*

Proof. By induction on i, with $0 \leqslant i \leqslant m$, we construct a valuation $(t_0, \ldots, t_i)$ for clock variables $(x_0, \ldots, x_i)$ such that all constraints D_{jk} for $0 \leqslant j, k \leqslant i$ are met.

To begin with, we set $t_0 = 0$. Then, trivially, $(v, x_0 \mapsto t_0) \models D_{00}$.

For the induction step, suppose that for some $i < m$ we have a valuation $(t_0, \ldots, t_i)$ for variables $(x_0, \ldots, x_i)$ such that all constraints D_{jk} for $0 \leqslant j, k \leqslant i$ are met. In order to extend this valuation to x_{i+1}, we have to find a value t_{i+1} such that the following simple guards hold for valuation $(v, x_0 \mapsto t_0, \ldots, x_{i+1} \mapsto t_{i+1})$:

$$D_{i+1,0} \quad \cdots \quad D_{i+1,i} \quad D_{0,i+1} \quad \cdots \quad D_{i,i+1} \quad D_{i+1,i+1}. \tag{1}$$

Here the first $i + 1$ simple guards give upper bounds for t_{i+1}, the second $i + 1$ simple guards give lower bounds for t_{i+1}, and the last simple guard is trivially met by any choice for t_{i+1}. We claim that each of the upper bounds is larger than or equal to each of the lower bounds. In particular, the minimum of the upper bounds is larger than or equal to the maximum of the lower bounds. This gives us a non-empty interval of possible values for t_{i+1} to choose from. Formally, we claim that, for all $0 \leqslant j, k < i + 1$, the following formula holds for valuation $(v, [x_0 \mapsto t_0, \ldots, x_i \mapsto t_i])$:

$$x_j - e_{j,i+1} (\prec_{j,i+1} \wedge \prec_{i+1,k}) x_k + e_{i+1,k}. \tag{2}$$

To see why (2) holds, observe that by induction hypothesis $(v, x_0 \mapsto t_0, \ldots, x_i \mapsto t_i) \models$

$$x_j - x_k \prec_{jk} e_{jk}. \tag{3}$$

Furthermore, since (C, D) is canonical and $v \in [\![C]\!]$, $v \models$

$$e_{jk}(\prec_{jk} \Rightarrow \prec_{j,i+1} \wedge \prec_{i+1,k})e_{j,i+1} + e_{i+1,k}. \tag{4}$$

Combination of (3) and (4), using Lemma 3.4(2), gives $(v, x_0 \mapsto t_0, \ldots, x_i \mapsto t_i) \models$

$$x_j - x_k (\prec_{j,i+1} \land \prec_{i+1,k}) e_{j,i+1} + e_{i+1,k},$$

which is equivalent to (2). This means that we can choose t_{i+1} in accordance with all the guards of (1). In particular, guard $D_{0,i+1}$ holds, which by the assumption that lower bounds on clocks are non-negative implies that t_{i+1} is non-negative. This completes the proof of the induction step and thereby of the lemma. $\square$

The following lemma essentially carries over from the non-parametric case too, see for instance [13]. As a direct consequence, semantic inclusion of constrained PDBMs is decidable for canonical PDBMs (using the oracle function).

Lemma 3.6. *Suppose $(C, D), (C', D')$ are constrained PDBMs and (C, D) is canonical. Then $[\![C, D]\!] \subseteq [\![C', D']\!] \Leftrightarrow ([\![C]\!] \subseteq [\![C']\!] \land \forall i, j : C \models e_{ij} (\prec_{ij} \Rightarrow \prec'_{ij}) e'_{ij}).$*

3.2. Operations on PDBMs

Our algorithm requires basically four operations to be implemented on constrained PDBMs: adding guards, canonicalization, resetting clocks and computing time successors.

3.2.1. Adding guards

In the case of DBMs, adding a guard is a simple operation. It is implemented by taking the conjunction of a DBM and the guard (which is also viewed as a DBM). The conjunction operation just takes the pointwise minimum of the entries in both matrices. In the parametric case, adding a guard to a constrained PDBM may result in a set of constrained PDBMs. We define a relation $\Leftarrow$ which relates a constrained PDBM and a guard to a collection of constrained PDBMs that satisfy this guard. For this we need an operation $\mathscr{C}$ that takes a PDBM and a simple guard, and produces a constraint stating that the bound imposed by the guard is weaker than the corresponding bound in the PDBM. Let $D^{ij} = (e_{ij}, \prec_{ij})$. Then

$$\mathscr{C}(D, x_i - x_j \prec e) = e_{ij} (\prec_{ij} \Rightarrow \prec) e.$$

Relation $\Leftarrow$ is defined as the smallest relation that satisfies the following rules:

(R_1) $\dfrac{\mathcal{O}(\mathscr{C}(D, f), C) = \text{yes}}{(C, D) \overset{f}{\Leftarrow} (C, D)},$

(R_2) $\dfrac{\mathcal{O}(\mathscr{C}(D, f), C) = \text{no}, \ f \text{ proper}}{(C, D) \overset{f}{\Leftarrow} (C, D[f])},$

(R_3) $\dfrac{\mathcal{O}(\mathscr{C}(D, f), C) = \text{split}}{(C, D) \overset{f}{\Leftarrow} (C \cup \{\mathscr{C}(D, f)\}, D)},$

(R_4) $\dfrac{\mathcal{O}(\mathscr{C}(D, f), C) = \text{split}, \ f \text{ proper}}{(C, D) \overset{f}{\Leftarrow} (C \cup \{\neg\mathscr{C}(D, f)\}, D[f])},$

(R_5) $\dfrac{(C, D) \overset{g}{\Leftarrow} (C', D') , \ (C'D') \overset{g'}{\Leftarrow} (C'', D'')}{(C, D) \overset{g \land g'}{\Leftarrow} (C'', D'')}.$

If the oracle replies "yes", then adding a simple guard will not change the constrained PDBM. If the answer is "no", then we tighten the bound in the PDBM. With the answer "split" there are two possibilities and two PDBMs with updated constraint systems are returned. Thus the result of the operation of adding a guard is a *set* of constrained PDBMs. The side condition "f proper" in R_2 and R_4 rules out guards of the form $x_i - x_i \prec e$ and thereby ensures that the diagonal bounds in the PDBM always remain equal to $(0, \leqslant)$. It is routine to check, using Lemma 3.4, that relation $\Leftleftarrows$ is well-defined in the sense that $(C, D) \overset{g}{\Leftleftarrows} (C', D')$ implies that (C', D') is a constrained PDBMs. In particular, the condition that clocks have non-negative lower bounds is met. Note that if we update a bound in D the semantics of the PDBM may become empty: a typical situation occurs when D contains a constraint $x \geqslant 5$ and we add a guard $x \leqslant 3$. Note however that $(C, D) \overset{g}{\Leftleftarrows} (C', D')$ and $[\![C]\!] \neq \emptyset$ implies $[\![C']\!] \neq \emptyset$. The following lemma characterizes $\Leftleftarrows$ semantically.

Lemma 3.7. $[\![C, D]\!] \cap [\![g]\!] = \bigcup \{[\![C', D']\!] \mid (C, D) \overset{g}{\Leftleftarrows} (C', D')\}.$

Proof. "$\subseteq$": Assume $(v, w) \in [\![C, D]\!] \wedge (v, w) \models g$. By structural induction on g we prove that there exists a constrained PDBM (C', D') such that $(C, D) \overset{g}{\Leftleftarrows} (C', D')$ and $(v, w) \in [\![C', D']\!]$.

For the induction basis, suppose g is of the form $x_i - x_j \prec e$. We consider four cases:
- $\mathcal{O}(\mathscr{C}(D, g), C) = $ yes. Let $C' = C$ and $D' = D$. Then trivially $(v, w) \in [\![C', D']\!]$ and, by rule R_1, $(C, D) \overset{g}{\Leftleftarrows} (C', D')$.
- $\mathcal{O}(\mathscr{C}(D, g), C) = $ no. By contradiction we prove that g is proper. Suppose g is not proper. Then, since $i = j$ and $v \models \neg e_{ij}(\prec_{ij} \Rightarrow \prec)e$, $v \models \neg(0 \prec e)$. By Lemma 3.4(6), $v \models e(\neg \prec)0$. But $(v, w) \models g$ implies $v \models 0 \prec e$. Hence, by Lemma 3.4(1), $v \models 0 < 0$, a contradiction. Let $C' = C$ and $D' = D[g]$. Then, by rule R_2, $(C, D) \overset{g}{\Leftleftarrows} (C', D')$. Since $v \in [\![C]\!]$ and $C' = C$, trivially $v \in [\![C']\!]$. Since $w \in [\![D]\!]_v$ and $(v, w) \models g$, easily $w \in [\![D[g]]\!]_v$. It follows that $(v, w) \in [\![C', D']\!]$.
- $\mathcal{O}(\mathscr{C}(D, g), C) = $ split and $v \models \mathscr{C}(D, g)$. Let $C' = C \cup \{\mathscr{C}(D, g)\}$ and $D' = D$. Then, by rule R_3, $(C, D) \overset{g}{\Leftleftarrows} (C', D')$. Since $v \in [\![C]\!]$ and $v \models \mathscr{C}(D, g)$, $v \in [\![C \cup \{\mathscr{C}(D, g)\}]\!]$. Since $w \in [\![D]\!]_v$ and $D' = D$, trivially $w \in [\![D']\!]_v$. It follows that $(v, w) \in [\![C', D']\!]$.
- $\mathcal{O}(\mathscr{C}(D, g), C) = $ split and $v \models \neg\mathscr{C}(D, g)$. By contradiction we prove that g is proper. Suppose g is not proper. Then, since $v \models \neg\mathscr{C}(D, g)$, $v \models \neg(0 \prec e)$. By Lemma 3.4(6), $v \models e\neg \prec 0$. But $(v, w) \models g$ implies $v \models 0 \prec e$. Hence, by Lemma 3.4(1), $v \models 0 < 0$, a contradiction. Let $C' = C \cup \{\neg\mathscr{C}(D, g)\}$ and $D' = D[g]$. Then, by rule R_4, $(C, D) \overset{g}{\Leftleftarrows} (C', D')$. Since $v \in [\![C]\!]$ and $v \models \neg\mathscr{C}(D, g)$, $v \in [\![C \cup \{\neg\mathscr{C}(D, g)\}]\!]$. Since $w \in [\![D]\!]_v$ and $(v, w) \models g$, easily $w \in [\![D[g]]\!]_v$. It follows that $(v, w) \in [\![C', D']\!]$.

For the induction step, suppose that g is of the form $g' \wedge g''$. Then $(v, w) \models g'$. By induction hypothesis, there exist C'', D'' such that $(C, D) \overset{g'}{\Leftleftarrows} (C'', D'')$ and $(v, w) \in [\![C'', D'']\!]$. Since $(v, w) \models g''$, we can use the induction hypothesis once more to infer that there exist C', D' such that $(C'', D'') \overset{g''}{\Leftleftarrows} (C', D')$ and $(v, w) \in [\![C', D']\!]$. Moreover, by rule R_5, $(C, D) \overset{g}{\Leftleftarrows} (C', D')$.

"$\supseteq$": Assume $(C, D) \overset{g}{\Leftleftarrows} (C', D')$ and $(v, w) \in [\![C', D']\!]$. By induction on the size of the derivation of $(C, D) \overset{g}{\Leftleftarrows} (C', D')$, we establish $(v, w) \in [\![C, D]\!]$ and $(v, w) \models g$.

There are five cases, depending on the last rule r used in the derivation of $(C, D) \overset{g}{\Longleftarrow} (C', D')$.

1. $r = R_1$. Then $C = C'$, $D = D'$ and $C \models \mathscr{C}(D, g)$. Let g be of the form $x_i - x_j \prec e$. Hence, $(v, w) \in [\![C, D]\!]$ and $v \models \mathscr{C}(D, g)$. By the first statement $(v, w) \models x_i - x_j \prec_{ij}^{D} e_{ij}^{D}$, and by the second statement $v \models e_{ij}^{D}$ $(\prec_{ij}^{D} \Rightarrow \prec)$ e. Combination of these two observations, using parts (2) and (4) of Lemma 3.4 yields $(v, w) \models g$.

2. $r = R_2$. Then $C = C'$, $D' = D[g]$ and $C \models \neg\mathscr{C}(D, g)$. Hence, $(v, w) \models g$ and $v \models \neg\mathscr{C}(D, g)$. Let g be of the form $x_i - x_j \prec e$. By Lemma 3.4(6), $v \models e \neg(\prec_{ij}^{D} \Rightarrow \prec) e_{ij}^{D}$. Using parts (2) and (4) of Lemma 3.4, combination of these two observations yields $(v, w) \models x_i - x_j \prec_{ij}^{D} e_{ij}^{D}$. Since trivially (v, w) is a model for all the other guards in D, $(v, w) \in [\![C, D]\!]$.

3. $r = R_3$. Then $C' = C \cup \{\mathscr{C}(D, g)\}$ and $D' = D$. Let g be of the form $x_i - x_j \prec e$. We have $(v, w) \in [\![C, D]\!]$. This implies $(v, w) \models x_i - x_j \prec_{ij}^{D} e_{ij}^{D}$. We also have $v \models e_{ij}^{D}$ $(\prec_{ij}^{D} \Rightarrow \prec)$ e. Combination of these two observations, using parts (2) and (4) of Lemma 3.4 yields $(v, w) \models g$.

4. $r = R_4$. Then $C' = C \cup \{\neg\mathscr{C}(D, g)\}$ and $D' = D[g]$. We have $v \models \neg\mathscr{C}(D, g)$ and $(v, w) \models g$. Let g be of the form $x_i - x_j \prec e$. By Lemma 3.4(6), $v \models e \neg(\prec_{ij}^{D} \Rightarrow \prec) e_{ij}^{D}$. Using parts (2) and (4) of Lemma 3.4 yields $(v, w) \models x_i - x_j \prec_{ij}^{D} e_{ij}^{D}$. Since trivially (v, w) is a model for all other guards in D, $(v, w) \in [\![C, D]\!]$.

5. $r = R_5$. Then g is of the form $g' \wedge g''$ and there are C'', D'' such that $(C, D) \overset{g'}{\Longleftarrow} (C'', D'')$ and $(C'', D'') \overset{g''}{\Longleftarrow} (C', D')$. By induction hypothesis, $(v, w) \in [\![C'', D'']\!]$ and $(v, w) \models g''$. Again by induction hypothesis, $(v, w) \in [\![C, D]\!]$ and $(v, w) \models g'$. It follows that $(v, w) \models g$. $\square$

3.2.2. Canonicalization

Each DBM can be brought into canonical form using classical algorithms for computing all-pairs shortest paths, for instance the Floyd–Warshall (FW) algorithm [11]. In the parametric case, we also apply this approach except that now we run FW *symbolically*, see Fig. 4. The algorithm repeatedly compares the difference between two clocks to the difference obtained by taking an intermediate clock into account (cf. the inequality in Definition 3.3). The symbolic FW algorithm contains a non-deterministic assignment, in which (C, D) non-deterministically gets a value from a set. This set may be empty, in which case the algorithm terminates unsuccessfully. We are interested in the (possibly empty, finite) set of all possible constrained PDBMs that may result when running the algorithm.

For the purpose of proving things we find it convenient to describe the computation steps of the symbolic FW algorithm in SOS style. In the SOS description, we use configurations of the form (k, i, j, C, D), where (C, D) is a constrained PDBM and $k, i, j \in [0, m+1]$ record the values of indices. In the rules below, k, i, j range over $[0, m]$:

$$\frac{(C, D) \overset{x_i - x_j \ (\prec_{ik} \wedge \prec_{kj}) \ e_{ik} + e_{kj}}{\Longleftarrow} (C', D')}{(k, i, j, C, D) \rightarrow_{\text{FW}} (k, i, j+1, C', D')},$$

$$(k, i, m+1, C, D) \rightarrow_{\text{FW}} (k, i+1, 0, C, D),$$

$$(k, m+1, 0, C, D) \rightarrow_{\text{FW}} (k+1, 0, 0, C, D).$$

We write $(C, D) \rightarrow_c (C', D')$ if there exists a sequence of $\rightarrow_{\text{FW}}$ steps leading from configuration $(0, 0, 0, C, D)$ to configuration $(m+1, 0, 0, C', D')$. In this case, we say

$\textsc{Floyd-Warshall}\ (C_0, D_0)$
$(C, D) := (C_0, D_0)$
for $k = 0$ **to** m
    do for $i = 0$ **to** m
        do for $j = 0$ **to** m
            $(C, D) := $ **choose** (C', D') **such that**
                $(C, D) \overset{x_i - x_j \ \prec_{ik} \wedge \prec_{kj}\ e_{ik} + e_{kj}}{\Longleftarrow}$
(C', D')
return (C, D)

Fig. 4. The Floyd–Warshall algorithm.

that (C', D') is an *outcome* of the symbolic Floyd–Warshall algorithm on (C, D). It is easy to see that the set of all outcomes is finite and can be effectively computed. If the semantics of (C, D) is empty, then the set of outcomes is also empty. We write $(C, D) \overset{g}{\Leftarrow}_c (C', D')$ iff $(C, D) \overset{g}{\Leftarrow} (C'', D'') \to_c (C', D')$, for some C'', D''.

The following lemma says that if we run the symbolic Floyd–Warshall algorithm, the union of the semantics of the outcomes equals the semantics of the original constrained PDBM.

Lemma 3.8. $[\![C, D]\!] = \bigcup \left\{ [\![C', D']\!] \mid (C, D) \to_c (C', D') \right\}.$

Proof. By an inductive argument, using Lemma 3.7 and the fact that, for any valuation (v, w) in the semantics of (C, D),

$(v, w) \models x_i - x_k \ \prec_{ik}\ e_{ik}$ and
$(v, w) \models x_k - x_j \ \prec_{kj}\ e_{kj}$ and therefore by Lemma 3.4(5)
$(v, w) \models x_i - x_j \ \prec_{ik} \wedge \prec_{kj}\ e_{ik} + e_{kj}.$ $\square$

Lemma 3.9. *Each outcome of the symbolic Floyd–Warshall algorithm is a constrained PDBM in canonical form.*

Proof. As in [11]. $\square$

Remark 3.10. Non-parametric DBMs can be canonicalized in $\mathcal{O}(n^3)$, where n is the number of clocks. In the parametric case, however, each operation of comparing the bound $D(x, x')$ to the weight of another path from x to x' may give rise to two new PDBMs if this comparison leads to a split. Then the two PDBMs must both be canonicalized to obtain all possible PDBMs with tightest bounds. Still, that part of these two PDBMs which was already canonical, does not need to be investigated again. So in the worst case, the cost of the algorithm becomes $\mathcal{O}(2^{n^3})$. In practice, it turns out that this is hardly ever the case.

3.2.3. Resetting clocks

A third operation on PDBMs that we need is resetting clocks. Since we do not allow parameters in reset sets, the reset operation on PDBMs is essentially the same as for DBMs, see [28]. If D is a PDBM and r is a singleton reset set $\{x_i := b\}$, then $D[r]$ is the PDBM

obtained by (1) replacing each bound D^{ij}, for $j \neq i$, by $(e_{0j} + b, \prec_{0j})$; (2) replacing each bound D^{ji}, for $j \neq i$, by $(e_{j0} - b, \prec_{j0})$. We generalize this definition to arbitrary reset sets by

$$D[x_{i_1} := b_1, \ldots, x_{i_h} := b_h] = D[x_{i_1} := b_1] \cdots [x_{i_h} := b_h].$$

It easily follows from the definitions that resets preserve canonicity. Note also that the reset operation is well-defined on constrained PDBMs: if (C, D) is a constrained PDBM, then $(C, D[r])$ is a constrained PDBM as well: since clocks can only be reset to natural numbers, lower bounds on clocks remain non-negative.

Lemma 3.11. *If (C, D) is canonical, then $(C, D[r])$ is canonical as well.*

The following lemma characterizes the reset operation semantically.

Lemma 3.12. *Let (C, D) be a constrained PDBM in canonical form, $v \in [\![C]\!]$, and w a clock valuation. Then $w \in [\![D[r]]\!]_v$ iff $\exists w' \in [\![D]\!]_v : w = w'[r]$.*

Proof. We only prove the lemma for singleton resets. Using Lemma 3.11, the generalization to arbitrary resets is straightforward. Let $r = \{x_i := b\}$ and $D' = D[r]$.

"$\Leftarrow$" Suppose $w' \in [\![D]\!]_v$ and $w = w'[r]$. In order to prove $w \in [\![D']\!]_v$, we must show that $(v, w) \models D'_{kj}$, for all k and j. There are four cases:

1. $k \neq i \neq j$. Then $D'_{kj} = D_{kj}$. Since $(v, w') \models D_{kj}$ and w and w' agree on all clocks occurring in D_{kj}, $(v, w) \models D'_{kj}$.
2. $k = i = j$. Then $D'_{kj} = D_{kj}$. Since $(v, w') \models D_{ii}$, $0 \prec_{ii} e_{ii}[v]$. Hence, $(v, w) \models D'_{kj}$.
3. $k \neq i = j$. Then $D'_{kj} = x_k - x_j \prec_{k0} e_{k0} - b$. Using that $(v, w') \models D_{k0}$, we derive $w(x_k) - w(x_j) = w'(x_k) - b \prec_{k0} e_{k0}[v] - b$. Hence, $(v, w) \models D'_{kj}$.
4. $k = i \neq j$. Then $D'_{kj} = x_k - x_j \prec_{0j} e_{0j} + b$. Using that $(v, w') \models D_{0j}$, we derive $w(x_k) - w(x_j) = b - w'(x_j) \prec_{0j} e_{0j}[v] + b$. Hence, $(v, w) \models D'_{kj}$.

"$\Rightarrow$" Suppose $w \in [\![D']\!]_v$. We have to prove that there exists a clock valuation $w' \in [\![D]\!]_v$ such that $w = w'[r]$. Clearly we need to choose w' in such a way that, for all $j \neq i$, $w'(x_j) = w(x_j)$. This means that, for any choice of $w'(x_i)$, for all $j \neq i \neq k$, $v, w' \models D_{jk}$. Using the same argument as in the proof of Lemma 3.5, we can find a value for $w'(x_i)$ such that also the remaining simple guards of D are satisfied. $\square$

3.2.4. Time successors

Finally, we need to transform PDBMs for the passage of time, notation $D \uparrow$. As in the DBMs case [13], this is done by setting the upper bounds $x_i - x_0$ to $(\infty, <)$, for each $i \neq 0$, and leaving all other bounds unchanged. We have the following lemma.

Lemma 3.13. *Suppose (C, D) is a constrained PDBM in canonical form, $v \in [\![C]\!]$, and w a clock valuation. Then $w \in [\![D \uparrow]\!]_v$ iff $\exists d \geqslant 0 \, \exists w' \in [\![D]\!]_v : w' + d = w$.*

Proof. "$\Leftarrow$" Suppose $d \geqslant 0$, $w' \in [\![D]\!]_v$ and $w' + d = w$. We claim that $w \in [\![D \uparrow]\!]_v$. For this we must show that for each guard f of $D \uparrow$, $(v, w) \models f$. Let f be of the form $x_i - x_j \prec e$. We distinguish between three cases:

- $i \neq 0 \land j = 0$. In this case, by definition of $D\uparrow$, f is of the form $x_i - x_0 < \infty$, and so $(v, w) \models f$ trivially holds.
- $i = 0$. In this case f is also a constraint of D. Since $w' \in \llbracket D \rrbracket_v$ we have $(v, w') \models f$, and thus $-w'(x_j) \prec e[v]$. But since $d \geqslant 0$, this means that $-w(x_j) = -w'(x_i) - d \prec e[v]$ and therefore $(v, w) \models f$.
- $i \neq 0 \land j \neq 0$. In this case f is again a constraint of D. Since $w' \in \llbracket D \rrbracket_v$ we have $(v, w') \models f$, and therefore $w'(x_i) - w'(x_j) \prec e[v]$. But this means that $w'(x_i) - w'(x_j) = (w(x_i) - d) - (w(x_j) - d) \prec e[v]$ and therefore $(v, w) \models f$.

"⇒" Suppose $w \in \llbracket D\uparrow \rrbracket_v$. If $m = 0$ (i.e., there are no clocks), then $D\uparrow = D$ and the right-hand side of the implication trivially holds (take $w' = w$ and $d = 0$). So assume $m > 0$. For all indices i, j with $i \neq 0$ and $j \neq 0$, $(v, w) \models D_{ij}$. Hence, $w(x_i) - w(x_j) \prec_{ij} e_{ij}[v]$. Thus, for any real number t, $w(x_i) - t - (w(x_j) - t) \prec_{ij} e_{ij}[v]$. But this means $(v, w - t) \models D_{ij}$. It remains to be shown that there exists a value d such that in valuation $(v, w - d)$ also the remaining guards D_{0i} and D_{i0} hold. Let

$$t_0 = \max(0, w(x_1) - e_{10}[v], \ldots, w(x_n) - e_{n0}[v]),$$
$$t_1 = \min(w(x_1) + e_{01}[v], \ldots, w(x_n) + e_{0n}[v]),$$
$$d = (t_0 + t_1)/2,$$
$$w' = w - d.$$

Intuitively, t_0 gives the least amount of time one has to go backwards in time from w to meet all upper bounds of D (modulo strictness), whereas t_1 gives the largest amount of time one can go backwards in time from w without violating any of the lower bounds of D (again modulo strictness). Value d sits right in the middle of these two. We claim that d and w' satisfy the required properties. For any i, since $(v, w) \models D_{0i}$, trivially

$$0 \prec_{0i} w(x_i) + e_{0i}[v]. \tag{5}$$

Using that D is canonical we have, for any i, j,

$$e_{ji}[v] \ (\prec_{ji} \Rightarrow \prec_{j0} \land \prec_{0i}) \ e_{j0}[v] + e_{0i}[v]$$

and, since $v, w \models D_{ji}$,

$$w(x_j) - w(x_i) \prec_{ji} e_{ji}[v].$$

Using these two observations, we infer

$$w(x_j) - e_{j0}[v] \ (\prec_{ji} \Rightarrow \prec_{j0} \land \prec_{0i}) \ w(x_j) - e_{ji}[v] + e_{0i}[v] \prec_{ji} w(x_i) + e_{0i}[v].$$

Hence,

$$w(x_j) - e_{j0}[v] \prec_{j0} \land \prec_{0i} w(x_i) + e_{0i}[v]. \tag{6}$$

By inequalities (5) and (6), each subterm of max in the definition of t_0 is dominated by each subterm of min in the definition of t_1. This implies $0 \leqslant t_0 \leqslant t_1$.

Now either $t_0 < t_1$ or $t_0 = t_1$. In the first case it is easy to prove that in valuation (v, w) the guards D_{0i} and D_{i0} hold, for any i:

$$w'(x_i) = w(x_i) - d < w(x_i) - t_0 \leqslant w(x_i) - (w(x_i) - e_{i0}[v]) = e_{i0}[v]$$

and thus $w'(x_i) < e_{i0}[v]$ and $v, w' \models D_{i0}$. Also

$$-w'(x_i) = -w(x_i) + d < -w(x_i) + t_1 \leqslant -w(x_i) + (w(x_i) + e_{0i}[v]) = e_{0i}[v]$$

and so $-w'(x_i) < e_{0i}[v]$ and $v, w' \models D_{0i}$.

In the second case, fix an i. If $w(x_i) - e_{i0}[v] < t_0$, then

$$w'(x_i) = w(x_i) - d = w(x_i) - t_0 < w(x_i) - (w(x_i) - e_{i0}[v]) = e_{i0}[v]$$

and thus $w'(x_i) < e_{i0}[v]$ and $v, w' \models D_{i0}$. Otherwise, if $w(x_i) - e_{i0}[v] = t_0$ observe that by $t_0 = t_1$, inequality (6) and the fact that, $t_1 = w(x_j) + e_{0j}[v]$, for some j, $\prec_{i0} = \leqslant$. Since

$$w'(x_i) = w(x_i) - d \leqslant w(x_i) - t_0 \leqslant w(x_i) - (w(x_i) - e_{i0}[v]) \leqslant e_{i0}[v]$$

and thus $w'(x_i) \leqslant e_{i0}[v]$ this implies $v, w' \models D_{i0}$. $\square$

3.3. Symbolic semantics

Having defined the four operations on PDBMs, we are now in a position to describe the semantics of a parametric timed automaton symbolically.

Definition 3.14 (*Symbolic semantics*). Let $\mathscr{A} = (Q, q_0, \rightarrow, I)$ be a PTA. The *symbolic semantics* of $\mathscr{A}$ is an LTS: the states are triplets (q, C, D) with q a location from Q and (C, D) a constrained PDBM in canonical form such that $[\![C, D]\!] \subseteq [\![I(q)]\!]$; the set of initial states is

$$\{(q_0, C, D) \mid (C_0, \mathsf{E}\uparrow) \stackrel{I(q_0)}{\Leftarrow}_c (C, D)\},$$

where $C_0 = \{p \geqslant 0 \mid p \in P\}$, E is the PDBM with $\mathsf{E}^{ij} = (0, \leqslant)$, for all i, j; the transitions are defined by the following rule:

$$\frac{q \xrightarrow{a,g,r} q', (C, D) \stackrel{g}{\Leftarrow}_c (C'', D''), (C'', D''[r]\uparrow) \stackrel{I(q')}{\Leftarrow}_c (C', D')}{(q, C, D) \rightarrow (q', C', D')}.$$

Observe that if (q, C, D) is a state in the symbolic semantics and $(v, w) \in [\![C, D]\!]$, then (q, w) is a state of the concrete semantics $[\![\mathscr{A}]\!]_v$. It is also easy to see that the symbolic semantics of a PTA is a finitely branching LTS. It may have infinitely many reachable states though.

In order to establish that each run in the symbolic semantics can be simulated by a run in the concrete semantics, we require two lemmas.

Lemma 3.15. *Suppose that* (q, C, D) *is an initial state of the symbolic semantics of* $\mathscr{A}$ *with* $(v, w) \in [\![C, D]\!]$. *Then the concrete semantics* $[\![\mathscr{A}]\!]_v$ *has an initial state* (q_0, w_0) *from which state* (q, w) *can be reached.*

Proof. Using the fact that $(v, w) \in [\![C, D]\!]$, the definition of initial states, Lemmas 3.8 and 3.7, we know that $q = q_0$, $(v, w) \models I(q_0)$ and $(v, w) \in [\![C_0, \mathsf{E}\uparrow]\!]$. By Lemma 3.13, we get that there exists a $d \geqslant 0$ and $w_0 \in [\![\mathsf{E}]\!]_v$ such that $w_0 + d = w$. Since $(v, w) \models I(q_0)$ and invariants in a PTA only give upper bounds on clocks, also $(v, w_0) \models I(q_0)$. It follows that (q_0, w_0) is a state of the concrete semantics $[\![\mathscr{A}]\!]_v$ and $(q_0, w_0) \xrightarrow{d} (q, w)$. Since $w_0 \in [\![\mathsf{E}]\!]_v$, w_0 is of the form $\lambda x.0$. Hence, (q_0, w_0) is an initial state of the concrete semantics. $\square$

Lemma 3.16. *Suppose that* $(q', C', D') \rightarrow (q, C, D)$ *is a transition in the symbolic semantics of* $\mathscr{A}$ *and* $(v, w) \in [\![C, D]\!]$. *Then there exists a pair* $(v, w') \in [\![C, D]\!]$ *such that in the concrete semantics* $[\![\mathscr{A}]\!]_v$ *there is a path from* (q', w') *to* (q, w).

Proof. By the definition of transitions in the symbolic semantics, Lemmas 3.8 and 3.7, we know that there is a transition $q' \xrightarrow{a,g,r} q$ in $\mathscr{A}$, and there are C'', D'' such that $(v, w) \models I(q)$, $(v, w) \in [\![C'', D''[r]\!\uparrow]\!]$ and $(C', D') \xleftarrow{g}_c (C'', D'')$. By Lemma 3.13, we get that there exists a $d \geqslant 0$ and $w'' \in [\![D''[r]]\!]_v$ such that $w'' + d = w$. Since $(v, w) \models I(q)$ and invariants in a PTA only give upper bounds on clocks, also $(v, w'') \models I(q)$. It follows that (q, w'') is a state of the concrete semantics $[\![\mathscr{A}]\!]_v$ and $(q, w'') \xrightarrow{d} (q, w)$. Using Lemma 3.12 we get that there exists a $w' \in [\![D'']\!]_v$ such that $w'' = w'[r]$. Using Lemmas 3.8 and 3.7 again, it follows that $(v, w') \models g$ and $(v, w') \in [\![C', D']\!]$. Since (q', C', D') is a state of the symbolic semantics, $(v, w') \models I(q')$. Hence, (q', w') is a state of the concrete semantics and $(q', w') \xrightarrow{a} (q, w'')$ is a transition in the concrete semantics. Combination of this transition with the transition $(q, w'') \xrightarrow{d} (q, w)$ gives the required path in the concrete semantics. □

Proposition 3.17. *For each parameter valuation v and clock valuation w, if there is a run in the symbolic semantics of $\mathscr{A}$ reaching state (q, C, D), with $(v, w) \in [\![C, D]\!]$, then this run can be simulated by a run in the concrete semantics $[\![\mathscr{A}]\!]_v$ reaching state (q, w).*

Proof. By induction on the number of transitions in the run.

As basis we consider a run with 0 transitions, i.e., a run that consists of an initial state of the symbolic semantics. So this means that (q, C, D) is an initial state. The induction basis now directly follows using Lemma 3.15.

For the induction step, assume that we have a run in the symbolic semantics, ending with a transition $(q', C', D') \rightarrow (q, C, D)$. By $(v, w) \in [\![C, D]\!]$ and Lemma 3.16, there exists a pair $(v, w') \in [\![C, D]\!]$ such that in the concrete semantics $[\![\mathscr{A}]\!]_v$ there is a path from (q', w') to (q, w). By induction hypothesis, there is a path in the concrete semantics leading up to state (q', w'). Extension of this path with the path from (q', w') to (q, w) gives the required path in the concrete semantics. □

Conversely, for each path in the concrete semantics, we can find a path in the symbolic semantics such that the final state of the first path is semantically contained in the final state of the second path.

Proposition 3.18. *For each parameter valuation v and clock valuation w, if there is a run in the concrete semantics $[\![\mathscr{A}]\!]_v$ reaching a state (q, w), then this run can be simulated by a run in the symbolic semantics reaching a state (q, C, D) such that $(v, w) \in [\![C, D]\!]$.*

Proof. In any execution in the concrete semantics, we can always insert zero-delay transitions at any point. Also, two consecutive delay transitions $(q, w) \xrightarrow{d} (q, w + d)$ and $(q, w + d) \xrightarrow{d'} (q, w + d + d')$ can always be combined into a single delay transition $(q, w) \xrightarrow{d+d'} (q, w + d + d')$. Therefore, without loss of generality, we only consider concrete executions that start with a delay transition, and in which there is a strict alternation of action transitions and delay transitions. The proof is by induction on the number of action transitions.

As basis we consider a run consisting of a single time-passage transition: $(q_0, w_0) \xrightarrow{d} (q_0, w_0 + d)$, where $w_0 = \lambda x.0$. By definition of the concrete semantics, $(v, w_0 + d) \models$

$I(q_0)$. Using Lemma 3.13, we have that $(v, w_0 + d) \in [\![C_0, \mathsf{E}\uparrow]\!]$ since $(v, w_0) \in [\![C_0, \mathsf{E}]\!]$. From $(v, w_0 + d) \in [\![C_0, \mathsf{E}\uparrow]\!]$ and $(v, w_0 + d) \models I(q_0)$, using Lemmas 3.7 and 3.8 we get that there exists C, D such that $(C_0, \mathsf{E}\uparrow) \overset{I(q_0)}{\Longleftarrow}_c (C, D)$ and $(v, w_0 + d) \in [\![C, D]\!]$. By definition, (C, D) is an initial state of the symbolic semantics. This completes the proof of the induction basis.

For the induction step, assume that the run in the concrete semantics of $[\![\mathscr{A}]\!]_v$ ends with transitions $(q'', w'') \overset{a}{\longrightarrow} (q', w') \overset{d}{\longrightarrow} (q, w)$. By induction hypothesis, there exists a run in the symbolic semantics ending with a state (q'', C'', D'') such that $(v, w'') \in [\![C'', D'']\!]$.

By definition of the concrete semantics, there is a transition $q'' \overset{g,a,r}{\longrightarrow} q'$ in $\mathscr{A}$ such that $(v, w'') \models g$ and $w' = w''[r]$. Moreover, we have $q' = q$, $w = w' + d$ and $(v, w) \models I(q)$. Using Lemmas 3.7 and 3.8 gives that there exist C', D' such that $(C'', D'') \overset{g}{\Longleftarrow}_c (C', D')$ and $(v, w'') \in [\![C', D']\!]$. By Lemma 3.12, $w' \in [\![D'[r]]\!]_v$. Moreover, by Lemma 3.13, $w \in [\![D'[r]\uparrow]\!]_v$. Using $(v, w) \models I(q)$, Lemmas 3.7 and 3.8, we infer that there exists C, D such that $(v, w) \in [\![C, D]\!]$ and $(C', D'[r]\uparrow) \overset{I(q)}{\Longleftarrow}_c (C, D)$. Finally, using the definition of the symbolic semantics, we infer the existence of a transition $(q'', C'', D'') \to (q, C, D)$. $\square$

Example 3.19. Fig. 5 shows the symbolic state-space of the automaton in Fig. 1 represented by constrained PDBMs. In the initial state the invariant $x \leqslant p$ limits the value of x, and since both clocks have the same value also the value of y. When taking the transition from S_0 to S_1 we have to compare the parameters p and q. This leads to a split where in the one case no state is reachable since the region is empty, and in the other (when $q \leqslant p$) S_1 can be reached. From then on, no more splits occur and only one new state is reachable.

3.4. Evaluating state formulas

We now define the predicate $\overset{\phi}{\Longleftarrow}$ which relates a symbolic state and a state formula ϕ (as defined in Definition 2.7) to a collection of symbolic states that satisfy ϕ.


Fig. 5. The symbolic state space of the PTA in Fig. 1.

In order to check whether a state formula holds, we break it down into its atomic sub-formulas, namely checking locations and clock guards. Checking that a clock guard holds relies on the definition given earlier, of adding that clock guard to the constrained PDBM. We rely on a special normal form of the state formula, in which all $\neg$ signs have been pushed down to the basic formulas.

Definition 3.20. State formula ϕ is in *normal form* if all $\neg$ signs in ϕ appear only in subformulae of the form $\neg q$.

Since each simple guard with a $\neg$ sign in front can be rewritten to equivalent simple guard without, for each state formula there is an equivalent one in normal form.

In the following, let f be a simple guard, and ϕ be in normal form.

$$(Q_1) \quad \frac{}{(q, C, D) \overset{q}{\Longleftarrow} (q, C, D)},$$

$$(Q_2) \quad \frac{q \neq q'}{(q, C, D) \overset{\neg q'}{\Longleftarrow} (q, C, D)},$$

$$(Q_3) \quad \frac{(C, D) \overset{f}{\Longleftarrow}_c (C', D')}{(q, C, D) \overset{f}{\Longleftarrow} (q, C', D')},$$

$$(Q_4) \quad \frac{(q, C, D) \overset{\phi_1}{\Longleftarrow} (q, C', D'), (q, C', D') \overset{\phi_2}{\Longleftarrow} (q, C'', D'')}{(q, C, D) \overset{\phi_1 \wedge \phi_2}{\Longleftarrow} (q, C'', D'')},$$

$$(Q_5) \quad \frac{(q, C, D) \overset{\phi_1}{\Longleftarrow} (q, C', D')}{(q, C, D) \overset{\phi_1 \vee \phi_2}{\Longleftarrow} (q, C', D')},$$

$$(Q_6) \quad \frac{(q, C, D) \overset{\phi_2}{\Longleftarrow} (q, C', D')}{(q, C, D) \overset{\phi_1 \vee \phi_2}{\Longleftarrow} (q, C', D')}.$$

The following lemma gives the soundness and completeness of relation $\overset{\phi}{\Longleftarrow}$.

Lemma 3.21. *Let ϕ be a state formula in normal form, q a location and (C, D) a con-strained PDBMs. Let $[\![q, \phi]\!]$ denote the set $\{(v, w) \mid (q, w) \models_v \phi\}$. Then*

$$[\![C, D]\!] \cap [\![q, \phi]\!] = \bigcup \left\{ [\![C', D']\!] \mid (q, C, D) \overset{\phi}{\Longleftarrow} (q, C', D') \right\}.$$

Proof. "$\subseteq$": Assume that $(v, w) \in [\![C, D]\!]$ and $(q, w) \models_v \phi$. We prove that there are C', D' such that $(v, w) \in [\![C', D']\!]$ and $(q, C, D) \overset{\phi}{\Longleftarrow} (q, C', D')$. We proceed by induction on the structure of ϕ.

- Base cases.
  - Suppose $\phi = q'$. As $(q, w) \models_v q'$, clearly, $q = q'$. Since, by rule Q_1, $(q, C, D) \overset{q}{\Longleftarrow}$ (q, C, D), we can take $C = C'$ and $D = D'$ and the result follows.

· Suppose $\phi = \neg q'$. Similar to the previous case, apply rule Q_2.

· Suppose $\phi = f$ with f a simple guard. Then $(v, w) \in [\![C, D]\!]$ and $(v, w) \models f$.

By Lemma 3.7 there exist C'', D'' such that $(C, D) \overset{f}{\Leftarrow} (C'', D'')$ and $(v, w) \in [\![C'', D'']\!]$. Lemma 3.8 yields the existence of C', D' with $(C'', D'') \to_c (C', D')$ and $(v, w) \in [\![C', D']\!]$. By application of rule Q_3 we have $(q, C, D) \overset{f}{\Leftarrow} (q, C', D')$.

• Induction step.

· Suppose $\phi = \phi_1 \wedge \phi_2$. Then $(q, w) \models_v \phi_1$ and $(q, w) \models_v \phi_2$. By applying the induction hypothesis on ϕ_1, we derive that there exist C'', D'' such that $(q, C, D) \overset{\phi_1}{\Leftarrow} (q, C'', D'')$ and $(v, w) \in [\![C'', D'']\!]$. Applying the induction hypothesis on ϕ_2 yields the existence of C', D' such that $(q, C'', D'') \overset{\phi_2}{\Leftarrow} (q, C', D')$ and $(v, w) \in [\![C', D']\!]$. Then by application of rule Q_4 also $(q, C, D) \overset{\phi_1 \wedge \phi_2}{\Leftarrow} (q, C', D')$.

· Suppose $\phi = \phi_1 \vee \phi_2$. Then $(q, w) \models_v \phi_1$ or $(q, w) \models_v \phi_2$. Suppose that $(q, w) \models_v \phi_1$. The induction hypothesis yields the existence of C', D' such that $(q, C, D) \overset{\phi_1}{\Leftarrow} (q, C', D')$ and $(v, w) \in [\![C', D']\!]$. Then, by application of rule Q_5, $(q, C, D) \overset{\phi_1 \vee \phi_2}{\Leftarrow} (q, C', D')$. The case $(q, w) \models \phi_2$ is similar (using rule Q_6).

"$\supseteq$": Assume $(q, C, D) \overset{\phi}{\Leftarrow} (q, C', D')$ and $(v, w) \in [\![C', D']\!]$. By induction on the structure of the derivation of $\overset{\phi}{\Leftarrow}$, we establish that $(v, w) \in [\![C, D]\!]$ and $(q, w) \models_v \phi$.

• Base cases. The derivation consists of a single step r.

· $r = Q_1$. Then $\phi = q$, $C = C'$, $D = D'$. Trivially $(v, w) \in [\![C, D]\!]$ and $(q, w) \models_v q$.

· $r = Q_2$. Similar to the previous case.

· $r = Q_3$. Suppose $\phi = f$ with f a simple guard. Then $(C, D) \overset{f}{\Leftarrow}_c (C', D')$. This means that there exist C'', D'' such that $(C, D) \overset{f}{\Leftarrow} (C'', D'')$ and $(C'', D'') \to_c (C', D')$. By Lemma 3.8 we have $(v, w) \in [\![C'', D'']\!]$. Then we have by Lemma 3.7 that $(v, w) \models f$ and $(v, w) \in [\![C, D]\!]$.

• Induction step. Consider the last rule r used in the derivation of $(q, C, D) \overset{\phi}{\Leftarrow} (q, C', D')$.

· $r = Q_4$. Then $\phi = \phi_1 \wedge \phi_2$ and $(q, C, D) \overset{\phi_1}{\Leftarrow} (q, C'', D'')$ and $(q, C'', D'') \overset{\phi_2}{\Leftarrow} (q, C', D')$ for some C'', D''. Applying the induction hypothesis to the derivation of $\overset{\phi_1}{\Leftarrow}$ yields $(q, w) \models_v \phi_2$ and $(v, w) \in [\![C'', D'']\!]$. Then applying the induction hypothesis to the derivation of $\overset{\phi_2}{\Leftarrow}$ yields $(q, w) \models_v \phi_1$ and $(v, w) \in [\![C, D]\!]$. Then also $(q, w) \models_v \phi_1 \wedge \phi_2$.

· $r = Q_5$. Then $\phi = \phi_1 \vee \phi_2$. Then $(q, C, D) \overset{\phi_1}{\Leftarrow} (q, C', D')$. By induction hypothesis we have $(q, w) \models_v \phi_1$ and $(v, w) \in [\![C, D]\!]$.

· $r = Q_6$. Similar to the previous case. $\square$

3.5. Algorithm

We are now in a position to present our model checking algorithm for parametric timed automata. The algorithm displayed in Fig. 6 describes how our tool explores the symbolic state space and searches for constraints on the parameters for which a reachability property $\exists \Diamond \phi$ holds in a PTA $\mathscr{A}$.

Reachable $(\mathcal{A}, \phi)$

    Result $:= \emptyset$, Passed $:= \emptyset$, Waiting $:= \{(q_0, C, D) \mid (C_0, \mathsf{E}\uparrow) \overset{I(q_0)}{\underset{c}{\Leftarrow}} (C, D)\}$

    while Waiting $\neq \emptyset$ **do**

      select (q, C, D) from Waiting

      Result $:=$ Result $\cup \{(q', C', D') \mid (q, C, D) \overset{\phi}{\Leftarrow} (q', C', D')\}$

      False $:= \{(q', C', D') \mid (q, C, D) \overset{\neg\phi}{\Leftarrow} (q', C', D')\}$

      for each (q', C', D') in False **do**

        if for all (q'', C'', D'') in Passed: $(q', C', D') \not\subseteq (q'', C'', D'')$ **then**

          add (q', C', D') to Passed

          for each (q'', C'', D'') such that $(q', C', D') \rightarrow (q'', C'', D'')$ **do**

            Waiting $:=$ Waiting $\cup \{(q'', C''', D'')\}$

    return Result

Fig. 6. The parametric model checking algorithm.

In the algorithm, we use inclusion between symbolic states defined by

$$(q, C, D) \subseteq (q', C', D') \overset{\triangle}{=} q = q' \wedge [\![C, D]\!] \subseteq [\![C, D']\!].$$

Note that whenever a triple (q, C, D) ends up in one of the lists maintained by the algorithm, (C, D) is a constrained PDBM in canonical form. This fact, in combination with Lemma 3.6, gives decidability of the inclusion operation. Our search algorithm explores the symbolic semantics in an "intelligent" manner, and stops whenever it reaches a state whose semantics is contained in the semantics of a state that has been encountered before. Despite this, our algorithm need not terminate.

If it terminates, the result returned by the algorithm is a set of satisfiable symbolic states, all of which satisfy ϕ, for any valuation of the parameters and clocks in the state.

Theorem 3.22. *Suppose (q, C, D) is in the result set returned by* Reachable *$(\mathcal{A}, \phi)$. Then (C, D) is satisfiable. Moreover, for all $(v, w) \in [\![C, D]\!]$, (q, w) is a reachable state of $[\![\mathcal{A}]\!]_v$ and $(q, w) \models_v \phi$.*

Proof. It is easy to see that all the symbolic states returned by the algorithm are satisfiable: the only operation that may modify the constraint set is adding a guard, but this will never lead to unsatisfiable constraint sets. Since all constrained PDBMs returned by the algorithm are in canonical form, they are all satisfiable by Lemma 3.5.

Suppose that $(v, w) \in [\![C, D]\!]$. By a straightforward inductive argument, using Lemmas 3.15, 3.16 and 3.21, it follows that (q, w) is a reachable state of $[\![\mathcal{A}]\!]_v$ and $(q, w) \models_v \phi$. $\square$

For invariance properties $\forall\Box\phi$, our tool runs the algorithm on $\neg\phi$, and the result is then a set of symbolic states, none of which satisfies ϕ. The answer to the model checking problem, stated in Section 2.2, is obtained by taking the union of the constraint sets from all symbolic states in the result of the algorithm; in the case of an invariance property we take the complement of this set.

A difference between the above algorithm and the standard timed model checking algorithm is that we continue the exploration until either no more new states are found or all paths end in a state satisfying the property. This is because we want to find all the possible constraints on the parameters for which the property holds. Also, the operations

on non-parametric DBMs only change the DBM they are applied to, whereas in our case, we may end up with a set of new PDBMs and not just one.

Some standard operations on symbolic states that help in exploring as little as possible, have also been implemented in our tool for parametric symbolic states. Before starting the state space exploration, our implementation determines the *maximal constant* for each clock. This is the maximal value to which the clock is compared in any guard or invariant in the PTA. When the clock value grows beyond this value, we can ignore its real value. This enables us to identify many more symbolic states, and helps termination. In fact, for unparameterized timed automata this trick *guarantees* termination [4,7].

4. Lower bound/upper bound automata

This section introduces the class of *lower bound/upper bound* (L/U) *automata* and describes several (rather intuitive) observations that simplify the parametric model checking problem for PTAs in this class. Our results use the possibility to eliminate parameters in certain cases. This is a relevant issue, because the complexity of parametric model checking grows very fast in the number of parameters. Moreover, our observations yield some decidability results for L/U automata, where the corresponding problems are undecidable for general PTAs. The applicability of the results is illustrated by the verification of Fischer's algorithm.

4.1. Lower bound/upper bound automata

Informally, each parameter in an L/U automaton $\mathscr{A}$ occurs either as a lower bound in the invariants and guards of $\mathscr{A}$ or as an upper bound, but never as both. For instance, p is an upper bound parameter in $x - y < 2p$. Lower bound parameters are for instance q and q' in $y - x > q + 2q'$ ($\equiv x - y < -q - 2q'$) and in $x - y < 2p - q - 2q'$. A PTA containing both the guards $x - y \leqslant p - q$ and $z < q - p$ is not an L/U automaton.

Definition 4.1. A parameter $p_i \in P$ is said to *occur* in the linear expression $e = t_0 + t_1 \cdot p_1 + \cdots + t_n \cdot p_n$ if $t_i \neq 0$; p_i *occurs positively* in e if $t_i > 0$ and p_i *occurs negatively* in e if $t_i < 0$. A *lower bound parameter* of a PTA $\mathscr{A}$ is a parameter that only occurs negatively in the expressions of $\mathscr{A}$ and an *upper bound parameter* of $\mathscr{A}$ is a parameter that only occurs positively in $\mathscr{A}$. We call $\mathscr{A}$ a *lower bound/upper bound (L/U) automaton* if every parameter occurring in $\mathscr{A}$ is either a lower bound parameter or an upper bound parameter.

From now on, we work with a fixed set $L = \{l_1, \ldots, l_K\}$ of lower bound parameters and a fixed set $U = \{u_1, \ldots, u_M\}$ of upper bound parameters with $L \cap U = \emptyset$ and $L \cup U = P$. Furthermore, we consider, apart from parameter valuations, also *extended parameter valuations*. Intuitively, an extended parameter valuation is a parameter valuation with values in $\mathbb{R}^{\geqslant 0} \cap \{\infty\}$, rather than in $\mathbb{R}^{\geqslant 0}$. Extended parameter valuations are useful in certain cases to solve the verification problem (over non-extended valuations) stated in Section 2.3. Working with extended parameter valuations may cause the evaluation of an expression to be undefined. For example, the expression $e[v]$ is not defined for $e = p - q$ and $v(p) = v(q) = \infty$. We therefore require that an extended parameter valuation does not assign the value ∞ to both a lower bound parameter and an upper bound parameter. Then

we can easily extend notions $e[v]$, $(v, w) \models e$ and $\mathscr{A}[v]$ (defined in Section 2) to extended valuations. Here, we use the conventions that $0 \cdot \infty = 0$, that $x - y \prec \infty$ evaluates to true and $x - y \prec -\infty$ to false. In particular, we have $[\![\mathscr{A}]\!]_v = [\![\mathscr{A}[v]]\!]$ for extended valuations v and L/U automata $\mathscr{A}$. Moreover, we extend the orders $\sim$ to $\mathbb{R} \cup \{\infty\}$ in the usual way and we extend them to extended parameter valuations via point wise extension (i.e. $v \sim v'$ iff $v(p) \sim v'(p)$ for all $p \in P$). We denote an extended valuation of an L/U automaton by a pair (λ, μ), which equals the function λ on the lower bound parameters and μ on the upper bound parameters. We write 0 and ∞ for the functions assigning, respectively, 0 and ∞ to each parameter.

The following proposition is based on the fact that weakening the guards in $\mathscr{A}$ (i.e., decreasing the lower bounds and increasing the upper bounds) yields an LTS whose reachable states include those of $\mathscr{A}$. Dually, strengthening the guards in $\mathscr{A}$ (i.e. increasing the lower bounds and decreasing the upper bounds) yields an LTS whose reachable states are a subset of those of $\mathscr{A}$. The result crucially depends on the fact that state formulae (by definition) do not contain parameters. The usefulness of this property (and of several other properties in this section) lies in the fact that the satisfaction of a property for infinitely many extended parameter valuations (λ', μ') is reduced to its satisfaction for a single valuation (λ, μ).

Proposition 4.2. *Let $\mathscr{A}$ be an L/U automaton and ϕ a state formula. Then*
1. $[\![\mathscr{A}]\!]_{(\lambda, \mu)} \models \exists \Diamond \phi \iff \forall \lambda' \leq \lambda, \mu \leq \mu' : [\![\mathscr{A}]\!]_{(\lambda', \mu')} \models \exists \Diamond \phi$.
2. $[\![\mathscr{A}]\!]_{(\lambda, \mu)} \models \forall \Box \phi \iff \forall \lambda \leq \lambda', \mu' \leq \mu : [\![\mathscr{A}]\!]_{(\lambda', \mu')} \models \forall \Box \phi$.

Proof (*Sketch*). The "$\Longleftarrow$" parts of both statements are trivial. The crucial observation for both "$\Longrightarrow$" parts is the following. For all linear expressions e in $\mathscr{A}$ and all extended parameter valuations (λ, μ), (λ', μ') with $\lambda' \leqslant \lambda$ and $\mu \leqslant \mu'$, we have that $e[\lambda, \mu] \leqslant e[\lambda', \mu']$. Therefore, if $((\lambda, \mu), w) \models x - y \prec e$, then $((\lambda', \mu'), w) \models x - y \prec e$. $\square$

The following example illustrates how Proposition 4.2 can be used to eliminate parameters in L/U automata.

Example 4.3. The PTA in Fig. 7 is clearly an L/U automaton: *min* is a lower bound and *max* is an upper bound parameter. Location S_1 is reachable irrespective of the parameter values. By setting the parameter *min* to ∞ and *max* to 0, one checks with a non-parametric model checker that $\mathscr{A}[(\infty, 0)] \models \exists \Diamond S_1$. Then Proposition 4.2(1) (together with $[\![\mathscr{A}]\!]_v = [\![\mathscr{A}[v]]\!]$) yields that S_1 is reachable in $[\![\mathscr{A}]\!]_{(\lambda, \mu)}$ for all extended parameter valuations $0 \leq \lambda, \mu \leq \infty$.

Clearly, $[\![\mathscr{A}]\!]_{(\lambda, \mu)} \models \exists \Diamond S_2$ iff $\lambda(min) \leq \mu(max) \wedge \lambda(min) < \infty$. We will see in this running example how we can verify this property completely by non-parametric model checking. Henceforth, we construct the automaton $\mathscr{A}'$ from $\mathscr{A}$ by substituting the parameter *max* by the parameter *min* yielding an (non L/U) automaton with one parameter, *min*. The next example shows that $[\![\mathscr{A}']\!]_v \models \exists \Diamond S_2$ for all valuations v, which essentially means that $[\![\mathscr{A}]\!]_{(\lambda, \mu)} \models \exists \Diamond S_2$ for all λ, μ such that $\mu(max) = \lambda(min) < \infty$. From this fact, Proposition 4.2(4.2) concludes that $[\![\mathscr{A}]\!]_{(\lambda, \mu)} \models \exists \Diamond S_2$ for all λ, μ with $\lambda(min) \leq \mu(max)$ and $\lambda(min) < \infty$.

The question whether there exists a (non-extended) parameter valuation such that a given location q is reachable, is known as the *emptiness problem* for PTAs. In [5], it is shown that the emptiness problem is undecidable for PTAs with three clocks or more.

Fig. 7. Reducing parametric to non-parametric model checking.

The following proposition implies that we can solve the emptiness problem for an L/U automaton $\mathscr{A}$ by only considering the timed automaton $\mathscr{A}[(0, \infty)]$. Since reachability for timed automata is decidable ([4]), the emptiness problem is decidable for L/U automata. Then it follows that the dual problem is also decidable for L/U automata. This is the *universality problem* for invariance properties, asking whether an invariance property holds for all parameter valuations.

Proposition 4.4. *Let $\mathscr{A}$ be an L/U automaton and q a location of $\mathscr{A}$. Then $\mathscr{A}[(0, \infty)] \models \exists \Diamond q$ if and only if there exists a (non-extended) parameter valuation (λ, μ) such that $[\![\mathscr{A}]\!]_{(\lambda, \mu)} \models \exists \Diamond q$.*

Proof. The "only if" part is an immediate consequence of Proposition 4.2(1) and the fact that $[\![\mathscr{A}[(0, \infty)]]\!] = [\![\mathscr{A}]\!]_{(0, \infty)}$. For the "if" part, assume that α is a run of $[\![\mathscr{A}[(0, \infty)]]\!]$ that reaches the location q. Let T' be the smallest constant occurring in $\mathscr{A}$ and let T be the maximum clock value occurring in α. (More precisely, if $\alpha = s_0 a_1 s_1 a_2 \ldots a_N s_N$ and $s_i = (q_i, w_i)$, then $T = \max_{i \leq N, x \in X} w_i(x)$; T' compensates for negative constants t_0 in expressions e of $\mathscr{A}$.) Now, take $\lambda(l_j) = 0$ and $\mu(u_j) = T + |T'| + 1$. Let $i \leqslant N$ and $g = x - y \prec e$ be the invariant associated with a state s_i occurring in α or the guard associated with the i^{th} transition taken by α. One easily shows that, since $w_i(x) - w_i(y) \prec e[0, \infty]$, also $w_i(x) - w_i(y) \prec e[\lambda, \mu]$, that is $((\lambda, \mu), w_i) \models g$. Hence, α is a run of $[\![\mathscr{A}]\!]_{(\lambda, \mu)}$, so $[\![\mathscr{A}]\!]_{(\lambda, \mu)} \models \exists \Diamond q$. $\square$

Corollary 4.5. *The emptiness problem is decidable for L/U automata.*

Definition 4.6. A PTA $\mathscr{A}$ is *fully parametric* if clocks are only reset to 0 and every linear expression in $\mathscr{A}$ is of the form $t_1 \cdot p_1 + \cdots + t_n \cdot p_n$, where $t_i \in \mathbb{Z}$.

The following proposition is basically the observation in [4], that multiplication of each constant in a timed automaton and in a system property with the same positive factor preserves satisfaction.

Proposition 4.7. *Let $\mathscr{A}$ be fully parametric PTA. Then for all parameter valuations v and all system properties ψ*

$$[\![\mathscr{A}]\!]_v \models \psi \iff \forall t \in \mathbb{R}^{>0} : [\![\mathscr{A}]\!]_{t \cdot v} \models t \cdot \psi,$$

where $t \cdot v$ denotes the valuation $p \mapsto t \cdot v(p)$ and $t \cdot \psi$ the formula obtained from ψ by multiplying each number in ψ by t.

Proof. It is easy to see that for all $t \in \mathbb{R}^{>0}$, $\alpha = s_0 a_1 s_1 a_2 \ldots a_N s_N$ with $s_i = (q_i, w_i)$ is a run of $[\![\mathscr{A}]\!]_v$ if and only if $s'_0 a_1 s'_1 \cdots a_N s'_N$ is a run of $[\![\mathscr{A}]\!]_{t \cdot v}$, where $s'_i = (q_i, t \cdot w_i)$ and $t \cdot w_i$ denotes $x \mapsto t \cdot w_i(x)$. $\square$

Then for fully parametric PTAs with one parameter and system properties ψ without constants (except for 0), we have $[\![\mathscr{A}]\!]_v \models \psi$ for all valuations v of P if and only if both $\mathscr{A}[0] \models \psi$ and $\mathscr{A}[1] \models \psi$. The need for a separate treatment of the value 0 is illustrated by the (fully parametric) automaton with a single transition equipped with the guard $x < p$. The target location of the transition is reachable for any value of p, except for $p = 0$.

Corollary 4.8. *For a fully parametric PTA $\mathscr{A}$ with one parameter, a constraint set C and a property ψ without constants (except 0), it is decidable whether $\forall v \in [\![C]\!] : [\![\mathscr{A}]\!]_v \models \psi$.*

Example 4.9. The PTA $\mathscr{A}'$ mentioned in Example 4.3 is a fully parametric timed automaton and the property $\exists \diamondsuit S_2$ is without constants. We establish that $\mathscr{A}'[0] \models \exists \diamondsuit S_2$ and $\mathscr{A}'[1] \models \exists \diamondsuit S_2$. Then Proposition 4.7 implies that $\mathscr{A}'[v] \models \exists \diamondsuit S_2$ for all v. As shown in Example 4.3, this implies that $[\![\mathscr{A}]\!]_{(\lambda,\mu)} \models \exists \diamondsuit S_2$ for all λ, μ with $\lambda(min) = \mu(max) < \infty$.

In the running example, we would like to use the same methods as above to verify that $[\![\mathscr{A}]\!]_{(\lambda,\mu)} \not\models \exists \diamondsuit S_2$ if $\lambda(min) > \mu(max)$. However, we can not take $min = max$ in this case, since the bound in the constraint is a strict one. The following definition and results allows us to move the strictness of a constraint into the PTA.

Definition 4.10. Let $P' \subseteq P$ be a set of parameters. Define $\mathscr{A}^{\leqslant}_{P'}$ as the PTA obtained from $\mathscr{A}$ by replacing every inequality $x - y \leqslant e$ in $\mathscr{A}$ by a strict inequality $x - y < e$, provided that e contains at least one parameter from P'. Similarly, define $\mathscr{A}^{\leqslant}_{P'}$ as the PTA obtained from $\mathscr{A}$ by replacing every inequality $x - y < e$ by a non-strict inequality $x - y \leqslant e$, provided that e contains at least one parameter from P'. For $\prec = <, \leqslant$, write $\mathscr{A}^{\prec}$ for $\mathscr{A}^{\prec}_P$. Moreover, define $v \prec_{P'} v'$ by $v(p) \prec v'(p)$ if $p \in P'$ and $v(p) = v'(p)$ otherwise.

Proposition 4.11. *Let $\mathscr{A}$ be an L/U automaton. Then for all extended valuations (λ, μ) of $\mathscr{A}$*
1. $[\![\mathscr{A}^{\leqslant}]\!]_{(\lambda,\mu)} \models \exists \diamondsuit \phi \Rightarrow \forall \lambda' < \lambda, \mu < \mu' : [\![\mathscr{A}]\!]_{(\lambda',\mu')} \models \exists \diamondsuit \phi$.
2. $[\![\mathscr{A}^{<}]\!]_{(\lambda,\mu)} \models \forall \square \phi \iff \forall \lambda < \lambda', \mu' < \mu : [\![\mathscr{A}]\!]_{(\lambda',\mu')} \models \forall \square \phi$.

Proof. 1. Let (λ, μ) be an extended valuation and assume that $[\![\mathscr{A}^{\leqslant}]\!]_{(\lambda,\mu)} \models \exists \diamondsuit \phi$. Let e be a linear expression occurring in $\mathscr{A}$. Then we can write $e = t_0 + e_1 + e_2$, where $t_0 \in \mathbb{Z}$, e_1 is an expression over the upper bound parameters and e_2 an expression over the lower bound parameters. Then we have

$$\mu \leqslant \mu' \Rightarrow e_1[\mu] \leqslant e_1[\mu'],$$
$$\lambda' \leqslant \lambda \Rightarrow e_2[\lambda'] \leqslant e_2[\lambda],$$
$$\lambda' \leqslant \lambda, \quad \mu \leqslant \mu' \Rightarrow e[(\lambda, \mu)] \leqslant e[(\lambda', \mu')].$$

If there is at least one parameter occurring, respectively, in e_1 or e_2 then, respectively,

$$\mu < \mu' \Rightarrow e_1[\mu] < e_1[\mu'],$$
$$\lambda' < \lambda \Rightarrow e_2[\lambda] < e_2[\lambda'].$$

Thus, if there is at least one parameter occurring in e, then

$$\lambda' < \lambda, \mu < \mu' \Rightarrow e[(\lambda, \mu)] < e[(\lambda', \mu')].$$

Now, let $g \equiv x - y \prec e$ be a simple guard occurring in $\mathscr{A}^{\leqslant}$ and let $g' \equiv x - y \prec' e$ be the corresponding guard in $\mathscr{A}$. Assume that $(w, (\lambda, \mu)) \models g$, i.e. $w(x) - w(y) \prec e[(\lambda, \mu)]$. We show that $(w, (\lambda, \mu)) \models g'$. We distinguish two cases.

Case 1. There exists a parameter occurring in e. Then $w(x) - w(y) \prec e[(\lambda, \mu)] < e[(\lambda', \mu')]$. Then certainly $((\lambda, \mu), w) \models g' \equiv x - y \prec' e$.

Case 2. The expression e does not contain any parameter. Then $g' \equiv g$ and hence $((\lambda, \mu), w) \models g'$.

It easily follows that every run of $[\![\mathscr{A}^{\leqslant}]\!]_{(\lambda, \mu)}$ is also a run of $[\![\mathscr{A}]\!]_{(\lambda', \mu')}$. Thus, if a state satisfying ϕ is reachable in $[\![\mathscr{A}^{\leqslant}]\!]_{(\lambda, \mu)}$ then it is also reachable in $[\![\mathscr{A}]\!]_{(\lambda', \mu')}$.

2. ($\Rightarrow$) This follows from statement (1) of this proposition: assume that $[\![\mathscr{A}^{<}]\!]_{(\lambda, \mu)} \models \forall \Box \phi$ and let λ', μ' be such that $\lambda < \lambda', \mu' < \mu$. Since $[\![\mathscr{A}^{<}]\!]_{(\lambda, \mu)} \nvDash \exists \Diamond \neg \phi$, we have

$$\neg \forall \lambda'' < \lambda', \mu' < \mu'' : [\![\mathscr{A}^{<}]\!]_{(\lambda'', \mu'')} \models \exists \Diamond \neg \phi.$$

Then contraposition of statement (1) together with $(\mathscr{A}^{<})^{\leqslant} = \mathscr{A}^{\leqslant}$ yields $[\![\mathscr{A}^{\leqslant}]\!]_{(\lambda', \mu')} \nvDash \exists \Diamond \neg \phi$. As $\mathscr{A}$ imposes stronger bounds than $\mathscr{A}^{\leqslant}$, also $[\![\mathscr{A}]\!]_{(\lambda', \mu')} \nvDash \exists \Diamond \neg \phi$, i.e. $[\![\mathscr{A}]\!]_{(\lambda', \mu')} \models \forall \Box \phi$.

2. ($\Longleftarrow$) Let (λ, μ) be an extended valuation. Assume that $[\![\mathscr{A}]\!]_{(\lambda'', \mu'')} \models \forall \Box \phi$ for all $\lambda'' > \lambda, \mu'' < \mu$ and that $\alpha = s_0 a_1 s_1 a_2 \ldots a_N s_N$ is a run of $[\![\mathscr{A}^{<}]\!]_{(\lambda, \mu)}$. We have to show that $s_N \models \phi$. Below, we construct $\lambda' > \lambda$ and $\mu' < \mu$ such that α is also a run of $[\![\mathscr{A}]\!]_{(\lambda', \mu')}$. Then we are done: since $[\![\mathscr{A}]\!]_{(\lambda', \mu')} \models \forall \Box \phi, s_N \models \phi$.

We use the following notation. For the run $\alpha = s_0 a_1 s_1 a_2 \ldots a_N s_N$ of $\mathscr{A}^{<}$, we write $s_k = (q_k, w_k), I(q_k) = \bigwedge_{i=0}^{J} I_{ik}, I_{ik} = x_i \prec_{ik} E_{ik}$, where J is the number of clocks in $\mathscr{A}$. As α is a run, we have that for all $k, 0 \leqslant k < N$, either $a_{k+1} \in \mathbb{R}^{\geqslant 0}$ or there exists a transition $q_k \xrightarrow{g_k, a_{k+1}, r_{k+1}} q_{k+1}$ in $\mathscr{A}^{<}$. We write the guard on this transition as $g_k = \wedge_{i, j \leqslant J} g_{ijk}$ with $g_{ijk} = x_i - x_j \prec_{ijk} e_{ijk}$. If $a_k \in \mathbb{R}^{\geqslant 0}$, then we put $\prec_{ijk} = <$ and $e_{ijk} = \infty$ for all $i, j \leqslant J$.

If for all i, j, k neither the guard g_{ijk} nor the invariant I_{ik} contains a parameter, then we can take λ' and μ' arbitrarily and we have that α is a run of $[\![\mathscr{A}]\!]_{(\lambda', \mu')}$. Therefore, assume that at least one of the guards g_{ijk} or invariants I_{ik} contains a parameter. Then, by definition of $\mathscr{A}^{<}$, this guard or invariant contains a strict bound. In this case, we construct $\lambda' > \lambda$ and $\mu' < \mu$ such that $w_k(x - y) < e[(\lambda', \mu')] < e[(\lambda, \mu)]$ for all $k < N$ and all expressions e occurring in the invariants I_{ik} or guard g_{ijk}. Informally, we use the minimum "distance" $e[(\lambda, \mu)] - w_k(x - y)$ occurring in α to slightly increase the lower bounds and slightly decrease the upper bounds yielding $\lambda < \lambda'$ and $\mu < \mu'$.

Formally, let

$$T_0 = \min_{k \leqslant N, i \leqslant J} \{E_{ik}[(\lambda, \mu)] - w_k(x_i) \mid \prec_{ik} = <\},$$
$$T_1 = \min_{k \leqslant N, i, j \leqslant J} \{e_{ijk}[(\lambda, \mu)] - (w_k(x_i) - w_k(x_j)) \mid \prec_{ijk} = <\},$$
$$0 < T < \min\{T_0, T_1\},$$

with the convention that $\min \emptyset = \infty$. At least one of the inequalities $\prec_{ijk}$ or $\prec_{ik}$ is strict, since at least one of the guards or invariants contains a parameter. Hence, either $T_0 < \infty$ or $T_1 < \infty$. Since $((\lambda, \mu), w_k) \models I_{ik} \wedge g_{ijk}$, we have that $T_0 > 0$ and $T_1 > 0$. Hence,

$0 < \min\{T_0, T_1\} < \infty$ and the requested T exists. The crucial property is that if $g_{ijk} \equiv x_i - x_j < e_{ijk}$ contains a parameter, then

$$w_k(x_i) - w_k(x_j) < e_{ijk}[(\lambda, \mu)] - T \tag{7}$$

and, similarly, if $I_{ik} \equiv x_i < E_{ik}$ contains a parameter, then $w_k(x_i) < E_{ik}[(\lambda, \mu)] - T$.

Now, we can distribute the value T over all parameters to obtain larger values for the lower bounds and smaller ones for the upper bounds. Let T' be the sum of the constants that appear in front of a parameter in one of the guards g_{ijk} or the invariants I_{ik}, i.e.

$$T' = \sum_{k \leqslant N, i \leqslant J} \text{sum\_of\_const}(E_{ik}) + \sum_{k \leqslant N, i, j \leqslant J} \text{sum\_of\_const}(e_{ijk}),$$

where $\text{sum\_of\_const}(t_0 + t_1 \cdot p_1 + \cdots + t_n \cdot p_n) = |t_1| + \cdots + |t_n|$. Since at least one of the guards or invariants contains a parameter, we have $T' > 0$.

Now, take $\lambda' = \lambda + T/T'$ and $\mu' = \mu - T/T'$. Let $i, j \leqslant J, k \leqslant N$ and consider the guard $\overline{g_{ijk}} \equiv x_i - x_j \prec^{\mathscr{A}}_{ijk} e_{ijk}$ in $\mathscr{A}$, which corresponds to the guard $g_{ijk} \equiv x_i - x_j \prec_{ijk} e_{ijk}$ in $\mathscr{A}^<$. We prove below that $((\lambda', \mu'), w_k) \models \overline{g_{ijk}}$. In a similar way, one can show that $((\lambda', \mu'), w_k) \models \overline{I_{ik}}$ for the invariant corresponding to I_{ik}. Then, α is a run of $[\![\mathscr{A}]\!]_{(\lambda', \mu')}$ and we are done.

Case 1. The expression $\overline{g_{ijk}}$ does not contain any parameter. Since $((\lambda, \mu), w_k) \models g_{ijk}$, $((\lambda', \mu'), w_k) \models \overline{g_{ijk}}$.

Case 2: There exists a parameter occurring in $\overline{g_{ijk}}$. Then $g_{ijk} \equiv x_i - x_j < e_{ijk}$ and we can write $e_{ijk} = t_0 + t_1 \cdot u_1 + \cdots + t_M \cdot u_M - t'_1 \cdot l_1 - \cdots - t'_K \cdot l_K$, with $t_i \geqslant 0, t'_i \geqslant 0$ for $i > 0$. Then

$$
\begin{aligned}
e_{ijk}[(\lambda', \mu')] &= \left(t_0 + \sum_{h=1}^{M} t_h \cdot u_h - \sum_{h=1}^{K} t'_h \cdot l_h \right) \left[\left(\lambda + \frac{T}{T'}, \mu - \frac{T}{T'} \right) \right] \\
&= t_0 + \sum_{h=1}^{M} t_h \cdot \left(\mu_h - \frac{T}{T'} \right) - \sum_{h=1}^{K} t'_h \cdot \left(\lambda_h + \frac{T}{T'} \right) \\
&= t_0 + \sum_{h=1}^{M} t_h \cdot \mu_h - \sum_{h=1}^{K} t'_h \cdot \lambda_h - \frac{T}{T'} \cdot \left(\sum_{h=1}^{M} t_h + \sum_{h=1}^{K} t'_h \right) \\
&\geqslant e_{ijk}[(\lambda, \mu)] - \frac{T}{T'} \cdot T' \quad \text{(by 7)} \\
&> w_k(x_i) - w_k(x_j).
\end{aligned}
$$

Thus, $((\lambda', \mu'), w_k) \models x_i - x_j < e_{ijk}$ and then also $((\lambda', \mu'), w_k) \models x_i - x_j \prec^{\mathscr{A}}_{ijk} e_{ijk}$. $\quad\square$

The previous result concerns the automaton that is obtained when *all* the strict inequalities in guards and invariants with parameters are changed into non-strict ones (or the other way around). Sometimes, we want to "toggle" only some of the inequalities. Then the following result can be applied.

Corollary 4.12. *Let $\mathscr{A}$ be an L/U automaton and $P' \subseteq P$.*

1. $[\![\mathscr{A}^{\leqslant}_{P'}]\!]_{(\lambda, \mu)} \models \exists\Diamond\phi \Rightarrow \forall \lambda' <_{P'} \lambda, \mu <_{P'} \mu' : [\![\mathscr{A}]\!]_{(\lambda', \mu')} \models \exists\Diamond\phi.$
2. $[\![\mathscr{A}^{<}_{P'}]\!]_{(\lambda, \mu)} \models \forall\Box\phi \iff \forall \lambda <_{P'} \lambda', \mu' <_{P'} \mu : [\![\mathscr{A}]\!]_{(\lambda', \mu')} \models \forall\Box\phi.$

Fig. 8. The converse of Proposition 4.11(1) does not hold.

Proof. Let (λ, μ) be an extended valuation. Let $\mathscr{A}_0$ be the automaton obtained from $\mathscr{A}$ by substituting p by $(\lambda, \mu)(p)$ for every $p \notin P'$. Then $[\![\mathscr{A}_{P'}^{\leqslant}]\!]_{(\lambda,\mu)} = [\![\mathscr{A}_0^{<}]\!]_{(\lambda,\mu)}$ and $[\![\mathscr{A}_{P'}^{\leqslant}]\!]_{(\lambda,\mu)} = [\![\mathscr{A}_0^{\leqslant}]\!]_{(\lambda,\mu)}$. Now the result follows by applying Proposition 4.11 to $\mathscr{A}_0$. $\square$

The following example shows that the converse of Proposition 4.11(4.11) does not hold.

Example 4.13. Consider the automaton $\mathscr{A}$ in Fig. 8. Recall that the clocks x and y are initially 0. Then $\mathscr{A} = \mathscr{A}^{\leqslant}$ and the location q is reachable if $max > 0$ but not if $max = 0$. This is so because if $max = 0$, then clock y is never augmented. Thus, $\forall \lambda' < 0, 0 < \mu'$: $[\![\mathscr{A}]\!]_{(\lambda',\mu')} \models \exists \Diamond \phi$, but not $[\![\mathscr{A}^{\leqslant}]\!]_{(0,0)} \models \exists \Diamond \phi$.

We believe that the class of L/U automata can be very useful in practice. Several examples known from the literature fall into this class, or can be modelled slightly differently to achieve this. We mention the root contention protocol [17], Fischer's mutual exclusion algorithm [19], the (toy) rail road crossing example from [5], the bounded retransmission protocol (when considering fixed values for the integer variables) and the biphase mark protocol (with minor adaptations). Moreover, the time constrained automata models of [24,23] can be encoded straightforwardly into L/U automata.

We expect that quite a few other distributed algorithms and protocols can be modelled with L/U automata, since it is natural that the duration of an event (such as the communication delay in a channel, the computation time needed to produce a result, the time required to open the gate in a rail road crossing) lies between a lower bound and an upper bound. These bounds are often parameters of the system.

The next section and Section 5 show that the techniques discussed in this section to eliminate parameters in L/U models reduce the verification effort significantly and possibly lead to a completely non-parametric model.

4.2. Verification of Fischer's mutual exclusion protocol

In this section, we apply the results from the previous section to verify the Fischer protocol described in Section 2.4. We establish the sufficiency of the protocol constraints completely by non-parametric model checking and the necessity of the constraints by eliminating three of the four parameters.

We also tried to use the prototype to verify the protocol model without any substitutions or changing of bounds, but this did not terminate within 20 h. Since we observed that the constraint lists of the states explored kept on growing, we suspected that this experiment would not terminate at all. (Recall that parametric verification is undecidable.) Verification of the reduced models took only 2 s.

Now, consider the Fischer protocol model from Section 2.4 again. In this section, we analyze a system $\mathcal{A}$ consisting of two parallel processes P_1 and P_2. It is clear that $\mathcal{A}$ is a fully parametric L/U automaton: *min_rw* and *min_delay* are lower bound parameters and *max_rw* and *max_delay* upper bound parameters.

The mutual exclusion property is expressed by the formula $\Phi_{ME} \equiv \forall\Box\neg(P_1.cs \wedge P_2.cs)$. In Section 2.4 we claimed that, when assuming the basic constraints $B_{ME} \equiv 0 \leqslant min\_rw < max\_rw \wedge 0 \leqslant min\_delay < max\_delay$, mutual exclusion is guaranteed if and only if $C_{ME} \equiv max\_rw \leqslant min\_delay$. To establish this formally, we will prove that $v \models B_{ME} \Rightarrow (\llbracket \mathcal{A} \rrbracket_v \models \Phi_{ME} \iff v \models C_{ME})$, for all valuations v.

4.2.1. Sufficiency of the constraints

We show that the constraints assure mutual exclusion, that is

$$\text{if } v \models C_{ME} \wedge B_{ME}, \text{ then } \llbracket \mathcal{A} \rrbracket_v \models \Phi_{ME}.$$

We perform the substitution

$$min\_rw \mapsto 0, \ max\_delay \mapsto \infty, \ min\_delay \mapsto max\_rw$$

to obtain a fully parametric automaton $\mathcal{A}'$ with one parameter, *max_rw*. We have established by non-parametric model checking that $\mathcal{A}'[0] \models \Phi_{ME}$ and $\mathcal{A}'[1] \models \Phi_{ME}$. Now Proposition 4.7 yields that $\llbracket \mathcal{A}' \rrbracket_v \models \Phi_{ME}$ for all valuations v (where only the value of *max_rw* matters). This means that $\llbracket \mathcal{A} \rrbracket_v \models \Phi_{ME}$ if $v(min\_rw) = 0$, $v(max\_rw) = v(min\_delay)$ and $v(max\_delay) = \infty$. Then Proposition 4.2(2) yields that the invariance property Φ_{ME} also holds if we increase the lower bound parameters *min_rw* and *min_delay* and if we decrease the upper bound parameter *max_rw*. More precisely, Proposition 4.2(2) implies that $\llbracket \mathcal{A} \rrbracket_v \models \Phi_{ME}$ for all v with $0 \leqslant v(min\_rw)$, $v(max\_rw) \leqslant v(min\_delay)$ and $v(max\_delay) \leqslant \infty$. Then, in particular, $\llbracket \mathcal{A} \rrbracket_v \models \Phi_{ME}$ if $v \models C_{ME} \wedge B_{ME}$.

4.2.2. Necessity of the constraints

We show that

$$v \models B_{ME} \wedge \neg C_{ME} \Rightarrow \llbracket \mathcal{A} \rrbracket_v \models \neg\Phi_{ME},$$

i.e. that if $v \models min\_rw < max\_rw \wedge min\_delay < max\_delay \wedge min\_delay < max\_rw$, then $\mathcal{A}[v] \models \neg\Phi_{ME} \equiv \exists\Diamond(P_1.cs \wedge P_2.cs)$. We consider the automaton $\mathcal{A}^{\leqslant}$ and proceed in two steps.

Step 1. Let v_0 be the valuation $v_0(min\_delay) = v_0(max\_delay) = 0$ and $v_0(min\_rw) = v_0(max\_delay) = 1$. By non-parametric model checking we have established that

$$\mathcal{A}^{\leqslant}[0] \models \neg\Phi_{ME}, \tag{8}$$

$$\mathcal{A}^{\leqslant}[v_0] \models \neg\Phi_{ME}. \tag{9}$$

We show that it follows that for all v

$$v \models 0 = min\_delay = max\_delay \leqslant min\_rw = max\_rw \Rightarrow \mathcal{A}^{\leqslant}[v] \models \neg\Phi_{ME}. \tag{10}$$

Assume $v \models 0 = min\_delay = max\_delay \leqslant min\_rw = max\_rw$. Consider $t = v(min\_rw)$. If $v(min\_rw) = 0$, then (8) shows that $\llbracket \mathcal{A}^{\leqslant} \rrbracket_v \models \neg\Phi_{ME}$. Therefore, assume $v(min\_rw) > 0$ and consider $\frac{v}{t} \equiv \lambda x.v(x)/t$. It is not difficult to see that

$$\frac{v}{t} \models 0 = min\_delay = max\_delay \leqslant min\_rw = max\_rw = 1.$$

Therefore, (9) yields $\llbracket \mathscr{A}^{\leqslant} \rrbracket_{v/t} \models \neg\Phi_{\mathrm{ME}}$. Since $\mathscr{A}^{\leqslant}$ is a fully parametric PTA, Proposition 4.7 yields that $\llbracket \mathscr{A}^{\leqslant} \rrbracket_{v} \models \neg\Phi_{\mathrm{ME}}$.

Step 2. Let $\mathscr{A}'$ be the automaton that is constructed from $\mathscr{A}^{\leqslant}$ by performing the following substitution *min_delay* $\mapsto$ 1, *max_delay* $\mapsto$ 1, *min_rw* $\mapsto$ *max_rw*. By parametric model checking we have established

$$v \models 1 \leqslant \mathit{max\_rw} \;\Rightarrow\; \llbracket \mathscr{A}' \rrbracket_{v} \models \neg\Phi_{\mathrm{ME}}. \tag{11}$$

This means that if

$$v \models \mathit{min\_delay} = \mathit{max\_delay} = 1 \leqslant \mathit{min\_rw} = \mathit{max\_rw} \;\Rightarrow\; \llbracket \mathscr{A}^{\leqslant} \rrbracket_{v} \models \neg\Phi_{\mathrm{ME}}.$$

By a argument similar to the one we used to prove (10), (where now the case v (*min_delay*) $= 0$ is covered by Equation (10) in Step 1), we can use Proposition 4.7 to show that

$$v \models \mathit{min\_delay} = \mathit{max\_delay} \leqslant \mathit{min\_rw} = \mathit{max\_rw} \;\Rightarrow\; \llbracket \mathscr{A}^{\leqslant} \rrbracket_{v} \models \neg\Phi_{\mathrm{ME}}.$$

Now, Proposition 4.2(1) yields that the reachability property $\neg\Phi_{\mathrm{ME}}$ also holds if the values for the lower bounds are decreased and the values for the upper bounds are increased. Note that we may increase *max_delay* as much as we want; $v(\mathit{max\_delay})$ may be larger than $v(\mathit{min\_rw})$. Thus we have

$$v \models \mathit{min\_rw} \leqslant \mathit{max\_rw} \wedge \mathit{min\_delay} \leqslant \mathit{max\_delay} \wedge \mathit{min\_delay} \leqslant \mathit{max\_rw}$$
$$\Rightarrow\; \llbracket \mathscr{A}^{\leqslant} \rrbracket_{v} \models \neg\Phi_{\mathrm{ME}}$$

and then Proposition 4.11 yields that

$$v \models \mathit{min\_rw} < \mathit{max\_rw} \wedge \mathit{min\_delay} < \mathit{max\_delay} \wedge \mathit{min\_delay} < \mathit{max\_rw}$$
$$\Rightarrow\; \llbracket \mathscr{A} \rrbracket_{v} \models \neg\Phi_{\mathrm{ME}}.$$

We have checked the result formulated in Eq. (11) with our prototype implementation. The experiment was performed on a SPARC Ultra in 2 s CPU time and 7.7 Mb of memory.

The substitutions and techniques used in this verification to eliminate parameters are ad hoc. Probably, more general strategies can be applied in this case, because the constraints are L/U-like (i.e., they can be written in the form $e \prec 0$ such that every p occurring negatively in e is a lower bound parameter and every p occurring positively in e is an upper bound parameter).

5. Experiments

5.1. A prototype extension of UPPAAL

Based on the theory described in Section 3, we have built a prototype extension of UPPAAL. In this section, we report on the results of experimenting with this tool.

Our prototype allows the user to give some initial constraints on the parameters. This is particularly useful when explorations cannot be finished due to lack of memory or time resources, or because a non-converging series of constraint sets is being generated. Often, partial results can be derived by observing the constraint sets that are generated during the exploration. Based on partial results, the actual solution constraints can be established in many cases. These partial results can then be checked by using an initial set of constraints.

model from	initial constraints	reduced	property	UPPAAL	time	memory
[27]	part of solution	no	safety	param	18 h	339 Mb
[27]	solution	no	safety	param	2.9 h	185 Mb
[27]	-	yes	safety	std	1 s	800 Kb
[25]	no	no	safety	param	40 m	38 Mb
[25]	solution	no	safety	param	1.6 m	36 Mb
[25]	solution	partly	safety	param	11 s	13 Mb
[25]	-	completely	safety	std	1 s	800 Kb
[25]	no	no	refinement	param	24 d	1.9 Gb
[25]	part of solution	no	refinement	param	8 d	1.4 Gb
[25]	solution	no	refinement	param	2.6 h	308 Mb

Fig. 9. Experimental results for the root contention protocol.

5.2. The root contention protocol

Description. The root contention protocol is part of a leader election protocol in the physical layer of the IEEE 1394 standard (FireWire/i-Link), which is used to break symmetry between two nodes contending to be the root of a tree, spanned in the network topology. The protocol consists of first drawing a random number (0 or 1), then waiting for some time according to the result drawn, followed by the sending of a message to the contending neighbor. This is repeated by both nodes until one of them receives a message before sending one, at which point the root is appointed.

Parametric approach. We use the UPPAAL models of [25,27], turn the constants used into parameters, and experiment with our prototype implementation (see Fig. 9 for results[5]). In both models, there are five constants, all of which are parameters in our experiments. The *delay* constant indicates the maximum delay of signals sent between the two contending nodes. The *rc_fast_min* and *rc_fast_max* constants give the lower and upper bound to the waiting time of a node that has drawn 1. Similarly, the *rc_slow_min* and *rc_slow_max* constants give the bounds when 0 has been drawn. It is reasonable to assume that initially, the constraints $rc\_fast\_min \leqslant rc\_fast\_max \leqslant rc\_slow\_min \leqslant rc\_slow\_max$ hold for each experiment.

We have checked for safety with the following property:

$$\forall\square \cdot (\neg(\text{Node}_1 \cdot \text{root} \wedge \text{Node}_2 \cdot \text{root}) \wedge \neg(\text{Node}_1 \cdot \text{child} \wedge \text{Node}_2 \cdot \text{child})).$$

Safety for [27]. The model in [27] consists of eight communicating processes, varying from three locations with six transitions to nine locations with 12 transitions, and has 4

---

[5] All experiments were performed on a 366 MHz Celeron, except the first experiment of safety for [27] and [25], and all the refinement experiments. These were performed on a 333 MHz SPARC Ultra Enterprise.

clocks in total. It is shown in [27], that the safety property holds (through a refinement relation), if the parameters obey the following relation: *delay < rc_fast_min*. We have checked that the error states, expressed in the safety property, are indeed unreachable when this parameter constraint is met. If we give no initial constraints, our experiments do not terminate. If we loosen the solution constraint to *delay ⩽ rc_fast_min*, we are able to check that no error states are reachable. In fact, it is argued in Remark 2 in [27], that the mentioned constraint is not *needed* for the correctness of the protocol. Rather than checking this on the parametric model without any initial constraints, which is a large task, we experiment with a non-parametric version of the model *without any timing constraints*. It turns out that this model satisfies the safety property, hence we deduce that the parametric model, in which guards and invariants have been added, satisfies the safety property for *any* valuation of the parameters.

Safety for [25]. A different model of the root contention protocol is proposed in [25], in which it is shown that the relation between the parameters for the safety property to hold, should obey: $2 * delay < rc\_fast\_min$. In fact, the model satisfies the safety property already when *delay < rc_fast_min*, but the stronger constraint is needed for proper behavior of the connecting wires. This model also consists of 8 communicating processes, varying from three locations with six transitions to 16 locations with 28 transitions, and has 6 clocks in total. The necessity and sufficiency of these constraints is shown in [25] by applying standard UPPAAL to several valuations for the parameters, and presented as an experimental result.

We have checked that the error states, expressed in the safety property, are indeed unreachable when either of these parameter constraints are met. We have also experimented without these initial constraints in an effort to generate constraints. This experiment terminates with a number of reachable error states. The union of the constraint sets of these states can be rewritten to the constraint *delay ⩾ rc_fast_min*.

Safety for [25] with L/U automata. Since the model used for safety is a L/U automaton, we can experiment with Proposition 4.2, as follows. We show that our invariant property is satisfied by a more general model of root contention, and deduce with part 2 of Proposition 4.2 that it holds for the constraints we are after. We first identify the sets $L = \{rc\_fast\_min, rc\_slow\_min\}$ and $U = \{delay, rc\_fast\_max, rc\_slow\_max\}$. We substitute infinity for both *rc_fast_max* and *rc_slow_max*, *rc_fast_min* for *rc_slow_min*. The new model, together with either the initial constraint *delay < rc_fast_min* or with $2 * delay < rc\_fast\_min$, satisfies the invariant property. This allows us to conclude that the original model satisfies the invariant property for any valuation of the parameters where *rc_fast_min ⩽ rc_slow_min*, and the given initial constraint are satisfied. This includes the special case *rc_fast_min ⩽ rc_fast_max ⩽ rc_slow_min ⩽ rc_slow_max*.

We can do even better by applying Proposition 4.11, if we first change each guard or invariant for *delay* to a *strict* version, and then substitute infinity for both *rc_fast_max* and *rc_slow_max*, and *rc_fast_min* for both *delay* and *rc_slow_min*. Now we have a model with only one parameter and no constants, which we can verify non-parametrically with standard UPPAAL, for two valuations of the parameter *rc_fast_min*, namely 0 and a nonzero value. The invariant property is satisfied, hence, by Proposition 4.7, we can deduce that it holds for all valuations of *rc_fast_min*, hence the original model satisfies the invariant property for any valuation of the parameters where *rc_fast_min ⩽ rc_slow_min*, and *delay < rc_fast_min*. Likewise, we can substitute *rc_fast_min*/2 for *delay*, and derive the other constraint. As can be seen in Fig. 9, the speed-up in terms of memory and time is drastic.

Finally, we can combine the results for initial constraints $delay < rc\_fast\_min$ and $delay = rc\_fast\_min$ with the fact that our model is a L/U automaton, and derive the *necessity* of constraint $delay < rc\_fast\_min$, as follows. Suppose that a parameter valuation for $delay$ and rc_fast_min exists, such that (1) the safety property holds, but (2) the constraint $delay < rc\_fast\_min$ is not satisfied. Assume this valuation assigns d to $delay$ and r to rc_fast_min. By our results, we know that $d \neq r$, so $d > r$. We now apply Proposition 4.2, and deduce that for each parameter valuation that assigns a value to upper bound parameter $delay$ which is *smaller* than d, and a value to lower bound parameter rc_fast_min which is *larger* than r, the safety property must hold. This includes valuations that satisfy constraint $delay = rc\_fast\_min$, which contradicts our results. We conclude that only for parameter valuations that satisfy constraint $delay < rc\_fast\_min$, the safety property holds.

Refinement for [25]. In [25], it is also shown that a refinement relation between the model of the most detailed level, and a model which is a bit more abstract, holds when the following relations are obeyed: $2 * delay < rc\_fast\_min$, and $2 * delay < rc\_slow\_min - rc\_fast\_max$. The refinement relation is such that it preserves both safety and liveness properties for the root contention protocol (which is proved in [25]). Again, the necessity and suffiency of the constraints is shown by experimenting with standard UPPAAL for several valuations for the parameters, and presented as an experimental result. Here, the most detailed model is put in parallel with a test automaton version of the more abstract model, and with a forward reachability exploration it is checked whether error states are reachable. If this is not the case, the refinement relation holds. This model consists of six communicating processes, varying now from four locations with five transitions to 11 locations with 87 transitions, and has 7 clocks in total.

We have checked for a completely parametric version of the system with the detailed model and the test automaton of the more abstract model, that error states in the test automaton are unreachable (i.e., the refinement relation holds), given both constraints initially. We have also experimented without these initial constraints in an effort to generate them. If we give no initial constraints, then the prototype consumes a huge amount of time and memory, but terminates successfully. It yields a number of reachable error states. The union of the constraint sets of these states can be rewritten to the constraint $2 * delay \geqslant rc\_fast\_min \vee 2 * delay \geqslant rc\_slow\_min - rc\_fast\_max$. This means that the refinement exists if $2 * delay < rc\_fast\_min$, and $2 * delay < rc\_slow\_min - rc\_fast\_max$.

When given one initial constraint: $delay \leqslant rc\_fast\_min$, this experiment terminates also successfully with a number of reachable error states. Again, the union of the constraint sets of these states can be rewritten to the constraint $2 * delay \geqslant rc\_fast\_min \vee 2 * delay \geqslant rc\_slow\_min - rc\_fast\_max$.

Since the models for refinement use constraints that fall outside the scope of L/U automata, we cannot apply Proposition 4.11 here.

5.3. The bounded retransmission protocol

Description. This protocol was designed by Philips for communication between remote controls and audio/video/TV equipment. It is a slight alteration of the well-known alternating bit protocol, to which timing requirements and a bound on the retry mechanism have been added. In [14] constraints for the correctness of the protocol are derived by hand, and some instances are checked using UPPAAL. Based on the models in [14], an automatic parametric analysis is performed in [1], however, no further results are given.

model from	initial constraints	property	UPPAAL	time	memory
[14]	yes	safety1	param	1.3 m	34 Mb
[14]	no	safety2	param	11 m	180 Mb
[14]	yes	safety2	param	3.5 m	64 Mb

Fig. 10. Experimental results for the bounded retransmission protocol.

Parametric approach. For our analysis, we use the timed automata models from [14]. These models typically consist of seven communicating processes, varying from two locations with four transitions to six locations with 54 transitions, and has 5 clocks and 9 non-clock variables in total. In [14] three different constraints are presented based on three properties which are needed to satisfy the safety specification of the protocol. We are only able to check two of these since one of the properties contains a parameter which our prototype version of UPPAAL is not able to handle yet (Fig. 10).

One of the constraints derived in [14] is that $TR \geqslant 2 \cdot MAX \cdot T_1 + 3 \cdot TD$, where TR is the timeout of the receiver, T_1 is the timeout of the sender, MAX is the number of resends made by the sender, and TD is the delay of the channel. This constraint is needed to ensure that the receiver does not time out prematurely before the sender has decided to abort transmission. The sender has a parameter $SYNC$ which decides for how long the sender waits until it expects that the receiver has realized a send error and reacted to it. In our parametric analysis we used TR and $SYNC$ as parameters and instantiated the others to fixed values. Using our prototype we did derive the expected constraint $TR \geqslant 2 \cdot MAX \cdot T_1 + 3 \cdot TD$. However, we also derived the additional constraint $TR - 2 \leqslant SYNC$ which was not stated in [14] for this property. The necessity of this constraint was verified by trying models with different fixed values for the parameters. The full set of constraints derived in [14] includes a constraint $TR \geqslant SYNC$ which is based on the property we cannot check. Therefore the error we have encountered is only present in an intermediate result, the complete set of constraints derived is correct. The authors of [14] have acknowledged the error and provided an adjusted model of the protocol, for which the additional constraint is not necessary.

The last constraint derived in [14] arises from checking that the sender and receiver are not sending messages too fast for the channel to handle. In this model we treat T_1 as a parameter and derive the constraint $T_1 > 2 \cdot TD$ which is the same as is derived in [14].

5.4. Other experiments

We have experimented with parametric versions of several models from the standard UPPAAL distribution, namely Fischer's mutual exclusion protocol, a train gate controller, and a car gear box controller.

In the case of Fischer's protocol (which is the version of the standard UPPAAL distribution, and not the one discussed in the rest of this paper), we parameterized a model with two processes, by turning the bound on the period the processes wait, before entering the critical section, into a parameter. We were able to generate the constraints that ensure the mutual exclusion within 2 s of CPU time on a 266 MHz Pentium MMX. Using these constraints as initial constraints and checking that now indeed the mutual exclusion is guaranteed, is done even faster. Fischer's protocol with two processes was also checked in [1], which took about 3 min.

5.5. Discussion

Our prototype handles parametric versions of bench-mark timed automata rather well. In some cases, the prototype will not generate a converging series of constraints, but in all cases we were able to get successful termination when applying (conjectures of) solution constraints as initial constraints in the exploration. The amount of time and memory used is then in many cases quite reasonable.

From our results it is not easy to draw clear-cut conclusions about the type of parametric model, for which our prototype can successfully generate constraints. It seems obvious from the case studies that the more complicated the model, the larger the effort in memory and time consumption. So it is worthwhile to have small, simple models. However, the danger of non-termination is most present in models which have a lot of behavioral freedom. The most promising direction, therefore, will be to experiment with conjectured solution constraints, and to combine this with the techniques for L/U automata.

6. Conclusions

This paper reports on a parametric extension to the model checker UPPAAL. This tool is capable of generating parameter constraints that are necessary and sufficient for a reachability or invariant property to hold for a linear parametric timed automaton. The semantics of the algorithms underlying the tool is given in clean SOS-style rules. Although the work [5] shows that parameter synthesis is undecidable in general, our prototype implementation terminates on many practical verification questions and the run time of the tool is acceptable. Significant reductions are obtained by parameter elimination in L/U automata.

There are several relevant and interesting topics for future research. First of all, serious improvements in the applicability of the tool can be obtained by improving the user interface. Currently, the tool generates many parameter equations whose disjunction is the desired constraint. Since the number of equations are can be quite large, it would be more convenient if the tool could simplify these set of equations. This could for instance be done with reduction techniques for BDDs.

Another relevant issue for parameter analysis is the theoretical investigation of the class of L/U automata. It would for instance be interesting to get more insight which types of problems are decidable for L/U automata and which are not. Furthermore, it would be interesting to investigate the use of L/U automata for synthesizing the constraints, rather than for analyzing given constraints as we did in this paper. On the practical side, the reduction techniques for L/U automata could be implemented.

Appendix A. Notational conventions

a	action
b	natural number
c	constraint
d	non-negative real number
e	linear expression
f	simple guard
g	guard
i, *j*	index
k	total number of actions

l	lower bound parameter
m	total number of clocks
n	total number of parameters
p	parameter
q	location
r	reset set
s	state
t, T	integer or real number
u	upper bound parameter
v	parameter valuation
w	clock valuation
x, y	clock
z	parametric zone
A	set of actions
C	set of constraints
D	parametric difference bound matrix
E	set of linear expressions
G	set of guards
I	invariant function
K	number of lower bound parameters
L	set of lower bound parameters
M	number of upper bound parameters
P	set of parameters
Q	set of locations
R	set of reset sets
S	set of states
U	set of upper bound parameters
X	set of clocks
$\mathscr{A}$	parametric timed automaton
E	unit PDBM
$\mathscr{L}$	labelled transition system
N	the natural numbers
R	the real numbers
Z	the integers
λ, μ	extended valuation of lower bound (upper bound) parameter, respectively
ϕ	state formula
ψ	system property

Acknowledgements

References

[1] A. Annichini, E. Asarin, A. Bouajjani, Symbolic techniques for parametric reasoning about counter and clock systems, in: E.A. Emerson, A.P. Sistla (Eds.), Proceedings of the 12th International Conference on Computer Aided Verification, Lecture Notes in Computer Science, vol. 1855, Springer, Berlin, 2000, pp. 419–434.

[2] A. Annichini, A. Bouajjani, M. Sighireanu, TReX: a tool for reachability analysis of complex systems, in: G. Berry, H. Comon, A. Finkel (Eds.), Proceedings of the International Conference on Computer Aided Verification (CAV'01), Paris, France, Lecture Notes in Computer Science, vol. 1855, Springer, Berlin, 2001.

[3] R. Alur, D.L. Dill, Automata for modeling real-time systems, in: M. Paterson (Ed.), Proceedings of the 17th ICALP, Warwick, Lecture Notes in Computer Science, vol. 443, Springer, Berlin, 1990, pp. 322–335 (Full version appeared as [4]).

[4] R. Alur, D.L. Dill, A theory of timed automata, Theoret. Comput. Sci. 126 (1994) 183–235.

[5] R. Alur, T.A. Henzinger, M.Y. Vardi, Parametric real-time reasoning, in: Proceedings of the 25th Annual Symposium on Theory of Computing, ACM Press, New York, 1993, pp. 592–601.

[6] M. Abadi, L. Lamport, An old-fashioned recipe for real time, in: J.W. de Bakker, C. Huizing, W.P. de Roever, G. Rozenberg (Eds.), Proceedings REX Workshop on Real-Time: Theory in Practice, Mook, The Netherlands, June 1991, Lecture Notes in Computer Science, vol. 600, Springer, Berlin, 1992, pp. 1–27.

[7] R. Alur, Timed automata, in: NATO-ASI Summer School on Verification of Digital and Hybrid Systems, Springer, Berlin, 1998, to appear.

[8] M. Bozga, C. Daws, O. Maler, A. Olivero, S. Tripakis, S. Yovine, Kronos: A model-checking tool for real-time systems, in: A.J. Hu, M.Y. Vardi (Eds.), Proceedings of the 10th International Conference on Computer Aided Verification, Vancouver, BC, Canada, Lecture Notes in Computer Science, vol. 1427, Springer, Berlin, 1998, pp. 546–550.

[9] G. Bandini, R.L. Lutje Spelberg, R.C.H. de Rooij, W.J. Toetenel, Application of Parametric Model Checking—The Root Contention Protocol using {LPMC}, in: Proceedings of the 34th Hawaii International Conference on System Sciences (HICCS'01), 2001.

[10] E.M. Clarke, O. Grumberg, D. Peled, Model Checking, MIT Press, Cambridge, MA, 1999.

[11] T.H. Cormen, C.E. Leiserson, R.L. Rivest, Introduction to Algorithms, McGraw-Hill, New York, 1991.

[12] A. Collomb–Annichini, M. Sighireanu, Parameterized reachability analysis of the IEEE 1394 Root Contention Protocol using TReX, in: P. Pettersson, S. Yovine (Eds.), Proceedings of the Workshop on Real-Time Tools (RT-TOOLS'2001), 2001.

[13] D. Dill, Timing assumptions and verification of finite-state concurrent systems, in: J. Sifakis (Ed.), Proceedings of the International Workshop on Automatic Verification Methods for Finite State Systems, Grenoble, France, Lecture Notes in Computer Science, vol. 407, Springer, Berlin, 1990, pp. 197–212.

[14] P.R. D'Argenio, J.-P. Katoen, T.C. Ruys, J. Tretmans, The bounded retransmission protocol must be on time!, in: E. Brinksma (Ed.), Proceedings of the Third Workshop on Tools and Algorithms for the Construction and Analysis of Systems, Enschede, The Netherlands, Lecture Notes in Computer Science, vol. 1217, Springer, Berlin, 1997, pp. 416–431.

[15] T.A. Henzinger, P.-H. Ho, H. Wong-Toi, HYTECH: A model checker for hybrid systems, in: O. Grumberg (Ed.), Proceedings of the 9th International Conference on Computer Aided Verification, Lecture Notes in Computer Science, vol. 1254, Springer, Berlin, 1997, pp. 460–463.

[16] T.S. Hune, J.M.T. Romijn, M.I.A. Stoelinga, F.W. Vaandrager, Linear parametric model checking of timed automata, in: T. Margaria, W. Yi (Eds.), Proceedings of the International Conference on Tools and Algorithms for the Construction and Analysis of Systems, Genova, Italy, Lecture Notes in Computer Science, vol. 2031, Springer, Berlin, 2001, pp. 189–203.

[17] IEEE Computer Society. IEEE Standard for a High Performance Serial Bus. Std 1394–1995, August 1996.

[18] K.J. Kristoffersen, F. Laroussinie, K.G. Larsen, P. Pettersson, W. Yi, A compositional proof of a real-time mutual exclusion protocol, in: M. Bidoit, M. Dauchet (Eds.), Proceedings TAPSOFT'97: Theory and Practice of Software Development, Lille, France, Lecture Notes in Computer Science, vol. 1214, Springer, Berlin, 1997, pp. 565–579.

[19] L. Lamport, A fast mutual exclusion algorithm, ACM Trans. Comput. Syst. 5 (1) (1987) 1–11.

[20] K.G. Larsen, F. Larsson, P. Pettersson, Wang Yi, Efficient verification of real-time systems: Compact data structures and state-space reduction, in: Proceedings of the 18th IEEE Real-Time Systems Symposium, IEEE Computer Society Press, Silverspring, MD, 1997, pp. 14–24.

[21] K.G. Larsen, P. Pettersson, Wang Yi, UPPAAL in a Nutshell, Int. J. Software Tools Technol. Transfer 1 (1–2) (1997) 134–152.

[22] R.F. Lutje Spelberg, W.J. Toetenel, M. Ammerlaan, Partition refinement in real-time model checking, in: A.P. Ravn, H. Rischel (Eds.), Proceedings of the Fifth International Symposium on Formal Techniques in Real-Time and Fault-Tolerant Systems (FTRTFT'98), Lyngby, Denmark, Lecture Notes in Computer Science, vol. 1486, Springer, Berlin, 1998, pp. 143–157.

[23] N.A. Lynch, Distributed Algorithms, Morgan Kaufmann, San Fransisco, CA, 1996.

[24] M. Merritt, F. Modugno, M. Tuttle, Time constrained automata, in: J.C.M. Baeten, J.F. Groote (Eds.), Proceedings CONCUR 91, Amsterdam, Lecture Notes in Computer Science, vol. 527, Springer, Berlin, 1991, pp. 408–423.

[25] D.P.L. Simons, M.I.A. Stoelinga, Mechanical verification of the IEEE 1394a Root Contention Protocol using Uppaal2k. Springer International Journal on Software Tools for Technology Transfer (STTT) 3 (2001) 469–485.

[26] M.I.A. Stoelinga, Fun with FireWire: a comparative study of formal verification methods applied to the IEEE 1394 Root Contention Protocol, Formal Aspects of Computing, to appear (2002).

[27] M.I.A. Stoelinga, F.W. Vaandrager, Root contention in IEEE 1394, in: J.-P. Katoen (Ed.), Proceedings of the 5th International AMAST Workshop on Formal Methods for Real-Time and Probabilistic Systems, Bamberg, Germany, Lecture Notes in Computer Science, vol. 1601, Springer, Berlin, 1999, pp. 53–74.

[28] S. Yovine, Model checking timed automata, in: G. Rozenberg, F.W. Vaandrager (Eds.), Lectures on Embedded Systems, Lecture Notes in Computer Science, vol. 1494, Springer, Berlin, 1998, pp. 114–152.