

Deciding the winner in parity games is in $UP \cap co-UP$

Marcin Jurdziński^{1,2}

*Basic Research in Computer Science, Centre of the Danish National Research Foundation, Department of Computer Science,
University of Aarhus, Ny Munkegade, Building 540, 8000 Aarhus C, Denmark*

Received 6 April 1998; received in revised form 24 August 1998

Communicated by L.A. Hemaspaandra

Abstract

We observe that the problem of *deciding the winner in mean payoff games* is in the complexity class $UP \cap co-UP$. We also show a simple reduction from *parity games* to mean payoff games. From this it follows that deciding the winner in parity games and the *modal μ -calculus model checking* are in $UP \cap co-UP$. © 1998 Elsevier Science B.V. All rights reserved.

Keywords: Games on graphs; Computational complexity

1. Introduction

Parity games are infinite duration two-player games played on graphs. From the results of Emerson et al. [4] it follows that the problem of *deciding the winner* in parity games is equivalent via linear time reductions to the *modal μ -calculus model checking*. (In fact, they show the equivalence of the modal μ -calculus model checking and the *non-emptiness problem for automata on infinite trees with parity acceptance conditions* [7,3]. It is an easy exercise to show the equivalence of the non-emptiness problem for parity automata and the problem of deciding the winner in parity games.)

The modal μ -calculus model checking problem is known to be in the complexity class $NP \cap co-NP$ [4], so it is considered unlikely to be NP -complete. On the other hand no polynomial time algorithms for it have been found so far. The modal μ -calculus is a powerful logic for specifying properties of finite

state distributed and reactive systems. The question of existence of a polynomial time algorithm for the model checking problem has relevance to the automatic verification of such systems.

A nondeterministic Turing machine is called *unambiguous* if for every input it has at most one accepting computation. The complexity class UP is the class of problems recognizable by unambiguous polynomial time nondeterministic Turing machines (e.g., see [9, pp. 283–284]). In other words, UP is the class of problems having *unique* short certificates (where *short* means “of length polynomial in the size of the instance of the problem”) checkable in polynomial time. In a standard fashion, by $co-UP$ we denote the class of problems being the complements of problems in UP .

The main result of this note is that the problem of deciding the winner in parity games has *unique short certificates*, i.e., it is in UP . To obtain the result claimed in the title we proceed in two steps. First we show a polynomial time reduction of parity games to the *mean payoff games* of Ehrenfeucht and Mycielski [2]. A polynomial time reduction of parity

¹ Email: mju@brics.dk.

² Partially supported by Polish KBN grant No. 8 T11C 002 11.

games to mean payoff games has been independently obtained by Puri [10] and Jerrum (see [11,12]). Next, using results of Zwick and Paterson [14], we argue that the problem of deciding the winner in mean payoff games is in $UP \cap co-UP$.

Before we proceed with the technical part of this note let us comment on the relationship between the complexity classes $UP \cap co-UP$ and $NP \cap co-NP$. Obviously

$$P \subseteq UP \cap co-UP \subseteq NP \cap co-NP.$$

There are relatively few (natural) problems known, which are in $NP \cap co-NP$, and are not known to be in P . Interestingly, to our best knowledge, all problems appearing in literature known to be in $NP \cap co-NP$ are in fact in $UP \cap co-UP$ as well (e.g., *primality* has been shown to belong to $UP \cap co-UP$ by Fellows and Koblitz [5]).

2. Games

We consider infinite duration games played by two players (player 0 and player 1) on finite graphs called arenas. An arena $A = (V, V_0, V_1, E)$ is a directed graph (V, E) with a partition $V_0 \cup V_1 = V$ of the set of vertices V . We require that the out-degree of each vertex is at least one. This ensures that every finite path in (V, E) can be prolonged.

A position in a game played on the arena A is a finite path in the underlying graph (V, E) . Note that although the arena is finite, the set of positions $Pos(V) \subseteq V^*$ in the game is infinite, because there are finite paths of arbitrary length in (V, E) . The initial position of a play starting from a vertex $v_0 \in V$ is the path consisting only of the vertex v_0 . A move in a game consists of extending the current position by one vertex. The result of a move from a position $\pi_i = \langle v_0, v_1, v_2, \dots, v_i \rangle$ is a new position $\pi_{i+1} = \langle v_0, v_1, v_2, \dots, v_i, v_{i+1} \rangle$, i.e., $(v_i, v_{i+1}) \in E$. Let $\pi_i = \langle v_0, v_1, v_2, \dots, v_i \rangle$ be the current position in a play. If $v_i \in V_0$ then it is the turn of player 0 to make a move, otherwise player 1 moves. The players make their moves indefinitely. Note that a play never gets “stuck”, because of our assumption that every vertex has at least one outgoing edge. The result of a play in the game is an infinite path $\pi = \langle v_0, v_1, v_2, \dots \rangle$ in the graph (V, E) . For brevity, we often just say a play, meaning the result of the play.

Formally, a game $G = (A, W)$ is an arena $A = (V, V_0, V_1, E)$, together with a winning condition $W \subseteq V^\omega$, specifying the set of winning plays for player 0. All the other plays are winning for player 1. A strategy for player 0 is a function $\zeta : Pos(V) \rightarrow V$, which given the current position in the game specifies what should be the next move to be made by player 0. A strategy is *memoryless* if it does not depend on the whole current position but only on the current vertex (i.e., the last vertex in the current position). Memoryless strategies are particularly simple and can be represented as functions $\sigma : V \rightarrow V$, i.e., they are finite objects. A winning strategy for player 0 from a vertex $v_0 \in V$ is a strategy such that every play starting in v_0 which is consistent with the strategy is winning for player 0. (A play is consistent with a strategy if player 0 makes all her moves along this play according to the strategy.)

The decision problem we are after is: given a game G and a starting vertex v_0 , decide if player 0 has a winning strategy in G from the vertex v_0 .

We deal with three types of infinite duration games played on arenas: *parity games*, *mean payoff games* and *discounted mean payoff games*. They differ in the way their winning conditions are represented.

Definition 1 (Parity game). A parity game [3,8] is a pair (A, p) where $A = (V, V_0, V_1, E)$ is an arena and $p : V \rightarrow \{0, 1, \dots, |V|\}$ is a function assigning a priority to every vertex of the arena. Let $\pi = \langle v_0, v_1, v_2, \dots \rangle$ be a play in a parity game. We define $Inf(\pi)$ to be the set of all numbers appearing infinitely often in the sequence $\langle p(v_0), p(v_1), p(v_2), \dots \rangle$. The play π is winning for player 0 if $\max(Inf(\pi))$ is even.

Definition 2 (Mean payoff game). A mean payoff game [2] is a quadruple (A, v, d, w) , where $A = (V, V_0, V_1, E)$ is an arena, v and d are natural numbers, and $w : E \rightarrow \{-d, \dots, -1, 0, 1, \dots, d\}$ is a function assigning an integer weight to every edge of the arena. A play $\pi = \langle v_0, v_1, v_2, \dots \rangle$ is winning for player 0 if

$$\liminf_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n w(v_{i-1}, v_i) \geq v.$$

Definition 3 (Discounted mean payoff game). Let $G = (A, v, d, w)$ be a mean payoff game and $0 < \lambda < 1$ be a real number called the discounting factor.

The *discounted mean payoff game* G_λ has the same arena and weight function as G . The only difference is the winning condition. A play $\pi = \langle v_0, v_1, v_2, \dots \rangle$ is winning for player 0 in G_λ if

$$(1 - \lambda) \sum_{i=1}^{\infty} \lambda^i \cdot w(v_{i-1}, v_i) \geq v.$$

The main aim of this note is to show that deciding the winner in parity games is in $\text{UP} \cap \text{co-UP}$. We have tried to provide explicit unique short certificates for winning strategies in parity games. One of the very promising candidates seemed to be *canonical signature assignments* studied by Walukiewicz [13]. We have failed, however, to devise a polynomial time algorithm to *check* that a signature assignment is indeed canonical. Instead, we relate parity games to mean payoff games by providing in Section 3 a simple reduction from the former to the latter. Then in Section 4 we argue that winning strategies for (discounted) mean payoff games have unique short certificates.

3. The reduction

In this section we show a simple reduction of parity games to mean payoff games. Essentially the same reduction has been given by Puri [10]. A reduction of parity games to mean payoff games has been also obtained by Jerrum [12].

Our reduction relies on *memoryless determinacy* of both parity and mean payoff games. A game is *determined* if for every starting vertex exactly one of the players has a winning strategy. Memoryless determinacy states that if a winning strategy for one of the players exists, then there is also a memoryless winning strategy for her. Memoryless determinacy for parity games was shown by Mostowski [8], and independently by Emerson and Jutla [3], and for mean payoff games by Ehrenfeucht and Mycielski [2], and by Gurvich et al. [6].

Theorem 4. *The problem of deciding the winner in a parity game reduces in polynomial time to the problem of deciding the winner in a mean payoff game.*

Proof. Let $G = (A, p)$ be a parity game, where $A = (V, V_0, V_1, E)$ is the arena. Given G and a starting

vertex $v_0 \in V_0$ we construct a mean payoff game $H = (A, v, d, w)$, such that player 0 has a winning strategy from v_0 in G if and only if she has a winning strategy from v_0 in H . Let $r = \max\{p(v) : v \in V\}$, and $m = |V|$. To define the mean payoff game H we set $v = 0$, $d = m^r$ and for every edge $e = (u, v) \in E$ we let $w(e) = (-m)^{p(u)}$.

Observe that the arenas of the games G and H are the same, hence a (memoryless) strategy in G is at the same time a (memoryless) strategy in H and vice versa. Let σ be a memoryless strategy for player 0 from vertex v_0 in game G (H). By G_σ (H_σ) we denote the subgame of G (H) obtained by first removing from the arena A all the edges coming out of vertices in V_0 except for those belonging to σ , and then by removing all the vertices unreachable from v_0 via the edges left. Plays in the game G (H) starting in v_0 and consistent with the strategy σ correspond exactly to the set of all plays in the game G_σ (H_σ). \square

Below, by a *simple cycle* we mean a cycle with no repeating vertices. The following lemma is used to prove that the above construction is indeed a reduction from parity games to mean payoff games.

Lemma 5. *Let σ be a memoryless strategy for player 0 in G (H). Then:*

- (1) *for every simple cycle c in G_σ the highest priority of a vertex on c is even if and only if the sum of the weights of the edges on c is nonnegative,*
- (2) *if σ is a winning strategy for player 0 from v_0 in G then for every simple cycle c in G_σ the highest priority of a vertex appearing on c is even,*
- (3) *if σ is a winning strategy for player 0 from v_0 in H then for every simple cycle c in H_σ the sum of the weights of the edges on c is nonnegative.*

Proof. The first clause is straightforward.

We prove the second clause; the proof of the third one is similar. Suppose that there exists a simple cycle c with the highest priority of a vertex on c being odd. Then player 1 can force the play from v_0 to c and also to stay in c indefinitely and thus win. This, however, contradicts our assumption that σ is a winning strategy for player 0.

Suppose that player 0 has a winning strategy from v_0 in the game G . Due to memoryless determinacy of parity games it follows that there is also a memoryless

winning strategy σ for player 0 from v_0 in G . We show that the memoryless strategy σ is also a winning strategy for player 0 from v_0 in H . In order to do that we have to argue that for every play $\langle v_0, v_1, v_2, \dots \rangle$ consistent with the strategy σ the following inequality holds:

$$\liminf_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n w(v_{i-1}, v_i) \geq 0.$$

We show now that every play $\pi = \langle v_0, v_1, v_2, \dots \rangle$ in H_σ can be essentially decomposed into simple cycles. This decomposition allows us to apply Lemma 5 and get the above inequality. Consider the following decomposition process of the play π in H_σ . We maintain a stack containing a sequence of distinct nodes forming a finite path u_0, u_1, \dots, u_h in H_σ , where h is the height of the stack. Whenever the next vertex from the path π to be considered happens to be already on the stack, we remove the vertices forming the cycle from the top of the stack. Otherwise we push the new vertex onto the stack. Observe that due to clauses (1) and (2) of Lemma 5 whenever we remove a cycle from the top of the stack the sum of the weights of the edges on the cycle is nonnegative. In this way only the weights of the edges which are on the stack may sum up to a negative value. Hence:

$$\begin{aligned} \liminf_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n w(v_{i-1}, v_i) \\ \geq \liminf_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^{h(n)} w(u_{i-1}, u_i) = 0, \end{aligned}$$

because $h(n) \leq |V|$, so the absolute value of the sum $\sum_{i=1}^{h(n)} w(u_{i-1}, u_i)$ is bounded by a constant.

In order to finish the proof of the theorem we have to show the reverse, i.e., if player 0 has a winning strategy from v_0 in H , then player 0 has also a winning strategy from v_0 in G . Again due to memoryless determinacy of mean payoff games we can restrict to memoryless strategies. Then it suffices to apply the same technique of decomposing plays in G_σ into simple cycles and use clauses (1) and (3) of Lemma 5. \square

Remark. In the above reduction the weight function w has values exponential with respect to the size of the parity game G , but their binary representations are of polynomial size and can be computed in polynomial

time. Note that having a reduction with polynomial values of the weight function would imply that parity games are in P, as a polynomial time algorithm for mean payoff games with polynomial weights is given by Zwick and Paterson [14].

4. The $UP \cap co-UP$ upper bound

Our aim now is to show that deciding the winner in mean payoff games is in $UP \cap co-UP$. Let $G = (A, v, d, w)$ be a mean payoff game. Ehrenfeucht and Mycielski [2] and independently Gurvich et al. [6] have shown that for every vertex v_0 of the arena A there exists a number $v(v_0)$, called the *value of G in v_0* , such that the following two conditions hold:

- (1) player 0 has a memoryless strategy such that for every play $\langle v_0, v_1, v_2, \dots \rangle$ consistent with this strategy

$$\liminf_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n w(v_{i-1}, v_i) \geq v(v_0),$$

- (2) player 1 has a memoryless strategy such that for every play $\langle v_0, v_1, v_2, \dots \rangle$ consistent with this strategy

$$\limsup_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n w(v_{i-1}, v_i) \leq v(v_0).$$

We call strategies satisfying those conditions *optimal*. Clearly, given the values of a game it is straightforward to decide the winner—it suffices to check whether $v(v_0) \geq v$. Hence the values of a mean payoff game may seem to be a plausible candidate for unique short certificates. Unfortunately, we have failed to devise a polynomial time algorithm to check whether a vector $\langle x_v \rangle_{v \in V}$ is indeed the vector of the *values* of the game. Moreover, whereas it is straightforward to extract a winning strategy in a parity game from the canonical signature assignment [13], here it is not at all clear to us how to do it efficiently given just the values of the mean payoff game.

In order to get the $UP \cap co-UP$ upper bound on the complexity of mean payoff games we use the following result of Zwick and Paterson [14].

Theorem 6 (Zwick and Paterson [14]). *The problem of deciding the winner in a mean payoff game reduces*

in polynomial time to the problem of deciding the winner in a discounted mean payoff game.

Hence it suffices to provide unique short certificates for winning strategies in *discounted* mean payoff games. Zwick and Paterson [14] have shown that *optimal memoryless strategies* for both players exist also in *discounted* mean payoff games. More precisely, for every vertex v_0 there exists a number $v_\lambda(v_0)$, called the *value of the discounted mean payoff game* G_λ in v_0 , satisfying the two conditions above, with the term $(1 - \lambda) \sum_{i=1}^n \lambda^i \cdot w(v_{i-1}, v_i)$ substituted for $1/n \sum_{i=1}^n w(v_{i-1}, v_i)$ and the number $v_\lambda(v_0)$ in place of $v(v_0)$.

The vector of the values of a discounted mean payoff game is the unique certificate we are after. The crucial property which allows us to quickly check whether some $(x_v)_{v \in V}$ is the vector of the values of a discounted mean payoff game is the following characterization due to Zwick and Paterson [14].

Theorem 7 (Zwick and Paterson [14]). *The vector $\bar{v} = \langle v_\lambda(v) \rangle_{v \in V}$ of the values of the discounted mean payoff game G_λ is the unique solution of the following system of equations*

$$x_v = \begin{cases} \max_{(v,u) \in E} \{ (1 - \lambda) \cdot w(v, u) + \lambda x_u \} & \text{if } v \in V_0, \\ \min_{(v,u) \in E} \{ (1 - \lambda) \cdot w(v, u) + \lambda x_u \} & \text{if } v \in V_1. \end{cases}$$

It only remains to show that the values of a discounted mean payoff games are *short*, i.e., can be written using a number of bits polynomial in $N = |G_\lambda|$, the size of the binary representation of the game G_λ . We prove this by applying a standard technique; the proof sketch below closely resembles, e.g., the proof of Lemma 2 in [1]. Let \bar{x} be the only solution of the system of equations from Theorem 7. Then

$$\bar{x} = (1 - \lambda) \cdot \bar{w} + \lambda \cdot Q \cdot \bar{x},$$

where \bar{w} is a vector of appropriate weights $w(u, v)$ and Q is a zero–one matrix with exactly one non-zero element in every row. We assume that the constant λ is a rational included in the binary representation of the game G_λ , hence it is a number a/b , with a and b being integers satisfying $\log a, \log b < N$. The above equation can be hence rewritten as

$$A \cdot \bar{x} = (b - a) \cdot \bar{w},$$

where $A = b \cdot I - a \cdot Q$, and I is the $|V| \times |V|$ identity matrix. Observe that A is a matrix with at most two non-zero elements in a row, and their absolute values are not bigger than 2^N . It is easy to show by induction on the size of the matrix that the absolute value of the determinant of A is not bigger than $4^{N|V|}$. According to Cramer's rule we have $x_v = \det A_v / \det A$, where A_v is obtained from A by replacing the v th column with the vector $(b - a) \cdot \bar{w}$. From this it follows that the elements of the vector \bar{x} can be written using a polynomial number of bits.

This settles the UP upper bound for the problem of deciding the winner in discounted mean payoff games. The co-UP upper bound follows easily, because by the existence of optimal strategies [14] player 0 does *not* have a winning strategy from a vertex v_0 in the game G_λ if and only if $v_\lambda(v_0) < v$, so it can be directly read from the vector of the values of the game.

Putting the UP \cap co-UP upper bound for discounted mean payoff games, together with the reductions of Theorems 4 and 6, we get our main result.

Theorem 8. *The problems of deciding the winner in parity, mean payoff and discounted mean payoff games are in UP \cap co-UP.*

Acknowledgements

I am indebted to Damian Niwiński for his guidance and encouragement. I thank Igor Walukiewicz for helpful comments, and an anonymous referee for pointing me to the work of Jerrum (see [11,12]) and Puri [10].

References

- [1] A. Condon, The complexity of stochastic games, Inform. and Comput. 96 (1992) 203–224.
- [2] A. Ehrenfeucht, J. Mycielski, Positional strategies for mean payoff games, Internat. J. Game Theory 8 (2) (1979) 109–113.
- [3] E.A. Emerson, C.S. Jutla, Tree automata, mu-calculus and determinacy (Extended abstract), in: Proceedings 32nd Annual Symposium on Foundations of Computer Science, IEEE Computer Society Press, 1991, pp. 368–377.
- [4] E.A. Emerson, C.S. Jutla, A.P. Sistla, On model-checking for fragments of μ -calculus, in: C. Courcoubetis (Ed.), Computer Aided Verification, 5th International Conference, CAV'93, Lecture Notes in Comput. Sci., Vol. 697, Springer, Berlin, 1993, pp. 385–396.

- [5] M.R. Fellows, N. Kobitz, Self-witnessing polynomial-time complexity and prime factorization, in: Proc. 7th Annual Conference on Structure in Complexity Theory, IEEE Computer Society Press, 1992, pp. 107–110.
- [6] V.A. Gurvich, A.V. Karzanov, L.G. Khachiyan, Cyclic games and an algorithm to find minimax cycle means in directed graphs, *USSR Comput. Maths. Math. Phys.* 28 (5) (1988) 85–91.
- [7] A.W. Mostowski, Regular expressions for infinite trees and a standard form of automata, in: A. Skowron (Ed.), *Computation Theory*, Lecture Notes in Comput. Sci., Vol. 208, Springer, Berlin, 1984, pp. 157–168.
- [8] A.W. Mostowski, Games with forbidden positions, Technical Report 78, University of Gdańsk, 1991.
- [9] C.H. Papadimitriou, *Computational Complexity*, Addison-Wesley, Reading, MA, 1994.
- [10] A. Puri, Theory of hybrid systems and discrete event systems, Ph.D. Thesis, University of California, Berkeley, CA, 1995.
- [11] C. Stirling, Local model checking games (Extended abstract), in: I. Lee, S.A. Smolka (Eds.), *CONCUR'95: Concurrency Theory*, 6th International Conference, Lecture Notes in Comput. Sci., Vol. 962, Springer, Berlin, 1995, pp. 1–11.
- [12] C. Stirling, Model checking and other games. Manuscript, available through <http://www.dcs.ed.ac.uk/home/cps>, 1996.
- [13] I. Walukiewicz, Pushdown processes: Games and model checking, in: T.A. Henzinger, R. Alur (Eds.), *Computer Aided Verification*, 8th International Conference, CAV'96, Lecture Notes in Comput. Sci., Vol. 1102, Springer, Berlin, 1996, pp. 62–74, Full version available through <http://zls.mimuw.edu.pl/~igw>.
- [14] U. Zwick, M. Paterson, The complexity of mean payoff games on graphs, *Theoret. Comput. Sci.* 158 (1996) 343–359.