



# Compact Data Structures and State-Space Reduction for Model-Checking Real-Time Systems

KIM G. LARSEN  
*BRICS, Aalborg University, Denmark*

kgl@cs.auc.dk

FREDRIK LARSSON  
*Department of Computer Systems, Uppsala University, Sweden*

fredrik@docs.uu.se

PAUL PETTERSSON  
*Department of Computer Systems, Uppsala University, Sweden*

paupet@docs.uu.se

WANG YI  
*Department of Computer Systems, Uppsala University, Sweden*

yi@docs.uu.se

**Abstract.** During the past few years, a number of verification tools have been developed for real-time systems in the framework of timed automata. One of the major problems in applying these tools to industrial-sized systems is the huge memory-usage for the exploration of the state-space of a network (or product) of timed automata, as the model-checkers must keep information about not only the control structure of the automata but also the clock values specified by clock constraints.

In this paper, we present a compact data structure for representing clock constraints. The data structure is based on an  $\mathcal{O}(n^3)$  algorithm which, given a constraint system over real-valued variables consisting of bounds on differences, constructs an equivalent system with a minimal number of constraints. In addition, we have developed an on-the-fly reduction technique to minimize the space-usage. Based on static analysis of the control structure of a network of timed automata, we are able to compute a set of symbolic states that cover all the dynamic loops of the network in an on-the-fly searching algorithm, and thus ensure termination in reachability analysis.

The two techniques and their combination have been implemented in the tool UPPAAL. Our experimental results demonstrate that the techniques result in truly significant space-reductions: for six examples from the literature, the space saving is between 75% and 94%, and in (nearly) all examples time-performance is improved. Noteworthy is also the observation that the two techniques are completely orthogonal.

**Keywords:** real-time systems, model checking, design tool, formal specification and verification, timed automata

## 1. Introduction

Reachability analysis has been one of the most successful methods for automated analysis of concurrent systems. Many verification problems, e.g., invariant checking can be solved by means of reachability analysis. It can in many cases also be used for checking whether a system described as an automaton satisfies a requirement specification formulated, e.g., in linear temporal logic, by converting the requirement to an automaton and thereafter checking whether the parallel composition of the system and requirement automata can reach certain annotated states (Vardi and Wolper, 1986; Holzmann, 1991; Aceto et al., 1998). However, the major problem in applying reachability analysis is the potential

combinatorial explosion of state spaces. To attack this problem, various symbolic and reduction techniques have been put forward over the last decade to efficiently represent state space and to avoid exhaustive state space exploration (e.g., Burch et al., 1990; Godefroid and Wolper, 1991; Valmari, 1990; Clarke et al., 1992, 1993; Emerson and Jutla, 1993; Andersen, 1995); such techniques have played a crucial role for the successful development of verification tools for finite-state systems.

In the last few years, new verification tools have been developed, for the class of infinite-state systems known as timed systems (Henzinger et al., 1995; Daws and Yovine, 1995; Bengtsson et al., 1996). Notably the verification engines of most tools in this category are based on reachability analysis of timed automata following the pioneering work of Alur and Dill (1990). A timed automaton is an extension of a finite automaton with a finite set of real-valued clock-variables. The foundation for decidability of reachability problems for timed automata is Alur and Dill's region technique, by which the infinite state space of a timed automaton due to the density of time, may effectively be partitioned into finitely many equivalence classes, i.e., regions in such a way that states within each class will always evolve to states within the same classes. However, reachability analysis based on the region technique is practically infeasible due to the potential state explosions arising from not only the control-structure (as for finite-state systems) but also the region space (Larsen et al., 1995).

Efficient data structures and algorithms have been sought to represent and manipulate timing constraints over clock variables (e.g., by difference bounded matrices (Bellman, 1957 and Dill, 1989)), or binary decision diagrams (Burch et al., 1990 and Asarin et al., 1997) and to avoid exhaustive state space exploration (e.g., by application of partial order reductions (Godefroid and Wolper, 1991; Valmari, 1990 and Pagani, 1996) or compositional methods (Andersen, 1995 and Larsen et al., 1995)). One of the main achievements in these studies is the symbolic technique (Dill, 1989; Yannakakis and Lee, 1993; Henzinger et al., 1994; Yi et al., 1994; Larsen et al., 1995), that converts the reachability problem to that of solving simple constraints. The technique can be simply formulated in an abstract reachability algorithm<sup>1</sup> as shown in Figure 1. The algorithm is to check whether a timed automaton may reach a state satisfying a given state formula  $\varphi$ . It explores the state space of the automaton in terms of symbolic states in the form  $(l, D)$  where  $l$  is a control-node and  $D$  is a constraint over clocks variables.

We observe that several operations of the algorithm are critical for efficient implementations. First, the algorithm depends heavily on the test operations for checking the inclusion  $DD'$  (i.e., the inclusion between the solution sets of  $D, D'$ ) and the emptiness of  $D_s$  in constructing the successor set  $SUCC$  of  $(l, D)$ . Clearly, it is important to design efficient data structures and algorithms for the representation and manipulation of clock constraints. One such well-known data structure is that of difference bounded matrix (DBM), which offers a canonical representation for constraint systems. It has been successfully used in several real-time verification tools, e.g., UPPAAL (Bengtsson et al., 1996) and KRONOS (Daws and Yovine, 1995). A DBM representation is in fact a weighted directed graph where the vertices correspond to clocks (including a zero-clock) and the weights on the edges stand for the bounds on the differences between pairs of clocks (Bellman, 1957; Dill, 1989; Yannakakis and Lee, 1993). As it gives an explicit bound for the difference between each pair of clocks, its space-usage is in the order of

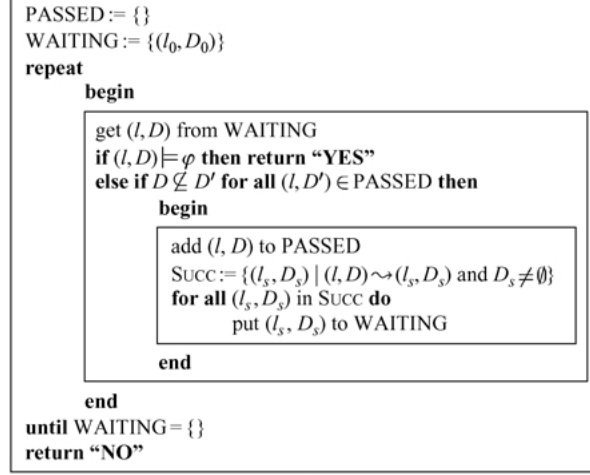


Figure 1. An algorithm for symbolic reachability analysis.

$\mathcal{O}(n^2)$  where  $n$  is the number of clocks. However, in practice it often turns out that most of these bounds are redundant.

In this paper, we present a compact data structure for DBM, which provides minimal and canonical representations of clock constraints and also allows for efficient inclusion checks. We have developed an  $\mathcal{O}(n^3)$  algorithm that given a DBM constructs a minimal number of constraints equivalent to the original constraints represented by the DBM (i.e., with the same solution set). The algorithm is essentially a minimization algorithm for weighted directed graphs, and hence solves a problem of independent interest. Note that the main global data structure of the algorithm in Figure 1 is the passed list (i.e., PASSED) holding the explored states. In many cases, it will store all the reachable symbolic states of the automaton. Thus, it is desirable that when saving a (symbolic) state in the passed list, we save the (often substantially smaller) minimal constraint system. The minimal representation also makes the inclusion-checking of the algorithm more efficient. Our experimental results demonstrate truly significant space-savings as well as better time-performance (see statistics in Section 5).

In addition to the local reduction technique above, which is to minimize the space-usage of each individual symbolic state, as the second contribution of this paper, we have developed a global reduction technique to reduce the total number of states to save in the global data structure, i.e. the passed list. It is completely orthogonal to the local technique. In the abstract algorithm of Figure 1, we notice the step of saving the new encountered state  $(l, D)$  in the passed list when the inclusion-checking for  $DD'$  fails (i.e.,  $DD'$ ). Its purpose is first of all to guarantee termination but also to avoid repeated exploration of states that have several predecessors. However, this is not necessary if all the predecessors of  $(l, D)$  are already present in the passed list. In fact, to ensure termination, it suffices to save only one state for each dynamic loop. An improved on-the-fly reachability algorithm according to the global reduction strategy has been implemented in UPPAAL (Bengtsson et al., 1996) based on static analysis of the control




Figure 2. A timed automaton.

structure of timed automata. Our experimental results demonstrate significant space-savings and also better time-performance (see statistics in Section 5).

The outline of this paper is as follows: In the next section we review the semantics of timed automata and the notion of DBM for clock constraints. Section 3 presents the compact data structure for DBM and the local reduction technique (i.e., the minimization algorithm for weighted directed graphs). Section 4 is devoted to develop the global reduction technique based on control structure analysis. Section 5 presents our experimental results for both techniques and their combination. Section 6 concludes the paper.

## 2. Preliminaries

### 2.1. Timed Automata

The model of timed automata was first introduced in Alur and Dill (1990) and has since then established itself as a standard model for real-time systems. For the reader not familiar with the notion of timed automata we give a short informal description.

Consider the timed automaton of Figure 2. It has two control nodes  $l_0$  and  $l_1$  and two real-valued clocks  $x$  and  $y$ . A state of the automaton is of the form  $(l, s, t)$ , where  $l$  is a control node, and  $s$  and  $t$  are non-negative reals giving the value of the two clocks  $x$  and  $y$ . A control node is labelled with a condition (the invariant) on the clock values that must be satisfied for states involving this node. Assuming that the automaton starts to operate in the state  $(l_0, 0, 0)$ , it may stay in node  $l_0$  as long as the invariant  $x \leq 4$  of  $l_0$  is satisfied. During this time the values of the clocks increase synchronously. Thus from the initial state, all states of the form  $(l_0, t, t)$ , where  $t \leq 4$ , are reachable. The edges of a timed automaton may be decorated with a condition (guard) on the clock values that must be satisfied in order to be enabled. Thus, only for the states  $(l_0, t, t)$ , where  $1 \leq t \leq 4$ , is the edge from  $l_0$  to  $l_1$  enabled. Additionally, edges may be labeled with simple assignments resetting clocks. For example, when following the edge from  $l_0$  to  $l_1$  the clock  $y$  is reset to 0 leading to states of the form  $(l_1, t, 0)$ , where  $1 \leq t \leq 4$ .

In general, a timed automaton is a standard finite-state automaton extended with a finite collection  $\mathcal{C}$  of real-valued clocks ranged over by  $x, y$ , etc. We use  $\mathcal{B}(\mathcal{C})$  ranged over by  $g$  (and latter  $D$ ), to stand for the set of formulas that can be an atomic constraint of the form:  $x \sim n$  or  $x - y \sim n$  for  $x, y \in \mathcal{C}$ ,  $\sim \in \{\leq, \geq\}^2$  and  $n$  being a natural

number, or a conjunction of such formulas. Elements of  $\mathcal{B}(\mathcal{C})$  are called clock constraints or constraint systems over  $\mathcal{C}$ .

**DEFINITION 1 (Timed automata)** A timed automaton  $A$  over clocks  $\mathcal{C}$  is a tuple  $\langle N, l_0, \rightarrow, I \rangle$  where  $N$  is a finite set of nodes (control-nodes),  $l_0$  is the initial node,  $\rightarrow \subseteq N \times \mathcal{B}(\mathcal{C}) \times 2^{\mathcal{C}} \times N$  corresponds to the set of edges, and finally,  $I : N \mapsto \mathcal{B}(\mathcal{C})$  assigns invariants to nodes. In the case,  $\langle l, g, r, l' \rangle \in \rightarrow$ , we write  $l \xrightarrow{g,r} l'$ .

Formally, we represent the values of clocks as functions (called clock assignments) from  $\mathcal{C}$  to the non-negative reals  $\mathbb{R}_+$ . We denote by  $\mathbb{R}_+^{\mathcal{C}}$  the set of clock assignments for  $\mathcal{C}$ . A semantical state of an automaton  $A$  is now a pair  $(l, u)$ , where  $l$  is a node of  $A$  and  $u$  is a clock assignment for  $\mathcal{C}$ , and the semantics of  $A$  is given by a transition system with the following two types of transitions (corresponding to delay- and edge-transitions):

- $(l, u) \rightarrow (l, u \oplus d)$  if  $I(l)(u)$  and  $I(l)(u \oplus d)$ ,
- $(l, u) \rightarrow (l', u')$  if there exist  $g$  and  $r$  such that  $l \xrightarrow{g,r} l'$ ,  $g(u)$  and  $u' = r[u]$ ,

where for  $d \in \mathbb{R}_+$ ,  $u \oplus d$  denotes the time assignment which maps each clock  $x$  in  $\mathcal{C}$  to the value  $u(x) + d$ , and for  $r \subseteq \mathcal{C}$ ,  $r[u]$  denotes the assignment for  $\mathcal{C}$  which maps each clock in  $r$  to the value 0 and agrees with  $u$  over  $\mathcal{C} \setminus r$ .

Clearly, the semantics of a timed automaton yields an infinite transition system, and is thus not an appropriate basis for decision algorithms. However, efficient algorithms may be obtained using a finite-state symbolic semantics based on symbolic states of the form  $(l, D)$ , where  $D \in \mathcal{B}(\mathcal{C})$  (Henzinger et al., 1994; Yi et al., 1994). We shall consider a clock constraint as a set of clock assignments and use  $u \in D$  to stand for  $u$  satisfied  $D$ .

The symbolic counterpart to the standard semantics is given by the following two (fairly obvious) types of symbolic transitions:

- $(l, D) \rightarrow (l, (D \wedge I(l))^\uparrow \wedge I(l))$ ,
- $(l, D) \rightarrow (l', r(g \wedge D))$  if  $l \xrightarrow{g,r} l'$ ,

where  $D^\uparrow = \{u \oplus d \mid u \in D \wedge d \in \mathbb{R}_+\}$  and  $r(D) = r[u] \cdot u \in D$ . It may be shown that  $\mathcal{B}(\mathcal{C})$  (the set of clock constraints) is closed under these two operations (and  $\wedge$ ) (Dill, 1989). Moreover, the symbolic semantics characterize the standard semantics in the sense that, whenever  $u \in D$  and  $(l, D) \rightarrow (l', D')$  then  $(l, u) \rightarrow (l', u')$  for  $u' \in D'$ .

Finally, we introduce the notion of networks of timed automata (Yi et al., 1994; Larsen et al., 1995). A network is the parallel composition of a finite set of automata for a given synchronization function. To illustrate the on-the-fly verification technique, we only need to study the case dealing with interleaving, i.e., the network of automata  $A_1, \dots, A_n$ , is the Cartesian product of  $A_i$ s. Assume a vector  $l$  of control nodes. We shall use  $l[i]$  to stand for the  $i$ -th element of  $l$  and  $l[l_i/l_i]$  for the vector where the  $i$ -th element  $l_i$  of  $l$  is replaced by  $l'_i$ . A control node (i.e., control vector)  $l$  of a network  $A_1, \dots, A_n$  is a vector where  $l[i]$  is a node of  $A_i$  and the invariant  $I(l)$  of  $l$  is the conjunction of  $I(l[1]), \dots, I(l[n])$ . The symbolic




Figure 3. Graph for  $E$  (a), its shortest-path closure (b), and shortest-path reduction (c).

semantics of networks is given in terms of control vectors by the following two types of symbolic transitions:

- $(l, D)(l, (D \wedge I(l))^{\uparrow} \wedge I(l))$
- $(l, D)(l[l'_i/l_i], r(g \wedge D))$  if  $l_i \xrightarrow{g,r} l'_i$

In the latter case, we shall say that the symbolic transition is derived by the edge  $l_i \xrightarrow{g,r} l'_i$ .

## 2.2. Difference Bounded Matrices and Shortest-Path Closure

To utilize the symbolic semantics of (networks of) timed automata algorithmically, as for example in the reachability algorithm of Figure 1, it is important to design efficient data structures and algorithms for the representation and manipulation of clock constraints.

One such well-known data structure is that of difference bounded matrices (DBM, see Bellman, 1957; Dill, 1989), which offers a canonical representation for constraint systems. A DBM representation of a constraint system  $D$  is simply a weighted, directed graph, where the vertices correspond to the clocks of  $C$  and an additional zero-vertex 0. The graph has an edge from  $x$  to  $y$  with weight  $m$  provided  $y - x \leq m$  is a constraint of  $D$ . Similarly, there is an edge from 0 to  $x$  with weight  $m$ , whenever  $x \leq m$  is a constraint of  $D$ .<sup>3</sup> As an example, consider the constraint system  $E$  over  $\{x_0, x_1, x_2, x_3\}$  being a conjunction of the atomic constraints  $x_0 - x_1 \leq 3$ ,  $x_3 - x_0 \leq 5$ ,  $x_3 - x_1 \leq 2$ ,  $x_2 - x_3 \leq 2$ ,  $x_2 - x_1 \leq 10$ , and  $x_1 - x_2 \leq -4$ . The graph representing  $E$  is given in Figure 3(a).

In general, the same set of clock assignments may be described by several constraint systems (and hence graphs). To test for inclusion between constraint systems  $D$  and  $D'$ ,<sup>4</sup> which we recall is essential for the termination of the reachability algorithm of Figure 1, it is advantageous if  $D$  is closed under entailment in the sense that no constraint of  $D$  can be strengthened without reducing the solution set. In particular, for  $D$  a closed constraint system,  $DD'$  holds if and only if for any constraint in  $D'$  there is a constraint in  $D$  at least as tight; i.e., whenever  $(x - y \leq m') \in D'$  then  $(x - y \leq m) \in D$  for some  $m \leq m'$ . Thus, closedness provides a canonical representation, as two closed constraint systems describe

the same solution set precisely when they are identical. To close a constraint system  $D$  amounts to derive the shortest-path closure for its graph and can thus be computed in time  $\mathcal{O}(n^3)$ , where  $n$  is the number of clocks of  $D$ . The graph representation of the closure of the constraint system  $E$  from Figure 3(a) is given in Figure 3(b). The emptiness-check of a constraint system  $D$  simply amounts to checking for negative-weight cycles in its graph representation. Finally, given a closed constraint system  $D$  the operations  $D^\dagger$  and  $r(D)$  may be performed in time  $\mathcal{O}(n)$ .

### 3. Minimal Constraint Systems and Shortest-Path Reductions

For the reasons stated above a matrix representation of constraint systems in closed form is an attractive data structure, which has been successfully employed by a number of real-time verification tools, e.g., UPPAAL (Bengtsson et al., 1996) and KRONOS (Daws and Yovine, 1995). As it gives an explicit (tightest) bound for the difference between each pair of clocks (and each individual clock), its space-usage is of the order  $\mathcal{O}(n^2)$ . However, in practice it often turns out that most of these bounds are redundant, and the reachability algorithm of Figure 1 is consequently hampered in two ways by this representation. First, the main data structure PASSED, will in many cases store all the reachable symbolic states of the automaton. Thus, it is desirable, that when saving a symbolic state in the PASSED-list, we save a representation of the constraint system with as few constraints as possible. Second, a constraint system  $D$  added to the PASSED-list is subsequently only used in checking inclusions of the form  $D'/D$ . Recalling the method for inclusion-check from the previous section, we note that (given  $D'$  is closed) the time-complexity of the inclusion-check is linear in the number of constraints of  $D$ . Thus, again it is advantageous for  $D$  to have as few constraints as possible.

In the following subsections we shall present an  $\mathcal{O}(n^3)$  algorithm, which given a constraint system constructs an equivalent reduced system with the minimal number of constraints. The reduced constraint system is canonical in the sense that two constraint systems with the same solution set give rise to identical reduced systems. The algorithm is essentially a minimization algorithm for weighted directed graphs. Given a weighted, directed graph with  $n$  vertices, it constructs in time  $\mathcal{O}(n^3)$  a reduced graph with the minimal number of edges having the same shortest path closure as the original graph. Figure 3(c) shows the minimal graph of the graphs in Figure 3(a) and (b), which is computed by the algorithm.

#### 3.1. Reduction of Zero-Cycle Free Graphs

A weighted, directed graph  $G$  is a structure  $(V, E_G)$ , where  $V$  is a finite set of vertices and  $E_G$  is a partial function from  $V \times V$  to  $\mathbb{Z}$  (the integers). The domain of  $E_G$  constitutes the edges of  $G$ , and when defined,  $E_G(x, y)$  gives the weight of the edge between  $x$  and  $y$ . We assume that  $E_G(x, x) = 0$  for all vertices  $x$ , and that  $G$  has no cycles with negative weight.<sup>5</sup>

Given a graph  $G$ , we denote by  $G^C$  the shortest-path closure of  $G$ , i.e.,  $E_{G^C}(x, y)$  is the length of the shortest path from  $x$  to  $y$  in  $G$ . A shortest-path reduction of a graph  $G$  is a graph  $G^R$  with the minimal number of edges such that  $(G^R)^C = G^C$ .

The key to reduce a graph is obviously to remove redundant edges, where an edge  $(x, y)$  is redundant if there exist an alternative path from  $x$  to  $y$  whose (accumulated) weight does not exceed the weight of the edge itself. For example, in the graph of Figure 3(a), the edge  $(x_1, x_2)$  is clearly redundant as the accumulated weight of path  $(x_1, x_0), (x_0, x_3), (x_3, x_2)$  has a weight (10) not exceeding the weight of the edge itself (also 10). The path  $(x_1, x_3), (x_3, x_2)$  makes also the edge  $(x_1, x_2)$  redundant. Being redundant, the edge  $(x_1, x_2)$  may be removed without changing the shortest-path closure. We shall use  $G \setminus (x_1, x_2)$  to denote the result of removing the edge  $(x_1, x_2)$  from the graph  $G$ .

Now, consider the edge  $(x_1, x_2)$  in the graph of Figure 3(b). Clearly, the edge is redundant as the path  $(x_1, x_3), (x_3, x_2)$  has equal weight. Similarly, the edge  $(x_3, x_2)$  is redundant as the path  $(x_3, x_1), (x_1, x_2)$  has equal weight. However, though redundant, we cannot just remove the two edges  $(x_1, x_2)$  and  $(x_3, x_2)$  as removal of one clearly requires the presence of the other. In fact, all edges between the vertices  $x_1, x_2$ , and  $x_3$  are redundant, but obviously we cannot remove them all simultaneously. The key explanation of this complicating phenomena is that  $x_1, x_2, x_3$  constitutes a cycle with length zero (a zero-cycle). However, for zero-cycle free graphs the situation is the simplest possible:

**Lemma 1** *Let  $G_1$  and  $G_2$  be zero-cycle free graphs such that  $G_1^C = G_2^C$ . If there is an edge  $(x, y) \in G_1$  such that  $(x, y) \notin G_2$ , then  $(G_1 \setminus \{(x, y)\})^C = G_1^C = G_2^C$ .*

**Proof:** Let  $\alpha$  denote the edge  $(x, y)$  and let  $m$  be the weight of  $\alpha$  in  $G_1$ . We will show that there is an alternative path in  $G_1$  not using  $\alpha$  with weight no more than  $m$ . From this fact the lemma obviously follows.

As  $G_1^C = G_2^C$ , the shortest path from  $x$  to  $y$  in  $G_2$  has weight no more than  $m$ . As  $\alpha \notin G_2$ , this path must visit some vertex  $z$  different from  $x$  and  $y$ . Now let  $m_1$  be the shortest path-weight from  $x$  to  $z$  and let  $m_2$  be the shortest path-weight from  $z$  to  $y$ ; note that  $G_1$  and  $G_2$  agrees on  $m_1$  and  $m_2$ , as they have the same shortest-path closure. Then clearly,  $m \geq m_1 + m_2$ .

Now assume that the shortest path in  $G_1$  from  $x$  to  $z$  uses  $\alpha = (x, y)$ . Then, as a sub-path,  $G_1$  will be a path from  $y$  to  $z$ . Since  $G_1$  also has a path from  $z$  to  $y$ , it follows that  $G_1$  will have a cycle from  $y$  via  $z$  back to  $y$ . The weight of this cycle can be argued to be no more than  $(m_1 - m) + m_2$ . However, as  $m \geq m_1 + m_2$  and there are no negative cycles, this cycle must have weight 0 contradicting the assumption that  $G_1$  is zero-cycle free.

Similarly, a contradiction with the zero-cycle free assumption of  $G_1$  is obtained, if the shortest path in  $G_1$  from  $z$  to  $y$  uses  $\alpha$ . Thus we can conclude that there is an path from  $x$  to  $y$  not using  $\alpha$  with length no greater than  $m$ . ■

From Lemma 1 it follows immediately that all redundant edges of a zero-cycle free graph may be removed without affecting the closure. On the other hand, removal of an edge which is not redundant will of course change the closure of the graph, and must be present in any graph with the same closure. Thus the following theorem follows:



**THEOREM 1** *Let  $G$  be a zero-cycle free graph, and let  $\{\alpha_1, \dots, \alpha_k\}$  be the set of redundant edges of  $G$ . Then  $G^R = G^C \setminus \{\alpha_1, \dots, \alpha_k\}$ .*

**Proof:** Follows from Lemma 1. ■

From an algorithmic point of view, redundancy of edges is easily determined given the closure  $G^C$  of a graph  $G$  as only path of length 2 needs to be considered: An edge  $(x, y)$  is redundant precisely when there is a vertex  $z$  ( $\neq x, y$ ) such that  $E_{G^C}(x, y) \geq E_{G^C}(x, z) + E_{G^C}(z, y)$ . Thus for zero-cycle free graphs computing  $G^R$  is  $\mathcal{O}(n^3)$ .

### 3.2. Reduction of Negative-Cycle Free Graphs

For general graphs (without negative cycles) our reduction construct relies on a partitioning of the vertices according to zero-cycles. We say that two vertices  $x$  and  $y$  are equivalent or zero-equivalent, if there is a zero-cycle containing them both. We write  $x \equiv y$  in this case. Given the closure  $G^C$  of a graph  $G$ , it is extremely easy to check for zero-equivalence:  $x \equiv y$  holds precisely when  $E_{G^C}(x, y) = -E_{G^C}(y, x)$ . Thus, in the graphs of Figure 3(a) and (b),  $\equiv$  partitions the vertices into the two classes  $\{x_0\}$  and  $\{x_1, x_2, x_3\}$ .

To obtain a canonical reduction, we assume that the vertices of  $G$  are ordered by assigning them indices as  $x_1, x_2, \dots, x_n$ . The equivalence  $\equiv$  now induces a natural transformation  $G_{\equiv}$  on the graph  $G$ :

**DEFINITION 2** *Given a graph  $G$ , the vertices of the graph  $G_{\equiv}$  are  $\equiv$ -equivalence classes, denoted  $E_k$ , of  $G$ . There is an edge between the classes  $E_i$  and  $E_j$  ( $i \neq j$ ) if for some  $x \in E_i$  and  $y \in E_j$  there is an edge in  $G$  between  $x$  and  $y$ . The weight of this edge is  $E_{G^C}(E_i^{\min}, E_j^{\min})$ , where  $E^{\min}$  is the vertex in  $E$  with the smallest index.*

Thus, the distance between  $E_i$  and  $E_j$  in  $G_{\equiv}$  is the weight of the shortest path in  $G$  between the elements of  $E_i$  and  $E_j$  with smallest index. It is obvious that  $G_{\equiv}$  is a zero-cycle free graph. It is also easy to see that  $G_{1\equiv} = G_{2\equiv}$  if  $G_1^C = G_2^C$ . Let  $H$  be the graph of Figure 3(a). Then  $H_{\equiv}$  will have vertices  $E_0 = \{x_0\}$  and  $E_1 = \{x_1, x_2, x_3\}$ . The two vertices are connected by two edges both having weight 3.

The following provides a dual to the operator of Definition 2:

**DEFINITION 3** *Let  $F$  be a graph with vertices being  $\equiv$ -equivalence classes with respect to a graph  $G = (V, E_G)$ . Then the expansion of  $F$  is a graph  $F^+$  with vertices  $V$  and with weight satisfying:*

- *For any multi-member equivalence class<sup>6</sup>  $\{z_1 < z_2 < \dots < z_k\}$  of  $F$ ,  $F^+$  contains a single cycle  $z_1, z_2, \dots, z_k, z_1$ , with the weight of the edge  $(z_i, z_{i+1})$  being the weight of the shortest path from  $z_i$  to  $z_{i+1}$  in  $G$ .*

- Whenever  $(E_i, E_j)$  is an edge in  $F$  with weight  $m$ , then  $F^+$  will have an edge from  $E_i^{\min}$  to  $E_j^{\min}$  with weight  $m$ .

We are now ready to state the main theorem giving the shortest-path reduction construct for arbitrary negative-cycle free graphs:

**THEOREM 2** *Let  $G$  be negative-cycle free graph. Then the shortest-path reduction  $G^R$  of  $G$  is given by the graph  $(G_{\equiv}^R)^+$ , i.e.,  $G^R = (G_{\equiv}^R)^+$ .*


**Proof:** We show: (1) that  $(G_{\equiv}^R)^+$  is a candidate for a shortest-path reduction of  $G$  in the sense that  $(G_{\equiv}^R)^+ = G^C$ , and (2) that  $(G_{\equiv}^R)^+$  is minimal.

1. We first prove that  $(G_{\equiv}^R)^+ = G^C$ . As all edges  $(x, y)$  of  $(G_{\equiv}^R)^+$  have weight of the form  $E_{G^C}(x, y)$ , it follows that for any path in  $(G_{\equiv}^R)^+$  there is a path in  $G$  with same weight.  
Now consider an edge  $(x, y)$  of  $G$ . We will demonstrate that there is a path in  $(G_{\equiv}^R)^+$  with no greater weight.

- If  $x = E_i^{\min}$  and  $y = E_j^{\min}$  for two  $\equiv$ -classes  $E_i$  and  $E_j$ , it follows that  $E_{G_{\equiv}}(E_i, E_j) \leq E_G(x, y)$ . Furthermore, due to the property of reduction construction, there is a path in  $G_{\equiv}^R$  between  $E_i$  and  $E_j$  with weight no greater than  $E_{G_{\equiv}}(E_i, E_j)$ . The same path, but now between the nodes with the minimal indices of the  $\equiv$ -classes, can be found in  $(G_{\equiv}^R)^+$ . Thus, there is a path in  $(G_{\equiv}^R)^+$  with weight no greater than  $E_G(x, y)$ .
- If  $x, y \in E_i$  for some  $\equiv$ -class  $E_i$ , an easy argument gives that  $E_{(G_{\equiv}^R)^+}(x, y) = E_{G^C}(x, y) \leq E_G(x, y)$ .
- Consider the case when  $x \in E_i$  and  $y \in E_j$  for two different  $\equiv$ -classes, and assume that  $E_G(x, y) = m$ .

Let  $m_1 = E_{(G_{\equiv}^R)^+}(x, E_i^{\min})$ ,  $m_2 = E_{(G_{\equiv}^R)^+}(E_i^{\min}, E_j^{\min})$ , and  $m_3 = E_{(G_{\equiv}^R)^+}(E_j^{\min}, y)$ . Note that by the reduction construction  $m_2 \leq E_{G^C}(E_i^{\min}, E_j^{\min})$ . Then there is a path in  $(G_{\equiv}^R)^+$  from  $x$  to  $y$  via  $E_i^{\min}$  and  $E_j^{\min}$  with weight  $m_1 + m_2 + m_3$ . Now, if  $m < m_1 + m_2 + m_3$ , there is a path in  $G$  from  $E_i^{\min}$  to  $E_j^{\min}$  of weight  $m - m_1 - m_3 < m_2$  contradicting that  $m_2$  is the weight of the shortest path in  $G$  between  $E_i^{\min}$  and  $E_j^{\min}$ . Thus the path  $x, E_i^{\min}, E_j^{\min}, y$  in  $(G_{\equiv}^R)^+$  has weight no greater than the edge  $(x, y)$  in  $G$ .

2. Next we prove that  $(G_{\equiv}^R)^+$  has minimal number of edges by showing that whenever  $H^C = G^C$  then  $H$  has at least as many edges as  $(G_{\equiv}^R)^+$ .  
As  $H^C = G^C$ ,  $H$  and  $G$  induces the same  $\equiv$ -equivalence relation on the same zero-length cycles. Obviously the fewest edges that will identify  $k$  ( $> 1$ ) vertices, with respect to  $\equiv$  is  $k$ . Hence,  $(G_{\equiv}^R)^+$  uses a minimal number of edges between vertices in the same  $\equiv$ -equivalence class.

Figure 4. A path in the graph  $H$ .

Now let  $(E_i^{\min}, E_j^{\min})$  be an edge in  $(G_{\equiv}^R)^+$  with weight  $m$ . We claim that  $H$  must have at least one edge from  $E_i$  to  $E_j$ .

Assume that this is not the case. Then, as  $H^C = G^C$ , there must be a path in  $H$  from  $E_i^{\min}$  to  $E_j^{\min}$  as shown in Figure 4 such that  $m = \sum_{i=0}^{k+2} v_i + \sum_{i=0}^{k+1} w_i$ .

Now let  $m_0 = E_{(G_{\equiv}^R)^+}(E_i^{\min}, E_0^{\min})$ ,  $m_1 = E_{(G_{\equiv}^R)^+}(E_0^{\min}, E_1^{\min})$ ,  $\dots$ ,  $m_{k+1} = E_{(G_{\equiv}^R)^+} \times (E_k^{\min}, E_i^{\min})$  (illustrated with dashed lines in Figure 4). Then  $m_0 \leq v_0 + w_0 + v'_1$ ,  $m_1 \leq v'_1 + w_1 + v'_2$ ,  $\dots$ ,  $m_{k+1} \leq v''_{k+1} + w_{k+1} + v_{k+2}$ , where  $v_1 = v'_1 + v''_1$ ,  $v_2 = v'_2 + v''_2$ ,  $\dots$ ,  $v_{k+1} = v'_{k+1} + v''_{k+1}$ .

It follows that  $\sum_{i=0}^{k+1} m_i \leq m$ . Hence  $(E_i, E_j)$  is redundant in  $(G_{\equiv}^R)^+$  and can be removed, contradicting Lemma 1. ■

First, note that the above construction of  $(G_{\equiv}^R)^+$  is well defined as  $G_{\equiv}$  is a zero-cycle free graph and the reduction construction of Theorem 1 thus applies. Given the closure  $G^C$  of  $G$  the constructions of Definitions 2 and 3 can be computed in  $\mathcal{O}(n^2)$ . Since  $G^R$  is computed from  $G$  in  $\mathcal{O}(n^3)$ , it follows that also  $(G_{\equiv}^R)^+$  can be constructed in  $\mathcal{O}(n^3)$ . Now applying the above construction to the graph  $H$  of Figure 3(a), we first note that  $H_{\equiv}^R = H_{\equiv}$  as  $H_{\equiv}$  has no redundant edges. Expanding  $H_{\equiv}$  with respect to the vertex ordering  $x_0 < x_1 < x_2 < x_3$  gives the graph of Figure 3(c), which according to Theorem 2 above is the shortest-path reduction of  $H$ .




Figure 5. Illustration of space-reduction.

Experimental results show that the use of minimal constrain systems (obtained by the above shortest-path reduction algorithm) as a compact data structure leads to truly significant space-savings in practical reachability analysis of timed systems: the space-savings are in the range 68–85%. We refer to Section 5 for more details.

#### 4. Global Reductions and Control Structure Analysis

The preceding section is about local reductions in reachability analysis in the sense that the technique developed is for each individual symbolic state. In this section, we shall develop a global reduction technique to reduce the total number of symbolic states to save in the global data structure, i.e., the passed list.

##### 4.1. Potential Space-Reductions

We recall the standard reachability analysis algorithm for finite graphs (see, for example, Papadimitriou, 1994). It is similar to the one in Figure 1, but simpler as no constraints but only control nodes are involved. The algorithm repeats three main operations: examining every new encountered node (to see if it is in the passed list), exploring the new encountered nodes (computing all their successors for further analysis), and saving the explored nodes in the passed list until all reachable nodes are present in the list (i.e., all new encountered nodes are already in the passed list).

Note that the saving of an explored node is to ensure termination and also to avoid repeated exploration of nodes with more than one incoming edge. However, it is not necessary to save all reachable nodes. Consider for example, the simple graph in Figure 5 with initial node  $l_0$ . Clearly, there is no need to save node  $l_2$ ,  $l_3$  or  $l_4$  as they will be visited only once if  $l_1$  is present in the passed list.

In fact, to guarantee termination on a finite graph, it is sufficient to save only one node for each cycle in the graph. For example, as  $l_1$  covers the two cycles of the graph in Figure 5, in addition to  $l_2$ ,  $l_3$ , and  $l_4$ , it is not necessary to save  $l_0$  either. In general, for a finite graph, there is a minimal number of nodes to save in the passed list in order to guarantee termination. However the trade-off of the space-saving strategy may be

increased time-consumption. Consider the same graph of Figure 5. If node  $l_0$  is not present in the passed list, it will be explored again whenever  $l_3$  is explored. This can be avoided by saving  $l_0$  when it is first visited. But the difference from saving  $l_1$  is that saving  $l_0$  is for efficiency and  $l_1$  for termination.

Now we again recall the abstract reachability algorithm in Figure 1 for timed systems. To ensure termination and also to avoid repeated exploration of states (that have more than one predecessors), it saves every new encountered state  $(l, D)$  in the passed list when the inclusion-checking for  $DD'$  fails (i.e.,  $DD'$ ). Obviously this is not necessary if all the predecessors of  $(l, D)$  already exist in the PASSED-list. Similar to the case for finite graphs, for termination, we need to save only one state for every dynamic loop of a timed automaton.

**DEFINITION 4 (Dynamic loops)** Assume a timed automaton with an initial state  $(l_0, D_0)$ . The set of symbolic states  $L_d = \{(l_1, D_1), \dots, (l_n, D_n)\}$  is a dynamic loop of the timed automaton if  $(l_1, D_1) \rightsquigarrow (l_2, D_2), \dots, (l_{n-1}, D_{n-1}) \rightsquigarrow (l_n, D_n)$  and  $(l_n, D_n) \rightarrow (l_1, D'_1)$  with  $D'_1 \subseteq D_1$ , and  $(l_0, D_0)$  is reachable in the sense that  $(l_0, D_0) \rightsquigarrow \dots \rightsquigarrow (l_1, D_1)$ . A symbolic state is said to cover a dynamic loop if it is a member of the loop.

We claim that to ensure termination, it is sufficient (but not necessary) to save a set of symbolic states that cover all the dynamic loops. Now, the problem is how to compute efficiently such a set.

#### 4.2. Control Structure Analysis and Application

We shall utilize the statical structure of an automaton to identify potential candidates of states to cover dynamic loops.

**DEFINITION 5 (Statical loops and entry Nodes)** A set of nodes  $L = \{l_1, \dots, l_n\}$  of a timed automaton is a statical loop if there is a sequence of edges  $l_1 \rightarrow l_2, l_{n-1} \rightarrow l_n$  and  $l_n \rightarrow l_1$  where  $l_i \rightarrow l_j$  denotes that  $l_i \xrightarrow{g,r} l_j$  for some  $g, r$  is an edge of the automaton. A node  $l_i \in L$  is an entry node of the statical loop  $L$  if it is an initial node of the automaton or there exists a node  $l \notin L$  (outside of the loop) and an edge  $l \rightarrow l_i$ . Further, we say that a vector of nodes (i.e. a node of a network) is an entry node if any of its components are entry nodes.

For example, nodes  $l_0, l_1, l_2$  and  $l_3$  in Figure 5 constitute a statical loop with entry nodes  $l_0$  and  $l_1$ ; another statical loop is nodes  $l_1$  and  $l_4$  with entry node  $l_1$ . In general, since the sets of control nodes and edges of a timed automaton are finite, the number of statical loops is finite and so is the set of entry nodes of all statical loops. In fact the set of entry nodes of a timed automaton can be easily computed by statical analysis using a stack or a slightly modified loop detecting algorithm (see, for example, Sedgewick, 1988).

Now note that according to Definition 4, a dynamic loop (a set of symbolic states) must contain a subset of symbolic states whose control nodes constitute a statical loop. As a statical loop always contains an entry node, we have the following fact.

**Proposition 1** Every dynamic loop of a timed automaton contains at least one symbolic state  $(l, D)$  where  $l$  is an entry node.

**Proof:** Standard proof by contradiction. ■

Following Proposition 1, to cover all the dynamic loops, we may simply save all the states whose control-nodes are an entry node, and ignore the others. Obviously, this will not give much reduction when dynamic loops include mostly entry nodes, which is the case when a network of automata contains a component whose nodes are mostly entry nodes, e.g., a testing automaton. For networks of automata, we adopt the strategy of saving the first derived states whose control nodes are an entry node, known as covering states in the following sense.

**DEFINITION 6 (Covering states)** Assume a network of timed automata with an initial state  $(l_0, D_0)$  and a given symbolic state  $(l, D)$ . We say that  $(l, D)$  is a covering state of the network if it is reachable in the sense that there exists a sequence of symbolic transitions  $(l_0, D_0) \rightsquigarrow (l_1, D_1), \dots, (l_n, D_n) \rightsquigarrow (l, D)$  and an  $i$  (standing for the  $i$ -th component of the network) such that  $l[i]$  is an entry node and  $(l_n, D_n) \rightsquigarrow (l, D)$  is derived by an edge  $l_n[i] \xrightarrow{g, f} l[i]$  for some  $g$  and  $r$ .

From the above definition, it should be obvious that we can easily decide whether a reachable symbolic state is a covering state by an on-the-fly algorithm when the entry nodes of all the component automata are known through statical analysis as discussed earlier.

Finally, we claim that the set of covering states of a network covers all its dynamic loops and therefore it suffices to keep them in the passed list for the sake of termination in reachability analysis.<sup>7</sup>

**THEOREM 3** Every dynamic loop of a network of timed automata contains at least one covering state.

**Proof:** Assume a dynamic loop  $L_d = (l_1, D_1) \rightsquigarrow \dots \rightsquigarrow (l_k, D_k)$  with no covering states. However, according to Proposition 1,  $L_d$  contains at least one entry node. Further, assume (without loss of generality) that the symbolic state  $(l, D) \in L_d$  is an entry node and the components  $l[1], \dots, l[m]$  of  $l$  are all in an entry node, and all the other components of  $l$ , i.e.,  $l[m+1], \dots, l[n]$ , are not.

Now, we claim that if  $L_d$  contains no covering states, the set of components  $l_i[1], \dots, l_i[m]$  will remain in an entry node in all symbolic states  $(l_i, D_i) \in L_d$ . Otherwise, if the set of local entry nodes changes, either grows or reduces, it will introduce a covering state. The case of growing is obvious due to the definition for covering states. The argument for the case of reducing is the same as the control nodes of all the

components will reach  $l_1$  again by the end of  $L_d$ , meaning that the set will sooner or later grows again.

In fact, the assumption that  $L_d$  contains no covering states, implies an even stronger property, that is, all symbolic transitions in  $L_d$  are derived by components in  $l_i[m+1], \dots, l_i[n]$ . A transition is derived by a local transition of a component in  $l[1], \dots, l[m]$ , means that the set of local entry nodes will either grow or reduce (discussed above) or the local transition leaves the current entry node and enters an another entry node. The later case implies that the new entry node is a covering state.

Now we construct  $L'_d$  by removing  $l_i[1], \dots, l_i[m]$  from all symbolic states  $(l_i, D_i) \in L_d$ , that is,  $L'_d$  contains only the components that are not in an entry nodes. Obviously, all the symbolic transitions of  $L_d$  are also in  $L'_d$ ; thus  $L'_d$  must be a loop by definition. However,  $L'_d$  contains no components that are in an entry node. This contradicts Proposition 1. ■

An improved reachability algorithm according to the saving strategy induced from Theorem 3 (i.e., saving only the covering sates in the passed list) has been implemented in UPPAAL. Our experimental results show that the space-reduction is between 13% and 72% (see Section 5, Tables 1 and 2).

## 5. Experimental Results

The techniques developed in preceding sections have been implemented and added to the tool UPPAAL (Bengtsson et al., 1996).<sup>8</sup> In this section, we present the results of an experiment where both the original version of UPPAAL and its extension were applied to verify the following six well-studied examples from the literature:

*Philips Audio Protocol* (Audio). The protocol was developed and implemented by Philips to exchange control information between components in audio equipment using Manchester encoding. The correctness of the encoding relies on timing delays between signals. It is first studied and manually verified in Bosscher et al. (1994).

We have verified that the main correctness property holds of the protocol, i.e., all bit streams sent by the sender are correctly decoded by the receiver (Larsen et al., 1995), if the timing error is  $\pm 5\%$ .

*Philips Audio Protocol with Bus Collision* (Audio w. collision). This is an extended variant of Philips audio control protocol with bus collision detection (Bengtsson et al., 1996). It is significantly larger than the version above since several new components (and variables) are introduced, and existing components are modified to deal with bus collisions.

In the experiment we checked that correct bit sequences are received by the receiver (i.e., Property 1 of Bengtsson et al., 1996), using the error tolerances set by Philips.

*Bang & Olufsen Audio/Video Protocol* (Bang & Olufsen). This is an audio control protocol highly dependent on real-time. The protocol is developed by Bang & Olufsen, to

Table 1. Space performance statistics: number of constraints (#) and percentage of current (%).

	Current	Local		Global		Local + global	
	#	#	%	#	%	#	%
Audio	828	219	26	774	93	206	25
Audio w. collision	646,092	198,178	31	370,800	57	111,632	17
Bang & Olufsen	778,288	249,175	32	642,752	83	204,795	26
Box sorter	625	139	22	175	28	36	6
Manufact. plant	92,592	27,042	29	50,904	55	14,933	16
Mutex 2	225	44	20	99	44	18	8
Mutex 3	3,376	621	18	1,360	40	240	16
Mutex 4	56,825	9,352	16	22,125	39	3,532	6
Mutex 5	1,082,916	158,875	15	416,556	38	59,720	6
Train crossing	464	130	28	384	83	114	25

Table 2. Time performance statistics: seconds (s) and percentage of current (%).

	Current	Local		Global		Local + global	
	s	s	%	s	%	s	%
Audio	0.44	0.43	98	0.44	100	0.47	107
Audio w. Collision	3,465.22	2,067.37	60	1,515.88	44	929.22	27
Bang & Olufsen	13,240.49	6,967.38	53	9,348.48	71	4,966.79	38
Box sorter	0.20	0.18	90	0.41	205	0.41	205
Manufact. plant	155.61	39.85	26	56.61	36	24.22	16
Mutex 2	0.13	0.14	108	0.15	115	0.14	108
Mutex 3	1.40	0.67	48	0.65	46	0.51	36
Mutex 4	102.49	24.48	24	25.97	25	12.14	12
Mutex 5	14,790.56	3,299.96	22	3,111.21	21	1,138.32	8
Train crossing	0.19	0.18	95	0.20	105	0.18	95

transmit messages between audio/video components over a single bus, and further studied in Havelund et al. (1997).

In the experiment we have verified the correctness criteria of the protocol. We refer the reader to Section 5.1 of Havelund et al. (1997) for more details.

*Box Sorter* (Box sorter). The example of Larsen et al., (1997) is a model of a sorter unit that sorts red and blue boxes. When the boxes moves down a lane they pass a sensor and a piston. The sorter reads the information from the sensor and sorts out the red boxes by controlling the position of the piston. We have shown, using UPPAAL, that only blue boxes pass the position of the piston and arrive at the end of the lane.

*Manufacturing Plant* (Manufact. plant). The example is a model of the manufacturing plant of Puri and Varaiya (1994) and Daws and Yovine (1995). It is a production cell with a 50 feet belt moving from left to right, two boxes, two robots and a service station. Robot



A moves boxes off the rightmost extreme of the belt to the service station. Robot B moves boxes from the service station to the left-most extreme of the belt.

Assuming an initial distance between the boxes on the belt we verified that no box will fall off the belt.

*Mutual Exclusion Protocol (Mutex 2–5).* It is the so-called Fischer’s protocol that has been studied previously in many experiments, e.g., Abadi and Lamport (1992); Shankar (1993). The protocol is to ensure mutual exclusion among several processes competing for a critical section using timing constraints and a shared variable. In the experiment we use the version of the protocol where a process may recover from failed attempts to enter the critical section, and also eventually leave the critical Section (Kristoffersen et al., 1997).

The protocol is shown to enjoy the invariant property: There is never more than one process existing in the critical section. The results for 2 to 5 processes are shown in Tables 1 and 2.

*Train Crossing Controller (Train crossing).* It is a variant of the train gate controller (Henzinger et al., 1995). An approaching train signals to the controller which reacts by closing the gate. When the train have passed the controller opens the crossing. We have verified that the gate is closed whenever a train is close to the crossing.

In Tables 1 and 2, we present the space (in number of timing constraints stored on the PASSED-buffer) and in the time requirements (in seconds) of the examples on a Sun SPARCstation4 equipped with 64 MB of primary memory. Each example was verified using the current algorithm of UPPAAL (Current), and using modified algorithms for: compact data structure for constraints (Local), control structure reduction (global), and their combination (Local + global).

As shown in Tables 1 and 2 both techniques give truly significant space savings: compact data structure for constraints saves 68–85% of the original consumed space while control structure reduction demonstrates more variation saving 13–72%. Both methods uniformly result in better time-performance on the examples consuming more than half a second, whereas the time-performance is worse on the smaller examples. Most significant is that the two techniques are completely orthogonal, witnessed by the numbers for the combined technique which shows a space-saving between 75% and 94%.

## 6. Conclusion

In this paper, we have two contributions to the development of efficient data structures and algorithms for memory-usage reduction in the automated analysis of timed systems.

First, we have presented a compact data structure, for representing the subsets of Euclidean space that arise during verification of timed automata, which provides minimal and canonical representations for clock constraints, and also allows for efficient inclusion checks between constraint systems. The data structure is based on an  $\mathcal{O}(n^3)$  algorithm which, given a constraint systems over real-valued variables consisting of bounds on

differences, constructs an equivalent system with a minimal number of constraints. It is essentially a minimization algorithm for weighted directed graphs, that extends the transitive reduction algorithm of Aho et al. (1972) to weighted graphs. Given a weighted, directed graph with  $n$  vertices, it constructs in time  $\mathcal{O}(n^3)$  a reduced graph with the minimal number of edges having the same shortest path closure as the original graph.

Second, we have developed an on-the-fly reduction technique to minimize the space-usage by reducing the total number of symbolic states to save in reachability analysis for timed systems. The technique is based on the observation that to ensure termination in reachability analysis, it is not necessary to save all the explored states in memory, but only certain critical states. Based on static analysis of the control structure of timed automata, we are able to compute a set of covering states that cover all the dynamic loops of a system. The set of covering states may not be minimal but sufficient to guarantee termination in an on-the-fly reachability algorithm.

The two techniques and their combination have been implemented in the tool UPPAAL. Our experimental results demonstrate that the techniques result in truly significant space-reductions: For a number of well-studied examples in the literature the space saving is between 75% and 94%, and in all large examples time-performance is improved. Noteworthy is also the observation that the two techniques are completely orthogonal.

As future work, we wish to further study the global on-the-fly reduction technique to identify the minimal sets of covering states that ensure termination and also avoid repeated explorations in reachability analysis for timed systems.

## Notes

1. Several verification tools for timed systems (e.g. UPPAAL, Bengtsson et al., 1996) have been implemented based on this algorithm.
2. For reasons of simplicity and clarity in presentation we have chosen only to consider the non-strict orderings. However, the techniques given extends easily to strict orderings.
3. We assume that  $D$  has been simplified to contain at most one upper and lower bound for each clock and clock-difference.
4. To be precise, it is the inclusion between the solution sets for  $D$  and  $D'$ .
5. This would correspond to constraint systems with empty solution set.
6. “ $<$ ” refers to the assumed ordering on the vertices of  $G$ .
7. Note that this is only a sufficient condition but not necessary.
8. For more information about the tool UPPAAL, see the web site <http://www.uppaal.com/>

## References

- Abadi, M., and Lamport, L. 1992. An old-fashioned recipe for real time. In *Proceedings of REX Workshop Real-Time: Theory in Practice*. LNCS No. 600.
- Aceto, L., Bergueno, A., and Larsen, K. G. 1998. Model checking via reachability testing for timed automata. In Bernard Steffen (ed.), *Proceedings of the 4th Workshop on Tools and Algorithms for the Construction and Analysis of Systems*. LNCS, No. 1384, Berlin: Springer-Verlag, pp. 263–280.
- Aho, A. V., Garey, M. R., and Ullman, J. D. 1972. The transitive reduction of a directed graph. *SIAM Journal on Computing* 1(2): 131–137.
- Alur, R., and Dill, D. 1990. Automata for modelling real-time systems. In *Proc. of Int. Colloquium on Algorithms, Languages and Programming*. LNCS, No. 443, Berlin: Springer-Verlag, pp. 332–335.

- Andersen, H. R. 1995. Partial model checking. In *Proceedings of Symp. on Logic in Computer Science*, pp. 398–407.
- Asarin, E., Maler, O., and Pnueli, A. 1997. Data-structures for the verification of timed automata. In *Proceedings of the International Workshop on Hybrid and Real-Time Systems*.
- Bellman, R. 1957. *Dynamic Programming*. Princeton University Press.
- Bengtsson, J., Griffioen, W. O. D., Kristoffersen, K. J., Larsen, K. G., Larsson, F., Pettersson, P., and Yi, W. 1996. Verification of an audio protocol with bus collision using UPPAAL. In R. Alur and T. A. Henzinger (eds), *Proceedings of the 8th International Conference on Computer Aided Verification*. LNCS, No. 1120, Berlin: Springer-Verlag, pp. 244–256.
- Bengtsson, J., Larsen, K. G., Larsson, F., Pettersson, P., and Yi, W. 1996. UPPAAL in 1995. In *Proceedings of the 2nd Workshop on Tools and Algorithms for the Construction and Analysis of Systems*. LNCS No. 1055, Berlin: Springer-Verlag, pp. 431–434.
- Bosscher, D., Polak, I., and Vaandrager, F. 1994. Verification of an audio-control protocol. In *Proceedings of Formal Techniques in Real-Time and Fault-Tolerant Systems*. LNCS, No. 863.
- Burch, J. R., Clarke, E. M., McMillan, K. L., Dill, D. L., and Hwang, L. J. 1990. Symbolic model checking:  $10^{20}$  states and beyond. In *Proceedings of IEEE Symposium on Logic in Computer Science*.
- Clarke, E. M., Filkorn, T., and Jha, S. 1993. Exploiting symmetry in temporal logic model checking. In *Proceedings of the 5th International Conference on Computer Aided Verification*. LNCS, No. 697.
- Clarke, E. M., Grumberg, O., and Long, D. E. 1992. Model checking and abstraction. In *Principles of Programming Languages*, pp. 450–462.
- Daws C., and Yovine, S. 1995. Two examples of verification of multirate timed automata with KRONOS. In *Proceedings of the 16th IEEE Real-Time Systems Symposium*, pp. 66–75.
- Dill, D. 1989. Timing assumptions and verification of finite-state concurrent systems. In J. Sifakis (ed.), *Proceeding of Automatic Verification Methods for Finite State Systems*. LNCS, No. 407, Berlin. Springer-Verlag, 197–212.
- Emerson E. A., and Jutla, C. S. 1993. Symmetry and model checking. In *Proceedings of the 5th International Conference on Computer Aided Verification*. LNCS, No. 697.
- Godefroid P., and Wolper, P. 1991. A partial approach to model checking. In *Proceedings of IEEE Symposium on Logic in Computer Science*, pp. 406–415.
- Havelund, K., Skou, A., Larsen, K. G., and Lund, K. 1997. Formal modeling and analysis of an audio/video protocol: An industrial case study using UPPAAL. In *Proceedings of the 18th IEEE Real-Time Systems Symposium*, pp. 2–13.
- Henzinger, T. A., Ho, P.-H., and W-Toi, H. 1995 A users guide to HYTECH. Technical Report, Department of Computer Science, Cornell University.
- Henzinger, T. A., Nicollin, X., Sifakis, J., and Yovine, S. 1994. Symbolic Model Checking for Real-Time Systems. *Information and Computation* 111(2): 193–244.
- Holzmann, G. 1991. *The Design and Validation of Computer Protocols*. Prentice Hall.
- Kristoffersen, K. J., Larroussinie, F., Larsen, K. G., Pettersson, P., and Yi, W. 1997. A compositional proof of a real-time mutual exclusion protocol. In *Proceedings of the 7th International Joint Conference on the Theory and Practice of Software Development*, pp. 565–579.
- Larsen, K. G., Pettersson, P., and Yi, W. 1995. Compositional and symbolic model-checking of real-time systems. In *Proc. of the 16th IEEE Real-Time Systems Symposium*, pp. 76–87.
- Larsen, K. G. Pettersson, P., and Yi, W. 1995. Diagnostic model-checking for real-time systems. In *Proceedings of Workshop on Verification and Control of Hybrid Systems III*. LNCS, No. 1066, Berlin, Springer-Verlag, 575–586.
- Larsen, K. G., Pettersson, P., and Yi, W. 1997. UPPAAL in a nutshell. *International Journal on Software Tools for Technology Transfer* 1(1–2): 134–152.
- Pagani, F. 1996. Partial orders and verification of real-time systems. In Bengt Jonsson and Joachim Parrow (eds), *Proceeding of Formal Techniques in Real-Time and Fault-Tolerant Systems*. LNCS, No. 1135, Berlin: Springer-Verlag, pp. 327–346.
- Papadimitriou, C. H. 1994. *Computational Complexity*. Addison-Wesley.
- Puri, A., and Varaiya, P. 1994. Verification of hybrid systems using abstractions. In *Hybrid Systems Workshop*. LNCS, No. 818, Berlin, Springer-Verlag.

- Sedgewick, R. 1993. *Algorithms*. 2nd edn, Addison-Wesley, 1988.
- Shankar, N. 1993. Verification of real-time systems using PVS. In *Proceedings of the 5th International Conference on Computer Aided Verification*. LNCS, No. 697, Berlin, Springer-Verlag.
- Valmari, A. 1990. A stubborn attack on state explosion. *Theoretical Computer Science* 3.
- Vardi, M. Y., and Wolper, P. 1986. An automata-theoretic approach to automatic program verification. In *Proceedings of Symposium on Logic in Computer Science*, pp. 322–331.
- Yannakakis, M., and Lee, D. 1993. An efficient algorithm for minimizing real-time transition systems. In *Proceedings of the 5th International Conference on Computer Aided Verification*, LNCS, No. 697, pp. 210–224.
- Yi, W., Pettersson, P., and Daniels, M. 1994. Automatic verification of real-time communicating systems by constraint-solving. In *Proceedings of the 7th International Conference on Formal Description Techniques*, pp. 223–238.



**Fredrik Larsson** studied engineering physics at Uppsala University, Sweden, and graduated as a M.Sc. in 1996. He then started as a Ph.D. student at the Department of Computer Systems also at Uppsala University. His main research interest was how to build tools for automatic analysis of real-time systems. He was one of the people behind UPPAAL, a tool for verification of systems modelled as timed automata, using reachability analysis and constraint solving techniques. He graduated as a Ph.Lic. in 2000 and is now working at Phoneticom, a company developing products for speech access to the web. He can be contacted by email at [fredrik.larsson@phoneticom.com](mailto:fredrik.larsson@phoneticom.com).



**Kim Guldstrand Larsen** received his Ph.D. in 1985 (from Edinburgh, supervised by Prof. Robin Milner). Since 1993 he holds a Professor position at Aalborg University, and in the period 2000 to 2002 he was a part-time industrial professor at Twente University. He is also Director of the BRICS Research Centre at Aalborg. From 1991 to 1997 Prof. Larsen was a CS member of the Danish Research Council for Natural Science. In 1999, Prof. Larsen was awarded an Honorary Doctorate (*Honoris causa*) from Uppsala University, and he is a member of the Royal Danish Academy of Sciences and Letters since 2000. His area of research is theoretical computer science, in particular validation and verification of embedded systems, real-time and hybrid systems. Prof. Larsen can be contacted by email at [kgl@cs.auc.dk](mailto:kgl@cs.auc.dk).



**Paul Pettersson** obtained his Ph.D. from Uppsala University, Sweden in 1999. His thesis, *Modelling and Verification of Real-Time Systems Using Timed Automata: Theory and Practice*, describes on-the-fly, symbolic, and compositional model-checking techniques for timed systems. From February 1999 to May 2000 he did a Post Doc at Aalborg University in Denmark. Dr. Pettersson now holds a senior lecturer position at the Department of information technology at Uppsala University. His current research concerns modelling, analysis, and testing of real-time and embedded systems. He can be contacted by email at [paupet@docs.uu.se](mailto:paupet@docs.uu.se).



**Wang Yi** is a professor in real time systems at Uppsala University. He obtained his Ph.D. (1991) in Computer Science at Chalmers University of Technology, Sweden. His main research interests and activities are in formal techniques and their application to the design and analysis of real time embedded systems. His main scientific contributions include the initiation and development of timed CCS (a calculus of real time systems) and UPPAAL (a tool for modelling and verification of timed systems).