

Program Termination Analysis in Polynomial Time

Chin Soon Lee*

Datalogisk Institut, Copenhagen University, Denmark

Abstract. Recall the size-change principle for program termination: An infinite computation is *impossible*, if it would give rise to an infinite sequence of size-decreases for some data values. For an actual analysis, *size-change graphs* are used to capture the data size changes in possible program state transitions. A graph sequence that respects program control flow is called a *multipath*. The set of multipaths describe possible state transition sequences. To show that such sequences are finite, it is sufficient that the graphs satisfy *size-change termination (SCT)*: Every infinite multipath has infinite descent, according to the arcs of the graphs. This is an application of the size-change principle.

SCT is decidable, but complete for PSPACE. In this paper, we explore efficient approximations that are sufficiently precise in practice.

For size-change graphs that can loop (i.e., they give rise to an infinite multipath), and that satisfy SCT, it is usual to find amongst them an *anchor*—some graph whose infinite occurrence in a multipath causes infinite descent. The SCT approximations are based on finding and eliminating anchors, as far as possible. If the remaining graphs cannot loop, then SCT is established.

An efficient method is proposed to find anchors among *fan-in free* graphs, as such graphs occur commonly in practice. This leads to a worst-case quadratic-time approximation of SCT. An extension of the method handles general graphs and can detect some rather subtle descent. It leads to a worst-case cubic-time approximation of SCT. Experiments confirm the effectiveness and efficiency of the approximations in practice.

1 Introduction

1.1 Background and Motivation of This Work

This work began as an investigation into the problem of non-termination for the offline partial evaluation of functional programs. To achieve finite function specialization, various techniques have been proposed to ensure that static variables assume boundedly many values during specialization (see [7] in this volume and references therein). The intuition behind these techniques is: Unbounded sequences of increases in a static variable are *impossible*, if they would give rise to unbounded sequences of size-decreases for some bounded-variable values.

* This research is supported by DAEDALUS, the RTD project IST-1999-20527 of the IST Programme within the European Union's Fifth Framework Programme.

The conditions for variable boundedness seen in the literature are complicated. To understand better the forms of descent detected by them, Neil Jones proposed studying a related, but simpler condition for *program termination*, based on the reasoning: Infinite program execution is *impossible*, if it would give rise to an infinite sequence of size-decreases for some data values.

In joint work with Jones and Ben-Amram, the *size-change termination (SCT)* condition, based on the above principle, was formulated and studied for a first-order functional language [12]. SCT is stated in terms of simple structures called *size-change graphs*. They capture the data size changes in possible program state transitions (in this case, function-call transitions), and are obtained by dataflow analysis. SCT is decidable, but complete for PSPACE. In [12], we argued that this called for “sufficiently strong” approximations. They form the topic of the present article.

Fast and scalable methods for termination analysis are not widely discussed in the literature. The need for complicated supporting analyses in a practical analyzer makes exponential-time procedures difficult to avoid. For example, the Prolog analyzers Termilog [13] and Terminweb [3] implement procedures similar to a decision method for SCT, except that they manipulate procedure-call descriptions that capture the instantiatedness of variables as well as argument size relations. The abstract information can be exponentially large in the program size.

Interestingly, when this is *not* the case, fast analysis still cannot be guaranteed. This is because SCT can be tested using Termilog or Terminweb by a straightforward encoding. Such generality is not required in practice. In this article, we develop effective and efficient approximations of SCT. By studying fast supporting analyses, it might then be possible to design practical analyzers based on the approximations. Below are different reasons for this research, connected to the study of automatic program manipulation.

- **Finite function specialization** requires static variables to be bounded. This can be ensured by adapting our techniques [9,10,7].
- **Termination of unfolding** during the offline partial evaluation of functional and logic programs is addressed by our techniques [9,15,7].
- **More aggressive program transformation** may be possible given the termination of certain computations. These computations need not be retained in the transformed program just for their termination behaviour.
- **Choice of evaluation strategy** for some Prolog compilers may be guided by the results of termination analysis [14].
- **Independent verification of machine-generated programs** that are not completely trusted is desirable.

1.2 Generality of SCT

The SCT condition subsumes natural ways to reason about the termination of programs that manipulate well-founded data. It is particularly suited to functional and logic programs, for which the size relations among *inputs* and *outputs*

of subcomputations can be approximated by size analysis. The following untyped functional programs, operating on lists and natural numbers, are all terminating because every infinite sequence of function-call transitions would give rise to a sequence of parameter values whose sizes decrease infinitely. We label the function calls for easy reference.

1. **Non-recursive calls** do not affect size-change termination, although they may cause descent to begin late. Consider call 1 below.

```
rev0(zs) = 1rev1(zs, [])
rev1(xs, acc) = if xs=[] then acc else
               2rev1(tl(xs), hd(xs)::acc)
```

2. **Indirect recursion**, as exhibited by the following tail-recursive program to multiply two numbers, is handled naturally.

```
mult(m,n) = 1loop1(m,0,n,0)
loop1(i1,j1,n1,a1) = if i1=0 then a1 else
                     2loop2(i1-1,n1,n1,a1)
loop2(i2,j2,n2,a2) = if j2=0 then 3loop1(i2,j2,n2,a2) else
                     4loop2(i2,j2-1,n2,a2+1)
```

3. **Nested calls** are handled by *size analysis*. Size analysis deduces that the return value of `sub(m,n)` below always has smaller size than the value of `m`.

```
quot(m,n) = if n=0 then false else
             if m<n then 0 else
             1quot(2sub(m,n),n)+1
sub(a,b) = if b=0 then a else
           3sub(a-1,b-1)
```

4. **Complex descent** in the *program states*: SCT induces a well-founded ordering on the program states in which every possible state transition decreases.

- a) *Lexical descent*. The parameters of `ack` below have sizes that descend lexically in every program state transition. To show termination, program state `ack(m,n)` is ordered below `ack(m',n')` whenever the sizes of `m` and `n` is a lexically smaller pair than the sizes of `m'` and `n'`.

```
ack(m,n) = if m=0 then n+1 else
            if n=0 then 1ack(m-1,1) else
            2ack(m-1, 3ack(m,n-1))
```

In Ex. 2, the sizes of `i` and `j`'s values also descend lexically in every sequence of program state transitions from `loop2` to `loop2`.

- b) *Descent in a sum of parameter sizes*. For the next program, the number `sx + sy + sz`, where `sx`, `sy` and `sz` are the sizes of `x`, `y` and `z`'s values, decreases on every program state transition. Plümer identifies this as an important form of descent [14].

```
p(x,y,z) = if z=[] then x else
            if y=[] then 1p(x,tl(z),y) else
            2p(z,tl(y),x)
```

- c) *Descent in the maximum over parameter sizes.* The number $\max(sx, sy)$, where sx and sy are the sizes of x and y 's values, decreases on every program state transition for the following program. This descent has been observed in a type inference algorithm [5,6].

$$\begin{aligned} q(x,y) = & \text{if } \text{hd}(\text{hd}(x)) \text{ or } \text{hd}(\text{hd}(y)) \text{ then true else} \\ & {}^1q(\text{hd}(\text{tl}(x)), \text{tl}(\text{tl}(x))) \text{ and} \\ & {}^2q(\text{hd}(\text{tl}(y)), \text{tl}(\text{tl}(y))) \end{aligned}$$

- d) *Combining different descent.* It is complicated to express the descent for the next program. Such descent is not expected in practice.

$$\begin{aligned} r(x,y,z,w) = & \text{if } w=[] \text{ then } x \text{ else} \\ & \text{if } y=[] \text{ or } z=[] \text{ then } {}^1r(x,x,1:z, \text{tl}(w)) \text{ else} \\ & {}^2r(\text{tl}(y), y, \text{tl}(z), 1:w) \end{aligned}$$

1.3 This Article

Following our earlier paper [12], we will work with a minimal first-order functional language with eager evaluation. In Sect. 2, we present its syntax and semantics. We define program state transitions (more precisely than in [12]), and relate the existence of an infinite state transition sequence to non-termination. *Size-change graphs* are bipartite graphs that describe the data size changes in possible program state transitions. In Sect. 3, we review the *size-change termination (SCT)* condition for a program's size-change graphs. As SCT is complete for PSPACE, we study viable PTIME approximations in Sect. 4. Section 5 contains some concluding remarks.

2 Programs and Abstraction of Their Behaviour

2.1 Syntax and Notations

Programs in our minimal subject language are generated by the following grammar, where $x, x_i \in \text{Par}$ and $f \in \text{Fun}$ are drawn from mutually exclusive sets of identifiers, and $op \in \text{Op}$ is a primitive operator.

$$\begin{aligned} p \in \text{Prog} & ::= d_1 \dots d_m \\ d, d_i \in \text{Def} & ::= f(x_1, \dots, x_n) = e^f \\ e, e_i, e^f \in \text{Expr} & ::= x \mid \text{if } e_1 \text{ then } e_2 \text{ else } e_3 \mid \\ & \quad op(e_1, \dots, e_n) \mid f(e_1, \dots, e_n) \end{aligned}$$

The *definition* of function f has the form $f(x_1, \dots, x_n) = e^f$, where e^f is called the *body* of f . The number $n \geq 0$ of parameters is its *arity*, written $ar(f)$. Notation: $\text{Param}(f) = \{f^{(1)}, \dots, f^{(n)}\}$ is the set of f parameters. In examples they are named by identifiers. Parameters are assumed to be in scope when they are used. Call expressions are sometimes labelled, e.g., ${}^c f(e_1, \dots, e_n)$. Without loss of generality, function names, parameter names and callsite labels are assumed to be distinct in a program. Constants are regarded as 0-ary operators.

2.2 Programs and Their Semantics

Programs in the subject language are untyped, and interpreted according to the call-by-value semantics in Fig. 4.1. The semantic operator \mathcal{E} is defined as usual for a functional language: $\mathcal{E}[\![e]\!]\mathbf{v}$ is the value of expression e in the environment $\mathbf{v} = (v_1, \dots, v_n)$, a tuple containing the values of parameters $f^{(1)}, \dots, f^{(n)}$.

\mathcal{E} has type $\text{Expr} \rightarrow \text{Val}^* \rightarrow \text{Val}^\sharp$, where Val^* is the set of finite value sequences. Domain $\text{Val}^\sharp = \text{Val} \cup \{\perp, \text{Err}\}$ includes values, plus Err to model runtime errors, and \perp to model non-termination. Function $\text{lift} : \text{Val} \rightarrow \text{Val}^\sharp$ is the natural injection. Any input \mathbf{v} for f is assumed to have length $\text{ar}(f)$.

Definition 1. (*Termination*) Function f of program p terminates on input \mathbf{v} if $\mathcal{E}[\![e^f]\!]\mathbf{v} \neq \perp$. Function f is terminating if it terminates on all inputs. A program p is terminating if all of its functions are terminating.

Definition 2. (*Program states and computations*)

1. A local program state is a triple $(e, f, \mathbf{v}) \in \text{Expr} \times \text{Fun} \times \text{Val}^*$ such that e is a subexpression of e^f and \mathbf{v} has length $\text{ar}(f)$.
2. For local program states (e, f, \mathbf{v}) and (e', g, \mathbf{u}) , write $(e, f, \mathbf{v}) \rightsquigarrow (e', g, \mathbf{u})$ to express a subcomputation. Define the relation \rightsquigarrow by case analysis of e .
 - a) Case $e \equiv \text{if } e_1 \text{ then } e_2 \text{ else } e_3$
 - $(e, f, \mathbf{v}) \rightsquigarrow (e_1, f, \mathbf{v})$
 - $(e, f, \mathbf{v}) \rightsquigarrow (e_2, f, \mathbf{v})$ if $\mathcal{E}[\![e_1]\!]\mathbf{v} = \text{lift } u$ where $u = \text{True}$.
 - $(e, f, \mathbf{v}) \rightsquigarrow (e_3, f, \mathbf{v})$ if $\mathcal{E}[\![e_1]\!]\mathbf{v} = \text{lift } u$ where $u \neq \text{True}$.
 - b) Case $e \equiv \text{op}(e_1, \dots, e_n)$
 - $(e, f, \mathbf{v}) \rightsquigarrow (e_k, f, \mathbf{v})$ if for $1 \leq i < k : \mathcal{E}[\![e_i]\!]\mathbf{v} \notin \{\perp, \text{Err}\}$.
 - c) Case $e \equiv g(e_1, \dots, e_n)$
 - $(e, f, \mathbf{v}) \rightsquigarrow (e_k, f, \mathbf{v})$ if for $1 \leq i < k : \mathcal{E}[\![e_i]\!]\mathbf{v} \notin \{\perp, \text{Err}\}$.
 - $(e, f, \mathbf{v}) \rightsquigarrow (e^g, g, \mathbf{u})$ if each $\mathcal{E}[\![e_i]\!]\mathbf{v} = \text{lift } u_i$ and $\mathbf{u} = (u_1, \dots, u_n)$.
3. A program state is a pair $(f, \mathbf{v}) \in \text{Fun} \times \text{Val}^*$ where \mathbf{v} has length $\text{ar}(f)$.
4. A program state transition (or function-call transition) is a pair of program states $(f, \mathbf{v}) \hookrightarrow (g, \mathbf{u})$ such that there is a sequence of local program states $(e_0, f, \mathbf{v}) \rightsquigarrow (e_1, f, \mathbf{v}) \rightsquigarrow \dots \rightsquigarrow (e_T, f, \mathbf{v})$, where $e_0 = e^f$, expressions e_0, \dots, e_{T-1} are not function calls, e_T is a function call to g , each $(e_t, f, \mathbf{v}) \rightsquigarrow (e_{t+1}, f, \mathbf{v})$ is a local state transition, and $(e_T, f, \mathbf{v}) \rightsquigarrow (e^g, g, \mathbf{u})$.
5. A state transition sequence has the form $(f_0, \mathbf{v}_0) \hookrightarrow (f_1, \mathbf{v}_1) \hookrightarrow (f_2, \mathbf{v}_2) \hookrightarrow \dots$ such that each $(f_t, \mathbf{v}_t) \hookrightarrow (f_{t+1}, \mathbf{v}_{t+1})$ is a program state transition.

Remark: Clearly, neither \rightsquigarrow nor \hookrightarrow are computable relations in general. However, we will be *approximating* properties of the \hookrightarrow relation for the subject program.

Proposition 1. If program p is not terminating, there exists an infinite state transition sequence: $(f_0, \mathbf{v}_0) \hookrightarrow (f_1, \mathbf{v}_1) \hookrightarrow (f_2, \mathbf{v}_2) \hookrightarrow \dots$

2.3 Abstraction of Program Behaviour

Definition 3. (*Data sizes*)

1. We will assume a size function $|\bullet| : Val \rightarrow \mathbb{N}$.
2. For e a subexpression of e^f , we describe e as non-increasing on $f^{(i)}$ if for every input $\mathbf{v} = (v_1, \dots, v_n)$ to f , $\mathcal{E}[\![e]\!]\mathbf{v} = \text{lift } u$ implies $|v_i| \geq |u|$; we describe e as decreasing on $f^{(i)}$ if $|v_i| > |u|$.

Typically, the size of a list is the number of constructors in the list, and the size of a natural number is its value. With this size function, reducing the head of `ls` by computing `hd(hd(ls)) : tl(ls)` is size-decreasing on `ls`.


We will assume, for our examples, that `hd(x)`, `tl(x)` and `x-1` are all size-decreasing on x . This means that `hd([])` and `tl([])` should evaluate to `Err`, rather than `[]`. A semantic consequence: Attempting to apply an unbounded number of size-decreasing operations to a value results in erroneous termination.

The composition of size-decreasing operations is size-decreasing. For instance, `hd(tl(x))` is size-decreasing on x . An expression consisting of just the variable reference x is clearly non-increasing on x .

Definition 4. (*Size-change graphs*)

1. A size-change graph Γ is formally a triple $f \xrightarrow{G} g$, where G is a bipartite graph $(Param(f), Param(g), A)$, and A is a set of arcs of the form $x \xrightarrow{r} y$ with $x \in Param(f)$, $y \in Param(g)$ and $r \in \{\downarrow, \Downarrow\}$. We also loosely write $x \xrightarrow{r} y \in G$. It is assumed that $x \xrightarrow{\downarrow} y$ and $x \xrightarrow{\Downarrow} y$ are not both in G .
2. The finite set of possible size-change graphs for program p is denoted SCG_p .

A set of size-change graphs for Ex. 4a is depicted below. Each `ack` $\xrightarrow{G_i}$ `ack` describes any program state transition due to the function call at c . (Note that for a transition at callsite 2, the nested call at callsite 3 must have terminated successfully.) The graphs `ack` $\xrightarrow{G_1}$ `ack` and `ack` $\xrightarrow{G_2}$ `ack` (formally the same graph) capture function-call transitions where the size of `m` is decreased, as indicated by the arc $m \xrightarrow{\downarrow} m$. The graph `ack` $\xrightarrow{G_3}$ `ack` captures function-call transitions where the size of `m` is not increased, as indicated by the arc $m \xrightarrow{\Downarrow} m$, and the size of `n` is decreased, as indicated by the arc $n \xrightarrow{\downarrow} n$.



Definition 5. A size-change graph describes the program state transitions for which it is safe.

1. Graph $f \xrightarrow{G} g$ is safe for $(f, (v_1, \dots, v_n)) \hookrightarrow (g, (u_1, \dots, u_m))$ if for each $f^{(i)} \xrightarrow{\downarrow} g^{(j)} \in G$, $|v_i| > |u_j|$, and for each $f^{(i)} \xrightarrow{\overline{\downarrow}} g^{(j)} \in G$, $|v_i| \geq |u_j|$.
2. A set of graphs \mathcal{G} is safe for the program p if for every program state transition $(f, \mathbf{v}) \hookrightarrow (g, \mathbf{u})$, there is a graph $f \xrightarrow{G} g$ in \mathcal{G} that is safe for it.

The size-change graphs seen earlier for **ack** are clearly safe for the program.

Sets and sequences of size-change graphs. We may regard a set of size-change graph as arcs of an *annotated digraph*.

Definition 6. (*Annotated digraph*)

1. An annotated digraph $D = (V, A)$ is a pair, consisting of vertex set V , and annotated-arc set A . An annotated arc has the form $u \xrightarrow{a} v$ where $u, v \in V$ and $a \in \Sigma$ for some set of annotations Σ . For any set A , let $\text{vert}(A) = \{u, v \mid (u, v) \in A\}$.
2. Let \rightarrow_D (or just \rightarrow) denote single-step reachability: $u \rightarrow v$ if some $u \xrightarrow{a} v$ is in A . The reachability relation \rightarrow^* is the reflexive transitive closure of \rightarrow .
3. A subgraph of D is some $D' = (V', A')$ such that $V' \subseteq V$ and $A' \subseteq A$.
4. Annotated digraph D is strongly-connected if for all $u, v \in V : u \rightarrow^* v$. A set of arcs A is strongly-connected if $(\text{vert}(A), A)$ is strongly-connected.
5. For a set of arcs A , define $\text{SCC}(A)$ to be the set of maximal strongly-connect subsets, or strongly-connected components (SCCs), of A .

Definition 7. An annotated call graph (ACG) for program p is an annotated digraph whose vertex set is Fun and whose arc set is some $\mathcal{G} \subseteq \text{SCG}_p$.

Definition 8. (*Multipaths*)

1. A multipath \mathcal{M} is a non-empty, finite or infinite sequence of size-change graphs: $f_0 \xrightarrow{G_1} f_1, \dots, f_t \xrightarrow{G_{t+1}} f_{t+1}, \dots$. If the graphs are members of \mathcal{G} , then \mathcal{M} is known as a \mathcal{G} -multipath. We also write $f_0 \xrightarrow{G_1} f_1 \xrightarrow{G_2} f_2 \dots$.
2. A thread th in \mathcal{M} is any finite or infinite connected sequence of size-change arcs: $x_0 \xrightarrow{r_1} x_1, \dots, x_t \xrightarrow{r_{t+1}} x_{t+1}, \dots$ such that there exists $T \geq 0$, for each $d = 1, 2, \dots : x_{d-1} \xrightarrow{r_d} x_d \in G_{T+d}$. We also write $th = x_0 \xrightarrow{r_1} x_1 \xrightarrow{r_2} x_2 \dots$.
3. Thread th has infinite descent if $r_d = \downarrow$ for infinitely many d .
4. Multipath \mathcal{M} has infinite descent if it has a thread with infinite descent.
5. A thread is complete if it starts at the beginning of the multipath, and has the same length as the multipath.

Referring to the picture seen earlier for the size-change graphs of **ack**, consider the multipath $\text{ack} \xrightarrow{G_1} \text{ack} \xrightarrow{G_2} \text{ack} \xrightarrow{G_3} \text{ack}$. The entire picture may be regarded as depicting this multipath. Observe the complete thread: $\mathbf{m} \xrightarrow{\downarrow} \mathbf{m} \xrightarrow{\downarrow} \mathbf{m} \xrightarrow{\overline{\downarrow}} \mathbf{m}$.

Role of size analysis. To obtain a safe set of size-change graphs, we could include a graph for each callsite c , safe for the possible function-call transitions originating there. Suppose that ${}^c g(e_1, \dots, e_n)$ occurs in e^f . We would include some graph $f \xrightarrow{G_\xi} g$. For the arcs of G_c , we include $f^{(i)} \xrightarrow{\downarrow} g^{(j)}$ if e_j is decreasing on $f^{(i)}$ and $f^{(i)} \xrightarrow{\Downarrow} g^{(j)}$ if e_j is non-increasing on $f^{(i)}$, as deduced by size analysis. Without global analysis, we could still obtain a safe set of graphs as follows.

Definition 9. *The natural graphs \mathcal{G}_p for subject program p contains a size-change graph $f \xrightarrow{G_\xi} g$ for each call ${}^c g(e_1, \dots, e_n)$ in e^f . The graph G_c contains the arc $f^{(i)} \xrightarrow{\Downarrow} g^{(j)}$ if e_j is $f^{(i)}$ and $f^{(i)} \xrightarrow{\downarrow} g^{(j)}$ if e_j is a sequence of destructive operators: hd , tl , \dots applied to $f^{(i)}$, e.g., $\text{hd}(\text{tl}(f^{(i)}))$.*




It is also *safe* to incorporate information from the conditions of **if** expressions. The code below has the structure of a program transformed by a binding-time improvement called The Trick [8]. Any transition $(\mathbf{f}, (v_1, v_2)) \hookrightarrow (\mathbf{g}, (u_1))$ is such that $|v_2| > |u_1|$ and $|v_1| = |u_1|$ (since for the transition to exist, $u_1 = \mathcal{E}[\text{hd}(\mathbf{s})](v_1, v_2) = \mathcal{E}[\mathbf{d}](v_1, v_2) = v_1$). This is a natural way for “fan-ins” (multiple arcs connected to the same parameter) to arise in our framework.

Example 5.

```

g(x) = if x=0 then 0 else
      1f(x-1, [0,1,2])+1

f(d,s) = if d=hd(s) then
          3g(hd(s)) else
          2f(d,tl(s))
    
```



3 Size-Change Termination

Definition 10. *A set of size-change graphs \mathcal{G} satisfies size-change termination (SCT) if every infinite \mathcal{G} -multipath has infinite descent.*

Theorem 1. *Let \mathcal{G} be safe for p . If \mathcal{G} satisfies SCT, then p is terminating.*

Proof. Suppose that \mathcal{G} satisfies SCT, but p is *not* terminating. By Proposition 1, there exists an infinite transition sequence $(f_0, \mathbf{v}_0) \hookrightarrow (f_1, \mathbf{v}_1) \hookrightarrow (f_2, \mathbf{v}_2) \dots$. By the safety of \mathcal{G} , there exists a \mathcal{G} -multipath \mathcal{M} : $f_0 \xrightarrow{G_1} f_1 \xrightarrow{G_2} f_2 \dots$ such that each $f_t \xrightarrow{G_{t+1}} f_{t+1}$ is safe for $(f_t, \mathbf{v}_t) \hookrightarrow (f_{t+1}, \mathbf{v}_{t+1})$. By SCT, \mathcal{M} has a thread th with infinite descent: $x_0 \xrightarrow{r_1} x_1 \xrightarrow{r_2} x_2 \dots$ such that for some T , and $d = 1, 2, \dots$, it holds that $x_{d-1} \xrightarrow{r_d} x_d \in G_{T+d}$. Define a value sequence u_0, u_1, \dots corresponding to th : Suppose x_i is $f_{T+i}^{(k)}$, and $\mathbf{v}_{T+i} = (v_1, \dots, v_n)$, then let $u_i = v_k$. By safety, the arc from x_{d-1} to x_d implies $|u_{d-1}| > |u_d|$ if $r_d = \downarrow$, and $|u_{d-1}| \geq |u_d|$ if $r_d = \Downarrow$. Since th has infinite descent, infinitely many of the r_d are \downarrow . This implies a sequence of data values whose sizes decrease infinitely, which is impossible. \square

Definition 11. The size-change graphs $f_0 \xrightarrow{G_1} f_1$ and $f_1 \xrightarrow{G_2} f_2$ can be composed to summarize the size changes among the variable values at the start and end of a state transition sequence modelled by $f_0 \xrightarrow{G_1} f_1 \xrightarrow{G_2} f_2$. Formally, the composition $f_0 \xrightarrow{G_1} f_1; f_1 \xrightarrow{G_2} f_2 = f_0 \xrightarrow{G} f_2$, where $G = G^\downarrow \cup G^\mp$ and

$$\begin{aligned} G^\downarrow &= \{x \xrightarrow{\downarrow} z \mid x \xrightarrow{r_1} y \in G_1, y \xrightarrow{r_2} z \in G_2, r_1 \text{ or } r_2 \text{ is } \downarrow\} \\ G^\mp &= \{x \xrightarrow{\mp} z \mid x \xrightarrow{\mp} y \in G_1, y \xrightarrow{\mp} z \in G_2, x \xrightarrow{\downarrow} z \notin G^\downarrow\} \end{aligned}$$

Definition 12. For $\mathcal{G} \subseteq SCG_p$, the composition closure of \mathcal{G} , denoted $\overline{\mathcal{G}}$, is the smallest set satisfying: $\overline{\mathcal{G}} = \mathcal{G} \cup \{f_0 \xrightarrow{G_1} f_1; f_1 \xrightarrow{G_2} f_2 \mid f_0 \xrightarrow{G_1} f_1, f_1 \xrightarrow{G_2} f_2 \in \mathcal{G}\}$.

Theorem 2. ([12]) \mathcal{G} satisfies SCT iff for every size-change graph $f \xrightarrow{G} f$ in $\overline{\mathcal{G}}$ such that $f \xrightarrow{G} f; f \xrightarrow{G} f = f \xrightarrow{G} f$, we have $x \xrightarrow{\downarrow} x \in G$ for some x .

It follows that SCT is decidable. Unfortunately, SCT is *intrinsically hard*.

Theorem 3. ([12]) SCT satisfaction is complete for PSPACE.

4 Polynomial-Time Approximations

The PSPACE-hardness result is motivation for studying viable polynomial-time approximations. The approximations we propose are centered around the ACG.

Definition 13. Let S be any finite set. The infinity set of an infinite sequence $s_1 s_2 \dots$ of S elements is the set $\{s \mid \{i \mid s_i = s\} \text{ is infinite}\}$.

Definition 14. Let \mathcal{H} be a set of size-change graphs. Call $f \xrightarrow{G} g \in \mathcal{H}$ an anchor for \mathcal{H} if every \mathcal{H} -multipath whose infinity set contains $f \xrightarrow{G} g$ has infinite descent. \mathcal{H} has an anchor means it contains some graph that is an anchor for it.

Lemma 1. Suppose that every non-empty strongly-connected $\mathcal{H} \subseteq \mathcal{G}$ has an anchor, then \mathcal{G} satisfies SCT.

Lemma 2. Suppose that \mathcal{H} is strongly-connected, and $\mathcal{A} \subseteq \mathcal{H}$ is a non-empty set of anchors for \mathcal{H} . If some non-empty strongly-connected subset \mathcal{H}^* of \mathcal{H} is without anchors, then \mathcal{H}^* is a subset of a member of $SCC(\mathcal{H} \setminus \mathcal{A})$.

Lemma 2 leads to the following procedure to approximate SCT for \mathcal{G} , provided we have some means to determine anchors for a strongly-connected \mathcal{H} . The procedure *SCP* (size-change polytime) below is called with each $\mathcal{H} \in SCC(\mathcal{G})$.

Procedure $SCP(\mathcal{H})$

1. Determine anchors \mathcal{A} for \mathcal{H} .
2. If \mathcal{A} is empty, fail.
3. Otherwise call SCP on each member of $SCC(\mathcal{H} \setminus \mathcal{A})$.

The procedure terminates, as \mathcal{H} is reduced on each recursive invocation. If it does not fail, argue that \mathcal{G} satisfies SCT as follows. If \mathcal{G} does *not* satisfy SCT, Lemma 1 guarantees a non-empty strongly-connected \mathcal{H}^* without anchors. At the start of the procedure, \mathcal{H} is equal to \mathcal{G} , so it includes \mathcal{H}^* . For each invocation of the procedure, if \mathcal{H} includes \mathcal{H}^* , the input to some recursive call of SCP is a superset of \mathcal{H}^* by Lemma 2. Therefore, if the procedure does not fail, \mathcal{H}^* is eventually discovered, and this would cause it to fail.

4.1 Deducing Anchors

For this, we introduce a few notions centered around the parameter set. Forward (backward) *thread preservers* for \mathcal{H} are parameters that are always forward (backward) connected to other thread preservers when they occur among the source (target) parameters in any graph of \mathcal{H} .

Definition 15. (*Thread preservers*) Let $P \subseteq \text{Par}$. For size-change graphs \mathcal{H} , define $\overrightarrow{TP}_P(\mathcal{H})$, $\overleftarrow{TP}_P(\mathcal{H})$ to be the largest sets such that:

$$\begin{aligned} \overrightarrow{TP}_P(\mathcal{H}) &= \{x \in P \mid \forall f \xrightarrow{G} g \in \mathcal{H} : \\ &\quad x \in \text{Param}(f) \Rightarrow \exists y \in \overrightarrow{TP}_P(\mathcal{H}) : x \xrightarrow{r} y \in G \text{ for some } r\} \\ \overleftarrow{TP}_P(\mathcal{H}) &= \{y \in P \mid \forall f \xrightarrow{G} g \in \mathcal{H} : \\ &\quad y \in \text{Param}(g) \Rightarrow \exists x \in \overleftarrow{TP}_P(\mathcal{H}) : x \xrightarrow{r} y \in G \text{ for some } r\} \end{aligned}$$

The subscript P is omitted if it is the set Par . $\overrightarrow{TP}_P(\mathcal{H})$ and $\overleftarrow{TP}_P(\mathcal{H})$ are called respectively the forward and backward preservers of \mathcal{H} with respect to P .

The *size-relation graph due to \mathcal{H}* is the digraph with vertex set Par and arcs corresponding to all the arcs in the size-change graphs of \mathcal{H} .

Definition 16. (*SRG*) The size-relation graph (SRG) due to \mathcal{H} , denoted $SRG_{\mathcal{H}}$, has the annotated arcs $\{x \xrightarrow{r}_{\Gamma} y \mid \Gamma = f \xrightarrow{G} g \in \mathcal{H}, x \xrightarrow{r} y \in G\}$.

Deducing anchors for fan-in free graphs.

Definition 17. (*Fan-in free*)

1. A size-change graph $f \xrightarrow{G} g$ is fan-in free if $x \xrightarrow{r} y, x' \xrightarrow{r'} y \in G \Rightarrow x = x'$.
2. A set of size-change graphs \mathcal{H} is fan-in free if all its graphs are fan-in free.

Experiments using Terminweb's size analysis indicate that both its polyhedron analysis and the simpler analysis using monotonicity and equality constraints frequently generate fan-in free graphs. This is despite the fact that equality constraints among variables (due to unifications) are common in Prolog.

Theorem 4. *Let \mathcal{H} be fan-in free and strongly-connected. If for $f \xrightarrow{G} g$ in \mathcal{H} , there exists $x \xrightarrow{\downarrow} y \in G$ with $x, y \in \overrightarrow{TP}(\mathcal{H})$, then $f \xrightarrow{G} g$ is an anchor for \mathcal{H} .*

Proof. As \mathcal{H} is strongly-connected, there is a multipath $f_0 \xrightarrow{G_1} f_1 \dots \xrightarrow{G_{n+1}} f_{n+1}$ with $f_0 = f_{n+1}$, containing every graph of \mathcal{H} . Let $P_i = \text{Param}(f_i) \cap \overrightarrow{TP}(\mathcal{H})$ for each i . By definition of $\overrightarrow{TP}(\mathcal{H})$ and the assumption that \mathcal{H} is fan-in free, $|P_i| \leq |P_{i+1}|$. Therefore $|P_0| \leq \dots \leq |P_{n+1}| = |P_0|$. Let $K = |P_0|$.

Consider any infinite \mathcal{H} -multipath $\mathcal{M} = f_0 \xrightarrow{G_1} f_1 \xrightarrow{G_2} f_2 \dots$. As before, let $P_t = \text{Param}(f_t) \cap \overrightarrow{TP}(\mathcal{H})$ for each t . Then each P_t has cardinality K . It is easy to see that the arcs from each P_t to P_{t+1} form K complete threads over \mathcal{M} . Suppose that $f \xrightarrow{G} g$ is in the infinity set of \mathcal{M} . By the premise, there exists some $x \xrightarrow{\downarrow} y \in G$ with $x, y \in \overrightarrow{TP}(\mathcal{H})$. This arc appears infinitely often among the K threads. Consequently, some thread has infinite descent. \square

Remark: The above proof gives insight into the form of descent inhibiting unbounded repetitions of $f \xrightarrow{G} g$ in any \mathcal{H} -multipath representing a state transition sequence. The sum of P_t parameter sizes is never increased, and strictly decreased for a transition described by $f \xrightarrow{G} g$, due to the decreasing arc in G .

Lemma 3. (Ben-Amram) *Let $\mathcal{H} \subseteq \mathcal{G}$ be non-empty, fan-in free and strongly-connected. And let N and M be the space requirement for \mathcal{G} and \mathcal{H} . Then $\overrightarrow{TP}(\mathcal{H})$ can be computed in time $O(M)$, with an $O(N^2)$ initialization performed on \mathcal{G} .*

The algorithm for computing $\overrightarrow{TP}(\mathcal{H})$ maintains a set of counts for the source parameters in each size-change graph of \mathcal{H} . These counts initially record the parameter out-degrees. Start by marking every parameter whose count is 0, and adding these parameters to a worklist. When processing x from the worklist, inspect those graphs with arcs connected to x . For such an arc $x' \xrightarrow{r} x$, if x' is not yet marked, reduce its count. If the count becomes 0, mark x' and insert it into the worklist. Terminate when the worklist is empty. We claim, without proof, that upon termination, the unmarked parameters form $\overrightarrow{TP}(\mathcal{H})$, and that the algorithm has $O(M)$ time-complexity, with an $O(N^2)$ initialization performed on \mathcal{G} . (Initialization is needed to set up the necessary indexing structures.)

The procedure *SCP* considers each graph of \mathcal{G} in no more than $|\mathcal{G}| \sim O(N)$ invocations. Each invocation of the procedure involves computing some $\overrightarrow{TP}(\mathcal{H})$, working out \mathcal{A} and computing $SCC(\mathcal{H} \setminus \mathcal{A})$. All these steps have linear time complexity. Therefore we have an $O(N^2)$ procedure to approximate SCT for fan-in free \mathcal{G} . We refer to the resulting approximation of SCT as SCP1.

An example. Consider the graphs \mathcal{G} for **ack** seen earlier. Now, \mathcal{G} is strongly-connected. We work out that $\overrightarrow{TP}(\mathcal{G}) = \{\mathbf{m}\}$ (\mathbf{m} is connected to \mathbf{m} in each graph of \mathcal{G}). We deduce that $\text{ack} \xrightarrow{G_1} \text{ack}$ and $\text{ack} \xrightarrow{G_2} \text{ack}$ are anchors for \mathcal{G} , as they have the arc $\mathbf{m} \xrightarrow{\downarrow} \mathbf{m}$. Intuitively, this means their occurrence in the infinity set of a \mathcal{G} -multipath would give the multipath infinite descent due to the decreasing arc.


Removing these graphs from consideration leaves $\mathcal{H} = \{\mathbf{ack} \xrightarrow{G_3} \mathbf{ack}\}$. We work out that $\overrightarrow{TP}(\mathcal{H}) = \{\mathbf{m}, \mathbf{n}\}$, and deduce that $\mathbf{ack} \xrightarrow{G_3} \mathbf{ack}$ is an anchor for \mathcal{H} , as it has the arc $\mathbf{n} \xrightarrow{\downarrow} \mathbf{n}$. We conclude that \mathcal{G} satisfies SCT.

In general, if \mathcal{G} satisfies SCP1, then in every function-call transition sequence of the form $(f, \mathbf{v}) \hookrightarrow^* (f, \mathbf{u})$ for the subject program, the parameter tuples \mathbf{v} and \mathbf{u} exhibit lexical descent in various sums of parameter sizes. SCP1 subsumes most practical forms of descent, including those of Ex. 1 to Ex. 4b in Sect. 1.2.

General graphs considered. The above method to deduce anchors suffers from obvious defects. To begin with, there does not appear to be any simple extension to cope with fan-ins. We consider the issue in greater detail.

Example 6.

$\mathbf{f}(\mathbf{x}, \mathbf{y}) = \text{if } \mathbf{y}=\mathbf{tl}(\mathbf{x}) \text{ then } {}^1\mathbf{f}(\mathbf{x}, \mathbf{y}) \text{ else } 0$



Now \mathcal{G} is strongly-connected, $\overrightarrow{TP}(\mathcal{G}) = \{\mathbf{x}, \mathbf{y}\}$, and in the graph $\mathbf{f} \xrightarrow{G_1} \mathbf{f}$, G_1 contains $\mathbf{x} \xrightarrow{\downarrow} \mathbf{y}$, but $\mathcal{M} = \mathbf{f} \xrightarrow{G_1} \mathbf{f} \xrightarrow{G_1} \mathbf{f} \dots$ has no infinite descent. Indeed, \mathcal{M} models the state transition sequence: $(\mathbf{f}, ([1], [])) \hookrightarrow (\mathbf{f}, ([1], [])) \hookrightarrow \dots$. Intuitively, when size-change graphs are not fan-in free, a descending arc between \overrightarrow{TP} parameters does not necessarily constitute “progress.” In the example, the arc $\mathbf{x} \xrightarrow{\downarrow} \mathbf{y}$ in G_1 indicates a size relationship between the values of \mathbf{x} and \mathbf{y} that persists throughout the state transition sequence.

It seems sensible to consider the thread behaviour in $\{\mathbf{x}\}$ and $\{\mathbf{y}\}$ separately. A little reflection reveals that arcs not in the same SCC of the SRG due to \mathcal{G} cannot give rise to anchor graphs. Therefore, we might define the following *critical parameter sets*, to be inspected independently for descent.

Definition 18. Define the critical parameter sets for size-change graphs \mathcal{H} as: $\text{Critical}(\mathcal{H}) = \{P \mid \overleftarrow{TP}_P(\mathcal{H}) = P, P \text{ strongly-connected in } \text{SRG}_{\mathcal{H}}, P \text{ maximal}\}$.

If P is a critical set of \mathcal{H} , then at the end of a state transition sequence described by any \mathcal{H} -multipath, the value of each P parameter has size no greater than the initial value of *some* P parameter. We will look for descent among such parameters. Observe that for fan-in free graphs, $\overleftarrow{TP}(\mathcal{H})$ at least includes $\overrightarrow{TP}(\mathcal{H})$. The use of a larger set of parameters makes it possible to detect more descent, but also means more complicated deductions (see Theorem 5).

Deducing anchors for general graphs.

Theorem 5. Let \mathcal{H} be strongly-connected. Then $f \xrightarrow{G} g$ in \mathcal{H} is an anchor for \mathcal{H} if there exists $P \in \text{Critical}(\mathcal{H})$ such that one of the following holds.

1. For $f_0 \xrightarrow{G_1} f_1$ in \mathcal{H} : $x \xrightarrow{r} y, x \xrightarrow{r'} y' \in G_1$ and $x, y, y' \in P$ implies $y = y'$ (i.e. no “fan-outs” among $x \in P$) AND there exists $x \xrightarrow{\downarrow} y \in G$ with $x, y \in P$.
2. The digraph with arcs $\{x \xrightarrow{\Gamma} y \in A \mid x, y \in P\}$, where A are $SRG_{\mathcal{H}}$ arcs, has no strongly-connected subgraph with an arc $x \xrightarrow{\Gamma} y$ where $\Gamma = f \xrightarrow{G} g$.

Refer to [11] for a proof. The first condition detects the same form of descent as the condition for fan-in free graphs in Theorem 4. When the backward thread-preserving P contains fan-outs, it is difficult to know whether the decreasing arcs among P parameters are significant for descent. The second condition implements a weak principle for identifying anchors: Any infinite \mathcal{H} -multipath has an infinite thread th remaining in just the P parameters (easily shown), so if there are no *non-descending* threads remaining infinitely among $x \in P$ in any multipath whose infinity set has $f \xrightarrow{G} g$, then th is infinitely descending in any multipath whose infinity set *does* contain $f \xrightarrow{G} g$, i.e., this graph is an anchor for \mathcal{H} . This reasoning is effective when there are sufficiently many decreasing arcs among $x \in P$. For example, the natural graphs of Ex. 4c in Sect. 1.2 are handled. Note the rather subtle descent in this example.

Lemma 4. *Let $\mathcal{H} \subseteq \mathcal{G}$ be non-empty and strongly-connected. Let N and M be the space requirement for \mathcal{G} and \mathcal{H} . Then $Critical(\mathcal{H})$ can be computed in time $O(M)$ for fan-in free \mathcal{H} and $O(M^2)$ in general, with an $O(N^2)$ initialization performed on \mathcal{G} .*

The critical sets of \mathcal{H} are computed as follows. Initialize \wp to $\{\overleftarrow{TP}(\mathcal{H})\}$. In the iterative step, replace each $P \in \wp$ with possibly several non-empty $\overleftarrow{TP}_{P'}(\mathcal{H})$, where P' are strongly-connected components in the SRG due to \mathcal{H} , restricted to P . This is repeated until no further change occurs. The calculation of $\overleftarrow{TP}_{P'}(\mathcal{H})$ is similar to the calculation of $\overrightarrow{TP}(\mathcal{H})$. The total \overleftarrow{TP} calculations for all P' in the current \wp can be accomplished in time $O(M)$, so the critical sets of \mathcal{H} can be worked out in time $O(M^2)$, as there are at most M iterative steps. The required initialization can be performed in time $O(N^2)$. To see that critical sets can be computed in time $O(M)$ for fan-in free graphs, we require the following.

Lemma 5. ([11]) *Let \mathcal{H} be non-empty, fan-in free and strongly-connected. If $P = \overleftarrow{TP}(\mathcal{H})$ and P' is a strongly-connected component in the SRG due to \mathcal{H} , restricted to P , then either $\overleftarrow{TP}_{P'}(\mathcal{H}) = \{\}$ or $\overleftarrow{TP}_{P'}(\mathcal{H}) = P'$.*

Thus no iteration of $\overleftarrow{TP}_{P'}(\mathcal{H})$ calculations is necessary for fan-in free graphs. The computation of critical sets is the dearest part of the procedure to deduce anchors. The conditions of Theorem 5 can be checked in time $O(M)$. It follows that the *SCP* procedure can be performed in time $O(N^2)$ for fan-in free graphs, and $O(N^3)$ in general. We refer to the resulting approximation of SCT as *SCP2*.

An example with fan-ins. Consider the graphs \mathcal{G} of Ex. 5 seen earlier. The initial step to determine anchors for \mathcal{G} finds the only critical set $\{\mathbf{x}, \mathbf{d}\}$. This blocks out \mathbf{s} , which cannot contribute to infinite descent in any multipath whose infinity set contains $\mathbf{g} \xrightarrow{G_1} \mathbf{f}$ or $\mathbf{f} \xrightarrow{G_3} \mathbf{g}$ (any thread visiting \mathbf{s} at some point in such a multipath is eventually continued to the $\{\mathbf{x}, \mathbf{d}\}$ component of the SRG, never to return to \mathbf{s}). By either condition in Theorem 5, $\mathbf{g} \xrightarrow{G_1} \mathbf{f}$ is determined to be an anchor for \mathcal{G} . Next, consider the strongly-connected $\{f \xrightarrow{G_2} f\}$. The critical parameter sets here are $\{\mathbf{d}\}$ and $\{\mathbf{s}\}$. The second critical set allows us to conclude that $f \xrightarrow{G_2} f$ is an anchor. It follows that \mathcal{G} satisfies SCT.

Domains	$u, v, v_i \in Val$ $w, w_i \in Val^\sharp = Val \cup \{\perp, Err\}$, where $\perp \sqsubseteq w$ for all w .
Types	$\mathcal{E} : Expr \rightarrow Val^* \rightarrow Val^\sharp$ $\mathcal{O} : Op \rightarrow Val^* \rightarrow Val^\sharp$ $lift : Val \rightarrow Val^\sharp$ (the natural injection) $strictapp : (Val^* \rightarrow Val^\sharp) \rightarrow (Val^\sharp)^* \rightarrow Val^\sharp$
Semantic operator	$\mathcal{E}[\![f^{(i)}]\!](v_1, \dots, v_n) = lift\ v_i$ $\mathcal{E}[\![\text{if } e_1 \text{ then } e_2 \text{ else } e_3]\!]\mathbf{v} = \mathcal{E}[\![e_1]\!]\mathbf{v} \rightarrow \mathcal{E}[\![e_2]\!]\mathbf{v}, \mathcal{E}[\![e_3]\!]\mathbf{v}$ $\mathcal{E}[\![op(e_1, \dots, e_n)]\!]\mathbf{v} = strictapp\ (\mathcal{O}[\![op]\!])\ (\mathcal{E}[\![e_1]\!]\mathbf{v}, \dots, \mathcal{E}[\![e_n]\!]\mathbf{v})$ $\mathcal{E}[\![f(e_1, \dots, e_n)]\!]\mathbf{v} = strictapp\ (\mathcal{E}[\![e^f]\!])\ (\mathcal{E}[\![e_1]\!]\mathbf{v}, \dots, \mathcal{E}[\![e_n]\!]\mathbf{v})$
Auxiliary	$w_1 \rightarrow w_2, w_3 =$ $\begin{cases} w_1, & \text{if } w_1 \in \{\perp, Err\}, \\ w_2, & \text{if } w_1 = lift\ u \text{ where } u = True, \\ w_3, & \text{if } w_1 = lift\ u \text{ where } u \neq True. \end{cases}$ $strictapp\ \psi(w_1, \dots, w_n) =$ $\begin{cases} \psi(v_1, \dots, v_n) & \text{if each } w_i = lift\ v_i; \text{ else} \\ w_i & \text{where } i = \text{least index such that } w_i \in \{\perp, Err\} \end{cases}$
Assume	$\mathcal{O}[\![op]\!]\mathbf{v} \neq \perp$

Fig. A.1. Program semantics. *True* is a distinguished element of *Val*

4.2 Assessment of the SCT Approximations

Lemma 6. ([11]) *Let \mathcal{H} be non-empty, fan-in free and strongly-connected. For $x, y \in \overrightarrow{TP}(\mathcal{H})$, if there exists $f \xrightarrow{G} g$ in \mathcal{H} such that $x \xrightarrow{r} y \in G$ for some r , there is a critical set P such that $x, y \in P$ and $\overrightarrow{TP}_P(\mathcal{H}) = P$.*

Ref	program	·	size	<i>SCT</i> (ms)	<i>SCP1</i> (ms)	<i>SCP2</i> (ms)	procedure(s)
[2]	read	ts	56 N	1103.53	-	- N	4.68 read_expt1 ...
[2]	qplan	ts	28 N	25.10	-	- N	2.35 schedule_plan ...
[2]	tictactoe	ts	27 N	12.37	-	- N	1.99 alpha_beta ...
[13]	vangelder	ts	24 Y	59.66	-	- N	2.37 q t r p
[2]	rdtok	ts	23 N	68.94	N	1.22 N	2.49 read_tokens ...
[2]	warplan	ts	22 N	13.25	-	- N	1.62 plan_achieve ...
[13]	sicstus3	ts	15 N	8.89	N	0.51 N	1.20 put_assoc
[2]	rdtok	ts	15 N	7.49	-	- N	1.17 read_string ...
[2]	aiakl	ts	15 Y	7.47	Y	0.04 Y	0.07 intersect
[2]	qplan	ts	14 Y	1.96	Y	0.17 Y	0.36 incorporate ...
[2]	serialize	ts	12 Y	1.66	Y	0.02 Y	0.05 split0
[2]	ann	ts	12 Y	9.67	-	- Y	0.08 collect_info
[13]	sicstus3	ts	11 N	7.00	N	0.45 N	0.95 get_assoc
[2]	browse	ts	11 N	20.19	N	0.23 N	1.10 get_pats
[2]	browse	ts	10 N	20.94	N	0.34 N	0.83 init
[2]	ann	ts	9 Y	2.57	Y	0.03 Y	0.09 intersect
[2]	ann	ts	9 Y	2.51	Y	0.03 Y	0.08 merge
[13]	yale_s_p	ll	8 Y	1.40	-	- Y	0.23 ab_holds
[2]	warplan	ts	8 Y	0.96	Y	0.03 Y	0.05 retract1
[2]	qplan	ts	8 Y	0.52	-	- Y	0.08 best_goal
[2]	hanoiapp	ts	8 Y	19.47	Y	0.03 Y	0.07 shanoi
[2]	ann	ts	8 N	2.49	N	0.39 N	0.88 un_number_vars_2
[13]	arit_exp	ts	7 Y	18.52	Y	0.14 Y	0.43 e g f
[14]	pl8.4.2	ts	7 Y	6.38	Y	0.17 Y	0.50 e n t
[2]	tak	ts	7 N	4.00	N	0.38 N	0.78 tak
[2]	rdtok	ts	7 N	7.93	N	0.29 N	0.70 read_digits
[2]	qplan	ts	7 N	5.86	N	0.38 N	0.88 variablise
[2]	qplan	ts	7 N	3.85	N	0.38 N	0.82 quantificate
[2]	qplan	ts	7 N	5.25	-	- N	0.78 cost
[2]	boyer	ts	7 N	1.72	N	0.38 N	0.81 tautology
[13]	ack	ts	6 Y	1.99	Y	0.02 Y	0.06 ack
[14]	mergesort_t	ll	6 Y	0.56	Y	0.09 Y	0.23 split2 split
[2]	warplan	ts	6 N	1.93	N	0.28 N	0.62 mkground ...
[2]	tictactoe	ts	6 N	0.60	N	0.16 N	0.74 choose_legal
[2]	qplan	ts	6 Y	2.74	Y	0.03 Y	0.08 mark
[2]	qplan	ts	6 N	5.92	N	0.37 N	0.85 variables
[2]	qplan	ts	6 Y	0.64	Y	0.04 Y	0.10 instantiate ...
[2]	peephole	ll	6 Y	2.56	Y	0.04 Y	0.20 popt3 popt31
[2]	peephole	ll	6 Y	1.01	Y	0.09 Y	0.21 popt1a popt1a1
[2]	peephole	ll	6 Y	1.01	Y	0.04 Y	0.09 popt2 popt21
[2]	hanoiapp	ts	6 N	16.62	N	0.20 N	0.73 shanoi
[2]	fib_t	ts	6 Y	1.70	Y	0.03 Y	0.07 add
[2]	ann	ts	6 N	5.99	N	0.39 N	0.87 numbervars_2
[2]	ann	ts	6 Y	0.94	Y	0.03 Y	0.09 subtract

Fig. B.1. Experiments conducted on a P3/800 running Redhat Linux 6.1

Lemma 7. *Let SCT , $SCP1$ and $SCP2$ denote those \mathcal{G} satisfying SCT , $SCP1$ and $SCP2$ respectively. Then $SCP1 \subset SCP2 \subset SCT$.*

Proof. It follows from Lemma 6 that any anchor deduced for some strongly-connected, fan-in free \mathcal{H} by the condition of Theorem 4 is also deduced to be an anchor by the first condition of Theorem 5. Therefore, $SCP1 \subseteq SCP2$. The other containment $SCP2 \subseteq SCT$ is clear. The containments are strict: $SCP1$ does not handle the natural graphs for Ex. 4c of Sect. 1.2, whereas $SCP2$ does; $SCP2$ does not handle the natural graphs for Ex. 4d, whereas SCT does. \square

Experimental results. Finally, we report some empirical findings. The size-change graphs used for our experiments have been generated using Terminweb’s size analysis [3], applied to the test suites in [14,1,2,4,13]. For each program, a set of size-change graphs \mathcal{G} is compiled using either the *list-length norm* (ll) or the *term-size norm* (ts). Instead of timing the analysis of each \mathcal{G} , we consider their SCCs, since our analyses scale linearly in this number. For each SCC, we time 1000 runs each of the SCT , $SCP1$ and $SCP2$ tests. The table in Fig. 4.1 shows the benchmark reference, program name, size norm used, total number of arcs in the tested component’s graphs (the problem size), results and timings of the various tests, and names of some procedures in the component.

A total of 103 programs with 281 components have been analyzed. For our experiments, the $SCP2$ test has given the same results as SCT in all but one instance, a contrived example due to Van Gelder. Apart from 19 components where the $SCP1$ test does not apply due to fan-ins (this situation is indicated by - in the table), it has also given the same results, suggesting that SCT is “overkill” in practice. Timing-wise, $SCP1$ is consistently and significantly faster than SCT , when it is applicable. For 19 small components, $SCP2$ is actually slower than SCT (probably due to its more difficult implementation), but it gives up quickly on complicated examples that eventually cause SCT to fail. For our table, we have simply ordered the experiments according to problem size, and listed the results at the top.

5 Concluding Remarks

We have reviewed the size-change condition for program termination. In a previous article, we established that SCT is complete for PSPACE. In this article, we have explored, theoretically and empirically, polynomial-time approximations of SCT . In particular, we have proposed two tests that may establish SCT . The first is simpler but restricted to fan-in free graphs. It has worst-case quadratic-time complexity and is very fast in practice. The general test has worst-case cubic-time complexity and accurately predicts SCT in practice.

There remain a number of obstacles to the analysis of “real programs.”

- **For functional programs:** We require a convincing extension of the basic principles to higher-order programs.

- **For Prolog programs:** Fast supporting analyses are needed. They have to be combined with SCP2 sensibly to produce a useful tool for Prolog.
- **For imperative programs:** Research is on-going. It is hard to deal with imperative programs, as they do not use well-founded data. Further, sequential programming encourages nested loops that maintain complex invariants among variable values. This makes the job of size analysis difficult.

References

1. K. R. Apt and D. Pedreschi. Modular termination proofs for logic and pure prolog programs. In *Advances in Logic Programming Theory*, pages 183–229. Oxford University Press, 1994.
2. F. Bueno, M. García de la Banda, and M. Hermenegildo. Effectiveness of global analysis in strict independence-based automatic program parallelization. In *International Symposium on Logic Programming*, pages 320–336. MIT Press, 1994.
3. Michael Codish and Cohavit Taboch. A semantic basis for termination analysis of logic programs and its realization using symbolic norm constraints. In Michael Hanus, Jan Heering, and Karl Meinke, editors, *Algebraic and Logic Programming, 6th International Joint Conference, ALP '97-HOA '97, Southampton, U.K., Sep 3–5, 1997*, volume 1298 of *LNCS*, pages 31–45. Springer, 1997.
4. D. De Schreye and S. Decorte. Termination of logic programs: The never-ending story. *The Journal of Logic Programming*, 19–20:199–260, 1994.
5. Carl C. Frederiksen. SCT analyzer. At <http://www.diku.dk/~xeno/sct.html>.
6. Carl C. Frederiksen. A simple implementation of the size-change principle. Technical Report D-442, Datalogisk Institut Københavns Universitet, 2001.
7. Neil Jones and Arne Glenstrup. Program generation, termination, and binding-time analysis. In *Principles, Logics, and Implementations of High-Level Programming Languages, PLI 2002 (invited paper)*. Springer, 2002.
8. Chin Soon Lee. Partial evaluation of the euclidean algorithm, revisited. *Higher-Order and Symbolic Computation*, 12(2):203–212, Sep 1999.
9. Chin Soon Lee. *Program Termination Analysis, and Termination of Offline Partial Evaluation*. PhD thesis, UWA, University of Western Australia, Australia, 2001.
10. Chin Soon Lee. Finiteness analysis in polynomial time. In M. Hermenegildo and G. Puebla, editors, *Proceedings of Static Analysis Symposium*, volume 2477 of *LNCS*. Springer, 2002.
11. Chin Soon Lee. Program termination analysis in polynomial time (with proofs). To be available as a technical report at DIKU, Denmark, 2002.
12. Chin Soon Lee, Neil D. Jones, and Amir Ben-Amram. The size-change principle for program termination. In *Proceedings of the ACM Symposium on Principles of Programming Languages, Jan 2001*. ACM, 2001.
13. Naomi Lindenstrauss and Yehohua Sagiv. Automatic termination analysis of logic program (with detailed experimental results). Article available at <http://www.cs.huji.ac.il/~naomil/>, 1997.
14. Lutz Plümer. *Termination Proofs for Logic Programs*, volume 446 of *LNAI*. Springer-Verlag, 1990.
15. Wim Vanhoof and Maurice Bruynooghe. Binding-time annotations without binding-time analysis. In *Logic for Programming, Artificial Intelligence, and Reasoning (LPAR), 8th International Conference, Havana, Cuba, Dec 3-7, 2001*, volume 2250 of *LNCS*, pages 707–722, 2001.

A Program Semantics

Programs are interpreted according to the semantic operator \mathcal{E} of Fig. A.1.

B Results of Experiments

As instantiatedness information is not accounted for, it is *ground termination behaviour* that is analyzed. A positive SCT result for a component does not indicate the termination of procedures in the component. It indicates that no looping occurs among these procedures. The ground termination of particular procedures can be recovered from this information and the call graph. SCT termination deductions are comparable with Termilog and Terminweb's.