

A modal μ perspective on solving parity games in quasi-polynomial time*

Karoliina Lehtinen

Kiel University

Kiel, Germany

kleh@informatik.uni-kiel.de

Abstract

We present a new quasi-polynomial algorithm for solving parity games. It is based on a new bisimulation invariant measure of complexity for parity games, called the *register-index*, which captures the complexity of the priority assignment. For fixed parameter k , the class of games with register-index bounded by k is solvable in polynomial time.

We show that the register-index of parity games of size n is bounded by $O(\log n)$ and derive a quasi-polynomial algorithm. Finally, we give a descriptive complexity account of the quasi-polynomial complexity of parity games: The winning regions of parity games with p priorities and register-index k are described by a modal μ formula of which the complexity, as measured by its alternation depth, depends on k rather than p .

CCS Concepts • **Theory of computation** \rightarrow **Complexity theory and logic**; *Logic and verification*; *Modal and temporal logics*;

Keywords parity games, modal μ calculus, algorithmic complexity, descriptive complexity

ACM Reference Format:

Karoliina Lehtinen. 2018. A modal μ perspective on solving parity games in quasi-polynomial time. In *LICS '18: LICS '18: 33rd Annual ACM/IEEE Symposium on Logic in Computer Science, July 9–12, 2018, Oxford, United Kingdom*. ACM, New York, NY, USA, 10 pages. <https://doi.org/10.1145/3209108.3209115>

1 Introduction

Parity games are two-player games of infinite duration, played on finite graphs labelled with integer priorities. They arise at the intersection of logic, games, and automata theory, with applications in reactive synthesis and verification. In particular, they are the model-checking games for the modal μ calculus, L_μ . Solving parity games – that is, deciding which player has a winning strategy – is the subject of much research, focused both on understanding their complexity and developing practical solvers [25].

Despite extended efforts, the exact complexity of solving parity games remains an open problem. The problem is in $\text{UP} \cap \text{coUP}$ [21] and admits a quasi-polynomial algorithm [9, 12, 16, 23]. While

a general polynomial algorithm is still elusive, many restricted classes of parity games are known to be solvable in polynomial, or even sub-polynomial time. This is the case for classes of games with a logarithmic number of priorities [9] as well as for classes of games where the structure of the underlying graph is restricted: parity games of bounded Kelly-width [19], DAG-width [3] and entanglement [4] are all solvable in polynomial time, while those with bounded tree-width and clique-width are even in LogCFL [15], the class of languages logspace reducible to context-free languages.

These structural measures on the graph of a parity game are orthogonal to the size of the priority assignment. However, the complexity of a parity game is not solely due to its structure nor to the size of its priority assignment, but to how both interact. For example, the cycles in a graph may be intertwined, causing high entanglement, yet unless the priorities are equally entangled within these cycles, this measure is not an accurate reflection of algorithmic complexity. The key insight of this paper is to capture the complexity of the priority assignment with a new measure which is logarithmically bounded in the size of the graph. This yields a new quasi-polynomial algorithm and a descriptive complexity account of the quasi-polynomial solvability of parity games.

A. A new measure of complexity

Our first contribution is to parametrise the complexity of the priority assignment of a parity game. We introduce the *register-index* as a new, L_μ describable, bisimulation invariant measure of complexity for parity games, which captures, roughly speaking, how many priorities the winner of a parity game needs to keep in memory in order to produce a witness of their victory. Like entanglement, register-index is defined in game-theoretic terms, using a parametrised k -register game that is played on a parity game arena.

We then define a k -parameterised polynomial-time algorithm which, for fixed k , solves parity games of register-index up to k . This adds parity games of bounded register-index to the classes of parity games known to be solvable in polynomial time. Notably, even the class of games of register-index 1 contains non-trivial parity games of arbitrarily high entanglement, tree-width, Rabin-index, Kelly-width, DAG-width and arbitrarily many priorities, as well as the known parity games exhibiting worst-case performance for recursive [2, 14], strategy improvement [13], and one of the existing quasi-polynomial algorithms [12].

We then show that the register-index of a parity-game of size n is $O(\log n)$. This yields a novel proof that parity games are solvable in quasi-polynomial time, based on reducing solving a parity game of size n with p priorities to solving one of size $O(p^{\log n} n \log n)$ with $O(\log n)$ priorities.

B. An L_μ account of register-index

*Work supported by BMBF project no. 01IS160253 "ARAMiS II"

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

LICS '18, July 9–12, 2018, Oxford, United Kingdom

© 2018 Copyright held by the owner/author(s). Publication rights licensed to Association for Computing Machinery.

ACM ISBN 978-1-4503-5583-4/18/07...\$15.00

<https://doi.org/10.1145/3209108.3209115>

Although the runtime of this algorithm does not improve over existing quasi-polynomial algorithms, it provides a novel, L_μ perspective on solving parity games in quasi-polynomial time. In the second half of the paper, we link the quasi-polynomial solvability of parity games to their descriptive complexity. It is well known that the winning regions of parity games with maximal priority p are described by a L_μ formula of which the alternation depth depends only on p . We show that the alternation depth of the L_μ formula describing the winning regions of parity games of register-index up to k and of maximal priority p depends only on k , rather than p . Its size is $O(kp^k)$. Since register-index is logarithmic in the size of the game, a L_μ formula of alternation depth k and size $O(kp^k)$ describes the winning regions of parity games of size up to 2^k with up to p priorities.

Finally, we show that register-index, when extended to L_μ formulas, is a decidable upper bound on the semantic complexity of formulas, and demonstrate that it is sensitive to non-trivial syntactic inefficiencies.

C. Overview

In brief, we provide new insights into the complexity of parity games and link their quasi-polynomial solvability to their descriptive complexity. Its main contributions are:

- Introducing a new, bisimulation invariant measure of complexity for parity games (Section 3);
- Showing that for fixed k solving parity games of register-index up to k is in P, thus adding parity games of bounded register-index to the classes of parity games known to be solvable in polynomial time (Theorem 3.6); Classes of games of different register-index are studied in Section 4;
- Bounding the register-index of parity games of size n with $O(\log n)$ (Theorem 4.7), resulting in a new quasi-polynomial parity game algorithm (Section 5);
- Relating register-index to complexity in L_μ (Section 6).

2 Background

2.1 Parity games

Definition 2.1 (Parity games). A parity game is an infinite-duration two-player zero-sum path-forming game, played between Even and her opponent Odd on a finite game graph $G = (V, V_e, V_o, v_i, E, \Omega)$ called the *arena*. The vertices V of the arena are partitioned into those belonging to Even, V_e , and those belonging to Odd, V_o . The *priority assignment* $\Omega : V \rightarrow I$ maps every vertex to a *priority* in I , a finite prefix of the non-negative integers. Starting at an initial state $v_i \in V$, a *play* proceeds with the owner of the current state v choosing a state v' among its successors in the directed edge relation $E \subseteq V \times V$. Thus the players collaboratively form a play, consisting of a potentially infinite path along the edges of the game graph. If the play is finite, the owner of its last position loses; Otherwise Even wins if the highest priority visited infinitely often is even, else Odd wins.

A (positional) *strategy* σ for a player $P \in \{\text{Even}, \text{Odd}\}$ maps every position v belonging to P in a parity game to one of its successors $\sigma(v)$. A play $\pi = v_0 v_1 \dots$ is said to agree with σ when for all i , if v_i belongs to P , then v_{i+1} is $\sigma(v_i)$. A strategy σ for player P is said to be winning for P if all plays starting at v_i that agree with σ are winning for P .

Theorem 2.2. [11] *In all parity games, one of the players has a positional winning strategy.*

It will often be convenient to assign priorities to edges instead of vertices, with $\Omega : E \rightarrow I$. A parity game with edge priorities can be converted into one with node priorities by introducing intermediate, priority-carrying nodes onto all edges and assigning a low priority to other vertices.

We denote the set $\{1, 2, \dots, n\}$ with $[n]$.

2.2 Entanglement

Register-index is inspired by and related to entanglement, which was introduced by Berwanger and Grädel [4] as a measure of how intertwined the cycles of a graph are.

Definition 2.3 (Entanglement). Entanglement is defined with a *detective game* in which a team of k detectives tries to catch a thief. The thief picks his initial position, while the detectives start outside of the graph. At each move, the detectives can either stay as they are or one of them can move to the current position of the thief. The thief then chooses a successor state not occupied by a detective. If such a state does not exist, the detectives win. If the thief can keep playing infinitely, he wins. Then, the entanglement of a graph is the minimal k for which k detectives have a winning strategy. For example, n -cliques have entanglement $n - 1$ and acyclic graphs have entanglement 0.

A positional strategy σ for player P on a parity game G induces a subgraph G_σ in which positions belonging to player P only have one successor. The entanglement of a parity game is the minimum entanglement of a subgame G_σ induced by a positional winning strategy σ .

Theorem 2.4. [4] *Parity games of bounded entanglement are solvable in polynomial time.*

The proof of this theorem uses a k -super-detective game, solvable in polynomial time for fixed k . For parity games on graphs of entanglement k , Even wins the parity game if and only if she wins the k -super-detective game. The k -register game, introduced next, can be seen as a bisimulation-invariant, infinite duration variation of this game¹.

Although not bisimulation invariant, entanglement is related to L_μ : the entanglement of a finite graph corresponds to the number of fixpoint variables needed to describe it up to bisimulation with a L_μ formula – a fact used to show the strictness of the L_μ variable hierarchy [6]. The winning regions of parity games of bounded entanglement are LFP-describable [10]. The relationship between entanglement and other structural measures is discussed in Berwanger et al. 2012 [5], while directed and undirected graphs of entanglement two are studied in Grädel et al. 2009 [17] and Belkhir and Santocanale 2007 [1] respectively.

3 The register-index of parity games

Like entanglement, the register-index of parity games is defined using a parametrised game: the k -register game. This game is played on a parity game arena, on which in addition to winning the parity game, player Even has produce a witness of her victory using a fixed amount of memory. This section introduces the k -register

¹Very roughly, instead of having to catch the thief just once, the super-detective needs to catch the thief infinitely often after a high even priority and only finitely often after a higher odd priority.

game and shows that it is an under-approximation of a parity game: if for any k Even wins the k -register game on a parity game G , she also wins in the parity game on G ; for every parity game G in which Even wins, there is k such that Even wins in the k -register game on G . We call the minimal such k the register-index of G . Solving a parity game G of bounded register-index then reduces to solving a k -register game on G which, for fixed k , can be done in polynomial time.

3.1 Register games

The k -register game consists of a normal parity game, augmented with k registers. Each register records the highest priority that has occurred in the parity game since it was last reset. The registers are ranked according to how long it has been since their last reset, with a newly reset register having rank 1. Initially, all registers contain 0 and the ranking is arbitrary. One player, let's say Even, is given control of the registers. At each turn, Even can choose to reset a register of any rank r . If the register contains the priority p , this produces output $2r$ if p is even and $2r + 1$ otherwise. As long as Even resets registers infinitely often, this produces an infinite sequence in $[2k + 1]^\omega$. To win, Even has to either win finitely in the underlying parity game, or produce an infinite sequence of outputs that satisfies her parity condition: the maximal priority output infinitely often must be even.

We now define formally the k -register game on a parity game G with priority domain I . Since the winning condition of the register game is a parity condition, we present the k -register game on a parity game arena G as a parity game on an arena $\mathcal{R}_e^k(G)$, of which the positions are positions of G paired with vectors in I^k that represent the contents of the registers. An additional binary variable t indicates whether the next move is a potential register reset ($t = 0$), or a move in the underlying parity game ($t = 1$).

Definition 3.1 (k -Register game). Let G be a parity game $(V^G, V_e^G, V_o^G, v_i^G, E^G, \Omega^G)$ and let I be the co-domain of $\Omega^G : V^G \rightarrow I$. For a fixed non-zero parameter $k \in \mathbb{N}$, the arena of the k -register game $\mathcal{R}_e^k(G)$ on G in which Even controls the registers, consists of $\mathcal{R}_e^k(G) = (V, V_e, V_o, v_i, E, \Omega)$ as follows. While G carries its priorities on its vertices, for the sake of clarity, $\mathcal{R}_e^k(G)$ carries them on its edges: $\Omega : V \rightarrow [2k + 1]$.

- V is a set of positions $(p, \bar{x}, t) \in V^G \times I^k \times \{0, 1\}$,
- V_o consists of $(p, \bar{x}, 1)$ such that $p \in V_o^G$,
- V_e consists of $V \setminus V_o$,
- v_i is $(v_i^G, \bar{0}, 0)$,
- E is the disjoint union of sets of edges E_{move} , E_{skip} and E_r for all $r \in [k]$ where:
 - E_{move} consists of edges $((p, \bar{x}, 1), (p', \bar{x}', 0))$ such that
 - $(p, p') \in E^G$ and
 - $x'_i = \max(\Omega^G(p'), x_i)$.
 - E_{skip} consists of edges $((p, \bar{x}, 0), (p, \bar{x}, 1))$.
 - For each $r \in [k]$, E_r consists of edges $((p, \bar{x}, 0), (p, \bar{x}', 1))$ such that:
 - $x'_i = x_i$ for $i > r$,
 - $x'_i = x_{i-1}$ for $1 < i \leq r$,
 - $x'_1 = 0$.
- Ω assigns priorities from $[2k + 1]$ to edges as follows:
 - Edges of E_{move} and E_{skip} have priority 1;
 - An edge $((p, \bar{x}, 0), (p, \bar{x}', 1)) \in E_r$ has priority $2r$ if x_r is even, and priority $2r + 1$ otherwise.




Figure 1. Parity game of register-index 2

Priority	1	–	2	1	–	0	2	–	1	–	2	–	1
Reset	–	2	–	–	1	–	–	2	–	1	–	2	–
Reg. 1	1	0	2	2	0	0	2	0	1	0	2	0	1
Reg. 2	1	1	2	2	2	2	2	2	2	2	2	2	2
Output	–	5	–	–	2	–	–	4	–	3	–	4	–

Figure 2. Fragment of a play in a 2-register game.

Terminology: Given a play in $\mathcal{R}_e^k(G)$, we call the *underlying play* its projection onto the first element of each visited position. At a position (p, \bar{x}, t) , we write that:

- a priority $q \in I$ occurs if $\Omega^G(p) = q$;
- a register of rank $r \in [k]$ holds priority $q \in I$ if $x_r = q$;
- Even resets the register of rank r and outputs $j \in [2k + 1]$ if the play follows an edge in E_r of priority j ;
- q is recorded at rank r if Even resets the register of rank r at a position $(p, \bar{x}, 0)$ where $x_r = q$;
- the interval of a position p at which Even resets a register a is the fragment of the play since the previous reset of a (not necessarily at the same rank).

A strategy for Odd in G induces a strategy for Odd in $\mathcal{R}_e^k(G)$. A strategy for Even in G , paired with a resetting strategy in $\mathcal{R}_e^k(G)$ from positions $(p, \bar{x}, 0)$ induces a strategy for Even in $\mathcal{R}_e^k(G)$.

The k -register game arena $\mathcal{R}_o^k(G)$ where Odd controls the registers is similar except that positions $(p, \bar{x}, 0)$ are in V_o , edges in $E_{\text{move}} \cup E_{\text{skip}}$ have priority 0, and edges $((p, \bar{x}, 0), (p, \bar{x}', 1))$ of E_r have priority $2r$ if x_r is even, and $2r - 1$ otherwise. The k -register game with Even (resp., Odd) in control of registers on a parity game arena G is the parity game on the arena $\mathcal{R}_e^k(G)$ (resp., $\mathcal{R}_o^k(G)$). Unless stated otherwise, the k -register game on G refers to $\mathcal{R}_e^k(G)$.

Example 3.2. Figure 1 shows an edge-labelled arena in which Even wins the parity game but loses the 1-register game. In the 1-register game, Odd's strategy is to stay at a position until Even resets the unique register and then repeat this in the other position. The unique register will then alternatively hold 1 and 2 (and briefly 0 between moves). If Even resets registers infinitely often, this strategy produces outputs 2 and 3 infinitely often, causing Even to lose.

The register contents during a play in the 2-register game is shown in Figure 2. Note how a reset of rank 2 causes the contents of the register of rank 1 to move into the register of rank 2. Even can win by resetting the register of rank 1 after seeing 1 or 0, and the one of rank 2 after seeing 2.

Dedicated readers may find it helpful to work out the register-index of the 4-node parity game in Figure 5.

For fixed k , the k -register game on G is a parity game on an arena of polynomial size and of priority domain $[2k + 1]$: it can be solved in polynomial time in the size of G .

3.2 Register-index

We first prove two fundamental properties of k -register games.

Lemma 3.3. *If Odd has a winning strategy in G , he also has a winning strategy in $\mathcal{R}_e^k(G)$, for all k . Dually, if Even has a winning strategy in G , she also has a winning strategy in $\mathcal{R}_o^k(G)$, for all k .*

Proof. We show that a winning strategy τ for Odd in G is also winning in $\mathcal{R}_e^k(G)$, for any k .

Odd obviously wins finite plays that agree with τ . Odd also wins if a play only contains a finite number of resets. For an infinite play with infinitely many resets that agrees with τ , consider the highest rank r at which registers are reset infinitely often. Observe that for a newly reset register to reach rank r , a register of rank r or higher has to be reset at least once. Therefore if a reset of rank r occurs at a position p , then the interval of p begins no later than at the previous position at which a reset of rank r or higher occurred. As a result, as long as there are infinitely many resets, the dominant odd priority of the play is recorded infinitely often at rank r . The maximal priority output infinitely often is therefore odd. The proof of the dual statement is similar. \square

The converse does not hold: Even can have a winning strategy in a parity game G but still lose the k -register game on G (see Example 3.2). However, for every parity game G , there is a k , bounded by the number of even priorities in G (and, as we will see later, by $1 + \log n$), such that Even has a winning strategy in G if and only if she has a winning strategy in $\mathcal{R}_e^k(G)$.

Lemma 3.4. *Let G be a parity game with k even priorities, in which Even has a winning strategy. Even has a winning strategy in $\mathcal{R}_e^k(G)$.*

Proof. Even's strategy in $\mathcal{R}_e^k(G)$ is to play a positional winning strategy in the underlying parity game G and reset registers as follows: when an even priority p occurs in the underlying game for the first time, she resets the register of highest rank, and whenever p occurs again, she resets that same register, regardless of its current rank.

A reset then records the highest priority seen between two occurrences of an even priority (excluding the first occurrence, but including the second). The rank of a register a reset at the occurrence of an even priority p is the number of distinct even priorities that have occurred since the last occurrence of p , plus one: every occurrence of a new even priority triggers the reset of a register that was last reset earlier than a , thus increasing the rank of a ; repeated occurrences of even priorities between consecutive occurrences of p trigger the reset of registers of smaller rank than a . Therefore, if a register is reset at v and another one at v' , and the interval of v is contained within the interval of v' , then the reset at v' is of strictly higher rank than the reset at v .

Let q be the highest priority that occurs in a play π infinitely often, q' the highest odd priority that is recorded infinitely often, and r the highest rank at which q' is recorded infinitely often. Now consider a suffix of π in which q is the first occurring priority and in which no odd priority larger than q' occurs. In this suffix, whenever at a position v a register is reset and q' is recorded, the interval at v is nested within two consecutive occurrences of q , as else q would

be recorded instead of q' . Therefore, if q' is recorded at rank r , the next reset at a position where q occurs records q at a higher rank than r . As a result, q is recorded infinitely often at a rank higher than r , so the maximal output produced infinitely often is even.

Similarly, if Odd wins a parity game G with k odd priorities, he has a winning strategy in $\mathcal{R}_o^k(G)$. \square

Definition 3.5 (Register-index). The *register-index* of a parity game G is the least integer k such that Even wins both G and $\mathcal{R}_e^k(G)$ or Odd wins both G and $\mathcal{R}_o^k(G)$. This amounts to the least k such that $\mathcal{R}_e^k(G)$, $\mathcal{R}_o^k(G)$ and G have the same winner.

Theorem 3.6. *Parity games of bounded register-index are solvable in polynomial time.*

Proof. For a fixed parameter k , let $G = (V, V_e, V_o, v_i, E, \Omega)$ be a parity game of register-index at most k with priority co-domain I : Even has a winning strategy in G if and only if she has a winning strategy in $\mathcal{R}_e^k(G)$. The game $\mathcal{R}_e^k(G)$ is itself a parity game, with state space $V \times I^k \times \{0, 1\}$ and priorities $[2k + 1]$. As k is fixed, this is PTIME-solvable. \square

4 Some register-index properties

This section studies register-index in more details. First, we compare register-index to other measures, then we prove one of the core theorems of this paper: register-index is logarithmic in the size of the game. In brief:

- The number of priorities in a parity game is an upper bound on its register-index (Lemma 3.4);
- Parity games in which a k -set of vertices intersects all cycles have register-index at most k (Lemma 4.1).
- Games of register-index 1 can have arbitrarily high entanglement and number of priorities;
- For all m there are parity games of register-index m with entanglement 2.
- Rabin index and register-index are orthogonal.
- A parity game of size n has register-index at most $1 + \log n$ (Theorem 4.7).

As with entanglement, the register-index depends only on the subgraphs induced by winning strategies. Therefore all examples are one-player games where Odd controls all positions but Even wins.

4.1 Register-index and entanglement

In some cases, entanglement and register-index coincide:

Lemma 4.1. *If there is a set of positions S of size k that intersects all cycles, then the parity game has register-index (and entanglement) at most k .*

Proof. The winning player can use a positional winning strategy in the underlying game and reset registers as follows: at the first occurrence of $v \in S$ she resets the register of maximal rank, and at each subsequent occurrence of v , she resets the same register, regardless of its rank. Each register that is reset at least twice records the highest priority that occurs in a cycle from a position in S to itself. As S intersects all cycles, registers get reset infinitely often. The underlying strategy forces all cycles to be dominated by a priority of the winning parity; all outputs are of this parity. \square




Figure 3. A parity game of register-index 1




Figure 4. A parity game of register-index 2

In general entanglement and register-index are orthogonal:

Lemma 4.2. *There are parity games of register-index 1 with both arbitrarily high entanglement and number of priorities.*

Proof. Consider the n clique C_n (Figure 3) where all nodes $v_i, i \in [n]$ belong to Odd, with the following edge-priorities: the edge from $v_i, i \in [n]$ to $v_j, j \in [n]$ if $i > j$ has priority $2i$; the edge from v_j to v_i has priority $2i - 1$. This parity game has entanglement n and $2n$ priorities. Even wins $\mathcal{R}_e^1(C_n)$: she resets the register whenever an even priority is seen. Any path with only odd priorities, followed by an even priority p , sees no priorities higher than p . Therefore, every recorded priority – and hence every output – is even. \square

Example 4.3. To see how changes in the order of priorities affect the register-index, reverse the edges of C_n (Figure 4) to obtain a family of parity games of register-index 2 and of size, entanglement and number of priorities $\Omega(n)$. Even's winning strategy in the 2-register game consists of resetting at each turn the register of rank 1 if it contains an odd priority, and resetting the register of rank 2 otherwise.

Lemma 4.4. *For all m , there exists a parity game of register-index at least m and entanglement 2.*

Proof. Let H_0 be the game arena consisting of a single node, belonging to Odd, with a self-loop of priority 0. This unique node is also the initial node of H_0 .

Then, for all $n > 0$, the arena H_n consists of two distinct copies of H_{n-1} with initial positions v_0 and v_1 respectively, with an edge (v_0, v_1) of priority $2n - 1$ and an edge (v_1, v_0) of priority $2n$. The position v_0 is also the initial position of H_n . See Figure 5.

Even wins these parity games H_n because all cycles are dominated by an even priority. H_0 is of entanglement 1, while H_i for $i > 0$ is of entanglement 2. We will show that for $n > 0$, Odd has a winning strategy in the n -register game $\mathcal{R}_e^n(H_n)$.

We reason inductively, and show that for each H_n , i) from the initial position in $\mathcal{R}_e^m(H_n)$ for $m \geq n$, with a register configuration in which register contents are bounded by $2n - 1$, Odd can force the game to output $2n + 1$ or a higher odd priority, before returning to the initial position, and ii) Even loses in $\mathcal{R}_e^n(H_n)$.

Base case. We first show that Odd has a winning strategy in $\mathcal{R}_e^1(H_1)$. His strategy is to always loop in the current position until Even resets the unique register, then move to the other position. This causes both 1 and 2 to be recorded infinitely often at rank 1, producing 2 and 3 infinitely often. Note that if Odd uses this same strategy in $\mathcal{R}_e^m(H_1)$ for $m \geq 1$, starting from a register configuration in which register contents are bounded by 1, although he can't necessarily win, he can force the game to either stay away from the initial position, or output $2n + 1$ or a higher odd priority before returning to the initial position.

Inductive step. Assume i) and ii) for H_n .

i) Consider the following strategy for Odd in $\mathcal{R}_e^m(H_{n+1})$ for $m \geq n + 1$. He first moves from v , the initial position of H_{n+1} onto the initial position v' of the second component of H_{n+1} , via the odd priority $2n + 1$. Then, he plays in H_n , which only contains priorities smaller than $2n + 1$, with a strategy that is winning in $\mathcal{R}_e^n(H_n)$. To counter this strategy, Even has eventually to reset a register of rank $n + 1$ or higher, after which Odd returns to the initial position. This is his strategy τ_{n+1} . Observe that if at the initial position all register contents are bounded by $2n + 1$, then Even will either lose in the second H_n component, or reset a register of rank $n + 1$ or higher when it contains the odd priority $2n + 1$, outputting $2n + 3$ or a higher odd priority.

ii) We show that he also has a winning strategy in $\mathcal{R}_e^{n+1}(H_{n+1})$. He begins by playing τ_{n+1} until $2n + 3$ is output and the play is back at the initial position. Note that the registers now all contain $2n + 2$, so he can not yet repeat τ_{n+1} . Instead, he plays a strategy that is winning in $\mathcal{R}_e^n(H_n)$ in the first H_n component of H_{n+1} . Again, Even will lose unless she resets register $n + 1$ infinitely often. After she has reset the register at rank $n + 1$ at least $n + 1$ times, all the registers hold values smaller than $2n + 1$. Odd can then return to the initial position, and again use τ_{n+1} to force output $2n + 3$. Thus by alternating between using τ_{n+1} to force the maximal odd output, and clearing the registers of higher priorities in order to be able to use τ_{n+1} again, Odd forces $2n + 3$ to be output infinitely often.

Hence H_n has register-index at least $n + 1$. \square




Figure 5. H_0, H_1, H_2, H_3 : A family of parity games of high register-index and entanglement at most 2

4.2 Rabin-index and register-index

Another related measure is the Rabin index [20] which, like register-index, considers both the structure of cycles in a parity game and its parity assignment. The idea of the Rabin index is to capture the minimal priority assignment that preserves winning regions and strategies regardless of the ownership partition of the game graph.

Definition 4.5 (Rabin index). Given a directed graph $G = (V, E)$, two priority assignments Ω, Ω' over V are equivalent if for all partitions of V into V_e and V_o , and all $v \in V$, the parity games (V, V_o, V_e, v, Ω) and $(V, V_o, V_e, v, \Omega')$ have the same winner and same set of winning strategies. The Rabin index of a parity game G with priority assignment Ω is then the size of the smallest co-domain of any equivalent priority assignment over the graph of G .

Lemma 4.6. For all k there are parity games of register-index at least k and Rabin index 1 and vice versa.

Proof. Since the Rabin index is agnostic to the ownership of nodes and register-index depends on the subgraphs induced by winning strategies, we can build games of low register-index and high Rabin index as follows: a game consisting of a starting position belonging to Even with one successor in a winning region of register-index 1, and another in a parity game of Rabin index k has register-index 1 and Rabin index at least k . Conversely, the graphs H_n of register-index at least $n + 1$ are all of Rabin index 1, as witnessed by the trivial colouring which assigns 0 to all positions. Rabin index and entanglement are therefore orthogonal. \square

4.3 A logarithmic bound on register-index

In this section we prove one of the main results of this paper:

Theorem 4.7. The register-index of parity games of size n is at most $1 + \log n$.

Before diving into the main proof, we define the following:

Definition 4.8. A strategy σ for Even in a k -register-game of maximal priority q is *defensive* if, from a configuration where the highest ranking register contains q or a higher even priority, plays which agree with σ never output $2k + 1$.

Definition 4.9. The *defensive* register-index of G where Even wins the parity game is the least integer k such that Even wins $\mathcal{R}_e^k(G)$ with a defensive strategy. The register-index of G is bounded by its defensive register-index.

Proof. Let $G = (V, V_e, V_o, v_i, E, \Omega)$ be the parity game induced by a positional winning strategy σ for Even in some parity game: nodes $v \in V_e$ have a unique successor, namely $\sigma(v)$. Let q be the least even priority larger or equal to all priorities appearing in G . We start by proving that if G has defensive register-index $r > 1$, it has two distinct subgames of defensive register-index $r - 1$.

If $r = 2$, this is trivial: all singleton games have register-index 1. We now consider the case of $r > 2$.

Let G_1, \dots, G_j be the maximal strongly connected subgames of G induced by vertices of priority up to $q - 2$, and let their defensive register-indices be k_1, \dots, k_j respectively: k_i is the number of registers Even needs in order to win with a defensive strategy in the k_i -register game from all positions of G_i . Let k_m be a maximum among these, or 1 if G has no strongly connected subgames of priority up to $q - 2$.

We show that $k_m \geq r - 1$. For contradiction, assume $k_m < r - 1$. Even's strategy in the $r - 1$ -register game on G is as follows. She resets the register of rank $r - 1$ only when q occurs. Then, in each subgame G_i , she simulates her winning strategy in the k_i -register game using only the k_i lowest ranking registers, leaving the register of rank $r - 1$ untouched. This forces Odd to eventually leave every subgame G_i , or lose. Since every cycle is dominated by an even priority, a play that eventually leaves every subgame G_i , or which eventually doesn't enter subgames any more, is either finite and winning for Even, or sees q infinitely often. This strategy is winning and defensive, since it resets the highest ranking register infinitely often when it contains q , and never when it contains an odd priority. G therefore has defensive register-index at most $r - 1$, rather than r , a contradiction.

Therefore $k_m \geq r - 1$. If $k_m = r$, we consider the strictly smaller subgame G_m instead of G and show that it too has two subgames of defensive register-index $r - 1$.

If only one subgame G_i has defensive register-index $k_m = k_i = r - 1$, we show that Even has a defensive winning strategy in the $r - 1$ register game on G , a contradiction. Her strategy begins by first resetting all registers that initially contain odd priorities (this might take several turns), regardless of the progress in the parity game. Note that if the register of rank $r - 1$ initially contains q or a higher even priority, this register does not get reset. After this initial series of resets, her strategy is as follows. She resets the register $r - 1$ whenever she sees q . Note that since q is the maximal priority and $r - 1 > 1$, immediately after this reset, the register of maximal rank

still contains q or a higher even priority. In subgames G_j where $j \neq i$, she simulates a winning strategy in $\mathcal{R}_e^{k_j}(G_j)$ without using the register of rank $r - 1$. When she enters G_i she plays a defensive winning strategy, this time using all $r - 1$ registers.

This strategy is winning: finite plays are trivially winning for Even; a play that eventually stays in a subgame agrees with a winning strategy within that subgame; a play that either leaves subgames infinitely often or eventually does not enter a subgame must see q infinitely often, and always between two visits to G_i , and therefore output $2(r - 1)$ but not $2(r - 1) + 1$ infinitely often. If the initial register-configuration contains q or a higher even priority in the register of rank $r - 1$, then this strategy never enters G_i without q or a higher even priority in the register of rank $r - 1$, and therefore never outputs $2(r - 1) + 1$. Hence it is defensive. The defensive register-index of G is therefore $r - 1$, a contradiction. We conclude that a game of defensive register-index $r > 1$ has at least two distinct subgames of defensive register-index $r - 1$.

Then, as the same argument applies to all subgames of G , a game of defensive register-index r has 2^{r-1} distinct subgames of defensive register-index 1. This bounds the defensive register-index of a parity game with $1 + \log n$ where n is the size of the parity game. The same bound applies to the register-index, which must be lower than the defensive register-index. \square

5 Algorithmic complexity

The algorithmic consequence of Theorem 4.7 is that one can always solve $\mathcal{R}_e^{1+\log |G|}(G)$ instead of G , which gives an alternative quasi-polynomial time algorithm for solving parity games:

Corollary 5.1 (Also from [9, 23]). *Parity games are solvable in quasi-polynomial time.*

Proof. Let n be the number of vertices in a parity game G with register-index k , with p distinct priorities, and m edges. The parity game $\mathcal{R}_e^k(G)$ has $O(np^k)$ positions, $2k + 1$ priorities and $O((m + kn)p^k)$ edges. Since $\mathcal{R}_e^k(G)$ has its priorities on its edges, for the complexity analysis we take the size of $\mathcal{R}_e^k(G)$ to be $O(knd^k)$, to account for an additional vertex for each of the $O(knd^k)$ edges of significant priority – that is, those in E_r for some $r \in [k]$.

The register-index of a parity game of size n is at most $1 + \log n$. Therefore solving parity games G reduces to solving $\mathcal{R}_e^k(G)$ where $k = 1 + \log n$. The $\mathcal{R}_e^k(G)$ game can then be solved with an algorithm exponential in the number of priorities of one's choice, say the small progress measure algorithm [22], to obtain a quasi-polynomial algorithm. \square

For a tighter complexity analysis, one can use Jurdiński and Lazic's quasi-polynomial succinct progress measure algorithm [23] which runs in $O(m\eta^{2.38})$ if the number of priorities is less than $\log \eta$, where η is the number of positions of odd priority. This is the case for $\mathcal{R}_e^k(G)$, which has over np^k positions of odd priority and $2k + 1$ priorities (as long as $p \geq 4$). With size $O(knp^k)$ where $k = 1 + \log n$, this yields a $O(n^{4.48+3.38 \log n} (\log n)^{2.38})$ algorithm. The space complexity remains quasi-polynomial.

Although this algorithm does not improve over the runtime of existing quasi-polynomial time algorithms, its parameterised versions, consisting of solving $\mathcal{R}_e^k(G)$ instead of G for a parameter k , provide polynomial algorithms for the rich classes of games of bounded register-index. Even the class of games of register-index 1 contains games which are otherwise of arbitrary complexity. In particular, the families of parity games which exhibit worst-case behaviour for strategy improvement, divide-and-conquer, and one of the existing quasi-polynomial algorithms [2, 12–14] all have constant register-index 1. In the former two cases, this is simply because although the games have complex features, the winner has a simple winning strategy, according to which all plays eventually only see one priority. The latter case is not as trivial, but every odd priority is still immediately followed by a larger even priority, so Even still has an easy 1-register game strategy consisting of resetting the unique register whenever an even priority occurs. However, the class of parity games of register-index 1 is not restricted to such simple games, as illustrated by Figure 3. Larger register-indices only occur if the winning strategy induces a connected region with the particular structure described in Lemma 3.2 and Theorem 4.7. Solving k -register games is therefore a complementary approach to existing parity game solving strategies. It seems plausible that solving $\mathcal{R}_e^k(G)$ and $\mathcal{R}_o^k(G)$ for a k as low as 2, or 3, and falling back onto a different algorithm if these don't have the same winner could be an effective approach to solving parity games in practice.

Beyond its potential for practical solving, this algorithm gives a descriptive complexity account of the quasi-polynomial solvability of parity games, as we will see in the next section.

6 Register-index and L_μ

The complexity of solving parity games is closely related to complexity in L_μ : The priorities in the model-checking parity games of an L_μ -formula reflect the complexity of the formula, as measured by the complexity of its fixpoint structure, known as its alternation depth [26]. Conversely, solving parity games is equivalent to model-checking a particular L_μ formula [18, Section 3.3.6], of which the alternation depth depends on the priorities in the parity game. In this section we argue that not only can we model-check this formula in quasi-polynomial time, but the formula itself can be chosen to be of logarithmic alternation depth and quasi-polynomial size.

We state the polynomial solvability of parity games of bounded register-index in terms of descriptive complexity: the winning regions of the k -register games on parity game arenas with a bounded number p of priorities can be described by a L_μ -formula of size $O(kp^k)$, and alternation-depth k .

After establishing the L_μ expressivity of register-games in Section 6.2 we consider some of its consequences:

- Register-index is bisimulation invariant (Corollary 6.6).
- The winning regions of parity games of bounded size are trivially described by a L_μ formula without fixpoints, however, this formula is exponential in the size bound. The logarithmic bound on the register-index with respect to the size of games means that the winning regions of parity games with p priorities of size up to 2^{k-1} are described by a L_μ formula of size $O(kp^k)$ with alternation depth k (Corollary 6.7)
- Finally, we extend the definition of register-index to L_μ -formulas and show that the register-index of a L_μ formula is a computable upper bound on its semantic alternation depth:

formulas of register-index k are equivalent to formulas of alternation depth k .

6.1 L_μ

For ease of representation we will use L_μ equational systems. Fix countably infinite sets $Prop = \{P, Q, \dots\}$ of propositional variables, and $Var = \{X, Y, \dots\}$ of fixpoint variables.

Definition 6.1. (L_μ) The syntax of unimodal *basic formulas* is given by:

$$\phi := P \mid X \mid \neg P \mid \phi \wedge \phi \mid \phi \vee \phi \mid \Diamond \phi \mid \Box \phi \mid \perp \mid \top$$

Conjunctions take precedence over disjunctions. The scope of modalities extends as little as possible to the right.

A L_μ sentence (S, X_i, Ω) is a finite set of S equations of the form $X_i = \psi_i$ where the left-hand sides are distinct fixpoint variables, of which one is the designated entry point X_i , coupled with a priority assignment $\Omega : Var \rightarrow I$ where I is a finite prefix of the non-negative integers.

The semantics of L_μ and its equational notation are standard – see for example [8]. In the context of this paper, the semantics are best described in terms of the model-checking parity games.

Definition 6.2. Given $S = (S, X_i, \Omega)$ a L_μ sentence and $\mathcal{M} = (V^M, E^M \subseteq V \times V, s_i \in V^M, Q^M : Prop \rightarrow \mathcal{P}(V^M))$ a structure we define the model-checking parity game $G(\mathcal{M}, S)$ as the structure $(V, E, V_e \subseteq V, P_0, \dots, P_q)$ where q is the maximal priority in the co-domain I of Ω , and:

- Positions in V consist of pairs (s, ϕ) where $s \in V^M$ and ϕ is a subformula of a basic formula in S ;
- There is an edge from $(s, \phi \wedge \phi')$ to (s, ϕ) and (s, ϕ') ; from $(s, \phi \vee \phi')$ to (s, ϕ) and (s, ϕ') ; from $(s, \Box \phi)$ to (s', ϕ) for each successor s' of s in E^M ; from $(s, \Diamond \phi)$ to (s', ϕ) for each successor s' of s in E^M ; from (s, X) to (s, ϕ_X) where $(X = \phi_X) \in S$;
- Positions $(s, \phi \vee \phi')$ and $(s, \Diamond \phi)$ satisfy V_e ;
- Positions (s, P) satisfy V_e if $s \notin Q^M(P)$;
- Positions $(s, \neg P)$ satisfy V_e if $s \in Q^M(P)$;
- Position (s, X) satisfies $P_{\Omega(X)}$ while other positions satisfy P_0 .
- (s_i, X_i) is the initial position.

Definition 6.3. A structure \mathcal{M} is said to satisfy a sentence S of L_μ , written $\mathcal{M} \models S$ if Even has a winning strategy in the parity game represented by $G(\mathcal{M}, S)$ when V_e represents positions belong to Even, $V \setminus V_e$ represents positions belonging to Odd and P_i represents positions of priority i .

Definition 6.4. The alternation depth of an L_μ system of equations consists of the number of even priorities in the co-domain I of Ω .

This presentation and the choice to only count even priorities is meant to emphasize the correspondence between alternation-depth, priorities in parity games and their register-index. A thorough discussion on how to define alternation depth, and comparison of definitions used in the literature can be found in Bradfield and Stirling 2007 [8].

6.2 Descriptive complexity

Theorem 6.5. There is a formula Win_I^k of L_μ with alternation-depth k that describes the class of parity games G of priority domain I for which Even has a winning strategy in $\mathcal{R}_e^k(G)$.

Proof. Let E_i stand for $V_e \wedge P_i$ and O_i for $\neg V_e \wedge P_i$; let $M_i(X) = (E_i \wedge \Diamond X) \vee (O_i \wedge \Box X)$. This subformula describes the owner of the current position (of priority i) taking a step in the parity game. Then, the formula for $Y_{\bar{x}, o}$ describes the existence of a winning strategy from a register configuration \bar{x} , while o indicates the last output value:

$$Y_{\bar{x}, o} = \bigvee_{j \in I} M_j(Y_{\bar{x}', 1}) \vee \bigvee_{j \in I, r \leq k} P_j \wedge M_j(Y_{\text{next}(\bar{x}', r, j), \text{out}(\bar{x}', r)})$$

Where $x'_i = \max(x_i, j)$ and $\text{next}(\bar{x}', r, j)$ is the new register contents after Even resets the register of rank r . That is to say $\text{next}(\bar{x}', r, j)_i = x'_i$ for $i > r$, but x_{i-1} for $1 < i \leq r$, and 0 for $i = 1$. Furthermore, $\text{out}(p, r)$ is $2r + 1$ for odd p and $2r$ otherwise.

Then, the formula Win_I^k is given by the system of equations consisting of $Y_{\bar{x}, p}$ where $\bar{x} \in I^k$ and $o \in [1..2k + 1]$, paired with $\Omega(Y_{m, \bar{x}, o}) = o$ and the entry point $Y_{0, 1}$.

In the formula $Y_{\bar{x}, o}$, the outermost disjunction gives Even the choice to proceed in the underlying parity game by playing to $M_j(Y_{\bar{x}', 1})$ where j is the current priority, or to reset a register. In either case, \bar{x} is updated. Ω guarantees that when a register is reset, a priority of the appropriate magnitude and parity is produced, as computed by $\text{out}()$. Hence by design, the model-checking parity game of Win_I^k on any structure G representing a parity game with priority domain I is functionally equivalent to the parity game $\mathcal{R}_e^k(G)$. \square

Win_I^k has alternation depth k and size $O(kd^k)$: the descriptive complexity of parity games of bounded register-index collapses, just like their algorithmic complexity.

Corollary 6.6. Register-index is bisimulation invariant.

Proof. Win_I^k holds in a parity game G with priority domain I whenever Even wins $\mathcal{R}_e^k(G)$. Dually, we can define OWin_I^k which holds in parity games G of priority domain I whenever Odd wins \mathcal{R}_o^k (note that this is not just the negation of Win_I^k). Recall that a parity game has register-index at most k if \mathcal{R}_o^k and \mathcal{R}_e^k have the same winner.

Then, parity games with priority domain I and of register-index at most k are exactly those that satisfy $(\neg \text{OWin}_I^k \wedge \text{Win}_I^k) \vee (\text{OWin}_I^k \wedge \neg \text{Win}_I^k)$. Since L_μ is bisimulation invariant, so is register-index. \square

Corollary 6.7. The winning regions of parity games of size up to 2^k with p priorities are described by an L_μ formula of size $O(kp^k)$ and alternation-depth k .

Note that the winning regions of parity games of bounded size n are also described by a formula without any fixpoints. However, that formula, obtained is by unfolding all the fixpoints of the standard parity game formula n times. For parity games with p priorities, this is a formula of size $O(2^{pn})$.

We extend the definition of register-index to L_μ -formulas:

Definition 6.8 (Register-index of a formula). A formula ψ of L_μ has register-index k if for all structures \mathcal{M} , the register-index of the model checking parity game for \mathcal{M} and ψ is bounded by k .

Using the L_μ description of the winner of k -register games, we then argue that a formula of register-index k is equivalent to a formula of alternation depth k . In other words, the alternation hierarchy is finite for the fragments of L_μ of bounded register-index. This is a simple consequence of the following theorem which relates the descriptive complexity of the model-checking games of a formula to the formula's semantic complexity.

Theorem 6.9. [24] *Given a formula ψ of L_μ , if for all structures \mathcal{M} it is the case that $\mathcal{M} \models \psi$ if and only if $G(\mathcal{M}, \psi) \models \phi$, then there exists a formula of the alternation depth of ϕ , constructed from ψ and ϕ that is equivalent to ψ .*

Corollary 6.10. *If ψ is an L_μ -formula of register-index k , then ψ is equivalent to a formula of alternation depth k .*

Proof. By definition of register-index for all structures \mathcal{M} , the formula ψ holds in \mathcal{M} exactly when Even wins $\mathcal{R}_e^k(G(\mathcal{M}, \psi))$. From Lemma 6.5, it is therefore the case that $\mathcal{M} \models \psi$ if and only if $G(\mathcal{M}, \psi) \models \text{Win}_I^k$ where I is the priority domain of $G(\mathcal{M}, \psi)$ which depends only on ψ . Then, from Theorem 6.9, ψ is equivalent to a formula of alternation depth k , built from ψ and Win_I^k . \square

Register-index is therefore a decidable upper bound on the semantic complexity of a formula. While the register-index can be as high as the syntactic alternation depth of a formula, it is less affected by syntactic inefficiencies: for example, if ϕ is unsatisfiable, but of high syntactic complexity, a formula $\psi(\phi)$ with subformula ϕ will have syntactic complexity at least as high as ϕ . However, since to win in the model-checking parity game of $\psi(\phi)$ a winning strategy for Even never needs to visit the subgame induced by ϕ , the register-index will only depend on $\psi(\perp)$.

Note that the entanglement of the model-checking games of a formula is unbounded even for simple formulas so entanglement does not lend itself to a similar analysis.

7 Discussion

A. Measuring the complexity of parity games

We introduced register-index, a measure of complexity of parity games that captures the complexity of the priority assignment. It is logarithmically bounded with respect to the size of a parity game.

Existing measures of complexity for parity games seem to be in one way or another agnostic to vertex ownership. Entanglement and register-index are measured on the one-player subgame induced by a winning strategy while the Rabin index and purely structural measures are completely independent of node-ownership. While register-index takes into account two sources of complexity – structure and priority assignment – it leaves the complexity arising from alternations between the two players untouched. However, since one-player games are known to be solvable in polynomial time, it seems plausible that node ownership could be incorporated into a more refined measure which would measure how cycles, priorities and alternations interact.

B. Solving parity games in practice

The algorithm induced by the logarithmic bound on the register-index of parity games uses different techniques from previous algorithms [9, 23] to prove yet again that parity games are solvable in quasi-polynomial time. It does not improve on the running time of state-of-the-art algorithms; however, in addition to its rich theory, it enjoys a family of parameterised versions, each of which solves parity games for a class of games of bounded register-index. The class of parity games of register-index as low as 1 already contains parity games of otherwise arbitrary complexity, including families of games which trigger the worst-case performance of recursive algorithms [14], strategy improvement algorithms [13], and one of the existing quasi-polynomial algorithms [12]. Conversely, single-player games with high register-index are trivial for algorithms like strategy-improvement which recognise a winning strategy quickly. Hence k -register games are a complementary solving method to state-of-the-art algorithms. Families of high register-index require a particularly intricate structure, as demonstrated in Lemma 4.7; how often such games are encountered in the wild is left as an open question. However, it seems likely that i) solving k -register games for small k could be a practical approach to solving parity games, and ii) existing solvers could potentially benefit from taking into account the register-index of games.

C. Parity games algorithms and L_μ

The final part of this paper discussed how register-index and the new quasi-polynomial algorithm relate to complexity in L_μ . This discussion connects two long-standing open problems: the decidability of the L_μ alternation hierarchy – that is, deciding, given a formula, whether it is equivalent to a formula of lower alternation depth – and the complexity of model-checking L_μ – that is, the complexity of solving parity games. There are trade-offs between formula size and alternation depth that echo the trade-offs between solving a smaller parity game with many priorities, and a larger game with fewer priorities. It remains open whether the other quasi-polynomial algorithms can also be understood in logical terms. Bojańczyk and Czerwinski present the original quasi-polynomial algorithm in terms of reachability automata [7], which is a step in this direction.

Acknowledgements. I thank Udi Boker and Patrick Totzke for their comments on drafts of this paper. I am also grateful to several anonymous reviewers for their insights and suggestions.

References

- [1] Walid Belkhir and Luigi Santocanale. 2007. Undirected graphs of entanglement 2. *FSTTCS 2007: Foundations of Software Technology and Theoretical Computer Science* (2007), 508–519.
- [2] Massimo Benerecetti, Daniele Dell'Erba, and Fabio Mogavero. 2017. Robust Exponential Worst Cases for Divide-et-Impera Algorithms for Parity Games. *Proceedings of International Symposium on Games, Automata, Logics, and Formal Verification* (2017), 121–135. <https://doi.org/10.4204/EPTCS.256.9>
- [3] Dietmar Berwanger, Anuj Dawar, Paul Hunter, and Stephan Kreutzer. 2006. DAG-width and parity games. In *STACS*, Vol. 6. Springer, 524–536. https://doi.org/10.1007/11672142_43
- [4] Dietmar Berwanger and Erich Grädel. 2005. Entanglement - a measure for the complexity of directed graphs with applications to logic and games. In *International Conference on Logic for Programming Artificial Intelligence and Reasoning*. Springer, 209–223. https://doi.org/10.1007/978-3-540-32275-7_15
- [5] Dietmar Berwanger, Erich Grädel, Lukasz Kaiser, and Roman Rabinovich. 2012. Entanglement and the complexity of directed graphs. *Theoretical Computer Science* 463 (2012), 2–25. <https://doi.org/10.1016/j.tcs.2012.07.010>
- [6] Dietmar Berwanger, Erich Grädel, and Giacomo Lenzi. 2007. The variable hierarchy of the μ -calculus is strict. *Theory of Computing Systems* 40, 4 (2007),

- 437–466.
- [7] Mikołaj Bojańczyk and Wojciech Czerwiński. 2018. *Automata-Toolbox*. University of Warsaw. <https://www.mimuw.edu.pl/~bojan/papers/toolbox-reduced-feb6.pdf> last checked April 2018.
 - [8] Julian Bradfield and Colin Stirling. 2007. Modal μ -calculi. *Studies in logic and practical reasoning* 3 (2007), 721–756.
 - [9] Cristian S. Calude, Sanjay Jain, Bakhadyr Khoussainov, Wei Li, and Frank Stephan. 2017. Deciding parity games in quasipolynomial time. In *Proceedings of the 49th Annual ACM SIGACT Symposium on Theory of Computing*. ACM, 252–263. <https://doi.org/10.1145/3055399.3055409>
 - [10] Anuj Dawar and Erich Grädel. 2008. The Descriptive Complexity of Parity Games. In *Computer Science Logic*. Springer, 354–368. https://doi.org/10.1007/978-3-540-87531-4_26
 - [11] E Allen Emerson and Charanjit S Jutla. 1991. Tree automata, mu-calculus and determinacy. In *Foundations of Computer Science, 1991. Proceedings., 32nd Annual Symposium on*. IEEE, 368–377.
 - [12] John Fearnley, Sanjay Jain, Sven Schewe, Frank Stephan, and Dominik Wojtczak. 2017. An ordered approach to solving parity games in quasi polynomial time and quasi linear space. In *Proceedings of the 24th ACM SIGSOFT International SPIN Symposium on Model Checking of Software*. ACM, 112–121. <https://doi.org/10.1145/3092282.3092286>
 - [13] Oliver Friedmann. 2009. An exponential lower bound for the parity game strategy improvement algorithm as we know it. In *Logic In Computer Science, 2009. LICS'09. 24th Annual IEEE Symposium on*. IEEE, 145–156. <https://doi.org/10.1109/LICS.2009.27>
 - [14] Oliver Friedmann. 2011. Recursive algorithm for parity games requires exponential time. *RAIRO-Theoretical Informatics and Applications* 45, 4 (2011), 449–457. <https://doi.org/10.1051/ita/2011124>
 - [15] Moses Ganardi. 2015. Parity Games of Bounded Tree-and Clique-Width.. In *FoSSaCS*. 390–404. https://doi.org/10.1007/978-3-662-46678-0_25
 - [16] Hugo Gimbert and Rasmus Ibsen-Jensen. 2017. A short proof of correctness of the quasi-polynomial time algorithm for parity games. *arXiv preprint arXiv:1702.01953* (2017).
 - [17] Erich Grädel, Łukasz Kaiser, and Roman Rabinovich. 2009. Directed Graphs of Entanglement Two. In *International Symposium on Fundamentals of Computation Theory*, Vol. 9. Springer, 169–181.
 - [18] Erich Grädel, Phokion G Kolaitis, Leonid Libkin, Maarten Marx, Joel Spencer, Moshe Y Vardi, Yde Venema, and Scott Weinstein. 2007. *Finite Model Theory and its applications*. Springer Science & Business Media.
 - [19] Paul Hunter and Stephan Kreutzer. 2007. Digraph measures: Kelly decompositions, games, and orderings. In *Proceedings of the eighteenth annual ACM-SIAM symposium on Discrete algorithms*. Society for Industrial and Applied Mathematics, 637–644. <https://doi.org/10.1016/j.tcs.2008.02.038>
 - [20] Michael Huth, Jim Huan-Pu Kuo, and Nir Piterman. 2015. The Rabin index of parity games: Its complexity and approximation. *Information and Computation* 245 (2015), 36–53. <https://doi.org/10.1016/j.ic.2015.06.005>
 - [21] Marcin Jurdziński. 1998. Deciding the winner in parity games is in $UP \cap co-UP$. *Inform. Process. Lett.* 68, 3 (1998), 119–124. [https://doi.org/10.1016/S0020-0190\(98\)00150-1](https://doi.org/10.1016/S0020-0190(98)00150-1)
 - [22] Marcin Jurdziński. 2000. Small progress measures for solving parity games. In *Annual Symposium on Theoretical Aspects of Computer Science*. Springer, 290–301. https://doi.org/10.1007/3-540-46541-3_24
 - [23] Marcin Jurdziński and Ranko Lazic. 2017. Succinct progress measures for solving parity games. In *2017 32nd Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)*, Vol. 00. 1–9. <https://doi.org/10.1109/LICS.2017.8005092>
 - [24] Karoliina Lehtinen. 2017. The descriptive complexity of modal μ model-checking parity games. *Proceedings of International Symposium on Games, Automata, Logics, and Formal Verification, 2017* (2017). <https://doi.org/10.4204/EPTCS.256.6>
 - [25] Tom van Dijk. 2018. Oink: an Implementation and Evaluation of Modern Parity Game Solvers. *CoRR abs/1801.03859* (2018). [arXiv:1801.03859](https://arxiv.org/abs/1801.03859) <http://arxiv.org/abs/1801.03859>
 - [26] Thomas Wilke. 2001. Alternating Tree Automata, Parity Games, and Modal μ -Calculus. *Bulletin of the Belgian Mathematical Society Simon Stevin* 8, 2 (2001), 359.