

# Iterative Context Bounding for Systematic Testing of Multithreaded Programs

Madan Msvathi   Shaz Qadeer  
Microsoft Research

Presented by Yuli Ye  
2008

# Concurrent software

- Operating systems
- Mail servers, web servers
- Databases
- Device drivers
- Games
- .....

# Concurrency is important

- Internet and multi-user environment require more and more applications to handle concurrency
- Hardware changes, e.g., multiple cores, require software to harness the hardware parallelism to improve performance

# Concurrency is a problem

- Windows 2000 hot fixes
  - Concurrency and synchronization errors are most common coding errors
- Windows Server 2003 late cycle defects
  - Synchronization errors are second in the list, next to buffer overruns
- Race conditions can lead to security vulnerability

# Concurrent programs are hard

- It is hard to write a correct concurrent program
  - People get more used to think sequentially than concurrently
- It is also hard to test a concurrent program
  - Thread interleaving may create subtle errors which are hard to catch
- Even when found, errors are hard to debug
  - An error may not repeat itself very often
  - An error may occur far away from its source

# Traditional testing methods

- Find interesting test scenario
  - Create some test cases that we think are “interesting”
- Stress testing
  - Run thousand threads for days
- Force scheduling variety
  - Use random() and sleep()

## **Disadvantages of the above three approaches**

- Many are heuristic based
- No guarantees on coverage
- Rely too much on the tester

# Testing with model checking

- Advantages
  - Systematically executes each thread schedule to control non-determinism
  - Capable of reproducing an error once found and hence easier to debug
- Disadvantages
  - State explosion: the number of possible program behaviors grow explosively with the size of the program
  - Almost infeasible for large concurrent program with limited resource of memory and time

# State explosion I

Thread 1

$x = 1;$   
 $y = 1;$

Thread 2


$x = 2;$   
 $y = 2;$

$x = 1;$

$y = 1;$


$x = 2;$

$y = 2;$





# State explosion II



## Theorem

With  $n$  threads and at most  $k$  steps at each thread, the total number of execution maybe as large as

$$(nk)! / (k!)^n \geq (n!)^k$$

# Iterative depth bounding

**Iterative depth bounding** limits the execution with a bounded number of steps

- Runs out of resource quickly as the depth is increased
- Most useful for program with small depth from the initial state, e.g., message-passing software
- Does not work well for multithread programs with fine-grained interaction through shared memory
- Usually have a very poor coverage of states explored

# CHESS: Iterative context bounding

A **context switch** occurs at a schedule points if the scheduler chooses a thread different from the current running thread


**There are two kinds of context switches**

- Preemptions – forced by the scheduler
  - e.g. time-slice expiration
- Non-preemptions – a thread voluntarily yields
  - e.g. blocking on an unavailable resource

In **context bounding**, we bound the number of preemptions but leave the number of non-preemptions unconstrained

# Benefits of context bounding 1

## Polynomial state space



## Theorem

If a program has at most  $c$  preemptions and  $n$  threads. Each thread has at most  $k$  steps of with at most  $b$  are potentially-blocking, the total number of execution is bounded by

$$n^k C_c \cdot (nb+c)! = O((n^2kb)^c \cdot (nb)!)$$

# Benefits of context bounding 2

## Possible deep exploration with small bounds

- The number of steps within each context remains unbounded, so we overcome the limitation of depth bounding
- The number of non-preemption within each context remains unbounded, therefore even a bound of zero may lead to complete termination executions

# Benefits of context bounding 3

## **Better coverage metric**

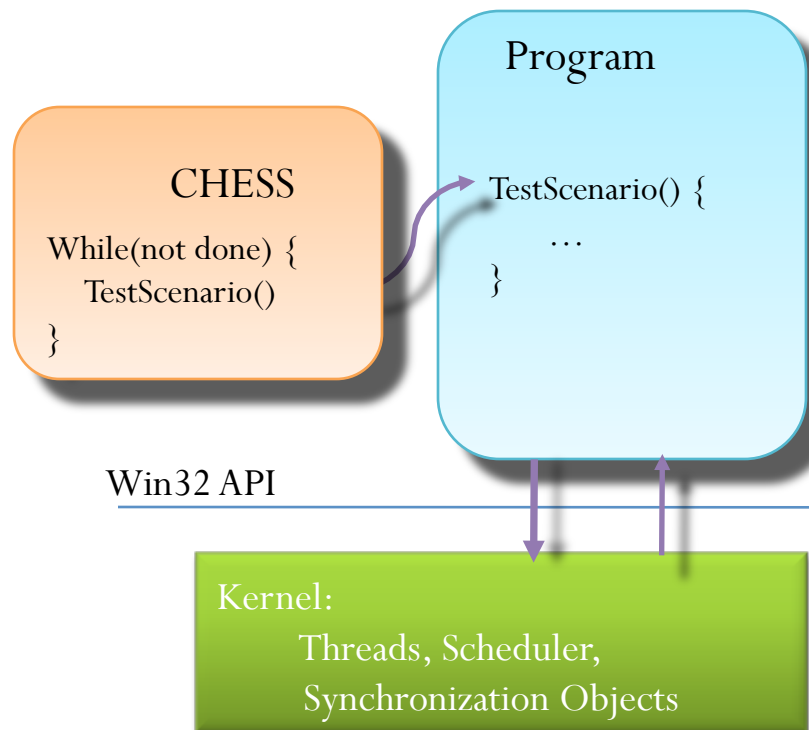
- Finds the smallest number of preemptions to the error
- Gives an estimate on the possible bugs remaining in the program and hence an estimate on the chance of their occurrence in practice

# Benefits of context bounding 4

## **Many bugs within small number of preemptions**

- Based on a non-blocking implementation of the work-stealing queue algorithm
  - Bounded circular buffers accessed concurrently by two threads
- A test harness and three bugs are given
  - Each bug found with at most 2 preemption
  - Although execution with 35 preemptions are possible

# Architecture of CHES



Tester Provides a Test Scenario

**CHES runs test scenario in a loop**

- Every run takes a different interleaving
- Every run is repeatable

**Intercept synchronization and threading calls**

- Control and schedule non-determinism

**Detects**

- Assertion violations
- Deadlock
- Livelock
- Data-races



# Conditions on TestScenario()

- TestScenario() should terminate in all interleavings
- TestScenario() should be idempotent
  - Free all resources
  - Reset global states
- TestScenario() should not interfere with other tasks in the program being tested

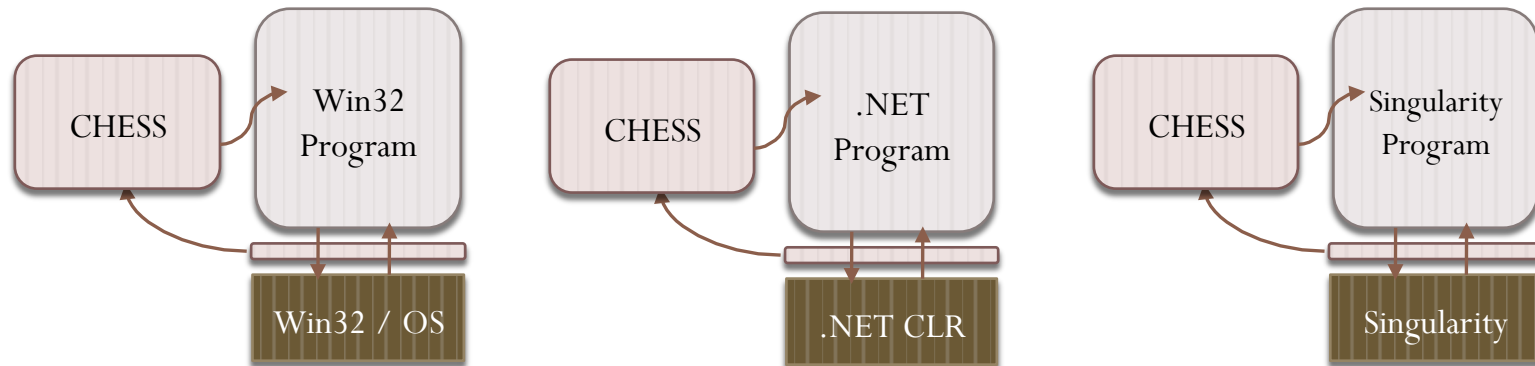
## **Observation:**

Existing stress tests usually satisfy these properties

# Perturb the system as little as possible

- Run the system as is
  - On the actual OS, hardware
  - Using system threads
  - Using system synchronization objects
- Advantages
  - Avoid reporting false errors
  - Easy to add to existing test frameworks
  - Use existing debuggers

# CHESS methodology generalizes



- CHESS works for
  - Unmanaged programs, such as code written in C and C++
  - Managed programs, such as code written in C#
  - Singularity applications
- With appropriate wrappers, can work for Java and Linux applications

# CHESS: The algorithm I

- Effectively search the state space of a program by systematically bounding the number of preemptions
- Assume the program is data-race free
- Context switch only at synchronization points
- Check for data-races in each execution

# CHESS: The algorithm II

**Input:** initial state  $s_0 \in \text{State}$  and context switch bound  $\text{csb}$

```
1 struct WorkItem { State state; Tid tid; int phase; }
2 Queue<WorkItem> workQueue;
3 WorkItem w;
4 int currPhase;
5 for  $t \in \text{Tid}$  do
6      $w.\text{state} := s_0$ ;
7      $w.\text{tid} := t$ ;
8      $w.\text{phase} := 0$ ;
9      $\text{workQueue.Add}(w)$ ;
10 end
11  $\text{currPhase} := 0$ ;
12 while  $\neg \text{workQueue.Empty}()$  do
13      $w := \text{workQueue.Front}()$ ;
14      $\text{workQueue.Pop}()$ ;
15     if  $\text{currPhase} < w.\text{phase}$  then
16         /* explored  $(\text{currPhase} + 1) * \text{csb} + \text{currPhase}$ 
17            preempting context switches */
18          $\text{currPhase} := w.\text{phase}$ ;
19     end
20     Search( $w, 0$ );
21 end
```

```
20 Search(WorkItem  $w$ , int  $\text{ncs}$ ) begin
21     if  $\neg w.\text{state.Enabled}(w.\text{tid})$  then
22         return;
23     end
24     WorkItem  $x$ ;
25      $x.\text{state} := w.\text{state.Execute}(w.\text{tid})$ ;
26      $x.\text{tid} := w.\text{tid}$ ;
27      $x.\text{phase} := w.\text{phase}$ ;
28     Search( $x, \text{ncs}$ );
29     for  $t \in \text{Tid} \setminus \{w.\text{tid}\}$  do
30          $x.\text{tid} := t$ ;
31         if  $\neg x.\text{state.Enabled}(w.\text{tid})$  then
32              $x.\text{phase} := w.\text{phase}$ ;
33             Search( $x, \text{ncs}$ );
34         else if  $\text{ncs} = \text{csb}$  then
35              $x.\text{phase} := w.\text{phase} + 1$ ;
36              $\text{workQueue.Push}(x)$ ;
37         else
38              $x.\text{phase} := w.\text{phase}$ ;
39             Search( $x, \text{ncs} + 1$ );
40         end
41     end
42 end
```

**Algorithm 1:** Iterative context bounding

# Why does this work?

## **Theorem**

To check a program, it is sufficient to insert a scheduling point before a synchronization operation in the program, provided that the algorithm also checks for data-races

**The strategy is essentially a partial-order reduction**

# Empirical evaluation

**Evaluation is done on a set of benchmark programs**

- Bluetooth
- File system model
- Work-stealing queue
- APE
- Dryad channels
- Transaction manager

# Characteristics of benchmarks

Programs	LOC	Num Threads	Max K	Max B	Max c
Bluetooth	400	3	15	2	8
File System Model	84	4	20	8	13
Work Stealing Q.	1266	3	99	2	35
APE	18947	4	247	2	75
Dryad Channels	16036	5	273	4	167

**Table 1.** Characteristics of the benchmarks. For each benchmark, this table reports the number of lines, the number of threads allocated by the test driver. For an execution, K is the total number of steps, B is the number of blocking instructions, and c is the number of preempting context switches. The table reports the maximum values of K, B, and c seen during our experiments.




## Bugs found with small context bound

Programs	Total Bugs	Bugs with Context Bound			
		0	1	2	3
Bluetooth	1	0	1	0	0
Work Stealing Queue	3	0	1	2	0
Transaction Manager	3	0	0	2	1
APE	4	2	1	1	0
Dryad Channels	3	1	2	0	0

**Table 2.** For a total of 14 bugs that our model checker found, this table shows the number of bugs exposed in executions with exactly  $c$  preempting context switches, for  $c$  ranging from 0 to 3. The 7 bugs in the first three programs was previously known. Iterative context-bounding algorithm found the 7 previously *unknown* bugs in Dryad and APE.


# Coverage vs. Context bound



# Dryad bugs

- Total of 7 bugs are found in spite of careful regression testing and months of production use
- The use-after-free bug has long error trace but requires only one preemption
  - Depth bounding is hard to find
- The error trace has 6 non-preempting context switches
  - Unrestricting non-preemption is important

# Coverage vs. time in Dryad



# Conclusion

- Currency is important but hard to get it right, building robust concurrency software remains a challenge
- Traditional testing and debugging methods are unsatisfying in providing guarantees of detecting and correcting errors
- CHES is a systematic testing tool that provides:
  - Good coverage without sacrificing the ability to go deep into the state space
  - Good integration capability with the existing test frameworks
  - Replay capability for debugging
- Iterative context bounding is a useful approach in designing concurrency testing tools

# Thank you!

- Musuvathi, M and Qadeer, S. Iterative context bounding for systematic testing of multithreaded programs. In *Proceedings of the 2007 ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI '07)*, pages 446 - 455. San Diego, California, USA, June 2007.
- <http://research.microsoft.com/projects/CHESS/>
- <http://research.microsoft.com/projects/CHESS/IterativeContextBounding.pdf>