

Can Message Buffers Be Axiomatized in Linear Temporal Logic?*

A. P. SISTLA

University of Massachusetts, Amherst, Massachusetts

E. M. CLARKE

Carnegie-Mellon University, Pittsburgh, Pennsylvania

N. FRANCEZ

The Technion, Haifa, Israel

AND

A. R. MEYER

Massachusetts Institute of Technology, Cambridge, Massachusetts

Message passing is one of the primary modes of interprocess communication in a distributed system. In this paper we investigate the possibility of *characterizing* and *axiomatizing* different message passing systems in temporal logic. Specifically, we consider FIFO buffers (queues), LIFO buffers (stacks) and unordered buffers (bags). We show that all bounded buffers are characterizable in propositional temporal logic (PTL) and so are axiomatizable. We prove that the theory of unbounded FIFO buffers is π_1^1 -complete and so is not axiomatizable. We also prove that the theories of unbounded LIFO and unordered buffers are decidable and hence are axiomatizable. © 1984 Academic Press, Inc.

1. INTRODUCTION

Exchange of information between executing processes is one of the primary reasons for process interaction. Many distributed systems implement explicit message passing primitives to facilitate intercommunication. Typically, a process executes a *write* command to pass a message to another process, and the target process accepts the message by

* This research was partially supported by NSF Grant MCS79-08365, and NSF Grant MCS80-10707.

executing a *read* command. The semantics of *write* and *read* may differ considerably depending on the methods used for storing or buffering messages that have been sent but not yet accepted by the receiving process.

Because message passing systems are so widely used, it is important to develop formal techniques for reasoning about them. In this paper we investigate the possibility (impossibility) of using *linear temporal logic* to characterize and axiomatize different message buffering mechanisms. This logic was originally introduced as a formal system for reasoning about sequences of events that are totally ordered in time. Recently, linear temporal logic has been proposed by Manna and Pnueli [2] and Owicki and Lamport [3] as an appropriate formal system for reasoning about parallel programs. The logic permits the description of a program's execution history without the explicit introduction of program states or of time. Moreover, important correctness properties such as mutual exclusion, freedom from deadlock and absence of starvation can be elegantly expressed in this system.

Specifically, we consider FIFO buffers (*queues*), LIFO buffers (*stacks*) and unordered buffers (*bags*). The set of distinct messages that can be written into the buffer is called the *message alphabet*. We specify a message buffer by the set of all valid infinite input/output message sequences.

Characterizing a message buffer in temporal logic consists of obtaining a formula that is true exactly on the set of sequences specifying the buffer. We show that it is possible to characterize bounded buffers over a finite alphabet in propositional linear temporal logic (PTL) and show how to obtain such a characterization. Although such bounded buffers can be characterized using ω -regular expressions, it is not obvious that they can be characterized in PTL since this logic is provably less expressive than ω -regular expressions [4]. Using first order temporal logic, we can give a uniform characterization of bounded buffers which is independent of the message alphabet. We also show that unbounded buffers are not characterizable in PTL.

Since the formulae we obtain for bounded buffers may be quite complicated, we introduce an extension of PTL in which certain atomic propositions are designated as auxiliary. The auxiliary propositions are not interpreted and are treated like existentially quantified monadic predicates. We give simple and succinct formulae in the extended logic which characterize bounded message buffers over a finite alphabet.

We also consider the problem of *axiomatizing* the various types of message buffers. A model of a message buffer is an infinite sequence of states denoting a series of legal read/write operations on the buffer. The theory of a message buffer is the set of all PTL formulae which are true in all models of the buffer. We say that a message buffer is axiomatizable iff the theory of the buffer is recursively enumerable. A simple and complete

axiom system for all the bounded buffers can be given using the characterization of these buffers in PTL. We show that, in general, unbounded FIFO buffers are not axiomatizable in PTL. We also show that unbounded LIFO buffers and unbounded unordered buffers are axiomatizable in PTL and in fact, we prove that the theories of these buffers are decidable.

One of our principal motivations for this work has been to study possible axiomatizations of computational behavior by temporal logic, and so we have formulated most of our results and indeed the title of this paper in terms of axiomatizability. Nevertheless, the detailed structure of axiom system actually plays no role in our results. All our non-axiomatizability results are in fact proofs that various theories are not recursively enumerable, and some of our axiomatizability results are proofs that the corresponding theories are recursively enumerable.

The paper is organized as follows: Section 2 defines the syntax and semantics of the linear temporal logic that we use in the remainder of the paper. In Section 3 we specify precisely those properties of message buffers that we would like to capture in temporal logic. Section 4 shows that bounded buffers can be characterized in the logic and describes how uninterpreted auxiliary proposition symbols can be added to simplify this construction. We also prove that it is impossible to give a characterization of unbounded message buffers in PTL. In Section 5 we consider axiomatization of message buffers in PTL. We show that unbounded FIFO message buffers are not axiomatizable in PTL while unbounded LIFO and unordered buffers are axiomatizable. The paper concludes in Section 6 with a summary and discussion of our results.

2. LINEAR TEMPORAL LOGIC

The language of PTL uses certain symbols called atomic propositions drawn from a finite set \mathcal{P} , the propositional connectives \wedge , \sim , and the temporal modalities **X** (“next-time”), **U** (“until”), **Y** (“last-time”), **S** (“since”), together with the parenthesis.

A *well-formed formula* in PTL is either an atomic proposition or is of the form $(\sim f_1)$, $(f_1 \wedge f_2)$, $(\mathbf{X} f_1)$, $(f_1 \mathbf{U} f_2)$, $(\mathbf{Y} f_1)$, $(f_1 \mathbf{S} f_2)$, where f_1, f_2 are well-formed formulae. We avoid parenthesis whenever the implied parsing of the formula is understood from the context. In addition, we use the following abbreviations:

$$(f_2 \vee f_1) \equiv \sim(\sim f_1 \wedge \sim f_2), (f_1 \supset f_2) \equiv \sim f_1 \vee f_2, \\ \mathbf{F} f \equiv (\mathbf{True} \mathbf{U} f), \mathbf{G} f \equiv \sim(\mathbf{F} \sim f).$$

F, **G** are the “sometimes,” “always” operators, respectively. A state is a mapping from \mathcal{P} into the set $\{\mathbf{True}, \mathbf{False}\}$. Let \mathcal{M} denote the set of all

states. Note that \mathcal{M} is finite. A model is a ω -sequence of states. An interpretation is an ordered pair (t, i) , where t is a model and $i \geq 0$ is an integer specifying the present state. We define the truth of a formula f in an interpretation (t, i) ($t, i \models f$) inductively as follows:

$t, i \models P$	where P is atomic iff $t_i(P) = \text{True}$;
$t, i \models f_1 \wedge f_2$	iff $t, i \models f_1$ and $t, i \models f_2$;
$t, i \models \sim f_1$	iff not $(t, i \models f_1)$;
$t, i \models \mathbf{X}f_1$	iff $t, i + 1 \models f_1$;
$t, i \models f_1 \mathbf{U}f_2$	iff there exists a $k \geq i$ such that $t, k \models f_2$ and for all j such that $i \leq j < k$, $t, j \models f_1$;
$t, i \models \mathbf{Y}f_1$	iff $i > 0$ and $t, i - 1 \models f_1$;
$t, i \models f_1 \mathbf{S}f_2$	iff there exists a $k \leq i$ $t, k \models f_2$ and for all j such that $k < j \leq i$ $t, j \models f_1$;

It is to be noted that $(t, i) \models \mathbf{F}(f)$ iff there exists $j \geq i$ such that $(t, j) \models f$, and $(t, i) \models \mathbf{G}(f)$ iff for all $j \geq i$ $(t, j) \models f$. A formula is *satisfiable* iff it is true in some interpretation and it is *valid* iff it is true in all interpretations. A consistent and complete axiomatization for the set of validities of PTL is presented in [14]. In [10] such an axiomatic system is presented for a restricted PTL that uses \mathbf{X} , \mathbf{U} as the only temporal modalities.

3. MODELLING MESSAGE BUFFERS

We characterize a message buffer by the set of legal *read/write* sequences allowed on the bugger. A *write* operation writes a message into the buffer; a *read* operation reads a message from the buffer and deletes it. At most one read or write operation is permitted at any instant of time. In the case of bounded buffers a write request will be rejected when the buffer is full; similarly, a read request on an empty buffer will be rejected. Rejected read/write requests are not included in the sequences of legal operations characterizing the buffer. We consider below three types of message buffers: FIFO buffers (*queues*), LIFO buffers (*stacks*), and unordered buffers (*bags*). In FIFO buffers the *earliest* written message currently in the buffer is the output for the next read request; with LIFO buffers the *latest* written message currently in the buffer is used; and with unordered buffers any message present in the buffer may be output.

Let Σ be the message alphabet and \mathcal{P}_Σ be the set of atomic propositions $\{R_\sigma \mid \sigma \in \Sigma\} \cup \{W_\sigma \mid \sigma \in \Sigma\}$ such that $\mathcal{P}_\Sigma \subseteq \mathcal{P}$. Let $ST = \{\phi \mid \phi: \mathcal{P} \rightarrow \{\text{True}, \text{False}\} \text{ such that } \phi(P) = \text{True} \text{ for at most one } P \text{ in } \mathcal{P}_\Sigma\}$. If $R_\sigma(W_\sigma)$ is true in a state, then it indicates that the message σ is read (written) from (into) the buffer in that state. Note that at most one operation (a read or a write) occurs in any state in ST .

Let $t \in ST^* \cup ST^\omega$ and $i_0 < i_1 < \dots$, be all the instances at which some messages $\sigma_0, \sigma_1, \dots$, are read from the buffer, i.e., $t_{i_k}(R_{\sigma_k}) = \text{True}$ for $k \geq 0$. Then $\pi_r(t)$ denotes the sequence $(\sigma_0, \sigma_1, \dots)$. Similarly, we define $\pi_w(t)$. Intuitively, $\pi_r(t)$ ($\pi_w(t)$) denotes the sequence of messages read from the buffer (written into the buffer) in t . Let $t[i]$ denote the sequence (t_0, t_1, \dots, t_i) , then $nb(i) = \text{length}(\pi_w(t[i])) - \text{length}(\pi_r(t[i]))$ is the number of messages in the buffer just after the instance i .

Let $\text{FIFO}_{\Sigma,k}$ be the set of all infinite sequences of states which denote legal series of read/write operations on a FIFO buffer of size k . Similarly, let $\text{LIFO}_{\Sigma,k}$, $\text{UNOR}_{\Sigma,k}$ be the corresponding sets of sequences for LIFO and unordered buffers, respectively. Unbounded buffers will be denoted in this scheme of notation by $k = \infty$. For $k \geq 0$,

$\text{FIFO}_{\Sigma,k} = \{t \in ST^\omega \mid \text{for all } i \geq 0, \pi_r(t[i]) \text{ is a prefix of } \pi_w(t[i]) \text{ and } nb(i) \leq k\}$;

$\text{LIFO}_{\Sigma,k} = \{t \in ST^\omega \mid \text{for all } i \geq 0, 0 \leq nb(i) \leq k \text{ and if for some } \sigma \in \Sigma, (t, i) \models R_\sigma \text{ then there exists a } j < i \text{ such that } (t, j) \models W_\sigma \text{ and } nb(j) = nb(i-1) \text{ and for all } l \text{ such that } j \leq l < i \text{ } nb(l) \geq nb(j)\}$;

$\text{UNOR}_{\Sigma,k} = \{t \in ST^\omega \mid \text{for all } i \geq 0, 0 \leq nb(i) \leq k \text{ and for all } \sigma \in \Sigma, \text{ the number of writes of the message } \sigma \text{ up to } i \geq \text{the number of reads of the message } \sigma \text{ up to } i\}$.

For a finite alphabet Σ , a formula f in PTL characterizes a FIFO message buffer of size k iff $\{t \mid (t, 0) \models f\} = \text{FIFO}_{\Sigma,k}$. Similarly, we define what it means to characterize LIFO and unordered buffers.

We say that a sequence t is a behavior of a FIFO (LIFO or unordered) buffer with *liveness* property iff t is a legal sequence as defined above and the buffer becomes empty infinitely often in t . That is, $nb(i) = 0$ in t for infinitely many values of i . This property guarantees that every message written into the buffer is eventually read which is a liveness property. As above we define what it means to characterize a message buffer with liveness property.

A *model* or *history* of a message buffer is an infinite sequence of states denoting a legal series of read/write operations on the buffer, as defined above. The *theory* of a message buffer is the set of all PTL formulae which are true in all interpretations of the form $(t, 0)$, where t is a model of the buffer. We say that a message buffer is *axiomatizable* in PTL iff the theory of the buffer is recursively enumerable.

4. CHARACTERIZING MESSAGE BUFFERS

4.1. Direct Characterization in PTL

In this section we show how we can characterize bounded buffers over a finite alphabet using PTL. We let fb_k , lb_k , ub_k , respectively, denote formulae in PTL characterizing FIFO, LIFO, and unordered message buffers of size k over the finite message alphabet Σ . First we describe how to obtain the formulae for the buffers with $k=1$ and $k=2$.

Let Σ be a finite message alphabet, and $\mathcal{P}_\Sigma = \{R_\sigma \mid \sigma \in \Sigma\} \cup \{W_\sigma \mid \sigma \in \Sigma\}$ be the set of atomic propositions. Throughout this section we use the following abbreviations:

$$\begin{aligned} W &= \bigvee_{\sigma \in \Sigma} W_\sigma \\ R &= \bigvee_{\sigma \in \Sigma} R_\sigma \\ Ex &= \left(\bigwedge_{\sigma_1 \neq \sigma_2} \sim (R_{\sigma_1} \wedge R_{\sigma_2}) \right) \wedge \left(\bigwedge_{\sigma_1 \neq \sigma_2} \sim (W_{\sigma_1} \wedge W_{\sigma_2}) \right) \wedge \sim (W \wedge R) \\ I &= \mathbf{G}(Ex). \end{aligned}$$

I asserts that at any instant at most one operation occurs on the buffer.

In the case of buffer size = 1, FIFO, LIFO, and unordered buffers are identical and the buffer behavior is as follows:

1. The *writes* and *reads* occur alternately;
2. The message read in each *read* operation is the message written by the previous *write* operation. Thus, $fb_1 = I \wedge f_a \wedge f_b$, where

$$\begin{aligned} f_a &= \mathbf{G}((W \wedge \mathbf{XF}(R \vee W)) \supset \mathbf{X}(\sim W \mathbf{U} R)) \\ &\quad \wedge \mathbf{G}((R \wedge \mathbf{XF}(R \vee W)) \supset \mathbf{X}(\sim R \mathbf{U} W)); \\ f_b &= \mathbf{G} \left(\bigwedge_{\sigma \in \Sigma} (R_\sigma \supset (\sim W \mathbf{S} W_\sigma)) \right). \end{aligned}$$

The first (second) conjunct in f_a asserts that every write (read) operation which is not the last operation on the buffer is followed by a read (write) operation before any other write (read) operation. It is easily seen that f_a and f_b assert properties (1) and (2), respectively.

Intuitively, the operation of a buffer of size two can be described as follows. Initially, *writes* and *reads* occur alternately. This continues until two writes occur successively without a *read* operation in between, and the buffer becomes full (formula l_2 expresses this). Subsequently, *reads* and

writes will again begin to alternate. After each *read* the buffer will have one message and after each *write* operation the buffer becomes full. This may continue forever, or until two *reads* occur successively without a *write* in between, making the buffer empty (r_2 expresses this); now the previous sequence repeats. This behavior is common for FIFO, LIFO, and unordered buffers of size two. The formulae l_2, r_2 are given below:

$$l_2 = W \wedge \mathbf{X}(\sim R \mathbf{U} W)$$

$$r_2 = R \wedge \mathbf{Y}(\sim W \mathbf{S} R).$$

In the remainder of this section we will frequently use the formula $\text{alt}(p, q, c)$ given below:

$$\text{alt}(p, q, c) = [(g \mathbf{U} c) \vee \mathbf{G}(g \wedge \sim c)] \wedge [(\sim c \mathbf{U} p) \supset (\sim q \mathbf{U} p)],$$

where

$$g = (p \supset \mathbf{X}(\sim p \mathbf{U} q)) \wedge (q \supset [\mathbf{X}(\sim q \mathbf{U} p) \vee \mathbf{X}(\sim(p \vee q) \mathbf{U} c)]).$$

The first conjunct in $\text{alt}(p, q, c)$ asserts that either there is a future instance at which c occurs and until this instance p, q occur alternately, or throughout the future p, q occur alternately without c occurring anywhere. The second conjunct asserts that if p occurs then it occurs before q . Thus, the previous intuitive description of the behavior of the buffer of size two is captured by the formula bv given below.

$$bv = \text{alt}(W, R, l_2) \wedge \mathbf{G}[l_2 \supset \mathbf{X} \text{alt}(W, R, r_2)] \wedge \mathbf{G}[r_2 \supset \mathbf{X} \text{alt}(W, R, l_2)];$$

bv asserts that l_2, r_2 occur alternately with alternating *read* and *writes* occurring in between. Any *read* after l_2 but before any following r_2 is on a buffer containing one message. Any *read* after an r_2 but before any following l_2 and any *read* before the first l_2 are on a buffer containing one element. The formulae *read-on-full*, *read-on-single* given below characterize *reads* on a full buffer and *reads* on a buffer with one message, respectively.

$$\text{read-on-full} = R \wedge (\sim r_2 \mathbf{S} l_2)$$

$$\text{read-on-single} = R \wedge [(\sim l_2 \mathbf{S} r_2) \vee \sim(\text{True} \mathbf{S} l_2)].$$

For FIFO buffers, a *read* on a full buffer reads the message written by the *write* before the previous *write*.

$$fb_2 = I \wedge bv \wedge g \wedge h,$$

where

$$g = \mathbf{G} \left(\text{read-on-full} \supset \bigwedge_{\sigma} (R_{\sigma} \supset [\sim W \mathbf{S} (W \wedge \mathbf{Y}(\sim W \mathbf{S} W_{\sigma}))]) \right),$$

$$h = \mathbf{G} \left(\text{read-on-single} \supset \bigwedge_{\sigma} [R_{\sigma} \supset (\sim W \mathbf{S} W_{\sigma})] \right).$$

The formula on the left side of \supset in g is true when *reads* occur on a full buffer, while the formula on the right side asserts that the message read at these instances is the message written by the last but one *write* operation; h asserts that *read* operations on a buffer containing a single message, read the message written by the previous *write* operation.

THEOREM 4.1. *For any ω -sequence of states t , $t, 0 \models fb_2$ iff $t \in \text{FIFO}_{\Sigma, 2}$.*

Let $t \in \text{LIFO}_{\Sigma, 2}$. If $t, i \models r_2$, then there exists $j < i$ such that $t, j \models l_2$. The message read at the instance i is the message written at the instance j . If $t, i \models R$ and $t, i \models \sim r_2$, then the message read at the instance i is the message written in the previous write operation. These properties are expressed by g' and h' , respectively,

$$g' = \mathbf{G}(r_2 \supset \bigwedge_{\sigma \in \Sigma} [R_{\sigma} \leftrightarrow \sim l_2 \mathbf{S} (l_2 \wedge W_{\sigma})])$$

$$h' = \mathbf{G} \left((\sim r_2 \wedge R) \supset \bigwedge_{\sigma \in \Sigma} [R_{\sigma} \leftrightarrow \sim W \mathbf{S} W_{\sigma}] \right).$$

Let $lb_2 = I \wedge bv \wedge g' \wedge h'$.

THEOREM 4.2. *For any ω -sequence of states t , $t, 0 \models lb_2$ iff $t \in \text{LIFO}_{\Sigma, 2}$.*

Let $t \in \text{UNOR}_{\Sigma, 2}$. Then for every $\sigma \in \Sigma$, for all $i \geq 0$ the number of messages of value σ written into the buffer up to the instance i is greater than or equal to the number of messages of value σ read from the buffer up to the instance i , and they do not differ by more than 2. For a given σ , we can obtain a formula bv_{σ} asserting the above property by replacing R by R_{σ} , W by W_{σ} in bv . Let $ub_2 = I \wedge bv \wedge \bigwedge_{\sigma \in \Sigma} bv_{\sigma}$.

The following theorem can be easily proved:

THEOREM 4.3. *For any ω -sequence of states t , $t, 0 \models ub_2$ iff $t \in \text{UNOR}_{\Sigma, 2}$.*

To characterize bounded buffers with the liveness property for buffers of sizes one and two, we add a conjunct to the formula bv by asserting that the buffer becomes empty infinitely often; this can be done by asserting that

whenever l_2 holds then there is a future instance where r_2 holds, and whenever there is a write operation it is eventually followed by a read operation.

The above approach can be extended to characterize bounded buffers of arbitrary size. However, this approach turns out to be complex and cumbersome. For an account of this the interested reader is referred to [9]. Below we take a different approach to prove that all bounded buffers are characterizable in PTL.

4.2. Star-free Regular Sets and Bounded Buffers

We assume that the reader is familiar with regular sets. Let \mathcal{A} be a finite alphabet. A *star-free* set $U \subseteq \mathcal{A}^*$ is inductively defined as follows:

The singleton $\{\delta\}$ is a star-free set over \mathcal{A} where $\delta \in \mathcal{A}$ or δ is the empty string. If V, W are star-free sets over \mathcal{A} , then $V \cup W$, \bar{V} (the complement of V relative to \mathcal{A}^*), and $V \cdot W$ (the concatenation of V and W) are star-free sets over \mathcal{A} . A set is star-free over \mathcal{A} only by implication from the preceding clauses.

A regular set $V \subseteq \mathcal{A}^*$ is a *non-counting* regular set over \mathcal{A} iff there is an integer $l \geq 0$ such that for all $x, y, z \in \mathcal{A}^*$,

$$xy^l z \in V \quad \text{iff} \quad xy^{l+1} z \in V.$$

It is proved in [7] that star-free sets are exactly the non-counting regular sets.

Let $FT_{\Sigma,k} = \{s \mid s \text{ is a prefix of some } t \in \text{FIFO}_{\Sigma,k}, \text{ i.e., } t \text{ is a model of the FIFO buffer}\}$. Similarly let $LT_{\Sigma,k}$, $UT_{\Sigma,k}$, respectively, be the sets of prefixes of sequences in $\text{LIFO}_{\Sigma,k}$, $\text{UNOR}_{\Sigma,k}$. We first prove that $FT_{\Sigma,k}$, $LT_{\Sigma,k}$, $UT_{\Sigma,k}$ are star-free regular sets by proving that they are non-counting regular sets. Remember that \mathcal{M} is the set of states and is finite.

LEMMA 4.4. *For $k \neq \infty$, $FT_{\Sigma,k}$, $LT_{\Sigma,k}$, $UT_{\Sigma,k}$ are non-counting regular sets over \mathcal{M} .*

Proof. Clearly all the above sets are regular sets since the corresponding buffers have a finite number of possible states. First we prove that $FT_{\Sigma,k}$ is a non-counting set. We want to prove that for all $x, y, z \in \mathcal{M}^*$,

$$xy^{k+1}z \in FT_{\Sigma,k} \quad \text{iff} \quad xy^{k+2}z \in FT_{\Sigma,k}.$$

For any $\sigma \in \Sigma$ if y contains more reads of σ than writes of σ or vice versa, then neither $xy^{k+1}z$ nor $xy^{k+2}z$ is in $FT_{\Sigma,k}$. Assume y has equal number of reads and writes for each $\sigma \in \Sigma$. Let $xy^{k+1}z \in FT_{\Sigma,k}$, n be the number of messages in the buffer after x , and α, β , respectively, be the sequences of

messages written and read from the buffer in y . Let $|\alpha| = |\beta| = m$. The non-trivial case occurs when $m, n > 0$. Let $p = n \bmod m$, $\alpha = \alpha_1 \alpha_2$, $\beta = \beta_2 \beta_1$, where $|\beta_2| = |\alpha_2| = p$. Clearly $n \leq k$. We see that all the messages in the buffer after xy^k are those written in y^k and it is easily seen that after xy^k the buffer contains the sequence $(\alpha_2 \alpha^q)$ for some $q \geq 0$. Since xy^{k+1} is a valid sequence of operations we see that $\alpha_1 = \beta_1$, $\alpha_2 = \beta_2$. Thus the buffer contains the same sequence after xy^k and after xy^{k+1} . From this it follows that xy^{k+2} is a valid sequence of operations and the contents of the buffer after xy^{k+2} is same as that after xy^{k+1} . Hence $xy^{k+2}z \in FT_{\Sigma,k}$. The same argument proves that if $xy^{k+2}z \in FT_{\Sigma,k}$ then $xy^{k+1}z \in FT_{\Sigma,k}$.

As above, if y contains unequal number of reads and writes for any $\sigma \in \Sigma$, then neither $xy^{k+1}z$ nor $xy^{k+2}z$ is in $LT_{\Sigma,k}$ or in $UT_{\Sigma,k}$. So assume y has equal number of reads and writes for each $\sigma \in \Sigma$. Let $xy^{k+1}z \in LT_{\Sigma,k}$. Since $k \geq 1$, it is seen that the messages read in y correspond with the messages written in y . Hence $xy^{k+2}z \in LT_{\Sigma,k}$. Similarly it is seen that if $xy^{k+2}z \in LT_{\Sigma,k}$ then $xy^{k+1}z \in LT_{\Sigma,k}$. Since y has equal number of reads and writes for every $\sigma \in \Sigma$, it easily follows that $xy^{k+1}z \in UT_{\Sigma,k}$ iff $xy^{k+2}z \in UT_{\Sigma,k}$. ■

From the results of [7] and the above lemma it follows that $FT_{\Sigma,k}$, $LT_{\Sigma,k}$, $UT_{\Sigma,k}$ are star-free regular sets.

Let L be the first-order languages of $(N, <)$ with monadic predicates, where N is the set of non-negative integers and $<$ is the usual less than relation. This language is shown to be expressively equivalent to PTL in [10]. Hence it is enough if we show that all bounded buffers are characterizable in L .

We assume that the predicate symbols in L are same as the propositional symbols in PTL, i.e., the elements of \mathcal{P} . Now we want to prove that for every star-free regular set, there is a formula in L which defines the star-free regular set. From this it follows that $FT_{\Sigma,k}$, $LT_{\Sigma,k}$, $UT_{\Sigma,k}$ are definable in L . Using this result we will prove that all bounded buffers are characterizable in L .

Let $f(x, y)$ be a formula in L with free integer variables x, y . Now we associate a language with f . Let M be an interpretation of x, y and the predicate symbols, that is, M is a function which associates non-negative integers with x, y and a subset of N with each predicate symbol. Let $m = M(x)$, $n = M(y)$. For any $i \geq 0$, let $\phi_i \in \mathcal{M}$ be the unique state such that for all $P_j \in \mathcal{P}$, $\phi_i(P_j) = \text{True}$ iff $i \in M(P_j)$. With M we associate a $t(M) \in \mathcal{M}^*$ defined as follows:

$t(M)$ is undefined if $n < m$,

$t(M)$ is the empty word if $n = m$,

For $m < n$, $t(M) = (\phi_m, \phi_{m+1}, \dots, \phi_{n-1})$.

Define $\mathcal{L}(f) = \{t(M) \mid M \models f\}$. $\mathcal{L}(f)$ is called the language associated with f . It can easily be shown that the successor function given by $y = x + 1$ is definable in L .

LEMMA 4.5. *For any star-free regular set E over \mathcal{M} there is a formula f_E in L such that $E = \mathcal{L}(f_E)$.*

Proof. By structural induction on E :

if E is a singleton say $\{t_0\}$ then

$$f_E = (y = x + 1) \wedge \left(\bigwedge_i P_i \right) \wedge \left(\bigwedge_j \sim P_j \right),$$

where the first conjunct is over all P_i such that $t_0(P_i) = \text{True}$, and the second conjunct is over all P_j such that $t_0(P_j) = \text{False}$;

If $E = V \cdot W$ then

$$f_E = \exists z (f_V(x, z) \wedge f_W(z, y));$$

If $E = V \cup W$ then $f_E = (f_V \vee f_W)$;

If $E = \bar{V}$ then $f_E = \sim(f_V)$.

It can easily be proved by induction that $E = \mathcal{L}(f_E)$. ■

THEOREM 4.6. *All bounded buffers with or without the liveness property are characterizable in PTL.*

Proof. First we prove that all bounded buffers without the liveness property are characterizable in PTL. It is proved in [7] that the non-counting regular sets are exactly the star-free sets. From this and Lemma 4.4 it follows that $FT_{\Sigma, k}$, $LT_{\Sigma, k}$, $UT_{\Sigma, k}$ are star-free sets. Using Lemma 4.5, we see that there is a formula $f(x, y)$ in L such that $\mathcal{L}(f) = FT_{\Sigma, k}$. Let $g(x) = \forall y f(x, y)$, and M be any interpretation of predicate symbols in L and the variable x such that $M(x) = 0$. Let $t_M = (t_0, t_1, \dots)$ be the model of PTL such that $t_i(P_j) = \text{True}$ iff $i \in M(P_j)$. It is easily seen that $M \models g$ iff $t_M \in \text{FIFO}_{\Sigma, k}$. This is because g asserts that every finite prefix of t_M is in $FT_{\Sigma, k}$. It is proved in [10] that L and PTL are expressively equivalent. From this it follows that there is a PTL formula, say h , such that for all interpretations M , such that $M(x) = 0$, $M \models g$ iff $(t_M, 0) \models h$. Hence for any $t \in \mathcal{M}^\omega$, $(t, 0) \models h$ iff $t \in \text{FIFO}_{\Sigma, k}$. Thus h characterizes bounded FIFO buffers of size k over the message alphabet Σ . Similarly it follows that bounded LIFO and unordered buffers are characterizable in PTL.

It can easily be seen that Lemma 4.4 also holds for bounded buffers with the liveness property. From this it follows that bounded buffers with the

liveness property are characterizable in PTL. The details are left to the reader. ■

In the Appendix we give a constructive method for obtaining a characterization of bounded buffers. This method also uses the translation from L to PTL. However it is known that this translation may cause a non-elementary blow up in the size of the formula. Thus the length of PTL formulae obtained may be non-elementary in the size of the buffer. Obtaining a characterization directly in PTL (i.e., without going through L) as given in [9], gives formulae of length $\exp(p(k))$, where p is a fixed polynomial in k .

Using first order temporal logic, a uniform characterization of bounded buffers can be given, that is, we can give a formula in first-order temporal logic characterizing a bounded buffer of size k , which is independent of the message alphabet Σ .

4.3. Using Auxiliary Propositions in PTL

Below we show that by introducing auxiliary propositions we can characterize bounded message buffers more elegantly and succinctly. The syntax of the well-formed formulae in this new logic is exactly the same as in PTL, except that some propositions are designated as auxiliary propositions and are not interpreted. Thus let $\mathcal{P} = \mathcal{P}_A \cup \mathcal{P}_1$ be the set of atomic propositions, where \mathcal{P}_A is the set of auxiliary propositions. As usual, an interpretation is a pair $\langle s', i \rangle$, where s' is an ω -sequence of states (s'_0, s'_1, \dots) , each state being a mapping from \mathcal{P}_1 into $\{\text{True}, \text{False}\}$, and $i \geq 0$ designates the present state. We define truth of a formula f in an interpretation $\langle s', i \rangle$ (denoted by $s', i \approx f$) as follows:

$s', i \approx f$ iff there exists a sequence $s = (s_0, s_1, \dots)$ such that $s, i \models f$, where for all $j \geq 0$, $s_j: \mathcal{P} \rightarrow \{\text{True}, \text{False}\}$ is an extension of s'_j .

In this new logic we can characterize bounded buffers more concisely. We show this for an FIFO buffer of size 2. An FIFO buffer of size 2 can be considered as two FIFO buffers each of size 1 in tandem as shown in Fig. 1.




FIGURE 1

External writes come into the left buffer while external reads are from the right buffer. Whenever the left buffer is full and the right buffer is empty the message in the left buffer is internally read and is written into the right buffer. We consider this internal reading and writing to be occurring

simultaneously and capture it by the auxiliary propositions I_σ for $\sigma \in \Sigma$. Let $fb_1(\mathbf{W}_\sigma, \mathbf{R}_\sigma)$ be the formula characterizing a buffer of size 1, where $\mathbf{W}_\sigma, \mathbf{R}_\sigma$ indicate vectors of propositions. The sequence of operations on the left buffer is characterized by $fb_1(\mathbf{W}_\sigma, \mathbf{I}_\sigma)$, and the sequence of operations on the right buffer is characterized by $fb_1(\mathbf{I}_\sigma, \mathbf{R}_\sigma)$. Let

$$fb_2 = fb_1(\mathbf{W}_\sigma, \mathbf{I}_\sigma) \wedge fb_1(\mathbf{I}_\sigma, \mathbf{R}_\sigma),$$

where the propositions in \mathbf{I}_σ are the auxiliary propositions.

LEMMA 4.7. $s, 0 \models fb_2$ iff $s \in \text{FIFO}_{\Sigma,2}$.

It is easily seen how we can extend the above approach to characterize bounded FIFO buffers of size k . For characterizing LIFO and unordered buffers we take a different approach. We use auxiliary propositions P_0, P_1, \dots, P_k . We will assert that P_j is true at an instance i iff the buffer has j messages before the operation of the i th instance:

$$h' = \mathbf{G} \left[\bigwedge_{0 \leq l < m \leq k} \sim (P_l \wedge P_m) \wedge \bigwedge_{0 \leq l < k} ((P_l \wedge W) \supset \mathbf{X}P_{l+1}) \wedge \bigwedge_{0 < l \leq k} ((P_l \wedge R) \supset \mathbf{X}P_{l-1}) \wedge (P_0 \supset \sim R) \wedge (P_k \supset \sim W) \right] \wedge P_0 \wedge g,$$

where $g = \mathbf{G} [\sim (R \vee W) \supset (\bigwedge_{0 \leq l \leq k} (P_l \leftrightarrow \mathbf{X}P_l))]$.

The first clause asserts that no more than one P_l is true at any instance, the second clause asserts that if P_l is true at an instance and the operation is a write operation then at the next instance P_{l+1} is true, the third clause asserts similar property for read operation, the last two clauses assert that there are no writes on a full buffer and no reads on an empty buffer. g asserts that if both R and W do not hold in the present state then the set of all P_l true in the present and next states are the same. Let

$$lb_k = I \wedge h' \wedge \mathbf{G} \left(\bigwedge_{0 < l \leq k} \left(P_l \supset \bigwedge_{\sigma} [R_\sigma \supset (\sim P_{l-1} \mathbf{S}(W_\sigma \wedge P_{l-1}))] \right) \right).$$

The last clause asserts that the message read at any instance when the buffer has l messages is same as the message written at the last instance when the buffer has $l-1$ messages. The following theorem can be easily proved.

THEOREM 4.8. $t, 0 \models lb_k$ iff $t \in LS_{\Sigma,k}$.

Similarly we can obtain a formula for unordered buffers. Note that the lengths of the formulae obtained above are only linear in k .

4.4. Characterizing Unbounded Buffers

Let \mathcal{P} be a finite set of atomic propositions and $s = (s_0, s_1, \dots)$ be an infinite sequence of states where each state is a mapping from \mathcal{P} into $\{\text{True}, \text{False}\}$. Let f be a formula in PTL and $SF(f)$ denote the set of subformulae of f . It is easily seen that $\text{card}(SF(f)) \leq \text{length}(f)$. For $i \geq 0$ let $[i]_{s,f} = \{g \in SF(f) \mid s, i \models g\}$.

LEMMA 4.9. *Let $0 \leq i \leq j$ be such that $[i]_{s,f} = [j]_{s,f}$. Then $s, 0 \models f$ iff $s', 0 \models f$, where $s' = (s_0, s_1, \dots, s_i, s_{j+1}, s_{j+2}, \dots)$.*

THEOREM 4.10. *Unbounded message buffers (unordered, FIFO, or LIFO) cannot be characterized in PTL.*

Proof. All the above unbounded buffers are identical in the case when the message alphabet is a singleton. We prove that there is no formula in PTL which characterizes an unbounded buffer over a message alphabet which is a singleton. Assume to the contrary f is a such formula. Let s be a model in which n messages are written in the first n states and n messages are read in the next n states. Choose $n > 2^{\text{length}(f)}$. Now it is easily seen that there are two integers i, j such that $0 \leq i < j < n$ such that $[i]_{s,f} = [j]_{s,f}$. Now by applying Lemma 4.9 we see that there is a sequence s' which has fewer than n writes followed by n reads such that $(s', 0) \models f$. Clearly s' is not a model of an unbounded buffer. This is a contradiction. ■

Indeed, we can show that there is no uniform characterization of unbounded buffers in first-order temporal logic. It can easily be seen that there are *partially interpreted* first-order temporal logics in which unbounded buffers can be characterized. For example, we can characterize unbounded FIFO buffers in a first order temporal logic that uses history variables ranging over sequences of messages, and the *prefix* relation (\leq) among the sequences.

5. AXIOMATIZING MESSAGE BUFFERS

Axiomatization of message buffers in PTL is a weaker notion than expressiveness. We show that in general unbounded FIFO buffers are not axiomatizable, while unbounded LIFO buffers and unbounded unordered buffers are axiomatizable even though they are not characterizable in PTL.

THEOREM 5.1. *Bounded FIFO, LIFO, and unordered buffers over any finite alphabet Σ are axiomatizable in PTL.*

Proof. Let fb_k be the formula in PTL characterizing the FIFO buffer of size k over a finite alphabet Σ . Consider a consistent and complete axiomatization for PTL as given in [14]. Let A be the set of axioms and I be the set of inference rules in this system. Then the system with $A \cup \{fb_k\}$ as the set of axioms and I as the set of inference rules forms a complete axiomatization for FIFO buffers of size k over Σ . Similarly, we can give an axiomatization for bounded LIFO and unordered buffers over a finite alphabet. ■

A finite state automaton M on infinite strings over an alphabet A is a 4-tuple (Q, δ, s_0, F) , where Q is the set of states, δ is the transition function, i.e., $\delta: Q \times A \rightarrow 2^Q$, $s_0 \in Q$ is the initial state, $F \subseteq Q$. The states in F are called final states. A run γ of M on an input $t \in A^\omega$, is an ω -sequence of states $(\gamma_0, \gamma_1, \dots)$ such that $\gamma_0 = s_0$, for all $i \geq 0$, $\gamma_{i+1} \in \delta(\gamma_i, t_i)$; γ is an accepting run iff some final state appears infinitely often in γ . An input t is said to be accepted by M iff there is an accepting run of M on t . For every PTL formula f there is a finite state automaton M_f over the alphabet \mathcal{M} such that $(t, 0) \models f$ iff t is accepted by M_f , where \mathcal{M} is the set of all mappings from \mathcal{P} into $\{\text{True}, \text{False}\}$ and \mathcal{P} is the set of all propositions appearing in f . Indeed, we can obtain such an automaton with number of states exponential in the length of f . A procedure to obtain such an automaton is given by [9]. We will be using this automaton frequently in our proofs.

We assume familiarity with the hierarchy notation of [11]. Σ_1^0 is the class of recursively enumerable sets and Π_1^0 is the class of sets which are complements of recursively enumerable sets. The class Σ_1^1 and its complement Π_1^1 reside low in the analytical hierarchy [11].

THEOREM 5.2. *For any Σ such that $\text{card}(\Sigma) \geq 2$, the theory of unbounded FIFO buffers over Σ is Π_1^1 -complete, the theory of unbounded FIFO buffers with liveness property is Π_1^0 -complete, and hence neither of the theories is axiomatizable.*

Proof. First we prove below that for $\Sigma = \{0, 1\}$ the theory of unbounded FIFO buffers is Π_1^1 -complete. From this result it automatically follows that the above theory is not axiomatizable.

We first prove that the set of PTL formulae satisfiable over some model of an unbounded FIFO buffer over $\{0, 1\}$ is Σ_1^1 -complete. We consider deterministic Turing machines on infinite strings with one read-only infinite input tape and one work tape. The set of encodings of all TMs on infinite strings which accept at least one input is known to be Σ_1^1 -complete. We will reduce this set to the set of PTL formulae which are true on some model of an unbounded FIFO buffer over Σ . A Turing machine on infinite strings works exactly like an ordinary Turing machine, but it takes infinite input

strings and it never halts. It accepts an input string by going through a final state infinitely often. Let $M = (\mathcal{A}, Q, \delta, H)$ be such a Turing machine, where \mathcal{A} is the alphabet (including both the input alphabet and the tape alphabet), Q is the set of states, $\delta: Q \times \mathcal{A} \times \mathcal{A} \rightarrow Q \times \mathcal{A} \times \{\text{left}, \text{right}\}$, H is the set of final states. After each step the input head of M moves right by one cell. If $\delta(q, \sigma_1, \sigma_2) = (q', \sigma'_2, \text{left})$, then whenever M is in state q and sees the symbol σ_1, σ_2 on the input and work tapes, respectively, then M moves into state q' , writes σ'_2 on the work tape and moves its head left, and it moves its input head right by one cell. We show below that given the encoding of M we can recursively obtain a formula f_M in PTL such that f_M is satisfiable on an unbounded FIFO buffer over $\{0, 1\}$ iff M accepts at least one input.

Let $C = (Q \times \mathcal{A}) \cup \mathcal{A}$. A partial id of M , is a sequence of values from C , containing exactly one symbol from $(Q \times \mathcal{A})$. A partial id denotes the contents of the work tape and the head position on the work tape and the state of finite control in the usual way. The formula f_M will be satisfied iff there is a ω -sequence of partial ids such that each succeeding partial id is obtained from the previous partial id by one move of M reading some input character, and there are infinitely many partial ids in this sequence containing a symbol of the form (q_f, σ) , where q_f is a final state. We call such a sequence an accepting sequence. Any such sequence denotes an accepting computation of M , and for every accepting computation of M there is such a sequence.

We fix a unary encoding of symbols from C using the character $0 \in \Sigma$. An encoding of a partial id is a sequence of encodings of the symbols in it separated by the symbol 1. An encoding of a sequence of partial ids is the sequence of encodings of the partial ids separated by two consecutive 1's. The formula f_M can easily be constructed from the following description.

The initial buffer history consists of a sequence of writes which places the encoding of the initial partial id followed by two consecutive 1's. After writing of the initial id, reading and writing of symbols occurs alternately (thus whenever a symbol is read, it is the symbol of the previous id). Each symbol written into the buffer is the value of the symbol in the new id assuming some input symbol on the input tape (f_M can express this because the value of a symbol in a new id depends only on the contents of that cell and its neighbors in the previous id, and the assumed value of input character). Two consecutive 1's are written at the end of each id.

Finally, f_M also asserts that there are infinitely many places where a symbol of the form (q_f, σ) is written into the buffer for some final state q_f . It is clearly seen that f_M is satisfiable on a model of an unbounded FIFO buffer over $\{0, 1\}$ iff M accepts at least one input.

Now we give a reduction in the other direction. Let f be a PTL formula and M_f be the finite state automaton on infinite strings corresponding to f .

From M_f we give a Turing machine M which operates as follows. M takes each symbol in its input to be an encoding of a function assigning truth values to the set of atomic propositions. M simulates M_f on the input, and at the same time it makes sure that the values of the propositions R_0, R_1, W_0, W_1 denote a valid FIFO buffer behavior. M accepts an input iff M_f accepts it and the input sequence denotes a valid FIFO buffer behavior. It is easily seen that M accepts at least one input iff f is satisfiable at the beginning of a model of an unbounded FIFO buffer over $\{0, 1\}$.

It can easily be shown that the set of encodings of TMs on infinite strings that accept at least one input is Σ_1^1 -complete. Hence the set of formulae in PTL that are satisfiable on a model of an unbounded FIFO buffer is Σ_1^1 -complete. From this it follows that the set of formulae, not satisfiable on any model of an unbounded FIFO buffer over $\{0, 1\}$ is Π_1^1 -complete. Hence the set of valid formulae is Π_1^1 -complete.

Now we prove that the set of PTL formulae satisfiable on a model of an FIFO buffer with liveness property is Σ_1^0 -complete. From this it follows that the set of PTL formulae valid on all models of FIFO buffers with liveness property is Π_1^0 -complete.

Given a Turing machine M on *finite* strings, analogous to before we can obtain a PTL formula f_M such that M accepts at least one input iff f_M is satisfiable on a model of an FIFO buffer with liveness property. Now we show that the set of PTL formulae satisfiable on an FIFO buffer with liveness property is in Σ_1^0 . Let f be satisfiable on a model of an FIFO buffer with liveness property and M_f be the automaton associated with f . Let $t \in \mathcal{M}^\omega$ be such a model. Let $\gamma = (\gamma_0, \gamma_1, \dots)$ be an accepting run of M_f on t . There are infinitely many values of $i \geq 0$, such that the buffer is empty before t_i in t . From this it follows that there are two instance i, j such that (1) $i < j$ and the buffer is empty at these instances in t , (2) $\gamma_i = \gamma_j$, and (3) there exists a k such that $i \leq k < j$ and γ_k is a final state of M_f . Now let $t' = \alpha \cdot \beta^\omega$, where $\alpha = (t_0, \dots, t_{i-1})$, $\beta = (t_i, t_{i+1}, \dots, t_{j-1})$. It is easily seen that M_f accepts t' , and hence f is true at the beginning of t' . Now we can easily give a Turing machine M on finite strings which takes a PTL formula f and checks if f is satisfiable on a model of an FIFO buffer with the liveness property. M guesses α, β and verifies that M_f accepts t' as given above. M halts iff M_f accepts t' . The details of M are left to the reader. ■

THEOREM 5.3. *The theory of unbounded LIFO buffers over a finite alphabet is decidable. Similarly the theory of LIFO buffers with the liveness property is decidable.*

Proof. Let f be a formula in PTL and M_f be the automaton associated with f as defined previously. From M_f we can obtain a push down automaton P_f operating on infinite strings. P_f uses its stack to make sure that the sequence of read/write operations represented by the input string is

a legal series of read/write operations on the buffer, while at the same time the finite state control of P_f makes state transitions exactly as M_f . P_f accepts an infinite string iff its finite state control goes through any of a set of final states infinitely often. P_f accepts an input t iff $t \in LS_{\Sigma, \infty}$ and $t, 0 \models f$. Thus f is satisfiable on a $t \in LS_{\Sigma, \infty}$ iff P_f accepts some input.

We now need to show that the question of whether P_f accepts some input is decidable. Assume that t is an input accepted by P_f . Consider an accepting computation C of P_f on the input. Let $nb(l)$, $r(l)$, respectively, be the number of messages in the stack and the state of the finite state control of P_f just before the l th step in C . It can easily be seen that there exists an infinite sequence of integers $i_0 < i_1 < i_2 < \dots < i_j < \dots$, such that for all $j \geq 0$, for all $m \geq i_j$, $nb(m) \geq nb(i_j)$, that is, from the i_j th step onwards the height of the stack never drops below $nb(i_j)$. Let l, m be integers such that $l < m$, $r(i_l) = r(i_m) = q$ and there exists a p such that $i_l \leq p \leq i_m$ and $r(p)$ is a final state of P_f . Clearly l, m exist since P_f goes through a final state infinitely often in the computation C . Note that P_f reads one input symbol on each step. Let $u = i_l$, $v = i_m$ and $\alpha = (t_0, t_1, \dots, t_{u-1})$, $\beta = (t_u, t_{u+1}, \dots, t_{v-1})$. We can also consider P_f as a push down automaton on finite strings.

Claim. P_f accepts some infinite string iff there exists a state q such that

- (a) there is a finite string α and there is a computation of P_f on input α reaching state q ; and
- (b) there is a finite string β and there is a computation of P_f starting with initial state q and empty stack and reaching q at the end but going through a final state during this computation.

Proof of claim. Assume P_f accepts an infinite input string. Then q, α, β as defined before satisfy (a) and (b). Assume (a) and (b) hold. Then it is easily seen that P_f accepts the infinite input $\alpha \cdot \beta^\omega$. ■

Now it is easily seen how to decide if P_f accepts at least one infinite input. That is check if there is a state q for which (a) and (b) of the above claim hold: (a) is the problem of deciding if a push down automaton on finite strings accepts at least one input which is known to be decidable; (b) can also be posed as such a problem and so is decidable. The above argument also shows that the theory of unbounded LIFO buffers with liveness property is decidable. ■

We will show that the problem of deciding if a given PTL formula is satisfiable on a model of an unbounded unordered buffer is reducible to certain problems in vector addition systems. A *vector addition system with states of dimension k* is a triple $G = (V, E, L)$, where (V, E) is a directed graph and $L: E \rightarrow N^k$, where N is the set of integers. A configuration is a pair (s, a) , where $s \in V$, $a \in N^k$. A path is a finite or infinite sequence of

length ≥ 2 of configurations c_0, c_1, \dots, c_i , where $c_i = (s_i, a_i)$ and for all $i \geq 0$ $(s_i, s_{i+1}) \in E$, $a_{i+1} = a_i + L(s_i, s_{i+1})$. Let $A \subseteq \{1, \dots, k\}$. A path is said to be positive with respect to A if for all $i \geq 0$, for all $j \in A$ $(a_i)_j \geq 0$. We say that a path is positive if it is positive with respect to $\{1, \dots, k\}$. We say that a configuration d is positively reachable with respect to A from the configuration c if there is a positive path with respect to A with initial configuration c and final configuration d ; and if $A = \{1, \dots, k\}$ then we simply say that d is positively reachable from c .

Let f be a PTL formula and M_f be the automaton associated with f . Let $G_f = (V, E, L)$ be the vector addition system associated with f defined as follows: V is the set of states of M_f and $(s_1, s_2) \in E$ iff there is an α (α is a function that assigns truth values to propositions) such that there is transition in M_f from s_1 to s_2 on input α , and $L(s_1, s_2) = (a_1, a_2, \dots, a_k)$, where for all i , $1 \leq i \leq k$,

$$\begin{aligned} a_i &= 1 && \text{if } \alpha(W_{\sigma_i}) = \text{True}, \\ a_i &= -1 && \text{if } \alpha(R_{\sigma_i}) = \text{True}, \\ a_i &= 0 && \text{otherwise.} \end{aligned}$$

It is easily seen that f is satisfiable on a model of an unordered buffer iff there exists an infinite positive path in G_f with initial configuration $(q_I, \vec{0})$ and containing infinitely many configurations of the form (q_f, a) , where q_I is the initial state and q_f is a final state in M_f .

Let $M = N \cup \{\omega\}$ and $<$ be the extension of "less than" relation to elements in M so that for each $i \in N$, $i < \omega$. Also we extend the usual $+$ operation so that for any $i \in M$, $i + \omega = \omega$. Let $a, b \in M^k$. Then $a \leq b$ iff $a_i \leq b_i$ for each i such that $1 \leq i \leq k$. Let c be any configuration. We define a labelled tree T_c as in [6]. Each node x in T_c is labelled with $l(x) \in M^k$ and is recursively defined as follows:

The root r of T_c is labelled with c , that is, $l(r) = c$. Let y be a node in T_c with $l(y) = (s, a)$:

(a) If there is a proper ancestor x of y in T_c such that $l(y) = l(x)$ then y is a leaf, that is, y does not have any children; otherwise

(b) For each $(s, s') \in E$ such that $L(s, s') + a = \gamma \geq \vec{0}$, there is a son z with $l(z) = (s', b)$, where b_i is given as follows:

If there is an ancestor x of y such that $l(x) = (s', d)$, where $d \leq \gamma$ and $d_i < \gamma_i$ then $b_i = \omega$, otherwise $b_i = \gamma_i$.

In [6] it is proved that T_c is finite.

LEMMA 5.4. *The following are equivalent:*

(a) *There is an infinite positive path in G with initial configuration c*

and containing infinitely many configurations of the form (q, d) , where q is a fixed state in G .

(b) There is a node in T_c with label (q, a) satisfying the following property:

Let A be the set of all i such that $a_i \neq \omega$ and $a' \in N^k$ be such that for all $i \in A$, $a'_i = a_i$ and for all i not in A , $a'_i = 0$. Then there exists $b \in N^k$ such that $a' \leq b$ and (q, b) is positively reachable with respect to A from (q, a') .

Proof. (b \Rightarrow a) Let q, a, b, A, a' be as given in (b). Let e_0, e_1, \dots, e_p be a positive path with respect to A , with $e_0 = (q, a')$, $e_p = (q, b)$. Let $e_i = (s_i, d_i)$, and Δ be the lower bound for all $(d_i)_j$ for $0 \leq i \leq p$ and for all j not in A . Since (q, a) is a node in T_c with $a_j = \omega$ for all j not in A , it can easily be shown that there is a configuration (q, f) such that $f_j = a_j$ for all $j \in A$, and $f_j \geq \Delta$ for all j not in A , and (q, f) is positively reachable from c , i.e., there is a positive path t with $c, (q, f)$ being the initial and final configurations, respectively. Let g be the sequence of states corresponding to t , and h be the sequence of states corresponding to e_1, e_2, \dots, e_p . Now consider the infinite sequence of states $g \cdot h^\omega$. It is clearly seen that this sequence gives us an infinite positive path satisfying the property given in (a).

(a \Rightarrow b) Let $C = c_0, c_1, \dots$, be an infinite positive path satisfying the property given in (a). Let $c_i = (s_i, a_i)$ for all $i \geq 0$. We define an infinite sequence $d_0 d_1 \dots d_i \dots$, as follows:

For all $i \geq 0$ $d_i = (s_i, b_i)$, where $b_i \in M^k$ is defined as follows: $b_0 = a_0$. For $i > 0$ we inductively define $b_i, \gamma_i \in M^k$ as follows:

For any j , $(\gamma_i)_j = \omega$ if $(b_{i-1})_j = \omega$, otherwise $(\gamma_i)_j = (a_i)_j$. For all j , $1 \leq j \leq k$, $(b_i)_j = \omega$ if there is an $l < i$ with $s_l = s_i$, $b_l \leq \gamma_i$, and $(b_l)_j < (\gamma_i)_j$; otherwise $(b_i)_j = (\gamma_i)_j$.

It is clearly seen that for each $i \geq 0$, d_i is the label of a node in T_c . Now the above sequence consists of an infinite subsequence $d_{j_0} d_{j_1} \dots$, such that for all $l \geq 0$, $s_{j_l} = q$. From this subsequence pick up an infinite subsequence with non-decreasing first coordinates of b_i . By repeatedly doing this for all coordinates, it can easily be shown that there is an infinite subsequence d_{i_0}, d_{i_1}, \dots , such that for all $l \geq 0$, $s_{i_l} = q$, $b_{i_l} \leq b_{i_{l+1}}$. But from the way we defined d_i , it has to be the case that there is an n , with the property that for all $l \geq n$, $b_{i_l} = b_{i_{l+1}}$. Since C is a positive path there exists l, m such that $l < m$ and $a_{i_l} \leq a_{i_m}$. Let A be the set of all j such that the j th coordinate of b_{i_l} is not ω . Clearly for all $j \in A$, the j th coordinate of $a_{i_l}, a_{i_m}, b_{i_l}, b_{i_m}$ are all equal.

Let $\alpha = b_{i_l}$ and $\beta, \gamma \in N^k$ be such that $\beta_j = \gamma_j = \alpha_j$ for all $j \in A$, and for all j not in A , $\beta_j = 0, \gamma_j = (j\text{th coordinate of } a_{i_m} - j\text{th coordinate of } a_{i_l})$. Now it is clearly seen that the configuration (q, γ) is positively reachable with respect to A from the configuration (q, β) . Also it is clear that there is a node in T_c with label (q, α) . Hence (b) is true. ■

Lemma 5.5 follows directly from the results of [6].

LEMMA 5.5. *Given a configuration $c = (q, a)$ it is decidable if there is a configuration $d = (q, b)$ with $a \leq b$ such that d is positively reachable from c in vector addition system G .*

The proof of the above lemma can easily be extended to the following lemma.

LEMMA 5.6. *Given $c = (q, a)$ and a set $A \subseteq \{1, 2, \dots, k\}$ it is decidable if there exists a configuration $d = (q, b)$ such that $a \leq b$ and d is positively reachable with respect to A from c . ■*

THEOREM 5.7. *The theory of unbounded unordered buffers over any alphabet Σ is decidable.*

Proof. It is enough if we show that the satisfiability problem is decidable. Let f be a PTL formula and M_f be the automaton associated with f , and G_f be the vector addition system with states of dimension k associated with f . Let $c = (s_0, \vec{0})$, where s_0 is the initial state in M_f . Then f is satisfiable iff there exists an infinite positive path in G_f with initial configuration c and containing infinitely many configurations of the form (q, d) , where q is a final state in M_f . Now to check this condition we use Lemma 5.4. We construct T_c and for each pair (q, a) , where q is a final state and some node in T_c is labelled with this pair, we do the following:

Let A be the set of all i such that $a_i \neq \omega$ and $a' \in N^k$ be such that for all $i \in A$, $a'_i = a_i$, and for all i not in A , $a'_i = 0$. Verify if there is some $b \geq a'$ such that (q, b) is positively reachable with respect to A from (q, a') . This can be done due to Lemma 5.6.

Now from Lemma 5.4, f is satisfiable on a model of an unbounded unordered buffer iff there exists a pair (q, a) which satisfies the above condition. Thus satisfiability is decidable and hence the theorem follows. ■

The following problem is known as the *reachability problem* for vector addition systems with states: Given a vector addition system G with states, and two configurations a, b of G , is b positively reachable from a in G ?

THEOREM 5.8. *The theory of unbounded unordered buffers with the liveness property over a finite alphabet Σ is axiomatizable (and also decidable) iff the reachability problem for vector addition systems with states of dimension $\text{card}(\Sigma)$ is decidable.*

Proof. First we want to show that the set of formulae satisfiable on a model of an unbounded unordered buffer with the liveness property is

recursively enumerable. To show this it is enough if we give a partial decision procedure to check if f is satisfiable on such a model. Let t be such a model at the beginning of which f is true, and let M_f be the automaton associated with f . Since the buffer becomes empty infinitely often in t , it easily follows that there is a model of an unbounded unordered buffer of the form $\alpha \cdot \beta^\omega$, where α, β are finite strings such that the buffer is empty after $\alpha \cdot \beta^i$ for all $i \geq 0$ and f is true at the beginning of the above model. Our partial decision procedure guesses α, β and verifies that they have the above properties. It is easily seen that we can easily verify if $\alpha \cdot \beta^\omega$ is accepted by M_f .

Now it easily follows that the set of valid PTL formulae in all models of unbounded unordered buffers with the liveness property is axiomatizable iff the set of satisfiable formulae over such models is decidable. We show that the set of satisfiable formulae is decidable iff reachability problem for vector addition systems with states is decidable.

First we reduce reachability problem to satisfiability problem. Let $G = (V, E, L)$ be a vector addition system with states of dimension k and it is required to determine if (t, b) is positively reachable from (s, a) . Let $\Sigma = \{\sigma_1, \sigma_2, \dots, \sigma_k\}$. We give a formula f such that there is a history $h \in \text{UNOR}_{\Sigma, \infty}$, in which the buffer becomes empty infinitely often and such that $(h, 0) \models f$ iff (t, b) is positively reachable from (s, a) in G . We use a proposition P_u for each $u \in V$. The formula f asserts the following:

- (i) For each $i, 1 \leq i \leq k$, initially a_i messages of value σ_i are written into the buffer; immediately after this P_s is true.
- (ii) The propositions P_u (for $u \in V$) are mutually exclusive. For $u \neq t$ if P_u is true at any instance i then the next proposition to be true in future at instance j will be P_v , where $(u, v) \in E$, and if $(c_1, c_2, \dots, c_k) = L(u, v)$ then between i and j , for all l such that $1 \leq l \leq k$ if c_l is positive (negative) then $|c_l|$ number of messages of value σ_l are written into (read from) the buffer.
- (iii) If P_t is true at any instance, either (ii) holds or the following condition is satisfied. Immediately after P_t is true, for all l such that $1 \leq l \leq k$, b_l messages of value σ_l are read from the buffer, and after this all propositions are false forever.
- (iv) There is a future instance from which point all propositions will be false forever.

It is easily seen how to obtain f . Thus if satisfiability problem is decidable then reachability problem is also decidable.

Now assume that the reachability problem is decidable. Let f be a PTL formula, and M_f, G_f be as defined before where $G_f = (V, E, L)$. Let q_i be the initial state in M_f . The following is easily seen:

There is a $t \in \text{UNOR}_{\Sigma, \infty}$ such that the buffer becomes empty infinitely

often in t and $(t, 0) \models f$ iff there is a $q \in V$ and a $q_f \in V$ which is a final state satisfying

- (i) $(q, \bar{0})$ is reachable from $(q_I, \bar{0})$ in G_f and,
- (ii) $(q, \bar{0})$ is reachable from $(q, \bar{0})$ by passing through q_f .

(ii) is not a direct reachability problem; however, we can put it as a reachability problem as follows: Introduce another copy of G_f , call it G'_f , and introduce a transition from q_f in G to q'_f in G'_f , which is labelled with $\bar{0}$. Now (ii) is satisfied in G_f iff $(q', \bar{0})$ is reachable from $(q, \bar{0})$ in the new vector addition system.

Since we assumed reachability is decidable, we can easily decide if there is a q satisfying (i) and (ii). ■

6. CONCLUSION

We have examined the possibility of using linear temporal logic to express the semantics of different message buffering systems. We have shown that it is possible to characterize bounded message buffers in PTL and that no such characterization is possible for unbounded buffers. We have also considered the possibility of axiomatization of the theory of various message buffer systems. We have shown that unbounded FIFO buffers are not axiomatizable in PTL. However it may be possible to give an axiomatization of these buffers in a logic weaker than PTL. This is still an open problem. We have also shown that unbounded LIFO and unordered buffers are axiomatizable in PTL. However to give a complete axiom system for these buffers in PTL is still an open problem. Some of these problems will be addressed in a future paper.

APPENDIX

Below we give a constructive method for characterizing bounded buffers.

A *Semi-automaton* A is a triple (Q_A, Δ_A, M_A) where Q_A is a finite set of states, Δ_A is a finite alphabet, and $M_A: Q_A \times \Delta_A \rightarrow Q_A$. Let $A = (Q_A, \Delta_A, M_A)$ and $B = (Q_B, \Delta_B, M_B)$, and $\delta: Q_A \times \Delta_A \rightarrow \Delta_B$. The *cascade product* of A, B with mapping δ is the semi-automaton $C = (Q_c, \Delta_c, M_c)$, where $Q_c = Q_A \times Q_B$, $\Delta_c = \Delta_A$ and for all $p \in Q_A$, $q \in Q_B$, $\sigma \in \Delta_c$, $M_c((p, q), \sigma) = (p', q')$, where $p' = M_A(p, \sigma)$, $q' = M_B(q, \delta(p, \sigma))$. The cascade product of three or more automata is defined by association to the left.

A *reset* is a semi-automaton $A = (Q, \Delta, M)$ where $Q = \{0, 1\}$, Δ is a disjoint union of 3 sets $\Delta_0, \Delta_1, \Delta_I$ such that for all $p \in Q$, $\sigma \in \Delta_0$, $M(p, \sigma) = 0$, for all $\sigma \in \Delta_1$, $M(p, \sigma) = 1$, and for all $\sigma \in \Delta_I$, $M(p, \sigma) = p$.

Let $A = (Q, A, M)$ be a semi-automaton and let M be extended in a natural way to the domain A^* . Now let $A_{p,q} = \{\sigma \in A^* \mid M(p, \sigma) = q\}$. It is proved in ([7, 8]) that if A is a cascade product of resets then $A_{p,q}$ is a star-free regular set. Indeed in [8] a constructive method is given to generate such a star-free regular set. We give a construction below to simulate bounded buffers by a cascade product of resets.

THEOREM. *Bounded FIFO, LIFO, and unordered buffers can be simulated by a cascade product of resets.*

Proof. Assume the message alphabet $\Sigma = \{0, 1\}$. It is easy to see how our construction can be extended to any finite Σ . Let k be the size of the buffer. We informally describe the construction. We use a cascade product of $2k$ resets. The number of resets in state 1 among the first k resets gives the number of messages in the buffer at any instance. The last k resets contain the contents of the buffer right justified. Number the resets left to right starting from 1. If the number of messages in the buffer is m , then the resets $2k - m + 1$ through $2k$ contain the buffer contents, the latest being contained in $(2k - m + 1)$ th reset and the oldest in $2k$ th reset.

Whenever a write operation occurs, then the first reset in state 0 among the first k resets is changed to state 1. Whenever a read operation occurs then the first reset in state 1 among the first k resets is changed to state 0. Overflow and underflow can be easily detected.

In addition to the above update, a write operation places the new message in $(2k - m)$ th reset if there are m messages already in the buffer.

In case of FIFO buffers, a read operation transfers the state of i th reset to $(i + 1)$ th reset for all i such that $2k - m + 1 \leq i < 2k$. In case of LIFO buffers decrementing of the counter in the first k resets is enough. In case of unordered buffers, a read operation on a message σ transfers the state of i th reset to $(i + 1)$ th reset for all i such that $2k - m + 1 \leq i < j$, where j is the first reset with state σ . The detailed designs of the resets are left to the reader. Finally, we can use another reset and handle error conditions. In this case if there is an error operation on the buffer then the $(2k + 1)$ th reset goes into state 1 indicating an error state. ■

Let A be the cascade product of resets as constructed in the above theorem. Let $\bar{0}$ be the state of A with all resets being in state 0. Let $B = \text{union of all } A_{p,q}$, where q is a non-error state of A and $p = \bar{0}$. Using the construction given in [8] we can obtain a star-free regular event corresponding to B , and using the translation given in Lemma 4.5 we can obtain a characterization of bounded buffers in the language L . From this and using the translation given in [10] we can obtain a characterization of bounded buffers in PTL.

ACKNOWLEDGMENT

The authors acknowledge the help of Professor Yuri Gurevich on an early version of Theorem 4.10.

RECEIVED: February 17, 1984; ACCEPTED: January 15, 1985

REFERENCES

1. SISTLA, A. P., CLARKE, E. M., FRANCEZ, N., AND GUREVICH, Y. (1982), Can message buffers be characterized in linear temporal logic? in "Proc. ACM Sympos. on Principles of Distributed Systems," Ottawa, Canada, August.
2. MANNA, Z., AND PNUELI, A. (1981), Verification of concurrent programs, in "The Correctness Problem in Computer Science," International Lecture Series in Computer Science, London.
3. OWICKI, S., AND LAMPORT, L. (1982), Proving liveness properties of concurrent programs, *TOPLAS* **4**, No. 3.
4. WOLPER, P. (1983), Temporal logic can be more expressive, in "Proc. 22nd IEEE Sympos. on Foundations of Computer Science, Nashville, Tenn. October, 1981; *Inform. and Control* **56**, 72-99.
5. KOSARAJU, S. R. (1982), Decidability of reachability in vector addition systems, in "Proc. 14th Annu. ACM Sympos. on Theory of Computing," May.
6. KARP, R. M., AND MILLER, R. E. (1969), Parallel program schemata, *J. Comput. System Sci.* **3**, No. 2, 147-195.
7. PAPERT, S., AND MCNAUGHTON, R. (1966), On topological events, in "Theory of Automata," University of Michigan Engineering Summer Conferences.
8. MEYER, A. R. (1969), A note on starfree events, *J. Assoc. Comput. Mach.* **16**, No. 2.
9. SISTLA, A. P., "Theoretical Issues in the Design and Verification of Distributed Systems, Ph.D. thesis, Harvard; Tech. Report CMU-CS-83-146, Carnegie-Mellon University.
10. GABBAY, D., PNUELI, A., SHEALAH, S., AND STAVI, J. (1969), Temporal analysis of fairness, "Seventh ACM Symposium on Principles of Programming Languages," Las Vegas, December.
11. ROGERS, H., (1967), "Theory of Recursive Functions and Effective Computability," McGraw-Hill, New York.
12. HAREL, D. (1984), Recurring dominoes: Making highly undecidable highly understandable, *Ann. Discrete Math.*
13. SISTLA, A. P., AND CLARKE, E. M. (1982), Complexity of propositional linear temporal logics, in "ACM Sympos. on Theory of Computing," San Francisco.
14. KAMP, H. (1968), "Tense Logic and Theory of Linear Order," Ph.D. Thesis, UCLA.