

Upper Bounds for Newton's Method on Monotone Polynomial Systems, and P-Time Model Checking of Probabilistic One-Counter Automata

ALISTAIR STEWART and KOUSHA ETESSAMI, University of Edinburgh
MIHALIS YANNAKAKIS, Columbia University

A central computational problem for analyzing and model checking various classes of infinite-state recursive probabilistic systems (including quasi-birth-death processes, multitype branching processes, stochastic context-free grammars, probabilistic pushdown automata and recursive Markov chains) is the computation of *termination probabilities*, and computing these probabilities in turn boils down to computing the *least fixed point* (LFP) solution of a corresponding *monotone polynomial system* (MPS) of equations, denoted $x = P(x)$.

It was shown in Etessami and Yannakakis [2009] that a decomposed variant of Newton's method converges monotonically to the LFP solution for any MPS that has a nonnegative solution. Subsequently, Esparza et al. [2010] obtained upper bounds on the convergence rate of Newton's method for certain classes of MPSs. More recently, better upper bounds have been obtained for special classes of MPSs [Etessami et al. 2010, 2012].

However, prior to this article, for arbitrary (not necessarily strongly connected) MPSs, no upper bounds at all were known on the convergence rate of Newton's method as a function of the encoding size $|P|$ of the input MPS, $x = P(x)$.

In this article, we provide worst-case upper bounds, as a function of both the input encoding size $|P|$, and $\epsilon > 0$, on the number of iterations required for decomposed Newton's method (*even with rounding*) to converge to within additive error $\epsilon > 0$ of q^* , for an *arbitrary* MPS with LFP solution q^* . Our upper bounds are essentially optimal in terms of several important parameters of the problem.

Using our upper bounds, and building on prior work, we obtain the first P-time algorithm (in the standard Turing model of computation) for quantitative model checking, to within arbitrary desired precision, of discrete-time QBDs and (equivalently) probabilistic 1-counter automata, with respect to any (fixed) ω -regular or LTL property.

Categories and Subject Descriptors: F.2.1 [Analysis of Algorithms and Problem Complexity]: Numerical Algorithms and Problems; G.1.5 [Numerical Analysis]: Roots of Nonlinear Equations; G.3 [Probability and Statistics]: Stochastic processes

General Terms: Algorithms, Theory

Additional Key Words and Phrases: Probability, model checking

ACM Reference Format:

Stewart, A., Etessami, K., and Yannakakis, M. 2015. Upper bounds for Newton's method on monotone polynomial systems, and p-time model checking of probabilistic one-counter Automata. *J. ACM* 62, 4, Article 30 (August 2015), 33 pages.

DOI: <http://dx.doi.org/10.1145/2789208>

Research was partially supported by the Royal Society, and by NSF grants CCF-1017955 and CCF-1320654. Authors' addresses: A. Stewart and K. Etessami, School of Informatics, University of Edinburgh; email: stewart.al@gmail.com; kousha@inf.ed.ac.uk; M. Yannakakis, Department of Computer Science, Columbia University; email: mihalis@cs.columbia.edu.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2015 ACM 0004-5411/2015/08-ART30 \$15.00

DOI: <http://dx.doi.org/10.1145/2789208>

1. INTRODUCTION

In recent years, there has been extensive work on the analysis of various classes of infinite-state recursive probabilistic systems, including recursive Markov chains, probabilistic pushdown systems, stochastic context-free grammars, multitype branching processes, quasi-birth-death processes and probabilistic 1-counter automata (e.g., Etessami and Yannakakis 2009, 2012; Esparza et al. 2006; Etessami et al. 2010, 2012; Brázdil et al. 2011). These are all finitely presentable models that specify an infinite-state underlying probabilistic system. These classes of systems arise in a variety of fields and have been studied by various communities. Recursive Markov chains (RMC), and the equivalent model of probabilistic pushdown systems (pPDS), are natural models for probabilistic programs with recursive procedures [Etessami and Yannakakis 2009; Esparza et al. 2006]. Quasi-birth-death (QBD) processes, which are essentially equivalent (in discrete-time) to probabilistic 1-counter automata (p1CA), are used in queueing theory and performance evaluation [Neuts 1981; Latouche and Ramaswami 1999]. Stochastic context-free grammars are a central model in natural language processing and are used also in biology [Durbin et al. 1999], and branching processes are a classical probabilistic model with many applications, including in population genetics [Harris 1963].

A central problem for the analysis and model checking of these systems is the computation of their associated *termination probabilities*. Computing these probabilities amounts to solving a system of fixed-point multivariate equations $x = P(x)$, where x is a (finite) vector of variables and P is a vector of polynomials with positive coefficients; such a system of equations is called a *monotone polynomial system (MPS)* because P defines a monotone operator from the nonnegative orthant to itself. Each of these classes of models has the property that, given a model M in the class, we can construct in polynomial time a corresponding MPS $x = P(x)$ such that the termination probabilities of M (for various initial states) are the *least fixed point (LFP)* solution of the system, that is, they satisfy the system, and any other nonnegative solution is at least as large in every coordinate. In general, a monotone polynomial system may not have any nonnegative fixed-point solution; consider for example $x = x + 1$. However, if it has a nonnegative fixed point, then it has a least fixed point (LFP). The systems constructed from probabilistic systems as previously described always have a LFP, which has values in $[0, 1]$ since its coordinates give the termination probabilities.

The equations are in general nonlinear, and their LFP solution (the vector of termination probabilities) is in general irrational even when all the coefficients of the polynomials (and the numerical input data of the given probabilistic model) are rational. Hence we seek to compute the desired quantities up to a desired accuracy $\epsilon > 0$. The goal is to compute them as efficiently as possible, as a function of the encoding size of the input (the given probabilistic model, or the MPS) and the accuracy ϵ . We first review some of the relevant previous work and then describe our results.

Previous Work. An algorithm for computing the LFP of MPSs, based on Newton's method, was proposed in Etessami and Yannakakis [2009]. Given a MPS, we can first identify in polynomial time the variables that have value 0 in the LFP and remove them from the system, yielding a new so-called *cleaned* system. Then, a dependency graph between the variables is constructed, the variables and the MPS are decomposed into strongly connected components (SCCs), and Newton's method is applied bottom-up on the SCCs, starting from the all-0 vector. It was shown in Etessami and Yannakakis [2009] that, for any MPS that has a (nonnegative) solution, the decomposed variant of Newton's method converges monotonically to the LFP. Optimized variants of decomposed Newton's method have by now been implemented in several

tools (see, e.g., Wojtczak and Etessami [2007] and Nederhof and Satta [2008]), and they perform quite well in practice on many instances.

Esparza et al. [2010] studied in detail the rate of convergence of Newton's method on MPSs (with or without decomposition). On the negative side, they showed that there are instances of MPSs $x = P(x)$ (in fact, even simple RMCs), with n variables, where it takes an exponential number of iterations in the input size to get even within just one bit of precision (i.e., accuracy $1/2$). On the positive side, they showed that after some initial number k_P of iterations in a first phase, Newton's method thereafter gains bits of precision at a linear rate, meaning that $k_P + c_P \cdot i$ iterations suffice to gain i bits of precision, where both k_P and c_P depend on the input, $x = P(x)$. For strongly connected MPSs, they showed that the length, k_P , of the initial phase is upper bounded by an exponential function of the input size $|P|$, and that $c_P = 1$. For general MPSs that are not strongly connected (and for general RMCs and pPDSs), they showed that $c_P = n2^n$ suffices, but they provided no upper bound at all on k_P (and none was known prior to the present article). Thus, they obtained no upper bounds, as a function of the size of the input, $x = P(x)$, for the number of iterations required to get to within even the first bit of precision (e.g., to estimate within $< 1/2$ the termination probability of a RMC) for general MPSs and RMCs. Proving such a general bound was left as an open problem in Esparza et al. [2010].

For special classes of probabilistic models (and MPSs), better results are now known. For the class of quasi-birth-death processes (QBDs) and the equivalent class of probabilistic 1-counter automata (p1CA), it was shown in Etessami et al. [2010] that the decomposed Newton method converges in a polynomial number of iterations in the size of the input and the bits of precision, and hence the desired termination probabilities of a given p1CA M can be computed within absolute error $\epsilon = 2^{-i}$ in a number of arithmetic operations that is polynomial in the size $|M|$ of the input and the number $i = \log(1/\epsilon)$ of bits of precision. Note that this is not polynomial time in the standard Turing model of complexity, because the numbers that result from the arithmetic operations in general can become exponentially long (consider n successive squarings of a number). Thus, the result of Etessami et al. [2010] shows that the termination problem for p1CAs can be solved in polynomial time in the unit-cost exact rational arithmetic model, a model in which arithmetic operations cost 1 time unit, regardless of how long the numbers are. It is not known exactly how powerful the unit-cost rational model is, but it is believed to be strictly more powerful than the ordinary Turing model. The question whether the termination probabilities of a p1CA (and a QBD) can be computed in polynomial time (in the standard model) was left open in Etessami et al. [2010].

Building on the results of Etessami et al. [2010] for computation of termination probabilities of p1CAs, more recently [Brázdil et al. 2011] showed how to do quantitative model checking of ω -regular properties (given by a deterministic Rabin automaton) for p1CAs, that is, compute within desired precision $\epsilon > 0$ the probability that a run of a given p1CA, M , is accepted by a given deterministic Rabin automaton, R , in time polynomial in $M, R, \log(1/\epsilon)$ in the unit-cost rational arithmetic model. The complexity in the standard Turing model was left open.

For the classes of stochastic context-free grammars, multitype branching processes, and the related class of 1-exit RMCs, we showed recently in Etessami et al. [2012] that termination probabilities can be computed to within precision ϵ in polynomial time in the size of the input model and $\log(1/\epsilon)$ (i.e., the # of bits of precision) in the standard Turing model [Etessami et al. 2012]. The algorithm is a variant of Newton's method, where the preprocessing identifies and eliminates (in P-time [Etessami and Yannakakis 2009]) the variables that have value 1 in the LFP (besides

the ones with value 0). Importantly, the numbers throughout the computation are not allowed to grow exponentially in length, but are always rounded down to a polynomial number of bits. The analysis then shows that the rounded Newton's algorithm still converges to the correct values (the LFP) and the number of iterations and the entire time complexity is polynomially bounded.

For general RMCs (and pPDSs) and furthermore for general MPSs, even if the LFP is in $[0, 1]^n$, there are negative results indicating that it is probably impossible to compute the termination probabilities and the LFP in polynomial time in the standard Turing model. In particular, we showed in Etessami and Yannakakis [2009] that approximating the termination probability of a RMC within *any* constant additive error $< \frac{1}{2}$, is at least as hard as the *square-root-sum* problem, a long-standing open problem that arises often in computational geometry, which is not even known to be in NP, and that it is also as hard as the more powerful problem, called PosSLP [Allender et al. 2009], which captures the essence of unit-cost rational arithmetic. Thus, if one can approximate the termination probability of a RMC in polynomial time, then it is possible to simulate unit-cost rational arithmetic in polynomial time in the standard model, something which is highly unlikely.

As we mentioned at the beginning, computing termination probabilities is a key ingredient for performing other, more general analyses, including model checking [Etessami and Yannakakis 2012; Esparza et al. 2006].

Our Results. We provide a thorough analysis of decomposed Newton's method and show upper bounds on its rate of convergence as a function of the input size and the desired precision, which holds for *arbitrary* monotone polynomial systems. Furthermore, we analyze a *rounded* version of the algorithm where the results along the way are not computed exactly to arbitrary precision but are rounded to a suitable number of bits (proportional to the number of iterations k of Newton's method that are performed), while ensuring that the algorithm stays well defined and converges to the LFP. Thus, the bounds we show hold for the standard Turing model and not only the unit-cost model. Note that all the previous results on Newton's method that we mentioned, except for Etessami et al. [2012], assume that the computations are carried out in *exact* arithmetic. To carry out k iterations of Newton's method with exact arithmetic can require exponentially many bits, as a function of k , to represent the iterates. In general, the fact that Newton's method converges with exact arithmetic does not even imply automatically that rounded Newton iterations will get anywhere close to the solution when we round to, say, only polynomially many bits of precision as a function of the number of iterations k , let alone that the same bounds on the convergence rate will continue to hold. We nevertheless show that suitable rounding works for MPSs.

In more detail, suppose that the given (cleaned) MPS $x = P(x)$ has a LFP $q^* > 0$. The decomposition into strongly connected components yields a DAG of SCCs with depth d , and we wish to compute the LFP with (absolute) error at most ϵ . Let q_{\min}^* and q_{\max}^* be the minimum and maximum coordinate of q^* . Then the rounded decomposed Newton's method will converge to a vector \tilde{q} within ϵ of the LFP, that is, such that $\|q^* - \tilde{q}\|_\infty \leq \epsilon$ in time polynomial in the size $|P|$ of the input, $\log(1/\epsilon)$, $\log(1/q_{\min}^*)$, $\log(q_{\max}^*)$, and 2^d (the depth d in the exponent can be replaced by the maximum number of *nonlinear* SCCs in any path of the DAG of SCCs). We also obtain bounds on q_{\min}^* and q_{\max}^* in terms of $|P|$ and the number of variables n , so the overall time needed is polynomial in $|P|$, 2^n and $\log(1/\epsilon)$. We provide actually concrete expressions on the number of iterations and the number of bits needed. As we shall explain, the bounds are essentially optimal in terms of several parameters. The analysis is quite involved and builds on the previous work. It uses several results and techniques from

Etessami and Yannakakis [2009], Esparza et al. [2010], and Etessami et al. [2012], and develops substantial additional machinery.

We apply our results then to probabilistic 1-counter automata (p1CAs). Using our analysis for the rounded decomposed Newton method and properties of p1CAs from Etessami et al. [2010], we show that termination probabilities of a p1CA M (and QBDs) can be computed to desired precision ϵ in polynomial time in the size $|M|$ of the p1CA and $\log(1/\epsilon)$ (the bits of precision) in the standard Turing model of computation, thus solving the open problem from Etessami et al. [2010].

Furthermore, combining with the results of Brázdil et al. [2011] and Etessami and Yannakakis [2012], we show that one can do quantitative model checking of ω -regular properties for p1CAs in polynomial time in the standard Turing model, that is, we can compute to desired precision ϵ the probability that a run of a given p1CA M satisfies an ω -regular property in time polynomial in $|M|$ and $\log(1/\epsilon)$ (and exponential in the property if it is given, for example, as a nondeterministic Büchi automaton or polynomial if it is given as a deterministic Rabin automaton).

The rest of the article is organized as follows. In Section 2, we give basic definitions and background, and discuss preliminaries. In Section 3, we consider strongly connected MPS, and we prove a central theorem which both provides upper bounds on Newton's method for strongly connected MPSs and also crucially provides bounds on how much the LFP of a strongly connected MPS would change if certain parameter values defining the MPS are decreased. In Section 4, we consider general MPS, building on the results we establish for the strongly connected case, we establish upper bounds on Newton's method for general MPSs. We then also discuss several senses in which our results are essentially optimal for MPSs. In Section 5, we analyze the special MPSs associated with probabilistic 1-counter automata (p1CAs), and we apply results from Section 4 to show that termination probabilities for these can be computed in P-time; we then use this to conclude that computing the model checking probability (to within desired precision) that a given p1CA satisfies a (fixed) LTL or ω -regular property can be done in P-time.

2. DEFINITIONS AND BACKGROUND

We first recall basic definitions about MPSs from Etessami and Yannakakis [2009]. A *monotone polynomial system of equations* (MPS) consists of a system of n equations in n variables, $x = (x_1, \dots, x_n)$, the equations are of the form $x_i = P_i(x)$, $i = 1, \dots, n$, such that $P_i(x)$ is a multivariate polynomial in the variables x , and such that the monomial coefficients and constant term of $P_i(x)$ are all nonnegative. More precisely, for $\alpha = (\alpha_1, \alpha_2, \dots, \alpha_n) \in \mathbb{N}^n$, we use the notation x^α to denote the monomial $x_1^{\alpha_1} x_2^{\alpha_2} \dots x_n^{\alpha_n}$. (Note that by definition $x^{(0, \dots, 0)} = 1$.) Then, for each polynomial $P_i(x)$, $i = 1, \dots, n$, there is some *finite* subset of \mathbb{N}^n , denoted C_i , and for each $\alpha \in C_i$, there is a positive (rational) coefficient $c_{i,\alpha} > 0$, such that $P_i(x) \equiv \sum_{\alpha \in C_i} c_{i,\alpha} x^\alpha$.

For computational purposes, we assume each polynomial $P_i(x)$ has rational coefficients,¹ and that it is encoded succinctly by specifying the list of pairs $\langle (c_{i,\alpha}, \alpha) \mid \alpha \in C_i \rangle$, where each rational coefficient $c_{i,\alpha}$ is represented by giving its numerator and denominator in binary, and each integer vector α is represented in *sparse representation*, by only listing its nonzero coordinates, i_1, \dots, i_k , by using a list $\langle (i_1, \alpha_{i_1}), \dots, (i_k, \alpha_{i_k}) \rangle$, giving each integer α_{i_j} in binary. (Proposition 2.1 in this section, from Etessami and Yannakakis [2009] and Etessami et al. [2012], shows that using such a sparse representation does not entail any extra computational cost.)

¹Although we also reason about MPSs with positive real-valued coefficients in our proofs.

We use vector notation, using $x = P(x)$ to denote the entire MPS. We use $|P|$ to denote the encoding size (in bits) of the MPS $x = P(x)$ having rational coefficients, using the succinct representation just described. Let $\mathbb{R}_{\geq 0}$ denote the non-negative real numbers. We shall often use 0 (respectively, 1) to refer to an all 0 (respectively, all 1) vector of appropriate dimension, where the dimension will be clear from the context. For vectors a and b , we use $a \leq b$ (respectively, $a < b$) to mean inequality (respectively, strict inequality) in every coordinate.

Note that $P(x)$ defines a monotone operator on the nonnegative orthant $\mathbb{R}_{\geq 0}^n$. In other words, $P : \mathbb{R}_{\geq 0}^n \rightarrow \mathbb{R}_{\geq 0}^n$, and if $0 \leq a \leq b$, then $P(a) \leq P(b)$. In general, an MPS need not have any real-valued solution: consider $x = x + 1$. However, because of monotonicity of $P(x)$, if there exists a solution $a \in \mathbb{R}_{\geq 0}^n$ such that $a = P(a)$, then there exists a *least fixed point* (LFP) solution $q^* \in \mathbb{R}_{\geq 0}^n$ such that $q^* = P(q^*)$, and such that $q^* \leq a$ for all solutions $a \in \mathbb{R}_{\geq 0}^n$. Indeed, if for $z \in \mathbb{R}^n$ we define $P^0(z) = z$, and define $P^{k+1}(z) = P(P^k(z))$, for all $k \geq 0$, then (as shown in Etessami and Yannakakis [2009]) value iteration starting at the all-0 vector converges monotonically to q^* : in other words $\forall k \geq 0$ $P^k(0) \leq P^{k+1}(0)$, and $\lim_{k \rightarrow \infty} P^k(0) = q^*$.²

Unfortunately, standard value iteration $P^k(0)$, $k \rightarrow \infty$, can converge very slowly to q^* , even for a fixed MPS with 1 variable, even when $q^* = 1$; specifically, $x = (1/2)x^2 + 1/2$ already exhibits exponentially slow convergence to its LFP $q^* = 1$ [Etessami and Yannakakis 2009]. It was shown in Etessami and Yannakakis [2009] that a decomposed variant of Newton's method also converges monotonically to q^* for an MPS with LFP solution q^* . More recently, in Etessami et al. [2012], a version of Newton's method with suitable rounding between iterations was studied. Rounding is necessary if one wishes to consider the complexity of Newton's method in the standard (Turing) model of computation, which does not allow unit-cost arithmetic operations on arbitrarily large numbers. In this article, we will apply a version of Newton's method to MPSs which uses both rounding and decomposition. Before describing it, we need some further background.

An MPS, $x = P(x)$, is said to be in *simple normal form* (SNF) if for every $i = 1, \dots, n$, the polynomial $P_i(x)$ has one of two forms: (1) Form_{*}: $P_i(x) \equiv x_j x_k$ is simply a quadratic monomial; or (2) Form₊: $P_i(x)$ is a *linear* expression $\sum_{j \in C_i} p_{i,j} x_j + p_{i,0}$, for some rational nonnegative coefficients $p_{i,j}$ and $p_{i,0}$, and some index set $C_i \subseteq \{1, \dots, n\}$. In particular, in any MPS in SNF form every polynomial $P_i(x)$ has multivariate degree bounded by at most 2 in the variables x . We will call such MPSs *quadratic MPSs*.

As shown in Etessami and Yannakakis [2009] and Etessami et al. [2012], it is easy to convert any MPS to SNF form, by adding auxiliary variables and equations.

PROPOSITION 2.1. [ETESSAMI AND YANNAKAKIS 2009, PROPOSITION 7.3; ETESSAMI ET AL. 2012, PROPOSITION 2.1]. *Every MPS, $x = P(x)$, with LFP q^* , can be transformed in P -time to an “equivalent” quadratic MPS $y = Q(y)$ in SNF form, such*

²Indeed, even if an MPS does not have a *finite* LFP solution $q^* \in \mathbb{R}_{\geq 0}^n$, it always does have an LFP solution over the extended nonnegative reals. Namely, we can define the LFP of any MPS, $x = P(x)$, to be the vector $q^* \in \overline{\mathbb{R}}_{\geq 0}^n$ over $\overline{\mathbb{R}}_{\geq 0} = (\mathbb{R}_{\geq 0} \cup \{+\infty\})$, given by $q^* := \lim_{k \rightarrow \infty} P^k(0)$. In general, it is PosSLP-hard to decide whether a given MPS has a finite LFP. (This follows easily from results in Etessami and Yannakakis [2009], although it is not stated there: it was shown there that it is PosSLP-hard to decide if $q_1^* \geq 1$ in an MPS with finite LFP $q^* \in \mathbb{R}_{\geq 0}^n$. Then, just add a variable x_0 , and an equation $x_0 = x_0 x_1 + 1$ to the MPS. In the new MPS, $q_0^* = +\infty$ if and only if $q_1^* \geq 1$.) However, various classes of MPSs, including those whose LFP corresponds to termination probabilities of various recursive probabilistic systems do have a finite LFP. Thus, in this article, we will only consider LFP computation for MPSs that have a finite LFP $q^* \in \mathbb{R}_{\geq 0}^n$. So when we say “ $x = P(x)$ is an MPS with LFP solution q^* ”, we mean $q^* \in \mathbb{R}_{\geq 0}^n$, unless specified otherwise.

that $|Q| \in O(|P|)$. More precisely, the variables x are a subset of the variables y , and $y = Q(y)$ has LFP p^* iff $x = P(x)$ has LFP q^* , and projecting p^* onto the x variables yields q^* .

Furthermore, for any MPS, $x = P(x)$, we can in P-time find and remove any variables x_i , such that the LFP solution has $q_i^* = 0$.³

PROPOSITION 2.2 [ETESSAMI AND YANNAKAKIS 2009, PROPOSITION 7.4]. *There is a P-time algorithm that, given any MPS³, $x = P(x)$, over n variables, determines for each $i \in \{1, \dots, n\}$ whether $q_i^* = 0$.*

Thus, for every MPS, we can detect in P-time all the variables x_j such that $q_j^* = 0$, remove their equation $x_j = P_j(x)$, and set the variable x_j to 0 on the right-hand side of the remaining equations. We obtain as a result a *cleaned* MPS, $x' = Q(x')$, which has an LFP $q^* > 0$.

Applying Propositions 2.1 and 2.2,

we assume without loss of generality in the rest of this article that every MPS is a cleaned quadratic MPS, with LFP $q^* > 0$.⁴

In order to describe *decomposed* Newton's method, for a *cleaned* MPS, $x = P(x)$ we need to define the *dependency graph*, $G_P = (V, E)$, of the MPS. The nodes V of G_P are the variables x_i , and the edges are defined as follows: $(x_i, x_j) \in E$ if and only if x_j appears in some monomial in $P_i(x)$ that has a positive coefficient.

We shall decompose the cleaned system of equation $x = P(x)$, into strongly connected components (SCCs), using the dependency graph G_P of variables, and we shall apply Newton's method separately on each SCC "bottom-up".

We first recall basic definitions for (a rounded-down version of) Newton's method applied to MPSs. For an MPS, $x = P(x)$, with n variables, we define $B(x) = P'(x)$ to be the $n \times n$ Jacobian matrix of partial derivatives of $P(x)$. In other words, $B(x)_{i,j} = \frac{\partial P_i(x)}{\partial x_j}$. For a vector $z \in \mathbb{R}^n$, assuming that the matrix $(I - B(z))$ is nonsingular, a single iteration of *Newton's method* (NM) on $x = P(x)$ at z is defined via the following operator:

$$\mathcal{N}_P(z) := z + (I - B(z))^{-1}(P(z) - z). \quad (1)$$

Let us now recall from Etessami et al. [2012] the rounded down Newton's method, with parameter h , applied to an MPS:

Definition 2.3. (*Rounded-Down Newton's Method (R-NM), with rounding parameter h .*) Given an MPS, $x = P(x)$, with LFP q^* , where $0 < q^*$, in the *rounded-down Newton's method* (R-NM) with integer rounding parameter $h > 0$, we compute a sequence of iteration vectors $x^{[k]}$, where the initial starting vector is $x^{[0]} := 0$, and such that for each $k \geq 0$, given $x^{[k]}$, we compute $x^{[k+1]}$ as follows.

- (1) First, compute $x^{[k+1]} := \mathcal{N}_P(x^{[k]})$, where the Newton iteration operator $\mathcal{N}_P(x)$ was defined in Eq. (1). (Of course, we need to show that all such Newton iterations are defined.)
- (2) For each coordinate $i = 1, \dots, n$, set $x_i^{[k+1]}$ to be equal to the maximum (non-negative) multiple of 2^{-h} which is $\leq \max(x_i^{[k+1]}, 0)$. (In other words, round down

³This proposition holds regardless whether the LFP q^* is *finite* or is over the *extended nonnegative reals*, $\mathbb{R}_{\geq 0}$. Such an extended LFP exists for any MPS. See footnote 2.

⁴For compatibility when quoting prior work, it will sometimes be convenient to assume quadratic MPSs, rather than the more restricted SNF form MPSs.

$x^{[k+1]}$ to the nearest multiple of 2^{-h} , while making sure that the result is non-negative.)

Now we describe the *Rounded-down Decomposed Newton's Method (R-DNM)* applied to an MPS, $x = P(x)$, with real-valued LFP $q^* \geq 0$. First, we use Proposition 2.2 to remove 0 variables, and thus we can assume we are given a cleaned MPS, $x = P(x)$, with real-valued LFP $q^* > 0$.

Let H_P be the DAG of SCCs of the dependency graph G_P . We work bottom-up in H_P , starting at bottom SCCs. For each SCC, S , suppose its corresponding equations are $x_S = P_S(x_S, x_{D(S)})$, where $D(S)$ denotes the union of the variables in “lower” SCCs, below S , on which S depends. In other words, a variable $x_j \in D(S)$ iff there is some variable $x_i \in S$ such that there is directed path in G_P from x_i to x_j . If the system $x_S = P_S(x_S, q_{D(S)}^*)$ is a linear system (in x_S), we call S a *linear SCC*, otherwise S is a *nonlinear SCC*. Assume we have already calculated (using R-DNM) an approximation $\tilde{q}_{D(S)}$ to the LFP solution $q_{D(S)}^*$ for these lower SCCs. We plug in $\tilde{q}_{D(S)}$ into the equations for S , obtaining the equation system $x_S = P_S(x_S, \tilde{q}_{D(S)})$. We denote the actual LFP solution of this new equation system by q'_S . (Note that q'_S is not necessarily equal to q_S^* , because $\tilde{q}_{D(S)}$ is only an approximation of $q_{D(S)}^*$.)

If S is a nonlinear SCC, we apply a chosen number g of iterations of R-NM on the system $x_S = P_S(x_S, \tilde{q}_{D(S)})$ to obtain an approximation \tilde{q}_S of q'_S ; if S is linear, then we just apply 1 iteration of R-NM, that is, we solve the linear system and round down the solution. Of course, we want to make sure our approximations are such that $\|q_S^* - \tilde{q}_S\|_\infty \leq \epsilon$, for all SCCs S , and for the desired additive error $\epsilon > 0$. We shall establish upper bounds on the number of iterations g , and on the rounding parameter h , needed in R-DNM for this to hold, as a function of various parameters: the input size $|P|$ and the number n of variables; the *nonlinear depth* f of P , which is defined as the maximum, over all paths of the DAG H_P of SCCs, of the number of nonlinear SCCs on the path; and the maximum and minimum coordinates of the LFP.

Bounds on the Size of LFPs for an MPS. For a positive vector $v > 0$, we use $v_{\min} = \min_i v_i$ to denote its minimum coordinate, and we use $v_{\max} = \max_i v_i$ to denote its maximum coordinate. Slightly overloading notation, for an MPS, $x = P(x)$, we shall use c_{\min} to denote the minimum value of all positive monomial coefficients and all positive constant terms in $P(x)$. Note that c_{\min} also serves as a lower bound for all positive constants and coefficients for entries of the Jacobian matrix $B(x)$, since $B(x)_{ij} = \frac{\partial P_i(x)}{\partial x_j}$.

We prove the following theorem in the appendix, establishing bounds on the maximum and minimum coordinates of the LFP q^* of an MPS $x = P(x)$.

THEOREM 2.4. *If $x = P(x)$ is a quadratic MPS in n variables, with LFP $q^* > 0$, and where $P(x)$ has rational coefficients and total encoding size $|P|$ bits, then*

- (1) $q_{\min}^* \geq 2^{-|P|(2^n-1)}$, and
- (2) $q_{\max}^* \leq 2^{2(n+1)(|P|+2(n+1)\log(2n+2)) \cdot 5^n}$.

3. STRONGLY CONNECTED MONOTONE POLYNOMIAL SYSTEMS

The following theorem is at the heart of this article. The theorem first establishes bounds on the distance between the LFP of a strongly connected MPS, and the LFP of another MPS obtained from it by decreasing values of some positive parameters. Second, the theorem establishes bounds on the number of iterations required by rounded down Newton's method to converge to within desired error $\epsilon > 0$, on the MPS with decreased parameter values. Because decomposed Newton's method works by

running Newton's method on strongly connected components sequentially, in a bottom-up fashion, we will later be able to use the bounds obtained in this theorem to analyze decomposed Newton's method on arbitrary MPSs in an inductive fashion.

THEOREM 3.1. *Let $P(x, y)$ be an n -vector of monotone polynomials with degree ≤ 2 in variables which are coordinates of the n -vector x and the m -vector y , where $n \geq 1$ and $m \geq 1$.*

Given nonnegative m -vectors y_1 and y_2 such that $0 < y_1 \leq 1$ and $0 \leq y_2 \leq y_1$, let $P_1(x) \equiv P(x, y_1)$ and $P_2(x) \equiv P(x, y_2)$. Suppose that $x = P_1(x)$ is a strongly connected MPS with LFP solution $0 < q_1^ \leq 1$.*

Let $\alpha = \min\{1, c_{\min}\} \min\{y_{\min}, \frac{1}{2}q_{\min}^\}$, where c_{\min} is the smallest nonzero constant or coefficient of any monomial in $P(x, y)$, where y_{\min} is the minimum coordinate of y_1 , and finally where q_{\min}^* is the minimum coordinate of q_1^* . Then*

- (1) *the LFP solution of the MPS $x = P_2(x)$ is q_2^* with $0 \leq q_2^* \leq q_1^*$, and*

$$\|q_1^* - q_2^*\|_{\infty} \leq \sqrt{4n\alpha^{-(3n+1)} \|P(1, 1)\|_{\infty} \|y_1 - y_2\|_{\infty}}. \quad (2)$$

Furthermore, if $x = P_1(x)$ is a linear system, then

$$\|q_1^* - q_2^*\|_{\infty} \leq 2n\alpha^{-(n+2)} \|P(1, 1)\|_{\infty} \|y_1 - y_2\|_{\infty}. \quad (3)$$

- (2) *Moreover, for every $0 < \epsilon < 1$, if we use $g \geq h - 1$ iterations of the rounded-down Newton's method with parameter*

$$h \geq \left\lceil 2 + n \log \frac{1}{\alpha} + \log \frac{1}{\epsilon} \right\rceil$$

applied to the MPS, $x = P_2(x)$, starting at $x^{[0]} := 0$, to approximate q_2^ , then the iterations are all defined, and $\|q_2^* - x^{[g]}\|_{\infty} \leq \epsilon$.*

The proof of Theorem 3.1 is rather involved, and will require a number of lemmas. Before starting to prove it, we first establish the following easy corollary of Theorem 3.1.

COROLLARY 3.2. *Let $x = P(x)$ be a strongly connected MPS with n variables, and with LFP q^* where $0 < q^* \leq 1$. Let $\alpha = \min\{1, c_{\min}\} \frac{1}{2}q_{\min}^*$, where c_{\min} is the smallest nonzero constant or coefficient of any monomial in $P(x)$.*

Then, for all $0 < \epsilon < 1$, if we use $g \geq h - 1$ iterations of R-NM with parameter $h \geq \lceil 2 + n \log \frac{1}{\alpha} + \log \frac{1}{\epsilon} \rceil$ applied to the MPS, $x = P(x)$, starting at $x^{[0]} := 0$, then the iterations are all defined, and $\|q^ - x^{[g]}\|_{\infty} \leq \epsilon$.*

PROOF OF COROLLARY 3.2. This follows by a trivial application of part (2) of Theorem 3.1, where we define y to be a dummy variable of dimension $m = 1$, and we define $y_1 = y_2 = y_{\min} = 1$, and where we define the n -vector of monotone polynomials $P(x, y)$, by replacing all constant terms $c > 0$ in every polynomial in $P(x)$ by cy . In this case, note that $P_1(x) = P_2(x) = P(x)$, and that since $y_{\min} = 1$, the α defined in the statement of this corollary is the same α as in the statement of Theorem 3.1. \square

Before giving the proof of Theorem 3.1, we will first state some useful known facts about MPSs and nonnegative matrices. We first recall some lemmas from Etessami et al. [2012].⁵

⁵In Etessami et al. [2012], the statements of Lemmas 3.3 and 3.4 assume that the MPS is in SNF form, but as noted in Etessami et al. [2012], the proofs of Lemmas 3.3 and 3.4 do not require that $x = P(x)$ is in SNF form, nor that it is an MPS, only that it is quadratic.

LEMMA 3.3 [ETESSAMI ET AL. 2012, LEMMA 3.3]. *Let $x = P(x)$ be a quadratic MPS, with n variables, and let $a, b \in \mathbb{R}^n$. Then*

$$P(a) - P(b) = B\left(\frac{a+b}{2}\right)(a-b) = \frac{B(a) + B(b)}{2}(a-b).$$

LEMMA 3.4. *Let $x = P(x)$ be a quadratic MPS. Let $z \in \mathbb{R}^n$ be any vector such that $(I - B(z))$ is nonsingular, and thus $\mathcal{N}_P(z)$ is defined. Then*

$$q^* - \mathcal{N}_P(z) = (I - B(z))^{-1} \frac{B(q^*) - B(z)}{2} (q^* - z).$$

We will also need the following lemma from Esparza et al. [2010].

LEMMA 3.5. [ESPARZA ET AL. 2010, LEMMA 5.4; ETESSAMI ET AL. 2012, LEMMA 3.7]. *Let $x = P(x)$ be a MPS, with polynomials of degree bounded by 2, with LFP, $q^* \geq 0$. Let $B(x)$ denote the Jacobian matrix of $P(x)$. For any positive vector $d \in \mathbb{R}_{>0}^n$ that satisfies $B(q^*)d \leq d$, any positive real value $\lambda > 0$, and any nonnegative vector $z \in \mathbb{R}_{\geq 0}^n$, if $q^* - z \leq \lambda d$, and $(I - B(z))^{-1}$ exists and is nonnegative, then*

$$q^* - \mathcal{N}_P(z) \leq \frac{\lambda}{2} d.$$

We next recall a number of basic facts from matrix analysis and Perron-Frobenius theory. For a square matrix A , let $\rho(A)$ denote the spectral radius of A . Recall that a nonnegative square matrix A is called *irreducible* if its underlying directed graph is strongly connected, where the adjacency matrix of its underlying directed graph is obtained by setting the positive entries of the matrix A to 1.

LEMMA 3.6 (SEE, E.G., HORN AND JOHNSON [1985], THEOREM 8.4.4). *If A is an irreducible nonnegative square matrix, then there is a positive eigenvector $v > 0$, such that $Av = \rho(A)v$. Such a vector v is called the Perron vector of A . It is unique up to rescaling by a positive factor.*

LEMMA 3.7. (SEE, E.G., LANCASTER AND TISMENETSKY [1985, THEOREM 15.4.1] AND EXERCISE 1, PAGE 540) *If A is an irreducible nonnegative square matrix and $0 \leq B \leq A$, but $B \neq A$, then $\rho(B) < \rho(A)$.*

LEMMA 3.8. (SEE, E.G., LANCASTER AND TISMENETSKY [1985, THEOREM 15.2.2, PAGE 531].) *If A is a square matrix with $\rho(A) < 1$, then $I - A$ is nonsingular and $(I - A)^{-1} = \sum_{i=0}^{\infty} A^i$.*

LEMMA 3.9. (SEE, E.G., LANCASTER AND TISMENETSKY [1985, SECTION 15.3 AND EXERCISE 11].) *If A is an irreducible nonnegative square matrix, and $v > 0$ is a positive eigenvector associated with some eigenvalue r , that is, such that $Av = rv$, then $r = \rho(A)$. Thus, $v > 0$ is the Perron vector (which is unique up to scaling).*

We proceed now to the proof of Theorem 3.1. We prove each of the two parts in turn.

PROOF OF PART 1 OF THEOREM 3.1. First note that, if $q_1^* = q_2^*$, then the result is trivial. So we assume henceforth that $q_1^* \neq q_2^*$. The proof is rather long and involved. We outline first the structure of the proof. There are four main steps. In step (i), we show that $q_1^* - q_2^* = [I - B_1(\frac{1}{2}(q_1^* + q_2^*))]^{-1} (P_1(q_2^*) - P_2(q_2^*))$, and in particular we show that the matrix $[I - B_1(\frac{1}{2}(q_1^* + q_2^*))]$ is nonsingular, where $B_1(x)$ is the Jacobian matrix of

$P_1(x)$. In step (ii), we bound $\| [I - B_1(\frac{1}{2}(q_1^* + q_2^*))]^{-1} \|_\infty$. In step (iii), we bound $\|P_1(q_2^*) - P_2(q_2^*)\|_\infty$. In step (iv), we combine the bounds from steps (ii) and (iii) to obtain the desired bound on $\|q_1^* - q_2^*\|_\infty$.

STEP (i). We observe first that $0 \leq q_2^* \leq q_1^*$. To see this, consider $P_1^k(0) = P_1(P_1(\dots P_1(0)\dots))$, that is, the k th iterate of P_1 applied to the vector 0, and $P_2^k(0)$. We know that for any MPS, $x = P(x)$ with LFP $q^* \in \mathbb{R}_{\geq 0}^n$, we have $\lim_{k \rightarrow \infty} P^k(0) = q^*$ [Etessami and Yannakakis 2009]. Thanks to the monotonicity of P , for any $x \geq 0$, we have $P_1(x) \geq P_2(x)$. By the monotonicity of P_1 and an easy induction, $P_1^k(0) \geq P_2^k(0)$. So $q_1^* = \lim_{k \rightarrow \infty} P_1^k(0) \geq \lim_{k \rightarrow \infty} P_2^k(0) = q_2^*$.

Because $x = P_1(x)$ is at most quadratic, we can apply Lemma 3.3 to get

$$B_1\left(\frac{1}{2}(q_1^* + q_2^*)\right)(q_1^* - q_2^*) = P_1(q_1^*) - P_1(q_2^*) = q_1^* - P_1(q_2^*). \quad (4)$$

Multiplying both sides of Eq. (4) by -1 , and then adding $(q_1^* - q_2^*)$ to both sides, we get

$$\begin{aligned} \left(I - B_1\left(\frac{1}{2}(q_1^* + q_2^*)\right)\right)(q_1^* - q_2^*) &= (q_1^* - q_2^*) - (q_1^* - P_1(q_2^*)) \\ &= P_1(q_2^*) - q_2^* \\ &= P_1(q_2^*) - P_2(q_2^*). \end{aligned} \quad (5)$$

Provided that $I - B_1(\frac{1}{2}(q_1^* + q_2^*))$ is nonsingular, we can multiply both sides of Eq. (5) by $\left[I - B_1\left(\frac{1}{2}(q_1^* + q_2^*)\right)\right]^{-1}$, to get

$$q_1^* - q_2^* = \left[I - B_1\left(\frac{1}{2}(q_1^* + q_2^*)\right)\right]^{-1} (P_1(q_2^*) - P_2(q_2^*)). \quad (6)$$

It remains to show that the matrix $I - B_1(\frac{1}{2}(q_1^* + q_2^*))$ is nonsingular. By Lemma 3.8, it suffices to show that the spectral radius $\rho(B_1(\frac{1}{2}(q_1^* + q_2^*))) < 1$. For this purpose, we will use the following lemma.

LEMMA 3.10. (This is a variant of Lemma 6.5 from Etessami and Yannakakis [2009].) *For any strongly connected MPS, $x = P(x)$, with LFP $q^* > 0$, and Jacobian $B(x)$, we have $\rho(B(q^*)) \leq 1$, and for all vectors y with $0 \leq y < q^*$, $\rho(B(y)) < 1$.*

PROOF. We will only show here that $\rho(B(q^*)) \leq 1$ if $x = P(x)$ is strongly connected, but in fact this holds for any MPS, $x = P(x)$, with LFP $q^* > 0$. We do so because we will only use the strongly connected case.

If we have $0 \leq z \leq y$ and $z \leq P(z)$, then Lemma 6.4 of Etessami and Yannakakis [2009] shows that for any $d \geq 1$, $B^d(z)(y - z) \leq P^d(y) - P^d(z)$. Let $x^i = P^i(0)$, for all $i \geq 1$. Recall that $\lim_{i \rightarrow \infty} x^i = q^*$. Also note that, because $x = P(x)$ is strongly connected, $x^i < q^*$ for all i .

Then, for all $i, d \geq 1$, $B^d(x^i)(q^* - x^i) \leq P^d(q^*) - P^d(x^i) = q^* - x^{i+d}$. But since $\lim_{d \rightarrow \infty} x^{i+d} = q^*$, we see that the right-hand side goes to 0. But since $(q^* - x^i) > 0$ for all i , it must be the case that $B^d(x^i) \rightarrow 0$, as d goes to infinity. But this is a necessary and sufficient condition for $\rho(B(x^i)) < 1$. Now notice that for any vector y such that $0 \leq y < q^*$, there is some i such that $y \leq x^i$. Thus, by monotonicity of $\rho(B(x))$ in $x \geq 0$, we must have $\rho(B(y)) < 1$.

Thus, also, since $\lim_{i \rightarrow \infty} x^i = q^*$, and by continuity of the spectral radius function, we get that $\rho(B(q^*)) \leq 1$. \square

We will apply this lemma to the system $x = P_1(x)$, which has LFP q_1^* , and $y = \frac{1}{2}(q_1^* + q_2^*)$. To apply the lemma, we need to show that $\frac{1}{2}(q_1^* + q_2^*) < q_1^*$, that is, that $q_2^* < q_1^*$. We will show in fact something stronger, which will be needed later on: we will bound from below the ratio of the minimum to the maximum entry of the vector $q_1^* - q_2^*$. For this, we will use the following lemma.

LEMMA 3.11. *If A is an irreducible, nonnegative $n \times n$ matrix with minimum nonzero entry a_{\min} , and $u \geq 0$ is a nonzero vector in \mathbb{R}^n with $Au \leq u$, then $a_{\min} \leq 1$ and if the minimum and maximum coordinates of u are denoted u_{\min} and u_{\max} , respectively, then we have $\frac{u_{\min}}{u_{\max}} \geq a_{\min}^n$. In particular, $u > 0$.*

PROOF. Let i, j be some coordinates with $u_i = u_{\min}$ and $u_j = u_{\max}$. Because A is irreducible and nonnegative, there is a power $0 \leq k \leq n$ with $(A^k)_{ij} > 0$. By matrix multiplication, for any $k \geq 1$, $(A^k)_{ij} = \sum \prod_l A_{i_l, i_{l+1}}$, where the sum is taken over all length $k+1$ sequences of indices i_1, \dots, i_{k+1} , with $i_1 = i$ and $i_{k+1} = j$, and with $i_l \in \{1, \dots, n\}$ for all l ranging from 1 to k . At least one of these products is nonzero and thus it is at least a_{\min}^k . That is $(A^k)_{ij} \geq a_{\min}^k$. Since $Au \leq u$, and A is nonnegative, a simple induction gives that $A^k u \leq u$. And since u is nonzero, $u_{\max} = u_j > 0$, so $0 < A^k u_j \leq u_i$. Since $u_i = u_{\min}$, this means $u > 0$. Also, $1 \geq \frac{u_{\min}}{u_{\max}} = \frac{u_i}{u_j} \geq A^k u_j \geq a_{\min}^k$. Note that since $1 \geq a_{\min}^k$, this implies $a_{\min} \leq 1$. We know that $1 \leq k \leq n$, so $a_{\min}^k \geq a_{\min}^n$. \square

We will apply this lemma to the matrix $A = B_1(\frac{1}{2}(q_1^* + q_2^*))$ and the vector $u = (q_1^* - q_2^*)$. We verify that the hypotheses of the lemma hold. Note that since $x = P_1(x)$ is a strongly connected system of equations, for any $x > 0$, the matrix $B_1(x)$ is nonnegative and irreducible. In particular, $B_1(\frac{1}{2}(q_1^* + q_2^*))$ (as well as $B_1(q_1^*)$) is nonnegative and irreducible. Regarding u , recall that $q_1^* \neq q_2^*$, thus u is nonzero. We verify next that $Au \leq u$. By (4), we have $B_1(\frac{1}{2}(q_1^* + q_2^*))(q_1^* - q_2^*) = q_1^* - P_1(q_2^*)$. Now $P_1(q_2^*) \geq P_2(q_2^*) = q_2^*$. Thus, $q_1^* - P_1(q_2^*) \leq q_1^* - q_2^*$. So $B_1(\frac{1}{2}(q_1^* + q_2^*))(q_1^* - q_2^*) \leq (q_1^* - q_2^*)$.

We verify finally that the minimum nonzero entry of $A = B_1(\frac{1}{2}(q_1^* + q_2^*))$ is at least α . Since each polynomial in $P(x, y)$ has degree no more than 2, each entry of $B_1(x)$ is a polynomial of degree no more than 1 in both x and in the entries of y_1 when these are treated as variables. In other words, each entry of $B_1(x)$ can be expressed in the form $(\sum_i c_i x_i) + (\sum_j c'_j y_j) + c''$, where c_i, c'_j , and c'' are all nonnegative coefficients and constants of $P(x, y)$ (possibly multiplied by 2 in the case where the term of $P(x, y)$ they originate from is of the form cx^2) for all indices i and j . The nonzero entries of $B_1(\frac{1}{2}(q_1^* + q_2^*))$ are $\geq \alpha$, because the coefficients c_i, c'_j , and c'' are all $\geq c_{\min}$, and the entries of $\frac{1}{2}(q_1^* + q_2^*)$ are $\geq \frac{1}{2}q_{\min}^*$.

Thus, we can apply Lemma 3.11, to matrix $A = B_1(\frac{1}{2}(q_1^* + q_2^*))$ and vector $u = (q_1^* - q_2^*)$, which yields that

$$\frac{(q_1^* - q_2^*)_{\min}}{(q_1^* - q_2^*)_{\max}} \geq \alpha^n. \quad (7)$$

In particular, we have

$$q_2^* < q_1^*. \quad (8)$$

Consequently, $\frac{1}{2}(q_1^* + q_2^*) < q_1^*$. Therefore, by Lemma 3.10, $\rho(B_1(\frac{1}{2}(q_1^* + q_2^*))) < 1$, which by Lemma 3.8 implies that $I - B_1(\frac{1}{2}(q_1^* + q_2^*))$ is nonsingular. This concludes Step (i).

STEP (ii). We distinguish two cases, depending on whether $x = P_1(x)$ is a linear or nonlinear system. We shall separately obtain upper bounds on the matrix norm $\left\| \left(I - B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right) \right)^{-1} \right\|_\infty$ in both these cases.

Case 1: $x = P_1(x)$ is a nonlinear system. Then, the matrix $B_1(x)$ is dependent on x , that is, a variable in x appears in some entry of $B_1(x)$ with nonzero coefficient.

As we showed earlier, $B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right)$ is a nonnegative irreducible matrix with spectral radius $\rho \left(B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right) \right) < 1$, and thus by Lemma 3.8, $\left(I - B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right) \right)^{-1}$ exists and $\left(I - B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right) \right)^{-1} = \sum_{i=0}^{\infty} \left(B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right) \right)^i \geq 0$. To bound the norm of the matrix, we will use the following result from Etessami et al. [2010].

LEMMA 3.12 [ETESSAMI ET AL. 2010, LEMMA 18]. *Let $A \in \mathbb{R}_{\geq 0}^{n \times n}$ and $b \in \mathbb{R}_{\geq 0}^n$ such that: $(I - A)^{-1} = \sum_{k=0}^{\infty} A^k$, $(I - A)^{-1}b \leq 1$, and A is an irreducible nonnegative matrix whose smallest nonzero entry is $c > 0$, and $b \neq 0$ and $p > 0$ is the largest entry of b . Then, $\|(I - A)^{-1}\|_\infty \leq \frac{n}{pc^n}$.*

We will take $A = B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right)$ and $b = \left(I - B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right) \right) v$ in this lemma, where v is the Perron vector of $B_1(q_1^*)$, (i.e., $B_1(q_1^*) \cdot v = \rho(B_1(q_1^*)) \cdot v$), normalized so that its maximum entry $v_{\max} = 1$. By definition, $(I - A)^{-1}b = v \leq 1$. Thus, all the conditions of the lemma are satisfied. As we showed earlier, the smallest nonzero entry of the matrix $B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right)$ is at least α . We need to bound the largest entry of b .

Note that, since $x = P_1(x)$ is strongly connected and $q_1^* > 0$, the matrix $B_1(q_1^*)$ is non-negative and irreducible, and by Lemma 3.10, $\rho(B_1(q_1^*)) \leq 1$. Thus, $(I - B_1(q_1^*))v \geq 0$. So

$$b \geq \left(B_1(q_1^*) - B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right) \right) v \geq 0. \quad (9)$$

To bound (from below) the maximum entry of b , it suffices to bound from below some entry of the matrix $B_1(q_1^*) - B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right)$, and the minimum entry v_{\min} of v . Since $q_2^* < q_1^*$ and $B_1(x)$ is dependent on x , there is some entry of $B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right)$, say $B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right)_{ij}$, which is strictly smaller than that of $B_1(q_1^*)_{ij}$. The entry $B_1(x)_{ij}$ must be of the form $(\sum_i c_i x_i) + (\sum_j c'_j y_j) + c''$, where for some k , $c_k > 0$ so that the term $c_k x_k$ depends on x_k . We must therefore have $B_1(q_1^*)_{ij} - \left(B \left(\frac{1}{2} (q_1^* + q_2^*) \right) \right)_{ij} \geq c_{\min} \frac{1}{2} (q_1^* - q_2^*)_k$, for some indices i, j, k . From inequality (7) we know that $\frac{(q_1^* - q_2^*)_k}{(q_1^* - q_2^*)_{\max}} \geq \alpha^n$, for all indices k . Thus, since $(q_1^* - q_2^*)_{\max} = \|q_1^* - q_2^*\|_\infty$, we have

$$\begin{aligned} B_1(q_1^*)_{ij} - B_1 \left(\frac{1}{2} (q_1^* + q_2^*) \right)_{ij} &\geq c_{\min} \frac{1}{2} (q_1^* - q_2^*)_k \\ &\geq c_{\min} \frac{1}{2} \alpha^n \|q_1^* - q_2^*\|_\infty \\ &\geq \alpha^{n+1} \frac{1}{2} \|q_1^* - q_2^*\|_\infty. \end{aligned} \quad (10)$$

Inequality (10) gives us a lower bound for a single entry of the non-negative matrix $\left(B_1(q_1^*) - B_1\left(\frac{1}{2}(q_1^* + q_2^*)\right)\right)$, namely the (i,j) th entry. In $\left(B_1(q_1^*) - B_1\left(\frac{1}{2}(q_1^* + q_2^*)\right)\right)v$, this (i,j) th entry is multiplied by a coordinate of v , which is at least v_{\min} . Thus, combining inequalities (10) and (9), we have $\|b\|_\infty \geq \alpha^{n+1} \frac{1}{2} \|q_1^* - q_2^*\|_\infty v_{\min}$. We can use Lemma 3.11 to bound v_{\min} .

LEMMA 3.13. *If v is the Perron vector of $B_1(q_1^*)$, normalized so that $v_{\max} = 1$, then $\frac{v_{\min}}{v_{\max}} = v_{\min} \geq \alpha^n$.*

PROOF. The vector v is the Perron vector of the (irreducible, nonnegative) matrix $B_1(q_1^*)$, it satisfies $B_1(q_1^*)v \leq v$ (since $\rho(B_1(q_1^*)) \leq 1$), and the smallest nonzero entry of $B_1(q_1^*)$ is at least α . Therefore, by Lemma 3.11, $v_{\min} \geq \alpha^n v_{\max} = \alpha^n$. \square

So $\|b\|_\infty \geq \alpha^{2n+1} \frac{1}{2} \|q_1^* - q_2^*\|_\infty$. Since the smallest nonzero entry of $A = B_1(\frac{1}{2}(q_1^* + q_2^*))$ is at least α , and since $\|b\|_\infty \geq \alpha^{2n+1} \frac{1}{2} \|q_1^* - q_2^*\|_\infty$, Lemma 3.12 now gives that

$$\left\| \left(I - B_1 \left(\frac{1}{2}(q_1^* + q_2^*) \right) \right)^{-1} \right\|_\infty \leq \frac{2n}{\alpha^{3n+1} \|q_1^* - q_2^*\|_\infty}. \quad (11)$$

Case 2: $x = P_1(x)$ is a linear system. Then, the matrix $B_1(x)$ is a constant matrix B_1 , independent on x . The matrix B_1 is nonnegative and irreducible, and $\rho(B_1) < 1$, because we already know from Lemma 3.10 that for all z such that $0 \leq z < q_1^*$, we have $\rho(B_1(z)) < 1$, but $B_1(z)$ is independent of z , because B_1 is a constant matrix.

Let us apply Lemma 3.3 with $a = q_1^*$, $b = 0$, and $P_1(x)$ in place of $P(x)$. We get $B_1 \cdot (q_1^* - 0) = P_1(q_1^*) - P(0)$. Multiplying both sides of this equation by -1 and then adding q_1^* to both sides, we get $(I - B_1)q_1^* = P_1(0)$, and thus $q_1^* = (I - B_1)^{-1}P_1(0)$. Since $q_1^* > 0$, we must have that $P_1(0) \neq 0$. But $P_1(0) \geq 0$. Indeed, $\|P_1(0)\|_\infty \geq c_{\min} \min\{1, y_{\min}^2\} \geq \alpha^2$. The smallest nonzero entry of B_1 is at least $c_{\min} \cdot \min\{1, y_{\min}\} \geq \alpha$. We now apply Lemma 3.12 to $A := B_1$ and $b := P_1(0)$, where we note that $(I - B_1)^{-1}P_1(0) = q_1^* \leq 1$. Lemma 3.12 thus gives

$$\left\| \left(I - B_1 \left(\frac{1}{2}(q_1^* + q_2^*) \right) \right)^{-1} \right\|_\infty \leq n\alpha^{-(n+2)}. \quad (12)$$

Since $\|q_1^* - q_2^*\|_\infty \leq 1$ ($q_1^* \leq 1$ and $q_2^* \geq 0$), and $0 < \alpha \leq 1$, and since $n \geq 1$, the upper bound (11) for the nonlinear case is worse than the upper bound (12) for the linear case, so the upper bound (11) holds in all cases.

This concludes Step (ii).

STEP (iii). To bound $\|P_1(q_2^*) - P_2(q_2^*)\|_\infty$, we use the following lemma, which holds for all points in $[0, 1]^n$, including q_2^* .

LEMMA 3.14. *If $0 \leq x \leq 1$, then $\|P_1(x) - P_2(x)\|_\infty \leq 2\|P(1, 1)\|_\infty \|y_1 - y_2\|_\infty$.*

PROOF. Since each entry of $P(x, y)$ is a quadratic polynomial, for each $b \in \{1, 2\}$ and each $d \in \{1, \dots, n\}$, the d th coordinate, $(P_b(x))_d$, of $P_b(x) = P(x, y_b)$ has the form

$$\sum_{i,j} a_{d,i,j} x_i x_j + \sum_{i,j} c_{d,i,j} y_b y_{b,j} + \sum_{i,j} c'_{d,i,j} x_i y_{b,j} + \sum_k a'_{d,k} x_k + \sum_k c''_{d,k} y_{b,k} + c'''_d,$$

where $y_{b,j}$ refers to the j th coordinate of the m -vector y_b , and where all the coefficients $a_{d,i,j}$, $c_{d,i,j}$, $c'_{d,i,j}$, $c''_{d,k}$ and c'''_d , are nonnegative. Also, recall $0 < y_1 \leq 1$ and $0 \leq y_2 \leq y_1$. Thus, we have the following series of inequalities. (For the first inequality, note that $y_{1,i}y_{1,j} - y_{2,i}y_{2,j} = y_{1,j}(y_{1,i} - y_{2,i}) + y_{2,i}(y_{1,j} - y_{2,j}) \leq (y_{1,i} - y_{2,i}) + (y_{1,j} - y_{2,j})$ and $x_i \leq 1$; the other inequalities are straightforward.)

$$\begin{aligned}
& \|P_1(x) - P_2(x)\|_\infty \\
&= \max_d \sum_{i,j} c_{d,i,j} (y_{1,i}y_{1,j} - y_{2,i}y_{2,j}) + \sum_{i,j} c'_{d,i,j} x_i (y_{1,j} - y_{2,j}) + \sum_k c''_{d,k} (y_{1,k} - y_{2,k}) \\
&\leq \max_d \sum_{i,j} c_{d,i,j} ((y_{1,i} - y_{2,i}) + (y_{1,j} - y_{2,j})) + \sum_{i,j} c'_{d,i,j} (y_{1,j} - y_{2,j}) + \sum_k c''_{d,k} (y_{1,k} - y_{2,k}) \\
&\leq \max_d \sum_{i,j} 2 \cdot c_{d,i,j} \cdot \|y_1 - y_2\|_\infty + \sum_{i,j} c'_{d,i,j} \cdot \|y_1 - y_2\|_\infty + \sum_k c''_{d,k} \|y_1 - y_2\|_\infty \\
&= \left(\max_d \sum_{i,j} 2c_{d,i,j} + c'_{d,i,j} + c''_{d,k} \right) \cdot \|y_1 - y_2\|_\infty \\
&\leq 2\|P(1, 1)\|_\infty \|y_1 - y_2\|_\infty
\end{aligned}$$

□

STEP (iv). We combine the results of the previous steps to conclude the proof. From Eq. (6) established in Step (i), taking norms on both sides, yields

$$\|q_1^* - q_2^*\|_\infty \leq \left\| \left(I - B_1 \left(\frac{1}{2}(q_1^* + q_2^*) \right) \right)^{-1} \right\|_\infty \|P_1(q_2^*) - P_2(q_2^*)\|_\infty. \quad (13)$$

Inserting our upper bound (11) for $\|(I - B_1(\frac{1}{2}(q_1^* + q_2^*)))^{-1}\|_\infty$, which holds in all cases (whether $x = P_1(x)$ is linear or nonlinear), gives

$$\|q_1^* - q_2^*\|_\infty \leq \frac{2n}{\alpha^{3n+1} \|q_1^* - q_2^*\|_\infty} \|P_1(q_2^*) - P_2(q_2^*)\|_\infty.$$

We now move the $\|q_1^* - q_2^*\|_\infty$ terms to the left and take square roots to obtain

$$\|q_1^* - q_2^*\|_\infty \leq \sqrt{2n\alpha^{-(3n+1)} \|P_1(q_2^*) - P_2(q_2^*)\|_\infty}. \quad (14)$$

Using Lemma 3.14, we have,

$$\|q_1^* - q_2^*\|_\infty \leq \sqrt{4n\alpha^{-(3n+1)} \|P(1, 1)\|_\infty \|y_2 - y_1\|_\infty},$$

which completes the proof of the first inequality of part 1 of Theorem 3.1.

If $x = P_1(x)$ is a linear system, then we can substitute the bound (12) into the inequality (13). This gives

$$\|q_1^* - q_2^*\|_\infty \leq n\alpha^{-(n+2)} \|P_1(q_2^*) - P_2(q_2^*)\|_\infty.$$

Again, applying Lemma 3.14, we get

$$\|q_1^* - q_2^*\|_\infty \leq n\alpha^{-(n+2)} 2\|P(1, 1)\|_\infty \|y_1 - y_2\|_\infty.$$

This completes the proof of part 1 of Theorem 3.1. □

We will next establish part 2 of Theorem 3.1.

PROOF OF PART 2 OF THEOREM 3.1. The proof consists of two steps: In step (i), we will show that, starting from $x^{[0]} := 0$, all the iterations of R-NM, applied to $x = P_2(x)$ are defined. In step (ii), we will prove the upper bound on the convergence rate of R-NM.

STEP (i). We will show, by induction on k , that $x^{[k]}$ is defined, $0 \leq x^{[k]} \leq q_2^*$ and $\rho(B_2(x^{[k]})) < 1$.

For the base case, it is clear that $x^{[0]} = 0 \leq q_2^*$. Furthermore, $B_2(0) \leq B_1(0)$, hence $\rho(B_2(0)) \leq \rho(B_1(0))$. Since $0 < q_1^*$, Lemma 3.10 implies that $\rho(B_1(0)) < 1$. Therefore, $\rho(B_2(0)) < 1$.

For the induction step, assume that $0 \leq x^{[k]} \leq q_2^*$ and $\rho(B_2(x^{[k]})) < 1$. We will show first that $\mathcal{N}_{P_2}(x^{[k]})$ is well defined and $0 \leq x^{[k+1]} \leq q_2^*$. Since $\rho(B_2(x^{[k]})) < 1$, by Lemma 3.8, $(I - B_2(x^{[k]}))$ is nonsingular and so $\mathcal{N}_{P_2}(x^{[k]})$ is well defined. Lemma 3.8 also gives that $(I - B_2(x^{[k]}))^{-1} = \sum_{i=0}^{\infty} B_2(x^{[k]})^i \geq 0$. Lemma 3.4 yields that

$$q_2^* - \mathcal{N}_{P_2}(x^{[k]}) = (I - B_2(x^{[k]}))^{-1} \frac{B_2(q_2^*) - B_2(x^{[k]})}{2} (q_2^* - x^{[k]}).$$

Note that $(q_2^* - x^{[k]}) \geq 0$, thus $B_2(q_2^*) - B_2(x^{[k]}) \geq 0$, and we have just shown that $(I - B_2(x^{[k]}))^{-1} \geq 0$. So all the terms on the right of this equation are nonnegative, and thus $q_2^* - \mathcal{N}_{P_2}(x^{[k]}) \geq 0$. That is, $q_2^* \geq \mathcal{N}_{P_2}(x^{[k]})$. The point $x^{[k+1]}$ is defined by rounding down $\mathcal{N}_{P_2}(x^{[k]})$ and maintaining nonnegativity, thus for all coordinates i , either $x_i^{[k+1]} = 0$, in which case trivially we have $x_i^{[k+1]} = 0 \leq (q_2^*)_i$, or else $0 \leq x_i^{[k+1]} \leq \mathcal{N}_{P_2}(x^{[k]})_i \leq (q_2^*)_i$. Thus, $x^{[k+1]} \leq q_2^*$.

It remains to show that $\rho(B_2(x^{[k+1]})) < 1$. We will prove $\rho(B_2(x^{[k+1]})) < 1$ by considering separately the cases where $P_1(x)$ contains nonlinear or only linear polynomials. We examine first the nonlinear case. We need the following lemma.

LEMMA 3.15. *If $x = P(x)$ is a strongly connected quadratic MPS with n variables, with LFP $q^* > 0$, and there is some nonlinear quadratic term in some polynomial $P_i(x)$, then if $0 \leq z < q^*$, then $\mathcal{N}_P(z)$ is defined and $\mathcal{N}_P(z) < q^*$.*

PROOF. Lemma 3.10 tells us that $\rho(B(q^*)) \leq 1$. Nonlinearity of $P(x)$ means that $B(x)$ does depend on x . That is, some entry of $B(x)$ contains a term of the form cx_i for some x_i with $c > 0$. So $B(z) \neq B(q^*)$, and $B(z) \leq B(q^*)$ since B is monotone. Since $x = P(x)$ is strongly connected and $q^* > 0$, Lemma 3.10 yields that $\rho(B(z)) < 1$. By Lemma 3.8, $(I - B(z))$ is nonsingular and so the Newton iterate $\mathcal{N}_P(z)$ is well defined. Consider the equation given by Lemma 3.4:

$$q^* - \mathcal{N}_P(z) = (I - B(z))^{-1} \frac{B(q^*) - B(z)}{2} (q^* - z).$$

We know that $q^* - z > 0$, and thus $B(q^*) - B(z) \geq 0$. Since $\rho(B(z)) < 1$, by Lemma 3.8, $(I - B(z))^{-1} = \sum_{k=0}^{\infty} B(z)^k \geq 0$. This and Lemma 3.4 is already enough to yield that $q^* - \mathcal{N}_P(z) \geq 0$, and we just need to show that this is a strict inequality.

We first show that if $P_i(x)$ contains a term of degree 2, then $\left(\frac{B(q^*) - B(z)}{2} (q^* - z) \right)_i > 0$. This term of degree 2 must be of the form $cx_j x_k$ for some j, k . Then, $B(x)_{i,j}$ has a term cx_k with $c > 0$ and so $(B(q^*) - B(z))_{i,j} \geq c(q^* - z)_k$. But then $\left(\frac{B(q^*) - B(z)}{2} (q^* - z) \right)_i \geq c(q^* - z)_k (q^* - z)_j > 0$.

Now we will show that for all $i \in \{1, \dots, n\}$, $(q - \mathcal{N}_P(z))_i > 0$. If $P_i(x)$ contains a term of degree 2, then we have just shown that $\left(\frac{B(q^*)-B(z)}{2}(q^* - z)\right)_i > 0$. But $(I - B(z))^{-1} = \sum_{k=0}^{\infty} B(z)^k \geq I$. So $(q - \mathcal{N}_P(z))_i \geq \left(\frac{B(q^*)-B(z)}{2}(q^* - z)\right)_i > 0$. If $P_i(x)$ does not contain a term of degree 2, there must be some other x_j with $P_j(x)$ containing a term of degree 2 and, since $x = P(x)$ is strongly connected, x_i depends on x_j , possibly indirectly. That is, there is a sequence of variables i_0, i_1, \dots, i_l with $l < n$, $i_0 = i$, $i_l = j$, and for each $0 < m \leq l$, the variable x_{i_m} appears in a term of $P(x)_{i_{m-1}}$. Let k be the least integer such that $P(x)_{i_k}$ contains a term of degree 2. Then, if $0 < m \leq k$, x_{i_m} appears in a degree 1 term in $P(x)_{i_{m-1}}$, that is one of the form $c_m x_{i_m}$ with $c_m > 0$. So $B(x)_{i_{m-1}, i_m}$ contains the constant term $c_m > 0$. So $B(z)_{i_{m-1}, i_m} \geq c_m > 0$. So $B^k(z)_{i, i_k} \geq \prod_{m=0}^{k-1} B(z)_{i_m, i_{m+1}} \geq \prod_{m=0}^{k-1} c_m > 0$. Since $P(x)_{i_k}$ contains a term of degree 2, from above $\left(\frac{B(q^*)-B(z)}{2}(q^* - z)\right)_{i_k} > 0$. So $\left(B^k(z) \frac{B(q^*)-B(z)}{2}(q^* - z)\right)_i > 0$. But $q^* - \mathcal{N}_P(z) = (I - B(z))^{-1} \frac{B(q^*)-B(z)}{2}(q^* - z) = \left(\sum_{m=0}^{\infty} B^m(z)\right) \frac{B(q^*)-B(z)}{2}(q^* - z) \geq B^k(z) \frac{B(q^*)-B(z)}{2}(q^* - z)$. So $(q^* - \mathcal{N}_P(z))_i > 0$ for all i , as required. \square

We will only actually need to apply Lemma 3.15 in the case when $q_2^* = q_1^*$ and $x = P_1(x)$ is nonlinear. Suppose that $q_1^* = q_2^*$ and some polynomial in $P_1(x)$ is nonlinear in x . We claim that then $P_1(x) \equiv P_2(x)$. That is, for all those variables in y , say $(y)_j$, that actually appear in some polynomials in $P(x, y)$, it must be the case that $(y_1)_j = (y_2)_j$. Otherwise, if there is some variable $(y)_j$ with $(y_2)_j < (y_1)_j$ such that $(y)_j$ appears in $P_i(x, y)$, then $(P_2(q_1^*))_i = (P(q_1^*, y_2))_i < (P(q_1^*, y_1))_i = (q_1^*)_i$, so q_1^* is not a fixed point of $P_2(x)$, contradicting that $q_1^* = q_2^*$. Thus, if $x = P_1(x)$ is nonlinear and $q_1^* = q_2^*$ then $x = P_2(x)$ is also nonlinear and $q_2^* = q_1^* > 0$, so we can use Lemma 3.15, which shows that if $0 \leq x^{[k]} < q_2^*$, then $\mathcal{N}_{P_2}(x^{[k]}) < q_2^*$ and so $0 \leq x^{[k+1]} < q_2^* \leq q_1^*$. Since $x^{[k+1]} < q_1^*$, we have $\rho(B_1(x^{[k+1]})) < 1$. Since $B_2(x^{[k+1]}) \leq B_1(x^{[k+1]})$, we also have $\rho(B_2(x^{[k+1]})) < 1$.

This leaves us with two cases remaining to show that $\rho(B_2(x^{[k+1]})) < 1$: first, the case where $x = P_1(x)$ is nonlinear and $q_2^* \neq q_1^*$, and second, the case where $x = P_1(x)$ is linear or constant. We have shown already that $0 \leq x^{[k+1]} \leq q_2^*$. It thus suffices to show that in these cases, for any $0 \leq z \leq q_2^*$, $\rho(B_2(z)) < 1$.

For the first case, suppose that $q_2^* \neq q_1^*$ and that $x = P_1(x)$ is nonlinear, and thus $B_1(x)$ depends on x . Then, we have previously argued that $q_2^* < q_1^*$ (see inequality (8)). But then $B_1(q_2^*) \neq B_1(q_1^*)$. For any $0 \leq z \leq q_2^*$, $B_2(z) \leq \tilde{B}_2(q_2^*) \leq B_1(q_2^*) \leq B_1(q_1^*)$ and because $B_1(q_2^*) \neq B_1(q_1^*)$, we have $B_2(z) \neq B_1(q_1^*)$. But $B_1(q_1^*)$ is irreducible, and Lemma 3.7 then tells us that $\rho(B_2(z)) < \rho(B_1(q_1^*))$. We know, by Lemma 3.10, that $\rho(B_1(q_1^*)) \leq 1$. So $\rho(B_2(z)) < 1$.

For the second case, suppose that $x = P_1(x)$ is linear. Then, $B_1(x)$ is a constant matrix. Thus, $B_1(z) = B_1(0)$ for all $0 \leq z$. But Lemma 3.10 tells us that, since $0 < q_1^*$, $\rho(B_1(0)) < 1$. Thus, $\rho(B_1(z)) < 1$ for all $0 \leq z \leq q_2^*$. Since $0 \leq B_2(z) \leq B_1(z)$, we have $\rho(B_2(z)) < 1$ for all $0 \leq z \leq q_2^*$. This concludes the proof of the inductive step.

Thus, the R-NM iterations applied to $x = P_2(x)$ are defined in all cases, and yield iterates $0 \leq x^{[k]} \leq q^*$, for all $k \geq 0$.

STEP (ii). We now prove the upper bound on the rate of convergence for R-NM applied to $x = P_2(x)$.

LEMMA 3.16. *Suppose an MPS, $x = P(x)$, with n variables has LFP $0 \leq q^* \leq 1$, and for some n -vector $v > 0$, we have $B(q^*)v \leq v$. Suppose we perform $g \geq h - 1$ iterations*

of R-NM with parameter $h \geq 2 + \lceil \log \frac{v_{\max}}{v_{\min} \cdot \epsilon} \rceil$ on the MPS $x = P(x)$, and suppose that, for all $k \geq 0$, every iteration $x^{[k]}$ is defined and $0 \leq x^{[k]} \leq q^*$. Then $\|q^* - x^{[g]}\|_\infty \leq \epsilon$.

PROOF. By induction on k , we claim that $\forall k \geq 0, q^* - x^{[k]} \leq (2^{-k} + 2^{-h+1}) \frac{1}{v_{\min}} v$. Note that this would indeed yield the lemma: for all $k, 0 \leq x^{[k]} \leq q^*$, and the claim would yield $q^* - x^{[g]} \leq (2^{-h+1} + 2^{-h+1}) \frac{1}{v_{\min}} v \leq 2^{-\log \frac{v_{\max}}{v_{\min} \cdot \epsilon}} \frac{1}{v_{\min}} v = \epsilon \frac{1}{v_{\max}} v \leq \epsilon 1$.

It remains to prove by induction on $k \geq 0$ that $q^* - x^{[k]} \leq (2^{-k} + 2^{-h+1}) \frac{1}{v_{\min}} v$. This is true for $k = 0$, because $q^* \geq 0 = x^{[0]}$, and $q^* - x^{[0]} = q^* \leq 1 \leq \frac{1}{v_{\min}} v$.

Lemma 3.5 then gives that $q^* - \mathcal{N}_P(x^{[k]}) \leq (2^{-(k+1)} + 2^{-h}) \frac{1}{v_{\min}} v$. Now, by definition of $x^{[k+1]}$, $\mathcal{N}_P(x^{[k]}) - x^{[k+1]} \leq 2^{-h} 1 \leq 2^{-h} \frac{1}{v_{\min}} v$. So $q^* - x^{[k+1]} \leq (2^{-(k+1)} + 2^{-h+1}) \frac{1}{v_{\min}} v$ as required. \square

To use Lemma 3.16 to get a bound on using R-NM on $x = P_2(x)$ to compute q_2^* , note that because $0 \leq B_2(q_2^*) \leq B_1(q_1^*)$, the Perron vector $v > 0$ of $B_1(q_1^*)$, which satisfies $B_1(q_1^*)v \leq v$, must also satisfy $B_2(q_2^*)v \leq v$.

Thus, if we perform $g \geq h - 1$ iterations of R-NM on $x = P_2(x)$, with parameter $h \geq \left\lceil 2 + \log(\alpha^{-n} \epsilon^{-1}) \right\rceil \geq \left\lceil 2 + \log \frac{v_{\max}}{v_{\min} \cdot \epsilon} \right\rceil$, we obtain that $\|q_2^* - x^{[h-1]}\|_\infty \leq \epsilon$. (Here we have used the fact, from Lemma 3.13, that for the Perron vector v , with $v_{\max} = 1$, $\frac{v_{\max}}{v_{\min}} = \frac{1}{v_{\min}} \leq \alpha^{-n}$.) This completes the proof of Theorem 3.1. \square

4. GENERAL MONOTONE POLYNOMIAL SYSTEMS

In this section, we use the rounded-down decomposed Newton's method (R-DNM), to compute the LFP q^* of general MPSs. First, we consider the case where $0 < q^* \leq 1$.

THEOREM 4.1. *For all ϵ , where $0 < \epsilon < 1$, if $x = P(x)$ is an MPS with LFP solution $0 < q^* \leq 1$, with $q_{\min}^* = \min_i q_i^*$, and the minimum nonzero coefficient or constant in $P(x)$ is c_{\min} , then rounded-down decomposed Newton's method (R-DNM) with parameter*

$$h \geq \left\lceil 3 + 2^f \cdot \left(\log\left(\frac{1}{\epsilon}\right) + d \cdot \left(\log(\alpha^{-(4n+1)}) + \log(16n) + \log(\|P(1)\|_\infty) \right) \right) \right\rceil,$$

using $g \geq h - 1$ iterations for every nonlinear SCC (and 1 iteration for linear SCC), gives an approximation \tilde{q} to q^ with $\tilde{q} \leq q^*$ and such that $\|q^* - \tilde{q}\|_\infty \leq \epsilon$.*

Here d denotes the maximum depth of SCCs in the DAG H_P of SCCs of the MPS $x = P(x)$, f is the nonlinear depth, and $\alpha = \min\{1, c_{\min}\} \cdot \frac{1}{2} q_{\min}^$.*

Before proving the theorem, let us note that we can obtain worst-case expressions for the needed number of iterations $g = h - 1$, and the needed rounding parameter h , in terms of only $f \leq d \leq n \leq |P|$, and ϵ , by noting that $\log(\|P(1)\|_\infty) \leq |P|$, and by appealing to Theorem 2.4 to remove references to q_{\min}^* in the bounds. Noting that $c_{\min} \geq 2^{-|P|}$, these tell us that $\alpha = \min\{1, c_{\min}\} \frac{1}{2} q_{\min}^* \geq 2^{-|P|} 2^{n-1}$. Substituting, we obtain that any

$$g \geq \left\lceil 2 + 2^f \cdot \left(\log\left(\frac{1}{\epsilon}\right) + d \cdot (|P| 2^n (4n + 1) + (4n + 1) + \log(16n) + |P|) \right) \right\rceil \quad (15)$$

iterations suffice in the worst case, with rounding parameter $h = g + 1$. Thus, for $i = \log(1/\epsilon)$ bits of precision, $g = k_P + c_P \cdot i$ iterations suffice, where $c_P = 2^f$ and $k_P = O(2^f 2^n n d |P|)$, with tame constants in the big-O.

PROOF OF THEOREM 4.1. For every SCC S , its *height* h_S (respectively, *nonlinear height* f_S) is the maximum over all paths of the DAG H_P of SCCs starting at S , of the number of SCCs (respectively, nonlinear SCCs) on the path. We show by induction on the height h_S of each SCC S that $\|q_S^* - \tilde{q}_S\|_\infty \leq \beta^{h_S} \delta^{2^{-f_S}}$ where $\beta = 16n\alpha^{-(3n+1)} \|P(1)\|_\infty$ and $\delta = \left(\frac{\epsilon}{\beta^d}\right)^{2^f}$. Note that since $n \geq 1$, $\epsilon < 1$, and $\alpha \leq c_{\min}$, we have $\beta \geq 1$ and $\delta \leq 1$, and thus also $\delta \leq \sqrt{\delta}$.

Let us first check that this would imply the theorem. For all SCCs, S , we have $1 \leq h_S \leq d$ and $0 \leq f_S \leq f$, and thus $\|q_S^* - \tilde{q}_S\|_\infty \leq \beta^{h_S} \delta^{2^{-f_S}} \leq \beta^d \delta^{2^{-f}} = \beta^d \left(\frac{\epsilon}{\beta^d}\right) = \epsilon$.

We note that h is related to δ by the following:

$$h \geq 2 + n \log \frac{1}{\alpha} + \log \frac{2}{\delta}. \quad (16)$$

This is because $\log \frac{2}{\delta} = 1 + \log \frac{1}{\delta} = 1 + 2^f (\log \frac{1}{\epsilon} + d \log \beta) = 1 + 2^f \left(\log \left(\frac{1}{\epsilon} \right) + d \log(16n\alpha^{-(3n+1)} \|P(1)\|_\infty) \right)$. Note that (16) implies that this inequality holds also for any subsystem of $x = P(x)$ induced by a SCC S and its successors $D(S)$ because the parameters n and $1/\alpha$ for a subsystem are no larger than those for the whole system.

We now prove by induction on h_S that $\|q_S^* - \tilde{q}_S\|_\infty \leq \beta^{h_S} \delta^{2^{-f_S}}$.

In the base case, $h_S = 1$, we have a strongly connected MPS $x_S = P_S(x)$. If S is linear, we solve the linear system exactly and then round down to a multiple of 2^{-h} . Then $f_S = 0$, and we have to show $\|q_S^* - \tilde{q}_S\|_\infty \leq \beta^{h_S} \delta^{2^{-f_S}} = \beta \delta$. But $\|q_S^* - \tilde{q}_S\|_\infty \leq 2^{-h} \leq \frac{\delta}{2} \leq \beta \delta$.

For the base case where S is nonlinear, Eq. (16) and Corollary 3.2 imply that $\|q_S^* - \tilde{q}_S\|_\infty \leq \frac{\delta}{2}$, which implies the claim since $\delta \leq 1$ and $\beta \geq 1$, hence $\frac{\delta}{2} \leq \beta^{h_S} \delta^{2^{-f_S}} = \beta^1 \delta^{2^{-1}}$.

Inductively, consider an SCC S with $h_S > 1$. Then S depends only on SCCs with height at most $h_S - 1$. If S is linear, it depends on SCCs of nonlinear depth at most $f_{D(S)} = f_S$, whereas if S is nonlinear, it depends on SCCs of nonlinear depth at most $f_{D(S)} = f_S - 1$. We can assume by inductive hypothesis that $\|q_{D(S)}^* - \tilde{q}_{D(S)}\|_\infty \leq \beta^{h_S-1} \delta^{2^{-f_{D(S)}}}$. Take q'_S to be the LFP of $x_S = P_S(x_S, \tilde{q}_{D(S)})$.

Suppose $x_S = P_S(x_S, q_{D(S)}^*)$ is linear in x_S . Then Theorem 3.1 with $y_1 := q_{D(S)}^*$ and $y_2 := \tilde{q}_{D(S)}$, yields

$$\|q_S^* - q'_S\|_\infty \leq 2n_S \alpha^{-(n_S+2)} \|P(1, 1)\|_\infty \|q_{D(S)}^* - \tilde{q}_{D(S)}\|_\infty.$$

But $2n_S \alpha^{-(n_S+2)} \|P(1, 1)\|_\infty \leq \frac{\beta}{2}$, so $\|q_S^* - q'_S\|_\infty \leq \frac{\beta}{2} \|q_{D(S)}^* - \tilde{q}_{D(S)}\|_\infty \leq \frac{\beta}{2} \beta^{h_S-1} \delta^{2^{-f_S}} = \frac{1}{2} \beta^{h_S} \delta^{2^{-f_S}}$. Since $\|q'_S - \tilde{q}_S\|_\infty \leq 2^{-h} \leq \frac{\delta}{2} \leq \frac{1}{2} \beta^{h_S} \delta^{2^{-f_S}}$, it follows that $\|q_S^* - \tilde{q}_S\|_\infty \leq \beta^{h_S} \delta^{2^{-f_S}}$.

Suppose that $x_S = P_S(x_S, q_{D(S)}^*)$ is nonlinear in x_S . Theorem 3.1, with $y_1 := q_{D(S)}^*$ and $y_2 := \tilde{q}_{D(S)}$, yields that

$$\|q_S^* - q'_S\|_\infty \leq \sqrt{4n\alpha^{-(3n+1)} \|P(1)\|_\infty \|q_{D(S)}^* - (\tilde{q})_{D(S)}\|_\infty}. \quad (17)$$

Note that the α from Theorem 3.1 is indeed the same or better (i.e., bigger) than the α in this theorem, because $y_{\min} = (q_{D(S)}^*)_{\min} \geq q_{\min}^*$ and $(q_S^*)_{\min} \geq q_{\min}^*$. Rewriting (17) in terms of β , we have $\|q_S^* - q'_S\|_\infty \leq \sqrt{\frac{1}{4} \beta \|q_{D(S)}^* - (\tilde{q})_{D(S)}\|_\infty}$. By inductive assumption,

$\|q_{D(S)}^* - \tilde{q}_{D(S)}\|_\infty \leq \beta^{h_S-1} \delta^{2^{-f_S}+1}$, and thus $\|q_S^* - q'_S\|_\infty \leq \sqrt{\frac{1}{4} \beta^{h_S} \delta^{2^{1-f_S}}} \leq \frac{1}{2} \beta^{h_S} \delta^{2^{-f_S}}$. Thus, to show that the inductive hypothesis holds also for SCC S , it suffices to show that, for the approximation \tilde{q}_S , we have $\|q'_S - \tilde{q}_S\|_\infty \leq \frac{1}{2} \beta^{h_S} \delta^{2^{-f_S}}$. But $\beta \geq 1$, $h_S \geq 1$, $2^{-f_S} \leq 1$ and $\delta \leq 1$, so $\frac{1}{2} \delta \leq \frac{1}{2} \beta^{h_S} \delta^{2^{-f_S}}$, so it suffices to show that $\|q'_S - \tilde{q}_S\|_\infty \leq \frac{1}{2} \delta$. Part 2 of Theorem 3.1 tells us that we will have $\|q'_S - \tilde{q}_S\|_\infty \leq \frac{1}{2} \delta$ if $g \geq h - 1$ and $h \geq \lceil 2 + n \log \frac{1}{\alpha} + \log \frac{2}{\delta} \rceil$. But (since h is an integer) we have already established this in Eq. (16), hence the claim follows. \square

Next, we want to generalize Theorem 4.1 to arbitrary MPSs that have an LFP, $q^* > 0$, without the restriction that $0 < q^* \leq 1$. The next lemma allows us to establish this by a suitable “rescaling” of any MPS which has an LFP $q^* > 0$. If $x = P(x)$ is a MPS and $c > 0$, we can consider the MPS $x = \frac{1}{c}P(cx)$.

LEMMA 4.2. *Let $x = P(x)$ be a MPS with LFP solution q^* , and with Jacobian $B(x)$, and recall that for $z \geq 0$, $\mathcal{N}_P(z) := z + (I - B(z))^{-1}(P(z) - z)$ denotes the Newton operator applied at z on $x = P(x)$. Then*

- (i) *the LFP solution of $x = \frac{1}{c}P(cx)$ is $\frac{1}{c}q^*$;*
- (ii) *the Jacobian of $\frac{1}{c}P(cx)$ is $B(cx)$;*
- (iii) *a Newton iteration of the “rescaled” MPS, $x = \frac{1}{c}P(cx)$, applied to the vector z is given by $\frac{1}{c}\mathcal{N}_P(cz)$.*

PROOF. From Etessami and Yannakakis [2009], we know that the value iteration sequence $P(0), P(P(0)), P(P(P(0))) \dots P^k(0)$ converges to q^* . Now note that for the MPS $x = \frac{1}{c}P(cx)$, the value iteration sequence is $\frac{1}{c}P(0), \frac{1}{c}P(c\frac{1}{c}P(0)) = \frac{1}{c}P(P(0)), \frac{1}{c}P(P(P(0))) \dots$ which thus converges to $\frac{1}{c}q^*$. This establishes (i).

For (ii), note that, by the chain rule in multivariate calculus (see, e.g., Apostol [1974, Section 12.10]), the Jacobian of $P(cx)$ is $cB(cx)$. Now (iii) follows because

$$z + (I - B(cz))^{-1} \left(\frac{1}{c}P(cz) - z \right) = \frac{1}{c}(cz + (I - B(cz))^{-1}(P(cz) - cz)) = \frac{1}{c}\mathcal{N}_P(cz). \quad \square$$

We use Lemma 4.2 to generalize Theorem 4.1 to MPSs with LFP q^* , where q^* does not satisfy $q^* \leq 1$.

THEOREM 4.3. *If $x = Q(x)$ is an MPS with n variables, with LFP solution $q^* > 0$, if c'_{\min} is the least positive coefficient of any monomial in $Q(x)$, then R-DNM with rounding parameter h' , and using g' iterations per nonlinear SCC (and one for linear), gives an approximation \tilde{q} such that $\|q^* - \tilde{q}\|_\infty \leq \epsilon'$, where*

$$g' = 2 + \lceil 2^f \cdot (\log(\frac{1}{\epsilon'}) + d \cdot (2u + \log(\alpha'^{-(4n+1)}) + \log(16n) + \log(\|Q(1)\|_\infty)) \rceil$$

and $h' = g' + 1 - u$, where $u = \max\{0, \lceil \log q_{\max}^* \rceil\}$, d is the maximum depth of SCCs in the DAG H_Q of SCCs of $x = Q(x)$, f is the nonlinear depth, and $\alpha' = 2^{-2u} \min\{1, c'_{\min}\} \min\{1, \frac{1}{2}q_{\min}^*\}$.

We can again obtain worst-case expressions for the needed number of iterations g' , and the needed rounding parameter h' , in terms of only $f \leq d \leq n \leq |Q|$, and ϵ' , by noting that $\log(\|Q(1)\|_\infty) \leq |Q|$ and by appealing to Theorem 2.4 to remove references to q_{\min}^* and q_{\max}^* in the bounds. Substituting and simplifying, we get that to guarantee additive error at most ϵ' , that is, for $i = \log(1/\epsilon')$ bits of precision, it suffices in the

worst-case to apply $g' = k_Q + c_Q \cdot i$ iterations of R-DNM with rounding parameter $h' = g' + 1$ (which is more accurate rounding than $h' = g' + 1 - u$), where $c_Q = 2^f$, and $k_Q = O(2^f 5^n n^2 d(|Q| + n \log n))$ (and we can calculate precise, tame, constants for the big-O expression).

COROLLARY 4.4. *If $x = P(x)$ is an MPS with LFP solution q^* with $0 < q_{\min}^* \leq q_i^* \leq q_{\max}^*$ for all i , with the least coefficient of any monomial in $P(x)$, c_{\min} , with f the nonlinear depth of the DAG of SCCs of $x = P(x)$ and with encoding size $|P|$ bits, we can compute an approximation \tilde{q} to q^* with $\|q^* - \tilde{q}\|_{\infty} \leq \epsilon$, for any given $0 < \epsilon \leq 1$, in time polynomial in $|P|, 2^f, \log \frac{1}{\epsilon}, \log \frac{1}{q_{\min}^*}$ and $\log q_{\max}^*$.*

PROOF OF COROLLARY 4.4. After preprocessing to remove all variables x_i with $q_i^* = 0$, which takes P-time in $|P|$, we use R-DNM as specified in Theorem 4.3. Calculating a Newton iterate at z is just a matter of solving a matrix equation and if the coordinates of z are multiples of 2^{-h} this can be done in time polynomial in $|P|$ and h . Theorem 4.3 tells us that the number of iterations and h are polynomial in $2^f, \log \frac{1}{\epsilon}, \log \frac{1}{q_{\min}^*}, \log q_{\max}^*, n, \log \frac{1}{c_{\min}}$ and $\log \|P(1)\|_{\infty}$. The last three of these are bounded by $|P|$. Together, these give the corollary. \square

PROOF OF THEOREM 4.3. If $q_{\max}^* \leq 1$, then Theorem 4.1 gives this immediately. So we assume that $q_{\max}^* > 1$. u is chosen so that $2^u \geq q_{\max}^*$. We rescale and use Lemma 4.2 with scaling parameter $c = 2^u$. This yields the “rescaled” MPS $x = 2^{-u}Q(2^u x)$, which has LFP $p^* = 2^{-u}q^* \leq 1$.

So we can apply Theorem 4.1 to this rescaled MPS $x = P(x)$, where $P(x) \equiv 2^{-u}Q(2^u x)$, and letting $\epsilon := 2^{-u}\epsilon'$. Then Theorem 4.1 gives us the needed number of iterations g and the rounding parameter $h = g + 1$, needed to obtain an approximation \tilde{p} of the LFP $p^* = 2^{-u}q^*$, such that $\|\tilde{p} - p^*\|_{\infty} \leq \epsilon$.

In the bounds specified for Theorem 4.1 for g and h , in place of q_{\min}^* we get $p_{\min}^* = 2^{-u}q_{\min}^*$, and in place of c_{\min} we get $2^{-u}c'_{\min}$. Thus α becomes the α' we have specified in the statement of this theorem. Furthermore, the $\|P(1)\|_{\infty}$ appearing in Theorem 4.1 is now $\|2^{-u}Q(2^u 1)\|_{\infty}$, but it is easy to verify that for a quadratic MPS, $\|2^{-u}Q(2^u 1)\|_{\infty} \leq 2^u \|Q(1)\|_{\infty}$.

Theorem 4.1 tells us that if we use R-DNM on $x = P(x)$ for g iterations per nonlinear SCC and a precision of $h = g + 1$ bits, we will obtain an approximation \tilde{p} to the LFP p^* of $x = P(x)$ with $\|\tilde{p} - p^*\|_{\infty} \leq \epsilon$ provided that $h \geq \lceil 3 + 2^f \cdot (\log(\frac{1}{\epsilon}) + d \cdot (\log(\alpha^{-(4n+1)}) + \log(16n) + \log(\|P(1)\|_{\infty}))) \rceil$. This condition is satisfied if we take $g = g'$ and $h = g' + 1$ because

$$\begin{aligned} & \lceil 3 + 2^f \cdot (\log(\frac{1}{\epsilon}) + d \cdot (\log(\alpha^{-(4n+1)}) + \log(16n) + \log(\|P(1)\|_{\infty}))) \rceil \\ & \leq 3 + 2^f (\log(\frac{1}{2^{-u}\epsilon'}) + d(\log(\alpha'^{-(4n+1)}) + \log(16n) + \log(2^u \|Q(1)\|_{\infty}))) \\ & = 3 + 2^f (u + \log(\frac{1}{\epsilon'}) + d(\log(\alpha'^{-(4n+1)}) + \log(16n) + u + \log(\|Q(1)\|_{\infty}))) \\ & \leq g' + 1 = h. \end{aligned}$$

Thus, applying R-DNM on $x = P(x)$ with parameters $g = g'$ and $h = g' + 1$ yields an approximation \tilde{p} to the LFP p^* of $x = P(x)$ with $\|\tilde{p} - p^*\|_{\infty} \leq \epsilon$ or, in terms of the original MPS, $\|\tilde{p} - 2^{-u}q^*\|_{\infty} \leq 2^{-u}\epsilon'$.

To obtain Theorem 4.3, we now show that if we apply R-DNM to $x = Q(x)$ with LFP q^* , using rounding parameter h' and using g' iterations per nonlinear SCC (where h'

and g' were specified in the statement of the theorem), we will obtain an approximation \tilde{q} to q^* that satisfies $\tilde{q} = 2^u \tilde{p}$. This would then give us that $\|q^* - \tilde{q}\|_\infty = \|2^u p^* - 2^u \tilde{p}\|_\infty = 2^u \|p^* - \tilde{p}\|_\infty \leq 2^u \epsilon = \epsilon'$, which is what we want to prove.

Since we are using the decomposed Newton's method, we will show that $\tilde{q}_S = 2^u \tilde{p}_S$ for every SCC S by induction on the depth of the SCC S . Suppose that for the variables $D(S)$ that S depends on (if any), we have that $\tilde{q}_{D(S)} = 2^u \tilde{p}_{D(S)}$. If we call the k th iterate of R-NM applied to $x_S = P_S(x_S, \tilde{p}_{D(S)})$ with parameter h , $x^{[k]}$ and the k th iterate of R-NM applied to $x_S = Q_S(x_S, \tilde{q}_{D(S)})$ with parameter h' , $x'^{[k]}$, then we aim to show by induction on k that $x'^{[k]} = 2^u x^{[k]}$.

The base case is $x'^{[0]} = 0 = 2^u x^{[0]}$. By abuse of notation, we will call the Newton iterate of $x_S = P_S(x_S, \tilde{p}_{D(S)})$, $\mathcal{N}_P(x_S)$ and that of $x_S = Q_S(x_S, \tilde{q}_{D(S)})$, $\mathcal{N}_Q(x_S)$. Note that because we assume that $\tilde{q}_{D(S)} = 2^u \tilde{p}_{D(S)}$, $x_S = P_S(x_S, \tilde{p}_{D(S)})$ is the result of scaling $x_S = Q_S(x_S, \tilde{q}_{D(S)})$ using $c = 2^u$. So Lemma 4.2 (iii) yields that $\mathcal{N}_P(x_S) = 2^{-u} \mathcal{N}_Q(2^u x_S)$. If $x'^{[k]} = 2^u x^{[k]}$, then $\mathcal{N}_Q(x'^{[k]}) = 2^u \mathcal{N}_P(x^{[k]})$.

If $(\mathcal{N}_P(x^{[k]}))_i \leq 0$, we would set $x_i^{[k+1]} := 0$. If so, $\mathcal{N}_Q(x'^{[k]})_i = 2^u \mathcal{N}_P(x^{[k]})_i \leq 0$, so we would set $x_i'^{[k+1]} := 0$.

If $(\mathcal{N}_P(x^{[k]}))_i > 0$, we set $x_i^{[k+1]}$ to be the result of rounding $(\mathcal{N}_P(x^{[k]}))_i$ down to a multiple of 2^h . But then $\mathcal{N}_Q(x'^{[k]}) = 2^u \mathcal{N}_P(x^{[k]}) > 0$ and we would set $x_i'^{[k+1]}$ to be the result of rounding $(\mathcal{N}_Q(x'^{[k]}))_i$ down to a multiple of $2^{-h'}$. Note that $h' = h - u$. So the result of rounding $2^u (\mathcal{N}_P(x^{[k]}))_i$ down to a multiple of $2^{-h'}$ is just 2^u times the result of rounding $(\mathcal{N}_P(x^{[k]}))_i$ down to a multiple of 2^{-h} . So $x_i'^{[k+1]} = 2^u x_i^{[k+1]}$.

This completes the induction showing that $x'^{[k]} = 2^u x^{[k]}$ for all $k \geq 0$. Note that $g = g'$. So $\tilde{q}_S = x'^{[g]} = 2^u x^{[g]} = 2^u \tilde{p}_S$. This in turn completes the induction on the SCCs, showing that $\tilde{q} = 2^u \tilde{p}$, which completes the proof. \square

4.1. How Good Are Our Upper Bounds for R-DNM on MPSs?

We have proved upper bounds on the number of iterations required by R-DNM to converge to within additive error $\epsilon > 0$ of the LFP q^* for an arbitrary MPS $x = P(x)$.

We now discuss some important parameters of the problem in which our upper bounds are essentially optimal and cannot be improved substantially.

To begin with, our upper bounds for the number of iterations required contain a term of the form $2^f \log \frac{1}{\epsilon}$. Here f denotes the nesting depth of nonlinear SCCs in the dependency graph G_P of the input MPS, $x = P(x)$.

It was already pointed out in Esparza et al. [2010, Section 7] that such a term is a lower bound on the number of iterations required, even for exact Newton's method (whether decomposed or not), even for rather simple MPSs. Namely, Esparza et al. [2010] provided a family of simple examples entailing this lower bound. Indeed, consider the following MPS, $x = P(x)$, which is a simpler variant of the bad MPSs noted in Esparza et al. [2010]. The MPS has $n + 1$ variables, x_0, \dots, x_n . The equations are

$$\begin{aligned} x_i &= \frac{1}{2}x_i^2 + \frac{1}{2}x_{i-1} \quad , \quad \text{for all } i \in \{1, \dots, n\} \\ x_0 &= \frac{1}{2}x_0^2 + \frac{1}{2}. \end{aligned}$$

note that
this is a PPS ⁽¹⁸⁾

The LFP of this MPS is $q^* = 1$, and it captures the termination probabilities of a (rather simple) stochastic context-free grammar, pBPA, or 1-exit Recursive Markov chain. Note that the encoding size of this MPS is $|P| = O(n)$.

As observed in Etessami and Yannakakis [2009], *exact* Newton's method, starting from $x^{(0)} := 0$, on the univariate equation $x_0 = \frac{1}{2}x_0^2 + \frac{1}{2}$ gains exactly one bit of precision per iteration. In other words, if $x^{(k)}$ denotes the k th iterate, then $1 - x^{(k)} = 2^{-k}$. It is an easy exercise to show this.

Suppose we perform m iterations of exact NM on the bottom SCC, $x_0 = \frac{1}{2}x_0^2 + \frac{1}{2}$, and suppose that by doing so we obtain an approximation $q'_0 = 1 - a_0$, where $a_0 = 2^{-m}$. Plugging the approximation q'_0 into the next higher SCC, the equation for x_1 becomes $x_1 = \frac{1}{2}x_1^2 + \frac{1}{2}q'_0$. For the rest of the argument, we do not need to appeal to Newton iterations: even *exact* computation of the LFPs for the remaining SCCs will yield bad approximations overall unless $1 - q'_0 \leq \frac{1}{2^{2^n}}$ (showing that the system of equations is terribly *ill-conditioned*).

Indeed, by induction on $i \geq 0$, suppose that the value obtained for LFP of x_i is $q'_i = (1 - a_i)$. Then, after plugging in q'_i in place of x_i in the SCC for x_{i+1} , the adjusted LFP, q'_{i+1} , of the next higher SCC: $x_{i+1} = (1/2)(x_i^2) + (1/2)(1 - a_i)$, becomes $q'_{i+1} = 1 - \sqrt{a_i}$. Thus, by induction on depth, the adjusted LFP of x_n becomes $q'_n = 1 - (a_0)^{2^{-n}}$. But $a_0 = 2^{-m}$. Thus, $q'_n = 1 - 2^{-m2^{-n}}$.

We would like to have error $1 - q'_n = 2^{-m2^{-n}} \leq \epsilon$. Taking logs, we get that we must perform at least $m \geq 2^n \log \frac{1}{\epsilon}$ NM iterations on the bottom SCC alone.

Note that n here is also the (nonlinear) depth f of SCCs in this example.

Other terms in our upper bounds on the number of iterations required to compute the LFP of a general MPS are $\log \frac{1}{q_{\min}^*}$, and $\log q_{\max}^*$.

Simple "repeated squaring" MPSs, with $x_i = x_{i-1}^2$, $x_0 = \{\frac{1}{2} \text{ or } 2\}$, show that we can have $q_{\min}^* \leq \frac{1}{2^{2^n}}$, and $q_{\max}^* \geq 2^{2^n}$, where n is the number of variables. In Theorem 2.4, we give explicit lower bounds on q_{\min}^* and explicit upper bounds on q_{\max}^* , in terms of $|P|$ and n , showing that linear-double-exponential dependence on n is indeed the worst case possible.

Consider first the dependence of our bounds on $\log q_{\max}^*$. Assume $q_{\max}^* \geq 1$. Let us note that if we use the *rounded* version of Newton's method then in order to compute q_{\max}^* within additive error $0 < \epsilon < 1$, we can not do with fewer than $\approx \log q_{\max}^*$ bits of precision, that is, we require space $\approx \log q_{\max}^*$ just to write down an additive approximation to the LFP, and thus also we require time at least $\approx \log q_{\max}^*$ in the standard Turing model of computation, irrespective of what algorithm we use.

What about our dependence on q_{\min}^* ? Do we really need a worst-case number of iterations that is polynomial in $\log \frac{1}{q_{\min}^*}$, even assuming that $q_{\max}^* \leq 1$? The answer is yes: we require at least $\approx \log \frac{1}{q_{\min}^*}$ iterations, starting from 0, already to converge to within more than a single bit of precision. To see this, consider the MPS, parametrized by $n \in \mathbb{N}$, that is described in Figure 1.

Let (p, p', q^*) denote the LFP of the MPS depicted in Figure 1, where p denotes the LFP in the x coordinates, p' denotes the LFP in the x' coordinates, and q^* denotes the LFP in the y coordinates. The following claims about this MPS and its LFP are not difficult to verify.

PROPOSITION 4.5.

- (1) $0 < p_j \leq 2 \cdot 2^{-2^j}$, for $j \in \{1, \dots, 2n\}$.
- (2) $(1 - 2^{-2^j}) \leq p'_j \leq 1$, for $j \in \{1, \dots, 2n\}$.

$$\begin{aligned}
y_i &= \frac{1}{2}y_i^2 + \frac{1}{2}y_{i-1} \quad , \quad \text{for all } i \in \{1, \dots, n\} \\
y_0 &= \frac{1}{2}y_0^2 + \frac{1}{2} \cdot x'_{2n} \\
x_j &= (x_{j-1})^2 \quad , \quad \text{for all } j \in \{1, \dots, 2n\} \\
x'_j &= x'_{j-1} + x'_{j-1}x_{j-1} \quad , \quad \text{for all } j \in \{1, \dots, 2n\} \\
x_0 &= \frac{1}{2} + \frac{1}{100} \cdot x_{2n} \cdot x'_{2n} \cdot y_n \\
x'_0 &= \frac{1}{2} + \frac{1}{100} \cdot x_{2n} \cdot x'_{2n} \cdot y_n
\end{aligned}$$

Fig. 1. Strongly connected MPS on which Newton requires exponentially many iterations.

$$(3) \quad (1 - 2^{-2^{2n-i}}) \leq q_i^* \leq 1$$

Thus, the LFP is ≤ 1 in all coordinates, and furthermore p'_{2n} is extremely close to 1. Now, notice that if we simply plug in $x'_{2n} := p'_{2n}$, and if we ignore the x and x' variables, then the remaining MPS on y variables looks very very “close” to MPS (18), for which we previously argued that Newton iteration would require exponentially many iterations to get within even a single bit of precision. Let $y_i^{(k)}$ denote the k th Newton iterate, in the coordinate corresponding to y_i , starting at 0, on the strongly connected MPS in Figure 1. By careful analysis of Newton’s method, we can in fact show the following.

$$\text{PROPOSITION 4.6. } |q_n^* - y_n^{(2^{n-2})}| > \frac{1}{2} - \frac{1}{2^{2^n}}.$$

We forgo a detailed proof of this bound, but let us explain in rough terms why it holds. We first note that Newton’s method does not converge too fast at coordinate y_0 . Indeed it converges “essentially linearly” for the initial 2^{n-2} iterations. Specifically, for all $k \leq 2^{n-2}$, we can show that $y_0^{(k)} \leq 1 - \frac{1}{2^{2k}}$. Next, an analysis similar to the “square rooting” analysis earlier in this section, adapted from Esparza et al. [2010], which was used to show that the nonstrongly connected MPS of (18) requires exponentially many Newton iterations to converge to within a single bit of precision, can be applied to analyze the errors in the variable y_n obtained by the Newton iterates, and yields that $y_n^{(k)} \leq 1 - 2^{-(2k \cdot 2^n)}$. So $y_n^{(2^{n-2})} \leq \frac{1}{2}$. We omit further details.

Thus, already for strongly connected MPSs, the polynomial dependence on $\log \frac{1}{q_{\min}^*}$ for the number of iterations of Newton’s method required to get within just one bit of precision, is unavoidable.

It should be noted however that the worst-case bounds on q_{\min}^* and q_{\max}^* are not representative of many important families of MPSs. In particular, note that MPSs whose LFP corresponds to (termination) probabilities must have $q_{\max}^* \leq 1$. Furthermore, for a number of classes of probabilistic systems, we can prove bounds of the form $\log \frac{1}{q_{\min}^*} \leq \text{poly}(|P|)$. Indeed, for MPSs corresponding to QBDs and probabilistic 1-counter automata, which we consider in the next section, such bounds were established in Etessami et al. [2010].

If the family of MPSs happens to have $\log \frac{1}{q_{\min}^*}, \log q_{\max}^* \leq \text{poly}(|P|)$, then our upper bounds show that the total number of iterations of R-DNM needed is only exponential

in ($f \leq d$), the (nonlinear) depth of SCCs, and thus if $f \leq \log |P|$, then for such MPSs R-DNM runs in P -time

in the encoding size of the input, $|P|$ and $\log \frac{1}{\epsilon}$, in the standard Turing model of computation,

to compute an approximation to the LFP q^* , within additive error $\epsilon > 0$.

It should be noted that for the case of *strongly connected MPSs* only, and only for *Exact Newton's Method*, without rounding, Esparza et al. [2010] obtained comparable result to ours in terms of worst-case dependence on $\log \frac{1}{q_{\min}^*}$ and $\log q_{\max}^*$.⁶ However, in Esparza et al. [2010], they did not obtain any constructive bounds in terms of $|P|$, q_{\min}^* or q_{\max}^* for MPSs that are not strongly connected, nor did they obtain any results for rounded versions of Newton's method. Using exact Newton's method of course entails the assumption of a unit-cost arithmetic model of computation, rather than the Turing model.

5. MPSs AND PROBABILISTIC 1-COUNTER AUTOMATA

A *probabilistic 1-counter automaton* (pICA), M , is a 3-tuple $M = (V, \delta, \delta_0)$ where V is a finite set of control states and $\delta \subseteq V \times \mathbb{R}_{>0} \times \{-1, 0, 1\} \times V$ and $\delta_0 \subseteq V \times \mathbb{R}_{>0} \times \{0, 1\} \times V$ are transition relations. The transition relation δ is enabled when the counter is nonzero, and the transition relation δ_0 is enabled when it is zero. For example, a transition of the form, $(u, p, -1, v) \in \delta$, says that if the counter value is positive, and we are currently in control state u , then with probability p , we move in the next step to control state v and we decrement the counter by 1. A pICA defines in the obvious way an underlying countably infinite-state (labeled) Markov chain, whose set of configurations (states) are pairs $(v, n) \in V \times \mathbb{N}$. A run (or trajectory, or sample path), starting at initial state (v_0, n_0) is defined in the usual way, as a sequence of configurations $(v_0, n_0), (v_1, n_1), (v_2, n_2), \dots$ that is consistent with the transition relations of M .

As explained in Etessami et al. [2010], pICAs are in a precise sense equivalent to discrete-time *quasi-birth-death processes* (QBDs), and to *1-box recursive Markov chains*.

Quantities that play a central role for the analysis of QBDs and pICAs (both for transient analyses and steady-state analyses, as well as for model checking) are their *termination probabilities* (also known as their *G-matrix* in the QBD literature, see, e.g., Latouche and Ramaswami [1999], Bini et al. [2005], and Etessami et al. [2010]). These are defined as the probabilities, $q_{u,v}^*$, of hitting counter value 0 for the first time in control state $v \in V$, when starting in configuration $(u, 1)$.

Corresponding to the termination probabilities of every QBD or pICA is a special kind of MPS, $x = P(x)$, whose LFP solution q^* gives the termination probabilities of the pICA. The MPSs corresponding to pICAs have the following special structure. For each pair of control states $u, v \in V$ of the pICA, there is a variable x_{uv} , which represents the probability of termination at configuration $(v, 0)$ starting at configuration $(u, 1)$. The equation for each variable x_{uv} has the following form:

$$x_{uv} = p_{uv}^{(-1)} + \left(\sum_{w \in V} p_{uw}^{(0)} x_{wv} \right) + \sum_{y \in V} p_{uy}^{(1)} \sum_{z \in V} x_{yz} x_{zv}, \quad (19)$$

where for all states $u, v \in V$, and $j \in \{-1, 0, 1\}$, the coefficients $p_{uv}^{(j)}$ are non-negative transition probabilities of the pICA, and such that for all states $u \in V$,

⁶Technically, their bounds are with respect to *relative error*, and their bounds for strongly connected MPSs do not depend at all on q_{\max}^* , but of course if q_{\max}^* is large, then in order to obtain absolute (additive) error $\epsilon > 0$, the relative error required is $\epsilon' = \frac{\epsilon}{q_{\max}^*}$, and since their bounds depend on $\log \frac{1}{\epsilon'}$, they depend (indirectly) on $\log q_{\max}^*$, with the same magnitude as ours.

$\sum_{j \in \{-1, 0, 1\}} \sum_{v \in V} p_{uv}^{(j)} \leq 1$. We can, of course, clean up this MPS in P-time (by Proposition 2.2), to remove all variables x_{uv} for which $q_{u,v}^* = 0$. In what follows, we assume this has been done, and thus that for the remaining variables $0 < q^* \leq 1$.

In Etessami et al. [2010], the decomposed Newton's method (DNM) is used *with exact arithmetic* in order to approximate the LFP for p1CAs using polynomially many arithmetic operations, that is, in polynomial time in the *unit-cost arithmetic model of computation*. However, Etessami et al. [2010] did not establish any result in the standard Turing model of computation. We establish instead results about R-DNM applied to the MPSs arising from p1CAs, in order to turn this method into a P-time algorithm in the standard model of computation.

It was shown in Etessami et al. [2010, Theorem 9] that in any path through the DAG of SCCs of the dependency graph for the MPS associated with a p1CA, M , there is at most one nonlinear SCC, that is, that the nonlinear depth f of these special MPSs is at most 1. Also, Etessami et al. [2010] obtained a lower bound on q_{\min}^* , the smallest positive termination probability. Namely, if c_{\min} denotes the smallest positive transition probability of a p1CA, M , and thus also the smallest positive constant or coefficient of any monomial in the corresponding MPS, $x = P(x)$, they showed the following.

LEMMA 5.1 [ETESSAMI ET AL. 2010, COROLLARY 6]. $q_{\min}^* \geq c_{\min}^{r^3}$, where r is the number of control states of the p1CA.

They used these results to bound the *condition number* of the Jacobian matrix for each of the linear SCCs, and to thereby show that one can approximate q^* in polynomially many arithmetic operations using decomposed Newton's method. Here, we get a stronger result, placing the problem of computing termination probabilities for p1CA in P-time in the standard Turing model, using the results from this article.

THEOREM 5.2. *Let $x = P(x)$ be the MPS associated with p1CA, M , let r denote the number of control states of M , and let m denote the maximum number of bits required to represent the numerator and denominator of any positive rational transition probability in M .*

Apply R-DNM, including rounding-down linear SCCs, to the MPS $x = P(x)$, using rounding parameter $h := 8mr^7 + 2mr^5 + 9r^2 + 3 + \lceil 2 \log \frac{1}{\epsilon} \rceil$ and such that for each nonlinear SCC we perform $g = h - 1$ iterations, whereas for each linear SCC we only perform 1 R-NM iteration.

This algorithm computes an approximation \tilde{q} to q^ , such that $\|q^* - \tilde{q}\|_{\infty} < \epsilon$. The algorithm runs in time polynomial in $|M|$ and $\log \frac{1}{\epsilon}$, in the standard Turing model of computation.*

PROOF OF THEOREM 5.2. We apply Theorem 4.1, which tells us that R-DNM with parameter

$$h \geq \left\lceil 3 + 2^f \cdot \left(\log\left(\frac{1}{\epsilon}\right) + d \cdot (\log(\alpha^{-(4n+1)}) + \log(16n) + \log(\|P(1)\|_{\infty})) \right) \right\rceil, \quad (20)$$

using $g = h - 1$ iterations for every SCC, gives an approximation \tilde{q} to q^* with $\tilde{q} \leq q^*$ and such that $\|q^* - \tilde{q}\|_{\infty} \leq \epsilon$. Here $f \leq 1$ since, by Theorem 9 of Etessami et al. [2010], there is at most 1 nonlinear SCC in any path through the dependency graph. Furthermore, $n = r^2$, since the variables in x are indexed by two states x_{uv} . Also, $d \leq n$, and so $d \leq r^2$. Also, $c_{\min} \geq 2^{-m}$ and so by Lemma 5.1, $q_{\min}^* \geq 2^{-mr^3}$. So $\alpha \geq 2^{-(mr^3+1)}$. To show that $\|P(1)\|_{\infty} \leq r$, by Eq. (19), $P(1)_{uv} = p_{uv}^{(-1)} + (\sum_{w \in V} p_{uw}^{(0)}) + \sum_{y \in V} p_{uy}^{(1)} r \leq r$. Plugging all this into Eq. (20), we get: $h \geq \lceil 3 + 2 \cdot (\log(\frac{1}{\epsilon}) + r^2 \cdot ((4r^2 + 1)(mr^3 + 1) + \log(16r^2) + \log r)) \rceil$.

Noting that $\log(16r^2) + \log r = \log(16r^3)$, and noting that $r \geq 1$ implies $\log(16r^3) \leq 4r$, we have

$$h \geq 3 + 8mr^7 + 2mr^5 + 9r^4 + \left\lceil 2 \cdot \log\left(\frac{1}{\epsilon}\right) \right\rceil.$$

Note that the rounding parameter h and the number of iterations $g = h - 1$ are both polynomials in the encoding size of the p1CA, and in $\log \frac{1}{\epsilon}$. Thus, each iteration of R-DNM can be computed in polynomial time, and we only do polynomially many iterations. Thus, the entire computation of \tilde{q} can be carried out in P-time in the Turing model of computation. \square

5.1. Application to ω -Regular Model Checking for p1CAs

Since computing termination probabilities of p1CAs (equivalently, the G -matrix of QBDs) plays such a central role in other analyses (see, e.g., Latouche and Ramaswami [1999], Bini et al. [2005], Etessami et al. [2010], and Brázdil et al. [2011]), the P-time algorithm given in the previous section for computing termination probabilities of a p1CA (within arbitrary desired precision) directly facilitates P-time algorithms for various other important problems.

Here we highlight just one of these applications: a P-time algorithm *in the Turing model of computation* for model checking a p1CA with respect to any ω -regular property. An analogous result was established by Brázdil et al. [2011] in the unit-cost RAM model of computation.⁷

THEOREM 5.3. *Given a p1CA, M , with states labeled from an alphabet Σ , and with a specified initial control state v , and given an ω -regular property $L(B) \subseteq \Sigma^\omega$, which is specified by a nondeterministic Büchi automaton, B , let $Pr_M(L(B))$ denote the probability that a run of M starting at configuration $(v, 0)$ generates an ω -word in $L(B)$. There is an algorithm that, for any $\epsilon > 0$, computes an additive ϵ -approximation, $\tilde{p} \geq 0$, of $Pr_M(L(B))$, that is, with $|Pr_M(L(B)) - \tilde{p}| \leq \epsilon$. The algorithm runs in time polynomial in $|M|$, $\log \frac{1}{\epsilon}$, and $2^{|B|}$, in the standard Turing model of computation.*

PROOF SKETCH. By Theorem 5.2, we know we can compute termination probabilities q^* for a p1CA, M , with additive error $\epsilon > 0$ in time polynomial in $|M|$ and $\log \frac{1}{\epsilon}$.

Let us first observe that if we do not insist on having the ω -regular property specified by a non-deterministic Büchi automaton B , and instead assume it is specified by a deterministic Rabin automaton R , then the analogous theorem follows immediately as a corollary of Theorem 5.2 and results established by Brázdil et al. [2011]. Specifically, in Brázdil et al. [2011], it was shown that, given a p1CA, M , and a deterministic Rabin automaton, R , and given $\epsilon > 0$, there is an algorithm that, first, decides in P-time whether $Pr_M(L(R)) > 0$, and if so computes a value \tilde{p} which approximates $Pr_M(L(R))$ with relative error $\epsilon > 0$, that is, such that $|Pr_M(L(R)) - \tilde{p}|/Pr_M(L(R)) < \epsilon$, and the algorithm runs in time polynomial in $|M|$, $|R|$, and $\log \frac{1}{\epsilon}$, in the unit-cost RAM model of computation.

The first observation we make is that, the results in Etessami et al. [2010] and Brázdil et al. [2011] together imply that, for p1CAs, there is no substantial difference in complexity between relative and absolute approximation, because the probability $Pr_M(L(R))$ can be bounded away from zero by $1/2^{\text{poly}(|M|, |R|)}$ if it is not equal to zero

⁷In the more recent journal version of their paper, Brázdil et al. [2014] use the results of this article (which appeared in the conference version [Stewart et al. 2013]), in order to similarly conclude that model checking can be carried out in P-time in the Turing model of computation.

(which can be detected in P-time). Thus, computing $Pr_M(L(R))$ with given relative error $\epsilon > 0$ is P-time equivalent to computing $Pr_M(L(R))$ with ϵ absolute error.

Second, a close inspection of Brázdil et al. [2011] shows that the *only* use made in their entire paper of the unit-cost RAM model of computation is for the purpose of computing termination probabilities for p1CAs, and specifically because they directly invoke the earlier result from Etessami et al. [2010] which showed that termination probabilities q^* for a p1CA can be ϵ -approximated in polynomial time in the unit-cost RAM model. Thus, the only thing needed in order to obtain an absolute error ϵ -approximation of $Pr_M(L(R))$ in P-time in the standard Turing model of computation is to appeal instead to Theorem 5.2 of this article for computation of termination probabilities in P-time in the standard Turing model, and apply the rest of the construction in Brázdil et al. [2011].

Next, let's first note that we can of course use Safra's construction to convert any non-deterministic Büchi automaton B to a deterministic Rabin automaton of size $2^{O(|B| \log |B|)}$. So, obtaining a complexity bound that is polynomial in $2^{|B| \log |B|}$ is no more difficult.

Let us now sketch why one can in fact obtain the (slightly) better complexity bound, polynomial in $2^{|B|}$, by using the algorithm for model checking of RMCs of Etessami and Yannakakis [2012], combined with Theorem 5.2 and Lemma 5.1, and with the key result of Brázdil et al. [2011], which establishes that nonzero *nontermination* probabilities for a p1CA are also bounded away from zero by $1/2^{\text{poly}(|M|)}$.

As shown in Courcoubetis and Yannakakis [1995] and Etessami and Yannakakis [2012], for probabilistic model checking a naive subset construction can be used (instead of Safra's construction) to obtain from a BA, B , a deterministic Büchi automaton, D , such that $|D| = 2^{|B|}$. (Although it need not be the case that $L(D) = L(B)$, it is shown in these references via an intricate combinatorial analysis that D is nevertheless sufficient to perform both qualitative and quantitative probabilistic model checking.) One then constructs the "product" $M \otimes D$, of the p1CA, M , with the deterministic Büchi automaton D . A key observation is that this "product" remains a p1CA. In terms of RMCs, p1CAs correspond to the subclass of 1-box RMCs, and the "product" of a 1-box RMC with a deterministic BA, D , remains a 1-box RMC.

It was shown in Etessami and Yannakakis [2012] that given a "product" (1-box) RMC $M \otimes D$, it is possible to construct a finite-state *conditioned summary chain*, \mathcal{M}' , which is a finite state Markov chain and whose transition probabilities are *rational expressions in positive termination and nontermination probabilities* of the (1-box) RMC. It is then possible to identify in P-time certain bottom strongly connected components \mathcal{T} of \mathcal{M}' , such that the probability $Pr_M(L(B))$ is equal to the probability that starting from a specific initial state of \mathcal{M}' , a run eventually hits a state in \mathcal{T} .

We describe here for concreteness the construction of the summary chain \mathcal{M}' in terms of the p1CA, and specify the set \mathcal{T} of "accepting" bottom SCCs. The product p1CA $\hat{M} = M \otimes D$ is a p1CA that has one state (u, S) for every state u of the p1CA M and every state S of D (i.e., every subset S of states of the automaton B), and has transition relations $\hat{\delta}$ and $\hat{\delta}_0$ corresponding to the transition relations δ and δ_0 of the p1CA M , defined in a straightforward way: If the labeling function of the p1CA M is λ , and D has state set Q_D and transition function σ , then $\hat{\delta} = \{((u, S), p, b, (v, \sigma(S, \lambda(v)))) \mid (u, p, b, v) \in \delta, S \in Q_D\}$ and $\hat{\delta}_0 = \{((u, S), p, b, (v, \sigma(S, \lambda(v)))) \mid (u, p, b, v) \in \delta_0, S \in Q_D\}$. Consider the system of Eqs. (19) for the product p1CA \hat{M} and let q^* be its LFP. That is, for every pair of states $(u, S), (w, T)$ of \hat{M} , the corresponding entry $q^*_{(u, S), (w, T)}$ is the probability that \hat{M} started at configuration $((u, S), 1)$ reaches the counter value 0 for the first time ("terminates") at state (w, T) . Let $ne(u, S) = 1 - \sum_{(w, T)} q^*_{(u, S), (w, T)}$ be the non-termination

probability from (u, S) , that is, the probability that, starting from $((u, S), 1)$, the counter never reaches 0. Let $ne'(u, S) = ne(u, S) + \sum_{(w, T)} q_{(u, S), (w, T)}^* \cdot ne(w, T)$; this is the probability that \hat{M} starting from configuration $((u, S), 2)$ never reaches counter value 0.

The set of states of the summary chain \mathcal{M}' is as follows: For every state (u, S) of \hat{M} , the chain \mathcal{M}' has states (u_0, S) , (u'_0, S) , and in addition it has a state (u_1, S) iff $ne(u, S) > 0$, and a state (u'_1, S) iff $ne'(u, S) > 0$. The initial state is $s_0 = (v_0, \{r_0\})$, where v is the initial state of the p1CA M and r_0 is the initial state of B .

The summary chain \mathcal{M}' has the following transitions. For every transition of $\hat{\delta}_0$ from a state (u, S) to a state (w, T) with probability p and counter-increment 0 (respectively, 1), \mathcal{M}' has a transition from (u_0, S) to (w_0, T) (respectively, (w'_0, T)) with the same probability p . For every transition of $\hat{\delta}$ from a state (u, S) to a state (w, T) with probability p and counter-increment 0 (respectively, 1), if (u_1, S) is a state of \mathcal{M}' (i.e., if $ne(u, S) > 0$), then \mathcal{M}' has a transition from (u_1, S) to (w_1, T) (respectively, (w'_1, T)), if the latter state exists in \mathcal{M}' , with probability $p \cdot ne(w, T)/ne(u, S)$ (respectively, $p \cdot ne'(w, T)/ne(u, S)$). (Note: \mathcal{M}' does not have any direct transitions corresponding to transitions of \hat{M} that decrement the counter.) Every state (u'_0, S) of \mathcal{M}' has a transition to (w_0, T) with probability $q_{(u, S), (w, T)}^*$ for every state (w, T) of \hat{M} such that $q_{(u, S), (w, T)}^* > 0$; in addition, (u'_0, S) has a transition with probability $ne(u, S)$ to (u_1, S) if $ne(u, S) > 0$. Every state (u'_1, S) of \mathcal{M}' has a transition to (w_1, T) with probability $q_{(u, S), (w, T)}^* \cdot ne(w, T)/ne'(u, S)$ for every state (w, T) of \hat{M} such that $q_{(u, S), (w, T)}^* > 0$ and $ne(w, T) > 0$; in addition (u'_1, S) has a transition with probability $ne(u, S)/ne'(u, S)$ to (u_1, S) if $ne(u, S) > 0$. This concludes the description of the summary chain \mathcal{M}' .

Once the summary chain \mathcal{M}' is constructed, a set \mathcal{T} of its bottom SCCs are characterized as *accepting*. The characterization of accepting bottom SCCs depends only on the underlying graph of \mathcal{M}' (i.e., it does not depend on the precise values of the transition probabilities), and can be done in polynomial time in the size of \mathcal{M}' . Specifically, a bottom SCC is accepting if it contains a state (z, S) , where S contains a state s of B , such that the pair (z, s) is a *special* pair of one of two types: A pair (z, s) is special of type 1 if s is an accepting state of B and the subgraph of \mathcal{M}' induced by the nodes reachable from $(z, \{s\})$ contains a bottom SCC that has a state (z, T) with $s \in T$. A pair (z, s) is special of type 2 if $z = u'_0$ or $z = u'_1$ for some state u of M , and the following graph $M'(z, s)$ has a bottom SCC that contains a state (z, T) with $s \in T$. If $z = u'_0$ (respectively, $z = u'_1$), then $M'(z, s)$ is the subgraph of \mathcal{M}' induced by all the nodes that are reachable from the nodes $(w_0, \{t\})$ (respectively, $(w_1, \{t\})$), such that the p1CA M has a path π from the configuration $(u, 1)$ that terminates (reaches for the first-time counter value 0) at $(w, 0)$, and the automaton B has a run on π from s to t that goes through an accepting state; this condition can be tested in polynomial time using the algorithms for (non-probabilistic) Recursive State Machines of Alur et al. [2005]. It is shown in Etessami and Yannakakis [2012] (more generally, for RMCs) that the probability $Pr_M(L(B))$ is equal to the probability that a run of \mathcal{M}' starting from the initial state s_0 eventually hits a state in \mathcal{T} .

In this way, the model checking problem is boiled down to the problem of computing hitting probabilities in a *finite-state* Markov chain \mathcal{M}' whose transition probabilities are simple rational expressions with numerators and denominators that are products of coefficients in a p1CA together with positive termination and non-termination probabilities of a p1CA.

It is well known that nonzero hitting probabilities for a finite-state Markov chain are the unique solution $(I-A)^{-1}b$, to a linear system of equations $x = Ax+b$, where the coefficients in A and b come from the transition probabilities of the Markov chain. The key remaining question is, *how well conditioned is this linear system of equations?* In other words, what happens to its unique solution if we only approximate the coefficients in A and b to within a small error? Now, the key is that applying Lemma 5.1 (which is from Etessami et al. [2010]), and applying the key result in Brázdil et al. [2011], together shows that both positive termination and positive nontermination probabilities of the product p1CA are bounded away from 0 by $1/2^{\text{poly}(|M|,|D|)}$.

Under these conditions, *exactly the same* known condition number bounds from numerical analysis that were used in Etessami et al. [2010], namely Theorem 17 of Etessami et al. [2010], which is a version of Theorem 2.1.2.3 of Isaacson and Keller [1966], also establish that the linear system of equations that one has to solve for hitting probabilities in the conditioned summary chain \mathcal{M}' derived from a p1CA are “polynomially well conditioned”, meaning that approximating their nonzero coefficients within suitable $1/2^{\text{poly}}$ additive error yields a linear system of equations whose unique solution is ϵ -close to the unique solution of the original system, for the chosen $\epsilon > 0$. Thus, we can first approximate the coefficients of the conditioned summary chain \mathcal{M}' in P-time, and we can then solve for the unique solution of the corresponding system of linear equations, in order to obtain the desired approximation of the probability $Pr_{\mathcal{M}}(L(B))$ in P-time. We omit further details. \square

APPENDIX: PROOF OF THEOREM 2.4

THEOREM 2.4. *If $x = P(x)$ is a quadratic MPS in n variables, with LFP $q^* > 0$, and where $P(x)$ has rational coefficients and total encoding size $|P|$ bits, then*

- (1) $q_{\min}^* \geq 2^{-|P|(2^n-1)}$, and
- (2) $q_{\max}^* \leq 2^{2(n+1)(|P|+2(n+1)\log(2n+2)) \cdot 5^n}$.

PROOF. We first prove (1), by lower bounding q_{\min}^* in terms of the smallest positive constant c_{\min} in $P(x)$.

LEMMA A.1. *If $x = P(x)$ has LFP $q^* > 0$, and least term c_{\min} , then $q_{\min}^* \geq \min\{1, c_{\min}\}^{2^n-1}$.*

PROOF. We first observe that, since $q^* > 0$, and there are n variables, it must be the case that $P^n(0) > 0$. To see this, for any $y \geq 0$, let us use $Z(y)$ to denote the set of zero coordinates of y . For any $k \geq 0$, $P^{k+l}(0) \geq P^k(0)$, for all $l \geq 0$, so $Z(P^{k+l}(0)) \subseteq Z(P^k(0))$. Thus, either $|Z(P^{k+1}(0))| = |Z(P^k(0))|$ or $|Z(P^{k+1}(0))| \leq |Z(P^k(0))| - 1$. Now $|Z(0)| = n$ and $|Z(P^k(0))| \geq 0$ for all k , so there must be some least $0 \leq k \leq n$ such that $|Z(P^k(0))| = |Z(P^{k+1}(0))|$ and such that $Z(P^k(0)) = Z(P^{k+1}(0))$.

Note that, for any $y \geq 0$, $Z(P(y))$ depends only on $Z(y)$ and on $P(x)$, but not on the specific values of nonzero coordinates of y .

So if for some $n \geq k \geq 0$, $Z(P^{k+1}(0)) = Z(P^k(0))$, then, by a simple induction $Z(P^{k+l}(0)) = Z(P^k(0))$ for all $l \geq 0$. So we must have $Z(P^k(0)) = Z(P^n(0)) = Z(P^{n+l}(0))$, for all $l \geq 0$. Now $\lim_{m \rightarrow \infty} P^m(0) = q^*$. Now if $P^n(0)_i = 0$, then $P^{n+l}(0)_i = 0$ for all $l \geq 0$, and so $q_i^* = 0$. This contradicts our assumption that $q^* > 0$. So $P^n(0) > 0$.

Let us use $P^k(0)_{@}$ to denote the minimum value of any nonzero coordinate of $P^k(0)$. First, $P(0) \neq 0$, that is, there is some nonzero constant in some polynomial $P(x)_i$. Thus $P(0)_{@} \geq c_{\min}$. We show by induction that for $k > 0$, $P^k(0)_{@} \geq \min\{1, c_{\min}\}^{2^k-1}$. This is true for $k = 1$. We assume that $P^k(0)_{@} \geq \min\{1, c_{\min}\}^{2^k-1}$. If for some

coordinate i , $P^{k+1}(0)_i = P(P^k(0))_i > 0$, there must be a term in $P(x)_i$ which is not zero in $P(P^k(0))_i$, this is either a constant c , or a linear term cx_j with $P^k(0)_j > 0$, or a quadratic term cx_jx_l with $P^k(0)_j > 0$ and $P^k(0)_l > 0$. In any of these three cases, this term is $\geq c_{\min} \min\{1, P^k(0)_{@}\}^2$. Since $P^k(0)_{@} \geq \min\{1, c_{\min}\}^{2^k-1}$, we now have that $P^{k+1}(0)_{@} \geq c_{\min} (\min\{1, c_{\min}\}^{2^k-1})^2 \geq \min\{1, c_{\min}\}^{2^{k+1}-1}$. So for all k , $P^k(0)_{@} \geq \min\{1, c_{\min}\}^{2^k-1}$. In particular $P^n(0)_{@} \geq \min\{1, c_{\min}\}^{2^n-1}$. But $P^n(0) > 0$ so $P^n(0)_{\min} \geq \min\{1, c_{\min}\}^{2^n-1}$. We know $q^* \geq P^n(0)$, so $q^*_{\min} \geq \min\{1, c_{\min}\}^{2^n-1}$. \square

To get our lower bound on q^*_{\min} in terms of $|P|$ and n , we just note that clearly $c_{\min} \geq 2^{-|P|}$. This and Lemma A.1 give the bound $q^*_{\min} \geq 2^{-|P|(2^n-1)}$ in part (1) of the theorem.

We now prove part (2). To prove the upper bound on q^*_{\max} , we need the following isolated root separation bound for systems of polynomial equations by Hansen et al. [2011]:

THEOREM A.2 [HANSEN ET AL. 2011, THEOREM 23]. *Consider a polynomial system of equations*

$$(\Sigma) \quad g_1(x_1, \dots, x_n) = \dots = g_m(x_1, \dots, x_n) = 0, \quad (21)$$

with polynomials of degree at most d and integer coefficients of magnitude at most 2^τ .

Then, the coordinates of any isolated (in Euclidean topology) real solutions of the system are real algebraic numbers of degree at most $(2d+1)^n$, and their defining polynomials have coefficients of magnitude at most $2^{2n(\tau+4n \log(dm))(2d+1)^{n-1}}$. Also, if $\gamma_j = (\gamma_{j,1}, \dots, \gamma_{j,n})$ is an isolated solution of (Σ) , then for any i , either

$$2^{-2n(\tau+2n \log(dm))(2d+1)^{n-1}} < |\gamma_{j,i}| \quad \text{or} \quad \gamma_{j,i} = 0. \quad (22)$$

To apply Theorem A.2, we now establish that q^* is an isolated solution of an MPS with $\text{LFP } q^* > 0$.

LEMMA A.3. *If $x = P(x)$ is a quadratic MPS with $\text{LFP } q^* > 0$, then q^* is an isolated solution of the system of equations $x = P(x)$.*

PROOF. First, we consider strongly connected MPSs. These can be divided into two cases, linear strongly connected MPSs, where $B(x) = B$ is a constant matrix and $P(x)$ is affine, and nonlinear strongly connected MPSs, where $B(x)$ is not a constant matrix and $P(x)$ is nonlinear.

For the linear case, the Jacobian is a constant $B(x) = B$, and $x = P(x) = Bx + P(0)$. We know that $\rho(B(q^*)) \leq 1$ from Lemma 3.10, and thus since $B = B(0) = B(q^*)$, from Lemma 3.10, we know that $\rho(B) < 1$, and thus $(I - B)$ is nonsingular, and there is a unique solution to $x = P(x) = Bx + P(0)$, namely $q^* = (I - B)^{-1}P(0)$. Being unique, this solution is isolated.

Now suppose, for contradiction, that $x = P(x)$ is a nonlinear strongly connected quadratic MPS but that $q^* > 0$ is not an isolated solution to $x = P(x)$. Because q^* is not isolated, there is another fixed-point q with $\|q^* - q\|_\infty \leq q^*_{\min}$ and $q \neq q^*$. Then $q \geq 0$ and, since q^* is the least nonnegative fixed-point, $q \geq q^*$. From Lemma 3.3, we have

$$P(q) - P(q^*) = B \left(\frac{1}{2}(q^* + q) \right) (q - q^*).$$

Because q^* and q are fixed points

$$q - q^* = B \left(\frac{1}{2}(q^* + q) \right) (q - q^*).$$

Lemma 3.11 now yields that since $q - q^* \geq 0$ but $q - q^* \neq 0$ and $B(\frac{1}{2}(q^* + q))$ is irreducible, $q > q^*$. Thus, $q - q^* > 0$ is a positive eigenvector of the irreducible matrix $B(\frac{1}{2}(q^* + q))$ associated with eigenvalue 1, thus $\rho(B(\frac{1}{2}(q^* + q))) = 1$ by Lemma 3.9.

We now again invoke the assumption of nonisolation of q^* , which implies there is a vector $q' \neq q^*$ such that $q' = P(q')$ and $\|q^* - q'\|_\infty \leq \min\{q_{\min}^*, \frac{1}{2}(q - q^*)_{\min}\}$. By the same reasoning as previously mentioned, we have that $q' > q^*$ and $\rho(B(\frac{1}{2}(q^* + q'))) = 1$. But now the condition $\|q^* - q'\|_\infty \leq \frac{1}{2}(q - q^*)_{\min}$ yields that $q' \leq q^* + \frac{1}{2}(q - q^*) < q$. We thus also have that $\frac{1}{2}(q^* + q) > \frac{1}{2}(q^* + q')$, and because $B(x)$ is nonconstant and monotone in x , we have $B(\frac{1}{2}(q^* + q)) \geq B(\frac{1}{2}(q^* + q'))$ and $B(\frac{1}{2}(q^* + q)) \neq B(\frac{1}{2}(q^* + q'))$. However, $\rho(B(\frac{1}{2}(q^* + q))) = 1 = \rho(B(\frac{1}{2}(q^* + q')))$. This contradicts Lemma 3.7. So q^* is also isolated in this case.

This establishes that for all strongly connected MPSs, with LFP $q^* > 0$, q^* is isolated.

Now suppose that $x = P(x)$ is not strongly connected. For each SCC S of $x = P(x)$, the MPS $x_S = P_S(x_S, q_{D(S)}^*)$ is strongly connected, so its LFP q_S^* is an isolated solution of $x_S = P_S(x_S, q_{D(S)}^*)$. That is, there is an $\epsilon_S > 0$ such that if q_S has $\|q_S - q_S^*\|_\infty \leq \epsilon_S$ and $q_S = P_S(q_S, q_{D(S)}^*)$, then $q_S = q_S^*$. Now take $\epsilon = \min_S\{\epsilon_S\}$. We claim that if $\|q - q^*\|_\infty \leq \epsilon$ and $P(q) = q$, then $q = q^*$. We can show this by induction on the depth of strongly connected components. If S is a bottom strongly connected component, then q_S has $\|q_S - q_S^*\|_\infty \leq \epsilon \leq \epsilon_S$ and $q_S = P_S(q_S)$. So $q_S = q_S^*$. If S is a SCC and for all variables $D(S)$ that variables in S depend on, directly or indirectly, $q_{D(S)} = q_{D(S)}^*$, then q_S has $q_S = P_S(q_S, q_{D(S)}^*) = P_S(q_S, q_{D(S)}^*)$. But this and $\|q_S - q_S^*\|_\infty \leq \epsilon \leq \epsilon_S$ are enough to establish $q_S = q_S^*$. This completes the induction showing that $q = q^*$. So q^* is isolated solution for any MPS with LFP $q^* > 0$. \square

For each x_i , let d_i be the product of the denominators of all coefficients of $P(x)_i$. Then, $d_i x = d_i P(x)_i$ clearly has integer coefficients which are no larger than $2^{|P|}$. Suppose $x = P(x)$ has LFP $q^* > 0$, and suppose that coordinate k is the maximum coordinate of q^* , that is, that $q_k^* = q_{\max}^*$. Now consider the system of $n + 1$ polynomial equations, in $n + 1$ variables (with an additional variable y), given by

$$d_i x_i = d_i P(x)_i, \quad \text{for all } i \in \{1, \dots, n\}; \quad \text{and } x_k y = 1. \quad (23)$$

Lemma A.3 tells us that $q^* > 0$ is an isolated solution of $x = P(x)$. If $z \in \mathbb{R}^n$ is any solution vector for $x = P(x)$, there is a unique $w \in \mathbb{R}$ such that $x := z$ and $y := w$ forms a solution to the Eqs. (23); namely let $w = \frac{1}{z_k}$. So, letting $x := q^*$, and letting $y := \frac{1}{q_k^*}$, gives us an isolated solution of the Eqs. (23). We can now apply Theorem A.2 to the system (23). For $y = \frac{1}{q_k^*}$, Eq. (22) in Theorem A.2 says that

$$2^{-2(n+1)(|P|+2(n+1)\log(2n+2))5^n} < \frac{1}{q_k^*} \quad \text{or} \quad \frac{1}{q_k^*} = 0.$$

Clearly $\frac{1}{q_k^*} \neq 0$, so $\frac{1}{q_{\max}^*} = \frac{1}{q_k^*} > 2^{-2(n+1)(|P|+2(n+1)\log(2n+2))5^n}$. So

$$q_{\max}^* < 2^{2(n+1)(|P|+2(n+1)\log(2n+2))5^n}. \quad (24)$$

\square

ACKNOWLEDGMENTS

We thank the anonymous referees for their helpful comments.

REFERENCES

- E. Allender, P. Bürgisser, J. Kjeldgaard-Pedersen, and P. B. Miltersen. 2009. On the complexity of numerical analysis. *SIAM J. Comput.* 38, 5, 1987–2006.
- R. Alur, M. Benedikt, K. Etessami, P. Godefroid, T. Reps, and M. Yannakakis. 2005. Analysis of recursive state machines. *ACM Trans. Program. Lang. Syst.* 27, 4, 786–818.
- T. Apostol. 1974. *Mathematical Analysis* (2nd ed.). Addison-Wesley.
- D. Bini, G. Latouche, and B. Meini. 2005. *Numerical Methods for Structured Markov Chains*. Oxford University Press.
- T. Brázdil, S. Kiefer, and A. Kucera. 2011. Efficient analysis of probabilistic programs with an unbounded counter. In *Proceedings of the 23rd International Conference on Computer Aided Verification (CAV)*. 208–224.
- T. Brázdil, S. Kiefer, and A. Kucera. 2014. Efficient analysis of probabilistic programs with an unbounded counter. *J. ACM* 61, 6.
- C. Courcoubetis and M. Yannakakis. 1995. The complexity of probabilistic verification. *J. ACM* 42, 4, 857–907.
- R. Durbin, S. R. Eddy, A. Krogh, and G. Mitchison. 1999. *Biological Sequence Analysis: Probabilistic Models of Proteins and Nucleic Acids*. Cambridge University Press.
- J. Esparza, S. Kiefer, and M. Luttenberger. 2010. Computing the least fixed point of positive polynomial systems. *SIAM J. Comput.* 39, 6, 2282–2355.
- J. Esparza, A. Kučera, and R. Mayr. 2006. Model checking probabilistic pushdown automata. *Log. Meth. Comput. Sci.* 2, 1, 1–31.
- K. Etessami, A. Stewart, and M. Yannakakis. 2012. Polynomial-time algorithms for multi-type branching processes and stochastic context-free grammars. In *Proceedings of the 44th ACM Symposium on Theory of Computing (STOC)*. (Full version is available at ArXiv:1201.2374.)
- K. Etessami, D. Wojtczak, and M. Yannakakis. 2010. Quasi-birth-death processes, tree-like QBDs, probabilistic 1-counter automata, and pushdown systems. *Perf. Eval.* 67, 9, 837–857.
- K. Etessami and M. Yannakakis. 2009. Recursive Markov chains, stochastic grammars, and monotone systems of nonlinear equations. *J. ACM* 56, 1.
- K. Etessami and M. Yannakakis. 2012. Model checking of recursive probabilistic systems. *ACM Trans. Comput. Log.* 13, 2, 12. (Conference versions in *TACAS'05* and *QEST'05*.)
- K. A. Hansen, M. Koucký, N. Lauritzen, P. B. Miltersen, and E. P. Tsigaridas. 2011. Exact algorithms for solving stochastic games: Extended abstract. In *Proceedings of STOC*. 205–214. (See full Arxiv version, arXiv:1202.3898.)
- T. E. Harris. 1963. *The Theory of Branching Processes*. Springer-Verlag.
- R. A. Horn and C. R. Johnson. 1985. *Matrix Analysis*. Cambridge University Press.
- E. Isaacson and H. B. Keller. 1966. *Analysis of Numerical Methods*. Wiley.
- P. Lancaster and M. Tismenetsky. 1985. *The Theory of Matrices* (2nd ed.). Academic Press.
- G. Latouche and V. Ramaswami. 1999. *Introduction to Matrix Analytic Methods in Stochastic Modeling*. ASA-SIAM Series on Statistics and Applied Probability.
- M.-J. Nederhof and G. Satta. 2008. Computing partition functions of PCFGs. *Res. Lang. Computation* 6, 2, 139–162.
- M. F. Neuts. 1981. *Matrix-Geometric Solutions in Stochastic Models: An algorithmic approach*. Johns Hopkins University Press.
- A. Stewart, K. Etessami, and M. Yannakakis. 2013. Upper bounds for Newton's method on monotone polynomial systems, and P-time model checking of probabilistic one-counter automata. In *Proceedings of the 25th International Conference on Computer Aided Verification (CAV)*, 495–510.
- D. Wojtczak and K. Etessami. 2007. PReMo: An analyzer for probabilistic recursive models. In *Proceedings of the 13th International Conference on Tools and Algorithms for the Construction and Analysis of Systems (TACAS)*. 66–71.

Received October 2014; revised April 2015; accepted June 2015