

# $O(m \log n)$ Time Algorithms for DFA Minimization and More

Antti Valmari

Tampere University of Technology

## Part I: DFA Minimization

- 1 Deterministic Finite Automata
- 2 Minimization of Deterministic ...
- 3 Block Splitting
- 4 Hopcroft's Ideas [1971]
- 5 Gries' Data Structures (Roughly)
- 6 More Recent ... Data Structure
- 7–9 DFA Min...in  $O(m \log n)$  Time

## Part II: and More


- 10 Bisimilarity
- 11 Minimal Bisimilar Graph
- 12 Paige-Tarjan  $O(m \log n)$  RCP Alg...
- 13 Extension to Bisimulation [2009, 2010]
- 14 Markov Chains and Markov Dec...
- 15 Conclusions

# Part I:

# DFA Minimization

# 1 Deterministic Finite Automata

- $D = (Q, \Sigma, \delta, \hat{q}, F)$ 
  - $Q$  = **states**
  - $\Sigma$  = **labels** (the alphabet)
  - $\delta$  = **transitions** (*partial function!*)
  - $\hat{q}$  = **initial state**
  - $F$  = **final states**



- Let  $n = |Q|$ ,  $m = |\delta| = |\text{defined transitions}|$ ,  $\alpha = |\Sigma| = |\text{available labels}|$ 
  - technical convenience assumption:  $n = O(m)$  (e.g.,  $n \leq 2m + 1$ )
- Let  $q \in Q$
- The **language** accepted by  $q$  is the set of strings of labels on the paths from  $q$  to final states
  - e.g., bottom middle state:  $\{aba, aaba, \dots, ba, baaba, \dots\}$
  - e.g., bottom right state:  $\{\varepsilon, aba, aaba, aaaba, \dots, \dots\}$
- Denote it with  $\mathcal{L}(q)$
- The language accepted by  $D$  is  $\mathcal{L}(D) = \mathcal{L}(\hat{q})$


## 2 Minimization of Deterministic Finite Automata

- **The minimization problem:**

Find the smallest DFA that accepts the same language as the given DFA.

- solution:



1. Remove **irrelevant** states and transitions. (textbook stuff)
  - those that are not reachable from  $\hat{q}$
  - those from which no final state can be reached (except  $\hat{q}$ )
2. Merge states that accept the same language.



- How can we test if  $\mathcal{L}(q_1) = \mathcal{L}(q_2)$ ?

### 3 Block Splitting

- States are partitioned into **blocks**
- $q_1$  and  $q_2$  go to different blocks only when it is certain that  $\mathcal{L}(q_1) \neq \mathcal{L}(q_2)$
- Initially blocks are  $F$  and  $Q \setminus F$  (or just one of them, if the other is  $\emptyset$ )
- Blocks are split as long as possible
- Reason for putting  $q_1$  and  $q_2$  to different blocks: for some label  $a$  and block  $B$ ,  $q_1$  has and  $q_2$  does not have an  $a$ -transition to a state in  $B$
- At most  $n - 1$  successful splittings ( $n = |\text{States}|$ )
- **Problem:** vulnerable to lots of work




–  $O(n^2)$  even if  $\alpha = 1$  and  $m = O(n)$

- **Cost of “useless” “little” work may be important**

## 4 Hopcroft's Ideas [1971]

- Clarified and improved by Gries [1973] (and, e.g., Knuutila [2001])
- Assumes that  $\delta$  is full
- **Idea:** Traverse transitions backwards
  - **splitter** = (block, label) =  $(B, a)$
  - process one splitter at a time
  - find the  $q$  such that  $\delta(q, a) \in B$ , move them to tentative new blocks
  - each block splits to backwards-encountered and others (if both non-empty) $\Rightarrow$  no futile scanning of states without relevant output transitions



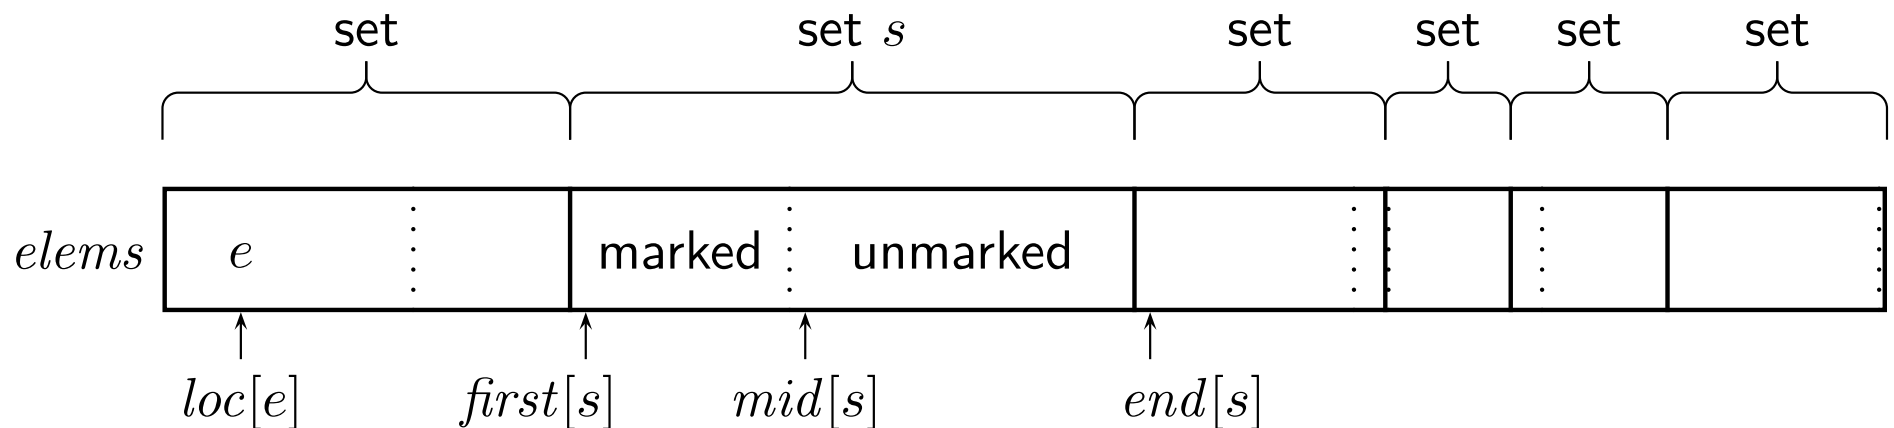
- **Idea:** If  $(B, a)$  has been used and  $B$  splits to  $B_1$  and  $B_2$ , then it suffices to use one of  $(B_1, a)$  and  $(B_2, a)$ 
  - use the “smaller” one (meaning of “smaller” is less trivial than it seems!) $\Rightarrow$  each state is used as a splitter state at most  $\log_2 n$  instead of  $n$  times
- Both ideas  $\Rightarrow$  running time in  $O(\alpha n \log n)$   $m \leq \alpha n$ , often  $m \ll \alpha n$

## 5 Gries' Data Structures (Roughly)

- Partition of  $Q$ 
  - for each block, there is a doubly linked list of the states in it
  - the block has a pointer to its main list and “tentative new” list
  - each state has a pointer to its block
  - the block knows its size (and the size of the “tentative new” list)
- Inverse transitions: for each  $q$  and  $a$ , the states  $q'$  such that  $\delta(q', a) = q$
- Worksets for temporary storage
  - unprocessed splitters, and pointers to them for each  $(B, a)$
  - backwards-encountered states — don't cut the branch on which you sit!
  - backwards-encountered blocks = **touched blocks**
- $\Theta(\alpha n)$  memory
  - while** there are unprocessed splitters  $(B, a)$  **do**
    - choose any and remove it from the workset
    - compute its backwards-encountered states
    - use backwards-encountered states to tentatively split blocks
    - for** each touched block  $B'$  **do**
      - split or reset back to earlier status
      - if**  $B'$  is split **then for**  $b \in \Sigma$  **do** update the  $(B', b)$

## 6 More Recent Refinable Partition Data Structure

- Maintains a partition of  $\{1, 2, \dots, N\}$ , for some  $N$
- Like Gries, constant time  $Mark(e)$  and amortized constant time  $Split(s)$



- Also  $sidx[e]$
- Of course, all arrays must be updated appropriately in each operation
- $Mark(e)$  swaps  $e$  with the first unmarked element and increments  $mid[s]$ 
  - **Trick 1 for the future:** returns set number iff all elements were unmarked
- **Trick 2 for the future:**  $Split(s)$  gives new block number to smaller half
  - earlier papers: marked states become the new block
  - does not affect amortized speed, as long as new  $\neq$  unmarked bigger half



## 7 DFA Minimization in $O(m \log n)$ Time 1/3

- Valmari & Lehtinen [2008], improvements Valmari [2010]
  - **Problem:** How to avoid spending excessive time scanning empty splitters?
    - empty  $(B, a) =$  no  $a$ -transition ends at  $B$
    - and how to avoid using and initializing  $\Theta(\alpha n)$  memory? $\Rightarrow$  cannot use, e.g., `in_trans[ state, label ]`
  - **Idea:** Non-empty splitters constitute a partition of transitions that can be maintained similarly to blocks
- $\Rightarrow$  Two refinable partitions in the same program
- $\mathcal{B} =$  blocks, partition of states
  - $\mathcal{C} =$  **cords**, partition of transitions
- Inverse transitions:  $In\_trans[q] = \{ (q_1, a, q_2) \in \delta \mid q_2 = q \}$ 
    - numbers of input transitions of  $q$  *in arbitrary order* $\Rightarrow$  no need for  $\Theta(\alpha n)$  data structures, easy to initialize
  - Worksets: only one,  $W$ 
    - e.g., array of integers used as a stack

## 8 DFA Minimization in $O(m \log n)$ Time 2/3

- Algorithm *in great detail*

```
 $c := 1; b := 2; W := \emptyset;$   
while  $c \leq |\mathcal{C}|$  do  
  use cord  $\#c$  to split blocks  
   $c := c + 1; W := \emptyset$   
  while  $b \leq |\mathcal{B}|$  do  
    use block  $\#b$  to split cords  
     $b := b + 1; W := \emptyset$ 
```

```
for  $\ell := \mathcal{C}.first[c]$  to  $\mathcal{C}.end[c] - 1$  do  
   $b' := \mathcal{B}.Mark( tail[ \mathcal{C}.elems[\ell] ] )$   
  if  $b' > 0$  then  $W := W \cup \{b'\}$   
for  $b' \in W$  do  $\mathcal{B}.Split(b')$ 
```

---

```
for  $\ell := \mathcal{B}.first[b]$  to  $\mathcal{B}.end[b] - 1$  do  
  for  $t \in In\_trans( \mathcal{B}.elems[\ell] )$  do  
     $c' := \mathcal{C}.Mark(t)$   
    if  $c' > 0$  then  $W := W \cup \{c'\}$   
for  $c' \in W$  do  $\mathcal{C}.Split(c')$ 
```

- Earlier Trick 1  
 $\Rightarrow b'$  and  $c'$  may be added to  $W$  without testing if they are already there
- Why is there no workset for unprocessed block-splitters and cord-splitters?
  - Hopcroft, Gries: either smaller or new half must be added to unprocessed
  - Trick 2: new number is given to smaller half  $\Rightarrow$  these cases are the same
  - unprocessed are chosen for processing in first in – first out order $\Rightarrow$  the unprocessed are always  $\{b, b + 1, \dots, |\mathcal{B}|\}$  and  $\{c, c + 1, \dots, |\mathcal{C}|\}$

## 9 DFA Minimization in $O(m \log n)$ Time 3/3


- **Problem:** Not enough time to initially sort the transitions!  $O(m \log m)$ 
  - we aim at so fast an algorithm that ordinary things become too slow
  - solution: counting sort + classification trick of Aho & al. [1974] exercise
  - practical average-time solution: hash table
  - engineer's: don't bother,  $O(m \log m)$  is not much worse than  $O(m \log n)$
- A prototype implementation
  - 590 lines of C++ (+ libraries for formatted i/o and error messages)
  - refinable partition data structure: 50 lines, other data structures: 20 lines
  - block splitting: 40 lines
  - Aho & al. [1974] & heapsorting as an alternative: 60 lines
  - removal of irrelevant states: 70 lines, other initialization: 80 lines
  - input and output: 130 lines
  - heapsort: 40 lines, range-checking array: 60 lines, main comment: 40 lines

⇒ Compared to general idea of the difficulty of programming fast DFA minimization, this is very simple

- 50 sec on a laptop when  $n = 10^5$ ,  $\alpha = 1000$ ,  $m = 5 \cdot 10^6$  (includes i/o)

# Part II: and More


# 10 Bisimilarity



- Weakest relation that satisfies:
  - $q_1 \sim q_2 \Rightarrow \text{label}(q_1) = \text{label}(q_2)$
  - $q_1 \sim q_2 \wedge (q_1, a, q'_1) \in \Delta_1 \Rightarrow \exists q'_2 : q'_1 \sim q'_2 \wedge (q_2, a, q'_2) \in \Delta_2$
  - $q_1 \sim q_2 \wedge (q_2, a, q'_2) \in \Delta_2 \Rightarrow \exists q'_1 : q'_1 \sim q'_2 \wedge (q_1, a, q'_1) \in \Delta_1$
  - $q_1^{\text{init}} \sim q_2^{\text{init}}$ , or  $\forall q_1 \in \text{Initials}_1 : \exists q_2 \in \text{Initials}_2 : q_1 \sim q_2$ , and ...
- Fundamental relation in concurrency theory
- Often much better “strong” equivalence than isomorphism
  - unifies states that only differ on  $x$ , if the next thing on  $x$  is  $x := 0$

# 11 Minimal Bisimilar Graph

- Nondeterministic version of DFA minimization
- Remove what is not reachable from initial states
- Fuse bisimilar states
- Applications
  - smaller graph for further processing
  - bisimilarity test
- **Problem** due to nondeterminism: **three-way splitting**



# 12 Paige–Tarjan $O(m \log n)$ RCP Algorithm [1987]



- **Relational coarsest partition problem:**  
Nondeterministic graph, **only one label** ( $\alpha = 1$ )
- **Idea: Compound blocks, “ $q$ - $B$ -counters”, and “ $q$ -counters”**

- Compound block  $\approx$   
union of blocks that has been used for splitting  
– let  $\hat{B}$  be the compound block that covers block  $B$
- Biggest block in a compound block  
need not be used in further splitting  
 $\Rightarrow$  each state is used in a splitter at most  $\log_2 n$  times

- **Problem:** How to implement three-way splitting?


$\Rightarrow$  Maintain, for each  $q$  and  $B$ , the number of transitions from  $q$  to  $B$

- Each  $B'$  has (at most) three kinds of sub-blocks
  - **left block:**  $\#(q, B) > 0 = \#(q, \hat{B} \setminus B)$
  - **middle block:**  $\#(q, B) > 0 < \#(q, \hat{B} \setminus B)$
  - **right block:**  $\#(q, B) = 0$



## 13 Extension to Bisimulation [2009, 2010]

- (Still RCP) cannot afford to represent counters that store 0  
⇒ data structure trickery!

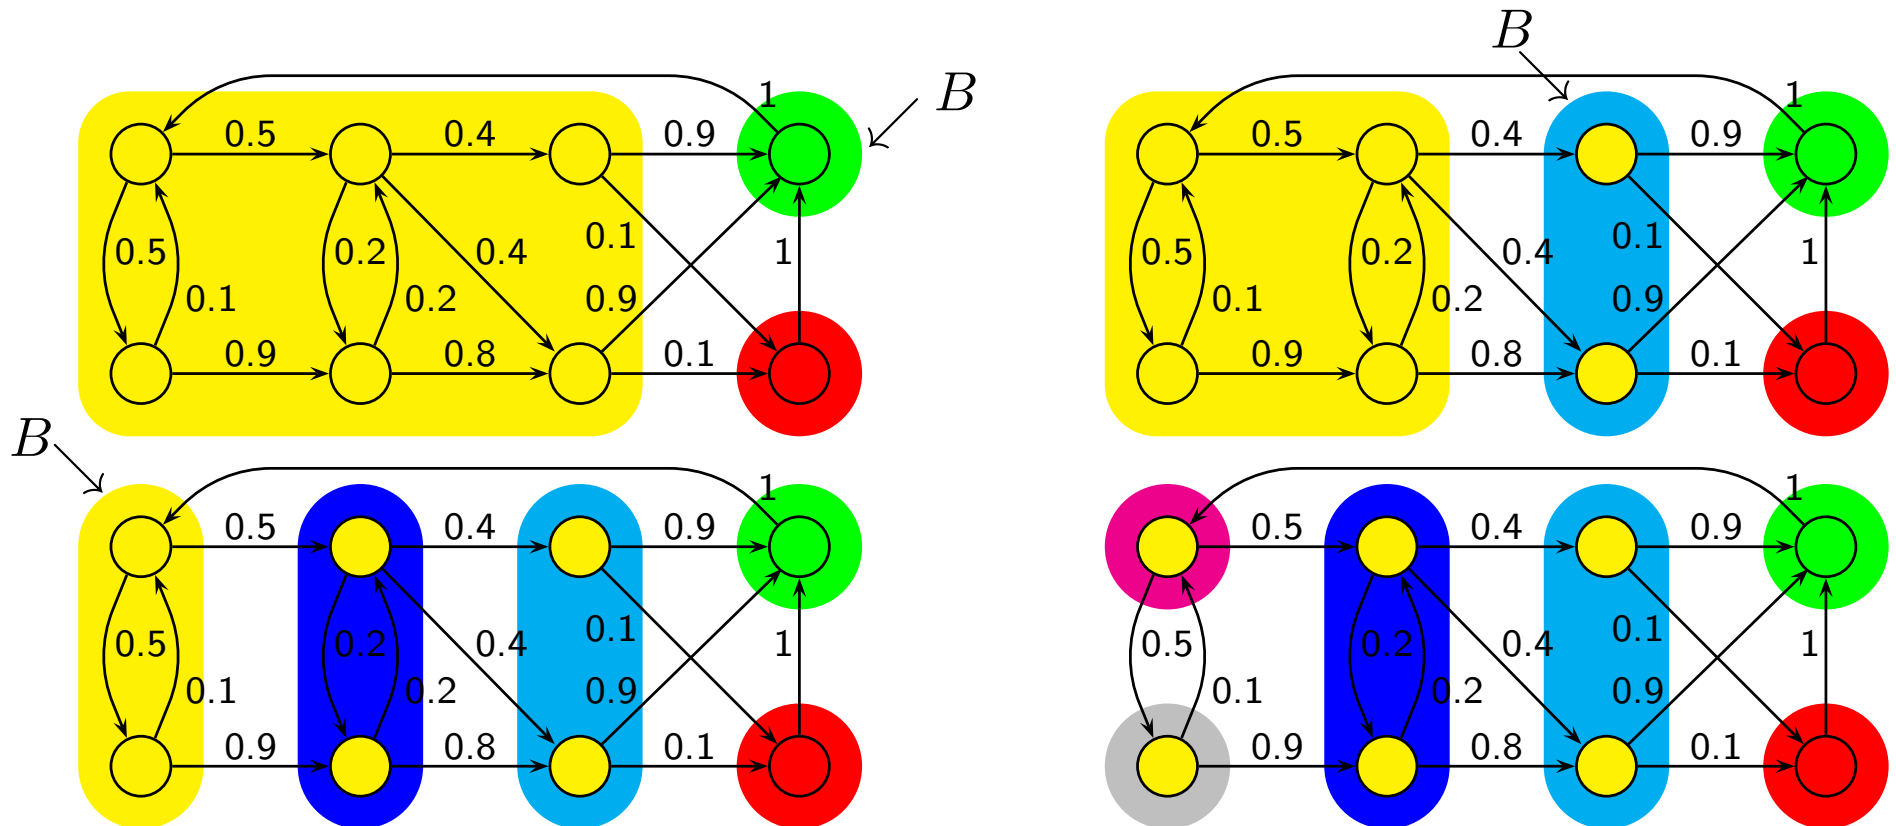


- Extend  $q$ - $B$ -counters to  $q$ - $a$ - $B$ -counters
- Use cords (and the year 2010 DFA tricks)
- Initialization: sort each cord according to start states
  - $O(m_a \log m_a)$ , where  $m_a \leq n^2$
- Splitting
  - distinguish left blocks from middle blocks by counter values
  - extract left blocks like with DFAs
  - update old and create new counter when extracting middle block
  - main loop and splitting of cords are like with DFAs
- Tricky details here and there and in the correctness proof, but it works!



# 14 Markov Chains and Markov Decision Processes

- These ideas have also been applied to state lumping of Markov Chains by Valmari and Franceschinis [2010]
  - would be a story of its own
  - finds use for the elegant but mostly useless majority candidate algorithm



- Everything has been put together in a program for minimizing Markov Decision Processes

# Part III:

# Conclusions

# 15 Conclusions

- $O(m \log n)$  time DFA, bisimulation, and MDP minimization are possible
  - such an algorithm for DFAs was found amazingly late, 1971  $\leftrightarrow$  2008
  - the other two problems are strictly more general
  - $O(m \log n)$  time Markov chain lumping was solved in 2003, but our results simplify it
- The breakthrough was the use of another partition, this time of transitions
- In the end the programs are relatively simple
  - we got rid of some data structures in earlier algorithms
- However, many tricky ideas had to be fine-tuned to make it all work
  - algorithms, correctness proofs, and programs
  - details **are** important for obtaining the promised performance!
- Hidden theme: representing a mapping where some result value is far more common than others
  - sparse mapping
  - avoid explicitly representing that value

## Part IV:

Thank you for attention.

Questions?