Revisiting the Equivalence Problem
for Finite Multitape Automata

James Worrell[*]

Department of Computer Science, University of Oxford, UK

Abstract. The decidability of determining equivalence of deterministic multitape automata (or transducers) was a longstanding open problem until it was resolved by Harju and Karhumäki in the early 1990s. Their proof of decidability yields a **co-NP** upper bound, but apparently not much more is known about the complexity of the problem. In this paper we give an alternative proof of decidability, which follows the basic strategy of Harju and Karhumäki but replaces their use of group theory with results on matrix algebras. From our proof we obtain a simple randomised algorithm for deciding equivalence of deterministic multitape automata, as well as automata with transition weights in the field of rational numbers. The algorithm involves only matrix exponentiation and runs in polynomial time for each fixed number of tapes. If the two input automata are inequivalent then the algorithm outputs a word on which they differ.

1 Introduction

One-way multitape finite automata were introduced in the seminal 1959 paper of Rabin and Scott [15]. Such automata (under various restrictions) are also commonly known as transducers—see Elgot and Mezei [6] for an early reference. A multitape automaton with k tapes accepts a k-ary relation on words. The class of relations recognised by deterministic automata coincides with the class of k-ary rational relations [6].

Two multitape automata are said to be equivalent if they accept the same relation. Undecidability of equivalence of non-deterministic automata is relatively straightforward [8]. However the deterministic case remained open for many years, until it was shown decidable by Harju and Karhumäki [9]. Their solution made crucial use of results about ordered groups—specifically that a free group can be endowed with a compatible order [13] and that the ring of formal power series over an ordered group with coefficients in a division ring and with well-ordered support is itself a division ring (due independently to Malcev [11] and Neumann [14]). Using these results [9] established the decidability of *multiplicity equivalence* of non-deterministic multitape automata, i.e., whether two non-deterministic multitape automata have the same number of accepting computations on each input. Decidability in the deterministic case (and, more

generally, the unambiguous case) follows immediately. We refer the reader to [16] for a self-contained account of the proof, including the underlying group theory.

Harju and Karhumäki did not address questions of complexity in [9]. However the existence of a **co-NP** *guess-and-check* procedure for deciding equivalence of deterministic multitape automata follows directly from [9, Theorem 8]. This theorem states that two inequivalent automata are guaranteed to differ on a tuple of words whose total length is at most the total number of states of the two automata. Such a tuple can be guessed, and it can be checked in polynomial time whether the tuple is accepted by one automaton and rejected by the other. In the special case of two-tape deterministic automata, a polynomial-time algorithm was given in [7], before decidability was shown in the general case.

A **co-NP** upper bound also holds for multiplicity equivalence of k-tape automata for each fixed k. However, as we observe below, if the number of tapes is not fixed, computing the number of accepting computations of a given nondeterministic multitape automata on a tuple of input words is **#P**-hard. Thus the guess-and-check method does not yield a **co-NP** procedure for multiplicity equivalence in general.

It is well-known that the equivalence problem for single-tape weighted automata with rational transition weights is solvable in polynomial time [18,19]. Now the decision procedure in [9] reduces multiplicity equivalence of multitape automata to equivalence of single-tape automata with transition weights in a division ring of power series over an ordered group. However the complexity of arithmetic in this ring seems to preclude an application of the polynomial-time procedures of [18,19]. Leaving aside issues of representing infinite power series, even the operation of multiplying a family of polynomials in two non-commuting variables yields a result with exponentially many monomials in the length of its input.

In this paper we give an alternative proof that multiplicity equivalence of multitape automata is decidable, which also yields new complexity bounds on the problem. We use the same basic idea as [9]—reduce to the single-tape case by enriching the set of transition weights. However we replace their use of power series on ordered groups with results about matrix algebras and Polynomial Identity rings (see Remark 1 for a more technical comparison). In particular, we use the Amitsur-Levitzki theorem concerning polynomial identities in matrix algebras. Our use of the latter is inspired by the work of [3] on non-commutative polynomial identity testing, and our starting point is a simple generalisation of the approach of [3] to what we call *partially commutative* polynomial identity testing.

Our construction for establishing decidability immediately yields a simple randomised algorithm for checking multiplicity equivalence of multitape automata (and hence also equivalence of deterministic automata). The algorithm involves only matrix exponentiation, and runs in polynomial time for each fixed number of tapes.

2   Partially Commutative Polynomial Identities

2.1   Matrix Algebras and Polynomial Identities

Let F be an infinite field. Recall that an F-algebra is a vector space over F equipped with an associative bilinear product that has an identity . Write $F\langle X\rangle$ for the free F-algebra over a set X. The elements of $F\langle X\rangle$ can be viewed as polynomials over a set of non-commuting variables X with coefficients in F. Each such polynomial is an F-linear combination of monomials, where each monomial is an element of X^*. The degree of a polynomial is the maximum of the lengths of its monomials.

Let A be an F-algebra and $f \in F\langle X\rangle$. If f evaluates to 0 for all valuations of its variables in A then we say that A satisfies the *polynomial identity $f = 0$*. For example, an algebra satisfies the polynomial identity $xy - yx = 0$ if and only if it is commutative. Note that since the variables x and y do not commute, the polynomial $xy - yx$ is not identically zero.

We denote by $M_n(F)$ the F-algebra of $n \times n$ matrices with coefficients in F. The Amitsur-Levitzki theorem [1,4] is a fundamental result about polynomial identities in matrix algebras.

Theorem 1 (Amitsur-Levitzki). *The algebra $M_n(F)$ of $n \times n$ matrices over a commutative ring F satisfies the polynomial identity*

$$\sum_{\sigma \in S_{2n}} x_{\sigma(1)} \ldots x_{\sigma(2n)} = 0 \,,$$

where the sum is over the $(2n)!$ elements of the symmetric group S_{2n}. Moreover $M_n(F)$ satisfies no identity of degree less than $2n$.

Given a finite set X of non-commuting variables, the *generic $n \times n$ matrix algebra $F_n\langle X\rangle$* is defined as follows. For each variable $x \in X$ we introduce a family of commuting indeterminates $\{t_{ij}^{(x)} : 1 \le i, j \le n\}$ and define $F_n\langle X\rangle$ to be the F-algebra of $n \times n$ matrices generated by the matrices $(t_{ij}^{(x)})$ for each $x \in X$. Then $F_n\langle X\rangle$ has the following universal property: any homomorphism from $F\langle X\rangle$ to a matrix algebra $M_n(R)$, with R an F-algebra, factors uniquely through the map $\Phi_n^X : F\langle X\rangle \to F_n\langle X\rangle$ given by $\Phi_n^X(x) = (t_{ij}^{(x)})$.

Related to the map Φ_n^X we also define an F-algebra homomorphism

$$\Psi_n^X : F\langle X\rangle \to M_n(F\langle t_{ij}^{(x)} \mid x \in X, 1 \le i, j \le n\rangle)$$

by

$$\Psi_n^X(x) = \begin{pmatrix} 0 & t_{12}^{(x)} & & \\ & \ddots & \ddots & \\ & & & t_{n-1,n}^{(x)} \\ & & & 0 \end{pmatrix}$$

where the matrix on the right has zero entries everywhere but along the super-diagonal.

2.2   Partially Commutative Polynomial Identities

In this section we introduce a notion of *partially commutative polynomial identity*. We first establish notation and recall some relevant facts about tensor products of algebras.

Write $A \otimes B$ for the tensor product of F-algebras A and B, and write $A^{\otimes k}$ for the k-fold tensor power of A. If A is an algebra of $a \times a$ matrices and B an algebra of $b \times b$ matrices, then we identify the tensor product $A \otimes B$ with the algebra of $ab \times ab$ matrices spanned by the matrices $M \otimes N$, $M = (m_{ij}) \in A$ and $N = (n_{ij}) \in B$, where

$$M \otimes N = \begin{pmatrix} m_{11}N & \cdots & m_{1a}N \\ \vdots & & \vdots \\ m_{a1}N & \cdots & m_{aa}N \end{pmatrix}$$

In particular we have $F^{\otimes k} = F$.

A *partially commuting set of variables* is a tuple $\boldsymbol{X} = (X_1, \ldots, X_k)$, where the X_i are disjoint sets. Write $F\langle \boldsymbol{X} \rangle$ for the tensor product $F\langle X_1 \rangle \otimes \cdots \otimes F\langle X_k \rangle$. We think of $F\langle \boldsymbol{X} \rangle$ as a set of polynomials in partially commuting variables. Intuitively two variables $x, y \in X_i$ do not commute, whereas $x \in X_i$ commutes with $y \in X_j$ if $i \neq j$. Note that if each X_i is a singleton $\{x_i\}$ then $F\langle \boldsymbol{X} \rangle$ is the familiar ring of polynomials in commuting variables $x_1, \ldots, x_k$. At the other extreme, if $k = 1$ then we recover the non-commutative case.

An arbitrary element $f \in F\langle \boldsymbol{X} \rangle = F\langle X_1 \rangle \otimes \cdots \otimes F\langle X_k \rangle$ can be written uniquely as a finite sum of distinct *monomials*, where each monomial is a tensor product of elements of X_1^*, $X_2^*, \ldots$, and X_k^*. Formally, we can write

$$f = \sum_{i \in I} \alpha_i (m_{i,1} \otimes \cdots \otimes m_{i,k}), \tag{1}$$

where $\alpha_i \in F$ and $m_{i,j} \in X_j^*$ for each $i \in I$ and $1 \leq j \leq k$. Thus we can identify $F\langle \boldsymbol{X} \rangle$ with the free F-algebra over the product monoid $X_1^* \times \ldots \times X_k^*$.

Define the *degree* of a monomial $m_1 \otimes \ldots \otimes m_k$ to be the total length $|m_1| + \ldots + |m_k|$ of its constituent words. The degree of a polynomial is the maximum of the degrees of its constituent monomials.

Let $\boldsymbol{A} = (A_1, \ldots, A_k)$ be a k-tuple of F-algebras. A *valuation* of $F\langle \boldsymbol{X} \rangle$ in $\boldsymbol{A}$ is a tuple of functions $\boldsymbol{v} = (v_1, \ldots, v_k)$, where $v_i : X_i \to A_i$. Each v_i extends uniquely to an F-algebra homomorphism $\widetilde{v}_i : F\langle X_i \rangle \to A_i$, and we define the map $\widetilde{\boldsymbol{v}} : F\langle \boldsymbol{X} \rangle \to A_1 \otimes \ldots \otimes A_k$ by $\widetilde{v} = \widetilde{v}_1 \otimes \ldots \otimes \widetilde{v}_k$. Often we will abuse terminology slightly and speak of a valuation of $F\langle \boldsymbol{X} \rangle$ in $A_1 \otimes \cdots \otimes A_k$. Given $f \in F\langle \boldsymbol{X} \rangle$, we say that $\boldsymbol{A}$ satisfies the partially commutative identity $f = 0$ if $\widetilde{\boldsymbol{v}}(f) = 0$ for all valuations $\boldsymbol{v}$.

Next we introduce two valuations that will play an important role in the subsequent development. Recall that given a set of non-commuting variables X, we have a map $\Phi_n^X : F\langle X \rangle \to F_n\langle X \rangle$ from the free F-algebra to the generic n-dimensional matrix algebra. We now define a valuation

$$\Phi_n^{\boldsymbol{X}} : F\langle \boldsymbol{X} \rangle \longrightarrow F_n\langle X_1 \rangle \otimes \cdots \otimes F_n\langle X_k \rangle \tag{2}$$

by $\Phi_n^{\boldsymbol{X}} = \Phi_n^{X_1} \otimes \cdots \otimes \Phi_n^{X_k}$. Likewise we define

$$\Psi_n^{\boldsymbol{X}} : F\langle\boldsymbol{X}\rangle \longrightarrow M_n(F\langle t_{ij}^{(x)} \mid x \in X_1, 1 \le i, j \le n\rangle) \otimes \cdots$$

$$\otimes M_n(F\langle t_{ij}^{(x)} \mid x \in X_k, 1 \le i, j \le n\rangle)$$

by $\Psi_n^{\boldsymbol{X}} = \Psi_n^{X_1} \otimes \cdots \otimes \Psi_n^{X_k}$. We will usually elide the superscript $\boldsymbol{X}$ from $\Phi_n^{\boldsymbol{X}}$ and $\Psi_n^{\boldsymbol{X}}$ when it is clear from the context.

The following result generalises (part of) the Amitsur-Levitzki theorem, by giving a lower bound on the degrees of partially polynomial identities holding in tensor products of matrix algebras.

Proposition 1. *Let $f \in F\langle\boldsymbol{X}\rangle$ and let L be a field extending F. Then the following are equivalent: (i) The partially commutative identity $f = 0$ holds in $M_n(L) \otimes_F \cdots_F \otimes M_n(L)$; (ii) $\Phi_n(f) = 0$. Moreover, if f has degree strictly less than n then (i) and (ii) are both equivalent to (iii) $\Psi_n(f) = 0$; and (iv) f is identically 0 in $F\langle\boldsymbol{X}\rangle$.*

Proof. The implication (ii) $\Rightarrow$ (i) follows from the fact that any valuation from $F\langle\boldsymbol{X}\rangle$ to $M_n(L) \otimes_F \cdots_F \otimes M_n(L)$ factors through Φ_n. To see that (i) $\Rightarrow$ (ii), observe that $\Phi_n(f)$ is an $n^k \times n^k$ matrix in which each entry is a polynomial in the commuting variables $t_{ij}^{(x)}$. Condition (i) implies in particular that each such polynomial evaluates to 0 for all valuations of its variables in F. Since F is an infinite field, it must be that each such polynomial is identically zero, i.e., (ii) holds.

The implications (ii) $\Rightarrow$ (iii) and (iv) $\Rightarrow$ (i) are both straightforward, even without the degree restriction on f.

Finally we show that (iii) $\Rightarrow$ (iv). Let $m_1 \otimes \ldots \otimes m_k$ be a monomial in $F\langle\boldsymbol{X}\rangle$, where $m_i = m_{i,1} \ldots m_{i,l_i} \in X_i^*$ has length $l_i < n$. Then $\Psi_n(m_1 \otimes \cdots \otimes m_k)$ is an $n^k \times n^k$ matrix whose first row has a single non-zero entry, which is the monomial

$$t_{12}^{(m_{1,1})} \cdots t_{l_1,l_1+1}^{(m_{1,l_1})} \cdots t_{12}^{(m_{k,1})} \cdots t_{l_k,l_k+1}^{(m_{k,l_k})} \tag{3}$$

at index $(1,\ldots,1), (l_1 + 1, \ldots, l_k + 1)$.

It follows that Ψ_n maps the set of monomials in $F\langle\boldsymbol{X}\rangle$ of degree less than n injectively into a linearly independent set of matrices, each with a single monomial entry over the commuting indeterminates $t_{ij}^{(x)}$. Condition (iv) immediately follows. $\qquad\square$

The hypothesis that f have degree less than n in Proposition 1 can be weakened somewhat, but is sufficient for our purposes.

2.3   Division Rings and Ore Domains

A ring R with no zero divisors is a *domain*. If moreover each non-zero element of R has a two-sided multiplicative inverse, then we say that R is a *division*

ring (also called a *skew field*). A domain R is a *(right) Ore domain* if for all $a, b \in R \setminus \{0\}$, $aR \cap bR \neq 0$. The significance of this notion is that an Ore domain can be embedded in a division ring of fractions [4, Corollary 7.1.6], something that need not hold for an arbitrary domain. If the Ore condition fails then it can easily be shown that the subalgebra of R generated by a and b is free on a and b. It follows that a domain R that satisfies some polynomial identity is an Ore domain [4, Corollary 7.5.2].

Proposition 2. *The tensor product of generic matrix algebras $F_n\langle X_1 \rangle \otimes \cdots \otimes F_n\langle X_k \rangle$ is an Ore domain for each $n \in \mathbb{N}$.*

Proof (sketch). We give a proof sketch here, deferring the details to Appendix A.

By the Amitsur-Levitzki theorem, $F_n\langle X_1 \rangle \otimes \cdots \otimes F_n\langle X_n \rangle$ satisfies a polynomial identity. Thus it suffices to show that $F_n\langle X_1 \rangle \otimes \cdots \otimes F_n\langle X_n \rangle$ is a domain for each n. Now it is shown in [4, Proposition 7.7.2] that $F_n\langle X \rangle$ is a domain for each n and set of variables X. While the tensor product of domains need not be a domain (e.g., $\mathbb{C} \otimes_{\mathbb{R}} \mathbb{C} \cong \mathbb{C} \times \mathbb{C}$), the proof in [4] can be adapted *mutatis mutandis* to show that $F\langle X_1 \rangle \otimes \cdots \otimes F\langle X_k \rangle$ is also a domain.

To prove the latter, it suffices to find central simple F-algebras $D_1, \ldots, D_k$, each of degree n, such that the k-fold tensor product $D \otimes_F \cdots \otimes_F D$ is a domain. Such an example can be found, e.g., in [17, Proposition 1.1]. Then, using the fact that $D \otimes_F L \cong M_n(L)$ for any algebraically closed extension field of F, one can infer that $F_n\langle X_1 \rangle \otimes \cdots \otimes F_n\langle X_k \rangle$ is also a domain. $\square$

3   Multitape Automata

Let $\Sigma = (\Sigma_1, \ldots, \Sigma_k)$ be a tuple of finite alphabets. We denote by S the product monoid $\Sigma_1^* \times \cdots \times \Sigma_k^*$. Define the length of $s = (w_1, \ldots, w_k) \in S$ to be $|s| = |w_1| + \ldots + |w_k|$ and write $S^{(l)}$ for the set of elements of S of length l. A *multitape automaton* is a tuple $A = (\Sigma, Q, E, Q_0, F)$, where Q is a set of *states*, $E \subseteq Q \times S^{(1)} \times Q$ is a set of *edges*, $Q_0 \subseteq Q$ is a set of *initial states*, and $Q_f \subseteq Q$ is a set of *final states*. A *run* of A from state q_0 to state q_m is a finite sequence of edges $\rho = e_1 e_2 \ldots e_m$ such that $e_i = (q_{i-1}, s_i, q_i)$. The *label* of ρ is the product $s_1 s_2 \ldots s_m \in S$. Define the *multiplicity* $A(s)$ of an input $s \in S$ to be the number of runs with label s such that $q_0 \in Q_0$ and $q_m \in Q_f$. An automaton is *deterministic* if each state reads letters from a single tape and has a single transition for every input letter. Thus a deterministic automaton has a single run on each input $s \in S$.

3.1   Multiplicity Equivalence

We say that two automata A and B over the same alphabet are *multiplicity equivalent* if $A(s) = B(s)$ for all $s \in S$. The following result implies that multiplicity equivalence of multitape automata is decidable.

Theorem 2 (Harju and Karhumäki). *Given automata A and B with n states in total, A and B are equivalent if and only if $A(s) = B(s)$ for all $s \in S$ of length at most $n - 1$.*

Theorem 2 immediately yields a **co-NP** bound for checking language equivalence of deterministic multitape automata. Given two inequivalent automata A and B, a distinguishing input s can be guessed, and it can be verified in polynomial time that only one of A and B accepts s. A similar idea also gives a **co-NP** bound for multiplicity equivalence in case the number of tapes is fixed. In general we note however that the *evaluation problem*—given an automaton A and input s, compute $A(s)$—is #**P**-complete. Thus it is not clear that the **co-NP** upper bound applies to the multiplicity equivalence problem without bounding the number of tapes.

Proposition 3. *The evaluation problem for multitape automata is #**P**-complete.*

Proof. Membership in #**P** follows from the observation that a non-deterministic polynomial-time algorithm can enumerate all possible runs of an automaton A on an input $s \in S$.

The proof of #**P**-hardness is by reduction from #SAT, the problem of counting the number of satisfying assignments of a propositional formula. Consider such a formula φ with k variables, each with fewer than n occurrences. We define a k-tape automaton A, with each tape having alphabet $\{0, 1\}$, and consider as input the k-tuple $s = ((01)^n, \ldots, (01)^n)$. The automaton A is constructed such that its runs on input s are in one-to-one correspondence with satisfying assignments of φ. Each run starts with the automaton reading the symbol 0 from a non-deterministically chosen subset of its tapes, corresponding to the set of false variables. Thereafter it evaluates the formula φ by repeatedly guessing truth values of the propositional variables. If the i-th variable is guessed to be true then the automaton reads 01 from the i-th tape; otherwise it reads 10 from the i-th tape. The final step is to read the symbol 1 from a non-deterministically chosen subset of the input tapes—again corresponding to the set of false variables. The consistency of the guesses is ensured by the requirement that the automaton have read s by the end of the computation. □

3.2 Decidability

We start by recalling from [9] an equivalence-respecting transformation from multitape automata to single-tape weighted automata.

Recall that a single-tape automaton on a unary alphabet with transition weights in a ring R consists of a set of *states* $Q = \{q_1, \ldots, q_n\}$, *initial states* $Q_0 \subseteq Q$, *final states* $Q_f \subseteq Q$, and *transition matrix* $M \in M_n(R)$. Given such an automaton, define the *initial-state vector* $\boldsymbol{\alpha} \in R^{1 \times n}$ and *final-state vector* $\boldsymbol{\eta} \in R^{n \times 1}$ respectively by

$$\alpha_i = \begin{cases} 1 & \text{if } q_i \in Q_0 \\ 0 & \text{otherwise} \end{cases} \quad \text{and} \quad \eta_i = \begin{cases} 1 & \text{if } q_i \in Q_f \\ 0 & \text{otherwise} \end{cases}$$

Then $\boldsymbol{\alpha} M^l \boldsymbol{\eta}$ is the weight of the (unique) input word of length l.

Consider a k-tape automaton $A = (\boldsymbol{\Sigma}, Q, E, Q_0, Q_f)$, where $\boldsymbol{\Sigma} = (\Sigma_1, \ldots, \Sigma_k)$, and write $S = \Sigma_1^* \times \cdots \times \Sigma_k^*$. Recall the ring of polynomials

$$F\langle \boldsymbol{\Sigma} \rangle = F\langle \Sigma_1 \rangle \otimes \cdots \otimes F\langle \Sigma_n \rangle,$$

as defined in Section 2. Recall also that we can identify the monoid S with the set of monomials in $F\langle \boldsymbol{\Sigma} \rangle$, where $(w_1, \ldots, w_k) \in S$ corresponds to $w_1 \otimes \cdots \otimes w_k$—indeed $F\langle \boldsymbol{\Sigma} \rangle$ is the free F-algebra on S.

We derive from A an $F\langle \boldsymbol{\Sigma} \rangle$-weighted automaton $\widetilde{A}$ (with a single tape and unary input alphabet) that has the same sets of states, initial states, and final states as A. We define the transition matrix M of $\widetilde{A}$ by combining the different transitions of A into a single matrix with entries in $F\langle \boldsymbol{\Sigma} \rangle$. To this end, suppose that the set of states of A is $Q = \{q_1, \ldots, q_n\}$. Define the matrix $M \in M_n(F\langle \boldsymbol{\Sigma} \rangle)$ by $M_{ij} = \sum_{(q_i, s, q_j) \in E} s$ for $1 \leq i, j \leq n$.

Let $\boldsymbol{\alpha}$ and $\boldsymbol{\eta}$ be the respective initial- and final-state vectors of $\widetilde{A}$. Then the following proposition is straightforward. Intuitively it says that the weight of the unary word of length l in $\widetilde{A}$ represents the language of all length-l tuples accepted by A.

Proposition 4. *For all $l \in \mathbb{N}$ we have $\boldsymbol{\alpha} M^l \boldsymbol{\eta} = \sum_{s \in S^{(l)}} A(s) \cdot s$.*

Now consider two k-tape automata A and B. Let the weighted single-tape automata derived from A and B have respective transition matrices M_A and M_B, initial-state vectors $\boldsymbol{\alpha}_A$ and $\boldsymbol{\alpha}_B$, and final-state vectors $\boldsymbol{\eta}_A$ and $\boldsymbol{\eta}_B$. We combine the latter into a single weighted automaton with transition matrix M, initial-state vector $\boldsymbol{\alpha}$, and final-state vector $\boldsymbol{\eta}$, respectively defined by:

$$\boldsymbol{\alpha} = (\boldsymbol{\alpha}_A \ \boldsymbol{\alpha}_B) \qquad M = \begin{pmatrix} M_A & 0 \\ 0 & M_B \end{pmatrix} \qquad \boldsymbol{\eta} = \begin{pmatrix} \boldsymbol{\eta}_A \\ -\boldsymbol{\eta}_B \end{pmatrix}$$

Proposition 5. *Automata A and B are multiplicity equivalent if and only if $\boldsymbol{\alpha} M^l \boldsymbol{\eta} = 0$ for $l = 0, 1, \ldots, n - 1$, where n is the total number of states of the two automata.*

Proof. Since S is a linearly independent subset of $F\langle \boldsymbol{\Sigma} \rangle$, from Proposition 4 it follows that A and B are multiplicity equivalent just in case $\boldsymbol{\alpha}_A (M_A)^l \boldsymbol{\eta}_A = \boldsymbol{\alpha}_B (M_B)^l \boldsymbol{\eta}_B$ for all $l \in \mathbb{N}$. The latter is clearly equivalent to $\boldsymbol{\alpha} M^l \boldsymbol{\eta} = 0$ for all $l \in \mathbb{N}$. It remains to show that we can check equivalence by looking only at exponents l in the range $0, 1, \ldots, n - 1$.

Suppose that $\boldsymbol{\alpha} M^i \boldsymbol{\eta} = 0$ for $i = 0, \ldots, n - 1$. We show that $\boldsymbol{\alpha} M^l \boldsymbol{\eta} = 0$ for an arbitrary $l \geq n$.

Consider the map $\Phi_l : F\langle \boldsymbol{\Sigma} \rangle \to F_l \langle \Sigma_1 \rangle \otimes \cdots \otimes F_l \langle \Sigma_k \rangle$, as defined in (2). Observe that $\boldsymbol{\alpha} M^l \boldsymbol{\eta}$ is a polynomial expression in $F\langle \boldsymbol{\Sigma} \rangle$ of degree at most l. Therefore by Proposition 1 ((ii) $\Leftrightarrow$ (iv)), to show that $\boldsymbol{\alpha} M^l \boldsymbol{\eta} = 0$ it suffices to show that

$$\Phi_l(\boldsymbol{\alpha} M^l \boldsymbol{\eta}) = 0. \tag{4}$$

Let us write $\Phi_l(M)$ for the pointwise application of Φ_l to the matrix M, so that $\Phi_l(M)$ is an $n \times n$ matrix, each of whose entries is an $n^k \times n^k$ matrix belonging to $F_l\langle \Sigma_1 \rangle \otimes \cdots \otimes F_l \langle \Sigma_k \rangle$. Since Φ_l is a homomorphism and $\boldsymbol{\alpha}$ and $\boldsymbol{\eta}$ are integer vectors, (4) is equivalent to:

$$\boldsymbol{\alpha}\,\Phi_l(M)^l\,\boldsymbol{\eta}\,=\,0\,. \tag{5}$$

Recall from Proposition 2 that the tensor product of generic matrix algebras $F_l\langle\Sigma_1\rangle\otimes\cdots\otimes F_l\langle\Sigma_k\rangle$ is an Ore domain and hence can be embedded in a division ring. Now a standard result about single-tape weighted automata with transition weights in a division ring is that such an automaton with n states is equivalent to the zero automaton if and only if it assigns zero weight to all words of length n (see [5, pp143–145] and [18]). Applying this result to the unary weighted automaton defined by $\boldsymbol{\alpha}$, M, and $\boldsymbol{\eta}$, we see that (5) is implied by

$$\boldsymbol{\alpha}\,\Phi_l(M)^i\,\boldsymbol{\eta}\,=\,0 \qquad i=0,1,\ldots,n-1\,. \tag{6}$$

But, since Φ_l is a homomorphism, (6) is implied by

$$\boldsymbol{\alpha}M^i\boldsymbol{\eta}=0 \qquad i=0,1,\ldots,n-1\,. \tag{7}$$

This concludes the proof.   □

Theorem 2 immediately follows from Proposition 5.

Remark 1. The difference between our proof of Theorem 2 and the proof in [9] is that we consider a family of homomorphisms of $F\langle\boldsymbol{\Sigma}\rangle$ into Ore domains of matrices—the maps Φ_l—rather than a single "global" embedding of $F\langle\boldsymbol{\Sigma}\rangle$ into a division ring of power series over a product of free groups. None of the maps Φ_l is an embedding, but it suffices to use the lower bound on the degrees of polynomial identities in Proposition 1 in lieu of injectivity. On the other hand, the fact that $F_l\langle\Sigma_1\rangle\otimes\cdots\otimes F_l\langle\Sigma_k\rangle$ satisfies a polynomial identity makes it relatively straightforward to exhibit an embedding of the latter into a division ring. As we now show, this approach leads directly to a very simple randomised polynomial-time algorithm for solving the equivalence problem.

3.3   Randomised Algorithm

Proposition 5 reduces the problem of checking multiplicity equivalence of multi-tape automata A and B to checking the partially commutative identities $\boldsymbol{\alpha}M^l\boldsymbol{\eta}=0$, $l=0,1,\ldots,n-1$ in $F\langle\boldsymbol{\Sigma}\rangle$. Since each identity has degree less than n, applying Proposition 1 ((iii) ⇔ (iv)) we see that A and B are equivalent if and only if

$$\boldsymbol{\alpha}\,\Psi_n(M)^l\,\boldsymbol{\eta}=0 \qquad l=0,1,\ldots,n-1\,. \tag{8}$$

Each equation $\boldsymbol{\alpha}\Psi_n(M)^l\boldsymbol{\eta}=0$ in (8) asserts the zeroness of an $n^k\times n^k$ matrix of polynomials in the commuting variables $t_{ij}^{(x)}$, with each polynomial having degree less than n. Suppose that $\boldsymbol{\alpha}\Psi_n(M)^l\boldsymbol{\eta}\neq 0$ for some l—say the matrix entry with index $((1,\ldots,1),(l_1+1,\ldots,l_k+1))$ contains a monomial with non-zero coefficient. By (3) such a monomial determines a term $s\in\Sigma_1^{l_1}\times\cdots\times\Sigma_k^{l_k}$ with non-zero coefficient in $\boldsymbol{\alpha}M^l\boldsymbol{\eta}$, and by Proposition 4 we have $A(s)\neq B(s)$.

We can verify each polynomial identity in (8), outputting a monomial of any non-zero polynomial, using a classical identity testing procedure based on the *isolation lemma* of [12].

Lemma 1 ([12]). *There is a randomised polynomial-time algorithm that inputs a multilinear polynomial $f(x_1, \ldots, x_m)$, represented as an algebraic circuit, and either outputs a monomial of f or that f is zero. Moreover the algorithm is always correct if f is the zero polynomial and is correct with probability at least $1/2$ if f is non-zero.*

The idea behind the algorithm described in Lemma 1 is to choose a weight $w_i \in \{1, \ldots, 2m\}$ for each variable x_i of f independently and uniformly at random. Defining the weight of a monomial $x_{i_1} \ldots x_{i_t}$ to be $w_{i_1} + \ldots + w_{i_t}$, then with probability at least $1/2$ there is a unique minimum-weight monomial. The existence of a minimum-weight monomial can be detected by computing the polynomial $g(y) = f(y^{w_1}, \ldots, y^{w_k})$, since a monomial with weight w in f yields a monomial of degree w in g. Using similar ideas one can moreover determine the composition of a minimum-weight monomial in f.

Applying Lemma 1 we obtain our main result:

Theorem 3. *Let k be fixed. Then multiplicity equivalence of k-tape automata can be decided in randomised polynomial time. Moreover there is a randomised polynomial algorithm for the function problem of computing a distinguishing input given two inequivalent automata.*

The reason for the requirement that k be fixed is because the dimension of the entries of the transition matrix M, and thus the number of polynomials to be checked for equality, depends exponentially on k.

The above use of the isolation technique generalises [10], where it is used to generate counterexample words of weighted single-tape automata. A very similar application in [2] occurs in the context of identity testing for non-commutative algebraic branching programs.

4 Conclusion

We have given a simple randomised algorithm for deciding language equivalence of deterministic multitape automata and multiplicity equivalence of nondeterministic automata. The algorithm arises directly from algebraic constructions used to establish decidability of the problem, and runs in polynomial time for each fixed number of tapes. We leave open the question of whether there is a deterministic polynomial-time algorithm for deciding the equivalence of deterministic and weighted multitape automata with a fixed number of tapes. (Recall that the 2-tape case is already known to be in polynomial time [7].) We also leave open whether there is a deterministic or randomised polynomial time algorithm for solving the problem in case the number of tapes is not fixed.

A Proof of Proposition 2

We first recall a construction of a *crossed product division algebra* from [17, Proposition 1.1]. Let $z_1, \ldots, z_k$ be commuting indeterminates and write $F =$

$\mathbb{Q}(z_1^n, \ldots, z_k^n)$ for the field of rational functions obtained by adjoining $z_1^n, \ldots, z_k^n$ to $\mathbb{Q}$. Furthermore, let K/F be a field extension whose Galois group is generated by commuting automorphisms $\sigma_1, \ldots, \sigma_k$, each of order n, which has fixed field F. (Such an extension can easily be constructed by adjoining extra indeterminates to F, and having the σ_i be suitable permutations of the new indeterminates.) For each i, $1 \leq i \leq k$, write K_i for the subfield of K that is fixed by each σ_j for $j \neq i$; then define D_i to be the F-algebra generated by K_i and z_i such that $az_i = z_i\sigma_i(a)$ for all $a \in K_i$. Then each D_i is a simple algebra of dimension n^2 over its centre F. It is shown in [17, Proposition 1.1] that the tensor product $D_1 \otimes_F \cdots \otimes_F D_k$ can be characterised as the localisation of an iterated skew polynomial ring—and is therefore a domain.

The following two propositions are straightforward adaptations of [4, Proposition 7.5.5.] and [4, Proposition 7.7.2] to partially commutative identities.

Proposition 6. *Let $f \in F\langle X_1 \rangle \otimes \cdots \otimes F\langle X_k \rangle$. If the partially commutative identity $f = 0$ holds in $D_1 \otimes_F \cdots \otimes_F D_k$ then it also holds in $(D_1 \otimes_F L) \otimes_F \cdots \otimes_F (D_k \otimes_F L)$ for any extension field L of F.*

Proof. Noting that the D_i are all isomorphic as F-algebras, let $\{e_1, \ldots, e_{n^2}\}$ be a basis of each D_i over its centre F. For each variable x appearing in f, introduce commuting indeterminates t_{xj}, $1 \leq j \leq n^2$, and write $x = \sum_{j=1}^{n^2} t_{xj}e_j$. Then we can express f in the form

$$f = \sum_{\nu \in \{1, \ldots, n^2\}^k} f_\nu \cdot (e_{\nu(1)} \otimes \cdots \otimes e_{\nu(k)}), \tag{9}$$

where $f_\nu \in F\langle t_{xj} : x \in X_1, 1 \leq j \leq n^2 \rangle \otimes_F \cdots \otimes_F F\langle t_{xj} : x \in X_k, 1 \leq j \leq n^2 \rangle$.

By assumption, each f_ν evaluates to 0 for all values of the t_{xj} in F. Since F is an infinite field it follows that each f_ν must be identically zero. Now we can also regard $\{e_1, \ldots, e_{n^2}\}$ as a basis for $D_i \otimes_F L$ over L. Then by (9), $f = 0$ also on $(D_1 \otimes_F L) \otimes_F \cdots \otimes_F (D_k \otimes_F L)$. □

Proposition 7. *$F_n\langle X_1 \rangle \otimes \cdots \otimes F_n\langle X_k \rangle$ is a domain.*

Proof. Recall that if L is an algebraically closed field extension of F, then we have $D_i \otimes_F L \cong M_n(L)$ for each i. By Proposition 6 it follows that an identity $f = 0$ holds in $D_1 \otimes_F \cdots \otimes_F D_k$ if and only if it holds in $M_n(L) \otimes_F \cdots \otimes_F M_n(L)$. But by Proposition 1 the latter holds if and only if $\Phi_n(f)$ is identically zero.

To prove the proposition it will suffice to show that the image of Φ_n contains no zero divisors, since the latter is a surjective map. Now given $f, g \in F\langle X_1 \rangle \otimes \cdots \otimes F\langle X_k \rangle$ with $\Phi_n(fg) = 0$, we have that $D_1 \otimes_F \cdots \otimes_F D_k$ satisfies the identity $fg = 0$. Since $D_1 \otimes_F \cdots \otimes_F D_k$ is a domain, it follows that it satisfies the identity $fhg = 0$ for any h in $F\langle X_1 \rangle \otimes \cdots \otimes F\langle X_k \rangle$. But now $M_n(L) \otimes_F \cdots \otimes_F M_n(L)$ satisfies the identity $fhg = 0$ for any h. Since h can take the value of an arbitrary matrix (in particular, any matrix unit) it follows that $M_n(L) \otimes_F \cdots \otimes_F M_n(L)$ satisfies either the identity $f = 0$ or $g = 0$, and so, by Proposition 1 again, either $\Phi_n(f) = 0$ or $\Phi_n(g) = 0$. □

References

1. Amitsur, S.A., Levitzki, J.: Minimal identities for algebras. Proceedings of the American Mathematical Society 1, 449–463 (1950)
2. Arvind, V., Mukhopadhyay, P.: Derandomizing the isolation lemma and lower bounds for circuit size. In: Goel, A., Jansen, K., Rolim, J.D.P., Rubinfeld, R. (eds.) APPROX and RANDOM 2008. LNCS, vol. 5171, pp. 276–289. Springer, Heidelberg (2008)
3. Bogdanov, A., Wee, H.: More on noncommutative polynomial identity testing. In: IEEE Conference on Computational Complexity, pp. 92–99. IEEE Computer Society (2005)
4. Cohn, P.M.: Further Algebra and Applications. Springer (2003)
5. Eilenberg, S.: Automata, Languages, and Machines, vol. A. Academic Press (1974)
6. Elgot, C.C., Mezei, J.E.: Two-sided finite-state transductions (abbreviated version). In: SWCT (FOCS), pp. 17–22. IEEE Computer Society (1963)
7. Friedman, E.P., Greibach, S.A.: A polynomial time algorithm for deciding the equivalence problem for 2-tape deterministic finite state acceptors. SIAM J. Comput. 11(1), 166–183 (1982)
8. Griffiths, T.V.: The unsolvability of the equivalence problem for ϵ-free nondeterministic generalized machines. J. ACM 15(3), 409–413 (1968)
9. Harju, T., Karhumäki, J.: The equivalence problem of multitape finite automata. Theor. Comput. Sci. 78(2), 347–355 (1991)
10. Kiefer, S., Murawski, A., Ouaknine, J., Wachter, B., Worrell, J.: On the complexity of equivalence and minimisation for Q-weighted automata. Logical Methods in Computer Science 9 (2013)
11. Malcev, A.I.: On the embedding of group algebras in division algebras. Dokl. Akad. Nauk 60, 1409–1501 (1948)
12. Mulmuley, K., Vazirani, U.V., Vazirani, V.V.: Matching is as easy as matrix inversion. In: STOC, pp. 345–354 (1987)
13. Neumann, B.H.: On ordered groups. Amer. J. Math. 71, 1–18 (1949)
14. Neumann, B.H.: On ordered division rings. Trans. Amer. Math. Soc. 66, 202–252 (1949)
15. Rabin, M., Scott, D.: Finite automata and their decision problems. IBM Journal of Research and Development 3(2), 114–125 (1959)
16. Sakarovich, J.: Elements of Automata Theory. Cambridge University Press (2003)
17. Saltman, D.: Lectures on Division Algebras. American Math. Soc. (1999)
18. Schützenberger, M.-P.: On the definition of a family of automata. Inf. and Control 4, 245–270 (1961)
19. Tzeng, W.: A polynomial-time algorithm for the equivalence of probabilistic automata. SIAM Journal on Computing 21(2), 216–227 (1992)